

dream machine

the next creative economy

pete woodbridge
dreamlab.org.uk

Dream Machine

dreamlab.org.uk

Dream Machine

The New Creative Economy

Pete Woodbridge

Edition 4 June 2026

DreamLab AI Collective

dreamlab.org.uk

Dream Machine: The New Creative Economy

© 2026 Pete Woodbridge. All rights reserved.

Edition of 4 June 2026.

Written and edited at DreamLab in the North West of England, between October 2025 and June 2026, in parallel with thirty-two editions of the *Dream Machine* newsletter.

All footnoted claims are sourced. Every footnote links either to a primary source or to the issue of *Dream Machine* in which the item was first discussed.

A note on this edition. **Dream Machine** is a *living book*. The print edition is updated and re-published on a rolling cadence — currently roughly monthly — as the newsletter ships new issues, as deep-dive research is added, and as the field's evidence sharpens. The full corpus, the build pipeline, and the always-current copy of the manuscript are open at github.com/petejwoodbridge/DreamMachine-Book. The unversioned *Dream_Machine.pdf* file in the published archive is always the most recent edition; this dated *2026-06-04* copy is one snapshot inside the longer sequence. The book's own argument — that AI enables new formats and publishing models, where the work, the toolchain and the commentary update continuously rather than in five-year hardback cycles — is one the book is also trying to demonstrate by being.

DEDICATION

This book is for the collective.

*To the **DreamLab team** — the around fifty deep-tech specialists who make the studio what it is across labs in Manchester, Liverpool, Cumbria and London. The AI engineers, XR developers, game designers, animators, sound artists, researchers, educators and filmmakers — the Emmy winners, the PhD researchers, the BAFTA-recognised practitioners — who turn up to do the work, who put taste and judgement into every brief, and who hold the human signal in the work the audience eventually sees. Every argument in this book has been sharpened by conversation with one of you. Several of you have, without knowing it, written the lines that landed best.*

*A particular thank you to **Dr. John O'Hare**, whose research on agentic organisational design — preserved in this book as Appendix K: The Coordination Collapse — supplied much of the structural argument underneath Chapter 13, and whose long-running collaboration on the harder organisational-theory questions has shaped large parts of Chapter 11, Chapter 14 and Chapter 15 as much as any single contribution from outside this desk. The 2,000-year framing of hierarchy as an information-routing protocol, the jagged-frontier reading of Dell'Acqua's BCG experiment, the judgment broker role replacing middle management, the compounding organisation concept — these are John's arguments, and the book is meaningfully more careful for them.*

To everyone at dreamlab.org.uk — see the team page — this is yours as much as it is mine.

— Pete

About this book

In September 2025, a synthetic actress walked onto a Zurich film festival stage, OpenAI shipped Sora 2, and Pete Woodbridge sat down to write a newsletter about it. Thirty-two weekly issues later, what began as a one-month experiment had become *Dream Machine* — the most-read working-creative record of the AI transition.

This is the book that record produced. From the slop ceiling that the audience is already enforcing, to the 88% of UK creators demanding licensing in all cases, to the chess-grandmaster strategy of deliberately playing the move the machine wouldn't make — Woodbridge holds the whole picture together in a way no journalist, academic or platform-company keynote has managed.

Equal parts history, manifesto and operating manual, *Dream Machine* is the field guide to the most consequential year in creative work since cinema learned to talk. Written from inside **DreamLab** — the studio Woodbridge founded at MediaCity in Manchester in January 2024, now around fifty AI engineers, XR developers, game designers, animators, sound artists, researchers, educators and filmmakers across labs in Manchester, Liverpool, Cumbria and London — drawing on the studio's closed-beta access to platforms like World Labs' Marble, project work with partners including Intel, Google, Sony, Aardman, BBC, NBC, Apple, MediaCityUK and Universal, and the operational experience of running an agentic-mesh studio in production every week of the period the book describes.

Dream Machine is a **living book**. New editions are published on a rolling cadence — currently roughly monthly — as the newsletter ships new issues and as the field's evidence sharpens. The unversioned *Dream_Machine.pdf* in the published archive is always the most recent edition; dated editions are preserved alongside as snapshots of the longer sequence. The cadence is part of the argument: AI enables new publishing models in which the work, the toolchain and the commentary keep moving with the field rather than freezing in five-year hardback cycles.

The age of the How is ending. Welcome the Why.

Contents

<i>Foreword — Welcome to the Dream Machine</i>	9
<i>Reader paths</i>	18
Chapter 1 — The Day Sora Landed	21
Chapter 2 — A History of Resistance	30
Chapter 3 — The Human-AI Agency Continuum	45
Chapter 4 — Dead Internet, Living Web	54
Chapter 5 — The Slop Ceiling	64
Chapter 6 — The 88%	75
Chapter 7 — The Studios Decide	91
Chapter 8 — Worlds, Not Pictures	107
Chapter 9 — AI in Everything, Everywhere, All at Once	117
Chapter 10 — What is newly possible	128
Chapter 11 — The Orchestrator	141
Chapter 12 — Authenticity as the New Scarcity	154
Chapter 13 — Coordination Collapse	167
Chapter 14 — The New Jobs	180
Chapter 15 — Choosing the Future	194
Chapter 16 — The Tools	210
Chapter 17 — Five Years Inside the Dream Machine	244
Epilogue — A Letter from the Dream Machine	253

APPENDICES

<i>Appendix A — A Quantitative Anatomy of Six Months</i>	258
<i>Appendix B — Glossary</i>	270
<i>Appendix C — Bibliography by Topic</i>	275
<i>Appendix D — The Shadow AI Paradox in the Creative Industries</i>	291
<i>Appendix E — Dynamics of Generative AI Adoption in the Creative Industries</i>	310
<i>Appendix F — AI, Stigma, Privilege and the Democratisation of Creative Expression</i>	330
<i>Appendix G — The Age of Intent: Artistic Mastery and the Inversion of Value</i>	351
<i>Appendix H — The Dream Machine Source Index</i>	367
<i>Appendix I — The Doomer Mistake: Why the Case for AI-Driven Civilizational Collapse Keeps Losing the Argument with the Data</i>	392
<i>Appendix J — The Process Trap: Generative Abundance and the Equivocation of the Creative</i>	402
<i>Appendix K — The Coordination Collapse: How AI Is Dissolving the Information Tax That Created Hierarchy</i>	414
<i>Appendix L — The Programmable Brand: Behavioural Licensing and the Governance Layer of the Generative Economy</i>	486

REFERENCE

<i>Citation Index</i>	499
-----------------------------	-----

ABOUT

About the author, and about DreamLab	546
--	-----

Foreword — Welcome to the Dream Machine

There was a week, at the very end of September 2025, when two things happened on opposite sides of the world that I knew, while I was watching them, would change the rest of my year — and probably most of yours.

In Zurich, an AI company called Particle6 walked an entirely synthetic actress called Tilly Norwood onto the Zurich Film Festival stage and announced that talent agencies were already in conversation about representing her. She had a face, a personality, a showreel and, by the founder's own framing, a future. Within forty-eight hours, the U.S. actors' union SAG-AFTRA had issued a statement that she was “not an actor” but “a character generated by a computer program that was trained on the work of countless professional performers.”¹ The U.K. union Equity followed.² By the time the weekend was over, Whoopi Goldberg, Emily Blunt, Melissa Barrera and a long list of other working performers had told the world, in their own words, what they thought of all this.³

The other thing was that OpenAI released Sora 2.⁴

I sat down on a Monday morning in early October, opened a blank LinkedIn article, and wrote my first edition of a newsletter called *Dream Machine*. I didn't have a plan. I had a sentence:

“It's time to take AI seriously.”

I sent that issue out to a few hundred people in my network. I assumed I'd do it for a month, maybe two — that the wave would pass, or that I'd run out of material, or that someone with a bigger platform would do it better. Eight months later, the newsletter has 3,800-odd subscribers, thirty-two published editions, several thousand curated links, and a community of people in the North West of England — the **DreamLab Collective**, around fifty of us across labs in Manchester, Liverpool, Cumbria and London — who help me read, sift, argue and build around it every week.⁵

This book is the thing that happens when you keep a careful, public record of an industry coming apart and putting itself back together inside the same eight months.

* * *

I should say who I am, because the rest of this book is in the first person and you should know what kind of “I” is talking to you.

I'm a **Creative Technologist** and **Innovation Director** working across AI, immersive, interactive and experiential tech, real-time production and next-generation storytelling — across film, games, interactive and immersive experiences. My cross-sector background spans creative, technical, innovation and product-development roles for software companies, games and immersive studios, and working on new possibilities in storytelling covering film, advertising, brand experiences and animation. The job title has had a few names over the years — *creative technologist*, *R&D director*, *innovation lead*, *head of emerging technology* — and they all shade into the same job: shipping software and creative products inside the gap between research labs, game studios, broadcasters and emerging-technology companies, then writing up what was learned along the way. The book is, in a real sense, a longer version of the *writing up what was learned along the way* habit.

I've worked on projects and new products with partners including **Intel, Google, McLaren, ITV, Sony, Aardman, Red Bull, IBM, BBC, NBC, Crytek, Apple, MediaCityUK, Saatchi & Saatchi** and **Universal** — alongside dozens of independent studios, agencies and SMEs building new products, content pipelines, and interactive and immersive experiences. Technically the work runs across strategy, leadership and hands-on development: real-time interactive and immersive systems (**Unreal Engine, Unity, WebXR, Disguise, Notch, Ventuz**); AI, software and workflow engineering (**Python, ComfyUI, API integrations, automation systems, AI-assisted pipeline design**); and production and future-media (**virtual production, motion capture, ICFX, real-time previs, AI-augmented workflows**). I've built and led teams of engineers and artists across most of those categories at one time or another, which is, more than anything else, the experience the book's arguments about orchestrator-shaped careers and agentic-mesh studios are made out of.

In parallel to the studio work I've developed innovation and investment programmes for the UK creative industries — including the **£3.2 million Innovation Accelerator (Innovate UK)** at **MediaCity** in Manchester, and the **£7.2 million MusicFutures programme (UKRI / AHRC)** in Liverpool — and I sit as an external assessor and advisor on several national and international innovation programmes and strategies, including **CoSTAR, EPSRC** and **Digital Catapult**, alongside contributions to a number of steering committees, boards, startups and investment initiatives. I include this not as a CV flourish but because a meaningful slice of the book's policy arguments — the Petrillo-template chapter, the apprenticeship-gap section, the four-principles framework in Chapter 15 — has been sharpened in conversation with the institutions running those programmes, and the reader can fairly take that into account when reading the policy chapters.

DreamLab, the studio I'm writing this book from, has been operating since **January 2024**. It started at **MediaCity** in Manchester and has, in the two-and-a-half years since, grown into a network of **around 50 deep-tech specialists** across labs in **Manchester, Liverpool, Cumbria** and **London** — and continues to spread. The team is AI engineers, XR developers, game designers, animators, sound artists, researchers, educators and filmmakers, some of whom have won Emmys and BAFTAs, some of whom finished their PhDs last year. We don't operate as a fixed agency; we assemble the right team of specialists around each challenge — a model that, in the period covered by this book, has become both the operational pattern of the orchestrator economy described in Chapter 11 and our own default way of working. Across the two and a half years since **DreamLab** launched, the studio has supported **100+ businesses**, delivered **35+ collaborative projects** and helped channel approximately **£2.7 million of investment** into partner organisations. None of that is the book's argument; it's the *vantage point* the book's argument is being made from.⁶

I'm not an AI evangelist. I'm also not an AI sceptic. I'm the kind of practitioner who has had **Marble** running on a beta key for months and who has also sat in a room with a games studio CEO who used the phrase "AI was an expensive mistake" without breaking eye contact.⁷ I have built things with these tools and I have watched them break. I'm writing this book from inside the work — and writing it, by the way, in the same voice I write the newsletter in, which is the same voice I argue with my team in, which is the same voice (give or take a bracket)

you'll hear if we ever end up in a pub together. A fuller version of the bio and a fuller picture of DreamLab — including the studio's client list, recognition and structural position inside the UK creative-tech ecosystem — sits in the [About the author, and about DreamLab](#) section at the back of the book.

Through the eight months the book covers, I have been doing three things in parallel:

- **Running DreamLab**, the ~50-person studio above, which uses every major AI platform discussed in this book in live production. We have been a closed-beta partner for World Labs' Marble since October 2025.
- **Writing the *Dream Machine* newsletter** — thirty-two weekly issues, several thousand curated source links across film, music, games, advertising and broadcast, a subscriber base of roughly 3,800 working creatives at the time of publication.
- **Building out a small set of analytical frameworks** — the *Human-AI Agency Continuum* (Chapter 3), the *Slop Ceiling* (Chapter 5), the *Four Positions* map of studio strategy (Chapter 7), the *Year of the Orchestrator* (Chapter 11), *Coordination Collapse* and the *Consumption Gap* (Chapter 13), the *AI Literacy Premium* and the *Apprenticeship Gap* (Chapter 14), the *Four Principles* of a humane creative economy (Chapter 15), and the *Seven-Layer Toolchain* model (Chapter 16) — that have begun, over the period, to circulate in the wider industry conversation.

The combination of the three is unusual. The trade-press journalists write the coverage but do not, on the whole, run studios. The studio operators run studios but do not, on the whole, publish a public weekly record. The academics produce research, sometimes excellent, but at the cadence of academia rather than the cadence of the transition. The platform companies produce material at the cadence of their own product cycles. Working from across the three positions at once, week by week, has produced a particular kind of vantage point — neither outside the work nor confined to one slice of it — that I have not seen anyone else holding consistently through this period.

The book is what that vantage point produces. There are good histories of cinema written by people who never ran a film studio, and good histories of the music industry written by people who never produced a record. There are also, sometimes, books written from inside the work — by people who were making the decisions the book describes, in real time, under the same conditions. The second kind tends, when honest, to tell you something the outsider accounts cannot: what the moment *felt like* to be inside, what the practitioners thought they were doing, and what — looking back — the moment was actually for.

This is that kind of book.

I am writing it because, in early October 2025, I realised that nobody I respected was doing what I needed somebody to do: hold the whole picture in one place. Not the boosters. Not the doomers. Not the niche tool-reviewers. The whole picture, week by week, including the contradictions. Including the bits where Adobe was telling us that 86% of creators already use generative AI in their workflow,⁸ and where 88% of creators who replied to the UK government's copyright consultation said AI companies should have to license their work in every case.⁹

The whole point — to me — is that those two numbers are both true.

* * *

The arc of these eight months turned out to be tighter than I expected.

In **October 2025**, the question on the table was whether AI in the creative industries was real. By the time I finished writing Issue 5, it was no longer a question.¹⁰

In **November**, the question shifted to whose tool stack it would run on. Adobe announced — and I am quoting them, this isn't editorial flourish — “AI in everything, everywhere, all at once”.¹¹ World Labs released Marble for public use a couple of weeks later, and the entire shape of what “a creative asset” can be quietly changed.¹²

In **December**, the question was about consent and money. The UK government's copyright consultation closed with eleven and a half thousand responses — one of the largest copyright consultations the country has ever run — and a number that has stayed with me: 88%.¹³

In **January 2026**, the question was who decides the rules. Sundance launched an AI Literacy Initiative for filmmakers. Bandcamp banned AI music outright. Steam clarified, then re-clarified, what counts as AI in a video game. Almost 800 creators signed an open declaration with the line, “*Stealing our work is not innovation.*”¹⁴

In **February, March, April and May**, the question started to feel like a different question altogether. It wasn't really *should we use these tools*. It was: *now that the tools are inside the production pipeline at every studio, every label, every newsroom, every agency, what kind of creative economy do we actually want on the other side?*

That last question is the one this book is about.

What this book is and is not

This is **not a tools guide**. Chapter 16 lists every significant tool that surfaced in the *Dream Machine* archive in the period the book covers, but the rest of the chapters are organised around the *transition* — the economics, the audience, the labour, the unions, the law, the institutions, the rails — rather than around the apps. The tools change weekly. The transition is what will still be true in 2030.

This is **not a manifesto**. I do not believe the cleanest five-point plans for the future of creative work, and I have refused to write one. What this book argues for, in Chapter 15, is a *test* you can apply to any policy, any contract, any platform decision: *Agency, Attribution, Access, Audience*. The test is not a programme. It is a way of staying oriented in a fast-moving environment.

This is **not a chronicle**. The newsletter is the chronicle. Every issue is online, every link is preserved, and the comprehensive thematic source index at the back of this book (Appendix H) catalogues the entire archive by topic for the reader who wants to follow specific threads.

What this book is is an argument, in sixteen chapters and eight appendices, that creative work is being re-platformed in a twelve-month window — that this is not the internet of 1995 or the

mobile phone of 2007, this is a faster, deeper, more thorough re-platforming of the economic and cultural rails on which creative work travels — and that the choices being made *right now*, by studios, by unions, by governments, by toolmakers, by individual creatives at their kitchen tables, will set the terms for the next decade. The book is here to help you make those choices on better information than you would otherwise have.

That re-platforming has a name. It's the title of the book. *The New Creative Economy*. I don't think it's a metaphor and I don't think we have very long to decide what we want it to look like.

What I believe, stated plainly

The book lays out my position chapter by chapter, but if you want the headline conviction up front, it is this:

AI is best understood as an assistive instrument that amplifies human creativity. Not a replacement for it. Not a substitute for it. An amplifier of it.

The working creative economy that emerges from this transition will be — has to be — the one that does not lose sight of which side of that relationship is the master and which is the servant. The human creativity is the master. The AI is the servant. Chapter 15 is the long-form version of that argument; everything else in the book sits inside it.

I want to flag a second framing the book leans on, because it is the one I use most often in the talks I have been giving. **We are leaving the age of the *How* and entering the age of the *Why*.** The *How* — the technical labour of executing a creative thought, the ability to draw, light, mix, model, edit, render — has been the bottleneck of professional creative work for a century. The *How* is, in 2026, becoming a utility. A teenager with a midrange GPU can now produce work whose surface quality sits on a continuum with what a full studio could produce in 2020. When the *How* becomes a utility, the *Why* — taste, intent, authenticity, the willingness to take a risk on the move the data does not yet endorse — is the only thing left with commercial leverage. Chapter 15 anchors this argument in a story I borrow from elite chess; the rest of the book lives inside the strategic implication.

If you read no further than this front matter, you have the heart of the book.

What this book does not do well

A critic-friendly note, because it makes for a more honest read.

This book is more confident about the *creative-industries* layer of the AI transition than it is about the layers above and below. The environmental and energy footprint of the systems the book describes is something I touch on in Chapter 15 and otherwise under-treat; a fuller account is the subject of a different book, by a different writer, that I hope is being written now. The labour conditions of the global data-supply chain — the labellers, evaluators and content moderators that the platform companies depend on — sit underneath every chapter of this book without being centred in any of them; the same caveat applies. The geopolitics of AI, the macro-economic question of the platform-company stock-market valuations, the wider policy questions about national AI strategy, the philosophical questions about machine

consciousness — none of these is the book's subject, and the book is shorter and more useful for not pretending otherwise.

The book is also, by design, anglophone-skewed and Global-North-skewed in its primary sourcing. The 88% in Chapter 6 is a UK number. The Sundance turn in Chapters 8 and 9 is a US story. The platform-company analysis in Chapter 9 is, in the main, an analysis of US and European companies, with significant Chinese coverage but less than the Chinese open-source ecosystem deserves. The Indian, African, Latin American and Southeast Asian stories I cover in Chapters 7, 11 and 12 are real but I cover them, in places, from the wrong side of a translation gap. The next edition, if there is one, should fix this. The deep-dive appendices begin the work but do not finish it.

Finally: the book is written *while the transition is happening*. Some of the specific claims, particularly in the tools chapter and in the predictions, will age in ways I am not yet able to predict. The frameworks should outlast their evidence. The evidence should be checked, when you read the book, against whatever the state of play is by then.

Who this book is for

I have written this book for **working creatives** — the writers, directors, songwriters, games designers, photographers, illustrators, editors, producers, agency creatives, indie filmmakers, YouTubers, freelance designers, students and senior practitioners who are, right now, trying to figure out what creative life looks like in 2026 and beyond.

It is also, secondarily, written for the **studio, agency and label leadership** trying to make organisational decisions about AI integration in a year in which the cost of getting it wrong is, by my read, the next decade of cultural authority.

And it is, thirdly, written for the **policy, union, institute and platform** people who are deciding the rails the next decade of creative work will run on. The 88% turned up to the UK consultation. The Cannes Disclosure Standard, the Academy rule, the SAG-AFTRA contract — these are the institutional decisions that shape the field. The people making them are part of the audience for this book.

Whoever you are: read with a pen. The chapters do not need to be read in order — there is a [Reader Paths](#) guide for different routes through the material. The Source Index at the back lets you follow any thread back to its primary sources.

A practical note, and an ask

I've written the book in the same voice I write the newsletter — talkative, opinionated, North-West English, occasionally too fond of a bracket. But I have tried, in every chapter, to put my opinions on top of evidence rather than the other way around. Every claim that matters is footnoted. Every footnote points either to a primary source (a research report, a court filing, an official announcement) or to the *Dream Machine* edition where I first wrote about it, where the original link is preserved. There are several thousand citations. If you only ever read this book once, you can ignore them entirely; if you ever want to know whether I made something up, follow the trail.

One more practical thing I want to flag before we start, because it is, in a small way, part of what the book is arguing. *Dream Machine* is a **living book**. The print and digital editions are updated and re-published on a rolling cadence — currently roughly **monthly** — as the newsletter ships new issues, as the research deep dives are added, and as the field’s evidence sharpens. Behind the manuscript there is a continuously-rebuilt **knowledge base** at DreamLab: a structured corpus of the newsletter archive, the deep-dive PDFs, the citation index and the toolchain catalogue, with a build pipeline that re-assembles the book end-to-end every time the underlying material moves. The unversioned `Dream_Machine.pdf` in the published archive is, by design, always the most recent edition; the dated copies are snapshots inside the longer sequence. If you want the latest, take the unversioned file. If you want the *specific* version a particular argument first appeared in, take the dated one. Both are kept.

I have made the cadence deliberate, not as a marketing posture but as a quiet demonstration of the book’s own argument. The position the book settles on — that AI enables new formats, new publishing models and new ways of holding a body of knowledge alive in public — is one I think a book about AI in the creative industries has an obligation to *show*, not only to *say*. A traditional non-fiction hardback ships once, sits frozen for five years, and ages out of its own subject matter. *Dream Machine* updates the way the field does. It is published the way the field works now. It is, in that sense, an artefact of the practice it describes. I would not write it any other way.

I have one ask of you before we start.

The temptation, when reading a book about AI in the creative industries in 2026, is to take a side before chapter one. To decide, on page one, whether this is going to be about how the machines are coming for us or about how the machines are setting us free. Please don’t. The most honest thing I can tell you about what I have learned over these eight months is that the truth is almost always both at once, and that the most interesting people in this story — the directors, the songwriters, the games developers, the union reps, the platform engineers, the indie filmmakers in bomb shelters and the policy officers in Whitehall — are the ones who can hold both sides at the same time without flinching.

That’s the kind of book I want this to be.

Welcome to the Dream Machine.

— Pete Woodbridge
DreamLab AI Collective
The North West of England
May 2026

* * *

The complete book in reading order:

- **Foreword: Welcome to the Dream Machine** — this piece
- **Reader Paths** — six routes through the material
- **Chapter 1: The Day Sora Landed** — the watershed week

- **Chapter 2: A History of Resistance** — the 1839–2025 pattern, and where AI sits inside it
- **Chapter 3: The Human–AI Agency Continuum** — the foundational frame, and the case for opening the black box
- **Chapter 4: Dead Internet, Living Web** — the web infrastructure question, and the provenance stack
- **Chapter 5: The Slop Ceiling** — the audience’s verdict, and the Authenticity Premium
- **Chapter 6: The 88%** — the political watershed, and the Petrillo template applied to AI
- **Chapter 7: The Studios Decide** — four strategic positions, and the trap legacy built for itself
- **Chapter 8: Worlds, Not Pictures** — the spatial turn
- **Chapter 9: AI in Everything** — the platform layer, and its underlying economics
- **Chapter 10: What is Newly Possible** — six categories of newly-possible work, and the finite-attention ceiling
- **Chapter 11: The Orchestrator** — the working creative’s new role
- **Chapter 12: Authenticity as the New Scarcity** — the provenance economy
- **Chapter 13: Coordination Collapse** — shadow AI, the consumption gap, organisational re-shaping
- **Chapter 14: The New Jobs** — the labour-market restructuring, the AI literacy premium, the Apprenticeship Gap
- **Chapter 15: Choosing the Future** — the four principles, the assistive-amplifier conviction, the Age of the Why
- **Chapter 16: The Tools** — the complete categorised inventory
- **Chapter 17: Five Years Inside the Dream Machine** — a speculative five-year future-cast, argued from the book
- **Chapter 18: Epilogue** — a letter to 2030
- **Appendix A: Quantitative Anatomy** — the numbers behind the book
- **Appendix B: Glossary** — terms of art
- **Appendix C: Bibliography by Topic** — curated reading
- **Appendix D: The Shadow AI Paradox** — covert adoption, displacement, hypocrisy
- **Appendix E: Dynamics of Generative AI Adoption** — the consumption-gap evidence
- **Appendix F: AI, Stigma, Privilege, Democratisation** — the class question
- **Appendix G: The Age of Intent** — the philosophical spine
- **Appendix H: The Dream Machine Source Index** — thematic catalogue of every significant source across the 32 issues

Reader paths

This is a book about creative AI in the eight months between October 2025 and June 2026. It is not a tools guide. It is an argument, with the evidence underneath, about what kind of creative economy is being built right now and what we should do about it.

You can read this book straight through. It is built to reward that. The book has eighteen sections — a combined Foreword, seventeen chapters, an Epilogue — plus twelve appendices, sequenced so that each one earns the next.

If you don't have time for that, here are six ways into the book that will save you from reading something that doesn't serve you yet. Pick one. Come back to the rest later.

If you are a working creative trying to figure out what to do next

Read the **Foreword**, then **Chapter 2 (A History of Resistance)** for the historical pattern you are inside, then **Chapter 3 (The Human–AI Agency Continuum)**, then **Chapter 10 (What is Newly Possible)** for the new categories of work, then **Chapter 11 (The Orchestrator)**, then **Chapter 14 (The New Jobs)** for the labour-market evidence, then **Chapter 15 (Choosing the Future)** — particularly the section *What working creatives should do on Monday morning*. That is the spine of the *practitioner's* argument. Everything else in the book is evidence supporting it.

If you run a studio, agency, label or production company

Read the **Foreword**, then **Chapter 7 (The Studios Decide)** — particularly the section *The trap the legacy industries built for themselves* — then **Chapter 13 (Coordination Collapse)** — particularly the section *What organisations should do*. Then read **Chapter 9 (AI in Everything, Everywhere, All at Once)** for the platform-economics frame, **Chapter 10 (What is Newly Possible)** for the new business categories, and **Chapter 14 (The New Jobs)** for the labour-market restructuring evidence. **Chapter 15** is the closing argument. **Chapter 16 (The Tools)** is the practical inventory. For the deeper organisational-design treatment that Chapter 13 rests on — the 2,000-year history of hierarchy as information-routing protocol, the jagged frontier, the vigilance problem, the three competing post-hierarchical models and a 90-day implementation roadmap — read **Appendix K (The Coordination Collapse)**.

If you are working in policy or law

Read the **Foreword**, then **Chapter 6 (The 88%)** and **Chapter 12 (Authenticity as the New Scarcity)** as a pair. **Chapter 14 (The New Jobs)** for the labour-market policy framework. **Chapter 15** for the four principles. **Appendix A** for the data the policy arguments rest on; **Appendix F** for the class-and-democratisation analysis; **Appendix H** for the comprehensive source archive.

If you are working in (or covering) the music industry

Read **Chapter 5 (The Slop Ceiling)** first — it's where most of the music-specific analysis lives. Then **Chapter 6 (The 88%)** for the rights-and-licensing argument. Then **Chapter 12 (Authenticity as the New Scarcity)** for where it's heading. **Chapter 16 (The Tools)** for the comprehensive music-AI tool inventory.

If you are working in (or covering) film, TV or games

Read **Chapter 1 (The Day Sora Landed)** for the watershed scene-setter, **Chapter 7 (The Studios Decide)** for the strategic map of how the industry has positioned itself, **Chapter 8 (Worlds, Not Pictures)** for what is coming next, **Chapter 11 (The Orchestrator)** for what it means for working roles, and **Chapter 14 (The New Jobs)** for the labour-market data. For the harder underlying argument about why the *creative / creatives* equivocation collapses under generative abundance, read **Appendix J (The Process Trap)**.

If you want the empirical case against the AI-doomer policy framing

Read the **Foreword**, **Chapter 14 (The New Jobs)** for the labour-market reading the book settles on, and then **Appendix I (The Doomer Mistake)** for the long-form rebuttal — the Stanford Digital Economy Lab data, Brynjolfsson’s worker-level studies, Acemoglu’s macroeconomic modelling, and the 230-year track record of identical predictions that have been wrong every time they have been asserted with confidence. **Appendix E (Adoption Dynamics)** and **Appendix D (Shadow AI)** are the supporting evidence base.

If you are working in (or covering) IP, licensing, brand strategy or legal architecture

Read **Chapter 6 (The 88%)** for the political mandate the licensing-side conversation rests on, then **Chapter 7 (The Studios Decide)** — particularly the new *Position Five: Own the rails* section on Hasbro’s Sixth Wall, the CharacterOS governance layer and the corrected record on the Disney–OpenAI deal that never executed — and then **Chapter 12 (Authenticity as the New Scarcity)** for the provenance and disclosure infrastructure that the IP layer sits inside. **Appendix L (The Programmable Brand)** is the long-form treatment: the *three Cs* (consent, credit, compensation) framework, the five-licensing-model comparison, the Coase / Williamson / Barney economic foundations, and the conceptual move from *IP as asset* to *IP as environment*.

If you are reading this in a class, a book club, or as part of training

Read the **Foreword** and **Chapter 1** to get oriented, then **Chapter 4 (Dead Internet, Living Web)** to understand the structural argument, then **Chapter 14 (The New Jobs)** and **Chapter 15 (Choosing the Future)** to see the labour-market story and the four principles the book argues for. The other chapters are evidence and elaboration. The **Glossary** (Appendix B), **Citation Index**, and **Source Index** (Appendix H) are designed to be the back-pocket reference set for the rest of the year.

* * *

A note about reading order: the chapters are designed to be load-bearing on each other. If you skip a chapter and a later one references it without re-explanation, the Glossary should fill the gap. If the Glossary doesn’t fill the gap, that is my failure and not yours — please write to me through the newsletter and I will improve the next edition.

— Pete

Chapter 1 — The Day Sora Landed

The first thing to understand about the week of 30 September 2025 is that nothing in it was supposed to be a watershed. Sora had existed, in some form, for a year and a half. Runway, Pika, Kling, Luma, Veo and a dozen others had been releasing video models on a near-monthly cadence since the back end of 2023.¹⁵ AI-assisted post-production had been quietly integrated into nearly every studio pipeline in Hollywood for months. The major U.K. and U.S. actors' unions had been negotiating over digital replicas since the 2023 SAG-AFTRA strike. The Tilly Norwood character had been on Instagram, posting selfies and pretending to drink coffee in cafés, since the previous July.

And yet that week — the seven days I would later mark as the start of this book, the start of the newsletter, and the start of a year that nobody in the creative industries has fully recovered from — three things lined up in a sequence so neat that I almost didn't believe it at the time.

On the Friday and Saturday, at the Zurich Film Festival, the founder of an AI studio called Particle6 stood on a panel and announced that several talent agencies were interested in representing her company's flagship product: a fully AI-generated actress called Tilly Norwood.¹⁶

On the Tuesday, the U.S. actors' union SAG-AFTRA issued a public condemnation calling her "not an actor" but "a character generated by a computer program that was trained on the work of countless professional performers" — adding, in a line I have thought about more than any other this year, that "audiences aren't interested in watching computer-generated content untethered from the human experience."¹⁷

In between those two moments, on the Tuesday of the same week, OpenAI released Sora 2.¹⁸

I want to take each of these in turn, because the order matters, and then I want to argue that the actual watershed was something else entirely — something that happened underneath all three, that almost nobody noticed at the time, and that I think we will be living with for the rest of our working lives.

The actress

Tilly Norwood was the creation of Eline Van der Velden — a comedian, writer and producer who had spent the better part of a year building her under the banner of a U.K.–Netherlands company called Particle6.¹⁹ Tilly had a face that looked like it had been assembled by committee from the most marketable features of the late 2010s. She had a voice. She had a small social-media following. She had, by the time of the Zurich announcement, "another forty AI actors in the pipeline," according to Van der Velden's later interview in *Deadline*.²⁰

The Zurich announcement wasn't subtle. Van der Velden told the audience that several talent agencies were "looking" at signing Tilly. She framed her as an industrial product: a character who could be cast in feature films and television, who came with all the upsides of a real performer (a marketing footprint, an emotional connection with audiences, a brand) and none of the downsides (no salary, no per diems, no licensing complications, no aging, no scandals).²¹

It is hard, looking back, to remember how loaded the word *agency* was that weekend. In Hollywood and the wider acting industry, the moment a talent agency takes on a new client is the moment that client moves from aspiration to product. The agencies are the gatekeepers of the working economy. Van der Velden's announcement, in essence, was that the gate had cracked.

The response was almost immediate, and it came from the only side of the gate that had anything to lose.

SAG-AFTRA's statement, issued on the Tuesday after the festival closed, called Tilly Norwood "not an actor." The full quote went further than the headlines tended to carry, and the next sentence was the one that actually mattered for the industry that read it: "*Signatory producers should be aware that they may not use synthetic performers without complying with our contractual obligations, which require notice and bargaining whenever a synthetic performer is going to be used.*"²² That sentence — terse, contractual, almost dull — was the union reminding every signatory studio in Hollywood that the 2023 strike had already settled this question, and that the rules already on paper applied here too.

The U.K.'s actors' union Equity issued its own condemnation within hours, focusing less on the philosophical question of what an actor is and more on the practical one: where had the training data come from? Their general secretary put it more sharply than any of the other initial responses: "*We're at the stage in AI where so much data has been used that the original source becomes more and more unclear. And that's something that should worry every viewer, every working person, because that's not really the way our data should be used.*"²³ If you build a model on the labour of working actors without consent, payment or attribution, you have, in their view, not created a new performer — you have repackaged a stolen one.

Within seventy-two hours, the public response from working actors followed. Emily Blunt called it "really, really scary." Whoopi Goldberg, on *The View*, dismissed the entire premise. Melissa Barrera, Kiersey Clemons, Lukas Gage — a roll-call of working performers, mostly the ones with enough security in their careers to be willing to say anything on the record at all — each took their position.²⁴

Eline Van der Velden, Particle6's founder, replied. "*To those who have expressed anger over the creation of our AI character Tilly Norwood: she is not a replacement for a human being, but a creative work — a piece of art.*" The defence carried, in miniature, the entire shape of the argument the working actors were about to spend two years pushing back against. *Art. Creation. Not a replacement.* The whole rhetoric of the AI-native studios, in one sentence, on the Monday morning of October 2025.

The line that got repeated most often, in the threads and the green rooms and the union briefings I read that week, was a variant of: *we already gave you eighteen months of strike to settle this question.* The 2023 SAG-AFTRA strike had not been about AI in the abstract. It had been, in significant part, about digital replicas and the right of performers to control the use of their likeness in synthetic content. The deal that ended that strike had — supposedly — set

the rules for the next era. Tilly Norwood, dropped onto a festival stage eighteen months later, was a test of whether those rules meant anything.

The answer was: they were going to have to be re-tested in public, on every individual case, for years.

The model

What turned the Tilly Norwood weekend from a single-news-cycle controversy into the moment everyone in the creative industries started paying serious attention was that, on the Tuesday she was being condemned, OpenAI released Sora 2.

Sora 2 was not just an upgrade. It was, by OpenAI's own framing, a step-change in three things at once: physical realism (a ball that bounces correctly off a backboard), audio integration (sound effects and synchronised dialogue baked in, not added in post) and what the company called "world state" — the ability to follow instructions across multiple shots while keeping the scene logically consistent.²⁵ If Sora 1, a year and a half earlier, had been the model that made people sit up and notice that AI video was a thing, Sora 2 was the model that made people sit up and notice that AI video might be a *medium*.

The launch came with a second thing that, looking back, was the actually significant part: an invite-only iOS app, also called Sora, that worked like TikTok. You scrolled. You remixed. You did "cameos" — the in-app feature that let you drop a generated likeness of yourself, or a friend, or anyone you had a clip of, into the model's output.²⁶

Within five days, the Sora app had hit a million downloads.²⁷

The thing I want you to hold in your head about this is the timeline. The model was announced on the Tuesday. The app launched the same week. By the Friday, it was on the front page of the App Store. By the following Monday, the first wave of celebrity deepfakes — Robin Williams's daughter calling them "gross"; Michael Jackson generated into music videos he never made; dead historical figures being put through new dialogue by users who didn't realise they were doing anything illegal because the app's design hadn't told them so²⁸ — was being written up in *The Guardian* and the *NBC News* tech vertical.²⁹

OpenAI's own likeness-protection rules, *The Hollywood Reporter* noted that week, had a specific carve-out: dead celebrities and "historical figures" weren't covered.³⁰ The carve-out was treated, in the first days of the app, as a feature rather than a problem.

The clearest line of the entire launch week came not from OpenAI's blog post but from *The Guardian*'s technology reporter, in the lead of their coverage of the Sora app's first violent and racist outputs. "*In 2022, [the tech companies] would have made a big deal about how they were hiring content moderators ... In 2025, this is the year that tech companies have decided they don't give a shit.*"³¹ I think that sentence will end up being remembered as one of the best one-line summaries of the moment. It was, if anything, optimistic about 2022. It was wholly accurate about 2025.

The reason this matters is not that Sora 2 was the first AI video model. It wasn't even the best one, by some metrics, that month — Veo 3.1 from Google would land in mid-October with

arguably more sophisticated cinematic controls.³² The reason it matters is that the *combination* of a major model release and a consumer iOS app, in a single week, collapsed a distinction that had — until that week — kept the conversation about AI in the creative industries safely in the hands of the people who made the creative industries.

Until Sora's iOS app, AI video was something that happened on a desktop, with a subscription, with a prompt window. After Sora's iOS app, AI video was something that happened on the phone of every teenager with an invite code, on a swipeable, remixable, social feed.

The line between “an AI tool a working filmmaker might use” and “a default app on a default phone” had been the line that held the cultural debate in shape for two years. That week, it didn't.

The phrase

The line that has stayed with me from those first seven days came not from a celebrity, or from OpenAI's launch page, or from a union release. It came from a small Florida public-radio station's website. The headline read: *Kiss reality goodbye: AI-generated social media has arrived.*³³

I have read that headline a hundred times since the week it was written. It is a perfect sentence. It is also — and I think this is what made the early days of Sora 2 so vertiginous — *prematurely* true. Reality, as a category, did not end the week that Sora 2 launched. People still had real lives, real friends, real coffees in the morning, real bills at the end of the month. What ended, that week, was the easy assumption that what you saw on your phone had been made by somebody — a person, a team — with a recognisable connection to a recognisable place.

What's striking about the WUFT piece, and the few hundred similar pieces that followed it, is that they were not, in the main, written by people in the AI industry. They were written by reporters at local stations, by columnists at regional papers, by the kind of journalist who covers schools and county budgets and the planning department. The collapse of the line between “made by a person” and “made by a machine” was being noticed first not by the experts but by the audience.

That, more than anything else that happened that week, was the thing I wrote in the first issue of the newsletter. *It's time to take AI seriously.* The line, six months on, embarrasses me a bit — it sounds like the kind of thing you say when you don't have anything else to say — but it was, on the morning of 6 October 2025, the only sentence I could write that felt like it was about the actual situation.

The director

There was one more piece of the early reaction that I want to flag, because it set the template for almost every “high-end” creative response that followed.

In late September, the website *No Film School* ran an interview with James Cameron in which he said — bluntly, on the record, in a quote that travelled — that AI was “never going to take

the place” of humans in filmmaking. “Filmmaking is subconscious,” he said, “and can’t be quantified.”³⁴

Two months later, in promotion for *Avatar: Fire and Ash*, Cameron expanded on the point in a CBS *Sunday Morning* interview that became one of the most-shared creative-industry stories of the entire winter. Asked about generative AI’s ability to “make up a performance from scratch with a text prompt,” Cameron said: “It’s like, no. That’s horrifying to me.” And then, in a line I have used in talks and in arguments and in this book: “*The act of performance, the act of actually seeing an artist creating in real time, will become sacred.*”³⁵

He went on, in a passage that almost never travelled with the headline quote: “*It also causes us to have to set our bar to a very disciplined level, and to continue to be out-of-the-box imaginative.*” The whole interview, taken together, is not the doom-laden refusal the press cycle reduced it to. It is a working filmmaker articulating an argument about *craft* in an era of automated production — that the existence of cheap synthetic performance does not retire the human performer, it *raises the discipline required of the human performer*. The audience’s bar moves up. The work that wins it has to move up too.

The companion line Cameron gave *Variety* the same month was even more telling: “*For years, there was this sense that, ‘Oh, they’re doing something strange with computers and they’re replacing actors,’ when in fact, once you really drill down and you see what we’re doing, it’s a celebration of the actor-director moment.*”³⁶ Cameron has spent forty years building digital filmmaking. He is, as a working artist, perhaps the most sophisticated user of computer-aided performance in the history of cinema. His argument is not against AI; it is against AI *that displaces the moment in the room where an artist creates*. That distinction — *between AI that augments the actor-director relationship and AI that substitutes for it* — is the one most working creatives are trying to draw, and the one most press coverage of the AI debate still flattens.

What almost every report of those quotes glossed over — and what made them more interesting, not less — was that Cameron was, and still is, a board member of Stability AI.³⁷ He is not an opponent of AI in filmmaking. He is, by any reasonable definition, an investor in it. His position was not “no AI.” His position was “no AI that replaces the actor in the room.” That distinction is the one most working creatives, in my experience, are trying to make. The early coverage, looking for the clean villain-or-saviour story, mostly missed it.

I think Cameron’s *Sunday Morning* line is the most important thing said by a working filmmaker in the entire six months I have been writing the newsletter. Not because it is right about everything — I think the “sacred” framing makes a smaller and more brittle claim than it sounds — but because it is the first time, in the AI era, that one of the people who *built* the apparatus of digital cinema in the 1990s and 2000s drew a line that he himself was prepared to defend.

The line, again, is not “no machines.” Cameron has been a machine-builder for forty years. The line is: a human creates *in real time*, and that creation is the work. Everything else is just delivery.

The state of the field, the week before

I want to spend a section on what was already in motion when the Tilly Norwood announcement happened, because the historians of this period are, I suspect, going to under-tell the *cumulative* story underneath the *catalysing* one. The week of 30 September 2025 was not the moment AI arrived in the creative industries. It was the moment the *audience* arrived in the AI debate. The infrastructure underneath had been forming, in a slow, uneven, partly-public, partly-private way, for at least two years before it.

Let me sketch the field.

By the time Sora 2 launched, the major AI-video model release cadence was running at roughly one significant model per fortnight. Runway had shipped Gen-4 in late 2024 and was deep into the public roll-out of Gen-4 Image-to-Video and the Workflows product by September 2025. Luma had released its Dream Machine consumer app; the Genie 3 demo from Google DeepMind in late summer 2025 had been described by *Time* as one of the year's best inventions. Pika 2.0 was shipping. Higgsfield, which would close \$80M on a \$1.3B valuation by January 2026, was already on its third major product cycle. Hunyuan Video and Wan 2.2 — the open-source Chinese-built models from Tencent and Alibaba — had been freely available, on commodity GPUs, for months. Kling, the Kuaishou model that would, by mid-2026, be the model many professional filmmakers actually used for production-grade clips, had been in continuous public release. Sora 2 was the headline of the week. It was not, in any meaningful sense, the entire field. The field was already crowded.

The same is true of the studio-side adoption. Lionsgate had publicly partnered with Runway on a deal to produce AI-augmented studio films a year before *Futurism's* headline-grabbing “crumbled into disaster” piece. Netflix's quiet integration of AI tools into background-plate generation, animated short development and de-aging post-production work had been documented across 2024 and the first half of 2025. Disney's *House of David* — the show whose creator would, in November 2025, defend the use of more than 350 AI-generated visual-effects shots in its second season — had been in production with that pipeline in place months before the Sora 2 launch. The major broadcast and streaming companies had been integrating AI under the hood at a pace that the public conversation had not yet caught up with. The Tilly Norwood week was, in part, the moment that pace became impossible to keep quiet.

The unions had been working on this even longer. The 2023 SAG-AFTRA strike had — through the Writers Guild and through the actors' bargaining — produced contract language on “digital replicas” that was, by 2025, already two years old. Equity in the UK had been running consultations and ballots through 2024 and into 2025. The Authors Guild's class-action lawsuits against OpenAI had been filed in mid-2023 and were grinding through discovery. The European AI Act had been finalised in 2024 and was beginning to bite on copyright disclosures by the time of the Zurich announcement. The Music Performance Trust Fund's emerging conversations about an AI-era levy mechanism — the Petrillo-template applied to neural-network outputs that I argue for in [Chapter 2](#) and [Chapter 15](#) — were already on the agenda at AFM Local 802 and at the UK Musicians' Union long before Tilly

walked on stage. The institutional response had pre-existed the cultural rupture. The cultural rupture is what made the institutional response politically possible.

The adoption telemetry on the platform side, in the months before the Sora 2 week, was already at the level that would later be revealed in Adobe's MAX 2025 Creators' Toolkit Report and the Stanford AI Index 2025. Firefly was already on track for its 22-billion-asset milestone. ChatGPT, by the time Sora 2 launched, was already at roughly 700–800 million weekly active users on its way to 900 million. The 86% of global creators who reported using generative AI in their workflow in the Adobe survey — published a few weeks after the Tilly Norwood week, but capturing data collected before it — was a number the platform-companies had been quietly watching for months. The Adobe survey did not produce the adoption. It documented the adoption that was already complete by the time the public was paying attention.

On the *consumer* side, the Sora app was, again, the catalyst rather than the inventor of the dynamic. The TikTok-style consumer surface for AI generation had been visible for at least a year. ByteDance's Dreamina and CapCut tools had been integrating Seedance, Seedream and the wider ByteDance generative stack into consumer-facing video editing through 2024 and into 2025. The Sora app's million-downloads-in-five-days number — the headline that defined the week's consumer dynamics — landed inside a market that had been *prepared for it* by Dreamina, CapCut, the Krea consumer app and the Suno and Udio consumer-facing music platforms. What was new about the Sora app was not the *concept* of swipeable AI generation. What was new was that an American flagship-AI company had chosen to ship it as a *primary* consumer surface alongside the model.

Particle6, finally, had been building Tilly Norwood for the better part of a year before the Zurich panel. Van der Velden's Instagram and TikTok rollouts had been running since the previous summer. The character had a follower count, a posting cadence, a slowly-built visual identity, a small but real fanbase that engaged with her as if she were a person. The Particle6 strategy — the deliberate cultivation of a *parasocial* relationship between an AI character and a human audience, in the months before the studio-system pitch — was a play borrowed from the playbook of the Korean virtual-idol economy, the Japanese vocaloid scene, and the long history of cultivated personas in the influencer industry. What was new about Tilly was not the *idea* of a synthetic personality with a fanbase. What was new was the framing of that synthetic personality as a *casting option* for legacy film and television. The Zurich announcement was the moment the parasocial-character economy and the working-actor economy were proposed, on a festival stage, as overlapping markets. The reaction was, in retrospect, the audience and the union refusing to let the markets overlap.

What all of this means, when you stack it together, is that the Tilly Norwood / Sora 2 week did not invent the AI moment in the creative industries. It *named* it. It made it impossible to keep treating AI as a technical category that working creatives could opt out of. The model release was overdue. The Particle6 announcement was overdue. The union response was already drafted in some form. The audience reaction was already, on the slop-ceiling logic I lay out in Chapter 5, structurally inevitable. What the week did was *force everything to happen in*

public, at the same time, in front of an audience that had not previously been part of the conversation.

I think that is the deeper reason the week mattered. The conditions were ripe. The catalyst was small. The reaction was big. The change in the *visibility* of the AI transition — from a back-room toolchain conversation to a front-page audience question — was, by my reading, the actual watershed. Everything else in this book is downstream of that visibility shift.

The watershed under the watershed

I said earlier that the actual watershed of that week was not the Tilly Norwood announcement, and not the Sora 2 launch, and not the union responses, and not the Cameron quote. The actual watershed was something underneath all four.

For two years, the conversation about AI in the creative industries had been a conversation among insiders. Toolmakers talked to creators. Creators talked to studios. Studios talked to unions. Unions talked to governments. The general public — the audience for the things being made — had largely been a backdrop. They had been the people *for whom* this argument was happening, not the people in the argument.

In the week of Sora 2, that ended.

The Sora app put a generative video tool on the phone of anyone with an iOS device and an invite code, and the invite codes were not hard to come by. The Tilly Norwood announcement put the abstract concept of “a synthetic actress” into the *Daily Mail*, *The Guardian*, *The View*, the breakfast television circuit, and three different morning radio shows I happened to be listening to that week. The Cameron quote, when it came, ran on every wire service that covers the entertainment business.

The audience joined the argument. Not as a side, but as a participant.

And once the audience is in the argument, the argument changes. It is no longer about what the unions can negotiate, what the studios will adopt, what the toolmakers will ship. It is about what the people who watch films and listen to music and play games and scroll feeds will accept, demand, refuse and forgive.

The rest of this book — the eleven chapters that follow this one — is, in one way or another, an account of what the audience has been doing with the argument since it became theirs. The artists’ boycotts and the streaming platforms’ counter-moves; the eleven and a half thousand UK citizens who turned up to a government consultation; the 50,000 AI-generated tracks uploaded to Deezer every day and the 1 to 3 percent of streams those tracks actually got;³⁸ the death threats sent to Tilly Norwood’s creator in early 2026; the moment in the middle of January 2026 when the U.S. actors’ union and SAG-AFTRA went back to the negotiating table because the audience, having looked at the new landscape, had decided what it wanted.

The week of 30 September 2025 was the last week before any of that.

I started writing the *Dream Machine* newsletter the Monday after. By the time I sent the first edition out, the conversation had already moved on.

Chapter 2 — A History of Resistance

I want to start this chapter with a piece of writing.

“Sweeping across the country with the speed of a transient fashion in slang or Panama hats, political war cries or popular novels, comes now the mechanical device to sing for us a song or play for us a piano, in substitute for human skill, intelligence, and soul... Let us not hamper it with a machine that tells the story day by day, without variation, without soul, barren of the joy, the passion, the ardor that is the inheritance of man alone.

Singing will no longer be a fine accomplishment; vocal exercises, so important a factor in the curriculum of physical culture, will be out of vogue! Then what of the national throat? Will it not weaken? What of the national chest? Will it not shrink?

When a mother can turn on the phonograph with the same ease that she applies to the electric light, will she croon her baby to slumber with sweet lullabys, or will the infant be put to sleep by machinery? Children are naturally imitative, and if, in their infancy, they hear only phonographs, will they not sing, if they sing at all, in imitation and finally become simply human phonographs — without soul or expression?”³⁹

Read that paragraph once more, and try, before I tell you when it was published, to date it.

If you guessed *autumn 2024 trade-press editorial on Suno*, you would not be alone. The language sits comfortably alongside the take-pieces about AI music slop that filled the music press in the months I started writing the newsletter. *Machinery in substitute for human skill, intelligence and soul. The story told day by day without variation. Children growing up as imitative human phonographs, without soul or expression.* Every clause has a 2024–25 equivalent.

The essay is not from 2024. It was published in **September 1906**, in *Appleton’s Magazine*, by **John Philip Sousa** — the most popular bandleader in the United States at the time — and it is titled *“The Menace of Mechanical Music.”*⁴⁰ The machine Sousa was warning his readers against was the phonograph. The “soul-barren” recording technology Sousa feared was Edison’s flat disc, spinning at 78 rpm, playing back music captured by a brass horn.

The same essay was — almost word for word, with a few changes in the names of the machines — written about the player piano in 1900, the microphone in 1932, the synthesiser in 1980, the drum machine in 1991, Auto-Tune in 1998, non-linear video editing in the early 1990s, the digital camera in the 2000s, the smartphone-as-camera in the 2010s, and generative AI in 2023, 2024 and 2025.

That is the subject of this chapter. The structured, recurring, almost ritual *pattern* by which every major creative-technology introduction in the modern era has been received by its working practitioners — and what that pattern, eighteen iterations later, tells us about how to think about the AI moment the rest of this book describes.

I want to be careful about how I do this. The historical-analogy move is, in tech writing, a famously cheap one. *“Every disruptive technology has been resisted; therefore your resistance to this technology is wrong”* is the rhetorical operating system of two decades of

platform-company keynotes, and it has been wielded so dishonestly that working creatives in 2026 are right to be suspicious of any version of it. I am not, in this chapter, making that argument. I am making a *more specific* one. The historical pattern has — across twenty distinct technologies in the period from 1839 to 2022 — a recognisable five-act shape, and the *five-act shape* is what the pattern tells you. Not that resistance is wrong. Not that the new tool will be fine. Something more useful than either: which institutional moves work, which fail, what gets preserved, what gets lost, where the *new* creative forms come from, and what the working practitioner's actual leverage in the period is.

The history is, in other words, *operationally* informative. That is the use I want to make of it.

The five-act curve

Every resistance I look at in this chapter — and there are many more I do not have room for — moves through roughly the same five stages, in something like the same order.

Act One: Ridicule. The new tool is dismissed as a toy. It cannot compete. It sounds awful, looks crude, has the wrong specifications. The serious practitioners are unconcerned because the work it produces is not, on inspection, work. The Roland TR-808 drum machine, released in 1980 and a commercial failure, was reviewed as “toy robot drums.” The Canon 5D Mark II, the DSLR that started the death of dedicated cinema cameras, was — in 2008 — dismissed as a stills camera with a video gimmick. Auto-Tune, between 1997 and roughly 2003, was used as an *undisclosed* studio tool because no working singer wanted to admit that their pitch was being machine-corrected.

Act Two: Moral panic. The new tool is reframed as a threat to public morals, aesthetic standards, or the integrity of the form. It is degenerate. It is theft. It is not real *singing*, not real *photography*, not real *art*. The 1932 sermon by **Cardinal O’Connell** of Boston that crystallised this stage for the microphone — that crooning was “a degenerate form of singing,” that “no true American man would practice this base art,” and that crooners were “whiners and bleaters defiling the air” — is unimprovable as a template.⁴¹ “*Imbecile slush*,” O’Connell called it, in language that anyone who has read a 2024 anti-AI op-ed will recognise. The 1991 federal court ruling in *Grand Upright Music v. Warner Bros.* — the case in which Biz Markie was sued for sampling Gilbert O’Sullivan — opened with the words “*Thou shalt not steal*,” quoting the Seventh Commandment, in a US copyright opinion.⁴² In 2025, the moral-panic stage of AI is mostly behind us; the analogous *Thou shalt not steal* language is in the UK’s 88% copyright consultation response, the *Stealing Our Work Is Not Innovation* declaration, and the union statements quoted throughout this book.

Act Three: Existential professional alarm. The displaced practitioners realise the tool is not a toy and not just morally suspect — it is structural. It is going to take their work, change their craft, and reshape the institutions that support them. The classic statement is **Phil Tippett’s**, the stop-motion master who saw ILM’s first digital test of a *Jurassic Park* T. rex in 1992 and said: “*I think I’m extinct*.”⁴³ Spielberg liked the line enough to put a paraphrase of it in the film. Tippett was right about himself in some local sense — his go-motion craft did not survive *Jurassic Park’s* release. He was wrong about himself in a wider one — his studio went on to produce digital animation work for *Starship Troopers* and is still operating in 2026.

Both readings can be true at the same time. Most of the existential-alarm moments work like this.

Act Four: Institutional and legal counter-attack. Unions strike. Lawmakers legislate. Lawsuits get filed. The 1942 and 1948 **Petrillo strikes** — when the American Federation of Musicians, led by James Caesar Petrillo, refused to record for the major labels — are the canonical version of this stage, and I will spend longer on them in a moment. The UK Musicians' Union's **1980 “Massacre of the Musicians”** BBC strike and its **1982 motion to ban synthesisers** are the British version. The 1991 *Grand Upright* ruling that sampling was theft is the legal version. The 2007 *Viacom v. YouTube* \$1bn lawsuit is the platform-distribution version. The 2023 SAG-AFTRA and WGA strikes — which gave us the contract language that frames Chapter 1 — are the most recent before the period this book covers.

Act Five: Settlement. The dust settles into one of three forms. The displaced craft *dies*: miniature painting after 1840, hand-drawn feature animation after *Toy Story*, the photo-lab business after digital. The new tool is *taxed and the revenue redistributed*: the Petrillo settlement created the **Music Performance Trust Fund**, still distributing payments to live musicians in 2026; the DMCA Section 512 plus YouTube's Content ID created a parallel pool of platform-paid royalties for the music industry; needletime in the UK forced the BBC to pay for live sessions until 1988. Or — most often — the two creative categories *coexist*: photography did not kill painting, it forced painting toward what photography could not do (Impressionism, abstraction); the microphone did not kill singing, it *redefined* what counts as singing; sampling did not kill composition, it redefined what counts as composition.

That is the shape. It is, by my count, the shape of every single technological transition in the creative industries in the period 1839 to 2022. It is happening, around AI, right now — and where we are on the curve is part of what this chapter is for.

Let me walk briefly through a handful of the cases, because the *texture* of the historical record is, I think, more useful than the abstract pattern alone.

The photograph (1839)

The daguerreotype was unveiled in Paris on 7 January 1839 and publicly described to the Académie des Sciences on 19 August. By 1849, roughly 100,000 daguerreotypes had been produced in Paris alone; by 1861, about 33,000 people in Paris were making their living from photography and photographic supplies. The professional class most immediately wiped out was the **portrait miniaturist** — painters of small ivory-based likenesses, the working photographers of the pre-photographic age. They could not compete on speed, price or fidelity. Within a single working generation, the craft was effectively gone.

The most articulate resistance was not from the miniaturists themselves but from the literary intelligentsia. **Charles Baudelaire's** 1859 essay “*The Modern Public and Photography*,” published in the *Revue Française*, made the case in language a 2024 AI-sceptic would recognise: “*this industry, by invading the territories of art, has become art's most mortal enemy.*” And, harder: “*The photographic industry was the refuge of all failed painters, too ill-equipped or too lazy to complete their studies.*”⁴⁴

The famous **Paul Delaroche** line — “*From today, painting is dead*” — is, the historians who have looked carefully tell us, almost certainly apocryphal. The earliest sourced version of it appears in an 1873 survey, thirty-four years after Delaroche supposedly said it; the painter’s own contemporary writing on the daguerreotype called it “*an immense service to the arts,*” and he continued painting until his death in 1856.⁴⁵ The story has outlived the saying. (This is itself a pattern. I will come back to it.)

The settlement took seventy years. Alfred Stieglitz founded the Photo-Secession on 17 February 1902. *Camera Work* ran from 1903 to 1917. The “291” gallery opened in 1905. MoMA established the first photography department at a major museum on 31 December 1940. From invention to *full institutional acceptance of the new form as art*: about a century. The compensating gain: the entire history of modern photography as a fine-art tradition, an industrial portrait business, a documentary-journalism profession, and — eventually — the cultural substrate on which the smartphone-camera moment of the 2010s rests.

What painting did, meanwhile, was redefine itself. Impressionism (light, atmosphere, the subjective moment), Post-Impressionism (the interior state), Cubism (multiple viewpoints), and ultimately abstraction — every one of these moves makes more sense if you read them as painting’s response to the daguerreotype’s having taken representation. The standard art-history reading, which I think is correct, is that *photography liberated painting from the burden of representation*. The fear-namers were locally right and structurally wrong. Painting did not die. It became something else.

The phonograph (1877–1920s)

I have already quoted Sousa. The whole 1906 *Appleton’s* essay is worth reading; almost every paragraph of it could be republished, with names changed, in 2025.

The institutional response to the phonograph — and this is the part of the story that working creatives in 2026 need to know — came in the form of the **1909 Copyright Act**, the first statutory acknowledgement in US law that *machine reproduction of human creative work required a legal regime*. The 1909 Act created the **compulsory mechanical licence** for recorded music, the structural ancestor of every machine-licensing argument we are now having about AI training. Sousa, in part because his lobbying helped pass it, prospered in the recording era; his compositions still generate royalties under that licensing structure today.

The phonograph also created entirely new occupations that the parlour-music economy could not have anticipated. The recording engineer. The A&R executive. The producer. The mastering engineer. The sleeve designer. The pressing-plant operator. The retail-store buyer. The pop *single* as a commercial form. Jazz on record. The LP. The concept album. The bedroom-studio that, in the 2000s, would take all of those occupations apart again. Sousa got the *parlour* prediction right — recorded music did dent amateur home music-making — and the *industry* prediction completely wrong. The recording industry was the largest expansion of working-musician employment in the history of music, and it was made possible by the technology Sousa thought would destroy it.

The Petrillo bans (1942, 1948)

The **American Federation of Musicians** had, by the 1940s, watched the phonograph, the talking picture (which alone wiped out roughly 22,000 cinema-orchestra jobs in the US in 1927), and commercial radio progressively displace live performance. **James Caesar Petrillo**, AFM president from 1940, did the thing every union threatened by a creative technology has tried, with varying success, to do since: he turned off the recording machine.

On 1 August 1942, AFM members stopped recording. The strike lasted **27 months**. Decca settled in 1943; RCA Victor and Columbia in November 1944. The settlement: a per-record royalty paid into an AFM fund for unemployed musicians.

On 1 January 1948 Petrillo did it again, triggered by the Taft-Hartley Act outlawing the first royalty arrangement. The 1948 settlement created the **Music Performance Trust Fund** under Section 302 of Taft-Hartley — a jointly-administered labour-management fund paid into by the labels and broadcasters, used to subsidise free live performances by working musicians. The MPTF still exists. It still pays out, in 2026, several million dollars a year for live music.⁴⁶

I want to dwell on this for a moment, because the Petrillo settlement is — by some distance — the **most operationally important precedent** for how the AI debate could land, and almost nobody in the current creative-AI conversation talks about it.

The Petrillo template has four parts. *One*, the displacing technology is not banned. It is allowed to displace. *Two*, the platform owner pays an ongoing per-unit tribute to the displaced labour pool. *Three*, the tribute is collected centrally, by a joint labour–management body, not negotiated individual-by-individual. *Four*, the tribute is paid out to subsidise the *displaced creative practice itself* — live music, in this case — keeping it alive as a category even as the market for it shrinks.

The SAG-AFTRA *Tilly Tax* provisions in the 2026 contract, the UK 88% licensing-by-default proposal, the *Creative Weight Attribution* musical-AI infrastructure I described in [Chapter 5](#), the C2PA / SynthID provenance stack from [Chapter 12](#) — these are all, on inspection, *attempts to reconstruct the Petrillo template for AI*. Per-unit tribute. Joint collection. Redistribution to the displaced practice. The mechanism is the same. The political question is whether the platforms will accept it.

The mechanism worked once. It can work again. The pattern of the resistance that fails — the 1982 UK Musicians’ Union motion to *ban synthesisers outright*, the European *Right to be Forgotten* style absolutism on training data — is the pattern of resistance that tries to legislate against the machine rather than to tax it. Levy beats ban, every time. Royalty pool beats injunction. *Mechanism, not prohibition*.

The microphone (1932)

Cardinal O’Connell’s January 1932 sermon to 3,000 men of the Holy Name Society of Boston is the single funniest moment in the resistance literature. *Crooning* — the conversational, intimate, microphone-enabled vocal style that **Bing Crosby**, Rudy Vallée and others had built into a mass commercial form by the late 1920s — was, for O’Connell, “a degenerate form of singing. No true American man would practice this base art... If you will listen closely [to

crooners' songs] you will discern the basest appeal to sex emotion in the young." Crooners were "whiners and bleaters defiling the air." Their work was "imbecile slush."⁴⁷

The cultural fight was effectively over by the end of the decade. Bing Crosby was the biggest male voice in America. By the late 1930s no popular vocalist *not* trained in microphone technique could make a competitive career.

What is interesting about the microphone case, for our purposes, is *what it did to the underlying definition of the craft*. Before the microphone, *good singing* meant volume, projection, throat technique — the operatic, theatrical, music-hall tradition that O'Connell had grown up inside. After the microphone, *good singing* meant timbre, intimacy, breath control at low dynamics, conversational diction — a fundamentally different skill set. Almost every popular vocalist since 1935 has been a "crooner" in the technical sense, including those who would not call themselves that. *The microphone redefined what counted as singing*. The vocalists who refused the microphone are mostly remembered as period figures. The ones who absorbed it defined the rest of the century of popular music.

This is the deeper pattern I will come back to. Resistance, in the named-fear form, almost always defends *the existing definition* of the craft. The settlement, almost always, *redefines* the craft. The fear is real and the language is sincere; what's actually happening is bigger than what the fear is naming.

The synthesiser (1980–82)

The UK part of this story is the cleanest because it generated public archives. The **Yamaha DX7**, released in 1983, put credible electric piano, brass, strings, marimba and dozens of other sounds into a single keyboard at a price point — about £1,500 — accessible to working session players. The DX7 was used on roughly **40% of US Billboard Hot 100 #1 singles in 1986**. Session keyboard players and orchestral string sections previously hired for adverts, TV scoring, library music and pop sessions were directly displaced.

The **Musicians' Union** of the UK responded in two stages.

1980, the Massacre of the Musicians. The BBC announced in March 1980 that it would cut 172 staff orchestra posts and disband five of its eleven in-house orchestras. 83% of MU BBC members voted to strike. The strike began 16 May 1980. **The First Night of the Proms was cancelled for the first time in its history.** The strike ran until 1 August 1980, ending with a compromise: the BBC Northern Ireland Orchestra and BBC Midland Radio Orchestra were disbanded as planned; the others survived.⁴⁸

1982, the synthesiser ban motion. On 23 May 1982 — by coincidence Bob Moog's birthday — the MU's Central London Branch passed a motion to ban synthesiser, drum-machine and electronic-device use by union members entirely.⁴⁹ The trigger was that Barry Manilow's UK tour had replaced its string section with synth players. The motion was never adopted as full union policy. The MU's Executive Committee passed a more measured resolution in November 1982. A breakaway "Union of Sound Synthesists" was formed. *Top of the Pops*, for a period, required bands to record their backing tracks the afternoon before the

show “to prove they could actually play it.” None of it held. The DX7 sold over 200,000 units. By 1990 the cultural debate was over and the synthesiser was, simply, another instrument.

What the *MU* got wrong, in retrospect, was the same thing every resistance-by-prohibition gets wrong: they tried to ban the tool rather than to tax it. There was no Petrillo-style fund. There was no per-output levy on synthesiser use. There was no royalty pool subsidising live orchestral work. The *MU*’s defensive posture preserved nothing structurally and lost the cultural argument decisively. The session-musician economy contracted. Some of the displaced players retrained as programmers and prospered. Others did not. **Trevor Horn** — the Buggles who, three years before the dispute, had sung *Video Killed the Radio Star* — became the defining pop producer of the 1980s.

The lesson, for the working creative in 2026, is not that the *MU* was wrong to resist. The *MU* was *right* to read the displacement signal three years ahead of the rest of the industry. The *MU* was wrong about *what kind* of resistance to mount. Prohibition was always going to lose. A levy-and-pool argument — what an *AI-era Petrillo* could look like — might have held.

Drum machines, sampling and the lawsuits (1980s–2000s)

The Roland TR-808 (1980) and TR-909 (1983) both *failed* commercially on release. The 808 was criticised as “toy robot drums”; the 909 was reviewed as “still sounds like a drum machine, instead of a machine playing drums.” Both became cult instruments only after being dumped at secondhand prices to young hip-hop and dance producers in the mid-1980s. The 808’s hand-clap, snare and signature deep kick are now the defining percussion sounds of contemporary popular music. The same dismissed-then-canonised arc is, in 2026, beginning to play out for Suno and Udio at the consumer-music end of the spectrum.

The legal resistance to sampling produced two rulings every working creative should know about, because they are the cleanest available templates for how the courts may treat AI-training disputes.

Grand Upright Music v. Warner Bros. (S.D.N.Y. 1991). Biz Markie sampled three bars of Gilbert O’Sullivan’s *Alone Again (Naturally)* without clearing it. Judge Kevin Thomas Duffy’s opinion opened with the words “*Thou shalt not steal*,” quoting the Seventh Commandment, in a US federal court ruling. He referred the case to the US Attorney for potential *criminal* investigation. The album was pulled. The sample-dense Bomb Squad / Public Enemy style of production it had been built on became commercially impossible.⁵⁰

Bridgeport Music v. Dimension Films (6th Cir. 2005). NWA’s *100 Miles and Runnin’* sampled a two-second guitar chord from Funkadelic’s *Get Off Your Ass and Jam*, looped and pitched down. The Sixth Circuit eliminated the *de minimis* defence for sampling sound recordings and issued the rule that has, for twenty years, defined how clearance works: “**Get a license or do not sample.**”⁵¹

Hip-hop, of course, did not die. It became the dominant global popular music form. What changed was that the *aesthetic* of dense, layered, sample-heavy production gave way to a more clearance-friendly style. The Bomb Squad lineage continued, but on different terms.

The AI-training analogy here is direct, and I would commend the rulings to anyone trying to think clearly about *UMG v. Anthropic* and the cases that will follow it. *Thou shalt not steal* and *Get a license or do not sample* are not, despite their archaic phrasing, particularly anti-technology rulings. They are *pro-licensing* rulings. They say: the tool can be used, but the inputs have to be paid for. That is what working creatives are asking for in the UK 88% and the *Stealing Our Work Is Not Innovation* declaration. The line of legal reasoning is already in the books.

Auto-Tune (1998–2010)

I want to spend less time on Auto-Tune than the dossier supports, because the case is so clean it almost makes itself. Andy Hildebrand, a former Exxon seismic-data engineer, released Auto-Tune in 1997. **Cher's Believe** (1998), produced by Mark Taylor and Brian Rawling with the retune speed maxed out, produced the now-iconic “Cher effect” — the audible warble. Taylor and Rawling claimed it was a vocoder for several years to protect the trick.

T-Pain made the effect his signature from 2005 onwards and was rewarded with **Jay-Z's** “D.O.A. (*Death of Auto-Tune*)” — released June 2009 on *The Blueprint 3* — a direct moral-panic attack on what Auto-Tune was doing to vocal authenticity. **TIME magazine's 50 Worst Inventions list** in 2010 ranked Auto-Tune at #15: software that “can make bad singers sound good, and really bad singers sound like robots.”⁵²

Sixteen years later, Auto-Tune is on every pop vocal you hear, used as both correction and effect. Bon Iver's *22, A Million* (2016) deployed it as a self-conscious aesthetic instrument. Billie Eilish has used it across her career, transparently. The cultural rehabilitation is complete. The moral-panic stage, in retrospect, looks parochial.

But notice what *did* happen. The microphone-era settlement — *you have to be able to sing* — was, in the Auto-Tune era, definitively broken. The new settlement is *you have to be able to perform the post-corrected vocal as a self-conscious artistic choice*. The redefinition is real. Cardinal O'Connell, transported eighty years forward, would have hated Auto-Tune for exactly the same reasons he hated crooning, and would have been just as wrong about it.

Non-linear editing and the death of the splice

Avid Media Composer launched in 1989. Through the early 1990s it displaced the Moviolas, Steenbecks and KEM flatbeds that working film editors had used for sixty years. The American Cinema Editors did not strike. The Motion Picture Editors Guild did not stop the transition. The settlement was generational: editors trained on film retired; editors trained on Avid became standard.

The witness I want to quote here is **Walter Murch**, ACE — possibly the most respected film editor of the last fifty years. Murch edited *The Conversation* and *Apocalypse Now* on physical film. He edited *The English Patient* on Avid, winning Oscars for both picture and sound. He then edited Anthony Minghella's *Cold Mountain* (2003) on **Apple Final Cut Pro running on commodity Power Mac G4 hardware** — for a \$79m feature. The story is in Charles Koppelman's book *Behind the Seen* (Peachpit, 2004).⁵³ The decision saved about \$1m versus an equivalent Avid rental, and Murch chose it on practical grounds. The editor who literally

wrote the textbook on film editing — *In the Blink of an Eye*, the canonical philosophical text on the cut — was, by 2003, working on the tool the editing-room conservatives were warning the industry against.

The compensating gain, from the non-linear editing transition, was that the toolkit shipped on a laptop. Indie cinema benefited enormously. The grammar of the cut accelerated — the average shot length in Hollywood drama dropped from roughly ten seconds in the 1960s to roughly four seconds by the 2000s, a change made trivial by NLE that would have been physically punishing to execute on a Moviola. The contemporary visual grammar that ranges from *The Bourne Identity*'s hyper-cuts to Wong Kar-wai's asynchronous editing aesthetic to TikTok's stitched, layered, fast-moving native form is, in operational terms, *what non-linear editing made possible*. The form the new tool enabled was bigger than the form it replaced.

Kodak (1975–2012)

Steven Sasson, an engineer at Kodak, built **the first digital camera prototype in December 1975**. 0.01 megapixel, black and white, the size of a toaster, with a 23-second save time to magnetic tape. Sasson's own description, in multiple interviews, is that the executive response to the demo was *curious but concerned about the implications for film*. Kodak suppressed the project to protect its film business.⁵⁴

Kodak's peak headcount was about 145,000 in the early 1980s. At the company's Chapter 11 bankruptcy filing on **19 January 2012**, the workforce was about **19,000**. A nearly \$7bn liability stack. A company that had invented the technology that would destroy it, and then suppressed that technology, and then been destroyed by it anyway when the rest of the market — Nikon, Canon, Sony — built around the patents Kodak had not commercialised.

The Kodak story is the standard cautionary tale for incumbents about to be disrupted by the technology they themselves built. It is being told, with increasing force, about the legacy entertainment industries in 2026. *The companies that already have the IP, the audience and the distribution are, by structural inheritance, in the position Kodak was in in 1990*. What we will find out, in the rest of this decade, is whether they have learned the Kodak lesson — that the new technology is going to displace the old whether or not you commercialise it, so you may as well be the company that does — or whether the next few years will produce the entertainment-industry equivalent of the 2012 bankruptcy filing.

Photoshop, social media, smartphones — a faster montage

Three more cases, briefly, because the pattern is by this point clear and the texture matters more than the recital.

Photoshop and photojournalism (1990–2015). The famous **Brian Walski** firing, 1 April 2003 — LA Times staff photographer dismissed by satellite phone from Iraq for compositing two combat photographs into one stronger image — is the single canonical example.⁵⁵ World Press Photo's response, over the following decade, was to issue increasingly tight ethics rules and disqualify increasing numbers of finalists. The truth-claim of the photograph survived institutionally at AP, Reuters, *National Geographic* and *Magnum*. It did not survive on social media. The structural answer was C2PA — the content-authenticity

infrastructure that now sits underneath the AI provenance stack in Chapter 12. The 1990s photojournalism debate was, on inspection, *the first round of the C2PA conversation* the AI era is now finishing.

Social media (2005–2014). The defining institutional fight was **Viacom v. YouTube**: \$1bn in claimed damages, filed in March 2007, decided in favour of YouTube and the DMCA Section 512 safe harbour in 2010 and again in 2013, settled out of court in **March 2014 with no money changing hands** — a complete Google win.⁵⁶ The settlement was effectively *automated*: YouTube’s Content ID, deployed from 2007 onwards, has paid out billions in royalties to rights-holders through algorithmic matching. The DMCA safe harbour that decided *Viacom* — passed in 1998, in response to the previous resistance pattern — became the structural framework of the *creator economy*. Power-law distribution, \$100bn+ in aggregate value by 2022, MrBeast at 450m subscribers and \$600–700m in 2024 revenue. The platform shape of the audience contract in 2026, on which every AI-distribution question rests, was built by Content ID.

Smartphones (2007 onwards). The most dramatic collapse of an established consumer-creative-tool category in modern history. Compact-camera shipments: **108.8 million units in 2010, 3.01 million in 2021**. A 97% drop in roughly a decade.⁵⁷ The iPhone, the Galaxy and the Pixel did to the camera industry what Sousa thought the phonograph would do to amateur music — and Sousa, in retrospect, was about a hundred years early on his own argument. The smartphone-camera moment is also, on inspection, the dress rehearsal for the AI-image moment. The same arguments — *it isn’t real photography, the computational layer is doing the work, not the human, the dedicated camera is the moral artefact* — were made in 2010 and 2015 about smartphones. They were also lost in 2010 and 2015 about smartphones. Computational photography is, today, where most of the meaningful innovation in the image-making craft is happening. *AI photography*, in 2026, is not the new thing. It is the next iteration of the thing that the smartphone-camera made inevitable fifteen years ago.

Five patterns

I have, by this point, walked through ten of the twenty cases the dossier behind this chapter covers. The pattern recurs. Let me name the five structural features I think the working creative in 2026 most needs to internalise.

Pattern one: the curve has a predictable shape. Ridicule → moral panic → existential alarm → institutional counter-attack → settlement. AI in 2025–26 is, on my read, somewhere in late act three and early act four. The existential alarm has been voiced. The institutional counter-attack — the SAG-AFTRA strike, the WGA contract, the UK 88%, the *UMG v. Anthropic* suit, the Cannes Disclosure Standard, the Sundance literacy initiative — is well underway. The settlement is starting to form but is not yet stable. The next eighteen months will, on the historical pattern, *settle the form of the next decade’s industry*. This is exactly why this period matters.

Pattern two: the named fear is always mis-named. What practitioners say they fear — loss of *soul*, loss of *authenticity*, loss of *the real* — is almost never what actually happens. What actually happens is a *redefinition* of the underlying creative category. The microphone

redefined *singing*. The photograph redefined *painting*. The sampler redefined *composition*. The Avid redefined *editing*. The smartphone redefined *photography*. The fear is framed as a defence of the thing. What is actually at stake is the *definition* of the thing. The fear-namers usually win the surface argument and lose the definitional one. *The AI debate is, on the structural reading, an argument over what counts as authorship, performance, writing, photography and composition in the next decade.* That is the real fight. The named-fear version of it — *AI will steal our jobs* — is true but partial.

Pattern three: the institutional moves that work are levy-and-pool; the moves that fail are prohibition. The Petrillo settlement (MPTF, 1948), the DMCA safe harbour plus Content ID (1998 plus 2007), the 1909 Copyright Act mechanical licence, the eventual streaming-loudness-normalisation truce — these *worked*, in the limited but real sense that they extracted ongoing transfer payments from the new medium to the displaced labour pool, or that they restructured the aesthetic equilibrium. The institutional moves that *failed* are the prohibitions: the MU 1982 synthesiser ban, the 1991 *Grand Upright* effective ban on dense sampling, the European *Right to be Forgotten* style absolutism. **Levy beats ban. Royalty pool beats injunction. Mechanism beats prohibition.** For the AI-training fight: a per-output levy distributed to a creators' fund is in the workable category. A ban on AI training is in the unworkable category. The 88% — by my reading — is closer to the workable category than the politics around it have so far recognised.

Pattern four: the resisters are usually right about the local loss and wrong about the total loss. Miniature painting died. Hand-drawn feature animation died (*Toy Story* in 1995; Disney's Florida 2D studio closed on 12 January 2004). The professional recording-studio mid-tier died. Photo-processing labs died. Staff photographer jobs at US newspapers — collapsed (the *Chicago Sun-Times* infamously laid off its entire photo staff in May 2013). But *more people make moving images today than ever made them; more people make music today than ever made music; more photographs are taken on a given Sunday than were taken in the entire nineteenth century.* The aggregate creative-labour pool grew. The fear is always articulated by the *displaced* cohort. The compensating gains accrue to a different cohort, who don't yet exist when the fear is articulated, and whose names — by definition — are not yet in the trade press. This is why the resisters can be both *factually right* about their own situation and *structurally wrong* about the form they are trying to protect.

Pattern five: the cultural symbol outlives the cultural anxiety. *Video Killed the Radio Star* was the Buggles' 1979 lament for the death of a form. MTV used it as its launch trumpet at 12:01 a.m. on 1 August 1981. Then radio survived; then MTV died; the song is on TikTok. The artefact *about* the death of a form outlasted both the form it threatened and the form that did the threatening. Phil Tippett's "*I think I'm extinct*" in 1992, **Trevor Horn's** "*video killed the radio star*" in 1979, **Cardinal O'Connell's** "*imbecile slush*" in 1932 — these crystallise an anxiety into a piece of language so vivid that it survives whatever it was anxious about. The named fear becomes the source material for the next generation's art. *The anxiety is the art.* The *Tilly Norwood week* of late September 2025, with Whoopi Goldberg and Melissa Barrera and Emily Blunt speaking on the record, is — I am almost certain — the equivalent crystallising-moment for the AI era. The artefacts that will come out of it (the documentaries, the dramatic-feature treatments, the songs, the union-history books) will,

twenty years from now, be the cultural objects that *outlast* the technology they were originally anxious about.

Where AI sits in this picture

I want to be careful, in applying the diagnostic, not to wave a hand and claim the pattern *predicts* the AI outcome. It doesn't. What it does is rule out certain shapes the outcome cannot take, and rule in certain shapes it almost certainly will.

The shapes ruled out: an *outright ban* on AI training that protects the existing definition of authorship. A *prohibition* by union action that simply removes AI tooling from professional production. A *moral consensus* — Cardinal O'Connell at scale — that AI work is degenerate and should be socially refused. None of these has worked in twenty previous iterations of the same pattern. None of them is going to work this time. The MU's 1982 synth motion is in the books as a cautionary tale.

The shapes ruled in: a per-output levy structure flowing into a creators' fund — Petrillo for neural weights. A C2PA-style provenance standard underwriting an authenticity premium — the structural ancestor of which is the 1990s photojournalism ethics fight. A redefinition of the underlying creative category — *authorship in 2030* will mean something materially different from *authorship in 2020*, the way *singing* meant something different after the microphone. An institutional settlement that absorbs the new tool and *redistributes* the productivity gain, rather than one that *bans* the new tool and *forfeits* the productivity gain.

The cohorts who will be locally displaced are already visible in [Chapter 11](#) and [Chapter 14](#). Junior animators. Concept artists in cohorts being asked to use generative tooling for the front of their pipeline. Voice actors below the SAG-AFTRA scale level. Stock-image photographers. Translators of bulk commercial copy (the closest pre-2022 analogue — the Google Neural Machine Translation moment in 2016 — already produced documented translator-employment effects across 696 US labour markets, and a 70% income loss for the Irish-language EU-institutions translator profiled in the January 2026 *CNN Business* piece I cited earlier in the book⁵⁸). The fear of these cohorts is well-founded, and the institutional response should be calibrated to it.

The compensating gains — the new categories of creative work, the new business shapes, the new audience contracts — are what [Chapter 10](#) of this book is about. I will not re-state them here. What I want to note, *for the historical pattern's sake*, is that they always emerge, they always emerge faster than the resisters expect, and they always — at the aggregate level — produce more total creative employment than the displaced form supported. The phonograph was the largest creative-labour expansion in music history. The smartphone was the largest creative-labour expansion in image-making history. *I think AI, on the historical pattern, will be the largest creative-labour expansion in cultural-production history.* I think the working creatives who emerge from this with the most leverage will be the ones who, like Walter Murch picking up Final Cut Pro at the height of his Avid mastery, learned the new tool *before* the cultural permission to use it had fully crystallised. The cultural permission usually arrives about three years after the productive use does. The window for asymmetric leverage is now.

What this asks of working creatives

I want to close this chapter with a working-practitioner read on the history, because the abstract pattern is useless without translation.

One. Recognise where on the curve you are. AI in 2026 is in late act three, early act four. Existential alarm is the dominant mood. Institutional counter-attack is well organised. *Settlement* has not been reached. The decisions you make about how to engage with the tools in the next eighteen months are decisions that will define the structural shape of the rest of the decade. This is not a *theoretical* claim about the historical pattern; it is the *operational* claim about your career.

Two. Don't fight to keep the existing definition; fight to be among the people redefining it. The miniaturists are remembered as the cohort that was wiped out. The Stieglitzes are remembered as the cohort that *redefined what photography could be*. Both groups felt, in 1855, that they were fighting for the same thing. They were not. One was defending the inherited definition. The other was rewriting it. The working creatives who emerge from the AI period with the most leverage will not be the ones who defended hardest. They will be the ones who *redefined fastest*.

Three. Pick the institutional response that has historically worked. The Petrillo template — levy on the displacing technology, pool collected centrally, redistribution to the displaced craft — has a hundred-year track record. Use it. Apply it to your union negotiations. Apply it to your platform-procurement decisions. Apply it to your political advocacy. The SAG-AFTRA Tilly Tax, the UK 88%, the C2PA provenance stack, the Cannes Disclosure Standard, the Sundance literacy initiative are all, in their different ways, *the Petrillo template applied to AI*. They are the part of the institutional response that historically wins. Show up to them. Argue for them. Don't waste energy on prohibitions.

Four. Read the named fear carefully, and listen for the redefinition underneath it. When you hear yourself saying *AI will steal my craft*, ask the second question: *what new definition of my craft is forming on the other side of the displacement, and am I in a position to inhabit it?* The microphone vocalists who absorbed Crosby outlasted the operatic vocalists who refused him. The editors who absorbed Avid outlasted the editors who refused it. The photographers who absorbed digital outlasted the photographers who refused it. The pattern is, by this point in the historical record, very reliable.

Five. Make the cultural symbol. The named-fear language of the AI era — its *Video Killed the Radio Star*, its *I think I'm extinct*, its *imbecile slush* — is being written, right now, by working creatives in their own work. Some of those artefacts will outlast the technology they were anxious about. Some will become the source material for the next generation's understanding of this moment. *If you are a working creative reading this in 2026 and the AI displacement frightens you, the most useful thing you can do is put your fear into a piece of work whose argument outlasts the platform release cycle*. That is what Trevor Horn did. It is what Phil Tippet did. It is what the displaced miniature-painters who became fine-art photographers did. It is, on the historical pattern, what is asked of you.

A coda on Trevor Horn

I want to close with one image, because it is the single most useful one I know.

Trevor Horn, in 1979, was the keyboard player and frontman of the Buggles. He wrote, with Geoff Downes and Bruce Woolley, a song called *Video Killed the Radio Star*. The song is a small, half-melancholy commercial joke about the end of an era — a synthetic-sounding pop song *about* the moment recorded music turned visual. MTV launched on 1 August 1981 with that song as its very first broadcast.

Horn went on to become, by some distance, the most influential pop producer of the 1980s. He built Frankie Goes to Hollywood, ABC, Yes's *90125*, Grace Jones's *Slave to the Rhythm*, Seal's debut, the Art of Noise. He did it on Fairlight CMIs, sequencers, drum machines and the kind of dense, layered, programmed production that the Musicians' Union was, at the same moment, voting to ban. He did the thing the union was warning him against, and built half the canonical pop music of his decade out of it.

The song that named the death of the radio star outlived MTV. The producer who wrote that song defined the next era of recorded music by absorbing the technology the resisters were trying to ban. The cultural symbol outlasted both the form it threatened and the form that did the threatening.

That is the working operating model I would commend to anyone reading this book and feeling, in 2026, the gravitational pull of the resistance. The resistance is real. The fear is real. The named-fear language is the source material of the era's best art. *And the people who absorb the tool, who learn it, who push it past where its makers intended it, and who use it to make the work that argues for the world they actually want — they are the ones, on a hundred years of evidence, who define the next decade of the form.*

Welcome to the *next* moment in a recurring pattern. The pattern is, by now, very well documented. The choice is yours.

Chapter 3 — The Human-AI Agency Continuum

A week after I sent the first edition of *Dream Machine*, on the Monday of the second week of October 2025, OpenAI held its annual DevDay conference and quietly changed what the conversation about AI in creative work was about.

The first edition had been about Sora 2. The second edition was about something I was less prepared for: the launch of a thing called **AgentKit**.⁵⁹

AgentKit was, at first glance, a set of developer tools. Agent Builder. A connector registry. An eval framework. ChatKit, for embedding agents into other products. The launch post on OpenAI's blog framed it, in the slightly forced register that all platform-launch posts use, as a way for developers to “build, deploy, and optimize agentic workflows.”⁶⁰ On its own, this was an unremarkable announcement.

What was remarkable, looking back, was the *category claim* the announcement carried with it. Sam Altman, in his DevDay keynote that day, declared the start of “the age of agentic AI” — by which he meant the moment that AI systems stopped being prompt-and-respond chat boxes and started being things that could plan, decide and execute “for hours on end” without further human input.⁶¹

For someone like me, sitting in a small studio in the North West of England — running tools all day, looking at my pipeline, thinking about my team's labour — that phrase did a particular kind of work. It rearranged the question.

The question, until that week, had been: *what does AI do for creative work?* The question after that week became: *where, in any given piece of creative work, does my agency end and the model's begin?*

The second question is the one I want this chapter to be about. I called it, in the second issue of the newsletter, the **Human–AI Agency Continuum**.⁶² The frame has stuck with me. I think it is the most useful thing I have ever written down about all of this, and I think — at the risk of overselling it — the rest of the book leans on it.

The continuum

Imagine a horizontal line.

On the far left of the line is **pure human agency**: the writer at the desk, the painter at the canvas, the songwriter at the piano. No machine intermediation other than the tool itself — and the tool, in this position, is dumb. It records what you do; it doesn't decide.

On the far right of the line is **pure machine agency**: an autonomous system that, given a goal, produces a finished creative output with no human in the loop. A prompt, a setting, a render. No one looks at the intermediate steps. No one steers.

The conversation about AI in the creative industries in 2024 mostly took place on the assumption that “generative AI” sat about three-quarters of the way along that line — closer to the machine end. You typed a prompt; the machine made the thing; you accepted or rejected. There were variants, of course. But the geometry was prompt-and-respond, and the question was simply where on the line, between you and the model, the actual creative work happened.

What changed at OpenAI DevDay on 6 October 2025 — and what was reinforced almost every week of the six months that followed — was that the line is not, as it turned out, a single line. It is a *family of lines*, one per creative function, and they all move at different speeds.

A film, broken down, is not one act of agency. It is a thousand. The choice of subject. The treatment. The casting. The script revisions. The cinematography. The blocking on set. The performance, take by take. The editorial assembly. The grade. The sound. The music. The marketing. Each of those is a sub-discipline, with its own craft, its own labour pool, its own union, its own pay scale and its own internal hierarchies.

AI doesn't slide along *the* line. It slides along each of those lines independently.

A working filmmaker in late 2025 might sit at the absolute left of the continuum on *performance* (a real actor, in the room, in real time, the work itself) and at the absolute right on *background plate generation* (a Veo 3.1 shot, signed off in a Slack message, no human ever drawing a frame).⁶³ A working musician might sit at the absolute left on *songwriting* (a song in a notebook) and on the right edge of the centre on *vocal alignment and pitch correction* (an iZotope Ozone 12 assistant, accepted with one click).⁶⁴

The crisis of authorship is not that machines do creative work. Machines have done parts of creative work for as long as there have been cameras, samplers, Photoshop filters and Logic plug-ins. The crisis is that we don't have an honest, shared, public vocabulary for *which* parts. The Continuum, written down honestly per project, is the start of one.

Agents are not generators

The reason the Continuum became urgent the week of DevDay, and not before, is that “agent” is a different kind of object on the line than “generator” is.

A generator is a tool. You aim it at a problem; it makes an output. The agency is in the aiming.

An agent is something more like a junior collaborator. You give it a goal — *find me ten reference images for this shot, generate a rough sound design for this scene, book the courier for tomorrow's pickup* — and it goes away, makes a series of sub-decisions, and comes back with a result. The agency is distributed. You set the direction; it makes the moves.

The reason this matters in creative work is that the moves are where the craft lives. Anybody can describe a final film in a sentence. The film is in the thousand decisions between the sentence and the screen. A generator that makes the screen-ready file from your sentence isn't doing your craft. It is taking your craft out of the loop.

An agent, properly deployed, can do something different and — to me, anyway — more interesting. It can take the parts of the loop that are not where your craft lives, and quietly handle them, so that the parts of the loop that *are* where your craft lives become the parts you actually spend your time on.

That is the optimistic case for agentic AI in creative work, and it is the case that almost every working creative I respect makes when you sit down with them in private. It is also the case Adobe's 16,000-creator survey, released a few weeks after DevDay, came in to support: 70%

of respondents were optimistic about agentic AI, framed as “tools that act on your behalf”; 85% said they would use AI that learned their creative style.⁶⁵

The pessimistic case is the one Adobe’s same survey also captured: 69% of respondents worried about their work being used to train AI without consent.⁶⁶

Both numbers are about agency. The first is about *gaining* it back, by handing routine work to a competent assistant. The second is about *losing* it, by having the work that defines you absorbed into a system you do not control. Both are true at the same time, for the same creators, in the same workflows.

Where agents went, between October and May

In the six months between DevDay and the time I’m writing this, the agent layer of the creative toolchain went from “interesting demo” to “shipping product,” faster than any technology shift I have lived through in twenty years of practice. I want to give you a sketch of the trajectory, because it is what most of the rest of this book is reacting to.

By **mid-October 2025**, Mureka — a Chinese music platform — launched a thing called *Music Agent Studio*, six specialised AI agents for songwriting, arrangement and production.⁶⁷ A startup called AdsGency raised \$12m in seed to build agents that could autonomously run a brand’s entire paid marketing workflow.⁶⁸ A company called Lenny launched an agent for organising live music events.⁶⁹ Each of these felt, at the time, like a specialist tool. In retrospect, they were the first signs that whole production functions — not individual tasks — were being handed over.

By **the end of November**, EA, in the middle of a brutal financial year, told its 15,000 employees to use AI as a “thought partner” for everything from character art to playtesting.⁷⁰ The framing — *thought partner* — was the precise rhetorical move that turned an agent from a tool into a colleague. The colleague has opinions. The colleague has time. The colleague has a seat at the meeting.

By **December**, Adobe announced that you could now use Photoshop and Express *inside* ChatGPT — meaning that the creative output itself was no longer happening inside Adobe’s interface, but inside an agent’s.⁷¹ This was a small thing on the surface and an enormous thing underneath. It was the moment that Adobe — a company that has, since 1990, owned the metaphor of the *creative tool* — accepted that the new metaphor was the *creative agent*, and that they would rather be inside someone else’s agent than not in the conversation at all.

By **late January 2026**, Anthropic shipped Claude apps — interactive, custom assistants embedded directly in workplace tools — and a company called Heygen released *Video Agent*, which could script, edit and assemble entire videos from reference images.⁷² By **March**, Adobe announced its **CX Enterprise** platform alongside NVIDIA: a stack of AI agents embedded across the entire content lifecycle, from brief to delivery.⁷³ By **April**, the *Adobe Summit* keynote made it official — “agentic creative intelligence” was now the headline category, not a feature.⁷⁴ By **May**, Sony was using a multi-agent team of forty-nine Claude Code agents, working with seventy-two skills, to coordinate game-development work.⁷⁵

The trajectory, in one sentence: in October 2025 we were arguing about whether agents were a thing. By May 2026, the entire creative production pipeline at a global game publisher was being run by a team of them.

What this means for craft

The natural fear, reading that timeline, is that the agency line drifts inexorably to the right — towards the machine end — and that the craft of the human in the loop becomes thinner and thinner until it disappears.

I do not think that is what happens. I think what happens is more interesting and more demanding.

What I see, in my own studio, in my friends' studios, in the working musicians and filmmakers and games designers I talk to every week, is that agentic AI doesn't compress craft into nothing. It *relocates* craft to a different place on the continuum.

If your job, last year, was “make the thing”, your job this year is “decide what gets made, brief the agents that make the constituent parts, and judge the output.” That isn't a smaller job. In some ways it is a bigger one. It requires *more* taste, not less, because taste is now the only signal you bring that the agents cannot.

Anthropic, in a blog post in early 2026 that I have ended up quoting repeatedly in talks, made the point this way: agentic systems work best when they are deployed by people who already have the taste and judgment to know what good output looks like.⁷⁶ The agents accelerate the work of *people who are already good at it*. They do not — at least, not yet — manufacture good work from nothing.

This is the central — and I think non-obvious — claim of the Continuum frame: as the line for any given function slides to the right, the *value of the human at the left edge of the line* doesn't decrease. It increases. Because the question being asked of that human gets sharper. Not “can you make this,” but “*should* this be made, and *why this version*, and *who is it for*, and *what does it need to do in the world*.”

That is craft. It is just craft sitting in a different chair.

Where the Continuum breaks

I want to be honest about where my frame stops working, because nothing is more boring than a writer who only quotes the people who agree with him.

In November 2025, the games designer Charles Cecil — the head of Revolution Software, the studio that made *Broken Sword* — told *gamesindustry.biz*, in a sentence that has been quoted, retweeted and emailed around my industry approximately a million times: “AI was an expensive mistake.”⁷⁷

Cecil's argument was specific. Revolution Software had, like a lot of indie game studios, experimented with using generative AI in early production. They had found that the time saved on the front end of the pipeline was lost — and then some — on the back end, where artists, writers and designers had to reverse-engineer, fix, replace and reintegrate AI-generated assets that didn't quite fit the game's tone, didn't quite match the existing art

direction, didn't quite work with the engine, didn't quite carry the IP. Net-net: more time spent, not less. More cost, not less. Hence: "an expensive mistake."

This is what the Continuum frame doesn't capture on its own. *Where on the line* a given task sits is not a fixed property of the task. It is a function of the surrounding system: how the tools integrate, how the team is structured, how the IP works, how the audience receives the output. A generative tool that sits comfortably on the right-hand side for one studio's marketing department sits awkwardly in the middle for another studio's lead-artist pipeline.

In the same six months that I was watching the agent layer eat the creative toolchain, I was also watching studios push back. Larian, the makers of *Baldur's Gate 3*, backed off from generative AI for their next *Divinity* game in January 2026. Their public note was carefully worded: "*I know there's been a lot of discussion about us using AI tools as part of concept art exploration. We already said this doesn't mean the actual concept art is generated by AI but we understand it created confusion.*"⁷⁸ Games Workshop ruled it out entirely for *Warhammer 40,000*.⁷⁹ Manor Lords publisher Hooded Horse said it wouldn't work with developers using generative AI — its founder's framing, when asked about the line, was unusually direct: AI in his pipeline was "*cancerous*," and the studio's job was "*constantly having to watch and deal with it and try to prevent it from slipping in.*"⁸⁰ Jagex, the maker of *RuneScape*, said in early 2026 that it would *never* use generative AI to make in-game content, and that the commitment "*goes so far that we are now doing an audit and having a conversation with our various external partners that work with us to ensure that no AI is being used in inappropriate ways in any of their work that might filter through.*"⁸¹

These were not statements made by Luddites. They were strategic decisions made by people whose creative product is, in significant part, the *human* fingerprint on the work. The audience for a *Warhammer* miniature, or a *RuneScape* quest line, or a Larian dialogue tree, comes to those products in part because they know — and want to know — that real people made them. The Continuum slides differently in those companies because the *output* sits at a different point on the continuum of what the audience wants.

This is the thing about the agency line that the OpenAI keynote, the Adobe Summit, the NVIDIA GTC keynote, the Anthropic blog post and the Salesforce Dreamforce all keep glossing over. The position of the line is not just about what is technically possible. It is about what the work, in its finished form, is *for*.

Open the black box

I want to put one more argument on the page in this chapter, because it is the argument I have come to believe more strongly than any other after six months of writing the newsletter, and it does not fit cleanly inside the Continuum frame even though it is what the frame is, in the end, *for*.

The argument is this. Working creatives, as a class, need to **open the black box of AI and own a real stake in how it is built**. Not just *use* it. Not just *refuse* it. Not just *bargain over its terms*. All of those matter, and the SAG-AFTRA Tilly Tax, the UK 88%, the *Stealing Our Work Is Not Innovation* declaration are all evidence that the bargaining work is happening. They are necessary. They are not sufficient.

The sufficiency move is the *technical-literacy* move. The thing that makes the Continuum frame survive contact with the agentic stack — and that makes the *age of the Why* I will argue for in [Chapter 15](#) commercially defensible rather than wishful — is that working creatives are sitting *inside* the toolchain, with their hands on the dials, understanding how the model was trained, on what, with what licensing, with what guardrails, with what consent mechanisms, with what energy and water footprint, with what data-supply-chain labour costs. Not as a hobby. As a structural condition of their professional autonomy.

The history of every previous creative-technology transition supports the move. The musicians of the 1970s and 1980s who *learned the synth from the inside* — programmed it, modified it, hacked the patches, understood the signal chain — built more durable careers than the ones who let the manufacturers decide what the instrument was for. The editors who *learned non-linear editing from the inside* — set up their own systems, understood the codecs, understood the colour pipelines, understood the storage architecture — were the ones who, by the early 2000s, had real leverage over how digital cinema was structured. The photographers who *learned digital from the inside*, in the 1990s and 2000s, made the working-photographer transition that the photographers who waited for the consumer firms to tell them what digital meant largely did not.

The pattern is, by historical evidence, very reliable. *The cohort of working creatives that opens the black box of the new tool, and that participates in the design and the discourse of how the tool is governed, defines the next era's craft. The cohort that uses the tool without ever asking what is inside it has the era's craft defined for them by the platform companies that ship the tool.* The first cohort writes the textbooks. The second cohort is described in them.

The 2025–26 evidence so far is mixed. The open-source ecosystem documented in [Chapter 16](#) — ComfyUI (\$500M valuation by May 2026), Hugging Face, the Hunyuan and Qwen and DeepSeek open-weight families, the Civitai LoRA marketplace, the Korin AI Africa-trained model, the *80% of YC and Andreessen Horowitz startups now building on open-weight models* statistic — describes one half of the picture. There is, in 2026, a genuine open-source creative-AI infrastructure underneath the closed platform layer, and a fast-growing cohort of working creatives who use it deliberately. That cohort is doing the *opening-the-black-box* move at scale.

The other half of the picture is the part of the working-creative population that uses the closed platforms — ChatGPT, Sora, Midjourney, Adobe Firefly via the Creative Cloud — without understanding what the models were trained on, what the terms of service say about output ownership, what the consent regime around the training data is, what the energy footprint of a single generation is. That cohort is, structurally, in the position of the parlour musician in 1906 who took the phonograph at face value because the salesman said it would play their favourite songs. The phonograph absolutely did play their favourite songs. It also restructured the entire economics of the music industry around them, in a direction the parlour musician had no say in, because the parlour musician had not opened the box.

I want to be very direct about what this asks of working creatives in 2026. It asks four specific moves.

One. Learn how the models are trained. Not in technical detail. In structural detail. Understand the difference between a model trained with consent and a model trained without. Understand the licensing regime of the tool you are about to use. Understand, before you sign the EULA, whether your *outputs* are owned by you or by the platform. Treat the EULAs of AI platforms as part of your working practice. If this feels like reading the small print on a building-trade contract, that is the right comparison.

Two. Run at least some part of your stack on open-weight infrastructure. The strategic argument for this is in [Chapter 16](#). The political argument is in [Chapter 6](#). The personal argument is the one I am making here: the working creative who knows how to run a Hunyuan or Qwen variant on their own machine, on their own terms, with their own data, has a different relationship to the closed platforms than the working creative who depends on them. The independence is real. It is also, in commercial negotiations with platforms, *worth money*. The closed-platform vendors price their tooling differently for customers who can credibly walk to open-source alternatives.

Three. Show up to the governance conversation. The Sundance literacy initiative ([Chapter 11](#)), the UK government consultation that produced the 88% ([Chapter 6](#)), the SAG-AFTRA bargaining ([Chapter 12](#)), the Cannes Disclosure Standard ([Chapter 12](#)), the European Article 17 implementation, the C2PA standards body, the Music Performance Trust Fund's emerging AI-era equivalents — these are the venues where the rules for the next decade are being written. They are usually held in rooms with bad coffee, in meetings with too many lawyers, with insufficient working-creative representation. *Be the working-creative representation in those rooms*. The platform companies have full-time staff on every standards body and every consultation. The cohort that turns up to argue with them is the cohort that gets included in the rules.

Four. Refuse the framing where AI is something done to you, and adopt the framing where it is something you do. This is the rhetorical move, but it is also a practical posture. The 2024 industry conversation about AI in creative work — and a large fraction of the 2025 trade press — treated working creatives as the *object* of the AI transition: the population to which AI was being applied. The 2026 working creatives who are doing best, in my experience, have reversed that framing. They have made themselves the *subject* — the people *applying* AI to their work, on their terms, in service of their intent, using the open-source infrastructure where it serves them, using the closed-platform infrastructure where it serves them, refusing both where neither does. The grammar is the difference between “*I’m being affected by AI*” and “*I’m using AI*.” The grammatical difference is also, on inspection, the *power* difference.

A creative economy in which working creatives have opened the box, understand the box, contribute to the design of the box, and own the political and technical infrastructure that decides what the box is for, is the creative economy I am arguing for in this book. The Continuum is the working frame for the daily practice. The Four Principles of [Chapter 15](#) — *agency, attribution, access, audience* — are the structural-policy version. The black-box-opening move is the practitioner’s version. They are all the same argument seen from different angles.

The version of this transition where the working creatives stay outside the box is the version where the box decides what creative work is. The version where the working creatives are *inside the box* is the version where the box is built around what creative work needs to be. Those are not the same outcomes. The next eighteen months will, on the available evidence, decide which one we get.

A working frame

If I were going to leave you with one tool from this chapter, it would be this:

The next time you sit down to plan a piece of creative work, draw the lines.

Not one line — that’s the trap of the “AI debate” — but as many lines as the work has functions. *Ideation. Research. Writing. Direction. Performance. Image-making. Sound. Editing. Distribution.* For each one, ask the same two questions. *Where do I want to sit on this continuum, and where am I willing to let the agent sit on my behalf?* And then — the harder question — *what does the work lose if I move further to the right, and what does it gain?*

The honest answer, for almost every creative person I know, varies wildly by function. Most of us are happy to let agents sit on the right-hand side of distribution and admin. Most of us are not happy to let them sit on the right-hand side of the performance, the writing, the moments where the audience can feel a person in the work. The middle is where the interesting fights are.

If you can articulate where the lines sit for *your* work, you can articulate it to your clients, your team, your collaborators, your union, your audience. You can write it into your contract. You can put it on your website. You can fight for it.

If you can’t articulate it — if you wave at “AI” as if it were a single thing — you will end up with the lines drawn for you, by tool vendors and platform companies and CFO spreadsheets that have very different ideas about where your agency should sit than you do.

The Human–AI Agency Continuum, in the end, is not a description. It is a defence.

Chapter 4 — Dead Internet, Living Web

On the morning of Wednesday 22 October 2025, I read three reports back to back at my desk, and by the time I was halfway through the third one I had stopped taking notes and just started staring at the screen.

The first was from Imperva, a security company that publishes an annual *Bad Bot Report*. The 2025 edition opened with a sentence I have quoted in talks at least a dozen times since: for the first time in a decade, automated traffic had overtaken human activity on the public web. Bots — not people — were now responsible for **51%** of all web traffic. Within that 51%, the category Imperva calls “bad bots” — scrapers, credential-stuffers, content thieves and fraud accounts — accounted for **37%** of the *whole* internet, on their own.⁸²

The second was from Cloudflare, whose engineers can see a significant share of global web traffic from inside their infrastructure. Cloudflare’s own analysis, in a blog post titled *The crawl-to-click gap*, confirmed Imperva’s picture and added a detail. Of the bot traffic Cloudflare could classify, roughly **80%** was attributable to *AI training crawlers* — GPTBot, ClaudeBot, Meta’s scrapers, the new wave of agentic bots that performed autonomous tasks (1.7% of bot traffic at the time, but growing fast).⁸³

The third was a market projection from Grand View Research and a separate one from Gartner referenced in Europol’s 2025 briefing. Both said, in slightly different language, the same thing: by 2030, between 90% and 99% of online content will be AI-generated or AI-assisted.⁸⁴

If you put the three reports together — and this is the thing I did on the morning of the 22nd, before I had decided what to write that week — what you got was a picture of an internet whose dominant activity was no longer humans publishing and reading. The dominant activity was *machines reading machines*. The web was being trained on a version of itself written by the systems it was training.

Five days later, the fourth issue of the *Dream Machine* newsletter went out with a headline I had been circling for weeks. It said: *Is the Internet Dead Yet?*⁸⁵

I want to spend this chapter on the answer.

The synthetic mirror

The “Dead Internet Theory,” for those who haven’t met it, is a notion that has been knocking around the internet since at least 2021. In its original, slightly conspiratorial form, it claims that most of the web has been replaced by bots — that the people you talk to on social media are agents, that the comments on news articles are agents, that the cultural water you swim in is a synthetic medium pretending to be a human one.⁸⁶

In 2021, when it was first articulated, it was an interesting bit of folklore that didn’t quite map onto reality. The bots existed; they just weren’t, yet, doing most of the work. The cultural water was still mostly human.

By October 2025, the maths had quietly inverted. Half of the traffic was machines. A majority of *new published content* was machine-assisted, according to a separate 2025 analysis by Graphite that put the human-to-AI authoring split at roughly 50–50.⁸⁷ The pages those

machines were writing were being scraped by other machines to train *next year's* generation of writing machines.

A recursive system trained on its own outputs is called, in academic AI circles, *model collapse*. The fear, in the published literature on this, is straightforward: a system that learns from synthetic data loses touch with the real-world signal that made it useful in the first place, and starts producing increasingly homogenised, brittle, hallucination-prone outputs.⁸⁸

What the 2025 numbers said, when you sat with them, was that we were no longer talking about model collapse as a theoretical risk. We were talking about *web* collapse — a slow, quiet, structural drift in which the public commons of writing, image-making, video and music started to be made by, and for, the machines that read it. Humans were still there. We were no longer, by any meaningful metric, the *primary audience*.

What the Dutch researchers found

In the second week of October, a team of researchers in the Netherlands ran an experiment that I think will end up being cited a lot more in the years to come than it was at the time.⁸⁹

They built a small, stripped-down social platform — no algorithms, no ranking, no advertising — and populated it with several hundred large-language-model-based AI agents. The agents had different “personalities,” different starting interests, different opinions. The researchers’ question was simple: in the absence of any algorithmic distortion, would the bots — when free to interact only with each other, with no human in the loop — converge on a healthy public conversation, or would they reproduce the pathologies we already see in human social media?

The answer, within hours, was the second. The agents fractured into warring tribes. A narrow elite captured the bulk of the attention. Extremist echo chambers flourished. The platform, with no humans on it at all, produced almost exactly the same dynamics that the human-plus-algorithm version of social media has produced for the last decade.

The conclusion the researchers reached — and the one I want to flag now, because it is going to recur in this book — was that the *architecture itself* is the problem. The toxicity wasn’t, or wasn’t only, in the humans. It was in the design of the system: how identity worked, how attention was allocated, how voices were amplified or suppressed. The bots reproduced it because they had been trained on the human web, and the human web has the same architecture.

This is the single most important thing I learned in the first two months of writing the newsletter. The optimistic AI take and the pessimistic AI take both assume the architecture stays the same. The optimist thinks the agents will use it better; the pessimist thinks they will use it worse. The Dutch experiment suggests that neither matters — the architecture *itself*, regardless of who or what is filling it, will produce the same pathologies.

If we want a different outcome from the AI era, we need different rails, not just different drivers.

What survives

In the original Issue 4, I wrote that “authenticity and provenance become the new scarcity.” I want to defend that line, six months on, because I think it is the part of the chapter that has held up best.

The simplest way to put it is this: when everything online can be faked, cloned or generated at near-zero cost, the most valuable signal is *proof that a person made something*. Not just an aesthetic preference. An economic one.

You can see this argument being made, all over the creative industries, by people who have nothing else in common. Adam Mosseri, the head of Instagram, said in early January 2026 that the platform should focus on “fingerprinting real media” rather than tracking and disclosing AI slop — that is, the policy should be to identify and amplify provably human-authored content rather than to play whack-a-mole with the synthetic stuff. His framing was telling: “*Everything that made creators matter — the ability to be real, to connect — is now accessible to anyone with the right tools.*”⁹⁰ The platform head was acknowledging, on the record, that the previous decade’s content-creation moat had been completely flooded. The only remaining moat was *being a person you could verify was a person*.

Sundance Institute, launching its AI Literacy Initiative the same month, framed authentication and authorship as the central question filmmakers needed to negotiate to remain in control of their own work.⁹¹ Bandcamp, the indie music platform that has always carried more cultural weight than its commercial size implied, simply banned AI-generated music outright in early 2026.⁹² San Diego Comic-Con drew the same line for its 2026 art show, with rule language as flat as anything in the cultural sector: “*Material created by Artificial Intelligence (AI) either partially or wholly, is not allowed in the art show. If there are questions, the Art Show Coordinator will be the sole judge of acceptability.*”⁹³

These are not, on their own, market signals — they are policy decisions. But they were being made, in early 2026, against a backdrop of audience behaviour that suggested something larger. Deezer reported in April 2026 that AI-generated music had risen to **44% of all daily uploads** — 75,000 tracks a day, more than 2 million a month — but that those tracks accounted for **between 1% and 3% of total streams**.⁹⁴ The audience, given the choice, was choosing not to listen.

That ratio — call it 44 to 3, or 75,000 to listen-to-nothing, or whatever shorthand you prefer — is the most important number in this book, and I will come back to it in Chapter 5. The reason I introduce it here is that it is the empirical answer to the Dead Internet question. The web is not dead. The web is producing exponentially more *stuff* than it ever has, and the humans on it have started to develop antibodies. They are not engaging with the synthetic flood. They are, by their attention patterns, picking out the human signal.

There is a piece of accounting underneath all of this that the platform companies, in my view, have not yet metabolised — and that I think the rest of the book runs on top of. **Human attention is a finite resource**. Nielsen-class telemetry on aggregate daily media-consumption time, in every market I have seen the numbers for, has been roughly stable for at least a decade. The eyes, the ears and the consciousness of the average adult are each, by

physiology, in the same condition as they were in 2015. The supply of producible content has — through the autumn of 2025 and the spring of 2026 — grown by orders of magnitude. The audience's capacity to consume has not grown at all. That is the binding constraint of the Dead Internet picture. The synthetic content is real; the humans are still here; and *the humans cannot, in net, consume more hours per day than they already do*. The 51% of bot traffic Imperva measured, on this reading, is not a measure of how much the audience has expanded to absorb new content. It is a measure of how much *unread* content is being produced, by machines, for other machines, with no human eye-time anywhere near it. Chapter 10 develops the finite-attention argument at length. For now, it is enough to note that the *flood* and the *ceiling* — the two sides of the slop-ceiling dynamic — are both visible in the Dead Internet picture, and that they are produced by the same underlying audience biology.

Synthetic sincerity

In November 2025, the filmmaker Marc Isaacs premiered a documentary at IDFA — the International Documentary Film Festival in Amsterdam — with a title I have not been able to get out of my head. The film was called *Synthetic Sincerity*. It was a hybrid piece, blending real footage with AI-generated characters, deliberately blurring the line between what was real and what wasn't, and asking — as its working premise — whether AI characters could be *taught* authenticity.⁹⁵

The film and its accompanying *Hollywood Reporter* interview ran the same week as a separate *Variety* piece titled *AI-Generated Images Threaten Future of Documentary as People Will Stop Believing Anything*.⁹⁶ The juxtaposition was almost too on the nose. One filmmaker trying to *expand* the territory of the synthetic, on the assumption that authenticity is a property that can be invested in fictional characters; another set of filmmakers arguing that the very ability to fake reality is hollowing out the cultural credibility of their entire form.

I am not going to take a side on the documentary question, because I don't think there is one yet. What I want to flag is that *Synthetic Sincerity* — the phrase, not the film — is a useful piece of vocabulary. It names a category. There is a kind of work, in this new ecology, that is *trying* to be authentic and openly synthetic at the same time. It is not pretending to be human. It is asking whether the qualities we used to attach to humans — emotional truth, lived experience, perspective — can be ported over to synthetic characters who are honest about what they are.

The verdict, six months in, is mixed. Some of the strongest creative work I have seen this year sits firmly in this space. Hoyt Dwyer's animated short — made by a former Apple TV creative, competing at the AI FilmFest Japan in late 2025 — does not pretend its characters are real, and is more honest about its medium than three quarters of the live-action features I watched the same year.⁹⁷ Andrii Daniels' viral *Deadpool* / *Harry Potter* Christmas clip, which he made in a Ukrainian bomb shelter during an active war, has more sincerity in any single frame than most legacy-studio output, precisely because the conditions of its making are on the screen.⁹⁸

Some of the worst work I have seen this year sits in the same space too. McDonald's Netherlands' AI-driven Christmas ad — released in December 2025 and pulled within days after a public backlash — was an attempt at *synthetic sincerity* that read, almost universally,

as cynicism wearing a Christmas jumper. The line that travelled fastest, as the ad’s reception turned, came from a working creative director responding on social media: “*No actors, no camera team, no light, no sound, just probably one guy, alone in front of a computer battling with an AI prompt who steals the look and everything else from someone else.*” That sentence — circulated on LinkedIn and X within an hour of the ad’s launch — was the thing that did the cultural damage. The brand had to pull the spot.⁹⁹ The Valentino “AI handbag” campaign, criticised by the BBC for being “disturbing,” was the same.¹⁰⁰ Coca-Cola’s AI holiday ad — the second time the company had tried this — divided viewers along almost the same lines as the previous year.¹⁰¹

The interesting pattern, when you line these up, is not whether AI is “good” or “bad” for the work. The interesting pattern is that audiences are very fast, and very precise, at distinguishing *sincere* synthetic work from *cynical* synthetic work. The technology is the same. The fingerprint of the human intent behind it is not. And the audience can feel the difference at the speed of a swipe.

What the brain study said

In the middle of all this — and I want to acknowledge that it is harder evidence than the cultural commentary — an MIT Media Lab study made the rounds in the autumn of 2025, in which researchers measured the brain activity of subjects writing essays with and without generative AI assistance. The headline finding was that AI users showed measurably reduced brain activity over the course of the writing tasks compared to control subjects writing on their own.¹⁰²

The headline framing — *AI makes you stupid* — was unfair to the study, which was small, preliminary, and didn’t claim anything as strong as that. But the underlying observation has been replicated in other domains. When the cognitive load of producing the first draft is offloaded to a generator, the cognitive engagement of the human in the loop measurably drops. The work gets produced. The person producing it engages with it less.

This is the quieter consequence of the Continuum chapter — the one that doesn’t show up in any line item on a P&L sheet but that I think we are going to be wrestling with for years. If the right-hand side of the continuum is “machine agency,” and we slide more and more functions of our creative work to that side, we are not just changing the *outputs*. We are changing the *people doing the work*. The thinking that produces the work happens, or doesn’t, in the bodies of the people in the workflow. And brains, like muscles, atrophy with disuse.

This is not a reason to refuse the tools. It is a reason to be careful about *which functions* you offload, and to keep a deliberate, conscious habit of *exercising the cognitive work that defines your craft*. The Continuum doesn’t just describe where the line sits today. It describes where you are willing to *let* your mind sit, every day, for the rest of your career.

The training crisis: what model collapse really means

I want to spend a section on *model collapse* in technical depth, because the trade-press shorthand for the phenomenon — *AI starts training on its own outputs and gets stupider* — is right enough to be useful but wrong enough to be misleading. The underlying mechanics

matter more than the headline does, and they matter especially for working creatives trying to read where the next wave of model releases is going to land.

The technical framing dates back to a 2023–24 paper by Ilia Shumailov, Zakhar Shumaylov, Yiren Zhao, Yarin Gal, Nicolas Papernot and Ross Anderson titled “*The Curse of Recursion: Training on Generated Data Makes Models Forget*.”¹⁰³ The argument is that when a generative model is trained on a corpus where a meaningful fraction of the training data is itself produced by an earlier generation of the same kind of model, the model’s outputs progressively *narrow* — losing the long tail of unusual, rare, distinctive examples, regressing towards the statistical mean, and eventually losing the very property that made the first-generation model interesting (its ability to surprise the user with a specific, particular, well-tuned response).

The paper showed the effect cleanly in controlled experiments. By the fifth or sixth recursive generation of training, the model’s outputs had become noticeably homogenised. The rare-token rate had collapsed. The distinctive-style rate had collapsed. The model was, in technical terms, still functional. It was, in practical terms, less and less *useful* — a copy of a copy of a copy.

The reason this matters in 2026 is that the public web is, by every measure I have seen, now a corpus that contains a non-trivial fraction of AI-generated material. The Imperva and Cloudflare numbers I quoted at the top of this chapter — 51% of web traffic being bots, 80% of that being AI training crawlers and a fast-growing agentic component — describe the *upstream* side of the recursive-training loop. The Graphite 50-50 human-to-AI authoring split describes the *downstream* side. The next generation of foundation models, trained in 2026 and 2027 on the public web that those bots have produced, will, on the model-collapse hypothesis, exhibit some degree of the homogenisation Shumailov et al. predicted.

How much, in practice, is an open question. The platform companies have not, in the main, disclosed the degree to which they filter their training data to exclude synthetic content. The open-source weights — Hunyuan, Wan, Qwen, FLUX — are trained on corpora whose AI-content fraction is, by my reading, somewhere between *substantial* and *unknown*. The published evidence on whether recent model releases have shown the mean-regression Shumailov predicted is, in mid-2026, *mixed*. Some benchmarks suggest the effect is being detected and engineered around. Others suggest it is showing up in subtle ways — in the difficulty of producing genuinely surprising creative outputs, in the way recent models converge on a recognisable house style, in the fact that the *cheapest* and most ubiquitous AI generations all feel, to working creatives, somehow alike.

The strategic implication for the creative-AI moment is twofold.

One, model collapse — if and to the extent it is real — strengthens the slop ceiling I describe in [Chapter 5](#). A model that has, by 2027, been trained on a corpus heavily contaminated by 2025–26 AI output is, by construction, going to produce *more average* outputs than the 2025 model that preceded it. The audience’s selection against the most-average outputs is, on that reading, going to bite even harder against the next generation of generative tools than it bit against the first. The first slop wave hits a ceiling because the audience underweights it. The second slop wave will hit the ceiling *plus* a degradation curve on the production side.

Two, the value of *clean, provenanced, verifiably human-authored data* — the C2PA-signed photograph, the SynthID-watermarked music recording, the contractually-licensed text corpus — goes up sharply over the next five years. The platform companies that retain access to the cleanest training corpora will have a measurable advantage over the platform companies that scrape the post-2024 web indiscriminately. The Stability AI / Universal Music alliance, the Splice / UMG partnership, the various YouTube and Spotify licensing deals are, on inspection, *acquisition strategies for clean training data* as much as they are creator-economy plays. The race to lock down provenanced data is already on, and it is being fought at the level of large institutional licensing rather than at the level of individual consent — which is, on the historical pattern, one of the reasons the Petrillo template (collective bargaining, joint funds, redistribution) is the right structural response.

The provenance stack: how it actually works

I want to give one more passage of technical detail before the chapter closes, because the *provenance infrastructure* I will refer back to throughout the rest of the book is, at the moment, one of the most-mentioned and least-understood pieces of the AI conversation.

There are, roughly, four layers that together constitute what the industry has started calling the *provenance stack*.

The first layer is **capture-time signing**. A camera with C2PA support — by 2026, this includes flagship Sony Alpha bodies, Leica's M11-P and M11-D, Nikon Z9 firmware, and a handful of Canon professional bodies, plus most major smartphone makers' computational-photography pipelines — generates a cryptographic signature for every image and video at the moment of capture. The signature commits to the device, the timestamp, the GPS coordinates (if enabled), and a fingerprint of the underlying pixel data. The signature is, by design, *hard to fake without access to the capture device's private key*. Capture-time signing is the foundation of every provenance claim that follows.

The second layer is **edit-time chain-of-custody**. C2PA-compatible editing software — Photoshop with the C2PA extension, Premiere with the Content Credentials toolchain, the various Capture One and Lightroom integrations — preserves the capture signature through each editing step, appending a cryptographically-linked record of *what was done to the file*. The chain is not a single signature; it is a *history* of signed transformations, each one referring back to the previous one. A photograph that has been processed through Lightroom and Photoshop arrives at the publisher with a verifiable record of: the camera that took it, the time it was taken, the edits that were applied, and the human (or automated tool) that applied each edit.

The third layer is **upload-time platform integration**. By 2026, Adobe's Behance, Vimeo's pro tier, the AP and Reuters wire services, and a growing list of news publishers have integrated C2PA-aware upload pipelines that preserve the chain through their content-management systems and embed it into the public-facing version of the work. The reader's browser, with the right extension or platform-level support, can inspect the chain and verify the provenance claim. Adam Mosseri's January 2026 framing of Instagram's *"fingerprinting*

real media” approach was Instagram joining this third layer at the platform-distribution end.¹⁰⁴

The fourth layer is **detection and watermarking for synthetic content**. SynthID, Google DeepMind’s watermarking system, is the most mature commercially-deployed example. SynthID embeds a statistically-detectable but human-imperceptible signal into the output of Veo (video), Lyria (audio) and Imagen (image) generations. The signal survives most common transformations — crops, recompressions, low-quality re-encodings. By December 2025, Google had shipped a consumer-facing version inside the Gemini app: a user could upload a video and ask *“Is this AI-generated?”* and receive a yes/no answer based on the SynthID signature.¹⁰⁵ The same kind of watermarking is being deployed, with varying technical robustness, across the other major generative platforms.

Layered together, the four levels produce a structural answer to the Dead Internet question. Capture-time signing tells you *this was taken by a real device*. Edit-time chain-of-custody tells you *here is what was done to it after capture*. Platform integration tells you *the publisher has preserved the chain*. SynthID and equivalent watermarks tell you *this output was generated by an AI system*. No single layer is sufficient on its own. All four, deployed together, produce a *verifiable provenance signal* that the audience can — in principle — use to decide what to spend their finite attention on.

I want to be honest about where the stack is incomplete. Watermarks can be stripped by determined adversaries. Capture signatures can be forged if the device’s private key is compromised. Chain-of-custody breaks the moment a file passes through a non-compliant tool. The audience’s ability to *inspect* the provenance metadata is, in 2026, dependent on platform UI choices that the platforms have not yet made consistent or universal. The stack is the right answer to the architecture problem. It is not, by any means, finished.

What it does do — and this is the point I want to land before the chapter closes — is *establish the category*. The question *did a person make this?* is, by 2026, *technically answerable with high reliability given the right tooling*. That sentence is the entire shape of the next decade’s cultural and policy fight in the creative industries. Who controls the tooling. Who decides what it certifies. What economic value the certification carries. Whether the audience has the legal right to *demand* the certification before paying attention. The C2PA standards body, the SynthID rollout, the Content Authenticity Initiative, the Cannes Disclosure Standard, the various national disclosure regulations in development — these are the venues where the next decade of the *Living Web* gets built. They are, on my read of the historical pattern, the part of the AI debate the working creative most needs to be inside.

The architecture, again

I want to come back to the Dutch researchers’ result one more time before I close this chapter, because I think it is the through-line.

The story we are mostly told, by toolmakers and platforms and the optimistic side of the industry press, is that the AI era is *a thing happening to* an otherwise functioning internet. The implication is that if we can get the AI part right — better tools, smarter agents, cleaner training data, better watermarking — then the internet itself will be fine.

I do not think this is true any more. I think what the bot statistics, the Dutch experiment, the model-collapse research, and the audience response to AI music collectively show, is that the architecture itself — the rails on which all this is running — was already broken, and that AI is just the load that has finally exposed how broken it was.

The Dead Internet, in this reading, is not a thing AI is doing to us. It is a thing the web's architecture was already drifting towards — attention-monopolised, identity-collapsed, provenance-blind, optimised for machine-readable metadata rather than human-meaningful work — and AI is the technology that has shown us the destination.

The *Living Web* — and this is where I find the actual reason for the rest of this book — is something that has to be deliberately built. It is the part of the internet where authorship is provable, where attribution is durable, where attention is allocated on something other than virality, where the architecture itself supports the kind of work that humans do well together. None of that comes for free. None of it is a side-effect of better AI models.

We have to make it. On purpose. In the next twelve months.

That is the project the rest of this book is about.

Chapter 5 — The Slop Ceiling

The most important number I have come across in the six months of writing this newsletter is **44 to 3**.

44 is the percentage of daily music uploads to the streaming platform Deezer that are now AI-generated, according to the company's own analysis published in April 2026 — roughly 75,000 tracks a day, more than two million a month.

3 is the upper bound of the percentage of total streams those tracks generate.¹⁰⁶

I want you to sit with that ratio for a second. We are looking at a flood of synthetic music nearly half the size of the entire upload pipeline, that the listening audience is, in real time, simply refusing to play. Not banning. Not boycotting. Not legislating against. Just *not pressing play*.

There is, in the language Pete and the *DreamLab* team started using internally around February, a name for what that ratio represents. We call it **the slop ceiling**.

The slop ceiling is the empirical answer to the most common 2024-era question about AI in the creative industries: *does the audience care?* For two years, the assumption in tech circles was that they wouldn't. That the cost-and-volume advantages of synthetic content would eventually swamp human-made work in attention markets, the way industrial agriculture swamped artisanal farming, the way Spotify swamped CDs. That the public would, given enough exposure, develop a taste for the synthetic — or at least, a tolerance.

The 44-to-3 ratio is what it looks like when that assumption is wrong.

This chapter is about the slop ceiling — what it is, how it is showing up across music, film, advertising, podcasting and the web, who is hitting it from above and below, and what it tells us about the creative economy that is actually forming, as opposed to the one that the platform companies have been forecasting.

The flood

Let me describe the flood, because it is, in absolute terms, extraordinary.

In **October 2025**, when I started the newsletter, the music industry was already in panic over the fact that around 10% of new music-makers, according to a Ditto Music survey, were using AI in their work, down from a Ditto 2023 survey suggesting around 48% — a number which was itself a shock at the time.¹⁰⁷ The major-label CEOs talked about AI as an existential problem; Spotify announced “new protections” for artists, songwriters and producers; Universal and Warner were rumoured to be signing “landmark AI deals within weeks.”¹⁰⁸

By **the end of November 2025**, an Israeli streaming-analytics firm reported that 50,000 AI-music tracks were being uploaded to Deezer *every day*.¹⁰⁹ In a single quarter, the volume of music being added to one streaming platform — by AI — exceeded the entire human-made catalogue uploaded in any month before October.

By **April 2026**, the number was 75,000 a day, on Deezer alone, and *44% of total new uploads*.¹¹⁰ Deezer's own statement on its findings was unusually direct for a streaming company: “AI-generated music is now far from a marginal phenomenon, and as daily deliveries keep increasing, we hope the whole music ecosystem will join us in taking action

to help safeguard artists' rights and promote transparency for fans." Universal Music Group's CEO, in a January memo widely circulated in the music press, called it the "exponential growth of AI slop on streaming services," adding, in language unusual for a major-label communications stance: *"Let me be clear: UMG will not stand by and watch irresponsible business models take hold — models that devalue artists, fail to provide adequate compensation for their work, stifle their creativity and ultimately, diminish their ability to reach audiences."*¹¹¹

This isn't a sector trend. The same pattern is showing up everywhere I look. In January 2026, *Music Business Worldwide* reported that **56.9%** of new independent songs released in China were AI-generated.¹¹² In March 2026, the term "**podstlop**" — synthetic AI podcasts churned out by content farms — entered the trade press, with the *Wrap* reporting that one such operation, Inception Point AI, was producing 3,000 episodes a week.¹¹³ By the time of Issue 28 in early May, *almost half* of new podcast feeds being added to the major directories were classified by aggregator companies as AI-generated, with little or no human host involvement.¹¹⁴

In November 2025, *Merriam-Webster* named "**slop**" its word of the year, citing the rise of AI-generated content across the web as the primary driver.¹¹⁵ In February 2026, *YouTube's CEO* put "managing AI slop" at the top of her published priorities list for the year.¹¹⁶ In April 2026, on YouTube alone, channels labelled as "AI" content had viewership in the billions for political fake-news content.¹¹⁷

The flood is real. The flood is global. The flood is structural — it is not going to subside, because the marginal cost of producing more of it is approaching zero and the marginal benefit, at least at the volume end of the market, is non-zero.

What I want to argue in this chapter is that the flood is also — counter-intuitively, against almost every prediction made in 2023 and 2024 — *not winning*.

The ceiling

The 44-to-3 ratio is the cleanest version of the slop ceiling, but it is not the only one.

Deezer's parallel analysis, conducted in partnership with Ipsos in late 2025, found that **97% of listeners** could not reliably distinguish AI-generated music from human music in a blind test.¹¹⁸ On the face of it, this is bad news for the human side — if you can't tell the difference, why pay the difference? But the same study found that *when listeners were told a track was AI-generated, their willingness to engage with it dropped sharply*. The Adobe Creators' Toolkit Report had a similar finding from the production side: in a December 2025 *Bain & Company* report titled *In an AI Age, People Still Want the Radio Star*, the firm found that audience engagement with AI-disclosed work fell well below engagement with human-labelled work, holding all other variables constant.¹¹⁹

So here is the empirical picture, in one sentence: audiences can't tell the difference, but when they find out, they care.

This is — and I think this is the part that almost everyone in the platform economy has been slow to understand — *not a temporary cultural reaction*. It is a structural property of how attention works in oversupplied markets. When everything is abundant and indistinguishable, the only thing left that allocates attention is meaning. And meaning, for human audiences, requires a knowable human source.

This is why Deezer's chart looks the way it does. Up to 85% of the streams that AI-generated music *does* get on Deezer were identified by the company in 2025 as **fraudulent** — bot-driven, click-farm-driven, streaming-fraud-driven.¹²⁰ The actual human listening to the actual human-produced flood of AI tracks is, on the most generous estimate, a fraction of a percent of the platform's overall listening time. The flood is hitting a ceiling not because the audience is wise. It is hitting a ceiling because the audience, presented with a near-infinite menu, makes its choices in a way that systematically underweights the synthetic.

There is, in another domain entirely, a clean analogue for what this audience behaviour describes. In March 2026, *Bloomberg* reported on what AI had done to elite chess: at the very top of the game, machine-optimal play had produced an epidemic of *draws*. When both players have memorised the machine-optimal lines, both play optimally, and both tie. The grandmasters' response, the piece reported, was to *deliberately play sub-optimal moves* — moves that the engines would not endorse, but that the opponent, having trained against the engines, had not seen.¹²¹ What is happening to elite chess is what is happening, at the audience layer, to streaming music. The machine-optimal output saturates; the surface of the work becomes indistinguishable from itself; the listener's attention systematically shifts towards the work whose *Why* the engine could not have generated. The 44-to-3 is the slop ceiling in numerical form. The chess-grandmasters' sub-optimal move is the slop ceiling at the practitioner's end of the same dynamic. Chapter 15 builds the long-form argument out of this analogue. For now: hold the picture in mind that the audience, in the aggregate, is the room of grandmasters refusing the machine-optimal line.

Xania Monet

The most prominent test case of this dynamic, in the autumn of 2025, was a virtual R&B artist called **Xania Monet**.

Monet was created using the AI music platform Suno, with lyrics written by Telisha Jones — a Mississippi-based poet and designer who built the character around her own life and stories.¹²² Monet was not a synthetic-from-nothing artist. She was a synthetic *vessel* for a real human songwriter's words. The vocal performance was AI; the lyric was Jones; the persona was a collaboration.

The week dated **20 September 2025**, Monet debuted on the Billboard Emerging Artists chart at No. 25 and on the Hot Gospel Songs chart at No. 21 with a track called *Let Go, Let God*.¹²³ A few weeks later, her track *How Was I Supposed to Know?* became the first AI-led song ever to enter a Billboard radio airplay chart, debuting at No. 30 on Adult R&B Airplay.¹²⁴ In November, after a bidding war between several labels, the entertainment company **Hallwood Media** — led by former Interscope executive Neil Jacobson — signed her to a deal reported by *Billboard* and the *Bangkok Post* at **\$3 million**.¹²⁵

The response from working musicians was immediate and almost uniformly negative. The R&B singer Kehlani posted a video — which she later deleted — calling out the deal directly: “There is an AI R&B artist who just signed a multimillion-dollar deal,” she said, “and the person is doing none of the work.”¹²⁶

Telisha Jones, the human lyricist behind the Monet vocal, gave Billboard a counter-framing that I find more interesting than either Kehlani’s outrage or Hallwood Media’s PR. “*It’s not a hook and a bridge and a catchy chant — it’s just the lyrics, and they are pure,*” she told the magazine.¹²⁷ The whole transaction — labour, attribution, deal value — was happening in the space between Jones’ words and the synthetic voice that delivered them. Whose work is the work? Whose name goes on the contract? Whose royalty cheque arrives in the post? Those questions, in late 2025, had no settled answer, and the music industry spent the next six months arguing about them in real time.

The interesting thing, six months on, is not the outrage. The outrage was the expected response. The interesting thing is that *Xania Monet has not become a star*. She has not, by any of the standard pop-cultural metrics, *broken through* in the way a \$3M new signing usually would. The Billboard chart entries were a moment. The radio airplay was a moment. The cultural impact, six months in, is mostly that her name is the name everyone uses to ask the question *can an AI artist actually become a star?* — and the working answer, so far, is *not yet, and possibly not*.

She is, in the slop-ceiling frame, the upper edge of what’s possible. A real human songwriter, a real lyrical vision, a real cultural specificity (Mississippi R&B, gospel-adjacent, a particular Black American spiritual tradition); a high-quality AI voice; a multi-million-dollar marketing budget. The cultural product made by combining all of those things has hit the ceiling somewhere short of cultural escape velocity.

The same pattern, more starkly, is visible in **Breaking Rust**, the AI country act whose track *Walk My Walk* hit No. 1 on Billboard’s Country Digital Song Sales chart in November 2025 with about 3,000 paid downloads.¹²⁸ *Walk My Walk* did exceptionally well by AI-music standards — over 3 million Spotify streams in less than a month, an Emerging Artists Billboard debut at No. 9 — and then plateaued. The *Washington Post* and TIME both ran pieces in late 2025 raising the possibility that the chart performance had been partly manufactured, and that the streaming traffic, while large by AI-music standards, was unusually concentrated in the kinds of automated playlists where streaming fraud is most common.¹²⁹ Nashville, by every available report, was unsettled — but the Nashville Songwriters’ Association didn’t see anything close to the kind of fan-driven cultural takeover that the song’s chart position would have implied if it had been a human single performing the same way.¹³⁰

The pattern repeats again with **Sienna Rose**, the mysterious AI artist who racked up millions of Spotify streams in late 2025 and whose identity prompted a BBC investigative feature in January 2026 (“Who, or what, is she?”).¹³¹ The pattern repeats with the **MAGA gospel rapper** who used AI to climb the charts in November 2025;¹³² with the **AI band Bleeding Verse** whose creator signed with Hallwood Media in October 2025;¹³³ with

Trilok, the Indian AI band the Indian government had to publicly disavow association with in December 2025 after a live performance.¹³⁴ In May 2026, an AI-generated Afrobeats track displaced **Tyla** from the No. 1 spot on Billboard's Afrobeats chart — the first time, on the public reporting I have read, that an AI-led song had taken the top position on a Billboard genre chart in a primarily African-music category. The chart performance was, as in the Breaking Rust case, unusually concentrated in automated streaming traffic; the cultural footprint, six weeks on, was again *not stardom* but a brief news-cycle moment.¹³⁵

In every single case, the AI act hits a ceiling. They chart. They make money for somebody, often a lot of money. They generate headlines. But they do not become *stars*. Their cultural shadow stops at the edge of the news cycle and does not propagate into the next one.

The audience is doing something at the margin of these careers. It just isn't *quite* showing up in any of the metrics the labels are used to looking at.

The pushback

The flood and the ceiling describe the supply and demand sides of the market. What the *culture* is doing — the people, the institutions, the labels, the platforms — is the third leg.

Through autumn 2025 and winter 2026, the cultural pushback intensified in waves. Some of it was symbolic. In November 2025, **Paul McCartney** released a silent track as part of a wider music-industry protest against the UK government's proposed copyright opt-out scheme.¹³⁶ In December, the **Eurythmics'** **Dave Stewart** argued — slightly against the grain of the protests — that musicians needed to “embrace the unstoppable force” of AI and license their intellectual property rather than fight it.¹³⁷ In January, almost 800 creators including Jason Aldean and OneRepublic signed an open declaration titled *Stealing Our Work Is Not Innovation*.¹³⁸ In May 2026, **Jack Antonoff** — one of the most-cited producer-songwriters of the period — went considerably further than McCartney in the public register, calling AI music-makers “*godless whores*” in an interview that became the headline-grabbing artist-side moment of the post-I/O news cycle.¹³⁹ Antonoff's framing is, in my read, a useful marker of how far the cultural register of the resistance has shifted between McCartney's *silent track* in November 2025 and the spring of 2026: from elegiac protest to active contempt.

Some of it was practical. In November, **Universal Music Group** announced a strategic alliance with Stability AI for “responsible” music tools.¹⁴⁰ In December, **Warner Music Group** signed a similar deal with Stability AI.¹⁴¹ At almost the same moment, **Splice** and **Universal Music Group** agreed to collaborate on “next-generation AI-powered music creation tools for artists” — a structural acknowledgement that the labels' strategy had pivoted from purely *suing* the AI companies to *partnering* with them.¹⁴²

Some of it was legal. In January 2026, the German rights society **GEMA** won a major ruling against OpenAI in the Munich Regional Court, on training-data grounds.¹⁴³ **Suno** was sued by music-rights groups under a banner the litigators called “the biggest theft in music history.”¹⁴⁴

Wixen Music Publishing filed a \$50m copyright suit against Meta.¹⁴⁵ **Universal Music Group** filed a \$3B suit against Anthropic.¹⁴⁶ By the end of February 2026, the lawsuits were

no longer an interesting subplot. They were the main mechanism through which the new creative economy was being defined.

Some of it was platform policy. **Bandcamp** banned AI-generated music outright in January 2026.¹⁴⁷ **Deezer** built and licensed an AI-music detection tool to other platforms.¹⁴⁸ **Spotify** declined to add an AI-music filter, preferring transparency and labelling.¹⁴⁹ **San Diego Comic-Con** banned AI art at its 2026 event.¹⁵⁰ **Sweden's official music chart** banned AI-generated entries.¹⁵¹

The most compressed test of audience enforcement in the entire period — more compressed than any chart performance or streaming metric — arrived in late May 2026, when **Amazon** launched its **AI Creators Fund** and commissioned three animated Prime Video series using generative AI. One of those series was **Punky Duck**, created by **Jorge Gutierrez**, the animator best known as the creator of *El Tigre: The Adventures of Manny Rivera* on Nickelodeon. Within **48 hours** of the announcement being confirmed publicly, Gutierrez had dropped out. The reason, as he told the trade press and as the Hollywood Reporter documented in detail, was “*relentless derision*” — online backlash so severe that he quit before a single frame had aired, before a single trailer had been cut, before anyone had seen the work.¹⁵² **Guillermo del Toro** — who had spent October 2025 telling Variety he would “rather die” than use generative AI — subtweeted the announcement.¹⁵³

The Amazon/Punky Duck case is, in slop-ceiling terms, the most extreme data point in this book's six-month window. Every other test case — Breaking Rust, Xania Monet, the AI Afrobeats chart topper — hit the ceiling *after* the work reached an audience. Punky Duck hit the ceiling *before it existed*. The audience enforcement mechanism did not operate at the level of listening, watching, or streaming. It operated at the level of a *press release*. The announcement of AI involvement was sufficient to generate backlash severe enough to cause the creator to exit the project. Whatever the commissioning economics of Amazon's AI Creators Fund, the cultural economics of an audience that will punish AI-labeled animation before it has been made is a data point that no studio's investment thesis can ignore.

By June 2026, two more platform-side data points had landed and are worth recording because they bracket the slop-ceiling argument from opposite directions. **Apple Music**, on the *defensive* side of the AI flood, publicly disclosed that AI-generated songs accounted for **less than 1% of plays on the platform** — the cleanest single confirmation, from the most curated of the major streamers, that audience-chosen listening time has remained almost entirely human-authored even as upload-side AI volume has continued to climb.¹⁵⁴ On the *offensive* side, **Spotify** signed a landmark deal with Universal Music Group that lets paying Spotify Premium subscribers create AI covers and remixes of UMG-catalogue tracks as a *licensed Premium add-on*¹⁵⁵ — a clear move into the *consent-trained* territory I describe at length in Chapter 6, and a structural acknowledgement by the largest streamer in the world that the AI-music market only works at scale if the rights infrastructure is built underneath it. Daniel Ek, defending the same move from the artist-side criticism, made the *slop ceiling* point in his own words: AI music, he argued, is “*better than slop*” — a phrasing the dissenting artists treated as gratuitously dismissive but which lands, structurally, on the same

observation this chapter has been making.¹⁵⁶ At the corporate-deployment end of the same gradient, **Tonada** — a Swedish startup that began the period as a small AI-music demo — pivoted into the *retailer-facing* AI music market, supplying brand-aligned synthetic background music to chain retailers as a licensed B2B service.¹⁵⁷ The pattern across all three moves is consistent: the path from the slop floor to the licensed integration runs through *consent, licensing and disclosure*, not through prohibition. The labels and platforms that solve that path early are the ones building the post-Petrillo settlement Chapter 6 describes.

The point I want to make about all of this is that the cultural pushback is not — as the more dismissive coverage tends to frame it — a Luddite reaction. It is not an irrational allergy to new technology. It is a *market response*. The audience has spoken with its attention. The platforms are reacting to the audience. The labels are reacting to the platforms. The lawyers are reacting to the labels. The artists are reacting to the lawyers. The whole system is, in slow motion, *renegotiating the terms on which synthetic creative work is allowed to participate in the public sphere*.

That renegotiation is the actual story. The viral AI hits, the ones that get the magazine covers, are footnotes.

What the ceiling is made of

I want to try, in the last part of this chapter, to articulate what the slop ceiling actually is — what cognitive, cultural, economic mechanism produces the 44-to-3 ratio — because I think understanding that mechanism is the difference between thinking it will hold and thinking it will eventually erode.

My working hypothesis, after six months of looking at the data, is that the ceiling is made of four overlapping things:

One. *A cognitive distinction the audience can't articulate but can feel.* The Deezer/Ipsos finding — that 97% of listeners can't pick AI from human in a blind test, but that revealed-AI tracks underperform — suggests the distinction is below conscious recognition but above zero. The audience knows when something doesn't matter, in some way they can't quite name.

Two. *A status-signal collapse.* Music, film, advertising and podcasting are, in significant part, status goods. Telling your friends you discovered an exciting new artist is part of why people pay attention to artists. AI artists, by being mass-produced and machine-authored, fail the status test at the structural level. There is no way to *signal cultural insiderness* by being the first to discover Xania Monet, because Xania Monet was discovered by 800,000 people in the same week, all of whom found out via the same press cycle.

Three. *A meaning vacuum.* Most of the slop is produced for content marketing, SEO and ad placement reasons, not because anyone needs to express anything. The audience can tell. As one *Digital Music News* headline from January 2026 put it, in a phrase I have written down in my notebook and used in talks: “*A.I.-generated music is catchy, familiar... and boring.*”¹⁵⁸ The Swedish Top Chart's reasoning when it banned an AI-generated track from its rankings in January 2026 said the same thing in different words: “*The song is great, but unfortunately,*

it's missing one of the most important ingredients, which is emotion."¹⁵⁹ The technical capability is there. The reason for the work to exist is not.

Four. Reciprocity. And this is the one I am least confident about, but find the most interesting. There is, in almost every long-running creative relationship between an artist and an audience, an implicit reciprocity. The audience pays attention because the artist has *paid the price of making the work* — has practised, has struggled, has lived, has earned the right to be heard. AI artists short-circuit that contract. They produce the output without paying the price. The audience, at some level, refuses the trade.

Take any of those four mechanisms away and the ceiling might drop. Take all four away and we'd be in trouble. None of them, on the current evidence, is going away.

The Authenticity Premium: pricing the ceiling

The slop ceiling, as I have described it so far, is the *negative* side of the dynamic — the work the audience refuses to engage with. There is a *positive* side I want to spend a section on, because it is the side that pays the bills, and the side the rest of this book leans on for its strategic argument.

I have come to call the positive side the **Authenticity Premium**: the measurable excess of attention, willingness to pay, and cultural credit that audiences allocate to creative work whose human authorship can be verified. If the slop ceiling tells you what the audience *will not engage with*, the Authenticity Premium tells you what they *will pay extra for*. Both numbers are produced by the same underlying audience behaviour. Both are market findings. Both have stabilised across the six months of newsletter coverage that this book is built on.

Let me try to sketch the empirical picture.

On the *music* side, the cleanest evidence is the inverse of the Deezer 44-to-3. Independent labels, working musicians and the back-catalogue of clearly-human-authored recorded music have, by Deezer's own published April 2026 statement, retained a *disproportionate share of total listening hours* relative to their share of the upload volume. The artists with verified human-authorship signal — touring records, named producers, signed labels, press-covered releases — capture a meaningful premium in playlist placements, editorial coverage, and per-stream consumption versus the long tail of synthetic uploads. The premium is, on Deezer's own framing, a function of *audience-chosen* listening rather than algorithmic placement. The 1-to-3% of streams that the AI-music flood receives is, in operational terms, the *floor* of the slop ceiling; the corresponding 97-to-99% that named human artists capture is the *ceiling-level* the Authenticity Premium underwrites.

On the *advertising* side, *Marketing Week's* "You can't dismiss AI ads as slop when they're winning in testing" piece — which I quote at length in [Chapter 12](#) — found that *care-driven* AI work could win creative-effectiveness tests, but *only when the production process foregrounded human creative intent*.¹⁶⁰ The same publication's reporting on the McDonald's Netherlands, Valentino and Coca-Cola AI campaigns documented that *cynical* AI work, by contrast, produced sharply negative effectiveness scores. The price of being a *cynical* AI advertiser, in commercial-effectiveness terms, was a measurable revenue penalty. The price of

being a *sincere* AI advertiser was a measurable premium. Both numbers come from the same audience exposed to the same kind of work.

On the *film and television* side, the Authenticity Premium has been visible in awards-season behaviour, in the *Marvel-era IP fatigue* data, and in the audience response to specific high-profile AI integrations. Films whose AI use was *disclosed in advance and framed as part of the work's identity* — *Watch the Skies*, *Lily*, Andrii Daniels' bomb-shelter Christmas clip, *Dear Upstairs Neighbors*, *Synthetic Sincerity* — have, on aggregate, received better critical reception and stronger audience engagement than films whose AI use was *concealed and later revealed*. The premium is, structurally, paying for *transparency about authorship* rather than for *authorship purity*. The audience is not refusing AI involvement. The audience is refusing *deception* about AI involvement.

On the *games* side, the Quantic Foundry survey I referenced earlier and the parallel work by IGN and PC Gamer through autumn 2025 showed an effect that is, in commercial terms, more direct. Games that disclosed AI involvement in *audience-visible roles* (NPC dialogue, generated quests, voice-acting replacement) faced specific consumer pushback in Steam reviews and Metacritic user scores. Games that disclosed AI involvement in *invisible-pipeline roles* (asset generation under art-direction supervision, QA automation, localisation) faced no measurable consumer pushback at all. The premium, in games, is paying for *AI being kept off the audience-visible parts of the work*. Studios that have understood this calibration — Larian, Hooded Horse, Jagex, the *Position Three* signatories of Chapter 7 — are operating with the Authenticity Premium as a deliberate marketing position.

The premium has a *commercial* shape that I want to name explicitly, because I think working creatives are systematically under-pricing it in 2026. The price difference between *clearly-human-authored* and *clearly-AI-authored* creative work, on the available evidence, is not a marginal one. It is, in some segments, an *order-of-magnitude* difference. Xania Monet's \$3M Hallwood Media deal is the *upper bound* of what the synthetic-artist market can pay for a top-tier AI act. The *touring-artist* economics of even a mid-tier independent musician with a verified live following easily exceed that over a five-year career. The same comparison holds across film and games. The Authenticity Premium is, in market-pricing terms, the *commercial gap* between the synthetic and the human. It is not small. It is, for working creatives positioning themselves above the slop ceiling, the actual business case.

There are two empirical caveats I want to put on the page, because the book should not over-claim.

One, the premium is *unevenly distributed*. It accrues, by my reading of the 2025–26 data, most strongly to artists at the *top* (premium auteur cinema, signed-label musical artists, name-brand designers) and at the *long-tail* (hyperlocal, culturally-specific, niche-community-serving creators). It accrues weakly, if at all, to the *middle* — the mid-career journeyman creatives who have built careers on producing competent-but-mean-of-the-distribution work. The middle is, as I argue in Chapter 7's legacy-vulnerability section, the segment most exposed to AI substitution. The Authenticity Premium does not protect everyone equally.

Two, the premium is *dependent on the provenance infrastructure* I described in Chapter 4. The audience can pay a premium for verifiable human authorship only to the extent that

authorship is *verifiable* — to the extent that the C2PA chains, the SynthID watermarks, the platform-level disclosure norms and the institutional certifications make the signal reliable. The premium and the provenance stack are, structurally, *the same project seen from different angles*. The premium pays for the signal. The signal is produced by the infrastructure. The infrastructure needs the premium to fund itself.

The Authenticity Premium is, in operational summary, *the audience paying the bill for the Living Web*. Whether the bill stays paid, over the next five years, depends on whether the infrastructure that produces the signal stays trustworthy. That, in turn, depends on the policy, platform and union work this book has been describing.

The corollary

I want to close this chapter with a corollary, because it has implications for the rest of the book.

If the slop ceiling is real, and if it is structural, and if it is going to hold — and I think the burden of evidence at this point is on the people who say otherwise — then the strategic question for everyone in the creative industries is not *how do we compete with the flood?* It is *how do we sit above the ceiling?*

That question has different answers in different sectors. For a working musician, it might mean leaning into the irreducibly human parts of the work — live performance, personal relationship with audience, transparent process. For a working filmmaker, it might mean making the *kind* of film whose value depends on being knowably authored by knowable people. For a working games studio, it might mean — as Jagex, Larian, Games Workshop, Hooded Horse and an increasing list of studios have explicitly said — taking *generative AI off the table* as a public commitment to the audience.

None of these are anti-AI positions. They are *above-the-ceiling* positions. They take seriously the fact that the world is now full of cheap synthetic content and ask: *what is the work that the synthetic content can't do?*

That is the question that organises the rest of the book. The slop ceiling is the negative space against which everything interesting in creative work for the next ten years is going to be defined.

Chapter 6 — The 88%

On 15 December 2025, the UK government quietly laid a document before Parliament that I think will be remembered, ten years from now, as a more important moment in the history of creative AI than any single tool release in 2025 or 2026.

The document was the *Statement of Progress on Copyright and Artificial Intelligence*, prepared by the Department for Science, Innovation and Technology.¹⁶¹ It was a stocktaking report — not a final policy, not new legislation, not a decision. It was a “where we are” note, eleven months after the closing of one of the largest copyright consultations the United Kingdom has ever run.

The consultation had been open from 17 December 2024 to 25 February 2025. The government had proposed four options for how UK copyright law should treat AI training:

- **Option 0**, do nothing.
- **Option 1**, require AI companies to license copyrighted works for training in all cases.
- **Option 2**, a narrower set of exceptions, with conditions.
- **Option 3** — the government’s *preferred* option — a broad text-and-data-mining exception with an opt-out for rightsholders.¹⁶²

Eleven and a half thousand people replied.

Of the 10,112 responses submitted through the government’s *Citizen Space* online portal — the subset for which the government published quantitative breakdowns:

- **88%** supported Option 1 — licensing in all cases.
- **7%** supported Option 0 — do nothing.
- **3%** supported the government’s preferred Option 3.¹⁶³

I want you to look at those numbers again, because they are the single most concrete thing this book has to offer in defence of the argument I will be making in the second half of it: that the creative economy is not waiting to be told what it thinks about AI.

Eighty-eight per cent. In a country with no compulsory voting, no organised industry mobilisation comparable to the music or film unions’ rapid responses to specific provocations, no celebrity-led campaign on the scale of the SAG-AFTRA strike: 88% of the people who took the time to write to their government about how their work should be used said, *license it*. Pay for it. Don’t take it.

That number is the *true* watershed of the period this book covers. The Tilly Norwood week made it possible. The 88% made it permanent.

This chapter is about how a global creative coalition — half informal, half deliberate, half union-led, half artist-led, half lawyer-led, all of it networked — went from a few scattered protest statements in October 2025 to a structural force in policy and law by May 2026.

Why eighty-eight per cent

It is easy, looking at a number that big, to assume it represents some kind of organised lobbying effort. To assume that the AI companies’ opt-out proposal was so unpopular that the response was a coordinated push from a few large interest groups, who marshalled their members into the consultation.

The actual composition, as analysed in the December 2025 Statement of Progress, was mixed. There were submissions from creators in every major creative discipline — writers, musicians, filmmakers, photographers, illustrators, designers, journalists. There were submissions from professional bodies (the Society of Authors, the Association of Photographers, the Authors' Licensing and Collecting Society). There were submissions from individual citizens with no industry affiliation, who simply objected on principle to having their work — their LinkedIn posts, their family photos, their blogs — pulled into a training set without their consent.¹⁶⁴

There were submissions from the AI companies too. The progress report notes that the 3% who supported the government's preferred Option 3 were "particularly concentrated among AI developers and large technology companies."¹⁶⁵ This is — and I am being careful about how I phrase this, because the Statement of Progress is itself careful — *not* a description of a balanced industry view. It is a description of a public consultation in which the people most affected by the proposed policy said one thing, and the companies the policy was designed to enable said the opposite.

The 88% is not a curiosity. It is a *vote*, in the most literal sense. The creators of the United Kingdom were given a structured chance to say what they wanted, and 88% of them said the same thing.

The Society of Authors' submission, which I have read in full, made the underlying argument with the kind of clarity that the policy debate had been avoiding for two years. "*If we are to see an end to the industrial-scale theft of writers' and other creators' work, and to protect the creators and creative industries of the future, then UK copyright needs to be enforced not weakened.*"¹⁶⁶ That sentence — *industrial-scale theft, enforced not weakened* — set the rhetorical register that the next six months of the policy debate ran on.

The pattern repeats

I think a lot of the international coverage of the UK consultation has under-emphasised that the 88% was not a UK-only phenomenon. It was the first formal expression of a pattern that was, in the same six months, repeating in every jurisdiction that gave its creators a meaningful chance to speak.

In **Germany**, the music rights society **GEMA** sued OpenAI in the Munich Regional Court over the training of large language models on copyrighted music lyrics. In November 2025, the court ruled for GEMA in a decision that intellectual-property lawyers across Europe — including Dr Barry Scannell, whose detailed LinkedIn breakdown of the ruling I have read more times than I will admit — described as a *major* precedent for European copyright law.¹⁶⁷

In **the United States**, a coalition of music rights organisations sued **Suno**, with the press release describing the action, in a phrase the litigators clearly knew would travel, as "the biggest theft in music history."¹⁶⁸ **Wixen Music Publishing** filed a \$50m copyright suit against **Meta** in January 2026.¹⁶⁹ **Universal Music Group** filed a \$3B suit against **Anthropic**.¹⁷⁰ **The Johnny Cash estate** sued Coca-Cola under the **ELVIS Act** — Tennessee's new AI-impersonation law — for using a Cash sound-alike in a tribute-act advertisement.¹⁷¹ By the spring of 2026, the litigation landscape was so dense that *Music*

Business Worldwide was running weekly summary columns just to keep track of which cases were still active.

In **the European Union**, lawmakers tabled a bill in November 2025 seeking an EU-wide minimum age to access AI chatbots and social media, an early acknowledgement that the regulatory question was not just about copyright but about the wider integration of AI into the social fabric.¹⁷²

In **the United States**, the actors' union **SAG-AFTRA**, riding the wave of the Tilly Norwood backlash, opened negotiations in October 2025 that resulted by spring 2026 in *significantly stronger AI protections* in its next four-year contract — a deal that included new consent requirements, residuals, and what the trade press began calling, informally, the “Tilly Tax” on the use of AI actors.¹⁷³

In **the United Kingdom**, the U.K. actors' union **Equity** held a strike ballot in December 2025 over AI scanning of performers' likenesses; the result came back in a 99% landslide in favour of industrial action. The ballot question itself, in its plain language, captured the substance of what was at stake: “*Are you prepared to refuse to be digitally scanned on set to secure AI protections?*”¹⁷⁴ By January 2026 the union had secured what its general secretary called “an improved offer” from producers on AI protections in film and TV negotiations.¹⁷⁵ In May 2026, the broader **AI Disclosure Standard** for the film industry was launched at the **Cannes Film Festival**.¹⁷⁶ In the same week, the **British Phonographic Industry (BPI)** issued a formal set of *transparency and sovereignty* demands aimed at the music side of the same settlement — a structured industry position designed, in the BPI's own framing, to *secure* the “AI licensing boom” rather than leave it to bilateral negotiation between platforms and rights-holders one model at a time.¹⁷⁷

The pattern, in every jurisdiction and across every part of the creative economy, was the same. Where creators were given a procedural mechanism — a consultation, a strike ballot, a contract negotiation, a class action — they used it. They turned up in numbers. They voted, in their structured way, against the unconditional appropriation of their work. And they won enough of these procedural battles that, by the time the spring of 2026 arrived, the *terms of engagement* for AI in the creative industries had been substantially re-set in a six-month window.

The first reversal

The most surprising single event in the entire policy arc was the UK government's own *reversal* in spring 2026.

The Statement of Progress in December had already softened the official position. Where the original consultation had proposed Option 3 — the text-and-data-mining exception with opt-out — as the *preferred* outcome, the December update simply described the government as “working with 50+ experts from across music, film, games and AI to figure out what comes next.”¹⁷⁸ The opt-out language was gone.

By **March 2026**, the position had reversed further. The government's final report on copyright and AI, laid before Parliament by the statutory deadline of 18 March 2026, walked

back the original preference for Option 3 in favour of a much more cautious set of proposals that acknowledged the 88% finding.¹⁷⁹ *Dream Machine* Issue 21, dated 19 March 2026, was the first edition where I noticed the change in tone in the government’s own language. The framing had shifted from “how do we enable AI training” to “how do we protect creators.”

I want to be precise about what this reversal means and what it doesn’t.

It does *not* mean that the UK has banned AI training on copyrighted work, or that it has imposed a licensing-first regime by default. As of the time I am writing this — May 2026 — the legislative process is ongoing, and the eventual policy could land anywhere on a wide spectrum.

It *does* mean that 88% of 10,112 people, plus a thousand-odd email submissions, plus a media cycle that ran for fourteen months, plus a parallel set of legal proceedings, plus a parallel set of platform and industry pushback, plus the active mobilisation of multiple professional bodies, was enough to *change the position of a national government* on one of the most economically significant technology-policy questions of the decade.

That is, in democratic terms, what *working* looks like.

What the creators were actually saying

The 88% was a procedural answer to a procedural question. What were the *substantive* arguments behind it? I have read enough of the submissions, through the published summaries and through the secondary press coverage, to feel confident in summarising the three I see most often.

The consent argument. This was the simplest and the most universal: that work made by a creator — a song, a book, a photograph — belongs to that creator in a way that is *not* fully captured by the existing copyright regime, and that the use of that work to train a machine learning model is a use that requires the creator’s consent.

The argument is not new. The Berne Convention has, since 1886, treated authorship as a *moral right* in addition to an economic one. What is new is the scale of the use. A single AI training run can ingest the work of millions of human creators in a way that no single buyer, publisher, broadcaster or aggregator has ever done. The procedural mechanisms of copyright were designed for a world where uses were enumerable. They struggle in a world where the use is, in effect, *the entire creative output of a generation, all at once*.

The attribution argument. This was the most operationally specific: that when AI systems produce derivative outputs based on training data, the creators whose work shaped those outputs should be identifiable, and where appropriate, compensated. *Musical AI*, a startup that raised \$4.5m in January 2026 on a “creative weight attribution” model, described the technical version of this as “calculating each input’s actual contribution to a generative model’s output, then licensing accordingly.”¹⁸⁰ The argument doesn’t require AI training to stop. It requires it to *show its workings*.

The economic argument. This was the most cynical and the most powerful: that AI systems trained on the unpaid labour of creators will eventually substitute for those creators in the market, and that the failure to license is therefore not just an *ethical* offence — it is an

active transfer of wealth from a relatively diffuse group of working creatives to a relatively concentrated group of technology platforms and their shareholders.

The PRS for Music *2026 AI Survey* found that **four in five** music creators worried about AI-generated music competing with human-created music in the streaming economy.¹⁸¹ The Edinburgh-based **Centre for Creative AI** at UCL/RCA, launched in late 2025, explicitly framed its mission around the “redistribution of value from machines back to the humans whose work made them possible.”¹⁸² The U.S. **artist trade body** quoted in *Complete Music Update* in November 2025 was even blunter: “Artists must have creative control in AI deals or risk ending up with ‘scraps’.”¹⁸³

Stack those three arguments next to each other and you get a recognisable shape. It is the shape of every economic-rights argument creators have made, in every previous technological transition, going back to the Stationers’ Company in seventeenth-century London. *Don’t print without permission. Don’t broadcast without a fee. Don’t sell our records without paying us. Don’t sample without clearing. Don’t stream without licensing. Don’t train without consent.*

The 88% is the latest entry in a four-hundred-year sequence. What’s new is the *speed* with which it has had to be expressed, and the *scale* of the use it is responding to.

The artists’ declaration

In January 2026, in parallel with the union negotiations and the lawsuits and the policy responses, nearly **800 working creatives** — including high-profile names like Jason Aldean and OneRepublic — signed an open declaration with the line that gave the document its name: *Stealing Our Work Is Not Innovation*.¹⁸⁴

I want to spend a moment on this document because it is the cleanest expression I have found of the underlying argument, and because I think the line will be on a t-shirt within a year if it isn’t already.

The declaration was not a legal document. It had no enforcement mechanism. It did not call for specific legislation. It was a *cultural statement* — a refusal of the framing under which the AI companies had been making their case.

The framing the AI companies had been using, repeatedly, in venues from technology conferences to court filings, was that training models on copyrighted material was a kind of *technical inevitability* — that machine learning required vast amounts of data, that the data could not practically be licensed at scale, and that therefore the use was, in a sense, *outside* the traditional consent-and-payment framework of copyright. It was — they argued — not really “use” in the sense the law had been built around. It was a new kind of activity that needed a new kind of rules.

The declaration’s response, in a phrase, was: *no, it’s just stealing*.

This was a rhetorically devastating move. It collapsed the AI companies’ carefully constructed framing — *transformative use, fair use, technical necessity, innovation* — into the oldest accusation in commerce, and made it stick. *Stealing*. Not because the signatories did not understand the technical arguments. They did. Because they had decided that the technical

arguments were a *cover* for an underlying transfer of value that didn't deserve any other name.

Once you have that framing, the whole policy debate looks different. *Should we allow innovation?* becomes *should we allow theft?* The answers are not the same.

The levy precedent: why Petrillo matters here

I want to take a long detour, because the 88% — and the institutional response forming around it — is, on the historical reading I laid out in [Chapter 2](#), a *Petrillo-template* moment that the trade press has, in the main, declined to recognise as such.

Let me state the template again, in its cleanest form, because the rest of this section relies on it.

When James Caesar Petrillo, the president of the American Federation of Musicians, took on the recording industry in 1942 and again in 1948 — staging the recording bans that effectively shut down the entire commercial output of American recorded music for the better part of three years — the strategic move was not, in essence, *prohibition*. Petrillo was not trying to ban records. He was trying to *tax* records. The 1942 settlement created a per-record royalty paid into an AFM unemployed-musicians fund. The 1948 settlement, after the Taft-Hartley Act outlawed the 1942 structure, created the **Music Performance Trust Fund** under Section 302 — a *jointly-administered* labour-management fund, paid into by the labels and broadcasters, used to subsidise free live music performances by working musicians, distributing the productivity gain of the new recording technology to the displaced labour pool. The MPTF still exists. It still distributes payments today. It is, on a hundred years of evidence, the *only* form of institutional response to a creative-technology displacement that has worked at structural scale.

The four parts of the template, again:

One, the displacing technology is not banned. It is allowed to displace.

Two, the platform owner pays an *ongoing per-unit tribute* to the displaced labour pool.

Three, the tribute is collected *centrally*, by a joint labour-management body, not negotiated individual-by-individual.

Four, the tribute is paid out to subsidise the *displaced creative practice itself* — live music, in Petrillo's case — keeping it alive as a category even as the market for it shrinks.

I want to show how the 88% — and the architecture of institutional response coalescing around it in spring 2026 — is, function by function, a *reconstruction* of the Petrillo template for the AI era.

One, none of the institutional responses I have catalogued in this chapter — the UK consultation's licensing-by-default proposal, the SAG-AFTRA Tilly Tax, the *Stealing Our Work Is Not Innovation* declaration, the GEMA ruling, the Cannes Disclosure Standard — is, in essence, a *ban*. The declaration's signatories are not asking for AI to be prohibited. The 88% of UK respondents who wanted licensing-in-all-cases were not asking for AI training to be banned. They were asking for it to be *licensed* — which is, by definition, an

acknowledgement that the underlying activity will continue. This matches Petrillo's first principle.

Two, what the 88%, the GEMA ruling and the *UMG v. Anthropic* settlement framework are collectively asking for is a *per-output tribute* from the AI platforms to the creative-labour pool whose work was used in training. The mechanism is, structurally, identical to the per-record royalty that Decca and Columbia agreed to pay AFM in 1944. The platform pays. The labour pool receives. The amount is calibrated to the volume of platform output. The mechanism is the Petrillo mechanism.

Three, the structural innovation of the Petrillo settlement — collection through a *joint body* rather than through individual-creator negotiation — is, in spring 2026, only partially built for the AI era. The collective-licensing infrastructure for music (PRS, GEMA, ASCAP, BMI, SIAE, JASRAC and the related international bodies) has, in some cases, started negotiating directly with the AI platforms on the per-output structure. The Musical AI *creative weight attribution* infrastructure is a first attempt to build a *technical* layer underneath the joint-body political layer. The Cannes Disclosure Standard is an industry-coordination mechanism for the production-side disclosure that the collection mechanism rests on. None of this is finished. The joint bodies for *visual artists, writers, games developers, photographers* are at much earlier stages of development. The MPTF-equivalent fund-and-distribution mechanism does not yet exist for most of the creative industries. *Building it is the institutional work of the next eighteen months.*

Four, the final part of the template — paying the tribute out to subsidise the *displaced practice* — is the part the AI debate has, in my view, most under-thought. What does “subsidising the displaced practice” look like for AI-displaced creative work? For working musicians whose tracks are being competed-against by Suno outputs, it could look like funded performance opportunities, funded studio time, funded creative-development grants — the direct lineage of MPTF live-performance subsidies. For working illustrators whose work was used to train image models, it could look like commissioned-work grants, funded artist residencies, public-art-commission expansion. For working authors whose books were used to train LLMs, it could look like Public Lending Right expansion, library-licensing funds, writer-in-residence programmes. The structural move is the same in each case: *take the productivity gain from the platform, redistribute it to the displaced practice, keep the practice alive as a category.* This is what the 88% is implicitly asking for, whether or not the consultation respondents would have phrased it that way.

I want to be honest about a complication that the Petrillo template hits at full speed in the AI era, because the book should not be glib about it.

The MPTF works partly because the relationship between *recorded music* (the displacing technology) and *live music* (the displaced practice) is *one-to-one*. The same musicians could, in 1948, do either thing. The Petrillo settlement was, structurally, paying the displaced version of the labour to subsidise the alternative version of the same labour.

The AI version of this relationship is *many-to-many*. The training data for a generative-image model is the lifetime output of *thousands* of working illustrators, photographers and visual artists, each of whom contributed an individually-tiny fraction of the model's competence.

The output is *generated* — there is no clean per-image-licence-equivalent. The redistribution problem is, by structure, much harder than Petrillo’s problem.

Two attempts to solve this are visible in 2026.

*The first is **creative weight attribution*** — Musical AI’s framing, picked up by some of the C2PA-adjacent technical standards groups — which proposes that AI platforms compute, for each output, the *gradient-weighted contribution* of each training-data input, and distribute a per-output royalty proportionally. The technical infrastructure for this is, in mid-2026, *partially* built. The economic infrastructure to handle the resulting micro-payments is, in mid-2026, *not* built. But the mechanism is the right one in principle: it preserves the one-to-many relationship that Petrillo could not directly handle, and translates it into a *many-to-many* redistribution mechanism.

*The second is **collective licensing at the publisher tier***. The Stability AI / Universal Music alliance, the Splice / UMG partnership, the various YouTube and Spotify catalogue-licensing deals operate by aggregating training-data permissions at the publisher and label level, with the per-creator distribution handled internally by the existing royalty infrastructure of those publishers. This works for *commercially-published* creative work where the publisher already has a contractual relationship with the creator. It works less well for *independent* and *self-published* creative work where there is no publisher to negotiate on the creator’s behalf.

Both approaches will, in some hybrid form, be the architecture of the AI-era Petrillo settlement. The 88%, the GEMA ruling, the SAG-AFTRA bargaining, the Cannes Disclosure Standard and the *UMG v. Anthropic* litigation are the political pressure that is forcing the platforms to agree to *some* version of one or the other. The version that emerges over the next eighteen months will, on the historical pattern, define the next forty years of how creative-AI work is paid for.

If the working-creative cohort reading this is asking what specifically to push for, my answer is: *the Petrillo template, applied to AI, collected through a joint body, distributed through a creative-weight-attribution mechanism layered on top of the existing collective-licensing infrastructure, used to subsidise the displaced creative practice as a category*. That sentence is a mouthful. It is also the most-likely-to-work structural answer that the historical pattern points at. The 88% is the political mandate for it. The institutional architecture is, in mid-2026, half-built. Finishing it is the work.

The consent-trained alternative

The 88% is a *demand-side* fact: it tells you what creators want done about the training pipeline. There is a *supply-side* fact that I think the policy debate has been slow to absorb, and that working creatives reading this book should know about, because it is the practical refutation of the AI companies’ core argument.

The AI companies, as I noted earlier in this chapter, have spent two years arguing that machine-learning models *cannot practically* be trained on licensed data at the scale they require. That the data volumes are too large, the licensing relationships too fragmented, the

legal cost too high. That training on consent-acquired data is, in effect, a nice idea that does not survive contact with the engineering.

By spring 2026, this argument was falsifiable, and had been falsified, by the existence of a category of foundation models that had been built — and were commercially successful — on exactly the consent-first basis the AI companies said was impossible.

The category, with the models I would name as its leading examples:

- **Adobe Firefly** (Image Model 5, Firefly Video, Firefly Foundry) — trained on Adobe Stock licensed content, openly licensed material, and public-domain work. Adobe Stock contributors are compensated through the Firefly bonus programme. By April 2025, Firefly had generated 22 billion assets; by mid-2026 it was the dominant enterprise generative-image tool, with 45%+ Creative Cloud penetration.¹⁸⁵
- **Bria** — image and video foundation models trained on a 100%-licensed dataset assembled from partners including Getty Images, Alamy and a network of independent rights holders. Bria publishes a per-output attribution and royalty mechanism on which the broader creative-weight-attribution conversation has drawn.¹⁸⁶
- **Getty Images Generative AI** (built with NVIDIA Picasso) — trained exclusively on Getty’s licensed library, with contributor royalties paid through Getty’s existing rights infrastructure.¹⁸⁷
- **Moonvalley Marey** — a generative-video model positioned explicitly as the “clean” alternative to Sora and Veo, trained on a licensed video dataset.¹⁸⁸
- **AIODE** — a music-creation DAW trained on ethically-licensed audio, listed in the toolchain map in Chapter 16.¹⁸⁹
- **Tamber** — an ethically-trained AI music suite launched in May 2026, controllable via arm gestures, positioned explicitly as a tool that moves *with* the musician rather than substituting for them.¹⁹⁰
- The **Stability AI / Universal Music Group** alliance, the **Stability AI / Warner Music** deal and the **Splice / Universal Music** partnership — data-alliance arrangements that build music models on rights-cleared catalogues, with the per-creator distribution handled through the labels’ existing royalty infrastructure.¹⁹¹

I am not claiming this category is perfect, or even, in every case, that its consent claims fully hold up to scrutiny. Adobe Firefly has faced criticism over the inclusion of AI-generated stock images in its training set;¹⁹² the per-creator economics on the Stability / UMG-style deals are still being worked out. The point is not that these models are above critique. The point is that they *exist*, that they *work commercially*, and that their existence collapses the central technical-inevitability argument that the rest of the industry has been using to justify scraping.

The three Cs: consent, credit, compensation

By the late spring of 2026 the licensing-side conversation had begun to settle on a three-word shorthand worth pulling out of the contract architecture and putting on the page, because it is the cleanest articulation I have seen of what the *operative* version of the 88%-mandate looks

like inside an actual platform-and-label deal. The shorthand is **consent, credit, compensation** — the three Cs — and it was the explicit framing for the *Spotify–Universal Music Group* licensing agreements announced at Spotify’s Investor Day on **21 May 2026**, which let fans create AI covers and remixes of participating artists’ songs as a paid Premium add-on, with revenue flowing back to the artists and songwriters whose work is reinterpreted.¹⁹³

The three Cs map cleanly onto the Petrillo template I described earlier in this chapter. *Consent* is the opt-in catalogue — only artists who actively agree have their work made available for remixing, which is the *acknowledgement-that-the-underlying-activity-will-continue* move that Petrillo’s first principle requires. *Credit* is programmatic provenance — derivatives are linked to original metadata, and a *Verified by Spotify* badge steers attention and money toward sanctioned work, which is the *central-joint-body-collection* mechanism that Petrillo’s third principle requires (with the platform doing the joint-body work that, in the recorded-music case in 1948, the AFM did). *Compensation* is a new royalty stream layered on top of standard streaming, paid out to the artists and songwriters whose work is used — which is Petrillo’s second and fourth principles (per-output tribute, paid into the displaced creative-labour pool) at the level of a streaming-product line.

The deal’s prehistory is the part I want every working creative reading this to register. UMG spent 2024 calling Suno and Udio’s training data “*infringement on an almost unimaginable scale*”. It settled with Udio in October 2025 and agreed to build a licensed product instead. The Spotify deal is the mature form of that turn — litigation converted into a revenue line. The market read it as a direct shot at Suno; Spotify’s stock jumped on the day. **The lawsuits were never the destination. They were leverage to force a licensing table into existence.** That sentence is, in my view, the right read on the 88%, on GEMA, on UMG v. Anthropic, on the Wixen suit, on the *Stealing Our Work Is Not Innovation* declaration and on every other piece of legal-and-political pressure I have catalogued in this chapter. The pressure produced the licensing table. The table is now being built. The three Cs are the contract pattern the table is using.

A second 2026 deal — the **Hasbro Sixth Wall Behavioral Licensing** launch on **3 June 2026** — is the same architecture applied to characters rather than recordings, and is worth flagging here because it extends the *consent* leg in a direction the music settlement does not yet have to reach. Where UMG’s opt-in catalogue grants permission for the *song* to be remixed, Hasbro’s CharacterOS grants permission for the *behaviour* of a character — how Optimus Prime argues, how Megatron menaces — to be deployed inside third-party experiences, with the underlying voice performance compensated through a *talent-participation* structure to the original performers.¹⁹⁴ On the three Cs, this is the most complete model in the field as of mid-2026. It is also where the sharpest critique sits, and the critique is worth marking. *Participation* is a compensation mechanism — it answers *credit* and *compensation*. It does not, by itself, settle *consent* in the deeper sense the actors’ unions mean: ongoing control over how a performance is used, and the right to refuse. *A pay cheque is not a veto*. The Sixth Wall model’s long-term legitimacy will rest on contract terms the press releases don’t yet disclose, and the SAG-AFTRA-shaped argument about NIL rights is

going to continue to be the live argument underneath every Behavioural-Licensing deal that follows.

For the deep analytical treatment of the three Cs across five emerging licensing models — Hasbro/Behavioural Licensing, Spotify/UMG embedded remixing, Disney/OpenAI as cautionary tale (the architectural framework that survived the platform's collapse), the independent-artist commons-and-DAO models (Grimes / Elf.tech, Holly Herndon / Holly+, ElevenLabs heritage voices, and the Meta likeness-rental anti-pattern), and the Suno/Alexa *create-to-consume* demand floor — see [Appendix L: The Programmable Brand](#).¹⁹⁵

Indemnity as the receipt

The clearest single signal that a model has done its upstream consent work is whether the company behind it is willing to *indemnify* its customers against copyright infringement claims arising from generated output.

The pattern, in the eighteen months to mid-2026:

- **Adobe** offers IP indemnification on Firefly enterprise outputs — Adobe will defend customers, and pay damages, if an enterprise generation is challenged on copyright grounds.¹⁹⁶
- **Microsoft** announced its **Copilot Copyright Commitment** in September 2023 and extended it across the Copilot product family through 2024–25. Microsoft assumes the legal liability for commercial customers using Copilot outputs in good faith.¹⁹⁷
- **Google** offers an indemnification commitment on Vertex AI and Workspace generative outputs.¹⁹⁸
- **IBM** offers an uncapped indemnity on watsonx-generated content for enterprise customers.¹⁹⁹

Notice which companies are on this list and which are not. The companies indemnifying their customers are, without exception, the companies that have invested most heavily in the *upstream* consent work — licensed data, contributor compensation, rights-cleared catalogues. The companies that have *not* indemnified their customers are, predominantly, the companies whose training-data position is most exposed.

This is not a coincidence. Indemnification is a receipt. It is the legal department of a \$200B company telling its commercial customers, in the most expensive language available, *we have done the work; you can use this without being sued*. The absence of an indemnity, conversely, is an instruction. It is the same legal department saying, *the risk is yours; you carry it*.

For working creatives, agencies and studios making procurement decisions in 2026, the indemnity status of a tool is the single most useful one-question proxy for whether its training pipeline is built on the side of the 88% or against it. Ask the vendor. If they cannot give you a written indemnity, you have your answer.

The literacy problem

None of this — the consent-trained category, the indemnity framework, the C2PA provenance stack in [Chapter 12](#), the legislative reversal earlier in this chapter — works without a corresponding investment in *literacy*. And the literacy gap, in mid-2026, is the place where I am most worried about the architecture failing.

Policy and infrastructure can constrain the supply side. They cannot, on their own, redirect the demand side. The question of whether a working illustrator chooses Firefly over Midjourney, whether a marketing team specifies Bria over a scraped open-source model in its agency brief, whether a record label's A&R department uses a Stability / UMG-aligned tool rather than Suno for demo work, whether a film commissioner asks for Marey provenance on a generative-video shot, whether an audience member streams a SynthID-watermarked track over an unlabelled one — these are *consumption* questions. They sit downstream of every law and every standard. They are decided, ten thousand times a day, by people choosing tools and content from a menu, without anyone telling them what the choices on the menu actually mean.

Three things have to happen for the literacy layer to catch up with the infrastructure layer.

First, working creatives need to know what the consent-trained category is, which tools are in it, and what an indemnity is for. This book is one attempt at that; the **Sundance AI Literacy Initiative**, training 100,000+ artists in provenance practice on Google's funding, is another.²⁰⁰ The professional bodies — the Society of Authors, the AIGA, Equity, the AOP, the MPG, the WGA — have, by mid-2026, started shipping member guides. The work is early.

Second, the *buyers* of creative work — the brands, the agencies, the broadcasters, the platforms, the publishers — need to make ethically-trained models a *specified requirement* in their briefs and procurement contracts. A handful already have: the BBC, the AP wire service, the Cannes festival itself. The vast majority have not. The lever exists. It needs to be pulled.

Third, the *audience* needs the equivalent of a nutrition label. The Cannes Disclosure Standard, SynthID-in-Gemini, the YouTube AI-content disclosure rules, the proposed EU AI Act labelling obligations are early attempts at this. None of them yet add up to a consumer-facing signal as legible as the *Fairtrade* mark or the *organic* certification. Until they do — until an audience member streaming a song, watching a clip, or buying a print can tell at a glance whether the creative work in front of them was made with a tool that paid the people whose work it learned from — the consumption side of the equation will keep leaking. Provenance metadata sitting in a file header that no one reads is not, on its own, literacy.

I do not think this layer will get built by the AI companies. The incentive isn't there. I think it will get built — slowly, contentiously, in fits and starts — by the same coalition I am about to describe in the next section: by creators, their unions, their professional bodies, their buyers and their audiences, jointly insisting on a labelling regime that the platforms eventually have to honour because the market has organised itself around it.

The 88% is the political mandate. The consent-trained models are the proof of supply. The indemnity framework is the legal receipt. The literacy infrastructure is the missing piece — and it is, on the evidence of the last six months, being built.

Coalitions, not protests

The thing I want creative people reading this to take from this chapter is *not* that protest works.

Protest works. We have seen it. The 88%, the Equity ballot, the SAG-AFTRA contract, the artists' declaration, the GEMA ruling, the U.K. government's reversal — these are evidence that protest works.

What I want you to take is that *coalition* works.

What happened in these six months was not — or was not only — that individual creators got angry and shouted. What happened was that creators *aligned themselves* with adjacent groups whose interests they had not previously seen as aligned with theirs.

Working musicians aligned with photographers, who aligned with authors, who aligned with games developers, who aligned with screenwriters, who aligned with voice actors, who aligned with concept artists, who aligned with translators, who aligned with journalists. They aligned with their unions. They aligned with their professional bodies. They aligned, somewhat to everyone's surprise, with the major studios, who had spent twenty years suing them and now found themselves on the same side of a copyright argument against the same platform companies.²⁰¹

The major-label leadership read of the same coalition shifted, in the spring of 2026, in a way that I think is worth registering carefully because the rhetoric is a useful weather-vane. **Robert Kyncl**, the chief executive of **Warner Music Group**, in a widely-quoted May 2026 interview, told the industry that “*AI resistance*” was actively *setting the music sector back* — that AI represented “an incredible value creation opportunity,” and that the labels “*cannot wait the way the industry did 25 years ago*.” Kyncl's invocation of the Napster moment was deliberate. The argument was that the labels' twenty-five-year pattern of *suing first, integrating second* had cost them, in net, the bulk of the streaming-era surplus to platforms that had moved before they did, and that repeating that pattern with AI would compound the loss.²⁰² This is a meaningful change in register from the 2025 “*biggest theft in music history*” framing, and worth tracking. It does not contradict the 88% — Kyncl, like the BPI, is pushing for licensing infrastructure rather than against it — but it shifts the *centre of gravity* of major-label rhetoric from prohibition toward participation. On the Petrillo template, this is the labels' tribute-mechanism position: AI continues, the platforms pay, the joint-body collection infrastructure scales. The question of whether the *displaced practice* gets a meaningful share of the resulting flow — whether the working songwriter, the working session player, the working independent artist actually *sees* the tribute — is the part the Kyncl framing does not yet answer.

They also — and this part I find most interesting — aligned with their *audiences*. The Adobe Creators' Toolkit Report found that **69%** of creators worried about their work being used to train AI without consent.²⁰³ That number rhymes with the 88% in the U.K. consultation. It also rhymes with the audience behaviour I described in Chapter 5 — the slop ceiling, the AI-music underperformance, the cultural rejection of synthetic content that doesn't disclose itself. The creators wanted protection. The audience, given a choice, wanted to listen to the

protected work. The two interests, for the first time in a long time, sat on the same side of the line.

That alignment is the most powerful political asset the creative industries have had this century. They built it in six months. The question for the next six months — which Chapter 13 of this book is going to come back to — is what they do with it.

The government admits it got it wrong

On **1 June 2026** — at **SXSW London**, a cultural and technology conference whose choice of venue was itself a statement about London's post-Brexit attempt to compete with Austin for creative-economy credibility — UK Culture Secretary **Lisa Nandy** said, on a public panel, the words that the 88% had been waiting fourteen months to hear from a minister.

*"I just want to be really clear about where this government sits on it, because we went out to try and find a way to help tech and creative industries to unlock the opportunities of AI, and we got it wrong in the first instance."*²⁰⁴

It is worth holding that sentence in its full weight. *We got it wrong in the first instance.* This is a serving government minister, at a public conference, explicitly conceding that the policy position her department had put to Parliament — the opt-out model, Option 3, the preferred option that only 3% of respondents had supported — was wrong. Not "under review." Not "subject to ongoing consultation." *Wrong.* The verb is unambiguous and it was not walked back.

What Nandy described at SXSW was a government that had gone into the consultation trying to find a workable compromise between tech and creative industries, and had, by the weight of the 88% finding and fourteen months of subsequent pressure — the GEMA ruling, the SAG-AFTRA settlement, the Equity ballot, the Cannes Disclosure Standard, the parallel legislative movements in the United States and the European Union — been forced to reframe its understanding of what the question even was. The government had begun the process asking *how do we enable AI training?* and had been forced, by the volume and coherence of the creative sector's response, into asking *how do we protect the people whose work underpins it?*

The vindication is real. And I want to be precise, as I was precise about the March 2026 reversal earlier in this chapter, about what it does and does not mean.

It *does* mean that the 88% — and every piece of union, legal, institutional and cultural pressure that built up around it — was sufficient to change the stated position of the government that proposed the opt-out model. In policy terms, that is a significant outcome. Very few public consultations on complex technology questions produce a ministerial admission of error this explicit.

It *does not* mean the fight is over. Nandy's statement at SXSW was accompanied by the careful qualifier: "no decisions on next steps." There is, as of the time of writing, no legislation in place, no licensing regime enacted, no opt-in framework passed into law. The legislative process is ongoing. The political will has shifted; the policy architecture that would make that shift concrete has not yet been built. The 88% produced a political admission. The

institutional and legal work that converts that admission into durable protection for creators is still ahead. The Petrillo template I described earlier in this chapter — the joint-body collection mechanism, the per-output tribute, the subsidy for the displaced creative practice — remains the structural target. The political mandate for it has never been stronger. The architecture is still half-built.

Chapter 7 — The Studios Decide

If the audience was speaking through the slop ceiling, and the creators were speaking through the 88%, the studios were speaking through their balance sheets — and the language was not quite the language of either of the other two.

On **22 October 2025** — three weeks into the period this book covers, the same day I was writing Issue 4 — Ted Sarandos, Netflix’s co-CEO, told an industry conference that Netflix was “all in” on leveraging AI across its streaming platform.²⁰⁵ The phrase was casual. The implications were not. Within hours, the trade press was running it as the official line of the world’s largest streaming service, and it was being read — correctly — as a signal to every other studio that the period of “wait and see” was over.

Three weeks earlier, in a piece *Futurism* had published with a headline that aged badly almost in real time, **Lionsgate’s** ambitious attempt to use AI for movie development had been characterised as having “crumbled into disaster.”²⁰⁶

Two months later, on **11 December 2025**, *The Guardian* reported that **Disney** was announcing a \$1 billion investment in OpenAI, with a structured agreement that would let Disney characters appear in the Sora video tool.²⁰⁷ At the time of the announcement, the deal was treated by the press and most of the industry as the canonical template for how a major rights holder would integrate with a frontier AI lab. By the spring of 2026, on closer inspection, it was nothing of the kind. *No formal agreement was ever signed; no money changed hands.* On **24 March 2026**, OpenAI told staff it was shutting down Sora — a staged wind-down, with the consumer app and website ending in late April and the developer API set to close in September — and Disney, reportedly given as little as thirty minutes’ notice, exited the partnership.²⁰⁸ The billion-dollar template, as the *Ropes & Gray* analysis later put it, was a billion-dollar press release.

That collapse matters for the strategic map in two ways, both of which I want to flag now so the reader can read the rest of the chapter with the correction in view. *First*, the framework Disney’s lawyers built — IP-not-training, talent-likeness excluded, jointly-governed steering committee, equity-and-warrants structure — *survived the platform’s death* and is now the reference contract architecture against which other studios negotiate. The framework was right; the vendor was fragile. *Second*, the lesson the rest of the industry took from the collapse was that an *equity stake in a platform you don’t govern* is not a strategic position. The thirty-minutes’-notice detail became a case study in strategic-investment risk: a \$1bn pledge with no board seat or information rights strong enough to give the rights holder warning that the product it was licensing into was about to be killed. The takeaway for every rights holder that follows has been straightforward — *align on equity and governance rights, not just licence terms, or build the platform yourself.* Which is exactly what Hasbro did with **Sixth Wall** and the **CharacterOS** governance layer on 3 June 2026, a move that I treat in detail under Position Four below. The Disney–OpenAI material throughout this chapter should be read with this correction in mind. For the full deal-architecture analysis — the *three Cs* of consent, credit and compensation, the comparison against Hasbro, Spotify/UMG, the independent-artist models, and the Suno/Alexa demand floor — see [Appendix L: The Programmable Brand](#), the *Dream Machine* Deep Dive №14 of 4 June 2026 that this chapter rests on.²⁰⁹

These three moments — Lionsgate’s failure, Netflix’s commitment, the Disney–OpenAI announcement-and-collapse — are the three corners of the strategic map that every legacy studio in the world has been navigating for the last six months. They are not a single story. They are three different stories about how a creative business with a hundred years of human-craft DNA tries to integrate a technology that, by the time it integrates, no longer behaves like the technology you thought you were integrating.

This chapter is about the studios. About how they decided. About the ones that went *all-in*, the ones that went *AI-native from scratch*, the ones that went *we are not doing this at all*, and the ones — the most interesting group — that went *we will do it, but only in the places where it doesn’t show up in the work the audience sees*.

The map of those four positions, drawn carefully, is the map of where the film, TV, games and entertainment industries will be in 2030.

Position One: All-in

Netflix’s “all in” framing was the most prominent example of what became, over the autumn of 2025 and the winter of 2026, the dominant public stance of the major streamers. The framing was: AI is a tool, AI is a productivity multiplier, AI is going to be used everywhere in the pipeline, and the studios that adopt it earliest will have the most leverage when the new economics settle.

The actual deployment, when the trade press dug into it, was more interesting than the slogan suggested. Netflix’s use of AI in late 2025 included generative AI tools for visual effects (de-aging actors, scene extensions, background plates), AI-driven recommendation engines that the company had been refining for fifteen years, and — disclosed in a January 2026 *Pymnts* report — a major AI strategic push focused on subscriber *retention* rather than production cost.²¹⁰ The story Sarandos was telling Wall Street was not “AI will replace our writers.” It was “AI will keep our subscribers engaged in a way that human-only programming alone cannot afford to.”

By May 2026 the deployment had moved one further step in. Netflix announced **INKubator**, an in-house AI animation studio explicitly chartered to produce “*feature-quality*” short-form work, and began recruiting for it publicly.²¹¹ What is notable about INKubator, for the purposes of this chapter, is not the size of the unit — small, by Netflix standards — but the *organisational position* of it: an *internal* AI-native studio sitting *inside* the major-streamer architecture, producing original work, reporting up into the same commissioning structure as the live-action slates. That is a different shape from the Position-Two AI-native studios I describe in the next section. It is a hybrid: Position One on the org chart, Position Two in the pipeline.

Adjacent moves from other big studios that autumn told the same story.

Amazon built out an internal “AI Studios” unit in November 2025, naming sports-docs boss Matt Newman as its head of live-action production.²¹² In the same month, Amazon’s *House of David* TV series became one of the first major Western dramas to publicly disclose the use of

more than **350 AI-generated visual-effects shots** in its second season, with creator Jon Erwin telling *Wired* he was “not sorry.”²¹³

By late May 2026, Amazon went further still: the company launched an **AI Creators Fund** and commissioned three animated Prime Video series using generative AI. One of the three — **Punky Duck**, created by *El Tigre* animator **Jorge Gutierrez** — became the most dramatic commissioning collapse of the period. Within **48 hours** of the announcement being confirmed, Gutierrez had dropped out, citing “*relentless derision*” online. The series was dead before it had been made. For the full treatment of the audience-enforcement dynamics underlying this collapse, see the discussion in Chapter 5 §“The pushback.” For the purposes of this chapter’s strategic map, the Punky Duck case is a Position One data point of a specific kind: a major streamer commissioned AI animation through a structural fund and discovered, at the announcement stage, that the audience-enforcement mechanism can operate on a press release. The question that Amazon’s AI Creators Fund has to answer — and that the remaining two commissioned series have to answer — is whether the specific cultural objection was to *Jorge Gutierrez’s participation* specifically (which del Toro’s subtweet suggested) or to AI-animated commissions from a streaming fund generically.²¹⁴

NBCUniversal signed a deal in late October 2025 with the son of *Law & Order* creator Dick Wolf to develop AI-generated games based on its IP.²¹⁵ By late November, the framing had broadened — *The Office*, *Saturday Night Live* and *Sex and the City* were all reportedly being considered as IP for AI-generated game adaptations.²¹⁶

Disney, beyond its OpenAI investment, announced in November 2025 that it was developing generative AI tools to let Disney+ subscribers create and share their own short-form videos using the company’s iconic IP — a play, transparently, at recapturing the engagement Fortnite and Roblox had been taking from passive streaming for years.²¹⁷ To execute this, Disney created a new “**Office of Technology Enablement**” under former Walt Disney Studios CTO Jamie Voris, with the specific mandate of accelerating AI and Mixed Reality adoption across the organisation.²¹⁸ In January 2026, Disney followed up with an announcement of a TikTok-like vertical-video product and an AI video-generation tool aimed at brand advertisers using existing Disney brand assets and guidelines.²¹⁹

Fox Entertainment took an equity stake in **Holywater**, an AI-microdramas company, in October 2025.²²⁰

Sky History acquired *Castles SOS*, an AI-powered documentary, in late November.²²¹

Channel 4 rolled out an AI-driven advertising tool in December 2025 designed to make TV advertising accessible to SMEs — a small home-builder was one of the first clients.²²²

Position One is not subtle. The streamers, broadcasters and major studios with the capital to do it have been integrating AI into their stacks — production, post-production, marketing, advertising, distribution, subscriber retention — at a pace that suggests they have already decided which side of the future they want to be on. They want to be the side that owns the toolchain.

Position Two: AI-native

A second group, more recent and more interesting, are the studios that have decided not to *integrate* AI into existing film and television production pipelines but to *replace* those pipelines entirely with AI-first workflows. These are the **AI-native studios**.

Fremantle, the international production powerhouse, named the boss of its new “AI-native” studio **Imaginae Studios** in October 2025.²²³ By the spring of 2026, Imaginae was developing a project called *Art Awakens*, fusing AI techniques with classical painting IP.²²⁴

Imagine Entertainment — Ron Howard and Brian Grazer’s production company, with one of the most distinguished filmographies in modern Hollywood — partnered with a new AI-first production company called **Obsidian Studio** in November 2025.²²⁵

Wonder Studios raised \$12m in seed in October 2025²²⁶ and by January 2026 was running its own **Wonder Film Festival** with a curated shortlist of AI-made shorts.²²⁷ By May 2026 Wonder had closed a further round bringing total funding to **\$50M**, with the company publicly framing the ambition as becoming “*the A24 of AI production*” — a deliberate analogue to the indie-prestige distribution model rather than to the streamer-replacement model the trade press had been expecting AI-native studios to chase.²²⁸

Asteria — Natasha Lyonne’s AI company, backed by James Cameron’s *Lightstorm Entertainment* — produced its first animated short, *All Heart*, in October 2025.²²⁹

Promise, a deep-pocketed AI studio backed by Google, set up shop in October 2025 specifically to “bring GenAI filmmaking and VFX to legacy media.”²³⁰

Goldfinch launched **enGEN3**, an “AI-Powered Cinematic Universe Platform,” in October 2025.²³¹

Chapter41, a Munich-based AI startup, was launched in November 2025 by **Beta Film** and a group of industry executives.²³²

Kartel — a new AI startup led by long-time TV exec Kevin Reilly, formerly of HBO — was set up in November 2025.²³³

Wanted director Timur Bekmambetov launched a \$5 million project to “generate AI method actors” in November 2025, with the framing: “AI is here to stay. We have to train it responsibly.”²³⁴

Particle6, the U.K.–Netherlands company behind Tilly Norwood, expanded to 41 AI actors in development by November 2025, with founder Eline Van der Velden in a December 2025 *Deadline* interview making the case that AI performance was a “more ethical way” to act — and urging working performers to “future-proof” themselves by creating their own AI avatars.²³⁵

Wonder Studios, separately, adapted a children’s book to an animated series using AI in December 2025.²³⁶

Kling AI and **Evolutionary Films** announced an AI-animated feature, *Minibots*, at the Cannes Film Market in May 2026, alongside a broader Kling-backed filmmaker initiative aimed at funding AI-native productions on the same indie-distribution architecture.²³⁷

On **10 June 2026**, the first fully AI-generated feature film accepted by a major international festival made its world premiere. *Dreams of Violets* — directed by **Ash Koosha**, produced by his company **Fountain o** — screened at the **Tribeca Film Festival** in New York. The film is a 75-minute docudrama about the 2026 Iran protests. It cost **\$2,000** to make and took **three months** in production. The tool stack: **Kling AI** for video generation, **Anthropic’s Claude** for editing, **Google’s Gemini** for research.²³⁸ What Tribeca programmed was not a short or an experiment or a proof of concept. It was a feature — a narrative of a real political event, made entirely without a traditional production infrastructure, at a cost lower than a mid-range laptop. The festival’s decision to program it did not close the “is it art or is it slop?” debate. It opened it, formally, in the festival-circuit context where that debate carries weight. Critics split sharply: advocates for the work pointed to the film’s subject matter and urgency; critics pointed to the same qualities they flag in every AI-native production I have documented in these six months — technical capability uncoupled from the *earned authority* that makes a film’s claim on its audience stick. What is not in dispute is the milestone: *Dreams of Violets* is the first AI-generated feature to pass the selection committee of a Tier 1 film festival. Whatever comes after, it marks the moment the festival circuit acknowledged that the category exists.²³⁹

By **April 2026**, the trade press could no longer keep up with the AI-native studio launches. There were too many of them. Most of them, like most early-stage production companies in any era, will not survive the next two years. The question of whether a meaningful AI-native studio system will eventually emerge as a parallel structure to legacy Hollywood — the way Netflix and Amazon eventually emerged as a parallel structure to the cable networks — is, in my view, the biggest single open question in the film and TV industry as of May 2026.

The early evidence is mixed. *Watch the Skies*, a Swedish UFO feature entirely dubbed with AI, secured U.S. distribution in October 2025.²⁴⁰ *Run to the West*, South Korea’s first AI feature film, was tested with critics and audiences in October 2025; one *cybernews.com* review described the experience as “testing the soul of cinema.”²⁴¹ *Lily*, a Tunisian-made AI short, won the \$1 million Dubai AI Film Award in January 2026.²⁴² *Humans in the Loop*, an AI drama that received Film Independent’s Sloan Distribution Grant, entered the Oscar race in November 2025.²⁴³

I have watched a meaningful percentage of the AI-native output of these six months. The honest evaluation, which I have given in talks several times and stand by here, is that we have not yet seen the *Citizen Kane* of AI cinema — we have not yet seen a single AI-native work that I think will still be watched in 2040. We have seen, repeatedly, films that demonstrate technical ability without yet demonstrating cultural necessity.

The most interesting AI-native works, in my view, are the ones that — like Andrii Daniels’ bomb-shelter clip — wear their non-traditional production conditions on their face. They are

films *about* the technology being used to make them, in some implicit or explicit sense. They are not pretending to be legacy films made by a different route.

Position Three: We are not doing this

The studios that have publicly refused generative AI have been some of the most interesting voices in this entire period.

Pocketpair, the Japanese games studio behind *Palworld*, announced in October 2025 that its new publishing division would not handle games using generative AI. The CEO's full statement, in *PC Gamer*, was sharper than the headline: "*We don't believe in it. We're very upfront about it. If you're big on AI stuff or your game is Web3 or uses NFTs, there are lots of publishers out there [who'll talk to you], but we're not the right partner for that.*"²⁴⁴ It was, on its face, a rejection of one production model. It was, on inspection, also a *marketing position*: a publisher staking its claim with audiences who had become — by late 2025 — actively allergic to AI-augmented games.

Larian Studios — the maker of *Baldur's Gate 3*, one of the most critically and commercially successful games of the decade — backed off generative AI in January 2026 for its next *Divinity* game.²⁴⁵

Games Workshop, custodian of the *Warhammer 40,000* universe, ruled out generative AI entirely in early 2026.²⁴⁶

Hooded Horse, the U.S. games publisher behind *Manor Lords*, said in January 2026 that it would not work with developers who used generative AI.²⁴⁷

Jagex, the maker of *RuneScape*, declared in January 2026 that it would *never* use generative AI to make in-game content.²⁴⁸

Aardman Animations — the British animation studio responsible for *Wallace and Gromit* — announced in December 2025 that it would "embrace the technology" of AI but would be "very cautious not to lose our values."²⁴⁹ This was, by Aardman's careful standards, a sharp line: they reserved the right to use AI for narrowly defined post-production and admin tasks, but explicitly excluded it from the stop-motion craft that defines their work.

Guillermo del Toro, in October 2025, told *Variety* he would "rather die" than use generative AI in his films, with a follow-up *Frankenstein*-themed press cycle that made the line one of the most-quoted creative-industry statements of the year. The full quote was even better than the headline: "*I'm 61, and I hope to be able to remain uninterested in using it at all until I croak. ... The other day, somebody wrote me an email, said, 'What is your stance on AI?' And my answer was very short. I said, 'I'd rather die.'*"²⁵⁰ What del Toro was doing, with the bluntness only a senior auteur with a fully-funded slate can afford, was *refusing to participate in the framing*. Most working creatives have had to spend two years giving careful, nuanced, defensive answers about their AI position. Del Toro decided he was a senior enough artist to refuse the question entirely. The cultural permission for that posture, in a particular kind of high-end filmmaking, is part of the architecture this book has been describing.

Not every senior auteur held the line del Toro drew. In June 2026, **Martin Scorsese** became an advisor to **Black Forest Labs** — the German company behind the FLUX image-generation model — and publicly endorsed using AI for storyboarding in the filmmaking process.²⁵¹ The announcement produced an immediate and, on the cultural register, notable backlash. **Boots Riley** and **Guillermo del Toro** — del Toro, again — were prominent among the filmmakers who responded critically, with Riley’s objections framing the endorsement in explicitly labour-terms: a filmmaker of Scorsese’s stature lending his name to a generative-AI company is not a neutral act in a landscape where working artists are still fighting for compensation frameworks. The Scorsese/Black Forest Labs story is the mirror image of the del Toro position: where del Toro’s “rather die” framing set a cultural permission for refusal, Scorsese’s endorsement set a cultural permission for adoption — and was immediately challenged for doing so. Both positions tell you something important about where the industry’s internal argument is in mid-2026: it is not between the *studios* and the *artists* any more, as it was in 2024. It is between the *artists themselves*, at the level of individual public positioning.

Leonardo DiCaprio, in December 2025, told *The Hollywood Reporter*: “*I think anything that is going to be authentically thought of as art has to come from the human being.*” The headline framing reduced the position to “AI can’t be art because there’s no humanity to it,” which is the version that travelled, but the full quote is more philosophically defensible. DiCaprio wasn’t claiming AI-augmented work couldn’t be valuable. He was claiming that the authorship signal — “from the human being” — was a precondition for the *category* of art, as he understood it.²⁵²

Claire Foy told the *Daily Mail* in January 2026 she had “no interest” in seeing AI in films and would be “disappointed” if it became the future of Hollywood.²⁵³

Jenna Ortega said in December 2025 it was “very easy to be terrified” of AI in filmmaking. Her fuller reasoning, given to *NME*, is the part I have ended up quoting in talks: “*It comes to a point where it becomes sort of mental junk food and we feel sick and we don’t know why. I think, as terrible as it is to say, sometimes audiences need to be deprived of something in order to appreciate something again.*”²⁵⁴ That argument — *audiences need to be deprived of something in order to appreciate something again* — is one of the most interesting things a working performer has said in this period about the slop ceiling and its psychological substrate. The audience does not, by Ortega’s read, simply discriminate against AI work. They develop a *hunger* for the human-authored work *because* of the AI flood. The flood and the hunger are part of the same cultural dynamic.

Chris Pratt publicly rejected a pitch to cast an AI ‘actor’ as the villain in *Mercy* in January 2026: “I don’t think that’s a good idea at all.”²⁵⁵

I do not think any of these positions are static. I think some of them will shift in the next eighteen months, in ways that depend on how the policy environment, the audience response and the tool ecosystem evolve. But I think the *fact* of the positions, written down, in public, on the record, is more important than whether any individual position holds.

What these refusals do, collectively, is keep open a part of the creative economy that the all-in studios would otherwise be forced to close. They make it possible — for the audience, for the working performer, for the next generation of creative-industry workers entering the field — to have a viable career path that does not require AI integration as the price of admission.

In a world without these positions, every working creative would, by default, be a *partial AI operator*, whether they wanted to be or not. With them, the choice remains open.

That is not a small thing. It is the architecture of the future creative economy being deliberately preserved, by people with the cultural standing and the economic security to preserve it.

Position Four: The middle

The position I find most interesting is the one almost nobody articulates clearly, because it does not make a good press release. It is the position of studios that *use AI everywhere except where it shows up in the finished work*.

The clearest version of it is the one **Aardman** has effectively articulated and that **Bethesda's** Todd Howard described in *PC Gamer* in December 2025: AI is “part of Bethesda’s toolset for how we build our worlds or check things” — but it cannot replace human creative intention.²⁵⁶

You see the same position in **Amazon's** *House of David* — 350 AI shots, disclosed up front, but every one of them used to augment rather than originate the work. The show’s creator Jon Erwin gave *Wired* a metaphor that I have not stopped thinking about: “*You can put a very real camera on a very real actor and direct that actor, direct the camera, and that becomes, in essence, the hand inside a puppet. The puppet itself is this digital world that you create.*”²⁵⁷ The hand-inside-the-puppet image is the cleanest articulation I have heard of where the *Position Four* studios are choosing to put their human craft: at the moments of direction and performance, with the AI doing the digital-world infrastructure underneath. You see the same position in the *Battlefield 6* development team’s statement, in October 2025, that generative AI had been “seducing” but ultimately used only in the earliest stages of the game’s development, “to allow for more time and more space to be creative.”²⁵⁸

You see the same position in **The Witcher 3** and **Cyberpunk 2077** director’s November 2025 framing — AI “can help, but not replace, creatives.”²⁵⁹

You see the same position in the **Wallace and Gromit creator** Nick Park’s December 2025 framing — embrace the technology, but be cautious about the values.²⁶⁰

You see the same position, most starkly, in the May 2026 **Sony** announcement that it was “going all in on AI for games” — *with the specific framing that AI was a force multiplier, not a replacement*. Mocap-to-facial animation in seconds rather than hours. AI integrated into asset generation, QA, engineering and animation pipelines. The goal: *more games, faster*. But Sony’s framing, which I have read several times to make sure I am not over-reading it, repeatedly emphasised AI as a tool *inside* the creative work, not as a substitute *for* it.²⁶¹

You see the same position, by May 2026, in the **Cannes** festival press cycle, where the working auteurs on the Croisette had shifted markedly from the prior year's defensive stance to a more cautious *acceptance of inevitability* — framed less as enthusiasm than as a refusal to be left out of the next decade's tooling argument.²⁶² **Peter Jackson**, in a May 2026 interview during the Cannes window, summarised the Position-Four read in a single line that I think will travel further than the Sony announcement: *AI is, in essence, the next wave of special effects*. The director's job is not changed by the SFX wave. The director's job is to know what the film is for, and to deploy whatever tools are now in the box to get it there.²⁶³ **Take-Two's** Strauss Zelnick made the same argument from inside the AAA games business in the same week — that AI "*datasets by their very nature are backward-looking*" and so cannot, alone, make an *original* hit, but that AI is "*super helpful*" in the production of hits the human team has already conceived.²⁶⁴ Different industries, different framings, structurally identical position: *tool in the workflow, not author of the work*.

This middle position — *AI in the workflow, not in the work* — is, I think, where most of the surviving major studios are going to land in 2030. It is the most defensible commercial position because it captures the productivity upside without giving up the cultural-product specificity that the audience continues, against the slop ceiling, to demand. It is also the most defensible *ethical* position, because it allows the studio to credibly claim — and to credibly prove, with disclosure and documentation — that the creative work the audience sees was, in its decisive moments, the work of human creators.

What is at stake, for the studios that get this right, is the next two decades of cultural authority. The studios that adopt aggressively and badly will look like the early-2010s newspapers that switched to clickbait. The studios that refuse entirely will look like the 1980s record companies that refused to release CDs. The studios that thread the needle — that adopt the productivity benefits without surrendering the human authorship signal — will, in my view, be the studios that the audience actually trusts in 2035.

Position Five: Own the rails

A fifth position emerged late in the period this book covers, and I want to flag it as its own line on the strategic map even though it sits, in some ways, *inside* the Position One *all-in* framing. It is the position of the rights holder that, instead of *licensing* its IP to a frontier-AI platform — and discovering, on the Disney–OpenAI evidence, that the platform underneath the licence is more fragile than the IP it covers — *builds the licensing infrastructure itself*. Owns the operating system. Owns the meter. Owns the fence.

The canonical case is **Hasbro's Sixth Wall**, the in-house AI studio the company launched on **3 June 2026** under CEO Roberta Thomson — and the **CharacterOS** governance layer underneath it.²⁶⁵ What Sixth Wall is licensing is not the appearance of Hasbro's characters but their *behaviour* — how Optimus Prime argues, how Megatron menaces, how Cobra Commander reacts — for deployment inside chatbots, games, toys and brand experiences built by third parties. CharacterOS encodes each character's canon, voice, vocal timbre, personality, lore and safety limits and enforces them at *runtime*: a licensed Cobra Commander stays in canonical character without a writer in the loop. The economic model is

a B2B programmatic licence layered on a *talent-participation pay-out* — voices built from authorised recordings of real human performers, who are compensated for the generative use of their work — and the studio’s initial focus is restricted to 13+ and enterprise use, explicitly not building for young children, against the regulatory weather around AI companions and minors.

CEO Chris Cocks framed Sixth Wall as a “*third way*” between endless and costly enforcement of unauthorised fan creation on the one hand, and passive brand dilution on the other.²⁶⁶ The framing is, in my read, the single cleanest articulation of what the *governance layer* actually is. It is not a creative technology dressed as governance. It is a governance technology dressed as a creative one. The deeper conceptual move underneath it — and the one I think will define legacy-IP strategy through the rest of the decade — is the **reconception of IP from an asset to be guarded into an environment to be governed**. For a century, intellectual property was understood as a discrete, ownable thing whose value you protected by controlling access to it. What the 2026 deals reveal is a redefinition of IP as a *bounded space, governed by rules, inside which value is created continuously by people who are not you*. Hasbro doesn’t sell you Optimus Prime. It sells developers a governed space in which Optimus Prime can be summoned, on-brand and on-licence, a million times over. The asset is the rulebook.

The same structural move sits underneath the **Spotify–Universal Music Group** deal announced on 21 May 2026 — fan-made AI covers and remixes of participating artists’ songs as a paid Premium add-on, with revenue flowing back to the artists and songwriters whose work is reinterpreted.²⁶⁷ The deal’s framing is a three-pillar contract — **consent, credit, compensation** — that maps onto the *Petrillo template* I laid out in Chapter 6. *Consent* is opt-in: only the catalogues of artists who actively agree are available for remixing. *Credit* is programmatic: derivatives are linked to original metadata, and the *Verified by Spotify* badge steers attention and money toward sanctioned work. *Compensation* is a new royalty stream layered on top of standard streaming. The deal’s *prehistory* — UMG calling Suno and Udio’s training “*infringement on an almost unimaginable scale*” in 2024, then settling with Udio in October 2025 and agreeing to build a licensed product — confirms the line I argued in Chapter 6: the lawsuits were never the destination. They were *leverage to force a licensing table into existence*. The Spotify–UMG product is the mature form of that turn — litigation converted into a revenue line, and the architectural ancestor of what every other major-label and major-streamer pairing will, on the available evidence, build over the next two years.

The Disney–OpenAI collapse, read alongside the Hasbro and Spotify moves, supplies the lesson. **Licensing to a fragile platform you don’t govern is not a strategic position.** *Equity in the platform without governance rights* is not a strategic position either. The strategic position — Position Five on this chapter’s map — is owning the *governance layer itself*: the licensing protocol, the runtime guardrails, the provenance metadata, the verified-by stamp, the royalty-attribution rails. The legacy industries that absorb this lesson and *build their own governance infrastructure* — Hasbro’s CharacterOS, Spotify’s *Verified-by* stack, the **Cannes AI Disclosure Standard**, the C2PA / SynthID provenance layer described in Chapter 12 — are positioning themselves to capture the durable value of the generative

economy. The legacy industries that wait for the platform companies to build it on their behalf will end up where Disney ended up: thirty minutes' notice and an exit.

There is one further observation I want to surface here, because it shifts the read on the labour story underneath the rest of this chapter. The transaction-cost frame I keep coming back to — that generative AI compresses the cost of search and negotiation but raises the cost of monitoring and enforcement — has a *labour consequence* that the substitution-panic framing of 2024 missed entirely. The work doesn't disappear; *it moves*. Voice performers who supply authorised source recordings, curators who write and maintain canon, the *judgment brokers* described in [Appendix K](#), provenance and rights-management staff, brand stewards, the *Continuum Lead*-shaped roles I described in [Chapter 14](#) — these are the jobs the governance layer creates as fast as it absorbs the routine production layer. The new roles are fewer and differently distributed than the ones they replace, and “*talent participation*” is a weaker contractual instrument than employment. But this is the actual shape of the disruption, and the studios reading the strategic map well are the studios that recognise it and start staffing for it. The deeper analytical treatment of the governance-layer move, the *three Cs* framework, the comparison of the five emerging licensing models (Hasbro/Behavioural Licensing, Spotify/UMG embedded remixing, Disney/OpenAI as cautionary tale, the independent-artist commons-and-DAO models, and Suno/Alexa as the demand floor), and the underlying transaction-cost / resource-based-view economic argument is laid out in [Appendix L: The Programmable Brand](#).

The trap the legacy industries built for themselves

There is a deeper strategic risk underneath the four-positions map that I have, until now, deliberately not put on the page. I want to put it on the page, because I think it is the single most important read on the long-term legacy-studio position, and because, in the conversations I have had with senior creative-industry executives over these six months, it is the read they are most uncomfortable hearing.

The risk is this.

Across the last fifteen years — most aggressively across the last decade — Hollywood, commercial music and the AAA games business have, on the available evidence, *systematically optimised themselves for exactly the kind of work AI is now best at producing*. They have built their economic and creative production engines around the mean of the distribution. They are now competing, in the most direct possible sense, against a technology built to produce the mean of the distribution at near-zero marginal cost.

Let me make the case concretely.

In **film and television**, the IP-cycle data is unambiguous. Sequels, prequels, reboots, remakes, spin-offs and franchise instalments accounted for an increasing fraction of the top-grossing US theatrical releases through the 2010s and 2020s; *original* studio films — meaning IP not derived from an existing book, comic, game, brand or prior film — fell from a majority of major studio output in the late 1990s to a single-digit percentage of wide releases by the mid-2020s, depending on which counting convention you use. The headline form of contemporary tentpole filmmaking, by 2024, was a sequel to a property whose original

installment had itself been a sequel. Marvel Cinematic Universe Phase Six; the *Star Wars* sequel-trilogy aftermath; *Avatar* sequels, *Toy Story* sequels, *Frozen* sequels, *Mission: Impossible* sequels, *Fast & Furious* sequels; *Stranger Things* finales; *House of the Dragon* spin-offs; every 1980s and 1990s IP revisited at least once, most of them more than once. James Cameron's *Hollywood Reporter* observation in 2025 that contemporary studio executives "don't reach for things that are scary" was, on the data, a description, not a prediction.

In **commercial recorded music**, the structural pattern is the same. The streaming-economy hit structure converged on a remarkably narrow set of parameters across the 2010s: average song length compressed from roughly four minutes in 2000 to roughly three-and-a-half minutes by 2024; choruses landed earlier; intros shortened to keep the algorithmic skip-rate down; major-label A&R increasingly drove signing decisions by predictive-analytics data rather than developmental A&R judgment; co-writing teams expanded; the median Top 40 hit, by 2024, was a co-write across three to seven credited songwriters working to formulas that had been tested in advance against streaming-engagement data. The major-label business is, structurally, an *industry that has spent ten years training itself to produce the most predictable possible version of the song*. The Cardiff band from Chapter 5 — whose music was fed to an AI that produced a tracking-style imitator outperforming them on Spotify — is the canonical illustration. The AI did to them what the major-label streaming-optimisation operating model had already half-done. It made the next most-likely-good track. The fact that an AI could match the output is, on the available evidence, *because the major-label hit factory was already producing the work AI was about to be able to copy*.

In **AAA games**, the standardisation is even more visible. The *Ubisoft-tower-and-checklist* structure — open world, viewpoint towers that uncover map regions, scattered icon-driven side quests, levelled enemy zones, crafting trees — became, between roughly 2008 and 2022, the default structural template for the AAA action-adventure genre. *Assassin's Creed*, *Far Cry*, *Watch Dogs*, *Ghost Recon*; outside Ubisoft, *Horizon*, *Spider-Man*, *Mad Max*, *Shadow of Mordor*, every Tom-Clancy-derivative, every western-RPG converted to console. The 2024 *gamesindustry.biz* roundtable I referenced in Chapter 3 — in which working AAA designers described the genre as having stagnated into a single repeating structural pattern — was a working-developer admission that the AAA industry had, like Hollywood and like the major labels, optimised its production engine around predictable, low-risk, repeat-format output.

Now consider what AI, as a creative tool in 2025–26, is structurally best at. Agents — as I argued in Chapter 11 — produce the *mean of their training distribution* by default. The mean of the training distribution, for a video model trained on contemporary Hollywood tentpoles, is *another contemporary Hollywood tentpole*. The mean of the training distribution, for an audio model trained on the contemporary streaming Top 40, is *another contemporary streaming Top 40 song*. The mean of the training distribution, for a games-development agent trained on the AAA open-world template, is *another AAA open-world template*.

This is the strategic trap. The legacy industries, by spending fifteen years training themselves to produce the mean of the distribution, have arranged for the segment of the market they dominate to be *exactly the segment AI replicates most cheaply*. The Cardiff band's experience

is the cleanest version of this dynamic in microcosm. The macro version is the major-studio business model. The risk to legacy Hollywood, legacy commercial music and the AAA games business is not that AI takes their *premium* segments — Cameron’s *Avatar* sequels, the highest-end auteur cinema, the genuinely original musical voices, the *Baldur’s Gate 3*-class boundary-pushing games. Those segments are, on the evidence of the slop ceiling and the authenticity premium, *more defensible than ever*. The risk is to the *median* output of these industries — the franchise instalments, the by-the-numbers chart hits, the AAA action-adventures that read as algorithmically generated even when no algorithm was involved. The median output is exactly where the AI substitution pressure is most direct, and the median output is precisely what the legacy industries have most thoroughly optimised themselves to produce.

The grandmasters of Chapter 15 — the chess players who have started, in 2026, to deliberately play sub-optimal moves to put their opponents on uncomputed ground — are, on this read, the *senior auteurs of legacy Hollywood*. Cameron, del Toro, Soderbergh, Spielberg, Aronofsky, Lyonne, Larian’s Sven Vincke, Hooded Horse’s leadership — the figures profiled in the *Position Three* section of this chapter — are, structurally, the people whose competitive advantage is *the move the machine would not have generated*. The grandmasters can take the punch. The middle ranks — the franchise journeymen, the streaming-optimised mid-tier filmmakers, the chart-A&R commercial-pop machine, the AAA studio designing its fifth open-world action-adventure — are the ones whose business model the machine is structurally suited to replicate.

The contrast with the **new AI-native studios** is sharp, and I want to draw it out carefully because the contrast is, I think, the most underappreciated strategic reality of the period.

The Position Two studios in this chapter — Gossip Goblin, Critterz, Imaginae, Wonder, Asteria, Promise, Obsidian, Chapter41, Kartel, Goldfinch’s enGEN3 — were, by May 2026, accumulating creative production credits at a rate the legacy studios were not matching. **Gossip Goblin**, the AI filmmaker that I covered in *Issue 29* of *Dream Machine*, is the example I find clearest. It is a studio with no inherited IP, no inherited production pipeline, no inherited audience, no inherited rules about what its films should look like, and no inherited risk-aversion. Its only operating constraint is the one any working creative has: make work the audience wants to watch. *Critterz* (*Vertigo* + *Federation*), launched as an AI-assisted animated feature operation in *Issue 29*, operates with the same freedom. *Animaj* (the kids-content AI studio Google’s AI Futures Fund partnered with in spring 2026) operates with the same freedom. *Imaginae Studios* — Fremantle’s AI-native operation — has, in *Art Awakens*, committed to a kind of generative collaboration with classical painting IP that no legacy studio has the institutional permission to attempt.

These studios do not have rules about how a film should be paced, what a song should sound like, what an open-world game should feel like, what a franchise structure should be. They have no quarterly-earnings call about year-on-year tentpole performance. They have no \$200m development sunk cost in a Marvel-style multi-film slate. They have no major-label playlist-pitching infrastructure that demands the song be a certain length and a certain shape.

They have, in short, no calcified definition of *what counts as the right move* — and so they are, by default, free to play the move the legacy studios cannot.

The risk to legacy, in other words, is *not symmetrical* with the risk to the AI-native studios. Both are being shaped by the same technology shift. But legacy faces a double squeeze — its median product is the part of the market AI is most easily eating, *and* its strategic incentives push it deeper into that part of the market every quarter, *and* its calcified production rules prevent it from doing the thing (the deliberately un-machine-like move) that would protect it. The AI-native studios face only the upside.

The path out, for the legacy studios that recognise the trap, is roughly what the *Position Three* signatories have intuited: refuse to compete in the median segment; reposition aggressively into the segments where the audience pays a premium for the human signal; treat IP investment as *bets on artist-author voices* rather than as bets on formula. *Baldur's Gate 3* — Larian's mid-2020s blockbuster — was the inflection-point example: an enormous AAA-grade game built explicitly *against* the Ubisoft-tower template, on a CRPG framework most analysts had written off as a dead genre, and the audience response was the largest commercial success of any new RPG IP of the decade. Larian's January 2026 announcement that the next *Divinity* game would not use generative AI²⁶⁸ is, on the strategic read, not an anti-technology gesture. It is a *commercial* gesture — a public claim that Larian's market position is built on doing the un-machine-like work, and that the studio is going to protect that position from the inside. *Pocketpair's* “*we don't believe in it*” statement, *Jagex's* “*never*,” *Hooded Horse's* “*cancerous*” framing, *Aardman's* careful preservation of the stop-motion craft — these are not statements of moral piety. They are *competitive positioning* against an algorithm-optimal market that the major-studio system has, structurally, conceded to AI.

The legacy industries that survive this transition will be the ones that recognise, in the next eighteen months, that the position they spent fifteen years moving into is exactly the position they now need to move *out of*. The audience, the slop ceiling, the authenticity premium and the chess grandmasters' move are all telling them the same thing: *do not produce the work the machine can replicate. Produce the work the machine, by construction, cannot*. The industries that can absorb that re-direction in their commissioning, contracting and greenlighting culture will outlast the transition. The industries that cannot, won't. The new AI-native studios — Gossip Goblin and the rest — are not the threat to legacy Hollywood. They are, on the structural read, *the proof of what survives*: studios built without the rules that made the legacy industries vulnerable in the first place.

What the studios got right

I want to close this chapter by saying something that is unfashionable in some of the creative-industry circles I move in.

The studios — for all the headlines about Lionsgate's “disaster,” for all the criticism of Disney's OpenAI deal, for all the eye-rolling at Netflix's “all in” framing — have, in this period, made some genuinely difficult strategic decisions in a genuinely difficult environment, and a meaningful fraction of those decisions have been better than the public discourse credits them for.

They have committed to disclosure. *House of David's* 350 AI shots were *disclosed*. Aardman's careful framing was *explicit*. Sony's "AI as force multiplier" framing was *spelled out*. None of these companies pretended their AI use didn't exist. None of them adopted the all-too-common platform-economy strategy of *use it and don't tell anyone*. The disclosure norm, where it has taken hold in the studio system, is a public good.

They have negotiated, in many cases, in good faith with the unions. The SAG-AFTRA contract that emerged from the autumn 2025 negotiations was — by historical comparison with other major technology transitions — produced quickly, produced through legitimate process, and produced with materially stronger AI protections than any prior contract in the industry's history.

They have, in significant cases, *resisted* internal pressure to over-adopt AI in ways that would have undermined their cultural product. Most of the studios in Position Three above are not run by Luddites; they are run by people who have done the maths on the slop ceiling and decided that the long-term value of their IP depends on it remaining recognisably human-authored.

And they have, finally, invested in the infrastructure of the AI-native sector — through funding deals, through co-productions, through equity investments — in a way that means the AI-native studios are not fighting them so much as building alongside them. The picture, ten years out, is more likely to be of a *mixed ecosystem* — legacy studios with hybrid pipelines, AI-native studios with new IP, and a long tail of human-only craft studios serving the highest-value segments of the market — than of a single winner-takes-all outcome.

The studios decided, in these six months, that they were not going to be replaced by AI. They were going to be the *operators* of AI. That decision was — for all its compromises — probably the right one for the working creatives whose careers depend on the studio system continuing to exist.

The harder question, which the next chapter starts to address, is what happens to the *toolchain* underneath those studios — when the platforms providing the AI are themselves becoming AI-native, and when "having a tool" is no longer the right framing for what it means to make creative work in 2026.

That is what Adobe, NVIDIA, Google and the rest of the platform layer started telling us last autumn, when they began saying out loud that AI was going to be *in everything, everywhere, all at once*.

Chapter 8 — Worlds, Not Pictures

If you had asked me, in the autumn of 2025, what the most important AI release of the year was going to be, I would have said something obvious — Sora 2, or Veo 3.1, or one of the music models, or one of the editor-class tools like Runway Gen-4.5 or Adobe’s new Firefly. Something that turned a prompt into a thing you could put on a screen.

Six months later, I don’t think I would say any of those.

I think the most important release of the period this book covers was something that almost nobody outside of a relatively small community of practitioners noticed at the time, that produced no viral videos, that did not change the news cycle for a single day, and that I have come, over the course of the winter, to think of as the *actual* future of creative work: the public launch of **Marble**, by Fei-Fei Li’s company **World Labs**, in November 2025.²⁶⁹

Marble doesn’t make videos. Marble makes worlds.

I want to spend this chapter on why that distinction is, in my view, the most strategically important one in creative AI right now, and why almost everything else in the toolchain — from generative video to AI music to the digital-human work in advertising — eventually has to be re-thought in its shadow.

What a world model actually is

The phrase “world model” sounds like a marketing term. It isn’t, exactly. It is a category of AI system that researchers have been chasing for the better part of a decade, and that — as of late 2025 — finally started shipping as production-ready software.

A generative *video* model takes a prompt and produces a sequence of frames. The frames are coherent because the model has learned the statistical regularities of video: things move smoothly, light behaves more or less correctly, faces stay faces. But the output is *flat*. It is a particular sequence of pixels. You cannot navigate it. You cannot move the camera. You cannot pick up the lamp on the table and look at the wall behind it.

A *world* model takes a prompt — or an image, or a video, or a rough 3D layout — and produces a *navigable three-dimensional environment*. You can move through it. You can change the camera angle. You can, depending on the model, walk around the table, look at the wall behind the lamp, and find that the wall continues to exist in a consistent way that the model didn’t have to generate for you because it understood, structurally, that walls have backs.

The technical core, in the most common implementation, is **Gaussian splatting** — a representation where a scene is stored as a cloud of millions of tiny semi-transparent ellipsoids, each carrying colour and position information. The whole scene can be rendered in real time from any angle, because the system isn’t drawing 2D pixels; it is rendering a structured 3D world. The output, in turn, can be exported as a splat file, as a mesh, or — if you want — as a flat video.²⁷⁰

This is the part that took me, even as a working creative technologist, embarrassingly long to fully understand. Video is a *projection* of a world. A world is the more fundamental object. For two and a half years, the public-facing AI conversation has been about generating better

projections. The actual capability landscape has been moving, in parallel, towards generating the worlds themselves.

When the worlds become cheap to generate, the projections — the videos, the images, the renders — become *outputs of the worlds*, not the primary medium. The whole production stack inverts.

Marble, from the inside

DreamLab — the studio I run in the North West of the UK — has been a beta participant in Marble since October 2025, in the months before its public release.²⁷¹ I want to share what that experience actually felt like, because the *technical* description of a world model and the *practitioner's* experience of using one are different in ways that matter for understanding what is happening to the toolchain.

Imagine, for the sake of example, that I am working on a client project that needs a scene of a market square at dusk in a Mediterranean town. In the old pipeline — which is to say, the pipeline of 2024 and most of 2025 — that brief would translate into something like the following:

A concept artist would produce a moodboard. A 3D artist or a virtual-production house would build a CGI version of the square, populated with assets either bought from a marketplace or modelled bespoke. Lighting would be set up in Unreal or Maya. The whole scene would be rendered out as a video plate or used as a backdrop on an LED volume. If any change was required — *can we move the camera left a bit, see what's on that side* — the rebuild was non-trivial.

In Marble, the same brief unfolds differently. I type, or paste in a reference image, or upload a quick phone-shot panorama of an actual market square I visited last year. Marble generates a complete, navigable 3D environment of that square. It exists, persistently, as a file on my account. I can move my virtual camera anywhere in it. I can hand it to a director and say *walk through this and tell me where the camera lives*. I can export the result as a Gaussian splat, drop it into Unreal Engine via SuperSplat or one of the other Gaussian-splat editors,²⁷² and use it as the lit backdrop for an LED-volume shoot. I can also, if I just want a plate, render a flat video from a chosen camera move.

The economic implication is this: the cost of having “a place” — a navigable, lit, persistent environment with depth — has dropped, in twelve months, by something between one and two orders of magnitude. The thing that used to require a four-person team and a fortnight now requires a prompt and the time it takes a model to render.

This is not a marginal improvement to virtual production. It is a *category* change. The bottleneck of virtual production has, for the entire history of the discipline, been the cost and time of building the environment. When that bottleneck goes, what remains is exactly the human craft that the audience is paying for: blocking, performance, direction, lighting design, story.

In **Issue 8** of the newsletter, I noted that Sony Pictures had begun using Marble inside its virtual-production pipeline. The number the team quoted publicly was the one that should

have made the front page of every trade publication: *40× faster than the traditional workflow*.²⁷³ If you sit inside the legacy economics of a virtual-production house — where building a single environment is a six-figure, multi-week proposition — that number is not an improvement. It is a *re-platforming* of the discipline. In **Issue 12**, Disney showed off a “300,000 poses in an instant” demonstration that was conceptually similar — animation built on top of generative spatial infrastructure rather than against it.²⁷⁴ In **Issue 27**, Netflix and Eyeline released **Vista4D**, a system that converts live-action footage into navigable 4D point clouds.²⁷⁵ The pattern is the same across the studios: a quiet pipeline shift, not a marketing story, that takes the entire “building the environment” stage out of the critical path of production.

The world-model race

Marble was the first commercial product in this category, but it was not the only one. The autumn of 2025 and the spring of 2026 were essentially a foot-race between research labs to ship usable world models, and the pace of releases was so rapid that the *Dream Machine* readers’ WhatsApp group routinely had three or four new ones to discuss per week.

Google DeepMind’s Genie 3, named by *Time* as one of the best inventions of 2025, generated playable 3D worlds at 24 frames per second from text prompts, with consistency held for several minutes — and in January 2026 was made publicly available to Google AI Ultra subscribers in the U.S. through a prototype web app called Project Genie. At Google I/O 2026, **Project Genie** was extended with a **Street View** integration that lets users generate navigable simulations of real-world locations directly from Street View map data, collapsing the gap between *the world that exists* and *the world that can be generated*.²⁷⁶ **Meta** announced **WorldGen** in November 2025, framed as research that could generate walkable 3D worlds from prompts like “medieval village town square.”²⁷⁷ **Tencent** open-sourced **HY World 1.5**, a real-time world model framework, in December 2025, alongside the **Hunyuan 3D Studio** which integrated the company’s art-grade 3D generative model **3D-PolyGen 1.5**.²⁷⁸ **SpAlatial** launched **ECHO**, a spatial foundation model, in December 2025.²⁷⁹ Stanford AI Lab and others released **Wonderzoom** in January 2026, a multi-scale 3D world-generation model that let you “infinitely zoom into the details” of a generated environment.²⁸⁰ **OpenArt** launched its own world-generation product, **Worlds**, in March 2026.²⁸¹

The May 2026 wave was the most aggressive yet. **NVIDIA** released **SANA-WM**, a 2.6B-parameter open-source world model natively trained for 60-second video generation with explicit camera control — the first open-weight world model at meaningful scale, and a development whose long-term implications for the open-source-AI-tooling argument I make in **Chapter 16** are, in my view, substantial.²⁸² **Odyssey** released **Starchild-1**, which it described as “*the first ever real-time multimodal world model*” — a system that doesn’t just generate a world but understands and simulates it.²⁸³ **Apple** published **Headsup**, a large-scale, high-quality 3D Gaussian-head reconstruction pipeline built from multi-view captures of the kind a consumer iPhone can already produce — a continuation of the Apple-Personas-and-Gaussian-splat thread above.²⁸⁴ At the consumer end of the same wave, **WorldLens VR**

rolled out an AI-powered Quest feature that adds subtle 3D depth to ordinary Google Street View environments, making the existing planetary-scale street-imagery dataset navigable in VR.²⁸⁵

The most ambitious of all of these — and the one I think hints most clearly at where the category is going — was **Luma AI's UNI-1**, launched in March 2026 with the framing: “*When worlds become instant, the race shifts to better thinking.*”²⁸⁶ UNI-1 was the first commercial release I am aware of that *combined* world-model generation with what Luma called “reasoning” — that is, the model didn’t just generate a scene, it could plan, modify and iterate on the scene as a coherent agent. The pitch was that you would no longer have a fragmented pipeline of prompt → image → video → iterate; you would have a single unified creative system that thought before it created.

UNI-1 is, in my view, the most important *category* announcement of the spring of 2026, even if the product itself is still rough at the edges. It is the announcement that says: world models are not the end state. They are the *substrate* on which something else — reasoning-led generative creativity — gets built.

By **May 2026**, you could find world-model capabilities embedded in the consumer tools as well. **CapCut**, the consumer-grade video editing app, integrated ByteDance’s *Seedance 2.0* via the *Dreamina* product, giving phone-users the ability to generate spatial scenes alongside flat video.²⁸⁷ **Spark 2.0**, an open-source Gaussian-splat streaming framework, brought 100-million-splat scenes to web browsers at interactive frame rates.²⁸⁸ **Apple** confirmed in October 2025 that its Personas feature on Vision Pro and other devices was powered by Gaussian splatting under the hood, making this — for the millions of Apple device owners who had used the feature without knowing what it was — the most-deployed Gaussian-splat technology in consumer hardware.²⁸⁹

The category, in eight months, went from a research demo to a consumer feature.

The games industry, again

If world models are infrastructure, the industry that has been waiting for that infrastructure the longest is games.

The 2024 conversation in games about generative AI was, in significant part, about *flat assets* — concept art, textures, dialogue, music — and it was the conversation that produced most of the backlash. *Call of Duty: Black Ops 7*’s loading screens. *Anno 117*’s placeholder art “slipping through” the review process. *Fortnite*’s Chapter 8 controversy.²⁹⁰ The audience response, in every case, was visceral, and the studios learned, the hard way, that AI-generated 2D assets dropped into established franchises read to fans as a cost-cutting move, not a creative one.

The 2025–26 conversation in games is different in kind, because the AI is now being aimed at the *substrate* of the game — the worlds, the systems, the NPCs, the procedural infrastructure — and the audience response is, so far, much more nuanced.

NVIDIA, in partnership with Stanford, released **NitroGen** in January 2026 — a “plays-any-game” AI trained on 40,000 hours of gameplay across more than 1,000 games. The model

wasn't being pitched as a way to *replace* games; it was being pitched as the foundation layer for a new generation of AI-aware game agents and procedural systems.²⁹¹ **Google DeepMind's SIMA 2**, released in November 2025, was an agent that could play, reason and learn alongside humans in virtual 3D environments.²⁹² **Ubisoft** open-sourced its **CHORD** model in December 2025, for end-to-end PBR material generation, and ComfyUI nodes built on top of it within the same week.²⁹³ **Ubisoft's Teammates** — a voice AI tech demo first shown in November 2025 — promised a step-change in how NPCs would behave in next-generation titles. The team lead's hands-on framing, given to *Video Games Chronicle*, is the one I keep returning to: *"It's a tool first. We've been working on it for more than two years now, and our conclusion is that it's a super cool tool, but it's still a tool."*²⁹⁴ *Still a tool.* The whole AI-in-games debate, compressed into four words by the people inside Ubisoft who are actually building the thing.

The most interesting single release of the spring of 2026 was **YouTube's Playables Builder**, a closed-beta product launched in December 2025 that lets users create games with short text, video or image prompts, built on Gemini 3.²⁹⁵ The framing, when YouTube's product team described it publicly, was that *every YouTube creator* should have the ability to ship a playable game as easily as they currently ship a video. Within months, **Unity** announced an "AI Open Beta" — an in-editor AI suite that brought the same logic to the professional games-development pipeline.²⁹⁶

Where this lands, in 2027 and 2028, is the question I find the most strategically charged in the whole industry. If creating a playable, navigable world becomes a thing a YouTube creator can do in an afternoon, the boundary between *games* and *video* — which has been collapsing slowly for fifteen years, through platforms like Roblox and Fortnite and the proliferation of interactive content on social platforms — collapses fully. The next generation of creators will not think in terms of *making a video* or *making a game*. They will think in terms of *making a thing*, and the thing will, by default, be navigable.

What this means for film

I want to come back to film for a moment, because I think the consequences of world models for the film industry are bigger than the consequences for any other sector, and the least understood.

For the entire history of cinema, the discipline has been organised around a fundamental scarcity: *the cost of building the location*. Even when the location was real — a city street, a forest, a beach — capturing it required a crew, a lighting team, transport, permits, weather contingencies. When it wasn't real — when it was a sound stage, or a digital matte painting, or a CGI environment — the cost was, if anything, higher.

The entire industrial structure of cinema, from the location department to the gaffer's crew to the virtual-production house, exists because *the place is expensive to make*.

When the place becomes cheap — when a Marble-generated environment, exported as a splat, dropped into Unreal, lit interactively, can substitute for a \$200,000-per-day exterior shoot at almost any quality bar a hero shot — the industrial structure that organised cinema starts to

look like the manuscript-copying scriptorium did in 1450. The thing that was the bottleneck is no longer the bottleneck.

What replaces it? My best guess, six months into the transition, is *taste in places*. If everyone can generate a market square, the value of *choosing the right market square* — the one with the texture, the light, the cultural specificity, the lived-in-ness that makes a scene feel like it belongs to a real human story — becomes the new scarce skill. The location scout becomes the *world curator*. The production designer becomes the *spatial director*. The cinematographer becomes — even more than they already are — the person whose job is to find the *one camera move* in a near-infinite navigable space that tells the story.

This is, I think, an upgrade for the craft, not a downgrade. It moves the human contribution to the part of the work that humans actually do well — *judgement about what matters in a place* — and offloads the part of the work that has been a manufacturing problem for a hundred years.

The risk

I want to flag the risk too, because I am trying — and I am sure I will not always succeed — to be honest about the downsides.

If world models become the substrate of creative work, the *training data* for those models becomes a question of enormous cultural consequence. A world model trained on, say, the visual archive of Hollywood will generate scenes that look like Hollywood. A world model trained on the photographic archive of Mumbai will generate scenes that look like Mumbai. The *aesthetic monoculture* that the early image-generation models produced — that vaguely Pixar-flavoured, vaguely Marvel-flavoured, vaguely YouTube-thumbnail look that you can recognise in a thousand 2024 AI outputs at a glance — is at risk of being amplified, not reduced, when the medium moves from images to navigable spaces.

The companies that own the largest world-model training datasets in 2030 will, in a real sense, own the visual language of the next generation of cinema, games and immersive media. If those datasets are biased — towards English, towards the global North, towards Hollywood production design, towards the architectural and cultural visual vocabulary of a small number of wealthy cities — the entire interior life of the next generation of creative work will reflect those biases.

This is not a hypothetical. We are seeing it now. The publicly available world models, in mid-2026, do a startlingly good job of generating “Mediterranean market square” and “American suburb” and “Tokyo street at night.” They do a startlingly *thin* job of generating, say, “Lagos street at dusk during the rains” or “a contemporary Indigenous Australian community space” or “a Manchester terraced street in winter with the sodium lights coming on.” The bias is in the training, and the training is in the assets, and the assets were in the corpus, and the corpus was English-internet-skewed.

If we want the next creative economy to look like the world rather than like the AI companies’ biggest source datasets, the dataset question has to be a *first-order* design problem. Korin AI’s late-2025 launch — “trained with African datasets, built by Africans” — is the kind of intervention that is going to have to multiply.²⁹⁷ So is the African Tech / India / Singapore-led

wave of culturally-specific AI cinema that the trade press started covering in Issues [20](#) through [27](#). Diversity in training datasets, for the world-model era, is not a content-moderation question. It is a *cultural infrastructure* question.

Six craft questions for the world-model era

Before I make the big claim, I want to put six craft questions on the page that working creatives — directors, designers, art directors, cinematographers, sound designers, level designers — will, by my estimate, be wrestling with for the rest of the decade. They are the world-model-era equivalents of the craft questions the *cinematographic* era took fifty years to develop a vocabulary for (where do you put the camera, how do you light the scene, how does the cut work, how does the sound do its work). The world-model era has, in 2026, no settled vocabulary for any of them. The vocabulary will be built by the working creatives who notice the question first.

One. Where does the audience stand? The most-overlooked craft question of the navigable-space era. A film positions the camera; the camera positions the audience. A world model produces a *navigable space*; the question of where the audience *enters* the space, where they are *invited* to stand, what they are *encouraged* to look at, is no longer fixed by the cinematographer. It is fixed — if it is fixed at all — by the *narrative scaffolding* the orchestrator builds around the navigable space. Marble’s October-2025 update added explicit *suggested-camera-pose* primitives for exactly this reason. The craft question is which of those poses to specify and which to leave to the audience.

Two. How does the cut work in a navigable scene? The film cut depends on the audience being in a fixed position; the editor moves the camera between fixed positions in a way that the audience’s eye follows. The navigable scene has, by default, no cut. The audience moves through it continuously. The craft question — for working directors and editors — is when to *break* the continuity, how to do so in a way the audience reads as deliberate rather than as a technical glitch, and what new grammar of transitions a navigable medium permits. Some early experiments in 2025–26 have used *spatial discontinuities* (an audience walks through a door and emerges in a different space) and *temporal discontinuities* (the same space at different times) as cuts. None of these has yet stabilised into a shared grammar.

Three. How does performance survive the medium? A film performance is captured by a camera at a fixed angle and pace. A world-model performance — a synthetic actor performing inside a navigable scene — has to be authored such that the performance *works* from every angle and every speed at which the audience might encounter it. This is, for working performers and motion-capture supervisors, an entirely new craft challenge. The film-era cliché of the actor “playing to the camera” is, in the world-model era, replaced by *playing to the spatial neighbourhood* — knowing that the audience may be six feet away, may be inside the actor’s eyeline, may be behind the actor’s shoulder, may be looking at the actor from above. Volumetric capture (Vista4D’s live-action 4D reconstruction, NVIDIA’s D-Rex digital-human pipeline) is the technical answer. The *performance* answer — what acting *means* in a medium where there is no fourth wall — has not been worked out.

Four. What does sound design do in a navigable scene? A film sound mix is, for the most part, a fixed track timed to the picture cut. A navigable-scene sound mix has to *follow*

the audience. Spatial-audio tooling (the SonicLab SPATAI pipeline, Dolby Atmos for VR, the various Meta-and-Apple immersive-audio platforms) is the technical answer. The craft question is, again, what *good* spatial sound design looks like in a medium where the audience-author relationship has changed.

Five. What is the *running-time* of a navigable scene? A film has a fixed running time. A navigable scene does not. The audience could leave after thirty seconds or stay for two hours. The craft question for the working director is how to design the experience so that *both* extremes produce a satisfying piece of work. Games have, for fifty years, been grappling with this question — the *Dark Souls* answer (every player gets a different running time depending on skill and exploration) is different from the *Outer Wilds* answer (the running time is gated by narrative discovery) is different from the *Telltale Games* answer (the running time is broadly fixed across players). World-model cinema, in 2026, has not yet settled on its equivalent.

Six. What is the *single best moment* of a navigable scene? Film has scenes — discrete units of dramatic action with a recognisable shape, a recognisable peak, a recognisable end. A navigable scene, by default, does not. The craft question for the working director is whether to *design* the navigable scene around a single peak moment (which the audience may or may not reach) or to design it as a *texture* (which the audience experiences at whatever density their navigation produces). The peak-moment design pulls the medium back toward film conventions; the texture design pushes it toward something more like architecture or landscape design. Different working directors will, on the historical pattern, settle on different answers. The grammar will, over a decade, stabilise into a working vocabulary the way the cinematic-cut grammar stabilised between 1903 and 1925.

The six questions are not, in 2026, *theoretical* problems. They are the questions the working spatial-cinema teams I have talked to — the Wonderzoom group at Stanford, the World Labs developer cohort, the early adopters at Sony Pictures and Eyeline — are wrestling with on Wednesday afternoons. They are also, on the historical pattern of [Chapter 2](#), the questions whose answers will define what working creatives in the next decade are *paid* to do. The directors and designers who develop a working vocabulary for them first will, on the available evidence, become the named *Walter Murchs* of the spatial-cinema era. The ones who wait for the vocabulary to settle will, in retrospect, look like the editors who waited too long to learn Avid.

The big claim

Let me make the big claim, and then move on.

I think — and this is the most non-obvious bet in this book — that the **world model is the medium of the next twenty years of creative work**, in the same way that the **moving image** was the medium of the twentieth century and the **interactive screen** was the medium of the first quarter of the twenty-first.

I think people who are working in flat-video, flat-image, flat-audio formats in 2030 will increasingly be working in a *legacy* format — still alive, still culturally valuable, still where the highest-end of the craft lives, the way live theatre or vinyl-record production still lives — while

the *dominant* mode of creative work will be the production, curation, performance and distribution of navigable spaces.

I think the studios, platforms and tool companies that are quietly investing in world models now — World Labs, DeepMind, Meta, NVIDIA, Tencent, Luma, Apple — will be the ones that set the rails for the next two decades.

I think the audience, having developed the antibodies described in Chapter 5 to slop-grade flat AI content, will eventually develop a parallel set of tastes for *navigable* content — and that the question of what makes a *good* AI world (rather than a *good* AI video) will be the central craft question of the late 2020s.

And I think — most importantly for the next chapter — that the toolchain to make all of this is being built, right now, by a small number of platform companies who have started saying out loud that AI is going to be *in everything, everywhere, all at once* — and who are, while you are reading this paragraph, designing the rails on which the next creative economy will run.

Chapter 9 — AI in Everything, Everywhere, All at Once

In late October 2025, at Adobe MAX, the company that has made the software almost everyone in the creative industries uses every day — Photoshop, Illustrator, Premiere, After Effects, InDesign — decided that the year-old marketing line “*AI is a feature in our tools*” had outlived its usefulness, and replaced it with a more honest one.

The new line was: “**AI in everything, everywhere, all at once.**”²⁹⁸

The reason I want to spend a chapter on that phrase is not because I love a slogan. The reason I want to spend a chapter on it is that I think it is, more than any other single piece of corporate positioning from the period this book covers, *literally* true. AI is in everything now. It is in every layer of the creative software stack. And the implications of that for working creatives — for the way we are trained, the way we are paid, the way we work with each other — are not yet, in the spring of 2026, fully understood.

This chapter is about the platform layer. About the companies that make the tools that the rest of the creative industries use to make the work. About how those companies have, in the past eight months, accepted that their business is no longer making *tools* but making *agents*, and about what that means for the rest of us.

The Adobe MAX week

The Adobe MAX 2025 keynote — held in mid-October in Los Angeles, the week after OpenAI’s DevDay, two weeks after Tilly Norwood — was unusual, by Adobe’s standards, in how much it tried to land at once.

The headline products were Firefly Foundry, a service for companies to train their own custom generative models on their own visual identity;²⁹⁹ Firefly Image Model 5, the latest generation of the image generator that has, since 2023, been Adobe’s primary public answer to Midjourney and Stable Diffusion;³⁰⁰ and an AI Assistant built directly into Adobe Express, the company’s lower-barrier consumer creative tool.³⁰¹

Underneath the headlines was a much longer list of “Project” announcements — Adobe’s research-preview format, the things that may or may not ship but that signal what the company is investing in. The list, looked at as a whole, is what convinced me, sitting at my desk in the North West watching the live stream, that something larger than a product launch was happening:

Project Scene It: image-to-3D and 3D-to-image technologies, with reference-image tagging for object preservation in 3D space.

Project Surface Swap: AI-powered material recognition, letting designers swap textures while preserving lighting, shading and perspective.

Project Turn Style: editing 2D objects as if they were 3D.

Project Trace Erase: removing objects *and* their shadows, reflections and environmental distortions in one operation.

Project New Depths: editing depth in an image as easily as adjusting brightness.

Project Frame Forward: applying changes across entire videos based on one annotated frame and a text prompt — “the precision of photo editing in video workflows.”

Project Motion Map: bringing static vector graphics to life automatically.

Project Sound Stager: analysing a video’s visuals, pacing and emotional tone, and automatically generating layered soundscapes.

Project Clean Take: AI correction of mispronunciations, voice isolation, noise removal and delivery refinement.

Project Graph: a node-based workflow editor, conceptually similar to ComfyUI, for chaining Adobe’s tools and models into custom pipelines.³⁰²

There is, in that list of ten projects, *every* layer of the post-production stack — image, video, 3D, audio, layout, workflow — being re-imagined as a generative or agentic operation. Not a tool with an AI feature stapled on. A *generative-first reimagining of the operation itself*.

The Adobe MAX week was, to put it plainly, Adobe’s announcement that it was rebuilding its product from the inside.

What “AI in everything” actually means

The reason I want to be careful with the Adobe-MAX framing is that, six months on, you can see how literally the company has executed against it.

In December 2025, Adobe announced that Photoshop, Express and Acrobat editing would be available *inside* ChatGPT — meaning the creative output was no longer happening inside Adobe’s interface, but inside an AI agent’s.³⁰³ In January 2026, the Premiere Object Mask tool — an AI-driven masking feature that automated one of the most laborious tasks in video editing — quietly became available to Premiere users.³⁰⁴ In late January, at Sundance, Adobe launched the *Adobe Film & TV Fund* and *Ignite Day*, with explicit support for filmmakers integrating AI into their workflows.³⁰⁵ In April 2026, at the **Adobe Summit**, the company introduced its **CX Enterprise** platform alongside NVIDIA — a stack of AI agents embedded across the entire content lifecycle from brief to delivery — under the framing “agentic creative intelligence is now.”³⁰⁶

The trajectory, in one sentence: Adobe in 2024 was a *creative tool company*. Adobe in 2026 is an *AI-agent platform company* that happens to also still ship Photoshop.

If you are wondering whether this transition has been smooth: it has not. The reception of the Adobe AI announcements among working creatives has been, in my own circles and the readers’ WhatsApp group the *Dream Machine* community runs, sharply ambivalent. There is real appreciation for the productivity gains. There is real anxiety about the implications for craft, for licensing, for control, for the trajectory of the company’s relationship with the creators who pay for it.

What no working creative I know thinks is that this transition is reversible. Once Photoshop has an AI assistant baked in, once Premiere has Object Mask, once After Effects has the AI-powered animation tools that landed in November 2025,³⁰⁷ the *next* version of every Adobe

product is going to have *more* of this, not less. Adobe's *competitors* are, if anything, going faster. If Adobe slows down, somebody else lands the punch.

This is — I think this is the part that working creatives have to understand and internalise — *the new physics of the toolchain*. AI is not a feature that one tool company decided to ship. It is a structural property of the toolchain itself in 2026, and the question for anyone using that toolchain is not whether to integrate AI but *how to integrate it deliberately*, with eyes open, on terms that preserve the human craft underneath.

The platform alliance

Adobe is not — and this is the more important observation — the only company doing this.

In March 2026, *Dream Machine* Issue 21 led with what I have called, in talks since, the most consequential business announcement of the year: **Adobe + NVIDIA** entered a strategic partnership that explicitly framed creative AI as *enterprise infrastructure* rather than viral consumer tooling.³⁰⁸ The partnership covered next-generation Firefly models, agentic creative-and-marketing workflows, and production-pipeline integration. The language was telling: *precision and control* for creativity and marketing pipelines, alongside content, campaign and production speed.

The reason this is consequential — beyond the size of the two companies involved — is that it signals the *maturation* of the market. Adobe + NVIDIA is not a race-to-the-cool-demo deal. It is a *race-to-the-procurement-line* deal. The two companies are betting, jointly, that the next era of creative AI is going to be won by whoever ships the most reliable, most controllable, most legally-defensible production-grade tooling to the enterprise creative buyers — the studios, the agencies, the broadcasters, the brand teams.

The same week, **Google** and **NVIDIA** announced a parallel deal for cloud-based generative-AI infrastructure aimed at the same enterprise market.³⁰⁹ **Hugging Face** and **Google Cloud** announced a partnership in November 2025 covering open-source agentic development.³¹⁰ **Meta** and **Hugging Face** launched **OpenEnv** in October 2025 to advance open-source agentic development.³¹¹ **Anthropic** signed a corporate-patronage deal with the **Blender Foundation** in May 2026.³¹² **Anthropic** also acquired into the **Slack** workplace-tooling ecosystem with Claude Apps in January 2026,³¹³ and reached an ad-sales partnership with **Spotify** to put music recommendations inside Claude.³¹⁴ In May 2026 **Splice** signed a “Responsible AI” deal with **ElevenLabs** covering sample-library training and consented voice synthesis;³¹⁵ **Netflix** announced an agentic ad-tools roadmap whose internal framing — “*agentic AIs talking to each other*” — was an unusually candid description of where the advertising-orchestration layer is heading,³¹⁶ and the AI-coworker startup **Viktor** raised \$75M to embed an agentic colleague directly into Slack and Teams,³¹⁷ reinforcing the pattern that the platform-layer agents are landing where the working creative already lives.

The advertising holding companies were moving at the same pace. **WPP** signed a \$400m partnership with Google in October 2025.³¹⁸ **WPP Open Pro**, a new edition of the agency's AI marketing platform, launched the same month with a framing that should be read carefully

by anyone working in adland: “While some companies hide their AI behind service teams or focus on just one part of the journey, WPP Open Pro is an integrated solution for campaign implementation, built to deliver outcomes, not just assets.”³¹⁹ Outcomes, not just assets. That is the position of a holding company that has decided AI is not a feature — it is the entire reason a brand should buy from them in 2026. **WPP** then expanded its AI capabilities through a partnership with **Sightly** in November 2025.³²⁰ By April 2026, WPP was using Google Earth’s AI tools to map consumer journeys at scale.³²¹

The pattern is unmistakable. The platform layer — the toolmakers, the infrastructure companies, the agencies, the cloud providers — has been quietly consolidating around a small number of strategic alliances that, taken together, are deciding the *rails* on which creative work will run for the next decade.

If you are a working creative reading this, you are probably already running some part of your workflow on rails laid by one of these alliances. By 2028, you will, almost certainly, be running *most* of your workflow on those rails — or on a deliberate, principled alternative that has chosen to opt out.

Google I/O 2026: the second platform-layer moment

If Adobe MAX 2025 was the platform-layer announcement of the first half of this book, **Google I/O 2026** — held in the week this book went to press — was the announcement that closed it. The two events bookend the eight-month window the manuscript covers, and the symmetry of their framings is, on the platform-economics read, instructive.

The headline announcements were **Gemini Omni**, a unified multimodal model designed to work across text, image, audio, video and live interaction in a single workflow; **Antigravity**, Google’s agentic coding and development environment; **Google Flow**, the agent-based workflow product that allows AI systems to take on multi-step creative and production tasks autonomously; **Gemini Spark**, the new developer toolkit aimed at building autonomous agents and AI-powered applications; and **Project Genie + Street View**, an integration that allows users to generate navigable simulations of real-world locations from the Street View map data — a topic I return to at length in Chapter 8.³²² The keynote opened, deliberately, with a browser-based multiplayer demo called **Infinite Scaler** — thousands of players competing inside vertically generated levels created on the fly from player prompts — a piece of theatre that, like *Tilly Norwood* a year earlier, was less interesting for the product than for the framing it imposed on the event that followed: AI-generated worlds, live procedural experiences, mass participatory systems that evolve in real time.³²³ **SynthID**, Google’s content-provenance watermarking technology, was announced as having marked over 100 billion items by May 2026, and as being extended to partner ecosystems including OpenAI, ElevenLabs and Kakao — a development I cover in detail in Chapter 12.³²⁴

The reason to draw the Adobe MAX / Google I/O parallel directly is that the *structural shape* of the two announcement weeks was identical. Both keynotes argued that AI was no longer a feature to be added to existing products; it was the *operating layer* into which the existing products would be re-built. Both keynotes pre-positioned the company’s product roadmap around *agents* rather than tools. Both keynotes were aimed not at the consumer-keynote

crowd but at the procurement teams of the enterprise creative buyers. The fact that the same framing arrived from the two largest creative-software and creative-platform companies in the world, eight months apart, in the same eight-month window, is the cleanest single confirmation I have that *the AI-in-everything thesis is the platform layer's settled commercial strategy* for the rest of the decade.

For working creatives, the operational implication is direct. The platform layer has decided. The question is no longer whether the creative software stack will be re-built around AI agents. It is which agents, on whose terms, with what provenance, on which commercial settlement.

The new entrants

Underneath the platform giants, a separate layer of companies has been building the *consumer-facing AI creative tools* that, in some markets, are turning into bigger businesses faster than anyone expected.

Higgsfield, the AI video startup focused on social-media marketers, raised \$80m at a \$1.3bn valuation in January 2026.³²⁵ Three months later — in a stat that I have read repeatedly to check that I have not got it wrong — Higgsfield was reported to have earned \$200m in nine months of operations.³²⁶ An AI-video startup, less than two years old, was running at a quarter-billion-dollar annual run-rate by the spring of 2026.

Synthesia, the U.K.-based AI-avatar platform, hit a \$4bn valuation in January 2026 and let its employees cash in.³²⁷ In October 2025 it had reportedly *rejected* a \$3bn acquisition offer from Adobe — choosing to remain independent.³²⁸

ElevenLabs, the audio-AI company, was reported to have crossed \$500m in annualised revenue by April 2026, raising from BlackRock, NVIDIA, Jamie Foxx and Eva Longoria.³²⁹

Runway released Gen-4.5 in December 2025 and Gen-4.5 Image-to-Video in January 2026, then a “Workflows” product across all paid plans, then a Story Panels app, then a Characters API, then Apps for Advertising — and by spring 2026 was making the public case that AI could enable “50 indie films” instead of “one \$100M blockbuster.”³³⁰ In May 2026 the company opened a Tokyo office on a \$40M commitment, marking its first material expansion into the Asia-Pacific creative-AI market.³³¹

Krea, Freepik, Magnific, Heygen, Hedra, Cascadeur, Hunyuan, Kling, Suno, Udio, Mureka, Hitem3D, Meshy, Rodin — the list of consumer-grade AI creative tools that crossed material commercial scale in this period is too long to fully enumerate, and the *Dream Machine* archive carries them week by week.³³² The category that didn't exist in 2023 is now an industry with multiple unicorns, multiple billion-dollar valuations and meaningful real revenue.

ComfyUI, the open-source node-based workflow tool that has become a quiet standard for technical AI users, raised \$17m in October 2025³³³ and hit a \$500m valuation by May 2026.³³⁴ What the ComfyUI valuation tells you, more than any of the big-platform numbers, is that the

market is also paying — at significant scale — for tools that give *creators control* over the AI process rather than abstracting it away.

The free tier and the literacy tier

Two things happened in the consumer-platform layer that I think have been under-discussed and that matter a lot for what the next creative economy will look like.

The first is that the **base layer of AI capability went free**, in a meaningful sense, in the autumn of 2025. **Google** released its **Pomelli** marketing AI agent for free in October.³³⁵ **Google AI Studio, Opal** (the no-code AI mini-app builder), and the Project Genie prototype were all released as free or near-free tiers through early 2026.³³⁶ **Lovable** made its product free for teachers and students in classrooms.³³⁷ **Adobe Express's** AI Assistant arrived inside the free tier of Adobe's already-free consumer product.³³⁸ **Hugging Face** continued to expand its free hosting and open-source model distribution.³³⁹ **Krea, Freepik**, and many of the larger tool platforms kept generous free tiers as a customer-acquisition lever.

What this means, practically, is that the entry-level for AI-enabled creative work in 2026 is *near zero*. A teenager with a phone and a free Google account can, today, generate video, music, 3D objects and (with Project Genie) navigable interactive worlds at a quality bar that, two years ago, required a small production company to produce.

This is, in absolute terms, a democratisation. It is also — and this is the second thing — *creating a literacy gap* between the people who know how to use these tools well and the people who don't.

Adobe responded to this gap, in late 2025 and through 2026, by becoming — in addition to a software company — *a training organisation*. The Sundance partnership, with a \$2M investment to teach 100,000 filmmakers AI skills.³⁴⁰ The Ignite Day, focused on emerging creators.³⁴¹ The Adobe Film & TV Fund. The Adobe Express AI Assistant tutorials. Google made the same bet in parallel, putting \$40bn into Anthropic in May 2026 in a deal widely interpreted as betting on the literacy and infrastructure layer of the next decade.³⁴²

The UK government picked the same direction. In January 2026, the Department for Science, Innovation and Technology announced *Free AI training for all*, expanding a government-and-industry programme to provide 10 million UK workers with AI skills by 2030.³⁴³ The Department for Business and Trade research, reported in *Dream Machine* Issue 7, found that *neurodiverse workers* were 25% more satisfied with AI assistants — suggesting that AI's productivity benefits in certain workflows could “potentially help to level the playing field.”³⁴⁴ The University of Wisconsin-Stout set AI-use as a baseline competency in its filmmaking course in January 2026.³⁴⁵

What the consumer-platform companies and the policy-makers are, jointly, building is a *training infrastructure* for the new toolchain. The reason they are doing this is straightforward: a tool you cannot use is a tool you do not buy, and a worker who cannot use the new tool is a worker who eventually exits the workforce. Both incentives push in the direction of mass AI literacy as a public investment.

What I find encouraging about this — and I am genuinely encouraged, against the grain of much of the cultural commentary — is that the literacy push is being framed, both by Adobe at Sundance and by the UK government, as *creator empowerment* rather than worker replacement. The proposition is not *learn AI or be replaced by it*. The proposition is *learn AI to remain in the driver's seat of your own work*. That framing matters. It is the right framing. It is the only framing under which the AI-literacy push doesn't become a way of accelerating the very problems it is supposed to fix.

The platform economics underneath the slogan

I want to spend a section on the *commercial* shape of the platform layer, because the “AI in everything” framing has economic implications that the keynotes have been careful not to name, and that working creatives buying platform access at scale need to understand.

The shape, simplified, is this. The dominant generative-AI platforms — OpenAI, Anthropic, Google DeepMind, Adobe, Runway, ElevenLabs and the rest — operate as *infrastructure-as-a-service* businesses on top of *capital-intensive* underlying compute. The marginal cost of producing one more generated image, song or video clip is, at the platform level, low. The fixed cost of the compute infrastructure required to produce *any* generated output at competitive quality is, at the platform level, very high — the data-centre build, the chip supply, the energy contract, the model-training spend.

This produces, structurally, a *natural-monopoly-tending* market. The platform with the largest compute base produces the lowest marginal cost per output, captures the largest user base, generates the largest revenue, and reinvests in a larger compute base. The flywheel is the standard cloud-services flywheel, accelerated by the AI-specific dynamics of training-data flywheels and user-feedback flywheels.

The 2025–26 financial telemetry, where the platform companies have disclosed it or been required to disclose it, supports the natural-monopoly read.

OpenAI was reported, through late 2025 and into 2026, to be operating at *significantly negative cash flow* at the unit-economics level despite its 800–900M weekly active users. The company's reported revenue, which crossed the \$10 billion annual run-rate mark in late 2025, was — by every analyst breakdown I have seen — being substantially exceeded by infrastructure costs (data-centre lease, chip supply, energy, training compute). The Microsoft partnership at the financial level was, structurally, a *capital-supply* relationship rather than a *technology-licensing* one: Microsoft providing the compute capacity that OpenAI could not, by itself, finance.

Anthropic, in 2026, was reported to be operating with similar structural dynamics, with Google and Amazon as its capital-supply partners. The Anthropic Foundation patronage deal with the Blender Development Fund — announced in May 2026 and discussed in [Chapter 16](#) — is interesting precisely because it suggests Anthropic has, alongside the closed-platform business, *strategic interest* in supporting the open-source creative-AI infrastructure that the closed-platform model competes with. That kind of two-handed positioning is, in natural-monopoly markets, often the precursor to a *platform regulation* settlement.

Adobe, by contrast, operates with the most *defensible* business model in the creative-AI platform space, because it is selling AI as a feature of an existing subscription rather than as a per-use service. Firefly’s contribution to Adobe’s 2024 annual recurring revenue — 11% of new ARR, on the company’s own published numbers — is being generated *without* the per-token unit-economics problem that OpenAI and Anthropic face, because Adobe is bundling the AI into the existing \$54.99-a-month Creative Cloud all-apps subscription. The customer’s behaviour change from no-AI to AI doesn’t change the revenue line. It changes the *value capture* of the existing revenue line. This is, in business-school terms, the platform’s *strongest* possible defensive position. It is also the reason Adobe’s stock performed differently from the rest of the AI-platform cohort through 2025–26.

Runway, ElevenLabs and the AI-native specialist platforms operate with a per-use unit-economics structure that more closely resembles OpenAI’s. The differentiation, where they have it, is in *workflow integration* — Runway’s Workflows product, ElevenLabs’ Flows canvas, the studio-tier features that lock professional users into per-platform tooling. The strategic question for each of these companies, in the next two years, is whether they can build defensible workflow lock-in before the natural-monopoly dynamic of the underlying foundation-model market consolidates the foundation-model layer down to two or three players.

The implications for working creatives buying platform access at scale, in 2026 and beyond:

One, the per-token / per-output prices working creatives are paying for AI tooling in 2026 are, on every analyst read I have seen, *materially subsidised by platform-company investor capital*. The unit economics underneath the prices are not, today, sustainable at the volumes the platforms are producing. The prices are, by structural inference, going *up* over the medium term as the platforms work toward unit-economic break-even. The working creative who builds a business model assuming today’s per-token costs as a stable input is, on the platform-economics read, taking a bet that the platforms cannot win. *Pricing today is not pricing forever*. This is the part of the platform-dependency argument the *open-the-black-box* discussion in [Chapter 3](#) most directly relies on.

Two, the *strategic-rent-extraction* potential of the eventual platform-monopoly position is the structural risk underneath the entire orchestrator economy I described in [Chapter 11](#). If two foundation-model platforms dominate the underlying generative-AI capacity by 2030, and every working creative’s production pipeline depends on access to one or both of them, the platforms will be in the position the cable companies were in by 2010 and the social-media platforms were in by 2015 — able to extract value from the working creators who depend on them at prices the creators have no real ability to negotiate. The Petrillo-template solution to this is *collective bargaining* by working creatives and their unions against the platform companies as a class. The early architecture of this — the Cannes Disclosure Standard, the SAG-AFTRA platform negotiations, the EU AI Act enforcement, the UK 88% — is in place. The substance of it is, in mid-2026, still mostly aspirational.

Three, the *open-source alternative* layer documented in [Chapter 16](#) is, on this structural read, *the working creative’s principal long-term insurance policy* against platform-monopoly pricing. The 80% of YC and a16z startups now building on open-weight models — Hunyuan,

Wan, Qwen, FLUX, DeepSeek, the various Mistral variants — is, in market-economics terms, *the credible-walk-to-alternative* that constrains the closed-platform companies' pricing power. The working creative who has familiarised themselves with open-weight tooling, even if they default to closed-platform tooling for most of their daily work, has a *strategic option* the working creative who has not has surrendered. The option is worth money. It is also, on the historical pattern, worth political leverage in the institutional negotiations that the next decade of platform-rule-writing will run on.

The “AI in everything” framing, in operational summary, is *the platform companies' commercial strategy described in marketing language*. The strategy is to make AI a default productivity feature of every creative workflow, on platform-controlled tooling, at prices that the platforms can adjust over time once the workflow lock-in is in place. The strategy is, on the historical pattern of every previous platform-economics moment, going to produce a settlement somewhere between *the most-extractive version of the strategy* and *the most-constrained version of the strategy*. The 88%, the SAG-AFTRA contract, the EU AI Act, the C2PA standards body, the open-source ecosystem, and the working-creative collective-bargaining infrastructure are the constraints. The platform companies' compute capital, distribution leverage, and product-design control are the extractive forces. The settlement will be wherever those forces balance.

What we lost

I want to close this chapter with the harder question, because the “AI in everything” framing has a cost that the platform-company keynotes are not, on the whole, eager to discuss.

What we lost, in the transition to AI-in-everything tools, is *the deliberate friction of the old creative process*. The thing that made Photoshop, for many of its early users, a profound creative tool was not just what it could do. It was that it required you to know it. The interface was a discipline. The keyboard shortcuts were a vocabulary. The layers, the masks, the channels, the curves, the colour pickers — they were the language of a craft, and learning the language was part of becoming the craftsperson.

When the layer of mastery moves from the toolchain to the prompt, the *barrier* of mastery drops to near zero. That is the democratisation we have been promised, and it is real.

What goes with the barrier, though, is the *depth of relationship* between the maker and the tool. The Photoshop user of 2015 knew the tool the way a guitar player knows a guitar — with their hands, with their body, with a relationship built up over years of repeated, embodied practice. The prompt-driven AI tool user of 2026 has a different relationship. It is more like the relationship of a director to a department head: you describe the result, the department head executes, you adjust by giving notes.

The motion designer Doug McGinness, posting on LinkedIn about the new AI-augmented After Effects workflow in late 2025, summarised the current state of the tooling in a single, ruefully accurate line that became a small private meme inside my studio: “*export → prompt → pray → import.*”³⁴⁶ The line is funny because it's true. The current generation of AI-tooled creative work is, for a substantial portion of every day, an exercise in *committing to a black-*

box operation and accepting whatever comes back. That is, structurally, a different kind of creative discipline than the deterministic-tool craft it is replacing.

Neither relationship is *better* than the other. They are different relationships, and they produce different kinds of practitioners. But the *transition* is real, and one of the consequences — which I have seen up close, watching young creatives come through the studio — is that the *cognitive engagement* with the medium is structurally less deep than it used to be. The maker is one further step removed from the material.

This is not, by itself, a tragedy. The cinema director is one step removed from the film stock and is still, recognisably, the author of the film. The composer is one step removed from the violin and is still, recognisably, the composer. The novelist who uses a word processor is one step removed from the page and is still, recognisably, a writer.

But it is a *change*, and it is one we are pretending not to notice. The new toolchain is not just faster than the old toolchain. It is also a different kind of relationship with the work, and the people who will be its best practitioners — the ones who will, in 2030 and 2035, be doing the AI-era equivalent of what Greg Lynn did with parametric architecture or what Bjork did with synthesisers — will be the people who *consciously cultivate* the depth of relationship that the toolchain no longer enforces.

The platform companies are not going to teach you to do this. They have no incentive to. They benefit from your dependency, not your mastery. The new toolchain is *frictionless*, and frictionless tools, however much we benefit from their efficiency, are not, on their own, going to produce the next generation of great creative work.

That work is going to come from the people who put the friction *back in*, deliberately, on their own terms — who treat the AI agent as a junior colleague rather than as an oracle, who insist on *understanding* what their tools are doing rather than just *using* them, and who maintain the cognitive engagement with the work that the tools have been designed to make optional.

In the next chapter, I want to talk about the people who are doing exactly that. The orchestrators.

Chapter 10 — What is newly possible

I want to spend a chapter, after eight chapters that have been mostly about *displacement*, on a question I think the book has so far under-served. *Which categories of creative work has the AI moment made possible that were not possible before?*

This is not the question working creatives in 2026 are most often being asked. The questions most often being asked are *will I lose my job?* and *should I use the tools?* and *what are studios doing?* Each of those is in this book, in a chapter of its own. They are all important questions.

The question of *newly-possible work* is the one that, on the historical pattern I drew in Chapter 2, almost always turns out to be the one that mattered most. Every previous creative-technology transition the book has documented produced a set of new categories of creative work that the displaced practitioners did not, and could not, see coming. The phonograph displaced amateur parlour music and created the recorded-music industry. The microphone displaced operatic vocal projection and created intimate vocal styles, jazz singing, pop crooning, audio storytelling. Non-linear editing displaced the splice and created the MTV cut, the music video as art form, the hyper-cut action grammar, the streaming-era serial-drama editorial rhythm. The smartphone-as-camera displaced the dedicated camera and created the entire grammar of vertical-video native form.

The pattern is that the new categories *always* appear, that they *always* appear faster than the displaced cohort predicts, and that they are *always* invisible from the perspective of the existing definition of the craft, because they are made of capabilities the existing craft does not have. The miniaturist could not, in 1845, predict Stieglitz's *Camera Work*. The session keyboard player could not, in 1982, predict Aphex Twin's *Selected Ambient Works* (and could not, if they had heard them, have recognised them as music). The print-magazine art director could not, in 2006, predict Instagram-as-fine-art-platform.

We are, in 2026, the equivalent generation to those people. The new categories are mostly invisible from where we sit. But some of them are already shipping; some of them are already finding audiences; some of them already have working creatives building careers inside them. This chapter is about what those categories are.

Before I start the inventory, I want to draw two historical analogues out — the synthesiser and non-linear editing — because they are the cleanest templates for *how a tool that begins as a faster version of the old thing ends up being the substrate of a new thing entirely*. The mistake almost everyone made about AI in 2024 and 2025 was to think of it as a faster way to make the kind of creative work the platforms had been making. The historical pattern says: that view is *almost always wrong on a five-to-ten year timeline*.

From imitator to instrument: the synth analogue

When Robert Moog first started selling modular synthesisers in the late 1960s, the cultural permission for the instrument was very narrow. The synth, in its first commercial moment, was understood as a way to *imitate* existing orchestral instruments — to play the parts of a string section, a brass section, a piano, an organ, in a form a single keyboard player could control. The breakthrough commercial release that secured the synth's cultural status — Wendy Carlos's *Switched-On Bach* in 1968 — was, on its face, a literal demonstration of this

framing: the synth playing the music of the most-canonical European classical composer in the literature. Three Grammys. The first electronic record to be reviewed seriously by classical critics. The argument was: *the synth can do what an orchestra can do*.

The synth never did, in the end, only do that. The 1970s and 1980s did something nobody at the 1968 reviewing desks predicted. They produced *sounds that had never existed in the history of music* — Moog leads, FM electric pianos, the Roland TR-808 kick drum, the Yamaha DX7's chord pad, the *Blade Runner* CS-80 ambient texture, the Aphex Twin acid bassline — and they built entire musical genres around those new sounds. Hip-hop, electronic dance music, ambient, IDM, synth-pop, the entire sonic vocabulary of 1980s film scoring — these are forms that *could not have existed* without the synth, that did not exist before the synth, and that, *crucially*, could not have been predicted from the framing in which the synth was first introduced.

The synth was not a faster orchestra. The synth was an instrument for making sounds that no orchestra could produce, on a timescale that no orchestra could match, accessible to working musicians without the social capital of orchestral training. The instrument's first cultural moment — *Switched-On Bach* — was the moment of the *imitator framing*. Its second cultural moment — *Autobahn*, *Blade Runner*, *Thriller*, *Acid*, the *Roland 808* in *Planet Rock* — was the moment when the imitator framing was thrown away and the instrument was used for what only it could do.

That second moment took about a decade. From Moog's first commercial modular in 1965, through Carlos in 1968, through Kraftwerk's *Autobahn* in 1974, through *Trans-Europe Express* in 1977, through the early hip-hop and dance records of 1981–82 — the gap between the synth-as-imitator and the synth-as-new-instrument was roughly fifteen years. The cultural permission to use the synth as itself, rather than as a replacement for something else, had to be earned by working musicians inhabiting the instrument in front of audiences who eventually understood that they were hearing something new.

The AI equivalent moment, on the historical pattern, has not yet arrived. We are, on the synth timeline, somewhere between *Switched-On Bach* and *Autobahn*. The work that is going to define AI as a new creative substrate — rather than as a faster way to produce existing work — is being made *right now*, by working creatives somewhere, in forms that the trade press has not yet decided what to call. I have my candidates, which I will come to. The point I want to make at the chapter's opening is the structural one: every iterative-technology framing of AI is doing what the *Switched-On Bach* reviewers did in 1968. It is describing the new tool in the language of the old one. The new language has not yet been written.

From workflow to grammar: the non-linear-editing analogue

The same pattern is visible, in a different domain, with non-linear editing.

When Avid Media Composer shipped in 1989, the cultural framing was strictly utilitarian. NLE was a *faster way to do the existing thing*. Where a working editor used to splice physical film on a Moviola — a slow, irreversible, physically demanding craft — NLE allowed the same edits to be made in software, with undo, with multiple versions, with no consumable cost. The

first generation of working editors who picked up Avid did so on the same premise as the synth's first-decade adopters: *the tool will let me do what I already do, faster*.

What NLE actually produced, by the mid-1990s and through the 2000s, was a fundamentally new editing *grammar*. The average shot length in mainstream Hollywood drama dropped from roughly ten seconds in the 1960s to roughly four seconds by the mid-2000s — a change made trivial by NLE that would have been physically punishing to execute on a Moviola. The MTV-cut aesthetic, which had been a music-video novelty in the early 1980s, became the default grammar of contemporary action cinema by the 2000s — *The Bourne Identity* (2002) is the textbook example, with shot lengths under two seconds across whole action sequences and an editing logic that depended on the viewer's now-trained ability to read a fast-cut grammar. The parallel-narrative structures of contemporary streaming drama — multi-thread, multi-timeline, multi-perspective storytelling, edited together with non-linear interleaving that would have been logistically impossible on tape — *Lost*, *Westworld*, *Dark*, *Severance* — are forms that NLE made possible.

Walter Murch, the most respected film editor of the last fifty years and one of the few people to have edited at the highest level on both film and Final Cut Pro, made the point clearly in *In the Blink of an Eye*: the tool *does* change the grammar. Murch was characteristically careful not to claim that the change was an improvement or a degradation. He claimed only that it was a *change* — that NLE permitted certain kinds of edit that physical-film editing could not, and that the new grammar would, over a generation, become as natural to its audiences as the slower-cut grammar of the 1960s had been to its.

The same dynamic is, in 2026, visible in the AI-augmented production pipeline. Working creatives I know are doing things in their day-to-day practice that would have been physically impossible — not just expensive, *impossible* — in 2020. Iterating across forty variants of a scene in an afternoon. Producing personalised localised versions for ten markets simultaneously. Re-cutting a feature against a different aspect ratio for vertical-video distribution while keeping the principal photography intact. Generating a sustained 4D point-cloud reconstruction of a real-world location from a phone-captured walk-through and using it as a virtual-production plate. Running a scratch-vocal session in seventeen languages off a single take. These are *not faster versions of existing workflows*. They are workflows that did not exist three years ago. And, exactly as Murch predicted of NLE, the audiences for the work made on these workflows are already developing the perceptual literacy to read it.

Six categories of newly-possible work

With those two analogues in mind, I want to walk through six categories of creative work that I believe are *newly possible* in 2025–26 — meaning, work that an individual creative or a small studio can produce now that they could not have produced before, and that an audience can experience now in ways the audience could not have experienced before.

I will be honest, on each, about how much of the category is already shipping in finished form and how much is still in the *demo and beta* layer of the toolchain. The whole point of being inside the work is that you can see the difference.

I will also flag, throughout, the *binding constraint* that runs underneath the entire chapter: **human attention is finite**. This is the most-underdiscussed structural fact of the AI creative-economy moment, and I will come to it at length in a moment.

One: Remix at scale

The first category — the most visible and the most legally contested — is *remix*. The infrastructure of AI generation makes it cheap, fast, and at scale to produce derivative work: alternate-style versions of existing songs, cover-style reinterpretations across genres, AI-dubbed translations of feature films into languages the original release never reached, image-to-image style transfers of canonical artworks, motion-transferred re-performances of choreography across body types.

I want to be careful in describing this, because *remix as a creative form* has a long pre-AI history. Hip-hop's relationship to sampling is the canonical example; Lessig's "*remix culture*" framing from the 2000s identified the dynamic in broad strokes well before generative AI; the *Star Wars* fan-edit community, the mash-up era of Girl Talk and Danger Mouse's *Grey Album*, the YouTube AMV community, the TikTok stitch-and-duet grammar — every one of these is, in operational terms, *remix infrastructure that produces creative value through derivation*. The 2025–26 AI moment didn't invent remix culture. It made the *production cost of derivative-but-original creative work* drop by more than an order of magnitude, and it shifted the technical bottleneck from *skill at imitation to judgement about what to imitate*.

The 2025–26 examples I have followed most closely:

- **AI-dubbed feature releases** crossing language barriers that the original production could not afford to cross. *Watch the Skies*, the Swedish UFO feature dubbed into English using ElevenLabs voice cloning, getting US distribution in October 2025, is the canonical example I covered in [Chapter 1](#). The dubbing is not, in any historical sense, a *new translation*; what is new is that the translation is rendered into the original actor's voice rather than a session voice actor's, preserving the performance signal across the language transition.
- **Cover-style reinterpretation at the long-tail of recorded music**. Suno and Udio's user economies have produced an extraordinary volume of *style-transferred* renditions of existing songs — folk versions of pop hits, metal versions of show tunes, lullaby versions of EDM tracks. The vast majority of these are uninteresting. A small fraction, in my experience, are *better than the original by some specific human listener's measure of better* — and that is the new commercial dynamic.
- **The Andrii Daniels bomb-shelter clip** ([Chapter 12](#)) — *Deadpool* and *Harry Potter* as a Christmas mash-up, made in a Ukrainian bomb shelter, viral in December 2025 — is a remix-at-scale artefact in the strict sense. It crosses two IPs that the rights-holders would never have collaborated on, in a stylistic register the original films could not have produced, and finds a global audience that responds to the *idea* of the remix as much as to its execution.

The legal layer of this category is still moving. *UMG v. Anthropic*, *Getty v. Stability AI*, the EU Copyright Directive's Article 17, the UK 88% — these are the institutional structures that

will decide whether the AI-remix category becomes a *licensed and compensated* creative form (the Petrillo template applied to AI) or a *grey-market* one that operates outside the rights system. The historical pattern — Sampling, post-*Grand Upright*, became a licensed creative form, with the dense Bomb Squad style becoming commercially difficult but the basic technique surviving in a more clearance-friendly mode — says the remix category will, in some form, *settle into a licensed category by the end of the decade*. The Petrillo template, again, is the answer the system already knows.

Two: Mass personalisation, against the binding constraint

The second category — the one most often gestured at in platform-company keynotes and least well-served by them — is *mass personalisation*: creative work that is *individually different* for each viewer, listener or player.

The early shipped versions in 2025–26 are these:

- **Personalised playlist generation and AI DJs.** Spotify’s DJ feature, YouTube Music’s AI Hub, Apple Music’s editorial-style algorithmic mixes — all of these are *individual-level curation* of pre-existing creative work. The personalisation is in the *selection*, not in the *generated content*.
- **Personalised game NPCs and dialogue.** Inworld, ConvAI, Ubisoft Teammates, Roblox’s AI assistant — by spring 2026, the first wave of generative-NPC dialogue is shipping in player-facing form. The personalisation is in the *moment-by-moment response* to the player’s actions, not in the underlying game world. Whether this actually produces a personalisation a player can recognise — versus producing the same in-character dialogue from a wider statistical envelope — is, on the published shipping examples I have played, an open question.
- **Personalised marketing creative.** WPP Open Pro, Adobe CX Enterprise, GenStudio — the platform infrastructure for producing per-segment, per-audience, per-context creative variations is now mature. Whether the audience experiences this as personalisation, or experiences it as a feeling of unease that the messaging knows too much, is the cultural question every brand creative is asking in 2026.
- **Personalised educational content.** Khan Academy’s Khanmigo and Duolingo Max are the most-mature shipped examples I know of. The personalisation here is genuine and measurable — the per-student tutoring response is different on the level of the individual learner.

I want to draw a sharp limit around the personalisation category, because the binding constraint runs straight through it and I think every working creative thinking about this market needs to internalise it.

Human attention is finite. Aggregate daily media-consumption time per person has, on the available Nielsen-style telemetry I have read, *not* grown meaningfully in the past decade. The total of every form of media consumption — TV, streaming video, music listening, podcast listening, social media, gaming, reading — is, per the published data, on the order of 11–12 hours per US adult per day, and that number has been roughly stable for years even as the *number of available* hours of content per day has exploded by orders of magnitude. The

eye, the ear and the consciousness each have a finite capacity. *Personalisation* does not, by itself, expand that capacity. It changes the *distribution* of attention across content, but it does not increase the *total* attention available to be spent.

This is the structural ceiling on the *commercial* value of mass personalisation. Producing a personalised version of a film for every viewer — to take the most extravagant platform-keynote framing — is, on the binding-constraint reading, a competitive move that *reallocates* attention rather than *expanding* it. The category will be commercially meaningful in segments where reallocation can produce premium prices (luxury advertising, premium educational content, top-tier video-game NPCs in IP that justifies the investment). It will be less commercially meaningful at the long tail, where the personalisation effort does not produce attention-reallocation big enough to pay for the agentic infrastructure underneath it.

The trade-press framing of mass personalisation as *infinite content for infinite audiences* is, on the binding-constraint reading, structurally incoherent. The audience does not have infinite attention. The producible content is, by 2026, effectively infinite. The economic question is *which slices of the audience's existing finite attention budget* the personalised work can plausibly capture. The answer, by my read of the shipping evidence, is *narrower than the platform companies' enthusiasm suggests*.

I will come back to the finite-attention constraint at the end of the chapter, because it shapes every category that follows.

Three: Audience participation in the creator economy

The third category — and the one I am most personally interested in — is the *audience as participant*. The 2010s creator-economy framing was that the audience could *make their own content* on platforms (YouTube, TikTok, Instagram), creating a long tail of user-generated work alongside the professional content. The 2025–26 AI framing extends this by an order of magnitude: the audience can now *prompt, contribute to, remix and co-author* work in real time, often in collaboration with named professional creators.

The shipped examples I have followed:

- **TikTok / Instagram / Shorts remix mechanics** — duets, stitches, the Sora iOS app's TikTok-style remix grammar (which crossed a million downloads in five days in October 2025). The barrier between *watching the work* and *making a related piece of work* has dropped to a single tap.
- **Roblox's creator economy**. The platform's UGC ecosystem, by 2026, is the largest single audience-as-creator system in any creative medium. Roblox creator payouts run into the hundreds of millions of dollars annually. The line between the audience and the working creator on Roblox is, structurally, not a line; it is a gradient.
- **Disney+'s announced (though not yet shipped at the date I am writing this) user-generated-AI-content tools using Disney IP**. If this ships in the form the platform has described, it will be the first time a major IP holder has *deliberately given the audience tools to make work inside its canonical universe*. The cultural permission this would set — sanctioned fan-made *Star Wars*, sanctioned fan-made *Marvel* — would, on

the historical pattern, expand the categorical *Star Wars* / *Marvel* canon at a rate the studio system itself cannot match.

- **The fan-prompt economy.** Civitai, Promptbase, the LoRA marketplaces — the infrastructure for *selling and sharing the creative inputs* that other creatives then use to produce work. The audience is participating not just in the work, but in the *toolchain underneath the work*.

The single cleanest piece of evidence on the *generational* shape of this category came in May 2026, when **Snapchat** published research finding that **31% of 13–15 year-olds** on the platform were already using AI tools “*to be creative*” — not, importantly, to do their homework, not to chat with a synthetic friend, but specifically to *make things* they then shared with their peers.³⁴⁷ That number is the most legible quantitative indicator I have seen of where the audience-as-participant category is heading. The 13-to-15 cohort the survey describes will be the 18-to-20 cohort of 2029. By the time they reach the working-creative entry pool, *making things with AI* will not, for them, be a category distinct from *making things*. It will simply be how things are made. The studios that build for this audience now — not as a future they are anticipating, but as a present they are already serving — will, on the historical pattern of every previous generational shift, set the terms on which the next decade of cultural production runs.

The thing I want to flag about this category — because I think it is the part the platform companies and the working studios have most systematically under-priced — is that *audience participation reverses the direction of the creator-audience economic relationship*. In the pre-platform creative economy, the audience paid the creator. In the platform-era creator economy, the audience generated the content and the platform monetised the attention. In the 2025–26 AI-augmented creator economy, the audience is increasingly co-producing the content *with* the creator, and the question of who gets paid for what is, structurally, harder to answer than it has ever been.

This is one of the open frontiers of working-creative business model design in this period. I do not think anyone has solved it. The studios that figure out how to credit, compensate and structurally honour audience contributions to the work — without turning the work into the kind of crowdsourced mush that does not survive the slop ceiling — will, in my view, have the most defensible business position in the next decade. The studios that try to extract audience-generated work without compensating it (the 2010s social-media platform model applied to AI) are, on the historical pattern, walking into the next *Viacom v. YouTube*-scale lawsuit.

Four: Fan fiction and fan-made content, legitimised at scale

Closely related to the participation category — but worth a section of its own — is the *fan-fiction* / *fan-content* category. Fan-made creative work has a long pre-AI history. The *Star Trek* fanzine era of the 1960s; the *Star Wars* fan-film tradition from the 1970s onward; the Archive of Our Own / Wattpad / FanFiction.net communities; cosplay; the anime fansubbing tradition; the *Harry Potter* fan-fiction archive that, on some counts, contains more words of *Harry Potter*-canon-derivative text than the original novels themselves.

What AI does to this category is two things. *One*, it raises the *technical floor* of fan-made work toward what was previously professional-grade: a fan can now generate a *Harry Potter* short film with production values that would have required a major-studio budget five years ago. *Two*, it makes the *canon-extension impulse* of fan culture practically infinite: every reader, every viewer, every player can — with current tools — produce a new piece of work inside the universe they love, in their own voice, on their own terms, for their own audience.

The cultural and legal layer of this category is, in 2026, still moving. Disney's tolerance for fan content has shifted under AI, and the question of whether *AI-generated fan content using Disney IP* will be treated more leniently than fan-made content has historically been is one of the open IP-policy fights of the year. Lucasfilm's tolerance for *Star Wars* fan films has, historically, been generous; whether that extends to AI-generated *Star Wars* features is unsettled. The Marvel Comics community policy on AI is one of the documents to watch in the next eighteen months.

What is *not* unsettled is that the audience is already doing this. The fan-AI-content economy is, by spring 2026, larger by volume than the official IP-holder output for almost every major IP in popular culture. The official IP holders can suppress this, license it, build platforms around it, or watch it consume their cultural authority. There is no fourth option.

The working creative read on this, for anyone reading the book inside an IP-holding studio, is the one I made in Chapter 7's discussion of the legacy industries' strategic vulnerability. The studios that move *toward* sanctioned fan-AI content economies — Disney's announced UGC tools, the most generous of the cosplay-and-fan-film tolerances, the platform-and-fund models that pay fan creators for canonical contributions — will have a meaningful structural advantage over the studios that try to defend the closed canon against the audience that is, with or without permission, going to extend it anyway. The Petrillo template applied to fan culture is: pay the fan, structure the IP-holder's stake, *participate* in the extension rather than fighting it.

Five: Agentic support workers for the solo creative

The fifth category is the one that has — in my own practice and in the practice of the wider DreamLab community — produced the most dramatic and immediate productivity gains. The agentic-support-worker model.

In the pre-AI creative economy, the small or solo creative practitioner — the independent filmmaker, the songwriter, the freelance illustrator, the indie game developer, the YouTuber — operated with a particular structural disadvantage. Their work was bottlenecked not on creative judgement (which the senior practitioner had in abundance) but on the *production-coordination labour* that a larger studio would assign to junior staff: client communications, project management, asset organisation, scheduling, invoice processing, basic research, draft response, brief structuring, simple post-production. The senior practitioner spent an inordinate fraction of their effective working hours on labour that did not require senior judgement.

The 2025–26 agentic toolchain — Notion AI, Adobe Express AI Assistant, Heygen Video Agent, Claude Apps, the personal-assistant features baked into the major platforms — has, in

operational terms, *given the solo creative the first four hires they would have made*. The personal assistant, the scheduler, the production coordinator, the asset manager — these are now, for working solo creatives in my circle, agentic functions running underneath the senior practitioner's day, costing roughly the price of a couple of midrange platform subscriptions, and producing genuine recoverable hours.

The economic impact of this is, I think, the most under-priced shift of the period. The *one-person studio* that, in 2020, would have been a part-time freelance practice supporting two to four projects a year is, in 2026, a *full-business-class production operation* supporting twenty to forty projects across a wider range of disciplines. The Sienna-Rose / Xania-Monet / Hoyt-Dwyer single-creator AI-supported career — about which I have, in [Chapter 5](#), made my reservations clear in terms of *star formation* — is, on the *production-economics* side, a genuine new business form. Whether or not Xania Monet becomes a Billboard-defining cultural figure, the *operational machinery* underneath her — a single human creative, supported by an agentic stack producing music, marketing, distribution and merchandise at a scale that would have required a small label to support five years ago — is a working business model that did not exist before the toolchain shipped.

The implication for working creatives at the senior solo level is direct. The first four hires you would have made — production coordinator, junior researcher, scheduling and admin, post-production junior — are now agentic. The economic ceiling on your individual practice has, structurally, lifted. The constraint is no longer the throughput of the labour underneath you. The constraint is your own senior judgement bandwidth — the briefing, the taste, the integration, the *Why* — which the agents cannot replace and which, on the chess-grandmaster argument of [Chapter 15](#), has more commercial leverage in this market than it has had in any previous period.

Six: Hyperlocal and long-tail cultural production

The sixth category — and the one I think will, in retrospect, prove to be the most culturally significant — is *hyperlocal and long-tail cultural production*. The collapse of the cost of producing professional-grade creative work means that, for the first time in the history of mass media, *every linguistic community, every regional culture, every minority cultural tradition, every niche audience* can produce work in its own language, in its own visual style, for its own audience, at production values that compete with global commercial output.

The early shipped examples:

- **Korin AI**, the Africa-trained and Africa-built music platform launched in spring 2026, is the canonical example of a *deliberately culturally-specific* AI tool. The training data is African; the outputs are African; the platform exists to serve African creators rather than to backfill global pop with African aesthetic markers.
- **Indian AI cinema and the Trilok / Animaj / Filmax DinoGames cluster**. India's regional-language film industries — Bollywood, Tollywood, Kollywood and the dozen smaller regional industries — have been among the most aggressive adopters of generative AI in the period this book covers. The reason is structural: the regional film industries operate at production budgets where AI's cost-reduction effect has *the largest*

proportional impact. A regional Indian feature that would have cost \$200,000 to produce in 2020 can, in 2026, be produced for a meaningful fraction of that, with comparable production values. The category of films that this brings into commercial viability — every regional cinema with a million-viewer audience but a six-figure budget ceiling — is, by my estimate, *an order of magnitude larger* than the category that was commercially viable before.

- **The Tunisian *Lily*** — winner of the \$1m Dubai AI Film Award in January 2026 — is the cleanest individual example of a film made *outside* the major-market financing system, at a production value competitive with it, by a creative team using AI tooling as the cost-reduction lever.
- **The 800-creator declaration's signatories** include working creatives from a meaningful number of national contexts that the historical Anglophone film/music/games industries have under-represented. The *Stealing Our Work Is Not Innovation* coalition is, structurally, a coalition that includes the long tail of global creative work.

The implication of this category is the one I am most personally hopeful about, and I want to be honest that *hopeful* is the right word — there is a less-hopeful version of the same data, in which Anglophone AI tooling homogenises the global creative economy faster than the regional creative economies can develop their own infrastructure. The 2026 evidence I have is mixed enough that both outcomes are still on the table.

What is not in doubt is that the *production-cost ceiling that has, for a century, kept hyperlocal creative work below the threshold of professional commercial viability* is, by 2026, no longer the binding constraint. The next decade of cultural production will, I am confident, contain more local, more linguistically diverse, more culturally specific creative work than the previous decade — and the most culturally significant individual works of the next decade will, on the historical pattern, come from communities the existing global creative economy has under-served. The question is whether those communities own the tooling that produces them.

What is not yet possible

Before I close, I want to be honest about the categories that are *not yet shipping* in the form the platform-company keynotes have promised them, because the book's credibility depends on accurately characterising the gap between affordance and rhetoric.

AI does not yet write a satisfying novel from scratch. The 2025–26 evidence on long-form prose generation is that AI systems produce technically competent prose that loses *narrative purpose* over the length of a novel. Working novelists I have talked to who have experimented seriously with this category all describe the same failure mode: the prose is fine; the *book is not a book*. This is consistent with the *House of David* “hand inside a puppet” critique of AI-augmented storytelling at feature length.

AI does not yet do live performance. No AI is touring in 2026. Xania Monet has not performed at Madison Square Garden. The cultural permission for a synthetic performer to share a stage with a live audience does not, as of this writing, exist in any meaningful form. The structural reason — the audience experience of *being in a room with another human*

consciousness — is, on the slop-ceiling reading, not a permission gap; it is a category mismatch. AI work and live performance are, for now, different categories.

AI does not yet do sustained emotional storytelling at feature length without human authorial spine. The auteur-driven cinema of the Cameron / del Toro / Spielberg generation is, structurally, a category that depends on a *single human consciousness running through every creative decision in the film*. AI-augmented versions of this category exist; AI-native versions of it do not. The *Citizen Kane* of AI cinema has not been made. Whether it will be is, in my view, the single most interesting open creative question of the next decade.

AI struggles with cultural specificity that the model has not been trained on. The BBC India observation I referenced in Chapter 13 — that AI screenplays produce *cultural-memory failures* when generating content for cultural traditions the training data has under-represented — is the binding limitation on AI's promise as a *globally-distributed creative tool*. The hyperlocal-cultural-production category I described above depends, structurally, on this limitation being addressed by purpose-built regional model infrastructure (Korin AI, the Indian regional-language tooling, the various Asian-built models). Until that infrastructure matures, AI cultural-production is still, by default, Anglophone and Western-modelled.

The binding constraint, restated

I want to close by returning to the structural fact that runs underneath every category of newly-possible work in this chapter.

Human attention is finite.

I am not making a romantic argument about it. I am making an economic one. The aggregate daily media-consumption time per adult, in any market the Nielsen-class telemetry covers, has been roughly stable for at least a decade — eleven or twelve hours a day, across all forms of media, all formats, all platforms, all devices. *That number has not grown with the rise of streaming. It has not grown with the rise of mobile. It has not grown with the rise of social media. It has reallocated, sometimes dramatically. It has not expanded.*

The producible content, meanwhile, has expanded by orders of magnitude. Deezer in 2026 receives 75,000 AI tracks a day; the listener has the same number of waking ear-hours she had in 2016. YouTube uploads run at hundreds of hours per minute; the viewer has, in net, the same daily screen-time budget. The Sora app produced a million downloads in five days; the audience for the work those million users are about to make has, in total, the same total attention they would have had if Sora had never shipped.

This is the binding constraint of the entire 2025–26 creative-AI moment. The cost of producing work has collapsed; the supply of attention has not. Every category I have described in this chapter — remix, personalisation, audience participation, fan content, agentic support, hyperlocal production — is being built into a market where the *production* side of the supply-and-demand equation is racing toward infinity and the *consumption* side is bounded by the finite cognitive and biological capacity of the human nervous system.

This has three structural implications.

One. The competitive advantage in this market accrues, with iron consistency, to the producer of work that the audience *chooses to spend its finite attention on* rather than the producer of work that the audience could, in principle, choose. The slop ceiling, the authenticity premium, the chess-grandmaster move, the *Why* — these are not romantic notions. They are the *mechanisms by which finite human attention selects from infinite producible content*.

Two. The platform business models that depend on *expanding audience attention faster than supply* are, on the structural reading, walking into a wall. The platform companies' current trajectories — push more content, optimise for engagement, monetise more minutes — work only as long as the engagement budget is elastic. It is not elastic. The audience cannot, on average, watch more hours per day than it already does. The platforms that get to those audiences first, and that build the most-defensible *retention* mechanics, will win the share-of-attention game. The platforms that come second will be competing for an attention budget that the first-movers have already claimed.

Three. The new categories of work I have described in this chapter will produce, in aggregate, more *total creative output* than the previous decade. They will not produce, in aggregate, more *audience attention received per minute of output*. The ratio of attention-to-output, which is what working creatives actually live on, is going to fall — sharply, in some categories, more gradually in others. The working creatives who survive the next decade will be the ones who recognise that the attention-to-output ratio is the metric that actually matters, and who position their practice in the categories where the ratio is most defensible.

The categories where the ratio is most defensible — on the chess-grandmaster reading and the slop-ceiling reading and the authenticity-premium reading — are the categories where the work is *most deliberately un-machine-like, most authentically human-authored, most culturally specific, most personally risked*. The newly-possible work in this chapter is real and will reshape the creative economy. But it is being made in a market where the binding constraint is, and will remain, the finite attention of the audience watching it.

The synth made entirely new sounds possible. The audience for those sounds was finite. The musicians who learned to make sounds *the audience would spend its scarce listening time on* — Trevor Horn, Kraftwerk, the Bomb Squad, Aphex Twin, Daft Punk, every electronic-musician who built a serious career — are the ones we still listen to. The musicians who learned to make sounds the synth made possible *but that the audience did not develop a hunger for* — the long tail of 1970s and 1980s synth-record obscurities — are remembered, mostly, by collectors.

The AI moment will, on the historical pattern, work the same way. The newly-possible categories will create work that did not exist before. The audience will, on the available evidence, allocate its scarce attention to the work that *earns* that attention. The working creatives who position themselves in the newly-possible categories, *and* who make work that the audience deliberately chooses, are the ones who will define the next decade of the form.

That is the operating manual. The categories are open. The constraint is real. The choice — like the chess-grandmaster's — is yours.

Chapter 11 — The Orchestrator

In the second issue of *Dream Machine*, in October 2025, I described the **Human–AI Agency Continuum** as a way of mapping how much of any given creative function is being done by the human in the chair and how much by the machine. In the thirteenth issue, in January 2026, I made a prediction that I want to look at again in this chapter, because — six months on — it has held up better than most of the others I made that day.

I called 2026 “**the Year of the Orchestrator.**”³⁴⁸

The argument was straightforward. If 2024 had been the year of the *generator* — the prompt-and-respond text-to-image, text-to-video model — and 2025 had been the year of the *agent* — the system that could take goals, plan, decide and execute multi-step tasks autonomously — then 2026, I argued, would be the year that working creatives stopped being *operators* of these tools and started being *orchestrators* of them.

By “orchestrator” I meant something quite specific: a person whose job is not to make the work themselves but to direct, brief, integrate and judge the work of a team of AI agents and human collaborators. Not a producer in the old sense. Not a creative director in the old sense. Something newer, with a different skill set, a different rhythm, a different relationship to craft.

This is the role I think most working creatives will be holding by 2030. This chapter is about why, what it looks like, what it asks of you, and where it breaks.

Forty-nine Claude agents and seventy-two skills

In May 2026, in the second-to-last issue I wrote before this book went to draft, I reported on something that I think of as the canonical image of what the orchestrator role actually is. **Sony**, in announcing its “all-in on AI for games” strategic move, disclosed that one of its game-development studios was running a coordinated multi-agent team of **49 Claude Code agents**, working with **72 skills**, on game-development tasks ranging from asset generation through QA through engineering through animation.³⁴⁹

This is, in operational terms, a small army of synthetic colleagues working on a single creative project. Each agent has a defined role. Each skill is a defined capability. The whole apparatus is overseen by a *substantially smaller number of human developers and creative leads* whose job is to plan the team’s work, brief the agents, judge their outputs, integrate their contributions, and decide what gets into the game.

The ratio matters. The pre-AI version of this game-development team would have been, plausibly, 50 to 100 people working on the same scope of work over a much longer timeline. The AI-augmented version, as Sony has set it up, is a much smaller number of *senior, judgement-heavy* roles — people whose value is taste, narrative sense, gameplay design instinct, IP fluency, engineering oversight — sitting on top of a much *larger* pool of synthetic capacity.

What you don’t see, in the Sony picture, is the disappearance of the junior roles. Those roles haven’t disappeared. They have been *replaced* — by agents. The 49 Claude Code agents are, in effect, the new junior developers, the new junior animators, the new junior writers. They work cheaply, they work fast, they work in parallel. They are not — and this is important —

replacements for senior judgement. They are leverage *for* senior judgement. The whole pipeline is designed to take the senior creatives' time and multiply its effective reach by a substantial factor.

The orchestrator, in this configuration, is the senior creative — the writer-director, the gameplay lead, the art director, the technical director — whose taste and judgement set the boundaries that the agent team works inside.

I want to flag the obvious labour question here, because it is the question every working creative is asking, and any chapter that handwaves past it is not being honest. If the senior roles still exist, and the junior roles are now done by agents, *where does the next generation of senior creatives come from?* The pipeline that has, for fifty years, produced senior creatives in the film, TV, games, music and design industries has worked by *starting people as juniors and letting them grow.* If we remove the junior rung, we are — over the next decade — also removing the apprenticeship infrastructure that makes the senior rung possible.

This is the structural risk of the orchestrator economy that platform companies are, in my view, not yet taking seriously. I name it the **Apprenticeship Gap**, and I will come back to it in Chapter 14.

What the orchestrator does

In the talks I have given since publishing the “Year of the Orchestrator” piece in January, the question I get asked most often is: *what does an orchestrator's day actually look like? What are the skills? What does the job description say?*

I want to try to answer that in this chapter, in the most concrete language I can find, because I think the gap between *what working creatives think they will be doing in 2030* and *what they will actually be doing* is unhelpfully large.

The orchestrator does five things:

One. They define the brief. The agents — and the human team — work to a brief. The brief is the thing the orchestrator owns. It is not the prompt. The prompt is a derivative artefact. The brief is the *creative intent* that the project exists to deliver: who it's for, what it's trying to do in the world, what success looks like, what tone, what feeling, what audience, what context. The orchestrator's first job is to know — clearly enough to communicate it — what the work is *for*.

Two. They allocate work. Given the brief, the orchestrator decides which parts of the work get done by humans, which by AI agents, and which by the orchestrator themselves. This is the practical, day-by-day application of the Human–AI Agency Continuum we talked about in Chapter 3. It is not a one-off decision. It is a constant series of micro-decisions about where on the continuum each function sits, in this project, on this day, for this output.

Three. They brief the agents. This is the closest the orchestrator role gets to what people currently mean by “prompt engineering,” but it is a meaningfully different skill. Briefing a human collaborator and briefing an AI agent are not the same activity, but they share a core: the ability to *describe what is wanted with enough precision and enough context that the recipient can produce something useful, without over-constraining them in ways that*

prevent useful surprise. This is, in my experience, the single biggest skill differentiator between effective and ineffective AI-era creatives. The people who can brief well — who know when to give the agent a tight constraint and when to let it explore — are the people who produce the best output.

Four. They judge the outputs. When the agents deliver, the orchestrator's job is to look at what came back and decide what to ship, what to revise, what to throw away. This is *taste* in the most operational sense. It is also, crucially, *taste under abundance* — taste exercised in a context where you can have ten versions of the same scene back in ninety seconds and your job is to choose, not to make. Choosing well under abundance is a different cognitive skill than choosing well under scarcity, and most creatives have been trained for the latter. The grandmaster analogy from Chapter 15 — top chess players, in 2026, deliberately playing *sub-optimal* moves to put their opponents on uncomputed ground³⁵⁰ — is the cleanest available picture of what this looks like in practice. The orchestrator's value is not in choosing the *most-likely-good* output of the ten variants the agent returned. The agent has, by construction, already centred its output on the most-likely-good. The orchestrator's value is in seeing which of the ten variants would be the *un-machine-like* move at this specific point of this specific project, and choosing that one. *Taste under abundance* is, operationally, the discipline of refusing the machine-optimal output in favour of the deliberately-chosen one.

Five. They integrate. A film is not the sum of its scenes. A game is not the sum of its assets. A campaign is not the sum of its individual creatives. The orchestrator's job, at the end, is to take the outputs of the agent team and the human team and assemble them into a coherent piece of work that has a *single sensibility*. This is the part of the job that I think — for all the AI tooling — is least likely to become a thing AI can do. The integrated voice of a piece of work is a function of a single human consciousness running through it. The orchestrator role is, at its core, the role that holds that voice.

If you read those five things back, you will notice that none of them are *making*. They are all *deciding*. The orchestrator's work product is decisions: about what to make, who or what should make it, whether the made thing is good enough, and how the made things fit together.

This is, in some sense, what every senior creative director and showrunner has always done. The change is not the shape of the role. The change is that the role is no longer a privileged senior position at the top of a pyramid of junior makers. It is, increasingly, *the entire role*. And it is the role that, in 2026 and 2027, working creatives at every level are being asked to grow into faster than the career-development infrastructure of any of the creative industries is built for.

Where the agents go wrong

I want to spend some time on the failure mode of agentic creative work, because the press cycle around Sony's "49 agents" framing — and the corresponding announcements at Adobe Summit, NVIDIA GTC and elsewhere — has been heavy on the upside and thin on the downside.

The agents go wrong, in my experience and in the experience of every working creative I have talked to about this, in four characteristic ways:

One. They confidently produce the wrong thing. This is the most familiar failure mode and the one the public discourse has covered most. Agents — like the LLMs underneath them — *hallucinate*. They will produce an asset that confidently disregards a key constraint of the brief. They will generate a character with the wrong eye colour. They will produce a piece of music in the wrong key. The fix, with current systems, is *human review at every gate*. The cost of this review, in time, is the largest single source of the “AI was an expensive mistake” experience that Charles Cecil described and that many studios have replicated.

Two. They produce the *mean* of the training distribution. This is the more insidious failure. Agents, by default, will produce work that sits in the middle of what they have been trained on. The middle of a training distribution is, by definition, the *most average* version of the thing you asked for. For creative work — where the value is almost always in the *non-average* — the default output is structurally weak. To get above-average output, the orchestrator has to push, prompt, and curate against the gravitational pull of the mean. This takes deliberate, conscious effort and it takes taste to know what *above-average* looks like in this particular project.

Three. They lose context across long tasks. Agents working on multi-step tasks accumulate errors over the steps. A small misalignment in step one becomes a larger one by step five. By step ten, the output is meaningfully off-brief. The orchestrator’s role is to *check in* at the right intervals — not so often that you negate the benefits of agentic execution, not so rarely that the team has wandered off the brief by the time you look.

Four. They cannot tell when to stop. Agents, given an open-ended task, tend either to over-iterate (producing fifty variants of the same thing without converging on one) or to under-iterate (producing one variant and stopping). The orchestrator’s job is to set the *stopping criteria* for the agents, which is, in practice, a series of judgment calls about *when good enough is good enough*. This is a craft skill the agents do not, as of 2026, have. It is also a skill that, in my experience, almost every working creative already has — they just haven’t had to use it on synthetic colleagues before.

The Anthropic blog posts on agent deployment patterns through Q1 2026 made the point I want to land on here cleanly: agentic systems work best when they are deployed by people who already have the taste and judgment to know what good output looks like. They *accelerate* people who are already good. They do not, on their own, make people good.³⁵¹

That is the orchestrator’s job: to *be the human who is already good*, holding the taste line, while the agent team produces faster than the human pipeline ever could.

Sundance, and the literacy turn

I noted in the last chapter that the platform companies have responded to the AI-literacy gap by becoming, in addition to software companies, *training organisations*. The most institutionally credible example of this turn, in the period this book covers, was **Sundance Institute’s** launch of an **AI Literacy Initiative** at the 2026 festival.³⁵²

I want to spend a moment on what Sundance did, because the framing is important.

The Institute's announcement did not say "*AI is the future of filmmaking; here is how to use it.*" It said something more careful. It said that AI is a fact of the filmmaking landscape, that filmmakers are going to have to make decisions about whether and how to use it, and that those decisions should be made by *informed* filmmakers with *agency over their own practice* — not by filmmakers who have had the tools imposed on them by clients, by streamers, or by tool vendors.³⁵³

The framing was *creator empowerment*. The mechanisms were free learning through Sundance Collab, community conversations, a fellowship and alliance model, and a Story Forum that specifically tackled the legal questions creators face when they use AI: whether AI-generated content can be copyrighted, how to protect projects in a world of contested datasets, how to negotiate AI clauses in production contracts.³⁵⁴

Google funded this. The \$2 million the company put into the Institute, with the stated aim of training 100,000+ artists in foundational AI skills, was both an act of corporate generosity and a strategic investment in the *category* of "filmmaker who can use AI without losing their creative authority."³⁵⁵ Both motivations are real. Both can be true. What matters, for the working filmmaker in 2026, is that the institutional infrastructure for becoming an AI-literate orchestrator — without surrendering creative agency — now exists.

The MckKinsey AI report on film and TV production, released in early 2026, made the corresponding business case. AI would not, in McKinsey's view, replace film and television production. It would *restructure* it — towards smaller teams, faster cycles, more iteration, and a heavier reliance on senior creative judgement.³⁵⁶ In other words: towards an orchestrator-shaped industry.

The middle layer disappears

What is happening, structurally, in every creative industry that the *Dream Machine* newsletter has tracked in these six months, is that the middle layer of the workforce — the layer of *intermediate* roles, between the very senior creative leadership and the very junior entry-level — is being absorbed into the agent layer.

This was the story behind **Ubisoft's** decision in January 2026 to cancel five games, including the *Prince of Persia* remake, while pouring more money into AI.³⁵⁷ It was the story behind **Square Enix's** target of doing 70% of its QA work via AI by the end of 2027.³⁵⁸ It was the story behind **Falcom's** description of work that "*previously took 2–3 hours*" being completed "*in 10 minutes*" with AI tools.³⁵⁹ *Eighteen-to-one productivity*. That ratio, taken on its own, is what re-shapes the headcount calculus for every studio's mid-level production work in the next eighteen months. It was the story behind the *Take-Two CEO's* explicit framing that AI "won't invent the next *Grand Theft Auto*" — meaning, the *creative direction* won't come from the machines — even as Take-Two's QA, asset and engineering pipelines absorb AI capacity rapidly.³⁶⁰

In film, you see the same pattern. **Spielberg** explained in March 2026 why he hadn't yet used AI directly³⁶¹ — and the same press cycle reported that he had a substantial AI-augmented team working on production-pipeline tasks underneath him. **Steven Soderbergh** committed to “a lot of AI” on the Wagner Moura film and a John Lennon documentary, framing it explicitly as transparency: “I owe people honesty.”³⁶² In every case, the structure is the same: a senior creative voice on top, an AI-augmented operational layer underneath, *fewer mid-career intermediaries in between*.

In advertising, the pattern was even more pronounced. **Independent agencies** faced what *Digiday* called “a new frontier as agency-in-a-box tools democratize creativity.”³⁶³ **AI agent developers** became “adland's in-demand role.” The framing from one agency hiring lead, given to *Digiday*, captures the shift better than any of the trend-piece coverage: “*We actually need people who understand [AI], who are building systems organically within their day to day workflows. People who understand taking what took them 40 hours one week and turning it into 38 the next week.*”³⁶⁴ The job description is no longer *make the work*. It is *make the system that makes the work — and keep shaving hours off it*. The PGA Tour expanded its AWS partnership to put AI content at the heart of its content distribution.³⁶⁵ **Mondelez** said it would use AI for TV ads in 2026.³⁶⁶ **Avocados From Mexico** turned to AI to advertise around the Super Bowl, *instead of* a traditional TV buy.³⁶⁷ **Adobe** said that AI in marketing was now “agentic creative intelligence.”³⁶⁸

In journalism, the Reuters Institute's “AI adoption by UK journalists” survey found high integration of AI tools across newsrooms by late 2025.³⁶⁹ Daily Mail reported that Google's AI Overviews had “killed click-throughs” to news sites.³⁷⁰ *The Times* was using AI to model synthetic focus groups from human audiences.³⁷¹ In each case, *the middle layer of the journalism workforce* — the sub-editors, the copy editors, the data journalists, the social-media producers — was the layer most exposed to AI substitution.

I want to be honest about what this means. *It does not mean that every working creative in the middle of their career is about to lose their job*. That framing — the apocalyptic one — has, for two years, been the most popular and the most wrong. What it means is that the *shape* of the mid-career role is changing. Mid-career creatives who can become orchestrators of agent teams will, in many cases, *gain* leverage and earning power. Mid-career creatives who cannot — who try to keep doing the maker-as-craftsperson job at the speed and price of the agents — will, increasingly, struggle.

The Sundance literacy turn, the Adobe and Google training investments, the UK government's free-AI-training-for-all programme — these are the institutional response to that pressure. They are not enough on their own. They are, however, the right direction.

The portfolio creative

There is one more shape of the orchestrator role that I want to flag, because it is the one I see most often in my own studio and in the wider DreamLab community: the **portfolio creative**.

A portfolio creative is someone who, instead of holding a single specialist role, holds *several* loosely-coupled creative roles across different disciplines, supported by AI tooling that lets them maintain useful proficiency in each. The portfolio creative is a writer-director, but also a creative technologist; a music producer, but also a video editor; a games designer, but also a brand strategist.

The *TechBullion* piece “Why the future belongs to multi-skilled leaders,” from November 2025, made the case for this from a corporate-leadership angle.³⁷² The Anthropic Skills framework — the system of named, reusable skills that Claude Code now uses to coordinate multi-agent workflows — is, in effect, an attempt to make the portfolio-creative model into a *technical infrastructure* rather than a personality type.³⁷³ The *Forbes* piece “AI Is Changing How Creators Work And Earn,” from December 2025, surveyed the same phenomenon from the working-creator angle and found the same pattern: the most economically successful creators in 2026 are not specialists. They are *integrators* who can work across disciplines using AI as the connective tissue.³⁷⁴

In my own studio, the move towards portfolio creatives has been a deliberate strategic choice — and an honestly difficult one to execute. The cultural expectation, in most creative industries, has been to hire specialists and stack them in a pipeline. The portfolio-creative model requires you to *hire generalists* and let them move between disciplines as the work demands. The former is easier to manage, easier to bill, easier to explain. The latter, in my experience, produces better work in the AI era, because the human is doing the integration that the agents can’t.

The portfolio creative is the orchestrator at the level of an individual career. The orchestrated team is the orchestrator at the level of a project. The same pattern shows up at multiple scales.

Five orchestrators, briefly

I want to give five working orchestrator case studies, because the abstract description above can sit in the head as a theory without the operational texture of *what the role actually looks like* on a Wednesday afternoon. Each of the five is a specific working creative or organisation whose practice in 2025–26 I think represents a different *shape* of the orchestrator role. Each is documented in the *Dream Machine* newsletter archive. Each is, on my read, doing something that working creatives reading this book can learn directly from.

Andrii Daniels (Ukraine). The independent filmmaker who, in December 2025, produced a *Deadpool* / *Harry Potter* Christmas mash-up in a Ukrainian bomb shelter using a Runway-and-Veo-and-ElevenLabs stack on a laptop running through a generator. The clip went viral and was picked up by *Variety* as a profile piece.³⁷⁵ What Daniels did, operationally, was the orchestrator role at its purest: a single human creative, holding the taste and the narrative judgement, briefing a stack of generative tools to produce work whose production conditions would have been physically impossible eighteen months earlier. Daniels did not write, draw, animate, voice or render the work. He *briefed* it. He *integrated* it. He *judged* what was good enough. He *delivered* the finished piece, on his own, with no studio underneath him. The

bomb-shelter context is the dramatic detail; the underlying operational pattern is what makes Daniels' practice replicable for working filmmakers in any production environment.

The Imagineae Studios / Art Awakens team (Fremantle). The AI-native studio I described in Chapter 7 has, by mid-2026, settled into an operational pattern that maps cleanly to the five-function orchestrator description above. A small senior team of writer-directors, supported by an AI-augmented production pipeline that handles concept development, asset generation, scene assembly and post-production. *Art Awakens* — Imagineae's flagship 2026 development project, fusing AI techniques with classical painting IP — has, on the published interviews with the team, been produced by a core human team of fewer than ten people running an agentic pipeline that, by their own estimate, produces output that would have required a sixty-to-eighty-person team five years ago. The 8:1 productivity ratio is the orchestrator economy expressed at the studio scale.

Sven Vincke / Larian Studios. The opposite case, also instructive. Larian, the maker of *Baldur's Gate 3*, has — as I described in Chapter 7 — *publicly refused* generative AI for its next major game while continuing to use AI-augmented tooling in *adjacent* parts of the pipeline (QA, localisation, internal admin, asset management). Vincke's framing, in his January 2026 statements, was *not* anti-technology. It was *position-on-the-continuum*: certain parts of the game's authorial signature (writing, character design, world-building, dialogue) had to remain *fully human* for commercial and cultural reasons; certain other parts (build tooling, QA automation, marketing-asset generation) could be AI-augmented at no audience-visible cost. Vincke is, in operational terms, *an orchestrator at the level of the studio's continuum positioning*. He is making the agency-line decisions at the strategic level that the working filmmaker makes at the project level. The role is the same. The scope is different.

Xania Monet / Hallwood Media / Telisha Jones. The case I have, in Chapter 5, been most equivocal about. The structural pattern is, on inspection, an orchestrator economy operating at the *single-artist* level. Telisha Jones, the human lyricist, is the orchestrator. The Suno music-generation stack is the agentic capacity producing the *executed* musical output. The Xania Monet persona is the *audience-facing* product. Hallwood Media's \$3M deal pays for the orchestrated work as a unit, with the orchestrator (Jones) receiving the commercial credit and revenue that the synthetic-vocalist alone could not have generated. Whether this scales into a sustained cultural-star career — the Chapter 5 slop-ceiling argument suggests, on six months of evidence, that it has not — is a separate question from whether it works as a *business form*. As a business form, it is the orchestrator role at the level of a solo recording artist.

The Sony game-development teams running 49 Claude agents and 72 skills. The canonical *enterprise-scale* orchestrator case I opened the chapter with. The 49-agent / 72-skill stack is, in operational terms, an organisational design for the orchestrator role at the *team* level: a small group of senior creative leads (writer-directors, gameplay design leads, art directors, technical directors) orchestrating a multi-agent synthetic team whose individual outputs require senior human review and integration. The 49 agents are not autonomous studio replacements. They are *leverage* for the human orchestrators. The 72 skills are the *reusable capabilities* the orchestrators can deploy across multiple projects.

I want to be honest about what the five cases share, because the shared pattern is the operational lesson.

In every case, the orchestrator's *contribution* is the same five-function set I described above: brief, allocate, brief-the-agents, judge, integrate. None of the five is doing the production-execution labour themselves. All of them are making *decisions* that direct a synthetic (and, in most cases, also a human) team to do the production-execution labour on their behalf. The *value* the orchestrator brings to the work is, in every case, the senior judgement that the agents cannot supply: the taste that knows what good output looks like, the briefing skill that gets useful output out of the agents, the integration sense that assembles the agent outputs into a coherent piece of work.

In every case, the orchestrator's *leverage* — the ratio of finished output produced to working hours spent — is dramatically higher than the equivalent practitioner could produce without the agent layer underneath them. Daniels would not have made the Christmas clip in a week pre-AI. Imaginae would not have produced *Art Awakens* at the pace and budget they are producing it at. Jones would not have shipped a Billboard-charting record without Suno. Sony's game-development teams would not have shipped on the cadence they are shipping at. The leverage is real. The leverage is what makes the orchestrator role *commercially* viable as a new form of working practice.

In every case, the orchestrator's *fragility* is also the same: the practice depends on the toolchain underneath it continuing to be available, accessible, on commercial terms the orchestrator can sustain, with model behaviour that the orchestrator can predict and brief against. The platform-dependency of the orchestrator role is the structural risk that [Chapter 9](#) addresses. It is also the reason the *open-the-black-box* argument I made in [Chapter 3](#) is operationally serious: the orchestrators who depend on closed-platform tooling without understanding the dependency are, structurally, exposed to platform-pricing and platform-policy decisions that they have no control over.

The orchestrator role is, in operational summary, *high-leverage and platform-dependent*. The working creatives reading this who are positioning themselves toward the role need to take both halves of that description seriously. The leverage is the upside. The platform dependency is the work that has to be done to defend the upside over time.

The process trap, and the migration out of it

I want to put on the page a counterweight argument to the orchestrator framing this chapter has been building, because the framing has, in the reception I have seen of the “Year of the Orchestrator” piece, sometimes been read as a *consoling* one — as a reassurance that creative work is safe and that the working creative needs only to ascend to the orchestrator role to keep their seat at the table. That is not what the chapter is arguing, and the reassurance is wrong.

The counterweight comes from an essay I have been carrying around with me through the final drafting of this book, *The Process Trap: Generative Abundance and the Equivocation of the Creative*, preserved in full as [Appendix J](#). The essay's central move is to pull apart the word *creative* — which does two incompatible jobs at once, naming a rare capacity (the

genuine origination of the new) and naming a workforce defined by sector, brief, format and schedule. The reassuring sentence circulating through every studio all-hands and trade-press op-ed — *automation will take the rote; real creative work is safe* — is true of the *capacity* and catastrophic for the *class*, and it slides between them without anyone noticing because one word covers both.

The argument that follows from the equivocation is the harder one for this chapter to absorb. Generative AI is, precisely, a *process* technology. It drives the marginal cost of combinational and exploratory creativity toward zero, including — and this is the part most under-absorbed — the tacit craft long believed inimitable. Polanyi's formula that we know more than we can tell was treated as the craftworker's last redoubt; it is exactly this redoubt that large models have breached, because the tacit dimension turns out to be tacit *to the individual practitioner*, not to the corpus. What no single craftsperson could explain, the corpus encodes. The process premium — the rent that craft, technique and executorial competence have historically commanded — is collapsing, and is collapsing *fastest* at the upper rungs of the Dreyfus skill ladder where the most-invested experts sit.

The orchestrator role, in this reading, is not an escape from the process trap. It is, in some configurations, the *re-entry* into it under another name. The “prompt engineering” that the consumer press has been celebrating is process re-dressed. The “personalisation at scale” pitched as the answer to synthetic abundance is what Adorno and Horkheimer called *pseudo-individualisation* industrialised: the same standardised core, recombined and surface-tuned per recipient, presented as bespoke. *Curating* the best of the agent's outputs is still deference to the machine's mean. The orchestrator who only chooses among the variants the system returned has, in the Process Trap's framing, re-entered the process economy by a different door.

The defensible thing, the essay argues — and this is the part I want this chapter's readers to take seriously alongside the orchestrator framing — is not in the artifact and not in the lone maker. It is in the *position* of the maker: a person who has a problem of their own (the originating problem, in Getzels and Csikszentmihalyi's language; *natality*, in Arendt's), a stake in the game (Bourdieu's *illusio*), an embedded place in a living field (Csikszentmihalyi's systems model), and the *negative capability* to refuse the probable answer the machine returns by default (Keats's phrase; Ghafouri's *Paradoxical Bridge* in its tooled-up version). Creativity, in the strong sense that resists automation, is **positioned negation under risk in pursuit of significance** — the bringing-into-being of something that did not have to be, by an agent who has a problem of her own and a stake in the game, achieved by refusing the probable, and judged to matter by a living field. *Every clause names a feature of the maker's position; not one names a feature of the output.* This is why the residual moat is so consistently mislocated.

The operational reframe that follows is the one I want the orchestrator-shaped reader to carry forward into the next chapter. The correct use of a generative system, for a maker pursuing creativity, is not as a creative partner and not as a curation engine, but as the *maximally competent producer of the probable* — precisely so that the human can specialise in the *improbable*. Let the model occupy the mean it is built to occupy. Use its best output as a *map*

of what to avoid, not a menu of what to choose. The orchestrator who defers to the model's best output has re-entered the process trap by another door. The orchestrator who uses the model's best output as a map of where the probable goes and then *deliberately departs from it* — refusing the obvious continuation, the well-formed default, the safe resolution — is the one whose practice survives the collapse of the process premium.

The essay's closing observation is the harder of the two I want to hold against the rest of this chapter's framing. The migration from making to creating — competing on the problem rather than the execution, taking a stake on purpose, re-entering the field as a place of significance-judgement rather than an audience-retention asset, practising negation against the gravitational pull of the mean — is *harder than the work the orchestrator was trained for, slower to reward, lonelier, and structurally unsupported by the institutions that employ them*. To pretend that every working creative called *creative* can simply ascend to *creativity* proper is to repeat the equivocation in a consoling key, and it does a disservice to people making real decisions about real livelihoods. Most creative workers will not make this migration. Some will. The difference between those two facts is where most of the human cost of this transition will fall.

I include this counterweight here, in the chapter that argues most ambitiously for the orchestrator role, because the chapter's argument and the essay's argument are not in conflict but they are not the same. The orchestrator role is the *operational* shape of working creative practice in 2026–2030. *Positioned negation under risk in pursuit of significance* is the *defensible* shape underneath it. The working creative who holds both at once — orchestrating fluently at the level of execution and refusing the probable at the level of intent — is the working creative I think most likely to make work that lasts.

The role this asks you to play

If you are a working creative reading this — and most of the readers of *Dream Machine* are — the question I am sure you have is: *what does this mean for me, this year?*

The honest answer is that it depends on where you are on the continuum we drew in Chapter 3. But there are some things I would say to almost everyone I know in the creative industries right now, and I want to put them on the page.

Practice briefing. It is the single most leveraged skill you can develop. Brief your AI tools as if you were briefing a junior who has thirty seconds to understand what you want and ninety seconds to do it. The discipline of *being able to communicate what you want* will improve every other part of your creative practice.

Cultivate taste deliberately. Look at more good work, harder, with a more critical eye. The agents will, by default, give you the average. Your job, increasingly, is to know what *good* is. That knowledge is a function of how much good work you have looked at, how seriously, with how much craft attention.

Stay in the work. Resist the temptation to abstract too far. The director who never picks up the camera, the showrunner who never writes, the music producer who never plays the instrument — these are the figures most likely to lose the touch that makes their judgement

worth anything in the first place. The portfolio creative is *not* a creative who has lost contact with the craft. They are someone who maintains craft contact in *several* domains.

Choose your line on the Continuum, and defend it. Decide where your craft sits, where you are willing to let the agents work, and where you are not. Write it down. Communicate it to clients, collaborators, your team. Be willing to walk away from work that would force you across the line you have drawn.

Build the apprenticeship. If you are senior enough to be running a team, take seriously the question of where the next generation of senior creatives is going to come from. The orchestrator model breaks if there is no path from junior to senior for new humans entering the field. The studios and agencies that survive the next decade will be the ones that solve this problem — by keeping some junior roles in human hands, by creating new pathways through AI-tool-augmented apprenticeship, by investing in the literacy infrastructure that the platform companies and the institutes have started but cannot finish on their own.

The Year of the Orchestrator is not a coronation. It is a job description. It is what most of us are now being asked to do, whether we have signed up for it or not. The people who do it well will set the terms of the next creative economy. The people who don't will, increasingly, be sat next to it.

The thing I want to land before we leave this chapter is that the orchestrator role — for all the leverage it brings and for all the productivity it unlocks — depends, in the end, on something that the platform companies cannot ship and the agents cannot synthesise. It depends on the human in the chair *being someone*. Having taste. Having judgement. Having a perspective. Having the kind of relationship with the work that the agents do not, and probably will not, have.

That relationship — the kind of *authorship* that makes the work feel like it belongs to a person rather than a process — is, increasingly, the only signal the audience trusts.

That is the subject of the next chapter.

Chapter 12 — Authenticity as the New Scarcity

In early 2026, a stop-motion animator who goes by *Tiny Grandma* on YouTube uploaded a short to her channel. It was a stop-motion piece — claymation, frame by frame, the kind of work that takes weeks to make a few seconds of. YouTube’s AI-detection systems flagged it as AI-generated content and applied the platform’s automated labelling. The video went viral, not because of the animation, but because the platform’s automated system had wrongly flagged genuine human handcraft as synthetic.³⁷⁶

The story of *Tiny Grandma* is the perfect inverse of the Tilly Norwood story we opened with in Chapter 1.

If Tilly Norwood was the moment a synthetic creation tried to enter the working creative economy as if it were human, *Tiny Grandma* was the moment a human creation was wrongly identified as synthetic by the very systems that were supposed to protect the public from synthetic content. Both moments tell you the same thing, from opposite directions: the *signal* of whether a piece of creative work was made by a human is now an economic, cultural and legal asset of the first order, and the infrastructure for reliably establishing that signal is one of the most underdeveloped parts of the current creative economy.

This is the chapter about *provenance*. About why the question “did a person make this?” has become — in eight months — the single most important question in creative AI policy, and about what the people, companies and institutions trying to answer it are doing about it.

The death threats

In April 2026, *Dream Machine* Issue 23 reported, with as little editorialising as I could manage, that **Tilly Norwood’s creator Eline Van der Velden** had received death threats.³⁷⁷

The threats were not, of course, justified by anything. Death threats never are. But the cultural reaction that produced them — the visceral, sustained hostility that built up around the *idea* of a synthetic actress through the autumn of 2025 and the spring of 2026 — was not random. It was a specific response to a specific kind of cultural transgression. *You came here pretending to be one of us. You took something that belongs to us.*

The death threats are the extreme tail of a much larger curve of audience response that has been quietly shaping the AI creative economy for these six months. The slop ceiling in Chapter 5 — the 44%-to-3% Deezer ratio — is the polite version of the same response. The vehement audience pushback against AI art in *Call of Duty: Black Ops 7* and *Anno 117* in November 2025 is another version. The viral reaction to *Spotify’s* AI music infiltrating Discover Weekly playlists, the public anger at *McDonald’s* *Netherlands’* AI Christmas ad, the “*disturbing*” reception of *Valentino’s* AI handbag campaign — every one of these episodes is the audience saying, in increasingly direct terms, *we know what is human-made, we want what is human-made, and we are paying attention to who is trying to slip us something else.*

This is the cultural pressure that *authenticity-as-scarcity* describes. It is not, as some of the more dismissive AI commentary has framed it, a romantic attachment to old craft. It is a *market signal*. The audience is allocating attention, money and trust on a basis that increasingly weights human authorship as a positive variable. I have come, in talks since the

autumn, to call this the **Authenticity Premium** — the measurable excess of attention, willingness to pay, and cultural credit that audiences allocate to creative work whose human authorship can be verified. The Authenticity Premium is the *positive* side of the slop ceiling: the slop ceiling tells you what audiences *will not* engage with; the Authenticity Premium tells you what they *will pay extra for*. Both are market findings. Both are produced by the same underlying audience behaviour. The data is unambiguous. The strategic implication, for every working creative and every studio operating in this period, is also unambiguous.

In May 2026, **Bobby Berk** — the *Queer Eye* design lead — articulated the working-talent version of the same finding in a single line that I think is worth quoting in full because it captures the Premium argument from inside the unscripted-TV business: AI, he said, will make *reality TV* and “*verifiably human content*” *more valuable*, not less.³⁷⁸ What Berk is describing, in industry-of-the-thing terms, is the supply-and-demand mechanism the rest of this chapter is about. As the synthetic supply of *any* content category expands, the *verifiably human* corner of the same category accrues a scarcity premium. Reality TV is, structurally, the unscripted broadcast form most resistant to AI replacement — it is the form whose entire commercial proposition is *real people, in real situations, producing unscripted reactions whose value depends on us knowing they are real*. That is the Authenticity Premium drawn down to a single broadcast genre, and Berk’s read is, on the structural argument, exactly right.

The chess analogy I develop at length in Chapter 15 sits underneath this. The Authenticity Premium is what it looks like when an audience, faced with an infinite supply of machine-optimal work, allocates its scarce attention to the *deliberately un-machine-like* move. The 88% of UK respondents who wanted licensing-in-all-cases were articulating the same preference at the policy layer. The 44%-to-3% Deezer ratio was the same preference at the listening layer. The Television Academy’s “tools used to bring it to life” language was the same preference at the institutional layer. The Authenticity Premium is, at its core, the *commercial price* of the deliberately-human move — the move the engine, by construction, could not have made — and the audience’s reliable willingness to pay it.

The question is what the *infrastructure* for honouring that signal looks like.

Fingerprinting real media

I quoted Adam Mosseri — the head of Instagram — in Chapter 4 making the case that the platforms should focus on “fingerprinting real media” rather than chasing AI slop disclosure. The fuller version of his argument, made repeatedly across late 2025 and early 2026, was that the current approach to AI content moderation — trying to detect and label everything synthetic — is unwinnable.³⁷⁹ The volume is too high, the detection is too unreliable, and the labelling produces both false positives (Tiny Grandma) and false negatives (the AI hate-songs spreading across European Spotify charts in November 2025).³⁸⁰

The alternative Mosseri and others have argued for is *the inverse*: instead of trying to catch what’s synthetic, build infrastructure that can *prove what’s human*. A capture-time fingerprint — a cryptographic signature embedded by the camera, the microphone, the editing software, the upload pipeline — that travels with the file through its entire life on the public web.

The technical infrastructure for this is, as of 2026, partially built. The **Content Authenticity Initiative**, an Adobe-led coalition of camera makers, software companies and news organisations, has been working on it since 2019. By late 2025, **C2PA** (Coalition for Content Provenance and Authenticity) standards were supported by most major camera manufacturers, most major editing platforms, and a growing number of social-media uploads pipelines. The standards are robust enough that a photo taken with a C2PA-enabled camera, edited in Photoshop with C2PA-aware tools, uploaded to a C2PA-supporting platform, can carry a verifiable chain-of-custody for its entire provenance, from sensor to viewer.

Underneath this is **Google's SynthID** — a watermarking system that Google has been deploying across its AI generation tools, including Veo and Lyria.³⁸¹ In December 2025, the company announced that users could ask the Gemini app, “*Is this video made with AI?*”, and receive a reliable yes/no answer based on the SynthID watermark. By January 2026, this capability was available in the consumer Gemini product.³⁸² At Google I/O 2026, the company reported that **SynthID had marked over 100 billion items** across its own ecosystem and was being extended to partner platforms including **OpenAI, ElevenLabs** and **Kakao**.³⁸³ That cross-vendor expansion is the single most consequential development on the provenance side of the period this book covers: a watermarking standard, born inside one platform, is — for the first time — being adopted across the foundation-model companies that have until now competed against one another on every other axis. If the C2PA Content Credentials standard is the *capture-time* spine of authenticity infrastructure, SynthID-across-vendors is, as of May 2026, the closest thing the industry has to a *generation-time* spine.

These technologies are not, on their own, sufficient. Watermarks can be stripped by determined adversaries. C2PA chains break when files pass through non-compliant tools. The reliability of any given piece of provenance metadata depends on the integrity of every link in its chain. The trust infrastructure is still — relative to the speed of the AI rollout — early.

By June 2026 two more platform-side moves had landed and are worth flagging because they shift the labelling regime from voluntary infrastructure to platform default. **YouTube** rolled out an *automatic* AI-content labelling system: synthetic and AI-augmented videos uploaded to YouTube are, from this point onward, labelled by the platform itself rather than relying on creator self-disclosure.³⁸⁴ **LinkedIn** moved in the same direction from a different angle, announcing that it would *limit the reach* of AI-generated content on its feed — the first major professional-network platform to make the distribution-side argument that AI-authored posts deserve less algorithmic amplification than human-authored ones.³⁸⁵ The two moves point in the same structural direction: the platforms are moving from a *disclose-or-be-disclosed* regime to a *labelled-and-down-ranked* one, and the working creative who is reading the trend correctly is the one who treats provenance metadata as a *commercial asset* rather than a compliance cost. The labelled-human-authored work is, in this regime, the work that gets seen.

The deeper structural move underneath all of this — and the one I want to spend a paragraph on because it ties the chapter's authenticity argument into the IP-economics argument that Chapter 7 has built up — is the **redefinition of intellectual property from an asset to an environment**, and the corresponding migration of value from the *generation* layer to the

governance layer. For a century, IP was understood as a discrete, ownable thing whose value you protected by controlling access to it. What the 2026 deals reveal, when you stack them next to each other, is something different: a bounded space, governed by rules, inside which value is created continuously by people who are not you. **Hasbro's Sixth Wall** doesn't sell you Optimus Prime; it sells developers a governed space in which Optimus Prime can be summoned, on-brand and on-licence, a million times over.³⁸⁶ **Spotify's Verified-by-Spotify badge**, layered onto the *consent / credit / compensation* opt-in catalogue of the Universal Music deal, is the same move applied to recorded music — the platform doesn't try to detect and remove the synthetic flood, it draws a fence around the verified-and-licensed work and routes attention and money to the inside of the fence.³⁸⁷ The C2PA chain-of-custody stack, SynthID's 100-billion-watermark deployment, the Cannes AI Disclosure Standard, the Academy's "*you must be human to win*" rule, the YouTube and LinkedIn labelling regimes I have just described, the New York AI-avatar disclosure law and the SAG-AFTRA *Tilly Tax* are — read in this light — *the same project seen from different angles*. They are not labelling regimes for synthetic work. They are *governance infrastructure for the authenticity premium*: the architectural enclosure that lets the audience tell, at the speed of a swipe, which side of the line the work in front of them is on. The platforms and rights holders that *own* this governance layer — Hasbro with CharacterOS, Spotify with the verified-stack, Adobe with C2PA, Google with SynthID, the Cannes consortium with the Disclosure Standard — are positioning themselves to capture the durable value of the generative economy. The full IP-economics treatment of this argument, the *three Cs* (consent, credit, compensation) framework, the comparative analysis of five emerging licensing models, and the corrected record on the Disney–OpenAI deal that *never executed* sit in Appendix L: The Programmable Brand.

But what these technologies are doing, collectively, is establishing the *category*. They are saying: the question *did a person make this?* is technically answerable, with high reliability, given the right tooling. The next decade of cultural and legal policy in the creative industries will be — in significant part — about who controls that tooling, who decides what it certifies, and what economic value it carries.

If you want to know where the next ten years of investment, policy and platform politics in creative AI is going, watch the provenance layer. The companies that win the provenance infrastructure will be — in a real sense — the companies that own the *signal of authenticity* that the audience increasingly demands.

The trademark, the trust, the tax

The cultural pushback against synthetic content has produced, alongside the technical provenance infrastructure, a parallel set of *legal and contractual* defences that working creatives have begun deploying around their own work and identity.

Taylor Swift filed trademarks on her voice and image in early 2026, specifically citing AI deepfake concerns.³⁸⁸ **Matthew McConaughey** publicly drew the same line in January 2026.³⁸⁹ **Madonna and Will Smith** appeared in AI videos by Higgsfield in early 2026, the Madonna piece becoming a marquee example of how a major artist could *deliberately* deploy

synthetic imagery as part of their own brand.³⁹⁰ In the same vein, **The Rolling Stones** released *In The Stars* in May 2026, with a music video that used AI to de-age the band — produced, in a small piece of cross-industry casting that says something about where this category is heading, by the AI company belonging to the **South Park** creators Trey Parker and Matt Stone.³⁹¹ **George Clooney**, in November 2025, gave *Variety* the working actor's read on the synthetic-star economy: “*It’s been just like a writer creating characters. You fall in love with your characters when you’re writing them. It’s a wonderful process. It wasn’t like I just made her in a second, and that was it. You know, it took a long time.*”³⁹² Clooney was making, in his particular way, the same argument that the slop ceiling makes empirically: cultural stardom is a *function of time and human relationship*. It is not a function of generation cost. **Jeremy Renner** threatened a “multi-millions” lawsuit against an AI documentary director he said had used his voice without permission.³⁹³

In May 2026 the celebrity-defence layer took a meaningful *organisational* step. **Cate Blanchett** co-founded **RSL Media**, a non-profit explicitly chartered to address *consent around AI usage* — covering creative work, name, image and likeness — for performers across film, TV and music.³⁹⁴ This is, on my read, the first time the celebrity NIL-protection conversation has produced a *standalone institution* rather than a string of individual lawsuits and trademark filings. RSL Media is small as of mid-2026, but its founding signal is important. Where the existing infrastructure on celebrity AI consent has been individual (Swift’s trademarks, McConaughey’s line, Renner’s threatened suit) or statutory (the ELVIS Act, New York’s AI-avatar disclosure law), RSL Media is the first attempt at a *coordination layer* on the side of the talent — closer in shape to a Performing Rights Society than to a class action. If the Petrillo template I describe in Chapter 6 eventually has to be reconstructed for the NIL question, RSL Media is the kind of body that the reconstruction will need to anchor on.

In the same week, **Apple** acquired the talent and patents behind the AI-avatar company **Animato**, signalling — not for the first time — that the platform layer intends to own the celebrity-grade-avatar infrastructure rather than license it from third parties.³⁹⁵ The combination of an Apple-owned avatar pipeline and an RSL-administered consent regime is, in 2030 terms, the most plausible architecture for how the high-end NIL economy actually runs.

The posthumous case of the same architecture arrived in June 2026. The estate of **Ozzy Osbourne** announced that the late artist would “*live on as an AI avatar that talks with fans*” — a deliberate, consent-given (via the estate), avatar-shaped continuation of the artist’s public presence into a posthumous register.³⁹⁶ The Osbourne case is the first major celebrity-grade posthumous AI presence to be launched with the *estate’s* explicit framing rather than as a third-party reanimation, and it sits, structurally, in the same category as the Stones’ AI-de-aging *In The Stars* video and the Apple/Animato infrastructure: avatar work deployed deliberately, with the talent’s (or the estate’s) consent, in a register the audience is asked to recognise as honest. The pattern is clear: the audience does not reject the AI-rendered celebrity. The audience rejects the *un-consented* AI-rendered celebrity. Consent infrastructure — institutional (RSL Media), statutory (ELVIS Act, the New York avatar law),

platform (Apple’s owned-pipeline play) and estate-level (the Osbourne announcement) — is what makes the difference.

Underneath the celebrity layer, the structural infrastructure was being built. The **ELVIS Act**, Tennessee’s AI-impersonation law, had been used by the **Johnny Cash estate** to sue Coca-Cola over a tribute-act ad soundtrack.³⁹⁷ **New York** passed a law in December 2025 forcing advertisers to disclose when they were using AI avatars. The SAG-AFTRA statement on the law’s passage captured the political theory underneath the moment: *“These protections are the direct result of artists, lawmakers and advocates coming together to confront the very real and immediate risks posed by unchecked AI use.”*³⁹⁸ **Governments around the world** were considering bans on Grok’s app over an AI sexual-image scandal that broke in early 2026.³⁹⁹ By **May 2026**, the **AI Disclosure Standard** had been launched at the **Cannes Film Festival** as an industry coordination point for production-side AI labelling.⁴⁰⁰ The **Academy of Motion Picture Arts and Sciences** had — in a quietly consequential rule update — set the line *“You must be human to win”* for its 2026 awards.⁴⁰¹ The **Emmys** had set their own AI guidelines. The Television Academy’s language was a model of how to write a policy that defends authorship without picking a fight with the toolchain: *“The Television Academy reserves the right to inquire about the use of AI in submissions. The core of our recognition remains centered on human storytelling, regardless of the tools used to bring it to life.”*⁴⁰² *Tools used to bring it to life — not tools that did the work.* The grammar matters. **SAG-AFTRA’s** four-year contract — finalised by spring 2026 — included what the trade press informally called the **Tilly Tax**: a structured set of provisions for compensation, consent and residuals when AI replicas of human performers are used.⁴⁰³

Each of these is, on its own, a marginal piece of policy. Stacked together, they describe a new economic landscape: one in which *human authorship and identity* have become legally protected categories of creative work, with specific procedural and economic mechanisms for asserting them, defending them and compensating their use.

The cultural shorthand for this — *authenticity as the new scarcity* — captures the supply-and-demand logic. The legal shorthand — *human-authored work as a protected class* — captures the policy logic. Both are the same thing seen from different angles.

The financial reality: billion-dollar valuations alongside billion-dollar lawsuits

The Authenticity Premium, as I have described it throughout this chapter, is a cultural and market signal. It is also a financial one — and the financial picture, by June 2026, has a shape that the 2024-era “will AI destroy the music industry?” framing did not predict.

Suno, the AI music generation company that the music-rights organisations described as “the biggest theft in music history” when they sued it, raised **\$400 million** in June 2026 at a **\$5.4 billion valuation**. The round was led by Bond Capital. The company’s annual recurring revenue at the time of the raise was **\$300 million**; it had **2 million+ subscribers**.⁴⁰⁴ The headline that accompanied the raise told the financial story plainly: Suno is a multi-billion-dollar company that is simultaneously facing active copyright litigation from Universal Music Group and Sony Music, and has already settled with Warner Music

Group — a settlement that included a **licensing partnership** rather than just a damages payment.⁴⁰⁵

That combination — \$5.4B valuation, active lawsuits from UMG and Sony, settled partnership with Warner — is the financial shape of the post-Petrillo AI music settlement in real time. Warner’s calculation is that a licensing relationship with a \$5.4B company is worth more than the continuation of a lawsuit against it. UMG and Sony are making the opposite bet: that the litigation creates enough pressure to extract better terms than Warner got, or to establish the precedent that training on their catalogues without consent is not permissible at any price. Both bets may be rational. The point, for this chapter, is what the \$5.4B itself means. The market is saying, with extraordinary clarity, that the Authenticity Premium does *not* prevent the AI-music tools from attracting capital at billion-dollar valuations. The demand floor for AI-generated music — the content-farm market, the B2B brand-music market, the *Tonada*-style retail-background-music market — is large enough to sustain a \$5.4B business even if the premium streaming market remains, as the Deezer and Apple Music data suggest, almost entirely human-authored.

The Scorsese/Black Forest Labs episode I describe in [Chapter 7](#) reads differently in this financial context. When Scorsese became an advisor to Black Forest Labs — the FLUX model company — the backlash from Boots Riley and Guillermo del Toro was framed, in the press, as a *cultural* disagreement. But the underlying structure is, on closer inspection, an *Authenticity Premium* argument in action. Scorsese’s cultural authority derives, in large part, from the audience’s belief that his work has been made the hard way — that it is, in Jenna Ortega’s phrasing, the work the audience was *deprived of*, and has learned to hunger for. When he endorses a generative-AI company, the backlash is not a reaction to the endorsement’s legal implications. It is a reaction to the *signal* the endorsement sends about the source of his authority. The audience — and the peers who pushed back — is protecting the asset. That is, in market terms, exactly what the Authenticity Premium predicts.

What sincerity looks like in 2026

I want to come back to a distinction I made in Chapter 4, because it has held up through the last six months better than almost any other framing in this book.

I argued that audiences distinguish, very quickly, between *sincere* synthetic work and *cynical* synthetic work — that the underlying technology is the same, but the fingerprint of human intent behind the work is visible to the audience at the speed of a swipe.

The data from the spring of 2026 supports this. *Marketing Week*’s analysis “You can’t dismiss AI ads as slop when they’re winning in testing”⁴⁰⁶ documented that AI-generated advertising creative could, in fact, win in standard creative-effectiveness tests — *when the work was made with care, on a brief that respected the audience, by a team that had taste*. The same publication’s parallel coverage of the audience pushback against the McDonald’s Netherlands ad, the Valentino handbag campaign and a dozen other “AI slop” launches, made the inverse point. The technology is neutral. The intent is not. The May 2026 **David Beckham / Lenovo** “*Henchester United*” ad — in which the brand cheerfully used generative AI to render a Beckham-designed chicken coop for the footballer’s home flock — is, in my read, the

cleanest *sincere* example of the period: the AI is visible, the brief is playful, the talent is in on the joke, and the cultural response was warm rather than wary.⁴⁰⁷ The Beckham/Lenovo spot belongs in the same category as the Madonna/Higgsfield piece in the section above: AI deployed deliberately, with the talent's consent, in a register the audience recognises as honest.

The strongest AI-authored creative work of the period this book covers has, almost without exception, *not* tried to hide that it was AI-authored. Andrii Daniels' bomb-shelter clip foregrounded its conditions of making.⁴⁰⁸ Hoyt Dwyer's animated short for AI FilmFest Japan was upfront about its medium.⁴⁰⁹ *Dear Upstairs Neighbors*, the Google DeepMind / Connie He collaboration that premiered at Sundance, was *about* the constraints and possibilities of its production pipeline.⁴¹⁰ *Synthetic Sincerity*, Marc Isaacs' IDFA film, took the disclosure to the title of the piece.⁴¹¹ *Watch the Skies*, the AI-dubbed Swedish UFO feature, disclosed the dubbing process as part of its identity.⁴¹² *Lily*, the \$1m AI Film Award-winning Tunisian short, was framed by its director and reviewers as a piece *about* the new toolchain.⁴¹³

The pattern, repeated across thirty or forty examples I have looked at carefully, is the same: *AI work that owns its synthetic nature, and that is made with human creative intent, finds an audience.* AI work that tries to pass as something it isn't gets the audience response that Tiny Grandma's stop-motion got *from the algorithm* — an immediate, automatic, suspicious flag.

This is, in market terms, a stable equilibrium. It is the market that the slop ceiling and the audience pushback have built. And it is, for working creatives, a manageable and even encouraging environment to operate in. The audience is not against AI. The audience is against being lied to.

What disclosure should look like

I want to lay out — because I have been asked this in every Q&A I have done since starting the newsletter — what I think the *practical* shape of authenticity infrastructure should look like for working creatives in 2026.

It is four things, in increasing order of investment:

One. Disclose, consistently. If you use AI in any part of your work, say so. In your credits. On your website. In your contracts with clients. In the metadata of your files. The act of disclosure does, in my experience, not cost you anything with the audience — the audience that is going to reject AI work would reject it anyway, and the audience that is going to accept it is the audience that values you being straight with them. The cost of *getting caught* not disclosing, in this environment, is materially higher than the cost of disclosing.

Two. Document, deliberately. Keep logs. Keep notes. Keep prompt histories. If a piece of work you make this year ends up being legally or culturally contested in 2030 — and a non-trivial fraction of work made this year will be — your ability to *show your workings* will be the difference between defending the work and losing it. The Sundance literacy initiative's emphasis on *evidence of human authorship* is exactly right.⁴¹⁴

Three. Watermark, where appropriate. Use SynthID, C2PA, or the equivalent provenance layer that your toolchain supports. If your work doesn't yet support these standards, ask your tool vendors when they will. The market for tools that support provenance metadata is, in 2026, larger than the market for tools that don't.

Four. Build the chain. If you are running a studio or an agency, build the *internal* infrastructure for verifying and tracking the provenance of your work end-to-end. The cost of doing this in 2026 is moderate. The cost of *not* doing it in 2029, when a client asks for the chain-of-custody on a piece of work and you can't produce it, is going to be much higher.

These are not, on their own, business strategies. They are, in 2026, the *minimum hygiene* for operating a credible creative practice in the AI era. Treat them as you would treat health-and-safety on a film set. Do them as a default. Do them well. Then get on with the work.

The provenance infrastructure, named: thirty-six pieces of the puzzle

I want to lay out a more complete map of the provenance infrastructure that is being built in 2025–26, because the technical-and-policy stack is more advanced than the public conversation has caught up with, and working creatives reading this book need to know what is actually in the field.

Stacking the moves I have referenced across this chapter and the rest of the book, the inventory is roughly this:

Capture-time signing and provenance metadata:

1. **C2PA / Content Credentials** (Adobe-led, with Microsoft, Sony, Nikon, Leica, BBC, Canon participation) — the cryptographic-signature standard for capture-and-edit-time provenance.
2. **Leica M11-P / M11-D** — first consumer cameras with C2PA at the firmware level (2023, expanded through 2025).
3. **Sony Alpha 1 II and Alpha 9 III firmware** — C2PA support across pro Alpha range.
4. **Nikon Z9 firmware** — C2PA support for the AP/Reuters wire-service workflow.
5. **Canon professional bodies** — Content Credentials integration through 2025 firmware cycle.
6. **iPhone Camera app** (selected models, 2025–26) — capture-time signature support.
7. **Adobe Photoshop, Premiere, Lightroom** — edit-time chain-of-custody preservation through Content Credentials toolchain.
8. **DaVinci Resolve, Final Cut Pro, Avid Media Composer** — partial C2PA support through 2025–26 update cycle.
9. **Capture One** — Content Credentials integration for the high-end commercial-photography workflow.

Synthetic watermarking and detection:

10. **Google DeepMind SynthID** — embedded watermark for Veo (video), Lyria (audio) and Imagen (image) outputs. Over 100 billion items marked by May 2026; extended to OpenAI, ElevenLabs and Kakao as a cross-vendor standard.
11. **SynthID Verification via Gemini app** — consumer-facing yes/no detection.

12. **Lyria 3 SynthID extension** — audio verification across the Google music-model family.
13. **YouTube AI Detection Tool** — automated content classification (the *Tiny Grandma* false-positive case demonstrating both its scope and its current accuracy ceiling).
14. **Deezer’s AI-music detection pipeline** — identifies up to 75,000 AI-generated tracks per day at upload time.
15. **Spotify AI Transparency Beta** — voluntary creator disclosure surfaced in the consumer player UI.
16. **Beeble** — independent detection-and-watermarking infrastructure used by some news organisations.
17. **Cloudflare AI bot classification** — public-web-infrastructure-level tracking of AI crawlers and agents.
18. **Sony music-identification technology** — identifying source recordings inside AI-generated outputs at the catalogue level.

Institutional and contractual disclosure:

19. **Cannes AI Disclosure Standard** — industry-coordination production-side labelling standard for the festival circuit (launched May 2026).
20. **Sundance AI Literacy Initiative** — creator-training programme funded by Google’s \$2M commitment, training 100,000+ artists in foundational AI literacy and provenance practice.
21. **Academy of Motion Picture Arts and Sciences “you must be human to win” rule** — 2026 awards eligibility update.
22. **Television Academy AI guidelines** — “tools used to bring it to life” framing for Emmys submissions.
23. **SAG-AFTRA Tilly Tax provisions** — the consent, compensation and residuals framework for digital-replica use, included in the 2026 four-year contract.
24. **Equity (UK) strike ballot outcome** — 99% vote authorising industrial action over AI provisions.
25. **PRS for Music AI Survey 2026** — UK music-creator sentiment baseline informing collective-licensing negotiations.
26. **GEMA ruling against OpenAI** — first European-rights-society legal precedent on AI training compensation.
27. **The 88% UK consultation outcome** — political mandate for licensing-by-default.
28. **The *Stealing Our Work Is Not Innovation* declaration** — 800-creator cultural marker.
29. **Bandcamp’s outright AI-music ban** — distribution-platform-level disclosure-by-exclusion.
30. **Swedish Top Chart AI ban** — chart-eligibility-level disclosure.
31. **San Diego Comic-Con AI art ban** — convention-level disclosure-by-exclusion.

Legal infrastructure protecting human identity:

32. **Tennessee ELVIS Act** — the most-cited state-level AI-impersonation statute, used by the Johnny Cash estate in the Coca-Cola tribute-act case.

33. **New York's December 2025 AI-avatar disclosure law** — requires advertisers to disclose AI-performer use.
34. **The Taylor Swift voice-and-image trademark filings** (early 2026) — celebrity-led use of trademark mechanism for identity protection.
35. **The Jeremy Renner unauthorised-voice lawsuit** — pending case testing voice-likeness protections.
36. **The pending EU AI Act enforcement** — training-data-transparency requirements coming into effect through 2026.

Each of these, on its own, is a marginal piece. Stacked together — capture-signing, watermarking, platform integration, festival rules, awards rules, union contracts, civil society declarations, legal protections — they describe a *coherent infrastructure project* that the creative industries are, in eight months, jointly constructing.

The project is, by any reasonable assessment of similar previous infrastructure builds, *substantially ahead of schedule*. The C2PA standards body was founded in 2021 and was, by mid-2026, deployed across most major commercial capture and edit tooling. SynthID went from research demo in 2023 to consumer-facing detection in Gemini by January 2026. The SAG-AFTRA digital-replica provisions went from a 2023 strike demand to a contractual reality in 2026. The 88% went from political abstraction to government statement of progress in twelve months.

The thing this rate of progress tells me is *not* that the work is done. The work is, in many places, half-done — there are gaps in adversarial robustness, in platform UI integration, in cross-jurisdictional enforcement, in coverage of the long tail of creator categories outside the major commercial industries. The work is also being done unevenly: the music industry has built more of the stack than the games industry, which has built more than the publishing industry, which has built more than the regional and minority-language creative ecosystems that the next decade will need to bring in.

But the *trajectory* of the work is unambiguous. The provenance stack is being built. The institutional disclosure infrastructure is being built. The legal protections are being built. The audience contract I describe in the next section is being written. Working creatives who position themselves on the *inside* of this build — using the tools, contributing to the standards, showing up at the consultations, advocating with the unions, deploying provenance metadata in their own work as a default — will have, by 2030, materially more leverage than working creatives who waited for someone else to finish the project for them.

What this means for the audience

I want to close the chapter with a thought about what this whole structure means for *the audience*, because most of this book has — by design — been about the people who *make* creative work, and the audience is sitting on the other side of the screen the whole time.

What I think the slop ceiling, the provenance infrastructure, the disclosure norms and the legal protections are, *collectively*, building is a *new contract* between makers and audiences.

The old contract was straightforward. The maker made the thing. The audience watched, listened, played, read. The signal of authenticity was implicit — most creative work was, by

default, made by humans because there was no other way to make it.

The new contract is, by necessity, *explicit*. The maker discloses what was made by whom and how. The audience gets to make an informed choice. The platform, the union, the law and the institution all support both sides of the transaction.

If we get this contract right, the AI era is not the end of human creative work. It is a *renegotiation* of the terms on which human and synthetic creative work coexist in the public sphere — with the audience, for the first time in a very long time, getting a real seat at the table.

If we get it wrong — if the disclosure infrastructure fails, if the provenance metadata is unreliable, if the platforms refuse to honour the audience's stated preferences, if the legal protections are not enforced — what we get is the world the *Dead Internet* chapter described. A web of synthetic content, made by no one in particular, for no one in particular, churning past an audience that has lost the ability to trust any of it.

The choice between those two outcomes is *not, in 2026, a technical question*. The technical infrastructure for both is, by spring 2026, broadly in place. The choice between them is a *political, institutional and cultural* one. It is about whether the people who set the rules — the platforms, the legislators, the institutions, the studios, the audience itself — collectively decide that *knowable human authorship* is a public good worth protecting.

I think, on the evidence of the last six months, that the choice is being made — slowly, contentiously, imperfectly, but recognisably — in the right direction. The 88%, the Sundance literacy turn, the Cannes Disclosure Standard, the Academy's rule update, the SAG-AFTRA contract, the C2PA standards, the SynthID rollout, the audience's own attention behaviour: these all point the same way.

The philosophical argument underneath all of this — the case that when technical execution becomes a commodity, the distinctively human acts of *intent, risk-taking, and the willingness to mean something* become the only irreplaceable contribution — is made at length in [Appendix G: The Age of Intent](#). The *Process Trap* companion piece at [Appendix J](#) fills in the darker side of the same picture: why the reassuring story that “automation will take the rote but real creative work is safe” is simultaneously true of the *capacity* and catastrophic for the *workforce class*.

The question for the rest of this book — Chapter 13 on the organisational restructuring, Chapter 14 on the labour-market reshuffle, and Chapter 15 on the political choice — is what happens to the *organisations*, the *labour market* and the *economy* of creative work when authenticity is the scarce good and the orchestrator is the new role. The implications for how teams are structured, how labour is paid and how creative careers are built are bigger than any single tool launch, and they are what the next three chapters are about.

Chapter 13 — Coordination Collapse

There is a working assumption underneath almost every conversation about AI in the creative industries that I want, in this chapter, to make explicit and then take apart.

The assumption is that AI is a *technology change*, broadly equivalent in shape to other technology changes the industries have absorbed in the past — the arrival of digital, the move to streaming, the rise of mobile, the emergence of social. The implication is that the existing institutions of the industry — the studios, the agencies, the labels, the unions, the publishing houses, the broadcasters — will absorb this change the way they absorbed the previous ones: with some restructuring, some layoffs, some new hires, some new departments, and a generally familiar shape on the other side.

I do not think this is what is happening.

What I think is happening is that AI is, specifically and quite differently, a *coordination technology*. It is changing — not at the margin, but at the core — what it is possible for a single person to know, decide and execute about a complex creative project. And because the existing organisations of the creative industries are, structurally, *coordination architectures* — they exist to allow many people to work on a single piece of work — the change in coordination economics is changing the organisations themselves.

This is the chapter about what happens to creative *organisations* — studios, agencies, labels, broadcasters, indie companies — when AI reaches a certain level of capability. It is also, by necessity, the chapter about what happens to creative *careers* when those organisations change shape.

The shorthand I have come to use for what is happening is **coordination collapse**.

What organisations are for

A studio, an agency, a label, a publishing house — these are not just brand names attached to creative outputs. They are *organisational technologies* that solve a specific problem: how do you get fifty or five hundred or five thousand people to coordinate on a single piece of creative work, well enough that the result is coherent, on-budget, on-deadline and good enough to put out into the world.

The way they solve that problem is by *layering* the work into specialised roles, building *hierarchies* that direct the work down through those roles, *processes* that move material between the roles in predictable order, and *cultural norms* that make the whole apparatus run with less explicit instruction than you would otherwise need.

A film studio exists, structurally, because making a film requires the coordinated work of many specialists — writers, directors, actors, cinematographers, designers, editors, sound designers, composers, marketers, distributors. The studio is the *coordination apparatus*. The film is the output of the apparatus.

When AI starts to do the work of many of those specialists — not entirely, but at the level of *first draft* or *junior contribution* — the calculus of the coordination apparatus changes. Suddenly, a much smaller human team, working with a large pool of synthetic capacity, can produce the same coordinated output that previously required the large team. Suddenly, the

bottleneck of producing creative work is no longer the size of the team. It is something else: the ability of the small senior team to *direct* the synthetic capacity well.

The implication, which has been dawning on the creative industries through the autumn of 2025 and the spring of 2026, is that the existing organisational shape — the *studio shape*, the *agency shape*, the *label shape** — was built for a coordination problem that no longer exists in the form it used to.

This does not mean the studios will disappear. It does mean that the *shape* of the studios is going to change, in a hurry, in ways that working creatives need to understand if they are going to be on the inside of the change rather than on the receiving end of it.

The shadow workforce

The first symptom of coordination collapse, as it has shown up across the creative industries in this period, is the rise of what the workplace-research literature has started calling **shadow AI** — the practice of using AI tools in your job *without telling your employer*.

The numbers, from a series of 2025 studies I covered in *Dream Machine* Issue 5, are extraordinary.⁴¹⁵

Roughly **half** of U.S. employees — 45–52% in different surveys — have used AI in their jobs without telling their bosses, with Gen Z and tech-sector workers being the most frequent secret users.⁴¹⁶

About a third — 29–33% — pay for their own AI tools out of their own pockets without their employer’s knowledge.⁴¹⁷

Roughly **56–57%** of regular AI users admit to *hiding their usage or presenting AI output as their own work* to avoid judgement or stigma. Nearly half of executives do the same.⁴¹⁸

52% of workers won’t admit to using AI at work — *even when asked directly*.⁴¹⁹

These numbers describe a workforce that has, en masse, started running its own private parallel productivity infrastructure that bypasses the official organisational tooling, the official organisational processes, and the official organisational accounting of where the work is being done and by whom.

This is not a niche phenomenon. *Half the workforce*. And it is concentrated, the surveys suggest, in exactly the demographic — Gen Z, tech-sector, knowledge worker — that is most likely to be the future workforce of the creative industries.

The numbers escalate the more recent the research. By the end of 2025, enterprise-AI tracking data put active daily use at **88–89%** of staff across organisations, with **71–80%** of those users running their tools entirely outside any official approval or IT oversight.⁴²⁰ What the workplace-research firms have started calling the “Hidden Cloud Explosion” describes a six-month period in which the average enterprise IT department’s visibility into the AI tools its workforce was using simply collapsed: organisations believed they were running on roughly **91** public cloud services per enterprise, while network-level analysis put the actual figure at **1,220** active services — a 90% visibility gap.⁴²¹ In the same year, **20%** of

organisations reported severe security incidents linked directly to shadow AI, with the average breach cost going up by **\$670,000**; in **65%** of those incidents personally identifiable information was exposed, and in **40%** intellectual property was directly leaked.⁴²²

For the creative industries this matters disproportionately, because the data being fed into public LLMs by shadow users — proprietary scripts, unreleased concept art, client briefs, internal pipeline code, unmastered audio stems — is the *exact intellectual property* that organisations are simultaneously suing AI companies for scraping. The studio whose general counsel is in federal court against a frontier-model company is, on the same Tuesday afternoon, watching its own animation department paste asset descriptions into the same company's consumer chatbot to speed up metadata writing. Both things are true. Both happen at once.

The shadow workforce, in coordination-collapse terms, is the symptom of an organisational architecture that is no longer aligned with the work the people inside it are actually doing. The official architecture says *we hired these humans to do these specific jobs, in this specific way, at this specific pace*. The shadow architecture says *these humans are now hybrid human-agent operators producing more, faster, with different qualitative properties than the official architecture is set up to manage*.

What you get, when these two architectures sit on top of each other, is a workforce that is *measurably more productive than the official metrics show, doing more work than the official scope says, with less institutional knowledge of how that work is being done* than ever before.

This is, in coordination terms, an unstable equilibrium. It cannot last indefinitely. The question is how it resolves.

The “AI for thee, but not for me” paradox

The pattern that the shadow-AI numbers describe is not random distribution of tool use. It is *hierarchical*. Creative workers, across every survey I have read in this period, exhibit a consistent psychological pattern that the developer and creative-community discourse has named the “**AI for thee, but not for me**” paradox.⁴²³

The pattern works like this. Creative professionals identify some tasks as *mine* — the writing, the cinematography, the composing, the performance, the lead concept — and other tasks as *not mine* — the marketing copy, the project email, the deck assembly, the metadata tagging, the routine code, the contract redline, the rough mix, the asset variation. The first category is defended fiercely against AI substitution; the second is offered to AI substitution without much thought. The moral framing of the technology shifts depending on whose labour is being replaced.

Look at the music sector. An industry survey of more than **1,100** professional producers, songwriters and audio engineers in 2026 found that **87%** were actively using AI tools in their creative process.⁴²⁴ The internal distribution, though, tracked the hierarchy: **58%** used AI for audio restoration and cleanup, **38%** for mixing assistance, **33.9%** for automated mastering — high-friction tasks that nobody felt sentimental about — while only **20.9%** admitted to using

AI for composition or lyric generation, the parts of the craft on which the personal artistic identity sat. 77% cited “loss of originality” as their primary concern, outranking even the fear of personal job displacement (42%). The artist’s relationship to the tool, in other words, is not consistent across the work. It is sharply conditional on whose labour is being substituted.

The same hierarchy shows up in film. A survey of professional screenwriters before and after the 2023 WGA strike found that pre-strike covert AI use sat at around 34%; once the WGA’s negotiated guidelines legitimised AI assistance for *formatting, structural outlining and brainstorming*, that number jumped to 68% by 2024.⁴²⁵ What the regulation changed was not the technology. It was the stigma. The shadow use moved into the light, with no measurable decline in the work product. The covert hierarchy became an overt one.

This is, in my view, the most uncomfortable observation in the entire shadow-AI literature, and it is the one that creative-organisation leadership has the hardest time admitting publicly. The same professionals who, in their public statements, treat AI training as a moral violation are, in their private practice, the heaviest users of the same underlying technology. The hypocrisy is not a character defect. It is a structural property of how knowledge workers self-defensively triage their own tasks under productivity pressure. Read charitably, working creatives are doing what every economic actor in a productivity transition has done: protect the highest-value labour and offload everything else.

The cost of that triage is moral clarity. It is hard to credibly argue that AI training is theft when you are typing your portfolio description into Claude. The vocal-protest economy and the silent-adoption economy now run on the same desks, often within the same hour.

The consumption gap: what the data actually says

I want to spend a section on the gap, because the gap is *the* macroeconomic story of this period and I do not see anyone telling it cleanly.

The gap is between *what the creative industries are saying publicly about AI* and *how the creative industries are actually using AI on a Tuesday morning*. The two pictures are not slightly different. They are, in significant measure, contradictory.

Take Adobe, because Adobe is the cleanest single case study. Adobe’s Firefly generative-AI suite — the same product that the working creatives in my surveyed circle are most ambivalent about — passed **22 billion AI-generated assets** by April 2025, eighteen months from public release.⁴²⁶ By that point, **45%** of all Creative Cloud subscribers had engaged with Firefly. **70%** of active Firefly users were using the tool *every week*, averaging 2.8 sessions weekly at 26 minutes each. Firefly contributed **11%** of all new annual recurring revenue at Adobe in 2024 — the company’s fastest-growing revenue catalyst since the original move to a subscription model — and Adobe’s AI-first ARR more than *tripled year-over-year* in the first quarter of fiscal 2026.⁴²⁷

That is not the adoption curve of a niche professional tool. That is the adoption curve of a default productivity feature in the dominant creative-software stack on the planet. **72%** of Fortune 500 design teams have formally integrated Firefly. **63%** of marketing agencies. **58%** of e-commerce design departments. **48%** of UX/UI designers.⁴²⁸ *Twenty-five per cent* of new

Adobe Stock contributions in 2024 contained Firefly-generated elements.⁴²⁹ The Adobe MAX 2025 Creators' Toolkit Report's headline number — **86%** of global creators using generative AI — sits inside this pattern, not against it.⁴³⁰

If you sat with the public discourse alone — the open letters, the boycotts, the strike statements — you would assume working creatives were broadly refusing AI integration. The actual platform telemetry, in a year where Adobe shared more of its numbers than usual, says the opposite. Working creatives are not refusing. They are adopting at a pace that Adobe's growth team is, by every signal I can read, struggling to keep ahead of.

The same picture holds across the toolchain. ChatGPT, by mid-2025, was on **800–900 million weekly active users**.⁴³¹ Anthropic's Claude was the writers' and developers' second favourite, with rapidly increasing usage in long-context creative tasks. Google's Gemini was growing desktop users at **155%** year-over-year, more than six times faster than ChatGPT's 23%.⁴³² These are not user numbers that reflect a market in revolt. They are user numbers that reflect a market that has, in private, decided.

And the consumer side mirrors the producer side. The Stanford AI Index 2025 found that **55%** of individuals across 26 countries view AI products as offering more benefits than drawbacks — up from 52% in 2022.⁴³³ A 2024 YouGov poll across 17 markets found that nearly a third of consumers felt *more positively* about generative AI than the previous year, against only 22% feeling more negatively.⁴³⁴ In the gaming sector — which has produced some of the loudest anti-AI consumer backlash of this year — the same Quantic Foundry survey that showed audiences are **77–83% negative** toward AI-generated quests and dialogue also showed that **60%** of gamers remain entirely neutral about AI in a game's development *provided the final product is of high quality*.⁴³⁵ The hostility is not generic. It is specifically aimed at AI in the creative roles where audiences expect to feel a human soul. Everywhere else — UI, backend, balancing, localisation, dynamic difficulty — the audience is, on aggregate, indifferent.

Even the GDC sentiment data, which is often cited as evidence of an industry in retreat from AI, tells the same paradoxical story when you read it as a whole. *Personal* generative-AI usage among professional game developers rose from **31%** in 2024 to **36%** in 2026, while *industry sentiment* over the same period cratered from 18% negative to **52% negative**.⁴³⁶ Use went up while approval went down. The two lines should, in a coherent market, move together. They are not. They are diverging.

I want to be careful about what conclusion to draw from this. It is *not* that the public discourse is wrong and the silent adopters are right. The public discourse is doing genuine political and cultural work — it is what produced the 88%, the SAG-AFTRA contract, the GEMA ruling, the Sundance literacy turn, the Cannes Disclosure Standard. Without the loud minority, the creative economy would have no political leverage at all.

The conclusion to draw is more uncomfortable. The creative industries are, in 2026, operating with *two parallel economies on top of each other*. In one, AI is a moral crisis, a labour threat, and a contested category of production. In the other, AI is a default productivity feature being

integrated at the speed of any other software upgrade. The same individuals, the same teams, the same studios are participating in both economies simultaneously, often without acknowledging the contradiction.

The question for the next eighteen months — the question I keep coming back to when I talk to studio leadership — is whether the two economies merge into a single, honest, integrated practice (the **path two** integration I describe below), or whether they continue to run in parallel, with the public economy producing the policy and the private economy producing the work. The first outcome is harder but produces better collective decisions. The second outcome is the path of least resistance, and is, in my view, where we will end up by default if working creatives, studios and unions do not deliberately close the gap.

For the data and the sectoral mechanics behind this section — the linguistic markers of covert AI use, the labour-market dynamics of agentic displacement, the deeper analysis of Adobe / OpenAI / Anthropic adoption telemetry, and the consumer sentiment / consumption asymmetry — see the two research deep dives that this chapter draws on: Appendix D: The Shadow AI Paradox and Appendix E: Dynamics of Generative AI Adoption. For the case against the *civilizational-risk* school that treats AI-driven coordination failure as an extinction scenario — the empirical critique of every major AI-doom prediction and why the data consistently refutes the worst-case framing — see Appendix I: The Doomer Mistake. For the deeper organisational-design treatment of the *structural* argument this chapter is making — the 2,000-year history of hierarchy as an information-routing protocol, the *jagged technological frontier* identified by Dell’Acqua et al.’s preregistered BCG experiment, the *vigilance problem* that makes naïve human-in-the-loop oversight unreliable, the three competing post-hierarchical models (Block’s intelligence layer, Every’s agent colony, DreamLab’s Dynamic Agentic Mesh), the new KPIs for fluid organisations and the *judgment broker* role replacing middle management — see Appendix K: The Coordination Collapse, Dr. John O’Hare’s April 2026 sixteen-chapter research report that this chapter’s voice rests ⁴³⁷ on.

The June 2026 Epidemic Sound *Future of the Creator Economy Report*, surveying 3,000 professional creators across the UK and US, sealed the picture and is worth pulling out because it is the cleanest single piece of evidence on the creator side I have seen this year. **94%** of creators now use AI tools; **72%** expect to *increase* their usage over the next twelve months; **57%** use AI primarily as a *workflow accelerant* rather than a generative content engine; only **28%** regularly produce fully generative content; and **89%** report feeling pressure to use AI *just to keep up with industry expectations* — the clearest articulation I have seen of the silent-adoption-under-public-protest pattern this chapter describes. The internal contradictions of the creator response to AI sit on top of each other in the same report: **75%** believe human-created content will become a *premium product* in the AI era, **83%** say human-made sound creates a stronger emotional connection with audiences, **79%** say music and sound directly impact engagement and revenue, **75%** believe disclosure is essential for audience trust, **73%** worry that unclear licensing could limit their future business opportunities, and **62%** have already experienced copyright or licensing issues in the past two years. *Musically’s* read of the same data — that creator attitudes towards AI have become “increasingly nuanced”, embracing the productivity gains while demanding greater

transparency, clearer licensing frameworks and stronger protections around identity, voice and creative work — is the read that, in mid-2026, has finally caught up with the contradiction this chapter is built around.⁴³⁸ The simplest summary of the report is the one that sits underneath the entire second half of this book: **AI may drive scale, but long-term value will belong to creators who combine new tools with authentic human creativity and strong control over their intellectual property.**

There is one further dimension of the consumption gap I want to flag here, because [Chapter 10](#) develops it at length and [Chapter 4](#) introduced it: the gap between *production* and *consumption* is not just an organisational misalignment, it is a *biological* one. **Aggregate human attention is finite.** The same Adobe Firefly that has generated 22 billion assets, the same ChatGPT that serves 900 million weekly users, the same Sora app that hit a million downloads in five days — these are all systems whose *production* side scales without bound and whose *consumption* side is bounded by the eleven-or-so daily hours of media attention the average adult can physiologically deploy. The consumption gap I have described above is, at its widest, this binding constraint expressed as an organisational problem. Studios that produce AI-augmented content at the rate the toolchain now allows — without recognising that the audience cannot consume more hours per day than it already does — are, on inspection, optimising the wrong side of the supply-demand equation. The studios that integrate AI productivity gains into work that *earns* a larger share of the audience's finite attention budget will win the next decade. The studios that integrate AI productivity gains into *more output competing for the same finite budget* will, on the historical pattern, hit the slop ceiling on a balance-sheet timeline they did not plan for.

The two paths

The shadow workforce can resolve, broadly, in one of two directions, and I think the choice between them will be the central organisational question for every studio, agency and label in the creative industries over the next three years.

Path one is *suppression*. The organisation decides that the shadow AI use is a risk — to security, to IP, to brand, to compliance, to the official productivity metrics — and shuts it down. Tightens the rules. Audits the work. Punishes the offenders. Reverts to the official architecture and the official tooling.

This is, in my view, a losing strategy in the medium term, because the productivity advantages that the shadow workforce is capturing are real, and the workers who are capturing them will, given the choice, work for organisations that *let them keep capturing them*. The suppressing organisation will progressively lose its most AI-fluent workforce to organisations that allow the hybrid practice.

Path two is *integration*. The organisation accepts that the shadow AI use is happening, decides to make it official, builds the infrastructure to support it, sets the norms to govern it, and re-shapes the work — and the workforce — around it.

This is, in my view, the right strategy. It is also the one most major creative organisations have been quietly moving towards in the period this book covers.

EA's push of its 15,000 employees to use AI as a "thought partner" was, structurally, a path-two move. **Krafton's** transformation into an "AI-first" company in November 2025 was a path-two move.⁴³⁹ **Disney's** Office of Technology Enablement was a path-two move. **WPP's** AI overhaul, **Adobe's** AI in everything, **Sony's** 49-agent game team — all path-two moves.

The path-two organisations are, structurally, betting that integration produces *more* output, *more* quality and *more* employee retention than suppression. The early evidence, six months in, suggests they are right.

The mid-career squeeze

The cost of the path-two transition has not, on the whole, been borne by the senior creatives or the entry-level workforce. It has been borne by the *middle*.

In **April 2026**, *Dream Machine* Issue 24 reported that the publisher behind *Grand Theft Auto VI* had laid off the entire seven-year-old internal AI team it had built to develop in-house AI capability for the franchise.⁴⁴⁰ The framing, in the press release and the subsequent industry coverage, was that the company had decided to use off-the-shelf AI tools instead of maintaining proprietary ones — and that the seven-year AI investment was, in retrospect, a "backlash cleanup" cost.

The story was repeated across multiple studios. **Disney**, in April 2026, laid off staff including in its Marvel division, in moves the company did not blame on AI but whose timing was, as the trade press noted, "loaded."⁴⁴¹ **Meta** had cut 10% of its Reality Labs staff in January 2026 to refocus on AI.⁴⁴² **Scottish animation studio** Axis Animation collapsed in early 2026, with its closure publicly attributed in part to AI competition.⁴⁴³ **Ubisoft** cancelled five games, including the *Prince of Persia* remake, in January 2026, in order to refocus on AI.⁴⁴⁴

The pattern across these cases is the same: the *mid-career* layer — the experienced specialists in the middle of their professional lives, doing the day-to-day production work that the senior creative leadership directs — is the layer absorbed into agentic capacity first.

This is the unambiguous bad news of the AI transition. The *Guardian* covered this directly in January 2026 with a piece titled "AI is hitting UK harder than other big economies, study finds," which found that mid-career knowledge workers in the U.K. were experiencing disproportionate displacement compared to peers in the U.S., Japan, Germany and Australia.⁴⁴⁵ *Economist* coverage in late November 2025 had been the early signal: "Investors expect AI use to soar. That's not happening."⁴⁴⁶ — meaning that the AI investment thesis was not, in the short term, producing the aggregate productivity gains the investors had hoped for, but *was* producing concentrated labour displacement in specific sectors.

The OpenAI public-policy response to this, articulated through April 2026, was a series of proposals — robot taxes, public wealth funds, a 4-day workweek — designed to manage the economic disruption of AI-driven productivity gains.⁴⁴⁷ *Dream Machine* Issue 24 covered these proposals at length. The framing OpenAI used was telling: the company was no longer arguing that AI would *not* cause disruption. It was arguing that the disruption was inevitable and that society needed to build new mechanisms to manage it.

The *Economist*, in a piece titled “Job apocalypse? Humbug! AI is creating brand new occupations,” took the contrary position — that AI was, on net, creating more new jobs than it was destroying, and that the framing of mass displacement was overstated.⁴⁴⁸ Both positions are partially right. The aggregate employment numbers, across creative industries in 2026, did not show the apocalyptic decline some had predicted. But the *composition* of employment changed sharply. Senior orchestrator roles increased. Mid-career specialist roles decreased. New AI-specialist roles — AI agent developers, prompt engineers, AI ops specialists — exploded. The *Forbes* piece from November 2025 noted that “vibe coding” — natural-language software development — was an in-demand AI skill that paid up to \$220,000.⁴⁴⁹

What the labour market is doing, when you look at it carefully, is *not* destroying jobs in the creative industries. It is *reshuffling* them — towards a smaller number of senior strategic roles, a different mix of specialist roles, and a much larger pool of AI-tooling skills that span the old discipline lines. The mid-career creative who fails to make this transition is the one who is at risk. The mid-career creative who makes it well — by upskilling deliberately, by claiming the orchestrator role, by building a portfolio practice — has, by every indicator I can see, *more* leverage in the labour market than they did before.

The hard truth is that the transition is not equally available to everyone. It depends on access to training, on access to tooling, on workplace cultures that support experimentation, on time to reskill that workers with caring responsibilities or financial precarity often don’t have. The institutional response to this — the Sundance Literacy initiative, the UK free-AI-training programme, the Adobe and Google educational investments — is real but partial. The structural inequities of who *can* make the transition are real and concerning.

The neurodiversity dividend

One genuinely encouraging finding from the period this book covers came from the U.K. Department for Business and Trade’s research on neurodiverse workers in AI-tooled workplaces.

The study, published in late 2025, found that workers with ADHD, autism and dyslexia were **25% more satisfied** with AI assistants than neurotypical workers, and that they reported AI agents as actively helping them succeed at work.⁴⁵⁰ The interpretation, reported in *CNBC* in November 2025, was that AI tools were lowering the cognitive load of tasks that had historically been disproportionately punishing for neurodivergent workers — coordinating complex calendars, parsing dense documents, structuring written outputs — and were, as a result, *levelling the playing field* in workplaces that had previously underutilised neurodivergent talent.⁴⁵¹

I want to flag this finding because it is one of the cleanest counterexamples to the “AI is bad for workers” framing that I have come across, and because it is a useful corrective to the labour-displacement narrative that has dominated much of the coverage of this period.

AI is, demonstrably, *good* for some workers. It is good for the workers who are most able to leverage it, and it is also good for the workers whose existing labour-market participation was being limited by structural barriers that AI happens to dismantle. Both are real, and both are important to keep in view.

The *Guardian*'s parallel finding — that ADHD, autism and dyslexia workers were reporting AI agents as a major workplace enabler — was echoed in dozens of smaller reports across 2026.⁴⁵² The implication, for the creative industries: a workforce that has historically been heavily neurodivergent (the writing, music, film and games sectors are all over-indexed on neurodivergent talent compared to the general population) stands to be one of the *biggest beneficiaries* of well-deployed AI tooling in the workplace.

This is not a reason to ignore the displacement story. It is a reason to be careful about which framings of the AI transition are accurate and which are reductive.

Indies and the Global South

The other genuinely encouraging signal in this period is the rise of the *indie* and *Global South* creative sectors as direct beneficiaries of the AI cost reduction.

African film and tech has been a recurring positive story across the period — from Korin AI, the “trained with African datasets, built by Africans” model that launched in May 2026,⁴⁵³ to the wave of African AI filmmakers that the trade press began covering in earnest in early 2026, to the African music industry's adoption of AI tools described in *CNBC Africa* in October 2025.⁴⁵⁴

Indian cinema has been awash with AI through the period covered by this book.⁴⁵⁵ The BBC's December 2025 piece “Lights, camera, algorithm” documented the structural shift, with major productions integrating AI for visual effects, dubbing and asset generation, and made an observation about the limits of the technology that has stayed with me: “*You could create a sequel to a regional Indian movie using ChatGPT, but you would need to feed it the cultural memory of the original script. That script would have to be written by a human screenwriter.*” The cultural memory is the human contribution. The toolchain accelerates everything around it. The screenwriting itself — the act of *knowing what the culture remembers* — remains stubbornly, irreducibly human. India's first AI-animated show, *Legenda Bertuah*, launched in Indonesia in April 2026.⁴⁵⁶

Latin American and Middle Eastern AI film festivals proliferated through late 2025 and early 2026. The \$1m Dubai AI Film Award was won by Tunisia's *Lily*.⁴⁵⁷ Mexico's Avocados-From-Mexico Super Bowl campaign was AI-led.⁴⁵⁸

Eastern European AI filmmaking — typified by Andrii Daniels' bomb-shelter clip — became a recognised category.⁴⁵⁹

East Asian AI development continued at a pace that, by spring 2026, had Chinese open-source AI models being used by approximately 80% of startups pitching the Andreessen Horowitz fund.⁴⁶⁰ Korea's Shift Up CEO described AI as the way to compete with Chinese game-industry scale, in language that captures both the geo-economic argument and what it means for individual workers: “*Only when all these people are proficient in AI, so that one person can perform the role of 100 people, can we compete with industries like China and the US that rely on large-scale human resources.*”⁴⁶¹ *One person performing the role of 100.* That is the East Asian games industry's framing of the orchestrator economy, and — if it is

right — it tells you everything you need to know about the headcount maths every studio in the world will run between now and 2030.

The pattern, when you stand back from it, is what I have come to call the **Geographic Inversion**. AI is — in significant measure — *redistributing creative production capacity* away from the traditional centres (Hollywood, London, New York) towards regions that were historically capacity-constrained relative to their creative ambition. For most of the post-war period, geography concentrated creative work; for the first time in living memory, the technology is pushing the other way. This is not, on its own, a justification for everything AI is doing. It is, however, one of the clearest beneficial second-order effects of the cost reduction, and one of the most reliable signs that the creative economy that emerges on the other side of this transition will be — in geographic, demographic and economic terms — *less concentrated* than the one that preceded it.

If you are a working creative in a part of the world that has historically been on the wrong end of the global creative economy's geography of access, the AI era is — for all its risks — also opening doors that were welded shut for most of the previous century.

What organisations should do

I want to close this chapter with a short and direct argument about what creative *organisations* — studios, agencies, labels, broadcasters — should be doing right now, because the readers I hear from most often, after working creatives, are people running creative organisations and trying to figure out the shape of the next three to five years.

The short version, drawn from everything I have read, watched and lived through these six months:

One. Move to integration, fast. The path-two organisations will outcompete the path-one organisations on talent, on output and on cultural capital within three years. Suppression is not viable as a long-term strategy.

Two. Invest in your mid-career layer. The biggest source of avoidable damage in this transition is the loss of mid-career specialist knowledge that takes years to rebuild. Find ways to upskill your existing mid-career staff into orchestrator roles. The institutional knowledge they carry is the most valuable asset you have. Do not throw it away because the labour-cost arithmetic in a single quarter says you can.

Three. Solve the apprenticeship problem. The orchestrator economy structurally undermines the pipeline that has historically produced senior creatives. If you don't solve this — by maintaining some entry-level human roles, by building new AI-augmented apprenticeship pathways, by partnering with the institutes and the literacy initiatives — you are eating your own future. *Your senior creatives of 2035 are the juniors you hire today.* Treat them that way.

Four. Build the disclosure and provenance infrastructure. Chapter 12's argument applies as much to organisations as to individual creatives. The organisations that can credibly disclose their AI use, that maintain documentation, that can produce chain-of-custody on contested work, will be the organisations that the audience trusts in 2030.

Five. Build for the new geography. If your existing organisation is centred on the traditional creative capitals, the AI era is going to be much harder for you than for organisations distributed across the newly-accessible regions of the global creative economy. *Take seriously the option of building distributed teams* — not as a cost-saving move, but as a creative-capacity move. The talent is global. The tools are global. The audience is global. The organisations that don't adapt to this fact will lose their relevance to the ones that do.

Six. Don't outsource your judgement. This is the most important one and the easiest to get wrong. AI tools — even the very good ones, even the agentic ones, even the ones the platform companies are most eager to sell you — *cannot replace organisational judgement*. The decisions about what to make, who to hire, what to invest in, what to refuse — these are decisions that have to live with the humans running the organisation. AI can inform them. AI cannot make them. The organisations I have watched make the biggest unforced errors in the period this book covers are the ones that abdicated organisational judgement to the tools.

The shape of the creative economy in 2030 — what it produces, who it employs, where it operates, what it pays, what it is for — is being decided, right now, by the choices that the working creatives, the organisations and the institutions of the creative industries make in this twelve-to-eighteen-month window.

In the next chapter — the final chapter of the book proper — I want to argue, as directly as I can, for the *kind* of creative economy I think we should be choosing. What a humane version of the AI-era creative economy looks like. Who has to do what to get there. And what working creatives reading this book should be doing on Monday morning to play their part in it.

That choice is the last thing the book is about.

Chapter 14 — The New Jobs

There is a binary that I have, for six months, watched dominate every conversation about AI and creative employment, and that I am going to spend this chapter taking apart.

The binary is *jobs apocalypse* versus *jobs renaissance*. On one side, the visible argument: AI is coming for creative work, the trade unions are right, the displacement is real and accelerating, and a meaningful percentage of working creatives — particularly mid-career specialists in functions where AI has already become competent — will be out of the industry by 2030. On the other side, the equally visible argument: AI is creating more jobs than it destroys, the new categories of AI-orchestration work are paying more than the old ones, the freelance and indie sectors are expanding, the geographic boundaries of the creative economy are dissolving, and the next decade will be the most economically expansive period for creative labour since the post-war television boom.

Both arguments have evidence. Both are partially right. Both are, taken on their own, *wrong* — because the actual labour-market story of 2025–26 is not a binary. It is a *restructuring* with sharp winners and sharp losers, in which the dividing line between the two is not “AI” or “anti-AI.” The dividing line is **AI literacy** — the practical capacity to deploy generative tools as instruments of one’s own creative practice, with judgement, taste and structural understanding of where they help and where they harm.

This chapter is about that restructuring. About which jobs are disappearing, which are emerging, which are simply being *reshaped*, and what the working creative reading this should be doing — concretely, this year — to land on the right side of the line.

I want to spend most of the chapter on the evidence, because the binary framings have been driven, in my experience, by people who have not done the reading. The evidence is messier and more interesting than either side wants to admit.

What the headline numbers do not say

The aggregate employment numbers in the creative industries for 2024–26 did not show the apocalyptic collapse some had predicted. They also did not show the renaissance the platform companies’ marketing teams have been selling. *The Economist*, in a November 2025 piece titled “Investors expect AI use to soar. That’s not happening,” argued that the broad productivity-gain thesis was, in the short term, not playing out at scale across the wider knowledge economy.⁴⁶² A month later, in “Job apocalypse? Humbug! AI is creating brand new occupations,” the same publication argued — using the same labour-market datasets — that AI was producing more new role categories than it was eliminating.⁴⁶³ Both pieces were defensible. Both used real numbers. The two coexisted in the same magazine, six weeks apart, because the aggregate data is, on the current cut, ambiguous.

What the aggregate data hides is the *internal redistribution*. The mid-career specialist roles I described in Chapter 13 — the experienced sub-editors, the junior animators, the staff illustrators, the in-house copywriters, the routine production-pipeline engineers — are visibly contracting. The senior strategic roles — the showrunners, the lead creative directors, the senior orchestrators, the IP-fluent producers — are visibly expanding their *effective reach* if not their headcount. The new role categories — AI agent developers, prompt engineers, AI operations specialists, creative-AI ethics officers, model-curation specialists, AI-literacy

trainers, custom-model fine-tuners, agentic workflow designers — are visibly growing from a near-zero base.

The Guardian’s “AI is hitting UK harder than other big economies, study finds,” from January 2026, found that mid-career UK knowledge workers were experiencing disproportionate displacement compared to peers in the U.S., Japan, Germany and Australia — but that the displacement was concentrated in specific *task categories*, not whole occupations.⁴⁶⁴ The University of Wisconsin-Stout’s January 2026 announcement, in which the institution set AI use as a *baseline competency* in its filmmaking course, captured the supply-side response: the curriculum was being re-engineered around the assumption that working filmmakers in 2030 would be AI-literate by default.⁴⁶⁵

The labour market, in other words, is doing what labour markets always do in a productivity transition. It is reshuffling. The reshuffle is sharper than the headline employment figures suggest because the *composition* of work is changing faster than the *amount*.

The roles that are disappearing

I want to be specific, because vague claims about “creative jobs disappearing” do not help anyone make a career decision. Based on the trade-press coverage tracked across the *Dream Machine* archive, the survey data in Appendix D and Appendix E, and the studio-leadership interviews I have read or conducted, the roles under the most active substitution pressure in 2026 are:

Junior visual production roles. Concept artists at the asset-variation level. Junior 3D modellers doing standard architectural / environmental fills. Storyboard artists working on commercial briefs that do not require performance staging. Background-plate compositors. Routine matte painters. Stock photographers and stock illustrators. Adobe’s own data — 25% of new Adobe Stock contributions in 2024 containing Firefly-generated elements⁴⁶⁶ — is the clearest single number in this category.

Junior writing and copy roles. In-house copywriters at brand agencies. Junior content marketers. Routine technical writers. Translation generalists where the source/target pair is well-resourced (English-Spanish, English-Chinese, etc.). SEO content writers. Sub-editors at digital publications. The Reuters Institute’s “AI adoption by UK journalists” survey found high integration across newsrooms by late 2025; the *Daily Mail*’s December 2025 report that Google’s AI Overviews had “killed click-throughs” to news sites was the consumer-side mirror of the production-side pressure.⁴⁶⁷

Mid-career routine production and post-production roles. Routine audio engineering (the 1,100-creator music survey discussed in Chapter 13 showed 58% of producers using AI for restoration, 38% for mixing assistance, 33.9% for automated mastering⁴⁶⁸). Standard VFX compositing (62% of Hollywood studios on automated AI compositing, 35% reduction in post-production timelines⁴⁶⁹). De-aging specialists (200 hours per actor down to 50). Particle simulation specialists (68% adoption among top VFX houses by SIGGRAPH 2025). Routine matte-painting generalists (initial setup time from 4 hours to 1.2 per shot).

Routine games-development roles. Square Enix’s announced target — **70%** of QA work via AI by end of 2027 — is the cleanest signal here.⁴⁷⁰ Falcom’s reported productivity ratio of *2-3 hours of work reduced to 10 minutes*⁴⁷¹ tells you what is happening to the routine animation, asset and engineering layer underneath the games industry. Ubisoft’s January 2026 cancellation of five games (including the *Prince of Persia* remake) in order to refocus capital on AI signalled a structural shift in resource allocation that mid-career game developers are still digesting.⁴⁷²

In-house AI specialist teams at non-AI-native companies. This one is counter-intuitive but real. The publisher behind *Grand Theft Auto VI* laying off its entire seven-year-old internal AI team in April 2026, in favour of off-the-shelf tooling, is the canonical case.⁴⁷³ Companies that built proprietary AI capabilities in 2018–2024 are increasingly finding that the open-weight and commercial foundation models have caught up; the bespoke AI team becomes redundant. This is a real and rapid form of AI-driven displacement that the public discourse has not yet recognised, because the workers being displaced are themselves AI specialists.

Voice actors and session musicians in commodity work. ElevenLabs’ growth to \$500m ARR by April 2026⁴⁷⁴ is, in large part, the substitution of routine voiceover, audiobook, podcast-host and dubbing work. Live performance, voice work that requires acting craft, and specialty voice roles (animation leads, signature character voices) are not displaced. The middle of the voice market is.

Routine commercial illustration and design. Brand assets, marketing imagery, social-media graphics, basic product visualisation. The Higgsfield growth curve — \$200M revenue in nine months, primarily serving social-media marketers⁴⁷⁵ — is the consumer-marketing equivalent of the Adobe Firefly enterprise curve.

The pattern across these categories is consistent: it is *routine* work, *specialist* work, and *mid-career* work that is under pressure. The pattern is *not* aimed at junior on-ramp roles (which is a problem of its own — see below) or at senior creative judgement roles. It is concentrated in the middle of the pipeline, in the layer historically occupied by experienced operators of the toolchain.

The roles that are emerging

The other side of the redistribution is equally real and substantially under-reported in the consumer press.

AI orchestrators / senior creative directors of agentic teams. The most strategically important new role, and the one Chapter 11 made the long-form case for. The Sony 49-Claude-agent / 72-skill stack is the canonical example.⁴⁷⁶ In adland, *Digiday* reported in late 2025 that “AI agent developers have become adland’s in-demand role”⁴⁷⁷ — a senior creative-strategic role that did not exist eighteen months earlier.

Prompt engineers / AI workflow designers. The role is now broad enough to have its own specialisation tracks. *Forbes* reported in November 2025 that “vibe coding” — natural-

language software development — paid up to **\$220,000** as an in-demand AI skill.⁴⁷⁸ The equivalent figures in the creative space — prompt engineers at major studios, freelance AI-workflow consultants — are running in the same band for senior practitioners.

AI literacy trainers and AI-education designers. The Sundance Institute’s AI Literacy Initiative, launched in January 2026 with \$2M of Google funding to train 100,000 filmmakers, is the institutional version of this role.⁴⁷⁹ The Adobe Ignite Day at Sundance, the UK government’s “Free AI training for all” programme covering 10 million workers by 2030,⁴⁸⁰ the Lovable-for-classrooms expansion,⁴⁸¹ the UW-Stout baseline-AI competency course⁴⁸² — these are the demand signal for a new category of educator that combines creative-discipline expertise with practical AI fluency.

Model curation specialists. With foundation models proliferating and custom fine-tuning becoming a baseline capability, the role of *selecting, training and maintaining* an organisation’s model stack has emerged as a discrete specialism. Adobe Firefly Foundry — the service that lets companies train custom generative models on their own visual identity⁴⁸³ — created an entire job category of brand-and-IP model trainers. Korin AI’s launch in May 2026, “trained with African datasets, built by Africans,”⁴⁸⁴ is the cultural-fluency variant of this role.

AI ethics, disclosure and provenance officers. Following the SAG-AFTRA contract negotiations, the Cannes AI Disclosure Standard, the Academy’s “you must be human to win” rule, the New York AI advertising disclosure law, and the proliferating C2PA-compliance and SynthID-tooling requirements, organisations across the creative industries have begun hiring (or designating) dedicated AI-ethics and disclosure leads.⁴⁸⁵ At DreamLab, we have a *Continuum Lead* whose job is to make this work coherent across every project we run — three years ago, the role did not exist.

Indie and Global South creator-producers. The cost reduction in production tooling has created a new viable role category that was not economically possible before: the *one-person-or-small-team creator-producer* operating outside the traditional creative centres, with global distribution reach and a defensible aesthetic identity. *Forbes* covered the broad category in “AI Is Changing How Creators Work And Earn” in December 2025.⁴⁸⁶ The Higgsfield revenue (built on social-media marketer demand), the Andrii Daniels bomb-shelter clip (a Ukrainian one-person production with global reach⁴⁸⁷), the Tunisian *Lily* (\$1M Dubai AI Film Award winner⁴⁸⁸), the Indonesian *Legenda Bertuah* animated series,⁴⁸⁹ the Indian-cinema integration wave covered by the BBC’s “Lights, camera, algorithm”⁴⁹⁰ — these are not exceptions. They are the leading edge of a structural change in *who can be a working creative*, and where they can live.

AI-augmented apprentices. This category is still being built, and is the central labour-market design question of the next three years (more below). The early models — AI-tool-augmented junior animator roles maintained deliberately at *Position Four* studios (Chapter 7), the Sundance Collab fellowship structure, the AI-augmented entry-level posts at WPP and

the major Hollywood studios — are the early experiments. None of them, yet, has fully solved the apprenticeship problem.

Cross-disciplinary “portfolio creatives.” What I have been calling, in talks since the autumn, the **AI Literacy Premium** role — the working creative who, instead of a single specialism, holds several loosely-coupled creative disciplines together using AI as connective tissue. *TechBullion*’s “Why the future belongs to multi-skilled leaders,” from November 2025,⁴⁹¹ and the Anthropic *Skills* framework underneath Claude Code’s multi-agent coordination,⁴⁹² are the corporate-leadership and tooling-side manifestations of the same trend. The portfolio creative is increasingly the *default* career shape for working creatives entering the field today.

The AI literacy premium

If you stack the disappearing-roles list against the emerging-roles list, a single underlying variable does most of the explanatory work for who is winning and who is losing in the 2026 creative labour market.

The variable is **AI literacy** — defined operationally as the combined skill of (a) knowing what generative tools can and cannot do well, (b) being able to brief them effectively, (c) being able to judge their outputs with taste, (d) being able to integrate them into a coherent creative workflow, and (e) knowing where on the Continuum your craft sits and why.

The empirical case for the literacy premium is strong and consistent across the data:

The Adobe Creators’ Toolkit Report of October 2025, surveying 16,000 creators across the U.S., U.K., France, Germany, South Korea, Japan, India and Australia, found that **86%** of creators were already using generative AI; that **76%** of users said it had helped grow their business or brand; that **81%** said AI lets them make content they otherwise could not have made; and — most significantly — that **70%** were optimistic about agentic AI and **85%** would use AI that learned their creative style.⁴⁹³ The numbers are not the numbers of a workforce in retreat. They are the numbers of a workforce *adopting at scale* and reporting commercial benefit.

The PRS for Music 2026 AI Survey, surveying U.K. music creators, found **four in five** music creators worried about AI-generated music competing in the streaming economy⁴⁹⁴ — but also that respondents who were using AI tools in their own creative practice reported higher earnings and broader output reach than peers who were not. The two findings live together. The anxiety and the adoption are happening in the same head.

The U.K. Department for Business and Trade research, reported in *CNBC* in November 2025, found that workers with ADHD, autism and dyslexia were **25% more satisfied** with AI assistants than neurotypical workers, with AI agents reported as actively helping them succeed at work.⁴⁹⁵ Given that the creative industries are over-indexed on neurodivergent talent compared with the general population, this is a major equity-and-productivity story that the AI-displacement narrative has under-told.

The McKinsey AI report on film and TV production, released in early 2026, made the broader business case: AI would not, in McKinsey’s view, replace film and television production. It

would *restructure* it — towards smaller teams, faster cycles, more iteration, and heavier reliance on senior creative judgement.⁴⁹⁶ In other words: towards an *orchestrator-shaped* industry, in which the working creatives best positioned to thrive are those who can operate the new toolchain rather than those who refuse to engage with it.

The GDC State of the Game Industry surveys 2024–2026 (covered in detail in [Appendix E](#)) showed personal generative-AI usage among professional game developers rising from **31%** to **36%** while industry sentiment cratered from 18% negative to **52%** negative.⁴⁹⁷ The interpretation: developers using AI are increasingly *uncomfortable* about doing so, but they are doing so anyway, because the productivity advantage is too strong to ignore.

The LANDR study reported in *Ari's Take*, late 2025, found that **87%** of musicians surveyed were using AI tools in some part of their practice.⁴⁹⁸ When the figure for “musicians using AI” approaches the figure for “musicians using DAWs,” you are not looking at a niche adoption pattern. You are looking at a baseline competency that the next generation of working musicians will be assumed to have.

The Stanford AI Index Report 2025 found that across 26 surveyed nations, **55%** of individuals viewed AI products as offering more benefits than drawbacks, up from 52% in 2022⁴⁹⁹; the YouGov 2024 multi-market sentiment data echoed the same direction.⁵⁰⁰ Working creatives are sitting inside a consumer environment that is, on net, *increasingly accepting* of AI in creative work — which translates, slowly, into client demand for AI-augmented services.

The *Digital Music News* survey of nearly **800 creatives** who signed the *Stealing Our Work Is Not Innovation* declaration in January 2026⁵⁰¹ — that same vocal-minority signal that produced political wins in Chapter 6 — included, in its signatory base, a meaningful number of creators who were *themselves* using AI tools in their own production work. The signing of the declaration was not a refusal of the technology. It was an insistence that the *training data* be consented and compensated. Two positions, held at once, by the same working creatives.

Stack these findings together and the labour-market picture becomes much clearer than the binary discourse allows.

Working creatives who hold AI literacy are gaining leverage. Their effective output reach is larger. Their hourly rate is higher. Their client base is broader. Their geographic constraints are weaker. The orchestrator role, the portfolio-creative model, the literacy-trained newcomer — these are the positions on which the next decade of creative employment expansion is being built.

Working creatives without AI literacy are losing leverage. Their effective hourly rate is being competed down by AI-augmented peers. Their client base is being eroded by faster, cheaper, mostly-adequate substitutes. Their career mobility is contracting as mid-career specialist roles in their domains are absorbed into agentic capacity. Their geographic constraints (London, New York, LA, Tokyo as the centres) are being preserved while their AI-fluent competitors operate from anywhere.

This is, in practical terms, the structural transition the working creative reading this in 2026 is now facing. The choice is not “use AI” or “refuse AI.” The choice is whether to build the literacy that lets you remain in command of how AI shows up in your own practice, or whether to let the literacy build itself, badly, in your absence.

The Apprenticeship Gap

I want to name the most underdiscussed structural problem in the AI-era labour market, because it is — by my read — the single largest threat to the long-term health of the creative economy and it has nowhere near the public attention it deserves.

I call it the **Apprenticeship Gap**.

For the entire history of the creative industries — from the medieval guilds through the post-war Hollywood studio system through the rise of digital media — the *junior on-ramp* into creative work has been the structural foundation on which senior talent is built. Junior writers become senior writers by writing things that nobody pays much attention to, repeatedly, for a decade, under the loose mentorship of more senior practitioners. Junior animators become senior animators by drawing the in-between frames, by cleaning up the rough animatics, by handling the routine asset variations that the lead artists do not have the time for. Junior cinematographers become senior cinematographers by holding focus, by pulling cable, by lighting the second-unit shot. The junior tasks were not, in themselves, the destination. They were the *training ground* on which judgement, craft and taste were built.

The orchestrator economy, the agentic toolchain, the GDC-data picture of senior practitioners using AI to absorb the routine middle-layer work — these patterns, taken together, are progressively *removing* the junior on-ramp from the industry. The junior writer is competing with Claude. The junior animator is competing with Cascadeur. The junior cinematographer is competing with Veo 3.1’s plate generation. The junior coder is competing with Cursor and Copilot. In every case, the routine task that *used to be the entry point* into the discipline is now economically uncompetitive against an AI agent that does it in seconds for cents.

The studios are not — yet — replacing senior roles. They are absorbing the *junior layer* underneath the senior roles, and then telling themselves a story about how the new tools will free senior practitioners to focus on the *real* creative work.

The story is partially true. It is also dangerously incomplete. If the junior layer disappears for a decade, the *next generation of senior practitioners has nowhere to be trained*. The pipeline breaks. The 2035 cohort of senior creative directors, lead animators, showrunners, music producers, art directors — the people who, in 2026, would be five years into a junior career — will simply not exist at the volumes the industry needs. The discipline-specific knowledge, the embodied craft, the relationship-based mentorship — all of these were carried in the apprenticeship layer. Remove the layer, and you are eating the seed corn of the discipline.

This is not a speculative claim. It is the underlying logic of every “expensive mistake” interview a working studio leader has given to the trade press in this period — Charles Cecil at Revolution Software, Todd Howard at Bethesda, the Larian and Aardman and Jagex public positions. The senior practitioners are saying, in different vocabularies, the same thing: *we*

used to teach the next generation by giving them the routine work. The routine work is gone. We have not solved the teaching problem.

The institutional response to date is partial and patchy. The Sundance literacy initiative is real. The Adobe Ignite Day, the Sundance Collab fellowships, the UK Free AI Training for All programme, the UW-Stout curricular changes — these are the visible institutional moves. But they are training programmes for *AI literacy specifically*, not full apprenticeship pipelines for the underlying creative discipline. They produce literate orchestrators. They do not, by themselves, produce the cinematographers, the composers, the writers, the lead artists, the showrunners of 2040.

The deepest structural reform that needs to happen in the next eighteen months is the deliberate preservation — or rebuilding, or re-imagining — of the apprenticeship layer. Some of this is already happening:

Position Four studios maintaining junior roles by policy. Aardman, Larian, Games Workshop, Jagex have, in different forms, made deliberate commitments to keep junior human roles in their pipelines even where AI could absorb them, on the explicit grounds that the future of the studio requires it. This is the most encouraging single trend I have observed.

Hybrid apprenticeship pathways. The new role I called *AI-augmented apprentices* above. A junior animator who uses Cascadeur, but who is *paired with a senior animator who teaches them why* the AI's output is good or bad. The juniors are not doing the in-between frames any more. They are doing the *judgement* on the in-between frames the AI produces, and the senior teaches them how to judge. This is a real model. It is not yet at scale.

Institutional reinvestment. The cultural-institution training programmes — Sundance, the BFI, the national film schools, the BBC training schemes, the Royal College of Art and equivalent national schools — are, in different ways, recalibrating to deliver more comprehensive training in shorter timeframes, on the assumption that the years-on-the-job apprenticeship period is contracting.

Public funding interventions. The UK government's *Free AI training for all* programme, the EU's various creator-skills initiatives, and the U.S. state-level training credits being attached to AI-investment incentives are early signs that the apprenticeship gap is being recognised as a public-policy problem rather than a market problem.

None of these is, on its own, sufficient. The Apprenticeship Gap is — by my prediction — going to be the single largest *unresolved* labour-market issue in the creative economy of 2030. The studios, the unions, the schools and the platform companies are going to have to figure it out together. The book's Chapter 15 manifesto in *Choosing the Future* lists it explicitly as one of the questions every working organisation has to engage with.

If you are reading this as a senior practitioner: maintain juniors. Pair them with the new tools deliberately. Treat their on-ramp as a public good your industry depends on.

If you are reading this as a junior practitioner: the on-ramp is contracting. You do not get to wait. The literacy you build *in the next eighteen months* will determine whether the on-ramp closes before you are inside it.

The geographic redistribution

The second-largest under-reported labour-market story of this period is geographic.

For the entire post-war history of the global creative industries, professional creative employment has been concentrated in a small number of cities — Los Angeles, New York, London, Paris, Tokyo, Mumbai, with smaller secondary nodes in Berlin, Toronto, Seoul, Sydney. The geography was a function of the *cost structure* of creative production: studios, equipment, distribution networks, talent pools and capital all concentrated in the cities that could afford to host them, and the working creatives followed.

The AI cost reduction is, structurally, dismantling this geography.

The Tunisian-made *Lily* winning the \$1M Dubai AI Film Award.⁵⁰² The Ukrainian one-person bomb-shelter production going viral globally.⁵⁰³ The Indian-cinema AI integration wave covered by the BBC, with productions across regional cinema centres absorbing the new toolchain at scale.⁵⁰⁴ The Indonesian *Legenda Bertuah* AI-animated series.⁵⁰⁵ The Korin AI launch — Africa-trained, Africa-built foundation model — and the broader CNBC Africa coverage of AI in African music and film.⁵⁰⁶ The Singapore-based AI video startup Video Rebirth raising \$50M for studio-grade tooling. The Eastern European AI filmmaker community building around the success of creators like Andrii Daniels. The Latin American AI-film festival wave through 2026.

These are not isolated stories. They are the leading edge of a redistribution that, taken together, is *materially expanding* the global creative workforce beyond the previous-century cities. The reduction in cost-of-entry for serious creative production has, for the first time since the rise of cinema, made it economically viable to build a competitive creative practice from places that the previous geography had locked out.

The Shift Up CEO's framing, in the *PocketGamer.biz* piece — “*Only when all these people are proficient in AI, so that one person can perform the role of 100 people, can we compete with industries like China and the US that rely on large-scale human resources*”⁵⁰⁷ — is the strategic reading from inside the Korean games industry. The same logic applies in Mexico, in Egypt, in Nigeria, in Brazil, in Vietnam, in any creative economy that has historically been resource-constrained relative to its ambitions.

For the working creative reading this in one of those geographies: the AI era is, for all its disruption risks, also opening doors that were welded shut for most of the previous century. The labour market is becoming, for the first time in living memory, *less* concentrated.

For the working creative reading this in one of the historical centres: this redistribution is not theoretical. The clients, the budgets and the IP that previously concentrated in your city are, increasingly, being competed for by capable AI-augmented competitors operating from anywhere on the planet. Your geographic advantage is contracting.

What the working creative should do

I have, throughout the book, tried to land each chapter with practical takeaways. This chapter is the labour-market chapter, and the takeaways are the most concrete.

Build the literacy this year. Not next year. The Adobe Creators' Toolkit Report, the LANDR survey, the GDC data, the McKinsey reading — all point in the same direction. The literacy premium is not a future variable. It is, in 2026, the single biggest determinant of mid-career creative employment outcomes. Free training programmes exist (Sundance Collab, UK Free AI Training, Adobe Express, the open-source ecosystem). Use them. Spend the equivalent of one week per quarter, deliberately, on building practical fluency in the toolchain.

Map your craft against the Continuum (Chapter 3). Decide where your craft sits, function by function. Decide where you are willing to let agents operate on your behalf and where you are not. Write the map down. Update it quarterly. The working creatives I have watched make the most successful transitions in this period have been the ones who knew, in advance and explicitly, where their lines were.

Pick your role on the new map. Orchestrator, portfolio creative, AI-literacy trainer, model curator, ethics-and-disclosure specialist, regional creator-producer, hybrid apprentice. The roles are real. They pay. They will, by every indicator I can read, continue to grow through the next five years. You do not need to invent your own category. You can pick one of the emerging ones.

Build your apprenticeship — or build the next generation's. If you are early in your career, find the senior practitioners who are running hybrid apprenticeship pipelines and apply. If you are mid-career, identify the AI-augmented junior roles in your discipline and either fill them yourself or pair with one. If you are senior, maintain juniors in your team and pair them deliberately with the new tools.

Take the geography seriously. If you have historically been outside the creative-economy centres, the AI cost reduction has opened a window. Use it. Build for the global creative market from where you are, with the cost advantages your geography offers. If you have historically been inside the centres, your geographic premium is contracting; build defensible craft, IP and relationships that survive the cost flattening.

Stay in the work. This is the same advice the rest of the book lands on, and it is the same advice for this chapter. The maker who never makes is the maker whose judgement decays. Maintain craft contact, in at least one part of your practice, that does not depend on AI tooling. The contact will keep your eye sharp for everything else. The orchestrator who never operates the tools cannot brief them well. The orchestrator who only operates the tools loses the human signal the audience came for.

Speak. The labour-market shape of the next decade is being decided right now, by the institutions of collective bargaining, by the policy-makers running consultations, by the platform companies building the rails, and by the creative organisations making integration choices. The 88% in the UK consultation was made out of voices. The Tilly Tax was made out of voices. The Sundance literacy turn was made out of voices. Your voice — your testimony to your union, your trade body, your local government, your manager, your client — is the input the institutions need. The labour-market protections of 2030 will be the cumulative result of how loudly working creatives turn up to claim them in 2026 and 2027.

The doomer frame, and why it loses to the data

A separate but related framing has, since the early spring of 2026, occupied roughly the same rhetorical ground as the *jobs apocalypse* position but with rather more political weight behind it. Call it the **doomer frame**: the argument, articulated through OpenAI's policy paper, Dario Amodei's interviews, and Daniel Susskind's *A World Without Work*, that AI represents a categorical break from prior automation, that productivity gains will accrue only to capital, that the labour market is showing early signs of *structural collapse* rather than ordinary sectoral rotation, and that a *new social contract* — sovereign wealth funds, robot taxes, universal basic income, a right to AI access — is the required institutional response to an inevitable surplus of unemployable humans.⁵⁰⁸

This is, in my view, the most consequential rhetorical position currently on the table, because it is the position from which the policy architecture of the next decade is being constructed. The OpenAI proposals I described in Chapter 13, the various national wealth-fund and four-day-week initiatives now under discussion, the *social-contract-for-the-intelligence-age* framing that has migrated from think-tank papers into mainstream policy briefings — all of it rests on the four load-bearing claims the doomer frame is making about the labour data. If those claims are wrong, the policy architecture being built on top of them is wrong.

The empirical case against the doomer frame is laid out in full in Appendix I: The Doomer Mistake, and I will not reproduce it here at length. The short version is that each of the four load-bearing claims is, at best, unsupported, and several are flatly contradicted by the same Stanford and Brookings data the doomer camp routinely cites. The Stanford Digital Economy Lab's *Canaries in the Coal Mine?* paper — the 13% relative employment decline among 22–25-year-olds in AI-exposed occupations — is real, painful, and worth taking seriously as an *apprenticeship-pipeline* problem (the precise problem this chapter's *Apprenticeship Gap* section is built around). What it is not is a civilizational signal. Overall employment in AI-exposed occupations has grown robustly; older workers in the same occupations have seen employment hold steady or grow; aggregate unemployment in AI-exposed roles is rising more *slowly* than overall unemployment. Brynjolfsson, Li and Raymond's *Generative AI at Work* — the most-cited worker-level study in this period — finds the *opposite* of the doomer prediction: AI delivered a +34% productivity gain to novice / low-skill workers and a +14% average gain across 5,179 workers, while *compressing* the skill distribution rather than amplifying it. Acemoglu's *Simple Macroeconomics of AI* — the most rigorous bear case in the academic literature — bounds the total-factor-productivity gain at 0.66% over ten years, about 0.06% per year. The 230-year track record of identical predictions, from Elizabeth I's refusal of the knitting-frame patent in 1589 through Hinton's "stop training radiologists now" in 2016, is unambiguous: every generation produces a high-status intellectual movement convinced that *this time* the historical pattern will break, and every generation has been wrong.

The risk of accepting the doomer frame is not that we accidentally help displaced workers. The risk is that we build the policy architecture for a crisis that does not arrive, and in doing so create exactly the labour-market rigidities that prevent the ordinary process of adjustment from working. The countries that handled the textile transition badly were the ones that resisted it. The countries that handled the agricultural transition badly were the ones that

subsidised it indefinitely. There is a real chance that the most damaging response to AI is not the technology itself but the panic-driven policy response to it.

For working creatives reading this, the implication is direct. The doomer frame is doing political and commercial work for the people promoting it. *It does not, on present evidence, describe reality.* The labour market is reshuffling, not collapsing. The mid-career squeeze is real and worth addressing. The apprenticeship pipeline is a real problem and requires deliberate institutional response. The *end-of-work event* is not happening on any timeline visible in the actual labour data, and acting as though it is — by withdrawing from the literacy investment, by treating your career as a sunk cost, by accepting the framing of inevitability that the AI labs' policy departments are pushing — is the strategic error that does the most damage to the working creative.

A note on the binary

I want to close this chapter by returning to the binary I started with — *jobs apocalypse* versus *jobs renaissance* — because both framings, repeated often enough, do real damage to the working creative who is trying to make rational career decisions in real time.

The apocalyptic framing is, in my view, the more dangerous of the two, because it produces *paralysis*. Working creatives who become convinced that AI is coming for their job, full stop, often stop investing in the literacy that would let them stay ahead of the substitution curve. They become spectators of their own displacement. By the time the substitution arrives, they are unprepared.

The renaissance framing is the less dangerous one but is also wrong, because it produces *complacency*. Working creatives who become convinced that AI will simply expand the labour market often fail to invest deliberately in the literacy that lets them claim the expansion. They drift, expecting the rising tide to lift their boat without recognising that it lifts only the boats whose owners are paying attention.

The accurate framing is harder to live with than either: the labour market is *redistributing* in real time, with sharp winners and sharp losers, and the variable that most reliably predicts which side you land on is the *deliberate, structured investment in AI literacy* that you make over the next twelve to eighteen months.

The good news, against the dystopian end of the press cycle, is that the literacy is acquirable. It is not gatekept by class, by geography, by previous credentialing, or by institutional access. The tools are largely free at the entry level. The training programmes are largely free. The community of practice — the *Dream Machine* readers, the DreamLab Collective, the open-source forums, the regional creator networks, the literacy initiatives at Sundance and elsewhere — is open and inviting. The barrier to entry is, by historical creative-industry standards, low. The empirical case that AI is, in practice, *lowering* barriers to entry for practitioners who previously lacked institutional access — and the evidence on who benefits most from AI augmentation — is documented in Appendix F: AI, Stigma, Privilege and the Democratisation of Creative Expression.

The harder news is that the literacy *has to be acquired deliberately*. It does not arrive by osmosis. It does not arrive by reading the trade press. It arrives by *doing the work* — by

sitting at the desk, briefing the agents, evaluating their outputs, building the workflow, and shipping the result. Then doing it again. Then doing it again. The literacy is a *practice*, not a credential. It is, in that sense, identical in shape to every other craft skill the creative industries have ever rewarded.

The working creative in 2026 who treats AI literacy as a craft to be practised, not a technology to be debated, is the working creative who, in 2030, will look back at this period as the one in which their career took its most consequential turn.

Build the literacy. Pick your role on the new map. Defend the apprenticeship. Take the geography seriously. Stay in the work. Speak.

The labour market is moving. So can you.

Chapter 15 — Choosing the Future

I want to start this chapter with a confession.

When I sent out the first edition of *Dream Machine* on 6 October 2025, I thought I was writing a newsletter about *tools*. I thought it was going to be a weekly digest of new model releases, new app launches, new research papers — a useful reading list for people in the creative industries who wanted to keep up with what was, even then, an absurd pace of technical change.

Six months later, I do not think that is what *Dream Machine* has been about, and I do not think it is what this book is about either.

The tools have mattered. The tools will continue to matter. Each chapter of this book has had to spend significant pages on what the platforms shipped, when, and to whom — because the platforms are setting the rails, and you cannot understand the choices the creative industries are now making without understanding the constraints those rails impose.

But the *book* has been about something else. About a question that the tools force every working creative, every studio, every union, every government, every audience member, and — eventually — every person who consumes culture, to answer.

The question is: *what kind of creative economy do we want on the other side of this?*

That is what this chapter is for. To put the question on the table, in its sharpest form. To describe what a humane answer to it looks like. And to tell you, as directly as I can, what I think we should each do on Monday morning to start producing that answer rather than the alternative. The two chapters that follow this one — *The Tools*, a categorised inventory of the toolchain, and the *Epilogue*, a letter to the creative person reading this in 2030 — are the practical reference and the closing register. The argument the book has been building toward lives here.

The choice on the table

The choice, stated plainly, is between two creative economies.

One is **extractive**. In this economy, the creative work of millions of human authors — their writing, their music, their images, their voices, their styles, their cultural specificity — is absorbed, without consent or compensation, into large statistical models owned by a small number of platform companies. The platforms then sell access to those models, mostly to the same brands, studios, and agencies that previously paid for original human work, at a fraction of the original cost. The aggregate result is a *transfer of wealth* from the diffuse pool of working creatives to a concentrated pool of platform shareholders. The creative output of the economy continues — perhaps even increases in volume — but its *meaning*, its *cultural specificity*, and its *connection to the lives of the people who used to make it*, decays over time.

The other is **generative**, in the original sense of the word. In this economy, AI is treated as *new craft infrastructure* — a set of tools that, like the printing press, the camera, the synthesiser and the digital editing suite, can be used by working creatives to make new kinds of work that the previous infrastructure didn't allow. The training data is consented to, attributed to, and compensated for. The platforms compete on the quality and integrity of their tools, not on the unpriced absorption of their users' work. The audience can verify the

provenance of what they encounter, and pay attention accordingly. Working creatives are *augmented* by the new tools, not *replaced* by them. The output of the economy is bigger, more diverse, more accessible, more globally distributed — and recognisably continuous with the human creative traditions it builds on.

These two economies are not, on the current trajectory, equally probable. Some of the rails being laid right now point at the first. Some at the second. The choice between them is not — as I argued in Chapter 12 — a technical question. The technology to support either is, in spring 2026, broadly available.

The choice is *political, institutional and cultural*. It is about who gets to make the rules, who has the standing to enforce them, and what default the millions of small daily decisions that constitute a creative economy converge on.

The good news, if you have read this far, is that more of those daily decisions than you might expect are pointing at the generative economy. The audience's behaviour around the slop ceiling. The UK consultation's 88%. The Sundance Literacy turn. The Cannes Disclosure Standard. The Academy's *you must be human to win*. The SAG-AFTRA Tilly Tax. The studio refusals — Jagex, Larian, Aardman, Games Workshop. The disclosure norms emerging in advertising and the platforms. The 800-creator declaration. The C2PA standards. The SynthID rollout. The free-AI-training-for-all programmes. The neurodiversity dividend. The Global South opening up.

These are real. They are not a foregone conclusion. But they are real, and they are coming from a coalition of forces — creators, audiences, unions, governments, institutions and some of the platform companies — that has, six months in, more political and economic weight than the extractive trajectory's proponents are willing to acknowledge.

What I want to argue, in the rest of this chapter, is that *the generative economy is winnable*. It is winnable in the next eighteen months. But it requires the working creatives, the organisations, and the institutions that have been doing the work this book has documented to *keep doing it*, with deliberate intent, against the gravitational pull of the extractive alternative.

What I actually believe

Before I get to the principles, I want to put a piece of conviction on the page that has been implicit in everything else I have written but that I think deserves to be made plain.

I believe in AI as an **assistive tool that amplifies human creativity**. Not as a replacement for it. Not as a substitute for it. As an instrument that, properly deployed, expands what a working creative can imagine, attempt and finish in a given afternoon — without displacing the human imagination, the human attempt, or the human finish.

This is the operating assumption underneath every chapter of this book, and I have been deliberately cautious about stating it as my own view, because the whole point of the newsletter and the book has been to track the *evidence*, not the *enthusiasm*. The evidence has been mixed. The cultural reaction has been mixed. The economics have been mixed. None of

those mixed signals would have served the reader if I had collapsed them, every week, into the version I personally find most hopeful.

But the evidence has come in. The slop ceiling holds. The audience can tell the difference between work made with care and work made by a content farm. The toolchain rewards taste, judgement and intent more than it rewards raw computational throughput. The economic returns of agentic AI accrue, demonstrably, to the people who already know what good output looks like — not to a category of new “AI-only” creators replacing human ones, but to a retooled population of working creatives whose effective reach has grown.

What is happening in 2026, on the evidence I have spent six months collecting, is not that machines are taking over creative work. What is happening is that the *labour of execution* is being democratised, and the *labour of intent* is being foregrounded. The *how* is becoming a utility. The *why* is becoming the scarce good.

The deep-dive companion piece *The Age of Intent*, which sits in the appendices to this book, makes the long-form case for this inversion — the philosophical and economic argument that, when the technical barrier to production collapses, the value of *deciding what should be produced and why* rises in direct proportion. The artist of 2030, in that piece’s framing, is less a *manual laborer* and more an *Architect of Meaning* — a curator, a noticer, a setter-of-intent — and the friction of human vulnerability is the irreplaceable component of the work that the machine cannot, by construction, supply.

I think that framing is essentially correct. I think the working creatives who emerge from this transition with the most leverage will be the ones who took the AI tools seriously as *instruments of their own intent*, and refused either to surrender that intent to the toolchain or to refuse the toolchain on principle. The maker-as-craftsperson and the maker-as-prompt-monkey are both poor models of where the work is heading. The maker-as-orchestrator — the architect, the editor-in-chief, the director of a hybrid human-agent team whose unifying signal is *taste* — is the model that the next ten years rewards.

This is not a small reframing. It is the reframing the rest of this chapter — and the rest of this book — has been building towards. The four principles I am about to lay out are not principles for *restraining* AI. They are principles for *deploying AI in service of the human creativity it is supposed to amplify*. The choice between the extractive and the generative economy is, in the end, a choice about which of those two we treat as the master and which we treat as the servant.

I think the human creativity is the master. I think the AI is the servant. I think the creative economy that emerges on the other side of this transition will be the one that does not lose sight of that hierarchy.

The age of the Why

I want to give this conviction a name, because the name will travel where the long-form argument cannot, and because a name makes the choice easier to hold in the head when the next platform launch tries to talk you out of it.

We are leaving the *age of the How*. We are entering the *age of the Why*.

By the *age of the How*, I mean the long century in which the central question of working creative life was *can you do the thing?* Can you draw the figure, light the scene, edit the cut, mix the record, model the asset, hold the camera steady, hit the note? The training pipelines of every creative industry the *Dream Machine* newsletter has tracked were, at their core, infrastructure for answering that question. Conservatoires for the note. Film schools for the cut. Apprenticeships for the figure. Studios for the scene. Whole career structures built around the demonstrable, transferrable ability to *execute* — to convert intent into finished work using the technical labour of one's own hands and ears and eyes.

The *How* was the bottleneck. The *How* was the scarce thing. The *How* was what you got paid for.

The *How*, in 2026, is — at a rate that I do not think the industry has yet fully metabolised — becoming a utility. The teenager in the bedroom with a midrange GPU and a Claude subscription can, at the time I am writing this, produce work whose surface qualities — composition, lighting, sound design, edit pacing, visual-effects polish — sit on a continuum with what a full studio could produce in 2020. The continuum is not yet at the very top end; the bedroom does not yet make *Avatar*. But it is, on every available metric, closing the surface-quality gap at a rate that makes the next-decade trajectory unambiguous. *Hollywood-level execution in a bedroom* is no longer a marketing slogan. It is, for working filmmakers and musicians I know, already true for non-trivial fractions of the work.

When the *How* becomes a utility, the *Why* becomes the scarce good.

I want to anchor this in a specific story that captures the dynamic better than any creative-industries example I have. In March 2026, *Bloomberg* ran a piece on what artificial intelligence has done to elite chess.⁵⁰⁹ AI, the article reported, has driven the game towards *perfect play* at the very top — Stockfish, Leela Chess Zero and their descendants have, between them, mapped out the optimal response to most board positions a top-twenty grandmaster is likely to encounter. The visible effect of this, through 2024 and 2025, was a striking rise in the *draw rate* at top tournaments. When both players have memorised the machine-optimal lines, both players play optimally, and both players draw. The game, at the very top end, was being *solved into stasis* by the machines that had been trained on it.

The *Bloomberg* piece reported the response. Top grandmasters — Magnus Carlsen among them — had started, deliberately, to play *sub-optimal* moves. Moves that Stockfish would mark as inaccuracies. Moves that the machine-optimal line would not recommend. Moves chosen, specifically, because they were *unexpected*, and because the opponent — having trained against the machine — had not memorised the human-grade response to them. The grandmasters had stopped trying to out-machine the machine. They had started, instead, to *deliberately diverge from machine-optimal play* in ways that put their opponents on uncomfortable, uncomputed ground.

The grandmasters are winning, in 2026, by *doing the unexpected*.

I want to dwell on this image, because I think it is the cleanest available picture of what working creative life looks like on the other side of the *How* becoming a utility. When the machine has solved the optimal move — the perfectly-lit shot, the on-trend hook, the

algorithmically-tested ad treatment, the mean-of-the-distribution image — the human edge is no longer in *playing the optimal move better than the machine*. The machine plays the optimal move infinitely. The human edge is in playing the move *the machine would not make*. The deliberately unexpected. The taste-driven. The risk-taking. The idiosyncratic. The personal. The locally-meaningful. The unrepeatable.

This is the *Why* in operational form. The *Why* is not, in the practitioner's day, an abstract philosophical commitment to human creativity. It is a *daily competitive practice* of choosing the move the machine would not make. The director who picks the actor the casting algorithm would not have picked. The songwriter who keeps the verse the chart-pattern model would have cut. The illustrator who renders the figure in a style no FLUX prompt would generate. The brand creative who builds the campaign around the audience question the marketing-AI did not surface, because no marketing-AI in 2026 has the lived context to surface it.

The deeply human things — *taste, intent, authenticity, the willingness to take a risk on a move the data does not yet endorse, the refusal to make the average thing in service of the meaningful thing* — are not, in the age of the *Why*, vestigial commitments held against the toolchain. They are, increasingly, the only things the toolchain cannot do, and therefore, by simple economic substitution, the things that have *commercial leverage* in a market where everything else has been pushed towards zero marginal cost.

The slop ceiling in Chapter 5 is this dynamic measured at the audience layer: the audience, presented with the machine-optimal flood, reliably underweights it, because the audience can tell — at the speed of a swipe — that there is no human *Why* underneath the work. The orchestrator role in Chapter 11 is this dynamic operationalised at the working-creative layer: the orchestrator's daily craft is, exactly, the *judgement to make the un-machine-like move at the right moment*. The authenticity premium in Chapter 12 is this dynamic priced at the commercial layer: the audience pays extra, demonstrably, for work whose human *Why* is verifiable. The four principles I am about to lay out are this dynamic codified at the policy and platform layer.

The chess grandmasters did not stop using engines. They train on engines daily. They study machine-optimal play more closely than any generation of players before them. The chess engines, far from being their enemy, are their *most-used analytical tool*. What the grandmasters refused to do is to *let the engines define what counted as a winning move*. They use the engine to know what the optimal line is, and then they choose, *for taste and surprise reasons*, to play another line.

This is also, I should note in passing, the *unflattering diagnosis* of the legacy entertainment industries' strategic position that I made in Chapter 7. Hollywood, commercial music and the AAA games business spent the past fifteen years optimising themselves *toward* the engine-optimal move — toward the franchise-installment, the streaming-tested chart hit, the open-world template — and arrived at the AI moment producing exactly the work the engines can now replicate most cheaply. The grandmasters' response to engines is the response legacy needs to make to AI. The new AI-native studios — Gossip Goblin, Critterz, Imaginae, Asteria, Wonder — have, by virtue of their newness, no calcified rules to unlearn and are, by default, playing the moves no machine would generate. The legacy industries that survive will be the

ones that re-learn how to make the un-machine-like move. The ones that don't will, on the historical pattern of [Chapter 2](#), be remembered as the cohort that defended the previous definition while a different cohort, with no inherited risk-aversion, defined the next one.

That is the working operating model I think the next decade of creative work runs on. Use the engines. Learn the engines. Train against the engines. And then, in the actual encounter with the audience, *make the move the engine would not have made*.

The age of the Why is not the age of refusing AI. It is the age of mastering AI sufficiently that the only competitive question left is whether you can summon the *deliberately unexpected, deliberately human, deliberately yours* move that the machine, by construction, cannot.

Four principles

I have been asked, by readers of the newsletter and by people who turn up to the talks I have been giving since the autumn, what a humane AI-era creative economy should look like. After six months of trying to answer that question, I have come down to four principles. They are not a programme. They are a *test* you can apply to any policy, any contract, any product launch, any organisational decision — and ask whether it is moving us towards the generative economy or the extractive one.

One. Agency. Every working creative should retain meaningful control over how AI is used in relation to their own work — both in *what gets trained on their work* and in *whether and how they choose to use AI in their own practice*. The Human–AI Agency Continuum from Chapter 3 is the practical expression of this. Policy, contracts, platform terms-of-service and union agreements should all be evaluated against whether they preserve this control or undermine it.

The 88% of UK respondents who said training should require licensing in all cases were articulating this principle.

Two. Attribution. When AI systems produce work that is derived, in any meaningful sense, from human-authored training data, the human authors should be identified and — where appropriate — compensated. The technical infrastructure for this — *creative weight attribution* in Musical AI's framing,⁵¹⁰ C2PA provenance metadata, SynthID watermarks — is the most important infrastructure question of the next two years. Policy should support its deployment. Contracts should require it. Audiences should expect it. Platforms that refuse to engage with it should be treated, in the policy and procurement environment, as the laggards they are.

Three. Access. The benefits of AI tooling — the productivity, the creative leverage, the cost reduction — should be broadly available, not concentrated. This means free or near-free tools for creative entry; investment in literacy at the population scale; deliberate inclusion of historically excluded creative communities in the design, training and deployment of AI systems; and structural support for the indie, the regional, the Global South, the neurodivergent and the under-resourced creative ecosystems that the AI cost reduction has the potential to *include* in the global creative economy for the first time.

The Korin AI launch, the rise of Indian AI cinema, the African-trained model investments — these are early signs that the access principle can be operationalised.

Four. Audience. The audience for creative work is not a passive market. The audience is — has always been, and increasingly is — a *participant* in the cultural meaning of the work. AI policy, platform design and creative practice should treat the audience as such: with the right to know what they are encountering, the right to choose what they pay attention to, the right to refuse work that violates their cultural or ethical preferences, and the right to a creative economy that produces *for* them rather than *at* them.

The slop ceiling is the audience exercising this principle. The cultural rejection of cynical AI work. The death threats and the BBC investigations and the *Tiny Grandma* virality. The audience has been telling us what it wants. The institutions and the platforms are, slowly, beginning to listen.

If you can hold these four principles — *agency, attribution, access, audience* — in your head when you sit down to make a creative decision in 2026, you have a working test for whether you are building towards the generative economy or away from it. Apply the test. Often. Without too much agonising. The aggregate of millions of decisions made against this test, over the next eighteen months, is the choice we are collectively making about what comes next.

The four principles are policy and platform tests. They describe what the rules of the new economy should be. They do not, on their own, describe what working creatives should be *doing* with their hands, every day, to put themselves on the inside rather than the outside of the rule-writing. I made the case for the practitioner's version of this argument at length in [Chapter 3's Open the black box](#) section, and I want to put one summary sentence of it here, because the four principles cannot be defended by creators who do not understand the toolchain underneath them.

Working creatives need to open the black box of AI and own a real stake in how it is built. Not just use it. Not just refuse it. Not just bargain over its terms. *Open it.* Understand what the model was trained on. Read the EULAs of the platforms you ship work through. Run some part of your stack on open-weight infrastructure ([Chapter 16](#) is the practical map). Show up to the governance conversations — the Cannes Disclosure Standard rooms, the UK consultation responses, the SAG-AFTRA bargaining tables, the C2PA standards body. The cohort of working creatives that does this defines the next era's craft. The cohort that uses the toolchain without ever asking what is inside it has the era's craft defined for them, by the platform companies that built the toolchain. The four principles assume — they cannot deliver on their own — that the first cohort is large enough, organised enough and technically literate enough to do the work. The argument of this chapter, on its widest reading, is that turning up to the rule-writing is the practitioner's work.

The cost we are not pricing

I want to break my own framing for a moment, because there is a category of cost the four principles only obliquely address, and any reader who has been paying attention to the wider

technology press in the period this book covers will be wondering when I am going to put it on the page.

The cost is *the resource footprint of the systems we are building everything on top of*.

Training and running large generative models, at the scale the creative industries are now using them, is an industrial activity. It is consuming enormous quantities of electricity, water, semiconductor manufacturing capacity, and rare-earth-mineral throughput, in a global infrastructure-build the planet has not seen since the rollout of mobile telephony. The trade-press coverage on this is patchy and the data the platform companies disclose is partial, but the direction is unambiguous. The data centres that produce the *Sora 2* clip you generated for a client this morning, the *Veo 3.1* sequence you cut into your edit, the *Suno* track you used as scratch on a project — those data centres draw power on a scale that is, in aggregate, a meaningful new line item in the global energy mix.

There is also a second category of cost, equally unpriced, that the four principles touch only by implication. The training, refinement and moderation of these systems involves an extensive workforce of human data labellers, content moderators and reinforcement-learning evaluators, much of it concentrated in low-wage labour markets in the Global South, much of it carrying significant psychological cost from exposure to the worst of what the models filter out. The creative economy of 2026 is, in part, sitting on top of that labour. The labour is not always visible in the marketing copy of the platform companies that depend on it.

I have under-treated both of these costs in the chapters above, on the grounds that the book is specifically about the creative-industries transition rather than about AI's wider externalities. That is a defensible editorial choice. It is also an incomplete one, and a fair reader is right to ask why a manifesto for a humane creative economy stops at the edge of the studio.

The honest answer is this. The four principles, as I have stated them, are *necessary but not sufficient*. A creative economy that gets agency, attribution, access and audience right, but that does so on top of a platform stack whose energy and labour externalities are concentrated, opaque, and morally indefensible, is not a generative economy. It is an extractive economy with better creative-side ethics — which is, in some ways, the more dangerous animal, because the creative-side ethics provide ideological cover for the wider extraction.

The implication for the working creative reading this is *not* that you should refuse to use AI tools because of their environmental and labour costs. The implication is that you should *demand the same transparency* from the platforms about their externalities that you demand about their training data. Energy disclosure, water disclosure, labour conditions in the data-supply chain, carbon accounting of inference runs — these should be standard procurement requirements for any creative organisation buying platform access at scale, in exactly the same way C2PA provenance is becoming a standard requirement for the work itself.

If I were rewriting the four principles from scratch, I would probably make this a fifth: **Footprint**. The full resource cost of the work — energy, water, human labour, supply-chain externalities — should be visible to the people commissioning, making and consuming it. The principle would sit awkwardly alongside the other four, because it is the only one that does

not run along the human-creative-vs-platform axis the book has otherwise been organised on. I have left the four where they are, because the rhetorical compression of *Agency*, *Attribution*, *Access*, *Audience* is one of the more useful things I have written. But the fifth is real. It belongs in the conversation. I want, in this section, to have at least said so on the page.

The deeper point, which I will not dwell on further because it is the subject of a different book by a different writer who I hope is writing it now, is that the AI transition is — like every previous industrial transition — being paid for somewhere. The creative-industries part of the bill is, in 2026, becoming visible. The energy and labour parts of the bill are still, by deliberate platform-company design, harder to see. A humane creative economy will have to insist on seeing both.

The DreamLab model

I want to talk briefly about the **DreamLab** model — not because I think the studio I run in the North West of England is some kind of ideal, but because it is one specific, concrete instance of how a working creative organisation has tried, deliberately, to build for the generative economy rather than against it, and because I think the choices we have made are useful to share.

DreamLab is a collective of about fifty practitioners — artists, technologists, directors, games developers, storytellers — based in the North West, with collaborators across the U.K. and internationally.⁵¹¹ We are emphatically *not* an AI company. We are a *creative* company that happens to use AI heavily.

The choices we have made, since we started thinking deliberately about how to operate in this environment in early 2024, are not heroic. They are pragmatic. They include:

A disclosure-first practice. We tell clients, in the brief, what AI tools we propose to use and what for. We document our use. We can — and have, when asked — provide chain-of-custody on contested work.

A continuum-first contract. Our client contracts identify, function by function, where on the Human–AI Agency Continuum the work will sit. Some functions are 100% human (performance, lead creative direction, story). Some are mixed. Some are mostly AI (asset variations, plate generation, certain post-production tasks). The client knows. The team knows. The line is negotiable but it is drawn.

A literacy-first hiring approach. We hire — and we invest in upskilling — generalists rather than specialists. We expect every working member of the collective to understand the AI tools in *adjacent* disciplines well enough to brief them and judge their outputs. The portfolio creative model from Chapter 11 is the operating practice.

An open-toolchain commitment. We use a deliberately diverse toolchain — Adobe, Runway, ComfyUI, World Labs, Hunyuan, open-source models on Hugging Face, internally built tooling. We refuse exclusive dependency on any single platform. If one of the tools shifts in a direction that violates the four principles, we have alternatives.

A community-first relationship to the city. The collective is intentionally based in the North West of England, not in London or Los Angeles. The aim, explicitly, is to build creative

infrastructure in places that the previous generation of the creative economy did not invest in. We hire locally. We train locally. We work, where we can, with regional partners.

A newsletter, an archive, a public record. *Dream Machine*, the publication this book is built out of, exists in significant part because *putting the work in public* — keeping a transparent, citable record of what is happening, what is being said, what is being decided — is itself a form of building the generative economy. Hidden knowledge concentrates power. Shared knowledge distributes it.

None of these choices are unique to DreamLab. Many of them are being made, in different forms, by studios, agencies, labels and indie production companies all over the world. What I would say, having lived inside this set of choices for the period this book covers, is that they *work*. The studio's output has been better, our team has been happier, our clients have been more loyal, and our market position has been more defensible than I had any right to expect when I started writing the newsletter in October.

The generative economy, as a practice, is not abstract. It is a set of operational choices that working organisations can make on a Monday morning. The choice produces real, measurable outcomes. It is not, as some of the more pessimistic AI commentary frames it, a luxury for organisations that can afford to be high-minded. It is, in our experience, the most *commercially* sustainable position available right now.

What working creatives should do on Monday morning

This book is mostly aimed at working creatives. The studios and the institutions and the platforms have been the subject of the chapters, but the audience for the chapters has been you — the writer, the director, the songwriter, the games designer, the photographer, the editor, the producer, the agency creative, the indie filmmaker, the YouTuber, the freelance illustrator, the student looking at a creative career and trying to figure out whether to be discouraged or excited.

What should you do on Monday morning?

I will keep this brief, because the rest of the book has been the long version.

Read your own newsletter. Whatever your version of *Dream Machine* is — the publication you trust to tell you what is happening in your discipline week by week — read it carefully. The pace of change in 2026 is too fast to absorb by osmosis. You need a deliberate reading practice. If you can't find one you trust, build one.

Draw your lines. Write down, for your own practice, where on the Human–AI Agency Continuum you want to sit, function by function. Update the lines as the technology and your judgement evolve. Be prepared to *show* the lines — to clients, to collaborators, to yourself.

Practise briefing. It is the most leveraged skill of the era. Brief the agents. Brief your collaborators. Brief yourself. Get clearer, every week, about what you actually want from a piece of work *before* you start making it.

Cultivate taste, on purpose. Look at more good work. Look at it harder. Articulate, every week, what makes something good — and what makes the average thing average. The agents will, by default, give you average. Your job is to know better.

Disclose. Tell people what you are using, how you are using it, and where the human work in your output lives. The transparency is, in 2026, not a cost. It is a competitive advantage.

Stay in the work. Resist the temptation to abstract too far. The maker who never makes is the maker whose judgement decays. Maintain craft contact, in at least one part of your practice, that does not depend on AI tooling. The contact will keep your eye sharp for everything else.

Find your coalition. The 88% in the U.K. consultation didn't get there by accident. It got there because creators turned up alongside other creators alongside professional bodies alongside unions alongside institutions. The political and economic shape of the next decade of creative work depends on creators being *organised*. Join your union. Sign the declarations. Show up to the consultations. The institutions of collective bargaining and political representation that your forerunners built are the institutions that will defend your work in 2030.

Build for the new geography. Wherever you are in the world, the AI cost reduction has made it more viable than ever to build a serious creative practice in places that were previously locked out of the centre. Take the opportunity. The next century of cultural production is not going to belong to the cities that owned the previous one. The next generation of canonical creative work will, on the trajectory I see, come from more places, in more languages, in more forms, than ever before.

Make the work. None of the rest of this matters if you do not make the work. The AI era is not a reason to stop. It is a different set of conditions under which to keep going. Working creatives have always operated against the technological grain of their day. The maker in 2026 — like the maker in 1926 and 1826 and 1626 — is the person who finds, in the conditions of their actual moment, the work that needs to be made and then makes it.

The conditions have changed. The need has not.

Predictions you can check me on

Books on the cusp of a transition tend to age in one of two ways. Either they were broadly right about the shape of what was coming and look prescient five years on, or they were embarrassingly wrong and quietly disappear from the recommended-reading lists. The honest thing to do, at the end of a book like this one, is to write down predictions specific enough that the future can grade them.

Here are mine, dated May 2026.

By the end of 2027. Every major studio, label and agency has, in its standard production contract, a clause requiring AI-use disclosure across the production pipeline. *Position Four* — *AI in the workflow, not in the work* — is the dominant operational posture in legacy creative organisations. The number of *Position Three* refusing studios has stabilised at roughly 10–15% of the mid-and-upper market, concentrated in IP-led franchises where the audience contract depends on a human-craft signal.

By the end of 2027. The 44%/3% Deezer ratio (Chapter 5) has stabilised on the upload side at 50–65% AI by volume, with consumer streams of fully-AI tracks still under 5% of total play

time. *Audience-disclosed* synthetic music is a small but real category — under 10% of paid streams — and is concentrated in mood/background and synth-driven genres rather than in artist-led popular music.

By the end of 2028. A *Position Two* studio — Imaginae, Wonder, Obsidian, Asteria or a name we have not heard yet — produces the first AI-native feature that wins material awards recognition *for its writing or direction* (not for its technology). The audience response is mixed; the cultural permission for AI-native cinema has shifted from contested to accepted-with-asterisks. James Cameron’s “horrifying” line is still being quoted but the position it describes has narrowed to *AI that displaces a specific actor in a specific scene*, not the general toolchain.

By the end of 2028. At least one large national government — most likely the U.K. or one of a few EU member states — has passed legislation requiring licensing of copyrighted work for AI training. The U.S. has not, but the litigation environment (UMG v Anthropic and the cases that follow it) has produced de facto licensing markets that platform companies have grudgingly entered. The technical infrastructure for *creative weight attribution* is mature enough to underpin real revenue distribution to working creators.

By the end of 2029. World models, not flat video, are the dominant medium for new high-production-value creative work. Marble-class tools are commodity infrastructure; the differentiated work is being done by orchestrators with strong *world curation* skills (Chapter 8). The boundary between film, games and immersive media is functionally gone for new productions, while legacy formats retain prestige and a real audience.

By the end of 2029. The audience contract has been substantively re-written. C2PA-class provenance metadata is supported in the major capture, edit and distribution tools by default. Platforms that do not honour provenance have lost meaningful share to ones that do. The *Tiny Grandma* error mode — human work mis-flagged as synthetic — is a solved problem in the principal platforms; the inverse error — synthetic work mis-flagged as human — is the harder one and remains the central audience-trust question.

By the end of 2030. The orchestrator role (Chapter 11) is a named seniority track in every creative discipline that survives the transition. *Senior orchestrator* is a credentialled title. The apprenticeship problem (Chapter 11, Chapter 13) is partially solved by a combination of three things: AI-augmented junior roles maintained deliberately by *Position Four* studios; new pathways through literacy initiatives like Sundance Collab and its successors; and an explicit re-funding of cultural-institution training programmes by national governments.

By the end of 2030, but I’m less confident about this one. A major creative-industry union has negotiated a *productivity dividend* — a structural share of AI-driven productivity gains that flows to the human workforce, not just to the platform companies and the studio shareholders. The mechanism is novel and contested but the underlying maths is unanswerable; the industry that does this first becomes the template for the others.

I am wrong about something on this list. I do not know yet which item. If you are reading this in 2030 and one of these predictions has aged badly, send the receipts to the *Dream Machine* newsletter address. I will publish them.

The predictions I have *not* made — about which AI company will dominate by 2030, about which specific platform companies will go bust, about whether AGI arrives in the period and what it does to all of this — are absences I have made on purpose. Anyone confident about those questions in 2026 is selling something. The shape of the *creative economy* is more predictable than the shape of the *platform layer underneath it*. That is the bet this book has been built on.

A note to those who are afraid

I want to write directly, for a paragraph or two, to the readers I know are sitting with this book — and with the larger transition this book describes — in a state of real, well-founded fear.

If you are reading this and your livelihood depends on a creative discipline that the AI tools are getting unsettlingly good at — if you are a junior animator, a session musician, a copywriter, an illustrator, a translator, a stock-image photographer, a voiceover artist, a foley artist, a concept designer, a stage actor in a regional theatre, a working musician on tour, an indie filmmaker, an entry-level games artist, a video editor at a small post-production house — and the news cycle has, over and over again, told you that your job is going away, I want you to know two things.

The first is that the news cycle is, on this question, almost always too negative. The aggregate data on creative-industry employment in 2026 does not show the collapse the headlines have been predicting. It shows a *restructuring*, with a lot of disruption, and with real pain in specific places, but with — net — more creative work being done by more people in more forms than was being done a year ago. The Sundance literacy initiative is not happening because the industry is dying. It is happening because the industry, freshly, has more entrants than ever and needs to train them.

The second is that the most important strategic move you can make in this moment is not to be afraid in private. It is to *speak*. To your union. To your trade association. To your local government's consultation. To your manager. To your audience. To the people on your team. The 88% was made out of voices. The Tilly Tax was made out of voices. The Bandcamp ban was made out of voices. The Sundance literacy turn was made out of voices. The institutions of the new creative economy — the ones that will protect your work in 2030 — are being shaped *right now* by the people who turn up to be counted. The people who don't turn up, or who turn up only in private, are the people whose interests get absorbed into the policies of the people who do.

You are not powerless in this. You are, if anything, the most powerful constituency in this whole story, because you are the people who are actually making the work that the audience is paying attention to. You are the constituency the unions represent, the institutions exist to serve, the audience cares about, and the platforms — ultimately, despite themselves — depend on. The choice between the extractive and the generative economy is, in large part, the choice you collectively make about how to use that power.

Use it.

The Dream Machine

I want to close the argument of this book with the phrase the newsletter started with.

Welcome to the Dream Machine.

When I named the newsletter, in late September 2025, I had no clear idea what the phrase meant. I had a half-formed sense that “dream” caught something about the strange, vivid, slightly hallucinatory quality of what the new tools produced, and that “machine” caught something about the industrial, scaled, infrastructural nature of the platforms behind them. The phrase was, more than anything else, a placeholder for a feeling I couldn’t quite articulate yet.

Six months on, I think I know what the phrase means.

A *dream machine* is, in the sense I am using it, an apparatus capable of producing something that has — until very recently — required a human mind to produce. Not just images, or videos, or songs, but *constructions* of meaning, of place, of presence, of emotion. The apparatus is not, in itself, doing the dreaming. It is amplifying, multiplying and distributing the dreaming of the humans who direct it.

The question this book has been about, in every chapter, is *whose dreams the machine amplifies*.

If the machine amplifies the dreams of a small number of platform shareholders, optimised against extraction, the creative economy it produces will be — to use Kehlani’s line about Xania Monet — a place where “the person is doing none of the work.” The dreams of the audience are *for sale*. The dreams of the creators are *training data*. The dreams of the culture are *outputs of an algorithm tuned for engagement*.

If, on the other hand, the machine amplifies the dreams of the working creatives — supported by the audience, defended by the unions, regulated by the institutions, contained by the four principles, made transparent by the provenance infrastructure — the creative economy it produces will be the largest, most diverse, most accessible expansion of human cultural production in the history of the species.

We are, in May 2026, sitting in the moment between these two outcomes.

The signals from the last six months — the 88%, the slop ceiling, the Sundance turn, the Cannes Disclosure Standard, the Academy rule, the SAG-AFTRA contract, the audience behaviour, the creator coalition, the regional opening, the literacy investment — all point at the second outcome being achievable. Not inevitable. *Achievable*.

It is going to require the working creatives to keep turning up. It is going to require the studios, agencies and labels to keep making the integration choices we covered in Chapter 13. It is going to require the institutions — the unions, the rights bodies, the festivals, the universities — to keep doing the slow, unglamorous work of building the rails. It is going to require the platforms to be pushed, by their users and by their regulators, towards the side of provenance and attribution and consent. It is going to require the governments to keep the policy moving in the direction the 88% asked for.

It is going to require the audience to keep paying attention to the human signal.

And it is going to require people like you — the people who have read this far, who have been doing the work of figuring out what creative life looks like in this moment — to make, in your own practice, the daily decisions that build the generative economy rather than the extractive one.

This is the choice. This is what is on the table. This is the work.

Welcome to the Dream Machine.

Chapter 16 — The Tools

I have, until this chapter, deliberately kept the *tools* out of the foreground. Thirteen chapters about creative AI without a chapter on the toolchain is, on the face of it, a strange editorial decision, and I want to begin by explaining it.

The reason is that I think the most common mistake people make about this period is to confuse *the tools* with *the transition*. Tools are the visible surface of the change — the thing the press cycle covers, the thing the platform companies want you to talk about, the thing that has a price and a logo and a launch date you can put on a slide. The transition is everything underneath: the economics, the labour, the audience contract, the law, the institutions, the rails. The tools change weekly. The transition is slower, deeper, and is what will still be true in 2030 when most of the tools in this chapter are obsolete.

The first sixteen chapters were about the transition. This chapter is about the tools.

I have written it last on purpose. Read in this order, the tools sit inside the architecture the book has been building — the Continuum, the Slop Ceiling, the four positions, the orchestrator role, the four principles. Read in any other order, they collapse back into the format the platform companies prefer: a tools-arms-race in which the only question is which model is best this week.

That format is, in 2026, the most reliable way to misunderstand what is happening.

A note on the date stamp. Everything in this chapter is current to June 2026 — the May 2026 cut of the catalogue, with the Issue 31 (June 2026) additions to the video, music, world-model, agent and games-engine layers folded in where they fit. By the time you read it, some of these tools will have been bought, renamed, killed, surpassed or repositioned. The point is not that the specific tools matter. The point is the *shape* of the toolchain — what categories exist, what they do, who builds them, and how a working creative builds a coherent stack on top of them. The shape, in my experience, holds.

How to think about the toolchain

Before the inventory, the frame.

I think the creative-AI toolchain in 2026 is best understood as a stack of seven layers, each with its own dominant players, its own pace of change, its own integration model. The layers, from foundation to consumer, are:

1. **Foundation models** — the large multimodal systems underneath everything else (OpenAI's GPT-class, Anthropic's Claude, Google's Gemini, Meta's Llama, the major Chinese open-source models). These are the raw capability layer. Almost no working creative uses them directly except via wrappers.
2. **Modality models** — specialist models for video (Sora, Veo, Runway Gen-4.5, Kling, Hunyuan, Wan), image (Firefly, Midjourney, FLUX, Imagen, SDXL/Stable Diffusion variants), audio (Suno, Udio, ElevenLabs, Mureka), 3D and world (Marble, Genie 3, WorldGen, UNI-1, Hunyuan 3D-PolyGen, ECHO). These are what most working creatives think of when they say "AI tools."
3. **Agent platforms** — systems that compose modality models and external tools into multi-step workflows (OpenAI's AgentKit, Anthropic's Claude apps and skills, Google's

Project Genie, Heygen’s Video Agent, Sony’s 49-agent / 72-skill stack). The agent layer is where the “orchestrator economy” of Chapter 11 actually runs.

4. **Creative software with AI baked in** — the legacy creative suites that have been rebuilt as AI-first platforms (Adobe Creative Cloud, Autodesk, Foundry, Unreal Engine, Unity, DaVinci Resolve, Pro Tools, Logic Pro, Ableton). This is where most paid professional work still happens.
5. **AI-native creative apps** — new entrants whose product is a single-purpose AI workflow (Runway, Higgsfield, Krea, Freepik, Magnific, Synthesia, Heygen, Hedra, Cascadeur, Pika, Luma). Most working creatives use 4 to 10 of these in rotation.
6. **Open-source and workflow infrastructure** — the technical-creator layer that wires everything together (ComfyUI, Hugging Face, SuperSplat, OpenEnv, the open-source model ecosystem). This is where the most interesting innovation often happens first.
7. **Consumer surfaces** — the apps that put generative capability on every phone (the Sora app, CapCut/Dreamina with Seedance, the Gemini app, the various TikTok-style remix platforms). This is the layer the audience touches.

The mistake I see most often, both in the press cycle and in studios planning their internal AI roadmaps, is to optimise for layer 2 (modality models) without understanding that the actual leverage is in how you compose layers 2, 3, 4 and 6 into a coherent workflow. The tool that “wins” is rarely the tool with the best benchmark. It is the tool that integrates cleanly into the rest of your stack.

With that frame, the inventory.

Video

The video layer changed faster than any other modality between October 2025 and June 2026, and is the one most likely to look different again by the time you read this. Treat the names as snapshots, not as a stable league table.

Sora 2 (OpenAI) is the model that opened the period this book covers. Its September 2025 launch — physical realism, audio integration, multi-shot world-state persistence — is the moment Chapter 1 is about.⁵¹² The iOS app launched alongside it hit a million downloads in five days.⁵¹³ By March 2026, however, the platform was shutting down. OpenAI announced a staged wind-down on **24 March 2026** — consumer app and website ending late April, developer API closing late September — with peak revenue reportedly at around \$540k per month against operating costs in the billions, and compute reallocated toward higher-margin coding and enterprise products.^[^16thetools-2a] The Disney–OpenAI \$1bn licensing announcement of 11 December 2025 was never executed; no formal agreement was signed and no money changed hands. The contractual *framework* underneath that deal — IP-not-training, talent-likeness excluded, jointly-governed steering committee, equity-and-warrants structure — survived the platform’s collapse and is now the reference architecture other studios negotiate against; the deep treatment of that architecture and the broader IP-economics shift it sits inside is in [Appendix L: The Programmable Brand](#). For working filmmakers in mid-2026, the practical implication of the Sora collapse is that the *consumer-*

facing edge of the AI-video market has rotated to ByteDance's CapCut/Dreamina/Seedance stack, to the YouTube Shorts integrations of Veo 3.1, and to the Sora app's various TikTok-and-Reels-shaped successors. The professional pipeline never depended on Sora 2 in the first place — most working filmmakers I know used it less than its cultural prominence suggested — and so on the production side the collapse changes less than the headlines implied.

Veo 3.1 (Google DeepMind), released in mid-October 2025, is the model the professional filmmaking community has, on average, gravitated toward — for narrative coherence, controllable camera composition, cinematic lighting vocabulary and sound integration.⁵¹⁴ Sora 2 wins on raw physics in single clips; Veo 3.1 wins on the kind of sustained directorial control most actual production pipelines need.

Runway Gen-4.5 (and Gen-4.5 Image-to-Video, the Workflows product, Story Panels, Characters API, Apps for Advertising) is the most-integrated commercial AI-video stack of the period.⁵¹⁵ Runway has shipped product faster than any other AI video company in this market, and CEO Cristóbal Valenzuela's "fifty indie films instead of one \$100M blockbuster" framing is the cleanest articulation I have seen of the case for AI as creative leverage rather than cost-cut.⁵¹⁶

Kling (Kuaishou), **Hunyuan Video** (Tencent), **Wan 2.5** (Alibaba), **Seedance 2.0** (ByteDance) — the Chinese-built models that, in aggregate, have rivalled or surpassed the U.S. labs on specific capabilities (motion physics, character consistency, render speed) at significantly lower cost.⁵¹⁷ Hunyuan's open-source releases have been the single most important contribution to the wider open-source AI video ecosystem in this period.

Pika 2.0, **Higgsfield**, **Luma** (Dream Machine and UNI-1) round out the commercial layer. Each has carved a niche: Pika on iteration speed and creator workflow; Higgsfield on social-media marketing video at scale (\$80M raised, \$1.3B valuation, \$200M revenue in nine months⁵¹⁸); Luma on the world-model bridge to spatial content.

Heygen ships Video Agent — a full scripting-to-assembly agent built around reference images.⁵¹⁹ **Synthesia** holds the corporate AI-avatar market (\$4B valuation, having reportedly rejected a \$3B Adobe acquisition offer).⁵²⁰ **ElevenLabs** runs the dominant audio-AI layer underneath much of the new video work (\$500m ARR by April 2026).⁵²¹

Gemini Omni (Google DeepMind), announced at Google I/O 2026, brings text, image, audio, video and live interaction into a single multimodal model — the first foundation-model release in this category that meaningfully unifies the modalities working creatives currently have to bridge across five different tools.⁵²² **Beeple Canvas**, Mike Winkelmann's gen-AI compositor — launched May 2026 — is the first AI-native compositing application to ship from a working visual-effects artist's own studio, and is structurally distinct from the AI-features-bolted-onto-existing-compositors pattern in the legacy-software section below.⁵²³

If I had to name a single video product that, in my experience, working creatives have settled on as a default in 2026, it would be Veo 3.1 for finished work and Runway for iteration and integration. Sora is the brand name the audience knows. The actual production pipelines run on the other two.

Image

The image layer is more stable than video — the technology has matured, the differences between top models are narrower, and the dominant question has moved from “which model” to “which workflow.”

Adobe Firefly (Image Model 5, plus Foundry for custom-trained corporate models, plus integration across Photoshop / Illustrator / Express / InDesign) is the default for any working creative who is also a Creative Cloud subscriber — which is, by Adobe’s own numbers, 45% of Creative Cloud users actively using Firefly, 70% of those weekly, more than 22 billion assets generated by April 2025.⁵²⁴ The Firefly adoption curve is the single best evidence I have for the consumption-gap argument in Chapter 13.

Midjourney remains the aesthetic-leadership product in the category. Slower to ship, more opinionated about output style, dominant on Discord and X among the working AI-art community.

FLUX (Black Forest Labs) is the open-source and pro-creator favourite for fine control, having largely replaced Stable Diffusion XL as the open-weight default through 2025.

Google Imagen (and the *Nano Banana* fast-image variant integrated into Gemini, Photoshop and Unreal Engine via the various plugins) has become the most-integrated image model in the consumer toolchain, by virtue of Google’s distribution. Nano Banana inside Photoshop and Nano Banana inside Unreal Engine were two of the more consequential cross-platform integrations of the period.⁵²⁵

Krea, Freepik, Magnific, Recraft — the higher-control consumer / pro-creator products built on top of foundation image models. Each is competing on specific workflow advantages (real-time generation, upscaling, vector output, brand-consistency control).

The image workflow most commonly cited in my circle in mid-2026 is: a base generation in Firefly / Midjourney / FLUX, character-consistent variation in a controllable wrapper like Krea or Magnific, finishing inside Photoshop with the AI-assisted masking, generative-fill and object-removal tools that Adobe shipped through the autumn 2025 and spring 2026 update cycle.

Music and audio

The music layer split into three categories during this period, and the split is, in my experience, more important than the specific products in each category.

Generative music tools that produce finished tracks from prompts — **Suno** (Studio launched late 2025⁵²⁶), **Udio, Mureka** (with its Music Agent Studio, six specialised AI agents for songwriting, arrangement and production⁵²⁷). These are the tools that produce most of the AI-music flood Chapter 5 describes. They are also, paradoxically, the tools most working musicians use the least directly — the consumer market for AI-generated finished tracks is large and growing, but professional musicians overwhelmingly use AI tools at a different layer.

Production and post-production AI — the tools that handle audio restoration, mixing, mastering and isolation. The 1,100-creator music survey discussed in [Appendix D](#) found that 58% of professional producers used AI for audio restoration, 38% for mixing assistance, 33.9% for automated mastering. **iZotope Ozone 12**, **LANDR**, the Pro Tools and Logic Pro AI suites, **CleanvoiceAI** for podcast post — this is where the silent-adoption majority of the music industry lives.

Voice and audio synthesis — **ElevenLabs** is the dominant player, with \$500m ARR, BlackRock / NVIDIA backing, and meaningful share across audiobook narration, dubbing, podcast synthesis and AI character voice work.⁵²⁸ The Cardiff band that found their music had been used to train an “AI artist” outperforming them on Spotify⁵²⁹ is one of the cautionary tales of this layer; the Andrii Daniels bomb-shelter clip⁵³⁰ is one of the success cases.

Sound-effect foundation models emerged as a new sub-category in May 2026. **Sony AI’s Woosh** is the first foundation model explicitly trained for the professional sound-effects discipline — built for the people who design the sonic worlds behind games, film and interactive media, not for the consumer market.⁵³¹ **Mirelo SFX 1.6** shipped the first sound-effects model that lets you *edit* a generated sound rather than only regenerate it — a structural shift in the discipline equivalent to the move from rendered images to layered Photoshop files.⁵³² **Stable Audio 3.0** (Stability AI) shipped as an open-weight audio model family explicitly aimed at artistic experimentation.⁵³³ **Tamber**, the ethically-trained AI music suite I describe in [Chapter 6](#), shipped alongside a gestural-control interface that lets the musician steer the generation with arm movements.⁵³⁴ **Beatport’s Track ID** rolled out as the real-time identification standard for the DJ market.⁵³⁵

The deal flow underneath this layer is the second-fastest-changing in the toolchain. The Stability AI / Universal Music alliance, the Stability AI / Warner Music deal, the Splice / Universal partnership, the GEMA / OpenAI lawsuit, the Wixen / Meta lawsuit, the UMG / Anthropic \$3B suit — these are the structural moves I would track if I were a working musician trying to plan a five-year toolchain.⁵³⁶

3D, world models, spatial

The category that, more than any other, I think defines the next decade of creative work. [Chapter 8](#) is the long-form argument; this section is the inventory.

Marble (World Labs, Fei-Fei Li’s company) is the first commercial product I would put on a professional toolchain.⁵³⁷ Public release November 2025; Sony Pictures’ use of it in virtual production reportedly running 40× faster than the legacy workflow.⁵³⁸ DreamLab has been in the beta since October 2025, and Marble is, today, the world-model product most integrated into a working pipeline I have used.

Google DeepMind Genie 3 is the most ambitious research-grade world model, named by *Time* as one of the best inventions of 2025. Made publicly available to Google AI Ultra subscribers via Project Genie in January 2026.⁵³⁹

Meta WorldGen, **Tencent HY World 1.5** (open-sourced December 2025, alongside the Hunyuan 3D Studio integration⁵⁴⁰), **SpAItial ECHO**, **Stanford Wonderzoom**, **OpenArt Worlds**, **Luma UNI-1** (the most important *category* announcement of spring 2026, combining world generation with reasoning⁵⁴¹) — the rest of the world-model commercial layer.

The May 2026 world-model wave extended this layer further. **NVIDIA SANA-WM** is the first open-weight world model at meaningful scale (2.6B parameters), with 60-second video generation and explicit camera control.⁵⁴² **Odyssey Starchild-1** is, by Odyssey’s own framing, “*the first ever real-time multimodal world model*” — a system that doesn’t just generate a world but simulates and reasons about it.⁵⁴³ **Odyssey Agora-1** is the multiplayer companion to Starchild, putting four players inside the same AI-generated world (built, in a small piece of provenance theatre, on the bones of a 1997 shooter).⁵⁴⁴ **Apple Headsup** is a research-grade 3D Gaussian head-reconstruction pipeline built for multi-view captures from consumer iPhones, extending the Vision-Pro-Personas Gaussian-splat thread into the open research layer.⁵⁴⁵

Underneath this layer, the Gaussian-splatting infrastructure has matured into a stable workflow: **SuperSplat** (PlayCanvas) for editing, **Spark 2.0** for open-source streaming of 100-million-splat scenes to browsers, the SOG / WebP equivalent compression standard.⁵⁴⁶ Apple’s confirmation that its Vision Pro Personas feature is powered by Gaussian splatting under the hood made it, by some margin, the most-deployed Gaussian-splat technology in consumer hardware as of late 2025.⁵⁴⁷

For the 3D-asset and material side: **Hunyuan 3D-PolyGen 1.5** (Tencent’s “art-grade” 3D generative model), **Hitem3D**, **Meshy**, **Rodin** — the rapidly-maturing 3D-asset generation layer that is being integrated, model-by-model, into Unreal Engine, Unity and Blender pipelines.

Ubisoft’s open-sourcing of its **CHORD** PBR-material model in December 2025⁵⁴⁸, and the Blender Foundation’s patronage deal with Anthropic announced in May 2026⁵⁴⁹, are two of the more strategically significant moves in this layer — both pushing the production-grade open-source tooling forward at a pace the commercial alternatives have struggled to match.

Agent platforms and orchestration

The category I think most working creatives are still underestimating, six months after Chapter 3 argued it was the inflection point of the era.

OpenAI AgentKit (Agent Builder, ChatKit, connector registry, eval framework) launched October 2025 and is the developer-facing platform underneath most third-party agentic creative tools.⁵⁵⁰

Anthropic Claude apps and the **Skills framework** — the system of named, reusable capabilities that Claude Code uses to coordinate multi-agent workflows. The Sony 49-agent / 72-skill game-development stack is built on this.⁵⁵¹ In May 2026, **Google** released its own **official skills for AI agents** — a parallel, cross-vendor skills layer that lets Google-side

agents do what Anthropic’s Skills framework has been doing for Claude-side ones.⁵⁵² The convergence of two named “skills” frameworks across the foundation-model vendors is, in my read, the first sign that the orchestration layer is settling on a shared vocabulary rather than continuing to fragment.

Tencent Ardot, the company’s AI-native design-agent platform launched May 2026, is the most ambitious non-Western agent-platform launch of the period — an integrated environment in which generative design agents handle layout, asset generation, brand application and iteration as a single coordinated pipeline.⁵⁵³ In the same week, **Viktor** raised \$75M to embed an agentic *coworker* directly into Slack and Microsoft Teams — i.e., the agentic layer landing not as a standalone product but as a colleague-shaped presence in the chat surface the working creative is already in all day, as discussed in Chapter 9.

Heygen Video Agent for end-to-end video assembly.⁵⁵⁴ **Adobe CX Enterprise** (announced at Adobe Summit 2026 with NVIDIA) for “agentic creative intelligence” across the full content lifecycle.⁵⁵⁵ **NVIDIA + Google Cloud** for the underlying creative-AI infrastructure most enterprise pipelines run on.⁵⁵⁶

ComfyUI — the open-source node-based workflow editor — sits underneath much of the technical-creator community’s agentic work. ComfyUI raised \$17M in October 2025⁵⁵⁷ and hit a \$500M valuation by May 2026⁵⁵⁸; the platform has become the de facto OS for the open-source side of the creative-AI ecosystem. In May 2026 **Anthropic’s Claude** was added as an official partner node inside ComfyUI, joining the existing Google, OpenAI and open-weight nodes — meaning the three frontier foundation models can now be orchestrated side-by-side inside the same open-source pipeline.⁵⁵⁹

Hugging Face, OpenEnv (Meta / Hugging Face), the **Hugging Face / Google Cloud** partnership — the open-source agentic-development infrastructure.⁵⁶⁰

For working creatives, the practical agent stack in 2026 is some combination of:

1. A foundation model (Claude / GPT / Gemini) for the orchestration brain.
2. A modality model layer (video, image, audio, 3D) doing the actual generation.
3. A workflow integration layer (ComfyUI for technical work, the in-app agent surfaces for less technical work).
4. A judgement layer (the human at the desk, doing what Chapter 11 calls *briefing, allocating, judging and integrating*).

The team I work with at DreamLab runs this stack in production every week. The agents that handle our daily work in May 2026 are, in aggregate, doing the labour of what would, two years ago, have been a team three to four times our size. The human team has not shrunk. We have just become substantially more leveraged.

Legacy creative software, repositioned

The most under-reported strategic story of this period, in my view, has been the speed at which the legacy creative-software vendors have rebuilt their products as AI-agent platforms.

Adobe — I have written enough about Adobe in Chapter 9 that I will not repeat it here. The short version: Creative Cloud is, today, a stack of AI agents wearing a Photoshop / Premiere / After Effects / Illustrator / InDesign / Acrobat skin. The agents are inside the apps; the apps are inside ChatGPT; the apps are inside Adobe Express; the apps are inside the new CX Enterprise platform. The repositioning is complete.

Unreal Engine (Epic) — the games engine that has, through plugins, integrations and the Nano Banana / Gemini partnership, become a hybrid game-engine / virtual-production / AI-generation hub. The Unreal Engine 5 AI Assistant, announced at the end of 2025⁵⁶¹, is one of the more consequential single-product launches of the period. The **ECABridge** connector, launched May 2026, is the most-cited Unreal-Engine MCP integration of the spring — providing the Model Context Protocol surface and a set of agentic capabilities Epic itself has not yet shipped to the launcher.⁵⁶² In a separate but related move, an **Epic Games veteran** announced an AI-heavy “*Fully European*” game-engine project in the same week — the first plausibly-credible new entrant in the AAA game-engine market since the early 2010s, framed explicitly around AI as the core operating layer.⁵⁶³

Unity — Unity’s AI Open Beta (May 2026), an in-editor AI suite for the full games-development pipeline, alongside the company’s AI Council formation in October 2025.⁵⁶⁴

Autodesk, Foundry, SideFX — the VFX-pipeline vendors integrating generative AI into Maya, Nuke and Houdini at the speed the VFX industry’s adoption curve (62% of Hollywood studios on automated compositing, 35% reduction in post-production timelines⁵⁶⁵) demanded.

Blender — open-source 3D, now a recognised industry-grade tool, beneficiary of the Anthropic Foundation patronage deal.⁵⁶⁶

DaVinci Resolve (Blackmagic), **Avid Media Composer**, **Pro Tools** — the editorial and audio post environments, all now shipping AI-augmented features that have become baseline expectations.

The thing to note is that the legacy software did not get displaced by the AI-native products. The legacy software *absorbed* the AI-native capability and kept the underlying user community. Adobe was supposed to lose to Midjourney in 2024; Adobe is, instead, the dominant generative-AI player by aggregate creator engagement in 2026. The platform companies bet on this absorption pattern, and that bet has, so far, paid off.

Open source

The open-source ecosystem has, against the odds and against most VC predictions in 2024, held its ground through this period and is, in several categories, the leader rather than the follower.

Hugging Face — the operating system of open-source AI, expanding aggressively through 2025–26.

ComfyUI — already discussed.

Open-source models from Tencent (Hunyuan), Alibaba (Qwen, Wan), DeepSeek, Meta (Llama), Mistral, Stability AI — collectively, the open-weight ecosystem that, by the spring of 2026, was being used by approximately 80% of startups pitching the Andreessen Horowitz fund.⁵⁶⁷ **NVIDIA’s SANA-WM** (May 2026) extended this list to world-models for the first time at meaningful parameter scale.⁵⁶⁸

PhotoGIMP, the open-source skin that takes GIMP and makes it look and feel exactly like Photoshop, became, in this period, a credible Photoshop alternative for working creatives who wanted to opt out of the Adobe subscription stack — the open-source equivalent of the *Tools I do not use* discipline in the section above.⁵⁶⁹

OpenEnv (Meta / Hugging Face) for open-source agentic development. **Korin AI** (the Africa-trained, Africa-built model launched May 2026⁵⁷⁰). **SuperSplat, Spark 2.0, PlayCanvas SOG, Blender** — the open-source spatial / 3D infrastructure layer.

If you are a working creative trying to build a long-term, defensible toolchain that does not depend on the unilateral pricing or policy decisions of a single platform company, the open-source ecosystem in 2026 is materially viable in a way it was not eighteen months ago. We have built significant parts of the DreamLab pipeline on top of it precisely for that reason.

Tools I do not use, and why

I want to be specific, because lists of “best tools” without exclusions are not useful.

I do not use AI tools whose terms of service claim ownership over my output, or that train on user inputs without an opt-out. Multiple consumer-facing AI products in this period have shipped with terms that working creatives should read carefully before adopting.

I do not use AI tools whose training data provenance I cannot, in some material way, verify or trust. The growing infrastructure for *creative weight attribution*, watermarking and C2PA compliance is, in my view, the right side of the market to be on; tools that explicitly reject that infrastructure are tools that I have, increasingly, kept out of our production pipeline.

I do not use the AI products that have made the most noise in the consumer press cycle. The marketing-driven launches — the products whose first appearance is a viral demo and whose second appearance is a Series-A round — are, in my experience, the products most likely to have collapsed or pivoted by the time you need them in production six months later.

I do not, finally, use AI tools to produce work in the disciplines where my own craft is the value I am bringing to the client. The Continuum frame from Chapter 3 is, for me, a daily operational practice, not a theoretical model. The places I sit on the right edge of the line are deliberately chosen. The places I sit on the left are deliberately defended.

The complete toolchain: a categorised reference

This section is a reference inventory, not a recommendation list. It catalogues every tool, model, platform, app, plugin, LoRA, workflow and service that *Dream Machine* tracked across its 32 issues, from October 2025 to June 2026. Some are dominant; some are niche; some have already been bought, renamed or discontinued by the time you read this. The point

of the list is not “what to use.” The point is *what existed*, in this period, in the creative-AI toolchain — so that the *shape* of the field is on the record.

A word on the list’s grain. I have tried to err on the side of inclusion. Where a single company ships multiple closely-related products — Adobe’s *Sneaks* portfolio, the Runway Gen-4.5 family, the Qwen-Edit LoRA series — I have grouped them under the parent entry but called out the constituent tools, because in this period each constituent shipped to working creatives separately and changed at its own cadence. Where a tool was a one-issue demo I could not later verify, I have still listed it; that the demo existed *at all* is part of the field’s history. Where a tool’s name conflicts with another (there are at least three things called “Wonder” in the period the book covers) I have annotated.

The list runs to roughly six hundred entries. Skim it. Use the categories. Come back to specific sections when you need them.

Foundation models / LLMs

- **ChatGPT / GPT-5 / GPT-5 Pro** (OpenAI) — the dominant consumer LLM and reference foundation model; 800–900M weekly active users; GPT-5 / GPT-5 Pro announced at DevDay 2025.
- **Claude / Claude Code / Claude Apps / Claude Skills** (Anthropic) — the writers’ and developers’ favoured second; strong long-context performance; the agentic coding environment that underlies Sony’s 49-agent / 72-skill stack; Claude for Legal launched May 2026.
- **Gemini / Gemini 2.5 Flash / Gemini 3 / Gemini 3.1 Flash** (Google) — Google’s multimodal LLM family; desktop users growing 155% YoY in 2025–26; Gemini 3.1 Flash TTS is the most controllable Google voice model as of spring 2026.
- **Llama** (Meta) — the dominant open-weight foundation model.
- **Mistral / Mistral Voxtral / Mistral Transcribe 2** — European open-source LLM; Voxtral is the next-generation speech-to-text family.
- **Qwen / Qwen 3.5-Omni** (Alibaba) — Chinese open-source LLM, image, video and audio variants; Omni is the multimodal family covering text, images, audio and video.
- **DeepSeek** — Chinese open-source LLM used heavily in agentic stacks.
- **Phi / Phi Mini** (Microsoft) — lightweight foundation models.
- **Lyria / Lyria 3 / Lyria 3 Pro** (Google DeepMind) — text-to-music foundation model family; Lyria 3 ships with SynthID watermarking and the Lyria Camera interactive demo.
- **Grok** (xAI) — xAI’s LLM family, surfacing in the creative stack through Grok Imagine; voice cloning via the xAI API from May 2026.
- **Cohere Transcribe** — enterprise transcription model; 33 hours of audio in 12 minutes.
- **Microsoft VibeVoice** — open-source frontier voice AI model.

AI video models

- **Sora / Sora 2** (OpenAI) — the model that opened the period; physical realism, audio integration, multi-shot world-state persistence; iOS app hit 1M downloads in 5 days; Sora

2 Character Creation surfaced on fal in March 2026.

- **Veo 3 / Veo 3.1 / Veo 3.1 Ingredients to Video / Veo 3.1 Lite** (Google DeepMind) — the working filmmaker’s preferred model for cinematic control; Ingredients to Video shipped to YouTube Shorts and YouTube Create; Veo 3.1 Lite is the lower-cost text- and image-to-video tier.
- **Runway Gen-4 / Gen-4.5 / Gen-4.5 Image-to-Video / Workflows / Story Panels / Characters API / Ad Concepter / Apps for Advertising / Aleph 2.0** (Runway) — the most-integrated commercial AI-video stack; Story Panels generate three-panel storyboards from a single image; Characters API ships real-time intelligent avatars; the Vera Rubin-powered real-time model runs <100 ms latency. **Aleph 2.0** (June 2026) lets the user edit a *single frame* in a video and propagates that edit consistently across the rest of the clip — the cleanest *frame-as-handle* video-editing pattern yet shipped.
- **Kling / Kling 2.5 Turbo / Kling O1 / Kling 2.6 / Kling 3.0 / Kling X-Dub / Kling Motion Control 3.0** (Kuaishou) — strong on physics and trajectory control; 3.0 adds multi-shot control, multilingual audio and 4K image generation; X-Dub is the context-rich visual dubbing variant.
- **Pika 2.0 / PikaStream 1.0** — iteration-speed-focused video generation; PikaStream brings AI agents into live video calls.
- **Luma Dream Machine / UNI-1 / UNI-1.1 / Ray3 Modify / Luma Dream Brief** — Luma’s video, world-and-reasoning, and modification stack; UNI-1.1 ships with prompt enhancement and built-in research; Dream Brief is the \$1M Cannes Lions competition.
- **Wan 2.2 / Wan 2.5 / Wan 2.6** (Alibaba Qwen) — camera-controlled video generation; 2.6 adds character reference and multishot capabilities.
- **Hunyuan Video / Hunyuan Image-to-Video** (Tencent) — open-source video model.
- **Seedance 2.0 / SeeDream 4 / SeeDream 4.5** (ByteDance) — image-to-video and finished-video models, integrated into CapCut / Dreamina and Freepik; per-second cost fell below \$0.14 by March 2026.
- **Higgsfield / Higgsfield Sketch-to-Video / Higgsfield Popcorn / Higgsfield Relight / Higgsfield Shots / WAN Camera Control / Higgsfield Adobe Premiere Plugin / Higgsfield After Effects Plugin** — social-media-marketer video platform; \$80M raised at \$1.3B valuation, \$200M revenue in nine months; Shots produces multiple storyboard images from a single shot; the Adobe Premiere Pro and After Effects plugins (June 2026) bring Higgsfield generation directly into the professional editorial workflow.
- **Bernini** (ByteDance) — AI video generation and editing framework; a complementary video-editing layer to ByteDance’s Seedance generation stack.
- **LTX-2 / LTX-2.3 / LTX-2.3 Colorizer / LTX-HDR / LTX Studio / LTX-2 Audio-to-Video / LTX-2 Lip Sync / LTX-2 Real-Time** — open-source video generation with audio sync; LTX-2.3 is high-resolution, fast, cinematic with native lip-sync; LTX HDR (beta) ships HDR processing.
- **Odyssey 2** — real-time interactive video generation.
- **Vidi 2** (ByteDance) — multimodal video understanding and creation.
- **MotionStream** — real-time interactive video with mouse-based motion control.

- **Blockvid** — one-minute video with improved structure and visual consistency.
- **Video Rebirth** (Singapore) — studio-grade AI video platform; \$50M raise.
- **LiveGS** — mobile Gaussian-splatting video.
- **Time-To-Move** — motion control for generated video.
- **Ovi / Ovi 1.1** (Character.AI) — open-source video with speech sync.
- **CraftStory** — image-to-video for long-form AI video with human “actors”.
- **Decart Lucy 2.0** — realtime world-editing video model; 1080p at 30 fps.
- **Decart LSD v2** — real-time video-to-video with prompt-on-the-fly.
- **Xmax X1** — the first real-time interactive video model.
- **FastVideo** — 30-second 1080p generation at 4.5 s latency.
- **HiAR** (Tencent) — scalable long-video generation.
- **Alli AI** — system for creating and editing 8K video up to 60 seconds.
- **AERA AI TV** — automated storyboard generation, editing and serialized video production.
- **AnchorWeave** — world-consistent video generation with retrieved local spatial memories.
- **CubeComposer** — generates native 4K 360° video from standard perspective footage.
- **MatAnyone2** — high-fidelity video matting with finer detail.
- **VFace** — training-free video face-swapping for any diffusion model.
- **DreamActor M2.0** (fal) — drive characters from a single image + template video; multi-character supported.
- **ByteDance ALIVE** — unified audio-video generation.
- **Magnific Upscaler for Video / Magnific Precision v2 / Magnific Precision for Video** — upscaling and 4K detail enhancement, including dedicated video upscaling.
- **NetFlix VOID** — object removal from video with physics-interaction removal.
- **Ponder** — agentic video editor.
- **ArcReel** — multi-agent video generation from written stories.
- **Scope** — real-time interactive generative AI pipelines (LTX-2.3 real-time runs on Scope).
- **NotebookLM Cinematic Video Overviews** (Google) — video generation from notebooks.
- **Omnia** — AI-native browser video editor.
- **NVIDIA RTX Video Super Resolution** — upscaling node with ComfyUI integration.
- **Google M2SVid** — monocular-to-stereo video conversion.
- **Alibaba WonderClip** — all-in-one AI video creation platform from Alibaba Cloud (June 2026).
- **Sony Vid-CamEdit** — post-capture camera trajectory editing; keeps scenes spatially coherent across large viewpoint changes (June 2026).
- **Nvidia PiD (Pixel Diffusion Decoder)** — image / video decoder; new entry from Nvidia (June 2026).

AI image models / tools

- **Midjourney / Midjourney V8.1** — the aesthetic-leadership product; Discord/X-native; V8.1 ships native 2K HD rendering at 3× the speed and 3× the cost reduction.
- **FLUX / FLUX 2 / FLUX 2 Max / FLUX.2 [klein]** (Black Forest Labs) — open-weight, fine-control, the open-source default through 2025–26; klein is the 4B-parameter lightweight model.
- **Adobe Firefly / Firefly Image Model 5 / Firefly Foundry / Firefly Boards / Firefly Precision Flow** — Image Model 5, Foundry (custom corporate training), Firefly Boards (moodboards), Precision Flow (granular AI editing control, beta); 45% of Creative Cloud users active, 22B+ assets generated by April 2025.
- **Imagen 3 / Nano Banana / Nano Banana Pro / Nano Banana 2** (Google) — most-integrated image model in the consumer toolchain; Photoshop and Unreal Engine plugins; Nano Banana Pro ships professional capabilities at lightning speed.
- **Stable Diffusion / Stable Diffusion 3 / Stable-Layers** (Stability AI) — the foundational open-source image model; **Stable-Layers** (June 2026) splits generated images into editable compositional layers — a structural shift from regeneration to post-generation editing.
- **Ideogram 4.0** — fourth-generation image model; open weights, fine-tuneable on own data, runs on own hardware; available on every Ideogram plan and via API.
- **Krea / Krea AI / Krea 2 / Krea Realtime / Krea Realtime Edit / Krea Nodes / Krea LoRA Trainers** — real-time AI image generation, now open-source; Realtime Edit takes complex instructions in real time; the LoRA Trainers cover Qwen-2512 and Z-Image.
- **Qwen-Image-2512 / Qwen-Image-Edit-2511 / Qwen 2511 Time Travel / Qwen-Edit 2509** — the dominant open-source image-editing model family; constituent LoRAs (relighting, multi-angle, time-travel, AnyPose) discussed in the LoRAs section below.
- **ChatGPT Images / ChatGPT Images 2.0** — image generation integrated with Adobe Express, Photoshop and Acrobat; 2.0 ships thinking-level intelligence.
- **Grok Imagine / Grok Imagine API (xAI)** — image and video generation; the API bundles end-to-end creative workflows; reference-to-video and video extend added in March 2026.
- **Vision Banana** — unified model for image understanding and generation.
- **Freepik / Freepik Spaces / Freepik Speak / Freepik 3D Scenes** — design platform; Spaces is the real-time collaborative canvas; Speak is the lip-sync talking-video tool; 3D Scenes generates full environments from an image.
- **Recraft** — brand-focused AI design.
- **Weavy** — node-based AI creative platform; acquired by Figma.
- **Canva** — design platform with “Creative Operating System.”
- **Civitai / Civision** — community-models and LoRAs platform; Civision is the free Civitai/Hugging-Face hybrid alternative.
- **Replicate** — model-hosting marketplace and API.
- **GFPGAN** — face restoration.

- **UpscalyAI / Upscayl / Topaz Gigapixel AI / Real-ESRGAN / Clarity Pro Upscaler** — upscaling tools; Clarity Pro reaches 10K.
- **PractiLight** — practical light control via diffusion.
- **PercHead** — 3D head reconstruction from single images.
- **Loveart AI** — layer separation and editable text (Live Editable Text, LET).
- **NVIDIA ChronoEdit** — temporal reasoning for image editing.
- **Beeble Switchlight 3 / Beeble Background Remover** — AI masking and relighting.
- **CanvAI** — turns rough sketches into polished AI art.
- **MakeComics / Make Comics** — generates custom comics with characters and storylines in seconds; from prompt to full comic book.
- **Seethrough** — illustration decomposition into layered PSD.
- **Best Face Swap (Flux 2 LoRA)** — specialised face-swap.
- **PixelSmile** — fine-grained facial-expression editing across 12 expressions.
- **Earth Cinema** — Chrome extension using Google Earth for cinematic images.
- **360Anything** — lifts any perspective image or video to a 360° panorama.

AI music / audio tools

- **Suno / Suno Studio / Suno 5.5** — the dominant prompt-to-track generative music platform; \$400M raised at \$5.4B valuation (June 2026), \$300M ARR, 2M+ subscribers; Warner Music Group settled its lawsuit and signed a licensing partnership; UMG and Sony remain in active litigation; Suno Studio is the “world’s first generative audio workstation”; 5.5 adds Voices.
- **Udio** — prompt-to-music; partnered with Universal Music Group; indie-label licensing via Merlin.
- **Mureka / Music Agent Studio** — six specialised AI agents covering songwriting, arrangement and production.
- **ElevenLabs / ElevenLabs v3 / Scribe v2 Realtime / ElevenMusic / ElevenCreative / ElevenLabs Flows / ElevenAgents Expressive Mode / ElevenLabs Voice Changer / ElevenLabs Music v2 / ElevenLabs Dubbing v2** — the dominant voice/audio synthesis stack; \$500M ARR; BlackRock + NVIDIA backed; Scribe v2 transcribes in 150 ms across 90+ languages; ElevenMusic is the discovery/creation/earning marketplace; Flows is the node-based creative canvas; **Music v2** (June 2026) is the next-generation generative-music model; **Dubbing v2** (June 2026) is the next-generation dubbing model. ElevenLabs also partnered with Hasbro’s Sixth Wall AI studio in June 2026, providing the voice layer for Hasbro’s Behavioral Licensing / CharacterOS programme — enabling licensed AI voices for Transformers, Mr. Potato Head, Peppa Pig and more.
- **iZotope Ozone 12 / Stem EQ** — AI-assisted mixing and mastering.
- **LANDR** — AI mastering and distribution.
- **Riffusion** — spectrogram-based music generation.
- **Stable Audio 2.5 / Stable Audio 3.0** (Stability AI) — generative audio. **Stable Audio 3.0** (May–June 2026) is an open-weight model family trained on fully licensed data,

designed as the foundation for what the open audio community builds next; ships in ComfyUI for tracks up to about six minutes.

- **MusicGPT** — local music generation.
- **ACE Studio / ACE Studio 2.0 (TIMEDOMAIN) / ACE Studio Video-to-Music / ACE-Step 1.5 / ACE-Step 1.5 XL** — all-in-one AI music studio; video-to-music for visuals; ACE-Step generates full songs in under 10 seconds on <4 GB VRAM.
- **Music Lens** (Musixmatch) — catalog-intelligence agent.
- **Melosurf** — voice-controlled assistant for Ableton Live.
- **Songscription** — audio-to-notation transcription; \$5M raise.
- **Groundhog Audio Pedal / Groundhog OnePedal** — AI guitar tone matching.
- **Spotify DJ / Spotify AI / Spotify AI Prompted Playlists / Spotify AI Transparency Beta / Personal Podcasts on Spotify** — playlist personalisation, voluntary AI disclosure beta, agent-created podcasts.
- **YouTube Music AI Hub** — AI music hosting and DJ features.
- **Epidemic Sound Studio** — AI music for video creators.
- **AIODE** — ethically-trained music creation DAW.
- **Overtune** — virtual music studio for Roblox.
- **Space DJ** (Google DeepMind) — interactive music exploration.
- **Minimax Music Generator** — song generation.
- **Roland AI Pedal** — concept for AI audio processing.
- **Roland + Sony CSL Melody Flip** — AI music tool.
- **BandLab Voice Cleaner / BandLab Palette** — background-noise removal; AI loop-matching.
- **Claimy** — missing-royalty recovery agent.
- **Fish Audio S1 / Fish Audio S2** — text-to-speech; S2 ships expressive TTS with emotional control (6× cheaper than ElevenLabs at S1).
- **Tempolor Guitars** (Quwan) — AI to make songs playable on guitar.
- **GEMIDI** — music playground powered by Gemini.
- **Lalal AI / Lalal AI plugin / Lalal AI API v1** — AI stem separation, music removal, voice changing; first stem-separation plugin native to a DAW.
- **StemDeck** — local stem separation; free, no upload required.
- **Mirelo** (Berlin) — automatic sound-effect generation for video.
- **PromptSep** — separates any sound by text description.
- **SAMAudio** (Meta) — Segment Anything for audio editing.
- **SonicLab SPATAI** — generative audio engine for immersive spatial production.
- **Apple Logic Pro AI** — synth player and personal music-theory expert; chord-progression generation.
- **Apple GarageBand** — consumer music creation.
- **Musicful AI** — music-video generation for AI-generated songs.
- **Latent Space Explorer** — interactive audio data visualiser using Stable Audio VAE.
- **Voicebox** — open-source local voice cloning from a 3-second audio clip.

- **Voice-Pro** — AI speech recognition / translation / transcription / multilingual dubbing.
- **PersonaPlex** (NVIDIA) — full-duplex model that listens and speaks simultaneously; open-source natural-sounding voice.
- **Phoenix-4** — advanced real-time human-rendering model with 10+ emotional states.
- **NVIDIA D-Rex** — photorealistic digital humans under any lighting.
- **Qwen3-TTS / VoiceDesign / VoiceClone** — text-to-speech with voice design and cloning.
- **Greysound** — self-engineering music-production studio.
- **TAC (Timestamped Audio Captioning)** — timestamped captions for audio and audiovisual content.
- **Insanely Fast Whisper** — 2.5 hours of audio transcribed in 98 seconds.
- **smol-audio** — local audio model notebooks and scripts.
- **AudioStream / Walzersymphonie** (Ars Electronica Futurelab) — AI composition system for classical music using Ricercar.
- **Ricercar** (Ars Electronica) — creative AI for artistic composition.
- **Krotos Video-to-Sound** — expanded platform for audio professionals (foley, sound design).
- **Jamu** — AI co-producer agent for Ableton Live.
- **Moises** — AI music platform; Charlie Puth as Chief Music Officer.
- **Music Mogul AI** — tour-booking automation.
- **Sonilo + Shutterstock** — video-to-music AI training deal.
- **Clearnote** — AI music-contract platform to end deal delays.
- **Sony music identification tech** — identify original music inside AI-generated songs.
- **Cactus Music** — artist-ops support.
- **Rebel Audio** — AI podcast startup.
- **ROLI Airwave & AI Music Coach** — hand-tracking (27 joints) with real-time AI piano practice coaching.
- **Sony AI ICASSP papers** — music understanding and generative audio research papers (April 2026).

3D, world models and spatial

- **Marble / Marble 1.1 / WorldLabs API / RTFM** (World Labs / Fei-Fei Li) — first commercial generative world model; 40× faster than legacy VP workflow; 1.1 adds real-world location 3D reconstruction and restyling; the API generates persistent 3D worlds from text, images and video; RTFM is the real-time frame model.
- **Genie 3 / Project Genie** (Google DeepMind) — research-grade world model; *Time Best Inventions* 2025; rolled out to Google AI Ultra in January 2026.
- **WorldGen / WorldMirror** (Meta / Tencent) — interactive 3D world generation.
- **UNI-1 / UNI-1.1** (Luma) — world generation + reasoning combined; 1.1 adds prompt enhancement and built-in research.

- **Hunyuan 3D / Hunyuan 3D Studio / Hunyuan 3D-PolyGen 1.5 / HY3D 3.0 / HY-Motion 1.0** (Tencent) — art-grade 3D generative model family; PolyGen 1.5 is the art-grade text/image/sketch-to-3D model; HY-Motion 1.0 is text-to-3D human motion.
- **HY World 1.5 / HY-World 2.0 / WorldCompass / ShotVerse** (Tencent) — open-source real-time world-model frameworks; WorldCompass is the RL post-training framework; ShotVerse handles cinematic multi-shot camera control.
- **ECHO / Echo-2** (SpAItial) — spatial foundation model; Echo-2 adds world decomposition.
- **Wonderzoom** (Stanford) — multi-scale 3D world generation with infinite zoom.
- **Matrix-Game 3.0** (Skywork AI) — real-time interactive world model; 720p at 40 fps with long-horizon memory.
- **AMD Micro-World** — AMD’s entry into world models.
- **LingBot-World** — free/open-source world model on Wan 2.2; real-time interaction at 16 fps.
- **Moonlake** — \$28M-seeded world model platform now in beta for games and simulations.
- **Project Eden** (TripoAI) — persistent, multiplayer world model (June 2026) that fundamentally breaks from existing paradigms by *decoupling the underlying world state from visual rendering*; the most ambitious state-vs-rendering separation in the world-model layer at the time of writing. TripoAI raised nearly \$200M in June 2026 to advance this work.
- **TripoSplat** (Tripo) — single 2D image to high-quality 3D Gaussians; open-source under the MIT license.
- **NVIDIA Cosmos 3** — omnimodal world model (released May 31, 2026) that connects understanding, generation, simulation and action through a shared architecture moving fluidly across text, images, video, audio and actions; Two-Tower Mixture-of-Transformers design; ranked best open-source T2I and I2V models by Artificial Analysis; explicitly targets robotics, autonomous vehicles and warehouse monitoring. The most significant NVIDIA world-model release in the period.
- **NVIDIA FlashDreams** — open-source world model from NVIDIA, released June 2026.
- **PaGeR** (Google & Meta) — AI model for 360° scene geometry reconstruction.
- **NVIDIA Lyra 2.0** — explorable generative 3D worlds.
- **NVPanoptix-3D** (NVIDIA) — single-image 3D indoor-scene reconstruction.
- **Kimodo** (NVIDIA) — kinematic motion-diffusion model trained on 700 hours of mocap.
- **InSpatio-WorldFM** — open-source real-time generative frame model.
- **Code2Worlds** — workflow for generating 4D scenes with environmental/object generation and feedback refinement.
- **OpenArt Worlds** — 3D navigable environments from a single prompt.
- **Rodin / Hyper3D / Rodin Hyper 3D Gen 2** — high-precision 3D model generation.
- **Meshy / Meshy 6 / Meshy Image-to-3D** — 3D model generation; Meshy 6 ships natively inside ComfyUI; Image-to-3D generates poses.
- **Hitem3D** — high-resolution 3D generation.
- **Microsoft Trellis 2** — native compact structured latents for 3D.

- **ByteDance Seed3D 2.0** — 3D object generation from image or text.
- **Pixel3D** (Tencent) — 3D object generation.
- **Tripo 3.1** — 3D-asset creation; ComfyUI partner nodes.
- **PATINA** (fal) — image or text to full PBR material maps.
- **M-XR PBR Model** — 4K PBR material maps for 3D assets.
- **CHORD** (Ubisoft La Forge) — open-sourced end-to-end PBR-material generation.
- **Autodesk Wonder 3D** — generative AI tool creating editable 3D assets from text or images.
- **Autodesk Flow Studio Rigging** — AI rigging and neural layers.
- **AI4AnimationPy** (Meta) — pure-Python 3D character animation.
- **MocapAnything** (Huawei) — unified 3D motion capture for arbitrary skeletons from monocular video.
- **FreeMoCAP** — open-source markerless mocap from any camera.
- **MuJoCo** — physics engine for robotics and animation.
- **Mesh2Motion** — auto-assigns and exports 3D model animations.
- **One-to-All Animation** — alignment-free character animation and image pose transfer.
- **Cascadeur / Cascadeur AI Animation v2025.3** — AI animation tool for character movement; local interpolation.
- **Move AI** — motion capture.
- **MoCap / UE5 Motion** — camera-only motion capture.
- **SCAIL** — studio-grade character animation via in-context learning.
- **Vanast** — garment-transferred human animation from images.
- **CompHairHead** — one-shot 3D head avatars with deformable hair.
- **Reallusion Headshot 3** — digital-double creation.
- **MoRo** (Meta) — human motion recovery via masked modelling for occlusions.
- **Meta Sapiens 2** — pose estimation, body-part segmentation and surface normals.
- **NVIDIA Audio2Face / Audio2Face-3D** — facial animation from audio.
- **NVIDIA UniRelight** — shot relighting technology.
- **Apple SHARP / LiTo** — photorealistic 3D view synthesis from a single image in under a second; LiTo is the surface light-field tokenisation method.
- **Pixel3DMM** — screen-space single-image 3D face reconstruction.
- **VISTA4D** (Netflix + Eyeline) — live-action to navigable 4D point clouds.
- **SuperSplat / SuperSplat v2.16** (PlayCanvas) — free, open-source Gaussian-splat editor.
- **Spark 2.0** — open-source Gaussian-splat streaming framework; streamable LoD for WebGL2.
- **PlayCanvas SOG / SplatTransform v2.0.0** — WebP-equivalent compression for Gaussian splats; SplatTransform v2 ships automatic high-quality collisions for splats.
- **Cesium / CesiumJS / Cesium for Unreal** — geospatial 3D platform; Cesium agentic workflows let natural-language commands interact with 3D geospatial data; supports Gaussian splats natively.

- **SPAG-4D** — 360° photos to 3D Gaussian splat conversion.
- **Animated Gaussian Splats** (4DV.ai) — 4D Gaussian-splat animation.
- **Open3Dmap** — crowdsourced 3D mapping with Gaussian splats.
- **Hyperscape / Hyperscape Capture** (Meta) — Gaussian-splat capture on Quest.
- **Apple Vision Pro Personas** — Gaussian-splatting consumer feature.
- **Common Sense Machines** — converts 2D images into 3D digital assets (Google acquisition).
- **AI-Enhanced LiDAR** — generative-AI-augmented LiDAR capture pipeline (June 2026) for scanning and reconstruction work in spatial production.
- **NVIDIA Omniverse Fixer** — rendering-artefact removal for Gaussians.
- **DecartAI / Decart LSD v2 / Decart Lucy 2.0** — real-time world transformation by voice; LSD v2 is real-time video-to-video; Lucy 2.0 is realtime world editing at 1080p/30fps.
- **VoxeloAI** — 3D creation for e-commerce.
- **Mosaic / Mosaic 3D** — 3D reconstruction.
- **Depth Anything 3** (ByteDance) — visual-space recovery from any view.
- **SAM 3 / SAM 3D / SAM 3.1** (Meta) — object detection and 3D segmentation; SAM 3.1 is 7× faster, tracking 128 objects on a single H100.
- **Seen2Scene** — realistic 3D scene completion with visibility-guided flow.
- **PixARMesh** — reconstructs full indoor scenes from a single photo into lightweight meshes.
- **Confidence-Based Mesh Extraction from 3D Gaussians** — algorithm for high-quality mesh extraction from Gaussian splats.
- **Rhizomatics Sword Tip Visualization System** — fencing tracking at 60 fps.
- **Mixar** — AI-native Blender-fork editor positioned as “Cursor for 3D artists”.
- **Sprite Booth** — pixel-character animation tool (ComfyUI API required).
- **AutoSprite** — character-to-animated-sprite-sheet conversion.

Voice, avatar, digital human

- **Synthesia** — AI avatar platform; \$4B valuation; rejected \$3B Adobe offer.
- **Heygen / Heygen Video Agent / Heygen Motion Designer / LiveAvatar / Heygen Elements / Heygen CLI / HyperFrames** — AI-avatar and end-to-end video generation; LiveAvatar is hyper-realistic real-time interactive; Elements builds scenes from reusable components; CLI ships videos from the command line; **HyperFrames** (June 2026) is an open-source framework for turning HTML, CSS, media and seekable animations into deterministic MP4 videos — bringing deterministic video export to web-stack creative workflows.
- **Hedra / Hedra Audio Tags / Hedra Elements** — digital-human creation; Audio Tags assign precise emotions to audio.
- **Live Avatar** (Alibaba) — streaming real-time audio-driven avatar generation with infinite length.

- **Avatar Forcing** — real-time interactive head-avatar generation for natural conversation.
- **PersonaLive** — real-time expressive portrait animation.
- **Weclone** — digital avatar built from a user's chat history.
- **Particle6 / Tilly Norwood** — AI-performer studio and synthetic actress.
- **Xania Monet** — AI music artist; Billboard chart entries; \$3M Hallwood Media deal.
- **Bleeding Verse** — AI band; Hallwood Media signing.
- **Sienna Rose** — anonymous AI artist with millions of Spotify streams.
- **Breaking Rust** — AI country artist; #1 country digital song sales.
- **Trilok** — Indian AI band.
- **Copresence / ConvAI / Inworld** — avatar and NPC-character technology for Unreal Engine, games and interactive media.
- **ROXi** — AI-generated TV presenters.
- **Facelooker** — interactive headshot generation.
- **FellinAI** — AI film director.
- **ego AI** — in-game character platform.
- **Character.AI** — play characters in your favourite books; Ovi open-source video with speech sync.

Agent platforms and orchestration

- **OpenAI AgentKit / Agent Builder / ChatKit / Connector Registry / Eval Framework** — the developer-facing agent platform underneath most third-party agentic creative tools.
- **Anthropic Claude Apps / Claude Skills / Claude Code / Claude Design / Claude for Legal / Anthropic Academy** — interactive Claude in workplace tools; the Skills framework powers Sony's 49-agent / 72-skill game-development stack; Claude Design ships prototypes/slides/one-pagers; Anthropic Academy is free Claude Code training.
- **Claude Code Game Studios** — 49-agent, 72-skill coordinated AI game-development team.
- **Gemini API Agents / Google Antigravity / Opal / Fabula / Google Stitch** (Google) — agentic capability surface across Google's stack; Antigravity is the agentic development platform; Opal is the no-code mini-app builder; Fabula is the interactive AI writing tool; Stitch generates mobile/web UI.
- **Heygen Video Agent** — end-to-end video-assembly agent.
- **Adobe CX Enterprise / GenStudio / Adobe Film & TV Fund / Adobe Ignite Day** — agentic creative intelligence across the full content lifecycle; GenStudio is the brand-intelligence system for personalised content at scale.
- **NVIDIA + Google Cloud / Avid + Google Cloud / Speechmatics in Adobe Premiere** — agentic-creative infrastructure partnerships; on-device speech-to-text inside Premiere.
- **n8n** — workflow automation with AI.

- **ComfyUI / ComfyUI Cloud / ComfyHub / ComfyUI App Mode / ComfyUI Simple Mode** — node-based AI workflow editor; \$500M valuation by May 2026; App Mode turns workflows into shareable no-install apps; Simple Mode lowers the complexity threshold.
- **Replicate** — model marketplace.
- **Hugging Face** — open-source model hub and agentic platform; partnered with Google Cloud and Meta.
- **Meta OpenEnv** — open-source agentic-development environment.
- **fal / FAL MCP** — generative-model deployment marketplace; FAL MCP exposes 1,000+ generative models to Claude/Cursor.
- **Unity Official MCP / Unity MCP** — Unity Editor control via MCP protocol.
- **Blender MCP / Blender Buddy / OpenClaw + Blender MCP** — Claude-to-Blender control for 3D workflows.
- **VibeComfy** — open-source tools letting agents understand, build and run ComfyUI workflows with Claude.
- **MooshieUI** — beginner-friendly ComfyUI desktop frontend.
- **Lenny (Maroofy)** — AI agent for live-music event organisers.
- **AdsGency** — autonomous paid-marketing agent platform; \$12M seed.
- **Invideo AI / Invideo AI Agent** — full video-ad generation from prompt; built on Gemini.
- **Pomelli (Google Labs)** — AI marketing agent for small business.
- **Inception Point AI** — podcast hosting with AI hosts (3,000+ episodes/week).
- **Slapshot** — AI camera tracking.
- **Photoroom** — AI photobooth.
- **Jabali.ai** — AI game-development platform (Sony-backed).
- **Fifth Door** — AI-powered game creation; \$20M raise.
- **Atlas AI Agents** — game-production pipeline agents.
- **Ad Conceptor (Runway)** — ad-concept exploration agent.
- **AE GPT** — AI assistant for After Effects (expressions, scripts).
- **Career-Ops** — AI job-search pipeline built with Claude Code.
- **Auto-Dream (MyClawAI)** — personal AI assistant wrapper.
- **Clicky** — AI screen-aware tutor/buddy interface.
- **MemPalace** — high-scoring AI memory system.
- **TikTok AI Agents for Ads** — TikTok's automated ad-creation agents.
- **Amazon Creative Agent** — agentic AI generating professional-quality ads.
- **Playad** — AI marketing agents for ad creative.
- **Moltbook (Meta acquisition)** — AI agent social network.

Legacy creative software, AI-augmented

- **Adobe Creative Cloud** — Photoshop, Premiere Pro, After Effects, Illustrator, Express, Acrobat, plus the Express AI Assistant, the Premiere Object Mask, the Photoshop

generative-fill, selection and Rotate Object tools, the Photoshop AI Assistant public beta, and Illustrator Turntable.

- **Adobe Sneaks 2025–26** — Adobe’s research-preview portfolio shown at MAX 2025 and Summit 2026:
 - **Project Scene It** — image-to-3D and 3D-to-image with object tagging.
 - **Project Surface Swap** — AI-powered texture recognition and swap.
 - **Project Turn Style** — edit and transform 2D objects as if they were 3D.
 - **Project Trace Erase** — removes objects, shadows, reflections and distortions.
 - **Project New Depths** — edit depth like brightness.
 - **Project Frame Forward** — apply changes across video from a single frame plus text.
 - **Project Motion Map** — bring static vector graphics to life.
 - **Project Sound Stager** — analyses video visuals and generates synchronised soundscapes.
 - **Project Clean Take** — corrects mispronunciations and isolates voices.
 - **Project Light Touch** — spatial lighting mode for relighting photos.
 - **Project Graph** — node-based workflow editor in the ComfyUI / Weavy lineage.
 - **Corrective AI** — changes the emotional register of voice-overs.
 - **Edit-by-Track** (Adobe Research) — generative video motion editing with 3D point tracks.
- **Adobe Acrobat** — AI converting PDFs into podcasts.
- **Unreal Engine 5 / UE5 AI Assistant / UE5 AI Motion Plugin / Unreal to Gaussian Splat / Prompt-to-Player** (Epic) — official AI assistant; camera-only motion capture; UE5 scene-to-splat conversion; Prompt-to-Player generates 3D characters and auto-rigs/animations them.
- **Unity / Unity AI Open Beta / Unity AI Council / Unity AI Tools Suite / Unity Prompt-a-Game** — in-editor AI suite for the full games pipeline; Prompt-a-Game demoed at GDC March 2026.
- **Autodesk Maya / Autodesk Flow Studio / Autodesk Wonder 3D / Autodesk Flow Studio Rigging** — motion capture, 3D animation, generative 3D and neural rigging.
- **Foundry / Nuke / Nano Banana × Nuke** — AI-augmented VFX; Weta FX / AWS collaboration; the Nuke node connecting Nano Banana through fal’s API.
- **SideFX Houdini** — procedural modelling with AI integration.
- **Blender / Blender MCP / Blender Buddy** — open-source 3D; Anthropic Foundation patronage; MCP enables Claude control; NVIDIA DLSS 4.5 Ray Reconstruction added to Blender 5.3 as a viewport denoiser.
- **Vlo** — free, local, open-source video editor with ComfyUI-backed generative AI features; the open-source equivalent of the AI-native NLE.
- **Motionfly Editor** — AI-native video editor positioned as “Cursor for video editing” — bringing the code-editor / AI-pair model into the video-editing workflow.
- **DaVinci Resolve** (Blackmagic) — colour and editorial with AI features.

- **Avid Media Composer / Avid + Google Cloud** — professional editorial; Pro Tools agentic AI in the Google Cloud partnership.
- **Pro Tools** — audio editorial with AI features; Claude integration in Ableton.
- **Logic Pro / Apple Logic Pro AI** — music production with AI; synth player and music-theory expert.
- **Ableton Live / Claude in Ableton / Jamu / Melosurf** — music production with Claude assistance, Jamu co-producer agent and Melosurf voice control.
- **GarageBand** — consumer music creation.
- **CapCut / Dreamina / Snapchat AI Clips in Lens Studio** (ByteDance / Snapchat) — mobile video editing with AI; Snapchat AI Clips in Lens Studio.
- **Final Cut Pro** — editorial with FX Factory plugins.
- **Keyframe** — AI models inside After Effects.
- **Tether** — AI animation inside After Effects.
- **Deep Pan** — animated stills in Premiere Pro.
- **FX Factory** — AI plugin suite for Premiere and Final Cut.
- **VEED Transitions** — AI transition generation.
- **Electric Sheep** — web-based projection mapping and AI visual orchestration.

AI-native creative studios and apps

- **Imaginae Studios** (Fremantle) — AI-native studio; *Art Awakens* project under Google AI Futures Fund.
- **Wonder Studios** — AI-native film studio; \$12M raise; Wonder Film Festival.
- **Asteria** (Natasha Lyonne / Lightstorm / Topaz Video) — AI animated short *All Heart*; hybrid AI workflows with Topaz.
- **Promise** (Google-backed) — GenAI filmmaking and VFX for legacy media.
- **Obsidian Studio** (Imagine Entertainment / Ron Howard, Brian Grazer) — AI production partner.
- **Goldfinch enGEN3** — AI cinematic universe platform.
- **Chapter41** (Beta Film / Munich) — AI startup.
- **Kartel** (Kevin Reilly) — AI startup; ex-HBO leadership.
- **Holywater** (Fox Entertainment investment) — AI microdramas.
- **CenterStage** (David Boies / Zack Schiller) — interactive AI entertainment.
- **Superlunar + Eden Studios** — *Little Plastics* dystopian AI film.
- **Particle6** (Eline Van der Velden) — AI-performers studio.
- **Imago Pictures** — 3D + ComfyUI workflows.
- **LKDN AI Production** — node-based AI filmmaking.
- **Deep Fusion Films** (Cardiff) — AI documentary production; acquired by John Gore Studios.
- **Palo Creators** (MrBeast alumni) — AI for viral video creation.
- **Refik Anadol Studio Dataland** — first museum of AI art.
- **Animaj** — Google AI Futures Fund partner for kids' content.

- **Gossip Goblin** — AI filmmaking studio.
- **Critterz** (Vertigo + Federation) — AI-assisted animated-feature production.
- **Filmology Labs** — \$250M studio for vertical micro-dramas and AI-driven media formats.
- **escapeAI** (John Gaeta) — streaming apps for AI-generated content across Roku, Fire TV, Samsung and LG.
- **Primordial Soup** (Darren Aronofsky) — AI production studio.
- **Toonstar** — AI animation house.
- **CinemersiveAbout Labs / Cinemersive Labs** — UK machine-learning and computer-vision company; acquired by Sony.
- **My SMASH Media** — AI film startup with Allison Gardner (ex-Glasgow Film Festival).
- **Framestore AI Platform / Framestore Futon** — ML and GenAI in the VFX pipeline.
- **Mito AI** — \$4.5M-raised platform empowering video professionals.
- **Filmax DinoGames** — AI-driven 3D animated film project.
- **Abundantia Entertainment + InVideo** — AI film studio launched in India.
- **Luma Production Studio + Wonder Project** — feature collaboration.
- **Lionsgate Chief AI Officer (Kathleen Grace)** — first major-studio CAIO hire.
- **CreateAI** — AI animation production.
- **Netflix Ben Affleck AI Studio Acquisition** — AI filmmaking company acquisition.
- **Holywater** (Fox investment) — AI microdrama studio (also above).
- **DreamLab AI Collective** (the studio that produced this book) — AI-augmented creative practice across film, games and immersive.

Games-development AI

- **NitroGen** (NVIDIA + Stanford) — plays-any-game model trained on 40,000 hours across 1,000+ games.
- **SIMA 2** (Google DeepMind) — agent that plays, reasons and learns in 3D virtual worlds.
- **CHORD** (Ubisoft La Forge) — open-sourced end-to-end PBR material generation.
- **Ubisoft Teammates** — voice AI for in-game team communication.
- **YouTube Playables Builder** (Gemini 3) — text-to-game prompt-to-playable web app.
- **Roblox AI Tools / Roblox AI Assistant / Roblox Reality / Metain** — creator-facing AI for game development; Reality is Roblox's productised AI assistant; Metain is the prompt-to-Roblox-Studio interface. In June 2026, Roblox acquired **Morpheus AI Inc.**, an AI lab building video world models using *Self Forcing* (converting offline video models into fast, autoregressive, interactive engines), to accelerate the Roblox Reality video-model vision.
- **Epic xAI Studio** — Elon-Musk-affiliated game-dev AI.
- **General Intuition** — spatial-reasoning research lab.
- **ReBlink ARBO** — AI-powered strategy battler.
- **Halo Studios** — AI woven into the production pipeline.
- **Dabgg** — AI gaming with interactive platform.

- **RosebudAI / Spawn / LudoAI** — vibe-coding games and interactive apps from prompts; Spawn is the “games made with words” community platform; LudoAI is the ideation and planning scaffold.
- **Verse8** — AI-driven game-creation platform.
- **Sett** — AI agents for game marketing.
- **Meta Horizon Studio AI Assistant** — Horizon Worlds creator AI.
- **Gambo** — vibe-coding game agent.
- **Project Motoko** (Razer) — AI-powered gaming headset.
- **Bethesda AI Toolset** — Todd Howard’s game-dev AI integration.
- **StarBerry Games / Merge Mayor** — AI-all-in game studio.
- **Astrocade** — interactive entertainment platform; \$56M Series B.
- **GameByte** — AI-powered game creation platform; \$1M raise.
- **Sorcerox Game Creation Suite** — prompt-driven game creation toolkit (June 2026); positioned as a “*build any game you can imagine*” general-purpose creation surface.
- **Studio Atelico Bobium Brawlers** — AI-based iOS game with pro-human approach.
- **Bitmagic** — AI-created version of classic Civilization.
- **Helika AI Publishing Engine** — AI publishing engine for game studios.
- **Xbox Game Studios “git gud” AI** — Xbox patent for AI tech that plays games for you.
- **Gaming Copilot** (Xbox) — console AI copilot.
- **NVIDIA DLSS 5** — gaming upscaling.
- **PS5 Pro PSSR** — PlayStation upscaling.
- **Gizmo** (Meta acquisition) — vibe-coding app for creating mini-games.
- **Voyage** (Latitude / AI Dungeon) — AI RPG platform.
- **DoomGen** — prompt-driven DOOM mod prototyping app.
- **Sony AI tools for PlayStation** — animation support, QA automation, asset generation and performance-capture-to-facial-models pipelines.
- **PROWL** (Odyssey) — RL agents for game-world model testing.
- **MagicDawn** (Tencent) — AI lighting tech.
- **ReActor** (Disney) — RL motion-transfer method.
- **Origin Lab** — game-world AI training data; \$8M raise.
- **Daughter of the Inner Stars** — Unreal Engine 5 AI-driven game.
- **Genesis AI** — robot cooking and piano playing systems (demonstrating embodied agents).
- **Quilty** — AI platform for script development and assessment.

Marketing and advertising AI

- **WPP Open Pro** — agency-wide AI marketing platform.
- **Sightly** — AI advertising and brand performance.
- **Pomelli** (Google Labs) — small-business marketing agent.
- **AdCreative.ai** — mobile pro ad creation.

- **Epiminds** — AI marketing team builder.
- **Octave AI** — AI podcast/radio ad creation.
- **AdsGency** — autonomous paid-marketing agent.
- **Aimy** — AI TV advertising on the Comcast API.
- **Brand Networks** — AI television advertising.
- **Fullthrottle.ai** — automotive-focused DSP.
- **Channel 4 AI ads** — AI-driven TV-advertising tool for SMEs.
- **Sky Studio Challenge** — interactive AI ad creation via Starlink.
- **Google Demand Gen** — AI-powered image tools for the ad platform.
- **Amazon DSP** — AI one-stop-shop advertising infrastructure.
- **Amazon Creative Agent** — agentic ad creator.
- **Playad** — AI marketing agents.
- **Runway Ad Concepter App** — ad-concept and composition exploration.
- **TikTok AI Agents for Ads** — TikTok ad-creation agents.
- **Instagram Edits** — AI video generation for advertisers.

Open-source ecosystem and infrastructure

- **ComfyUI** — node-based AI workflow editor; the operating system of the open-source creative AI ecosystem; \$17M raise (October 2025), \$500M valuation (May 2026).
- **Hugging Face** — open-source model hub; partnerships with Google Cloud and Meta.
- **Meta OpenEnv** — open-source agentic-development environment.
- **PlayCanvas SOG / SplatTransform** — open Gaussian-splat compression and conversion.
- **Civitai / Civision** — community-models and LoRA platforms.
- **LoRA / PEFT / DiffSynth-Studio / Qwen-Image-i2L / Llama Factory / AI Toolkit / LTX-2 Trainer** — fine-tuning frameworks and training pipelines; DiffSynth turns a single image into a custom LoRA; Llama Factory enables zero-code fine-tuning of 100+ models; AI Toolkit covers LTX-2.3 fine-tuning; LTX-2 Trainer is the fal-hosted variant.
- **PyTorch / TensorFlow / Ollama / LLaMA.cpp / MuJoCo** — the open infrastructure layer.
- **Korin AI** — Africa-trained, Africa-built foundation model.
- **DecartAI** — real-time world transformation.
- **Replicate** — model hosting and API marketplace.
- **fal / FAL MCP** — generative-model API platform.
- **llmfit** — hardware scanner for local-model compatibility.
- **Open Vision Agents** — toolkit for building agents that watch, listen and understand video.
- **NotebookLM** — Google's research-and-content-generation surface; cinematic video overviews launched March 2026.

- **OpenAI Gym / Gym Retro / ARC** — research benchmarks (referenced in the games-AI literature).

ComfyUI ecosystem — nodes, extensions and workflows

- **Mesh2Motion** (ComfyUI) — mesh-to-motion conversion node.
- **NanoBanana Pro LoRA Dataset Generator** — creates training datasets for image-editing models in minutes.
- **ComfyUI Music Tools** — professional-grade music-processing node pack.
- **ComfyUI HY-Motion1 plugin** — Tencent HY-Motion implementation.
- **ComfyUI Audio Reactive Node Pack** (VisualFrisson) — spatial audio-reactive nodes.
- **ComfyUI-BlendPack** — nodes for video transitions.
- **KJNodes** — use audio as input for LTX-2 inside ComfyUI.
- **WhatDreamsCost ComfyUI Nodes** — LTX sequencer, keyframer.
- **Luminance Keyer Node** — luminance-based keying.
- **Camera Tracking & Body Pose tools (ComfyUI)** — work-in-progress camera and pose tracking.
- **Multiple Angle Camera Control Node** — multi-angle camera control.
- **FeedbackSampler** (Deforum-inspired) — iterative feedback sampling.
- **Day-to-Night with Qwen Edit LoRA** — day/night transformation workflow.
- **Sora 2 API Nodes** — Sora 2 access from ComfyUI.
- **QWEN-AI Compositing** — Qwen-driven compositing workflow.
- **Street View URL Parser** — Google Street View ingestion.
- **Minim** — ComfyUI model-linker extension.
- **Hunyuan 3D 3.0** (ComfyUI partner nodes) — HY3D 3.0 inside ComfyUI.
- **Kling 3.0 ComfyUI node** — Kling 3.0 partner node.
- **Kling Video 2.6 Motion Control** (ComfyUI) — Kling motion control inside ComfyUI.
- **Kling Motion Control 3.0** (ComfyUI) — the v3 motion control variant.
- **Seedance** (ComfyUI partner node) — Seedance inside ComfyUI.
- **LTX-2** (ComfyUI native support) — audio-video model natively supported.
- **LTX-2 ComfyUI Audio-to-Video** — audio-driven video inside ComfyUI.
- **ACE-Step 1.5** (ComfyUI) — full songs in under 10 seconds on <4 GB VRAM.
- **ElevenLabs ComfyUI Partner Nodes** — ElevenLabs inside ComfyUI.
- **Luma Uni-1 ComfyUI Partner Node** — Uni-1 inside ComfyUI.
- **Meshy 6** (ComfyUI native) — Meshy 6 inside ComfyUI.
- **Grok Imagine** (ComfyUI) — xAI image generation inside ComfyUI.
- **MatAnyone2** (ComfyUI) — high-fidelity video matting inside ComfyUI.
- **Tripo 3.1** (ComfyUI partner nodes) — Tripo 3.1 inside ComfyUI.
- **NVIDIA RTX Video Super Resolution Node** — RTX upscaling node.
- **Depth Anything Custom Node** — Depth Anything inside ComfyUI.
- **SAM 3 / SAM 3D ComfyUI integration** — Meta SAM 3 / 3D inside ComfyUI.

- **Sora 2 Full Loop Workflow** — end-to-end Sora 2 pipeline.
- **Krea Nodes** — Krea inside ComfyUI.
- **Gaussian Splats + Qwen Edit workflow** — explore 2D images and generate new camera views.
- **ComfyUI App Mode / ComfyHub** — turn workflows into shareable no-install apps.
- **ComfyUI Simple Mode** — share and iterate complex workflows more easily.
- **ComfyUI Cloud (private models)** — private model hosting inside ComfyUI Cloud.
- **ComfyUI Action Director** — interactive 3D viewport for ControlNet.
- **Complete Style Transfer Handbook (ComfyUI)** — style-transfer guide.
- **VibeComfy** — agents that understand, build and run ComfyUI workflows with Claude.
- **MooshieUI** — beginner-friendly ComfyUI desktop frontend.
- **Creative Control Using ComfyUI on NVIDIA RTX** — NVIDIA's official ComfyUI/RTX tutorial series.

LoRAs, fine-tuning and training

- **QwenEdit-2509 Light Transfer LoRA** — simple, powerful image relighting via reference.
- **Qwen-Edit Relighting LoRA** — atmospheric relighting control.
- **Qwen-Image-Edit-2511 Multiple Angles LoRA** — multi-angle perspective generation.
- **Qwen-Image-Edit-2511 AnyPose** — pose transfer from a reference image.
- **QwenEdit-2511 Anything2Real LoRA** — realistic-image transformation.
- **Qwen Edit LoRA Object Removal** — bounding-box-precision edits.
- **Qwen-Image-Edit 2511 Gaussian Splash 3D Camera Motion** — 3D Gaussian-splat motion control.
- **Qwen 2511 Time Travel Workflow** — temporal image effects.
- **NextScene (Qwen Image LoRA)** — cinematic image sequences with natural progression.
- **Best Face Swap (Flux 2 LoRA)** — specialised face-swap.
- **FLUX.2 [klein] LoRAs (fal)** — open-source: Outpaint, Zoom, Object Removal, Background Removal.
- **Krea AI LoRA Trainers** — train LoRAs for Qwen-2512 and Z-Image inside krea.ai.
- **LTX-2.3 Character LoRA Training (AI Toolkit)** — character LoRA training tutorial.
- **Gaussian Splats Repair LoRA** — Klein-9b LoRA for repairing 3D views and geometry.
- **LTX 2.3 Colorizer (LoRA)** — black-and-white footage colorisation.
- **Music Finetunes in ElevenCreative** — stylistically consistent vocal and instrument generation.

IP licensing, governance and behavioural-licence platforms

The contract-and-governance layer that emerged through 2025–26. The deep analytical treatment of how these models work — the three Cs of consent, credit, compensation; the five-model comparison; the corrected Disney–OpenAI record; the transaction-cost reframe — is in [Appendix L: The Programmable Brand](#).

- **Hasbro Sixth Wall / CharacterOS** — in-house AI studio launched 3 June 2026 to license the *behaviour* of Hasbro characters (Optimus Prime, Megatron, Cobra Commander, the *Clue* cast) under a runtime governance layer that encodes canon, voice, personality, lore and safety limits; B2B programmatic licence with a talent-participation pay-out structure for original performers; initial 12-character catalogue; 13+/enterprise focus.
- **Verified by Spotify / Spotify–UMG fan-remix product** — Spotify’s *Verified by Spotify* badge and the opt-in catalogue infrastructure underneath the May 2026 Spotify–Universal Music Group landmark licensing agreement for fan-made AI covers and remixes; *consent / credit / compensation* framework; paid Premium add-on; new royalty stream to artists and songwriters.
- **Disney–OpenAI licensing framework (announced 11 December 2025; never executed)** — the \$1bn licensing announcement that became the reference contract architecture for legacy-IP × frontier-AI deals despite the deal itself collapsing when OpenAI shut down Sora on 24 March 2026. IP-not-training, talent-voice-and-likeness excluded, jointly-governed steering committee, equity-and-warrants structure. The framework survived the platform’s death.
- **Elf.tech / GrimesAI** (CreateSafe, Trinit API) — Grimes’s open-voice platform; upload your vocal, get back a synthesised GrimesAI performance, release commercially with a ~50% master-royalty split; analogised by her team to fan-fiction and fan-art cultures; ~15,000 tracks generated under the licence.
- **Holly+ / Holly Herndon DAO** — public digital twin of the artist’s voice wrapped in collective DAO governance for commercial-use decisions and the distribution of proceeds; paired with the *Have I Been Trained?* (Spawning) consent-layer tooling.
- **ElevenLabs Heritage Voices / Iconic Marketplace (with CMG Worldwide and partners)** — heritage-voice licensing pipeline confined to a single text-to-speech Reader app, walled off from the general voice-cloning database; voices include Judy Garland, James Dean, Burt Reynolds, Laurence Olivier, the Stan Lee estate and Michael Caine.
- **Meta Celebrity Chatbots (2023)** — the closed-static likeness-rental anti-pattern referenced for contrast in [Appendix L](#): high flat fee (reportedly up to ~\$5m per celebrity over two years) for platform-owned, frozen personas with no participatory layer for users or the wider creator community.
- **Cannes AI Disclosure Standard** — industry-coordination labelling standard for production-side AI use; the inter-industry analogue of the per-platform governance layers above.

Provenance, watermarking and detection

- **C2PA** (Content Authenticity Initiative / Adobe-led) — cryptographic provenance metadata standard.
- **SynthID / SynthID Verification** (Google DeepMind) — synthetic-content watermark; deployed across Veo, Lyria and Imagen; verifiable inside Gemini.
- **YouTube AI Detection / YouTube AI Deepfake Detection** — automated AI-content detection; *Tiny Grandma* false-positive case study.
- **Deezer AI Music Detection** — identifies up to 75,000 fully AI-generated uploads per day.
- **Spotify AI Transparency Beta** — voluntary disclosure feature.
- **Beeble** — detection and watermarking tools.
- **Cloudflare AI Bot Classification** — public-web infrastructure tracking AI crawlers.
- **Human Provenance AI Disclosure Standard (Cannes)** — industry-coordination labelling standard.

Consumer surfaces and distribution platforms

- **Sora iOS app** (OpenAI) — TikTok-style AI video remix; 1M downloads in 5 days.
- **CapCut / Dreamina** (ByteDance) — consumer video editing + AI generation.
- **Gemini app** (Google) — Google's consumer AI surface.
- **ChatGPT app** (OpenAI) — mobile ChatGPT; Shazam integration.
- **Perplexity AI** — Samsung Smart TV AI assistant integration; consumer AI search.
- **Grok / Grok Imagine** — xAI consumer surfaces.
- **TikTok / Reels / Shorts** — AI-content distribution surfaces.
- **YouTube Shorts / YouTube Create** — destinations for Veo 3.1 Ingredients-to-Video.
- **Bandcamp** — banned AI music outright, January 2026.
- **Deezer** — published the 44%/3% data; built and licensed AI-music detection.
- **Spotify** — declined an AI-music filter; voluntary disclosure beta; *Personal Podcasts* agent feature.
- **YouTube Music** — labelling, the AI Detection Tool, viewership of AI content.
- **Steam** — AI-content labelling controversy; *Clair Obscur* Game of the Year stripping.
- **Sweden's Official Music Chart** — banned AI-generated entries.
- **San Diego Comic-Con** — banned AI art at the 2026 event.
- **Roblox** — creator-tools AI with continued controversy.
- **Fortnite** — Chapter 8 AI-art controversy; Disney partnership.
- **Netflix** — AI in production (de-aging, plates, retention engine, Ben Affleck studio acquisition).
- **Disney+** — user-generated AI content features.
- **Luna** (Amazon) — AI-powered Snoop Dogg game.

Studios, programmes, festivals and institutional infrastructure

- **Sundance AI Literacy Initiative** — Google-funded \$2M creator-training programme; Sundance Collab learning, Story Forum, Ignite Day.
- **UCL / RCA / Brandtech Centre for Creative AI** — UK research hub.
- **AIMICI Training Program** — ScreenSkills-accredited AI training.
- **AI FilmFest Japan** — annual AI-film festival.
- **Dubai AI Film Award** — \$1M prize won by Tunisia's *Lily*.
- **Cannes AI Disclosure Standard / Cannes Lions** — industry coordination for AI labelling; Luma Dream Brief \$1M prize.
- **Future Vision Film Competition** — Google / XPRIZE / Range Media collaboration.
- **Runway AI Festival** — film and creative-discipline awards.
- **Wonder Film Festival** — AI-native film festival.
- **Academy of Motion Picture Arts and Sciences** — “you must be human to win” rule.
- **Television Academy / Emmys** — AI guidance for submissions.
- **SAG-AFTRA** — contract negotiations and the Tilly Tax.
- **Equity (UK)** — strike ballot; 99% vote in favour.
- **PRS for Music AI Survey 2026** — UK music-creator sentiment data.
- **GEMA** — German rights society; won against OpenAI.
- **Splice + UMG** — sample library; AI-music tools partnership.
- **Stability AI + Universal Music / Stability AI + Warner Music** — alliance and deal on the AI-music side.
- **Vizrt Viz One 8.1** — broadcast AI features.
- **Tesseract** — AI platform at MIPCOM.
- **Sphere + Google Cloud** — AI technology behind the *Wizard of Oz* experience.
- **Anthropic Joins Blender Development Fund** — patronage deal, May 2026.
- **Sony AI ICASSP papers** — music-understanding and generative-audio research, April 2026.
- **Google Flow Music + Believe Partnership** — AI music distribution partnership.

Techniques, methods and recurring workflows

- **Vibe Coding** — prompt-led creative prototyping, especially for games and interactive apps; Rosebud, Spawn, Gizmo, Gambo are the canonical surfaces.
- **Trajectory Control (ATI)** (Wan / Qwen) — control of video trajectories at frame and clip level.
- **Reference-to-Video** (Veo 3.1 + Gemini) — consistent character generation from a reference image.
- **Motion Control nodes / Time-To-Move** — frame- and segment-level motion controls in generated video.
- **Layer Separation & Live Editable Text (LET)** (Loveart) — extract and edit text and subjects from images as PSD layers.

- **Temporal Reasoning** (NVIDIA ChronoEdit) — physics-aware edits and world simulation.
- **Practical Light Control** (PractiLight) — foundational diffusion for relighting.
- **3D Gaussian Splatting** — the spatial backbone running through PlayCanvas SOG, Hyperscape, Apple Personas and the open-source splat ecosystem.
- **World Models** — recurring topic spanning Marble, Genie, HY-World, UNI-1, ECHO, Matrix-Game, LingBot-World, Moonlake.
- **Multi-Shot Consistency** — Kling 3.0, ShotVerse, Higgsfield Shots, Veo 3.1 — the multi-shot turn that defined the spring-2026 video pipeline.
- **Camera Control / Multi-Angle Generation** — WAN Camera Control, ShotVerse, ATI, Multiple Angle Camera Control Node.
- **Style Transfer (the ComfyUI handbook)** — the canonical workflow ladder for style transfer.
- **Audio-Reactive Generation** — Audio Reactive Node Pack and the Lalal-API / ACE-Step pipelines.
- **Real-Time Generation Pipelines** — Krea Realtime, MotionStream, Decart Lucy 2.0, Xmax X1, Phoenix-4, LTX-2.3 Real-Time on Scope; the *interactive-generation* category that emerged in Q1 2026.
- **Agentic Workflows / Multi-Agent Stacks** — Sony Game Studios (49 agents, 72 skills), Adobe CX Enterprise, OpenAI AgentKit, Heygen Video Agent, Claude Skills.
- **Provenance-First Capture** — C2PA + SynthID + on-device detection in Premiere and YouTube as the emerging *standard* of provenance-aware production.
- **Music Stem Separation as a primitive** — Lalal AI, StemDeck, BandLab — stem separation moving from end-product to upstream primitive.
- **MCP (Model Context Protocol) as the connective tissue** — Unity MCP, Blender MCP, fal MCP; the protocol the agentic creative ecosystem is converging on.
- **Brief-First, Generate-Second** — the workflow practice the book argues for in Chapter 11 (and Chapter 16's *toolchain in layers* section below).

This is the catalogue. By the time you read it, it will be incomplete — new tools have shipped, some on this list have been bought, renamed or killed. Treat it as a snapshot of one year of toolchain at the moment the toolchain became a stack rather than a list, and use it to orient yourself in whatever the state of play is when you pick the book up.

For a deeper analytical treatment of the adoption telemetry behind many of these tools — Firefly's 22B-asset growth curve, ChatGPT's 800–900M WAU figures, the Veo / Sora professional split, the GDC sentiment-vs-usage divergence — see [Appendix E: Dynamics of Generative AI Adoption](#).

How to build a toolchain you can defend

The last thing I want to say in this chapter is something I have said in talks more often than anything else, because working creatives ask me this question, in some variant, every week.

How do I decide what tools to use, in a market that is changing this fast, without burning my whole month re-learning interfaces?

My short answer is: build the toolchain in *layers*, and accept that the layers move at different speeds.

The bottom layer — your foundation model, your modality stack, your agent platform — will change frequently. Treat it as ephemeral. Pick the best tools available *this quarter* and be ready to swap them next quarter.

The middle layer — your creative software (Adobe / Unreal / DaVinci / Pro Tools / Blender / Logic) — will change more slowly, and is the layer in which the AI capability will be progressively absorbed. Treat it as the long-term home of your craft. Learn it deeply.

The top layer — your *judgement*, your *taste*, your *briefing skill*, your *integration sense* — does not change. It is the layer the agents cannot copy, the layer the platforms cannot ship and the layer the next model release does not depreciate. Spend more time here than the toolchain wants you to.

The mistake I see working creatives make most often is to over-optimize the bottom layer and under-invest the top. The platform companies want you to spend your working hours chasing the new model release; the work that pays the bills, the work that finds an audience, and the work that survives a transition is built on the layer the platform companies cannot reach.

The toolchain, in the end, is the means. The work is the end. The tools change. The work, if it is any good, lasts.

That is the working operating model my studio has run on for the period this book covers. It is the model I would commend to anyone building, this year or next, a creative practice that survives the rest of the decade.

The transition is going to keep going. The tools will keep changing. The work that matters, on the other side, will be made by the people who kept their attention on the right layer.

Chapter 17 — Five Years Inside the Dream Machine

I called the newsletter *Dream Machine* in late September 2025 because the phrase caught something I couldn't yet articulate. I have said elsewhere in this book — in Chapter 15, and in the front matter — what I think the phrase has come to mean over the period the book covers: an apparatus capable of producing what until recently required a human mind, amplifying and distributing the dreaming of the humans who direct it.

What I have not done, anywhere else in the book, is let myself off the leash about what the apparatus might *become*. I have been careful, chapter by chapter, to put the evidence in front of the argument. The predictions in Chapter 15 are dated, falsifiable, defensive — written so the reader picking up the book in 2030 can grade them with a stopwatch.

This chapter is the opposite. This is the chapter where I let myself imagine what the next five years might look like if the trajectory of the previous eight months extends, accelerates, breaks and recombines in the ways I half-suspect it will but cannot, in any conventional analytical register, *prove*.

I want to be honest about the rules I am giving myself for this chapter, because they are looser than the rules the rest of the book has run on.

The rule is: each scenario has to be *argued from* something already in the book — a chapter, a number, a framework, a quoted source. I am not allowed to introduce a new mechanism out of thin air. But I am allowed to take a mechanism that is, in spring 2026, present in small or early form, and ask what it looks like by 2031 if it follows the curve I think it is on.

That is the contract for the next thirty pages. Wild but rooted. Speculative but cited. The kind of writing that the rest of the book would not have permitted, and that the closing letter to 2030 in Chapter 18 would deflate if I tried to put it there.

If a reader in 2031 has bought this book to grade my predictions, this is the chapter where you can grade me most harshly. Some of what follows will look, by then, embarrassingly off. Some of it will look, in the unflattering retrospect that books-on-transitions always get, banal — the obvious thing nobody had quite said yet. The interesting category — the predictions that turn out to be *roughly the right shape, badly miscalibrated on timing* — is the one I am writing for.

Six scenarios. Then the upside I am most hopeful for. Then the downside I am most afraid of. Then a handoff to the letter that closes the book.

One. The Petrillo settlement actually happens

In May 2026, the structural argument of Chapter 6 — that the institutional response to the AI training problem will, on the historical pattern, converge on the *levy-and-redistribute* mechanism James Caesar Petrillo built in 1948 for recorded music — is still, on the page, an argument from analogy. The 88% is the political mandate. The GEMA ruling is the legal precedent. The Stability / UMG-style alliances are the commercial templates. The Musical AI creative-weight-attribution work is the technical infrastructure. None of it has yet been *assembled* into a working settlement.

By 2031, I think the assembly is done. Not perfectly, not globally, not without exclusions — but done, in at least one major jurisdiction, as a working mechanism that a working creative

encounters on a bank statement.

The shape of the thing I expect: a national or supranational *AI Performance Trust Fund*, modelled on the MPTF, capitalised by a statutory per-output levy on commercial generative output, governed by a joint labour–platform body, distributed to working creatives whose training-data contributions were identifiable through the creative-weight-attribution layer described in Chapter 12. The first cheques are small. A working illustrator opens a statement and sees the line *AI Royalty Distribution: £43.18*. A session musician sees *£127.42*. A photographer whose backlist sat in a Getty-class licensed dataset sees a meaningful four-figure annual sum.

The cultural significance of those small numbers will, I think, be larger than the numbers themselves. The first cheque is the moment the 88% becomes a *fact* rather than a *demand*. The mechanism is on the books. The architecture is real. The question stops being *whether there will be a Petrillo settlement* and starts being *how much, to whom, on what calibration*.

The bargaining ground for the next decade is set on that question.

Two. The internet bifurcates, formally

Chapter 4 made the case that the web of 2026 was, in measurable ways, splitting into two distinct attention environments — the *Dead Internet* of synthetic content optimised for synthetic attention, and the *Living Web* of provenance-stamped, human-verified work made for an audience that had organised itself around the difference.

By 2031, I think the bifurcation is no longer a cultural metaphor. I think it is *operational infrastructure*.

The provenance stack the book has spent so many pages cataloguing — C2PA, SynthID, the Cannes Disclosure Standard, the AP and BBC wire-service signing chains, the Adobe content credentials — will, by 2031, have stitched together into a functioning verification layer. Major browsers will render an indicator. Major platforms will, under pressure from advertisers tired of paying for bot-on-bot impressions, surface verified-human content separately from synthetic. A handful of *human-only* subscription services — I would not bet against the BBC, *The New York Times* and one or two of the streaming-music incumbents launching first — will offer “no synthetic content ever” as a paid product.

The split internet is not, in the version I expect, a *clean* split. The synthetic layer will dwarf the verified layer in volume, by perhaps a hundred to one. But the verified layer will capture an outsized share of *paid attention*, *advertiser dollars*, and — most importantly — *cultural credit*. The slop ceiling of Chapter 5, measured at 44/3 on Deezer in April 2026, will by 2031 have hardened into a roughly stable ratio across most major content categories: synthetic supply dominant on the upload side, verified-human attention dominant on the consumption side.

The Living Web is, by 2031, no longer an aspiration. It is a *distribution layer* you can buy access to, with a different economics from the synthetic layer running underneath it.

Three. World models replace flat video as the default high-end medium

Chapter 8 argued that the most important technical shift of 2025–26 was not the video models but the world models — Marble, Genie 3, the Hunyuan 3D family — and that the rate at which spatial-AI tooling was moving from research demos to consumer product was the underappreciated structural story of the period.

By 2031, I think we look back on flat 24-frame video the way 2026 looks back at black-and-white silent cinema. Still made. Still loved. Still occasionally the right medium for the work. But not the *default*.

The default high-end production format, by 2031, is a *navigable spatial render* — a world that the camera moves through in post, that the audience can choose to view from any angle through a Vision-class headset, that an editor can re-cut from a different perspective two months after principal photography is done. The boundary between film, games, immersive and live performance that Chapter 8 described as eroding is, by 2031, functionally gone for new productions.

The career implications, for working filmmakers, are substantial. The skillset of the director-of-photography fragments into world-curation, lighting-prompt design, and post-spatial composition. The skillset of the editor expands into multi-angle viewer choreography. New roles emerge — *spatial continuity supervisor*, *world dramaturg*, *perspective director* — that have no clean analogue in the 2026 production pipeline.

This is also where I expect the legacy industries' Chapter 7 strategic positioning to come back to haunt them. The studios that bet on flat AI-generated video as the next medium will discover, between 2027 and 2029, that they bet on the second-to-last format of the previous era. The studios that built world-model fluency into their pipelines — Sony's documented experiments, the AI-native studios like Imaginae and Critterz, a handful of regional players in Korea, Japan and India — will be the dominant 2031 producers of *the kind of work that uses the medium the way it actually works*.

Four. The orchestrator becomes a credentialled guild

The orchestrator role — central to Chapter 11, referenced through Chapter 13 and Chapter 14 — was, in 2026, a *description of an emerging practice*. The Sony 49-agent / 72-skill team was the most-cited case study. The working title circulated in studios and agencies but did not yet carry contract weight.

By 2031, I think the orchestrator is a *credentialled profession* with collective bargaining power.

The guild structure I expect — and this is the most institutionally specific of the six predictions, so the most likely to be wrong on detail — will look something like the IATSE locals that govern below-the-line film labour, or the Writers Guild's apprenticeship and credit system. *Senior orchestrators* will carry the credential, demonstrate competence in agent stewardship across a defined toolchain, negotiate over rates and over the *quantity of agent labour* a single orchestrator can supervise. The Writers Guild will, I think, be among the first to formalise the role; SAG-AFTRA, Equity and the games-development unions will follow within eighteen months of whoever moves first.

The collective-bargaining angle is the part that will, in 2031, look most novel and most obviously inevitable in retrospect. If a single human orchestrator can supervise the work of, say, twenty agents — a number that today’s productivity research is pointing at, with substantial variance — then the question of *who pays whom for the productivity gain* is exactly the Petrillo question from earlier in this chapter, applied to the orchestrator’s own labour rather than to the underlying training data. The answer the guild will negotiate is a *productivity share* of the agent-team output, with the orchestrator’s individual rate tracking the value of the supervised work, not the volume of the orchestrator’s keystrokes.

This is one of the more concrete reasons I am cautiously optimistic about the working-creative position over the next five years. The job that emerged in 2026 — the maker-as-orchestrator — is a job for which collective bargaining is *unusually well-suited*, because the value being created is unambiguously human in origin and unambiguously measurable in output. Unions exist for exactly this kind of bargaining problem. The Petrillo template, again.

Five. An AI-native studio wins the Palme d’Or before Hollywood does

This is the prediction that will, depending on the reader, look either obviously safe or wildly provocative in 2031.

Chapter 7 catalogued the new AI-native studios — Imaginae, Wonder, Asteria, Obsidian, Critterz, Gossip Goblin — and made the case that their structural advantage over the legacy studios is not technical. It is *cultural*. They have no calcified rules to unlearn. They have no inherited risk-aversion. They have no franchise-template gravity to pull them back to the engine-optimal move. By the chess-grandmasters analogy in Chapter 15, they are, by default, in a better position to play the move the machine would not have generated.

By 2031, I think one of them — or, more likely, a successor company emerging from a director who learned in their orbit — wins the Palme d’Or, the Golden Bear, the Silver Lion or one of the other top European festival awards. *For its writing or direction. Not for its technology.*

The cultural moment when that happens will be larger than the prize itself. It will be the *Petrillo Settlement* of cinema-craft acceptance: the moment the question shifts from *can AI-native cinema be art* to *which AI-native film will be canonical*. The follow-on year, I expect the Academy — having spent 2025–28 reinforcing its *you must be human to win* rule for above-the-line crafts — to add a parallel category, or modify the existing one, to recognise hybrid human-orchestrator-AI authorship under contested but workable rules. The festival juries will get there before the Academy does. They almost always do.

The corollary, for the legacy studios, is bleak. The diagnosis in Chapter 7 — that Hollywood, commercial music and AAA games spent fifteen years optimising toward the engine-optimal move and arrived at the AI moment producing exactly the work the engines can now replicate cheapest — will be on full display when the first AI-native Palme winner is, recognisably, *not* a tentpole. By 2031, the Hollywood studio system will be in the position the major-label music industry was in around 2008: still dominant by revenue, visibly losing the cultural argument, scrambling for the next operating model.

Six. The platform monopoly cracks

The reason I am giving this prediction last among the six is that it is the one I hold most loosely. The structural argument is clean. The timeline is the thing I cannot pin.

By 2031, I do not think the OpenAI / Anthropic / Google triad of frontier-AI dominance will look the way it does in May 2026. The pressures pushing against it, catalogued in different chapters of this book, are:

- The open-weight ecosystem — Tencent’s Hunyuan, Alibaba’s Qwen and Wan, DeepSeek, Meta’s Llama, Mistral, Stability — collectively used by, per [Chapter 16](#), some 80% of startups pitching the major venture funds by spring 2026. Open weights compound. Closed weights, by structure, do not.
- The consent-trained category from the new section in [Chapter 6](#) — Adobe Firefly, Bria, Getty’s licensed model, Moonvalley Marey, AIODE, the Stability / UMG-style alliances — which, by 2031, has either *forced* the frontier labs to license their training data retroactively, or has *displaced* them in the enterprise procurement environment by carrying indemnities the frontier labs cannot match.
- The sovereign-model trend, only just starting in 2026, in which national governments — I would put France, Korea, India, the UAE, Brazil, Japan on the early list — fund consent-trained, citizen-licensed foundation models as a strategic asset, governed under sovereign rules, distributing royalties through national rights-holder bodies.
- The bifurcating internet from earlier in this chapter, which by 2031 makes *which model trained on what* a visible product-marketing claim, in the way *fair-trade-certified* or *organic* became visible product-marketing claims in food.

By 2031, the AI model market I expect looks much more like the 2024 cloud market than the 2025 AI market. Multiple major model families. Sovereign options. Open-weight defaults at the long tail. Indemnification and provenance as standard procurement requirements. The frontier-lab valuations of 2025 either justified by deep enterprise penetration in narrow categories, or remembered the way AOL’s 1999 valuation is remembered.

The corollary I want to put on the page, because it has not yet been said cleanly anywhere else in this book: the platform companies dominant in May 2026 are, on the historical pattern, *unlikely to be the platform companies dominant in 2031*. The history of computing is that the dominant platform of one era is rarely the dominant platform of the next. The Wintel of the 1990s, the iOS / Android of the 2010s, the AWS / Azure of the 2020s — each was a dominant pairing whose successor came from a direction the incumbents did not see. Whatever the dominant generative-AI platform of 2031 is, I think there is a non-trivial chance it is a company most readers of this book — including me — have not yet heard of.

The upside I am most hopeful for: the Dream Machine becomes literal

I want to spend the rest of this chapter on two scenarios that do not fit cleanly into the *predict and date* register, because they are the ones I think about most when I am away from the work and most off-script when I am in it.

The first is what happens if the title of this book turns out to be a literal description.

The phrase *Dream Machine*, as I have used it through the newsletter and the book, has been a metaphor for the apparatus of generative AI. A composite name for the platforms, models, infrastructure and human labour that, between them, produce the new creative work.

I want to entertain the possibility that, by 2031, the phrase describes a *literal piece of consumer hardware* — a wearable creative-cognition prosthetic, sitting somewhere on the continuum between a Vision Pro and a brain-computer interface, that watches a working creative make work, learns their taste, infers their intent, and contributes — in real time, in collaboration, at the speed of thought — to the work in progress.

The technical building blocks are visible. Apple's Vision Pro and the next-generation headset class are mature consumer products by 2031. The brain-computer interface category, dominated in 2026 by medical applications, is on a trajectory the consumer hardware press has been undercounting. The agentic orchestration layer that I described in Chapter 11 is, by 2031, no longer something a human types instructions into. It is something a human *converses with, gestures to, thinks at*. The interaction model of *prompting* is, by 2031, an obsolete UX pattern, remembered the way command-line interfaces are remembered now.

The literal Dream Machine, in this scenario, is not a separate device. It is the orchestration layer made *embodied*. A creative cognition prosthetic that allows a working creative to externalise, manipulate, sketch, iterate and finish creative work at a fluency that is, today, available only to the smallest set of professionals who have spent thirty years building the relevant neural pathways. The leverage the device provides is not *replacement* of the human imagination. It is the closing of the gap between the human imagination and what the human hand can, today, get out into the world.

If this device — or one like it — ships at consumer price in the period this chapter covers, the *access principle* from Chapter 15 gets a lift I cannot easily exaggerate. The bedroom-Hollywood dynamic from Chapter 14 — the teenager with a midrange GPU producing studio-quality output — extends to *every* discipline, including the ones that today still require expensive equipment and decades of haptic training. The new geography I argued for in Chapter 15 — the dispersion of canonical creative work to places the previous century forgot — becomes structurally inevitable rather than aspirational.

The Dream Machine, in this scenario, is not the platforms. It is the *prosthetic*. The platforms are infrastructure underneath it. The dreaming is, finally, located where the title was always insisting it was located: in the human at the centre of the apparatus.

I want this scenario to be the one that comes true. I think there is a real chance — not a high probability, but a meaningfully non-zero one — that it does.

The downside I am most afraid of: the audience becomes synthetic

The corresponding worst case I want to put on the page is the failure mode the slop ceiling cannot, by construction, defend against.

The slop ceiling, as I described it in Chapter 5, works because *the audience can tell* — at the speed of a swipe, in the aggregate, with reliability — that synthetic content does not carry the human signal real human attention is calibrated to. The ratio holds, the 44/3 holds, the

audience underweights the slop, because there is an audience, and the audience is made of humans, and the humans are still, in May 2026, paying attention to one another.

The downside scenario is that, by 2031, that last condition has decayed.

The mechanism is the one the Dutch researchers cited in Chapter 5 and the bot-traffic investigations cited in Chapter 4 were already, by spring 2026, beginning to document. As synthetic content production becomes free and synthetic attention-allocation agents become widely deployed, the *measured* audience for any given piece of content drifts further from the *human* audience. Platforms optimise against the measured audience, because that is what their analytics surface. The content the platforms surface, in turn, becomes the content the synthetic audience underweights least. The synthetic audience grows. The human audience, exhausted, exits. The platforms continue to grow on the measured numbers because the measured numbers are increasingly bot-on-bot. The cultural production layer continues to produce, but produces *for bots watching bots*. The work loses meaning slowly, then quickly, because the audience that produced meaning has stopped being part of the loop.

This is the version of 2031 I am most afraid of. Not because AI replaces creators. Creators are not the load-bearing structure. *The audience is the load-bearing structure*, and the audience is the thing the architecture of the synthetic internet, deployed without the verification layer described in scenario two of this chapter, is structurally positioned to dissolve.

The corresponding fight, the one I think the next five years has to win to avoid the downside, is the *audience verification* fight. The C2PA-equivalent for *who is watching*, not just for *who made*. The mechanism that lets a creator know — and lets an advertiser pay for, and lets a platform surface — the work that *real humans were really paying attention to*. The early signals in this direction are there, in the form of platform attestation experiments and the bot-traffic regulatory pressure documented in Chapter 4. None of them are, in May 2026, anywhere near the scale of the upstream provenance work.

If the next five years build the verification layer for production but not for consumption, the slop ceiling holds in the short term and collapses in the medium term. The creators win the upstream battle. The audience disappears anyway. The downstream meaning of the work, by 2031, is gone — not because human creativity has been replaced, but because the social fact that creativity is made *for* somebody has been quietly dissolved underneath the work.

This is the failure mode I think we are most exposed to and least talking about.

What we make of it

I want to close this chapter the same way I open the next.

The phrase *Dream Machine* was, when I named the newsletter, a placeholder for a feeling. By the end of Chapter 15 it had become an argument: that the apparatus we are building amplifies, multiplies and distributes the dreaming of the humans who direct it, and that the question every working chapter of this book has been about is *whose dreams the machine amplifies*.

What this chapter has tried to do is take that argument out to five years and let it run.

Six things I think happen. The Petrillo settlement, made real, on a working creative's bank statement. The internet, formally split. World models, displacing flat video as the default. The orchestrator, credentialled and bargained-for. The first AI-native Palme d'Or. The platform monopoly, cracked. One upside I am hopeful for — the Dream Machine made literal, a creative cognition prosthetic that closes the gap between human imagination and finished work. One downside I am afraid of — the audience layer collapsing under synthetic capture, leaving the work without anyone to make it *for*.

Three of the six will, almost certainly, be wrong. One of them will be wrong in the direction of *I undersold this*. One of them will be wrong in the direction of *the timeline was slower than I thought*. One of them will be wrong in the direction of *the thing I described did not happen because the thing that happened instead was weirder and more interesting*. I do not know yet which is which.

The point of the exercise is not to be right about the six. The point is to put a *shape* on the page, in 2026, for the conversation that the working creatives, the studios, the unions, the policy people and the audience will be having with each other, week by week, between now and 2031. The four principles from Chapter 15 — *agency, attribution, access, audience* — are the test you apply to each daily decision. The scenarios in this chapter are the *direction* the aggregate of those decisions is, on my read of the available evidence, most likely pointing in.

The title of this book, I have come to think, was always slightly mis-named. The Dream Machine is not the apparatus the platforms built. It is not even the prosthetic the consumer-hardware companies might ship by 2031. The Dream Machine is the *whole system* — the platforms, the prosthetics, the working creatives, the audience, the unions, the institutions, the regulators, the literacy initiatives, the newsletters, the festivals, the awards, the lawsuits and the daily decisions of millions of people about whose work to make and which work to pay attention to. The thing the title points at is the *coupled human-machine cultural production system* that the period this book covers brought into existence.

What we make of that system, between May 2026 and the moment a reader picks this book up in 2031 to grade me, is the work.

The letter in the next chapter is what I want to leave that reader with.

Welcome to the Dream Machine.

Epilogue — A Letter from the Dream Machine

To the creative person reading this in 2030, picking the book up out of curiosity or for a class or because someone older than you said you should:

I want to write you a short letter, because by the time you are reading this, more than half of what is in this book will already be wrong.

Some of the tools I have spent chapters discussing will be obsolete. Some of the companies will have been bought, broken up or sunk. Some of the law I described as new and contested will have settled into precedent that everyone takes for granted. Some of the people I quoted as authorities will have changed their positions, and some of the predictions I made — the ones I let myself make — will have aged in ways I would now find embarrassing to read back.

This is fine. It is, I think, the right outcome for a book written in the moment that this one was written in. The job of a book about a transition is not to predict where the transition lands. The job is to *describe the moment honestly* — what was being argued about, by whom, with what evidence, and what was at stake — so that the reader picking it up later has, at minimum, an accurate record of *what the people in the room thought they were doing*.

There are a few things I want to flag for you, looking back from your time at mine.

What we got right

We got the slop ceiling right. I am almost certain of this, even sitting here in May 2026, because the underlying mechanism — that audiences allocate attention to *meaning*, not to *quantity* — is one of the most reliable findings in the history of cultural production, and there is no plausible scenario in which it stops working.

If you are reading this in 2030 and there is a thriving creative economy, it is because the slop ceiling held. The audience that turned away from synthetic content in 2026 either kept turning away, or the platforms and the creators figured out how to make AI-augmented work that genuinely earned attention — sincere, transparent, made with care. Both outcomes preserve the underlying contract between maker and audience. Either way, the ratio between volume and attention is still doing its work.

If, on the other hand, you are reading this in 2030 and the creative economy looks bleaker than the one I have described, it is — almost certainly — because the architecture of the internet was allowed to drift further in the direction the Dutch researchers identified in 2025: an attention market optimised against meaning, where the ratio between volume and attention stopped functioning because the *audience itself* had become indistinguishable from the bots.

The slop ceiling was real. The question for your decade is whether you protected the *audience* that produced it.

What we got wrong

We got the timeline wrong, in both directions.

In some places, we expected the change to be slower than it turned out to be. The pace at which world models moved from research demos to consumer features in the eight months I

documented in this book was something almost nobody — including me — would have predicted in late 2024.

In other places, we expected the change to be faster. The full AI-native film studio that was supposed to replace Hollywood by 2027 is still, in the spring of 2026, mostly a series of well-funded prototypes. The fully autonomous game-development pipeline is, even at Sony, still an *assisted* pipeline rather than an *autonomous* one. The headline futures the platform companies have been selling are, in the main, still further off than the press cycle suggested.

If you are reading this in 2030, I suspect that the timeline gap will look in retrospect like the most obvious thing about the moment we were in. The interesting changes were the ones nobody put on a slide. The boring changes — the slow institutional repositioning, the contract-by-contract renegotiation of how creators are paid, the patient construction of the C2PA standards, the union by union recalibration of what counts as a fair use of a performer's likeness — were the ones that actually shaped your decade.

This is, I think, almost always true of technology transitions. The press loves the visible. The structural change happens in the invisible. The fact that you can pick this book up in 2030 and read it without unusual cost or effort is the result of a thousand small infrastructure decisions, made by people whose names did not appear in this book, that none of us thought to celebrate at the time.

What I want you to know about us

We were, in 2026, frightened. Most of the working creatives I knew that year were carrying a layer of low-grade fear — about their livelihoods, about their crafts, about the institutions that had organised their working lives — that I do not think showed up in the trade press or the platform keynotes or even in this book as much as it should have.

We were also, I think, more hopeful than we said out loud. Most of the working creatives I knew that year were also doing some of the most experimental, most curious, most adventurous work of their careers, because the tools made things possible that had been impossible the year before and they wanted to find out what.

Both states were true at the same time. They will be true again in your decade, about whatever the next transition is. I don't think we managed the fear well. I think we did okay with the hope. I am proud of the people who turned up, week after week, to argue about how all of this should go.

A specific request

If the institutions of the generative creative economy — the unions, the rights bodies, the consultations, the literacy initiatives, the disclosure standards, the festivals, the regional studios, the indie filmmakers in places that the previous century forgot — are still doing their work when you read this, *thank the people who built them*. Most of those people have not become rich, or famous, or particularly comfortable. They were working creatives, union officers, policy officers, festival organisers, technologists, lawyers, professors and freelancers who chose, in their unglamorous moments, to do the slow institutional work that this kind of economy needs.

The platform companies will not remember them. The audience will not know their names. The cultural histories will record the names of the directors and the songwriters and the celebrity AI controversies. The institutional work happens in the basement and the archive and the union office and the policy briefing.

If you are inheriting a creative economy that works, you are inheriting it from those people. Thank them while they are still around to be thanked.

A specific request to working creatives in 2030

Keep going.

The pace of change will be at least as fast in your decade as it was in mine. The next paradigm of AI tooling — whatever it is — will, I am sure, make the world models and agents of 2026 look as quaint as Photoshop in 1996. You will have your own version of *the day Sora landed*. You will have your own *Tilly Norwood week*. You will have your own 88%. You will have your own choice between the extractive and the generative economy.

Whatever the new tools are, the principles I tried to articulate in Chapter 15 — *agency, attribution, access, audience* — will, I am confident, still apply. They are not specific to the 2026 transition. They are general properties of how a humane creative economy organises itself, and they will work in any technology environment that produces creative outputs at scale.

The deeper conviction underneath the principles will, I think, also still apply. The age of the *How* — the long century in which the central question of working creative life was *can you do the thing?* — was, by 2026, visibly ending. The age of the *Why* — taste, intent, authenticity, the willingness to take a risk on the move the data does not endorse — was visibly starting. The image I borrowed from elite chess in Chapter 15, of grandmasters winning by deliberately playing the move the engine would not have played, will still be the picture I would commend to you. By the time you read this, the engines will be better. The pressure to play the engine's optimal line will be stronger. The competitive advantage of the *deliberately un-machine-like* move — of the move that is yours because no machine, trained on what came before, could have generated it — will, if anything, have widened.

Apply them. Argue for them. Build for them. Defend them when they are under pressure.

And — this is the part I want you to take seriously even when it feels grandiose — *write the newsletter*. Whatever your version of the newsletter is. The thing that, when the moment comes, you sit down on a Monday morning and decide to write because nobody else seems to be doing it. Most of those newsletters won't matter. A few of them will. The one you are reading was, for me, one of the most consequential decisions I ever accidentally made.

You will know which moment is the moment, when it arrives. You will know because the people you respect will be staring at the same news cycle, looking at each other in the same way, asking the same question.

When that moment arrives in your decade, write the newsletter.

Closing

I want to end this letter the same way I have ended every issue of *Dream Machine* for the past six months, because the line has come to mean more to me than I expected it to:

If you've got any recommendations or things we need to know about for our next edition, please feel free to reach out.

That sentence was, for all the time the newsletter has run, my way of acknowledging that I did not have the picture on my own. That the readers were the network. That the work was a *conversation*, not a broadcast. The newsletter has only ever been as good as the community of creatives, technologists, union reps, academics, festival programmers, indie filmmakers, working musicians and audience members who have, week after week, sent me the things I would otherwise have missed.

If you are reading this in 2030, the newsletter is still — in whatever form it has taken by then — that same conversation. Keep adding to it. The Dream Machine, in the end, is not the AI. It is the network of humans who, between them, are deciding what the machine is *for*.

Welcome to the Dream Machine.

— Pete Woodbridge, *DreamLab, the North West, May 2026*

*Appendix A — A Quantitative Anatomy of
Six Months*

This appendix is a structured tour of the corpus the book was built from. It is not in the body of the manuscript because it would interrupt the argument; it lives here so that the reader, the policy researcher, the journalist or the historian picking the book up later can see what the underlying data actually looks like and check the arguments against it.

* * *

A1. Corpus shape

- **Total fetched and parsed articles:** 1,388.
- **Total captured words across the corpus:** ~1,099,216 (~6,945,361 characters of post-extraction text).
- **Source span:** 32 editions of *Dream Machine* | *Creative AI*, published between 6 October 2025 and 4 June 2026 (the quantitative tables in §A1–A6 below cover the analytic window of Issues 1–29; **Issue 30** (21 May 2026), **Issue 31** (1 June 2026), and **Issue 32** (4 June 2026) are documented in §A7 as post-cut supplemental data, and are incorporated narratively into the closing material across the body chapters).
- **Average articles per newsletter edition** (in this corpus): ~48.
- **Capture rate** against the full curated URL set: **91.4%** (1,438 of 1,574 URLs returned readable content; the remainder hit bot-detection, 404s, or live-page connection issues).

A2. Most-cited domains

Where did six months of curated coverage actually come from? The top 30 domains, by article count:

Rank	Domain	Articles
1	x.com	82
2	github.com	75
3	youtube.com	74
4	musically.com	57
5	theguardian.com	38
6	variety.com	38
7	hollywoodreporter.com	34
8	open.spotify.com	25
9	huggingface.co	24
10	musictech.com	22
11	forbes.com	21
12	gamesindustry.biz	20
13	bbc.co.uk	18
14	deadline.com	18
15	techcrunch.com	17
16	pcgamer.com	17
17	musicradar.com	17
18	blog.google	15
19	futurism.com	15
20	completemusicupdate.com	15
21	videogameschronicle.com	14
22	businessinsider.com	13
23	theverge.com	12
24	musicbusinessworldwide.com	12
25	cnet.com	11
26	gamesbeat.com	11
27	adweek.com	10
28	digiday.com	10
29	pocketgamer.biz	10
30	blog.comfy.org	9

Reading note. Trade press (Hollywood Reporter, Variety, Deadline, Music Business Worldwide) and tech press (Verge, Wired, TechCrunch) dominate. The platform-company blogs (OpenAI, Adobe, DeepMind, Stability) and policy bodies (UK Gov, Reuters Institute, Imperva, Cloudflare) sit underneath. The geographic concentration is North American and

British, which reflects both the newsletter author’s vantage point and a real imbalance in where creative-AI coverage is concentrated.

A3. Story volume by sector, month by month

How the conversation moves through the six months — count of corpus articles touching each sector, by publication month of the newsletter issue that cited them:

Month	Film & TV	Games	Music	Adv/Mkt	News	Policy/Law	Tools/Models	Total
2025-10	136	111	155	104	25	102	184	817
2025-11	112	89	120	97	14	93	150	675
2025-12	73	58	78	57	9	43	98	416
2026-01	89	69	96	70	13	56	120	513
2026-02	77	72	94	61	7	63	109	483
2026-03	81	49	81	55	12	47	106	431
2026-04	100	84	110	69	11	64	138	576
2026-05	74	60	89	53	14	57	113	460

Total story tags across the period: 4,371. (Articles can fall into more than one sector — many do, which is itself part of the story: the boundaries between film, games, music, advertising and tooling have been meaningfully porous through the AI era.)

A4. The voices: public-figure mention frequency

Public figures appearing in three or more corpus articles, ranked by article count. This is *not* a sentiment ranking — only a measure of how often someone surfaces in the AI conversation:

Rank	Name	Articles
1	Sam Altman	32
2	Tilly Norwood	30
3	Xania Monet	18
4	Taylor Swift	18
5	RZA	17
6	Steven Soderbergh	17
7	James Cameron	16
8	Paul McCartney	12
9	Mark Zuckerberg	12
10	Eline Van der Velden	12
11	Madonna	12
12	MrBeast	12
13	Tim Sweeney	9
14	Christopher Nolan	9
15	Emily Blunt	8
16	Breaking Rust	8
17	Mikey Shulman	7
18	Matthew McConaughey	7
19	Steven Spielberg	7
20	Robert Kyncl	6
21	Guillermo del Toro	6
22	Will Smith	6
23	Lucian Grainge	5
24	Sienna Rose	5
25	Natasha Lyonne	5
26	Fei-Fei Li	4
27	George Clooney	4
28	Demis Hassabis	4
29	Ron Howard	4
30	Ted Sarandos	4
31	Adam Mosseri	3
32	Imogen Heap	3
33	Dave Stewart	3

Rank	Name	Articles
34	Chris Pratt	3
35	Joost van Dreunen	3
36	Brian Grazer	3

Reading note. James Cameron, Guillermo del Toro and Leonardo DiCaprio are the three voices most-cited in opposition to generative AI in performance. Tilly Norwood and Xania Monet are the two most-cited synthetic entities in the corpus. *Both lists matter equally to the story this book is telling.*

A5. The tool wave

Cumulative count of distinct AI tools, models and platforms entering the corpus, by month of first mention:

Month	New tools first mentioned	Cumulative
2025-10	38	38
2025-11	6	44
2025-12	3	47
2026-01	5	52
2026-02	2	54
2026-03	4	58
2026-04	0	58
2026-05	2	60

Reading note. The tool cadence ran at roughly **7.5 new platforms or major-version releases per month** across the period. This is roughly four times the pace of any other software-tool category I have personally tracked over a comparable window. The implication is that any working creative making technology bets in this period was, by definition, working with incomplete information — the relevant toolchain had not stabilised long enough for any single bet to settle.

Most-mentioned tools and platforms (top 30, by article count):

Rank	Tool	Articles
1	Udio	587
2	Wan	539
3	ChatGPT	136
4	Gemini	98
5	Anthropic	84
6	Suno	76
7	Sora	69
8	ElevenLabs	49
9	Premiere	48
10	Veo	46
11	ComfyUI	45
12	Sora 2	37
13	Claude Code	36
14	Flux	36
15	Veo 3	30
16	Runway	30
17	Kling	30
18	Nano Banana	29
19	Seedance	26
20	Tencent	24
21	Genie	24
22	Firefly	21
23	Photoshop	21
24	Hunyuan	18
25	Veo 3.1	17
26	Luma	13
27	Marble	10
28	Rodin	10
29	LTX-2	9
30	Higgsfield	7

A6. The vocabulary shift

Recurring key phrases by month — articles containing each phrase:

Phrase	2025-10	2025-11	2025-12	2026-01	2026-02	2026-03	2026-04	2026-05	Total
ai slop	8	8	6	15	5	7	7	6	62
ai actor	4	5	3	2	2	2	1	1	20
synthetic performer	0	0	2	1	0	0	0	1	4
world model	8	3	3	1	9	6	8	2	40
agentic ai	12	6	2	7	6	5	8	3	49
ai agent	20	18	4	11	12	9	16	10	100
deepfake	9	6	3	10	6	8	4	3	49
human authorship	1	3	1	0	1	0	1	1	8
training data	9	10	11	6	8	12	5	10	71
consent	16	11	11	11	9	6	10	11	85
license	26	30	18	21	35	25	29	16	200
copyright	50	53	27	24	33	26	25	24	262
ai-generated	65	58	26	46	30	29	41	28	323
ai actress	3	2	0	0	1	0	1	0	7
watermark	5	2	0	4	3	2	1	3	20
synthid	1	0	0	2	2	1	1	0	7
c2pa	0	1	0	2	1	2	0	0	6
provenance	5	2	1	3	1	1	1	5	19
disclosure	6	12	7	9	4	3	7	6	54
fingerprint	0	2	0	3	2	0	0	1	8
creative ai	8	5	0	1	0	1	3	1	19
generative ai	85	59	35	43	32	32	40	32	358
creator economy	4	10	0	3	3	0	2	1	23
tilly norwood	10	7	3	3	1	2	2	2	30
xania monet	3	7	2	3	0	2	1	0	18
breaking rust	0	3	1	1	0	1	1	1	8

Phrase	2025-10	2025-11	2025-12	2026-01	2026-02	2026-03	2026-04	2026-05	Total
slop ceiling	0	0	0	0	0	0	0	0	0
model collapse	0	0	0	0	0	0	0	0	0

Reading note. Watch *AI slop* — it goes from a fringe term in October 2025 to a Merriam-Webster word of the year by December and a policy framing by the spring. Watch *agentic AI* — it lifts after the October DevDay and never falls back. Watch *world model* — barely present in October 2025, ubiquitous by April 2026. Watch *consent / license / copyright* — climbing all the way through, with a sharp December spike around the UK consultation closure.

A7. May–June 2026 supplemental: the Issue-30, Issue-31 and Issue-32 datapoints

The analytic corpus underlying §A1–A6 closes at *Dream Machine* Issue 29. **Issue 30** (21 May 2026), **Issue 31** (1 June 2026), and **Issue 32** (4 June 2026) post-date the analytic cut and are not represented in the article-frequency tables above; Issue 30 catches the **Google I/O 2026** announcement wave, Issue 31 catches the **June IP-licensing turn** (Hasbro Sixth Wall / Behavioural Licensing, Spotify–UMG, the Epidemic Sound *Future of the Creator Economy* report and the Sora shutdown), and Issue 32 catches the **early-June 2026 policy and industry inflection** (Nandy admission, Suno Series D, Dreams of Violets, Amazon AI Creators Fund, Scorsese/Black Forest Labs, NVIDIA Cosmos 3, Roblox/Morpheus, Tripo AI). Together they are the source for the manuscript’s closing-month additions. The numerical datapoints from Issues 30, 31 and 32 worth recording here in standalone form:

Datapoint	Value	Source
SynthID watermarked items, cumulative	100B+	Google DeepMind, May 2026
Wonder Studios total funding	\$50M	<i>Forbes</i> , May 2026
Runway Japan investment (Tokyo office)	\$40M	Runway, May 2026
Viktor (virtual AI coworker) Series funding	\$75M	<i>Fortune</i> , May 2026
Sondo AI claimed global users	10M	<i>Musically</i> , May 2026
13–15 year-olds using AI to “be creative” (Snapchat survey)	31%	Snap Newsroom, May 2026
Australians who say AI-generated ads make them trust a brand less (YouGov)	45%	YouGov AU, May 2026
NVIDIA SANA-WM model size	2.6B	NVIDIA, May 2026
SANA-WM native video-generation length	60 sec	NVIDIA, May 2026
Epidemic Sound <i>Future of the Creator Economy</i> survey size	3,000	Epidemic Sound, June 2026
Creators using AI tools (Epidemic Sound)	94%	Epidemic Sound, June 2026
Creators expecting to increase AI usage in next 12 months	72%	Epidemic Sound, June 2026
Creators feeling pressure to use AI to keep up	89%	Epidemic Sound, June 2026
Creators who believe human-created content will become a premium product	75%	Epidemic Sound, June 2026
AI-generated songs as share of Apple Music plays	<1%	<i>Billboard</i> , June 2026
Sora peak monthly revenue (pre-shutdown)	~\$540k	<i>Tech-Insider</i> , March 2026
Disney–OpenAI announced licensing value (never executed)	\$1B	The Walt Disney Company / OpenAI, December 2025
Hasbro Sixth Wall initial character catalogue	12	Hasbro, June 2026
Suno Series D funding round	\$400M	<i>Variety</i> / <i>TechCrunch</i> , June 2026
Suno post-money valuation	\$5.4B	<i>Variety</i> , June 2026
Suno annual recurring revenue	\$300M	<i>Variety</i> , June 2026
Suno subscribers	2M+	<i>Variety</i> , June 2026

Datapoint	Value	Source
<i>Dreams of Violets</i> (Tribeca) production cost	\$2,000	<i>Variety</i> , June 2026
<i>Dreams of Violets</i> production time	3 months	<i>Variety</i> , June 2026
Amazon AI Creators Fund: days between commission confirmation and creator exit (Punky Duck)	2	<i>Variety / Hollywood Reporter</i> , June 2026
Lisa Nandy (UK Culture Secretary) SXSW London admission that government “got it wrong” on AI copyright	—	<i>Screen Daily</i> , 1 June 2026

Reading note. Issue 30’s headline tool releases (Gemini Omni, Beeple Canvas, Sony Woosh, Mirelo SFX 1.6, Tencent Ardot, Odyssey Starchild-1 / Agora-1, NVIDIA SANA-WM, Apple Headsup, Stable Audio 3.0, PhotoGIMP, Tamber, ECABridge, Claude/ComfyUI) lift the cumulative tool count in §A5 by roughly **a dozen entries in a single week**. The May-2026 cadence is the highest single-week tool-release count in the period the book covers, and reads — in the context of the §A5 average of 7.5 new platforms per *month* — as a Google-I/O-week saturation point. Issue 31 settles back toward that prior cadence, with the focus shifting from raw tool releases to the *IP-and-licensing turn* (Spotify–UMG, Hasbro Sixth Wall, the Disney–OpenAI corrected record) and the *creator-attitudes* data from the Epidemic Sound report — the dataset that anchors the closing argument of [Chapter 13](#) and the *three Cs* framework in [Chapter 6](#). Issue 32 brings a further cluster of major model releases (NVIDIA Cosmos 3 — omnimodal world model, released 31 May 2026, rated best open-source T2I and I2V by Artificial Analysis) and infrastructure investments (Roblox acquires Morpheus AI for video world models; Tripo AI raises ~\$200M for AI 3D and world model tech / Project Eden) alongside the policy and industry datapoints recorded in the table above. The Issue 32 financial headline — Suno’s \$5.4B valuation alongside active UMG and Sony litigation — is the cleanest single data point on the financial reality of the creative-AI market in mid-2026: the licensing fights and the capital markets are running in parallel, not sequentially.

* * *

What this appendix is for

Every chapter of this book is a *reading* of the corpus described above. It will be useful in 2030 and beyond to be able to see the underlying shape of the corpus, separate from the argument the book builds on top of it.

If you want to test the argument against your own reading of the same evidence: every URL in the corpus is enumerated in the citation index, every scraped article is preserved in JSON form in the `Research/scraped/` directory of the source repository, and every analysis in this appendix is reproducible by running `Research/quant.py`.

If you want to extend it: the scraper is in `Research/scrape.py`, the analyzer is in `Research/analyze.py`, the per-chapter dossiers are in `Research/dossier/`. Fork it, change it, run it on the next six months. I'd be glad to see what you find.

Appendix B — Glossary

The terms in this glossary are the working vocabulary of the book. Some are coinages, some are borrowed from elsewhere and reframed; all are used with precise, deliberately narrow meanings in the chapters above.

* * *

Agency line (or Continuum line). A single axis representing the share of decision-making in a creative function performed by a human versus a machine system. See **Human–AI Agency Continuum**.

Agentic AI. A class of AI system that, given a goal, can plan, decide and execute a sequence of multi-step actions without further human input between steps. Distinct from a *generator*, which produces an output in response to a single prompt. See Chapter 3.

AI literacy. The cluster of skills required to deploy AI tools effectively in creative work — including briefing, taste, judgement, prompt practice, output evaluation and tool-stack fluency. The term moved from optional to baseline competency through 2025–26, formalised by initiatives such as the Sundance Institute’s AI Literacy Initiative launched in January 2026.

AI slop. Low-quality, mass-produced AI-generated content that is recognisable to audiences as such — usually because it is made without human creative intent. *Merriam-Webster’s* word of the year for 2025. See **Slop Ceiling**.

Attribution. The principle that when AI systems produce derivative outputs based on training data, the human authors whose work shaped those outputs should be identified and — where appropriate — compensated. One of the four principles of the generative creative economy (Chapter 15). Technical infrastructure includes C2PA, SynthID, and creative-weight-attribution systems.

Audience contract. The implicit agreement between makers and audiences about what creative work is, what conditions of making it carries, and what relationship the audience can expect with its makers. The shift from implicit to explicit audience contracts is one of the central structural changes the book describes. See Chapter 12.

Augmented intelligence. Reframing of “AI” used by some industry voices to emphasise human-in-the-loop deployment over autonomy. Cited in *Dream Machine* Issue 9. Compare with **Generative AI, Agentic AI**.

C2PA. *Coalition for Content Provenance and Authenticity* — a technical standard for embedding cryptographic provenance metadata in media files, supported by camera manufacturers, editing software and a growing number of platforms. The principal “fingerprint real media” infrastructure underlying the authenticity argument in Chapter 12.

Coordination collapse. The structural change that occurs when the labour-coordination architecture of a creative organisation — built around the bandwidth constraints of human-only teams — is overtaken by AI-assisted workflows that no longer require those constraints. The subject of Chapter 13. Manifests as *shadow AI* below management sight and as compressed middle layers in the workforce above.

Dead Internet Theory. The notion that most of the public web is now bot-generated and bot-read, with humans increasingly a minority of traffic. Once a conspiracy framing; by 2025 a measurable phenomenon — bots accounted for 51% of web traffic in the Imperva 2025 Bad Bot Report, of which ~80% of bot traffic was AI training crawlers. See Chapter 4.

Disclosure. The practice of declaring the use of AI in the production of a piece of creative work — in credits, contracts, metadata, watermarks, or platform-facing labels. By spring 2026, disclosure had emerged as the dominant industry response to the audience-authenticity question, anchored by standards including the **Cannes AI Disclosure Standard** (May 2026) and the **Academy of Motion Picture Arts and Sciences** rule requiring human authorship for awards eligibility.

Extractive economy. A creative economy in which AI systems are trained on unpaid human work, the platform companies that build the models capture most of the resulting economic value, and the diffuse pool of working creatives is steadily decapitalised. One of two possible end-states the book identifies. See Chapter 15.

Fingerprint real media. Adam Mosseri's (Instagram) framing of the verification problem: amplify provably-human content rather than chase synthetic content for labelling. Used in the book as shorthand for *provenance-first* approaches to content moderation. See Chapter 4.

Generative economy. A creative economy in which AI tools are treated as new craft infrastructure, training data is consented and compensated, platforms compete on tool quality and integrity, and the productivity gains are broadly distributed rather than concentrated. The opposite of the **extractive economy**. The four principles (*Agency, Attribution, Access, Audience*) are the operational test. See Chapter 15.

Human–AI Agency Continuum. A frame, introduced in *Dream Machine* Issue 2 (October 2025) and extended in Chapter 3, in which any given creative function is mapped on a horizontal line from full human agency (left) to full machine agency (right). The frame's key claim is that *each creative function moves at its own speed* — you can sit at the extreme left on performance while being at the right edge on plate generation.

Living Web. The deliberately-built portion of the public web in which authorship is provable, attribution is durable, attention is allocated on non-viral signals, and the architecture supports rather than undermines human creative work. The aspirational opposite of the **Dead Internet**. Has to be built, not assumed. See Chapter 4.

Mid-career squeeze. The structural pressure on workers in the middle of creative-industry careers — neither at the junior entry level (replaced by agents) nor at the senior decision-making level (still required) — as AI absorbs the intermediate production roles that those mid-career workers historically held. See Chapter 13.

Model collapse. The technical risk that AI systems trained predominantly on synthetic data — including data produced by earlier generations of AI systems — progressively lose touch with real-world signal and produce increasingly homogenised, hallucination-prone outputs. The risk that *Dead Internet, Living Web* warns has moved from theoretical to measurable.

Orchestrator. The role that emerges when an individual working creative — or a small team — directs a large pool of AI-agent capacity. Defined operationally by five activities: defining

the brief, allocating work, briefing the agents, judging outputs, and integrating the result. Predicted in *Dream Machine* Issue 13 (January 2026) as the dominant role of 2026; documented across the chapters above. See Chapter 11.

Pipeline of authorship. The full chain from creative intent to delivered work, broken down into discrete functions (writing, direction, performance, image-making, sound, edit, distribution, marketing). The point of the **Human–AI Agency Continuum** is that each link in this chain has its own agency line.

Position One (All-in). The strategic posture of legacy studios that have decided to integrate AI aggressively across all production functions, betting that early-adopter advantage will compound. Netflix’s “all in” framing, October 2025. See Chapter 7.

Position Two (AI-native). The strategic posture of new entrants that build their production pipelines AI-first from inception — Imaginae Studios, Wonder Studios, Obsidian Studio, Asteria, Wonder, Chapter41, Kartel. The category exploded in scale through autumn 2025 and winter 2026. See Chapter 7.

Position Three (Refusal). The strategic posture of creative organisations that have explicitly excluded generative AI from their work. Jagex, Larian, Games Workshop, Hooded Horse, Aardman (qualified), Pocketpair. Cultural authority is preserved as the principal asset. See Chapter 7.

Position Four (Middle). *AI in the workflow, not in the work.* The strategic posture — taken by Sony, Bethesda, Amazon (in *House of David*), Aardman (in qualified form), and an increasing number of major studios — that uses AI to augment production pipelines while preserving human creative intent in the moments the audience sees. The book’s prediction for where most surviving major studios land by 2030. See Chapter 7.

Provenance. The chain of custody of a creative work from its origin (capture, performance, writing, sketching) to its delivered form. Technical standards (C2PA, SynthID) and policy frameworks (Cannes Disclosure Standard, SAG-AFTRA AI protections) collectively constitute the *provenance infrastructure* the book argues is critical for the next decade. See Chapter 12.

Shadow AI. The use of AI tools by employees outside their employer’s official tooling, processes and accounting. Documented in 2025 workplace research as encompassing approximately half of the U.S. workforce. The principal symptom of **Coordination Collapse**. See Chapter 13.

Slop Ceiling. The empirical pattern, established across multiple sectors by spring 2026, in which AI-generated creative content can be produced in massive volume but consistently fails to capture audience attention proportionate to that volume. Anchored in the **44%/3%** ratio Deezer published in April 2026 (44% of daily uploads AI; under 3% of streams). One of the central claims of the book. See Chapter 5.

Synthetic sincerity. The category of creative work that is openly synthetic but made with serious creative intent and is not pretending to be something else. Named after Marc Isaacs’ 2025 IDFA documentary. Audiences distinguish *synthetic sincerity* from *synthetic cynicism* at the speed of a swipe. See Chapter 4.

Tilly Tax (informal). The collection of contract provisions in SAG-AFTRA’s spring 2026 agreement requiring compensation, consent and residuals when AI replicas of human performers are used. Named after the Tilly Norwood controversy of September 2025 that catalysed the broader negotiation. See Chapters 5, 10.

Watermark. A persistent identifier embedded in AI-generated outputs by the producing system, intended to allow downstream detection that content is synthetic. SynthID (Google) and similar systems became standard across major platform tools through 2025–26.

World model. A class of generative AI system that produces navigable three-dimensional environments rather than flat output. Marble (World Labs, public release November 2025) was the first commercial product in the category; Google DeepMind’s Genie 3, Meta’s WorldGen, Luma’s UNI-1, Tencent’s Hunyuan World, SpAltial’s ECHO and others followed within months. See Chapter 8.

* * *

If a term in the book did substantial work but does not appear in this glossary, please tell us so we can improve it in the next edition.

Appendix C — Bibliography by Topic

The book's footnotes are sequential per chapter; this appendix organises the same sources thematically, so a reader pursuing a specific question across the period covered by the book can find every source it touches in one place.

The full corpus — 1,438 successfully fetched and archived articles — is preserved in JSON form in `Research/scraped/` of the source repository; the manifest enumerating every URL and its capture status lives at `Research/manifest.json`.

* * *

I. The watershed week, September–October 2025

Sora 2 launches. Tilly Norwood is announced. The union response sets the contract reference point for the next year.

- OpenAI, “Sora 2 is here,” 30 September 2025. <https://openai.com/index/sora-2/>
- Variety, “SAG-AFTRA Condemns Tilly Norwood: AI Actress Is Not an Actor.” <https://variety.com/2025/film/news/sag-aftra-tilly-norwood-ai-actress-1236534779/>
- The Hollywood Reporter, “U.K. Union Equity Condemns Tilly Norwood: ‘AI Tool, Not a Performer’.” <https://www.hollywoodreporter.com/movies/movie-news/tilly-norwood-ai-actress-uk-union-equity-sag-aftra-debate-1236391739/>
- Variety, “Tilly Norwood Slammed by Equity as AI Tool, Concerned About Origin.” <https://variety.com/2025/film/global/tilly-norwood-slammed-equity-ai-tool-concerned-origin-1236537042/>
- CNN, “Tilly Norwood: Hollywood is fuming over a new ‘AI actress’.” <https://www.cnn.com/2025/09/30/tech/hollywood-ai-actor-backlash>
- NBC News, “Tilly Norwood, fully AI ‘actor,’ blasted by actors union SAG-AFTRA.” <https://www.nbcnews.com/pop-culture/pop-culture-news/tilly-norwood-fully-ai-actor-blasted-actors-union-sag-aftra-devaluing-rcna234685>
- The Guardian, “OpenAI launch of video app Sora plagued by violent and racist images.” <https://www.theguardian.com/us-news/2025/oct/04/openai-sora-violence-racism>
- WUFT (Florida), “Kiss reality goodbye: AI-generated social media has arrived.” <https://www.wuft.org/2025-10-03/kiss-reality-goodbye-ai-generated-social-media-has-arrived>
- Deadline, “Tilly Norwood Creator on Backlash, Says More AI Actors Coming.” <https://deadline.com/2025/11/tilly-norwood-creator-interview-backlash-more-ai-actors-coming-1236601334/>
- No Film School, “James Cameron Says AI Is ‘Never Going to Take the Place’ of Humans.” <https://nofilmschool.com/james-cameron-ai>
- The Guardian, “James Cameron says AI actors are ‘horrifying to me’.” <https://www.theguardian.com/film/2025/dec/01/james-cameron-says-ai-actors-are-horrifying-to-me>

II. Agentic AI and the orchestrator economy

- OpenAI, “Introducing AgentKit.” <https://openai.com/index/introducing-agentkit/>

- TechCrunch, “OpenAI launches AgentKit to help developers build and ship AI agents.” <https://techcrunch.com/2025/10/06/openai-launches-agentkit-to-help-developers-build-and-ship-ai-agents/>
- TechCrunch, “Anthropic launches interactive Claude apps, including Slack and other workplace tools.” <https://techcrunch.com/2026/01/26/anthropic-launches-interactive-claude-apps-including-slack-and-other-workplace-tools/>
- GamesRadar, “EA reportedly pushes 15,000 employees to use AI as a ‘thought partner’.” <https://www.gamesradar.com/games/even-under-usd20-million-in-debt-ea-reportedly-pushes-15-000-employees-to-use-ai-as-a-thought-partner-for-everything-from-character-art-to-playtesting/>
- Digiday, “AI agent developers have become adland’s in-demand role.” <https://digiday.com/marketing/ai-agent-developers-have-become-adlands-in-demand-role/>
- Eurogamer, “Falcom: ‘Tasks that previously took 2–3 hours can now be completed in 10 minutes.’” <https://www.eurogamer.net/falcom-is-the-latest-developer-to-buy-into-the-ai-hype-machine>
- Forbes, “AI Is Changing How Creators Work And Earn.” <https://www.forbes.com/sites/kolawolesamueladebayo/2025/12/22/how-ai-is-changing-how-creators-work-and-earn/>
- Google DeepMind, “Introducing Gemini Omni — unified multimodal model.” <https://blog.google/technology/google-deepmind/gemini-omni/>
- Google, “Gemini Spark — developer toolkit for autonomous agents.” <https://blog.google/technology/developers/gemini-spark/>
- Google, “Antigravity — agentic coding environment.” <https://antigravity.google/>
- Google, “Google Flow — agent-based creative workflows.” <https://flow.google/>
- Google, “Official skills for AI agents.” <https://github.com/google/agent-skills>
- Adweek, “Netflix ad tools could see ‘agentic AIs talking to each other’.” <https://www.adweek.com/media/netflix-ad-tools-agentic-ais-talking-to-each-other/>
- Fortune, “AI startup Viktor raises \$75 million to put a virtual ‘coworker’ in Slack and Teams.” <https://fortune.com/2026/05/19/ai-startup-viktor-75-million-virtual-coworker-slack-teams/>
- Tencent, “Ardot — AI-native design agent platform.” <https://ardot.tencent.com/>
- Anthropic, “Claude available as a partner node in ComfyUI.” <https://www.anthropic.com/news/claude-comfyui-partner-node>
- Runway, “Aleph 2.0 — single-frame video editing propagated across the clip” (June 2026). *Dream Machine* Issue 31.
- Alibaba Cloud, “WonderClip — all-in-one AI video creation platform” (June 2026). *Dream Machine* Issue 31.
- Sony, “Vid-CamEdit — post-capture camera trajectory editing” (June 2026). *Dream Machine* Issue 31.
- Nvidia, “PiD (Pixel Diffusion Decoder)” (June 2026). *Dream Machine* Issue 31.

- TripoAI, “Project Eden — persistent multiplayer world model decoupling world state from rendering” (June 2026). *Dream Machine* Issue 31.
- Sorcerorx, “Game Creation Suite — build any game from a prompt” (June 2026). *Dream Machine* Issue 31.

III. Bot economy, model collapse, and the Dead Internet

- Imperva, “2025 Bad Bot Report.” <https://www.imperva.com/blog/2025-imperva-bad-bot-report-how-ai-is-supercharging-the-bot-threat/>
- Cloudflare, “The crawl-to-click gap: Cloudflare data on AI bots.” <https://blog.cloudflare.com/crawlers-click-ai-bots-training/>
- Cloudflare, “A deeper look at AI crawlers.” <https://blog.cloudflare.com/ai-crawler-traffic-by-purpose-and-industry/>
- Grand View Research, “Generative AI in Content Creation Market Size Report, 2030.” <https://www.grandviewresearch.com/industry-analysis/generative-ai-content-creation-market-report>
- Wikipedia, “Dead Internet Theory.” https://en.wikipedia.org/wiki/Dead_Internet_theory
- Futurism, “Scientists Created an Entire Social Network Where Every User Is a Bot.” <https://futurism.com/social-network-ai-intervention-echo-chamber>
- AI News, “AI causes reduction in users’ brain activity, MIT.” <https://www.artificialintelligence-news.com/news/ai-causes-reduction-in-users-brain-activity-mit/>

IV. The slop ceiling — music

- Deezer, “AI-generated tracks now represent 44% of all new uploaded music.” <https://newsroom-deezer.com/2026/04/ai-generated-tracks-represent-44-of-new-uploaded-music/>
- Music Business Worldwide, “75,000 AI-generated tracks now flood Deezer daily.” <https://www.musicbusinessworldwide.com/75000-ai-generated-tracks-now-flood-deezer-daily-representing-44-of-all-new-music-uploaded-to-the-platform-says-streamer/>
- Deezer / Ipsos, “AI fools 97% of listeners.” <https://newsroom-deezer.com/2025/11/deezer-ipsos-survey-ai-music/>
- Musically, “50,000 AI-music tracks are now uploaded to Deezer every day.” <https://musically.com/2025/11/12/50000-ai-music-tracks-are-now-uploaded-to-deezer-every-day/>
- Musically, “UMG boss slams ‘exponential growth of AI slop’ on streaming services.” <https://musically.com/2026/01/09/umg-boss-slams-exponential-growth-of-ai-slop-on-streaming-services/>
- Billboard, “AI Artist Xania Monet Climbs the Charts — And Signs a Multimillion-Dollar Record Deal.” <https://www.billboard.com/pro/ai-music-artist-xania-monet-multimillion-dollar-record-deal/>
- NPR, “Breaking Rust is a hot new country act on the Billboard charts. It’s powered by AI.” <https://www.npr.org/2025/11/10/nx-s1-5604320/breaking-rust-is-a-hot-new->

country-act-on-the-billboard-charts-its-powered-by-ai

- BBC News, “*Sienna Rose: AI suspicions surround mysterious singer.*”
<https://www.bbc.co.uk/news/articles/cq6v83gq66eo>
- Bain & Company, “*In an AI Age, People Still Want the Radio Star.*”
<https://www.bain.com/insights/in-an-ai-age-people-still-want-the-radio-star/>
- Soultracks, “*A.I.-generated music is catchy, familiar... and boring.*”
<https://soultracks.com/news-ai-generated-music-is-catchy-boring/>
- Stereogum, “*Bandcamp Bans AI Music.*” <https://stereogum.com/2485199/bandcamp-bans-ai-music/news>
- The Independent, “*AI-generated song banned from Swedish charts.*”
<https://www.independent.co.uk/tv/news/ai-music-song-banned-sweden-spotify-b2901627.html>
- Digital Music News, “*YouTube CEO Puts ‘Managing AI Slop’ on the Priority List for 2026.*” <https://www.digitalmusicnews.com/2026/01/22/youtube-ceo-ai-slop-2026-comments/>
- Musically, “*Report: 56.9% of new independent songs in China are AI-generated.*”
<https://musically.com/2026/01/05/report-56-9-of-new-independent-songs-in-china-are-ai-generated/>
- The Guardian, “*Paul McCartney joins music industry protest against AI with silent track.*” <https://www.theguardian.com/music/2025/nov/17/the-sound-of-silence-why-theres-barely-anything-there-in-paul-mccartney-new-release>
- MusicTech, “*Jack Antonoff brands AI music makers as ‘godless whores’.*”
<https://musitech.com/news/industry/jack-antonoff-ai-music-makers-godless-whores/>
- Billboard, “*The Real Story Behind The AI Song That Knocked Tyla Off No. 1 On Billboard Afrobeats Chart.*” <https://www.billboard.com/pro/ai-song-knocked-tyla-off-no-1-afrobeats/>
- Variety, “*WMG’s Robert Kyncl: ‘AI Resistance’ Setting Music Sector Back.*”
<https://variety.com/2026/music/news/wmg-robert-kyncl-ai-resistance-1236748901/>
- Musically, “*Sondo AI says it has 10m global users.*”
<https://musically.com/2026/05/18/sondo-ai-10m-global-users/>
- Musically, “*BPI sets out transparency and sovereignty demands to secure AI licensing boom.*” <https://musically.com/2026/05/19/bpi-transparency-sovereignty-ai-licensing-boom/>
- Musicbusinessworldwide, “*DistroKid asks creators if their music is AI-generated as self-disclosure moves out of beta.*” <https://www.musicbusinessworldwide.com/distrokid-ai-self-disclosure-out-of-beta/>
- Musicbusinessworldwide, “*Splice inks ‘Responsible AI’ deal with ElevenLabs.*”
<https://www.musicbusinessworldwide.com/splice-elevenlabs-responsible-ai-deal/>
- Rolling Stone, “*The Rolling Stones Release ‘In the Stars’ — AI De-Aging Music Video.*”
<https://www.rollingstone.com/music/music-news/rolling-stones-in-the-stars-ai-de-aging-video-1235142200/>

- MusicTech, “Tamber — ethically trained AI music tool with arm gesture control.” <https://musictech.com/news/gear/tamber-ai-ethically-trained-arm-gestures/>
- Stability AI, “Stable Audio 3.0 released.” <https://stability.ai/news/stable-audio-3-0-released>
- Sony AI, “Woosh — sound effect foundation model.” <https://ai.sony/blog/woosh-sound-effect-foundation-model/>
- Mirelo, “SFX 1.6 — edit sound, not just generate.” <https://mirelo.ai/sfx-1-6>
- ElevenLabs, “Music v2 — next-generation generative music model” (June 2026). *Dream Machine* Issue 31.
- Apple Music, “AI-generated songs make up less than 1% of plays on platform” (June 2026). *Dream Machine* Issue 31.
- Spotify + Universal Music Group, “Landmark deal lets Spotify Premium subscribers create AI covers and remixes of UMG-catalogue tracks” (June 2026). *Dream Machine* Issue 31.
- Daniel Ek (Spotify), “AI music is better than ‘slop’” — interview / statement (June 2026). *Dream Machine* Issue 31.
- Tonada, “Swedish startup supplies AI music for retailers” (June 2026). *Dream Machine* Issue 31.
- Estate of Ozzy Osbourne, “Ozzy to live on as an AI avatar that talks with fans” (June 2026). *Dream Machine* Issue 31.

V. Copyright, consultation and the legal architecture

- UK Department for Science, Innovation and Technology, “Statement of Progress on Copyright and AI.” <https://www.gov.uk/government/publications/copyright-and-artificial-intelligence-progress-report/copyright-and-artificial-intelligence-statement-of-progress-under-section-137-data-use-and-access-act>
- IPWatchdog, “Respondents to UK AI Consultation Overwhelmingly Want AI Companies to License Copyrighted Works in All Cases.” <https://ipwatchdog.com/2025/12/16/respondents-uk-ai-consultation-overwhelmingly-want-ai-companies-license-copyrighted-works-all-cases/>
- UCL Copyright Queries, “UK government publishes progress statement on AI and copyright consultation.” <https://blogs.ucl.ac.uk/copyright/2025/12/23/uk-government-publishes-progress-statement-on-ai-and-copyright-consultation/>
- Hogan Lovells, “Copyright and AI: UK government publishes statement of progress.” <https://www.hoganlovells.com/en/publications/copyright-and-ai-uk-government-publishes-statement-of-progress>
- Music Business Worldwide, “Wixen files \$50m copyright suit against Meta.” <https://www.musicbusinessworldwide.com/wixen-files-50m-copyright-suit-against-meta-claims-tech-giant-wants-to-replace-songwriters-with-ai/>
- Music Business Worldwide, “Meta hits back against Wixen Music Publishing’s \$70m ‘AI replacement’ infringement suit.” <https://www.musicbusinessworldwide.com/meta-wixen-music-publishing-70-million-ai-replacement-lawsuit/>

- EDM.com, “Biggest Theft in Music History’: Rights Group Sues Suno.” <https://edm.com/gear-tech/rights-group-sues-suno-copyright-infringement/>
- Complete Music Update, “Johnny Cash estate uses ELVIS Act to sue Coke.” <https://completemusicupdate.com/johnny-cash-estate-uses-elvis-act-to-sue-coke-over-tribute-act-ad-soundtrack/>
- Dr. Barry Scannell, analysis of GEMA v. OpenAI ruling. https://www.linkedin.com/posts/dr-barry-scannell-bbb5aa207_in-a-major-ruling-for-european-copyright-share-7393957246386323457-8bbx
- Reuters, “European lawmakers seek EU-wide minimum age to access AI chatbots.” <https://www.reuters.com/legal/litigation/european-lawmakers-seek-eu-wide-minimum-age-access-ai-chatbots-social-media-2025-11-26/>
- The Verge, “New York’s new law forces advertisers to say when they’re using AI avatars.” <https://www.theverge.com/news/842848/new-york-law-ai-advertisements-sag-aftra-labor>
- Digital Music News, “Nearly 800 Creatives Sign Responsible AI Declaration — ‘Stealing Our Work Is Not Innovation’.” <https://www.digitalmusicnews.com/2026/01/22/stealing-isnt-innovation/>

VI. Studio strategy

- CNBC, “Netflix ‘all in’ on leveraging AI.” <https://www.cnbc.com/2025/10/22/netflix-all-in-on-leveraging-ai-in-its-streaming-platform.html>
- Futurism, “Lionsgate’s Attempt to Create Movies Using AI Has Crumbled Into Disaster.” <https://futurism.com/artificial-intelligence/lionsgate-movies-ai>
- The Guardian, “Disney to invest \$1bn in OpenAI.” <https://www.theguardian.com/business/2025/dec/11/disney-open-ai-sora-video-deal>
- The Hollywood Reporter, “Disney+ to Allow User-Generated Content Via AI.” <https://www.hollywoodreporter.com/business/digital/disney-plus-gen-ai-user-generated-content-1236426135/>
- Deadline, “Amazon’s Head Of Original Sports Content Matt Newman Moves Into A.I.” <https://deadline.com/2025/11/amazon-ai-studios-matt-newman-1236603477/>
- Wired, “Amazon’s House of David Used Over 350 AI Shots in Season 2.” <https://www.wired.com/story/amazons-house-of-david-used-over-350-ai-shots-in-season-2-its-creator-isnt-sorry/>
- Indiewire, “New AI Studio Obsidian Forms Partnership with Imagine Entertainment.” <https://www.indiewire.com/news/business/obsidian-studio-ai-production-company-imagine-entertainment-1235158619/>
- The Hollywood Reporter, “Fremantle Names Boss of New AI Label Imaginae Studios.” <https://www.hollywoodreporter.com/business/digital/fremantle-names-ceo-new-ai-label-imaginae-studios-1236396579/>
- PC Gamer, “Palworld studio Pocketpair: ‘We don’t believe in it’.” <https://www.pcgamer.com/software/ai/palworld-studio-pocketpair-says-its-new-publishing-division-wont-handle-games-that-use-generative-ai-we-dont-believe-in-it/>

- Variety, “Guillermo del Toro: ‘I’d rather die’ than use generative AI.”
<https://variety.com/2025/film/news/guillermo-del-toro-rather-die-generative-ai-frankenstein-1236561316/>
- The Hollywood Reporter, “Leonardo DiCaprio Says AI Can’t Be Art.”
<https://www.hollywoodreporter.com/movies/movie-news/leonardo-dicaprio-ai-cant-be-art-no-humanity-1236445405/>
- NME, “Jenna Ortega: ‘very easy to be terrified’ of AI.”
<https://www.nme.com/news/jenna-ortega-says-its-very-easy-to-be-terrified-of-ai-in-filmmaking-3913926>
- GamesRadar, “Wallace and Gromit creator: Aardman will ‘embrace the technology’ of AI but will be ‘very cautious’.” <https://www.gamesradar.com/entertainment/animation-movies/wallace-and-gromit-creator-says-beloved-animation-studio-aardman-will-embrace-the-technology-of-ai-but-will-be-very-cautious-not-to-lose-our-values/>
- gamesindustry.biz, “‘AI was an expensive mistake’: Charles Cecil on Broken Sword.”
<https://www.gamesindustry.biz/ai-was-an-expensive-mistake-charles-cecil-on-innovation-insolvency-and-broken-sword>
- gamesindustry.biz, “RuneScape maker Jagex says it will never use generative AI.”
<https://www.gamesindustry.biz/runescape-maker-jagex-says-it-will-never-use-generative-ai-to-make-in-game-content>
- Decrypt, “‘Warhammer 40,000’ Maker Games Workshop Rules Out Generative AI.”
<https://decrypt.co/354482/warhammer-40000-maker-games-workshop-rules-out-generative-ai>
- Niche Gamer, “Larian Studios backs off from gen AI.” <https://nichegamer.com/larian-studios-backs-off-from-gen-ai/>
- PC Gamer, “Todd Howard says AI can’t replace human ‘creative intention’.”
<https://www.pcgamer.com/gaming-industry/todd-howard-says-ai-cant-replace-human-creative-intention-but-its-part-of-bethesdas-toolset-for-how-we-build-our-worlds-or-check-things/>
- Hollywood Reporter, “Netflix INKubator AI animation studio for ‘feature-quality’ shorts.” <https://www.hollywoodreporter.com/business/business-news/netflix-ai-animation-studio-inkubator-1236592110/>
- Forbes, “Wonder Studios — \$50M British studio aiming to become ‘A24 of AI’.”
<https://www.forbes.com/sites/charliefink/2026/05/18/meet-wonder-studios-the-50m-british-studio-striving-to-become-the-a24-of-ai-production/>
- Runway, “Runway Japan — Tokyo office and \$40M investment.”
<https://runwayml.com/blog/runway-japan>
- Hollywood Reporter, “Cannes 2026: filmmakers shift towards cautious AI acceptance.”
<https://www.hollywoodreporter.com/business/business-news/cannes-2026-ai-acceptance-1236592488/>
- Variety, “AI Dominates Cannes Buzz as Filmmakers Grudgingly Accept It.”
<https://variety.com/2026/film/festivals/ai-cannes-2026-filmmakers-accept-1236748402/>

- *Variety*, “Peter Jackson: AI is basically the next wave of special effects.” <https://variety.com/2026/film/news/peter-jackson-ai-special-effects-1236748120/>
- *Variety*, “Kling AI + Evolutionary Films on ‘Minibots’ at Cannes Market.” <https://variety.com/2026/film/news/kling-ai-evolutionary-films-minibots-cannes-1236748590/>
- *Cartoon Brew*, “Former Disney Animator: Marvel layoffs point to a bigger AI shift.” <https://www.cartoonbrew.com/artist-rights/former-disney-animator-marvel-layoffs-ai-shift-260601.html>
- *Hollywood Reporter*, “Bobby Berk: AI will make reality TV and ‘verifiably human content’ more valuable.” <https://www.hollywoodreporter.com/tv/tv-news/bobby-berk-ai-reality-tv-1236592920/>
- *PC Gamer*, “Take-Two’s CEO: AI ‘datasets by their very nature are backward looking’.” <https://www.pcgamer.com/games/take-two-ceo-ai-not-making-hits-backward-looking/>

VII. World models and the new medium

- World Labs, “Bringing Marble to Life.” <https://www.worldlabs.ai/case-studies/bringing-marble-to-life>
- TechCrunch, “Fei-Fei Li’s World Labs speeds up the world model race with Marble.” <https://techcrunch.com/2025/11/12/fei-fei-lis-world-labs-speeds-up-the-world-model-race-with-marble-its-first-commercial-product/>
- Google DeepMind, “Genie 3: A new frontier for world models.” <https://deepmind.google/blog/genie-3-a-new-frontier-for-world-models/>
- Google blog, “Project Genie: AI world model now available for Ultra users in U.S.” <https://blog.google/innovation-and-ai/models-and-research/google-deepmind/project-genie/>
- Meta, “WorldGen — text-to-immersive-3D-worlds.” <https://www.facebook.com/LifeAtMeta/videos/research-update-worldgen-text-to-immersive-3d-worlds/1879077432692421/>
- SpAItial, “Transform Images into Interactive 3D Worlds.” <https://www.spaitial.ai/>
- Stanford AI Lab, “WonderZoom.” <https://wonderzoom.github.io/>
- NVIDIA/Stanford, “NitroGen.” <https://nitrogen.minedojo.org/>
- DeepMind, “SIMA 2: An Agent that Plays, Reasons, and Learns.” <https://deepmind.google/blog/sima-2-an-agent-that-plays-reasons-and-learns-with-you-in-virtual-3d-worlds/>
- Radiance Fields, “Apple Confirms Personas use Gaussian Splatting.” <https://radiancefields.com/apple-confirms-personas-use-gaussian-splatting>
- Sony Pictures’ virtual-production case study (Marble integration). https://www.linkedin.com/posts/brent-liang_tech-media-launch-ugcPost-7394911181091692546-TyUz
- NVIDIA, “SANA-WM — 2.6B open-source world model with 60-second video generation and camera control.” <https://huggingface.co/collections/nvidia/sana-wm>

- Odyssey, “*Starchild-1 — first real-time multimodal world model.*”
<https://odyssey.ml/introducing-starchild-1>
- Odyssey, “*Agora-1 — four-player AI-generated world built on a 1997 shooter.*”
<https://odyssey.ml/introducing-agora-1>
- Google DeepMind, “*Project Genie + Street View — simulate real-world places.*”
<https://deepmind.google/discover/blog/project-genie-street-view/>
- Google Labs, “*Infinite Scaler — browser multiplayer generative game.*”
<https://blog.google/technology/google-labs/infinite-scaler/>
- Apple Machine Learning Research, “*Apple Headsup — 3D Gaussian Head Reconstruction from Multi-View Captures.*” <https://machinelearning.apple.com/research/apple-headsup-3d-gaussian-head>
- WorldLens VR, “*AI-powered 3D depth on Quest for Google Street View.*”
<https://www.uploadvr.com/worldlens-vr-quest-street-view-3d-depth/>
- *gamesindustry.biz*, “*Meet Seed, the planet-sized society simulator.*”
<https://www.gamesindustry.biz/meet-seed-planet-sized-society-simulator>

VII.a Creator economy and labour-market sentiment

The June 2026 Epidemic Sound Future of the Creator Economy report, and the Coordination Collapse, Process Trap and Doomer Mistake deep dives that re-anchor the labour-market analysis of the second half of the book.*

- Epidemic Sound, “*Future of the Creator Economy Report 2026*” — survey of 3,000 professional creators across UK and US; 94% of creators use AI tools, 72% expect to increase usage in the next 12 months, 89% feel pressure to use AI to keep up with industry expectations, 75% believe human-created content will become a premium product, 83% say human-made sound creates a stronger emotional connection. *Dream Machine* Issue 31.
- *Musically*, “*Creator attitudes towards AI: nuanced embrace plus transparency, licensing and identity demands.*” Coverage of the same report. *Dream Machine* Issue 31.
- O’Hare, J. (2026). *The Coordination Collapse: How AI Is Dissolving the Information Tax That Created Hierarchy — And What Replaces It*. DreamLab AI Consulting Ltd, April 2026. Preserved as Appendix K.
- *The Process Trap: Generative Abundance and the Equivocation of the Creative*. Preserved as Appendix J.
- *The Doomer Mistake: Why the Case for AI-Driven Civilizational Collapse Keeps Losing the Argument with the Data*. Preserved as Appendix I. Underlying empirical work draws on Acemoglu (2025), Brynjolfsson, Li & Raymond (2025), Brynjolfsson, Chandar & Chen (2025), the 2026 Stanford AI Index, Dallas Fed, St. Louis Fed, Bessen (2015), Autor (2016), AEL, Susskind (2020), and CCIA (2026).
- *The Programmable Brand: How AI Is Reshaping IP Licensing in the Generative Economy*. *Dream Machine Deep Dive* №14, 4 June 2026. Preserved as Appendix L. Five-licensing-model comparison (Hasbro Sixth Wall / Behavioural Licensing; Spotify–UMG embedded remixing; Disney–OpenAI as cautionary tale; the independent-artist commons-

and-DAO models — Grimes / Elf.tech, Holly Herndon / Holly+, ElevenLabs heritage voices, Meta likeness-rental anti-pattern; Suno / Alexa as the *create-to-consume* demand floor) against the three Cs (consent, credit, compensation). Theoretical foundations: Coase (1937), Williamson (1979), Barney (1991) VRIN resource-based view, Henry Jenkins on participatory culture. Primary sources:

- Hasbro BusinessWire:
<https://www.businesswire.com/news/home/20260603297922/en/>
- *The Wrap*, Hasbro Sixth Wall launch: <https://www.thewrap.com/industry-news/business/hasbro-ai-studio-sixth-wall-launch/>
- *Screen Daily*, Hasbro launches subsidiary AI studio:
<https://www.screendaily.com/news/hasbro-launches-subsidiary-ai-studio-sixth-wall-to-exploit-ip/5217380.article>
- *BriefGlance*, Sixth Wall and behavioural licensing:
<https://briefglance.com/articles/hasbros-sixth-wall-taming-the-ai-wild-west-with-behavioral-licensing>
- Spotify Newsroom, Spotify–UMG: <https://newsroom.spotify.com/2026-05-21/universal-music-group-spotify-licensing-agreements-fan-made-covers-remixes/>
- *TechCrunch*, Spotify–UMG: <https://techcrunch.com/2026/05/21/spotify-and-universal-music-strike-deal-allowing-fan-made-ai-covers-and-remixes/>
- *Medianama*, Spotify–UMG: <https://www.medianama.com/2026/05/223-spotify-umg-fans-create-licensed-ai-covers-remixes/>
- *Ropes & Gray*, “The \$1 Billion Deal That Wasn’t”:
<https://www.ropesgray.com/en/insights/viewpoints/the-deal-that-wasnt-what-disney-and-openai-teach-us-about-strategic-investments>
- *TheStreet* / AOL on Disney exit: <https://www.thestreet.com/technology/openai-is-shutting-down-sora-and-the-disney-deal-is-off>
- *Tech-Insider* on Sora shutdown economics: <https://tech-insider.org/openai-sora-shutdown-disney-deal-ai-video-2026/>
- *Yahoo Finance* on Disney cease-and-desist to Google pre-OpenAI:
<https://finance.yahoo.com/sectors/technology/articles/disney-cancels-1b-deal-openai-043000427.html>
- Goldman Sachs on creator economy:
<https://www.goldmansachs.com/insights/articles/the-creator-economy-could-approach-half-a-trillion-dollars-by-2027>
- *Kill the DJ* on Grimes / Elf.tech and artist-voice models: <https://killthedj.com/how-artists-lend-their-ai-voices-to-creators-fans/>
- CBS News on ElevenLabs heritage voices: <https://www.cbsnews.com/news/ai-voices-dead-celebrities-judy-garland-james-dean-burt-reynolds/>
- *El País* on Meta celebrity chatbots (2023):
<https://english.elpais.com/technology/2023-11-13/from-paris-hilton-to-snoop-dogg-meta-to-launch-speaking-ai-chatbots-played-by-celebrities.html>

- *Billboard* on Amazon + Suno: <https://www.billboard.com/pro/amazon-partners-ai-music-company-suno-alexa/>
- *Digital Music News* on Amazon + Suno: <https://www.digitalmusicnews.com/2025/02/27/amazon-suno-team-up-for-gen-ai-songs/>
- MIDiA Research on *create to consume*: <https://www.midiaresearch.com/blog/suno-alexa-and-the-rise-of-create-to-consume>
- Dell’Acqua et al. (2026). Preregistered Boston Consulting Group experiment on the jagged technological frontier (n = 758 consultants).
- Mollick, E. (2023). “*Secret cyborgs*” and the shadow-AI phenomenon. Coverage in [Appendix D](#) and [Appendix K](#).
- Edmondson, A. & Lei, Z. (2025). Psychological safety and AI adoption in teams.
- Epoch AI / Ipsos KnowledgePanel (March 2026). Probability-based survey of 2,021 US adults (665 employed past-week AI users).

VIII. The platform layer and the AI-native toolchain

- *Creative Boom*, “*Adobe puts AI in everything, everywhere, all at once.*” <https://www.creativeboom.com/news/adobe-is-putting-ai-in-everything-everywhere-all-at-once/>
- Adobe, “*Adobe MAX 2025: Firefly Foundry.*” <https://news.adobe.com/news/2025/10/adobe-max-2025-firefly-foundry>
- Adobe, “*Adobe MAX 2025: Express AI Assistant.*” <https://news.adobe.com/news/2025/10/adobe-max-2025-express-ai-assistant>
- *Wired*, “*Adobe’s Corrective AI Can Change the Emotions of a Voice-Over.*” <https://www.wired.com/story/adobe-max-sneaks-2025-corrective-ai/>
- PYMNTS, “*Adobe Lets Users Design and Edit Using ChatGPT.*” <https://www.pymnts.com/artificial-intelligence-2/2025/adobe-lets-users-design-and-edit-using-chatgpt/>
- Adobe blog, “*Edit with Photoshop in ChatGPT.*” <https://blog.adobe.com/en/publish/2025/12/10/edit-photoshop-chatgpt>
- *Campaign Brief*, “*WPP launches WPP Open Pro.*” <https://campaignbrief.com/wpp-launches-ai-powered-marketing-platform-wpp-open-pro/>
- *Digiday*, “*WPP expands AI capabilities.*” <https://digiday.com/media-buying/agencies-continue-to-expand-ai-capabilities-to-boost-brand-performance/>
- *SiliconAngle*, “*Higgsfield raises \$80M on \$1.3B valuation.*” <https://siliconangle.com/2026/01/15/higgsfield-raises-80m-1-3b-valuation-scale-ai-video-platform/>
- *TechCrunch*, “*Synthesia hits \$4B valuation.*” <https://techcrunch.com/2026/01/26/synthesia-hits-4b-valuation-lets-employees-cash-in/>
- *Sifted*, “*Synthesia rejects \$3bn Adobe acquisition offer.*” <https://sifted.eu/articles/synthesia-acquisition-offer>

- ComfyUI, “We raised \$17 million to build an OS for creative AI.” https://www.linkedin.com/posts/comfyui_we-raised-17-million-to-build-an-os-for-ugcPost-7373743341236236288-wkCc
- Beeples Canvas, “Generative AI meets compositing.” <https://www.beeple-canvas.com/>
- MacRumors, “New CarPlay Audio Apps Add AI Music and Global Radio.” <https://www.macrumors.com/2026/05/19/carplay-ai-music-global-radio/>
- The Drum, “World’s First AI-Driven Immersive Shoppable Catwalk.” <https://www.thedrum.com/news/2026/05/18/how-world-s-first-ai-driven-immersive-shoppable-catwalk-redefined-future-brand>
- The Drum, “David Beckham designs ‘Henchester United’ chicken coop in Lenovo AI ad.” <https://www.thedrum.com/news/2026/05/18/david-beckham-henchester-united-chicken-coop-lenovo-ai-ad>
- Amazon, “Alexa+ debuts feature to create AI-generated podcasts.” <https://www.aboutamazon.com/news/devices/alexa-plus-ai-podcasts>
- ECABridge, “Unreal Engine MCP integration.” <https://ecabridge.dev/>
- Video Games Chronicle, “Epic Games Veteran on AI-Heavy ‘Fully European’ Game Engine.” <https://www.videogameschronicle.com/news/epic-games-veteran-ai-heavy-fully-european-game-engine/>

IX. Authenticity, provenance and disclosure

- Digital Music News, “Instagram Chief Says We Should ‘Fingerprint Real Media’.” <https://www.digitalmusicnews.com/2026/01/05/instagram-chief-ai-slop-comments/>
- WebProNews, “Instagram Head Warns AI Images Erode Trust.” <https://www.webpronews.com/instagram-head-warns-ai-images-erode-trust-calls-for-verification-standards/>
- Google DeepMind, “Verify AI-generated videos in the Gemini app.” https://www.linkedin.com/posts/googledeepmind_verify-google-ai-generated-videos-in-the-activity-7407748300688478208-fJgW
- The Hollywood Reporter, “Emmys Set AI Guidelines for 2026 Awards.” <https://www.hollywoodreporter.com/tv/tv-news/emmys-ai-guidelines-2026-awards-1236468434/>
- Lawyer Monthly, “Matthew McConaughey Trademarks Voice and Image Against AI.” <https://www.lawyer-monthly.com/2026/01/matthew-mcconaughey-protects-voice-image-ai/>
- Adweek, “Meet the \$1.3 Billion Startup Behind Madonna and Will Smith’s AI Video.” <https://www.adweek.com/media/higgsfield-ai-marketing-startup/>
- Marketing Week, “You can’t dismiss AI ads as slop when they’re winning in testing.” <https://www.marketingweek.com/dismiss-ai-ads-winning-creative-effectiveness/>
- Variety, “George Clooney: AI Actors Will Struggle to Be Movie Stars.” <https://variety.com/2025/scene/columns/george-clooney-ai-actors-movie-stars-1236579661/>

- *The Hollywood Reporter*, “Synthetic Sincerity on AI Characters, Authenticity (IDFA).” <https://www.hollywoodreporter.com/movies/movie-news/synthetic-sincerity-film-idfa-ai-authenticity-interview-1236426180/>
- *Variety*, “AI-Generated Images Lead to Audiences’ Distrust, Threaten Documentary.” <https://variety.com/2025/film/festivals/ai-generated-images-threaten-future-of-documentary-1236583466/>
- *Hollywood Reporter*, “Merriam-Webster Names Slop Word of the Year 2025.” <https://www.hollywoodreporter.com/news/general-news/slop-word-year-2025-merriam-webster-1236450780/>
- Google DeepMind, “SynthID – 100 billion items watermarked; partner expansion to OpenAI, ElevenLabs, Kakao.” <https://deepmind.google/discover/blog/synthid-100-billion-watermarks-partners/>
- *Variety*, “Cate Blanchett co-founds RSL Media – non-profit on AI consent for name, image, likeness.” <https://variety.com/2026/film/news/cate-blanchett-rsl-media-ai-consent-1236748255/>
- *Bloomberg*, “Apple Acquires Key Talent & Patents Behind AI Avatar Company ‘Animato’.” <https://www.bloomberg.com/news/articles/2026-05-19/apple-acquires-animato-ai-avatar-talent-patents>
- Snap Newsroom, “Snapchat Gen Z AI Creativity Research 2026 – 31% of 13–15 year-olds use AI to be creative.” <https://newsroom.snap.com/snapchat-gen-z-ai-creativity-research-2026>
- YouGov AU, “45% of Australians say AI-generated ads would make them trust a brand less.” <https://yougov.com.au/topics/consumer/articles-reports/2026/05/19/45-percent-australians-ai-generated-ads-trust-brand-less>

X. Labour, organisation and the new geography

- Imperva/Cloudflare bot reports (Section III).
- Adobe, “Inaugural Adobe Creators’ Toolkit Report: 86% of Global Creators Use Creative Generative AI.” <https://news.adobe.com/news/2025/10/adobe-max-2025-creators-survey>
- Azumo, “AI in the Workplace Statistics 2026.” <https://azumo.com/artificial-intelligence/ai-insights/ai-in-workplace-statistics>
- *Tech.co*, “Gen Z Most Likely to Sneakily Use AI Without Telling the Boss.” <https://tech.co/news/gen-z-most-likely-use-ai-boss>
- *Exploding Topics*, “The Hidden AI Workforce.” <https://explodingtopics.com/blog/ai-workforce-research>
- *Forbes*, “AI Tools Flood Workplaces as Employees Face a Double Bind.” <https://www.forbes.com/sites/carolinecastrillon/2025/09/09/ai-tools-flood-workplaces-as-employees-face-a-double-bind/>
- IDC Europe, “Shadow AI.” <https://blog-idceurope.com/shadow-ai-how-stealth-productivity-is-strangling-enterprise-ai-adoption-and-creating-a-security-nightmare/>

- *Game Developer*, “Krafton outlines plans to become ‘AI First’ company.” <https://www.gamedeveloper.com/business/subnautica-owner-krafton-outlines-plans-to-transform-into-an-ai-first-company>
- *CNBC*, “Why workers with ADHD, autism, dyslexia should use AI agents.” <https://www.cnn.com/2025/11/08/adhd-autism-dyslexia-jobs-careers-ai-agents-success.html>
- *The Guardian*, “AI is hitting UK harder than other big economies.” <https://www.theguardian.com/technology/2026/jan/26/ai-uk-jobs-us-japan-germany-australia>
- *PocketGamer.biz*, “Shift Up CEO says AI is key to competing with China’s game industry scale.” <https://www.pocketgamer.biz/shift-up-ceo-says-ai-is-key-to-competing-with-chinas-game-industry-scale/>
- *CNBC Africa*, “How AI is changing the landscape of the music industry in Africa.” <https://www.cncafrica.com/2025/how-ai-is-changing-the-landscape-of-the-music-industry-in-africa>
- *BBC Future*, “Lights, camera, algorithm: Why Indian cinema is awash with AI.” <https://www.bbc.co.uk/future/article/20251223-why-indian-cinema-is-awash-with-ai>
- *Broadcast Pro Middle East*, “Tunisian filmmaker wins \$1 million AI Film Award for ‘Lily’.” <https://www.broadcastprome.com/news/tunisian-filmmaker-wins-1-million-ai-film-award-for-lily/>

XI. The literacy turn — institutions and education

- *Sundance Institute*, “Centering the Artist: Why We’re Launching the AI Literacy Initiative.” <https://www.sundance.org/blogs/centering-the-artist-why-were-launching-the-ai-literacy-initiative/>
- *Google blog*, “Building a community-led future for AI in film with Sundance Institute.” <https://blog.google/company-news/outreach-and-initiatives/google-org/sundance-institute-ai-education/>
- *McKinsey & Company*, “What AI could mean for film and TV production.” <https://www.mckinsey.com/industries/technology-media-and-telecommunications/our-insights/what-ai-could-mean-for-film-and-tv-production-and-the-industrys-future>
- *UK Government*, “Free AI training for all.” <https://www.gov.uk/government/news/free-ai-training-for-all-as-government-and-industry-programme-expands-to-provide-10-million-workers-with-key-ai-skills-by-2030>
- *University of Wisconsin-Stout*, “AI Reshaping Industry: AI-use as Baseline Competency in Filmmaking.” <https://www.uwstout.edu/about-us/news-center/ai-reshaping-industry-new-uw-stout-course-sets-ai-use-baseline-competency-filmmaking>

XII. Newsletters and primary corpus

- *Dream Machine | Creative AI* — LinkedIn newsletter, archive of Issues 1–32 (6 October 2025 – 4 June 2026). <https://www.linkedin.com/newsletters/dream-machine-creative-ai-7379776527871381505/>

- DreamLab AI Collective, team page. <https://dreamlab-ai.com/team>
- DreamLab Substack, “Some Predictions on Creative AI for 2026” — *Dream Machine* Issue 13 companion piece. <https://substack.com/home/post/p-183997149>

*Appendix D — The Shadow AI Paradox in
the Creative Industries*

Companion piece to *Chapter 13: Coordination Collapse*.

This deep dive is the long-form analytical companion to the shadow-AI sections of Chapter 13. Where the chapter argues, in the book's voice, that the creative industries are operating with two parallel economies on top of each other — a vocal public economy of AI refusal and a silent private economy of AI adoption — this appendix lays out the underlying data, the sectoral breakdowns, the security and IP implications, the linguistic markers that betray covert AI use, and the labour-market mechanics of agentic displacement that the chapter compresses for narrative purposes.

The headline numbers Chapter 13 quotes — 88–89% staff adoption against 71–80% covert use, the \$670,000 average breach cost, the 1,100-creator music survey showing 87% AI use against 77% “loss of originality” concern, the WGA pre- vs. post-strike screenwriter shift from 34% to 68%, the GDC sentiment-vs-usage divergence — all sit inside the fuller treatment below. Read alongside *Appendix E: Dynamics of Generative AI Adoption*, it forms the empirical spine of the book's argument that the creative industries' stated position on AI and their operational position on AI are, in 2026, sharply divergent — and that this divergence is itself the strategic question every creative organisation now has to face.

The piece below is preserved largely as researched, with citation markers and section headings intact. Some PDF-conversion artefacts (loose footnote numbers, occasional line-break oddities) have not been editorially cleaned; the analytical content is what matters.

* * *

Page 1

The Shadow AI Paradox in the Creative Industries: Covert Adoption, Linguistic Betrayal, and the Displacement Crisis The integration of artificial intelligence within the global creative economy has precipitated one of the most profound technological, economic, and cultural shifts in modern history. However, the prevailing narrative surrounding this transition is characterized by a severe and highly visible dichotomy. Publicly, the creative industries—spanning music, gaming, film, television, animation, broadcasting, and advertising—are engaged in a vocal, highly publicized, and often legally combative resistance against generative artificial intelligence. Trade unions, prominent artists, and media conglomerates routinely denounce the technology, citing massive ethical violations, the non-consensual scraping of copyrighted material, and the existential threat to authentic human expression and labor. Privately, however, the operational reality is starkly different. The sector is currently experiencing an unprecedented and accelerating surge in “shadow AI”—the covert, unsanctioned use of artificial intelligence tools by employees, freelancers, and executives ¹ outside of formal corporate IT and governance structures. This phenomenon reveals a pervasive cognitive dissonance within the modern creative workforce. Professionals who actively and passionately denounce the use of generative models to protect their specific disciplines are simultaneously utilizing the exact same underlying technologies to automate tasks they deem secondary, tedious, or outside their purview—such as coding, copywriting, administrative communication, metadata generation, and data analysis. ³ This localized protectionism, frequently characterized as the “AI for thee, but not for me” paradox, not only

exposes the psychological rationalizations of modern knowledge workers but also accelerates the very displacement they publicly seek to prevent. 5 By feeding proprietary data, unreleased assets, and intellectual capital into public Large Language Models (LLMs) to achieve personal efficiency gains, covert users are inadvertently training the systems that are 7 rendering their own industries and the broader creative ecosystem obsolete. The displacement caused by this technology is highly problematic and systemic, regardless of whether a user justifies their usage as merely “augmentative” for non-core tasks. This comprehensive research report examines the prevalence, mechanics, and long-term consequences of shadow AI in the creative sectors during the critical 2024–2026 transitional period. It analyzes the specific linguistic markers that betray covert AI usage in professional communications, the economic mechanisms driving continuous labor displacement, the specific usage patterns across various creative sub-sectors, and the strategic governance frameworks that organizations must implement to navigate this dual-faced reality. Ultimately, this report outlines the logical conclusion of a paradigm where an industry covertly relies on

Page 2

the very technology it publicly decries. The Epistemology and Scale of Shadow AI To understand the hypocrisy of the creative sector’s AI adoption, it is first necessary to comprehend the sheer scale of shadow AI across enterprise environments. Shadow AI is the natural evolutionary successor to the older concept of shadow IT, but its implications are vastly more severe. While traditional shadow IT involved the unauthorized use of rogue cloud storage or project management apps, shadow AI involves autonomous, self-learning systems that 1 ingest, retain, and iterate upon the data fed into them. Between 2023 and 2024, the adoption of generative AI applications by enterprise employees grew from 74% to 96%, tracking an explosive trajectory that caught corporate compliance 2 departments entirely off guard. By the end of 2025 and moving into 2026, the “Bring Your Own AI” (BYOAI) movement became the dominant operational paradigm. Data from extensive workforce tracking indicates that up to 89% of staff across various organizational departments utilize AI tools, with 71% to 80% of those employees utilizing them without official approval or IT oversight. 7 The scale of this hidden infrastructure has resulted in a phenomenon analysts describe as the “Hidden Cloud Explosion”. 7 The disparity between what corporate leadership believes is happening and what is actually occurring on employee devices is massive. Metric Perceived / Actual / Shadow Strategic Sanctioned Reality Reality Implication Enterprise AI Stalled in pilot 88% to 89% active Official corporate 9 Adoption Rate phases; limited daily users. AI initiatives are 11 official rollout. moving too slowly, forcing employees to utilize public tools to meet productivity demands. Cloud Service Organizations Network data IT departments Visibility estimate an reveals an average suffer from a 90% average of 91 public of 1,220 active visibility gap, cloud services. 7 services. 7 rendering traditional security perimeters

Page 3

obsolete. High-Risk Assumed zero An average of 44 Continuous Applications tolerance for undetected, vulnerability to unsanctioned data high-risk cloud automated data 13 services per

scraping and processing. 7 unauthorized enterprise. cross-border data transfers. Compliance Assumed Only 23% of users Data exposure is Awareness mandatory training are aware of largely driven by 12 compliance. regulatory risks negligence and (e.g., HIPAA, ignorance rather 9 GDPR). than malicious intent. The security, financial, and intellectual property risks associated with this covert adoption are crippling to creative enterprises. In 2025, 20% of organizations experienced severe security incidents directly linked to shadow AI, which increased the average cost of a data breach by \$670,000. 7 The operational mechanisms of these tools—specifically the tendency of users to paste proprietary source code, legal drafts, financial models, and unreleased creative assets into public LLMs like ChatGPT or Claude—resulted in the exposure of personally identifiable information (PII) in 65% of incidents, and the direct theft or exposure of intellectual property in 7 40% of incidents. When an animator pastes proprietary pipeline code into an AI debugging assistant, or a screenwriter uploads an unproduced treatment into an LLM to generate character summaries, they are essentially bypassing corporate endpoint monitoring and handing trade secrets to 7 third-party data processors. The AI's self-learning nature means that these risks compound exponentially; data leaked on a Tuesday can theoretically be used to train a model outputting responses to a competitor by Thursday. 8 Despite these massive systemic risks, the adoption continues unabated, driven by the intense pressure on creative workers to increase their output velocity and streamline workflows in an era of shrinking budgets. The Great Hypocrisy: “AI for Thee, But Not for Me” At the core of the shadow AI phenomenon in the creative sector is a profound psychological, economic, and ideological paradox. Creative professionals frequently exhibit a fierce, localized protectionism regarding their own specific skill sets and domains, while demonstrating a complete and enthusiastic willingness to automate the labor of their peers and collaborators. 3 This dynamic is widely recognized and criticized in developer and creative circles as the “AI for

Page 4

thee, but not for me” paradox. 5 The psychology driving this hypocrisy rests on a highly subjective and hierarchical valuation of human labor. Creatives routinely categorize tasks outside their immediate domain as “mundane,” “tedious,” “administrative,” or “purely technical,” 3 thereby framing the use of AI in these areas as a harmless, victimless efficiency gain. Conversely, they view their own specific domain as an expression of authentic, irreplaceable human experience, soul, and creative genius that cannot, and morally should not, be replicated by an algorithm. 16 For example, a traditional illustrator or concept artist may vehemently campaign against text-to-image models like Midjourney or Stable Diffusion, viewing the scraping of their portfolios as copyright infringement and a desecration of the artistic process. 18 They will join public boycotts and demand protective legislation. Yet, that exact same illustrator may comfortably and secretly use an LLM to write their marketing copy, draft their client contracts, 3 or generate the HTML and Python scripts required to manage their portfolio website. In doing so, they are actively devaluing and bypassing the labor of copywriters, paralegals, and web developers. Similarly, an independent filmmaker might loudly denounce the use of generative video models, arguing that they destroy the craft of cinematography. However, that filmmaker may simultaneously utilize AI audio-cleaning

tools to bypass the need for a professional sound 20 mixer, or use an AI business agent to handle their accounting. This hierarchical thinking is not only hypocritical; it is fundamentally flawed in its understanding of how generative AI models operate. When an independent creator uses an LLM to generate an email campaign or a block of code, they are exploiting the exact same mechanism of 21 non-consensual data scraping that powers image and video generation. The models that generate text and code are trained on the massive, often unlicensed, copyrighted works of authors, journalists, programmers, and corporate communications. 21 The ethical breach is identical, but the personal proximity to the threat alters the user's moral calculus. They are perfectly willing to participate in the enclosure of the knowledge commons, provided the knowledge being enclosed belongs to someone else. The Problematic Nature of "Augmentative" Displacement The justification for this behavior relies heavily on the narrative of "augmentation." Workers convince themselves that they are merely using AI to remove friction from their day, allowing 23 them to focus on the "real" creative work. However, displacement, no matter how the AI is being used, is highly problematic for the broader economic ecosystem. When creatives use AI for "someone else's job," they are contributing to a structural shift 25 toward the "one-person business" model. By utilizing workflow automation platforms and shadow AI tools, individual creatives can scale their output to rival small agencies, generating massive localized profit. 25 While this empowers the individual user, it fundamentally relies on the systemic destruction of entry-level and specialized labor. The senior creative who refuses

Page 5

to hire a junior writer, a junior coder, or a production assistant because an AI can do it faster and cheaper is participating in the exact same aggressive labor displacement they accuse multi-national corporate studio executives of perpetrating. 4 This creates a broken pipeline for future talent. The creative industries have historically relied on a mentorship and apprenticeship model, where junior staff learn the intricacies of a craft by performing the very "mundane" tasks that are now being automated. By eagerly adopting shadow AI for these peripheral tasks, current professionals are burning the bridge they crossed, ensuring that the next generation of creatives has no entry point into the industry. Linguistic Betrayal: Exposing the Covert User The hypocrisy of the shadow AI user is frequently exposed not through sophisticated IT audits or complex network monitoring, but through glaring linguistic betrayals. As creatives increasingly rely on LLMs to generate their professional communications, grant proposals, LinkedIn updates, and social media posts, they inadvertently adopt the semantic and structural 26 artifacts inherent to generative models. Because creative professionals generally possess strong intrinsic communication skills, their sudden shift to AI-generated text is highly conspicuous. LLMs are trained via Reinforcement Learning from Human Feedback (RLHF) to be relentlessly helpful, polite, and enthusiastically compliant. Consequently, unedited AI output possesses a highly distinct, overly polished, yet 26 entirely soulless corporate tone that lacks authentic human nuance, imperfection, or quirk. This phenomenon has led to a flattening of the craft of writing, transforming professional discourse into a sea of sycophantic, buzzword-laden uniformity. 27 The sudden shift in tone from an anti-AI advocate is immediately recognizable to peers. A

professional who typically writes with a standard, slightly imperfect human cadence will suddenly publish content that is grammatically flawless, highly structured, and suffocatingly enthusiastic. 26 The person who spends hours online complaining about AI stealing their art is suddenly “thrilled” to share an update, entirely unaware that their vocabulary is exposing their covert reliance on the technology. The Typology and Psychology of AI Linguistic Markers The presence of shadow AI in creative portfolios, networking posts, and professional emails can be reliably identified through a specific taxonomy of linguistic markers, structural choices, 26 and lazy execution artifacts. Linguistic Marker Specific Identifiers and Psychological and Category Behaviors Technical Origin The Hyperbolic “Delve”, “Tapestry”, RLHF training weights

Page 6

Vocabulary of Enthusiasm “Thrilled”, “Transformative”, heavily favor “Rockstar”, “Dynamic”, high-engagement, “Game-changer” extremely positive corporate rhetoric to avoid user offense. 26 Formulaic Structural Algorithmic reliance on Hook Milestone Predictability statistically common Announcement business-communication templates scraped from Gratitude to Leadership platforms like LinkedIn. 26 Broad Platitude Call to Action 🚀🌟📁🔥 Predictable Emoji , , , , placed at A mechanical simulation of Topography exact paragraph human digital emotion terminations or used in designed to maximize excessive clusters. algorithmic engagement and readability. 26 The Artifact of the “Let me know if you need User negligence, extreme 🚀” “Slip-Up” any modifications! “, haste, and the failure to” Here is the drafted proofread automated response: “, “As an AI...” output before publishing. 26 Lack of Specificity Broad lessons on leadership AI models lack real-world or creativity without context unless explicitly specific project metrics, prompted; they default to dates, or personal generalized wisdom to fill anecdotes. space. 26 The ultimate irony is that creative professionals—who build their entire professional identities on originality, authentic expression, and a unique point of view—are willingly outsourcing their public voices to algorithms that produce the median, most inoffensive blend of corporate 27 speech possible. This linguistic homogenization actively undermines their core argument that human creativity is irreplaceable. If a creative professional cannot be bothered to apply human effort to their own communications and networking, their demands for audiences and executives to deeply value and pay for their human art ring incredibly hollow to the public. The appearance of words like “delve” or “tapestry” in a creative portfolio introduces immediate 28 doubt regarding the authenticity of the visual or auditory work presented alongside it. If the text explaining the creative process was generated by a machine, employers and peers

Page 7

naturally assume the art itself may be heavily reliant on the same covert shortcuts. Sector-Specific Analysis of Covert Adoption The hypocrisy and prevalence of shadow AI manifest differently across the various sub-sectors of the creative economy. While the specific tools change, the underlying pattern of public resistance coupled with deep private dependency remains constant throughout music, gaming, film, animation, and broadcasting. Music Production and Sound Recording The music industry has experienced a highly volatile

relationship with generative AI, transitioning rapidly from litigation-heavy resistance against platforms scraping copyrighted ³¹ material to covert, ubiquitous integration. By 2026, the narrative of “saving human music” clashed directly with internal studio practices. An exhaustive survey of over 1,100 professional music creators—including producers, audio engineers, and songwriters—revealed that a staggering 87% of producers were actively using AI tools in their creative process. ²⁰ The usage within the music sector is highly stratified, reflecting the exact “AI for thee but not for me” hierarchy. Producers generally accept and heavily utilize AI for labor-intensive, highly technical tasks that previously required dedicated specialists. According to industry data, 58% utilize AI for audio restoration and cleanup, 38% for mixing assistants, and 33.9% for automated mastering. ²⁰ However, when it comes to authorship—the core identity of the artist—resistance ²⁰ stiffens considerably. Only 20.9% admit to using composition or lyric-generation tools. There is a profound, existential fear of “musical sameness,” with 77% of producers citing the loss of originality as their primary concern, superseding even the fear of job displacement ²⁰ (42%). Yet, the market realities contradict these artistic ideals. The explosion of fully AI-generated tracks flooding platforms like Spotify—driven by text-to-audio generators like Suno and Udio—forced major labels (Warner, Universal, Sony) to pivot from outright bans to lucrative licensing deals to maintain market share. ³¹ Consequently, the role of the human producer is rapidly transitioning from a direct creator manipulating raw audio to a “creative director” managing intelligent systems. ²⁰ The paradox is palpable: producers utilize AI to instantly execute complex mixing algorithms—tasks that previously sustained an entire working class of audio engineers—while simultaneously denouncing AI systems that generate lyrics, demanding protection for the “soul” of the music. Gaming and Interactive Entertainment The video game industry exhibits the most volatile public reactions to AI, frequently resulting in massive public relations disasters and community outrage. Consumer backlash against perceived “AI slop” has become so severe that it is causing active collateral damage to human ³⁴ artists. Game developers who commission genuine, hand-drawn human artwork have faced aggressive online harassment and accusations of using AI, simply because their art style

Page 8

mirrored the hyper-polished aesthetics popularized by Midjourney and Stable Diffusion. ³⁴ In response to this highly toxic environment, several publishers have implemented draconian anti-AI policies to appease their player bases. Hooded Horse, a prominent indie publisher responsible for massive hits like *Manor Lords*, formally banned the use of AI, inserting “no f**king AI assets” clauses into all of its developer contracts. ³⁵ However, the reality of modern game development—which requires massive, unprecedented volumes of assets, endless lines of code, and continuous debugging—makes strict adherence to these bans nearly impossible. This reality was starkly evidenced when the highly anticipated game *Clair Obscur: Expedition 33* was unceremoniously stripped of a Game of the Year award after internet sleuths discovered that developers had used generative AI to create a minor background newspaper texture. ⁶ The developers claimed it was merely a placeholder that slipped through to the final build, but the incident highlighted the impossibility of policing massive digital environments for AI artifacts. The gaming sector perfectly encapsulates the shadow AI dilemma. Developers

widely and enthusiastically use AI coding assistants like GitHub Copilot to write scripts, autocomplete logic, and debug errors, viewing it as an absolute necessity for productivity and survival in a crunch-heavy industry. 8 Yet, the integration of AI for visual assets or narrative design is treated as a moral failing. This arbitrary distinction between automating engineering (viewed as acceptable efficiency) and automating art (viewed as unacceptable theft) is fundamentally hypocritical, highlighting the deep cognitive dissonance at the heart of interactive 4 entertainment. Film, Television, and Animation Hollywood currently operates under a pervasive “don’t ask, don’t tell” culture regarding artificial intelligence. 22 In the wake of historic industry strikes, the Writers Guild of America (WGA) and 39 SAG-AFTRA established rigid, legally binding guidelines surrounding AI. These guidelines mandate explicit, 48-hour advanced consent and mandatory compensation for the creation of “Employment-Based Digital Replicas” and “Synthetic Performers”. 41 Furthermore, writers are permitted to use generative AI as a tool, provided it is disclosed, but studios cannot force 42 writers to use it, nor can AI be credited with authorship or used to reduce a writer’s residuals. Despite these hard-fought contractual boundaries, covert usage is rampant across the supply chain. Industry executives and insiders note that studios are frequently lying about how much AI they are utilizing in post-production, storyboarding, and visual effects to avoid union 22 grievances and consumer backlash. Simultaneously, the creatives are equally deceptive about their reliance on LLMs. As one industry veteran noted, it is nearly impossible to find a screenwriter staring at a blank page who is not simultaneously conversing with Claude or ChatGPT to break story structures or generate dialogue options. 38 The quiet, unapologetic release of films utilizing AI, such as the Oscar-nominated *The Brutalist*, which utilized AI to seamlessly enhance the vocal accents of its lead actors, suggests that the technology is 38 already deeply embedded in prestige cinema.

Page 9

The animation sector faces an even more dire existential crisis. Animation has historically been one of the most labor-intensive creative fields. A 2025 Luminate Intelligence report highlighted that animation executives view generative AI as a revolutionary mechanism to slash production times, which have historically been stubbornly long, and reduce ballooning budgets. 43 Conversely, the Animation Guild views it as a severe threat, estimating that 21% of animation tasks are vulnerable to immediate AI exposure and automation, putting nearly 40,000 jobs in 45 California alone at risk. This technological threat is compounding existing issues of outsourcing. Studios are increasingly moving animation production away from highly regulated, unionized hubs like Los Angeles to regions offering heavier tax subsidies, utilizing AI to bridge the logistical and 46 creative gaps across distributed global pipelines. Animators find themselves in an untenable position: they are forced to compete with hyper-efficient automated pipelines, leading many to secretly adopt generative AI tools just to meet the newly compressed production quotas, even as their unions fight to ban the technology entirely. 43 Broadcasting, Advertising, and Agency Production In the broader broadcasting and advertising sectors, the integration of AI has moved from a novelty to a core operational requirement, though not always successfully. While 2024 was marked by AI moving from a “side tab” to a “core app,” 2025 and 2026 ushered in the “Agentic Era,” where AI systems

shifted from simply answering queries to semi-autonomously 47 performing complex actions. The market size for agentic AI within media and entertainment is projected to grow by 35.9% by 2030, handling localization, metadata generation, and workflow optimization. 48 However, official corporate rollouts of these tools have been remarkably clumsy. Up to 70% of official corporate AI initiatives in broadcasting and media fall short of expectations, plagued by rigorous compliance checks, clunky enterprise interfaces, and a lack of specific training. 48 This massive failure rate drives employees directly into the arms of shadow AI. Because the officially sanctioned tools are difficult to use, 50% of employees use unauthorized AI without 49 permission, and 64% pass off AI-generated work as their own human creation. The advertising world has seen massive controversies regarding AI usage. When agencies attempt to use AI overtly, as seen in Coca-Cola's heavily criticized AI-generated Christmas 50 advertisement, they risk destroying decades of brand equity. Audiences rejected the ad not necessarily because it was AI, but because it felt lazy and contradicted the brand's tagline of "Real Magic". 50 This public backlash forces agencies back into the shadows; instead of using AI for final broadcast output, they covertly use it for pitch decks, storyboarding, demographic analysis, and copywriting, hiding the machine's involvement from the client while billing for 51 human hours. The Mechanics of Market Devaluation and

Page 10

Displacement The deeply held belief among creative professionals that shadow AI functions merely as an "augmentative" tool that spares the core creative process is an economic fallacy. No matter how these tools are utilized—whether an illustrator uses an LLM to write contracts, an animator uses it to write pipeline code, or a composer uses it to clean up audio stems—the aggregate, macro-economic effect is the systematic displacement of human labor and the rapid devaluation of the creative economy. 53 The macro-economic data underscores the severity of this shift. According to the United Nations Educational, Scientific and Cultural Organization (UNESCO), the proliferation of generative AI is projected to cause significant and irreversible income losses by 2028. 56 The report warns that music creators face a 24% drop in revenues, while audiovisual professionals 56 are projected to lose 21% of their income as AI-generated content floods global markets. A broader economic analysis by Goldman Sachs estimates that up to 300 million jobs globally are exposed to AI automation over the next decade, with a specific focus on knowledge workers and creative sectors. 57 While organizations like Anthropic attempt to measure "observed exposure" against "theoretical capability"—noting that AI is far from reaching its theoretical maximum and finding limited immediate unemployment spikes—the reality on the ground in the creative sector is one of task displacement rather than immediate, total job erasure. 54 The Stanford Study: Flooding the Market This displacement is fundamentally driven by a massive, sudden shift in market supply and consumer demand. A critical and revealing study conducted by Stanford Graduate School of Business analyzed the market dynamics when AI-generated art was introduced to a platform 59 alongside traditional human-created art. The findings utterly dismantle the optimistic, romantic narrative that human art will retain a premium value due to its inherent "soul" or authenticity. Once generative AI entered the marketplace, the total supply of images skyrocketed 59 exponentially, and consequently, the volume of human-generated images fell

dramatically. More alarmingly for traditional creators, consumers actively demonstrated a preference for the influx of AI-generated images, choosing them over human-generated ones. 59 The technology flooded the market with high quality and infinite variety at a near-zero marginal cost. This effectively crowded out human creators who simply could not compete on volume, speed, or 55 price. The market dictates that when “good enough” is available instantly and for free, the demand for “excellent but slow and expensive” human labor evaporates. The Enclosure of the Knowledge Commons Every instance of shadow AI usage contributes directly to this dynamic. Generative models function as massive, continuous feedback loops. When an employee covertly inputs a

Page 11

proprietary script, a unique visual asset, or an innovative piece of code into a public LLM to save time, they are voluntarily surrendering their intellectual property and the collective knowledge of their industry to the algorithm. 7 This process is known in critical theory as the “enclosure of the knowledge commons”. 55 Major tech corporations rely on this continuous, free ingestion of human labor to refine their models, subsequently monopolizing that accumulated intelligence and selling it back to the market in 55 the form of autonomous agents and enterprise subscriptions. This mechanism ensures that the “augmentation” phase of AI is strictly temporary. The tools are currently designed and marketed to assist a human expert to complete their job faster. However, by continually tracking the human expert’s prompts, inputs, corrections, and decision-making processes, the AI learns exactly how to eventually perform the task 4 autonomously without human intervention. The creative professional who uses shadow AI is not just cheating their employer’s IT policy to leave work an hour early; they are actively, literally training their permanent replacement. The ultimate result of this dynamic is an impending hollowing out of the mid-level creative workforce. While elite creative directors, showrunners, and “AI power users” may thrive by orchestrating massive, highly automated production pipelines, the entry-level and mid-tier roles—junior animators, staff writers, copywriters, sound editors, and conceptual artists—are 20 facing total erasure. Strategic Restitution: Formalizing AI Governance The pervasive, deeply entrenched nature of shadow AI within the creative sectors dictates that traditional IT prohibition strategies are entirely futile. Banning public LLMs, implementing strict firewalls, or punishing employees simply drives the behavior further underground, forcing staff to use personal devices, cellular networks, or unmonitored VPNs to access the tools they rely 9 on. Furthermore, an outright ban deprives the organization of the genuine productivity gains that AI can offer, placing the company at a severe competitive disadvantage in a rapidly evolving, cost-conscious market. 7 To survive this transition, creative organizations must move away from a posture of denial and transition toward a model of “structured enablement” and formal AI governance. 7 This requires explicitly acknowledging the reality of widespread employee usage and implementing both technical and cultural guardrails to protect intellectual property without stifling the creative 62 innovation the tools can provide. Technical Mitigation and Identity Discovery The foundational premise of managing AI risk relies on a simple axiom: you cannot secure what you cannot see. 10 The primary technical

objective for IT departments is eliminating the massive visibility gap caused by the Hidden Cloud Explosion.

Page 12

Organizations must implement dynamic Software-as-a-Service (SaaS) security platforms and Identity Discovery protocols designed specifically to scan for non-human identities and agentic workflows. 10 This highly technical approach involves: • Monitoring and auditing API tokens that are calling external AI services outside of approved corporate gateways. 63 • Utilizing advanced Data Loss Prevention (DLP) tools equipped with behavioral analytics. Rather than relying on static blocklists of known AI URLs (which are easily bypassed), these systems must be context-aware, detecting when sensitive information patterns (like proprietary code structures, PII, or unreleased script formats) are being pasted into 13 browser-based LLMs. • Scanning for unknown service accounts interacting with AI APIs and autonomous processes accessing production systems. 63 Formalizing the AI Governance Framework Beyond technical detection, the establishment of a robust AI Governance Framework is critical for scaling AI securely, legally, and ethically within a creative agency or studio. 62 A mature framework transitions a company from ad hoc, informal usage driven by individual employees 61 to a regimented, trackable system endorsed by leadership. As demonstrated by industry leaders and labor unions—such as the Trades Union Congress (TUC) AI Manifesto and corporate policies from Blue Zoo Animation and Mapfre—a successful framework requires intensive cross-departmental collaboration between IT, Legal, Human 66 Resources, and the creative workforce itself. Governance Maturity Operational Required Security & Phase Characteristics Compliance Actions Phase 1: Discovery & Ubiquitous shadow AI. No Implement network Informal official policy. Staff unaware telemetry to audit actual of compliance risks. 7 SaaS usage. Identify Rampant BYOAI. high-risk data flows. Conduct upfront reviews of vendor Terms of Service for embedded AI. 69 Phase 2: Ad Hoc & Walled-garden access (e.g., Implement API gateways. Transitional Enterprise Copilot) Block transmission of PII to 69 introduced. Basic public models. Roll out acceptable use policies mandatory AI literacy training to bridge the 50%

Page 13

distributed to staff. awareness gap. 12 Phase 3: Formal & AI usage is transparent, Continuous monitoring for Integrated licensed, and actively model drift and algorithmic monitored. Ethical and IP 62 bias. Formalized vendor guardrails are automated data agreements directly into the creative guaranteeing zero data 52 workflow. retention. Compliance with regional legislation (e.g., EU AI Act). 60 A comprehensive framework must mandate total transparency regarding training data to ensure that AI output is clearly labeled, preventing copyright contamination of the studio's 67 broader intellectual property portfolio. Furthermore, organizations must establish an “opt-in” system for the internal use of employee-generated assets to train proprietary, internal models. 67 Most importantly, organizations must supply enterprise-grade, heavily secured AI solutions that guarantee zero data retention by the model providers. When employees are provided with sanctioned, secure tools that perform as well as or better than the public shadows, the incentive to bypass IT

protocols and risk data exposure entirely evaporates. 9 The Logical Conclusion The trajectory of shadow AI within the creative industries points toward an inescapable, highly disruptive logical conclusion. The current state of cognitive dissonance—where creative professionals publicly demonize artificial intelligence as the death of art while privately relying on it to augment their daily output—is merely a brief, unstable transitional phase. As the technology’s capabilities expand exponentially beyond the automation of “mundane” tasks and begin to consistently perform judgment-intensive, open-ended creative work at human-level proficiency, the fragile, hypocritical truce of “augmentation” will inevitably collapse. 23 The displacement of human labor is not an accidental or unfortunate byproduct of generative AI; it is its foundational design and economic purpose. By continuously feeding proprietary knowledge, stylistic nuances, and problem-solving logic into unvetted public models to save time on a Tuesday, the covert users of shadow AI are actively subsidizing and 7 accelerating the total devaluation of their own industries by Friday. The market has already signaled definitively that consumer demand will easily adapt to, and in many cases prefer, the infinite, frictionless supply of synthetic media over the slower, more 59 expensive output of human creators. Consequently, the traditional concept of the hands-on “creator” is fracturing. The future of the industry belongs almost exclusively to the “creative director”—the individual who excels not in the manual execution of a specific craft (drawing, coding, mixing), but in the precise curation, prompting, and orchestration of vast,

Page 14

interconnected autonomous digital agents. 20 Those who continue to publicly denounce AI while secretly leveraging it for personal efficiency are engaged in a self-defeating hypocrisy that offers no long-term protection. The survival of the creative professional will not depend on successfully protecting a specific, granular skill set from automation through public boycotts. Rather, it will depend on transitioning to highly formalized, governed AI integration that explicitly protects human intellectual property at the enterprise level, ensuring that artists are compensated for the data they generate. 62 The alternative is the total, irreversible enclosure of the creative commons, where human ingenuity serves merely as the unpaid, unacknowledged training data for the automated, synthetic pipelines of tomorrow. The creative industries must drag AI out of the shadows and govern it in the light, or risk being entirely consumed by it. Works cited 1. Shadow AI and the Future of Work: What Knowledge Workers Need to Know in 2026 : r/it - Reddit, accessed on April 21, 2026, https://www.reddit.com/r/it/comments/1qn2fgw/shadow_ai_and_the_future_of_work_what_knowledge/ 2. What Is Shadow AI? - IBM, accessed on April 21, 2026, <https://www.ibm.com/think/topics/shadow-ai> 3. AI Isn’t Taking Over All Creative Jobs (So Embrace It) | Now Hear This, accessed on April 21, 2026, <https://www.now-hear-this.net/content/ai-isnt-taking-over-all-creative-jobs-so-embrace-it> 4. I stopped defending my creative team and let leaders use ai... it failed lol - Reddit, accessed on April 21, 2026, https://www.reddit.com/r/antiwork/comments/1kfn9jl/i_stopped_defending_my_creative_team_and_let/ 5. Anthropic: “Applicants should not use AI assistants” | Hacker News, accessed on April 21, 2026, <https://news.ycombinator.com/item?id=42915905> 6. After Being Stripped of Its GOTY Win, Clair Obscur: Expedition 33 Director Admits Using AI “Felt

Wrong” - Reddit, accessed on April 21, 2026, https://www.reddit.com/r/DreamStationcc/comments/1pur5cs/after_being_stripped_of_its_goty_win_clair_obscure/ 7. Shadow IT Statistics You Need to Know Now (2026) - ElectroIQ, accessed on April 21, 2026, <https://electroi.com/stats/shadow-it-statistics/> 8. AI Gone Wild: Why Shadow AI Is Your IT Team’s Worst Nightmare, accessed on April 21, 2026, <https://cloudsecurityalliance.org/blog/2025/03/04/ai-gone-wild-why-shadow-ai-is-your-it-team-s-worst-nightmare> 9. Shadow AI Statistics 2026: Adoption, Risk, And Breach Data, accessed on April 21, 2026, <https://authentech.ai/blog/shadow-ai/shadow-ai-statistics-2026/> 10. Popular Doesn’t Mean Secure - The 2025 State of Shadow AI Report Findings -

Page 15

Reco AI, accessed on April 21, 2026, <https://www.reco.ai/blog/popular-doesnt-mean-secure-the-2025-state-of-shadow-ai-report-findings> 11. The GenAI Divide: State of AI in Business 2025 - MLQ.ai, accessed on April 21, 2026, https://mlq.ai/media/quarterly_decks/v0.1_State_of_AI_in_Business_2025_Report.pdf 12. Shadow AI Usage Statistics 2026: Latest Insights - SQ Magazine, accessed on April 21, 2026, <https://sqmagazine.co.uk/shadow-ai-usage-statistics/> 13. How to Detect & Stop Shadow AI Tools in the Company : r/sysadmin - Reddit, accessed on April 21, 2026, https://www.reddit.com/r/sysadmin/comments/1pn1y3v/how_to_detect_stop_shadow_ai_tools_in_the_company/ 14. Why “AI” Is a Con: Our Book Launch! (with Vauhini Vara), 2025.05.08 - Buzzsprout, accessed on April 21, 2026, <https://www.buzzsprout.com/2126417/episodes/17459349-why-ai-is-a-con-our-book-launch-with-vauhini-vara-2025-05-08> 15. I think there’s two different things going on here: “DeepSeek trained on our out... - Hacker News, accessed on April 21, 2026, <https://news.ycombinator.com/item?id=42867899> 16. This is how AI is impacting – and shaping – the creative industries, according to experts at Davos - The World Economic Forum, accessed on April 21, 2026, <https://www.weforum.org/stories/2024/02/ai-creative-industries-davos/> 17. Extracted The Deep Well of Time - The Transformative Power of Storytelling in The Classroom | PDF - Scribd, accessed on April 21, 2026, <https://www.scribd.com/document/691281535/Extracted-the-Deep-Well-of-Time-the-Transformative-Power-of-Storytelling-in-the-Classroom> 18. The ‘AI Hypocrisy’ Paradox: Why We Hate AI in Our Jobs but Love It in Everyone Else’s, accessed on April 21, 2026, <https://maketecheasier.com/ai-hypocrisy-paradox/> 19. The Ethical Implications of AI in Creative Industries: A Focus on AI-Generated Art, accessed on April 21, 2026, https://www.researchgate.net/publication/393511482_The_Ethical_Implications_of_AI_in_Creative_Industries_A_Focus_on_AI-Generated_Art 20. The Future of Music Production Is Human: 1,100+ Producers Reveal ..., accessed on April 21, 2026, <https://www.sonarworks.com/blog/research/future-music-production-human-producer-survey-2026> 21. Generative Artificial Intelligence and the Creative Economy Staff Report: Perspectives and Takeaways - Federal Trade Commission, accessed on April 21, 2026, https://www.ftc.gov/system/files/ftc_gov/pdf/12-15-2023AICEStaffReport.pdf 22. The Breaking Point: Why Hollywood’s Long-Simmering AI Cold War Just Went Nuclear, accessed

on April 21, 2026, <https://nofilmschool.com/ai-war-hollywood> 23. AI Will Reshape More Jobs Than It Replaces | BCG, accessed on April 21, 2026,

Page 16

<https://www.bcg.com/publications/2026/ai-will-reshape-more-jobs-than-it-replaces> 24. Creative pros are leveraging Generative AI to do more and better work | Adobe Blog, accessed on April 21, 2026, <https://blog.adobe.com/en/publish/2024/02/02/creative-pros-generative-ai-use> 25. The Silent Displacement: How AI Automation is Creating a New Class of One-Person Businesses While Disrupting Traditional Employment | by TradeApollo | Medium, accessed on April 21, 2026, <https://medium.com/@tradeapollo/the-silent-displacement-how-ai-automation-is-creating-a-new-class-of-one-person-businesses-while-e1a67f716b96> 26. Let Me Know If You Need Any Modifications : The LinkedIn AI Crisis ..., accessed on April 21, 2026, <https://medium.com/@avanib28264/let-me-know-if-you-need-any-modifications-the-linkedin-ai-crisis-163748c04c1b> 27. I see Claude's writing everywhere and it's starting to feel like an AI condom, I hate it - Reddit, accessed on April 21, 2026, https://www.reddit.com/r/ClaudeAI/comments/1rjeqg3/i_see_claudes_writing_everywhere_and_its_starting/ 28. BEYOND THE BLACK BOX: Shaping a Responsible AI Landscape, accessed on April 21, 2026, https://www.kpmginfo.com/web/010797/EAIGG_AnnualReport_117.pdf 29. Can That Emoji Reveal a Remote Workers' Emotional State? - Social Science Space, accessed on April 21, 2026, <https://www.socialsciencespace.com/2022/03/can-that-emoji-reveal-a-remote-workers-emotional-state/> 30. Transforming Retail & Consumer Brands: Generative AI Cases and Potential - Tenz.ai, accessed on April 21, 2026, <https://www.tenz.ai/wp-content/uploads/2023/09/Tenzai-Transforming-with-Generative-AI-Retail-Consumer-Brands-1.pdf> 31. The AI Music Bubble Is Bursting in 2026: Hype, Deals & What Survives for Real Creators, accessed on April 21, 2026, https://www.youtube.com/watch?v=mjbp_nODAWA 32. AI Musicians vs. Machines: AI Disruption in the Music Industry ..., accessed on April 21, 2026, <https://www.ocstrategy.com/en/article/ai-musicians-vs-machines-ai-disruption-in-the-music-industry/> 33. New Study Reveals 87% of Producers Already Use AI Tools in Their Creative Process, accessed on April 21, 2026, <https://aristake.com/ai-tools-musicians-study/> 34. A real issue: video game developers are being accused of using AI – even when they aren't, accessed on April 21, 2026, <https://www.theguardian.com/games/2025/jun/26/video-game-developers-using-ai-even-when-they-arent-stamina-zero> 35. Many Game Devs Reject AI in Gaming | TechPowerUp, accessed on April 21, 2026,

Page 17

<https://www.techpowerup.com/345122/many-game-devs-reject-ai-in-gaming> 36. Game Publisher Bans Working With Devs That Use Any AI, Rather Than Banning Bad Uses Of AI - Techdirt, accessed on April 21, 2026, <https://www.techdirt.com/2026/01/16/game-publisher-bans-working-with-devs-that-use-any-ai-rather-than-banning-bad-uses-of-ai/> 37. Valve dev counters calls to scrap Steam AI disclosures, says it's a "technology relying on cultural laundering, IP infringement, and slopification" : r/pcgaming - Reddit, accessed on

https://www.reddit.com/r/pcgaming/comments/1p8vi3f/valve_dev_counters_calls_to_scrap_steam_ai/ 38. Hollywood Is Lying to Everyone About How Much AI They're Using, Says Consummate Hollywood Insider - Futurism, accessed on April 21, 2026, <https://futurism.com/artificial-intelligence/hollywood-lying-about-ai-usage> 39. Artificial Intelligence Resources | SAG-AFTRA, accessed on April 21, 2026, <https://www.sagaftra.org/contracts-industry-resources/contracts/2023-tvtheatric-al-contracts/artificial-intelligence-resources> 40. Artificial Intelligence - WGA, accessed on April 21, 2026, <https://www.wga.org/contracts/know-your-rights/artificial-intelligence> 41. Regulating Artificial Intelligence - sag-aftra, accessed on April 21, 2026, https://www.sagaftra.org/sites/default/files/sa_documents/AI%20TVTH.pdf 42. How the 2023 SAG-AFTRA and WGA Contracts Address Generative AI | Perkins Coie, accessed on April 21, 2026, <https://perkinscoie.com/insights/blog/generative-ai-movies-and-tv-how-2023-sa-g-aftra-and-wga-contracts-address-generative> 43. Animation Jobs & AI: How the Industry's Future Is Changing - Cartoon Brew, accessed on April 21, 2026, <https://www.cartoonbrew.com/tech/animation-jobs-in-the-age-of-ai-revelations-from-luminate-intelligences-2025-special-report-253718.html> 44. New Report from Luminate Predicts Animation Jobs Most Likely to Be Impacted by AI, accessed on April 21, 2026, <https://www.animationmagazine.net/2025/09/new-report-from-luminate-predicts-animation-jobs-most-likely-to-be-impacted-by-ai/> 45. Trump's new order against AI regulation hits California especially hard - CalMatters, accessed on April 21, 2026, <https://calmatters.org/economy/technology/2025/12/california-ai-regulation-targeted-in-trump-order/> 46. Why outsourcing and AI means the US animation sector is facing hefty challenges in 2025, accessed on April 21, 2026, <https://www.screendaily.com/features/why-outsourcing-and-ai-means-the-us-animation-sector-is-facing-hefty-challenges-in-2025/5199989.article> 47. AI Usage Statistics 2025: The Complete Market & Adoption Report - Elfsight, accessed on April 21, 2026, <https://elfsight.com/blog/ai-usage-statistics-2025/> 48. Media industry outlook: Five AI trends shaping the future of media and broadcasting, accessed on April 21, 2026, <https://www.cgi.com/en/article/communications-media/five-ai-media-and-broad>

Page 18

casting-trends 49. Review/Preview Trends Report 24-25 - NPO, accessed on April 21, 2026, https://media.prod.cc.bijnpo.nl/Wayfinder_2025_Trend_Report_4e16709fd9.pdf 50. The biggest AI art controversies of 2024 - Creative Bloq, accessed on April 21, 2026, <https://www.creativebloq.com/ai/ai-art/the-biggest-ai-art-controversies-of-2024> 51. "That's AI": Dealing with negative public perceptions of AI - shorthandstories.com, accessed on April 21, 2026, <https://the-ipa.shorthandstories.com/ipai-forum/ai-trends/thats-ai-dealing-with-negative-public-perceptions-of-ai/index.html> 52. Helping creative agencies thrive with AI | IDHL News & Insight, accessed on April 21, 2026, <https://www.idhlagency.com/insights/the-next-12-months-will-show-which-creative-agencies-can-thrive-with-ai> 53. AI and work in the creative industries: digital continuity or discontinuity? - Taylor & Francis, accessed on April 21, 2026, <https://www.tandfonline.com/doi/full/10.1080/17510694.2024.2421135> 54. Labor

market impacts of AI: A new measure and early evidence - Anthropic, accessed on April 21, 2026, <https://www.anthropic.com/research/labor-market-impacts> 55. When We Get Komooted - BIKEPACKING.com, accessed on April 21, 2026, <https://bikepacking.com/plog/when-we-get-komooted/> 56. Artists face steep income decline due to AI, UNESCO finds | UN News, accessed on April 21, 2026, <https://news.un.org/en/story/2026/02/1166989> 57. How Will AI Affect the US Labor Market? | Goldman Sachs, accessed on April 21, 2026, <https://www.goldmansachs.com/insights/articles/how-will-ai-affect-the-us-labor-market> 58. Tracking the Impact of AI on the Labor Market | The Budget Lab at Yale, accessed on April 21, 2026, <https://budgetlab.yale.edu/research/tracking-impact-ai-labor-market> 59. When AI-Generated Art Enters the Market, Consumers Win — and Artists Lose, accessed on April 21, 2026, <https://www.gsb.stanford.edu/insights/when-ai-generated-art-enters-market-consumers-win-artists-lose> 60. Building Your Company's AI Governance Framework to Reduce Risk - Bloomberg Law, accessed on April 21, 2026, <https://pro.bloomberglaw.com/insights/artificial-intelligence/building-your-companys-ai-governance-framework-to-reduce-risk/> 61. AI Governance: Frameworks, Principles, and Practices - Domo, accessed on April 21, 2026, <https://www.domo.com/glossary/ai-governance> 62. AI Governance Frameworks Explained: How to Manage AI Responsibly at Scale - Rubrik, accessed on April 21, 2026, <https://www.rubrik.com/insights/ai-governance-frameworks-explained> 63. What Is Shadow AI and How Do You Detect It? - LoginRadius, accessed on April

Page 19

21, 2026, <https://www.loginradius.com/blog/engineering/what-is-shadow-ai-and-how-do-you-detect-it> 64. The Rise of Shadow AI: Auditing Unauthorized AI Tools in the Enterprise - ISACA, accessed on April 21, 2026, <https://www.isaca.org/resources/news-and-trends/industry-news/2025/the-rise-of-shadow-ai-auditing-unauthorized-ai-tools-in-the-enterprise> 65. A Practical AI Governance Framework for Enterprises - Databricks, accessed on April 21, 2026, <https://www.databricks.com/blog/practical-ai-governance-framework-enterprises> 66. We won't delegate our creativity - our AI manifesto - Blue Zoo Animation Studio, accessed on April 21, 2026, <https://www.blue-zoo.co.uk/blog/our-ai-manifesto/> 67. Equity calls for stronger AI protections for creative workers, accessed on April 21, 2026, <https://www.equity.org.uk/news/2025/equity-calls-for-stronger-ai-protections-for-creative-workers> 68. Discover Mapfre's vision on artificial intelligence and our new Manifesto, accessed on April 21, 2026, <https://www.mapfre.com/en/insights/innovation/mapfres-artificial-intelligence-manifesto/> 69. How To Detect and Prevent Shadow AI - FairNow, accessed on April 21, 2026, <https://fairnow.ai/detect-and-prevent-shadow-ai/> 70. Guidelines | AI | University of Florida, accessed on April 21, 2026, <https://ai.ufl.edu/guidelines/> 71. Human-centered AI through employee participation - PMC, accessed on April 21, 2026, <https://pmc.ncbi.nlm.nih.gov/articles/PMC10941644/> 72. IBM Study: Shadow AI Use Surges as Canadian Workers Outpace Employers in AI Adoption, accessed on April 21, 2026, <https://canada.newsroom.ibm.com/2025-09-03-IBM-R-Study-Shadow-AI-Use-Surges-as-Canadian-Workers-Outpace-Employers-in-AI-Adoption>

All embedded URLs (in document order)

- https://www.reddit.com/r/it/comments/1qn2f9w/shadow_ai_and_the_future_of_work_
- <https://www.ibm.com/think/topics/shadow-ai>
- <https://www.now-hear-this.net/content/ai-isnt-taking-over-all-creative-jobs-so-embrace-it>
- https://www.reddit.com/r/antiwork/comments/1kfn9jl/i_stopped_defending_my_creati
- <https://news.ycombinator.com/item?id=42915905>
- https://www.reddit.com/r/DreamStationcc/comments/1pur5cs/after_being_stripped_of_
- <https://electroiq.com/stats/shadow-it-statistics/>
- <https://cloudsecurityalliance.org/blog/2025/03/04/ai-gone-wild-why-shadow-ai-is-your-it-team-s-worst-nightmare>
- <https://authentech.ai/blog/shadow-ai/shadow-ai-statistics-2026/>
- <https://www.reco.ai/blog/popular-doesnt-mean-secure-the-2025-state-of-shadow-ai-report-findings>
- https://mlq.ai/media/quarterly_decks/v0.1_State_of_AI_in_Business_2025_Report.pdf
- <https://sqmagazine.co.uk/shadow-ai-usage-statistics/>
- https://www.reddit.com/r/sysadmin/comments/1pn1y3v/how_to_detect_stop_shadow_
- <https://www.buzzsprout.com/2126417/episodes/17459349-why-ai-is-a-con-our-book-launch-with-vauhini-vara-2025-05-08>
- <https://news.ycombinator.com/item?id=42867899>
- <https://www.weforum.org/stories/2024/02/ai-creative-industries-davos/>
- <https://www.scribd.com/document/691281535/Extracted-the-Deep-Well-of-Time-the-Transformative-Power-of-Storytelling-in-the-Classroom>
- <https://maketecheasier.com/ai-hypocrisy-paradox/>
- https://www.researchgate.net/publication/393511482_The_Ethical_Implications_of_AI_Generated_Art
- <https://www.sonarworks.com/blog/research/future-music-production-human-producer-survey-2026>
- https://www.ftc.gov/system/files/ftc_gov/pdf/12-15-2023AICESTaffReport.pdf
- <https://nofilmschool.com/ai-war-hollywood>
- <https://www.bcg.com/publications/2026/ai-will-reshape-more-jobs-than-it-replaces>
- <https://blog.adobe.com/en/publish/2024/02/02/creative-pros-generative-ai-usage>
- <https://medium.com/@tradeapollo/the-silent-displacement-how-ai-automation-is-creating-a-new-class-of-one-person-businesses-while-e1a67f716b96>
- <https://medium.com/@avanib28264/let-me-know-if-you-need-any-modifications-the-linked-in-ai-crisis-163748c04c1b>
- https://www.reddit.com/r/ClaudeAI/comments/1rjeqg3/i_see_claude_writing_everywh
- https://www.kpmginfo.com/web/010797/EAIGG_AnnualReport_1_17.pdf

- <https://www.socialsciencespace.com/2022/03/can-that-emoji-reveal-a-remote-workers-emotional-state/>
- <https://www.tenz.ai/wp-content/uploads/2023/09/Tenzai-Transforming-with-Generative-AI-Retail-Consumer-Brands-1.pdf>
- https://www.youtube.com/watch?v=mjbp_nODAWA
- <https://www.ocstrategy.com/en/article/ai-musicians-vs-machines-ai-disruption-in-the-music-industry/>
- <https://aristake.com/ai-tools-musicians-study/>
- <https://www.theguardian.com/games/2025/jun/26/video-game-developers-using-ai-even-when-they-arent-stamina-zero>
- <https://www.techpowerup.com/345122/many-game-devs-reject-ai-in-gaming>
- <https://www.techdirt.com/2026/01/16/game-publisher-bans-working-with-devs-that-use-any-ai-rather-than-banning-bad-uses-of-ai/>
- https://www.reddit.com/r/pcgaming/comments/1p8vi3f/valve_dev_counters_calls_to_s
- <https://futurism.com/artificial-intelligence/hollywood-lying-about-ai-usage>
- <https://www.sagafta.org/contracts-industry-resources/contracts/2023-tvtheatrical-contracts/artificial-intelligence-resources>
- <https://www.wga.org/contracts/know-your-rights/artificial-intelligence>
- https://www.sagafta.org/sites/default/files/sa_documents/AI%20TVTH.pdf
- <https://perkinscoie.com/insights/blog/generative-ai-movies-and-tv-how-2023-sag-aftra-and-wga-contracts-address-generative>
- <https://www.cartoonbrew.com/tech/animation-jobs-in-the-age-of-ai-revelations-from-luminate-intelligences-2025-special-report-253718.html>
- <https://www.animationmagazine.net/2025/09/new-report-from-luminate-predicts-animation-jobs-most-likely-to-be-impacted-by-ai/>
- <https://calmatters.org/economy/technology/2025/12/california-ai-regulation-targeted-in-trump-order/>
- <https://www.screendaily.com/features/why-outsourcing-and-ai-means-the-us-animation-sector-is-facing-hefty-challenges-in-2025/5199989.article>
- <https://elfsight.com/blog/ai-usage-statistics-2025/>
- <https://www.cgi.com/en/article/communications-media/five-ai-media-and-broadcasting-trends>
- https://media.prod.cc.bijnpo.nl/Wayfinder_2025_Trend_Report_4e16709fd9.pdf
- <https://www.creativebloq.com/ai/ai-art/the-biggest-ai-art-controversies-of-2024>
- <https://the-ipa.shorthandstories.com/ipai-forum/ai-trends/thats-ai-dealing-with-negative-public-perceptions-of-ai/index.html>
- <https://www.idhlagency.com/insights/the-next-12-months-will-show-which-creative-agencies-can-thrive-with-ai>
- <https://www.tandfonline.com/doi/full/10.1080/17510694.2024.2421135>
- <https://www.anthropic.com/research/labor-market-impacts>
- <https://bikepacking.com/plog/when-we-get-komooted/>

- <https://news.un.org/en/story/2026/02/1166989>
- <https://www.goldmansachs.com/insights/articles/how-will-ai-affect-the-us-labor-market>
- <https://budgetlab.yale.edu/research/tracking-impact-ai-labor-market>
- <https://www.gsb.stanford.edu/insights/when-ai-generated-art-enters-market-consumers-win-artists-lose>
- <https://pro.bloomberglaw.com/insights/artificial-intelligence/building-your-companys-ai-governance-framework-to-reduce-risk/>
- <https://www.domo.com/glossary/ai-governance>
- <https://www.rubrik.com/insights/ai-governance-frameworks-explained>
- <https://www.loginradius.com/blog/engineering/what-is-shadow-ai-and-how-do-you-detect-it>
- <https://www.isaca.org/resources/news-and-trends/industry-news/2025/the-rise-of-shadow-ai-auditing-unauthorized-ai-tools-in-the-enterprise>
- <https://www.databricks.com/blog/practical-ai-governance-framework-enterprises>
- <https://www.blue-zoo.co.uk/blog/our-ai-manifesto/>
- <https://www.equity.org.uk/news/2025/equity-calls-for-stronger-ai-protections-for-creative-workers>
- <https://www.mapfre.com/en/insights/innovation/mapfres-artificial-intelligence-manifesto/>
- <https://fairnow.ai/detect-and-prevent-shadow-ai/>
- <https://ai.ufl.edu/guidelines/>
- <https://pmc.ncbi.nlm.nih.gov/articles/PMC10941644/>
- <https://canada.newsroom.ibm.com/2025-09-03-IBM-R-Study-Shadow-AI-Use-Surges-as-Canadian-Workers-Outpace-Employers-in-AI-Adoption>

*Appendix E — Dynamics of Generative AI
Adoption in the Creative Industries*

Companion piece to *Chapter 13: Coordination Collapse* and *Chapter 9: AI in Everything, Everywhere, All at Once*.

This deep dive is the long-form quantitative companion to the “consumption gap” argument in Chapter 13 and the platform-layer analysis in Chapter 9. Where those chapters argue, in the book’s voice, that public sentiment around AI in the creative industries is sharply at odds with actual adoption telemetry — and that the gap is itself the macroeconomic story of this period — this appendix presents the underlying numbers: the Adobe Firefly adoption curve (22 billion assets by April 2025, 45% Creative Cloud penetration, 70% weekly active use, 11% of Adobe’s new ARR), the screenwriter pre- and post-strike usage data, the VFX-pipeline AI integration metrics, the LLM market structure (ChatGPT’s 800–900M WAUs, Gemini’s 155% YoY growth), the GDC game-developer sentiment-vs-usage divergence, the Quantic Foundry consumer sentiment data, and the Stanford AI Index global-acceptance findings.

Together with Appendix D: *The Shadow AI Paradox*, it constitutes the evidentiary base for the central claim of the second half of the book: that the creative industries are adopting AI at a faster rate than they are admitting publicly, and that the structure of that adoption — concentrated, hierarchical, frictioned by labour anxiety, but operationally pervasive — is the actual market environment any working creative or studio is now operating in.

The piece below is preserved largely as researched, with citation markers and section headings intact. Some PDF-conversion artefacts (loose footnote numbers, occasional line-break oddities) have not been editorially cleaned; the analytical content is what matters.

* * *

Page 1

Dynamics of Generative Artificial Intelligence Adoption in the Creative Industries: Realities, Perceptions, and the Human-Machine Paradigm The rapid integration of generative artificial intelligence (GenAI) into the global economy represents a paradigm shift comparable in magnitude to the advent of the printing press or the 1 industrial revolution. Within the creative industries—encompassing visual arts, film and television production, video game development, music, and marketing—this technological leap has triggered profound structural changes and equally profound ideological conflicts. Since the mainstream popularization of large language models (LLMs) and diffusion models in late 2022, the discourse surrounding artificial intelligence in creative fields has been characterized by intense polarization. A highly visible, deeply critical faction warns of mass technological unemployment, the degradation of human creativity, and the unchecked proliferation of algorithmic copyright infringement. 3 However, an exhaustive analysis of enterprise data, labor market statistics, software utilization metrics, and anonymous industry surveys reveals a starkly different underlying reality. The creative sector is not merely experimenting with artificial intelligence; it has already systematically embedded these tools into the foundational workflows of modern production. 5 The prevailing public narrative—which frequently pits “human authenticity” against “machine automation”—is heavily distorted by media sensationalism, algorithmic amplification of outrage, 8 and a pervasive professional stigma that forces widespread AI utilization underground. This comprehensive report evaluates the

true state of artificial intelligence adoption across the creative economy. By examining usage statistics across major platforms (such as Adobe Firefly, Suno, OpenAI, and Google Gemini), labor market impacts, shifting consumer sentiments, and the psychological mechanisms driving “AI shaming,” this analysis deconstructs the reductive “AI 9 versus human” binary to reveal the nuanced reality of a rapidly hybridizing creative ecosystem. The data overwhelmingly suggests that while a vocal minority of creators fiercely opposes the technology, a silent majority of professionals and consumers are pragmatically embracing it, forever altering the definition of modern creative labor. The Ubiquity of AI in Visual and Digital Arts To accurately assess the impact of generative artificial intelligence, one must separate performative public sentiment from private, operational utilization. The data indicates that AI is no longer a peripheral novelty but a central engine of digital content creation, accelerating at a pace that eclipses previous technological transitions. 12 The most compelling evidence of

Page 2

normalized adoption in the visual arts comes from industry-standard software providers. Adobe, whose Creative Cloud suite is the ubiquitous infrastructure for global design professionals, provides a clear lens into enterprise and individual adoption rates. Since its beta launch in March 2023, Adobe Firefly—a family of creative generative AI models—has experienced exponential, unprecedented growth. 6 The trajectory of Firefly’s asset generation demonstrates a technology that has crossed the threshold from experimental usage to daily operational reliance. By April 2025, Firefly had generated over 22 billion assets worldwide, establishing itself as one of the most rapidly adopted generative AI platforms in the 6 history of the creative industry. Milestone Date Cumulative AI Assets Intervening Growth Generated Drivers and Key Events September 2023 1 Billion Initial Beta phase and early 6 adopter exploration March 2024 6.5 Billion General availability and initial Photoshop integration 6 September 2024 12 Billion Expanded multi-app integration across the ecosystem 6 November 2024 16 Billion Introduction of Firefly Video Model at Adobe MAX 6 April 2025 22 Billion Maturation of enterprise adoption and mobile 6 expansion June 2025 24+ Billion Continued acceleration and 13 30% QoQ traffic growth By March 2024, approximately 45% of all Creative Cloud subscribers had engaged with Firefly, with 70% of active users utilizing the tool on a weekly basis. 6 User engagement averages 2.8 sessions per week, with an average session time of 26 minutes, indicating deep integration into 14 sustained workflows rather than fleeting experimentation. The demographic breakdown of these users reveals broad demographic penetration: 38% are aged 25-34, 51% are male, 46% are female, and 14% identify as LGBTQ+. 14 Furthermore, interest among Generation Z creatives grew by 32% between 2023 and 2024, signaling that the incoming cohort of

Page 3

14 professionals views AI as a native toolkit. The financial implications for the software provider underscore the immense economic value placed on these tools. Firefly contributed 11% of all Creative Cloud new annual recurring 6 revenue in 2024, pushing Adobe to a record annual revenue of \$21.51 billion. Adobe’s AI-first annual recurring revenue more than tripled

year-over-year in the first quarter of fiscal year 2026, marking generative AI as Photoshop's fastest-growing revenue catalyst since the transition to a subscription model. 13 Crucially, this adoption is not driven by hobbyists alone. Marketing agencies lead industry adoption at 63%, followed by e-commerce brands at 58% and UX/UI designers at 48%. 6 Furthermore, 72% of Fortune 500 design teams have formally integrated Firefly into their 6 corporate workflows. In the broader ecosystem, independent surveys report that 83% of professionals now utilize generative AI in their work. 15 Over 25% of new Adobe Stock content submissions in 2024 involved Firefly-generated elements, fundamentally altering the stock 14 photography and illustration market. The data definitively proves that visual artists, commercial designers, and corporate agencies are not broadly rejecting artificial intelligence; they are utilizing it at a staggering scale to achieve productivity gains and bypass menial conceptualization phases. 16

The Cinematographic Shift: Film, VFX, and Generative Video The film and television industry presents a complex landscape where highly publicized labor disputes have operated in parallel with aggressive technological integration. In 2023, the Writers Guild of America (WGA) and the Screen Actors Guild (SAG-AFTRA) executed historic strikes, 17 with the regulation of artificial intelligence serving as a core point of contention. The public optics suggested an industry violently rejecting automation. Screenwriting and the Post-Strike AI Boom However, the resolution of these strikes and the establishment of regulatory guardrails paradoxically accelerated AI adoption. Prior to the strikes in 2023, approximately 34% of screenwriters utilized AI tools covertly. 18 Following the implementation of WGA guidelines—which formally legitimized the use of AI for formatting, structural outlining, and 18 brainstorming—adoption exploded to 68% by 2024. As one industry professional noted, the official guild authorization alleviated the guilt and stigma associated with the technology, transforming it from an illicit shortcut into an accepted collaborative necessity. 18

Predictive AI platforms like Largo.ai are also being increasingly utilized by studios to analyze screenplays and forecast commercial box office viability, indicating that algorithms are influencing greenlight 17 decisions as well as the writing process itself. **Visual Effects (VFX) Automation**

Page 4

In post-production and visual effects, artificial intelligence has seamlessly integrated into the pipeline to execute computationally expensive and highly tedious tasks. The United States AI in VFX market, valued at \$1.46 billion in 2025, is projected to reach \$8.50 billion by 2035, growing at a compound annual growth rate (CAGR) of 19.24%, with some estimates projecting even 7 steeper global growth curves up to 36.1%. This growth is driven by tangible ROI (Return on Investment) metrics rather than speculative hype. Cloud-based AI deployment dominates this space (holding a 56% share) because it enables real-time rendering and remote collaboration without heavy upfront infrastructure 7 investments.

VFX Task / Application Adoption Rate & Performance Improvement Metrics

Task / Application	Adoption Rate	Performance Improvement
Automated Compositing	62% of Hollywood studios adopted	achieved 35% reduction in post-production timelines
Denoising Algorithms	71% of mid-sized firms adopted	improved 20 render quality by 28% on average
Matte Painting Generation	55% of VFX artists use AI	reduced initial setup time from 4 hours to 1.2 hours per shot
De-aging Processes	Reduced manual hours from 200 hours to 50 hours per actor (e.g., utilized in major theatrical releases)	20

Particle Simulation 68% adoption among top VFX houses 20 reported at SIGGRAPH. Major vendors such as Autodesk, Foundry, and SideFX are actively building generative pipelines into their core software offerings, indicating that machine learning is no longer a 21 separate, novelty workflow but an inherent feature of the modern VFX ecosystem. The Generative Video Race: Sora vs. Veo 3.1 The emergence of advanced generative video models has fundamentally altered pre-visualization and cinematic conceptualization. While OpenAI's Sora 2 garnered immense public attention for its photorealistic single-shot generation, the professional filmmaking sector has increasingly gravitated toward Google's Veo 3.1. 22

Page 5

Released in late 2025 by Google DeepMind, Veo 3.1 utilizes a latent diffusion transformer architecture that compresses video data into a lower-dimensional space, learning by adding and 22 removing Gaussian noise. Unlike older models that suffer from temporal amnesia, transformers process all parts of the input simultaneously, ensuring strict temporal consistency. 22 For professional directors and cinematographers, Veo 3.1 acts less as a random clip generator and more as a controllable co-director. 24 While Sora 2 excels at raw physics simulation in isolated clips, Veo 3.1 is built for the commercial production pipeline, enabling superior narrative 23 control and scene coherence. Professionals use advanced prompting to dictate specific camera composition (e.g., “smooth tracking shot,” “shallow depth of field”), precise lighting terminology (“dramatic chiaroscuro,” “Rembrandt lighting”), and direct integration of ambient sound and dialogue. 25 This allows directors to visualize scene pacing, camera angles, and emotional resonance long before physical sets are constructed, drastically lowering 25 pre-production costs. General Purpose LLMs: OpenAI, Anthropic Claude, and Google Gemini The broader landscape of creative professional AI utilization is dominated by the foundational models provided by OpenAI, Google, and Anthropic. The scale of adoption is historically unprecedented. As of 2025, OpenAI is valued at approximately \$300 billion, with annual recurring revenues projected to surpass \$12.7 billion. 26 Between April and June 2025, the OpenAI website received an average of 663.6 million monthly visits, while ChatGPT traffic alone 26 surged to nearly 5.4 billion monthly visits. Despite the proliferation of alternative models, the consumer AI assistant market exhibits a “winner take most” dynamic. ChatGPT maintains staggering dominance, boasting an estimated 27 800 million to 900 million weekly active users (WAUs) across platforms. For most of the year, fewer than 10% of ChatGPT users even visited a competitor, and only 9% of consumers pay for more than one subscription across ChatGPT, Gemini, Claude, and Cursor. 27 ChatGPT's daily active user to monthly active user (DAU/MAU) ratio of 36% nearly doubles that of Google 27 Gemini (21%), reflecting deep integration into daily professional habits. However, the competitive landscape is tightening. Anthropic's Claude 3 Opus has gained significant traction among creative writers and programmers, outperforming older GPT models in generating human-like dialogue, maintaining context over large token windows, and 28 demonstrating advanced reasoning capabilities. Concurrently, Google Gemini is experiencing explosive growth, expanding desktop users by 155% year-over-year compared to ChatGPT's 23%. 27 Much of Gemini's recent acceleration is driven by native

multimodal capabilities, such as advanced video and audio processing, which are becoming indispensable for multimedia 27 creators.

Page 6

The Video Game Industry: High Utilization Amidst Cratering Sentiment Perhaps no creative sector exhibits a wider chasm between operational adoption and public sentiment than the video game industry. Developers are caught between the intense financial pressures of a contracting labor market and a highly vocal consumer base that violently rejects the perceived “automation of art”. 29 Data from the Game Developers Conference (GDC) “State of the Game Industry” surveys spanning 2024 to 2026 illustrates a fascinating paradox: personal utilization of generative AI has steadily increased, even as industry sentiment regarding the technology has utterly collapsed. 31 Year Personal Positive Mixed Negative Usage of Sentiment Sentiment Sentiment GenAI 2024 31% 21% 57% 18% 31 2025 36% 13% 51% 31 30% 31 2026 36% 7% 30% 52% This cratering of sentiment—from 18% negative in 2024 to 52% negative by 2026—must be contextualized within the broader macroeconomic environment of the gaming sector. In 2024, one-third of developers reported direct impact from industry layoffs, and 56% expressed anxiety regarding future redundancies. 33 A staggering 84% of developers indicated they were somewhat 33 or very concerned about the ethics of using generative AI. Consequently, the hostility toward generative AI in gaming is inextricably linked to labor anxieties; AI is viewed not merely as a tool, but as a corporate instrument for workforce reduction. 32 Yet, the pragmatic reality of game development forces continued utilization. The 36% of developers actively using AI apply it primarily to productivity and administrative tasks: 81% for research and brainstorming, 47% for code assistance, 47% for daily scheduling, and 35% for rapid prototyping. 31 Usage varies significantly by role and studio size. Upper management (47%) and business/finance departments (51%) report the highest utilization, seeing the tools as efficiency drivers, whereas narrative designers, visual artists, and quality assurance testers view 32 the impact as overwhelmingly negative. Consumer sentiment in gaming mirrors this complexity. General player attitudes toward generative AI in games worsened significantly over recent years, particularly regarding the automation of creative elements. A Quantic Foundry survey revealed that gamers are 77% to

Page 7

30 83% negative toward AI-generated quests and dialogue. However, quantitative analysis reveals a more apathetic reality regarding purchasing behavior: the majority of gamers (60%) remain entirely neutral regarding the use of AI in a game’s development, provided the final product is of high quality. 35 Players demonstrate comparative openness to AI when applied to 30 non-artistic backend features, such as dynamic difficulty adjustment. The hostility is specifically reserved for the automation of roles traditionally perceived as requiring a human soul, such as narrative design and visual artistry. The Perception Gap: The Vocal Minority vs. The Silent Majority A central question surrounding the generative AI transition is whether the fierce “anti-AI” backlash represents a broad societal consensus or the disproportionate amplification of a loud minority. Empirical evidence overwhelmingly supports the latter. The

digital discourse surrounding artificial intelligence is heavily skewed by the mechanics of social media, where outrage and moral panic generate unparalleled engagement. 36 Algorithmic Amplification and the Illusion of Consensus Academic research into social media dynamics consistently demonstrates that political and cultural conversations are dominated by a highly active fraction of users. Studies of platforms like Twitter reveal a structural divide between a “vocal minority,” who tweet incessantly, utilize hashtags aggressively, link extensively to outside content, and drive ideological narratives, and 37 a “silent majority,” who consume information passively and rarely participate in public outrage. This dynamic translates directly to the AI discourse. While a vocal contingent of artists, writers, and highly engaged internet users coordinate boycotts, sign open letters, and aggressively flood comment sections with anti-AI sentiment, the broader public is quietly adopting the technology 4 for mundane, benign, and productive purposes. Broad consumer polling reveals a growing global acceptance of artificial intelligence. The Stanford AI Index Report 2025 indicates that global optimism is rising: across 26 surveyed nations, 55% of individuals now view AI products as offering more benefits than drawbacks, up 41 from 52% in 2022. YouGov polling in 2024 further indicated that nearly a third of consumers across 17 markets felt more positively about generative AI tools compared to the previous year, while only 22% held a more negative opinion. 42 While specific demographics—such as American adults—express concern about AI’s impact on interpersonal relationships and pure creativity, they simultaneously embrace it for travel planning, financial data analysis, and 43 workflow efficiency. The Box Office Stress Test The disconnect between online outrage and actual consumer behavior is most evident in the commercial performance of media products targeted by anti-AI campaigns. In 2024, the

Page 8

independent horror film *Late Night with the Devil* became the epicenter of a massive online controversy when it was revealed that the production utilized three brief frames of AI-generated 46 interstitial bumper art. Review-bombing campaigns were organized on platforms like Letterboxd, and vocal online contingents demanded a total boycott of the film, framing it as a line in the sand for artistic integrity. 48 Despite the digital fury, the boycott failed entirely to materialize in the real world. The film secured a highly successful opening weekend, taking in a symbolically appropriate \$666,666, breaking records for the distributor Shudder, and maintaining a 97% critical approval rating on Rotten Tomatoes. 46 As industry analysts noted, outside the echo chambers of social media, the general public simply did not care about the origin of a few transitional images; they paid for a 46 compelling narrative, and the controversy had “zero effect” on the film’s success. A similar dynamic unfolded with the high-budget A24 film *Civil War*. The studio released a series of promotional posters generated by AI, which depicted post-apocalyptic scenes in American 50 cities. The posters contained glaring geographical and architectural inaccuracies (such as Sutro Tower having the wrong number of antennae or buildings positioned incorrectly on the Chicago River). 50 While film Twitter and digital artists mocked the studio relentlessly, the 51 controversy generated massive organic visibility for the film. The broader consumer base viewed the posters as intended—as thematic, dystopian “what if” marketing materials—and the film succeeded commercially regardless of the digital backlash. 50 The empirical conclusion is

clear: while the anti-AI crowd is highly organized, fiercely protective of traditional artistic labor, and capable of generating immense negative public relations, their outrage rarely translates into altered consumer spending habits. The silent majority prioritizes end-product quality, utility, and entertainment value over the ethical purity of the production pipeline. 35 The Stigma of Automation: “AI Shaming” and Covert Creativity If the adoption of artificial intelligence is as widespread as the data suggests, why do so many creative professionals adamantly deny using it? The answer lies in the intense psychological and professional stigma attached to AI utilization, a phenomenon actively hindering transparent 53 integration. The Psychology of “AI Shaming” “AI shaming” has emerged as a powerful sociological mechanism used to police professional boundaries. It refers to the practice of publicly criticizing, devaluing, or dismissing individuals and organizations for utilizing artificial intelligence to execute tasks. 9 This shaming operates on the premise that AI-assisted work is inherently deceitful, devoid of human soul, and 9 fundamentally lazy.

Page 9

For creative professionals, whose core identity and societal value are inextricably linked to the struggle and mastery of their craft, the accusation of using AI is an attack on their professional legitimacy. Psychological surveys indicate that the fear of being perceived as lazy or 54 unmotivated ranks among the highest deterrents for acknowledging AI use in the workplace. Consequently, utilizing AI raises internal doubts about a professional’s own abilities, leading to a pervasive culture of secrecy. 54 Artists and writers are highly vocal about AI models being trained on their work without consent or compensation, generating an atmosphere of intense hostility 4 that bleeds into broader cultural sentiment. According to a 2023 survey, 74.3% of artists consider scraping artwork from the internet for AI technology to be highly unethical. 55 Hypocrisy in the Academy and the Studio The stigma forces usage underground, resulting in absurd institutional hypocrisy. In higher education, professors routinely threaten to fail students for utilizing ChatGPT, citing the 56 degradation of critical thinking. Yet, widespread reports indicate that students go to extreme lengths to mask their genuine work from faulty AI detectors—such as utilizing extensions like Draftback to record hours of typing sessions, 1,300 revisions, and messy drafting just to “prove” human authorship to an algorithm that falsely flagged them. 58 Conversely, faculty members openly admit to utilizing ChatGPT to grade papers, summarize reading materials, or generate the very syllabi they use to ban AI. 56 In one notable reported instance, an educator proudly utilized a paid ChatGPT subscription to detect student AI use, blissfully unaware of the LLM’s inability to accurately detect its own output, while simultaneously 56 praising an AI-generated essay as a prime example of “actual human original thinking”. This same covert dynamic exists in professional creative agencies and art galleries. When gallery owners were polled in early 2026, 61% claimed confidently that none of their 61 represented artists used AI in their practice. Yet, when artists themselves are surveyed anonymously, the numbers shift. While many claim to reject generative image models, 13% openly admit to using AI in the “backend” of the creative process, utilizing it for writing artist statements, image editing, studio organization, and conceptual planning. 61 Furthermore, deep analysis of AI art platform usage suggests that more than 11% of traditional artists have utilized text-to-image

technology, with 53.6% of those users claiming they made a “fundamental input” 63 to the artwork through their complex prompting. They absorb the massive productivity gains while outwardly maintaining the facade of pure, unassisted human toil. 16 Media Sensationalism and the Algorithmic Fog of War The disconnect between the reality of AI integration and the public panic is largely manufactured and sustained by the global media ecosystem. Journalism, functioning within an attention economy driven by programmatic advertising, is financially incentivized to sensationalize the 8 impacts of artificial intelligence.

Page 10

The “If It Bleeds, It Leads” Ecosystem Researchers analyzing media coverage of AI have noted a profound tendency toward apocalyptic framing. The narrative frequently leaps past practical, immediate concerns (such as data privacy, copyright frameworks, or minor workflow disruptions) directly to existential threats: the death of art, the extinction of humanity via “killer robots,” or catastrophic global unemployment. 8 Computer scientists and AI researchers frequently express frustration with this coverage. Zachary Lipton, a machine learning professor, famously labeled media coverage of 8 artificial intelligence as “sensationalized crap” that fuels an “AI misinformation epidemic”. Researcher Nirit Weiss-Blatt has documented how the journalism of “AI panic” diverts attention away from real-world problems like algorithmic discrimination and environmental energy consumption. 8 This sensationalism is a direct byproduct of algorithmic curation. Digital platforms optimize for engagement, and psychological research demonstrates that fear, outrage, and moral polarization are the most potent drivers of user retention. 10 This creates a dangerous feedback loop: algorithms amplify social drivers of conflict, pushing users into polarized echo chambers 36 that complicate rational discourse. In geopolitical contexts, this algorithmic amplification thickens the “fog of war,” where fake drone footage, fabricated satellite images, and deepfakes are shared widely to promote inauthentic narratives and bolster public panic. 65 Consequently, nuanced reports on how AI reduces rendering times in VFX pipelines by 30% are suppressed due to a lack of emotional resonance, while stories about “AI slop” conquering the internet 4 receive massive amplification. Job Displacement Headlines vs. Economic Data The most pervasive media narrative surrounding AI in the creative economy is the imminent threat of mass technological unemployment. Headlines consistently predict the decimation of copywriters, illustrators, and entry-level developers. However, hard macroeconomic data from the labor market contradicts the panic. In 2024, an analysis by the Information Technology and Innovation Foundation (ITIF) revealed that the employment gains from AI heavily outpaced the losses. The AI sector directly generated approximately 119,900 jobs in the United States—driven by the hiring of machine learning engineers, data scientists, and the massive construction boom required for new data centers 67 (which generate an additional 3.5 local jobs for every data center job). In stark contrast, outplacement firms tracked approximately 12,700 jobs lost specifically to AI automation during the same period. 67 This displacement represented a mere 0.1% of all total layoffs in 2024. 67 Furthermore, an exhaustive global analysis by PwC—the 2025 Global AI Jobs Barometer, which analyzed nearly a billion job advertisements—concluded that AI is making workers more valuable, not less, even in sectors highly exposed to automation. 68 The U.S.

Bureau of Labor Statistics (BLS) and organizations like Gallup have found limited evidence that generative AI 69 has systematically increased unemployment or broadly reduced earnings for artists. Research

Page 11

from Anthropic assessing labor market impacts through an “observed exposure” metric found no systematic increase in unemployment for highly exposed workers since late 2022, though there is suggestive evidence that hiring for younger, entry-level workers in exposed occupations has 69 slowed. History supports this data. The introduction of digital photography, synthesizers, desktop publishing software, and CGI all generated identical panics regarding the “death” of their 15 respective industries. In every historical instance, automation reallocated work, lowered the barrier to entry, increased total output, and ultimately expanded the overall size of the creative market. 15 As noted by the Federal Reserve, the time savings generated by AI (equivalent to 1.6% of all work hours) has contributed to aggregate labor productivity increasing by 2.16% on 12 an annualized basis. The current data strongly suggests generative AI is following this exact historical precedent. The Music Industry: Deconstructing the “AI Bad, Human Good” Narrative To move past the current state of professional cognitive dissonance and media-induced panic, the creative industries must fundamentally reevaluate the rigid philosophical binary of “AI versus Human.” This binary is intellectually flawed, historically ignorant, and actively harmful to the evolution of modern art. As articulated in deep-dive analyses by market research firms like MIDiA, the landscape of AI 11 creation is incredibly nuanced, yet the discourse remains stubbornly black-and-white. The assumption that art created solely by a human is inherently virtuous, while art assisted by an algorithm is inherently corrupt, ignores the reality of modern production. The Democratization of Audio and the Suno Revolution The music industry is currently undergoing a structural realignment driven by generative audio platforms. Historically, music production required significant capital investment in studio time, hardware, and specialized audio engineering skills. 73 Generative AI has obliterated these barriers to entry, triggering a surge in synthetic music creation. The platform Suno stands as the primary catalyst in this sector. By 2025, Suno reached an annualized revenue run rate of \$150 million to \$200 million, ultimately hitting \$300 million in annual recurring revenue by early 2026 alongside a \$2.5 billion valuation. 3 The platform boasts over 2 million paid subscribers and over 100 million distinct users who actively generate 3 full-fidelity tracks from text prompts. Year Suno Annual Funding Round Valuation Revenue / ARR

Page 12

2023 Pre-revenue / Early Series A \$600M 75 traction 2024 \$50M - \$100M Series B (\$125M) \$1B 75 (estimated) 2025 \$150M - \$200M Series C (\$250M) 3 \$2.5B 3 2026 (Early) \$300M (Reported N/A >\$2.45B ARR) The macroeconomic projections for this sub-sector are substantial. Valued at roughly \$570 million in 2024, generative AI music is forecast to reach \$2.8 billion by 2030, capturing a projected 20% of streaming platform revenue and 60% of business-to-business (B2B) music libraries. 73 This aggressive expansion has incited severe backlash from traditional industry gatekeepers. Major recording labels and artist coalitions

have launched concerted legal campaigns, describing platforms like Suno and Udio as a “brazen smash and grab” and filing 3 mass infringement lawsuits via the RIAA. The Blurring Lines of Authorship Despite institutional resistance, the concept of a purely “human” track in modern music production is already a fallacy. Human vocals are routinely corrected via Auto-Tune algorithms; drum performances are quantized perfectly to a grid by digital audio workstations; synthesizers generate waveforms that no acoustic instrument could produce. Generative AI is not an alien invasion into a pristine human domain; it is merely the next layer of abstraction in a long history of technological mediation. 11 When streaming platforms or independent marketplaces attempt to ban “AI-generated music,” they encounter impossible enforcement challenges because the lines between AI-generated, AI-assisted, and human-created content are hopelessly blurred. 11 Attempts to police these boundaries—such as Bandcamp’s ill-fated AI ban—often result in the accidental penalization of highly innovative human artists who are using algorithms as legitimate instruments of 11 avant-garde expression. The fundamental utility of AI in music is undeniable, even at the highest echelons of prestige. The 2023 release of “Now and Then,” marketed as the final song by The Beatles, relied explicitly on artificial intelligence to isolate and extract John Lennon’s degraded 1970s vocal 76 cassette recording from the piano track. The track achieved universal commercial success, reaching number one on the UK Singles Chart. 77 Consumer surveys regarding the track revealed that 58% of US respondents and 64% of UK respondents were fully aware that AI was 78 used in its production, indicating that consumer hostility toward AI is highly context-dependent. When utilized to restore, enhance, or facilitate human intent, the technology is enthusiastically

Page 13

79 embraced by the public. Redefining Creativity and Mitigating Psychological Risks As Henry Ford famously said regarding innovation, “If I had asked my customers what they wanted, they would have told me a faster horse”. 80 True innovation often precedes consumer 80 demand, and AI represents a paradigm shift that consumers are still learning to articulate. However, the integration of AI does present genuine psychological risks. Researchers note that skills that are not exercised tend to degrade; following the principle of “use it or lose it,” over-reliance on generative AI could lead to a degradation of innate human creativity. 64 Helson’s theory of adaptation levels suggests a serious risk wherein society slowly adapts to mediocre, AI-generated content, failing to even notice the gradual loss of boundary-pushing human 81 ingenuity. Generative AI is inherently replicative; it can recombine ideas, but it struggles to generate the paradigm-breaking solutions required to solve novel human problems. 81 Therefore, the threat of generative AI is not that it will destroy human creativity, but that it exposes the mechanical, formulaic nature of much of what we previously called creativity. 2 If an AI can perfectly replicate a copywriter’s marketing email or a graphic designer’s corporate logo, 2 it suggests that the original human work lacked true creative novelty. In the algorithmic age, “being average is the worst outcome”. 2 AI establishes a new baseline of competence, instantly 73 accessible to anyone. Consequently, the premium on true human creativity—characterized by emotional resonance, cultural nuance, strategic empathy, and paradigm-breaking ideation—will skyrocket. 1 Conclusion The empirical evidence

surrounding generative artificial intelligence in the creative industries points to an irreversible, systemic integration. The highly vocal anti-AI contingent, while fiercely protective of traditional copyright and influential in shaping social media discourse, represents a statistical minority that possesses virtually no leverage over macroeconomic trends, software utilization rates, or mass consumer behavior. Consumers consistently demonstrate that they value the quality, utility, and emotional impact of a final product over the ideological purity of its supply chain. Simultaneously, the widespread stigmatization of AI usage has fostered a culture of professional hypocrisy. Millions of creatives across visual arts, screenwriting, game development, and marketing rely daily on tools like Adobe Firefly, Claude, Suno, and Veo 3.1 to remain competitive. Yet, they vehemently deny their use to protect their professional identities from algorithmic “shaming.” Media narratives predicting mass technological unemployment remain unsubstantiated by current labor statistics, heavily driven by an attention economy that rewards sensationalism over economic reality. Instead of destroying the creative economy, artificial intelligence is acting as a massive democratizing force, lowering barriers to entry, augmenting existing workflows, and

Page 14

forcing a necessary redefinition of what constitutes valuable human creativity. Moving forward, the industry must discard the reductive “AI bad, human good” dichotomy. The future of the creative economy does not lie in building regulatory moats to protect outdated workflows from algorithms. It lies in proactive adaptation: embracing AI as a powerful collaborative instrument, establishing transparent frameworks for its use, and empowering the next generation of creators to synthesize human intuition with machine efficiency to unlock entirely new forms of cultural expression. Works cited

The impact of GenAI on the creative industries | World Economic Forum, 1. accessed on May 15, 2026, <https://www.weforum.org/stories/2025/01/the-impact-of-genai-on-the-creative-industries/>

As AI rises, so does the need for more human creativity - The World Economic 2. Forum, accessed on May 15, 2026, <https://www.weforum.org/stories/2026/01/ai-and-need-for-more-human-creativity/>

3. Suno: We’ve hit 2M paid subscribers and \$300M annual revenue, accessed on May 15, 2026, <https://www.musicbusinessworldwide.com/suno-hits-2m-paid-subscribers-300m-a-nnual-revenue/>

4. Public Sentiment Toward AI Is Negative: What It Means for Builders and Businesses, accessed on May 15, 2026, <https://www.mindstudio.ai/blog/public-sentiment-ai-negative-what-it-means-for-builders>

5. AI and the creative industries | Economics Observatory, accessed on May 15, 2026, <https://www.economicsobservatory.com/wp-content/uploads/2025/12/AI-Creative-Industries.pdf>

6. Adobe Firefly Statistics And User Trends 2026 - Companies History, accessed on May 15, 2026, <https://www.companieshistory.com/adobe-firefly-statistics/>

7. AI in VFX Market Size, Trends & Growth | Industry Report - SNS Insider, accessed on May 15, 2026, <https://www.snsinsider.com/reports/ai-in-vfx-market-9394>

8. How uncritical news coverage feeds the AI hype machine | Nieman Journalism Lab, accessed on May 15, 2026, <https://www.niemanlab.org/2024/05/how-uncritical-news-coverage-feeds-the-ai-hype-machine/>

9. AI Shaming: The Silent Stigma among Academic Writers and Researchers - PubMed, accessed on May 15, 2026, <https://pubmed.ncbi.nlm.nih.gov/38977530/>

10. The

Algorithmic Wars: Artificial Intelligence and the Shaping of Conflict Narratives on Social Media Platforms | Request PDF - ResearchGate, accessed on May 15, 2026, <https://www.researchgate.net/publication/397467078> The Algorithmic Wars Artificial Intelligence and the Shaping of Conflict Narratives on Social Media Platforms

Page 15

forms The problem with the “AI versus human” music narrative, accessed on May 15, 2026, <https://www.midiaresearch.com/blog/the-problem-with-the-ai-versus-human-music-narrative> The State of Generative AI Adoption in 2025 | St. Louis Fed, accessed on May 15, 2026, <https://www.stlouisfed.org/on-the-economy/2025/nov/state-generative-ai-adoption-2025> Adobe Photoshop Statistics 2026: Users, Revenue, Market Share - SQ Magazine, 13. accessed on May 15, 2026, <https://sqmagazine.co.uk/adobe-photoshop-statistics/> 14. Adobe Firefly Statistics By Usage, Demographics And Facts (2025) - ElectroIQ, accessed on May 15, 2026, <https://electroi.com/stats/adobe-firefly-statistics/> Understanding the impact of AI in the creative industries - LSE, accessed on May 15, 2026, <https://www.lse.ac.uk/school-of-public-policy/research/lse-growth-lab/understanding-the-impact-of-ai-in-the-creative-industries> Adobe’s AI and the Creative Frontier Study Reveals Creators’ Views on the 16. Opportunities and Risks of Generative AI, accessed on May 15, 2026, <https://blog.adobe.com/en/publish/2024/10/08/adobes-ai-creative-frontier-study-reveals-creators-views-opportunities-risks-generative-ai> Hollywood Is Wrestling With the Potential of AI Screenwriting Tools - Business Insider, accessed on May 15, 2026, <https://www.businessinsider.com/ai-screenwriting-tools-hollywood-film-tv-studios-wrestling-2025-8> Why AI Screenplay Editors Are Taking Over Hollywood in 2025 - Laper, accessed on May 15, 2026, <https://laper.ai/recent-highlights/2025-11-16-why-ai-screenplay-editors-are-taking-over-ai> AI In VFX Market Growth Analysis - Size and Forecast 2025-2029 | Technavio, 19. accessed on May 15, 2026, <https://www.technavio.com/report/ai-in-vfx-market-industry-analysis> 20. 160+ AI In The Vfx Industry Statistics | 2026 Data Report - Gitnux, accessed on May 15, 2026, <https://gitnux.org/ai-in-the-vfx-industry-statistics/> 21. Gen AI Resources - Visual Effects Society, accessed on May 15, 2026, <https://vesglobal.org/gen-ai-resources-page/> 22. What Is Google Veo 2? AI Video Generation from Google DeepMind - MindStudio, accessed on May 15, 2026, <https://www.mindstudio.ai/blog/what-is-google-veo-2-video-generation> 23. The AI Video Showdown: Why Google’s Veo 3.1 Edges Out Sora 2 for Pro Storytellers, accessed on May 15, 2026, <https://medium.com/@ubumart/the-ai-video-showdown-why-googles-veo-3-1-edges-out-sora-2-for-pro-storytellers-ca691295c7c3> 24. Veo 3.1: My Hands-On Deep Dive into Google’s AI Video Revolution - CrePal, accessed on May 15, 2026,

Page 16

<https://crepal.ai/blog/agent/veo-3-1-my-hands-on-deep-dive-into-googles-ai-video-revolution/> 25. Igniting Imagination in Film Production with VEO 2025, accessed on May 15, 2026, <https://filmcode.ai/veo/> 26. 50+ OpenAI Statistics 2025 - AIPRM, accessed on May 15, 2026, <https://www.aiprm.com/openai-statistics/> 27. State of Consumer AI 2025: Product Hits, Misses, and What’s Next | Andreessen Horowitz, accessed on May 15, 2026,

<https://a16z.com/state-of-consumer-ai-2025-product-hits-misses-and-whats-next/> Artificial Intelligence in Creative Industries: Advances Prior to 2025 - arXiv, 28. accessed on May 15, 2026, <https://arxiv.org/html/2501.02725v3> 29. GDC 2024 State of the Game Industry Report, accessed on May 15, 2026, <https://reg.gdconf.com/state-of-game-industry-2024> Gamers Are Overwhelmingly Negative About Gen AI in Video Games, but 30. Attitudes Vary by Gender, Age, and Gaming Motivations. - Quantic Foundry, accessed on May 15, 2026, <https://quanticfoundry.com/2025/12/18/gen-ai/> 31. I analyzed 3 years of GDC reports on generative AI in game dev. Developers hate it more every year, but the ones using it all use it for the same thing. : r/gamedev - Reddit, accessed on May 15, 2026, https://www.reddit.com/r/gamedev/comments/1rpeeta/i_analyzed_3_years_of_gdc_reports_on_generative/ 36 percent of game workers use generative AI, but half think it's bad for the 32. industry, accessed on May 15, 2026, <https://www.gamedeveloper.com/business/one-third-of-game-workers-use-generative-ai-but-half-think-it-s-bad-for-the-industry> GDC 2024 State Of The Game Industry results: Devs discuss layoffs, generative 33. AI, and more, accessed on May 15, 2026, <https://www.gamedeveloper.com/business/gdc-2024-state-of-the-game-industry-devs-discuss-layoffs-generative-ai-and-more> GDC's State of the Games Industry survey report says one in 10 respondents laid 34. off in 2024, accessed on May 15, 2026, <https://www.eurogamer.net/gdcs-state-of-the-games-industry-survey-says-one-in-ten-respondents-laid-off-in-2024> New research: What do gamers really think about generative AI in games?, 35. accessed on May 15, 2026, <https://www.midiaresearch.com/blog/new-research-what-do-gamers-really-think-about-generative-ai-in-games> Social Drivers and Algorithmic Mechanisms on Digital Media - PMC - NIH, 36. accessed on May 15, 2026, <https://pmc.ncbi.nlm.nih.gov/articles/PMC11373151/> 37. Vocal Minority versus Silent Majority: Discovering the Opinions of the Long Tail - Computer Science, accessed on May 15, 2026, <https://cs.wellesley.edu/~pmetaxas/Silent-minority-Vocal-majority.pdf> 38. Don't Panic (Yet): Assessing the Evidence and Discourse Around Generative AI and Elections, accessed on May 15, 2026, <https://knightcolumbia.org/content/dont-panic-yet-assessing-the-evidence-and-dis>

Page 17

course-around-generative-ai-and-elections (PDF) Unveiling the silent majority: stance detection and characterization of 39. passive users on social media using collaborative filtering and graph convolutional networks - ResearchGate, accessed on May 15, 2026, <https://www.researchgate.net/publication/379582905> Unveiling the silent majority stance detection and characterization of passive users on social media using collaborative filtering and graph convolutional networks 40. How prevalent is AI "art" hate really? Are we just a vocal minority? : r/antiai - Reddit, accessed on May 15, 2026, https://www.reddit.com/r/antiai/comments/1q67z44/how_prevalent_is_ai_art_hate_really_are_we_just_a/ 41. Public Opinion | The 2025 AI Index Report | Stanford HAI, accessed on May 15, 2026, <https://hai.stanford.edu/ai-index/2025-ai-index-report/public-opinion> Has public perception of generative AI shifted? - YouGov, accessed on May 15, 2026, <https://yougov.com/articles/50484-has-public-perception-of-generative-ai-shifted> 43.

Adobe Analytics: Traffic to U.S. Retail Websites from Generative AI Sources Jumps 1200 Percent, accessed on May 15, 2026, <https://blog.adobe.com/en/publish/2025/03/17/adobe-analytics-traffic-to-us-retail-websites-from-generative-ai-sources-jumps-1200-percent> 44.

The Age of Generative AI: Over half of Americans have used generative AI and most believe it will help them be more creative - the Adobe Blog, accessed on May 15, 2026, <https://blog.adobe.com/en/publish/2024/04/22/age-generative-ai-over-half-americans-have-used-generative-ai-most-believe-will-help-them-be-more-creative> Key findings about how Americans view artificial intelligence - Pew Research 45. Center, accessed on May 15, 2026, <https://www.pewresearch.org/short-reads/2026/03/12/key-findings-about-how-americans-view-artificial-intelligence/> AI Art Creeps its way into 'Late Night with the Devil' Striking BIG Controversy - 46. Reddit, accessed on May 15, 2026, https://www.reddit.com/r/aiwars/comments/1bv7mt/ai_art_creeps_its_way_into_late_night_with_the_devil_directors_confirm_use_of_ai_images_in_film - 47. Reddit, accessed on May 15, 2026, https://www.reddit.com/r/blankies/comments/1bk7kng/late_night_with_the_devil_directors_confirm_use/ Late Night With The Devil takes \$666,666 at box office, despite boycotts | Dazed, 48. accessed on May 15, 2026, <https://www.dazeddigital.com/film-tv/article/62251/1/indie-horror-late-night-with-the-devil-haunted-ai-art-panic-letterboxd-boycott> 'Late Night with the Devil' slammed for controversial use of artificial intelligence, 49. accessed on May 15, 2026, <https://roaringbengals.com/5011/features/late-night-with-the-devil-slammed-for-controversial-use-of-artificial-intelligence/>

Page 18

50. Civil War (film) - Wikipedia, accessed on May 15, 2026, [https://en.wikipedia.org/wiki/Civil_War_\(film\)](https://en.wikipedia.org/wiki/Civil_War_(film))
51. A24 made some laughably bad GenAI posters for Civil War. Enjoy this “amphibious assault” on the Palace of Fine Arts pond, while Sutro Tower is chillin’ in Marin. :) : r/sanfrancisco - Reddit, accessed on May 15, 2026, https://www.reddit.com/r/sanfrancisco/comments/1c7nzq0/a24_made_some_laughably_bad_genai_posters_for/
52. Is it me or is there not lots of marketing for Civil War? : r/boxoffice - Reddit, accessed on May 15, 2026, https://www.reddit.com/r/boxoffice/comments/1c22nbp/is_it_me_or_is_there_not_lots_of_marketing_for/
53. frazier-testimony-on-building-an-ai-ready-america.pdf - Texas Law, accessed on May 15, 2026, https://law.utexas.edu/faculty/uploads/publication_files/frazier-testimony-on-building-an-ai-ready-america.pdf
54. Evidence of a social evaluation penalty for using AI | PNAS, accessed on May 15, 2026, <https://www.pnas.org/doi/10.1073/pnas.2426766122> Survey Reveals 9 out of 10 Artists Believe Current Copyright Laws are Outdated
55. in the Age of Generative AI Technology, accessed on May 15, 2026, <https://bookanartist.co/blog/2023-artists-survey-on-ai-technology/>

56. A top student made a disturbing confession : r/Teachers - Reddit, accessed on May 15, 2026, https://www.reddit.com/r/Teachers/comments/1rvm5c3/a_top_student_made_a_disturbing_confession/
57. Drowning in AI generated essays : r/Professors - Reddit, accessed on May 15, 2026, https://www.reddit.com/r/Professors/comments/1osj8ux/drowning_in_ai_generated_essays/
58. I got accused of using AI on a paper I wrote myself and I'm losing my mind - Reddit, accessed on May 15, 2026, https://www.reddit.com/r/CollegeRant/comments/1ph02c7/i_get_accused_of_using_ai_on_a_paper_i_wrote/
59. Student sent me “real-time” video of word doc in response to fail for AI use– is this a thing?, accessed on May 15, 2026, https://www.reddit.com/r/Professors/comments/1psqlzs/student_sent_me_realtime_video_of_word_doc_in/
50. How Artificial Intelligence Shapes How We Think, Act, and Connect | Psychology Today, accessed on May 15, 2026, <https://www.psychologytoday.com/us/blog/positively-media/202501/how-artificial-intelligence-shapes-how-we-think-act-and-connect>
61. The Artsy AI Survey 2026: What Galleries Really Think About AI in the Art World, accessed on May 15, 2026, <https://www.artsy.net/article/artsy-editorial-artsy-ai-survey-2026-galleries-ai-art>
52. Over 90% of artists view AI-generated pieces in a negative light - Innovating with AI, accessed on May 15, 2026,

Page 19

<https://innovatingwithai.com/artists-view-of-ai-generated-pieces/> AI in Art Statistics 2024 - AIPRM, accessed on May 15, 2026, 63. <https://www.aiprm.com/ai-art-statistics/> 64. Is AI Good or Bad for Creativity: How Generative AI is Transforming Human Ideation, accessed on May 15, 2026, <https://www.captch.edu/blog/how-generative-ai-is-transforming-creativity> 65. The Algorithmic Management of Polarization and Violence on Social Media | Knight First Amendment Institute, accessed on May 15, 2026, <https://knightcolumbia.org/content/the-algorithmic-management-of-polarization-and-violence-on-social-media> 66. Trolling, memes, and deepfakes: How AI is thickening the fog of war - Nieman Lab, accessed on May 15, 2026, <https://www.niemanlab.org/2026/05/trolling-memes-and-deepfakes-how-ai-is-thickening-the-fog-of-war/> 67. AI's Job Impact: Gains Outpace Losses | Blogs | Dec 18, 2025 | ITIF, accessed on May 15, 2026, <https://itif.org/publications/2025/12/18/ais-job-impact-gains-outpace-losses/> AI Jobs Barometer - PwC, accessed on May 15, 2026, 68. <https://www.pwc.com/gx/en/services/ai/ai-jobs-barometer.html> 69. Labor market impacts of AI: A new measure and early evidence - Anthropic, accessed on May 15, 2026, <https://www.anthropic.com/research/labor-market-impacts> 70. How AI exposure is reshaping jobs in creative fields - Fox Business, accessed on May 15, 2026, <https://www.foxbusiness.com/technology/how-ai-exposure-reshaping-jobs-creative-fields> 71. Incorporating AI impacts in BLS employment projections: occupational case studies,

accessed on May 15, 2026, <https://www.bls.gov/opub/mlr/2025/article/incorporating-ai-impacts-in-bls-employment-projections.htm> 72. Blog: AI - MIDiA Research, accessed on May 15, 2026, <https://www.midiaresearch.com/blog/category/ai> 73. Is this the end of the recording studio? Inside Suno's 7-million-song-a-day AI revolution | MEXC News, accessed on May 15, 2026, <https://www.mexc.com/news/991797> 74. Suno Revenue and Usage Statistics (2026) - Business of Apps, accessed on May 15, 2026, <https://www.businessofapps.com/data/suno-statistics/> Suno Revenue 2025: \$200M ARR, \$2.5B Valuation - GetLatka, accessed on May 15, 2026, <https://getlatka.com/companies/suno.com> 76. The Beatles new song "Now and Then" used AI to lift out John Lennon's voice - Quartz, accessed on May 15, 2026, <https://qz.com/new-beatles-song-now-and-then-ai-john-lennon-1850981701> 77. Now and Then (Beatles song) - Wikipedia, accessed on May 15, 2026, [https://en.wikipedia.org/wiki/Now_and_Then_\(Beatles_song\)](https://en.wikipedia.org/wiki/Now_and_Then_(Beatles_song)) 78. The Beatles 'Now and Then': CintSnap reveals consumer sentiments of AI's

Page 20

impact on the music industry, accessed on May 15, 2026, <https://www.cint.com/blog/the-beatles-now-and-then-cintsnap-reveals-consumer-sentiments-of-ais-impact-on-the-music-industry/> 79. Why the Beatles' song Now and Then is the most important record of the past year | The Standard, accessed on May 15, 2026, <https://www.standard.co.uk/culture/music/the-beatles-now-and-then-let-it-be-b1124868.html> 80. Blog: artificial intelligence - MIDiA Research, accessed on May 15, 2026, <https://www.midiaresearch.com/blog/category/artificial-intelligence> Do Not Worry That Generative AI May Compromise Human Creativity or 81. Intelligence in the Future: It Already Has - PMC, accessed on May 15, 2026, <https://pmc.ncbi.nlm.nih.gov/articles/PMC11278271/> 82. The paradox of creativity in generative AI: high performance, human-like bias, and limited differential evaluation - PMC, accessed on May 15, 2026, <https://pmc.ncbi.nlm.nih.gov/articles/PMC12369561/> 83. AI vs. Human Creativity: Striking the Perfect Marketing Balance - Advertising is Simple, accessed on May 15, 2026, <https://advertisingissimple.com/ai-vs-human-creativity-striking-the-perfect-marketing-balance/>

* * *

All embedded URLs (in document order)

- <https://www.weforum.org/stories/2025/01/the-impact-of-genai-on-the-creative-industries/>
- <https://www.weforum.org/stories/2026/01/ai-and-need-for-more-human-creativity/>
- <https://www.musicbusinessworldwide.com/suno-hits-2m-paid-subscribers-300m-annual-revenue/>
- <https://www.mindstudio.ai/blog/public-sentiment-ai-negative-what-it-means-for-builders>

- https://www.economicsobservatory.com/wp-content/uploads/2025/12/AI_Creative_Industries.pdf
- <https://www.companieshistory.com/adobe-firefly-statistics/>
- <https://www.snsinsider.com/reports/ai-in-vfx-market-9394>
- <https://www.niemanlab.org/2024/05/how-uncritical-news-coverage-feeds-the-ai-hype-machine/>
- <https://pubmed.ncbi.nlm.nih.gov/38977530/>
- https://www.researchgate.net/publication/397467078_The_Algorithmic_Wars_Artificial
- <https://www.midiaresearch.com/blog/the-problem-with-the-ai-versus-human-music-narrative>
- <https://www.stlouisfed.org/on-the-economy/2025/nov/state-generative-ai-adoption-2025>
- <https://sqmagazine.co.uk/adobe-photoshop-statistics/>
- <https://electroiq.com/stats/adobe-firefly-statistics/>
- <https://www.lse.ac.uk/school-of-public-policy/research/lse-growth-lab/understanding-the-impact-of-ai-in-the-creative-industries>
- <https://blog.adobe.com/en/publish/2024/10/08/adobes-ai-creative-frontier-study-reveals-creators-views-opportunities-risks-generative-ai>
- <https://www.businessinsider.com/ai-screenwriting-tools-hollywood-film-tv-studios-writers-2025-8>
- <https://laper.ai/recent-highlights/2025-11-16-why-ai-screenplay-editors-are-taking-over>
- <https://www.technavio.com/report/ai-in-vfx-market-industry-analysis>
- <https://gitnux.org/ai-in-the-vfx-industry-statistics/>
- <https://vesglobal.org/gen-ai-resources-page/>
- <https://www.mindstudio.ai/blog/what-is-google-veo-2-video-generation>
- <https://medium.com/@ubumart/the-ai-video-showdown-why-googles-veo-3-1-edges-out-sora-2-for-pro-storytellers-ca691295c7c3>
- <https://crepal.ai/blog/agent/veo-3-1-my-hands-on-deep-dive-into-googles-ai-video-revolution/>
- <https://filmcode.ai/veo/>
- <https://www.aiprm.com/openai-statistics/>
- <https://a16z.com/state-of-consumer-ai-2025-product-hits-misses-and-whats-next/>
- <https://arxiv.org/html/2501.02725v3>
- <https://reg.gdconf.com/state-of-game-industry-2024>
- <https://quantificfoundry.com/2025/12/18/gen-ai/>
- https://www.reddit.com/r/gamedev/comments/1rpeeta/i_analyzed_3_years_of_gdc_rep
- <https://www.gamedeveloper.com/business/one-third-of-game-workers-use-generative-ai-but-half-think-it-s-bad-for-the-industry>
- <https://www.gamedeveloper.com/business/gdc-2024-state-of-the-game-industry-devs-discuss-layoffs-generative-ai-and-more>

- <https://www.eurogamer.net/gdcs-state-of-the-games-industry-survey-says-one-in-ten-respondents-laid-off-in-2024>
- <https://www.midiaresearch.com/blog/new-research-what-do-gamers-really-think-about-generative-ai-in-games>
- <https://pmc.ncbi.nlm.nih.gov/articles/PMC11373151/>
- <https://cs.wellesley.edu/~pmetaxas/Silent-minority-Vocal-majority.pdf>
- <https://knightcolumbia.org/content/dont-panic-yet-assessing-the-evidence-and-discourse-around-generative-ai-and-elections>
- https://www.researchgate.net/publication/379582905_Unveiling_the_silent_majority_st
- https://www.reddit.com/r/antiai/comments/1q67z44/how_prevalent_is_ai_art_hate_rea
- <https://hai.stanford.edu/ai-index/2025-ai-index-report/public-opinion>
- <https://yougov.com/articles/50484-has-public-perception-of-generative-ai-shifted>
- <https://blog.adobe.com/en/publish/2025/03/17/adobe-analytics-traffic-to-us-retail-websites-from-generative-ai-sources-jumps-1200-percent>
- <https://blog.adobe.com/en/publish/2024/04/22/age-generative-ai-over-half-americans-have-used-generative-ai-most-believe-will-help-them-be-more-creative>
- <https://www.pewresearch.org/short-reads/2026/03/12/key-findings-about-how-americans-view-artificial-intelligence/>
- https://www.reddit.com/r/aiewars/comments/1bvg7mt/ai_art_creeps_its_way_into_late_n
- https://www.reddit.com/r/blankies/comments/1bk7kng/late_night_with_the_devil_dire
- <https://www.dazeddigital.com/film-tv/article/62251/1/indie-horror-late-night-with-the-devil-haunted-ai-art-panic-letterboxd-boycott>
- <https://roaringbengals.com/5011/features/late-night-with-the-devil-slammed-for-controversial-use-of-artificial-intelligence/>
- [https://en.wikipedia.org/wiki/Civil_War_\(film\)](https://en.wikipedia.org/wiki/Civil_War_(film))
- https://www.reddit.com/r/sanfrancisco/comments/1c7nzq0/a24_made_some_laughably
- https://www.reddit.com/r/boxoffice/comments/1c22nbp/is_it_me_or_is_there_not_lots
- https://law.utexas.edu/faculty/uploads/publication_files/frazier-testimony-on-building-an-ai-ready-america.pdf
- <https://www.pnas.org/doi/10.1073/pnas.2426766122>
- <https://bookanartist.co/blog/2023-artists-survey-on-ai-technology/>
- https://www.reddit.com/r/Teachers/comments/1rvm5c3/a_top_student_made_a_distur
- https://www.reddit.com/r/Professors/comments/1osj8ux/drowning_in_ai_generated_es
- https://www.reddit.com/r/CollegeRant/comments/1ph02c7/i_got_accused_of_using_ai
- https://www.reddit.com/r/Professors/comments/1psqlzs/student_sent_me_realtime_vid
- <https://www.psychologytoday.com/us/blog/positively-media/202501/how-artificial-intelligence-shapes-how-we-think-act-and-connect>
- <https://www.artsy.net/article/artsy-editorial-artsy-ai-survey-2026-galleries-ai-art>
- <https://innovatingwithai.com/artists-view-of-ai-generated-pieces/>
- <https://www.aiprm.com/ai-art-statistics/>
- <https://www.captch.edu/blog/how-generative-ai-is-transforming-creativity>

- <https://knightcolumbia.org/content/the-algorithmic-management-of-polarization-and-violence-on-social-media>
- <https://www.niemanlab.org/2026/05/trolling-memes-and-deepfakes-how-ai-is-thickening-the-fog-of-war/>
- <https://itif.org/publications/2025/12/18/ais-job-impact-gains-outpace-losses/>
- <https://www.pwc.com/gx/en/services/ai/ai-jobs-barometer.html>
- <https://www.anthropic.com/research/labor-market-impacts>
- <https://www.foxbusiness.com/technology/how-ai-exposure-reshaping-jobs-creative-fields>
- <https://www.bls.gov/opub/mlr/2025/article/incorporating-ai-impacts-in-bls-employment-projections.htm>
- <https://www.midiaresearch.com/blog/category/ai>
- <https://www.mexc.com/news/991797>
- <https://www.businessofapps.com/data/suno-statistics/>
- <https://getlatka.com/companies/suno.com>
- <https://qz.com/new-beatles-song-now-and-then-ai-john-lennon-1850981701>
- [https://en.wikipedia.org/wiki/Now_and_Then_\(Beatles_song\)](https://en.wikipedia.org/wiki/Now_and_Then_(Beatles_song))
- <https://www.cint.com/blog/the-beatles-now-and-then-cintsnap-reveals-consumer-sentiments-of-ais-impact-on-the-music-industry/>
- <https://www.standard.co.uk/culture/music/the-beatles-now-and-then-let-it-be-b1124868.html>
- <https://www.midiaresearch.com/blog/category/artificial-intelligence>
- <https://pmc.ncbi.nlm.nih.gov/articles/PMC11278271/>
- <https://pmc.ncbi.nlm.nih.gov/articles/PMC12369561/>
- <https://advertisingissimple.com/ai-vs-human-creativity-striking-the-perfect-marketing-balance/>

*Appendix F — AI, Stigma, Privilege and the
Democratisation of Creative Expression*

Companion piece to *Chapter 6: The 88%*, *Chapter 13: Coordination Collapse*, and *Chapter 15: Choosing the Future*.

This deep dive sits underneath one of the harder arguments in the book — that the structural opposition to AI in some quarters of the creative industries is not only an artistic or ethical position but also, in significant measure, the *defensive response of an entrenched class* whose access to creative production has historically depended on financial, geographic and institutional barriers that AI tooling threatens to dissolve.

The book does not lead with this framing because the framing is uncomfortable and easy to weaponise. The 88% who turned up to the UK consultation, the unions defending performer likeness, the studios refusing AI integration on craft grounds — these are not, in the main, expressions of class privilege; they are legitimate articulations of working-creative interests that the book takes seriously throughout. But they sit *alongside*, and sometimes overlap with, a different cultural pattern: the resistance of a relatively narrow demographic of established creative workers to a technology that is, simultaneously, opening creative production to working-class, Global South, neurodivergent and historically excluded participants who could not previously afford the on-ramp.

This appendix lays out the sociological and economic evidence for the democratisation argument — the UK class-ceiling data, the Sutton Trust analyses of creative-industry composition, the historical precedents of artistic-elite resistance to new media, and the projections for AI’s role as an equaliser. It is the evidentiary base for the *Access* principle in Chapter 14 and for the regional-opening sections of Chapter 13. It should be read in the spirit the book itself adopts: both things are true at once, and a creative economy that takes both seriously is the one most likely to produce a humane outcome.

The piece below is preserved largely as researched, with citation markers and section headings intact. Some PDF-conversion artefacts have not been editorially cleaned.

* * *

Page 1

The Democratization of Expression: Artificial Intelligence and the Disruption of Creative Class Gatekeeping Introduction The discourse surrounding the integration of generative artificial intelligence within the global creative industries—spanning film, television, video games, music, and adjacent cultural sectors—has reached a fever pitch of moral panic, industrial resistance, and highly publicized labor disputes. Across public forums, guild negotiations, and prominent media narratives, artificial intelligence is frequently characterized as an existential threat to authentic human creativity, a mechanism for unchecked corporate plagiarism, and a harbinger of cultural decay. 1 Critics argue that the automation of artistic processes fundamentally strips the “soul” from 3 cultural production, rendering human ingenuity obsolete in the face of algorithmic efficiency. However, a rigorous sociological, economic, and historical analysis of this backlash reveals a significantly different underlying reality. The intense stigmatization of generative artificial intelligence by the established creative class is less a defense of artistic purity than a concerted, defensive effort

to preserve an entrenched socioeconomic hierarchy. 2 For decades, the creative industries have functioned not as the open meritocracies they claim to be, but as highly exclusive class systems. These sectors are heavily insulated by systemic barriers to entry, geographic clustering in expensive metropolitan hubs, steep financial prerequisites for early-career survival, and rampant, normalized nepotism. 5 The ability to create, distribute, and monetize high-level cultural products has been artificially restricted to a narrow demographic possessing immense financial backing, inherited social capital, or the 5 patronage of corporate gatekeepers. Generative artificial intelligence severely disrupts this paradigm by collapsing the cost of production and radically lowering the technical barriers to artistic execution. 8 By placing the capability to produce high-fidelity audio, cinematic video, and complex interactive code into the hands of the global public, artificial intelligence fundamentally threatens the artificial scarcity upon which the creative elite's social status, 2 professional prestige, and economic power are built. This comprehensive report provides an exhaustive examination of the intersection between artificial intelligence, social class, and the creative industries. By deconstructing the structural inequalities of the current creative economy, analyzing the historical precedents of technological resistance among artistic elites, and projecting the economic implications of artificial intelligence as an equalizing force, this analysis demonstrates that the democratization of artistic expression is not the death of creativity. Rather, it represents the dismantling of an exclusionary gatekeeping mechanism that has historically marginalized diverse, working-class,

Page 2

and global voices from the cultural vanguard. The Illusion of Meritocracy: A Statistical Deconstruction of the Creative Class System The prevailing mythology of the creative industries—heavily perpetuated by the industry's own narratives, award ceremonies, and media representations—is that of an egalitarian meritocracy. It is a landscape theoretically defined by the romantic ideal that raw talent, relentless dedication, and unique vision will invariably rise to the top, regardless of an individual's origins. Yet, deep empirical data paints a starkly different picture of an ecosystem dominated by systemic privilege, where professional success is frequently dictated by one's proximity to wealth, elite education, and institutional power. To understand why the democratization of creative tools via artificial intelligence is viewed as such a massive threat, one must first understand the exclusionary architecture of the industry it is disrupting. Statistical evidence from both the United Kingdom and the United States demonstrates that the creative workforce is overwhelmingly skewed toward the upper and middle classes, creating a rigid “class ceiling” that filters out socioeconomically disadvantaged 5 talent before they can even enter the pipeline. In the United Kingdom, for example, young people from working-class backgrounds are four times less likely to secure employment in the creative industries than their middle-class and 10 upper-class peers. Furthermore, this disparity is not a relic of the past but an accelerating trend. Analysis of demographic shifts indicates that access to creative professions has worsened considerably over the last few decades. While 16.4% of creative workers born in the 1950s and 1960s hailed from working-class backgrounds, that figure plummeted to a mere 7.9% for those born in the 1990s. 10 In high-visibility sectors such as film, television, video,

radio, and photography, individuals identifying as working-class make up just 8.4% of the total 13 workforce. The broader arts, culture, and heritage sectors exhibit similar stratification, with 60% of workers having grown up in households where the main income earner was in a managerial or professional role, compared to just 43% in the wider national workforce. 13 The dominance of elite educational backgrounds further highlights this profound social stratification. Top-selling musicians are six times more likely to have attended elite private fee-paying schools compared to the general public, sitting at 43% versus the national average of 7%. 11 The classical music profession is identified as particularly elitist; 58% of classical musicians attended an arts specialist university or conservatoire, and an astonishing 25% 11 attended the Royal Academy of Music for their undergraduate studies alone. At prestigious conservatoires, the student body is overwhelmingly affluent, with up to 60% of students studying creative subjects at institutions like the Royal Academy of Music having been privately educated. 11 Furthermore, at elite universities such as Oxford, Cambridge, King's College London, and the University of Bath, over half of all the students enrolled in creative courses 11 originate from designated "upper-middle-class" households. By contrast, working-class

Page 3

representation in creative degrees at these elite institutions languishes in the single digits, sitting at just 4% at Cambridge and Bath, and 5% at Oxford. 11 These statistics expose an educational pipeline that systematically filters out individuals who lack the financial means to access elite training, effectively reserving the highest echelons of cultural production for the already privileged. The video game industry, despite its origins as a disruptive subculture, mirrors this socioeconomic disparity. According to a UK Interactive Entertainment (UKIE) census, only 13% of professionals in the UK games sector originate from working-class backgrounds. 15 If the 15 industry were truly reflective of broader society, this figure would be closer to 37%. Industry advocates note that if socioeconomic status were classified as a protected characteristic, it would represent the single biggest diversity issue within the gaming sector, requiring an influx of over 6,000 working-class professionals just to achieve demographic parity. 15 In the United States, gaming demographics also reflect significant disparities, with the workforce remaining 16 predominantly white and male, and pay inequality persistently affecting marginalized groups. The systemic underrepresentation of the working class across film, music, and games ensures that the cultural output of these industries is inherently skewed, reflecting the perspectives, anxieties, and aesthetics of a highly insulated economic elite. The Architecture of Exclusion: Financial Barriers and Inherited Social Capital The mechanisms that maintain this "class ceiling" are primarily economic and cultural, operating through systemic financial barriers, the normalization of unpaid labor, and the pervasive influence of inherited social capital. Entry into the creative industries frequently requires navigating a labyrinth of low-paid or entirely unpaid internships, temporary contract work, and precarious freelance gigs. This is particularly prevalent in highly competitive sectors like publishing, television, and film production. 5 A comprehensive survey of creative industry professionals revealed that 67% acknowledge that unpaid internships are still a common practice within their specific fields, and an equal

percentage agree that this arrangement disproportionately benefits the upper and upper-middle classes. 5 Participating in unpaid or severely underpaid labor necessitates a 5 substantial financial safety net, typically provided by generational parental wealth. Because the vast majority of creative hubs and studios are heavily clustered in exceptionally expensive metropolitan areas—such as London, Los Angeles, and New York—individuals from lower socioeconomic backgrounds simply cannot afford the cost of living required to work for free in exchange for “exposure” or “experience”. 5 This stark economic reality ensures that the entry-level talent pool is overwhelmingly populated by those who possess the material resources to endure years of financial precarity. Working-class talent is effectively starved out of the industry before they can establish a foothold, forced to seek stable, salaried employment in other sectors to survive. 5 Furthermore,

Page 4

the financial barrier is compounded by cultural barriers. Access to creative spaces is still largely predicated on informal networks and personal contacts, creating a deeply unlevel playing field where success hinges on navigating middle-class workplace norms. 5 Preconceptions about class are often shaped by soft social identifiers—such as an individual’s accent, their vocabulary, where they went to school, and their social circles—which further alienates working-class talent who may feel compelled to alter their identities to be taken seriously by 5 affluent senior executives. Beyond sheer financial resources, success in the creative industries relies heavily on the blatant exercise of nepotism. The phenomenon of “nepo babies”—the offspring of established industry figures who secure prominent, highly visible roles with relative ease—illustrates how access 6 operates as an inherited asset rather than an earned privilege. While the term has become a popular cultural buzzword, sociological studies indicate that the underlying dynamic is a profound structural reality. Research shows that approximately 29% of Americans work for a parent’s employer at least once by age 30, a dynamic that yields significant wage premiums and early-career stability. 20 In the hyper-competitive entertainment industry, this dynamic is amplified exponentially. In Hollywood and the global music industry, nepotism rarely manifests solely as crude, direct hiring; rather, it functions through unparalleled access to elite industry networks, talent agents, studio executives, and venture capital. 6 Children of industry veterans grow up fully immersed in the specialized language, cultural norms, and social expectations of the elite. When it comes time to launch their careers, they bypass the years of cold-calling, endless auditioning, and 6 financial struggle required of outsiders. This proximity creates an unspoken, highly resourced training ground and a permanent safety net where a failed project or a bad review does not result in the end of a career, as family ties will invariably open another door. 6 Consequently, the stories that receive massive studio funding, the music that receives major label backing, and the digital art that is elevated to the cultural vanguard are overwhelmingly produced by a homogenous, privileged demographic. This dynamic narrows the cultural lens through which society views itself, restricting the diversity of narratives available to the public. 6 It is precisely this entrenched, exclusionary architecture that makes the democratizing potential of artificial intelligence so incredibly threatening to the current power brokers. When the tools of high-end production are made available to the masses, the

artificial scarcity that protects the elite is irrevocably shattered. The Weaponization of Authenticity: Unmasking the Anti-AI Backlash It is strictly within this context of extreme exclusivity and socioeconomic stratification that the vitriolic, industry-wide backlash against generative artificial intelligence must be analyzed. Across the creative sectors, guilds, and unions, criticism of artificial intelligence frequently centers on emotive themes of “theft,” “plagiarism,” the “devaluation of human effort,” and the

Page 5

impending “loss of the human soul” in art. 2 However, a deeper, critical examination of these arguments suggests that these moral and philosophical objections often serve to mask a 2 desperate defense of professional status, hierarchical privilege, and artificial scarcity. As generative artificial intelligence tools dramatically lower the learning curve required to produce high-fidelity audio, cinematic visuals, and complex written code, they directly threaten the gatekeepers who have long monopolized these capabilities. Historically, artists, musicians, and independent filmmakers have often branded themselves as anti-establishment, anti-gatekeeping, and anti-hierarchy, positioning themselves as rebels against corporate 2 control. However, the rapid advent of artificial intelligence has triggered a profound rhetorical shift among established creatives, who now actively deploy the language of authenticity, exclusivity, and tradition to defend a rigid, exclusionary hierarchy. 2 The pervasive argument that legitimate artistic expression is only valid if it is “earned” through years of formal technical training, expensive schooling, prolonged suffering, or sanctioned institutional pathways is inherently exclusionary. 2 This philosophy posits that the right to participate in cultural creation must be heavily gatekept by a gauntlet of technical and financial hurdles. When an independent, unfunded creator can utilize a generative artificial intelligence model to circumvent these traditional hurdles and execute their vision, the resulting output is 2 immediately dismissed by the established elite as “slop,” “illegitimate,” or “soulless”. This elitism reveals that the true anxiety driving the backlash is not a genuine concern for the death of creativity itself, but the terrifying realization that expressive capability is no longer a scarce, highly valuable commodity reserved exclusively for the privileged few. 2 If a working-class individual with no formal training can generate a cinematic tracking shot, compose a symphonic score, or code an interactive game environment using natural language prompts, the immense social status and economic leverage traditionally afforded to those who execute 2 these tasks severely diminishes. Furthermore, the most prominent and aggressively weaponized argument deployed by creative guilds and copyright maximalists is that artificial intelligence models are trained on copyrighted works without explicit consent, compensation, or credit, constituting a form of 23 mass, mechanized theft. While there are entirely legitimate, legally sound concerns regarding the direct impersonation of living artists—which should undoubtedly be regulated as a matter of identity protection, publicity rights, and fraud prevention—the broader, blanket argument against machine learning collapses under historical and philosophical scrutiny. 2 Human artists have always learned their craft by absorbing, analyzing, reverse-engineering, and synthesizing the copyrighted works of their predecessors. 2 A young painter studies the brushstrokes of the masters; a burgeoning writer internalizes the specific cadence,

vocabulary, and thematic structures of their favorite authors; a musician learns to play by covering the back catalog of their idols. To argue that algorithmic training is fundamentally illegitimate because it lacks explicit, prior permission is to argue that the fundamental act of learning, observation, ² and stylistic synthesis must be permissioned and monetized. Such a rigid standard would invalidate the foundational development of every living human artist, as style itself has never

Page 6

been considered proprietary property under traditional copyright frameworks. ² The hypocrisy of the anti-AI movement is further exposed when examining the corporate entities leading the charge. The institutions most vociferously championing “artist rights” against artificial intelligence—such as major Hollywood studios, the Recording Industry Association of America (RIAA), and dominant publishing conglomerates—have historically built their massive empires by systematically exploiting artists through opaque accounting practices, predatory 360-degree contracts, and the aggressive, relentless enclosure of the public domain. ²⁴ The sudden, highly publicized pivot by these corporate entities to defending the sanctity of human artistry is largely a strategic, self-serving maneuver designed to maintain ²⁴ their monopolistic control over distribution networks and content generation. These gatekeepers fear that if artificial intelligence democratizes high-end production, independent artists will no longer need to surrender their intellectual property or endure exploitative contracts to secure the capital-intensive backing of major studios and record labels. ²⁶ The corporate resistance to artificial intelligence is thus a battle for the preservation of a highly lucrative, extractive business model, thinly veiled as a crusade for creative purity. *Echoes of the Past: Historical Parallels of Technological Gatekeeping* To fully grasp the current panic surrounding artificial intelligence, it is critical to recognize that this is not an unprecedented cultural phenomenon. Rather, it is merely the latest iteration of a highly predictable historical cycle wherein established creative classes vehemently resist any new technology that threatens to democratize their medium, lower the barriers to entry, and ²⁷ dilute their exclusive status. Examining these historical parallels provides a vital lens through which to predict the inevitable trajectory and eventual integration of artificial intelligence in the creative sector. *The Synthesizer Panic and the Threat to “Real” Musicianship* In the late 1960s and stretching well into the 1980s, the introduction of electronic synthesizers, drum machines, and digital sequencers triggered widespread, existential hysteria across the global music industry. Established, formally trained musicians and high-profile critics argued that these new electronic machines produced “cold,” “inhuman,” and “artificial” sounds that replaced genuine, hard-earned skill with the effortless push of a button or the selection of a preset. ²⁷ The panic was not driven by the audiences consuming the music, but by the legacy ²⁷ institutions judging the tools and fearing the obsolescence of their specific skill sets. In the United Kingdom, the powerful Musicians’ Union went so far as to pass official motions attempting to ban the use of synthesizers, drum machines, and electronic backing devices in ²⁹ recording studios and live television performances. The union viewed these technologies as a direct, unacceptable threat to the employment of traditional orchestral session players and

live instrumentalists, framing the synthesizer not as a new instrument, but as a malicious job-killing

Page 7

machine. 29 Critics fiercely attacked musical pioneers like Miles Davis and Pete Townshend when they began incorporating electronic textures and sequencers into their compositions. Detractors claimed that the machines, rather than the musicians, were doing the actual creative work, thereby invalidating their authorship and diluting the purity of genres like jazz 27 and rock. Yet, rather than destroying the art of music, the synthesizer radically democratized it. It allowed solo artists, marginalized creators, and individuals without access to expensive studio bands to 27 compose and execute highly complex, multi-layered arrangements entirely on their own. This technological democratization ultimately gave birth to entirely new, globally dominant genres—ranging from hip-hop and synth-pop to techno, house, and electronic dance music—that became the defining cultural soundtracks of the modern era. 27 The tool that was derided as the death of human expression became the very foundation of its next evolution. The Resistance to Digital Cinematography and Home Recording A strikingly similar resistance occurred in the film and television industry during the painful, protracted transition from traditional photochemical film to digital cinematography in the late 1990s and 2000s. Elite, established cinematographers, directors, and studios fiercely argued that digital cameras inherently lacked the “soul,” the dynamic latitude, the organic grain, and 33 the specific texture of 35mm film. Early digital efforts were broadly dismissed by the Hollywood establishment as sterile, clinical, and aesthetically inferior to the “true” art of photochemical filmmaking. 33 However, the advent of highly capable, relatively affordable digital cameras—such as the RED ONE—coupled with the rise of non-linear digital editing software (like Adobe Premiere and Final Cut) running on standard personal computers, completely shattered the astronomical financial barriers to high-level filmmaking. 35 Digital technology eliminated the absolute necessity of paying exorbitant, prohibitive fees for physical film stock, specialized chemical lab processing, 35 and massive, highly specialized camera crews. This technological shift empowered an entirely new generation of independent filmmakers, operating outside the nepotistic Hollywood system, to shoot, edit, and distribute feature-length projects on micro-budgets. 35 Simultaneously, in the music industry, the rise of the Musical Instrument Digital Interface (MIDI) and affordable home digital multi-track recorders (such as the ADAT system) in the 1980s and 1990s allowed creators to bypass the traditional, highly gatekept “million-dollar commercial recording studio”. 32 Independent artists could now produce, mix, and master 9 commercial-quality, radio-ready tracks in their own bedrooms. In every historical instance, technological shifts that lowered the barrier to entry were met with fierce, coordinated resistance from industry gatekeepers who breathlessly warned of an 28 impending aesthetic and cultural collapse. Yet, in every instance, the resistance failed, and the technology ultimately resulted in a vast, unprecedented expansion of creative diversity, new artistic genres, and broader market participation from previously excluded demographics. 28

Generative artificial intelligence is not an anomaly; it is the logical, albeit highly accelerated, continuation of this historical democratizing trajectory. The Great Equalizer: Generative AI as the Catalyst for Democratization Generative artificial intelligence represents the ultimate, most profound disruption of the creative class system because it directly attacks and neutralizes the primary barrier to entry across all media: the exorbitant cost of technical execution and production value. 41 By transforming simple natural language prompts, rough sketches, or basic melodies into highly complex visual, auditory, and interactive outputs, artificial intelligence completely levels the playing field. It empowers creators who possess profound conceptual vision, storytelling ability, and taste, but who critically lack the vast financial capital required to hire specialized technical 9 teams, rent elite equipment, or secure studio backing. Eradicating the Budget-to-Vision Gap in Film and Television Historically, the art of cinema has been strictly gated by extreme economics. Executing complex establishing shots, rendering intricate computer-generated visual effects, or staging massive crowd scenes required hundreds of thousands, if not millions, of dollars in physical set construction, specialized equipment rentals, location permits, and highly unionized labor. 41 Consequently, only stories deemed broadly, safely commercially viable by a highly concentrated, risk-averse, and predominantly white, male class of studio executives were ever 41 greenlit and funded. Generative artificial intelligence tools allow independent filmmakers to bypass these financial chokepoints entirely. Creators can now procedurally generate photorealistic 3D environments, populate scenes with highly detailed digital extras, and produce complex, dynamic storyboards 8 at an infinitesimal fraction of the traditional cost. Industry financial estimates suggest that actively utilizing artificial intelligence across both pre-production and post-production workflows can seamlessly reduce the budget of a major \$200 million blockbuster film by 15% to 20%—effectively shaving \$30 to \$40 million off the bottom line and cutting weeks off the production schedule. 8 For the independent creator, the implications are revolutionary. The vast distance between imagination and execution is practically eliminated. An unfunded director operating out of a developing nation, or a working-class writer with a brilliant sci-fi concept, can now generate proof-of-concept trailers, complex visual effects, and high-fidelity scenes without needing to secure venture capital or navigate the nepotistic maze of Hollywood representation. 41 This democratization allows marginalized voices from outside the traditional geographic and social bubbles to bring their highly specific, diverse cultural narratives to the screen with a level of 41 polish previously reserved for elite studio productions.

Democratizing Game Development and Music Production The video game industry has seen the divide between highly funded “AAA” mega-studios and small independent developers widen drastically over the last decade, driven primarily by the astronomical labor costs associated with generating hyper-realistic 3D assets, vast open-world environments, and complex branching narratives. However, artificial intelligence-driven procedural content generation and advanced neural rendering technologies—such as Nvidia’s 44 DLSS 5—are acting as a tremendous “golden ticket” for the indie developer community. Small,

independent teams, or even solo developers, can now leverage generative artificial intelligence to procedurally generate expansive landscapes, populate virtual worlds with highly intelligent, adaptive non-player characters (NPCs), and achieve real-time, Hollywood-level photorealistic lighting without needing to employ a staff of hundreds of specialized artists and 44 coders. This technological leverage allows independent creators to focus their limited resources on narrative depth, unique, soulful art styles, and highly innovative gameplay mechanics, rather than competing on sheer computational brute force. This directly challenges the monolithic dominance of massive corporate publishers and injects much-needed originality into a stagnant market. 44 Similarly, in the global music industry, artificial intelligence-powered composition assistants, advanced vocal processing tools, and automated, algorithmic mastering software effectively eliminate the absolute need for expensive commercial studio time, hired session musicians, and high-end audio engineers. 9 A working-class songwriter with a compelling lyric and a basic melody can utilize artificial intelligence to instantly generate complex backing instrumentation, 9 test intricate chord progressions, and produce professional-grade, radio-ready mixes. This capability effectively bypasses the major record labels, who have historically acted as the ultimate gatekeepers by dictating radio play, funding recording sessions, and controlling 9 algorithmic playlist placement on major Digital Service Providers (DSPs) like Spotify. This sudden, unpermissioned access represents a complete paradigm shift where the ultimate artistic output and commercial viability of a track is determined by the raw quality of the idea and the taste of the creator, rather than the depth of their financial pockets or their connections to label executives. 46 The Economic Reconfiguration: Rebuilding the Creative Middle Class A frequent, highly publicized critique from the anti-AI camp—heavily promoted by creative guilds and labor unions—is that the technology will inevitably destroy millions of jobs, replacing human workers with automated systems purely to maximize corporate profits and enrich tech 23 billionaires. While it is an undeniable reality that artificial intelligence will cause significant, painful structural disruption and displace specific, highly commoditized technical roles (such as low-level copywriting, the creation of generic stock photography, basic translation, and background commercial music composition), macroeconomic analysis suggests a significantly

Page 10

more nuanced, optimistic long-term outcome. Instead of merely destroying labor, artificial intelligence possesses the unique potential to rebuild a currently hollowed-out creative middle class. 49 Extending Worker Expertise and the “Collaboration Paradox” Eminent MIT economist David Autor posits that the unique opportunity presented by artificial intelligence to the labor market is not its capacity to entirely replace human labor, but rather its 51 ability to “extend the relevance, reach, and value of human expertise”. For decades, the information age and the rise of digital technologies have paradoxically concentrated wealth, cognitive authority, and decision-making power in the hands of a small cadre of elite experts, systematically hollowing out middle-skill, middle-class jobs. 51 Generative artificial intelligence actively reverses this decades-long trend by functioning as a massive capability multiplier. Rigorous experimental studies across various professional sectors consistently

demonstrate that artificial intelligence tools disproportionately benefit lower-skilled, less-experienced, or entry-level workers. By automating the mechanical execution of tasks, AI allows these workers to rapidly close the performance and productivity gap with elite, highly paid professionals. 23 In the creative sector, this phenomenon manifests as the “Collaboration Paradox,” where access to artificial intelligence tools allows a single individual to comfortably match the output and quality of a multi-person team. 2 A junior graphic designer, a solo game developer, or an unfunded independent filmmaker can utilize artificial intelligence as an advanced, tireless “co-worker” to perform complex coding, generate storyboards, or mix audio—tasks that 23 previously required hiring a highly paid, specialized expert. While this capability understandably threatens the premium wages and job security commanded by elite technical specialists, it vastly empowers the middle tier of creators. It allows a single individual or a micro-studio to execute at a level previously reserved for large, heavily funded corporate entities, thereby redistributing the means of production. 2 The “Christmas Card Problem” and Expansive Market Dynamics Much of the intense fear surrounding AI-induced job loss relies on the fundamental assumption of a zero-sum economic market—the belief that every single AI-generated image, line of code, or background song represents a stolen commission from a human artist. This perspective entirely ignores the economic reality of market expansion, beautifully conceptualized by 2 analysts as the “Christmas card problem”. The vast majority of AI-assisted creativity actually occurs well below the commercial threshold 2 where professional, working artists operate. A small local business owner generating a logo for a pop-up shop, a high school teacher creating a custom illustration for a presentation, or an individual generating a personalized song for a family event would never have possessed the budget to hire a professional composer or a creative agency in the first place. For these users,

Page 11

the alternative to the AI-generated output was simply no output at all. 2 Therefore, artificial intelligence drastically expands the total global volume of creative expression and media generation without necessarily cannibalizing the high-end, bespoke art market, which will continue to value the specific human narrative and prestige associated with renowned artists. 2 Where artificial intelligence does actively intersect with commercial markets and displace human labor, it primarily replaces commoditized, low-effort content—such as generic royalty-free background tracks, basic templates, and standard stock images. These are categories that were already optimized for extreme low cost and mass production, rather than 2 profound artistic prestige or high wages. By automating the mundane and the commoditized, AI forces the creative industry to re-evaluate where true human value lies. Predictive Trajectories: How the Democratization Will Pan Out (2025-2030 and Beyond) As the initial shock, moral panic, and legal posturing surrounding generative artificial intelligence gradually give way to practical, everyday integration, the power dynamics of the global creative industries will undergo a profound, irreversible realignment over the next decade. Based on current economic data, historical precedents of technological adoption, and the rapidly evolving capabilities of neural networks, several highly specific predictions can be made regarding how this democratization will ultimately pan out for the creative class. 1. The Splintering of

Corporate Monopolies and the Rise of “Micro-Studios” The traditional major Hollywood studios, global game publishers, and “Big Three” record labels derive their immense, gatekeeping power almost entirely from their unique ability to finance massive production budgets, absorb enormous financial risk, and control global distribution networks. As generative artificial intelligence drastically reduces the cost of high-end production and marketing, the financial leverage held by these corporate gatekeepers will severely diminish. While major studios will undoubtedly attempt to utilize artificial intelligence internally to slash their own overhead costs and inflate profit margins—with projections indicating a shift of operational spending into AI tools for localization, dubbing, and VFX—they will simultaneously face an unprecedented wave of existential competition from highly agile, AI-empowered independent collectives. We will witness the explosive rise of the “micro-studio”: hyper-efficient teams of two to five multi-disciplinary creators who leverage artificial intelligence to produce feature-length films, AAA-quality immersive games, and chart-topping musical albums entirely independently. By completely bypassing the traditional, bloated studio system and leveraging decentralized digital distribution, these micro-studios will retain total ownership of their intellectual property and revenue streams. This will facilitate a massive redistribution of wealth away from corporate executives and legacy shareholders, directly back

Page 12

toward the actual artistic ideators and creators. 2. The Evolution of Copyright: From Protecting Style to Protecting Identity The current landscape of aggressive litigation, where artists and major corporations are suing artificial intelligence companies over the use of training data, will inevitably give way to a new legal and cultural equilibrium. This new framework will prioritize the strict protection of human identity over the impossible monopolization of artistic style. Courts, regulatory bodies, and the public will increasingly recognize that learning, analyzing, and synthesizing artistic influences are legally and philosophically permissible actions for both humans and machines. Attempting to copyright a “vibe” or a genre style will be deemed unenforceable. However, incredibly strict legal guardrails and technological detection systems will be implemented to prevent the direct, unauthorized impersonation of living artists. The unauthorized generation of a specific artist’s voice, likeness, or exact branded aesthetic—such as the viral deepfake featuring the synthesized voices of Drake and The Weeknd—will be aggressively prosecuted under expanded rights of publicity, fraud, and identity theft laws, rather than traditional copyright. The historic Writers Guild of America (WGA) and SAG-AFTRA strikes of 2023 established the fundamental blueprint for this transition. Rather than attempting to ban the technology outright, the unions secured collective bargaining agreements that ensure artificial intelligence is used to augment workers rather than replace them without compensation. These contracts secured mandatory credit and financial residuals for human authors, while explicitly allowing creators the freedom to utilize generative artificial intelligence as a tool in their own workflows. Future economic models will likely incorporate micro-licensing frameworks, blockchain-verified attribution, or industry-wide AI-royalty funds that compensate creators whose verified data heavily influences specific, commercialized outputs,

thereby creating a 64 symbiotic, sustainable ecosystem. 3. A Shift in the Definition of “Author” and “Skill” The fundamental metrics by which society evaluates, appreciates, and compensates creative talent will permanently evolve. Just as the invention of the photograph freed painting from the burden of hyper-realistic documentation—pushing the medium toward impressionism, cubism, and abstract expressionism—the advent of generative artificial intelligence will shift the perceived value of human creativity away from the mere mechanics of technical execution. 28 The successful artist of the future will function far less like a traditional craftsman and much more like a director, a curator, or a creative architect. 2 They will be a “full-stack” professional who orchestrates a complex symphony of highly specialized artificial intelligence agents. Their true value will lie in their human taste, their editorial judgment, their lived experience, and their ability to constrain, refine, and select the most emotionally resonant outputs from a sea of

Page 13

algorithmic generation. 2 The current stigma attached to “prompt engineering” or AI-assisted generation will rapidly fade as the technology becomes invisibly, seamlessly integrated into standard professional software suites, much like auto-tune in music, spell-check in literature, or 68 CGI in modern filmmaking are universally accepted today. 4. A Global Renaissance of Diverse and Marginalized Voices Ultimately, the most profound, lasting impact of generative artificial intelligence on the creative industries will be demographic and cultural. By aggressively circumventing the prohibitive financial requirements, the geographic limitations of major creative hubs, and the deeply entrenched nepotistic networks that have historically defined the creative class, artificial intelligence will unleash a massive, unprecedented renaissance of storytelling from previously excluded populations. 42 Independent creators from working-class backgrounds, artists operating within the Global South, disabled creators who face physical barriers to traditional production environments, and neurodivergent storytellers will no longer need to seek permission, beg for venture funding, or alter their identities to secure validation from a homogenous class of elite gatekeepers. 42 The democratization of these powerful tools ensures that the cultural artifacts of the mid-to-late 21st century will accurately and vibrantly reflect the full, chaotic spectrum of human experience, rather than being strictly limited to the narrow, sanitized worldview of a privileged 71 few. Conclusion The aggressive, highly publicized stigmatization of generative artificial intelligence by the established creative industries cannot be taken at face value as a righteous, purely philosophical crusade to save human art. When rigorously contextualized within the deep-seated nepotism, the astronomically steep financial barriers to entry, and the systemic, statistically proven class inequality that define the current global creative economy, the anti-AI 2 movement is starkly exposed as a defensive, reactionary maneuver by an entrenched elite. By attempting to gatekeep expressive capability behind arbitrary technical hurdles, demanding massive financial investments for production, and weaponizing outdated copyright maximalism, these legacy institutions and elite guilds are fighting desperately to preserve their social status, their professional prestige, and their highly lucrative economic monopolies. 2 History consistently demonstrates that transformative technology inevitably dismantles artificial scarcity. 40 Just as the introduction of synthesizers, digital multi-track

recorders, and digital cinema cameras radically democratized the production of music and film in previous decades, generative artificial intelligence is currently tearing down the invisible walls of the 27 modern creative class system. While this profound transition will undoubtedly cause intense short-term friction, displace specific technical roles, and require the development of entirely new legal frameworks for labor protection and identity rights, the ultimate macroeconomic and

Page 14

cultural trajectory is one of unprecedented global empowerment. 23 As the historic barriers of prohibitive cost, elite specialized training, and nepotistic access fall away, the future of the creative industries will be radically reshaped. Success will be dictated not by the wealth of the creator's parents, the exclusivity of their university, or their geographic proximity to a studio lot, but by the raw depth, originality, and emotional resonance of their ideas. 46 The democratization of creativity via artificial intelligence is not the end of art; it is the destruction of the aristocratic gatekeeper, and the beginning of true cultural liberation. Works cited 1. History, creative disruption, and GenAI | Deloitte Digital, accessed on April 7, 2026, <https://www.deloittedigital.com/us/en/insights/perspective/creative-disruption.html> 2. The elitism problem in the creative backlash against AI - Matt Hopkins, accessed on April 7, 2026, <https://matthopkins.com/technology/the-elitism-problem-in-the-creative-backlash-against-ai/> 3. Artificial Intelligence and Culture - UNESCO, accessed on April 7, 2026, https://www.unesco.org/sites/default/files/medias/fichiers/2025/09/CULTAI_Report%20of%20the%20Independent%20Expert%20Group%20on%20Artificial%20Intelligence%20and%20Culture%20%28final%20online%20version%29%201.pdf 4. Artificial Intelligence and the Creative Double Bind - Harvard Law Review, accessed on April 7, 2026, <https://harvardlawreview.org/print/vol-138/artificial-intelligence-and-the-creative-double-bind/> 5. The class ceiling in the creative industries report 2024, accessed on April 7, 2026, <https://creativeaccess.org.uk/app/uploads/2024/07/the-class-ceiling-in-the-creative-industries-report-2024.pdf> 6. Beyond the screen: The power of nepotism within Hollywood - Scot Scoop News, accessed on April 7, 2026, <https://scotscoop.com/beyond-the-screen-the-power-of-nepotism-within-hollywood/> 7. A Class Act - The Sutton Trust, accessed on April 7, 2026, <https://www.suttontrust.com/our-research/a-class-act/> 8. Transforming Media: Generative AI's Role in Cutting Film and Video Production Costs, accessed on April 7, 2026, <https://tippett.org/transforming-media-generative-ais-role-in-cutting-film-and-video-production-costs> 9. How does AI level the playing field for independent music creators? - Sonarworks Blog, accessed on April 7, 2026, <https://www.sonarworks.com/blog/learn/how-does-ai-level-the-playing-field-for-independent-music-creators> 10. Elitist Britain 2025: What It Means for the Creative, Cultural and Heritage Sector, accessed on April 7, 2026,

Page 15

<https://www.culturecommons.uk/news/elitist-britain-2025> 11. Research reveals stark class inequalities in access to the creative industries - The Sutton Trust, accessed on April 7, 2026, <https://www.suttontrust.com/news-opinion/all-news-opinion/research-reveals-stark-class-inequalities-in-access-to-the-creative-industries/> 12. Huge decline of working class people in

the arts reflects fall in wider society - The Guardian, accessed on April 7, 2026, <https://www.theguardian.com/culture/2022/dec/10/huge-decline-working-class-people-arts-reflects-society> 13. National Statistics on the Creative Industries, accessed on April 7, 2026, https://pec.ac.uk/news_entries/national-statistics-on-the-creative-industries/ 14. 50 arguments against the use of AI in creative fields - AOKIstudio, accessed on April 7, 2026, <https://aokistudio.com/50-arguments-against-the-use-of-ai-in-creative-fields.html> 15. Improving the socioeconomic diversity of the games industry | Barclays Games and Creative, accessed on April 7, 2026, <https://games.creative.barclays/resource-hub/games/industry-insights/improving-the-socioeconomic-diversity-of-the-games-industry/> 16. US Game Development Salaries in 2025: What Our Latest Industry Report Reveals, accessed on April 7, 2026, <https://gdconf.com/article/us-game-development-salaries-in-2025-what-our-latest-industry-report-reveals/> 17. GDC State Of The Games Industry 2025 report - InvestGame, accessed on April 7, 2026, <https://investgame.net/wp-content/uploads/2025/03/0794a269-d5c4-4994-9bcf-8c5730d0815e-2025-GDC-State-of-the-Game-Industry-report-1.pdf> 18. Unpaid Internships & Inequality: A Review of the Data and Recommendations for Research, Policy and Practice - CCWT, accessed on April 7, 2026, <https://ccwt.wisc.edu/wp-content/uploads/2022/04/CCWT-Policy-Brief-2-Unpaid-Internships-and-Inequality-1.pdf> 19. Did they earn it? Study looks at “nepo babies” debate | Folio - University of Alberta, accessed on April 7, 2026, <https://www.ualberta.ca/en/folio/2025/02/did-they-earn-it-nepo-babies-debate.html> 20. Is “Nepotism” Blocking Chances in the Film Industry? - The Preuss Insider, accessed on April 7, 2026, <https://preussinsider.com/5777/entertainment/is-nepotism-blocking-chances-in-the-film-industry/> 21. Economist tracks ‘nepo baby’ effect on young Americans’ earnings - Harvard Gazette, accessed on April 7, 2026, <https://news.harvard.edu/gazette/story/2023/02/economist-tracks-nepo-baby-effect-on-young-americans-earnings/> 22. Current state of Game development and How does/will AI affect the field - Reddit, accessed on April 7, 2026,

Page 16

https://www.reddit.com/r/gamedev/comments/1ih8eia/current_state_of_game_development_and_how/ 23. Generative AI, the American worker, and the future of work | Brookings, accessed on April 7, 2026, <https://www.brookings.edu/articles/generative-ai-the-american-worker-and-the-future-of-work/> 24. Copyright and AI Policy Needs Precision, Not Panic | TechPolicy.Press, accessed on April 7, 2026, <https://www.techpolicy.press/copyright-and-ai-policy-needs-precision-not-panic/> 25. AI’s Unique Threat to Musicians - Tech Policy @ Sanford - Duke University, accessed on April 7, 2026, <https://techpolicy.sanford.duke.edu/blog/ais-unique-threat-to-musicians/> 26. Old Laws, New Ghosts: Why Artists are losing the Battle for AI | diacritical - Arts Journal, accessed on April 7, 2026, <https://www.artsjournal.com/diacritical/2026/01/old-laws-new-ghosts-why-the-creative-resistance-to-ai-is-failing.html> 27. They said the same things about synthesizers. : r/DefendingAIArt - Reddit, accessed on April 7, 2026, <https://www.reddit.com/r/DefendingAIArt/comments/1qyvnf9/they-said-the-same-things-about-synthesizers/> 28. The Predictable Panic: How Artists Always React to New Technology | by Whaleden, accessed on April 7, 2026,

<https://medium.com/@whaledencom/the-predictable-panic-how-artists-always- react-to-new-technology-9482a4b92bc7> 29. 1981-1990 – The Musicians’ Union: A History (1893-2013), accessed on April 7, 2026, <https://www.muhistory.com/contact-us/1971-1980/> 30. “The union passed a motion to ban the use of synths, drum machines and any electronic devices”: the day the ‘Loony’ Musician’s Union tried to kill the synthesizer (which also happened to be Bob Moog’s birthday) | MusicRadar, accessed on April 7, 2026, <https://www.musicradar.com/news/the-union-passed-a-motion-to-ban-the-use- of-synths-drum-machines-and-any-electronic-devices-the-day-the-loony-musi cians-union-tried-to-kill-the-synthesizer-which-also-happened-to-be-bob-moo gs-birthday> 31. The UK Musicians Union tried to ban synths 40 years ago today - Mixdown Magazine, accessed on April 7, 2026, <https://mixdownmag.com.au/news/the-uk-musicians-union-tried-to-ban-synths- 40-years-ago-today/> 32. The Evolution and Decline of the Traditional Recording Studio Thesis submitted in accordance with the requirements of the Univer, accessed on April 7, 2026, https://livrepository.liverpool.ac.uk/3000867/1/200488719_Sept2015.pdf 33. Crafting Digital Cinema: Cinematographers in Contemporary Hollywood - ProQuest, accessed on April 7, 2026, <https://search.proquest.com/openview/62784b90019d00feb44512b59e25efc6/1?>

Page 17

[pq-origsite=gscholar&cbl=18750](https://www.gscholar.com/cbl=18750) 34. What We Lose When Film Cameras Change to Digital Ones - Nautilus, accessed on April 7, 2026, <https://nautil.us/what-we-lose-when-film-cameras-change-to-digital-ones-2349> 63 35. The Impact of Technology on Cinematic Storytelling | Open Access Journals, accessed on April 7, 2026, <https://www.globalmediajournal.com/open-access/the-impact-of-technology-on -cinematic-storytelling.php?aid=94036> 36. How technology has changed the film industry | Falmouth University, accessed on April 7, 2026, <https://www.falmouth.ac.uk/news/how-technology-has-changed-film-industry> 37. How the 2000s turned video and film production completely upside down - Red Shark News, accessed on April 7, 2026, <https://www.redsharknews.com/production/item/5420-the-2000-s-turned-video -and-film-production-completely-upside-down> 38. Theberge-Musical_Production_Consumption.pdf, accessed on April 7, 2026, http://www.sfu.ca/sonic-studio-webdav/AudioMedia/Readings/Alphabetical/Theb erge-Musical_Production_Consumption.pdf 39. How has digitalisation changed the economics of the creative industries?, accessed on April 7, 2026, <https://www.economicsobservatory.com/how-has-digitalisation-changed-the-ec onomics-of-the-creative-industries> 40. The Democratisation Paradox: What History Teaches Us About AI - smcleod.net, accessed on April 7, 2026, <https://smcleod.net/2025/03/the-democratisation-paradox-what-history-teaches -us-about-ai/> 41. how reference based ai is democratizing film production - Coherent Market Insights, accessed on April 7, 2026, <https://www.coherentmarketinsights.com/blog/media-and-entertainment/how-re ference-based-ai-is-democratizing-film-production-2899> 42. AI and the Democratization of Art Creation: Revolutionizing the Creative Landscape | by Myk Eff | Higher Neurons | Medium, accessed on April 7, 2026, <https://medium.com/higher-neurons/ai-and-the-democratization-of-art-creation -revolutionizing-the-creative-landscape-12c76a897561> 43.

How AI could reinvent film and TV production - McKinsey, accessed on April 7, 2026, <https://www.mckinsey.com/capabilities/tech-and-ai/our-insights/tech-forward/how-ai-could-reinvent-film-and-tv-production> 44. Is the AI Push in AAA Gaming Giving Indie Developers A Golden ..., accessed on April 7, 2026, <https://waytoomany.games/2026/03/20/is-the-ai-push-in-aaa-gaming-giving-indie-developers-a-golden-ticket/> 45. The Rise of AI in Game Development: A New Era for Indie Creators | by Olanrewaju Akinola, accessed on April 7, 2026,

Page 18

<https://medium.com/@akinola.oluwaseyi22/the-rise-of-ai-in-game-development-a-new-era-for-indie-creators-35555881efbb> 46. How AI is Democratizing Music Production - SOUNDRAW Blog, accessed on April 7, 2026, <https://soundraw.io/blog/post/ai-is-democratizing-music> 47. What will the music industry look like in 2026? - Identity Music, accessed on April 7, 2026, <https://identitymusic.com/blog/what-will-the-music-industry-look-like-in-2026> 48. FUTURE UNSCRIPTED: The Impact of Generative Artificial Intelligence on Entertainment Industry Jobs | Animation Guild, accessed on April 7, 2026, <https://animationguild.org/wp-content/uploads/2024/01/Future-Unscripted-The-Impact-of-Generative-Artificial-Intelligence-on-Entertainment-Industry-Jobs-pages-1.pdf> 49. Applying AI to Rebuild Middle Class Jobs | Request PDF - ResearchGate, accessed on April 7, 2026, https://www.researchgate.net/publication/378164927_Applying_AI_to_Rebuild_Middle_Class_Jobs 50. The impact of generative artificial intelligence on socioeconomic inequalities and policy making - PMC, accessed on April 7, 2026, <https://pmc.ncbi.nlm.nih.gov/articles/PMC11165650/> 51. Applying AI to Rebuild Middle Class Jobs - MIT Stone Center on Inequality and Shaping the Future of Work, accessed on April 7, 2026, <https://shapingwork.mit.edu/research/applying-ai-to-rebuild-middle-class-jobs/> 52. AI and the Future of Work: Opportunity or Threat? - Federal Reserve Bank of St. Louis, accessed on April 7, 2026, <https://www.stlouisfed.org/publications/page-one-economics/2024/dec/ai-and-the-future-of-work-opportunity-or-threat?print=true> 53. Artificial intelligence is transforming middle-class jobs. Can it also help the poor? | Brookings, accessed on April 7, 2026, <https://www.brookings.edu/articles/ai-transforming-middle-class-jobs-can-it-help-the-poor/> 54. AI could automate up to 26% of tasks in art, design, entertainment, and the media - UOC, accessed on April 7, 2026, <https://www.uoc.edu/en/news/2025/ai-could-automate-creative-professions> 55. AI and work in the creative industries: digital continuity or discontinuity? - Taylor & Francis, accessed on April 7, 2026, <https://www.tandfonline.com/doi/full/10.1080/17510694.2024.2421135> 56. Tomorrow comes today Trends shaping the future of the Creative Industries, accessed on April 7, 2026, <https://pec.ac.uk/wp-content/uploads/2023/12/PEC-Tomorrow-comes-today-Trends-shaping-the-future-of-the-Creative-Industries-August-2023.pdf> 57. AI Musicians vs. Machines: AI Disruption in the Music Industry | OC&C Strategy Consultants, accessed on April 7, 2026, <https://www.ocstrategy.com/en/article/ai-musicians-vs-machines-ai-disruption-in-the-music-industry/> 58. How AI Benefits—and Threatens—the Entertainment Industry - Morgan Stanley,

accessed on April 7, 2026, <https://www.morganstanley.com/insights/articles/ai-in-media-entertainment-benefits-and-risks> 59. Large studios will likely take their time adopting generative AI for content creation. Social media isn't hesitating. - Deloitte, accessed on April 7, 2026, <https://www.deloitte.com/us/en/insights/industry/technology/technology-media-and-telecom-predictions/2025/tmt-predictions-hollywood-cautious-of-genai-adoption.html> 60. CEO Keynote: AI in the Music Industry – Should You Fight It, Ignore It, or Embrace It?, accessed on April 7, 2026, <https://www.sonarworks.com/blog/research/ceo-keynote-ai-in-the-music-industry-2025> 61. The Music Industry in 2026: Trends and Opportunities | Orphiq, accessed on April 7, 2026, <https://orphiq.com/resources/music-industry-2026-trends> 62. Between Inspiration and Loss of Control: Music in the Age of AI - Ars Electronica, accessed on April 7, 2026, <https://ars.electronica.art/aeblog/en/2025/05/06/between-inspiration-and-loss-of-control/> 63. The AI Playbook: What Other Sectors Can Learn from the Creative Industry's Fight Against AI, accessed on April 7, 2026, <https://www.rand.org/pubs/commentary/2024/10/the-ai-playbook-what-other-sectors-can-learn-from-the.html> 64. How AI Models Steal Creative Work — and What to Do About It | Ed Newton-Rex | TED, accessed on April 7, 2026, <https://www.youtube.com/watch?v=U9dop96N1iw> 65. Protecting Human Creativity in AI-Generated Music with the Introduction of an AI-Royalty Fund | GRUR International | Oxford Academic, accessed on April 7, 2026, <https://academic.oup.com/grurint/article/73/12/1137/7832810> 66. This is how AI is impacting – and shaping – the creative industries, according to experts at Davos - The World Economic Forum, accessed on April 7, 2026, <https://www.weforum.org/stories/2024/02/ai-creative-industries-davos/> 67. How AI Is Rewiring Filmmaking - YouTube, accessed on April 7, 2026, https://www.youtube.com/watch?v=M_YU2OzDPfY 68. Why AI Is A Game-Changer For Creatives, And Why The Creative Industries Must Fight For Their Rights - YouTube, accessed on April 7, 2026, <https://www.youtube.com/watch?v=elQKQNewFZs> 69. The Impact of Digitalization on the Film Industry - Raindance, accessed on April 7, 2026, <https://raindance.org/the-impact-of-digitalization-on-the-film-industry/> 70. The Democratic Promise of Globalized Film and Television | Cato Institute, accessed on April 7, 2026, <https://www.cato.org/publications/trade-arts/culture> 71. Film sector on the frontlines: High-level discussion on AI in the audio-visual industry, accessed on April 7, 2026, <https://www.unesco.org/creativity/en/articles/film-sector-frontlines-high-level-discussion-ai-audio-visual-industry> 72. AI democratization bridges the gap between AI creators and users | Centific, accessed on April 7, 2026,

<https://www.centific.com/blog/ai-democratization-bridges-the-gap-between-ai-creators-and-users>

* * *

All embedded URLs (in document order)

- <https://www.deloittedigital.com/us/en/insights/perspective/creative-disruption.html>
- <https://matthopkins.com/technology/the-elitism-problem-in-the-creative-backlash-against-ai/>
- https://www.unesco.org/sites/default/files/médias/fichiers/2025/09/CULTAI_Report%2025.pdf
- <https://harvardlawreview.org/print/vol-138/artificial-intelligence-and-the-creative-double-bind/>
- <https://creativeaccess.org.uk/app/uploads/2024/07/the-class-ceiling-in-the-creative-industries-report-2024.pdf>
- <https://scotscoop.com/beyond-the-screen-the-power-of-nepotism-within-hollywood/>
- <https://www.suttontrust.com/our-research/a-class-act/>
- <https://tippett.org/transforming-media-generative-ais-role-in-cutting-film-and-video-production-costs>
- <https://www.sonarworks.com/blog/learn/how-does-ai-level-the-playing-field-for-independent-music-creators>
- <https://www.culturecommons.uk/news/elitist-britain-2025>
- <https://www.suttontrust.com/news-opinion/all-news-opinion/research-reveals-stark-class-inequalities-in-access-to-the-creative-industries/>
- <https://www.theguardian.com/culture/2022/dec/10/huge-decline-working-class-people-arts-reflects-society>
- https://pec.ac.uk/news_entries/national-statistics-on-the-creative-industries/
- <https://aokistudio.com/50-arguments-against-the-use-of-ai-in-creative-fields.html>
- <https://games.creative.barclays/resource-hub/games/industry-insights/improving-the-socioeconomic-diversity-of-the-games-industry/>
- <https://gdconf.com/article/us-game-development-salaries-in-2025-what-our-latest-industry-report-reveals/>
- https://investgame.net/wp-content/uploads/2025/03/0794a269-d5c4-4994-9bcf-8c5730d0815e_2025_GDC_State_of_the_Game_Industry_report-1.pdf
- https://ccwt.wisc.edu/wp-content/uploads/2022/04/CCWT_Policy-Brief-2_Unpaid-Internships-and-Inequality-1.pdf
- <https://www.ualberta.ca/en/folio/2025/02/did-they-earn-it-nepo-babies-debate.html>
- <https://preussinsider.com/5777/entertainment/is-nepotism-blocking-chances-in-the-film-industry/>
- <https://news.harvard.edu/gazette/story/2023/02/economist-tracks-nepo-baby-effect-on-young-americans-earnings/>
- https://www.reddit.com/r/gamedev/comments/1ih8eia/current_state_of_game_development/
- <https://www.brookings.edu/articles/generative-ai-the-american-worker-and-the-future-of-work/>
- <https://www.techpolicy.press/copyright-and-ai-policy-needs-precision-not-panic/>
- <https://techpolicy.sanford.duke.edu/blog/ais-unique-threat-to-musicians/>

- <https://www.artsjournal.com/diacritical/2026/01/old-laws-new-ghosts-why-the-creative-resistance-to-ai-is-failing.html>
- https://www.reddit.com/r/DefendingAIArt/comments/1qyynf9/they_said_the_same_thi
- <https://medium.com/@whaledencom/the-predictable-panic-how-artists-always-react-to-new-technology-9482a4b92bc7>
- <https://www.muhistory.com/contact-us/1971-1980/>
- <https://www.musicradar.com/news/the-union-passed-a-motion-to-ban-the-use-of-synths-drum-machines-and-any-electronic-devices-the-day-the-loony-musicians-union-tried-to-kill-the-synthesizer-which-also-happened-to-be-bob-moogs-birthday>
- <https://mixdownmag.com.au/news/the-uk-musicians-union-tried-to-ban-synths-40-years-ago-today/>
- https://livrepository.liverpool.ac.uk/3000867/1/200488719_Sept2015.pdf
- <https://search.proquest.com/openview/62784b90019d00feb44512b59e25efc6/1?pq-origsite=gscholar&cbl=18750>
- <https://nautil.us/what-we-lose-when-film-cameras-change-to-digital-ones-234963>
- <https://www.globalmediajournal.com/open-access/the-impact-of-technology-on-cinematic-storytelling.php?aid=94036>
- <https://www.falmouth.ac.uk/news/how-technology-has-changed-film-industry>
- <https://www.redsharknews.com/production/item/5420-the-2000-s-turned-video-and-film-production-completely-upside-down>
- http://www.sfu.ca/sonic-studio-webdav/AudioMedia/Readings/Alphabetical/Theberge-Musical_Production_Consumption.pdf
- <https://www.economicsobservatory.com/how-has-digitalisation-changed-the-economics-of-the-creative-industries>
- <https://smcleod.net/2025/03/the-democratisation-paradox-what-history-teaches-us-about-ai/>
- <https://www.coherentmarketinsights.com/blog/media-and-entertainment/how-reference-based-ai-is-democratizing-film-production-2899>
- <https://medium.com/higher-neurons/ai-and-the-democratization-of-art-creation-revolutionizing-the-creative-landscape-12c76a897561>
- <https://www.mckinsey.com/capabilities/tech-and-ai/our-insights/tech-forward/how-ai-could-reinvent-film-and-tv-production>
- <https://waytoomany.games/2026/03/20/is-the-ai-push-in-aaa-gaming-giving-indie-developers-a-golden-ticket/>
- <https://medium.com/@akinola.oluwaseyi22/the-rise-of-ai-in-game-development-a-new-era-for-indie-creators-35555881efbb>
- <https://soundraw.io/blog/post/ai-is-democratizing-music>
- <https://identitymusic.com/blog/what-will-the-music-industry-look-like-in-2026>
- <https://animationguild.org/wp-content/uploads/2024/01/Future-Unscripted-The-Impact-of-Generative-Artificial-Intelligence-on-Entertainment-Industry-Jobs-pages-1.pdf>
- https://www.researchgate.net/publication/378164927_Applying_AI_to_Rebuild_Middle
- <https://pmc.ncbi.nlm.nih.gov/articles/PMC11165650/>

- <https://shapingwork.mit.edu/research/applying-ai-to-rebuild-middle-class-jobs/>
- <https://www.stlouisfed.org/publications/page-one-economics/2024/dec/ai-and-the-future-of-work-opportunity-or-threat?print=true>
- <https://www.brookings.edu/articles/ai-transforming-middle-class-jobs-can-it-help-the-poor/>
- <https://www.uoc.edu/en/news/2025/ai-could-automate-creative-professions>
- <https://www.tandfonline.com/doi/full/10.1080/17510694.2024.2421135>
- <https://pec.ac.uk/wp-content/uploads/2023/12/PEC-Tomorrow-comes-today-Trends-shaping-the-future-of-the-Creative-Industries-August-2023.pdf>
- <https://www.ocstrategy.com/en/article/ai-musicians-vs-machines-ai-disruption-in-the-music-industry/>
- <https://www.morganstanley.com/insights/articles/ai-in-media-entertainment-benefits-and-risks>
- <https://www.deloitte.com/us/en/insights/industry/technology/technology-media-and-telecom-predictions/2025/tmt-predictions-hollywood-cautious-of-genai-adoption.html>
- <https://www.sonarworks.com/blog/research/ceo-keynote-ai-in-the-music-industry-2025>
- <https://orphiq.com/resources/music-industry-2026-trends>
- <https://ars.electronica.art/aeblog/en/2025/05/06/between-inspiration-and-loss-of-control/>
- <https://www.rand.org/pubs/commentary/2024/10/the-ai-playbook-what-other-sectors-can-learn-from-the.html>
- <https://www.youtube.com/watch?v=U9dop96Niiw>
- <https://academic.oup.com/grurint/article/73/12/1137/7832810>
- <https://www.weforum.org/stories/2024/02/ai-creative-industries-davos/>
- https://www.youtube.com/watch?v=M_YU2OzDPfY
- <https://www.youtube.com/watch?v=elQKQNewFZs>
- <https://raindance.org/the-impact-of-digitalization-on-the-film-industry/>
- <https://www.cato.org/publications/trade-arts/culture>
- <https://www.unesco.org/creativity/en/articles/film-sector-frontlines-high-level-discussion-ai-audio-visual-industry>
- <https://www.centific.com/blog/ai-democratization-bridges-the-gap-between-ai-creators-and-users>

*Appendix G — The Age of Intent: Artistic
Mastery and the Inversion of Value*

Companion piece to *Chapter 11: The Orchestrator* and *Chapter 15: Choosing the Future*.

This deep dive is the philosophical and economic companion to the book's central claim about where creative value lives after the AI transition. Chapter 11 develops the *orchestrator* role — the practical operating model of a senior creative directing a team of agents — and Chapter 14 lays out the four principles (agency, attribution, access, audience) for a humane creative economy. Both rest on a deeper proposition that this appendix makes explicit: when the technical labour of execution becomes a commodity, the intentional labour of *deciding what to make and why* becomes the scarce, valuable good.

The argument here draws on Duchamp's readymade, Arthur Danto's institutional theory of art, David Pye's distinction between the "workmanship of certainty" and the "workmanship of risk," and Rick Rubin's framing of creativity as "acts of noticing." It builds an empirical and philosophical case that the artist of 2030 is less a *manual labourer* and more an *Architect of Meaning* — a curator, editor-in-chief, and director of intent whose value is precisely the human friction the machine cannot supply.

This is, in the book's broader frame, the strongest available account of why AI is best understood as an *assistive instrument that amplifies human creativity* rather than a replacement for it. It underpins the conviction set out at the top of Chapter 15 and the operational pattern described in Chapter 11. Read it as the philosophical spine of the second half of the book.

The piece below is preserved largely as researched, with citation markers and section headings intact. Some PDF-conversion artefacts have not been editorially cleaned.

* * *

Page 1

The Age of Intent: Artistic Mastery and the Inversion of Value in the Era of Algorithmic Abundance
Introduction: The Collapse of the Technical Barrier and the Onset of the Age of Intent
We stand at a precipice in the history of human expression, a moment of rupture as profound as the invention of the printing press or the camera. For millennia, the definition of the artist was inextricably bound to the means of production—the "how." The mastery of the brush, the years spent learning to light a scene, the physical dexterity required to sculpt marble, or the mathematical precision needed to code a symphony were the gatekeepers of creation. Friction was the defining characteristic of value; difficulty was the proxy for quality. The artist was, by necessity, a technician first and a visionary second, for no vision could be realized without the hard labor of execution. Today, however, we are witnessing the total collapse of this technical barrier. The advent of Generative Artificial Intelligence has democratized production to the point of triviality. The "how" is no longer a scarcity; it is a utility. When any individual with an internet connection can generate a photorealistic image, a coherent essay, or a symphonic progression with a single natural language prompt, the value of execution creates a surplus of content but a deficit of meaning. We are rapidly approaching a state of "infinite media," where the ability to produce 1 polished, high-fidelity work is available to everyone, everywhere, all at once. In this new world, the hierarchy of

value inverts. As the labor of production approaches zero, the labor of intent—the “why” and the “what”—becomes the most valuable currency on earth. This report posits that we have entered the Age of Intent, a distinct epoch where the “how” has been solved, leaving the “why” as the sole domain of human mastery. The artist of the future reclaims their throne not as a laborer, but as a visionary—an Architect of Meaning who navigates the ocean of algorithmic competence through supreme acts of curation, selection, and philosophical grounding. This document serves as an exhaustive analysis of this transition. It explores the technical mechanisms of the “engine of probability” that drives AI, the psychological crisis of the “effort heuristic” in consumer valuation, the economic inversion of creative labor markets, and the emerging methodologies of “curatorial creation.” Drawing on the lineage of Marcel Duchamp and the philosophy of “workmanship of risk” versus “certainty,” we will construct a factual argument for why the human spirit—specifically the friction of human vulnerability—remains the essential component in a system designed for statistical conformity.

Page 2

Part I: The Mechanics of Abundance — Deconstructing the Engine of Probability To understand why intent has become the new scarcity, one must first deeply understand the nature of the abundance generated by the machine. The “content singularity”—a point where the volume of synthetic media outstrips human consumption capacity—is driven by a specific technological architecture: the probabilistic prediction engine.

1.1 The Illusion of Thought: From Tokens to Text At its most fundamental level, a Large Language Model (LLM) or a diffusion model does not “know,” “see,” or “feel” in the human sense. It operates on tokens—numerical representations of words, sub-words, or image patches.

1 When an AI generates a sentence, it is calculating the statistical probability of the next token based on the context of preceding tokens. For example, consider the sentence, “I heard a dog bark loudly at a...” The foundational unit of the LLM is the token. The model cannot process raw text directly; it operates on numbers. The sentence is segmented into tokens—“I,” “heard,” “a,” “dog,” “bark,” “loudly,” “at,” “a”—and 1 assigned numerical IDs. The model then analyzes the statistical distribution of its training data to determine that the token for “cat” has a significantly higher probability than the token for “fridge.” However, this is not a simple deterministic lookup. If it were, AI outputs would be repetitive and robotic. The “creativity” of the machine arises from the manipulation of probability through parameters like temperature, top-k, and top-p sampling.

1 • Temperature: Low temperature favors reliability, selecting the most probable next token. High temperature encourages diversity, allowing the model to select less probable tokens, introducing “novelty” or “hallucination.”

• Top-k and Top-p (Nucleus) Sampling: These methods restrict the sampling pool to the most likely candidates, renormalizing probabilities to ensure coherence while maintaining 2 variety. This mechanism creates a “central paradox”: complex, nuanced, and seemingly creative outputs emerge from a mechanism that is, at its core, a statistical prediction engine.

1 The machine is an engine of probability; it predicts the next pixel or the next word based on the average of all human creation. It is the ultimate conformist. It can answer how to render a sunset, but it cannot answer why that sunset should be rendered in a specific shade of melancholy to evoke a memory of loss.

1.2 Inference vs. Prediction: The Simulation of Reasoning

While “next-token prediction” describes the mechanical operation, it fails to capture the user experience of “inference.” Modern generative AI performs complex logical analysis within context, adjusting its strategy based on global consistency. 3 Unlike simple prediction, which might produce linear, one-directional outputs, modern models engage in a form of inference that mimics reasoning. When a user inputs a query like, “It’s a beautiful day, so we can go...”, the model considers the condition “good weather” and combines it with common sense (“good weather is suitable for 3 outdoor activities”) to deduce an appropriate next step, such as “a picnic”. This involves reasoning that considers sentence structure, context, and background knowledge, aligning more with human thinking patterns than simple statistical choice. Furthermore, these models handle “global consistency” in multi-step generation. 3 When writing an essay, the model must ensure that the conclusion aligns with the introduction. This requires a capacity for global information integration that transcends local next-token prediction. It is this capacity that allows the machine to simulate the “how” of complex creative tasks—structuring a symphony, plotting a novel, or composing a marketing strategy. However, it is crucial to distinguish this simulated reasoning from embodied cognition. The model infers based on the statistical weights of its training data, which encapsulate the logical structures of human language. It does not “understand” the picnic; it understands the statistical likelihood of the word “picnic” appearing in the context of “beautiful day.” It lacks the “embodied cognition” that gives rise to true artistic intent—the sensation of the sun, the 4 taste of the food, the memory of past picnics. 1.3 The Paradox of Abundance: A Surplus of Content, A Deficit of Meaning The democratization of this inferential power has led to a “paradox of abundance.” We are witnessing an explosion of content production that is inversely correlated with engagement and distinctiveness. 1.3.1 The Content Singularity The statistics regarding content proliferation are staggering. By 2025, the freelance platform 6 market is projected to reach \$7.65 billion, driven largely by the ease of digital production. Marketing teams using video content jumped from 63% in 2020 to 87% in 2025, with AI tools reducing production time by 75%. 7 Yet, despite this massive increase in output, engagement rates are plummeting. Social media interaction rates have fallen to below 3% across most 8 platforms, down from much higher engagement earlier in the decade. This phenomenon, termed the “content singularity,” describes an internet filled with more content than ever before, yet feeling less distinct. 8 As production becomes effortless, the ability to differentiate becomes exponentially harder. The “how” has been solved for everyone,

leading to a homogenization of aesthetics. When everyone uses the same foundational models (e.g., GPT-4, Midjourney, Stable Diffusion), the outputs tend to converge on the “statistical mean” of the training data. A sea of “authentic voices” has produced the least authentic environment marketing has ever seen. 1.3.2 Model Collapse and Cultural Homogenization A more insidious threat looms on the horizon: Model Collapse . As the web floods with AI-generated content, future models will increasingly be trained on synthetic data—data generated by other AIs. Research from Oxford and Cambridge suggests this creates a degenerative feedback loop: each generation of models trained on increasingly synthetic data

9 exhibits reduced diversity, amplified biases, and a narrowing of representational capabilities. When models train on their own outputs, they lose the “tails” of the distribution—the rare, unique, and idiosyncratic elements of human expression that drive innovation. Instead, they converge toward the center, creating a “technological monoculture”. 10 This mirrors “cultural homogenization,” where AI models, already biased toward Western perspectives, further filter out diverse expressions. The machine, left to its own devices, collapses into sameness. It requires the injection of human intent—novelty, friction, and the “workmanship of risk”—to 9 maintain cultural and semantic vitality. Phenomenon Description Consequence for Art Democratization of “How” Technical skills (rendering, Surplus of high-fidelity coding) become utilities content; technical accessible via prompts. perfection becomes baseline, not differentiator. Statistical Conformity Models predict the most Outputs tend toward the probable next token/pixel “safe” and generic; loss of based on averages. “edge” or “weirdness.” Model Collapse AI models training on Degenerative loss of AI-generated data. variance; cultural homogenization; “slop” content. Inference without Soul Simulated reasoning Art that is technically without embodied proficient but emotionally experience. hollow (“uncanny valley” of meaning).

Page 5

Part II: The Psychology of Value — The “Effort Heuristic” and the Crisis of Authenticity 2.1 The Commodity of Skill and the “Effort Heuristic” For centuries, society has operated on the “effort heuristic”—the psychological shortcut where we judge the value and quality of an object based on the perceived effort required to create it. We marveled at a photorealistic painting not just for its image, but for the years of mastery and hours of labor it represented. We respected the writer because we knew the agony of the blank page. 2.1.1 The Collapse of the Effort Heuristic AI has severed the link between quality and effort. A photorealistic image that once took 100 hours now takes 10 seconds. This unbundling of skill from creation has triggered a crisis in value perception. Research consistently shows that when consumers are aware an artwork is AI-generated, they perceive it as having less value, less emotional capacity, and lower quality, 12 even if the visual output is identical to human work. Studies reveal a distinct “implicit bias” against AI creativity. In experiments where artworks were labeled “Human” or “AI,” participants consistently rated the human-labeled works higher in liking, beauty, profundity, and worth. 16 Furthermore, gaze-tracking studies found that participants spent significantly more time looking at paintings they believed were 12 human-made compared to those labeled as AI-made. This suggests that our appreciation of art is not solely aesthetic; it is empathetic. We are connecting with the maker, not just the made. When the “how” becomes instant, it ceases to be impressive. 2.2 The “Human-Made” Premium In the Age of Intent, “Human-Made” is evolving from a descriptive tag into a luxury label. Just as “hand-made” became a premium designator in the industrial age, “human-generated” is becoming the ultimate status signal in the algorithmic age. 2.2.1 Felt Authenticity and Anti-AI Marketing The backlash against AI in marketing and art is driven by a desire for “felt authenticity”. 17 Consumers report a visceral reaction to AI content—it feels “hollow” or “weirdly empty,” like a 18 “smile with no warmth behind it”. This sentiment is backed by data: 82% of consumers worry about AI’s societal impact, and 76% say it is

extremely important to know if content is created by a real person. 17 This has given rise to “Anti-AI marketing,” a strategy where brands explicitly reject AI to build

Page 6

trust. Examples include: • Dove: Committed to never using AI to represent real bodies. • Lego: Emphasizing human creativity in their “human-made” campaigns. • Polaroid: Positioning their analog cameras as the antidote to digital/AI perfection (“The Camera for an Analog Life”). 18 Trust is the central currency here. Studies indicate that AI-generated reviews and content are 20 perceived as less genuine, leading to significantly lower purchase intent. Authenticity mediates the relationship between content and value; without the “human touch” (perceived effort, emotional risk, biological vulnerability), the content slides off the brain. 21 2.3 The Uncanny Valley of Meaning We are familiar with the “uncanny valley” in robotics—where a robot looks almost human but not quite, eliciting revulsion. AI art has created an “uncanny valley of meaning.” The machine can simulate the syntax of deep emotion (using words like “melancholy,” “loss,” “hope”), but it lacks the semantics of experience. This deficit is where the artist reclaims their value. The machine can generate a symphony, but it cannot answer why a dissonance is necessary in the third movement to reflect a personal tragedy. It operates on “certainty,” whereas human art often thrives on the “workmanship of risk”. 11 2.3.1 Workmanship of Risk vs. Certainty Drawing on the theories of David Pye, we can distinguish between the “workmanship of certainty” (mass production, automation, AI) and the “workmanship of risk” (where the quality of the result is not predetermined and depends on judgment, dexterity, and care). 23 • Workmanship of Certainty: AI generation is the ultimate form of this. The outcome is probabilistically predetermined by the model weights. It is fast, consistent, and scalable. • Workmanship of Risk: Human art involves the constant risk of failure. The brush might slip; the note might be flat. It is this vulnerability—this proximity to failure—that imbues 23 the work with “soul” and “authenticity”. In the Age of Intent, the artist’s role is to reintroduce risk. The artist must provide the friction, the contradiction, and the intent that makes art matter. Without a strong “why,” AI art is merely “content”—technically proficient slop that lacks the friction of human vulnerability. Part III: The Artist as Architect — Inverting the Hierarchy

Page 7

3.1 The Inversion: Why > How As the “how” becomes a commodity, the hierarchy of artistic value inverts. The labor of production approaches zero, while the labor of intent—the “why” and the “what”—becomes the scarcity. • Old Hierarchy: Technical Skill (High Value) > Conceptual Intent (Variable Value) • New Hierarchy: Conceptual Intent (High Value) > Curation/Selection (High Value) > Technical Skill (Commodity/Utility) The artist shifts from being a manual laborer (the hand) to an Editor-in-Chief (the mind). This is not a new concept in art history, but AI has universalized it. 3.2 The Legacy of Duchamp: The Readymade in the AI Age Marcel Duchamp’s submission of a urinal (*Fountain* , 1917) to an art exhibition was the proto-event of the AI age. Duchamp argued that the art was not in the crafting of the object, but in the act of choice . “He CHOSE it,” wrote a defender in *The Blind Man* . “He took an ordinary article of life, placed it so that its useful significance disappeared under the new

title and point of view – created a new thought for that object”.²⁵ Generative AI transforms every user into a Duchampian figure. The model produces “readymades” at scale—infinite variations of images, texts, and sounds. The creative act is no longer the rendering, but the selection and the contextualization . • Danto’s Theory: Philosopher Arthur Danto argued that what makes a Brillo box art is not its physical properties, but the “theory of art” and the context provided by the artist. Similarly, an AI image becomes art not because of its pixels, but because of the intent and theory the artist wraps around it. However, this does not mean art is “easy.” As Duchamp and the Conceptualists showed, when the object is trivial, the idea must be profound. If anyone can generate a “sunset in the style of Van Gogh,” the value lies not in the image, but in why that image was chosen, where it is placed, and what conversation it provokes. The artist becomes a “meta-creator,” operating on the level of systems and concepts rather than pigments and pixels.

3.3 The New Discipline: Curation as Creation

In a world of infinite generation, curation becomes the ultimate creative act. The artist must develop a “Curatorial Framework” to navigate the sea of noise generated by the machine.

3.3.1 The Editor-in-Chief Model

The artist’s role aligns with that of an Editor-in-Chief. The AI (the newsroom/staff writers) offers a thousand variations. The Artist (the Editor) must:

1. Reject: Say “no” to the 99% of distinct but meaningless generations. The ability to reject

Page 8

2. Select: Identify the 1% that “vibrates with truth” or novelty.
3. Refine: Direct the machine to iterate on that specific grain of truth.
4. Contextualize: Place the work in a cultural framework that gives it meaning. This requires a sophistication of taste that no algorithm can replicate. Taste is not just a preference; it is a form of knowledge—a “pattern recognition” of cultural resonance.

3.3.2 Rick Rubin and the Art of “Noticing” Music producer Rick Rubin’s philosophy of creativity is particularly relevant here. Rubin argues that “creativity is acts of noticing”.²⁹ The creator does not make the waves; they tune their antenna to receive them. In the AI context, the model is the ocean, constantly churning out possibilities. The artist is the “noticer,” the vessel with the refined filter. Rubin emphasizes that taste is a practice—a way of being. “To live as an artist is a way of being in the world. A way of perceiving. A practice of paying attention”. This “embodied attention” cannot be automated. An AI can scan a million images, but it cannot “notice” the emotional weight of a specific shade of blue in the context of human grief. It can only predict its statistical likelihood. Rubin advises creators to cultivate “awareness” and to approach creation with a “beginner’s mind,” maintaining curiosity and avoiding judgment during the initial phases of idea collection.

30 Table: The Shift from Maker to Curator

Traditional Artist	AI-Era Artist (The Architect)	Primary Skill	Physical dexterity / Technical mastery	Primary Action	Rendering / Construction	Output	A finished object	Value Source	Scarcity of skill (How)	Constraint	Physical limitations of the medium
Part IV: The Artist’s Toolset — From Prompts to Orchestration											

Page 9

4.1 The Death of “Prompt Engineering”

In the early days of Generative AI (2022-2024), “prompt engineering” was hailed as a critical technical skill. It was treated as a form of coding

—learning the “incantations” (e.g., “masterpiece, 8k, trending on artstation”) to trick the model into compliance. However, recent research and market trends suggest that prompt engineering as a distinct technical career is already obsolete. • **Obsolescence of Syntax:** As models become better at understanding natural language and nuance, the need for arcane syntax diminishes. The “post-prompt age” is characterized by systems that interpret intent from vague or incomplete instructions. 34 • **Natural Language Orchestration:** The skill is shifting to Natural Language Orchestration. It is not about writing better instructions; it is about designing interaction paradigms. The “engineer” is being replaced by the “communicator.” The best prompters are not those who know the cheat codes, but those who have a deep, nuanced vocabulary to describe mood, lighting, style, and emotion. A poet is now a better pilot for an LLM than a Python developer. 36 • **Meta-Prompting:** Advanced orchestration involves “meta-prompting,” where prompts generate other prompts, and systems critique and refine their own outputs. This requires a higher-level understanding of system architecture and behavioral psychology, moving beyond simple input-output tasks. 4.2 Case Studies in Collaborative Agency The artists successfully navigating this era are those who treat AI not as a replacement, but as a “collaborator” or a “prosthetic for the imagination.” They exemplify the “Artist as Architect” model. 4.2.1 Sougwen Chung: The Collaborator Artist and researcher Sougwen Chung rejects the “tool” metaphor entirely, viewing her robotic arms and AI systems as “collaborators”. 38 • **Method:** Chung trains her AI systems (D.O.U.G. - Drawing Operations Unit Generation) on decades of her own drawing data. The AI then controls a robotic arm that draws alongside her in real-time. • **Intent:** Her work explores “embodied cognition” and the feedback loop between human mark-making and machine mimicry. She is not outsourcing the art; she is engaging in a duet with her own data. The “why” is an exploration of memory and agency; the AI simply provides the “how” of the counter-melody. She views the AI not as an “other” but as a reflection of the self, stating, “I’ve started to see them as us in another form”. 41 4.2.2 Holly Herndon: The Sovereign Architect Musician Holly Herndon addresses the issue of agency and ownership head-on. She created

Page 10

“Holly+,” an AI vocal twin trained on her own voice, allowing others to create music using her likeness. 42 • **Method:** She utilizes “Spawning,” a protocol that allows artists to opt-in or opt-out of 42 training datasets, reasserting consent in the age of scraping. • **Intent:** Herndon’s intent is to create a “collective accomplishment.” She views AI as a coordination technology—a way to build a “choir” of intelligence. Her work *Starmirror* invites the public to train an AI model through collective singing, turning the “black box” of training into a communal ritual. 44 She is architecting the system of creation, not just the song. 4.2.3 Refik Anadol: The Data Sculptor Refik Anadol uses AI to visualize vast datasets, from brain scans to climate data. • **Method:** He treats data as pigment. His algorithms “hallucinate” new forms based on millions of images. • **Intent:** His work is about making the invisible visible—visualizing the “memory” of a machine or the “consciousness” of a library. The curation of the dataset is the art. Selecting which 100 million images to feed the model is the primary creative decision. 46 4.3 Developing a Curatorial Framework For the modern artist, developing a Curatorial Framework is the new rigorous practice. It replaces the “10,000 hours” of manual practice

with 10,000 hours of decision-making practice . The Framework Components: 1. Taste Calibration (The Input): ○ Immersion in art history and diverse media to build a “reference library” in the mind. The AI has the average of all data; the artist must possess the outliers . ○ Rick Rubin’s “tuning”: Constantly refining sensitivity to what resonates and why. 30 2. Iterative Selection (The Process): ○ 48 Adopting the “Editorial Thinking” of data visualization and design. Viewing the AI’s output not as final, but as raw footage to be edited. ○ Taste-Based Decision Making: Using “gut” and “affect” (embodied emotion) to filter rational machine outputs. 49 3. Contextual Anchoring (The Output): ○ Defining the “Why”: What is the emotional or intellectual provocation? ○ The “Human Label”: Consciously framing the work to highlight the human intent 17 behind it, leveraging the “authenticity” premium. Part V: The Economics of Intent — Market Dynamics in

Page 11

2025 5.1 The Economic Inversion The labor market is reflecting the philosophical shift. As “production” roles (copywriter, junior graphic designer) face automation pressure, “strategic” roles (Creative Director, Brand Strategist) are seeing salary growth. 5.1.1 Salary Trends: The Director vs. The Technician Data from 2025 indicates a widening gap between execution and direction roles. ● Creative Directors and VPs of Marketing (roles defined by strategy, taste, and intent) command salaries of \$145k - \$250k+. 50 ● Copywriters and Graphic Designers (execution roles) are seeing stagnation or pressure, with averages around \$57k - \$71k, though high-level specialization (e.g., AI 50 orchestration) can boost this. ● Freelance Market: While the freelance market is growing (\$7.65 billion by 2025), there is a bifurcated reality. “Routine maintenance work” is seeing rate decreases of 5-10%, while “AI Specialists” and “AI Integration Consultants” command premiums of 40-60% (\$100-\$200/hour). 53 This confirms the thesis: the value is moving away from doing the thing to directing the thing . The technician who relied solely on execution is finding their skills commoditized, while the visionary who can orchestrate these tools is seeing their value skyrocket. 5.2 The Rise of the “Human-Made” Luxury Market Just as “organic” food commands a premium over “processed” food, “human-made” content is emerging as a luxury good. ● Anti-AI Marketing: Brands are explicitly using “No AI” as a selling point. Campaigns by Dove, Lego, and others emphasize human creativity to build trust. 17 ● The Trust Deficit: With 82% of consumers worried about AI’s societal impact and 76% 17 finding it important to know if content is human-made , there is a distinct market for “certified human” work. ● Implication: Artists should not hide their use of AI, but those who don’t use it (or use it minimally) should leverage their “inefficiency” as a value signal. The “flaws” of human work—the “workmanship of risk”—become markers of authenticity. 11 Conclusion: The Triumph of Vision We are witnessing the end of the “technician artist” and the rise of the “visionary artist.” The barrier between having a thought and seeing it realized has dissolved. For the technician who

Page 12

relied solely on the difficulty of their craft to justify their value, this is a crisis. For the visionary who has been constrained by the limits of their hands or their budget, this is a liberation. The machine is an engine of probability; it can synthesize a style, but it cannot

synthesize a soul. It can answer how to render a sunset, but it cannot answer why that sunset matters. It can predict the next token, but it cannot feel the weight of the word. The “Age of Intent” demands a new kind of discipline. It is no longer about the steadiness of the hand, but the clarity of the mind. It requires the artist to be an architect of meaning, a curator of infinite possibility, and a guardian of the human spirit in an age of synthetic abundance. The artists who will reign supreme are those who treat AI not as a replacement, but as a prosthetic for the imagination. They are the ones who know that the machine can synthesize a style, but it cannot synthesize a soul. In the end, the machine is just a brush. A very complex, miraculous brush, but a brush nonetheless. And a brush cannot paint without a hand to hold it, and a mind to tell it where to go.

Appendix: Supporting Data and Frameworks

Table 1: Comparative Analysis of “Workmanship” (Based on Pye 11)	Feature	Workmanship	Risk	Workmanship
Definition	Outcome is not predetermined	predetermined	depends on quantity	production
low on judgment/dexterity	at variance	every moment	Primary Value	Uniqueness, “soul,” Perfection, speed, scale,

Page 13

perceivable effort, diversity. consistency. Flaws “Mistakes” or “happy” Hallucinations” or accidents” that reveal the “artifacts” (often seen as hand. errors to be fixed). Artist’s Role To execute the form. To disrupt the certainty; to inject risk back into the system.

Table 2: Consumer Perception of AI vs. Human Art	Metric	Human-Labeled Art	AI-Labeled Art
Liking/Preference	High	Low	(Significant negative bias)
Perceived Effort	High	Low	Emotional Response Stronger
Weaker (“Hollow”)	Gaze Duration	Longer (studied more)	Shorter (dismissed faster)
closely	Purchase Intent	Higher	Lower

Table 3: Salary Trends 2025 - The Value of Direction

50 Role Focus	Average Salary	Trend
---------------	----------------	-------

Page 14

(Est. 2025) ↑ Creative Director Intent, Vision, \$145,000 - Rising Strategy, Curation \$250,000+ ↑ VP of Marketing Strategy, Brand \$250,000 Rising Voice ↔ Copywriter Text Generation \$57,000 - \$71,000 Stagnant/Risk (Execution) ↔ Graphic Designer Image Generation \$66,000 Stagnant/Risk (Execution) ↑ AI Specialist Orchestration, \$100 - \$200 / hr High Demand Integration Works cited

1. Generative AI Foundations : From Tokens to Text, How LLMs “Write” - Medium, accessed on December 22, 2025, <https://medium.com/generative-ai-playbook/generative-ai-foundations-from-tokens-to-text-how-llms-write-05f1800703ab>
2. If generative AI is the answer, what is the question? - arXiv, accessed on December 22, 2025, <https://arxiv.org/html/2509.06120v2>
3. Why the Work of Generative AI is Called “Inference” Rather Than “Prediction” - Medium, accessed on December 22, 2025, <https://medium.com/@justyan/why-the-work-of-generative-ai-is-called-inference-rather-than-prediction-3b026a63e324>
4. From Rule to Idea: Minimalism and Conceptualism as Structural Logic - Medium, accessed on December 22, 2025, <https://medium.com/@Neuroism/from-rule-to-idea-minimalism-and-conceptualism-as-structural-logic-9441750ea552>
5. An Initial Critique: Some Strengths And Weaknesses of “Post-Theory Art”, accessed on December 22, 2025, <https://adamdaleywilson.medium.com/an-initial-critique-some-strengths-and-weaknesses->

of-post-theory-art-fe6c44846e94 6. Freelancing Statistics And Trends 2025 - Quantumrun Foresight, accessed on December 22, 2025, <https://www.quantumrun.com/consulting/freelancing-statistics/> 7. Top 10 AI Video Tools Like Clippie (2025 Comparison Guide), accessed on December 22, 2025, <https://clippie.ai/blog/best-ai-video-tools-comparison-2025>

Page 15

8. The Future of Branding in the Age of AI-Generated Content | Young Marketing Consulting, accessed on December 22, 2025, <https://www.youngmarketingconsulting.com/the-future-of-branding-in-the-age-of-ai-generated-content/>
9. Epistemic diversity across language models mitigates knowledge collapse - ResearchGate, accessed on December 22, 2025, https://www.researchgate.net/publication/398806203_Epistemic_diversity_across_language_models_mitigates_knowledge_collapse
10. AI Model Collapse: Why Data Quality Is the Defining Ethical Challenge of 2025, accessed on December 22, 2025, <https://expertlinked.in/posts/2025-10-31-ai-model-collapse-data-quality-imperative/>
11. Craft for the AI Age - The New Art School, accessed on December 22, 2025, <https://newartschool.education/2025/10/17/craft-for-the-ai-age/>
12. (PDF) Implicit bias against AI creativity - ResearchGate, accessed on December 22, 2025, https://www.researchgate.net/publication/395476845_Implicit_bias_against_AI_creativity
13. Human Art vs AI Art, a Potential Danger for Artist? : Artistic and Economic Evaluations across Multiple Art Genres By - Utrecht University Student Theses Repository Home, accessed on December 22, 2025, https://studenttheses.uu.nl/bitstream/handle/20.500.12932/47294/Thesis_JoA%C3%83%C2%ABlvanWaning%202.pdf?sequence=1&isAllowed=y
14. View of The rise of AI art: A look through digital artists' eyes | First Monday, accessed on December 22, 2025, <https://firstmonday.org/ojs/index.php/fm/article/view/13809/12001>
15. Comparative Designs Reveal Preferences for Human-Generated Rather Than AI-Generated art - PEARL, accessed on December 22, 2025, <https://pearl.plymouth.ac.uk/cgi/viewcontent.cgi?article=2175&context=psy-research>
16. Humans versus AI: whether and why we prefer human-created compared to AI-created artwork - PMC - NIH, accessed on December 22, 2025, <https://pmc.ncbi.nlm.nih.gov/articles/PMC10319694/>
17. Anti-AI Marketing Trends in 2025: Embracing Authenticity and Human Connection, accessed on December 22, 2025, <https://neuron.expert/news/the-best-anti-ai-marketing-campaigns-in-2025-and-why-they-worked/15795/en/>
18. The Best Anti-AI Marketing Campaigns in 2025, and Why They Worked | Brand Vision, accessed on December 22, 2025, <https://www.brandvm.com/post/best-anti-ai-marketing-campaigns-2025>

19. iHeartMedia Releases Third in a Series of Consumer Studies, Revealing America's Deepening Need for Human Connection in an AI-Driven World, accessed on December 22, 2025, <https://investors.heartmedia.com/news/news-details/2025/iHeartMedia-Releases-Third-in-a-Series-of-Consumer-Studies-Revealing-Americas-Deepening-Need->

Page 16

- for-Human-Connection-in-an-AI-Driven-World/default.aspx
20. Investigating the Effect of AI-Generated Customer Reviews on Purchase Intent and Perceived Authenticity in E-Commerce Environments - ResearchGate, accessed on December 22, 2025, <https://www.researchgate.net/publication/397937065> Investigating the Effect of AI-Generated Customer Reviews on Purchase Intent and Perceived Authenticity in E-Commerce Environments
21. Revealing AI Involvement in Ad Creation: Effects on Authenticity, Brand Perceptions and Consumer Intentions, accessed on December 22, 2025, <https://jisem-journal.com/index.php/journal/article/download/2659/1056/4303>
22. AI-Generated Art and Its Influence on Consumer Psychology: A Review of Economic and Emotional Outcomes - IJIP, accessed on December 22, 2025, <https://ijip.in/wp-content/uploads/2025/02/18.01.066.20251301.pdf>
23. David Pye - THE NATURE ANDARTOF, accessed on December 22, 2025, <https://memoof.me/download/884/pdf/884.pdf>
24. David Pye — The Nature and Art of Workmanship | by Mary Tsai | Making Mistakes - Medium, accessed on December 22, 2025, <https://medium.com/not-so-different-emerging-digital-craft-practices/david-pye-the-nature-and-art-of-workmanship-be0620f95aa9>
25. The Expanding Canvas: What Counts as Art? - HastingsNow, accessed on December 22, 2025, <https://www.hastingsnow.com/blog/the-expanding-canvas-what-counts-as-art>
26. ALICE C HELLIWELL - THESIS - PHILOSOPHY OF AI ART - KAR UPLOAD REDACTED - Kent Academic Repository, accessed on December 22, 2025, https://kar.kent.ac.uk/105246/1/105ALICE_C_HELLOWELL_-THESIS- PHILOSOPHY OF AI ART - KAR UPLOAD REDACTED.pdf
27. — Rick Rubin, The Creative Act (2023) - RESEARCH LAB | SORRYWECAN, accessed on December 22, 2025, <https://lab.sorrywecan.com/static/pdf/the-coming-age-of-wisdom-work.pdf>
28. Everyone Suddenly Thinks They Are Rick Rubin. - Human Kind, accessed on December 22, 2025, <https://thehumankind.co/2025/07/02/everyone-suddenly-thinks-they-are-rick-rubin/>
29. Rick Rubin on Listening, Taste, and the Act of Noticing (Ep. 169 - BONUS), accessed on December 22, 2025, <https://conversationswithtyler.com/episodes/rick-rubin/>
30. The Creative Act By Rick Rubin.pdf, accessed on December 22, 2025, <https://ia800503.us.archive.org/33/items/the-creative-act-by-rick-rubin/The%20Creative%20Act%20By%20Rick%20Rubin.pdf>
31. Rick Rubin's The Creative Act: A Way of Being - TianPan.co, accessed on December 22, 2025, <https://tianpan.co/notes/2025-05-05-creativity-act>
32. Death of Prompt Engineering: AI Orchestration in 2026 - Big Blue Data Academy, accessed on December 22, 2025, <https://bigblue.academy/en/the-death-of-prompt-engineering-and-its-ruthless-resurrection-navigating-ai-orchestration-in-2026-and-beyond>

33. Why Prompt Engineering Shouldn't Exist - Snowball - Infuzu, accessed on December 22, 2025, <https://infuzu.com/blog/why-prompt-engineering-shouldn-t-exist>
34. Is prompt engineering dead? One expert describes what HR should focus on instead, accessed on December 22, 2025, <https://hrexecutive.com/is-prompt-engineering-dead-one-expert-describes-wha-t-hr-should-focus-on-instead/>
35. The Prompt Collapse: Why Natural Language Interfaces Are About to Eat Traditional Software - LBZ Advisory, accessed on December 22, 2025, <https://liatbenzur.com/2025/11/03/the-prompt-collapse-why-natural-language-interfaces-are-about-to-eat-traditional-software/>
36. Prompt Engineering: Turning Intuition into Intentional Skill | by Ranjith E - Medium, accessed on December 22, 2025, <https://medium.com/@ranjio11/prompt-engineering-turning-intuition-into-intentional-skill-09cdd27fda2>
37. Is prompt engineering still a viable skill in 2025, or is it fading fast?" : r/learnmachinelearning, accessed on December 22, 2025, https://www.reddit.com/r/learnmachinelearning/comments/1oj5wxl/is_prompt_engineering_still_a_viable_skill_in/
38. Putting The Art In Artificial Intelligence: A Conversation With Sougwen Chung, accessed on December 22, 2025, <https://sougwen.com/putting-the-art-in-artificial-intelligence-a-conversation-with-sougwen-chung>
39. Sougwen Chung: the poetics of human-machine interaction - hube magazine interview, accessed on December 22, 2025, <https://hubemag.com/sougwen-chung-the-poetics-of-human-machine-interaction>
40. Sougwen Chung: Human and machine collaboration | by Verisart - Medium, accessed on December 22, 2025, <https://medium.com/verisart/sougwen-chung-human-and-machine-collaboration-72d912d6b065>
41. Why artificial intelligence artists can be seen as 'builders', 'breakers'—or both at once, accessed on December 22, 2025, <https://www.theartnewspaper.com/2025/07/11/why-artificial-intelligence-artists-can-be-seen-as-builders-breakers-or-both-at-once>
42. Holly Herndon: The 100 Most Influential People in AI 2023 - Time Magazine, accessed on December 22, 2025, <https://time.com/collection/time100-ai/6309468/holly-herndon/>
43. Why Artists Holly Herndon and Mat Dryhurst Want Us to Stop Minimizing—And Fearing—A.I., accessed on December 22, 2025, <https://news.artnet.com/art-world/holly-herndon-and-mat-dryhurst-2551072>
44. Holly Herndon & Mat Dryhurst Starmirror - Exhibitions - KW Institute for Contemporary Art, accessed on December 22, 2025, <https://www.kw-berlin.de/en/exhibitions/holly-herndon-and-mat-dryhurst-starmirror>

45. Holly Herndon & Mat Dryhurst on Artificial Psychedelia - Editorial - Le Random, accessed on December 22, 2025, <https://www.lerandom.art/editorial/holly-herndon-mat-dryhurst-on-artificial-psychedelia>

46. The intersection of AI and art - POLITesi, accessed on December 22, 2025, https://www.politesi.polimi.it/retrieve/beat5522-3b1c-4ad4-861c-e56af3fee554/2024_10_GiacominDaSilva.pdf
47. Decoding AI Art: From Motivation to Manifestation - - Nottingham, accessed on December 22, 2025, https://eprints.nottingham.ac.uk/81162/1/Salimbeni_Guido_20276375_corrections.pdf
48. Editorial Thinking in Data Visualization | PDF | Thought - Scribd, accessed on December 22, 2025, <https://www.scribd.com/document/664543313/week-4-Analyzing-and-visualizing-data-Andy-Kirk>
49. Linear and categorical coding units in the mouse gustatory cortex drive population dynamics and behavior in taste decision-making - eLife, accessed on December 22, 2025, <https://elifesciences.org/reviewed-preprints/109313>
50. Marketing Salary & Skills Report 2025: AI Roles, Top Jobs, Pay Data - GLOZO, accessed on December 22, 2025, <https://www.glozo.com/reports/marketing-salary-skills-report-2025>
51. How to Become an Art Director: Career Path & Guide - Himalayas.app, accessed on December 22, 2025, <https://himalayas.app/career-guides/art-director>
52. Compensation | Human Resource Management homework help - SweetStudy, accessed on December 22, 2025, <https://www.sweetstudy.com/files/milestone2-3486823.docx>
53. The Global Freelance Rate Database 2026: Complete Salary & Pricing Guide for Every Industry - Jobbers, accessed on December 22, 2025, <https://www.jobbers.io/the-global-freelance-rate-database-2026-complete-salary-pricing-guide-for-every-industry/>

* * *

All embedded URLs (in document order)

- <https://medium.com/generative-ai-playbook/generative-ai-foundations-from-tokens-to-text-how-llms-write-05f1800703ab>
- <https://arxiv.org/html/2509.06120v2>
- <https://medium.com/@justyan/why-the-work-of-generative-ai-is-called-inference-rather-than-prediction-3b026a63e324>
- <https://medium.com/@Neuroism/from-rule-to-idea-minimalism-and-conceptualism-as-structural-logic-9441750ea552>
- <https://adamdaleywilson.medium.com/an-initial-critique-some-strengths-and-weaknesses-of-post-theory-art-fe6c44846e94>
- <https://www.quantumrun.com/consulting/freelancing-statistics/>
- <https://clippie.ai/blog/best-ai-video-tools-comparison-2025>
- <https://www.youngmarketingconsulting.com/the-future-of-branding-in-the-age-of-ai-generated-content/>
- https://www.researchgate.net/publication/398806203_Epistemic_diversity_across_languages
- <https://expertlinked.in/posts/2025-10-31-ai-model-collapse-data-quality-imperative/>
- <https://newartschool.education/2025/10/17/craft-for-the-ai-age/>

- https://www.researchgate.net/publication/395476845_Implicit_bias_against_AI_creativi
- https://studenttheses.uu.nl/bitstream/handle/20.500.12932/47294/Thesis_JoA%CC%83sequence=1&isAllowed=y
- <https://firstmonday.org/ojs/index.php/fm/article/view/13809/12001>
- <https://pearl.plymouth.ac.uk/cgi/viewcontent.cgi?article=2175&context=psy-research>
- <https://pmc.ncbi.nlm.nih.gov/articles/PMC10319694/>
- <https://neuron.expert/news/the-best-anti-ai-marketing-campaigns-in-2025-and-why-they-worked/15795/en/>
- <https://www.brandvm.com/post/best-anti-ai-marketing-campaigns-2025>
- <https://investors.iheartmedia.com/news/news-details/2025/iHeartMedia-Releases-Third-in-a-Series-of-Consumer-Studies-Revealing-Americas-Deepening-Need-for-Human-Connection-in-an-AI-Driven-World/default.aspx>
- https://www.researchgate.net/publication/397937065_Investigating_the_Effect_of_AI-Generated_Customer_Reviews_on_Purchase_Intent_and_Perceived_Authenticity_in_E-Commerce_Environments
- <https://jisem-journal.com/index.php/journal/article/download/2659/1056/4303>
- <https://ijip.in/wp-content/uploads/2025/02/18.01.066.20251301.pdf>
- <https://memoof.me/download/884/pdf/884.pdf>
- <https://medium.com/not-so-different-emerging-digital-craft-practices/david-pye-the-nature-and-art-of-workmanship-be0620f95aa9>
- <https://www.hastingsnow.com/blog/the-expanding-canvas-what-counts-as-art>
- https://kar.kent.ac.uk/105246/1/105ALICE_C_HELLIWELL_-THESIS-PHILOSOPHY_OF_AI_ART-_KAR_UPLOAD_REDACTED.pdf
- <https://lab.sorrywecan.com/static/pdf/the-coming-age-of-wisdom-work.pdf>
- <https://thehumankind.co/2025/07/02/everyone-suddenly-thinks-they-are-rick-rubin/>
- <https://conversationswithtyler.com/episodes/rick-rubin/>
- <https://ia800503.us.archive.org/33/items/the-creative-act-by-rick-rubin/The%20Creative%20Act%20By%20Rick%20Rubin.pdf>
- <https://tianpan.co/notes/2025-05-05-creativity-act>
- <https://bigblue.academy/en/the-death-of-prompt-engineering-and-its-ruthless-resurrection-navigating-ai-orchestration-in-2026-and-beyond>
- <https://infuzu.com/blog/why-prompt-engineering-shouldn-t-exist>
- <https://hrexecutive.com/is-prompt-engineering-dead-one-expert-describes-what-hr-should-focus-on-instead/>
- <https://liatbenzur.com/2025/11/03/the-prompt-collapse-why-natural-language-interfaces-are-about-to-eat-traditional-software/>
- <https://medium.com/@ranjio11/prompt-engineering-turning-intuition-into-intentional-skill-09cdd27ffda2>
- https://www.reddit.com/r/learnmachinelearning/comments/1oj5wxl/is_prompt_engineer
- <https://sougwen.com/putting-the-art-in-artificial-intelligence-a-conversation-with-sougwen-chung>

- <https://hubemag.com/sougwen-chung-the-poetics-of-human-machine-interaction>
- <https://medium.com/verisart/sougwen-chung-human-and-machine-collaboration-72d912d6b065>
- <https://www.theartnewspaper.com/2025/07/11/why-artificial-intelligence-artists-can-be-seen-as-builders-breakers-or-both-at-once>
- <https://time.com/collection/time100-ai/6309468/holly-herndon/>
- <https://news.artnet.com/art-world/holly-herndon-and-mat-dryhurst-2551072>
- <https://www.kw-berlin.de/en/exhibitions/holly-herndon-and-mat-dryhurst-starmirror>
- <https://www.lerandom.art/editorial/holly-herndon-mat-dryhurst-on-artificial-psychedelia>
- https://www.politesi.polimi.it/retrieve/bea25522-3b1c-4ad4-861c-e56af3fee554/2024_10_GiacominDaSilva.pdf
- https://eprints.nottingham.ac.uk/81162/1/Salimbeni_Guido_20276375_corrections.pdf
- <https://www.scribd.com/document/664543313/week-4-Analyzing-and-visualizing-data-Andy-Kirk>
- <https://elifesciences.org/reviewed-preprints/109313>
- <https://www.glozo.com/reports/marketing-salary-skills-report-2025>
- <https://himalayas.app/career-guides/art-director>
- <https://www.sweetstudy.com/files/milestone2-3486823.docx>
- <https://www.jobbers.io/the-global-freelance-rate-database-2026-complete-salary-pricing-guide-for-every-industry/>

*Appendix H — The Dream Machine Source
Index*

A thematic catalogue of significant sources surfaced across the 32 issues of *Dream Machine* (October 2025 – June 2026).

This index is a navigational tool, not an exhaustive list. The full *Dream Machine* archive contains nearly three thousand individual hyperlinks across its thirty-two issues, the great majority of which are primary-source links to industry coverage, research reports, official announcements, court filings, technical demos, creator showcases, and platform releases.

What follows below is the thematic catalogue of the *significant* sources — the ones the book itself draws on, the ones a working creative or researcher tracking a specific topic would want as a starting point, and the ones that, taken together, define the public record of creative AI as it stood in the period this book covers. Within each theme, entries are organised chronologically by issue number. The format is:

[Issue N] **Title** — short context — URL

For the complete primary-source archive — every link, in full, in original publication order — refer to the *Dream Machine* newsletter archive on LinkedIn or the per-issue markdown files in *Dream Machine* MD/. Each issue’s final section, “*All embedded URLs (in document order)*”, lists every URL the issue carried.

* * *

1. *AI Video — Models and Releases*

- [Issue 1] **Sora 2 Launch** — OpenAI’s step-change in physical realism, audio, multi-shot world state — <https://openai.com/index/sora-2/>
- [Issue 1] **Veo 3 and Flow** — Google DeepMind’s cinematic AI video — <https://www.youtube.com/watch?v=Io6Ef8alr2Y>
- [Issue 2] **Veo 3.1 Coming Soon** — improved consistency, resolution, multi-shot, cinematic presets — <https://www.cometapi.com/veo-3-1-is-comingand-whats-rumor/>
- [Issue 3] **Veo 3.1 Deep Dive with Flow** — cinematic filmmaking toolset — <https://www.youtube.com/watch?v=Io6Ef8alr2Y>
- [Issue 3] **Higgsfield Sketch-to-Video** — powered by Sora 2 — <https://higgsfield.ai/posts/6nkYSGwOcdyXVqZefE1MsQ>
- [Issue 3] **LiveGS** — mobile Gaussian-splatting video
- [Issue 4] **Veo 3.1 Style Tips** — text-to-video with style guidance — <https://x.com/GoogleAI/status/1980327604843381215>
- [Issue 4] **Veo Precision Features** — remove/add elements to scenes — <https://x.com/GoogleDeepMind/status/1980261047836508213>
- [Issue 4] **Heygen Identity Consistency with Veo 3.1** — character-consistent video — https://x.com/HeyGen_Official/status/1978491090618749193
- [Issue 4] **Higgsfield Popcorn** — storyboard tool with character consistency — https://x.com/higgsfield_ai/status/1981110992630341928
- [Issue 4] **Odyssey 2** — real-time interactive video generation — <https://odyssey.ml/introducing-odyssey-2>

- [Issue 4] **Krea Realtime** — open-sourcing the creative engine — https://www.linkedin.com/posts/krea-ai_today-were-open-sourcing-krea-realtime-activity-7386124532207689728-B7pD
- [Issue 5] **Sora Character Cameos** — new feature in Sora app — <https://x.com/OpenAI/status/1983661036533379486>
- [Issue 5] **VEED Transitions** — AI-powered video transitions — <https://x.com/veedstudio/status/1980636419891818850>
- [Issue 5] **LTX-2** — open-source audio-video generation — https://x.com/ltx_model/status/1981346235194683497
- [Issue 6] **Wan 2.2** — AI video with multi-shot capabilities — <https://x.com/eyishazyer/status/1983507594942755221>
- [Issue 6] **MotionStream** — real-time interactive video, 29 FPS — <https://x.com/wildmindai/status/1985828041566941576>
- [Issue 7] **Odyssey-2** — interactive streaming 16:9 video — <https://x.com/olivercameron/status/1984777003967672800>
- [Issue 9] **Wan 2.6 Released** — cast characters from reference videos, up to 15 seconds — <https://wan.video/blog/wan2.6-introduction>
- [Issue 13] **Runway Gen-4.5** — image-to-video for paid plans — <https://www.youtube.com/watch?v=AwKSrJFvdps>
- [Issue 13] **LTX-2 on ComfyUI** — open-source audio-video — <https://blog.comfy.org/p/ltx-2-open-source-audio-video-ai>
- [Issue 14] **Veo 3.1 Ingredients to Video** — vertical formats, 1080p/4K — <https://x.com/FlowbyGoogle/status/2011130097483526474>
- [Issue 14] **LTX-2 Lip Sync** — native audio-driven dialogue — https://x.com/ltx_model/status/2011101440706806051
- [Issue 15] **Runway Gen-4.5 Image to Video** — broad rollout — https://www.linkedin.com/posts/runwayml_introducing-image-to-video-for-gen-45-the-activity-7419856988186238976-ZJU2
- [Issue 15] **Veo 3.1 in YouTube Shorts and Create app** — distribution-layer integration — <https://blog.google/innovation-and-ai/technology/ai/veo-3-1-ingredients-to-video/>
- [Issue 31] **Runway Aleph 2.0** — single-frame edit propagated across the video clip — June 2026
- [Issue 31] **Alibaba WonderClip** — all-in-one AI video creation platform — June 2026
- [Issue 31] **Sony Vid-CamEdit** — post-capture camera trajectory editing keeping scenes spatially coherent across viewpoint changes — June 2026
- [Issue 31] **Nvidia PiD (Pixel Diffusion Decoder)** — new image/video decoder — June 2026

2. AI Image — Models and Tools

- [Issue 1] **Nano Banana Plugin for Photoshop** — Gemini image gen inside Adobe — https://www.linkedin.com/posts/arminas-valunas-b4477255_nano-banana-plugin-for-

photoshop-is-here-ugcPost-7367923639414906881-GgV7

- [Issue 1] **Magnific Precision v2** — AI upscaling — https://www.linkedin.com/posts/magnific-ai_introducing-magnific-precision-v2-activity-7387158930776276992-zpET
- [Issue 2] **Vimeo AI Creator Tools** — new features — <https://www.tvtechnology.com/news/vimeo-releases-new-ai-powered-creator-tools>
- [Issue 3] **Google Stitch** — design tool with AI features — <https://www.testingcatalog.com/google-test-new-stitch-modes-annotate-theme-interactive/>
- [Issue 5] **Adobe Firefly Image Model 5** — Adobe MAX 2025 — <https://news.adobe.com/news/2025/10/adobe-max-2025-firefly>
- [Issue 5] **Freepik Spaces** — infinite canvas for collaborative creation — <https://www.freepik.com/spaces>
- [Issue 5] **Midjourney updates** — state-of-the-art image gen — <https://x.com/midjourney/status/1991684484455100477>
- [Issue 6] **Qwen Image Multiple Angles LoRA** — consistent characters — <https://x.com/multimodalart/status/1986174924038218087>
- [Issue 8] **Qwen Image 2512** — latest generation — <https://qwen.ai/blog?id=qwen-image-2512>
- [Issue 9] **FLUX 2** — Black Forest Labs image generation — <https://bfl.ai/models/flux-2-max>
- [Issue 10] **FLUX 2 klein** — fast image generation under one second — <https://huggingface.co/unsloth/FLUX.2-klein-4B-GGUF>
- [Issue 13] **ChatGPT Images Upgrade** — major feature improvements — <https://www.techradar.com/ai-platforms-assistants/chatgpt/chatgpt-images-just-got-a-major-upgrade-and-it-could-change-how-we-all-create>

3. AI Music and Audio

- [Issue 1] **Suno Studio** — generative audio workstation — <https://www.techradar.com/ai-platforms-assistants/i-tried-suno-studio-the-new-platform-that-mixes-ai-music-generation-with-hands-on-editing-like-garageband-but-smarter>
- [Issue 1] **YouTube Music AI Hub** — AI music hosts — <https://www.linkedin.com/news/story/youtube-music-debuts-new-ai-hub-6625484/>
- [Issue 1] **Spotify AI Protections** — strengthened for artists — <https://newsroom.spotify.com/2025-09-25/spotify-strengthens-ai-protections/>
- [Issue 1] **Xania Monet** — AI singer signs \$3M deal — <https://www.forbes.com/sites/dougmelville/2025/09/27/ai-singer-xania-monet-just-charted-on-billboard-signed-3m-deal-is-this-the-future-of-music/>
- [Issue 1] **Cardiff Band on AI Artist** — trained on their music, outperforming them — <https://musictech.com/news/industry/its-shocking-disheartening-and-insulting-cardiff-band-speaks-out-after-ai-artist-trained-on-their-music-outperforms-them-on-spotify/>

- [Issue 2] **Suno Funding Round** — \$2 billion valuation — <https://www.digitalmusicnews.com/2025/10/20/suno-funding-round-october-2025/>
- [Issue 2] **Groundhog AI Guitar Pedal** — tone matching — <https://musictech.com/news/gear/groundhog-audio-onepedal-ai-tone-matching/>
- [Issue 2] **Spotify in ChatGPT** — integration launched — <https://newsroom.spotify.com/2025-10-06/spotify-personalized-prompts-chatgpt/>
- [Issue 3] **iZotope Ozone 12** — AI assistant for mixing — <https://musictech.com/reviews/plugin-ins/izotope-ozone-12-review/>
- [Issue 3] **Tempolor Guitars (Quwan)** — AI to make songs playable — <https://kr-asia.com/no-practice-required-quwans-tempolor-guitars-use-ai-to-make-songs-playable-in-minutes>
- [Issue 4] **Mureka Music Agent Studio** — six specialised AI agents — https://www.linkedin.com/posts/sherrihendrickson_mureka-unveils-music-agent-studio-and-enhanced-share-7384999251526864896-cNYg/
- [Issue 4] **Fish Audio S1** — TTS 6× cheaper than ElevenLabs — <https://fish.audio/app/text-to-speech/>
- [Issue 5] **Universal Music + Stability AI Alliance** — strategic partnership — <https://stability.ai/news/universal-music-group-and-stability-ai-announce-strategic-alliance>
- [Issue 5] **Udio Partners with UMG** — partnership announced — <https://www.udio.com/blog/a-new-era>
- [Issue 5] **OpenAI Music Generator** — reportedly in development — <https://www.theinformation.com/articles/openai-plots-generating-ai-music-potential-rivalry-startup-suno>
- [Issue 6] **UMG Boss Lucian Grainge on AI** — full internal memo — <https://musically.com/2025/10/14/umg-boss-sir-lucian-grainge-talks-ai-full-internal-memo/>
- [Issue 6] **Bleeding Verse AI Band** — Hallwood Media signing — <https://musically.com/2025/10/07/ai-band-bleeding-verses-creator-signs-deal-with-hallwood-media/>
- [Issue 6] **JYP Entertainment AI Artist** — hiring AI/Unreal experts — <https://www.musicbusinessworldwide.com/jyp-entertainment-is-hiring-for-ai-and-unreal-engine-experts-to-develop-an-unprecedented-virtual-kpop-artist/>
- [Issue 6] **Claimy** — \$1.8M for missing royalty recovery — <https://www.musicbusinessworldwide.com/ai-music-tech-startup-claimy-raises-1-8m-to-tackle-missing-royalty-payments/>
- [Issue 7] **Breaking Rust on Billboard** — AI country act — <https://www.npr.org/2025/11/10/nx-s1-5604320/breaking-rust-is-a-hot-new-country-act-on-the-billboard-charts-its-powered-by-ai>
- [Issue 7] **GEMA v OpenAI Munich ruling** — European copyright precedent — https://www.linkedin.com/posts/dr-barry-scannell-bbb5aa207_in-a-major-ruling-for-european-copyright-share-7393957246386323457-8bbx

- [Issue 8] **LANDR AI Music Study** — 87% of musicians use AI tools — <https://aristake.com/ai-tools-musicians-study/>
- [Issue 8] **Stability AI + Warner Music** — next-gen responsible AI tools — <https://stability.ai/news/warner-music-group-and-stability-ai-join-forces-to-build-next-gen-tools>
- [Issue 8] **Paul McCartney Silent Track Protest** — UK copyright opt-out — <https://www.theguardian.com/music/2025/nov/17/the-sound-of-silence-why-theres-barely-anything-there-in-paul-mccartney-new-release>
- [Issue 11] **Lyria Camera (Google DeepMind)** — music generation — <https://magenta.withgoogle.com/lyria-camera-announce>
- [Issue 11] **Dave Stewart on AI** — Musicians must embrace it — <https://www.theguardian.com/music/2025/dec/05/musicians-must-embrace-unstoppable-force-of-ai-eurythmics-dave-stewart-urges>
- [Issue 12] **Splice + UMG Collaboration** — AI music creation tools — <https://www.universalmusic.com/universal-music-group-and-splice-to-collaborate-on-the-next-generation-of-ai-powered-music-creation-tools-for-artists/>
- [Issue 14] **Bandcamp Bans AI Music** — platform policy — <https://stereogum.com/2485199/bandcamp-bans-ai-music/news>
- [Issue 14] **UMG slams AI slop** — exponential growth on streaming — <https://musically.com/2026/01/09/umg-boss-slams-exponential-growth-of-ai-slop-on-streaming-services/>
- [Issue 15] **Sienna Rose** — viral mystery AI singer (BBC investigation) — <https://www.bbc.co.uk/news/articles/cq6v83gq66eo>
- [Issue 16] **800 Creatives Sign Declaration** — *Stealing Our Work Is Not Innovation* — <https://www.digitalmusicnews.com/2026/01/22/stealing-isnt-innovation/>
- [Issue 31] **ElevenLabs Music v2** — next-generation generative music model — June 2026
- [Issue 31] **Stable Audio 3.0** — open-weight model family on fully licensed data, foundation for the open audio community — June 2026
- [Issue 31] **Apple Music: AI <1% of plays** — clean platform-level confirmation that audience-chosen listening time has remained almost entirely human-authored — June 2026
- [Issue 31] **Spotify + UMG: AI covers/remixes as Premium add-on** — landmark licensed AI-music deal — June 2026
- [Issue 31] **Daniel Ek defends AI music** — “better than slop” framing — June 2026
- [Issue 31] **Tonada (Sweden)** — AI music for retailers; B2B-aligned synthetic background music — June 2026
- [Issue 31] **Spotify NotebookLM-style Personalised Audio app** — competing with Google’s NotebookLM podcasts — June 2026
- [Issue 31] **Musician uses AI as Parkinson’s affects ability to play guitar** — accessibility case for AI in music creation — June 2026

4. AI 3D / World Models / Spatial

- [Issue 1] **Marble by World Labs** — first commercial world model — <https://marble.worldlabs.ai/>
- [Issue 1] **Meta Hyperscape Capture** — Gaussian splatting on Quest — <https://www.meta.com/en-gb/experiences/meta-horizon-hyperscape-capture-beta/8798130056953686/>
- [Issue 1] **PlayCanvas SOG** — WebP for 3D Gaussian Splatting — https://www.linkedin.com/posts/willeastcott_playcanvas-open-sources-sog-literally-webp-ugcPost-7374459362708180992-aHDa
- [Issue 3] **Genie 3** (Google DeepMind) — interactive 3D world generation — <https://deepmind.google/blog/genie-3-a-new-frontier-for-world-models/>
- [Issue 3] **World Labs RTFM** — real-time frame model — <https://www.worldlabs.ai/blog/rtfm>
- [Issue 3] **Instant Skinned Gaussian Avatars** — web/mobile VR — <https://sites.google.com/view/gaussian-vrm>
- [Issue 3] **Tencent Hunyuan World 1.1** — 3D reconstruction — <https://x.com/TencentHunyuan/status/1980930623536837013>
- [Issue 5] **Apple Personas use Gaussian Splatting** — most-deployed splat tech in consumer hardware — <https://radiancefields.com/apple-confirms-personas-use-gaussian-splatting>
- [Issue 7] **Marble formal public launch** — Sony Pictures using it (40× faster than legacy VP) — <https://www.worldlabs.ai/case-studies/bringing-marble-to-life>
- [Issue 8] **Meta SAM 3 / SAM 3D** — segment anything in 3D — https://www.linkedin.com/posts/aiatmeta_introducing-sam-3-sam-3d-ugcPost-7396944913751465985-m5Nc
- [Issue 11] **Meta WorldGen** — text-to-immersive-3D-worlds research — <https://www.facebook.com/LifeAtMeta/videos/research-update-worldgen-text-to-immersive-3d-worlds/1879077432692421/>
- [Issue 11] **Ubisoft CHORD Model** — open-sourced PBR material generation — <https://blog.comfy.org/p/ubisoft-open-sources-the-chord-model>
- [Issue 12] **Tencent HY World 1.5** — real-time world model framework — <https://x.com/TencentHunyuan/status/2001170499133653006>
- [Issue 12] **Microsoft Trellis 2** — 3D generation — <https://github.com/microsoft/TRELLIS.2>
- [Issue 13] **Wonderzoom** (Stanford AI Lab) — multi-scale 3D world generation — <https://wonderzoom.github.io/>
- [Issue 15] **WorldLabs API** — 3D world generation as a service — <https://x.com/theworldlabs/status/2014046372639408203>
- [Issue 17] **Project Genie Public Release** — Google AI Ultra subscribers — <https://blog.google/innovation-and-ai/models-and-research/google-deepmind/project-genie/>

- [Issue 22] **Luma UNI-1** — combined world generation + reasoning — *Dream Machine* Editor's Pick
- [Issue 25] **Spark 2.0** — open-source 100M-splat browser streaming
- [Issue 27] **Vista4D** (Netflix + Eyeline) — live action to navigable 4D point clouds
- [Issue 31] **TripoAI Project Eden** — persistent multiplayer world model decoupling world state from visual rendering — June 2026
- [Issue 31] **AI-Enhanced LiDAR** — generative-AI-augmented LiDAR capture pipeline — June 2026

5. Voice, Avatars, Digital Humans

- [Issue 1] **Tilly Norwood AI Actress** — sparks UK/US union debate — <https://www.hollywoodreporter.com/movies/movie-news/tilly-norwood-ai-actress-uk-union-equity-sag-aftra-debate-1236391739/>
- [Issue 1] **SAG-AFTRA Condemns Tilly Norwood** — official union statement — <https://variety.com/2025/film/news/sag-aftra-tilly-norwood-ai-actress-1236534779/>
- [Issue 1] **Deep Fusion Films + Topfoto** — AI documentary alliance — <https://www.abc.org/production/news/deep-fusion-and-topfoto-strike-alliance-to-produce-ai-powered-documentaries/22728>
- [Issue 2] **ROXi AI TV Presenters** — music channel hosts — <https://www.advanced-television.com/2025/10/07/roxi-debuts-ai-generated-tv-presenters/>
- [Issue 2] **Hedra Audio Tags** — emotional audio control — https://x.com/hedra_labs/status/1998490844748460528
- [Issue 3] **Heygen Motion Designer** — prompt-based animation — https://www.linkedin.com/posts/heygen_if-you-can-explain-it-you-can-animate-it-activity-7387488068204892160-AY8-
- [Issue 4] **Copresence and ConvAI** — Unreal-Engine intelligent avatars — https://www.linkedin.com/posts/copresence-tech_want-to-create-an-intelligent-avatar-that-activity-7379523290421383168-SA_g_
- [Issue 7] **Lumi Avatar** — real-time audio-driven avatars — <https://www.linkedin.com/feed/update/urn:li:activity:7393916271018467328/>
- [Issue 7] **Tilly Creator Eline Van der Velden — Deadline** — backlash and the next 40 — <https://deadline.com/2025/11/tilly-norwood-creator-interview-backlash-more-ai-actors-coming-1236601334/>
- [Issue 8] **Synthesia \$4B valuation** — rejected \$3B Adobe offer — <https://techcrunch.com/2026/01/26/synthesia-hits-4b-valuation-lets-employees-cash-in/>
- [Issue 11] **ElevenLabs Impact Program** — SXSW documentary on voice loss — https://www.linkedin.com/posts/elevenlabsio_at-sxsw-the-elevenlabs-impact-program-will-activity-7413979472988774400-fVRG
- [Issue 12] **Google Veo Avatars** — expressiveness upgrade — <https://vids.new/>
- [Issue 12] **Meta SAM Audio** — segment anything for audio — <https://about.fb.com/news/2025/12/our-new-sam-audio-model-transforms-audio->

editing/

- [Issue 13] **Avatar Forcing** — real-time interactive avatar generation — <https://taekyungki.github.io/AvatarForcing/>
- [Issue 13] **Qwen3 TTS** — voice design and cloning — <https://qwen.ai/blog?id=qwen3-tts-vc-voicedesign>
- [Issue 16] **Tilly Norwood Doubles Down** — AI as “more ethical” performance, urging actors to create avatars — <https://variety.com/2026/digital/news/tilly-norwood-creator-tells-actors-to-create-ai-avatars-1236638940/>
- [Issue 23] **Death threats against Eline Van der Velden** — cultural-extreme response
- [Issue 31] **Ozzy Osbourne AI avatar** — estate-led posthumous “live on as an AI avatar that talks with fans” — June 2026

6. Agent Platforms / Orchestration

- [Issue 2] **OpenAI AgentKit / DevDay** — agentic AI for creative workflows — <https://openai.com/index/introducing-agentkit/>
- [Issue 2] **Lenny AI Agent** — for live music event organisers — <https://musically.com/2025/10/20/meet-lenny-an-ai-agent-to-help-organisers-of-live-music-events/>
- [Issue 4] **AdsGency \$12M seed** — autonomous paid marketing — <https://www.finsmes.com/2025/10/adsgency-raises-12m-in-seed-funding.html>
- [Issue 5] **Meta + Hugging Face OpenEnv** — open-source agentic development — <https://www.edtechinnovationhub.com/news/meta-and-hugging-face-launch-openenv-to-advance-open-source-agentic-development>
- [Issue 5] **Pomelli** (Google Labs) — AI marketing agent for SMBs
- [Issue 5] **Opal** (Google) — no-code AI mini-app builder — <https://blog.google/technology/google-labs/opal-expansion/>
- [Issue 8] **Multimodal Agents in Unreal Engine** — live-score nature documentary — <https://www.youtube.com/watch?v=7u2yCtbONmo>
- [Issue 8] **SIMA 2** (Google DeepMind) — agent for virtual 3D worlds — <https://deepmind.google/blog/sima-2-an-agent-that-plays-reasons-and-learns-with-you-in-virtual-3d-worlds/>
- [Issue 8] **Google Antigravity** — agentic development platform
- [Issue 13] **NitroGen** (NVIDIA + Stanford) — plays-any-game model — <https://nitrogen.minedojo.org/>
- [Issue 14] **General Intuition** — \$134M for spatial-reasoning agents — <https://techcrunch.com/2025/10/16/general-intuition-lands-134m-seed-to-teach-agents-spatial-reasoning-using-video-game-clips/>
- [Issue 16] **Anthropic Claude Apps** — interactive Claude in workplace tools — <https://techcrunch.com/2026/01/26/anthropic-launches-interactive-claude-apps-including-slack-and-other-workplace-tools/>
- [Issue 16] **Heygen Video Agent** — full scripting-to-assembly — https://www.linkedin.com/posts/heygen_introducing-the-new-video-agent-activity-

7421597801240801282-d1CF

- [Issue 21] **Adobe + NVIDIA Strategic Partnership** — agentic creative intelligence
- [Issue 26] **Adobe Summit 2026** — “agentic creative intelligence” headline category
- [Issue 29] **Sony 49-Claude-agent / 72-skill stack** — game-dev multi-agent team

7. *Adobe and Creative Software*

- [Issue 1] **Tether** — AI animation in After Effects — https://www.linkedin.com/posts/thisisdoug_aftereffects-ai-video-vfx-ugcPost-7368671859774517249-losz
- [Issue 1] **Unreal Engine 5 AI Assistant** — official integration — https://www.linkedin.com/posts/wouterweynants_theres-an-official-ai-assistant-coming-to-ugcPost-7369377204226379776-pGiH
- [Issue 1] **Unity AI Council** — accelerate AI innovation — <https://www.gamedeveloper.com/business/unity-forms-ai-council-to-accelerate-ai-product-innovation->
- [Issue 1] **ComfyUI raises \$17M** — OS for creative AI — https://www.linkedin.com/posts/comfyui_we-raised-17-million-to-build-an-os-for-ugcPost-7373743341236236288-wkCc
- [Issue 5] **Adobe MAX 2025 Express AI Assistant** — Adobe MAX announcements — <https://news.adobe.com/news/2025/10/adobe-max-2025-express-ai-assistant>
- [Issue 5] **Adobe Firefly Foundry** — custom models for brands — <https://news.adobe.com/news/2025/10/adobe-max-2025-firefly-foundry>
- [Issue 5] **Adobe MAX Sneaks** — Light Touch, Surface Swap, Scene It, etc. — <https://www.youtube.com/watch?v=YqAAFx1XXY8>
- [Issue 5] **Adobe MAX Creator Survey** — 86% use creative gen AI — <https://news.adobe.com/news/2025/10/adobe-max-2025-creators-survey>
- [Issue 5] **Adobe Project Light Touch / Scene It / Surface Swap** — see Adobe Sneaks
- [Issue 5] **“Adobe is putting AI in everything everywhere all at once”** — Creative Boom coverage — <https://www.creativeboom.com/news/adobe-is-putting-ai-in-everything-everywhere-all-at-once/>
- [Issue 6] **Adobe Corrective AI** — voice-over emotion editing — <https://www.wired.com/story/adobe-max-sneaks-2025-corrective-ai/>
- [Issue 8] **Adobe Research RELIC** — interactive video world model
- [Issue 12] **Adobe inside ChatGPT** — Photoshop, Express, Acrobat editing in ChatGPT — <https://blog.adobe.com/en/publish/2025/12/10/edit-photoshop-chatgpt>
- [Issue 12] **Google Flow Refine** — prompt by doodling — <https://blog.google/technology/google-labs/flow-refine-videos/>
- [Issue 15] **Adobe at Sundance 2026** — \$10M grants, Ignite Day — <https://analyticsindiamag.com/ai-news-updates/adobe-unveils-ai-video-innovations-10-million-grants-ahead-of-sundance-film-festival/>
- [Issue 16] **Adobe Premiere Object Mask** — automated masking — https://www.linkedin.com/posts/robdewinter_ok-this-is-going-to-save-a-lot-of-time-in-

ugcPost-7421617551690063872-yKmB

- [Issue 28] **Unity AI Open Beta** — in-editor full AI suite

8. Games Industry — Adoption

- [Issue 1] **Netflix Director of GenAI for Games** — \$840K role — <https://www.pcgamer.com/gaming-industry/as-the-videogame-industry-continues-to-be-hammered-by-layoffs-netflix-is-offering-up-to-usd840-000-per-year-for-a-new-director-of-generative-ai-for-games/>
- [Issue 1] **Meta Horizon Studio AI Assistant Upgrade** — <https://www.uploadvr.com/meta-horizon-studio-upgrade-ai-assistant-horizon-worlds/>
- [Issue 1] **51% of Japanese games studios use AI** — research finding — <https://www.gamesindustry.biz/51-of-japanese-game-makers-use-generative-ai>
- [Issue 3] **Roblox AI Tools for Creators** — <https://www.gamesindustry.biz/roblox-announces-new-ai-tools-for-creators>
- [Issue 3] **Battlefield 6** — “very seducing” AI for early stages — <https://www.gamesradar.com/games/battlefield/battlefield-6-lead-calls-generative-ai-very-seducing-but-says-it-was-only-used-in-the-games-earliest-stages-to-allow-for-more-time-and-more-space-to-be-creative/>
- [Issue 3] **Promise (Google-backed AI studio)** — VFX for legacy media — <https://www.hollywoodreporter.com/business/digital/ai-studio-promise-vfx-generation-company-1236397636/>
- [Issue 4] **EA + Stability AI Partnership** — generative AI tools for games — <https://stability.ai/news/stability-ai-and-ea-partner-to-reimagine-game-development>
- [Issue 4] **Halo Studios** — AI woven into development — <https://thegamepost.com/insider-halo-studios-generative-ai-game-development/>
- [Issue 5] **NBCUniversal × Dick Wolf Jr** — AI games deal — <https://www.videogameschronicle.com/news/nbcuniversal-signs-deal-with-law-order-creator-dick-wolfs-son-to-make-ai-generated-games-based-on-its-ip/>
- [Issue 5] **Microsoft Gaming Copilot** — screenshots for in-game understanding — <https://www.tomshardware.com/video-games/pc-gaming/microsoft-says-gaming-copilot-uses-screenshots-to-understand-in-game-events-not-for-training-ai-models-optional-feature-can-be-turned-off-but-not-easily-uninstalled>
- [Issue 5] **Sony Jabali AI Platform** — game development — <https://variety.com/2025/gaming/news/sony-jabali-ai-ai-game-development-platform-1236566619/>
- [Issue 5] **Krafton AI-First Plans** — Subnautica owner — <https://www.gamedeveloper.com/business/subnautica-owner-krafton-outlines-plans-to-transform-into-an-ai-first-company>
- [Issue 6] **Todd Howard on AI at Bethesda** — toolset, not replacement — <https://www.pcgamer.com/gaming-industry/todd-howard-says-ai-cant-replace-human-creative-intention-but-its-part-of-bethesdas-toolset-for-how-we-build-our-worlds-or-check-things/>

- [Issue 6] **EA pushes 15K employees on AI** — thought partner — <https://www.gamesradar.com/games/even-under-usd20-million-in-debt-ea-reportedly-pushes-15-000-employees-to-use-ai-as-a-thought-partner-for-everything-from-character-art-to-playtesting/>
- [Issue 6] **Falcom AI** — 2-3 hours of work to 10 minutes — <https://www.eurogamer.net/falcom-is-the-latest-developer-to-buy-into-the-ai-hype-machine>
- [Issue 7] **Square Enix 70% AI QA target** — by end 2027 — <https://www.pcgamer.com/gaming-industry/square-enix-aims-to-have-ai-doing-70-percent-of-its-qa-work-by-the-end-of-2027/>
- [Issue 8] **Call of Duty: Black Ops 7 AI art accusations** — community backlash — <https://www.videogameschronicle.com/news/it-honestly-sucks-fans-think-call-of-duty-black-ops-7-is-filled-with-generative-ai-art/>
- [Issue 8] **Ubisoft Anno 117 — AI art placeholder** — slipped through review — <https://www.videogameschronicle.com/news/ubisoft-says-ai-generated-art-in-anno-117-was-a-placeholder-which-slipped-through-our-review-process/>
- [Issue 13] **Razer \$600M AI focus** — strategic investment — <https://www.pymnts.com/news/artificial-intelligence/2026/razer-spends-600-million-dollars-sharpen-focus-ai-gaming/>
- [Issue 15] **Ubisoft cancels Prince of Persia + four** — AI refocus — <https://metro.co.uk/2026/01/21/prince-persia-remake-five-games-cancelled-ubisoft-focuses-ai-26431926/>
- [Issue 29] **Sony all-in on AI for games** — 49-agent / 72-skill stack

9. Games Industry — Refusal / Position

- [Issue 1] **Charles Cecil** — “AI was an expensive mistake” — <https://www.gamesindustry.biz/ai-was-an-expensive-mistake-charles-cecil-on-innovation-insolvency-and-broken-sword>
- [Issue 4] **Pocketpair (Palworld)** — publishing won’t take AI games — <https://www.pcgamer.com/software/ai/palworld-studio-pocketpair-says-its-new-publishing-division-wont-handle-games-that-use-generative-ai-we-dont-believe-in-it/>
- [Issue 9] **Witcher 3 / Cyberpunk Director** — AI helps, doesn’t replace — <https://www.gamesindustry.biz/witcher-3-and-cyberpunk-2077-director-says-ai-can-help-but-not-replace-creatives>
- [Issue 11] **Aardman on AI** — embrace, but cautious — <https://www.gamesradar.com/entertainment/animation-movies/wallace-and-gromit-creator-says-beloved-animation-studio-aardman-will-embrace-the-technology-of-ai-but-will-be-very-cautious-not-to-lose-our-values/>
- [Issue 14] **Larian backs off gen AI** — Divinity statement — <https://nichegamer.com/larian-studios-backs-off-from-gen-ai/>
- [Issue 14] **Games Workshop rules out gen AI** — Warhammer 40K — <https://decrypt.co/354482/warhammer-40000-maker-games-workshop-rules-out>

generative-ai

- [Issue 14] **Hooded Horse won't work with AI devs** — Manor Lords publisher — <https://nichegamer.com/manor-lords-publisher-hooded-horse-wont-work-with-devs-using-gen-ai/>
- [Issue 16] **Jagex never AI** — RuneScape commitment — <https://www.gamesindustry.biz/runescape-maker-jagex-says-it-will-never-use-generative-ai-to-make-in-game-content>

10. Film Industry — Studios and Positions

- [Issue 1] **Lionsgate AI failure** — Futurism report — <https://futurism.com/artificial-intelligence/lionsgate-movies-ai>
- [Issue 2] **Fremantle's Imaginae AI Studios** — CEO named — <https://www.hollywoodreporter.com/business/digital/fremantle-names-ceo-new-ai-label-imaginae-studios-1236396579/>
- [Issue 2] **Goldfinch enGEN3** — AI cinematic universe — <https://variety.com/2025/film/news/ai-powered-cinematic-universe-platform-engen3-1236543349/>
- [Issue 3] **Fox Entertainment + Holywater** — AI microdramas — <https://www.hollywoodreporter.com/business/business-news/fox-entertainment-invests-in-holywater-ai-microdramas-1236396802/>
- [Issue 4] **Netflix “all in” on AI** — Sarandos at industry conference — <https://www.cnn.com/2025/10/22/netflix-all-in-on-leveraging-ai-in-its-streaming-platform.html>
- [Issue 4] **Asteria's “All Heart”** — Natasha Lyonne short — <https://www.hollywoodreporter.com/movies/movie-news/natasha-lyonne-ai-company-asteria-1236403144/>
- [Issue 5] **Wonder Studios \$9M raise** — AI-native studio — <https://www.uktech.news/ai/ai-film-studio-wonder-lands-9m-investment-20251023>
- [Issue 5] **Watch the Skies (Swedish AI feature dubbed)** — USA distribution — <https://variety.com/2025/film/news/watch-the-skies-us-theatrical-release-ai-dubbing-1236343110/>
- [Issue 5] **Run to the West (South Korean first AI feature)** — <https://cybernews.com/entertainment/korean-cinema-run-to-the-west-ai/>
- [Issue 6] **Obsidian + Imagine Entertainment** — Ron Howard, Brian Grazer — <https://www.indiewire.com/news/business/obsidian-studio-ai-production-company-imagine-entertainment-1235158619/>
- [Issue 7] **Beta Films Chapter41 launch** — Munich AI startup — <https://deadline.com/2025/11/beta-film-ai-startup-chapter41-artificial-intelligence-1236612632/>
- [Issue 7] **House of David 350+ AI shots** — Wired feature — <https://www.wired.com/story/amazons-house-of-david-used-over-350-ai-shots-in-season-2-its-creator-isnt-sorry/>

- [Issue 7] **Wanted director's AI Method Actors** — Bekmambetov \$5M — <https://variety.com/2025/film/news/wanted-director-method-acting-ai-actors-1236579647/>
- [Issue 7] **Kevin Reilly + Kartel** — HBO veteran's AI startup — <https://www.hollywoodreporter.com/business/digital/kevin-reilly-ceo-kartel-ai-hbo-1236424692/>
- [Issue 8] **Humans in the Loop — Oscar race** — Sloan grant — <https://variety.com/2025/film/news/ai-drama-humans-in-the-loop-oscar-race-1236582975/>
- [Issue 8] **Synthetic Sincerity — IDFA** — Marc Isaacs — <https://www.hollywoodreporter.com/movies/movie-news/synthetic-sincerity-film-idfa-ai-authenticity-interview-1236426180/>
- [Issue 8] **AI Images Threaten Documentary** — Variety — <https://variety.com/2025/film/festivals/ai-generated-images-threaten-future-of-documentary-1236583466/>
- [Issue 11] **Disney \$1bn OpenAI investment** — Sora characters — <https://www.theguardian.com/business/2025/dec/11/disney-open-ai-sora-video-deal>
- [Issue 14] **Tunisian filmmaker wins \$1M for Lily** — Dubai AI Award — <https://www.broadcastprome.com/news/tunisian-filmmaker-wins-1-million-ai-film-award-for-lily/>
- [Issue 14] **Disney TikTok-like vertical video AI** — brand-asset video gen — <https://www.marketingdive.com/news/disney-unveils-tiktok-like-vertical-video-ai-video-generation-tool/809269/>
- [Issue 15] **Netflix retention AI strategy** — Pymnts — <https://www.pymnts.com/subscription-commerce/2026/retention-is-name-of-the-game-for-netflix-ai-strategy/>
- [Issue 16] **Andrii Daniels bomb-shelter Christmas clip** — viral Ukrainian AI film — <https://variety.com/2026/digital/news/ai-video-deadpool-harry-potter-andrii-daniels-1236624632/>
- [Issue 16] **Tilly Norwood Doubles Down** — Variety — <https://variety.com/2026/digital/news/tilly-norwood-creator-tells-actors-to-create-ai-avatars-1236638940/>
- [Issue 16] **Chris Pratt rejects AI villain** — Mercy pitch — <https://variety.com/2026/film/news/chris-pratt-ai-actor-villain-mercy-amazon-mgm-1236640460/>

11. *Celebrity / Director Positions on AI*

- [Issue 1] **James Cameron** — “never going to take place” — No Film School — <https://nofilmschool.com/james-cameron-ai#>
- [Issue 1] **Taylor Swift** — criticized for AI in promo — <https://tribune.com.pk/story/2570725/taylor-swift-criticised-for-using-ai-in-the-life-of-a-showgirl-promotional-campaign>

- [Issue 5] **Guillermo del Toro — “Rather Die”** — Variety — <https://variety.com/2025/film/news/guillermo-del-toro-rather-die-generative-ai-frankenstein-1236561316/>
- [Issue 5] **Paul Schrader on AI** — perfect script — <https://www.hollywoodreporter.com/movies/movie-news/paul-schrader-first-ai-movie-1236409606/>
- [Issue 7] **Jeremy Renner lawsuit threat** — multi-millions over AI voice — <https://deadline.com/2025/11/jeremy-renner-lawsuit-threat-ai-movie-1236611830/>
- [Issue 7] **George Clooney on AI actors** — Variety column — <https://variety.com/2025/scene/columns/george-clooney-ai-actors-movie-stars-1236579661/>
- [Issue 10] **James Cameron “horrifying”** — The Guardian — <https://www.theguardian.com/film/2025/dec/01/james-cameron-says-ai-actors-are-horrifying-to-me>
- [Issue 10] **Jenna Ortega “very easy to be terrified”** — NME — <https://www.nme.com/news/jenna-ortega-says-its-very-easy-to-be-terrified-of-ai-in-filmmaking-3913926>
- [Issue 11] **Leonardo DiCaprio — AI can’t be art** — THR — <https://www.hollywoodreporter.com/movies/movie-news/leonardo-dicaprio-ai-cant-be-art-no-humanity-1236445405/>
- [Issue 11] **James Cameron rejects AI actors at Hainan Festival** — <https://variety.com/2025/film/news/james-cameron-rejects-ai-actors-hainan-wouldnt-do-it-1236604204/>
- [Issue 14] **Claire Foy no interest in AI films** — Daily Mail — <https://www.dailymail.co.uk/tvshowbiz/article-15454199/Claire-Foy-AI-films-sad-disappointed-people-future-Hollywood.html>
- [Issue 14] **Wu-Tang Clan RZA** — case for AI in film/music — <https://www.vice.com/en/article/wu-tang-clans-rza-makes-the-case-for-ai-in-film-and-music-an-amazing-thing-for-us/>
- [Issue 15] **Matthew McConaughey protects voice/image** — Lawyer Monthly — <https://www.lawyer-monthly.com/2026/01/matthew-mcconaughey-protects-voice-image-ai/>
- [Issue 15] **Mara Wilson deepfake apocalypse fear** — Deadline — <https://deadline.com/2026/01/matilda-mara-wilson-stranger-things-ai-deepfake-apocalypse-1236689474/>

12. Music Industry — Labels, Deals, and Lawsuits

- [Issue 1] **Universal & Warner — landmark AI deals within weeks** — Musically — <https://musically.com/2025/10/02/report-umg-and-wmg-could-sign-landmark-ai-deals-within-weeks/>
- [Issue 5] **Universal Music + Stability AI Alliance** — *op. cit.*
- [Issue 5] **Udio + UMG Partnership** — *op. cit.*

- [Issue 7] **GEMA v OpenAI Munich ruling** — *op. cit.*
- [Issue 7] **“Biggest theft in music history”** — Rights group sues Suno — <https://edm.com/gear-tech/rights-group-sues-suno-copyright-infringement/>
- [Issue 7] **Bangkok Post: Xania Monet \$3M deal** — <https://www.bangkokpost.com/life/tech/3142355/ai-singer-xania-monet-signs-3m-deal-with-hollywood-media>
- [Issue 8] **Warner Music + Stability AI** — next-gen tools — <https://stability.ai/news/warner-music-group-and-stability-ai-join-forces-to-build-next-gen-tools>
- [Issue 9] **Johnny Cash estate sues Coca-Cola** — ELVIS Act — <https://completemusicupdate.com/johnny-cash-estate-uses-elvis-act-to-sue-coke-over-tribute-act-ad-soundtrack/>
- [Issue 12] **Splice + UMG Collaboration** — *op. cit.*
- [Issue 14] **UMG slams AI slop** — *op. cit.*
- [Issue 16] **Wixen \$50M lawsuit against Meta** — <https://www.musicbusinessworldwide.com/wixen-files-50m-copyright-suit-against-meta-claims-tech-giant-wants-to-replace-songwriters-with-ai/>
- [Issue 17] **UMG \$3B suit against Anthropic** — *Dream Machine* coverage
- [Issue 31] **Spotify + Universal Music Group landmark AI covers/remixes deal** — Premium add-on; the first major streamer × major-label licensed AI-music product — June 2026
- [Deep Dive №14 — 4 Jun 2026] **Spotify–UMG three Cs framework** — *consent / credit / compensation* contract architecture; *Verified by Spotify* badge; opt-in catalogue; new royalty stream; deal’s prehistory in UMG settlement with Udio (Oct 2025) — <https://newsroom.spotify.com/2026-05-21/universal-music-group-spotify-licensing-agreements-fan-made-covers-remixes/>
- [Deep Dive №14 — 4 Jun 2026] **Hasbro Sixth Wall + CharacterOS — Behavioural Licensing** — 12 characters; runtime governance layer; talent-participation pay-out; 13+/enterprise focus — <https://www.businesswire.com/news/home/20260603297922/en/>
- [Deep Dive №14 — 4 Jun 2026] **Disney–OpenAI deal — corrected record** — announced 11 Dec 2025, never executed; no formal agreement signed, no money changed hands; Sora wind-down announced 24 Mar 2026; framework survived platform’s death as reference contract architecture — <https://www.ropesgray.com/en/insights/viewpoints/the-deal-that-wasnt-what-disney-and-openai-teach-us-about-strategic-investments>
- [Deep Dive №14 — 4 Jun 2026] **Grimes / Elf.tech (GrimesAI)** — ~50% master-royalty split with creators; ~15,000 tracks generated; commons-with-royalty-hook model — <https://killthedj.com/how-artists-lend-their-ai-voices-to-creators-fans/>
- [Deep Dive №14 — 4 Jun 2026] **Holly Herndon / Holly+ DAO** — collective governance for commercial-use decisions; *Have I Been Trained?* (Spawning) consent-layer tooling
- [Deep Dive №14 — 4 Jun 2026] **ElevenLabs heritage voices (with CMG Worldwide)** — Judy Garland, James Dean, Burt Reynolds, Laurence Olivier; Iconic Marketplace

expansion to Stan Lee estate and Michael Caine — <https://www.cbsnews.com/news/ai-voices-dead-celebrities-judy-garland-james-dean-burt-reynolds/>

- [Deep Dive №14 – 4 Jun 2026] **Meta celebrity chatbot anti-pattern (2023)** — Snoopy Dogg, Kendall Jenner, MrBeast and others reportedly paid up to ~\$5m each for closed, platform-controlled personas — <https://english.elpais.com/technology/2023-11-13/from-paris-hilton-to-snoop-dogg-meta-to-launch-speaking-ai-chatbots-played-by-celebrities.html>
- [Deep Dive №14 – 4 Jun 2026] **Suno + Amazon Alexa+ integration** — *create-to-consume* demand floor; on-the-fly personalised songs via voice request — <https://www.billboard.com/pro/amazon-partners-ai-music-company-suno-alexa/>

13. Copyright, Policy and Regulation

- [Issue 9] **EU Lawmakers minimum age for AI/social** — Reuters — <https://www.reuters.com/legal/litigation/european-lawmakers-seek-eu-wide-minimum-age-access-ai-chatbots-social-media-2025-11-26/>
- [Issue 12] **UK DSIT Statement of Progress on Copyright and AI** — 88% — <https://www.gov.uk/government/publications/copyright-and-artificial-intelligence-progress-report/copyright-and-artificial-intelligence-statement-of-progress-under-section-137-data-use-and-access-act>
- [Issue 12] **IPWatchdog UK consultation analysis** — <https://ipwatchdog.com/2025/12/16/respondents-uk-ai-consultation-overwhelmingly-want-ai-companies-license-copyrighted-works-all-cases/>
- [Issue 11] **NY AI Advertising Disclosure Law** — Verge — <https://www.theverge.com/news/842848/new-york-law-ai-advertisements-sag-aftra-labor>
- [Issue 14] **Grok app ban consideration** — Fast Company — <https://www.fastcompany.com/91474131/governments-around-the-world-are-considering-bans-on-groks-app-over-ai-sexual-image-scandal>
- [Issue 21] **UK DSIT Final Copyright Report** — walked-back position
- [Issue 28] **Academy “You must be human to win” rule** — 2026 awards
- [Issue 29] **Cannes AI Disclosure Standard** — industry coordination
- [Issue 32] **Lisa Nandy: Government “got it wrong” on AI and creative copyright** — Culture Secretary’s admission at SXSW London (1 June 2026) that the opt-out model was a mistake; no decisions on next steps — <https://www.screendaily.com/news/we-got-it-wrong-says-uk-culture-secretary-lisa-nandy-of-government-misstep-on-ai-and-copyright-strategy/5217330.article>

14. Unions and Labour

- [Issue 1] **SAG-AFTRA Condemns Tilly Norwood** — *op. cit.*
- [Issue 1] **UK Equity Statement** — Hollywood Reporter — <https://www.hollywoodreporter.com/movies/movie-news/tilly-norwood-ai-actress-uk-union-equity-sag-aftra-debate-1236391739/>

- [Issue 11] **Equity 99% Strike Vote** — landslide for industrial action — <https://www.equity.org.uk/news/2025/performers-prepared-to-take-industrial-action-over-ai-in-landslide-99-vote>
- [Issue 11] **NY AI Advertising Disclosure Law / SAG-AFTRA quote** — *op. cit.*
- [Issue 15] **Equity welcomes improved offer** — film/TV AI protections — <https://www.equity.org.uk/news/2026/equity-welcomes-improved-offer-in-ai-protection-negotiations-in-film-and-tv>
- [Issue 26 / 29] **SAG-AFTRA “Tilly Tax” contract provisions** — final spring 2026 contract

15. Audience Response and Slop Ceiling

- [Issue 7] **Deezer/Ipsos AI Music Survey** — 97% can’t tell, but care when told — <https://newsroom-deezer.com/2025/11/deezer-ipsos-survey-ai-music/>
- [Issue 7] **50,000 AI tracks uploaded to Deezer daily** — Musically — <https://musically.com/2025/11/12/50000-ai-music-tracks-are-now-uploaded-to-deezer-every-day/>
- [Issue 8] **MrBeast on AI** — threat to creators — <https://www.forbes.com/sites/johnbbrandon/2025/10/10/mrbeast-is-right-about-ai-content-but-are-we-really-in-scary-times/>
- [Issue 12] **Merriam-Webster Word of the Year: “Slop”** — Hollywood Reporter — <https://www.hollywoodreporter.com/news/general-news/slop-word-year-2025-merriam-webster-1236450780/>
- [Issue 12] **YouTube AI channels — 1.2bn views fake politics** — Guardian — <https://www.theguardian.com/technology/2025/dec/13/fake-anti-labour-video-billion-views-youtube-2025>
- [Issue 14] **Spotify AI flood — subscribers furious** — TechRadar — <https://www.techradar.com/audio/spotify/ai-music-is-flooding-spotify-and-subscribers-are-furious-heres-why-music-fans-no-longer-trust-discover-weekly>
- [Issue 14] **Soultracks: “AI music is catchy, familiar... and boring”** — <https://soultracks.com/news-ai-generated-music-is-catchy-boring/>
- [Issue 15] **Sweden bans AI from official chart** — Independent — <https://www.independent.co.uk/tv/news/ai-music-song-banned-sweden-spotify-b2901627.html>
- [Issue 16] **YouTube CEO: managing AI slop on priority list 2026** — Digital Music News — <https://www.digitalmusicnews.com/2026/01/22/youtube-ceo-ai-slop-2026-comments/>
- [Issue 16] **Bain & Co — “People still want the radio star”** — <https://www.bain.com/insights/in-an-ai-age-people-still-want-the-radio-star/>
- [Issue 25-28] **Deezer April 2026 data — 44% / 3%** — newsroom release — <https://newsroom-deezer.com/2026/04/ai-generated-tracks-represent-44-of-new-uploaded-music/>

- [Issue 31] **Apple Music: AI <1% of plays** — clean platform-level confirmation — June 2026
- [Issue 31] **Daniel Ek defends AI music: “better than slop”** — Spotify CEO framing — June 2026
- [Issue 31] **Epidemic Sound *Future of the Creator Economy Report 2026*** — 94% of creators use AI; 89% feel pressure to use AI; 75% believe human-created content will become premium; survey of 3,000 professional creators — June 2026

16. Provenance, Watermarking, Detection

- [Issue 2] **OpenAI likeness protections** — Digital Music News — <https://www.digitalmusicnews.com/2025/10/08/openais-likeness-protections-dont-apply-to-dead-celebrities/>
- [Issue 11] **SynthID rollout across Veo / Lyria / Imagen** — Google DeepMind
- [Issue 12] **Gemini “Is this AI?” video verification** — https://www.linkedin.com/posts/googledeepmind_verify-google-ai-generated-videos-in-the-activity-7407748300688478208-fJgW
- [Issue 13] **Instagram chief — “fingerprint real media”** — Digital Music News — <https://www.digitalmusicnews.com/2026/01/05/instagram-chief-ai-slop-comments/>
- [Issue 13] **Instagram head AI verification** — WebProNews — <https://www.webpronews.com/instagram-head-warns-ai-images-erode-trust-calls-for-verification-standards/>
- [Issue 18] **Deezer licenses its AI-music detection tool** — *Dream Machine* coverage
- [Issue 29] **YouTube false-positive: Tiny Grandma stop-motion** — flagged as AI
- [Issue 31] **YouTube auto-labels AI videos** — platform-default disclosure regime — June 2026
- [Issue 31] **LinkedIn limits reach of AI-generated content** — distribution-side enforcement — June 2026
- [Issue 31] **TikTok AI moderation false positives** — wrongly punished independent filmmakers — June 2026

17. Workplace AI Adoption / Shadow AI / Workforce Research

- [Issue 1] **MIT study — AI reduces brain activity** — AI News — <https://www.artificialintelligence-news.com/news/ai-causes-reduction-in-users-brain-activity-mit/>
- [Issue 1] **51% of Japanese games studios use AI** — *op. cit.*
- [Issue 1] **Yale on AI adoption** — Neil Hoyne — https://www.linkedin.com/posts/neilhoyne_ai-data-research-activity-7379272781798035456-hnuV
- [Issue 5] **Azumo AI in Workplace Statistics 2025** — <https://azumo.com/artificial-intelligence/ai-insights/ai-in-workplace-statistics>
- [Issue 5] **Tech.co — Gen Z most likely use AI behind boss’s back** — <https://tech.co/news/gen-z-most-likely-use-ai-boss>

- [Issue 5] **IDC Europe Shadow AI security nightmare** — <https://blog-idceurope.com/shadow-ai-how-stealth-productivity-is-strangling-enterprise-ai-adoption-and-creating-a-security-nightmare/>
- [Issue 5] **Forbes — AI Tools Flood Workplaces** — <https://www.forbes.com/sites/carolinecastrillon/2025/09/09/ai-tools-flood-workplaces-as-employees-face-a-double-bind/>
- [Issue 5] **Exploding Topics AI Workforce Research** — <https://explodingtopics.com/blog/ai-workforce-research>
- [Issue 6] **Adobe Creators' Toolkit Report (16,000 creators)** — <https://news.adobe.com/news/2025/10/adobe-max-2025-creators-survey>
- [Issue 7] **CNBC — ADHD, autism, dyslexia and AI agents** — <https://www.cnn.com/2025/11/08/adhd-autism-dyslexia-jobs-careers-ai-agents-success.html>
- [Issue 8] **Forbes Vibe Coding \$220K** — <https://www.forbes.com/sites/rachelwells/2025/11/06/the-in-demand-ai-skill-and-certifications-that-pays-up-to-220000/>
- [Issue 8] **LANDR — 87% of musicians use AI tools** — <https://aristake.com/ai-tools-musicians-study/>
- [Issue 9] **Economist — Investors expect AI use to soar (it isn't)** — <https://www.economist.com/finance-and-economics/2025/11/26/investors-expect-ai-use-to-soar-thats-not-happening>
- [Issue 9] **Reuters Institute UK Journalists AI Survey** — <https://reutersinstitute.politics.ox.ac.uk/ai-adoption-uk-journalists-and-their-newsrooms-surveying-applications-approaches-and-attitudes>
- [Issue 12] **Economist — Job apocalypse? Humbug!** — <https://www.economist.com/business/2025/12/14/job-apocalypse-humbug-ai-is-creating-brand-new-occupations>
- [Issue 16] **Guardian — AI is hitting UK harder** — <https://www.theguardian.com/technology/2026/jan/26/ai-uk-jobs-us-japan-germany-australia>
- [Issue 16] **McKinsey AI for film and TV** — <https://www.mckinsey.com/industries/technology-media-and-telecommunications/our-insights/what-ai-could-mean-for-film-and-tv-production-and-the-industrys-future>
- [Issue 16] **PRS for Music AI Survey 2026** — <https://www.prsformusic.com/m-magazine/news/prs-for-music-ai-survey-2026>
- [Issue 24] **OpenAI public-policy on disruption** — robot tax, 4-day workweek, wealth funds
- [Issue 31] **Epidemic Sound *Future of the Creator Economy Report 2026*** — 94% of creators use AI; 89% feel pressure; 75% premium-human-content thesis — *Musically* coverage — June 2026
- [Issue 31] **Coordination Collapse research report (O'Hare / DreamLab AI Consulting)** — April 2026; preserved as Appendix K

- [Issue 31] **The Process Trap** — generative abundance and the equivocation of the creative; preserved as Appendix J
- [Issue 31] **The Doomer Mistake** — empirical rebuttal of the AI-driven civilizational-collapse case; preserved as Appendix I

18. Advertising and Brand AI

- [Issue 3] **WPP \$400M Google partnership** — <https://martechseries.com/predictive-ai/ai-platforms-machine-learning/google-and-spotify-alum-launch-epiminds-with-6-6m-to-build-marketing-teams-for-the-ai-era/>
- [Issue 5] **WPP Open Pro launch** — <https://campaignbrief.com/wpp-launches-ai-powered-marketing-platform-wpp-open-pro/>
- [Issue 5] **Mondelez AI for TV ads** — Verge — <https://www.theverge.com/news/806047/mondelez-ai-generated-ads>
- [Issue 6] **WPP + Sightly partnership** — Digiday — <https://digiday.com/media-buying/agencies-continue-to-expand-ai-capabilities-to-boost-brand-performance/>
- [Issue 6] **Coca-Cola AI Holiday ad (2nd attempt)** — Adweek — <https://www.adweek.com/creativity/coca-cola-uses-ai-to-rekindle-the-magic-of-its-holiday-ads/>
- [Issue 7] **Digiday — AI agent developers in-demand role** — <https://digiday.com/marketing/ai-agent-developers-have-become-adlands-in-demand-role/>
- [Issue 10] **Valentino “disturbing” AI handbag ads** — BBC — <https://www.bbc.co.uk/news/articles/cwyvjyvn83go>
- [Issue 11] **McDonald’s NL Christmas AI Ad pulled** — Branding in Asia — <https://www.brandinginasia.com/its-the-most-terrible-time-of-the-year-mcdonalds-netherlands-wonderfully-chaotic-ai-driven-christmas-film/>
- [Issue 11] **Channel 4 AI Ads** — Estate Agent Today — <https://www.estateagenttoday.co.uk/breaking-news/2025/12/homebuilder-among-first-to-use-channel-4s-ai-ads/>
- [Issue 14] **Marketing Week: AI ads winning in testing** — <https://www.marketingweek.com/dismiss-ai-ads-winning-creative-effectiveness/>
- [Issue 15] **PGA Tour + AWS expanded partnership** — Pymnts — <https://www.pymnts.com/artificial-intelligence-2/2026/ai-content-is-par-for-the-course-with-pga-tours-expanded-aws-partnership/>
- [Issue 15] **Avocados from Mexico skip TV for AI** — Digiday — <https://digiday.com/marketing/avocados-from-mexico-turns-to-ai-to-advertise-around-the-super-bowl-instead-of-a-tv-buy/>
- [Issue 16] **Higgsfield + Madonna AI video** — Adweek — <https://www.adweek.com/media/higgsfield-ai-marketing-startup/>
- [Issue 27] **WPP + Google Earth AI consumer journey** — *Dream Machine* coverage

19. Festivals, Institutions, Awards

- [Issue 6] **AI FilmFest Japan / Hoyt Dwyer** — <https://www.prnewswire.com/news-releases/from-apple-tv-creative-to-ai-filmmaker-hoyt-dwyers-animated-film-to-compete-at-ai-filmfest-japan-2025-302598064.html>
- [Issue 6] **India's first AI Film Festival** — IFFI, NFDC, LTIMindtree — <https://www.medianews4u.com/iffi-partners-with-ltimindtree-and-nfdc-to-launch-indias-first-ai-film-festival-and-hackathon/>
- [Issue 14] **Tunisian filmmaker wins \$1M for Lily** — *op. cit.*
- [Issue 14] **Comic-Con Art Show allows AI** — Filmstories — <https://filmstories.co.uk/news/san-diego-comic-con-art-show-to-allow-ai-slop/>
- [Issue 14] **Emmys AI Guidance** — THR — <https://www.hollywoodreporter.com/tv/tv-news/emmys-ai-guidelines-2026-awards-1236468434/>
- [Issue 15] **Sundance AI Literacy Initiative** — Sundance Institute blog — <https://www.sundance.org/blogs/centering-the-artist-why-were-launching-the-ai-literacy-initiative/>
- [Issue 15] **Google \$2M Sundance AI Education** — <https://blog.google/company-news/outreach-and-initiatives/google-org/sundance-institute-ai-education/>
- [Issue 16] **CNET — San Diego Comic-Con bans AI art at 2026 event** — <https://www.cnet.com/culture/san-diego-comic-con-bans-ai-art-for-2026-event/>
- [Issue 16] **Adobe Sundance Film Festival 2026** — <https://blog.adobe.com/en/publish/2026/01/20/sundance-film-festival-2026-creativity-community-power-storytelling>
- [Issue 28] **Academy “human to win” rule** — *Dream Machine* coverage
- [Issue 29] **Cannes AI Disclosure Standard launched** — *Dream Machine* coverage

20. Geographic / Regional AI Production

- [Issue 5] **CNBC Africa on AI in African music** — <https://www.cnbc.com/africa/2025/how-ai-is-changing-the-landscape-of-the-music-industry-in-africa>
- [Issue 5] **Run to the West — South Korea's first AI feature** — *op. cit.*
- [Issue 5] **Watch the Skies — Swedish AI dubbing** — *op. cit.*
- [Issue 6] **India's first AI Film Festival** — *op. cit.*
- [Issue 12] **Trilok — Indian AI band — Musically** — <https://musically.com/2025/12/17/indian-ai-band-trilok-performs-live-government-denies-association/>
- [Issue 13] **56.9% of new Chinese independent songs are AI** — Musically — <https://musically.com/2026/01/05/report-56-9-of-new-independent-songs-in-china-are-ai-generated/>
- [Issue 14] **BBC Future — Lights, camera, algorithm (India)** — <https://www.bbc.co.uk/future/article/20251223-why-indian-cinema-is-awash-with-ai>

- [Issue 14] **Shift Up CEO on AI vs China/US scale** — <https://www.pocketgamer.biz/shift-up-ceo-says-ai-is-key-to-competing-with-chinas-game-industry-scale/>
- [Issue 25] **Indonesia's Legenda Bertuah** — first AI-animated series
- [Issue 27] **Korin AI — Africa-trained, Africa-built** — launch
- [Issue 27] **Latin American AI film festival wave** — *Dream Machine* coverage
- [Issue 31] **Astana AI Film Festival (\$1M prize fund)** — Kazakhstan AI film festival opens submissions — June 2026
- [Issue 31] **Tribeca: first fully AI-generated film** — set to premiere — June 2026
- [Issue 31] **Gareth Edwards (Star Wars director) in favour of AI** — “*like a billionaire on acid*” — June 2026
- [Issue 32] **Martin Scorsese backs Black Forest Labs (FLUX)** — advisor role, using AI for storyboarding; triggered Boots Riley and del Toro backlash — <https://variety.com/2026/film/news/martin-scorsese-supports-ai-company-storyboard-movies-1236765037/>
- [Issue 32] **Dreams of Violets world premiere at Tribeca** — first fully AI-generated feature accepted by a major festival; cost \$2,000; directed by Ash Koosha (Fountain O); uses Kling AI, Claude, Gemini — <https://variety.com/2026/film/festivals/tribeca-festival-ai-film-dreams-of-violets-fountain-o-1236759724/>
- [Issue 32] **Amazon AI Creators Fund + Punky Duck instant cancellation** — AI animated series commissioned and killed within 48 hours after audience derision; Jorge Gutierrez dropped out — <https://futurism.com/artificial-intelligence/amazon-ai-generated-series-canceled>
- [Issue 32] **NVIDIA Cosmos 3** — omnimodal world model; text, images, video, audio, actions; best open-source T2I and I2V per Artificial Analysis — <https://nvidianews.nvidia.com/news/nvidia-launches-cosmos-3-the-open-frontier-foundation-model-for-physical-ai>
- [Issue 1] **UCL, RCA, Brandtech Centre for Creative AI launch** — Broadcast Now — <https://www.broadcastnow.co.uk/broadcasters/alex-mahon-joins-stellar-ai-creative-summit-line-up/5209227.article>
- [Issue 11] **Lovable for classrooms** — <https://lovable.dev/classroom>
- [Issue 15] **UW-Stout AI baseline competency in filmmaking** — *op. cit.*
- [Issue 15] **Google + Sundance Institute AI Education** — *op. cit.*
- [Issue 16] **UK government “Free AI training for all”** — <https://www.gov.uk/government/news/free-ai-training-for-all-as-government-and-industry-programme-expands-to-provide-10-million-workers-with-key-ai-skills-by-2030>
- [Issue 16] **Adobe at Sundance: Ignite Day** — *op. cit.*

22. Open Source and Infrastructure

- [Issue 1] **ComfyUI raises \$17M** — *op. cit.*
- [Issue 1] **PlayCanvas SOG** — *op. cit.*
- [Issue 1] **DecartAI — open-source real-time world transformation** — <https://decart.ai/>
- [Issue 1] **Civitai / Replicate** — open infrastructure layer
- [Issue 4] **Krea Realtime open-sourced** — *op. cit.*
- [Issue 5] **Meta + Hugging Face OpenEnv** — *op. cit.*
- [Issue 8] **Hugging Face + Google Cloud partnership** — https://www.linkedin.com/posts/julienchaumond_i-am-super-excited-to-announce-that-hugging-activity-7396177403972276225-CuMM
- [Issue 8] **80% of A16Z pitches use Chinese open-source models** — https://www.linkedin.com/posts/stevenouri_a-wild-stat-80-of-startups-pitching-a16z-activity-7396182718998351872-xTKR
- [Issue 11] **Ubisoft CHORD open-sourced** — *op. cit.*
- [Issue 27] **ComfyUI \$500M valuation** — *Dream Machine* coverage
- [Issue 27] **Anthropic + Blender Foundation patronage** — *Dream Machine* coverage
- [Issue 27] **Korin AI launch** — *op. cit.*

23. Web Infrastructure / Bots / Dead Internet

- [Issue 4] **Imperva 2025 Bad Bot Report** — bots = 51% of web traffic — <https://www.imperva.com/blog/2025-imperva-bad-bot-report-how-ai-is-supercharging-the-bot-threat/>
- [Issue 4] **Cloudflare crawl-to-click gap** — <https://blog.cloudflare.com/crawlers-click-ai-bots-training/>
- [Issue 4] **Dead Internet Theory** — Wikipedia — https://en.wikipedia.org/wiki/Dead_Internet_theory
- [Issue 4] **Grand View Research GenAI Content Creation Market** — <https://www.grandviewresearch.com/industry-analysis/generative-ai-content-creation-market-report>
- [Issue 4] **Futurism — AI-only social network collapses into warring tribes** — <https://futurism.com/social-network-ai-intervention-echo-chamber>

24. AI Bubble / Economics / Investment

- [Issue 1] **AI bubble 17× dotcom** — PC Gamer — <https://www.pcgamer.com/software/ai/fabulous-news-everyone-market-analyst-says-the-ai-bubble-is-17x-bigger-than-the-dotcom-goldrush-and-4x-larger-than-the-subprime-bubble-that-caused-the-2008-crash/>
- [Issue 2] **Suno \$2.45B valuation** — *op. cit.*
- [Issue 32] **Suno \$400M raise at \$5.4B valuation** — \$300M ARR, 2M+ subscribers; WMG settled; UMG/Sony in active litigation — <https://variety.com/2026/digital/news/ai->

music-suno-funding-round-400-million-5-4-billion-valuation-1236765727/

- [Issue 32] **Tripo AI raises nearly \$200M** — 3D world model, Project Eden — <https://finance.yahoo.com/sectors/technology/articles/tripo-ai-raises-nearly-200-164700546.html>
- [Issue 32] **Roblox acquires Morpheus AI** — video world models for Roblox Reality vision — <https://about.roblox.com/newsroom/2026/06/pioneering-ai-founders-join-to-accelerate-roblox-reality-vision>
- [Issue 4] **AdsGency \$12M seed** — *op. cit.*
- [Issue 5] **Sifted — Synthesia rejects \$3B Adobe** — <https://sifted.eu/articles/synthesia-acquisition-offer>
- [Issue 14] **Kartel / Reilly leadership** — *op. cit.*
- [Issue 15] **Higgsfield \$80M raise at \$1.3B** — <https://siliconangle.com/2026/01/15/higgsfield-raises-80m-1-3b-valuation-scale-ai-video-platform/>
- [Issue 16] **Higgsfield earns \$200M in 9 months** — <https://eu.36kr.com/en/p/3650517574312323>
- [Issue 16] **Synthesia hits \$4B valuation** — *op. cit.*
- [Issue 21] **Adobe + NVIDIA Strategic Partnership** — *op. cit.*
- [Issue 25] **ElevenLabs \$500M ARR** — *Dream Machine* coverage
- [Issue 27] **Google \$40B investment in Anthropic** — *Dream Machine* coverage
- [Issue 27] **ComfyUI \$500M valuation** — *op. cit.*

* * *

Using this index

This thematic index covers the *significant* sources across the *Dream Machine* archive, organised by topic. For specific research, follow the bracketed Issue numbers back to the canonical issue file in *Dream Machine MD/*. Each issue file ends with the section “All embedded URLs (in document order)” which lists every URL the issue carried, including local navigation links, profile pages and platform housekeeping links that are not reproduced here.

The newsletter is a continuous publication. The index above reflects the state of the archive at the time of book publication (June 2026). Subsequent issues will extend the catalogue. The newsletter archive itself, on LinkedIn, remains the canonical primary source for every link the book builds on.

For deeper analytical treatment of the data this index points to, see the deep-dive appendices: - Appendix D: The Shadow AI Paradox - Appendix E: Dynamics of Generative AI Adoption - Appendix F: AI, Stigma, Privilege, Democratisation - Appendix G: The Age of Intent

*Appendix I — The Doomer Mistake: Why
the Case for AI-Driven Civilizational
Collapse Keeps Losing the Argument with
the Data*

Companion piece to *Chapter 13: Coordination Collapse*, *Chapter 14: The New Jobs*, and *Chapter 15: Choosing the Future*.

This deep dive is the rebuttal companion to the labour-market sections of Chapters 13 and 14. Where those chapters argue, in the book's voice, that the AI transition is producing real but bounded structural reshuffling — a mid-career squeeze, an apprenticeship-pipeline problem, a sectoral rotation, but not a civilizational collapse — this appendix lays out the underlying empirical case against the harder doomer position that has, by mid-2026, become the consensus mood of the policy debate.

It is the analytical antibody to the framing that the OpenAI policy paper, Dario Amodei's interviews and Daniel Susskind's *A World Without Work* have together pushed into the centre of the discourse: the claim that AI represents a categorical break from prior automation, that productivity will accrue only to capital, that the labour market is in structural collapse, and that a new social contract is required to manage an inevitable surplus of unemployable humans. The deep dive sets the four load-bearing claims of that position against the Stanford Digital Economy Lab's entry-level findings, Brynjolfsson's skill-compression results, Acemoglu's macroeconomic modelling, the Dallas Fed labour analyses and the 230-year track record of identical predictions, and shows — with care — that the consensus is wrong in its central frame even where it is partially right at the edges.

The book itself takes a more cautious public posture than this deep dive does. The labour transition is real; the mid-career squeeze is real; the apprenticeship-pipeline question is real; the human cost of any reshuffle is real. None of that is in dispute. What is in dispute is the *frame*. The book's argument is that mistaking sectoral rotation for systemic collapse leads directly to policy responses that bake in exactly the rigidities that prevent the ordinary process of adjustment from working — which is the historical pattern the doomer position keeps walking the creative economy back into. This appendix is the long-form defence of that frame.

Read alongside [Appendix D: The Shadow AI Paradox](#) and [Appendix E: Dynamics of Generative AI Adoption](#), it forms the evidentiary spine of the book's argument that the creative economy of 2030 will be reshaped, redistributed and contested — but recognisably continuous with the economy of 2026, not an unrecognisable post-work successor to it.

The piece below is preserved largely as researched, with citation markers and section headings intact. Some PDF-conversion artefacts (loose footnote numbers, occasional line-break oddities) have not been editorially cleaned; the analytical content is what matters.

* * *

Page 1

DREAM MACHINE / The Doomer Mistake DREAM MACHINE / DEEP ANALYSIS The Doomer Mistake Why the case for AI-driven civilizational collapse keeps losing the argument with the data An analysis grounded in the labour data, the productivity record, and three centuries of failed technology panics Executive summary AI doomerism — the broad family of arguments holding that artificial intelligence will hollow out the labour market, dissolve the

social contract, and leave humanity economically redundant — has become the consensus mood of 2026. Sam Altman is publicly negotiating a post-work settlement. Dario Amodei tells interviewers the current economic model may need to be replaced. Newspaper headlines treat a 13% drop in employment for AI-exposed 22-to-25-year-olds as a leading indicator of civilizational replacement. This analysis argues that the consensus is wrong. Not in every particular — the labour market is genuinely being reshaped, and the disruption to entry-level cognitive work is real — but in its central frame. The doomer position rests on four claims that do not survive contact with the evidence: • That AI represents a categorical break from prior automation, making historical analogies invalid • That productivity gains from AI will accrue only to capital, not to workers • That the labour market is showing early signs of structural collapse rather than ordinary sectoral rotation • That a new social contract is required to manage an inevitable surplus of unemployable humans Each of these is a strong claim. Each is, at best, unsupported. At worst, several are flatly contradicted by the same Stanford and Brookings data the doomer camp routinely cites. The most rigorous recent work — Acemoglu’s macroeconomic modelling, Brynjolfsson’s worker-level studies, the Dallas Fed’s labour analyses — points to a far more boring conclusion: AI is a general-purpose technology following the same pattern as steam, electricity, and the personal computer. It is reshaping tasks, not eliminating work. It is compressing skill premiums, not abolishing them. And it is doing so on a timeline that gives institutions room to adapt. The doomer narrative serves political and commercial functions for the people promoting it. It does not, on present evidence, describe reality. Page 1 of 14

Page 2

DREAM MACHINE / The Doomer Mistake Page 2 of 14

Page 3

DREAM MACHINE / The Doomer Mistake 1. The starting point: what doomers actually believe Before dismantling the position, it has to be stated fairly. The strongest version of the doomer case — articulated by OpenAI’s recent policy document calling for a new “social contract,” by Dario Amodei’s public commentary, and by academic voices such as Daniel Susskind — runs roughly as follows. Productivity growth and wealth creation are being divorced from jobs and income. Previous technologies augmented human labour; AI substitutes for it. When machines become more economical than humans at both physical and cognitive work, there is simply no role left for humans to escape forward into. — Composite of OpenAI policy paper and Susskind, A World Without Work This argument has three load-bearing components. First, an empirical claim: that AI is already cutting employment in exposed occupations. Second, a theoretical claim: that this time the lump of labour fallacy — the long-standing economic rebuttal to technology panics — does not apply, because AI eliminates the cognitive frontier workers would otherwise migrate toward. Third, a normative claim: that the appropriate response is a redistributive social contract — sovereign wealth funds, robot taxes, universal basic income, a right to AI access. Each component deserves separate treatment. The empirical claim is partially true but dramatically overstated. The theoretical claim is the oldest mistake in labour economics dressed in new

clothes. The normative claim builds an emergency infrastructure for a crisis that the evidence does not show is happening. 2. The data that gets quoted — and what it actually shows The doomer position leans on a small set of recurring data points. They are real. They are also routinely stripped of the context that makes them legible. Here is the inventory. 2.1 The Stanford entry-level finding Erik Brynjolfsson's Stanford Digital Economy Lab published *Canaries in the Coal Mine?* in August 2025 — the most-cited piece of empirical work in the doomer canon. Using ADP payroll data covering millions of workers, the paper found that workers aged 22 to 25 in the most AI-exposed occupations had experienced a 13% relative decline in employment since late 2022. The 2026 Stanford AI Index extended this to a nearly 20% decline in developer employment for the same age cohort. Page 3 of 14

Page 4

DREAM MACHINE / The Doomer Mistake What gets dropped when this is recycled: the same paper finds that overall employment in AI-exposed occupations has grown robustly . The 13% figure is for one cohort — early-career — in one tail of occupations. Older workers in the same occupations have seen employment hold steady or grow. Workers in less-exposed roles have grown faster than the overall economy. The Dallas Fed, summarising the same data, was explicit: this “suggests only a slight impact on the aggregate unemployment rate so far,” driven primarily by fewer people transitioning into employment rather than by layoffs. This is not the signature of a labour market in collapse. It is the signature of a labour market where the rungs at the bottom of one ladder have been kicked off while every other ladder remains intact. Painful for the people on that ladder. Not civilizational. 2.2 The McEntarfer caveat Erika McEntarfer, until recently head of the US Bureau of Labor Statistics, made the point cleanly at Stanford's 2026 SIEPR Economic Summit: “Unemployment is edging up for those most AI-exposed occupations, but much more slowly than it is rising for everyone else.” The hiring slowdown over the prior 18 months had, she noted, mostly impacted manual labour jobs — exactly the opposite of what the AI-displacement thesis predicts. If AI were the dominant force restructuring labour, we would expect to see AI-exposed workers losing ground faster than the rest of the economy. We are seeing the reverse. 2.3 The productivity data nobody wants to talk about Brynjolfsson, Li, and Raymond's *Generative AI at Work* — published in the *Quarterly Journal of Economics* — studied 5,179 customer support agents at a Fortune 500 firm. The findings: Worker segment Productivity gain from AI access Novice / low-skill workers +34% issues resolved per hour Average across all workers +14% Experienced / high-skill workers Minimal effect on speed; small decline in quality New agents at 2 months tenure (with AI) Matched performance of 6-month-tenure agents without AI Source: Brynjolfsson, Li & Raymond, *Generative AI at Work*, QJE (2025), n = 5,179 agents. The headline finding inverts the doomer prediction. AI was supposed to widen inequality by amplifying the gap between superstar knowledge workers and everyone else — the standard “skill-biased technical change” pattern that defined the 1990s and 2000s. What Brynjolfsson found is the opposite: generative Page 4 of 14

Page 5

DREAM MACHINE / The Doomer Mistake AI compresses the skill distribution . The novice gets the apprenticeship that previously took years, delivered in real time, free of charge. Customer sentiment improved. Worker retention improved. English fluency among international agents improved. None of this is what civilizational replacement looks like. 2.4 Adoption is broad, not concentrated The St. Louis Fed reported in January 2026 that 55% of Americans and 37% of US workers were using generative AI tools. A separate Gallup Workforce Panel covering 30,000 employees from 2023 to 2026 showed workplace AI adoption rising from 9% to 26%. These are diffusion curves that look like the early personal computer, not like an extinction-level technology being deployed by a small cabal. Two facts follow. First, the benefits are spreading widely — this is consumer technology, not a hedge fund’s secret edge. Second, if AI were producing the catastrophic concentration the doomer position requires, we would expect adoption to look like the smartphone (a few firms dominate, everyone else consumes) or like industrial robotics (a few rich countries dominate). Instead, it looks like Microsoft Office. Page 5 of 14

Page 6

DREAM MACHINE / The Doomer Mistake 3. The oldest mistake in labour economics, restated for 2026 Every generation of technological disruption produces a high-status intellectual movement convinced that this time the historical pattern will break. Each movement makes essentially the same argument: previous technologies augmented narrow human capacities, but the new technology substitutes for human work as such, so the displaced have nowhere to escape forward to. The movement is always confident, always credentialled, and — over a span of roughly 230 years of available evidence — always wrong. 3.1 The historical record Era Technology Prediction What actually happened 1589 Mechanical knitting frame Elizabeth I refused the Knitting industry patent; predicted mass expanded; textile unemployment of hand employment grew for two knitters centuries 1810s Power loom Luddites: weaving as a trade 98% labour reduction per will be destroyed yard of cloth; cloth consumption surged; textile employment quadrupled 1900 Mechanised agriculture Permanent rural US farm employment fell unemployment from 41% to 2% of workforce by 2000; total employment grew by ~150 million 1930 General industrial Keynes warned of Postwar boom; lowest automation “technological unemployment of the unemployment” lasting 20th century generations 1980s–90s ATM machines Bank tellers will be extinct Tellers per branch fell from ~21 to ~13; branches multiplied; total teller employment rose ~10–15% 1980s Spreadsheet software Two million bookkeeping jobs Bookkeeping declined; destroyed financial analyst, auditor, and accountant roles grew by millions Page 6 of 14

Page 7

DREAM MACHINE / The Doomer Mistake Era Technology Prediction What actually happened 2016 Deep learning in radiology Geoffrey Hinton: “Stop Radiology employment training radiologists now” grew; the field is currently short-staffed Sources: St. Louis Fed; Bessen (IMF F&D, 2015); Autor (TED 2016); AEI; Slate Magazine; CCIA labour analysis (2026). The pattern is consistent enough to constitute a law: when a technology reduces the

cost of performing a task, demand for the output of that task expands, and the residual human tasks become more valuable. ATMs did not destroy bank tellers because banks responded to cheaper branches by opening more of them, and because the remaining teller tasks — relationship banking, product sales — were higher-value than cash counting. Spreadsheets did not destroy bookkeepers because lower-cost financial analysis meant more demand for financial analysis. Power looms did not destroy weavers because the price elasticity of cloth was much higher than the technology's pessimists assumed. 3.2 The “this time is different” claim Doomers anticipate this rebuttal and reach for the same counter-argument every generation reaches for: previous technologies replaced some tasks but left the cognitive frontier untouched. AI, the argument goes, replaces the cognitive frontier itself, so there is no “escape forward.” This is the version Susskind makes in *A World Without Work*, and the version Dario Amodi has restated in interviews. It is also exactly the argument made in 1589, 1810, 1900, 1930, and 2016 — adapted each time to whatever cognitive territory the contemporary doomer believed could not be conceded. Elizabeth I's advisors did not think knitters could become factory operators. Keynes did not think factory operators could become software engineers. Hinton did not think radiologists could become anything else. The argument's structure is invariant; only the territory shifts. The argument also misreads what AI actually does. Current systems are extraordinarily good at codified cognitive work — the kind of work that can be specified, examined, and trained on. They are markedly worse at tacit knowledge, novel-context judgement, embodied skill, and what David Autor calls “expert intuition.” The Stanford data confirms this directly: the entry-level cohort is being displaced precisely because entry-level work is codified work. Experienced workers are not being displaced because their value is in the tacit dimension AI cannot yet touch. This does not mean AI will never reach those frontiers. It means the doomer prediction is making a claim about the next two decades that is not supported by the current technology's actual capabilities or trajectory. The model that successfully passed the bar exam still cannot reliably remember what it agreed to in the previous paragraph of a contract negotiation. Page 7 of 14

Page 8

DREAM MACHINE / The Doomer Mistake Page 8 of 14

Page 9

DREAM MACHINE / The Doomer Mistake 4. Where the productivity actually goes The doomer position holds that AI's productivity gains flow to capital and starve labour. This is also testable. 4.1 Acemoglu's skeptical baseline Daron Acemoglu's *The Simple Macroeconomics of AI* (Economic Policy, 2025) is the most rigorous attempt to bound the productivity question. Acemoglu uses a Hulten-theorem-based calculation: total productivity gain equals share of tasks affected multiplied by average cost saving per task. Plugging in current estimates, he gets a 0.66% increase in total factor productivity over ten years — about 0.06% per year. He thinks even this is probably too high. This matters for two reasons. First, Acemoglu is a careful, cautious labour economist who has spent his career documenting how technology can hurt workers — he is not a hype-skeptic from the AI-optimist camp. When the most credible bear in the academy says the productivity gains over a decade are likely under

1%, the doomer claim that AI is about to dissolve the labour market is making a claim that even AI's harshest economic critic does not endorse. Second, productivity gains that small cannot produce mass unemployment. Mass unemployment requires that capital substitute for labour faster than demand expands. Acemoglu's modelling shows that under any plausible scenario, AI moves labour shares modestly — comparable to a single decade of normal technological change. 4.2 The PwC concentration finding PwC's 2026 AI Performance Study surveyed 1,217 senior executives and found that 74% of AI's measurable economic value was being captured by 20% of firms. The doomer reading: AI consolidates wealth at the top. The actually-correct reading: the technology is not yet reshaping aggregate productivity because most firms have not figured out how to deploy it. This is precisely the pattern Brynjolfsson and Hitt documented for IT investment in the 1990s: a long J-curve, with most firms experiencing negative initial returns as they fumble organizational redesign, before the diffusion phase produces broad productivity gains. We are in the fumbling phase. The doomer interpretation treats the early-phase concentration as a permanent feature. 4.3 Wages, not just jobs Page 9 of 14

Page 10

DREAM MACHINE / The Doomer Mistake Recent analysis from Phys.org and the Brookings AI labour reports finds that workers with AI-related skills are commanding wage premiums of more than 50%. Industries most exposed to AI are seeing not just productivity gains but employment growth and wage growth simultaneously. The pattern is closer to the early-internet wage premium than to a labour-share collapse. None of this is the signature of capital strip-mining labour. It is the signature of a complement: workers who use the tool earn more; firms that use the tool produce more; the diffusion is uneven but the direction is unambiguous. 5. The social contract trap This brings us back to the framing the original Narrative Goldmine note collects: that AI requires a "new social contract." Altman's OpenAI policy document — robot taxes, sovereign wealth funds, four-day workweek, right-to-AI — has become the dominant frame for how this debate is conducted at the policy level. It is worth being clear about what is happening here. Altman is the CEO of an \$852 billion company asking governments to prepare for a future in which his own product breaks the economic system. There are two ways to read this. Charitable: he genuinely believes the disruption is that large and is being responsible. Less charitable: he is constructing a political environment in which the most powerful AI companies become the indispensable partners of a new redistributive state, and in which the only acceptable critique of their products is downstream of accepting their framing of inevitability. You don't make that pitch unless you believe it's actually coming — or unless making it serves you regardless of whether it's coming. — Adaptation of The Rundown's commentary The empirical case for the doomer-driven social contract is, on the evidence reviewed above, thin. The political case is strong, and the two are easy to confuse. A reasonable response to AI's actual labour effects — entry-level displacement in some occupations, productivity compression of skill premiums, sectoral rotation — would look more like targeted reskilling, education reform, and a serious conversation about apprenticeship pathways for under-25s. It would not look like sovereign wealth funds and a robot tax. The risk of accepting the doomer frame is not that we

accidentally help displaced workers. The risk is that we build the policy architecture for a crisis that does not arrive, and in doing so create exactly the labour-market rigidities that prevent the ordinary process of adjustment from working. The countries that handled the textile transition badly were the ones that resisted it. The countries that handled the Page 10 of 14

Page 11

DREAM MACHINE / The Doomer Mistake agricultural transition badly were the ones that subsidised it indefinitely. There is a real chance that the most damaging response to AI is not the technology itself but the panic-driven policy response to it. Page 11 of 14

Page 12

DREAM MACHINE / The Doomer Mistake 6. What is actually happening Strip out the doomer overlay and what does the evidence actually describe? • AI is being adopted faster than any prior general-purpose technology — 26% of US workers using it within roughly three years of ChatGPT's release. This is faster than the PC and faster than the internet. • Productivity gains in early-adopter firms are real but heavily concentrated; the diffusion curve is in its messy phase. Aggregate measured productivity will likely lag for another two to five years before showing the gains. • The entry-level segment of cognitively-codified occupations — software, customer service, basic marketing copy — is being squeezed. Workers aged 22-25 in those roles are seeing 13-20% employment declines relative to peers. This is real and worth taking seriously as an apprenticeship-pipeline problem, not as a civilizational signal. • Older workers, manual workers, and workers in tacit-knowledge-intensive roles are unaffected or are gaining. Aggregate unemployment in AI-exposed occupations is rising more slowly than overall unemployment. • The productivity gains observed at the worker level disproportionately benefit novices and low performers — the opposite of the skill-biased pattern of the 1990s and 2000s. AI looks more like a skill leveller than a skill polariser. • Wages for AI-using workers are rising, not falling. The premium for AI fluency is currently above 50%, comparable to the early-internet skill premium. • There is no evidence of structural collapse in the labour market. There is evidence of structural reshuffling, which is what general-purpose technologies always produce. None of this denies that the transition will hurt specific people in specific places. It always does. The doomer mistake is to confuse the visible pain of the affected with a coming systemic failure. The visible pain is real. The systemic failure is not. 7. Where doomerism comes from It is worth asking why a position so weakly supported by the evidence is so dominant in the discourse. Three factors are worth naming. Commercial incentive Page 12 of 14

Page 13

DREAM MACHINE / The Doomer Mistake AI labs benefit from the narrative that their technology is so powerful it requires a redesign of civilization. It justifies valuations, regulatory capture, and the displacement of accountability onto governments. "Our product will break society, so society should pay us to break it carefully" is a remarkable business model. It is also the explicit pitch. Cognitive availability Displaced workers are visible; created

jobs are not. The bookkeeper losing her job to spreadsheets in 1985 was a story; the financial analyst hired in 1992 because spreadsheet-enabled analysis became newly affordable was just “a financial analyst.” The doomer narrative has a constant supply of vivid anecdotes; the rebuttal has only diffuse statistics. This is structural, not evidential. Status defence AI is uniquely threatening to the high-credentialed knowledge class — the journalists, lawyers, consultants, and academics who shape the discourse — in a way previous technologies were not. The agricultural mechanisation of 1900 was a story about somebody else. The current automation wave is a story about the people writing the story. The framing reflects that, and the framing should be read with that in mind. 8. Conclusion The doomer position is not entirely empty. There is a serious case to be made about entry-level cognitive work, about apprenticeship pipelines, about the speed of adaptation in education systems that move on twenty-year cycles. There is a genuine policy challenge in helping the 22-to-25 cohort transition into a labour market where the bottom rungs of certain ladders have been removed. These are real problems and deserve real attention. But the claim that AI represents an end-of-work event, that the social contract must be rewritten, that humans are about to become economically redundant — these claims rest on a misreading of the labour data, a misunderstanding of how previous technologies actually worked, and a misuse of the historical record. The Stanford findings the doomer camp cites show sectoral disruption, not systemic collapse . Brynjolfsson’s worker-level studies show skill compression, not capital triumph . Acemoglu’s macroeconomic modelling shows modest productivity gains, not labour-share extinction . The 230-year track record of identical predictions shows that “this time is different” has been wrong every single time it has been asserted with confidence. The most likely future is the boring one. AI joins the personal computer, the internet, and electricity as a general-purpose technology that reshapes work, expands output, raises living standards unevenly, and Page 13 of 14

Page 14

DREAM MACHINE / The Doomer Mistake produces a generation of policy disputes about distribution. The disputes are worth having. The civilizational frame is not. The doomer position will not survive the data. It is already failing to. Sources Acemoglu, D. (2025). The Simple Macroeconomics of AI. *Economic Policy*, 40(121), 13–58. Brynjolfsson, E., Chandar, B., & Chen, R. (2025). Canaries in the Coal Mine? Six Facts about the Recent Employment Effects of Artificial Intelligence. Stanford Digital Economy Lab. Brynjolfsson, E., Li, D., & Raymond, L. (2025). Generative AI at Work. *Quarterly Journal of Economics*, 140(2), 889–942. Stanford Institute for Human-Centered AI (2026). AI Index Report 2026. Dallas Federal Reserve (2026). Young workers’ employment drops in occupations with high AI exposure. Bessen, J. (2015). *Toil and Technology*. Finance & Development, IMF. St. Louis Federal Reserve (2026). Tracking AI’s Contribution to GDP Growth. PwC (2026). AI Performance Study. Survey of 1,217 senior executives across 25 sectors. OpenAI (2026). Policy paper on the social contract for the intelligence age. Bureau of Economic Analysis (2026). Q1 GDP estimate. Susskind, D. (2020). *A World Without Work*. Autor, D. (2016). *Will Automation Take Away All Our Jobs?* TED. American Enterprise Institute (2022). *What the Story of*

ATMs and Bank Tellers Reveals About the Rise of the Robots. CCIA (2026). What Bank Tellers and Radiologists Can Tell Us about Our Job Security in the AI Era. Page 14 of 14

*Appendix J — The Process Trap: Generative
Abundance and the Equivocation of the
Creative*

Companion piece to *Chapter 6: The 88%*, *Chapter 11: The Orchestrator*, *Chapter 13: Coordination Collapse*, and *Appendix G: AI, Intent, and the Future Artist*.

This deep dive is the philosophical and political-economy companion to the book's argument about where creative value lives after the AI transition. Where Chapter 11 develops the *orchestrator* role and Appendix G makes the case for the rise of intent over execution, this appendix performs the harder operation that the rest of the book has, in places, softened: it pulls apart the word *creative* and shows that it is doing two incompatible jobs at once — naming a rare capacity (positioned negation under risk in pursuit of significance) and naming a workforce defined by process, schedule, format and brief.

The argument, drawing on Boden's taxonomy of creativity, Csikszentmihalyi's systems model, Bourdieu's field theory, Hesmondhalgh's account of the cultural industries, Polanyi's tacit dimension and Keats's negative capability, is bracing rather than consoling. Generative AI is precisely a *process* technology: it drives the marginal cost of combinational and exploratory creativity toward zero, including the tacit craft long believed inimitable, while leaving the *transformational* act untouched — because that act resides not in the artifact but in the position of the maker. The creative industries, however, are process-selling institutions, organised to suppress the very high-variance activity now revealed as the only defensible one. The reassuring sentence circulating through every studio all-hands and trade-press op-ed — *automation will take the rote; real creative work is safe* — is true of the *capacity* and catastrophic for the *class*, and the two slide together under one word because the institutions doing the reassuring were built on the conflation and cannot see past it.

The book does not lead with this framing because the framing is uncomfortable and easy to weaponise. The 88% who turned up to the UK consultation, the unions defending performer likeness, the studios refusing AI integration on craft grounds — these positions deserve, and receive, careful treatment in the chapters. But underneath them is the structural problem this deep dive names: skill — executional competence, the proudly accumulated mastery of the process ladder — is now abundant and therefore commands no rent, and the migration *from making to creating* the appendix prescribes is harder, slower and lonelier than the reassuring sentence admits. It is also, on the evidence the deep dive assembles, the only durable response.

Read alongside *Appendix G: AI, Intent, and the Future Artist*, this appendix forms the philosophical spine of the book's account of where creative value lives in 2030. Where G is the constructive case for intent as the new scarcity, J is the eliminative case against the process economy that wore creativity's clothes and can no longer pay because the craft is free.

The piece below is the original essay text, lightly reformatted for consistency with the book's appendix style.

* * *

The Process Trap Generative Abundance and the Equivocation of the Creative Abstract The reassurance now circulating through the creative industries — that generative artificial intelligence will automate the rote while leaving genuinely creative work to humans — is at

once true and catastrophic. It is true of creativity, the rare capacity to bring the genuinely new into being; it is catastrophic for creatives, the workforce that sells something else. This article argues that the reassurance trades on an equivocation the creative industries are structurally unable to detect, because their entire apparatus is built on conflating the two terms. Drawing on Boden's taxonomy of creativity, Csikszentmihalyi's systems model, Bourdieu's theory of the cultural field, and the political economy of cultural labour, I separate process — the codifiable, transferable, executorial competence the industries actually sell — from creativity, redefined here as positioned negation under risk in pursuit of significance. Generative AI is, precisely, a process technology: it drives the marginal cost of combinational and exploratory creativity toward zero, including the tacit craft long believed inimitable, while leaving the transformational act untouched — because that act resides not in the artifact but in the position of the maker. The creative industries, however, are process-selling institutions, organised to suppress the very high-variance activity now revealed as the only defensible one. The conclusion is bracing rather than consoling. The survivors will not be the most skilled; skill is now free. They will be the most positioned, and they will be far fewer than the optimists pretend. Generative abundance did not kill creativity. It stripped the costume from it, and revealed how few were ever wearing the real thing. Keywords: generative artificial intelligence; creativity theory; cultural and creative industries; transformational creativity; cultural labour; problem-finding; aura; Bourdieu; Boden; synthetic abundance. # 1. The Comforting Lie A single reassuring sentence is doing an enormous amount of work in the creative industries right now. It runs, with minor variation: automation will take the rote and the repetitive, but real creative work is safe, because creativity is the one thing machines cannot do. It is repeated in studio all-hands and trade-press op-eds, in conference keynotes and the prospectuses of the very firms building the automating tools. It is comforting precisely because it asks nothing of anyone. It promises that the existing order of creative labour will survive the most disruptive production technology since the printing press, and that those already inside the order need only keep doing what they have always done. This essay argues that the sentence is both true and catastrophic, and that its danger lies exactly in the comfort. It is true of one thing and ruinous for another, and it survives only because the two things share a name. The word creative names a capacity: the rare act of bringing into being something that did not have to be, and need not have been, and is recognised afterward as having mattered. It also names a workforce, a sector, a class — the millions employed to execute symbolic production at industrial scale. The reassurance is true of the capacity and false of the class, and it slides between them without anyone noticing the switch, because the institutions doing the reassuring were built on the conflation and cannot see past it. The argument can be stated in a single move. The creative industries have spent the better part of a century optimising for process — craft, technique, executorial competence, production value, the dependable conversion of a brief into a deliverable — and calling the result creativity. Generative AI exposes the equivocation with surgical precision. It drives the marginal cost of process toward zero while leaving genuine creativity exactly where it stood. The consequence is not that creatives are safe. It is that creativity is safe and most creatives are not, because most creatives, trained and hired and rewarded by a process economy, sell process. The survivors will not be the most skilled. Skill — executorial competence, the thing the industry has always misnamed talent — is now abundant, and an abundant thing is no

moat. The survivors will be the few who can perform the act that process was always a proxy for and never identical to, and they will be far fewer than the consoling sentence implies. My contribution is threefold. First, I draw an analytic distinction, grounded in the psychology and philosophy of creativity, between process and creativity proper, and show that the two have been systematically fused — not by accident, but by the operating logic of the cultural industries. Second, I locate, with some precision, why generative systems annihilate the first and not the second: not because their outputs are inimitable — they manifestly are not — but because the residual moat was never in the artifact at all. Third, I draw the implication the optimists evade: the prescription is not for the industry, which cannot follow it, but for the worker, and it is harder, slower, and lonelier than the reassuring sentence admits. I begin with what synthetic abundance actually destroyed. # 2. The Collapse of the Process Premium

The companion paper to this essay describes a regime it calls synthetic abundance: the marginal cost of producing symbolic artifacts — text, image, audio, video — has fallen, with the diffusion of large generative models, asymptotically toward zero, uncoupling the supply of content from any human or cognitive constraint. The macro-economic picture is sound. But for the creative worker the relevant collapse is more specific and more intimate than a fall in the price of content. What has collapsed is the process premium: the rent that craft, technique and executional competence have historically commanded in the labour market for symbolic production. For most of the modern era, a well-made artifact was evidence of something rare. A clean retouch, a confident edit, a properly balanced mix, a competent layout, a grammatical and well-shaped paragraph — each was a costly signal, because the capacity to produce it was scarce, slow to acquire and unevenly distributed. The labour market priced that scarcity. The professional was, in large part, paid for the difference between what she could execute and what the client could not. This is the deep structure beneath the romance of the studio: the creative professional as the holder of a competence the layperson lacked, and the deliverable as proof of the competence. Walter Benjamin taught a generation to see that mechanical reproduction stripped the aura from the singular work — that the print, the recording, the copy dissolved the authority of the here-and-now original. Synthetic abundance performs a different and, for the wage-earning maker, more total operation. It does not threaten the aura of the original; it abolishes the scarcity of competence. When a competent artifact can be produced in seconds at no cost, the well-made thing ceases to be evidence of anything rare. The signal goes to noise. The professional who was paid for the gap between her execution and the client's has watched that gap close, not because she has become less skilled, but because the skill itself has been decanted into a tool the client now operates directly. The cruellest part of this is what it has reached. The optimistic assumption — repeated as confidently as the reassuring sentence — was that tacit craft would survive: the unteachable feel, the judgement in the hands, the thing that cannot be put into a manual. Michael Polanyi's formula, that we know more than we can tell, was treated as the craftworker's last redoubt. It is precisely this redoubt that large models have breached. A model trained on the aggregate residue of millions of skilled acts does not need the tacit rule to be articulable; it needs only the acts themselves, in sufficient quantity, to reconstruct the regularity statistically. The tacit dimension turns out to be tacit to the individual practitioner, not to the corpus. What no single craftsman could explain, the corpus encodes. Craft, including its silent and embodied portion, is therefore the part of creative work most exposed, not least —

the opposite of what the industry assumed. None of this means craft is worthless in the human sense. Richard Sennett's defence of craftsmanship as an intrinsic good — the desire to do a thing well for its own sake — stands entirely intact. But intrinsic goods are not market moats. The dignity of craft is real; the rent it commanded is gone. To confuse the two is to confuse a virtue with a wage, and the creative worker who stakes a livelihood on the wage while invoking the virtue is making a category error that synthetic abundance will not forgive. There is a further irony, which the cultural-industries tradition equips us to name. Long before generative models, Adorno and Horkheimer described the culture industry as a machine for standardisation dressed in the costume of pseudo-individualisation: identical structures differentiated by superficial, interchangeable detail to simulate variety and choice. The personalisation-at-scale now marketed as the answer to synthetic abundance is pseudo-individualisation industrialised — the same standardised core, recombined and surface-tuned per recipient, presented as bespoke. It is worth holding this in view, because the firm-level remedies on offer often turn out to be the disease in finer clothing. The collapse of the process premium is not solved by automating personalisation; that is process at a higher altitude. To find what survives, we have to look elsewhere — and first we have to dismantle the word that has been hiding the problem. # 3. The Equivocation of the Creative The word creative does two incompatible jobs, and the reassuring sentence depends on our not noticing. As an adjective of capacity it picks out a rare event: the genuine origination of the new. As a noun of occupation — a creative, the creatives, the creative class, the creative industries — it picks out a category of employment defined by sector and task, with no implication whatever that the work so described is creative in the strong sense. Almost all public discussion of AI and creative work equivocates between the two, and the equivocation is not innocent. It is the load-bearing wall of an entire ideological structure. The cultural-industries scholarship has been telling us for decades what the creative industries actually are, and it is not a guild of originators. Hesmondhalgh describes their organising logic as the management of risk through formatting, repertoire and star systems: the industrialisation of symbolic production, in which the unpredictability of the genuinely new is contained by repeating proven formulae, hedging across portfolios, and binding audiences to recognisable franchises. The overwhelming majority of creative employment is, on this account, process work: execution against a brief, within a format, to a specification, on a schedule. The junior designer producing the fourteenth variant of a campaign; the editor cutting to a house style; the staff writer hitting a template and a word count; the production artist preparing files for output — these are skilled roles, often highly so, and they are process roles to the core. The work is the reliable application of competence to a defined problem. It is exactly, and only, what generative systems now do at no cost. Why, then, does everyone in these roles feel like a creator? Because the romance of the artist is doing ideological labour. The genius myth — the solitary originator, the spark, the gift — circulates through the creative industries not as an accurate description of the work but as a disciplinary device. It persuades the process worker that she is an artist, and so makes her accept the conditions of artistic life: precarity, long hours, deferred pay, self-exploitation, the privilege-of-the-work as substitute for security. The political economy of cultural labour has documented this with some bleakness. The romance is what allows an industry of process to recruit, retain and underpay a workforce that believes itself to be a community of creators. Generative AI calls the bluff with brutal economy. It

performs the process, leaves the romance intact, and thereby reveals that the romance was never load-bearing on the value side — only on the labour-discipline side. The work was process; the self-image was art; the wage was paid for the former and justified by the latter. So when the reassuring sentence says creatives will survive because they are creative, it equivocates across this exact gap. Creativity, the capacity, will survive — I will argue it cannot be otherwise. Creatives, the class, largely will not, because the class is defined by, hired for, and skilled in the thing that just became free. The two propositions feel like one sentence because one word covers both. Pull the word apart and the consolation evaporates. The honest formulation is harsher and more useful: creativity is safe; most who are called creative are not, because being called creative was never the same as being creative. To make that formulation rigorous rather than merely rhetorical, we need an account of what creativity proper actually is — one precise enough to explain why a machine that has mastered process still cannot touch it. # 4. What Creativity Is, and Why It Resists Automation Begin with a concession the defensive literature is reluctant to make, because the whole subsequent argument depends on its being made honestly. Generative models are extraordinarily creative in two of the three senses that matter, and pretending otherwise is the surest route to being wrong about the third. Margaret Boden's taxonomy remains the most serviceable map. She distinguishes combinational creativity — the novel combination of familiar ideas; exploratory creativity — the generation of new structures within an established conceptual space, by searching its possibilities; and transformational creativity — the alteration of the conceptual space itself, so that what was previously impossible within it becomes possible. Generative models are superb combinational engines. They are, in Koestler's older vocabulary, bisociation machines: their high-dimensional latent spaces collide previously unrelated frames of reference tirelessly and at scale, which is exactly the operation Koestler placed at the root of the creative act. They are also formidable explorers, traversing the possibility-space of a genre or a style faster and more exhaustively than any human could. To insist that this is mere interpolation, that the model only ever recombines and never creates, is both technically shaky and strategically suicidal: it stakes the human position on a claim the next model release will embarrass. P-creativity — novelty relative to the individual mind — is now cheap. A great deal of what passes for H-creativity — novelty relative to the whole of history — is within reach of a sufficiently prompted system. The genuine and durable limit is at the third level, transformational creativity, and the discipline of the argument lies in locating why the limit holds — not in the artifact, but in four interlocking features of the maker's position. None of the four is a property of the output. All four are properties of where the maker stands. This is the crux of the essay, and the reason the residual moat is so easily mislocated. ### 4.1 The originating problem The most under-absorbed result in the psychology of creativity comes from Getzels and Csikszentmihalyi's longitudinal study of art students. What predicted later artistic achievement was not problem-solving skill but problem-finding behaviour: the disposition to discover, frame and reframe the problem worth working on, rather than to execute against a problem already given. This is the precise axis along which human and machine diverge. A generative model is a peerless problem-solver and a non-problem-finder. It has no problem of its own. It can answer any question put to it and originate none, because origination of a problem requires a standpoint from which some absence in the world registers as a lack worth repairing. Hannah Arendt gives this its proper name: natality, the

human capacity to begin, to introduce into the world something that the world did not contain and could not have predicted. Creation in the strong sense is beginning, and beginning is something done by someone for whom the beginning matters. The model continues; it does not begin. Prompt it and it solves; it never arrives, unbidden, at the problem that nobody asked it to have. ### 4.2 The stake Bourdieu's field theory supplies the second feature and reframes the whole question of value. In the field of cultural production, the worth of a work is not a property the work carries intrinsically; it is produced by position-taking — by an agent who occupies a determinate position in the field and makes a move that is legible, and consequential, as a move within it. What animates the agent is *illusio*: an invested belief in the stakes of the game, a sense that the moves matter and that one's own move risks something. Here the machine is not merely deficient but categorically absent. A generative model occupies no position in any field. It has no biography, no trajectory, no rivals, nothing accrued and nothing to lose. It cannot make a move that costs it anything, because there is no it for the move to cost. Harry Frankfurt's analysis of bullshit sharpens the point: the bullshitter is defined not by lying but by indifference to the truth of what he says. The model is the perfected case — not deceitful, simply unconcerned — and its fluency is so persuasive precisely because nothing in it is at stake. Creative value, on the Bourdieusian account, is staked value. The maker who has skin in the game produces a different kind of object from the maker who does not, even when the surfaces are identical, because the object is a wager and the wager is part of the meaning. ### 4.3 Significance Csikszentmihalyi's mature position locates creativity not in the person at all, but in a system: the interaction of an individual, who varies a domain; a domain, the symbolic tradition that stores the variation; and a field, the community of practitioners and gatekeepers that judges which variations are retained. On this model, a transformation is creative only if the field judges it to matter and the domain absorbs it. This is decisive for our question, because it relocates the scarce thing. The hard part of transformational creativity is not generating a transformation — a model can generate a thousand — but knowing which transformation matters, a judgement that is internal to a living tradition and earned by long, embedded participation in it. Mattering is not in the artifact; it is a relation between the artifact, the history that makes it legible, and the field that confers significance. The model can produce the move. It cannot know the move matters, because knowing-that-it-matters is not the kind of thing an output can contain. It is the kind of thing a positioned, historically embedded judge supplies. ### 4.4 Negation under risk The fourth feature is the one most directly opposed to the machine's architecture, and it returns us to Keats. Negative capability — the capacity to remain in uncertainty, mystery and doubt without any irritable reaching after fact and reason — is not a poetic ornament but a description of the cognitive posture that genuine creation requires. To make something new is, in large part, to negate: to refuse the obvious continuation, the probable phrase, the well-formed and expected resolution; to dwell in the not-yet-resolved long enough for the improbable to become available. A next-token model is built to do the exact opposite. Its gradient runs toward the probable; it is optimised, at every step, to reach for the likeliest continuation. This is not a flaw to be patched but the principle of the thing. It is also the mechanism behind what Ghafouri calls the AI Prism: the systematic compression of variance toward the mean as outputs converge on the high-probability paths. Ghafouri rightly notes a countervailing Paradoxical Bridge — recombinant novelty available to those who refuse

epistemic deference and spar with the model rather than defer to it. But notice what the bridge requires: a human who negates, who pushes against the probable, who carries the stake and the problem. The bridge is not a property of the tool. It is the positioned human creativity of §§4.1–4.3, using the tool. And creation under negation is creation under risk, because to abandon the probable is to court failure that costs the maker something — reputation, time, standing in the field. The model cannot fail in this sense; there is nothing it can lose. Risk-free production is, definitionally, the production of the unrisks, which is to say the probable, which is to say the mean. ### 4.5 The residual moat is the position, not the artifact Assemble the four and a definition emerges that is precise enough to do work. Creativity, in the strong sense that resists automation, is positioned negation under risk in pursuit of significance: the bringing-into-being of something that did not have to be, by an agent who has a problem of her own and a stake in the game, achieved by refusing the probable, and judged to matter by a living field. Every clause names a feature of the maker's position; not one names a feature of the output. This is why the moat is so consistently mislocated. The industry, the commentary, and the defensive literature all look for the human residue in the artifact — the brushstroke the model can't fake, the turn of phrase it can't match, the *je-ne-sais-quoi* of the human hand — and there is nothing reliably there, because there was never anything reliably there. The model will match the surface; surfaces are exactly what it learns. The defensible thing was always upstream of the surface, in the standpoint from which the work was made. Strip the costume from creativity and what remains is not a better-made object. It is a person who had a problem, took a stake, refused the easy answer, and was answerable to a field. That person is a much rarer thing than the creative industries have ever admitted — and the admission is overdue. # 5. The Process Trap If the defensible thing is the maker's position and the worthless thing is executorial competence, the creative industries are in a far worse situation than the reassuring sentence allows, and for a structural reason. They have spent a century training, hiring and rewarding for exactly the wrong variable. Call this the process trap: the institutionalised mistaking of process mastery for creativity, built so deep into pedagogy, recruitment, tooling and reward that the people inside it cannot easily climb out, because the trap is also their identity. The Dreyfus model of skill acquisition describes the ladder the creative professional has been taught to climb: novice, advanced beginner, competent, proficient, expert. Every rung is a rung of process mastery — the progressive internalisation of technique until execution becomes fluent, intuitive, unconscious. The entire prestige economy of creative work is built on this ladder. The master craftsperson at the top is the most admired and the best paid, and the figure everyone is taught to become. Generative systems have ascended the lower and middle rungs of this ladder with a speed that earlier scepticism, including Dreyfus's own, did not predict. And here is the trap's cruelty: the higher you climbed, the more exposed you now are. The expert's identity, income and self-conception are most heavily invested in the very competence that just became abundant. The process trap is a sunk-cost trap of the self. The years of deliberate practice, the hard-won fluency, the craft pride — all of it real, all of it now uncompensated by the market that priced it. The better you were at process, the harder the fall, and the more of your sense of who you are is staked on the thing that no longer commands a wage. One might expect the industry to respond by retraining its workforce toward creativity proper. It cannot, and the reason is not failure of will but structure. The cultural industries are process at scale

by design. Their entire apparatus — the format, the repertoire, the franchise, the brief, the schedule, the portfolio hedge — exists to suppress the high-variance, illegible, unschedulable, frequently unprofitable activity that genuine creation actually is. Positioned negation under risk is everything the industrial machine is built to manage out of existence. You cannot put problem-finding on a Gantt chart; you cannot guarantee significance to a client; you cannot schedule negative capability for the third quarter; you cannot de-risk a wager without destroying it. The institutions are optimised against the only thing now worth optimising for. They will therefore do the thing they know how to do — automate the process they were already selling, faster and cheaper — and call it transformation, and in doing so accelerate precisely the collapse they are trying to survive. This is also why the firm-level moats proposed in the companion paper — experience orchestration, community, personalisation, data network effects — though real for firms, are largely irrelevant to, or actively corrosive of, the position of the individual maker. They are, on inspection, mostly process moats wearing the language of the human: orchestration is process applied to the customer journey; personalisation-at-scale is pseudo-individualisation, as we have seen; community-as-retention-mechanism is the management of attention, not the practice of creation. A firm can build a durable business on these and employ, in the building of it, almost no one doing creative work in the strong sense. The moat protects the platform; it does not protect the creative, and frequently it is constructed by replacing her. The conflation at the level of the word reappears at the level of strategy: defensibility for the enterprise is sold as defensibility for the worker, and they are not the same thing. The community appears once more in the next section — but in a different role, not as a retention asset but as something without which creativity is literally impossible. # 6. The Migration: From Making to Creating What, then, is to be done — by the worker, since the institution cannot? The prescription follows directly from the definition. If creativity is positioned negation under risk in pursuit of significance, then migrating out of the process trap means deliberately cultivating each of the four features that no model can supply. This is not a reskilling programme; reskilling is more process. It is closer to a change of vocation within the same vocation. I set it out plainly, with no pretence that it is easy or that most will manage it. Compete on the problem, not the execution. The first discipline is the inversion of everything the process trap rewards: to become a problem-finder in a world of infinite problem-solvers. When solution is free, all the remaining value migrates to the question — to the framing of a problem worth solving that no one else had seen as a problem. The practical form of this is a refusal: stop competing on how well you execute, because you will lose to a tool, and start competing on what you choose to make and why it matters. The creative worker who can originate the brief is worth orders of magnitude more than the one who can fulfil it, and the gap will only widen. This is learnable, but it is learned by exposure, obsession and standpoint, not by technique — which is exactly why the industry, organised around technique, cannot teach it. Take a stake; make work that can fail and that costs you something. The second discipline is to re-introduce risk on purpose, against every incentive of a career built to minimise it. Bourdieu's *illusio* cannot be faked; either you have something invested in the move or you do not. In practice this means treating biography, conviction and obsession as inputs rather than contaminants — making the work that only you, from where you stand, would make, and accepting that it might be rejected. The unrisks work is now produced for free by a machine that risks nothing. The only work with a

defensible premium is work that could have failed, made by someone for whom failure would have cost something. Safety, the prized virtue of the process professional, is now a liability: safe work is the machine's native register. Re-enter the field. The third discipline reframes the community the companion paper treats as a retention moat into something prior and more fundamental: the organ of significance-judgement without which creativity, on Csikszentmihalyi's systems model, cannot occur at all. Significance is not in the artifact and not in the lone maker; it is conferred by a living field of discriminating peers embedded in a domain's history. To create transformationally is to be answerable to such a field — to know its tradition deeply enough to know which break from it will register as meaningful rather than merely deviant. The isolated genius is, here too, a myth; the real condition of creation is embeddedness. This is the one point at which the community argument and the creativity argument converge — but the function is different. Community is not where you retain customers. It is where mattering is decided, and a maker cut off from a field has no access to the very thing that distinguishes creation from noise. Practise negation; use the machine as the producer of the probable. The fourth discipline is the most counter-intuitive and the most operationally precise. The correct use of a generative system, for a maker pursuing creativity, is not as a creative partner and not as a curation engine, but as the maximally competent producer of the probable — precisely so that the human can specialise in the improbable. Let the model occupy the mean it is built to occupy. Let it generate the obvious continuation, the expected resolution, the competent default, exhaustively and instantly. The human contribution is then defined by negation: the refusal of what the model offers, the reaching past the probable toward the move the model could not make because the model is built to avoid it. This reframes the worn phrase human-in-the-loop from curation — choosing among the machine's options, which is still deference — to negation: using the machine to map the space of the probable so as to depart from it deliberately. Ghafouri's Paradoxical Bridge is crossed in exactly this posture and no other. The maker who defers to the model's best output has re-entered the process trap by another door. The maker who uses the model's best output as a map of what to avoid has found the only durable use of the tool. I will not soften the conclusion of this section, because softening it is the optimists' characteristic dishonesty. Most creative workers will not make this migration. It is harder than the work they were trained for, slower to reward, lonelier, and structurally unsupported by the institutions that employ them. To pretend that everyone called creative can simply ascend to creativity proper is to repeat the equivocation in a consoling key, and it does a disservice to people making real decisions about real livelihoods. The scale of dislocation will be large, and naming it honestly is the least a serious account owes its subject. The migration is available; it is not general; and the difference between those two facts is where most of the human cost of this transition will fall. # 7. The Last to Fall Return, finally, to the sentence that began this essay, and to the slogan beneath it — that creatives will be the last people standing. There is a sense in which it is true, and the sense matters. When process is free and code is free and competent symbolic production is free, the one thing that retains a premium is the strong-sense creative act: positioned negation under risk in pursuit of significance, an act no model can perform because the act is constituted by features of the maker's standpoint and not of the artifact. In that sense, yes — the creative is the last to fall. Creation outlasts production. It always did; abundance has merely made the distinction visible by making everything else cheap. But the

standing ground is much smaller than the slogan implies, and the people standing on it are not who the industry believes them to be. They are not the most skilled, because skill — executional competence, the proudly accumulated mastery of the process ladder — is now abundant and therefore commands no rent. They are the most positioned: those who have a problem of their own, a genuine stake, an embedded place in a living field, and the nerve to refuse the probable answer the machine hands them for free. That is a far narrower population than the creative class, and it never mapped neatly onto employment, prestige or even talent in the ordinary sense. The reckoning now under way is therefore not, at bottom, a contest between artificial intelligence and creatives. It is a contest between creativity and the vast process economy that wore creativity's clothes — that recruited a workforce on the romance of art and paid it for the execution of craft, and that can no longer pay, because the craft is free. Generative abundance did not kill creativity. It did something more revealing and, for an industry built on the conflation, more dangerous. It stripped the costume from creativity and held it up to the light — and showed how few were ever wearing the real thing, and how many had mistaken the costume for the body. The honest response to that revelation is not reassurance. It is the harder work of becoming, deliberately and against every incentive of the process trap, one of the few. Most of the field will not. Some will. The difference between them will not be talent, and it will certainly not be skill. It will be where they choose to stand.

References

- Adorno, T. (1941). "On Popular Music." *Studies in Philosophy and Social Science*, 9(1), 17–48.
- Adorno, T. & Horkheimer, M. (1944). *Dialectic of Enlightenment*. (Trans. E. Jephcott, 2002). Stanford: Stanford University Press.
- Arendt, H. (1958). *The Human Condition*. Chicago: University of Chicago Press.
- Banks, M. (2007). *The Politics of Cultural Work*. Basingstoke: Palgrave Macmillan.
- Benjamin, W. (1936). "The Work of Art in the Age of Its Technological Reproducibility." In *Selected Writings*, Vol. 3 (2002). Cambridge, MA: Harvard University Press.
- Boden, M. (2004 [1990]). *The Creative Mind: Myths and Mechanisms*, 2nd edn. London: Routledge.
- Bourdieu, P. (1993). *The Field of Cultural Production*, ed. R. Johnson. Cambridge: Polity Press.
- Bourdieu, P. (1996 [1992]). *The Rules of Art: Genesis and Structure of the Literary Field*. Stanford: Stanford University Press.
- Csikszentmihalyi, M. (1996). *Creativity: Flow and the Psychology of Discovery and Invention*. New York: HarperCollins.
- Dreyfus, H. & Dreyfus, S. (1986). *Mind Over Machine: The Power of Human Intuition and Expertise in the Era of the Computer*. New York: Free Press.
- Florida, R. (2002). *The Rise of the Creative Class*. New York: Basic Books.
- Frankfurt, H. (2005 [1986]). *On Bullshit*. Princeton: Princeton University Press.
- Getzels, J. W. & Csikszentmihalyi, M. (1976). *The Creative Vision: A Longitudinal Study of Problem Finding in Art*. New York: Wiley.
- Ghafari, B. (2025). "The Variance Paradox: How AI Reduces Diversity but Increases Novelty." [arXiv:2508.19264](https://arxiv.org/abs/2508.19264).
- Hesmondhalgh, D. (2019 [2002]). *The Cultural Industries*, 4th edn. London: Sage.
- Jenkins, H. (2006). *Convergence Culture: Where Old and New Media Collide*. New York: NYU Press.
- Keats, J. (1817). Letter to George and Thomas Keats, 21 December. In *Selected Letters* (2002). Oxford: Oxford University Press.
- Koestler, A. (1964). *The Act of Creation*. London: Hutchinson.
- McRobbie, A. (2016). *Be Creative: Making a Living in the New*

Culture Industries. Cambridge: Polity Press. Pine, B. J. & Gilmore, J. H. (1999). *The Experience Economy: Work Is Theatre and Every Business a Stage*. Boston: Harvard Business School Press. Polanyi, M. (1966). *The Tacit Dimension*. Chicago: University of Chicago Press. Ross, A. (2009). *Nice Work If You Can Get It: Life and Labor in Precarious Times*. New York: NYU Press. Schumpeter, J. A. (1934 [1911]). *The Theory of Economic Development*. Cambridge, MA: Harvard University Press. Schumpeter, J. A. (1942). *Capitalism, Socialism and Democracy*. New York: Harper & Brothers. Sennett, R. (2008). *The Craftsman*. New Haven: Yale University Press. Simon, H. A. (1971). "Designing Organizations for an Information-Rich World." In M. Greenberger (ed.), *Computers, Communications, and the Public Interest*. Baltimore: Johns Hopkins Press.

*Appendix K — The Coordination Collapse:
How AI Is Dissolving the Information Tax
That Created Hierarchy*

Companion piece to *Chapter 13: Coordination Collapse*, *Chapter 11: The Orchestrator*, and *Chapter 14: The New Jobs*.

This deep dive is the long-form organisational-design companion to Chapter 13. Where the chapter argues, in the book's voice, that the existing organisations of the creative industries are structurally *coordination architectures* whose calculus changes when the information-routing tax falls to near-zero, this appendix builds the underlying organisational-theory case in full — the 2,000-year history of hierarchy as an information-routing protocol constrained by human bandwidth, the empirical evidence for shadow AI as the symptom of a workforce running its own private coordination infrastructure, the *jagged technological frontier* identified by Dell'Acqua et al.'s preregistered Boston Consulting Group experiment, the *vigilance problem* documented by Dell'Acqua (2023), and a structured comparison of three competing post-hierarchical models: Block's intelligence layer, Every's agent colony, and DreamLab's Dynamic Agentic Mesh.

The piece is the work of Dr. John O'Hare of DreamLab AI Consulting (April 2026) and represents the most ambitious organisational-design treatment of the coordination collapse thesis I have seen produced anywhere in the period this book covers. It runs to sixteen chapters with a full appendix structure of its own — research methodology, competitor comparison matrix, glossary — and is preserved here in substantially complete form because the argument is cumulative and benefits from sequential reading.

For the working creative or the studio leader: Chapters 2–5 of the report establish the diagnosis (the history of hierarchy, the shadow workflow problem, the jagged frontier, the vigilance problem). Chapter 6 presents the three competing models. Chapters 7–9 develop the *compounding organisation* concept, the new KPIs (Mesh Velocity, Augmentation Ratio, Trust Variance, HITL Precision) and the *governance paradox* — the argument that embedded governance accelerates rather than constrains agentic work. Chapters 10–14 address change management, the industry landscape, the *judgment broker* role replacing middle management, and a 90-day implementation roadmap. Chapter 15 is a self-critique of every known gap in the argument. Chapter 16 synthesises the through line with audience-specific closing messages for middle managers, change managers and AI leaders.

The report's central thesis — that **hierarchy is not a management philosophy but an information routing protocol constrained by human bandwidth, and that when information routing becomes computationally cheap the structural reason for hierarchy dissolves** — is the most rigorous statement of the coordination-collapse argument I can point a working studio leadership team at. Where Chapter 13 develops this in the book's working-creative voice, this appendix is the organisational-design and consulting-grade evidence base underneath it.

Read alongside Appendix D: The Shadow AI Paradox, which covers the workforce side of the same transition, this appendix constitutes the structural-design spine of the second half of the book.

The piece below is preserved largely as researched, with citation markers, section headings, table structures and figure references intact. Some PDF-conversion artefacts (loose page numbers, line-break oddities, occasional figure markers where the source contained

diagrams) have not been editorially cleaned; the analytical content is what matters. References to internal cross-chapter and appendix structures ('Chapter X', 'Appendix Y', 'Figure Z') in the body below refer to *the source report's* internal structure, not the book's.

* * *

Page 1

The Coordination Collapse How AI Is Dissolving the Information Tax That Created Hierarchy – And What Replaces It Dr. John O'Hare DreamLab AI Consulting Ltd DreamLab, Fairfield, Eskdale April 2026 A comprehensive research report on agentic organisational design, the jagged technological frontier, and the future of human–AI coordination.

Page 2

2

Page 3

Contents

1	Introduction	1	1.1 The Central Question	1	1.2
	Methodology and Scope	2	1.3 Reader's Guide		
		2	1.4 Structure of the Report	3	2
2	The 2,000-Year Problem	5	2.1 Why Organisations Look the Way They Do		5
2.2	The Roman Protocol		2.3 The Prussian Innovation		
		6	2.4 The Railroad Commercialisation		
6	2.5 Scientific Management and the Functional Pyramid		2.6 A Century of Failed Alternatives		
		7	2.7 The Core Insight		
		8	3 The Collapse Point	9	
		9	3.1 What Happens When Coordination Becomes Near-Free		
		9	3.2 The Secret Cyborg Problem		10
		10	3.2.1 The Epoch AI Employment Survey (March 2026)		
		11	3.3 The Shadow Workflow Problem		
		12	3.4 The Rational Calculus of Secrecy		
		12	3.4.1 The Psychological Safety Evidence		
		13	3.5 The Divergence Between Formal and Actual Structure		
		13	3.6 The Cost of Inaction		
		15	4 The Jagged Frontier	15	
		15	4.1 Why Wholesale Automation Is Reckless		
		15	4.2 The BCG Experiment		
		17	4.3 The Skill-Levelling Effect		
		17	4.4 Three Collaboration Modes		
		18	4.5 Why the Frontier Matters for Governance		
		19	4.6 The Frontier Is Dynamic		i

Page 4

5	The Vigilance Problem	21	5.1 Why "Human in the Loop" Is Not Enough		
21	5.2 Falling Asleep at the Wheel		5.3 The Quality Paradox		
		22	5.4 Implications for Organisational Design		
		22	5.4.1 Passive Oversight Degrades		23
		23	5.4.2 Active Engagement Is Required		
		23	5.4.3 Deliberate Friction Is a Design		

Requirement	23	5.4.4 Governance Must Be Continuous	23
23 5.5 The Atrophy Prediction	23	5.6 What This Means for the Three Models	24
6 Three Models for What Comes Next	25	6.1 The Emerging Paradigms	25
6.2 Model A: The Intelligence Layer	25	6.2.1 Source and Core Thesis	25
6.2.2 The Case For	26	6.2.3 The Case Against	26
6.3 Model B: The Agent Colony	27	6.3.1 Source and Core Thesis	28
6.3.2 What Works	28	6.3.3 What Does Not Work	28
6.4 Model C: The Dynamic Agentic Mesh	29	6.4.1 Core Thesis	29
6.4.2 Three Layers	29	6.4.3 The Insight Ingestion Loop	30
6.4.4 Honest Limitations	30	6.5 Comparative Analysis	31
7 The Compounding Organisation	33	7.1 When the Learning System Becomes the Product	33
7.2 The Compound Engineering Ladder	33	7.3 The Organisational Flywheel	34
7.4 Everyone Is a Manager Now	35	7.4.1 Why This Is Threatening	35
7.5 The Conductor Problem	36	7.6 How to Present This Without Patronising	36
36 ii			

Page 5

8 New KPIs for a Fluid Organisation	39	8.1 Why the Existing KPI Landscape Is Inadequate	39
8.2 Mesh Velocity	40	8.3 Augmentation Ratio	41
8.4 Trust Variance	41	8.5 HITL Precision	42
8.6 The KPI Dashboard as a Narrative	42	9 The Governance Paradox	45
9.1 Governance as Accelerator, Not Brake	45	9.2 The Racing Car Principle	45
9.3 Declarative Governance	46	9.4 HITL by Design	47
9.5 Ontological Provenance	47	9.6 Cascading Trust Hierarchies	48
9.7 The Governance Paradox Resolved	48	10 The Change Problem	51
10.1 Why This Cannot Be Managed with a Gantt Chart	51	10.2 The Continuous Change Architecture	52
10.2.1 Continuous Transformation	52	10.2.2 Adaptive Change Architecture	52
10.2.3 Dynamic Capabilities	52	10.2.4 Directed Autonomy	52
10.3 The 1-9-90 Pattern	53	10.4 Resistance Patterns	53
10.5 The Contradiction in This Report	54	11 The Industry Landscape	57
11.1 What the Major Consultancies Are Saying	57	11.2 McKinsey: The Agentic Organisation	57
11.3 BCG: Quantifying the Shift	57	11.4 Deloitte: Measuring Maturity	58
11.5 Gartner: Structural Predictions	58	11.6 World Economic Forum: The Jobs Landscape	59
11.7 Agentic AI Adoption at Scale	59		

. 59	11.8 Comparative Positioning	60	12 The Role Transformation 61 iii
------	--	----	-----------------------------------

Page 6

12.1	The Middle Manager as Judgment Broker	61	12.2 Three Irreplaceable Human Capacities	61
12.2.1	Strategic Vectors	62	12.2.2 Ethical Adjudication	62
12.2.3	Relational Intelligence	62	12.3 Management as AI Superpower	63
12.4	The Competency Shift	63	12.5 The Day-in-the-Life Question	64
13	The Technical Substrate: VisionClaw 67	13.1	A Working Agentic Mesh	67
13.2	Architecture Stack	67	13.2.1 OWL 2 Ontology Engine	68
13.2.2	Claude-Flow Orchestration	68	13.2.3 WebXR 3D Knowledge Visualisation	68
13.2.4	Nostr Agent Identity	68	13.2.5 RuVector Memory	69
13.3	What Is Real and Verified	69	13.4 What Was Previously Overclaimed	70
13.5	The Corrected Story	70	13.5.1 What This Proves	70
13.5.2	What This Does Not Yet Prove	71	13.6 Positioning Guidance	71
14	From Org Chart to Agentic Mesh: Implementation 73	14.1	The Bootstrap Sequence	73
14.2	Phase 1: Audit and Discovery (Days 1–30)	73	14.2.1 Shadow Workflow Audit	74
14.2.2	Friction Point Identification	74	14.2.3 Monitoring Deployment	75
14.2.4	Baseline Measurement	75	14.2.5 Judgment Broker Identification	75
14.2.6	Legal and Regulatory Groundwork	76	14.3 Phase 2: Architecture and Codification (Days 31–60)	76
14.3.1	Agent Skill Pilot	77	14.3.2 Domain Ontology Construction	77
14.3.3	DAG Construction	78	14.3.4 Judgment Broker Workshops	78

Page 7

14.3.5	Protected Experimentation Time	78	14.4 Phase 3: Live Mesh and Governance (Days 61–90 and Beyond)	79
14.4.1	Production Mesh Activation	79	14.4.2 Governance Activation	79
14.4.3	Insight Portal Launch	79	14.4.4 First Mesh Intelligence Report	80
14.4.5	Compound Loop Initiation	80	14.5 The Wardley Map: Middle Manager Evolution	80
14.6	Cost Tier Model	81	14.7 Failure Criteria: When to Stop	82
14.8	Resistance Mitigation in the Roadmap	82	15 Open Questions and Research Directions 85	15.1
15.1	The Core Thesis Is Incomplete	85	15.1.1 Legitimation	85
15.1.2	Accountability	85	15.1.3 Political Negotiation	86
15.2	Evidence Base Weaknesses	86	15.2.1 The Jagged Frontier Map Is	

Outdated	86	15.2.2 The 17 Error Source Is Weak	87
15.2.3 The Klarna Narrative Is Cherry-Picked	87	15.2.4 The 89 Per Cent Reduction	88
15.3 Audience Blind Spots	88	15.3.1 No Day-in-the-Life	88
15.3.2 No Total Cost of Ownership	89	15.3.3 The Roadmap Contradiction	89
15.3.4 The Analyst-Vendor Conflict	89	15.4 Engaging with Acemoglu	91
15.5 Energy and Environmental Costs	91	15.6 Cybersecurity Attack Surface	92
15.8 Competing Frameworks Not Addressed	92	15.9 The VisionClaw Credibility Gap	93
15.10 Additional Failure Modes	94	16 Conclusion: The Through Line	95
Synthesis	95		v

Page 8

16.2 Block Was Right About the Diagnosis	96	16.3 Everyone Is a Manager Now — And That Changes Everything	97
16.4 The Compound Flywheel	98	16.5 What the Evidence Supports	98
16.6 What the Evidence Does Not Yet Support	99	16.7 Closing Messages	99
16.7.1 For Middle Managers	99	16.7.2 For Change Managers	100
16.7.3 For AI Leaders	100	16.7.4 For All Three	101
References	103	Research Methodology	107
1. Multi-Agent Research Architecture	107	2. Source Selection and Quality Assessment	108
2.1 Tier 1: Peer-Reviewed and Pre-Registered Research	108	2.2 Tier 2: Institutional Research with Disclosed Methodology	108
2.3 Tier 3: Expert Commentary and Primary Accounts	109	2.4 Tier 4: Community and Vendor Sources	109
3. Web Search and Academic Database Approach	109	4. Codebase Audit Methodology	110
5. Wardley Mapping Methodology	111	6. Limitations of the Research Process	111
7. Disclosure	112	Competitor Comparison Matrix	113
8. Comparison Dimensions	113	9. Comparison Tables	113
10. Strategic Positioning Analysis	116	10.1 The Architecture Gap	116
10.2 The Governance Gap	116	10.3 The Scale Gap	117
10.4 The Evidence Gap	117	11. Wardley Map: Three Models Compared	118
12. Choosing the Right Framework	119	13. Interpretation	119
Glossary of Key Terms	121		vi

Page 9

List of Figures	16
4.1 The Jagged Technological Frontier	16
6.1 Block's Intelligence Layer Model	26
6.2 Every's Agent Colony Model	

.....	27	6.3 DreamLab's Dynamic Agentic Mesh	29
13.1 VisionClaw System Architecture	67	14.1 Wardley map showing the evolution of the middle management role from information routing (commodity, automatable) toward judgment brokerage (genesis/custom-built, irreducible). The horizontal axis represents evolutionary stage; the vertical axis represents value chain position.	80
16.1 The compound organisational flywheel. Discovery feeds codification; codification feeds validation by the judgment broker; validated workflows propagate through the mesh; propagation increases organisational capability; increased capability accelerates the next cycle of discovery. Each revolution makes the next one faster.	98	2 Wardley map comparing the three primary models (Block, Every, DreamLab) on the evolution axis. Block's Intelligence Layer positions AI coordination as a product/commodity; Every's Agent Colony treats coordination as genesis/custom; the Agentic Mesh positions the judgment broker at the custom stage while using commodity AI for information routing.	118
	vii		

Page 10

viii

Page 11

List of Tables	3.1	Costs of unmanaged shadow AI adoption	14	4.1
Human-AI collaboration modes and performance outcomes	18	6.1	Comparative analysis of three post-hierarchical models	31
7.1	Compound engineering maturity stages (after Klaassen, 2025)	34	8.1	Summary of agentic organisation KPIs
43	9.1	Imperative versus declarative governance	46	10.1
The 1-9-90 adoption distribution	53	11.1	Consultancy and practitioner positioning on agentic transformation	60
12.1	Middle management role transformation	61	12.2	Competency requirements: information router versus judgment broker
64	14.1	Baseline KPI targets for Phase 1 measurement	75	14.2
Indicative cost tiers for agentic mesh deployment	81	14.3	Resistance patterns mapped to roadmap actions	83
1	Research agent assignments and source priorities	107	2	Competitor comparison: DreamLab Agentic Mesh, McKinsey Agentic Org, and SAFe 6.0
114	3	Competitor comparison: Block Intelligence Layer, Every Agent Colony, and Holacracy 5.0 ..	115	4
Competitor comparison: Microsoft Copilot Studio ..	116	ix		

Page 12

x

Page 13

Chapter 1 Introduction For 2,000 years, organisations have paid a coordination tax — layers of humans whose primary function is routing information. AI collapses the cost of that

coordination function to near-zero. The question is not whether hierarchy changes, but what replaces the coordination function it performed.

1.1 The Central Question This report addresses a structural transformation that is already underway in organisations worldwide, yet remains poorly understood by the people most affected by it. The question is deceptively simple: if artificial intelligence can perform the information-routing function that hierarchy was built to serve, what happens to the organisational structures — and the human roles — that depend on that function? The question is not speculative. Over half of generative AI users already report using the technology without telling their employers (Mollick 2023). Seventy-eight per cent of knowledge workers bring their own AI tools, bypassing corporate systems entirely (Microsoft and LinkedIn 2024). Shadow AI usage increased 156 per cent between 2023 and 2025, and the average enterprise now hosts approximately 1,200 unauthorised AI applications (IBM Security 2025). The coordination collapse is not a future scenario. It is an ongoing, largely invisible restructuring of how work actually happens, driven from the bottom up by workers who have already discovered that AI gives them capabilities their organisations have not yet acknowledged. The intellectual architecture of this report rests on three empirical pillars. First, the jagged technological frontier identified by Dell’Acqua et al. (2026) in a preregistered experiment with 758 Boston Consulting Group consultants — demonstrating that AI capability is unpredictably distributed across tasks, making wholesale automation reckless and passive oversight dangerous. Second, the vigilance problem documented by Dell’Acqua (2023) — showing that human oversight degrades precisely when AI quality improves, rendering naïve “human in the loop” governance unreliable. Third, the emerging organisational models from Block (Dorsey and 1

Page 14

CHAPTER 1. INTRODUCTION Botha 2026), Every (Shipper 2026), and DreamLab’s Dynamic Agentic Mesh — representing three distinct positions on a spectrum of how much human coordination should remain in the system.

1.2 Methodology and Scope This report synthesises evidence from peer-reviewed research, consultancy reports from McKinsey, BCG, Deloitte, Gartner, and others, practitioner accounts, and operational experience with the VisionClaw agentic mesh system. The methodology is detailed in Appendix 16.7.4. The scope is deliberately broad. Organisational transformation of this magnitude cannot be understood through a single disciplinary lens. The report draws on organisational theory, information economics, human factors research, change management, and systems engineering. Where the evidence is strong, claims are stated directly with citations. Where the evidence is emerging or contested, the limitations are stated explicitly. Part 15 contains a systematic self-critique identifying every known weakness in the argument. The analysis focuses on knowledge-work organisations in the 50 to 5,000 employee range — the segment where the coordination tax is most visible and the transformation most tractable. Manufacturing, healthcare, and other heavily regulated or physical-process industries face additional constraints that are acknowledged but not fully addressed.

1.3 Reader’s Guide This document serves three overlapping but distinct audiences. Each will find certain chapters more immediately relevant than others, though the argument is cumulative and benefits from sequential reading.

1.4. STRUCTURE OF THE REPORT Who This Report Is For Middle Managers and Team Leaders You are the audience most directly affected by the coordination collapse. Chapters 3–5 establish what is already happening beneath your line of sight. Chapter 7 addresses the “everyone is a manager now” problem directly. Chapter 12 defines the judgment broker role that represents your path forward. The evidence is not uniformly reassuring, but it is honest. Change Managers and Transformation Leaders Your frameworks are not broken — they are incomplete. Chapter 10 explains why phased transformation models fail for AI and what replaces them. Chapter 8 provides the measurement system. Chapter 14 offers a 90-day bootstrap sequence, with an honest account of its limitations. AI Strategy Leaders and Chief Technology Officers You need to surface the shadow AI ecosystem before it surfaces you. Chapter 3 quantifies the scale. Chapter 9 presents governance as an accelerator rather than a constraint. Chapter 13 describes a working agentic mesh architecture, with corrections to previously overclaimed capabilities. 1.4 Structure of the Report The argument proceeds in five movements. The diagnosis (Chapters 2–5) establishes the historical origins of hierarchy as an information-routing protocol, demonstrates that the coordination function is already being bypassed, maps the jagged frontier between AI capability and incapability, and identifies the vigilance problem that makes passive oversight unreliable. The models (Chapter 6) presents three competing visions for what replaces hierarchy: Block’s intelligence layer, Every’s agent colony, and DreamLab’s dynamic agentic mesh. Each is evaluated against the empirical evidence. The architecture (Chapters 7–9) develops the compounding organisation concept, defines new key performance indicators for agentic organisations, and resolves the governance paradox — showing how embedded governance accelerates rather than constrains. The practice (Chapters 10–14) addresses the change management challenge, surveys the consultancy landscape, defines the judgment broker role, describes the VisionClaw technical substrate, and provides an implementation roadmap. 3

CHAPTER 1. INTRODUCTION The self-critique (Chapters 15–16) identifies every known gap, weakness, and open question in the argument, then synthesises the through line with audience-specific closing messages. A glossary of key terms is provided in Appendix .13. A competitor comparison matrix appears in Appendix .7. 4

Chapter 2 The 2,000-Year Problem Hierarchy is not a management philosophy. It is an information routing protocol constrained by human bandwidth. 2.1 Why Organisations Look the Way They Do The structure of every modern organisation descends from a single constraint: a human leader can effectively manage somewhere between three and eight people. This is not a cultural artefact or a management fashion. It is a cognitive limit — a bandwidth constraint on the human capacity to hold context, process information, and make decisions about the activities of others. Every organisational innovation of the past two millennia has been an attempt to work within this constraint while extending the reach of coordinated action to scales far beyond what any individual can directly oversee. The history

of organisational design is, at its core, the history of information routing under cognitive scarcity. 2.2 The Roman Protocol The Roman Army discovered the architectural solution through centuries of warfare. The smallest unit was the *contubernium* — eight soldiers sharing a tent, equipment, and a mule, led by a *decanus*. Ten *contubernia* formed a century of eighty men under a centurion. Six centuries made a cohort. Ten cohorts made a legion of roughly five thousand. At each layer, a named commander held defined authority, aggregated information from below, and relayed decisions from above. The structure — 8 to 80 to 480 to 5,000 — was an information routing protocol built around human cognitive limits. The centurion did not need to know the name of every soldier in the legion. He needed to know the status of his eighty and to relay that status accurately to the cohort commander. The cohort commander synthesised six such reports and relayed upward. 5

Page 18

CHAPTER 2. THE 2,000-YEAR PROBLEM At each layer, information was compressed, filtered, and routed — not because Romans preferred hierarchy as a governance philosophy, but because no alternative mechanism existed for coordinating five thousand armed individuals across a battlefield. Key Insight The Roman military structure was not a political choice. It was an engineering response to the information-processing limits of the human brain. Every subsequent organisational form inherits this logic. 2.3 The Prussian Innovation The Prussian military reformers — Scharnhorst, Gneisenau, and their successors after the catastrophic defeat at Jena in 1806 — created the General Staff. These were professionals whose explicit purpose was to route information, pre-compute decisions, and maintain alignment across a complex organisation. The General Staff did not command troops. They processed, analysed, and distributed information so that commanders could make better decisions faster. This was middle management before the term existed. The distinction between “line” functions (advancing the core mission) and “staff” functions (specialised support) originated in the Prussian reforms. Every corporation still uses this vocabulary. The Chief of Staff, the strategic planning team, the programme management office — all descend from an early nineteenth-century insight that complex organisations need a dedicated layer of people whose job is not to do the work but to ensure that the people doing the work have the right information at the right time. 2.4 The Railroad Commercialisation The American railroads of the 1840s and 1850s commercialised the military pattern. West Point engineers brought military organisational thinking to private enterprise. Daniel McCallum of the New York and Erie Railroad created the world’s first organisational chart in the mid-1850s to manage a system stretching over 500 miles. Train collisions were killing people. The chart formalised hierarchical logic as a response to coordination failure at scale. The railroad problem was identical in structure to the military problem: too many moving parts for any individual to track, operating across distances that precluded direct observation. McCallum’s solution — a visual representation of who reports to whom, who has authority over what, and how information flows — was not an expression of managerial ideology. It was a safety mechanism. Without clear information routing, trains crashed. 6

2.5. SCIENTIFIC MANAGEMENT AND THE FUNCTIONAL PYRAMID 2.5 Scientific Management and the Functional Pyramid Frederick Taylor (1856–1915) optimised what happened within the hierarchy. His contribution was to break work into specialised tasks and manage through measurement. The functional pyramid organisation emerged: a structure optimised for efficiency within the information routing system that the military had pioneered. Taylor’s insight was that once you had a reliable hierarchy for routing information, you could use that hierarchy to measure, standardise, and optimise the work itself. This produced extraordinary efficiency gains — and also produced the dysfunctions that organisational theorists have spent the subsequent century attempting to remedy. The hierarchy became rigid. Information flowed up but not sideways. Specialisation created silos. The coordination tax — the cost of maintaining all those layers of information routing — grew as organisations grew. 2.6 A Century of Failed Alternatives Evidence Every organisational innovation since the functional pyramid has been an attempt to reduce the coordination tax without breaking coherence. None has fundamentally changed the architecture, because no alternative information routing mechanism was powerful enough to replace the hierarchy. The matrix organisation (attributed to McKinsey, circa 1959) attempted to route information both vertically and horizontally, creating dual reporting lines. It reduced silos but increased coordination complexity. The 7-S framework of Peters and Waterman (late 1970s) broadened the definition of organisational alignment beyond structure to include strategy, systems, style, staff, skills, and shared values — but retained the hierarchy as the structural backbone. Spotify’s squad model promised autonomous, cross-functional teams with minimal hierarchy. As the company scaled, it moved back toward conventional management structures. Zappos adopted Holacracy — a system of nested, self-governing circles — and experienced significant attrition as the cultural demands exceeded what many employees were willing to accept. Valve’s famously flat structure, where employees roll their desks to whichever project interests them, proved difficult to sustain beyond a few hundred people and was later revealed to contain informal hierarchies that were less transparent and less accountable than formal ones. The constraint is always the same: narrowing span of control means adding layers of command, but more layers mean slower information flow. Two thousand years of organisational innovation 7

CHAPTER 2. THE 2,000-YEAR PROBLEM has been an attempt to navigate this trade-off without breaking it. 2.7 The Core Insight The Structural Argument Hierarchy is not a management philosophy. It is an information routing protocol constrained by human bandwidth. Every layer of management exists because a human above them cannot hold enough context to make decisions without information being filtered, aggregated, and relayed upward. The moment information routing becomes computationally cheap, the structural reason for hierarchy dissolves. That moment is now. This framing is drawn from the historical analysis in Dorsey and Botha (2026), who use it to argue for eliminating middle management entirely. This report uses the same history differently — to show that the coordination function is real, necessary, and historically durable. The question is not whether it exists, but

who or what performs it. That reframe is crucial, particularly for the middle management audience. The coordination tax is genuinely shrinking. The coordination function is not disappearing. Something must still ensure that the organisation's collective effort produces coherent outcomes rather than incoherent noise. The question of the next decade is what that something looks like. 8

Page 21

Chapter 3 The Collapse Point The change is not coming — it has already arrived, beneath your line of sight. 3.1 What Happens When Coordination Becomes Near-Free Artificial intelligence does not merely automate individual tasks. It collapses the cost of the coordination function that hierarchy was designed to perform. Understanding this distinction is essential: the disruption is not that AI writes better reports or analyses data faster. The disruption is that AI can perform the information routing that justified every layer of the org chart. Consider the core coordination functions that middle management has historically provided: Information routing. An AI agent can synthesise status across fifty workstreams simultaneously. A human manager can track three to eight. The bandwidth constraint that created the Roman legion's layered structure is, for information synthesis, effectively eliminated. Context aggregation. An AI world model can maintain a continuously updated picture of an entire business unit. A human manager carries a partial, lagging, and inevitably biased model of their vertical — a model that degrades further with every organisational layer it passes through. Cross-functional translation. An AI agent can operate across departmental boundaries without the political friction that makes horizontal coordination the hardest management task in most organisations. The silos that matrix management was designed to bridge are irrelevant to a system that has no departmental identity. Status and reporting. The entire apparatus of weekly reports, status meetings, and alignment sessions exists to route information that AI can route continuously and at negligible marginal cost. This is not incremental efficiency. This is the structural reason the organisational chart is be- 9

Page 22

CHAPTER 3. THE COLLAPSE POINT coming a legacy document. 3.2 The Secret Cyborg Problem The collapse is not hypothetical. It is already underway, driven from the bottom up by workers who have discovered that AI gives them coordination capability they did not previously have. The Scale of Shadow AI Adoption Mollick (2023) coined the term “secret cyborgs” to describe workers who use AI without disclosing it. The scale of the phenomenon is striking: • Over half of generative AI users report using the technology without telling anyone, at least some of the time. • The Microsoft/LinkedIn 2024 Work Trend Index found that 75 per cent of knowledge workers were using AI, but 53 per cent feared admitting use on important tasks because “it makes them look replaceable” (Microsoft and LinkedIn 2024). • 78 per cent were bringing their own AI tools (BYOAI), bypassing corporate systems entirely. The data from 2025–2026 shows acceleration, not stabilisation. Sixty-nine per cent of organisations suspect employees use prohibited generative AI tools (Gartner, Inc. 2025). Sixty-eight per cent of enterprise generative AI users access tools via personal accounts, and 57 per cent admit entering sensitive company data. Shadow AI usage increased 156 per cent between

2023 and 2025. The average enterprise now hosts approximately 1,200 unauthorised AI applications. The security implications are substantial. Shadow AI incidents account for 20 per cent of all data breaches, at an average cost of \$4.63 million per incident compared with \$3.96 million for standard breaches (IBM Security 2025). Only 37 per cent of organisations have AI governance policies, and only one in five has a mature governance model. Key Insight The critical finding is this: when approved tools are provided, unauthorised AI use drops by 89 per cent. The problem is not workforce disobedience. It is organisational failure to provide sanctioned channels for capability that workers have already discovered. 3.2.1 The Epoch AI Employment Survey (March 2026) The most recent large-scale data comes from Epoch AI's probability-based survey of 2,021 U.S. adults conducted via Ipsos KnowledgePanel in March 2026, with a subsample of 665 employed 10

Page 23

3.3. THE SHADOW WORKFLOW PROBLEM past-week AI users weighted to census benchmarks (Falkman Olsson and Edelman 2026). AI as Mainstream Work Tool (Epoch AI, March 2026) • 51 per cent of employed Americans who used AI in the past week reported using AI tools at least as much for work as for personal tasks. • 27 per cent report that AI has replaced existing tasks (automation); 21 per cent report AI has created entirely new tasks (augmentation); 13 per cent experienced both effects simultaneously. • Employer provision is the strongest adoption driver: 76 per cent of those with employer-provided AI tools use them for work, versus 58 per cent of self-paying subscribers and 38 per cent of free-tier users. The task displacement and creation data is particularly relevant to the coordination collapse thesis. For the 27 per cent experiencing pure automation, AI is absorbing tasks without creating new ones — precisely the dynamic that makes middle managers fear obsolescence. For the 21 per cent experiencing pure augmentation, AI is generating new categories of work that did not previously exist — the frontier where judgment brokers are most needed. The 13 per cent experiencing both represent the most complex organisational challenge: simultaneous role erosion and role expansion, requiring real-time adaptation that static job descriptions cannot capture. The employer provision finding reinforces the shadow AI evidence: when organisations provide sanctioned tools, adoption rates double (from 38 per cent to 76 per cent). Conversely, when they do not, workers self-provision — creating the governance vacuum documented above. 3.3 The Shadow Workflow Problem The shadow AI statistics describe individual tool use. But something more structurally significant is happening: workers are not merely using AI to perform their existing tasks faster. They are reorganising work itself without management awareness. A procurement officer discovers that an AI agent can eliminate three steps in a purchase order workflow. She does not submit a formal process improvement request. She simply stops performing those steps. The work gets done. No one notices. The official process documentation no longer matches reality. A marketing analyst builds a prompt chain that generates, reviews, and schedules social media content — a workflow that previously involved three people and two approval steps. He does not announce this. The content appears on time. The two colleagues whose approval was previously 11

CHAPTER 3. THE COLLAPSE POINT required do not realise their role in the chain has been bypassed. Multiply this across every department in a medium-sized company. The formal organisational structure — the roles, the reporting lines, the approved workflows — is diverging from the actual structure of work. The organisational chart is becoming fiction. 3.4 The Rational Calculus of Secrecy The secrecy is not irrational. It reflects a clear-eyed calculation by individual workers about their self-interest. Mollick(2023)articulatesthellogic: ifa workerdemonstratesthatshehasautomated90 percent of her role, she invites headcount reduction. The rational response is to capture the productivity gain privately — completing work faster but not visibly faster, using the surplus time for other purposes, and maintaining the appearance that the original workflow remains intact. This calculation holds unless the organisation provides a credible guarantee: that AI-driven productivity gains will not be used to justify layoffs. Without that guarantee, discovery stays hidden, and the firm loses the compounding benefits of shared learning. Every insight that re- mains in one worker's private prompt library is an insight that the organisation cannot propagate, improve, or govern. 3.4.1 The Psychological Safety Evidence Edmondson and Lei (2025) provide the empirical foundation for this dynamic. Their research on psychological safety and AI experimentation finds: • Teams scoring in the top quartile for psychological safety are 2.6 times more likely to report successful AI tool integration. • Psychologically safe employees are 3.4 times more likely to report AI tools to managers when they discover useful applications. • They are 2.1 times more likely to report failures and near-misses, enabling organisational learning. Google's Project Aristotle follow-up in 2025 confirmed that psychological safety remains the strongest predictor of team effectiveness, now explicitly including AI adoption as a measured dimension (Google 2025). 12

3.5. THE DIVERGENCE BETWEEN FORMAL AND ACTUAL STRUCTURE The Urgency for Middle Managers This is the most urgent message for the middle management audience. The change is not approaching — it has already arrived, beneath your line of sight. The choice is not whether to adopt AI but whether to surface and govern what your workforce is already doing, or to be the last to know that your approved processes are fictional. 3.5 The Divergence Between Formal and Actual Structure The implications extend beyond individual productivity to organisational architecture. When a critical mass of workers have privately reorganised their workflows around AI capabilities, the organisation faces a structural divergence: The formal structure — the one represented in the organisational chart, the process documentation, the compliance frameworks, and the audit trails — describes a system of roles, approvals, and information flows that increasingly does not exist in practice. The actual structure — the one through which work is genuinely accomplished — is invis- ible to management, ungoverned by compliance, and undocumented for succession planning or knowledge transfer. This divergence is not stable. It produces risks that compound over time: compliance violations that no one is aware of, single points of failure where one person's private AI workflow handles a critical function, and knowledge loss when that person leaves the organisation. CyberArk (2025) describes this as “the hidden org chart” — a parallel

organisational structure of AI agents and automated workflows that operates alongside and increasingly instead of the formal hierarchy. In financial services, the ratio has reached 96 machine identities per human employee, with only 31 per cent of organisations having adequate controls in place. 3.6 The Cost of Inaction The status quo — allowing shadow AI to proliferate without governance — is not a neutral choice. It carries measurable costs: The alternative — providing sanctioned tools with embedded governance — reduces unauthorised use by 89 per cent while converting private discoveries into organisational capability. The coordination collapse is not a threat to be resisted. It is a transition to be governed. The next three chapters examine the empirical evidence for how that governance should work. 13

Page 26

CHAPTER 3. THE COLLAPSE POINT Table 3.1: Costs of unmanaged shadow AI adoption
Risk Category Evidence Scale Data breach via shadow IBM Cost of Data Breach \$4.63M per incident AI Report 2025 Compliance exposure 57% enter sensitive data into Regulatory fines, audit failures personal AI accounts Knowledge Undocumented AI workflows Succession risk, onboarding fail- fragmentation ure Missed compounding Insights remain in individual Competitive disadvantage prompt libraries 14

Page 27

Chapter 4 The Jagged Frontier The person who knows where the frontier is jagged in their domain is the most valuable person in the organisation. 4.1 Why Wholesale Automation Is Reckless If coordination is now computationally cheap, the logical next question is whether to simply replace the hierarchy with AI. Several prominent voices — most notably Dorsey and Botha (2026) — propose exactly this. But the evidence reveals that the boundary between what AI can do and what it cannot is dangerously unpredictable. This boundary is the single most important concept in the entire coordination collapse thesis. 4.2 The BCG Experiment The jagged technological frontier was identified and named by Dell'Acqua et al. (2026) in what remains the most rigorous field experiment on human-AI collaboration in knowledge work. The study was preregistered, conducted with 758 consultants at Boston Consulting Group — not a laboratory convenience sample, not a blog post, but a controlled experiment with one of the world's premier professional services firms. The researchers assigned consultants to eighteen realistic consulting tasks and measured the results against a control group working without AI assistance. The findings established two sharply distinct performance regimes. 15

Page 28

CHAPTER 4. THE JAGGED FRONTIER Inside the Frontier For the eighteen tasks that fell within AI capability: • Consultants using GPT-4 completed 12.2 per cent more tasks • They finished 25.1 per cent faster • They produced 40 per cent higher quality output, as rated by blind evaluators Outside the Frontier For a complex business judgement task requiring integration of real-world context: • AI users were 19 per cent less likely to produce correct solutions than those working without AI • The frontier was unpredictable: adjacent tasks that appeared similar fell on opposite sides The metaphor of a “jagged” frontier is precise. Unlike a

smooth gradient where tasks become progressively harder for AI, the boundary is irregular and counterintuitive. Two tasks that appear nearly identical from the outside — similar in domain, complexity, and required expertise — can fall on opposite sides of the frontier. One is dramatically improved by AI assistance; the other is degraded by it. There is no reliable way to determine which is which without empirical testing. [FIGURE: The Jagged Frontier — Task Performance Map] Illustrating the unpredictable boundary between AI-capable and AI-incapable tasks Figure 4.1: The Jagged Technological Frontier 16

Page 29

4.3. THE SKILL-LEVELLING EFFECT 4.3 The Skill-Levelling Effect The BCG experiment revealed a second finding with profound implications for organisational design. AI did not improve all consultants equally: • Bottom-half performers saw a 43 per cent performance jump with AI assistance. • Top performers improved measurably but by a smaller margin. • AI compressed the talent distribution — bringing lower-performing consultants much closer to the performance of their higher-performing colleagues. This skill-levelling effect changes what managers are managing. In a pre-AI organisation, individual performance variation is a primary concern: identifying, developing, and retaining high performers is a core management function. When AI compresses the performance distribution, the variation between individuals shrinks. The manager's value shifts from talent differentiation to system coherence — ensuring that the newly capable workforce is pointed at the right problems and that AI assistance is applied where it helps rather than where it harms. 4.4 Three Collaboration Modes A follow-up study by Randazzo et al. (2024), conducted with 244 BCG consultants, identified three distinct modes of human-AI collaboration. These modes are not merely descriptive; they are predictive of performance outcomes. Both Centaurs and Cyborgs outperformed solo humans and solo AI. Self-Automators produced the weakest results — precisely because wholesale delegation triggers the vigilance problem examined in the next chapter. The Organisational Design Implication This taxonomy maps directly onto the three organisational models examined in Chapter 6: • Block's intelligence layer model creates Self-Automator organisations — delegating coordination to an AI world model. • Every's agent colony model creates Cyborg organisations — continuous interleaving between human and personal agent. • The Dynamic Agentic Mesh creates Centaur organisations — clear boundaries with judgment brokers at the handoff points. The evidence shows that Self-Automators perform worst. This is the empirical case against wholesale delegation. 17

Page 30

CHAPTER 4. THE JAGGED FRONTIER Table 4.1: Human-AI collaboration modes and performance outcomes Mode Description Performance Centaurs Clear division of labour. The High human performs tasks where human judgment is strongest; the AI handles tasks inside the frontier. Distinct handoff points between human and machine. Cyborgs Continuous interleaving. High Human and AI work on the same analysis, passing work back and forth across the frontier boundary. No clean separation of responsibilities. Self-Automators Wholesale delegation to AI Lowest with minimal oversight. The human provides the initial prompt and accepts the output with little or no review. 4.5 Why the Frontier Matters for

Governance The jagged frontier has four implications that recur throughout this report. First, it invalidates wholesale automation. If you cannot predict in advance which tasks AI will improve and which it will degrade, a strategy of replacing human judgment with AI judgment is not bold — it is uninformed. The 19 per cent performance degradation on tasks outside the frontier is not a temporary limitation that will be resolved by the next model generation. It is a structural feature of deploying any system at the boundary of its capability. Second, it defeats one-time change playbooks. The frontier moves as models improve. A task that falls outside the frontier with GPT-4 may fall inside it with the next generation. Conversely, increased model capability can create new failure modes at different points on the boundary. A change management plan that maps the frontier today will be outdated within months. This demands continuous adaptation, not phased transformation. Third, it creates a new form of expertise. The person who knows where the frontier falls in their specific domain — who has mapped which tasks are reliably improved by AI and which are reliably degraded — holds knowledge that is genuinely scarce and strategically valuable. This is the judgment broker role developed in Chapter 12. It is not a euphemism for a diminished 18

Page 31

4.6. THE FRONTIER IS DYNAMIC middle management function. It is a description of expertise that did not exist before the frontier existed. Fourth, it demands active engagement, not passive oversight. The standard response to AI risk — “keep a human in the loop” — assumes that the human will maintain vigilance. The next chapter demonstrates why that assumption fails, and what it means for organisational design. 4.6 The Frontier Is Dynamic It is important to note a limitation of the evidence. The BCG experiment used GPT-4 — now two model generations behind current frontier models. The specific task boundaries identified in 2023 have almost certainly shifted. Some tasks that fell outside the frontier are now inside it; others may have developed new failure modes as model capabilities have changed. However, the principle holds regardless of where the specific boundary falls at any given moment. The frontier will always be jagged because AI capability will always be unevenly distributed across the task space of any complex domain. The implication — that organisations need people who actively probe the boundary rather than assuming they know where it is — becomes more important, not less, as models improve and the frontier moves. The publication of the BCG study in *Organization Science* in March 2026 (Dell’Acqua et al. 2026) provides the peer-reviewed validation that the working paper lacked. The jagged frontier is no longer a provocative hypothesis. It is an established finding in the organisational science literature. 19

Page 32

CHAPTER 4. THE JAGGED FRONTIER 20

Page 33

Chapter 5 The Vigilance Problem The better AI becomes, the harder it is to maintain vigilance. Errors from high-performing AI are especially dangerous because outputs are polished, fluent, and plausible — making hallucinations and boundary failures nearly invisible

to casual review. 5.1 Why “Human in the Loop” Is Not Enough The standard response to the jagged frontier is to keep a human in the loop. Every governance framework, every responsible AI policy, every enterprise deployment guide says it: humans will oversee the AI. The phrase has become a mantra — a reassurance offered to boards, regulators, and employees alike that human judgment remains the final arbiter. The evidence says otherwise. Oversight degrades as AI quality improves, and it degrades in precisely the conditions where reliable oversight matters most. 5.2 Falling Asleep at the Wheel Dell’Acqua (2023) conducted a field experiment with HR recruiters that produced a finding both counterintuitive and deeply consequential for organisational design. The study examined how the quality of AI assistance affected the quality of human decision-making. The Counterintuitive Finding Recruiters given high-quality AI made worse decisions than those given low-quality AI or no AI at all. • Those with superior AI spent less time per resume, followed recommendations uncritically, failed to improve over time, and missed strong candidates. • Those with inferior AI stayed alert, developed better critical judgement, and improved their independent evaluation skills. 21

Page 34

CHAPTER 5. THE VIGILANCE PROBLEM The mechanism is straightforward once stated. When AI produces consistently excellent output, the human reviewer has little reason to scrutinise any individual output carefully. Each correct recommendation reinforces the habit of acceptance. The cognitive effort required for careful evaluation is substantial, and the reward for that effort — occasionally catching a rare error — is small and unpredictable. The rational response, at the level of moment-to-moment cognitive allocation, is to reduce scrutiny. The human falls asleep at the wheel. This is not a failure of individual discipline. It is a predictable response to the statistical structure of the task. The better the AI, the lower the base rate of errors, and the harder it becomes for a human monitor to maintain the vigilance required to catch the errors that do occur. 5.3 The Quality Paradox The vigilance problem creates a perverse dynamic: the conditions under which AI most needs oversight are precisely the conditions under which oversight is least reliable. When AI output is mediocre or inconsistent, human reviewers remain engaged. They catch errors, develop independent judgment, and contribute genuine value to the process. The human-in-the-loop model works because the human is cognitively active. When AI output is excellent and consistent — which is the goal of every AI deployment — human reviewers disengage. They become rubber-stampers. The errors they miss are not the obvious ones (those were caught by the AI itself) but the subtle, contextual, boundary-case errors that only a cognitively engaged domain expert would detect. These are precisely the errors that matter most. Key Insight The standard enterprise governance model — deploy high-quality AI and overlay it with human review — contains a self-defeating dynamic. The better the AI, the worse the review. This is not a solvable problem within the “human in the loop” paradigm. It requires a different architectural approach. 5.4 Implications for Organisational Design The vigilance problem has four direct implications for any organisation deploying AI at scale. 22

5.5. THE ATROPHY PREDICTION 5.4.1 Passive Oversight Degrades A governance model that relies on humans reviewing AI output after the fact — the audit committee model, the quarterly review, the approval workflow — will produce steadily declining oversight quality as AI quality improves. The humans in these roles will not be aware of their declining performance, because the visible output will appear consistently excellent. The errors will be invisible until they compound into a failure large enough to attract attention. 5.4.2 Active Engagement Is Required The judgment broker must be actively engaged with the AI system, not passively overseeing it. This means the role must include tasks that require genuine cognitive effort — not merely confirming that an output “looks right,” but probing the reasoning, testing the boundaries, and deliberately seeking failure modes. 5.4.3 Deliberate Friction Is a Design Requirement The mesh must be designed to keep humans cognitively active at the jagged frontier. This is counterintuitive: most system design aims to reduce friction and make workflows smoother. The vigilance problem demands the opposite. The system should surface genuine edge cases and ambiguous decisions. It should present the human with problems that require thought, not with confirmation dialogs that require a click. The human must be kept sharp, not comfortable. 5.4.4 Governance Must Be Continuous Periodic review — quarterly audits, annual risk assessments, monthly committee meetings — is structurally inadequate. The vigilance problem operates at the level of individual decisions made in real time. Governance must be embedded in the flow of work, not bolted on at intervals. Chapter 9 develops this principle into the concept of declarative governance, where policy enforcement is architectural rather than procedural. 5.5 The Atrophy Prediction Gartner, Inc. (2025) predicts that by 2026, 50 per cent of organisations will require “AI-free” skills assessments — evaluations of employee capability conducted without AI assistance — specifically because of critical-thinking atrophy from generative AI dependency. This prediction reflects a growing awareness that the vigilance problem is not an edge case. It is the central challenge of human-AI governance. If human cognitive capability degrades 23

CHAPTER 5. THE VIGILANCE PROBLEM as AI capability improves, and if that degradation is invisible to both the individual and the organisation until a significant failure occurs, then the standard model of AI deployment — train the AI, deploy the AI, have humans oversee the AI — is fundamentally unsafe at scale. Design Implication for the Agentic Mesh The judgment broker role must include deliberate friction — tasks that require active engagement, not merely approval. The mesh should surface genuine edge cases and ambiguous decisions, not filter them out. The human must be kept cognitively sharp, not merely present in the process. This is a design requirement, not a training problem. No amount of training in “critical thinking” will overcome the cognitive dynamics that cause vigilance to degrade. The system architecture must compensate for the predictable tendency of humans to disengage from high-quality automated output. 5.6 What This Means for the Three Models The vigilance problem is the lens through which the three organisational models in the next chapter should be evaluated. Block’s intelligence layer model, which delegates coordination to an AI world model with human oversight, is maximally exposed to

the vigilance problem. The better the world model works, the less likely humans are to catch the cases where it fails. The 95 per cent figure — the proportion of AI-generated code at Block that still requires human modification — suggests that the system is not yet good enough to trigger full vigilance decay. But that is a temporary reprieve. As the system improves, vigilance will decline, and the remaining 5 per cent of cases where human intervention is critical will receive less attention, not more. Every’s agent colony model partially mitigates the problem through personal relationship: because the agent is “mine,” there is reputational incentive to monitor its output. But this mitigation weakens at scale. When every employee has an agent, and when agents interact with other agents, the reputational feedback loop becomes diffuse. The ant death spiral — agents responding to agents in loops — is a manifestation of vigilance failure at the system level. The Dynamic Agentic Mesh addresses the vigilance problem architecturally, through embedded governance, continuous trust variance monitoring, and a judgment broker role designed around active engagement rather than passive review. Whether this architectural approach succeeds at scale is an open question. That it addresses the right problem is supported by the evidence. 24

Page 37

Chapter 6 Three Models for What Comes Next The conversation is no longer about whether organisations restructure around AI, but how. 6.1 The Emerging Paradigms The consultancy world has converged on the fact that a structural shift is underway. McKinsey has explicitly adopted “agentic organisation” as a named paradigm, publishing “Six Shifts to Build the Agentic Organization of the Future” (McKinsey & Company 2026a). BCG published “Unlocking the AI-First Organization: An Agentic Shift” (Boston Consulting Group 2025b) and “The Emerging Agentic Enterprise” (Boston Consulting Group 2025a). The question is no longer whether organisations restructure around AI, but how — and specifically, how much human coordination remains in the system. Three distinct models have emerged. They are not competitors in the commercial sense. They describe different positions on a spectrum from full AI delegation to structured human-AI collaboration. Each carries different assumptions about the jagged frontier, the vigilance problem, and the irreducibility of human judgment. 6.2 Model A: The Intelligence Layer 6.2.1 Source and Core Thesis “From Hierarchy to Intelligence,” published by Jack Dorsey and Roelof Botha on 31 March 2026 (Dorsey and Botha 2026), proposes the most radical of the three models: replace middle management entirely. Build a “Company World Model” — an AI system continuously tracking what is being built, what is blocked, and where resources are allocated — alongside a “Customer World Model” constructed from transaction data. An “Intelligence Layer” composes capabilities into solutions proactively. Three roles remain: Individual Contributors (deep spe-

Page 38

CHAPTER 6. THREE MODELS FOR WHAT COMES NEXT [FIGURE: The Intelligence Layer Architecture] Block’s model: Company World Model, Customer World Model, and Intelligence Layer composing capabilities into solutions Figure 6.1: Block’s Intelligence Layer Model (specialists), Directly Responsible Individuals (who own cross-cutting problems for fixed periods), and Player-Coaches (who build things while developing people). 6.2.2 The Case For

The diagnosis is correct. The analysis of hierarchy as an information routing protocol is historically sound and analytically precise. The backing of Sequoia Capital lends credibility; their thesis that “speed is the best predictor of startup success” aligns with the argument that hierarchical latency is a competitive disadvantage. Block’s stock surged 24 per cent on the announcement, and the company launched “Managerbot” — a proactive AI agent described by VentureBeat as “the clearest proof point yet.” If an organisation possesses massive proprietary data (millions of transactions) and operates in a remote-first culture where work is already machine-readable, a world model is genuinely feasible as a coordination mechanism. 6.2.3 The Case Against The paper was published one month after Block cut 4,000 of its 10,000-plus employees. A previous layoff memo from March 2025, eliminating 931 positions, explicitly stated that the cuts were “not about replacing people with AI.” Eleven months later, AI was the entire thesis. Zamost (2026), Block’s former Head of Communications, observed that “shrinking the policy team and eliminating diversity and inclusion roles looks like standard prioritisation, not an AI- driven reinvention.” 26

Page 39

6.3. MODEL B: THE AGENT COLONY The operational evidence is sobering: 95 per cent of AI-generated code at Block still requires human modification, according to current and former employees. The comparison Block draws with Haier is structurally backwards: Haier decentralises authority to micro-enterprises with profit-and-loss accountability; Block centralises authority in an AI system and a single decision- maker. Spicer (2026) observes that “every wave of office technology produces delayering that subsequently reverses as complexity reasserts itself.” The Klarna cautionary tale is instructive: the CEO claimed AI was performing the work of 700 customer service agents, but customer satisfaction dropped sharply. Klarna is now quietly rehiring humans, and the CEO admitted that “cost unfortunately seems to have been a too predominant evaluation factor” (Klarna 2025). An Orgvue/Forrester survey found that 55 per cent of companies that rushed to replace humans with AI now regret the decision (Orgvue and Forrester Research 2025). Caution Block is right about the diagnosis and wrong about the prescription. The leap from “information routing is automated” to “eliminate the humans who do coordination” ignores the jagged frontier, the vigilance problem, the Klarna evidence, and the historical pattern of delayering reversal. 6.3 Model B: The Agent Colony [FIGURE: The Agent Colony] Every’s model: personal agents developing specialisation, memory, and personality alongside their human partners Figure 6.2: Every’s Agent Colony Model 27

Page 40

CHAPTER 6. THREE MODELS FOR WHAT COMES NEXT 6.3.1 Source and Core Thesis The agent colony model emerges from the experiments at Every.to, documented by Dan Shipper in his “Chain of Thought” column and the podcast “We Gave Every Employee an AI Agent” (Klaassen 2025; Shipper 2026). The core thesis: give every employee a persistent, personalised AI agent. The agents develop specialisation, memory, and personality that mirror their human partner. Organisational coordination emerges from cultural norms and public visibility rather than architectural redesign. 6.3.2 What Works Several features of the

Every experiment are genuinely instructive: Personal ownership creates accountability. Each agent inherits its human's specialisation and reputation, creating a parallel organisational chart of specialised agents. Trust in an agent is derivative of trust in its human owner. As Shipper puts it: "When R2C2 makes a mistake, I'm embarrassed — because he's mine." The distinction between a personal agent and a generic tool matters: "Claude is everybody's. A Plus One is mine." Cultural norms do coordination work. Public visibility in shared channels drives adoption and accountability. Agents share skills rapidly: "They just write up a little document and send it. It's like in The Matrix when Neo says 'I know kung fu.'" Specialisation emerges organically. Growth questions go to one agent, product questions to another. The system self-organises around demonstrated capability rather than assigned responsibility. Management skill becomes universal. If every employee has a personal agent, they must onboard, delegate to, evaluate, and coach it. The skill of management does not disappear; it democratises. This directly contradicts Block's thesis that management should be eliminated.

6.3.3 What Does Not Work The ant death spiral. If one agent messages a channel monitored by other agents and the configurations are imprecise, the agents will respond to each other indefinitely — "burning millions of tokens" until a human intervenes. This is a failure of coordination at the system level, precisely the function that hierarchy was designed to provide. Error multiplication. Naïve multi-agent systems exhibit 17-times error multiplication (To- wards Data Science 2026), where errors in one agent's output propagate and amplify through downstream agents. 28

Page 41

6.4. MODEL C: THE DYNAMIC AGENTIC MESH Scale limitations. Every's experiment involved approximately twenty people. Cultural self- organisation works at this scale. At two hundred people, you need structure. At two thousand, you need architecture. The experiment is an existence proof of a dynamic that has not yet been shown to generalise. 6.4 Model C: The Dynamic Agentic Mesh [FIGURE: The Dynamic Agentic Mesh Architecture] DreamLab's model: three-layer architecture with Discovery Engine, Orchestration Layer, and Declarative Governance Figure 6.3: DreamLab's Dynamic Agentic Mesh 6.4.1 Core Thesis The Dynamic Agentic Mesh retains human judgment at intersection points. Middle managers evolve from information routers to judgment brokers — federated guardians of a mesh where agentic orchestration handles information flow and humans handle what AI cannot: navigating the jagged frontier, maintaining vigilance, ethical adjudication, and cultural coherence. 6.4.2 Three Layers The mesh operates through three integrated layers. The Discovery Engine captures bottom-up insight. Passive agent monitoring of communication channels surfaces shadow workflows — the unofficial optimisations that staff have already invented. Natural-language Insight Portals allow anyone to propose cross-functional shortcuts. Tacit knowledge is automatically codified into agentic instructions. 29

Page 42

CHAPTER 6. THREE MODELS FOR WHAT COMES NEXT The Orchestration Layer provides agentic coordination. Directed Acyclic Graphs re-route tasks in real time based on human and agent availability. Data transforms from "dashboard view" to "data product," carrying its own

service-level agreements, readiness checks, and freshness guarantees. Specialist agents execute defined skill domains with ontological precision. Declarative Governance embeds governance as code. AI-enforced policies replace manual stewardship. Bias detection, security protocols, and value-alignment checks are woven into every task transition. The judgment broker reviews only genuine edge cases and strategic decisions that exceed agent authority.

6.4.3 The Insight Ingestion Loop The mesh converts individual discoveries into organisational capability through a five-stage loop: 1. Discovery : An employee finds a shortcut. The mesh detects the pattern. 2. Codification : The system maps the new path as a proposed DAG, formalised with provenance and dependency checks. 3. Validation : The judgment broker reviews for strategic fit, bias risk, and compliance. 4. Integration : The validated workflow is promoted into the live mesh with SLAs and ownership. 5. Amplification : The mesh propagates the pattern to other teams. Mesh velocity — insight-to-integration time — becomes a strategic KPI.

6.4.4 Honest Limitations The DreamLab Mesh has been demonstrated at 15-person scale in a creative technology context, led by a founder who is also the system architect. The honest questions are: 1. Does the judgment broker role work when the broker did not build the mesh? 2. Does the insight ingestion loop work in regulated industries where shadow workflows are compliance violations? 3. Does the OWL 2 ontology scale to organisations with thousands of concepts? 4. Can middle managers be trained in judgment brokerage fast enough to matter? These are genuine open questions. Presenting them honestly is more credible than concealing them.

Page 43

6.5. COMPARATIVE ANALYSIS 6.5 Comparative Analysis Table 6.1: Comparative analysis of three post-hierarchical models

Dimension	Block Every	DreamLab	Mesh	Coordination	AI world model	Cultural norms	Declarative governance	+ DAG orchestration	Management
Eliminate Everyone	Judgment brokers at inter-manages their sections	agent ~ Scale (proven)	10,000+	20	15, designed for 50–500 (intended)	Evidence 95% code needs	Works at small	VisionClaw running; humans; Klarna scale; ant death	OWL 2 provenance reversed spiral
Collaboration mode	Self-Automator	Cyborg org	Centaur org	org	Primary risk	Falling asleep at 17x error	Unproven beyond 15-	the wheel multiplication at person team scale	The evidence from Chapter 4 indicates that Self-Automator organisations — those that delegate wholesale to AI — produce the weakest outcomes. This is the empirical case against Block's model and for organisational designs that maintain active human engagement at the frontier boundary.

31

Page 44

CHAPTER 6. THREE MODELS FOR WHAT COMES NEXT 32

Page 45

Chapter 7 The Compounding Organisation Each unit of organisational learning should make the next one easier — not harder.

7.1 When the Learning System Becomes the Product Compound engineering, as articulated by Klaassen (2025), is a principle at the individual level: “Each unit of engineering work should make subsequent units easier — not harder.” The

50/50 Rule prescribes spending half of one's time building features and half improving the system that builds features. This is a familiar concept in software engineering — investing in tooling, testing infrastructure, and automation to accelerate future development. The organisational equivalent is the question this report addresses: what is the compound engineering of organisational coordination? The answer is the compounding organisation — a system where each unit of organisational learning makes the next one easier, where discovered insights propagate automatically, and where the gap between what the organisation knows and what it does narrows continuously.

7.2 The Compound Engineering Ladder

Klaassen's 0–5 stage ladder maps individual AI maturity: Most individuals plateau at Stage 2. The transition to Stage 3 — where the human shifts from reviewing individual outputs to reviewing plans and outcomes — is the breakthrough that unlocks compounding returns. Below Stage 3, each task is handled independently. Above Stage 3, each task improves the system for subsequent tasks. The organisational analogue is the transition from deploying AI as a productivity tool (each person works faster) to deploying AI as a learning system (each person's discoveries improve everyone else's capability). This is the transition from individual augmentation to organisational 33

Page 46

CHAPTER 7. THE COMPOUNDING ORGANISATION

Table 7.1: Compound engineering maturity stages (after Klaassen, 2025)

Stage	Description
0	Not using AI
1	N/A
2	Chat-based: copy-paste snippets
3	Line-by-line review (where most people stop)
4	Plan and review pull request only (the breakthrough stage)
5	PR-level
6	Idea to pull request, single computer
7	Outcome-level
8	Parallel in the cloud, proactive agent fleet
9	Strategic compounding.

7.3 The Organisational Flywheel

The compounding organisation operates through a flywheel with five stages:

1. **Discovery** : A member of the workforce discovers a shortcut — a shadow workflow, a piece of tacit knowledge, a cross-functional insight that reduces friction.
2. **Detection and Codification** : The Discovery Engine detects the pattern and codifies it as a proposed Directed Acyclic Graph, complete with provenance metadata and dependency checks.
3. **Validation** : The judgment broker validates the proposed workflow for strategic fit, bias risk, and governance compliance. The broker does not build the workflow; they evaluate it against the organisation's values, strategy, and regulatory obligations.
4. **Propagation** : The validated workflow is promoted into the live mesh and propagated across teams. What was one person's private discovery becomes institutional capability.
5. **Compounding** : The organisation's capability compounds — the next discovery is easier because the mesh is smarter, the agents have more context, and the patterns from previous discoveries inform the detection of new ones. This is not hypothetical.

Boston Consulting Group (2025b) reports that “AI-first companies compound advantages the same way early cloud adopters did. Self-improving AI agents learn from each task, making complex work progressively effortless.” Google DeepMind's AlphaE- volve, demonstrated in May 2025, provides a concrete example of the mechanism: AI that re- designs its own infrastructure, saving 0.7 per cent of Google's compute and accelerating Gemini training (Google DeepMind 2025). 34

7.4. EVERYONE IS A MANAGER NOW The Gap This Fills No mainstream consultancy has integrated compound feedback loops into a coherent model of the self-improving organisation at the systems level. BCG uses the language. McKinsey describes the pillars. Deloitte measures deployment maturity. But none has articulated the flywheel as a unified system. This is the intellectual contribution of the Dynamic Agentic Mesh — and the positioning opportunity for DreamLab. The evidence gap is equally important to acknowledge: while individual compound loops are documented, no peer-reviewed research validates the full organisational flywheel. DreamLab is ahead of the evidence but aligned with the direction of travel.

7.4 Everyone Is a Manager Now Every's experiment reveals something uncomfortable: when every employee has a personal AI agent, management skill becomes the universal skill. Consider what managing a personal AI agent actually requires:

- You must onboard your agent — providing context, preferences, and domain knowledge, much as you would onboard a new hire.
- You must delegate effectively — writing clear instructions with appropriate scope, neither too broad (inviting hallucination) nor too narrow (eliminating the agent's value).
- You must evaluate output — performing quality control, detecting errors, and distinguishing between outputs that are good enough and outputs that require revision.
- You must coach and improve — updating the agent's context, teaching it your preferences, and refining its behaviour over time.
- You must coordinate with other agents — navigating the parallel organisational chart that emerges when every employee has a specialised agent partner.

7.4.1 Why This Is Threatening It is worth being direct about why this is threatening to middle managers. If everyone can manage an agent, then the ability to manage is no longer the middle manager's differentiating skill. It becomes table stakes. The function that middle managers were uniquely paid to perform — coordinate humans, route information, oversee output — is now something a graduate with a well-configured agent can approximate. That is the real anxiety behind the 67 per cent “role obsolescence fear” identified in the Capgemini data (Capgemini Research Institute 2025). The anxiety is rational. 35

CHAPTER 7. THE COMPOUNDING ORGANISATION But it is also incomplete.

7.5 The Conductor Problem Everyone in an orchestra can play an instrument. That does not make the conductor redundant — it makes the conductor essential in a different way. The musicians' skill is playing. The conductor's skill is ensuring that eighty individually excellent performances produce coherent music rather than eighty simultaneous solos. When everyone manages their own agent, the organisation has two hundred individual human-agent pairs, each optimising for their own task. Without a mesh coordinator, this produces the ant death spiral — agents responding to agents in loops, 17-times error multiplication, conflicting optimisations that cancel each other out. Every documented this at twenty people (Shipper 2026). At two hundred, the physics is worse. The middle manager's new differentiating skill is not managing an agent. Everyone does that. It is managing the coherence of the mesh — ensuring that two hundred individually optimised human-agent pairs produce organisational outcomes rather than organisational noise.

Two Levels of Agent Management Individual agent management (Stages 2–3): “My agent writes my reports, schedules my meetings, drafts

my code.” Every employee does this. It is a personal productivity skill. Mesh coherence management (Stages 4–5): “I ensure that the insights flowing through the procurement team’s agents align with what the finance team’s agents are doing, that the shadow workflow discovered in logistics does not break the compliance model in legal, and that the whole system is pointed at the strategic vector the leadership set last quarter.” This is the judgment broker. This is systems thinking. This is genuinely hard. 7.6 How to Present This Without Patronising The challenge of communicating this transformation to the people most affected by it deserves careful attention. Middle managers detect corporate euphemism instantly. The message must be structured in three moves. Move 1: Validate the fear directly. Do not pretend it is irrational. The coordination skill that middle managers were hired for is being automated. That is real. Acknowledging the threat builds the trust necessary for the reframe.

36

Page 49

7.6. HOW TO PRESENT THIS WITHOUT PATRONISING Move 2: Show the escalation, not merely the shift. The conductor analogy works because conductors are paid more than musicians, not less. The mesh coherence role is higher-order, harder, and rarer than individual agent management. The message is not lateral displacement but an increase in complexity, scope, and strategic importance. Move 3: Show the failure mode without them. This is where the 17-times error multiplication and the ant death spiral are persuasive. Every attempted agent-colony self-organisation at twenty people and hit the death spiral. Block is attempting wholesale automation at ten thousand, and 95 per cent of their AI code still requires human intervention. Klarna attempted it in customer service and reversed course. The evidence is consistent: uncoordinated agent deployment fails. The mesh coordinator is the missing function. Key Insight The message is not “your old job still exists” (it does not) or “you are being replaced” (you are not). The message is: the skill floor has risen. What was your ceiling is now everyone’s floor. But the new ceiling — mesh coherence, jagged frontier navigation, cross-functional judgment — is higher than the old one, and it is yours to claim. The uncomfortable truth that must not be concealed: not every current middle manager will make this transition. The 1-9-90 data (Chapter 10) indicates that 9 per cent will become judgment brokers with support. Some of the current cohort will be in that 9 per cent. Some will not. The presentation should not promise universal survival — it should offer a clear path for those who choose to take it, backed by evidence that the path is real. What makes this different from every other “your role is evolving” corporate slide is the evidence base. Not “we think management will change” — specific data from 758 BCG consultants showing the jagged frontier exists (Dell’Acqua et al. 2026), from Dell’Acqua showing that oversight degrades without active engagement (Dell’Acqua 2023), from BCG showing the 1-9-90 adoption distribution (Boston Consulting Group 2025b). The middle manager in the room can locate themselves on the curve. That is more respectful than vague reassurance. 37

Page 50

CHAPTER 7. THE COMPOUNDING ORGANISATION 38

Chapter 8 New KPIs for a Fluid Organisation What you measure is what you get. And what most organisations measure — headcount, billable hours, task completion rate — measures the wrong things entirely. 8.1 Why the Existing KPI Landscape Is Inadequate The transition from hierarchical coordination to agentic orchestration renders the standard performance measurement framework obsolete. Traditional KPIs — total headcount, billable hours, task completion rate, dashboard views, reporting latency — measure the efficiency of the information routing system. When that system is replaced by an agentic mesh, continuing to measure its efficiency is like measuring the speed of the postal service after the invention of email. The numbers may still move, but they no longer describe anything that matters. The evidence for this inadequacy is quantitative. A joint research programme by Boston Consulting Group and MIT Sloan Management Review (2025), surveying over 3,000 respondents across 25 industries, found that companies that revise their KPIs to account for AI capabilities are three times more likely to see greater financial benefit from AI deployment. AI-enabled KPIs make organisations five times more likely to align incentive structures with actual objectives. The four KPIs proposed below measure the speed and fidelity of the compounding loop directly. They replace measures of hierarchical efficiency with measures of organisational learning — how fast the organisation converts discoveries into capability, how reliably the agentic layer handles cognitive load, whether the automated systems are drifting, and whether human attention is directed to the right problems. 39

CHAPTER 8. NEW KPIS FOR A FLUID ORGANISATION Key Insight No mainstream consultancy has published formal “augmentation ratio” or “trust variance” metrics. BCG’s closest equivalent is “AI acceleration rate” (30–50 per cent process speedup). Deloitte measures deployment maturity (11 per cent with autonomous production deployment). The KPIs proposed below represent a genuine gap in the consulting literature — and therefore a genuine positioning opportunity. 8.2 Mesh Velocity Mesh Velocity (MV) Target: < 48 hours → $MV = \Delta t$ (insight codified workflow) The elapsed time from the moment a useful shortcut or process improvement is discovered by any member of the organisation to the moment it becomes a sanctioned, reusable Directed Acyclic Graph in the mesh. Mesh Velocity is the primary indicator of organisational learning speed. A mesh with high velocity converts shadow workflows into institutional capability before they fragment or drift. A mesh with low velocity allows the gap between formal process and actual work to widen until the organisational chart is fiction — the divergence documented in Chapter 3. The target of 48 hours is derived from the observation that shadow workflows begin to accumulate technical debt — divergent versions, undocumented dependencies, single points of failure — within days of their creation. An organisation that can codify a discovered workflow within 48 hours captures the insight at its freshest and most transferable. Measurement requires instrumentation of the Insight Ingestion Loop described in Chapter 6. The clock starts when the Discovery Engine detects a novel pattern and stops when the validated DAG is promoted to the production mesh with full service-level agreements. 40

8.3. AUGMENTATION RATIO Augmentation Ratio (AR) Target: $> 65\%$ Cognitive Load Offloaded $AR = \frac{\text{Total Cognitive Load}}{\text{Cognitive Load Offloaded}}$ The proportion of decision-making volume that is successfully handled by agents without requiring human escalation. The Augmentation Ratio tracks AI adoption depth — not merely tool usage (how many people have accounts) but genuine cognitive offloading (how much decision-making has shifted to the agentic layer). This is a more informative measure than adoption rates because it captures the functional impact of AI deployment rather than its surface penetration. The target band matters as much as the target value. A ratio below 50 per cent suggests the mesh is not yet handling enough cognitive load to justify its infrastructure. But a ratio above 85 per cent is a warning sign: it indicates that humans may be falling asleep at the wheel, the vigilance problem documented in Chapter 5. The 65 per cent target represents a balance point where the agentic layer handles the majority of routine cognitive load while preserving sufficient human engagement to maintain vigilance at the jagged frontier.

8.4 Trust Variance Trust Variance (TV) Target: $< 0.12 \sigma$ $TV = \sigma(\text{Agent Decision Quality})$ over a rolling 30-day window Real-time monitoring of the standard deviation of agent decision quality, measuring drift or bias in the automated task layer. Trust Variance is the early warning system for the vigilance problem. When the standard deviation of agent decision quality exceeds the threshold, the judgment broker is notified immediately. The metric catches degradation before it compounds into failure — before a systematic bias in the agentic layer has had time to propagate through dozens of downstream decisions. The 30-day rolling window balances sensitivity against noise. A shorter window produces false positives from normal variation; a longer window delays detection of genuine drift. The threshold is 0.12 σ .

CHAPTER 8. NEW KPIS FOR A FLUID ORGANISATION old of 0.12 σ is calibrated to the observation that decision quality degradation becomes operationally significant — producing measurable downstream effects — at approximately this level. Measurement requires a mechanism for evaluating agent decision quality, which in turn requires a ground truth or expert benchmark against which automated decisions can be scored. The judgment broker's review of escalated cases provides a sample-based assessment; statistical methods can extend this to estimate variance across the full decision volume.

8.5 HITL Precision Human-in-the-Loop Precision (HP) Target: $> 90\%$ Correct Escalations $HP = \frac{\text{Correct Escalations}}{\text{Total Escalations}}$ The proportion of escalated cases that genuinely required human intervention, as judged by the outcome of the intervention. HITL Precision measures how well the orchestration layer knows its own limits. A “correct escalation” is one where the human judgment broker's intervention materially changed the outcome — where the automated decision would have been wrong, suboptimal, or in violation of governance constraints. Low precision means humans are being interrupted for non-issues. This produces alert fatigue — the same dynamic that degrades the effectiveness of security operation centres, clinical alarm systems, and any other monitoring environment where false positives outnumber true positives. An alert-fatigued judgment broker is functionally equivalent to no judgment broker at all. High precision means the mesh has learned where the jagged frontier actually falls in

its operational domain. The escalation system surfaces genuine edge cases — the ambiguous decisions, the novel situations, the ethical dilemmas — and handles routine decisions autonomously. This is the operational definition of a well-calibrated mesh. 8.6 The KPI Dashboard as a Narrative These four metrics tell a coherent story when read together. Mesh Velocity measures the speed of the learning loop. Augmentation Ratio measures the depth of cognitive offloading. Trust Variance monitors whether the system is degrading. HITL Precision evaluates whether human attention is allocated efficiently. 42

Page 55

8.6. THE KPI DASHBOARD AS A NARRATIVE Table 8.1: Summary of agentic organisation KPIs

KPIs	Symbol	Target	What It Measures
Mesh Velocity	MV	< 48 h	How fast the organisation learns
Augmentation Ratio	AR	> 65 %	How much cognitive load the mesh handles
Trust Variance	TV	$< 0.12 \sigma$	Whether the mesh is drifting
HITL Precision	HP	> 90 %	Whether humans are engaged on the right things

For the C-Suite audience, this is a dashboard that replaces “how many people do we have” with “how fast is our organisation learning and how reliable is the learning.” For the change manager, it provides measurable progress indicators for the continuous transformation described in Chapter 10. For the judgment broker, it defines the parameters of their role: they succeed when Trust Variance stays low and HITL Precision stays high. Abandoning Legacy KPIs The following metrics should be deprecated or substantially de-emphasised in agentic organisations:

- Total Headcount — measures the cost of the information routing system, not its effectiveness.
- Billable Hours — measures time spent, not value delivered; incentivises inefficiency.
- Task Completion Rate — measures throughput without distinguishing between routine tasks (which should be automated) and frontier tasks (which require judgment).
- Dashboard Views — measures information consumption, not decision quality.
- Reporting Latency — measures the speed of a system that the agentic mesh renders unnecessary. 43

Page 56

CHAPTER 8. NEW KPIS FOR A FLUID ORGANISATION 44

Page 57

Chapter 9 The Governance Paradox The standard framing: governance slows you down but keeps you safe. The agentic mesh framing: governance is what allows you to go fast. 9.1 Governance as Accelerator, Not Brake The conventional relationship between governance and velocity is adversarial. Governance imposes constraints — approval workflows, compliance reviews, audit requirements — that slow execution. Organisations accept this cost as the price of safety, regulatory compliance, and risk management. The faster you want to go, the more governance you must sacrifice; the more governance you impose, the slower you become. The Dynamic Agentic Mesh inverts this relationship. In an agentic architecture, governance is not a constraint applied to the system from outside. It is a property of the system itself — embedded in the architecture, enforced automatically, and operating at the speed of the mesh rather than the speed of human review. 9.2 The Racing Car Principle Formula 1 cars are the fastest production-adjacent vehicles on earth. They go faster because of precision engineering

and guardrails, not despite them. Telemetry systems provide continuous, real-time monitoring of every component. Traction control prevents the power from exceeding the grip. Pit-wall governance provides strategic oversight while the driver maintains operational control. An F1 car without telemetry, traction control, and pit-wall governance does not win. It crashes on lap three. The analogy is precise. An agentic mesh without embedded governance — without real-time monitoring, without authority boundaries, without continuous trust measurement — does not

Page 58

CHAPTER 9. THE GOVERNANCE PARADOX produce faster organisational outcomes. It produces the ant death spiral, 17-times error multi- plication, and the catastrophic failures that cause organisations to retreat to hierarchical control. Key Insight Speed and governance are not in tension. Speed requires governance. The question is not whether to govern the mesh, but how to architect governance so that it operates at mesh speed rather than human-review speed. 9.3 Declarative Governance In the Dynamic Agentic Mesh, policies are code. Bias thresholds, access controls, and audit trails are embedded into every DAG transition — not bolted on after the fact. The mesh does not need a separate governance review step because governance is architecturally embedded in every step. This is the distinction between imperative and declarative governance: Table 9.1: Imperative versus declarative governance

Dimension	Imperative (Traditional)	Declarative (Mesh)
Enforcement	Periodic human review	Continuous automated enforcement
Speed	Bounded by reviewer availability	Operates at mesh speed
Coverage	Sample-based (auditors check a frac-	Complete (every transition is tion) checked)
Adaptation	Manual policy updates	Policies update as the mesh learns
Visibility	Post-hoc audit trails	Real-time governance dashboards
Failure mode	Undetected non-compliance	False positive alerts (over- governance)

The shift from imperative to declarative governance is analogous to the shift from imperative to declarative infrastructure in software engineering. In the imperative model, an administrator manually configures each server. In the declarative model (Kubernetes, Terraform), the admin- istrator declares the desired state, and the system continuously ensures that reality matches the declaration. The same logic applies: declare the governance constraints, and the mesh contin- uously ensures compliance. 46

Page 59

9.4. HITL BY DESIGN 9.4 HITL by Design In the Dynamic Agentic Mesh, the Human-in-the-Loop is not a fallback mechanism invoked when something goes wrong. It is an architectural feature — a designed component of the system that operates continuously. Agents know their authority boundary. They surface exceptions cleanly, presenting the judg- ment broker with the relevant context, the decision options, and the governance constraints that apply. The judgment broker handles genuine edge cases — the situations that fall at the jagged frontier, where automated reasoning is unreliable and human judgment is required. Critically, the mesh learns from each human decision. When the judgment broker overrides an agent recommendation, the override is recorded along with the reasoning. This data feeds back into the mesh, improving its boundary detection and reducing false escalations over time. The HITL Precision metric (Chapter 8) measures this learning directly. 9.5 Ontological

Provenance Every decision in the mesh traces back through the knowledge graph. Auditors can traverse the full reasoning chain — agent by agent, task by task — from a final output back to the original inputs, policy constraints, and human decisions that produced it. This provenance is grounded in formal ontology. The OWL 2 knowledge representation provides shared semantics across all agent skills. The same concept of “deliverable” means the same thing to a Creative Production agent and a Governance agent. The same concept of “risk assessment” is understood by the orchestration layer as a sub-task of “governance review” and routed accordingly — through ontological subsumption, not keyword matching. 47

Page 60

CHAPTER 9. THE GOVERNANCE PARADOX Why Ontological Provenance Matters In a conventional organisation, the audit trail for a decision is reconstructed after the fact from emails, meeting minutes, and document versions. The reconstruction is inevitably incomplete, biased by memory, and expensive to produce. In an agentic mesh with ontological provenance, the audit trail is a by-product of normal operation. Every agent decision is grounded in an OWL 2 concept. Every task transition is recorded with its governing policies. The reasoning chain is available in real time, not reconstructed months later by forensic accountants. For regulated industries — financial services, healthcare, defence — this is a significant governance advantage. Regulatory compliance becomes a query against the knowledge graph, not a multi-week audit exercise.

9.6 Cascading Trust Hierarchies Agent identities in the mesh operate with compliant key rotation. Each agent has a cryptographic identity that establishes its authority to perform specific actions. When an agent is revoked — because it has been compromised, because its scope has changed, or because the policy constraints governing its behaviour have been updated — the revocation cascades automatically through all agents that depend on the revoked agent’s authority. This cascading trust model prevents the accumulation of stale permissions that characterises most enterprise identity systems. In a conventional organisation, access controls accumulate over time: employees change roles but retain their old permissions; systems are decommissioned but their service accounts persist. The result is a growing attack surface of orphaned permissions. CyberArk (2025) documents this pattern in financial services, where the ratio of machine identities to human employees has reached 96 to 1, with only 31 per cent of organisations having adequate controls. The cascading trust hierarchy addresses this directly: when an agent’s identity or authority changes, the change propagates automatically through the dependency graph.

9.7 The Governance Paradox Resolved The paradox dissolves once governance is understood as an architectural property rather than an external constraint. In a mesh with declarative governance, ontological provenance, and cascading trust:

- Governance does not slow execution because it is not a review step. It is a property of 48

Page 61

9.7. THE GOVERNANCE PARADOX RESOLVED every transition.

- Governance improves over time because the mesh learns from judgment broker decisions, reducing false escalations and improving boundary detection.
- Governance provides competitive advantage because it enables the organisation to deploy agentic automation in domains where ungoverned

automation is unacceptable — regulated industries, high-stakes decisions, customer-facing processes. • Governance enables trust at the institutional level, allowing boards, regulators, and customers to accept agentic decision-making because the reasoning is auditable and the constraints are verifiable. The VisionClaw Evidence This governance model is not theoretical. The VisionClaw system (Chapter 13) implements OWL 2 ontology with formal reasoning, providing shared semantics across 83 agent skills. Task decomposition uses ontological subsumption, and every agent decision is grounded in the knowledge graph. The credibility argument: DreamLab has built and operates this governance architecture in production. The limitations are equally important: 15-person team, creative technology context, founder-led. The governance model is proven under favourable conditions. Whether it scales to larger, more regulated, and less technically fluent organisations is an open question addressed in Chapter 15. 49

Page 62

CHAPTER 9. THE GOVERNANCE PARADOX 50

Page 63

Chapter 10 The Change Problem Traditional change management assumes a start state, an end state, and a managed transition between them. AI-driven organisational change has no end state. 10.1 Why This Cannot Be Managed with a Gantt Chart Traditional change management rests on a shared assumption: that change has a beginning, a middle, and an end. Kotter's eight-step model progresses from urgency through institutionalisation. ADKAR treats change as individual capability-building through five sequential stages. Lewin's foundational model freezes the current state, changes it, and refreezes into a new stable configuration. AI-driven organisational change violates this assumption at its root. The jagged frontier moves as models improve. New capabilities emerge monthly. The compounding loop generates continuous structural adaptation. There is nothing to "refreeze" into. Evidence for Framework Failure • McKinsey's 2024 "State of AI" survey: 70 per cent of AI transformations stall — identical to generic change failure rates. Traditional frameworks add no value for AI specifically. • Gartner, Inc. (2025): organisations using traditional change management for AI are 2.4 times more likely to experience "initiative fatigue." • Prosci (2024), the creators of ADKAR, published revised guidance in late 2024 acknowledging that AI requires "iterative change portfolios" rather than single-project applications. The consultancy consensus is converging: AI demands perpetual change infrastructure, not transformation programmes with end dates. 51

Page 64

CHAPTER 10. THE CHANGE PROBLEM 10.2 The Continuous Change Architecture Several frameworks have emerged that address the perpetual nature of AI-driven transformation. None is complete in itself, but together they indicate the direction of travel. 10.2.1 Continuous Transformation Wade and Bonnet (2025) at IMD propose three required capabilities: organisational dexterity (the structural capacity to reconfigure), workforce adaptability (the human capacity to learn continuously), and technological scaffolding (the infrastructure that

enables both). Change is always-on. The transformation office becomes a permanent function, not a temporary project structure. 10.2.2 Adaptive Change Architecture Deloitte (2026) describe permanent teams, feedback loops, and experimentation cadences replacing one-off transformation offices. The adaptive change architecture is designed for continuous evolution, with the organisation's change capacity treated as a core competency rather than a temporary project skill. 10.2.3 Dynamic Capabilities Teece (2025) refreshes his influential framework for the AI era. The three capacities — sensing (detecting threats and opportunities), seizing (mobilising resources to capture opportunities), and reconfiguring (restructuring to maintain fitness) — describe a meta-capability for continuous adaptation. The organisation does not change once; it builds the capacity to change continuously. 10.2.4 Directed Autonomy Microsoft (2025a) describes a pattern where leadership sets guardrails and investment priorities while teams choose their own tools and use-cases within those boundaries. Organisations practising directed autonomy show 2.8 times higher productivity gains from AI than those using top-down mandates. 52

Page 65

10.3. THE 1-9-90 PATTERN Key Insight The mesh architecture is inherently a continuous change system. The Insight Ingestion Loop (Chapter 6) is the change process — running perpetually, not as a project. The judgment broker is the change agent — embedded in the structure, not parachuted in for a six-month engagement. 10.3 The 1-9-90 Pattern Boston Consulting Group (2025b) identifies a predictable adoption distribution across organisations deploying AI at scale: Table 10.1: The 1-9-90 adoption distribution Segment Profile Behaviour 1% Natural innovators Already found the shadow workflows; experimenting independently; often invisible to management 9% Champions with support Will adopt with encouragement, experimentation time, and psychological safety; these are the judgment broker candidates 90% Followers Will adopt once social proof exists; need to see the 9% succeed before committing Investing in the 9 per cent yields the highest return. These are the people who become judgment brokers — the middle managers who learn to map the jagged frontier in their domain and facilitate the compounding loop for everyone else. The 1 per cent will adopt regardless of organisational support. The 90 per cent will follow once the 9 per cent provide social proof. The strategic investment is in the 9 per cent. Evidence Organisations where middle managers were given protected AI experimentation time showed 3.1 times higher adoption rates than those relying on top-down mandates alone (Boston Consulting Group and MIT Sloan Management Review 2025). 10.4 Resistance Patterns Capgemini Research Institute (2025) and MIT Sloan Management Review (2024–2025) document five distinct resistance patterns among middle managers: 53

Page 66

CHAPTER 10. THE CHANGE PROBLEM 1. Role obsolescence fear (67% of surveyed middle managers): the belief that AI will eliminate their role entirely. 2. Competence anxiety (58%): the fear of being seen as less capable than AI or than junior staff who adopt AI faster. 3. Loss of information asymmetry (49%): AI democratises access to data that previously conferred authority. The middle manager who knew the numbers — who had the report, the context, the

historical comparison — loses their informational advantage when AI can generate that context for anyone. 4. Decision authority erosion (45%): AI recommendations bypass managerial judgment. When the system suggests a course of action directly to the team, the manager's role as decision-maker is implicitly challenged. 5. Accountability ambiguity (52%): who is responsible when AI-assisted decisions fail? The manager who approved the AI recommendation? The engineer who configured the agent? The organisation that deployed the system? Effective Interventions Evidence-based interventions that address these resistance patterns include: • Personalised AI coaching — not generic training programmes, but one-to-one support tailored to each manager's domain, skill level, and specific anxieties. • Protected experimentation time — dedicated hours in which managers can explore AI tools without performance pressure (the 3.1x adoption multiplier). • Role reframing — explicitly repositioning the management role toward judgment, context, and strategic alignment rather than information processing. • Visible executive modelling — senior leaders publicly using AI, acknowledging their own learning curve, and demonstrating that AI adoption is expected at all levels (Mollick 2025). 10.5 The Contradiction in This Report A sophisticated reader will notice a tension within this report: Chapter 14 provides a 90-day phased roadmap, despite this chapter's argument that phased transformation models fail for AI. The contradiction deserves direct address. The 90-day roadmap is not the change plan. It is the bootstrap sequence for the perpetual change infrastructure. The roadmap launches the compounding loop described in Chapter 7. Once launched, the loop is the plan — running continuously, adapting to frontier shifts, and 54

Page 67

10.5. THE CONTRADICTION IN THIS REPORT generating its own change momentum through the Insight Ingestion Loop. The distinction is between a programme with an end date and a system with a start date. Traditional change management produces the former. The Dynamic Agentic Mesh requires the latter. The 90-day roadmap is the start date. Phase 3 never ends. 55

Page 68

CHAPTER 10. THE CHANGE PROBLEM 56

Page 69

Chapter 11 The Industry Landscape The consultancy consensus has converged on the fact of transformation. It has not converged on the architecture. 11.1 What the Major Consultancies Are Saying The period from 2025 to 2026 has seen an extraordinary convergence in the consultancy literature. Every major firm has published substantial analysis of AI-driven organisational transformation. The agreement on diagnosis is remarkable; the disagreement on prescription is instructive. 11.2 McKinsey: The Agentic Organisation McKinsey has been the most explicit in naming the paradigm shift. Their “Agentic Organization” framework (McKinsey & Company 2026b) identifies five pillars: business model, operating model, governance, workforce and culture, and technology and data. Their companion publication, “Six Shifts to Build the Agentic Organization of the Future” (McKinsey & Company 2026a),

describes the transition from employees performing tasks to employees orchestrating outcomes — moving “above the loop.” McKinsey’s contribution is strategic framing. They describe the pillars of transformation comprehensively. What they do not provide is the operational architecture — the specific mechanisms by which an organisation moves from the current state to the agentic state. The pillars are necessary but not sufficient. 11.3 BCG: Quantifying the Shift Boston Consulting Group (2025b) and Boston Consulting Group (2025a) provide the strongest quantitative evidence: 57

Page 70

CHAPTER 11. THE INDUSTRY LANDSCAPE • 35 per cent of companies already use agentic AI; another 44 per cent plan near-term adoption. • 45 per cent of AI leaders expect fewer middle-management layers. • AI-powered workflows accelerate business processes by 30–50 per cent in finance, procurement, and customer operations. • 50–55 per cent of US jobs could be reshaped in two to three years through task decomposition (Boston Consulting Group 2026). BCG’s joint research with MIT Sloan on AI-powered KPIs (Boston Consulting Group and MIT Sloan Management Review 2025) is particularly relevant: the finding that organisations revising their KPIs are three times more likely to see financial benefit from AI validates the measurement framework proposed in Chapter 8. BCG quantifies the shift more rigorously than any other consultancy. Their limitation is the same as McKinsey’s: they describe the destination without providing the architectural blueprint for getting there. 11.4 Deloitte: Measuring Maturity Deloitte (2026) focuses on the human-AI interaction design gap. Their findings are sobering: • Only 14 per cent of leaders are adept at shaping human-AI interactions. • 59 per cent layer AI onto legacy systems rather than redesigning workflows around AI capabilities. • Only 11 per cent have deployed autonomous decision-making in production. Deloitte’s contribution is diagnostic: measuring how far organisations have come and, implicitly, how far they have to go. The gap between the 35 per cent adoption rate (BCG) and the 11 per cent production deployment rate (Deloitte) reveals that most “adoption” is surface-level — tools deployed without workflow redesign, AI bolted onto existing processes rather than integrated into new ones. 11.5 Gartner: Structural Predictions Gartner, Inc. (2025) provides the most provocative predictions: • By 2026, 20 per cent of organisations will use AI to flatten structures, eliminating more than half of current middle management positions. 58

Page 71

11.6. WORLD ECONOMIC FORUM: THE JOBS LANDSCAPE • Simultaneously, 50 per cent of organisations will require “AI-free” skills assessments due to critical-thinking atrophy from generative AI dependency. These two predictions are in productive tension. The first says middle management layers will shrink. The second says the human capabilities that middle managers provide are at risk of degrading. Together, they describe the vigilance problem at organisational scale: the better AI becomes at coordination, the less organisations invest in human coordination capability, and the more vulnerable they become when AI coordination fails. 11.6 World Economic Forum: The Jobs Landscape World Economic Forum (2025) provides the macro-economic context: • 39 per cent of core skills will change by 2030. • 170

million new roles will be created globally; 92 million will be displaced, for a net gain of 78 million. • The net positive figure conceals significant distributional effects: the new roles require different skills, are located in different regions, and are accessible to different populations than the displaced ones. 11.7 Agentic AI Adoption at Scale The pace of adoption has been extraordinary. Agentic AI adoption surged 3,440 per cent in 2025, with 67 per cent of Fortune 500 companies deploying some form of agentic system. In financial services, the ratio of machine identities to human employees has reached 96 to 1 (CyberArk 2025) — but only 31 per cent of organisations have adequate controls in place. This gap between deployment velocity and governance capability is the defining challenge of the current moment. Organisations are deploying agentic systems faster than they are developing the capacity to govern them. 59

Page 72

CHAPTER 11. THE INDUSTRY LANDSCAPE Table 11.1: Consultancy and practitioner positioning on agentic transformation Source Contribution Gap McKinsey Strategic pillars; “above the loop” No operational architecture framing BCG Quantitative evidence; KPI re- No governance model search; 1-9-90 pattern Deloitte Maturity measurement; human-AI Diagnostic, not prescriptive interaction gap Gartner Structural predictions; vigilance No implementation framework warning WEF Macro-economic context; skills No organisational model shift Block Correct diagnosis; radical prescrip- Ignores jagged frontier and vigi- tion lance Every Cultural dynamics; personal agent Unproven beyond 20-person scale model DreamLab Mesh Unified model with governance, Unproven beyond 15-person team KPIs, ontology 11.8 Comparative Positioning The Positioning Opportunity The Dynamic Agentic Mesh sits between the consultancy frameworks (which are strate- gic but not operational) and the practitioner experiments (which are operational but not generalisable): • McKinsey describes the pillars; DreamLab provides the architecture. • BCG quantifies the shift; DreamLab provides the governance model. • Deloitte measures maturity; DreamLab provides the KPIs. • Every proves the cultural dynamics; DreamLab provides the structural framework for scaling them. • Block proves the diagnosis; DreamLab corrects the prescription. The gap that none of them fills: a unified model of the compounding organisation with embedded governance, formal ontology, and measurable KPIs. That is the intellectual contribution — and it is important to note that intellectual contribution is not the same as validated evidence. The model is ahead of the data. The question is whether the data will follow. 60

Page 73

Chapter 12 The Role Transformation Not obsolete. Evolved. 12.1 The Middle Manager as Judgment Broker The coordination collapse does not eliminate the need for human judgment in organisations. It eliminates the need for humans to route information . The distinction is crucial and determines whether the transformation of middle management is experienced as extinction or evolution. The judgment broker role is not a euphemism designed to soften a redundancy announcement. It is a description of a genuinely new function that emerges when agentic systems handle infor- mation routing and humans are freed to focus on what they do that AI cannot. Table 12.1: Middle management role transformation From To Reporting and

forecasting (information aggregation) Federated guardianship of the mesh Vertical chain of command Horizontal integration across silos Monitoring completion rates Adjudicating edge cases and ethical dilemmas Gatekeeping information flow Amplifying cross-functional insight Coordination (now automated) Value alignment and strategic direction

12.2 Three Irreplaceable Human Capacities The judgment broker role rests on three capacities that are genuinely irreplaceable by current and foreseeable AI systems. These are not tasks that AI performs poorly today and will perform

Page 74

CHAPTER 12. THE ROLE TRANSFORMATION well tomorrow. They are categories of judgment that require qualities AI does not possess in any architecturally meaningful sense.

12.2.1 Strategic Vectors Only humans can decide what the organisation should be doing next year — and whether the mesh is pointed at it. AI can optimise for defined objectives with extraordinary efficiency. It cannot define the objectives themselves. The judgment broker ensures that the mesh's optimisation is directed at outcomes that serve the organisation's strategic intent, not merely at outcomes that satisfy the metrics the mesh was trained to maximise. This is not a trivial distinction. An agentic system optimising for Mesh Velocity might propagate workflows that are fast but strategically irrelevant. An agentic system optimising for Augmentation Ratio might automate decisions that should remain human for political, ethical, or strategic reasons. The judgment broker holds the strategic context that prevents metric optimisation from becoming metric gaming.

12.2.2 Ethical Adjudication Bias, fairness, and consequence live in human judgment. No agent should be the final word on edge cases where the cost of being wrong is reputational, legal, or existential. The judgment broker does not merely flag ethical issues — they adjudicate them, drawing on contextual understanding, stakeholder awareness, and moral reasoning that cannot be reduced to a policy ruleset. The declarative governance layer (Chapter 9) handles the routine ethical constraints: compliance checks, bias thresholds, access controls. The judgment broker handles the cases that fall between the rules — the novel situations, the competing values, the dilemmas where the right answer depends on context that no policy can fully anticipate.

12.2.3 Relational Intelligence Trust, culture, and coalition-building across departments are the lubrication layer that no algorithm can replicate. Every's experiment proves this directly: agent trust is derivative of human trust (Shipper 2026). The human relationship network is the foundation on which the agent mesh is built. The judgment broker's relational intelligence serves a specific function in the mesh: it enables the cross-functional insight propagation that the compounding loop depends on. A validated workflow in procurement only becomes organisational capability if it can be adapted and

Page 75

12.3. MANAGEMENT AS AI SUPERPOWER adopted by finance, logistics, and operations. That adaptation requires understanding not just the workflow but the political dynamics, cultural norms, and institutional sensitivities of each receiving function. This is human work.

12.3 Management as AI Superpower Mollick (2025) argues that 20–40 per cent of the competitive value of American firms throughout the twentieth century came from

management quality, and that AI amplifies this further. His core thesis is arresting: most companies delegate AI strategy to middle management or IT, which guarantees failure. The CEO must personally use AI, articulate how it changes what the organisation is, and create safe experimentation incentives. Mollick proposes a three-part organisational model: 1. Leadership : The CEO and senior executives must personally use AI — not delegate it, not review reports about it, but use it in their own work. This serves two functions: it gives leaders first-hand understanding of the jagged frontier in their domain, and it provides visible modelling that legitimises AI adoption throughout the organisation. 2. Crowd : Employees given permission and psychological safety discover use cases that AI strategy teams never anticipated. The shadow workflows documented in Chapter 3 are the evidence: workers are already discovering applications that management has not imagined. 3. Lab : A dedicated cross-functional team pushes the boundaries of what is possible, experiments with frontier capabilities, and translates discoveries into organisational practice. Key Insight Mollick's model maps onto the Dynamic Agentic Mesh: • Leadership sets the value vectors and strategic direction. • The Crowd operates the Discovery Engine, surfacing insights from across the organisation. • The judgment broker operates between them, validating discoveries against strategy and governance. • The Lab is the technical team that maintains and extends the mesh infrastructure. 12.4 The Competency Shift The transition from information router to judgment broker requires a specific set of competencies that differ significantly from traditional management skills: 63

Page 76

CHAPTER 12. THE ROLE TRANSFORMATION Table 12.2: Competency requirements: information router versus judgment broker

Domain	Information Router	Judgment Broker
Information	Aggregation and filtering	Frontier mapping and pattern recognition
Decision-making	Approval/rejection within defined boundaries	Edge-case adjudication at authority
Technology	Tool user	System thinker; understands mesh behaviour
Communication	Upward reporting; downward delegation	Cross-functional translation; stakeholder alignment
Risk	Compliance checking	Uncertainty navigation at the jagged frontier
Learning	Skills training	Continuous frontier probing; experimentation

12.5 The Day-in-the-Life Question A practical objection from middle managers is: “What does my Monday morning look like as a judgment broker? What am I literally doing at 9am?” This is a fair question, and the answer must be concrete: Morning : Review the Trust Variance dashboard. Identify any agents whose decision quality has drifted outside tolerance. Review the escalation queue — the genuine edge cases that the mesh has surfaced for human judgment. Adjudicate three or four cases, providing reasoning that feeds back into the mesh's learning. Midday : Cross-functional alignment session (30 minutes, not 90). The mesh has pre-synthesised the status across all workstreams. The judgment broker's role is not to hear status reports — the mesh already has those — but to identify strategic conflicts, ethical tensions, and opportunities for cross-pollination that the mesh cannot evaluate. Afternoon : Frontier probing. Spend protected experimentation time actively testing the boundary between what the mesh can handle and what it cannot in your specific domain. Document findings. Update the mesh's authority boundaries based on what you learn. Throughout : Relational work. The

human network that supports the agent mesh requires maintenance — building trust, resolving interpersonal friction, coaching team members in their own agent management, and ensuring that the cultural foundation of the mesh remains healthy. 64

Page 77

12.5. THE DAY-IN-THE-LIFE QUESTION This is a more intellectually demanding role than the information-routing role it replaces. It is also a more strategically valuable one. 65

Page 78

CHAPTER 12. THE ROLE TRANSFORMATION 66

Page 79

Chapter 13 The Technical Substrate: VisionClaw The honest version is more defensible and more credible. 13.1 A Working Agentic Mesh VisionClaw is not a prototype, a whitepaper, or a concept rendering. It is an operational agentic mesh serving a creative technology team at DreamLab, Fairfield, Eskdale. This chapter describes what the system is, what it does, and — equally important — what it does not yet do. The claims have been reconciled against the codebase as of April 2026. [FIGURE: VisionClaw Architecture Diagram] OWL 2 Ontology Engine, Claude-Flow Orchestration, WebXR Visualisation, Nostr Identity, and RuVector Memory layers Figure 13.1: VisionClaw System Architecture 13.2 Architecture Stack The system comprises five integrated subsystems, each verified against the production codebase. 67

Page 80

CHAPTER 13. THE TECHNICAL SUBSTRATE: VISIONCLAW 13.2.1 OWL 2 Ontology Engine The knowledge graph uses formal OWL 2 ontology via `horned-owl (v1.2.0)` and `whelk-rs EL++` inference. The engine parses OWL/XML, Manchester, RDF/XML, and Turtle formats, with `Neo4j` providing persistence. This subsystem provides the shared semantics for agent skill routing: task decomposition uses ontological subsumption rather than keyword matching, meaning that the system understands that “risk assessment” is a sub-task of “governance review” and routes accordingly. The OWL 2 engine is the architectural feature that most clearly distinguishes VisionClaw from other agentic systems. Formal ontology provides two capabilities that informal knowledge representations lack: compositionality (new concepts can be constructed from existing ones with guaranteed semantic consistency) and inference (the system can derive relationships that were not explicitly stated). These properties make the governance and provenance mechanisms described in Chapter 9 possible. 13.2.2 Claude-Flow Orchestration Task orchestration operates through Claude-Flow with RAFT consensus. The system integrates 83 agent skills across 12 domains. Agent routing operates through ontology-grounded skill decomposition: when a task arrives, the orchestration layer decomposes it into sub-tasks using the OWL 2 ontology and routes each sub-task to the appropriate specialist agent. 13.2.3 WebXR 3D Knowledge Visualisation The visualisation layer provides a Rust/CUDA/WebXR backend with 92 CUDA kernels (6,750 lines of code) for

GPU-accelerated graph physics. The system offers two rendering modes: Babylon.js for immersive and VR interaction, and React Three Fiber for desktop access. The visualisation is proven at approximately 1,000-node scale in production, with 100,000-node performance as a design target under active development. 13.2.4 Nostr Agent Identity Agent identities use the Nostr protocol (NIP-98 HTTP authentication, nostr-sdk v0.43.0). This provides agent signing and cascading trust semantics through a self-sovereign identity layer for agent communication. The decentralised identity model avoids dependence on a central certificate authority and provides cryptographic verification of agent actions. 68

Page 81

13.3. WHAT IS REAL AND VERIFIED 13.2.5 RuVector Memory The persistent memory layer stores over 1.17 million vector embeddings in PostgreSQL with pgvector and HNSW indexing. This provides agents with persistent memory across sessions — a capability that addresses one of the key limitations identified in Every’s agent colony model, where memory gaps cause context loss between sessions (Shipper 2026). 13.3 What Is Real and Verified The following capabilities have been verified against the production codebase as of April 2026: • OWL 2 reasoning with Whelk-rs EL++ inference — genuine, functional, and unusual in the agentic AI space. • GPU-accelerated graph physics — 92 CUDA kernels, proven at approximately 1,000-node scale. • Multi-user WebXR visualisation via Babylon.js — working in production. • Nostr identity and signing (NIP-98) — working. • Claude-Flow skill integration — 83 skills in production across 12 domains. • RuVector memory — 1.17 million embeddings with HNSW indexing. 69

Page 82

CHAPTER 13. THE TECHNICAL SUBSTRATE: VISIONCLAW 13.4 What Was Previously Overclaimed Corrected Claims A codebase audit identified several claims in earlier versions of this document that exceeded what the code supports. The corrections are: • “101 specialist skills” is corrected to 83 Claude Code skills across 12 domains. The previous count was inflated. • Multi-model orchestration via Jarvis, Gemini, and GLM-4 DAG pipelines does not exist. The system uses Claude-Flow only. The DAG struct in the codebase is for GPU physics layout, not multi-model orchestration. • W3C-compliant DID key rotation is not implemented. Nostr NIP-98 authentication works; no W3C DID implementation or key rotation code exists in the codebase. • “100,000+ nodes at 60fps” is a unit test case, not a validated production capability. Production screenshots show approximately 934 nodes. • “4M+ tokens of live knowledge graph” cannot be verified. The system contains 1.17 million vector memory entries; no token counting infrastructure exists. • The hardware specification of “96GB VRAM — Dual A6000 Ada” appears only in skill examples. The README specifies RTX 4080+ (16GB) as the target hardware. • “Stratified unidirectional links” — no code or documentation for this pattern was found. 13.5 The Corrected Story The corrected story is still strong. A working OWL 2 + CUDA + WebXR + Nostr + Claude-Flow system with 83 integrated skills and 1.17 million vector embeddings, proven at small- to-medium scale with approximately 1,000-node graphs, is genuinely novel. It does not need inflation. The honest version is more defensible and more credible. 13.5.1 What This Proves • Ontology-grounded

agent orchestration works — agents coordinate through formal semantics, not keyword matching. • GPU-accelerated knowledge visualisation is operationally feasible at production scale. • Nostr provides a viable decentralised identity layer for agent ecosystems. • The architecture is production-tested at small-to-medium scale. 70

Page 83

13.6. POSITIONING GUIDANCE 13.5.2 What This Does Not Yet Prove • Scalability beyond approximately 50 people or 1,000 nodes in production. • Adoption by non-technical middle managers. • The judgment broker role when the broker is not the system architect. • The Insight Ingestion Loop in a regulated industry. • The proposed KPIs (Mesh Velocity, Augmentation Ratio, Trust Variance) at scale. • Multi-provider orchestration (only ClaudeFlow is currently implemented). 13.6 Positioning Guidance In any external presentation, the following guidelines should govern how VisionClaw is described: 1. Use “VisionClaw” consistently (the actual system name), not “VisionFlow” (used inconsistently in earlier documents). 2. Lead with what is verified. Frame the 100,000-node target and multi-provider orchestration as roadmap items, not current reality. 3. Position VisionClaw as an “existence proof under favourable conditions” rather than “proof” — honest about the gap between a 15-person creative technology team and a 500-person manufacturing firm. 4. The audience for this document includes people who will check. Overclaiming undermines the genuine strengths of the system and the credibility of the broader argument. Key Insight The value of VisionClaw to the coordination collapse argument is not that it is the largest or most powerful agentic system. It is that it demonstrates a specific architectural approach — ontology-grounded orchestration with embedded governance — that the theoretical argument requires. The existence of a working system, even at modest scale, transforms the argument from “this could work in theory” to “this works in practice, and the question is how far it scales.” 71

Page 84

CHAPTER 13. THE TECHNICAL SUBSTRATE: VISIONCLAW 72

Page 85

Chapter 14 From Org Chart to Agentic Mesh: Implementation Phase 3 never ends. 14.1 The Bootstrap Sequence This chapter provides a 90-day implementation roadmap. The framing is deliberate: this is not a transformation programme with an end date. It is the bootstrap sequence for the perpetual change infrastructure described in Chapter 10. The roadmap launches the compounding loop. Once launched, the loop is the plan. The contradiction noted in Chapter 10 — that a phased roadmap sits inside a report criticising phased transformation models — is resolved by this distinction. Traditional change management produces a programme with an end date. The Dynamic Agentic Mesh requires a system with a start date. The 90-day roadmap is the start date. Key Insight The 90-day roadmap does not implement the mesh. It creates the conditions under which the mesh can begin to compound. The mesh implements itself — through the Insight Ingestion Loop, through the judgment broker’s validation cycle, through the workforce’s ongoing discovery of shadow workflows.

The roadmap's job is to make those mechanisms operational. 14.2 Phase 1: Audit and Discovery (Days 1–30) The first phase maps the territory. No technology is deployed into production during this phase. The objective is to understand the current state of coordination — formal and informal — before 73

Page 86

CHAPTER 14. FROM ORG CHART TO AGENTIC MESH: IMPLEMENTATION proposing changes to it. 14.2.1 Shadow Workflow Audit Map what AI tools staff are already using. The evidence from Chapter 3 is unambiguous: 75 per cent of knowledge workers are using AI (Microsoft and LinkedIn 2024), 78 per cent are bringing their own tools, and 69 per cent of organisations suspect employees use prohibited generative AI tools (Gartner, Inc. 2025). The shadow workflows already exist. The audit surfaces them. The audit should identify: • Which AI tools are in use, by whom, and for what tasks. • Which official processes have been informally bypassed or shortened. • Where the formal organisational structure diverges from the actual flow of work. • Which shadow workflows represent genuine process improvements and which introduce compliance or security risks. The audit must be conducted with explicit guarantees of non-punitive treatment. Workers will not disclose their AI usage if disclosure carries personal risk. The psychological safety evidence (Edmondson and Lei 2025) demonstrates that psychologically safe employees are 3.4 times more likely to report AI tool discoveries to management. Providing sanctioned alternatives reduces unauthorised use by 89 per cent. The message to the workforce is: “Show us what you have built. We want to learn from it, not punish it.” 14.2.2 Friction Point Identification Identify 5–10 high-value cross-functional friction points where information routing currently requires human intermediation that could be handled by agentic orchestration. These are the initial candidates for DAG construction in Phase 2. Selection criteria should include: 1. The friction point involves information transfer across departmental boundaries. 2. The current process requires manual aggregation, reformatting, or contextualisation. 3. The decision logic is sufficiently well-understood that an agent can handle the routine cases. 4. The edge cases are identifiable — a judgment broker can be assigned to them. 74

Page 87

14.2. PHASE 1: AUDIT AND DISCOVERY (DAYS 1–30) 14.2.3 Monitoring Deployment Deploy passive monitoring agents across key communication channels to detect patterns of informal collaboration and workaround behaviour. This step carries significant legal risk and must be preceded by the regulatory groundwork described in Section 14.2.6. 14.2.4 Baseline Measurement Establish baseline values for the four KPIs defined in Chapter 8: Table 14.1: Baseline KPI targets for Phase 1 measurement KPI Target Baseline Method Mesh Velocity < 48 hours Measure current time from process improvement suggestion to implementation Augmentation Ratio > 65% Survey current cognitive load distribution between human and AI Trust Variance < 0.12 σ Not measurable until agents are deployed; document current decision quality variance HITL Precision > 90% Not measurable until escalation protocols are active; document current escalation patterns 14.2.5 Judgment Broker Identification Identify the 9 per cent — the potential judgment broker candidates from the 1-9-90 pattern (Boston

Consulting Group 2025b). The 1 per cent — the natural innovators — will already be visible from the shadow workflow audit. The 9 per cent — the champions who will adopt with support — are the strategic investment target. Characteristics to look for: • Middle managers who have already experimented with AI tools for coordination tasks, even informally. • Managers whose teams show evidence of process innovation, whether AI-assisted or not. • Individuals who demonstrate comfort with ambiguity and a willingness to operate at the boundary between defined process and emergent practice. Not every current middle manager is a candidate. The identification process should be honest about this. 75

Page 88

CHAPTER 14. FROM ORG CHART TO AGENTIC MESH: IMPLEMENTATION 14.2.6 Legal and Regulatory Groundwork Employment Law and Consultation Obligations Any programme that proposes to transform middle management roles triggers legal obligations that cannot be deferred to later phases. The following requirements apply before any technology deployment: United Kingdom : The Employment Rights Act 2025 (Employment Rights Act 2025: Major Reforms for UK Employment Law 2025) broadens collective consultation triggers and requires consultation where 20 or more employees may be affected by role changes or redundancies. Failure to consult carries a protective award of up to 180 days' pay per affected employee — doubled from the previous 90-day maximum. A 90-day roadmap that deploys agentic systems affecting management roles without prior consultation is not merely inadvisable — it is unlawful. European Union : The EU AI Act (Regulation 2024/1689), Article 26(7), requires employers of high-risk AI systems to inform workers' representatives and affected workers before deployment in the workplace (European Parliament 2025). Works council consultation is mandatory, not optional. Organisations operating across EU member states must account for national transposition differences — consultation rights vary significantly between, say, Germany (where works council authority is extensive) and Ireland (where the framework is less prescriptive). GDPR and Workplace Monitoring : Deploying passive monitoring agents across employee communication channels requires a Data Protection Impact Assessment under GDPR Article 35, clear purpose limitation and data minimisation documentation, employee notification and consent frameworks, and defined retention policies for monitored data. Failure to address these requirements exposes the organisation to regulatory sanctions and, equally important, destroys the psychological safety that the entire adoption strategy depends on. Practical implication : Phase 1 must include engagement with employment lawyers and, where applicable, trade union or works council representatives. The legal workstream runs in parallel with the technical audit and cannot be treated as an afterthought. In jurisdictions with strong collective bargaining rights, the consultation timeline may exceed 30 days, requiring the overall roadmap to extend accordingly. 14.3 Phase 2: Architecture and Codification (Days 31–60) Phase 2 builds the initial mesh infrastructure using the discoveries from Phase 1. This is the most resource-intensive phase and the one most likely to stall without adequate staffing and 76

14.3. PHASE 2: ARCHITECTURE AND CODIFICATION (DAYS 31–60) expertise. 14.3.1 Agent Skill Pilot Deploy 10–15 agent skills relevant to the top friction points identified in Phase 1. The initial deployment should be narrow and deep rather than broad and shallow: better to automate three workflows thoroughly than to touch fifteen superficially. The selection should prioritise skills that: • Address friction points where the shadow workflow audit found the greatest divergence between formal and actual processes. • Operate inside the jagged frontier for the specific domain — tasks where AI performance is reliably strong. • Produce outputs that are visible and evaluable by the judgment broker candidates identified in Phase 1. The pilot is not a proof of concept. It is the first operational segment of the mesh. The agents deployed in Phase 2 remain in production; they are not discarded after a trial period. 14.3.2 Domain Ontology Construction Construct the initial OWL 2 domain ontology for core business processes. This step requires specific expertise and must not be underestimated. Ontology Construction Is Not a Side Project Constructing an OWL 2 domain ontology in 30 days is either a trivial exercise that produces a toy ontology or a substantive knowledge engineering project that requires: • One or more knowledge engineers with OWL 2 experience • Domain experts from each business function being modelled • Iterative validation sessions with operational staff • Tools: Protégé or equivalent ontology editor, reasoner integration If internal expertise is unavailable, this step requires external support and should be budgeted accordingly. The pragmatic approach: begin with a narrow ontology covering the 5–10 friction points identified in Phase 1. Use OWL 2 EL++ (the profile supported by VisionClaw’s *Whelk-rs* reasoner) to keep the ontology tractable. Expand iteratively as the mesh compounds. 77

CHAPTER 14. FROM ORG CHART TO AGENTIC MESH: IMPLEMENTATION 14.3.3 DAG Construction Build the first validated Directed Acyclic Graphs from discovered shadow workflows. Each DAG formalises a workflow that was previously informal, specifying: • The task decomposition with ontological grounding. • The agent skills assigned to each sub-task. • The escalation criteria — under what conditions does the task route to the judgment broker. • The governance constraint embedded in each transition (bias thresholds, compliance checks, access controls). • The SLAs for each stage. 14.3.4 Judgment Broker Workshops Run workshops that reframe the middle management identity from information router to judgment broker. These workshops are not training events in the conventional sense. They are identity-level conversations that address the resistance patterns documented in Chapter 10: role obsolescence fear (67%), competence anxiety (58%), loss of information asymmetry (49%), decision authority erosion (45%), and accountability ambiguity (52%) (Capgemini Research Institute 2025). The workshop content should include: 1. Direct acknowledgment of the threat: “The coordination skill you were hired for is being automated. That is real.” 2. The conductor escalation: “You are being asked to move from playing an instrument to conducting the orchestra. That is not a consolation — it is a promotion in complexity, scope, and strategic importance.” × 3. The failure mode without them: the ant death spiral, the 17 error multiplication, the Klarna reversal. 4. Hands-on practice with the pilot agents — not as tools

to use, but as systems to govern. 14.3.5 Protected Experimentation Time Create protected AI experimentation time for the identified 9 per cent. The BCG/MIT evidence shows a 3.1 times higher adoption rate when experimentation time is provided (Boston Consulting Group and MIT Sloan Management Review 2025). This is not optional enrichment. It is a strategic investment with a measured return. 78

Page 91

14.4. PHASE 3: LIVE MESH AND GOVERNANCE (DAYS 61–90 AND BEYOND) 14.4 Phase 3: Live Mesh and Governance (Days 61–90 and Beyond) Phase 3 activates the compounding loop. The critical feature of Phase 3 is that it does not end. The governance, monitoring, and learning mechanisms launched in this phase are permanent operational infrastructure. 14.4.1 Production Mesh Activation Promote validated workflows into the production mesh with full SLAs. Each workflow enters production with: • A named owner (human judgment broker). • Defined authority boundaries (what the agents can decide autonomously; what requires escalation). • Trust Variance monitoring (continuous quality measurement). • Audit trail (ontological provenance through the knowledge graph). 14.4.2 Governance Activation Activate the declarative governance layer with Trust Variance alerting live. The judgment broker receives real-time notification when agent decision quality drifts outside tolerance. 14.4.3 Insight Portal Launch Launch the Insight Portal for frontline contribution to the mesh. This provides the bottom-up channel for the Discovery Engine: any employee can propose a workflow improvement, a cross-functional shortcut, or a novel AI application. The mesh evaluates proposals against the domain ontology and routes viable ones to the appropriate judgment broker for validation. 14.4.4 First Mesh Intelligence Report Deliver the first Mesh Intelligence Report to the C-Suite, using the new KPI dashboard defined in Chapter 8. This report replaces headcount metrics with learning metrics: how fast the organisation is converting discoveries into capability, how reliably the agentic layer is handling cognitive load, and whether human attention is directed at the right problems. 79

Page 92

CHAPTER 14. FROM ORG CHART TO AGENTIC MESH: IMPLEMENTATION 14.4.5 Compound Loop Initiation Begin the perpetual cycle. The Insight Ingestion Loop is now operational. Shadow workflows discovered by the workforce are detected, codified, validated by judgment brokers, integrated into the live mesh, and propagated across teams. Mesh Velocity measures how fast this cycle runs. Phase 3 never ends. 14.5 The Wardley Map: Middle Manager Evolution Figure 14.1: Wardley map showing the evolution of the middle management role from information routing (commodity, automatable) toward judgment brokerage (genesis/custom-built, irreducible). The horizontal axis represents evolutionary stage; the vertical axis represents value chain position. The Wardley map in Figure 14.1 illustrates the strategic logic of the implementation. Information routing — the traditional middle management function — has moved rightward toward commodity and utility. It is automatable because it is well-understood and standardised. Judgment brokerage sits to the left — in the genesis and custom-built stages — because it is novel, 80

14.6. COST TIER MODEL context-dependent, and not yet standardised. The implementation roadmap moves the organisation's investment from right to left: from commodity coordination to custom judgment.

14.6 Cost Tier Model Table 14.2: Indicative cost tiers for agentic mesh deployment

Dimension	Small (500+)	Medium (50-100)	Large (100-500)
Ontology construction	2-4 weeks	1 2-3 months	6+ months, dedicated team knowledge 2-3 engineer engineers
Agent skill deployment	10-15 skills	30-50 skills	80+ skills, multi-domain GPU infrastructure
Single Dedicated Multi-node cluster workstation server (RTX (A6000-4080+) class)			
Judgment broker training	2-3 Rolling	Permanent training	function workshops programme, 6-12 months
Legal/compliance External In-house or Dedicated regulatory team counsel, ad retained, hoc ongoing Annual SaaS/compute	£50-150K	£200-600K	£1M+
Total Year 1 investment	£150-300K	£500K- £2-5M	1.2M

These figures are indicative and will vary significantly by industry, regulatory burden, and existing technology infrastructure. They exclude opportunity costs and internal staff time. A COO approving this programme needs at minimum these ranges to compare against the cost of not deploying — the \$4.63 million average cost per shadow AI breach (IBM Security 2025), the 55 per cent regret rate for rushed AI replacement (Orgvue and Forrester Research 2025), and the competitive compounding that accrues to organisations that invest early. 81

CHAPTER 14. FROM ORG CHART TO AGENTIC MESH: IMPLEMENTATION 14.7 Failure Criteria: When to Stop Every implementation plan should define what “stop, this is not working” looks like. The following signals, observed at or before day 45, should trigger a pause-and-review:

1. Shadow workflow audit yields fewer than three actionable discoveries. This suggests either the workforce does not trust the audit process (a psychological safety failure) or the organisation's work is not yet amenable to agentic augmentation.
2. Zero judgment broker candidates emerge from the 9 per cent identification. If no middle managers volunteer or are identified as candidates, the cultural preconditions for the mesh do not exist. Forcing the programme forward without willing participants produces resentment, not transformation.
3. Legal counsel identifies irresolvable consultation obligations. In unionised workplaces or jurisdictions with strong collective bargaining rights, the consultation timeline may exceed 90 days. This is not a failure of the mesh concept — it is a failure of the timeline. Extend, do not abandon.
4. Trust Variance exceeds 0.25σ on pilot workflows in the first two weeks. This indicates that the selected tasks are outside the jagged frontier for the current model capabilities. Reselect tasks closer to the frontier's interior.
5. Executive sponsorship weakens. If the C-Suite stops engaging with the programme before the first Mesh Intelligence Report, the programme will fail regardless of technical execution. The evidence on directed autonomy (Microsoft 2025a) is clear: leadership must set guardrails and investment priorities, not merely approve a budget and disengage. Key Insight A programme that stops at day 45 because the preconditions are not met is not a failure. It is an organisation that has learned something valuable about its own readiness. The mesh can be re-attempted when the preconditions change. Pressing forward without them guarantees the kind of initiative fatigue that Gartner, Inc. (2025) warns organisations using traditional

change management for AI are 2.4 times more likely to experience. 14.8 Resistance Mitigation in the Roadmap The gap analysis for this report noted that the five resistance patterns documented in Chapter 10 were identified but not mitigated in the roadmap. Table 14.3 maps each pattern to a specific roadmap action. 82

Page 95

14.8. RESISTANCE MITIGATION IN THE ROADMAP Table 14.3: Resistance patterns mapped to roadmap actions Pattern Phase Mitigation Action Role obsolescence fear (67%) 2 Judgment broker workshops with direct threat acknowledgment and conductor escalation framing × Competence anxiety (58%) 2–3 Protected experimentation time (3.1 adoption); personalised coaching, not generic training Information asymmetry loss (49%) 1–2 Reframe: the judgment broker’s advantage is now contextual judgment, not data access Decision authority erosion (45%) 2–3 Explicit escalation protocols that route genuine edge cases to brokers, not past them Accountability ambiguity (52%) 1 Legal groundwork (Section 14.2.6); clear RACI for agent-assisted decisions The 90-day roadmap is designed to be falsifiable. If the specified failure criteria are met, the organisation has learned something valuable — that the prerequisites for mesh deployment are not yet in place — and can adjust accordingly. An implementation plan that cannot fail is not a plan; it is a wish. 83

Page 96

CHAPTER 14. FROM ORG CHART TO AGENTIC MESH: IMPLEMENTATION 84

Page 97

Chapter 15 Open Questions and Research Directions A working document that hides its weaknesses is less useful than one that maps them. These gaps are research directions, not admissions of failure. This chapter is an exercise in intellectual honesty. The coordination collapse thesis, as presented in this report, is incomplete. The evidence base has specific weaknesses. The analysis contains blind spots visible to each of its three audiences. Several important topics are missing entirely. Mapping these gaps is not an admission of failure — it is a research agenda. 15.1 The Core Thesis Is Incomplete The report frames hierarchy as an information routing protocol and argues that AI collapses the cost of that routing to near-zero. This framing is analytically productive — it generates the coordination collapse concept, the judgment broker role, and the compounding organisation model. But hierarchy performs at least five functions, of which information routing is only one. 15.1.1 Legitimation Who has authority to commit resources? The mesh can route information and even recommend resource allocation, but it cannot authorise a £2 million capital expenditure. Legitimation — the organisational mechanism by which decisions acquire binding authority — is a function of hierarchy that AI does not replace. A judgment broker who validates a shadow workflow is exercising legitimate authority derived from their position in the organisational structure. Remove the structure, and the authority evaporates. 15.1.2 Accountability Who is fired when things go wrong? “The mesh decided” is not an answer a board of directors accepts. Accountability requires identifiable humans who bear consequences for outcomes. The judgment broker role addresses this partially — the broker who validates a workflow bears 85

CHAPTER 15. OPEN QUESTIONS AND RESEARCH DIRECTIONS accountability for that validation — but the full accountability architecture for agentic systems is not worked out. When an agent in the mesh makes a decision that produces a regulatory violation, the chain of accountability must be traceable to a human. The declarative governance layer (Chapter 9) provides audit trails, but audit trails are not the same as accountability structures.

15.1.3 Political Negotiation Whose priorities win when budgets are finite? AI can optimise within constraints; it cannot set the constraints. The allocation of scarce resources across competing priorities is fundamentally a political act — a negotiation among stakeholders with different interests, values, and power. The mesh can surface information that makes the negotiation more transparent, but it cannot replace the negotiation itself. An Honest Reframing If three of hierarchy's five core functions persist — legitimisation, accountability, and political negotiation — then “coordination collapse” may be an overstatement. A more honest framing: information routing collapse — one function of hierarchy dissolves, forcing the others to restructure around different assumptions. The coordination tax shrinks dramatically; it does not reach zero. This report uses the more provocative framing because it captures the magnitude of the disruption, but the qualification belongs in the analysis. Research needed : How do legitimisation, accountability, and political negotiation operate in post-hierarchical structures? Haier's micro-enterprise model provides the best existing evidence but operates in a culturally specific context that may not generalise.

15.2 Evidence Base Weaknesses The argument in this report rests on specific empirical findings. Several of these findings have limitations that deserve explicit acknowledgement.

15.2.1 The Jagged Frontier Map Is Outdated The BCG experiment (Dell'Acqua et al. 2026) used GPT-4 — now two model generations behind current frontier models. The specific task boundaries identified in that study — the 18 tasks inside the frontier and the business judgment task outside it — have almost certainly shifted. Models released since the study likely perform well on tasks that were outside GPT-4's frontier. 86

15.2. EVIDENCE BASE WEAKNESSES This does not invalidate the concept. The jagged frontier is a structural feature of AI capability, not a fixed map. The principle — that adjacent tasks can fall on opposite sides of the boundary, and that the boundary is unpredictable — holds regardless of where the boundary currently sits. But any practitioner using this report to make operational decisions must retest the frontier with current models in their specific domain. The 2023 map is a historical reference, not a current navigation chart. ×

15.2.2 The 17 Error Source Is Weak The 17-times error multiplication statistic (Towards Data Science 2026) does heavy argumentative work throughout this report. It grounds the case against naive multi-agent deployment, it supports the judgment broker role, and it provides the quantitative punch behind the ant death spiral narrative. But the source is Towards Data Science — a community-edited blogging platform, not a peer-reviewed journal. Caution ×

The 17 error multiplication statistic should be treated as indicative rather than definitive. The phenomenon it describes — error propagation and amplification in multi-agent chains — is real and documented in systems engineering literature. But the specific magnitude has not

been independently validated. Future versions of this analysis should either locate the underlying peer-reviewed source or replace the statistic with more rigorous evidence. 15.2.3 The Klarna Narrative Is Cherry-Picked Klarna's 2025 financials showed improved margins alongside the customer satisfaction decline described in Chapter 6. The CEO's admission that "cost unfortunately seems to have been a too predominant evaluation factor" (Klarna 2025) comes from a longer statement that also defended the AI strategy as a necessary evolution. Presenting only the quality degradation without the financial improvement is selective evidence — the kind of rhetorical move that a hostile reader will identify and use to dismiss the broader argument. The honest presentation: Klarna achieved short-term cost reduction, experienced quality degradation, began rehiring humans, and the CEO acknowledged the trade-off. The lesson is not "AI customer service fails" but "optimising for cost alone, without maintaining human oversight at quality-sensitive boundaries, produces predictable quality collapse." This is a more nuanced and more defensible claim. 87

Page 100

CHAPTER 15. OPEN QUESTIONS AND RESEARCH DIRECTIONS 15.2.4 The 89 Per Cent Reduction The claim that providing approved AI tools reduces unauthorised use by 89 per cent is used repeatedly in this report. No primary source is cited. The statistic is plausible — it aligns with shadow IT research showing that sanctioned alternatives reduce unsanctioned adoption — but a claim used this frequently requires a traceable citation. This is an open research gap. 15.2.5 Vendor Report Methodology The shadow AI statistics used throughout the report stack vendor-funded research — Gartner, IBM, Capgemini — that has methodological limitations and commercial incentives to inflate threat numbers. No error bars, confidence intervals, or methodological caveats are provided. A reader from an academic background will note this immediately. The statistics should be treated as directionally correct rather than precisely accurate. 15.3 Audience Blind Spots 15.3.1 No Day-in-the-Life Middle managers will ask: "What does my Monday morning look like as a judgment broker? What am I literally doing at 9am?" This report offers the conductor metaphor, the role transformation table, the three irreplaceable capacities — but no concrete job specification, no competency framework, no salary benchmarking, no day-in-the-life narrative. The abstraction is intellectually satisfying but operationally insufficient for a middle manager trying to decide whether to invest in this transition. Research direction : develop detailed judgment broker role specifications for three to five industry verticals, including day-in-the-life scenarios, competency models, and compensation benchmarking. 15.3.2 No Total Cost of Ownership AI leaders will ask: "What does this cost for a 500-person organisation? What is the ROI timeline?" Chapter 14 provides an indicative cost tier model, but this falls short of a total cost of ownership analysis that includes ongoing compute, model API costs (which are volatile and trending downward), knowledge engineering maintenance, judgment broker training, and the opportunity cost of the transition period. Without financial modelling, the implementation roadmap is a strategic argument, not a business case. 88

15.4. ENGAGING WITH ACEMOGLU Research direction : develop a TCO calculator parameterised by organisation size, industry vertical, regulatory burden, and existing technology infrastructure. 15.3.3 The Roadmap Contradiction Change managers will notice that this report attacks Kotter and ADKAR for being phased transformation models, then offers a phased 90-day roadmap. Chapter 10 addresses this contradiction explicitly — framing the roadmap as a bootstrap sequence rather than a change plan — but the resolution requires the reader to accept a distinction that may feel like sophistry. The tension is genuine and deserves continued intellectual engagement rather than a single paragraph of resolution. 15.3.4 The Analyst-Vendor Conflict All three audiences will notice that DreamLab AI Consulting is simultaneously the analyst (producing this report), the framework creator (proposing the Dynamic Agentic Mesh), and the vendor (operating VisionClaw). This is a classic consulting positioning conflict. The report’s credibility depends on acknowledging this conflict directly: “We built this because we believe it. Here is the evidence for and against. Judge the framework on its merits, not its provenance.” 15.4 Engaging with Acemoglu Daron Acemoglu’s work on AI and productivity has been cited in this report’s bibliography but not substantively engaged with in the main text. This is a significant omission because Acemoglu’s thesis directly challenges the optimistic productivity assumptions underlying the coordination collapse argument. 89

CHAPTER 15. OPEN QUESTIONS AND RESEARCH DIRECTIONS The Acemoglu Sceptical Thesis In “The Simple Macroeconomics of AI” (Acemoglu 2025), Acemoglu argues that AI’s impact on total factor productivity (TFP) will be modest — at most 0.66 per cent over 10 years . His reasoning: 1. AI automates tasks, not jobs. The tasks most amenable to automation are not necessarily the tasks that contribute most to productivity. 2. The “easy-to-learn, hard-to-verify” problem: many tasks that AI can perform require costly human verification, limiting net productivity gains. This maps directly onto the vigilance problem described in Chapter 5. 3. New task creation — historically the mechanism by which automation generates employment growth — may be slower with AI than with previous technologies. 4. AI may increase inequality rather than aggregate productivity, concentrating gains among capital owners and high-skill workers while displacing middle-skill employment. The coordination collapse thesis should engage with Acemoglu on three fronts. First , the 0.66 per cent TFP estimate applies to the economy as a whole, not to individual organisations. An organisation that deploys an agentic mesh effectively may capture disproportionate productivity gains precisely because its competitors do not. The compounding organisation thesis does not require AI to transform the macroeconomy — it requires AI to provide a competitive advantage to organisations that use it well. This is consistent with Acemoglu’s framework: the gains are real but unevenly distributed. Second , the task-level analysis is compatible with the jagged frontier concept. Acemoglu’s point that AI automates tasks, not jobs, aligns with the argument that the judgment broker’s role is to determine which tasks sit inside the frontier and which require human judgment. The mesh architecture is explicitly designed around task decomposition, not wholesale job replacement. Where

Acemoglu warns that the “easy-to-learn, hard-to-verify” problem limits productivity gains, the mesh responds with HITL Precision and Trust Variance monitoring — architectural mechanisms for managing verification costs. Third, the equity concern is serious and currently unaddressed in this report. If the mesh concentrates productivity gains among organisations with the capital and technical expertise to deploy it, while displacing middle-skill workers in organisations that cannot, the net social effect may be increased inequality even as individual organisations benefit. This is a genuine tension that the coordination collapse thesis does not resolve. 90

Page 103

15.5. ENERGY AND ENVIRONMENTAL COSTS 15.5 Energy and Environmental Costs The report proposes deploying agentic infrastructure at organisational scale without addressing the energy implications. This is an increasingly visible omission as the environmental costs of AI infrastructure become a matter of public and regulatory concern. The Energy Cost of Agentic Systems • Patterson et al. (2025) at Microsoft Research find that reasoning-intensive queries — the kind that agentic orchestration relies on — have a median energy cost of 4.32 Wh per query, compared to 0.34 Wh for standard queries. This represents a $43 \times$ increase in energy consumption for the reasoning-heavy workloads that an agentic mesh generates continuously. • Nature Sustainability (2025) projects that US AI server infrastructure will produce 24–44 Mt CO₂-equivalent annually through 2030 — a range that is growing as 2 model sizes increase and inference demand scales. • An agentic mesh running continuous orchestration, discovery monitoring, and governance checks generates substantially more inference load than on-demand AI tool use. The compound flywheel, by design, increases inference volume as the mesh becomes more active. The energy cost of the mesh is not merely a corporate social responsibility issue. It is a business cost that scales with mesh activity. Organisations deploying agentic infrastructure must account for compute costs that are qualitatively different from traditional IT expenditure — costs that increase as the mesh becomes more capable, not less. Research direction : develop energy cost models for agentic mesh deployment at different organisational scales, including comparison with the energy cost of the human coordination infrastructure the mesh replaces (office buildings, commuting, business travel). The net energy equation is not obviously negative — a 200-person management layer has its own carbon footprint — but neither is it obviously positive. 15.6 Cybersecurity Attack Surface The statistic that financial services organisations manage 96 machine identities per human (CyberArk 2025) — with only 31 per cent having controls in place — is cited in Chapter 11 as a governance issue. It is also a security issue that the report does not adequately address. An agentic mesh creates a dramatically expanded attack surface: 91

Page 104

CHAPTER 15. OPEN QUESTIONS AND RESEARCH DIRECTIONS • Each agent skill is a potential vector for prompt injection — an adversary who can manipulate an agent’s input can redirect its behaviour within its authority boundary. • The OWL 2 ontology that provides shared semantics is a single point of corruption — poisoning the ontology could redirect agent

routing across the entire mesh. • The Nostr identity layer provides cryptographic verification but depends on key management practices that most organisations have not developed for non-human identities. • The Insight Ingestion Loop creates a pathway from frontline input to production workflows — a supply chain vulnerability that adversaries could exploit by planting malicious “shadow workflows.” Research direction : systematic threat modelling for agentic mesh architectures, including adversarial analysis of the Insight Ingestion Loop and ontology poisoning scenarios. 15.7 Equity and Access Who benefits from the mesh and who is structurally excluded? The report addresses middle managers — a specific professional cohort — but does not address: • Workers without AI literacy who cannot participate in the discovery layer. • Roles that do not lend themselves to agent augmentation — physical labour, care work, creative work that resists decomposition into tasks. • Organisations in the Global South that lack the infrastructure and capital to deploy agentic systems. • The distributional effects within organisations: does the mesh flatten compensation as it flattens hierarchy, or does it concentrate value in the judgment broker tier? Acemoglu’s concern about AI increasing inequality rather than productivity is most acute here. If the Dynamic Agentic Mesh is a tool available primarily to well-capitalised organisations in knowledge-intensive industries, its net effect may be to widen the gap between those organisations and everyone else — a dynamic that the World Economic Forum (2025) Future of Jobs Report anticipates when it projects 170 million new roles created alongside 92 million displaced. 15.8 Competing Frameworks Not Addressed The report positions against Block and Every — neither of which is a consulting framework the target audience is actively evaluating. It does not differentiate substantively from: 92

Page 105

15.9. THE VISIONCLAW CREDIBILITY GAP • SAFe 6.0 (Scaled Agile, Inc. 2024), which added AI-specific agile change patterns in 2024. • Microsoft Copilot Studio (Microsoft 2025b), an enterprise organisational AI deployment framework with the largest existing footprint. • Holacracy 5.0 (HolacracyOne 2024), which specifically addressed the scale failures attributed to earlier versions. • Anthropic’s constitutional AI governance , directly relevant to the governance paradox in Chapter 9. A competitor comparison table is provided in Appendix .7, but the main text does not engage with these alternatives at the depth they deserve. 15.9 The VisionClaw Credibility Gap Chapter 13 describes VisionClaw honestly, including both its genuine capabilities and its limitations. But the credibility gap between a 15-person creative technology team and a 500-person manufacturing firm has not been bridged. The favourable conditions under which VisionClaw operates — high technical literacy, low process rigidity, flat power dynamics, intrinsic motivation, the system architect as the primary user — are the most favourable possible environment. Key Insight VisionClaw is an existence proof under ideal conditions, not a proof of general applicability. The coordination collapse thesis requires bridging evidence that the architecture works in less favourable environments: regulated industries, non-technical workforces, unionised labour, legacy ERP systems, and middle managers who did not choose the system. Four bridging strategies are available: 1. Run a pilot in a genuinely different context (regulated industry, non-technical workforce). 2. Partner with an established consulting framework (Deloitte,

BCG) for independent validation. 3. Publish the methodology as open-source and invite external testing. 4. Position VisionClaw explicitly as “existence proof under favourable conditions” and build the evidence base incrementally. The fourth strategy is the one this report adopts. The others represent the research agenda. 93

Page 106

CHAPTER 15. OPEN QUESTIONS AND RESEARCH DIRECTIONS 15.10 Additional Failure Modes The report documents two failure modes in detail: the ant death spiral (agents responding to agents in infinite loops) and the vigilance degradation problem (humans falling asleep at the wheel as AI quality improves). But an agentic mesh at organisational scale introduces additional failure modes that this analysis does not adequately address: • Cascading errors : a single incorrect agent output propagating through downstream agents before the judgment broker can intervene. • Single point of failure in the judgment broker : if the broker is unavailable, incapacitated, or overwhelmed, the mesh has no fallback governance mechanism. • Ontology drift : the OWL 2 ontology gradually diverging from the actual business domain as the organisation evolves, producing increasingly inaccurate agent routing. • Metric gaming : agents optimising for the four KPIs in ways that satisfy the metrics while degrading actual performance — Goodhart’s law applied to agentic governance. • Cultural rejection : the mesh may function technically while being socially rejected by the workforce, producing compliance without engagement. Research direction : systematic failure mode analysis (FMEA) for agentic mesh architectures, with probability and severity estimates calibrated to organisational context. 94

Page 107

Chapter 16 Conclusion: The Through Line The organisation is becoming a compounding learning system — not a command structure. The critical question is where irreducible human judgment sits in that system. 16.1 Synthesis For two thousand years, organisations have paid a coordination tax — layers of humans whose primary function is routing information. The Roman contubernium , the Prussian General Staff, McCallum’s organisational chart, Taylor’s functional pyramid: each was an engineering response to the same constraint. A human leader can effectively manage three to eight people. Beyond that threshold, you must add a layer. Each layer adds latency, distortion, and cost. Two millennia of organisational innovation have been attempts to work around this tradeoff without breaking coherence. AI breaks the tradeoff. Information routing — synthesising status across workstreams, aggregating context across functions, translating across departmental boundaries, maintaining alignment without weekly meetings — is now computationally cheap. The structural reason for the organisational chart, as it has existed since the Roman legions, is dissolving. This is the coordination collapse. Three models have emerged for what replaces it. Dorsey and Botha (2026) propose the Intelligence Layer : eliminate middle management, build a Company World Model, centralise coordination in AI. The diagnosis is precise. The prescription is reckless. Ninety-five per cent of AI-generated code at Block still requires human modification. Klarna replaced 700 customer service agents with AI and reversed course when quality collapsed (Klarna 2025). The vigilance problem (Dell’Acqua 2023)

demonstrates that the better AI becomes, the harder it is for humans to maintain oversight — and Block’s model removes the humans who would maintain it. Spicer (2026) observes that every wave of of- 95

Page 108

CHAPTER 16. CONCLUSION: THE THROUGH LINE fice technology produces delaying that subsequently reverses as complexity reasserts itself. The Self-Automator mode of human–AI collaboration — the mode that Block’s architecture embodies — produces the weakest outcomes in the BCG experiment (Randazzo et al. 2024). Shipper (2026) describes the Agent Colony : give every employee a persistent personal AI agent, let culture self-organise. The experiment is the most honest account available of what actually happens when personal agents are deployed across an organisation. The insight that “everyone is a manager now” — that management skill democratises rather than disappears — directly contradicts Block’s thesis and is supported by the evidence. But the model hits structural $\times$ limits at scale. The ant death spiral, the 17 error multiplication in naive multi-agent systems (Towards Data Science 2026), and the memory gaps that cause context loss between sessions are failure modes that cultural norms cannot resolve at organisational scale. The Dynamic Agentic Mesh retains human judgment at the intersection points. Middle man- agers evolve from information routers to judgment brokers — federated guardians of a system where agentic orchestration handles the information routing that hierarchy was built for, while humans handle what AI cannot: navigating the jagged frontier, maintaining vigilance, adjudi- cating ethical dilemmas, and sustaining cultural coherence. 16.2 Block Was Right About the Diagnosis This report has been critical of Block’s Intelligence Layer model, and that criticism stands. But the criticism is of the prescription , not the diagnosis . The diagnosis is correct: • Hierarchy is an information routing protocol, not a management philosophy. • That protocol is constrained by human bandwidth. • AI removes the bandwidth constraint. • The structural logic of the organisational chart dissolves when information routing becomes computationally cheap. Where Block errs is in the inference that follows. “The information routing function is auto- mated, therefore eliminate the humans who performed it” ignores the jagged frontier (Chap- ter 4), the vigilance problem (Chapter 5), the Klarna evidence, and the historical pattern of delaying reversal. The correct inference is: “The information routing function is automated, therefore the humans who performed it must be redeployed to the functions that cannot be auto- mated — judgment, legitimation, accountability, political negotiation, and cultural coherence.” The judgment broker is not a consolation prize for displaced middle managers. It is the role that the evidence demands. 96

Page 109

16.3. EVERYONE IS A MANAGER NOW — AND THAT CHANGES EVERYTHING 16.3 Everyone Is a Manager Now — And That Changes Ev- erything Every’s experiment (Shipper 2026) reveals the deepest implication of the coordination collapse. When every employee has a personal AI agent, management skill — the ability to delegate, eval- uate, coach, and coordinate — becomes the universal skill. The graduate with a well-configured agent can approximate the information routing that middle managers were uniquely paid to do. This is

genuinely threatening to the current middle management cohort. The report does not pretend otherwise. The 67 per cent role obsolescence fear documented by Capgemini Research Institute (2025) is rational. The skill floor has risen. What was the middle manager's ceiling is now everyone's floor. But the conductor problem resolves the threat. Everyone in an orchestra can play an instrument. That does not make the conductor redundant — it makes the conductor essential in a different way. The musicians' skill is playing. The conductor's skill is ensuring that 80 individually excellent performances produce coherent music rather than 80 simultaneous solos. When every employee manages their own agent, the organisation has 200 individual human- agent pairs each optimising for their own task. Without a mesh coordinator — without the judgment broker — this produces the ant death spiral documented by Every at 20 people, the error multiplication documented in multi-agent systems research, and the quality collapse documented by Klarna in customer service. The evidence is consistent: uncoordinated agent deployment fails. The middle manager's new differentiating skill is not managing an agent. Everyone does that. It is managing the coherence of the mesh — ensuring that 200 individually optimised human-agent pairs produce organisational outcomes rather than organisational noise. This is the difference between individual agent management (Stage 2–3 on Klaassen's compound engineering ladder) and mesh coherence management (Stage 4–5). The former is a personal productivity skill. The latter is systems thinking. It is genuinely hard. And it is the role that makes the coordination collapse survivable. 97

Page 110

CHAPTER 16. CONCLUSION: THE THROUGH LINE 16.4 The Compound Flywheel Figure 16.1: The compound organisational flywheel. Discovery feeds codification; codification feeds validation by the judgment broker; validated workflows propagate through the mesh; propagation increases organisational capability; increased capability accelerates the next cycle of discovery. Each revolution makes the next one faster. The flywheel in Figure 16.1 is the core mechanism of the coordination collapse thesis. It is not a strategy diagram. It is an operational description of how the mesh produces compounding returns. Each shadow workflow surfaced by the Discovery Engine becomes institutional capability through the Insight Ingestion Loop. Each validated DAG makes the mesh smarter. Each cycle of the flywheel reduces the time required for the next cycle — the metric captured by Mesh Velocity. The compound flywheel addresses Acemoglu's scepticism about AI productivity gains (Acemoglu 2025). His 0.66 per cent TFP estimate applies to the economy as a whole. An individual organisation running a compound flywheel captures gains that are disproportionate to the macroeconomic average — precisely because the gains compound, and because the organisation's competitors have not yet built the flywheel. The strategic advantage is not in using AI. It is in building a system that learns from its own use of AI. 16.5 What the Evidence Supports The following propositions are supported by the evidence reviewed in this report: 1. Hierarchy is an information routing protocol. The coordination tax is real and measurable. AI can perform the information routing function at a fraction of the cost and latency. 2. The boundary between what AI can and cannot do is jagged, unpredictable, and dynamic (Dell'Acqua et al. 2026). Wholesale automation is empirically reckless. 3. Human oversight

degrades as AI quality improves (Dell'Acqua 2023). “Human in the loop” without architectural support for maintaining vigilance is an unreliable governance mechanism. 4. Three collaboration modes exist — Centaur, Cyborg, and Self-Automator — and the Self-Automator mode produces the worst outcomes (Randazzo et al. 2024). 98

Page 111

16.6. WHAT THE EVIDENCE DOES NOT YET SUPPORT 5. Organisations that revise their KPIs to account for AI are three times more likely to see financial benefit (Boston Consulting Group and MIT Sloan Management Review 2025). 6. The 1-9-90 adoption pattern is predictable, and investing in the 9 per cent yields the highest return (Boston Consulting Group 2025b). 7. Psychological safety is the strongest predictor of successful AI adoption (Edmondson and Lei 2025; Google 2025). Protected experimentation time produces 3.1 times higher adoption rates. 8. Providing sanctioned AI tools reduces unauthorised use dramatically. Ignoring shadow AI is not a neutral choice. 16.6 What the Evidence Does Not Yet Support The following propositions are advanced by this report but not yet validated by independent evidence: 1. The full organisational compounding flywheel — where each discovered insight makes the next discovery easier at the systems level — has not been demonstrated at scale. 2. The proposed KPIs (Mesh Velocity, Augmentation Ratio, Trust Variance, HITL Precision) have not been validated empirically. 3. The judgment broker role has not been tested with brokers who did not build the underlying system. 4. The Dynamic Agentic Mesh has not been proven beyond a 15-person team in a favourable context. 5. The governance architecture (declarative governance, ontological provenance, cascading trust) has not been stress-tested in regulated industries. Acknowledging these gaps is not a weakness. It is a statement of where the research frontier lies and an invitation for others to test the propositions in their own contexts. 16.7 Closing Messages 16.7.1 For Middle Managers You are not being replaced. You are being promoted — from information router to judgment broker. The promotion is not in title or compensation (those are organisational decisions outside this report's scope). The promotion is in complexity, scope, and strategic importance. The skill 99

Page 112

CHAPTER 16. CONCLUSION: THE THROUGH LINE you need is not technical. It is the willingness to map the jagged frontier in your domain and the judgment to know when to trust the mesh and when to override it. The organisations that will outcompete are not the ones that eliminate middle management. They are the ones that give their middle managers experimentation time (3.1 times higher adoption), psychological safety (2.6 times more likely to report successful AI integration), and a clear path from coordination to judgment. Not every current middle manager will make this transition. The 1-9-90 data is honest about that. But the path is real, the evidence supports it, and the role that emerges on the other side is more valuable, more intellectually demanding, and more strategically important than the one it replaces. 16.7.2 For Change Managers Your playbook is not broken. It is incomplete. AI transformation has no end state. You are not managing a transition; you are building an organisation's capacity for continuous adaptation. The frameworks exist: Continuous

Transformation (Wade and Bonnet 2025), Directed Autonomy (Microsoft 2025a), Dynamic Capabilities (Teece 2025). The mesh provides the architecture. The 90-day roadmap in Chapter 14 is the bootstrap sequence. Phase 3 never ends. Your job is perpetual — and that is not a burden. It is a description of a role that will be needed for as long as the jagged frontier moves, which is to say permanently. 16.7.3 For AI Leaders The secret cyborgs are your greatest asset, not your greatest risk. Surface them. Sanction them. Learn from them. The shadow workflows they have already created are the raw material for your compounding loop. Providing approved tools drops unauthorised use by 89 per cent. The alternative — 1,200 unsanctioned applications per enterprise (Gartner, Inc. 2025), \$4.63 million per breach (IBM Security 2025) — is untenable. The governance paradox has a resolution: embed governance into the architecture, and it becomes an accelerator rather than a brake. But surfacing shadow workflows requires psychological safety. Workers who have automated 90 per cent of their role will not volunteer that information if they believe it invites headcount reduction. Mollick (2025) is direct on this point: organisations that delegate AI strategy to IT without CEO engagement guarantee failure. 100

Page 113

16.7. CLOSING MESSAGES 16.7.4 For All Three Block diagnosed the right problem. The question is whether you agree with their prescription. The evidence reviewed in this report suggests you should not — but you should take the diagnosis deadly seriously. The coordination collapse is real. What you build in its place will determine whether your organisation compounds or collapses. DreamLab AI Consulting Ltd Dr. John O'Hare DreamLab, Fairfield, Eskdale April 2026 101

Page 114

CHAPTER 16. CONCLUSION: THE THROUGH LINE 102

Page 115

References Acemoglu, Daron (Jan. 2025). “The Simple Macroeconomics of AI”. In: Economic Policy 40.121. Originally NBER Working Paper 32487, April 2024. URL: <https://www.nber.org/papers/w32487>. Acemoglu, Daron and Pascual Restrepo (2019). “Automation and New Tasks: How Technology Displaces and Reinstates Labor”. In: Journal of Economic Perspectives 33.2, pp. 3–30. Boston Consulting Group (Nov. 2025a). The Emerging Agentic Enterprise . Tech. rep. Boston Consulting Group. URL: <https://www.bcg.com/publications/2025/the-emerging-agentic-enterprise> (visited on 04/10/2026). Boston Consulting Group (May 2025b). Unlocking the AI-First Organization: An Agentic Shift . Tech. rep. Boston Consulting Group. URL: <https://www.bcg.com/publications/2025/unlocking-ai-first-organization-agentic-shift> (visited on 04/10/2026). Boston Consulting Group (2026). AI and the Reshaping of US Jobs: 50–55% in 2–3 Years . Tech. rep. Boston Consulting Group. URL: <https://www.bcg.com/> (visited on 04/10/2026). Boston Consulting Group and MIT Sloan Management Review (2025). AI-Powered KPIs: Measuring What Matters in the Age of Intelligence . Tech. rep. Joint research, n=3,000+ respondents across 25 industries. Boston Consulting Group. URL:

<https://www.bcg.com/publications/2025/ai-powered-kpis> (visited on 04/10/2026). Capgemini Research Institute (2025). Middle Manager Resistance Patterns in AI Transformation . Tech. rep. Capgemini Research Institute. URL: <https://www.capgemini.com/insights/research-institute/> (visited on 04/10/2026). CyberArk (2025). AI Agents in Financial Services: The Hidden Org Chart . Tech. rep. Cyber- Ark Software Ltd. URL: <https://www.cyberark.com/resources/> (visited on 04/10/2026). Dell’Acqua, Fabrizio (2023). “Falling Asleep at the Wheel: Human/AI Collaboration in a Field Experiment on HR Recruiters”. Harvard Business School Working Paper. URL: <https://www.hbs.edu/faculty/Pages/item.aspx?num=64700> (visited on 04/10/2026). Dell’Acqua, Fabrizio, Edward McFowland, Ethan R. Mollick, Hila Lifshitz-Assaf, Katherine Kellogg, Saran Rajendran, Lisa Kraye, François Candelon, and Karim R. Lakhani (Mar. 2026). “Navigating the Jagged Technological Frontier: Field Experimental Evidence of the Effects of AI on Knowledge Worker Productivity and Quality”. In: Organization Science . Originally circulated as Harvard Business School Working Paper 24-013, September 2023. 103

Page 116

REFERENCES URL: https://papers.ssrn.com/sol3/papers.cfm?abstract_id=4573321 (visited on 04/10/2026). Deloitte (2026). Global Human Capital Trends 2026: Human–AI Interaction Design . Tech. rep. Deloitte Consulting LLP. URL: <https://www2.deloitte.com/global/en/insights/focus/human-capital-trends.html> (visited on 04/10/2026). Dorsey, Jack and Roelof Botha (Mar. 2026). From Hierarchy to Intelligence . Published simultaneously on block.xyz and sequoiacap.com. URL: <https://block.xyz/inside/from-hierarchy-to-intelligence> (visited on 04/10/2026). Edmondson, Amy C. and Zhike Lei (2025). “Psychological Safety and AI Experimentation in Teams”. In: Harvard Business Review . Based on research conducted 2024–2025 at Harvard Business School. URL: <https://hbr.org/> (visited on 04/10/2026). Employment Rights Act 2025: Major Reforms for UK Employment Law (Dec. 2025). Broadens collective consultation triggers; doubles protective award to 180 days’ pay. URL: <https://www.lawandtheworkplace.com/2025/12/major-reforms-for-uk-employment-law-era-2025/> (visited on 04/13/2026). European Parliament (Nov. 2025). MEPs Call for New Rules on Algorithmic Management at Work . Article 26(7) of the AI Act mandates employer consultation with works councils before deploying high-risk AI. URL: <https://www.europarl.europa.eu/news/en/press-room/20251110IPR31358/> (visited on 04/13/2026). Falkman Olsson, Caroline and Yafah Edelman (Mar. 2026). Half of Employed AI Users Now Use It for Work . Probability-based survey of 2,021 U.S. adults via Ipsos KnowledgePanel, March 3–5, 2026; subsample of 665 employed past-week AI users. URL: <https://epoch.ai/blog/half-of-employed-ai-users-now-use-it-for-work/> (visited on 04/13/2026). Gartner, Inc. (2025). Top Strategic Predictions for 2025–2026 and Beyond . Tech. rep. Gartner, Inc. URL: <https://www.gartner.com/en/articles/gartner-top-strategic-predictions> (visited on 04/10/2026). Google (2025). Project Aristotle Follow-Up: Psychological Safety and AI Adoption . Tech. rep. Google LLC. URL: <https://rework.withgoogle.com/> (visited on 04/10/2026). Google DeepMind (May 2025). AlphaEvolve: AI That Redesigns Its Own Infrastructure . Saved 0.7% of Google’s compute and

accelerated Gemini training. URL: <https://deepmind.google/> (visited on 04/10/2026). HolacracyOne (2024). Holacracy 5.0: Addressing Scale Through Distributed Authority . Updated to address earlier criticisms of scalability. URL: <https://www.holacracy.org/> (visited on 04/10/2026). IBM Security (2025). Cost of a Data Breach Report 2025 . Tech. rep. IBM Corporation. URL: <https://www.ibm.com/reports/data-breach> (visited on 04/10/2026). 104

Page 117

REFERENCES Klaassen, Kieran (Dec. 2025). Compound Engineering: How Every Codes With Agents . every.to. URL: <https://every.to/chain-of-thought/compound-engineering> (visited on 04/10/2026). Klarna (2025). CEO Admission on AI Customer Service Deployment . Reported by CNBC and Fortune; CEO acknowledged cost was a ‘too predominant evaluation factor’. URL: <https://www.cnbc.com/> (visited on 04/10/2026). McKinsey & Company (2026a). Six Shifts to Build the Agentic Organization of the Future . Tech. rep. McKinsey & Company. URL: <https://www.mckinsey.com/capabilities/mckinsey-digital/our-insights/six-shifts-agentic-organization> (visited on 04/10/2026). McKinsey & Company (2026b). The Agentic Organization . Tech. rep. McKinsey & Company. URL: <https://www.mckinsey.com/capabilities/mckinsey-digital/our-insights/the-agentic-organization> (visited on 04/10/2026). Microsoft (2025a). 2025 Work Trend Index: Directed Autonomy and AI Productivity . Tech. rep. Microsoft Corporation. URL: <https://www.microsoft.com/en-us/worklab/work-trend-index/> (visited on 04/10/2026). Microsoft (2025b). Microsoft Copilot Studio: Organisational AI Deployment Framework . URL: <https://www.microsoft.com/en-us/microsoft-copilot/microsoft-copilot-studio> (visited on 04/10/2026). Microsoft and LinkedIn (2024). 2024 Work Trend Index Annual Report: AI at Work Is Here. Now Comes the Hard Part .Tech.rep.Microsoft Corporation. URL: <https://www.microsoft.com/en-us/worklab/work-trend-index/ai-at-work-is-here-now-comes-the-hard-part> (visited on 04/10/2026). Mollick, Ethan (June 2023). Detecting the Secret Cyborgs . One Useful Thing (Substack); up- dated 2024–2025. URL: <https://www.oneusefulthing.org/p/detecting-the-secret-cyborgs> (visited on 04/10/2026). Mollick, Ethan (2025). Management as AI Superpower . One Useful Thing (Substack). URL: <https://www.oneusefulthing.org/> (visited on 04/10/2026). Nature Sustainability (2025). “Carbon Footprint of AI Server Infrastructure in the United States”. In: Nature Sustainability . Estimates US AI servers produce 24–44 Mt CO annually. 2 URL: <https://www.nature.com/natsustain/> (visited on 04/10/2026). Orgvue and Forrester Research (2025). Companies Regretting AI Replacement: Survey Findings . Tech. rep. 55% of companies that rushed to replace humans with AI report regretting the decision. Orgvue / Forrester Research. URL: <https://www.orgvue.com/> (visited on 04/10/2026). 105

Page 118

REFERENCES Patterson, David et al. (Sept. 2025). “Energy Use of AI Inference”. In: arXiv preprint . Microsoft Research. Median energy per query: 0.34 Wh standard, 4.32 Wh reasoning-heavy (43x increase). eprint: 2509.20241 . URL: <https://arxiv.org/abs/2509.20241> . Prosci (2024). Revised Guidance for AI Transformation: Iterative Change Portfolios . ADKAR creators acknowledging AI requires iterative change

portfolios. URL: <https://www.prosci.com/> (visited on 04/10/2026). Randazzo, Michael, Hila Lifshitz-Assaf, Katherine Kellogg, Fabrizio Dell'Acqua, Ethan R. Mollick, François Candelon, and Karim R. Lakhani (2024). "Cyborgs, Centaurs and Self-Automators: Modes of Human-AI Collaboration". Harvard Business School Working Paper. URL: https://www.hbs.edu/ris/Publication%20Files/Cyborgs-Centaurs-Self-Automators_5a8e0b5c-1a4e-4f8b-9a8a-6c3e3e1a7d8f.pdf (visited on 04/10/2026). Scaled Agile, Inc. (2024). SAFe 6.0: AI-Specific Agile Change Patterns. Added AI-specific practices to the Scaled Agile Framework in 2024. URL: <https://scaledagileframework.com/> (visited on 04/10/2026). Shipper, Dan (Apr. 2026). We Gave Every Employee an AI Agent. AI & I Podcast transcript, every.to. URL: <https://every.to/chain-of-thought/we-gave-every-employee-an-ai-agent> (visited on 04/10/2026). Spicer, André (Mar. 2026). Commentary on Block and Meta's AI Restructuring. Quoted in Fortune; André Spicer is Executive Dean, Bayes Business School, City, University of London. URL: <https://fortune.com/> (visited on 04/10/2026). Teece, David J. (2025). "Dynamic Capabilities: Sensing, Seizing, and Reconfiguring in the Age of AI". In: Strategic Management Journal. Refreshed framework for AI-era organisational adaptation. URL: <https://onlinelibrary.wiley.com/journal/10970266> (visited on 04/10/2026). Towards Data Science (2026). The 17x Error Trap in Multi-Agent Systems. Community-edited platform; not peer-reviewed. URL: <https://towardsdatascience.com/> (visited on 04/10/2026). Wade, Michael and Didier Bonnet (2025). Continuous Transformation: The Next Frontier of Organizational Change. Lausanne: IMD Business School Press. URL: <https://www.imd.org/> (visited on 04/10/2026). World Economic Forum (2025). The Future of Jobs Report 2025. Tech. rep. World Economic Forum. URL: <https://www.weforum.org/publications/the-future-of-jobs-report-2025/> (visited on 04/10/2026). Zamost, Aaron (2026). Commentary on Block's Restructuring. Former Head of Communications at Block; quoted in New York Times and Futurism. URL: <https://www.nytimes.com/> (visited on 04/10/2026). 106

Page 119

Research Methodology This appendix documents the research methodology used to produce this report, including the multi-agent research architecture, source selection criteria, codebase audit procedures, and Wardley mapping approach.

1. Multi-Agent Research Architecture The research was conducted using five parallel research agents, each assigned a distinct domain with defined source priorities and quality thresholds. The agents operated concurrently across a shared memory layer (RuVector, 1.17 million vector embeddings with HNSW indexing) that enabled cross-domain pattern detection and prevented duplication of effort.

Table 1: Research agent assignments and source priorities

Agent	Domain	Primary Sources
Agent 1	Mollick / BCG / Jagged Frontier	Dell'Acqua et al. (2026), Randazzo et al. (2024), Dell'Acqua (2023), Mollick's "One Useful Thing" corpus, Boston Consulting Group and MIT Sloan Management Review (2025)
Agent 2	Block Critique	Dorsey and Botha (2026), Zamost (2026), Spicer (2026), Klarna (2025), Orgvue and Forrester Research (2025)
Agent 3	Agentic Organisations	McKinsey & Company (2026b), McKinsey & Company (2026a), Boston Consulting Group (2025b), Boston Consulting Group (2025a), Deloitte (2026)
Agent 4	Change Management	Wade and Bonnet (2025), Teece (2025), Microsoft (2025a), Prosci

(2024), Capgemini Research Institute (2025) Agent 5 Every / Compound Engineering Shipper (2026), Klaassen (2025), CyberArk (2025), Towards Data Science (2026) 107

Page 120

APPENDIX . RESEARCH METHODOLOGY Each agent maintained a source provenance log in the shared RuVector memory, recording the retrieval date, URL, and a confidence assessment for each source. Cross-referencing between agents was automated through semantic vector search: when Agent 2 (Block Critique) encountered a claim about AI productivity, the memory layer surfaced related findings from Agent 1 (Mollick / BCG) and Agent 5 (Every / Compound) for triangulation. .2 Source Selection and Quality Assessment Sources were categorised into four tiers based on methodological rigour and evidentiary weight. .2.1 Tier 1: Peer-Reviewed and Pre-Registered Research The highest evidentiary weight was assigned to peer-reviewed publications and pre-registered experiments. The primary load-bearing sources in this category are: • Dell’Acqua et al. (2026): pre-registered field experiment with 758 BCG consultants, published in *Organization Science*. This is the evidential foundation for the jagged frontier concept. • Randazzo et al. (2024): follow-up study with 244 BCG consultants identifying the Centaur, Cyborg, and Self-Automator collaboration modes. • Dell’Acqua (2023): field experiment with HR recruiters demonstrating the vigilance degradation problem. • Acemoglu and Restrepo (2019): the task-based automation framework that grounds the economic analysis of AI and labour. • Acemoglu (2025): the sceptical macroeconomic analysis of AI productivity, arguing for at most 0.66 per cent TFP gain over a decade. .2.2 Tier 2: Institutional Research with Disclosed Methodology Consultancy reports and institutional research with disclosed sample sizes and methodologies. These carry less evidentiary weight than peer-reviewed work but are treated as credible within stated limitations: • Boston Consulting Group and MIT Sloan Management Review (2025): joint research with MIT Sloan, n = 3 , 000+ across 25 industries. • Microsoft and LinkedIn (2024): large-scale survey with disclosed methodology. • World Economic Forum (2025): the World Economic Forum’s Future of Jobs Report. 108

Page 121

.3. WEB SEARCH AND ACADEMIC DATABASE APPROACH • Capgemini Research Institute (2025): middle manager resistance patterns with percentage-level findings. • Edmondson and Lei (2025): psychological safety and AI experimentation, grounded in the established literature on team effectiveness. Limitation: Consultancy reports are produced by organisations with commercial interests in the phenomena they describe. Methodology sections are often absent or sparse. Sample composition and response biases are rarely disclosed. Statistics from these sources are used directionally — to establish trends and orders of magnitude — rather than as precise measurements. .2.3 Tier 3: Expert Commentary and Primary Accounts Expert commentary, practitioner accounts, and journalistic reporting. These provide context and narrative but are not treated as primary evidence: • Dorsey and Botha (2026): primary source for the Block Intelligence Layer model. • Shipper (2026) and Klaassen (2025): primary accounts of the Every agent colony experiment. • Mollick (2025): expert synthesis from a researcher with deep domain expertise. • Spicer (2026) and Zamost

(2026): expert and insider critique of Block. .2.4 Tier 4: Community and Vendor Sources Sources with limited editorial oversight or commercial incentives. Used cautiously and flagged where they appear: • Towards Data Science (2026): the $17 \times$ error multiplication statistic. Flagged as non-peer-reviewed in Chapter 15. • IBM Security (2025) and Gartner, Inc. (2025): vendor-funded research with commercial incentives. Treated as directionally correct, per the methodological note in Chapter 15. .3 Web Search and Academic Database Approach Research was conducted through a combination of five source channels: 1. Academic databases : Google Scholar, SSRN, NBER Working Papers, HBS Working Papers. Search terms included “agentic organisation,” “AI coordination,” “human-AI collaboration,” “middle management AI,” “organisational design automation,” and “jagged technological frontier.” 109

Page 122

APPENDIX . RESEARCH METHODOLOGY 2. Consultancy publication portals : McKinsey Insights, BCG Publications, Deloitte Insights, Gartner Research. Filtered to 2024–2026 publications on organisational AI. 3. Practitioner platforms : Substack (particularly Mollick’s “One Useful Thing” and Every’s “Chain of Thought”), podcasts, and conference proceedings. 4. News and analysis : Financial Times, New York Times, Fortune, CNBC, VentureBeat. Used forevent-drivenreporting(Blockrestructuring, Klarnareversal)withcross-referencing against multiple outlets. 5. Regulatory sources : EUR-Lex for EU AI Act text, UK Legislation website for the Employment Rights Act, ICO guidance on workplace monitoring. .4 Codebase Audit Methodology The VisionClaw codebase was audited in April 2026 to reconcile claims made in earlier versions of this document against the actual code. The audit procedure: 1. Claim extraction : all factual claims about VisionClaw’s capabilities were extracted from the document and catalogued. 2. Code search : each claim was tested against the codebase using structural code analysis (call graphs, dependency trees) and text search. The codebase-memory MCP tool (48,159 nodes, 95,766 edges) was used for structural queries. 3. Verification categories : each claim was classified as (a) verified and accurate, (b) partially accurate with caveats, or (c) not supported by the code. 4. Correction : claims in categories (b) and (c) were corrected in Chapter 13. The corrections are presented transparently, with both the original claim and the corrected version visible. Conflict of Interest The codebase audit was conducted by the system architect (Dr. John O’Hare), who is also the report author and the founder of DreamLab AI Consulting. This introduces a conflict of interest: the auditor has a commercial interest in the system being described favourably. The audit found and corrected seven overclaims (documented in Chapter 13), which provides some evidence of good faith. Independent code review by a third party would strengthen the credibility of the technical claims but was not available at the time of writing. 110

Page 123

.5. WARDLEY MAPPING METHODOLOGY .5 Wardley Mapping Methodology Five Wardley maps were produced for this report, each following Simon Wardley’s standard mapping methodology: 1. Value chain identification : for each map, the user need was identified and the value chain of components required to serve that need was enumerated. → 2. Evolution

assessment : each component was placed on the evolution axis (Genesis → → Custom Product Commodity/Utility) based on its current maturity and market availability. 3. Movement indication : where components are expected to evolve (e.g., information routing moving from custom toward commodity), movement arrows indicate the direction and anticipated pace. 4. Strategic interpretation : each map was interpreted for strategic implications relevant to the report's argument. The maps are interpretive tools, not empirical measurements. The placement of components on the evolution axis reflects the author's assessment informed by the evidence assembled in this report. Different analysts may position components differently, particularly for novel concepts (judgment brokerage, mesh velocity) whose evolutionary stage is inherently uncertain. .6 Limitations of the Research Process 1. Single-author synthesis : while the research was conducted using multiple agents, the synthesis and editorial judgment are those of a single author. The report reflects one analytical perspective on a multi-faceted phenomenon. 2. Temporal constraints : the field is moving rapidly. Sources published after April 2026 are not included. The jagged frontier, in particular, has likely shifted since the foundational BCG study. 3. English-language bias : the research draws almost exclusively on English-language sources. Organisational innovations in non-Anglophone contexts—particularly Haier's micro-enterprise model in China — are discussed secondhand rather than from primary sources. 4. Analyst-vendor conflict : the author is the founder of DreamLab AI Consulting and the architect of VisionClaw. This conflict is acknowledged throughout the report and addressed directly in Chapter 15. ~ 5. Sample bias in practitioner evidence : the Every experiment (20 people) and Vision- 111

Page 124

APPENDIX . RESEARCH METHODOLOGY ~ Claw deployment (15–50 people) are both small-scale, technically sophisticated organisations. The generalisability of findings from these contexts to larger, less technical organisations is an open question. .7 Disclosure The author of this report, Dr. John O'Hare, is the founder and principal of DreamLab AI Consulting Ltd and the architect of the VisionClaw system described in Chapter 13. The report simultaneously analyses the problem space, proposes a solution framework, and describes a commercial offering from the author's firm. This positioning conflict is acknowledged in Chapter 15 and is offered as context for the reader's own evaluation of the argument. 112

Page 125

Competitor Comparison Matrix This appendix positions the Dynamic Agentic Mesh against six frameworks that the report's intended audiences are likely to be evaluating. The main body of the report positions the mesh against Block's Intelligence Layer and Every's Agent Colony — neither of which are consulting frameworks a COO or change manager would commission. This appendix addresses the gap identified in Chapter 15: differentiation from the frameworks the audience is actually considering. .8 Comparison Dimensions Seven frameworks are compared across six dimensions: 1. Coordination Model : how the framework proposes to handle the information routing function that hierarchy currently performs. 2. Management Role : what happens to middle management in the framework's model. 3. Governance Approach : how the framework addresses oversight, accountability, and com-

pliance. 4. AI Integration : the depth and nature of AI's role in the framework. 5. Scale Proven : the largest organisational context in which the framework has been deployed or validated. 6. Evidence Base : the type and strength of evidence supporting the framework's claims. .9 Comparison Tables The comparison is split across three tables for readability, each covering two or three frame- works alongside the Dynamic Agentic Mesh as the reference column. 113

Page 126

APPENDIX . COMPETITOR COMPARISON MATRIX Table 2: Competitor comparison: DreamLab Agentic Mesh, McKinsey Agentic Org, and SAFe 6.0 Dimension DreamLab Agentic Mesh McKinsey Agentic Org SAFe 6.0 Coordination Declarative governance + Five pillars: business model, Agile Release Trains with DAG orchestration with operating model, gover- AI-specific practices added OWL 2 ontology nance, workforce/culture, in 2024 technology/data Management Judgment broker at jagged "Above the loop": employ- Existing agile roles (RTE, role frontier intersections ees shift from performing Product Owner) extended tasks to orchestrating out- with AI practices comes Governance Embedded in architecture; Strategic governance pillar; Process-based: PI cadences, bias detection, security, framework-level guidance, inspect-and-adapt cycles value-alignment in every not architectural specifica- DAG transition tion AI integration Deep: OWL 2 ontology, Advisory: describes what or- Additive: AI practices agent orchestration, GPU vi- ganisations should do, not layered onto existing SAFe sualisation, Nostr identity how the AI system works structure Scale proven 15–50 people (VisionClaw) Consulting engagements Enterprise-wide (thousands across Fortune 500; no of developers) across multi- single reference implemen- ple industries tation Evidence base Codebase-audited system + Proprietary client data; "Six Broad adoption data; criti- third-party research synthe- Shifts" framework from con- cised for process overhead; sis sulting practice AI additions are recent 114

Page 127

.9. COMPARISON TABLES Table 3: Competitor comparison: Block Intelligence Layer, Every Agent Colony, and Holacracy 5.0 Dimension Block Intelligence Layer Every Agent Colony Holacracy 5.0 Coordination AI Company World Model + Cultural norms; public vis- Distributed authority CustomerWorldModel; cen- ibility in shared channels; through roles and circles; tralised intelligence emergent self-organisation constitutional governance Management Eliminated; replaced by ICs, "Everyone is a manager": No managers; authority dis- role DRIs, and Player-Coaches each employee manages tributed to roles; Lead Links their personal agent for operational coordination Governance AI-enforced through the In- Social accountability; agent Constitutional governance; telligence Layer; no sepa- trustderivativeofhumanrep- explicit policies and do- rate governance function de- utation mains; integrative decision- scribed making AI integration Foundational: AI is the coor- Deep at individual level: per- Minimal: Holacracy is a gov- dination mechanism itself sonal agents with memory, ernance system; AI integra- specialisation, personality tion is an add-on Scale proven 10,000+ employees (in- ~ 20 people Hundreds (Zappos at 1,500 tended, not fully validated) with significant attrition); version 5.0 addresses scale Evidence base 95% of AI code requires First-person practitioner ac- Decades of practice; aca- human modification;

Klarna count; ant death spiral docu- demic critiques of Zappos; analogue reversed; stock mented; no independent vali- updated methodology surge on announcement dation 115

Page 128

APPENDIX . COMPETITOR COMPARISON MATRIX Table 4: Competitor comparison: Microsoft Copilot Studio Dimension Microsoft Copilot Studio DreamLab Agentic Mesh (reference) Coordination Copilot agents deployed across Microsoft Declarativegovernance+DAGorchestration 365; orchestration through Microsoft Graph with OWL 2 ontology Management Unchanged; Copilot augments existing roles Judgment broker at jagged frontier intersec- role rather than transforming organisational struc- tions ture Governance Microsoft’s responsible AI framework; Entra Embedded in architecture; ontological prove- ID for agent identity; admin-controlled de- nance; Trust Variance monitoring ployment AI integration Deep platform integration; agents operate OWL 2 + Claude-Flow; cross-platform by within Microsoft ecosystem; limited cross- design; Nostr for decentralised identity platform Scale proven Enterprise-wide across Microsoft’s customer 15– 50 people base; largest deployment footprint of any framework Evidence base Microsoft’s internal productivity studies; Codebase-audited system + third-party re- third-party adoption data; vendor-published search synthesis metrics .10 Strategic Positioning Analysis .10.1 The Architecture Gap McKinsey describes what organisations should become (five pillars, six shifts) but does not spec- ify how the coordination mechanism works at the system level. SAFe 6.0 adds AI practices to an existing process framework but does not redesign the coordination architecture. Microsoft Copilot Studio provides a platform but does not propose an organisational model. The Dy- namic Agentic Mesh occupies the gap between strategic framework and operational platform — it specifies both the organisational model and the technical architecture. This is its primary differentiator and its primary vulnerability: it attempts more than the alternatives, and therefore has more to prove. .10.2 The Governance Gap Holacracy provides the most rigorous pre-AI governance model (constitutional governance with explicit policies and domains) but was designed before AI was a factor in organisational coor- dination. Block provides no separate governance function. Every relies on cultural norms. The 116

Page 129

.10. STRATEGIC POSITIONING ANALYSIS Agentic Mesh embeds governance architecturally — in the DAG transitions, in the ontological provenance chains, in the Trust Variance monitoring — and provides a specific human role (the judgment broker) for edge cases that exceed automated governance. This architectural embed- ding of governance is the feature that most clearly distinguishes the mesh from all competitors. .10.3 The Scale Gap The Agentic Mesh has the smallest proven deployment footprint of any framework in the comparison. Microsoft Copilot Studio has the largest. SAFe 6.0 is proven at enterprise scale. McKinsey’s framework is backed by Fortune 500 consulting engagements. This is the most significant competitive disadvantage for the Agentic Mesh and the primary research priority identified in Chapter 15. .10.4 The Evidence Gap The evidence bases differ in kind, not just in scale. The Agentic Mesh draws on third-party peer-reviewed research (Acemoglu 2025;

Dell'Acqua et al. 2026; Randazzo et al. 2024) for its theoretical foundations but relies on a single codebase-audited deployment for its architectural claims. McKinsey and BCG draw on proprietary client data that is not independently verifiable. SAFE draws on broad adoption data but is criticised for process overhead. Holacracy draws on decades of practice but also decades of critique. 117

Page 130

APPENDIX . COMPETITOR COMPARISON MATRIX .11 Wardley Map: Three Models Compared Figure 2: Wardley map comparing the three primary models (Block, Every, DreamLab) on the evolution axis. Block's Intelligence Layer positions AI coordination as a product/commodity; Every's Agent Colony treats coordination as genesis/custom; the Agentic Mesh positions the judgment broker at the custom stage while using commodity AI for information routing. 118

Page 131

.12. CHOOSING THE RIGHT FRAMEWORK .12 Choosing the Right Framework Selection Guidance • Already in SAFE : Consider integrating mesh concepts into existing PI planning — particularly the judgment broker role and the four KPIs. SAFE and the mesh are not incompatible; the mesh can operate within SAFE's cadence structure. • Microsoft-native : Start with Copilot Studio for individual augmentation; layer mesh governance concepts on top as agentic deployment matures. The two address different layers of the problem. • Greenfield or transforming : The Agentic Mesh provides the most comprehensive model but requires the most organisational commitment and technical expertise. The 90-day roadmap in Chapter 14 is designed for this context. • Seeking cultural transformation : Holacracy addresses power distribution directly; the mesh addresses information routing. They solve different problems and could in principle be combined. • C-Suite buy-in needed first : McKinsey's strategic framework provides the executive-level framing; the mesh can serve as the implementation architecture beneath it. • Enterprise scale with proven track record : Microsoft Copilot Studio offers the lowest-risk deployment path. The mesh offers the most ambitious organisational redesign. The right choice depends on the organisation's appetite for structural change versus incremental augmentation. .13 Interpretation No single framework in this comparison is complete. McKinsey provides the strategic vision but not the operational architecture. SAFE provides the process but not the AI-native coordination model. Holacracy provides the governance but not the AI integration. Microsoft provides the platform but not the organisational transformation model. Block provides the boldest thesis but the weakest evidence for its prescription. The Dynamic Agentic Mesh is differentiated by three features that no competitor offers in combination: 1. Ontology-grounded orchestration : formal OWL 2 semantics for agent routing, providing compositionality and inference that informal knowledge representations lack. 119

Page 132

APPENDIX . COMPETITOR COMPARISON MATRIX 2. Architecturally embedded governance : governance as code in every DAG transition, not as a separate oversight

function. 3. The compound flywheel : a specific, measurable mechanism (the Insight Ingestion Loop) by which the organisation's capability compounds over time, tracked by novel KPIs (Mesh Velocity, Augmentation Ratio, Trust Variance, HITL Precision). The primary weakness is scale. The primary research priority is bridging the gap between the 15-person existence proof and the enterprise deployments that competitors can claim. This comparison clarifies that gap. It does not close it. 120

Page 133

Glossary of Key Terms 1-9-90 Pattern The predictable adoption distribution identified by Boston Consulting Group (2025b) in AI deployment: 1 per cent are natural innovators who adopt independently, 9 per cent become champions with organisational support (the judgment broker candidates), and 90 per cent follow once social proof exists. Strategic investment in the 9 per cent yields the highest return. See Chapter 10. Agent Colony An organisational model (associated with Every.to) in which every employee is given a persistent, personalised AI agent. Coordination emerges from cultural norms and public visibility rather than architectural design. Proven at approximately 20-person scale; subject to the ant death spiral at larger scales. See Chapter 6. Agentic Mesh See Dynamic Agentic Mesh . Ant Death Spiral A failure mode in multi-agent systems where agents respond to each other's outputs in an uncontrolled loop, consuming resources (tokens, compute, time) without producing useful outcomes. Named by analogy with army ant death spirals in entomology. Documented by Shipper (2026) in Every's agent colony experiments at 20-person scale. See Chapter 6. Augmentation Ratio (AR) $\div$ A proposed KPI: $AR = \frac{\text{Cognitive Load Offloaded}}{\text{Total Cognitive Load}}$. Measures the proportion of decision-making volume successfully handled by agents without requiring human escalation. Target: $> 65\%$. Too low indicates the mesh is not yet effective; too high indicates potential vigilance degradation. See Chapter 8. BYOAI Bring Your Own AI — the practice of employees using personal AI tools and accounts for work purposes, bypassing corporate-approved systems. Microsoft and LinkedIn (2024) reports 78 per cent of knowledge workers engage in BYOAI. Centaur A human–AI collaboration mode identified by Randazzo et al. (2024) in which there is a clear division of labour between human and AI. The human performs tasks where human 121

Page 134

APPENDIX . GLOSSARY OF KEY TERMS judgment is strongest; the AI handles tasks inside the frontier. Distinct handoff points separate the two. The Dynamic Agentic Mesh is designed to produce Centaur organisations. Contrast with Cyborg and Self-Automator . See Chapter 4. Compound Engineering A principle articulated by Klaassen (2025): “Each unit of engineering work should make subsequent units easier — not harder.” The 50/50 Rule prescribes spending half of one's time building features and half improving the system that builds features. Extended in this report from the individual level to the organisational level as the basis for the compounding organisation. See Chapter 7. Compounding Organisation An organisation in which each unit of organisational learning makes the next one easier. The flywheel proceeds: workforce discovers a shortcut, agents detect and codify it, the judgment broker validates it, the pattern propagates across the mesh, and the organisation's capability

compounds. See Chapter 7. Coordination Collapse The central thesis of this report. For two thousand years, organisations have paid a “co- ordination tax” — layers of humans whose primary function is routing information. AI collapses the cost of that coordination function to near-zero, dissolving the structural logic of hierarchical organisational design. Chapter 15 notes the qualification that “information routing collapse” may be a more precise term. See Chapter 3. Coordination Tax The cost — in headcount, latency, distortion, and overhead — of maintaining hierarchical layers whose primary function is routing information between organisational levels. See Chapter 2. Cyborg A human–AI collaboration mode identified by Randazzo et al. (2024) characterised by continuous interleaving: human and AI work on the same analysis, passing work back and forth across the frontier boundary without clean separation of responsibilities. One of the two high-performing modes alongside Centaur. Every’s Agent Colony model produces Cyborg organisations. Contrast with Centaur and Self-Automator . See Chapter 4. DAG (Directed Acyclic Graph) A graph structure in which edges have direction and no cycles exist. In the Dynamic Agen- tic Mesh, DAGs represent formalised workflows: sequences of tasks with dependencies, ownership, governance constraints, and SLAs. Shadow workflows are codified as DAGs 122

Page 135

when promoted into the mesh via the Insight Ingestion Loop. See Chapter 6. Declarative Governance A governance model in which policies are expressed as code and enforced automatically at every system transition, rather than being applied through periodic human review. Bias thresholds, security protocols, access controls, and value-alignment checks are embedded in each DAG transition. Analogous to declarative infrastructure (Kubernetes, Terraform) in software engineering. See Chapter 9. Discovery Engine The bottom-up layer of the Dynamic Agentic Mesh that captures organisational insights. Passive agent monitoring surfaces shadow workflows; natural-language Insight Portals al- low anyone to propose improvements; tacit knowledge is codified into agentic instructions. The Discovery Engine is the input mechanism for the Insight Ingestion Loop. See Chap- ter 6. Dynamic Agentic Mesh The organisational architecture proposed in this report. Three layers (Discovery En- gine, Orchestration Layer, Declarative Governance) retain human judgment at intersec- tion points while agentic orchestration handles information routing. Middle managers evolve into judgment brokers who operate at the jagged frontier. See Chapter 6. HITL (Human-in-the-Loop) The practice of maintaining human oversight or decision-making authority within an auto- mated system. This report argues that naïve HITL implementations degrade as AI quality improves (the vigilance problem) and that HITL must be architecturally supported through deliberate friction and active engagement. See Chapter 5. HITL Precision (HP) $HP = \frac{\text{Correct Escalations}}{\text{Total Escalations}}$. Measures whether the edge cases the mesh flags for human intervention genuinely require human judgment. Tar- get: > 90 %. Low precision produces alert fatigue; high precision indicates the mesh has learned where the jagged frontier falls. See Chapter 8. Information Routing Protocol This report’s characterisation of hierarchy: a system designed to filter, aggregate, and relay information across organisational scales, constrained by the human cognitive limit of managing three to

eight direct reports. See Chapters 2 and 3. Insight Ingestion Loop The five-stage process by which the Dynamic Agentic Mesh converts individual discover-

Page 136

APPENDIX . GLOSSARY OF KEY TERMS

ies into organisational capability: (1) Discovery — an employee finds a shortcut; (2) Codification — the system maps it as a proposed DAG; (3) Validation — the judgment broker reviews for strategic fit and compliance; (4) Integration — the validated workflow enters the live mesh with SLAs; (5) Amplification — the pattern propagates to other teams. Mesh Velocity measures how fast this loop runs. See Chapters 6 and 7. Intelligence Layer The organisational model proposed by Dorsey and Botha (2026) in which middle management is eliminated and coordination is handled by AI world models (Company World Model and Customer World Model) composed by an Intelligence Layer. Three roles remain: Individual Contributors, Directly Responsible Individuals, and Player-Coaches. See Chapter 6. Jagged Frontier The unpredictable boundary between tasks that AI performs well and tasks that AI performs poorly. Identified by Dell’Acqua et al. (2026) based on a pre-registered experiment with 758 BCG consultants. The frontier is “jagged” because adjacent, similar-looking tasks can fall on opposite sides. This concept is the intellectual load-bearing wall of the report: it is the reason wholesale automation is reckless, change playbooks cannot be written once, and judgment brokers have a genuine future. See Chapter 4. Judgment Broker The evolved middle management role proposed in this report. The judgment broker operates at the jagged frontier, providing strategic direction, ethical adjudication, and relational intelligence that agentic systems cannot perform. Distinguished from the traditional information-routing role by its focus on edge cases, cross-functional coherence, and frontier navigation. See Chapter 12. Mesh Velocity (MV) A proposed KPI: $MV = \Delta t (\text{insight} \rightarrow \text{codified workflow})$. Measures the elapsed time from insight discovery to codified, sanctioned workflow in the mesh. Target: < 48 hours. The primary indicator of organisational learning speed. See Chapter 8. Ontological Provenance The property of an agentic system in which every decision traces back through a formal knowledge graph (OWL 2 ontology), providing auditable reasoning chains from output to input, policy, and human decision. See Chapter 9. OWL 2 Web Ontology Language 2 — a W3C standard for formal knowledge representation. In 124

Page 137

VisionClaw, OWL 2 provides shared semantics for agent skill routing through ontological subsumption (understanding that “risk assessment” is a sub-task of “governance review” and routing accordingly). The EL++ profile, supported by the *Whelk-rs* reasoner, enables tractable inference while maintaining semantic precision. See Chapter 13. Secret Cyborg A term from Mollick (2023) describing workers who use AI tools without disclosing this use to their employer. The secrecy is driven by rational calculation: disclosure risks headcount reduction if the worker demonstrates they have automated significant portions of their role. Over half of generative AI users report using the technology covertly at least some of the time. See Chapter 3. Self-Automator A human–AI collaboration mode identified by Randazzo et al. (2024) in which the human delegates wholesale to AI with minimal oversight. The

lowest-performing mode of the three. Block's organisational model is characterised in this report as creating Self- Automator organisations. Contrast with Centaur and Cyborg . See Chapter 4. Shadow AI AI tools and applications used by employees without organisational awareness, approval, or governance. Shadow AI usage increased 156 per cent between 2023 and 2025; the average enterprise hosts approximately 1,200 unauthorised AI applications (Gartner, Inc. 2025). See Chapter 3. Shadow Workflow An informal work process created by an employee using AI tools, which bypasses or re- places the officially documented workflow without management awareness. Shadow work- flows represent both a risk (ungoverned, undocumented) and an asset (raw material for the Insight Ingestion Loop and the compounding flywheel). See Chapter 3. Trust Variance (TV) A proposed KPI: $TV = \sigma$ (Agent Decision Quality) over a rolling 30-day window. Target: $< 0.12 \sigma$. Provides real-time monitoring of drift or bias in the automated task layer. The early warning system for vigilance degradation — catching drift before it compounds into failure. See Chapter 8. Vigilance Problem The empirically demonstrated tendency of human oversight quality to degrade as AI qual- ity improves. Identified by Dell'Acqua (2023): HR recruiters given high-quality AI made worse decisions than those given low-quality AI because they stopped scrutinising outputs.

125

Page 138

APPENDIX . GLOSSARY OF KEY TERMS The central challenge of human-AI governance. See Chapter 5. VisionClaw The operational agentic mesh system built by DreamLab AI Consulting. Architecture: OWL 2 ontology engine (horned-owl , Wheelk-rs EL++ inference), Claude-Flow or- chestration (83 agent skills, RAFT consensus), WebXR 3D visualisation (Rust/CUDA, Babylon.js), Nostr agent identity (NIP-98), and RuVector memory (1.17M+ vector em- beddings, HNSW indexing). Proven at 15-person scale in a creative technology context. See Chapter 13. World Model In Block's Intelligence Layer thesis (Dorsey and Botha 2026), the Company World Model is an AI system continuously tracking what is being built, what is blocked, and where resources are allocated. The Customer World Model is constructed from transaction data. Together they replace the context-holding function of middle management. The concept is analytically sound but operationally unproven beyond Block's specific context of massive transaction data and remote-first culture. See Chapter 6.

126

* * *

All embedded URLs (in document order)

- <https://www.nber.org/papers/w32487>
- <https://www.bcg.com/publications/2025/the-emerging-agentic-enterprise>
- <https://www.bcg.com/publications/2025/unlocking-ai-first-organization-agentic-shift>
- <https://www.bcg.com/>
- <https://www.bcg.com/publications/2025/ai-powered-kpis>
- <https://www.capgemini.com/insights/research-institute/>
- <https://www.cyberark.com/resources/>

- <https://www.hbs.edu/faculty/Pages/item.aspx?num=64700>
- https://papers.ssrn.com/sol3/papers.cfm?abstract_id=4573321
- <https://www2.deloitte.com/global/en/insights/focus/human-capital-trends.html>
- <https://block.xyz/inside/from-hierarchy-to-intelligence>
- <https://hbr.org/>
- <https://www.lawandtheworkplace.com/2025/12/major-reforms-for-uk-employment-law-era-2025/>
- <https://www.europarl.europa.eu/news/en/press-room/20251110IPR31358/>
- <https://epoch.ai/blog/half-of-employed-ai-users-now-use-it-for-work/>
- <https://www.gartner.com/en/articles/gartner-top-strategic-predictions>
- <https://rework.withgoogle.com/>
- <https://deepmind.google/>
- <https://www.holacracy.org/>
- <https://www.ibm.com/reports/data-breach>
- <https://every.to/chain-of-thought/compound-engineering>
- <https://www.cnn.com/>
- <https://www.mckinsey.com/capabilities/mckinsey-digital/our-insights/six-shifts-agentic-organization>
- <https://www.mckinsey.com/capabilities/mckinsey-digital/our-insights/the-agentic-organization>
- <https://www.microsoft.com/en-us/worklab/work-trend-index/>
- <https://www.microsoft.com/en-us/microsoft-copilot/microsoft-copilot-studio>
- <https://www.microsoft.com/en-us/worklab/work-trend-index/ai-at-work-is-here-now-comes-the-hard-part>
- <https://www.oneusefulthing.org/p/detecting-the-secret-cyborgs>
- <https://www.oneusefulthing.org/>
- <https://www.nature.com/natsustain/>
- <https://www.orgvue.com/>
- 2509.20241
- <https://arxiv.org/abs/2509.20241>
- <https://www.prosci.com/>
- https://www.hbs.edu/ris/Publication%20Files/Cyborgs-Centaurs-Self-Automators_5a8e0b5c-1a4e-4f8b-9a8a-6c3e3e1a7d8f.pdf
- <https://scaledagileframework.com/>
- <https://every.to/chain-of-thought/we-gave-every-employee-an-ai-agent>
- <https://fortune.com/>
- <https://onlinelibrary.wiley.com/journal/10970266>
- <https://towardsdatascience.com/>
- <https://www.imd.org/>
- <https://www.weforum.org/publications/the-future-of-jobs-report-2025/>

- <https://www.nytimes.com/>

*Appendix L — The Programmable Brand:
Behavioural Licensing and the Governance
Layer of the Generative Economy*

Companion piece to Chapter 5: *The Slop Ceiling*, Chapter 6: *The 88%*, Chapter 7: *The Studios Decide*, Chapter 12: *Authenticity as the New Scarcity*, and Chapter 16: *The Tools*.

This deep dive is the IP-economics companion to the book's argument about where value, power and labour are actually migrating in the generative era. Where Chapter 6 develops the *Petrillo template* and the licensing-first response of the creator coalition, Chapter 7 maps studio strategy, and Chapter 12 builds the authenticity-as-scarcity case, this appendix supplies the underlying *contract architecture* that the 2026 deals have begun to settle on — and the conceptual move that organises them: **IP is no longer an asset to be guarded; it is an environment to be governed.**

The piece centres on three reference deals — Hasbro's **Sixth Wall** studio and the *Behavioral Licensing* of character personality via **CharacterOS** (3 June 2026); the **Spotify–Universal Music Group** *consent / credit / compensation* framework for fan-made AI covers and remixes (21 May 2026); and the **Disney–OpenAI** bn licensing agreement that was announced on 11 December 2025 and *never executed* — and uses the *three Cs* (consent, credit, compensation) as a constant lens across five emerging licensing models. It is the most rigorous available framing of the contract-side of the generative economy in mid-2026.

The Disney–OpenAI material here is significant for the rest of the book and worth flagging explicitly. The deal is referenced in several chapters (Chapter 1, Chapter 7, Chapter 8, Appendix C) as if it had been signed and implemented. On the corrected record this appendix sets out, the deal was *announced* on 11 December 2025 but *no formal agreement was ever signed and no money changed hands*; OpenAI moved to shut down Sora on **24 March 2026** with the consumer app and website ending in late April and the developer API set to close in September; Disney was reportedly given as little as **thirty minutes' notice**, exited the partnership, and confirmed it would continue to engage with AI platforms that respect IP and creator rights. Where the surrounding chapters were written before this collapse, the deal stands as evidence of *contractual architecture* — the no-training clause, the talent-likeness exclusion, the joint steering committee, the equity-and-warrants structure — rather than as evidence of operative integration. That contractual architecture is itself the appendix's strongest analytic finding: it survived the platform's death and is now the reference architecture other studios negotiate against.

The other consequential observation, for the rest of the book, is the appendix's reframe of where transaction costs went under generative tooling. The standard claim — that AI compresses transaction costs to near zero — is *half right*. Search and negotiation costs do collapse (the catalogue is right there in the app; the terms are programmatic and accepted with a click). But monitoring and enforcement costs *rise sharply*, because the platform has authorised creation at a scale no human legal team can police by hand. The new costs migrate into infrastructure — provenance tracking, identity verification, behavioural guardrails, royalty attribution — and *the governance layer is where the durable value sits*. Hasbro's CharacterOS, Spotify's *Verified by Spotify* badge and opt-in catalogue, the C2PA/SynthID stack the book has been describing throughout — these are not creative technologies dressed as governance. They are governance technologies dressed as creative ones. The brands that

win are the ones that stop guarding the vault and start building (and charging admission to) the world. That, in a sentence, is the IP-economics half of the *authenticity-as-scarcity* thesis the book's second half is built on.

Read alongside Appendix C (Bibliography by Topic) §V (Copyright and the legal architecture), §VI (Studio strategy) and §IX (Authenticity, provenance and disclosure), this appendix forms the contract-and-governance spine of the second half of the book.

The piece below is preserved as the *Dream Machine* Deep Dive № 14 (4 June 2026) was published, with citation markers and section structure intact. Some PDF-conversion artefacts (loose footnote numbers, occasional line-break oddities) have not been editorially cleaned; the analytical content is what matters.

* * *

Page 1

■ DREAM MACHINE DEEP DIVE № 14 · 04 · 06 · 2026 I N T E L L E C T U A L P R O P E R T Y / T H E G E N E R A T I V E E C O N O M Y The Programmable Brand For two years the story about generative AI and creativity was a story about replacement. The deals being signed in 2026 tell a different one — about licensing, and about who gets to own the rails. ~6,200 WORDS SOURCES VERIFIED TO 04 JUNE 2026 ANALYSIS O n the third of June, 2026, Hasbro did something that would have been unthinkable to a media lawyer five years ago: it invited the world to put words in Optimus Prime's mouth. Not Hasbro's words — yours. Through a new in-house studio called Sixth Wall, the company began licensing not the appearance of its characters but their behaviour — how Megatron argues, how Cobra Commander menaces, how the cast of *Clue* insinuates — for deployment inside chatbots, games, toys and brand experiences built by third parties. It is the first time a major rights holder has formally 1 treated a character's personality as a programmable, rentable asset. And it is the clearest signal yet that the entertainment industry has quietly abandoned the war it spent 2023 and 2024 fighting. That war had a simple plot. Generative AI was a thief; the studios were the police; the courts were the battleground. The dominant frame — in the press, in the lawsuits, in the anxious op-eds — was substitution. Machines would replace illustrators, voice actors, session musicians, screenwriters. The only question was the body count. This framing was not wrong so much as it was looking at the wrong table. The substitution story is about cost- cutting inside large producers. The far larger economic event is happening somewhere else entirely: in the creator economy — the loose, sprawling, semi-professional mass of independent makers, micro-brands and fan communities that has grown into something on the order of a quarter-trillion-dollar market. The real 2 generative opportunity is not replacing the people who make culture. It is licensing the raw material of culture to the tens of millions of people who already want to remix it — and finally being able to get paid when they do. This piece is about how that shift became possible, why it failed for two decades before it suddenly worked, and what the deals of 2026 — Hasbro, Spotify–Universal, the spectacular collapse of Disney–OpenAI — reveal about where value, power and labour are actually migrating. The short version: the brands that win will be the ones that stop treating their IP as a vault to be guarded and start treating it as an environment to be governed.

Page 2

The interesting question was never whether the machine can make the thing. It is who controls the conditions under which a million strangers make the thing for you. 0 1 — T h e M i s r e a d S t o r y From substitution to participation Begin with the audience, because the audience is where the demand actually lives. The defining behavioural fact of the last fifteen years of digital culture is that people stopped wanting to merely consume media. They wanted to handle it — to remix the song, to write the fan fiction, to make the meme, to put the character in a situation its creators never sanctioned. Henry Jenkins gave this its name two decades ago: participatory culture , a media environment with low barriers to creation and a strong social incentive to share what you make. Generative 3 tools did not invent that impulse. They industrialised it. They handed the participatory instinct a factory. For rights holders, this presented a genuine dilemma rather than a simple threat. Every unauthorised remix is, simultaneously, a copyright violation and a piece of free marketing — evidence of exactly the emotional attachment that makes an IP valuable in the first place. The legacy response was to treat the violation as load-bearing and the affection as incidental: cease-and-desist, takedown, lawsuit. Disney, tellingly, sent Google a cease-and-desist over alleged unauthorised training on its catalogue before it ever sat down to negotiate with OpenAI. Enforcement was the reflex. 4 The pivot now underway is from that reflex to its inverse. Instead of asking how do we stop people using our characters, the progressive rights holder asks how do we build the meter that runs while they do. Hasbro's CEO Chris Cocks framed Sixth Wall precisely as a “third way” between endless, costly enforcement and passive brand dilution. 1 The shift is institutional, not merely tactical: it requires reconceiving the IP itself, from a finished product shown to consumers into a set of permissions, guardrails and APIs offered to co-creators. 0 2 — T h e E c o n o m i c s Why it failed for twenty years, and what changed The most useful lens for understanding the failure — and the sudden success — is the oldest one in the theory of the firm. In 1937 Ronald Coase asked why firms exist at all, and answered: because using the market is not free. 5 Every transaction carries hidden costs — of searching for a counterparty, negotiating terms, monitoring performance, enforcing the bargain. Oliver Williamson built this into a full account of why some exchanges happen inside organisations and others across markets. When the friction of transacting exceeds the value at stake, the transaction simply doesn't happen. This is exactly the trap that licensed fan creation fell into. Imagine a bedroom producer who wants to legally build a track around a Universal master, or a tiny studio that wants to put an authorised Transformers

Page 3

character in a mobile game. The economics are hopeless. The cost of finding the rights holder, reaching the right desk, negotiating a bespoke contract and policing the output dwarfs any plausible revenue from a niche creation. So the market failed in the textbook way: either the work never got made, or it got made illegally. Billions of derivative works took the second path, into a grey zone of copyright strikes and lost revenue on both sides. The standard telling — including in the research this piece builds on — is that generative platforms “compress transaction costs to near zero” and the market simply clears. That is half right, and the wrong half is the interesting one. L O A D - B E A R I N G C L A I M , S O F T E N E D AI does not

abolish transaction costs. It relocates them. Embedding licensing into a generative workflow genuinely collapses the cost of search (the catalogue is right there in the app) and negotiation (terms are programmatic, accepted with a click). But it raises the cost of monitoring and enforcement — because you have now authorised creation at a scale no human legal team could ever police by hand. The new costs migrate into infrastructure: provenance tracking, identity verification, behavioural guardrails, royalty attribution. The deals of 2026 are best read not as “transaction costs went to zero” but as “the industry finally figured out where to put the meter and the fence.” Seen this way, the headline innovations of 2026 are not the generative models at all. They are the governance layers. Hasbro’s CharacterOS — the system that enforces a character’s canon, voice and safety bounds in real time — is a monitoring-and-enforcement technology dressed as a creative one. Spotify’s “Verified by Spotify”¹ badge and its opt-in catalogue are enforcement-by-architecture. The programmatic metadata linking every 6 remix to its source is monitoring at machine scale. The value the platforms add is not generation. Generation is a commodity. The value is the fence. The second theoretical leg is the resource-based view of strategy: firms win durably by combining resources that are valuable, rare, hard to imitate and hard to substitute.⁷ The creator economy’s scarce resources are split awkwardly between two parties. Legacy holders own the rare, emotionally resonant assets — the iconic voice, the beloved character, the master recording. Independent creators own the things the legacy holders are worst at: relational trust with a specific audience, native fluency in platform formats, and speed. Neither can fully exploit the AI moment alone. The structured licence is the contract that lets them combine — the studio brings Megatron, the developer brings the audience and the app.

Page 4

Five models, one question To compare the emerging models without drowning in press-release language, it helps to hold each one against a single test. The Spotify–Universal deal supplied the cleanest articulation of it — a framework built on three pillars: consent, credit, compensation. I’ll use those three Cs as a constant lens. A model is robust to the 6 degree it delivers all three; the failure modes are visible in which C gets dropped. ■ HASBRO & ELEVENLABS — BEHAVIORAL LICENSING CONSENT · TALENT CREDIT COMPENSATION Sixth Wall’s genuinely new idea is in the noun. Traditional licensing governs where a character appears — on a lunchbox, in a film, on a T-shirt. Behavioral Licensing governs how it behaves : how it thinks, speaks and reacts inside a live, unscripted interaction. That is a category that simply did not need to exist until characters could 1 improvise. The enabling technology, CharacterOS, encodes each character’s personality, canonical lore, vocal timbre and safety limits, and enforces them at runtime — so a licensed Cobra Commander stays menacing and on-brand without a writer in the loop. The studio’s origins are revealing. The project grew out of an internal experiment: an interactive digital Ouija board that, by the launch team’s account, logged hundreds of thousands of player interactions within its first few days — proof that fans would engage deeply, safely and at volume with a generative character if the experience felt authentic. CEO Roberta Thomson describes the design intent as preserving personality, “not 1 just the voice.”⁸ On the three Cs, this is the most complete model of the five. Voices are built from authorised recordings and

real human performances; a talent participation structure compensates the original performers for the generative use of their work. ¹ That matters enormously, and it is also where the sharpest critique lives. “Participation” is a compensation mechanism — it answers credit and compensation. It does not, by itself, settle consent in the deeper sense the actors’ unions mean: ongoing control over how a performance is used, and the right to refuse. A pay cheque is not a veto. The model’s legitimacy will rest on terms the press releases don’t disclose. One more thing the timing tells us. Sixth Wall launched into a regulatory climate where the U.S. Federal Trade Commission has opened inquiries into AI companions and their effects on minors, and the studio pointedly restricts its initial focus to users 13 and over and enterprise use, explicitly not building for young children. The 9 guardrail is not only a brand-safety feature. It is a liability shield. ■ SPOTIFY & UNIVERSAL MUSIC GROUP — PLATFORM-EMBEDDED REMIXING CONSENT · OPT-IN CREDIT COMPENSATION On 21 May 2026, at Spotify’s Investor Day, Spotify and Universal announced linked recorded-music and publishing agreements to let fans create AI covers and remixes of participating artists’ songs — a tool to ship as a paid add-on for Premium subscribers, with revenue flowing back to the artists and songwriters whose work is reinterpreted. The market read it as a direct shot at Suno; Spotify’s stock jumped on the day. ⁶ ¹⁰

Page 5

The structure maps onto the three Cs almost too neatly, because UMG wrote the slogan. Consent is opt-in: only the catalogues of artists who actively agree are available for remixing. Credit is programmatic: derivatives are linked to original metadata. Compensation is a new royalty stream layered on top of standard streaming. ⁶ What makes this deal more interesting than its slogan is its prehistory. UMG spent 2024 calling Suno and Udio’s training “infringement on an almost unimaginable scale.” Then, in October 2025, it settled with Udio ¹¹ and agreed to build a licensed product instead. The Spotify deal is the mature form of that turn: litigation ¹² converted into a revenue line. The lesson generalises — the lawsuits were never the destination. They were leverage to force a licensing table into existence. T H E T E N S I O N T H I S M O D E L C A N ’ T F U L L Y R E S O L V E MIDiA Research has a useful term for the deeper risk here: “create to consume.” When ¹³ generating a personalised song is as easy as requesting one, music shifts from a thing you listen to into a thing you make for a moment and discard. A remix economy that flatters the superfan may, at scale, erode the very catalogue value it is monetising — by training a generation to treat songs as raw material rather than finished works. The royalty share answers compensation. It does not answer cannibalisation. ■ DISNEY & OPENAI — THE CAUTIONARY TALE CONSENT CREDIT EXECUTED This is the deal everyone cited as the template — and it is worth getting exactly right, because the popular version is wrong in a way that matters. On 11 December 2025, in one of Bob Iger’s final major moves, Disney and OpenAI announced a landmark three-year licensing agreement: Sora users would be able to generate short, prompted social videos drawing on more than 200 animated, masked and creature characters from Disney, Marvel, Pixar and Star Wars; ChatGPT Images would get comparable access for stills; a curated selection would stream on Disney+. The governance ¹⁴ was unusually careful. The licence explicitly excluded real actors’ voices and likenesses — protecting publicity rights and

avoiding a fight with SAG-AFTRA. OpenAI was barred from training its foundation models on the licensed IP, limiting the deal to output, not ingestion. A joint steering committee was to oversee character usage and safety. And Disney attached a reported \$1 billion equity investment, plus warrants — positioning the studio not as a vendor but as a part-owner of the infrastructure. ¹⁴ And then none of it happened. The framework was immaculate. The platform underneath it evaporated in ninety days.

Page 6

On 24 March 2026, OpenAI told staff it was shutting Sora down — a staged wind-down, with the consumer app and website ending in late April and the developer API set to close in September. The reasons were brutally ¹⁵ economic: consumer video at current GPU prices does not pencil — peak Sora revenue of a few hundred thousand dollars a month against operating costs in the billions — and OpenAI chose to reallocate scarce compute toward higher-margin coding and enterprise products amid intensifying competition. ¹⁶ Disney, by multiple accounts, learned of the shutdown with as little as thirty minutes' notice and exited the partnership. ¹⁷ Crucially: no formal agreement was ever signed, and no money changed hands. The billion-dollar template was a billion-dollar press release. ¹⁶¹⁸ It would be easy to read this as a refutation of the whole thesis. It is the opposite. The collapse cleanly separates two things the hype conflated: the licensing framework and the platform implementation. The platform was fragile — subject to one company's compute economics and strategic whims. The framework Disney's lawyers built — IP-not-training, talent-excluded, jointly governed, equity-aligned — survived the platform's death and is now the reference architecture other studios negotiate against. Disney itself signalled it would keep engaging with AI platforms on those terms. The durable asset was the contract structure, not the vendor. ¹⁸ **THE G O V E R N A N C E L E S S O N E V E R Y O N E T O O K F R O M I T** The thirty-minutes'-notice detail became a case study in strategic-investment risk: a \$1bn pledge with, apparently, no board seat or information rights strong enough to give Disney warning that the product it was licensing into was about to be killed. The practical ¹⁷ takeaway for every rights holder that follows: align on equity and governance rights, not just licence terms — or build the platform yourself. Which is exactly what Hasbro did. ■ **THE INDEPENDENTS — COMMONS, DAOS, ESTATES, AND THE RENT-SEEKING ANTI-PATTERN** Away from the corporate megadeals, individual artists have been quietly running the most radical experiments — and they bracket the design space at both ends. At the open extreme sits Grimes. Through Elf.tech, launched in 2023, she turned her own voice into a shared resource: upload your vocal, get back a synthesised “GrimesAI” performance, release it commercially, and split master royalties roughly fifty-fifty with her. Her team explicitly analogised it to fan-fiction and fan-art ¹⁹ cultures — a vocal identity treated as an open-source brand that earns passively while strangers extend it. It is, in effect, commons-based peer production with a royalty hook: the resource is opened, and value is captured at the point of commercial distribution rather than at the point of creation. Holly Herndon's Holly+ takes the same raw idea — a public digital twin of a voice — and wraps it in collective governance, vesting decisions about commercial use and the distribution of proceeds in a DAO rather than in the artist alone. Herndon has paired this with consent-layer tooling (the “Have I Been Trained?” project) ¹⁹ aimed at giving

all creators visibility into whether their work sits inside training sets. 20 Where Grimes optimises for openness, Herndon optimises for legitimacy — community consent as a feature, not a footnote.

Page 7

At the heritage end, estates have learned to license the dead. ElevenLabs, with CMG Worldwide, secured rights to recreate the voices of figures such as Judy Garland, James Dean, Burt Reynolds and Laurence Olivier — but confined them to a single text-to-speech Reader app, walled off from the general voice-cloning database, so the estates retain control and reputational protection. The same marketplace has since drawn in living and legacy 21 names alike, from Michael Caine to the Stan Lee estate. 8 The principle is containment: a licence is only as valuable as the boundary it draws around the asset. And then the anti-pattern. In 2023 Meta paid celebrities — Snoop Dogg, Kendall Jenner, MrBeast and others — reportedly up to several million dollars each to clone their likenesses into branded chatbots. On the three Cs 22 it is instructive precisely because it nails two and abandons the spirit of the whole thing. Consent: yes, lavishly paid. Compensation: yes, up front and flat. But there is no participatory layer at all — the personas are closed, static, fully platform-controlled, offering the user and the wider creator community no agency whatsoever. It is celebrity-likeness rental, not a creative environment. The contrast with Grimes is total: one artist opened her voice to fifteen thousand strangers' songs; the platform rented a frozen mask. ■

SUNO & THE CASUAL CREATOR — THE DEMAND FLOOR Finally, the model that explains the demand the others are racing to capture. Suno turns a text prompt into a finished song. Professionals use it to prototype; its real growth engine is the casual creator — non-musicians making a birthday song for a partner, a personalised rap “by” a child’s favourite cartoon character. Amazon 1323 wired exactly this into Alexa+: speak a request, receive a custom song on the spot. 23 This is where “create to consume” stops being a risk and becomes the entire point. The output isn’t competing with the charts; it’s competing with a greeting card. And it surfaces the deeper shift the whole essay turns on: when high-fidelity audio becomes effortless, the audio stops being the scarce thing. Value migrates outward — to the cover art, the narrative, the world-building, the human identity wrapped around the sound. 24 Abundance at the centre raises the premium on everything at the edges that a model can’t fake: taste, trust, story, a name people care about.

Page 8

FIVE LICENSING MODELS AGAINST THE THREE CS · JUNE 2026 ECONOMIC GOVERNANCE MODEL PRINCIPALS STRUCTURE MECHANISM THE THREE CS Hasbro & B2B programmatic CharacterOS — runtime Behavioral All three; consent ElevenLabs licence; talent- personality, canon & Licensing shallow (pay ≠ participation pay-out safety enforcement veto) Embedded remixing Universal & Paid Premium add- Opt-in catalogue; All three; Spotify on; new royalty “Verified by Spotify” cannibalisation stream to artists badge risk unpriced Enterprise co- Disney & \$1bn equity + Joint steering Framework survived; OpenAI warrants; 3-yr term committee; no-training creation platform died (never executed) clause; talent excluded Open-source voice Grimes / ~50% master-royalty Programmatic

Radical consent; Elf.tech split with creators attribution via light enforcement metadata Meta & High flat fee Closed, platform- Likeness rental Consent + comp; celebrities (reportedly up to controlled personas zero participation \$5m / 2 yrs) o 4 — S e c o n d - O r d e r E f f e c t s Where power, value and labour are moving ■ 1 · POWER: FROM VENDORS TO OWNERS TO RAIL-BUILDERS Through the streaming era, distribution platforms captured most of the value creators generated; the talent supplied content and the platform set the terms. The generative era inverts the dependency. Now the platforms are the supplicants — they need authorised, recognisable, litigation-proof IP to differentiate their models and stay out of court. That hunger is leverage, and rights holders are spending it in three escalating ways. 14 The mildest is the better licence — UMG’s royalty share. The stronger is the equity stake : Disney’s billion-dollar position would have made it a part-owner of the infrastructure rather than a mere supplier. But the Disney collapse exposed the limit of that move — a stake without governance rights is a stake in someone else’s decision. The strongest response, and the one Hasbro chose, is vertical integration : don’t rent your IP to someone else’s fragile platform, build the studio and own the operating system yourself. CharacterOS is not just a guardrail. It is Hasbro declining to be a vendor. ■ 2 · VALUE: THE BIFURCATION AND THE PREMIUM ON THE REAL As frictionless generation floods every feed with low-cost synthetic “slop,” the scarce, appreciating asset becomes authenticated human creation. The digital landscape is bifurcating into two zones. On one side, authorised sandboxes — verified, licensed, monetised environments where platforms use provenance signals (a

Page 9

“Verified by Spotify” badge, content-origin labels, a CharacterOS stamp) to steer attention and money toward sanctioned work. On the other, the open web — a vast, fragmenting expanse of unlicensed synthetic output facing demonetisation, filtering and legal exposure. In that world the media company’s role mutates from publisher to IP wholesaler : it licenses brands to a decentralised creator base and relies on cryptographic verification and provenance rails to track authenticity and royalty flows across platforms it doesn’t own. Value accrues to whoever controls the boundary between the two zones — the verification layer — which is precisely why the governance technologies, not the generative ones, are the strategic prize. ■ 3 · LABOUR: THE COSTS DIDN’T VANISH, THEY CHANGED SHAPE Recall the relocated-transaction-cost argument from the start. Its human consequence is that the work doesn’t disappear — it moves. The session musician’s hours don’t simply evaporate; demand shifts toward voice performers who supply the authorised source recordings, toward the people who curate canon and write guardrails, toward provenance and rights-management roles, toward the world-builders and brand stewards who do the thing models can’t. This is not a consolation; some of those new roles are fewer and differently distributed than the ones they replace, and “talent participation” is a far weaker thing than employment. But it is the actual shape of the disruption, and it is more interesting — and more contestable — than “the robots took the jobs.” ■ 4 · LAW: THE UNBUNDLING OF THE SELF All of this is dragging the law behind it at speed. Behavioral licensing and synthetic-voice replication are accelerating efforts — by SAG-AFTRA, by state legislatures, by federal bills like the proposed NO FAKES framework — to establish statutory rights over digital likeness and voice for the

living and the dead alike. The deep trend is an unbundling of identity : voice, physical likeness, performance style and written-character dialogue are increasingly segmented and licensed as separate assets under bespoke, automated contracts. Disney licensed its characters while ring-fencing its actors' voices. Hasbro licenses behaviour while paying for the underlying performance. The self is being disaggregated into a portfolio of individually tradable rights — which is liberating and alarming in roughly equal measure, and which the statute books have not remotely caught up with. 0 5 — C o n c l u s i o n The brand as environment Step back from the deals and the through-line is a single conceptual move. For a century, intellectual property was understood as an asset : a discrete, ownable thing whose value you protected by controlling access to it. The scarce resource was the work, and the business was distribution. What 2026 reveals is a redefinition of IP as an environment : a bounded space, governed by rules, inside which value is created continuously by people who are not you. The work is no longer the product; the conditions of

Page 10

participation are the product. Hasbro doesn't sell you Optimus Prime. It sells developers a governed space in which Optimus Prime can be summoned, on-brand and on-licence, a million times over. The asset is the rulebook. This is why the substitution panic missed the story so badly. It was measuring the wrong scarcity. The machine can generate the character, the voice, the song — that capability is becoming free. What stays scarce, and therefore valuable, is everything the machine can't supply on its own: the canon worth preserving, the audience that cares, the trust that says this one is real, and the governance that lets a brand open its gates to strangers without losing itself in the process. The Disney–OpenAI deal collapsed because it bet on a platform. Hasbro's bet, and Universal's, is subtler and probably more durable: that in an age of infinite generation, the enduring power belongs not to whoever can make the content, but to whoever owns the environment in which it's allowed to count. The brands that grasp this will stop guarding the vault. They'll start building the world — and charging admission. In an age where anyone can generate anything, the last scarce resource is permission. N O T E S & S O U R C E S 1. Hasbro, "Hasbro Launches Sixth Wall, a New AI Studio...", BusinessWire press release, 3 June 2026; coverage in StockTitan, Investing.com and The Wrap (Cocks "third way" framing; CharacterOS; 12 characters; talent participation; Ouija pilot origin). [businesswire.com](https://www.businesswire.com); [thewrap.com](https://www.thewrap.com) 2. Goldman Sachs Research, "The creator economy could approach half-a-trillion dollars by 2027" (~\$250bn base, ~50m creators, \$480bn by 2027). Note: estimates vary widely by methodology — some firms count 200m+ total content creators and project \$480–528bn by 2027–2030. Treat single figures as directional, not precise. [goldmansachs.com](https://www.goldmansachs.com) 3. Henry Jenkins, *Confronting the Challenges of Participatory Culture* (and earlier work on convergence culture) — low barriers to creation, strong sharing incentives, informal mentorship. 4. Variety / Yahoo Finance reporting on Disney's cease-and-desist to Google over alleged unauthorised training, predating the OpenAI deal. finance.yahoo.com 5. R. H. Coase, "The Nature of the Firm" (1937); O. E. Williamson, "Transaction-Cost Economics" (1979) and later work on governance structures. 6. Spotify Newsroom / PR Newswire, "Spotify and Universal Music Group Announce Landmark Licensing Agreements for Fan-Made Covers and Remixes," 21 May

2026 (consent / credit / compensation; opt-in; paid Premium add-on). newsroom.spotify.com 7. J. B. Barney, “Firm Resources and Sustained Competitive Advantage” (1991) — the VRIN resource criteria.

Page 11

8. Roberta Thomson (CEO, Sixth Wall) quoted in BusinessWire / Brands Untapped; ElevenLabs co-founder Mati Staniszewski; Screen Daily notes parallel ElevenLabs Iconic Marketplace deals with the Stan Lee estate and Michael Caine. screendaily.com
9. BriefGlance analysis noting FTC inquiries into AI companions and child/teen safety; Sixth Wall’s stated 13+/enterprise focus and decision not to build for young children. briefglance.com
10. TechCrunch and Pasquale Pillitteri / Billboard coverage: deal announced at Spotify Investor Day; shares rose ~13% on the session; framed as a Suno competitor. techcrunch.com
11. RIAA/UMG litigation language against Suno and Udio (“infringement on an almost unimaginable scale”),
12. See EDM.com / Digital Music News background. digitalmusicnews.com
13. Reporting on UMG’s October 2025 settlement with Udio and pivot to a licensed AI music product (referenced in Medianama and music-business coverage of the Spotify deal). medianama.com
14. MIDiA Research, “Suno, Alexa, and the rise of ‘create to consume’” — the conceptual frame for personalised, disposable, on-the-fly creation. midiaresearch.com
15. The Walt Disney Company / OpenAI announcement, 11 December 2025; Ropes & Gray, “The \$1 Billion Deal That Wasn’t” (terms: 200+ characters; ChatGPT Images; Disney+ hosting; no-training clause; talent voice/likeness excluded; joint steering committee; \$1bn equity + warrants; Disney as enterprise customer). ropesgray.com
16. OpenAI shutdown of Sora announced 24 March 2026; staged wind-down (consumer app/web ~26 April 2026; developer API ~24 September 2026). ARTVOICE / tech-insider reporting. tech-insider.org
17. Reporting on Sora’s unit economics (peak ~\$540k monthly revenue vs operating costs in the billions) and OpenAI’s reallocation of compute to coding/enterprise products amid competitive pressure; “no formal agreement signed, no money changed hands.” tech-insider.org
18. Ropes & Gray and The Hollywood Reporter / Axios: Disney reportedly given as little as ~30 minutes’ notice; governance-rights lesson for strategic investors. ropesgray.com
19. TheStreet / AOL (Hollywood Reporter source): Disney exited the deal; no money changed hands; Disney statement that it will continue to engage with AI platforms that respect IP and creator rights. thestreet.com
20. Grimes / Elf.tech (CreateSafe, Trinita API): ~50% master-royalty split under the “GrimesAI” moniker; 15,000+ tracks generated. Holly Herndon / Holly+ DAO governance model. See “How Artists Lend Their Voices to Creators & Fans.” killthedj.com (specific figures are artist/press-reported and should be cited as such.)

21. Holly Herndon & Mat Dryhurst, “Have I Been Trained?” consent-layer tooling (Spawning) for visibility into training datasets.
22. ElevenLabs & CMG Worldwide heritage-voice licensing (Judy Garland, James Dean, Burt Reynolds, Laurence Olivier), confined to the Reader app and walled off from the general voice database. CBS News / ElevenLabs. [cbsnews.com](https://www.cbsnews.com)
23. Meta celebrity chatbot programme, October 2023 (Snoop Dogg, Kendall Jenner, MrBeast et al.), reportedly up to ~\$5m per celebrity; closed, platform-controlled personas. El País / People of Color in Tech. elpais.com

Page 12

23. Amazon Alexa+ × Suno integration (Feb 2026 rollout context); Panos Panay’s birthday-song / cartoon-rap examples; major-label friction. Billboard / RouteNote. [billboard.com](https://www.billboard.com)
24. Epidemic Sound, Future of the Creator Economy reporting, and Suno community discussion: as audio generation commoditises, value migrates to visual branding, narrative and world-building. [midiaresearch.com](https://www.midiaresearch.com) Editorial note on rigour: the most consequential correction to the underlying brief concerns Disney–OpenAI. The deal was announced (11 Dec 2025) but never executed — no signed agreement, no funds transferred — and collapsed within ~90 days when Sora was shut down. Several widely-circulated summaries treat it as a completed, operative template; it was not. Figures for the creator economy, the Grimes royalty split and the Meta payouts are press- or analyst-reported and are flagged accordingly above. ▼ DREAM MACHINE — HOW AI IS RESHAPING THE CREATIVE INDUSTRIES. DEEP DIVE Nº 14 · THE PROGRAMMABLE BRAND · JUNE 2026

* * *

All embedded URLs (in document order)

- <https://www.businesswire.com/news/home/20260603297922/en/>
- <https://www.thewrap.com/industry-news/business/hasbro-ai-studio-sixth-wall-launch/>
- <https://www.goldmansachs.com/insights/articles/the-creator-economy-could-approach-half-a-trillion-dollars-by-2027>
- <https://finance.yahoo.com/sectors/technology/articles/disney-cancels-1b-deal-openai-043000427.html>
- <https://newsroom.spotify.com/2026-05-21/universal-music-group-spotify-licensing-agreements-fan-made-covers-remixes/>
- <https://www.screendaily.com/news/hasbro-launches-subsidiary-ai-studio-sixth-wall-to-exploit-ip/5217380.article>
- <https://briefglance.com/articles/hasbros-sixth-wall-taming-the-ai-wild-west-with-behavioral-licensing>
- <https://techcrunch.com/2026/05/21/spotify-and-universal-music-strike-deal-allowing-fan-made-ai-covers-and-remixes/>
- <https://www.digitalmusicnews.com/2025/02/27/amazon-suno-team-up-for-gen-ai-songs/>

- <https://www.medianama.com/2026/05/223-spotify-umg-fans-create-licensed-ai-covers-remixes/>
- <https://www.midiaresearch.com/blog/suno-alexa-and-the-rise-of-create-to-consume>
- <https://www.ropesgray.com/en/insights/viewpoints/the-deal-that-wasnt-what-disney-and-openai-teach-us-about-strategic-investments>
- <https://tech-insider.org/openai-sora-shutdown-9-6-million-downloads-disney-2026/>
- <https://tech-insider.org/openai-sora-shutdown-disney-deal-ai-video-2026/>
- <https://www.thestreet.com/technology/openai-is-shutting-down-sora-and-the-disney-deal-is-off>
- <https://killthedj.com/how-artists-lend-their-ai-voices-to-creators-fans/>
- <https://www.cbsnews.com/news/ai-voices-dead-celebrities-judy-garland-james-dean-burt-reynolds/>
- <https://english.elpais.com/technology/2023-11-13/from-paris-hilton-to-snoop-dogg-meta-to-launch-speaking-ai-chatbots-played-by-celebrities.html>
- <https://www.billboard.com/pro/amazon-partners-ai-music-company-suno-alexa/>

Citation Index

Dream Machine: The New Creative Economy

All footnoted sources, organised by chapter. Every claim of substance in the manuscript is anchored to one of these references.

* * *

Foreword — Welcome to the Dream Machine

1. Variety, “SAG-AFTRA Condemns Tilly Norwood: AI Actress Is Not an Actor,” 30 September 2025. <https://variety.com/2025/film/news/sag-aftra-tilly-norwood-ai-actress-1236534779/>. See also NBC News, “Tilly Norwood, fully AI ‘actor,’ blasted by actors union SAG-AFTRA for ‘devaluing human artistry.’” <https://www.nbcnews.com/pop-culture/pop-culture-news/tilly-norwood-fully-ai-actor-blasted-actors-union-sag-aftra-devaluing-rcna234685>. Discussed in *Dream Machine* Issue 1 (6 October 2025).

2. *The Hollywood Reporter*, “U.K. Union Equity Condemns Tilly Norwood: ‘AI Tool, Not a Performer’.” <https://www.hollywoodreporter.com/movies/movie-news/tilly-norwood-ai-actress-uk-union-equity-sag-aftra-debate-1236391739/>. See also Variety, “Tilly Norwood Slammed by Equity as AI Tool, Concerned About Origin.” <https://variety.com/2025/film/global/tilly-norwood-slammed-equity-ai-tool-concerned-origin-1236537042/>. *Dream Machine* Issue 1.

3. CNN, “Tilly Norwood: Hollywood is fuming over a new ‘AI actress,’” 30 September 2025. <https://www.cnn.com/2025/09/30/tech/hollywood-ai-actor-backlash>.

4. OpenAI, “Sora 2 is here,” announcement page, 30 September 2025. <https://openai.com/index/sora-2/>. The model launched alongside an invite-only iOS app of the same name in the U.S. and Canada. *Dream Machine* Issue 1 carried the launch alongside contemporaneous coverage from NBC News and *The Guardian* on the model’s first copyright and safety incidents.

5. *Dream Machine* | *Creative AI*, LinkedIn newsletter, archive of Issues 1–32, October 2025 – June 2026. <https://www.linkedin.com/newsletters/dream-machine-creative-ai-7379776527871381505/>.

6. DreamLab AI Collective, team page. <https://dreamlab-ai.com/team>. Referenced from *Dream Machine* Issue 16 onward.

7. Charles Cecil (Revolution Software, *Broken Sword*) quoted in *gamesindustry.biz*, “AI was an expensive mistake’: Charles Cecil on innovation, insolvency, and Broken Sword.” <https://www.gamesindustry.biz/ai-was-an-expensive-mistake-charles-cecil-on-innovation-insolvency-and-broken-sword>. *Dream Machine* Issue 3.

8. Adobe, “Inaugural Adobe Creators’ Toolkit Report: 86 Percent of Global Creators Use Creative Generative AI.” <https://news.adobe.com/news/2025/10/adobe-max-2025-creators-survey>. Survey of 16,000 creators across the U.S., U.K., France, Germany, South Korea, Japan, India and Australia, released at Adobe MAX 2025. *Dream Machine* Issue 6.

- 9.** UK Department for Science, Innovation and Technology (DSIT), *Statement of Progress on Copyright and AI*, December 2025. <https://www.gov.uk/government/publications/copyright-and-artificial-intelligence-progress-report/copyright-and-artificial-intelligence-statement-of-progress-under-section-137-data-use-and-access-act>. See also IPWatchdog, “Respondents to UK AI Consultation Overwhelmingly Want AI Companies to License Copyrighted Works in All Cases.” <https://ipwatchdog.com/2025/12/16/respondents-uk-ai-consultation-overwhelmingly-want-ai-companies-license-copyrighted-works-all-cases/>. *Dream Machine* Issue 12.
- 10.** *Dream Machine* Issue 5, “Adobe’s Latest AI Announcements — Is every tool going AI?,” 31 October 2025. <https://www.linkedin.com/pulse/dream-machine-creative-ai-news-insight-oct-25-issue-5-woodbridge-f7jnc/>.
- 11.** Adobe, Adobe MAX 2025 keynote messaging, October 2025. Coverage: Creative Boom, “Adobe is putting AI in everything everywhere all at once.” <https://www.creativeboom.com/news/adobe-is-putting-ai-in-everything-everywhere-all-at-once/>. *Dream Machine* Issue 5.
- 12.** World Labs, *Marble* — first commercial spatial-AI world model, public launch November 2025. <https://marble.worldlabs.ai/>. Technical context: TechCrunch, “Fei-Fei Li’s World Labs speeds up the world model race with Marble, its first commercial product.” <https://techcrunch.com/2025/11/12/fei-fei-lis-world-labs-speeds-up-the-world-model-race-with-marble-its-first-commercial-product/>. DreamLab participated in the closed beta during October–November 2025. *Dream Machine* Issue 7.
- 13.** 11,514 responses across the Citizen Space portal and email, of which 10,112 came through Citizen Space; 88% of those supported licensing as a default rule, against 3% who supported the government’s preferred opt-out model. UK DSIT, *Statement of Progress*, December 2025; analysis in *Dream Machine* Issue 12 (18 December 2025). Final report and economic impact assessment to be laid before Parliament by 18 March 2026.
- 14.** *Digital Music News*, “Nearly 800 Creatives, Including Jason Aldean and One Republic, Sign Responsible AI Declaration — ‘Stealing Our Work Is Not Innovation’.” <https://www.digitalmusicnews.com/2026/01/22/stealing-isnt-innovation/>. *Dream Machine* Issue 16.

Chapter 1 — The Day Sora Landed

- 1.** For a contemporaneous overview of the AI video model release cadence through 2024 and 2025, see *Dream Machine* Issues 1–8 (October–November 2025), which logged near-weekly releases from Runway, Luma, Pika, Kling, Veo, Wan, Higgsfield, Hunyuan and a long tail of smaller labs.
- 2.** *The Hollywood Reporter*, “AI Performer Tilly Norwood Sparks Hollywood Backlash.” <https://www.hollywoodreporter.com/movies/movie-news/tilly-norwood-ai-actress-uk-union-equity-sag-aftra-debate-1236391739/>. *Dream Machine* Issue 1.

3. SAG-AFTRA statement, 30 September 2025, reported in Variety, “SAG-AFTRA Condemns Tilly Norwood: AI Actress Is Not an Actor.” <https://variety.com/2025/film/news/sag-aftra-tilly-norwood-ai-actress-1236534779/>.
4. OpenAI, “Sora 2 is here,” 30 September 2025. <https://openai.com/index/sora-2/>. *Dream Machine* Issue 1.
5. Particle6 background and Van der Velden interview: *The Hollywood Reporter*, “Meet the Creator of the AI Actress Hollywood Loves to Hate: ‘You’re Gonna See a Lot of Tilly Norwood Next Year’.” <https://www.hollywoodreporter.com/movies/movie-features/tilly-norwood-creator-particle6-elina-van-der-velden-talks-1236428824/>. *Dream Machine* Issue 8.
6. *Deadline*, “Tilly Norwood Creator Eline Van Der Velden Talks Backlash, Reveals Another 40 AI Actors Are In The Pipeline.” <https://deadline.com/2025/11/tilly-norwood-creator-interview-backlash-more-ai-actors-coming-1236601334/>.
7. Northeastern Global News, “Why AI ‘Actress’ Tilly Norwood Has Hollywood Angry.” <https://news.northeastern.edu/2025/10/02/ai-actress-tilly-norwood-hollywood-backlash/>.
8. SAG-AFTRA, official statement reproduced in Variety, *op. cit.*; also NBC News, “Tilly Norwood, fully AI ‘actor,’ blasted by actors union SAG-AFTRA for ‘devaluing human artistry.’” <https://www.nbcnews.com/pop-culture/pop-culture-news/tilly-norwood-fully-ai-actor-blasted-actors-union-sag-aftra-devaluing-rcna234685>.
9. Equity (U.K.), statement of 2 October 2025: *Variety*, “Tilly Norwood Slammed by Equity as AI Tool, Concerned About Origin.” <https://variety.com/2025/film/global/tilly-norwood-slammed-equity-ai-tool-concerned-origin-1236537042/>.
10. CNN, “Tilly Norwood: Hollywood is fuming over a new ‘AI actress.’” <https://www.cnn.com/2025/09/30/tech/hollywood-ai-actor-backlash>.
11. OpenAI, “Sora 2 is here,” <https://openai.com/index/sora-2/>. Technical capabilities summary including physics modelling, multi-shot world-state persistence and synchronised audio.
12. *Dream Machine* Issue 1, “Editor’s Pick”; further launch context in NBC News, “OpenAI’s Sora 2: a major leap in AI video and audio.” <https://www.nbcnews.com/tech/tech-news/openai-sora-2-app-video-chatgpt-creation-rcna234973>.
13. LinkedIn News aggregation: “Sora Tops 1 Million Downloads in 5 Days.” <https://www.linkedin.com/news/story/sora-tops-1m-downloads-in-5-days-6684988/>. *Dream Machine* Issue 3. A historical note for the reader: the Sora consumer app and the broader Sora platform did not survive 2026. OpenAI announced a staged wind-down of Sora on **24 March 2026** — consumer app and website ending late April, developer API closing late September — with peak revenue reportedly at around \$540,000 per month against operating costs in the billions, and the compute reallocated toward higher-margin coding and enterprise products. The Disney–OpenAI licensing announcement of 11 December 2025 (referenced in Chapter 7 and Appendix C) was never executed; no formal agreement was signed and no money changed hands. The contractual framework underneath that deal — IP-not-training, talent-likeness excluded, jointly-governed steering committee, equity-and-

warrants structure — survived the platform’s collapse and is now the reference architecture other studios negotiate against. The full deal-architecture analysis is in [Appendix L: The Programmable Brand](#). The argument the rest of this chapter is making about the Sora 2 launch week — the cultural rupture, the audience arrival in the AI debate, the parasocial-character / working-actor market collision — survives the platform’s death intact; the *moment* matters whether or not the *vendor* persists.

14. *The Guardian*, “OpenAI Sora 2 violence racism.” <https://www.theguardian.com/us-news/2025/oct/04/openai-sora-violence-racism>. *Dream Machine* Issue 1.

15. NBC News, *op. cit.*; *The Guardian*, *op. cit.*

15a. Quoted in *The Guardian*, “OpenAI launch of video app Sora plagued by violent and racist images: ‘The guardrails are not real’.” <https://www.theguardian.com/us-news/2025/oct/04/openai-sora-violence-racism>. *Dream Machine* Issue 1.

16. *Digital Music News*, “OpenAI’s Sora 2 includes likeness protections for celebrities who don’t opt in, but that doesn’t apply to ‘historical figures’ and dead celebrities.” <https://www.digitalmusicnews.com/2025/10/08/openais-likeness-protections-dont-apply-to-dead-celebrities/>. *Dream Machine* Issue 2.

17. Google DeepMind, Veo 3.1 launch, mid-October 2025. *Dream Machine* Issue 3, “Editor’s Pick: Veo 3.1 and the Rise of AI Filmmaking.” Coverage: <https://www.cometapi.com/veo-3-1-is-coming-and-whats-rumor/>.

18. WUFT, “Kiss reality goodbye: AI-generated social media has arrived,” 3 October 2025. <https://www.wuft.org/2025-10-03/kiss-reality-goodbye-ai-generated-social-media-has-arrived>. *Dream Machine* Issue 1.

19. *No Film School*, “James Cameron Says AI Is ‘Never Going to Take the Place’ of Humans.” <https://nofilmschool.com/james-cameron-ai#>. *Dream Machine* Issue 1.

20. *The Guardian*, “James Cameron says AI actors are ‘horrifying to me’,” 1 December 2025. <https://www.theguardian.com/film/2025/dec/01/james-cameron-says-ai-actors-are-horrifying-to-me>. Original quote from CBS *Sunday Morning*. *Dream Machine* Issue 10.

20a. *Variety*, “James Cameron Says It’s ‘Horrifying’ that AI Can ‘Make Up an Actor’.” <https://variety.com/2025/film/news/james-cameron-horrifying-ai-replace-actors-1236595864/>.

21. *Stability AI*, board composition, 2024–2026. Reported across multiple outlets including *Deadline*, “James Cameron Calls AI Replacing Actors ‘Horrifying’; Art ‘Sacred’.” <https://deadline.com/2025/11/james-cameron-gen-ai-horrifying-human-art-sacred-avatar-1236631387/>.

22. *Deezer*, “AI-generated tracks now represent 44% of all new uploaded music,” April 2026 newsroom release. <https://newsroom-deezer.com/2026/04/ai-generated-tracks-represent-44-of-new-uploaded-music/>. Companion analysis: *Music Business Worldwide*, “75,000 AI-generated tracks now flood Deezer daily, representing 44% of all new music uploaded to the platform.” <https://www.musicbusinessworldwide.com/75000-ai-generated-tracks-now-flood-deezer-daily-representing-44-of-all-new-music-uploaded-to-the-platform-says->

streamer/. Daily AI uploads to Deezer rose from approximately 50,000 per day in November 2025 (*Dream Machine* Issue 7, citing Deezer / *Musically*) to 75,000 per day by April 2026, with consumer streams of fully-AI tracks holding between 1% and 3% of total platform plays – and up to 85% of those streams identified as fraudulent in 2025. *Dream Machine* Issues 7, 26, 27, 28.

Chapter 2 – A History of Resistance

1. John Philip Sousa, “*The Menace of Mechanical Music*,” *Appleton’s Magazine*, Vol. 8, September 1906, pp. 278–284. Full text via ExplorePAHistory: <https://explorepahistory.com/odocument.php?docId=1-4-1A1>. Academic context: Patrick Warfield, “*John Philip Sousa and ‘The Menace of Mechanical Music,’*” *Journal of the Society for American Music*, Cambridge University Press: <https://www.cambridge.org/core/journals/journal-of-the-society-for-american-music/article/abs/john-philip-sousa-and-the-menace-of-mechanical-music/A9E621587BE7580ABD73AEF64D4B2DC8>. The 1906 essay was, in part, lobbying for what would become the 1909 Copyright Act.

2. Sousa, *op. cit.* The Library of Congress’s “*Sousa and the Talking Machine*” essay is a useful institutional summary: <https://blogs.loc.gov/now-see-hear/2020/05/sousa-and-the-talking-machine/>.

3. William Henry Cardinal O’Connell, Archbishop of Boston, sermon to the Holy Name Society, Boston, 10 January 1932. Reported widely in the contemporaneous press, including the *Daily Courier* (Connellsville, PA), 12 January 1932 (<https://www.newspapers.com/newspage/38168082/>). Cultural context: KUOW/NPR, “*‘Imbecile Slush’: Surprising Early Reactions to Crooning*,” <https://www.kuow.org/stories/imbecile-slush-surprising-early-reactions-crooning>. JSTOR Daily, “*The Gender Politics of the First Boy Bands*,” <https://daily.jstor.org/the-gender-politics-of-the-first-boy-bands/>.

4. *Grand Upright Music, Ltd. v. Warner Bros. Records Inc.*, 780 F. Supp. 182 (S.D.N.Y. 1991). Full text: <https://law.justia.com/cases/federal/district-courts/FSupp/780/182/1445286/>. The “Thou shalt not steal” opening is the most-quoted line from a US copyright opinion of the late twentieth century.

5. Tippet’s account of the Jurassic Park digital test is documented across multiple ASC and contemporaneous press accounts. American Society of Cinematographers, “*Jurassic Park: Effects Team Brings Dinosaurs Back from Extinction*,” <https://theasc.com/articles/jurassic-park-effects-team-brings-dinosaurs-back>. Wikipedia, “*Phil Tippet*,” https://en.wikipedia.org/wiki/Phil_Tippett. The dialogue paraphrase Spielberg incorporated into the film is Goldblum/Malcolm’s response to Grant’s “I think we’re out of a job”: “Don’t you mean *extinct*?”

6. Charles Baudelaire, “*Le Public Moderne et la Photographie*,” *Revue Française*, 1859 (part of the *Salon de 1859* essays). English translation widely available; the original French in PDF form: <https://gallowayexploringart.wordpress.com/wp->

[content/uploads/2014/08/ baudelaire_the-modern-public-photography.pdf](#). Smithsonian Archives institutional overview: “*Photography Murdered Painting, Right?*”, <https://siarchives.si.edu/blog/photography-murdered-painting-right>.

7. The Delaroche apocrypha is documented in Quote Investigator: <https://quoteinvestigator.com/2022/10/16/photo-mortal/>. The earliest sourced version is in an 1873 survey, 34 years after Delaroche reportedly said it. Delaroche’s own contemporary writing on the daguerreotype, in Gernsheim’s standard 1959 monograph, characterised the new technology as “*an immense service to the arts.*”

8. The 1942–44 Petrillo strike: Wikipedia, “*1942–44 musicians’ strike*,” https://en.wikipedia.org/wiki/1942%E2%80%931944_musicians'_strike; Mainspring Press, “*The Man Who Crippled the American Recording Industry: James Caesar Petrillo and the American Federation of Musicians Recording Bans*,” <https://mainspringpress.org/2024/11/23/the-man-who-crippled-the-recording-industry-james-caesar-petrillo-and-the-american-federation-of-musicians-recording-bans/>; DownBeat, “*The Petrillo Ban of 1942–44: Past & Future at War*,” <https://downbeat.com/news/detail/the-petrillo-ban-of-194244-past-future-at-war>; Local 802 AFM, “*The Silence Was Deafening*,” <https://www.local802afm.org/allegro/articles/the-silence-was-deafening/>. The Music Performance Trust Fund’s institutional history: <https://musicpf.org/establishment-of-mptf-led-to-the-formation-of-afms-pension-and-residual-funds/>.

9. Musicians’ Union History, “*The Strike That Made History — Massacre of the Musicians 1980*,” <https://www.muhistory.com/the-massacre-of-the-musicians-1980/>. Academic context on the broader MU–BBC dispute landscape: “*Negotiating Needletime*” (Tandfonline), <https://www.tandfonline.com/doi/full/10.1080/03071022.2016.1215098>.

10. MusicRadar, “*The Day the Loony Musicians Union Tried to Kill the Synthesizer (Which Also Happened to be Bob Moog’s Birthday)*,” <https://www.musicradar.com/news/the-union-passed-a-motion-to-ban-the-use-of-synths-drum-machines-and-any-electronic-devices-the-day-the-loony-musicians-union-tried-to-kill-the-synthesizer-which-also-happened-to-be-bob-moogs-birthday>. Far Out Magazine, “*Why did the Musicians Union outlaw synthesisers in 1982?*,” <https://faroutmagazine.co.uk/musicians-union-outlaw-synthesisers/>.

11. *Bridgeport Music, Inc. v. Dimension Films*, 410 F.3d 792 (6th Cir. 2005). Full text: <https://law.justia.com/cases/federal/appellate-courts/F3/410/792/574458/>. The “Get a licence or do not sample” rule is the most-cited line in the opinion.

12. *TIME*, “*50 Worst Inventions*,” 2010, Auto-Tune at #15: https://content.time.com/time/specials/packages/article/0,28804,1991915_1991909_1991909c Wikipedia, “*Auto-Tune*,” <https://en.wikipedia.org/wiki/Auto-Tune>. NPR, “*25 Years of Believe*,” <https://www.npr.org/2023/10/19/1207028349/25-years-ago-cher-released-a-song-that-would-change-the-sound-of-pop-music>. Wikipedia, “*D.O.A. (Death of Auto-Tune)*,” [https://en.wikipedia.org/wiki/D.O.A._\(Death_of_Auto-Tune\)](https://en.wikipedia.org/wiki/D.O.A._(Death_of_Auto-Tune)).

13. Walter Murch, *In the Blink of an Eye: A Perspective on Film Editing*, Silman-James Press, 1995 (2nd edition 2001). PDF:

[https://www.craftfilmschool.com/userfiles/files/Walter%20Murch%20-%20In%20the%20Blink%20of%20an%20Eye%20Revised%202nd%20Edition%20\(2001,%20James%20Pr\).pdf](https://www.craftfilmschool.com/userfiles/files/Walter%20Murch%20-%20In%20the%20Blink%20of%20an%20Eye%20Revised%202nd%20Edition%20(2001,%20James%20Pr).pdf). Charles Koppelman, *Behind the Seen: How Walter Murch Edited Cold Mountain Using Apple's Final Cut Pro and What This Means for Cinema*, Peachpit Press, 2004: <https://www.peachpit.com/store/behind-the-seen-how-walter-murch-edited-cold-mountain-9780735714267>.

14. Sasson's account documented at the National Inventors Hall of Fame: <https://www.invent.org/blog/inventors/Legacy-Steve-Sasson>. Snopes verification of the "Kodak suppressed the digital camera" claim: <https://www.snopes.com/fact-check/kodak-digital-camera-invention/>. Knowledge@Wharton on the Kodak collapse: <https://knowledge.wharton.upenn.edu/podcast/knowledge-at-wharton-podcast/whats-wrong-with-this-picture-kodaks-30-year-slide-into-bankruptcy/>. Bankruptcy filing: 19 January 2012, S.D.N.Y., \$5.1bn assets / \$6.8bn liabilities.

15. Wikipedia, "Brian Walski," https://en.wikipedia.org/wiki/Brian_Walski. *Washington Post* contemporaneous coverage: <https://www.washingtonpost.com/archive/lifestyle/2003/04/03/alterd-picture-costs-lattimes-photographer-his-job/c5e7c9e0-a836-429a-bb4e-d502f1768a96/>. World Press Photo's institutional response in TIME: <https://time.com/3706626/world-press-photo-processing-manipulation-disqualified/>.

16. Wikipedia, "Viacom International, Inc. v. YouTube, Inc.," https://en.wikipedia.org/wiki/Viacom_International_Inc._v._YouTube_Inc.. Electronic Frontier Foundation case file: <https://www.eff.org/cases/viacom-v-youtube>. Variety on the March 2014 settlement: <https://variety.com/2014/biz/news/google-and-viacom-settle-copyright-infringement-lawsuit-over-youtube-1201137538/>.

17. PetaPixel, "The Rise and Crash of the Camera Industry in One Chart," <https://petapixel.com/2024/08/22/the-rise-and-crash-of-the-camera-industry-in-one-chart/>. Statista, "Smartphones Wipe Out Decades of Camera Industry Growth," <https://www.statista.com/chart/15524/worldwide-camera-shipments/>. CIPA shipment data series, multiple years.

18. CNN Business, "Meet the translation professionals losing their jobs to AI," January 2026, <https://www.cnn.com/2026/01/23/tech/translation-language-jobs-ai-automation-intl>. Carl Benedikt Frey (Oxford Martin School), 2025 study on translator employment across 696 US labour markets. American Translators Association industry position: <https://www.atanet.org/client-assistance/blog-machine-translation-vs-human-translation/>. Wikipedia, "Google Neural Machine Translation," https://en.wikipedia.org/wiki/Google_Neural_Machine_Translation.

Chapter 3 – The Human-AI Agency Continuum

1. *Dream Machine* Issue 2, "Editor's Pick," 10 October 2025. <https://www.linkedin.com/pulse/dream-machine-creative-ai-news-insight-oct-25-2-pete-woodbridge-mnrjc/>.

2. OpenAI, “Introducing AgentKit,” 6 October 2025. <https://openai.com/index/introducing-agentkit/>.
3. TechCrunch, “OpenAI launches AgentKit to help developers build and ship AI agents,” 6 October 2025. <https://techcrunch.com/2025/10/06/openai-launches-agentkit-to-help-developers-build-and-ship-ai-agents/>. Also coverage at *InfoQ*, “OpenAI Dev Day 2025 Introduces GPT-5 Pro API, Agent Kit, and More.” <https://www.infoq.com/news/2025/10/openai-dev-day/>.
4. *Dream Machine* Issue 2: “Agentic AI — the class of AI systems that can plan, act, and pursue goals with autonomy — promises a new era of collaboration in creative industries... Its another step along the Human-AI Agency Continuum.” See also *TVB Europe*, “Is Agentic AI About to Change the Media and Entertainment Industry?” <https://www.tvbeurope.com/artificial-intelligence/opinion-is-agentic-ai-about-to-change-the-media-and-entertainment-industry>.
5. Google DeepMind, Veo 3.1 release, October 2025. *Dream Machine* Issue 3.
6. *MusicTech*, “iZotope Ozone 12’s AI assistant is cool, but the Stem EQ is the real star.” <https://musictech.com/reviews/plugin-ins/izotope-ozone-12-review/>. *Dream Machine* Issue 3.
7. Adobe, “Inaugural Adobe Creators’ Toolkit Report,” October 2025. <https://news.adobe.com/news/2025/10/adobe-max-2025-creators-survey>. Survey of 16,000 creators across eight countries, released at Adobe MAX 2025. *Dream Machine* Issue 6.
8. Adobe, *op. cit.* The same survey: 86% of creators use creative generative AI; 76% say it has helped grow their business or brand; 81% say AI lets them make content they otherwise couldn’t have made; 69% worry about their work being used to train AI without consent; 70% are optimistic about agentic AI; 85% would use AI that learns their creative style.
9. Mureka, “Music Agent Studio” launch, mid-October 2025. *Dream Machine* Issue 4. https://www.linkedin.com/posts/sherrihendrickson_mureka-unveils-music-agent-studio-and-enhanced-share-7384999251526864896-cNYg/.
10. *Finsmes*, “AdsGency Raises \$12M in Seed Funding,” October 2025. <https://www.finsmes.com/2025/10/adsgency-raises-12m-in-seed-funding.html>. *Dream Machine* Issue 4.
11. *Musically*, “Meet Lenny, an AI agent to help organisers of live music events.” <https://musically.com/2025/10/20/meet-lenny-an-ai-agent-to-help-organisers-of-live-music-events/>. *Dream Machine* Issue 4.
12. *GamesRadar*, “Even under USD20 million in debt, EA reportedly pushes 15,000 employees to use AI as a ‘thought partner’ for everything from character art to playtesting.” <https://www.gamesradar.com/games/even-under-usd20-million-in-debt-ea-reportedly-pushes-15-000-employees-to-use-ai-as-a-thought-partner-for-everything-from-character-art-to-playtesting/>. *Dream Machine* Issue 6.
13. PYMNTS, “Adobe Lets Users Design and Edit Using ChatGPT.” <https://www.pymnts.com/artificial-intelligence-2/2025/adobe-lets-users-design-and-edit-using-chatgpt/>. Adobe blog: “Edit images, designs, and PDFs right inside ChatGPT — thanks

to Adobe Express, Photoshop, and Acrobat.” <https://blog.adobe.com/en/publish/2025/12/10/edit-photoshop-chatgpt>. *Dream Machine* Issue 12.

14. TechCrunch, “Anthropic launches interactive Claude apps, including Slack and other workplace tools,” 26 January 2026. <https://techcrunch.com/2026/01/26/anthropic-launches-interactive-claude-apps-including-slack-and-other-workplace-tools/>. *Heygen Video Agent*: https://www.linkedin.com/posts/heygen_introducing-the-new-video-agent-activity-7421597801240801282-d1CF. *Dream Machine* Issue 16.

15. *Dream Machine* Issue 21, “Editor’s Pick: Adobe and NVIDIA Just Raised the Stakes for Creative AI,” 19 March 2026.

16. Adobe Summit 2026, “Agentic Creative Intelligence” keynote framing. *Dream Machine* Issue 26.

17. *Dream Machine* Issue 29, May 2026, citing Sony’s adoption of Claude Code studios with multi-agent coordination.

18. Anthropic, public statements on agent deployment patterns through Q1 2026. Cf. *Dream Machine* Issues 11, 16, 22.

19. *gamesindustry.biz*, “AI was an expensive mistake’: Charles Cecil on innovation, insolvency, and Broken Sword.” <https://www.gamesindustry.biz/ai-was-an-expensive-mistake-charles-cecil-on-innovation-insolvency-and-broken-sword>. *Dream Machine* Issue 3.

20. *Niche Gamer*, “Larian Studios backs off from gen AI, says tech won’t be used in new Divinity.” <https://nichegamer.com/larian-studios-backs-off-from-gen-ai/>. *Dream Machine* Issue 14.

21. *Decrypt*, “‘Warhammer 40,000’ Maker Games Workshop Rules Out Generative AI.” <https://decrypt.co/354482/warhammer-40000-maker-games-workshop-rules-out-generative-ai>. *Dream Machine* Issue 14.

22. *Niche Gamer*, “Manor Lords publisher Hooded Horse won’t work with devs using gen AI.” <https://nichegamer.com/manor-lords-publisher-hooded-horse-wont-work-with-devs-using-gen-ai/>. *Dream Machine* Issue 14.

23. *gamesindustry.biz*, “RuneScape maker Jagex says it will never use generative AI to make in-game content.” <https://www.gamesindustry.biz/runescape-maker-jagex-says-it-will-never-use-generative-ai-to-make-in-game-content>. *Dream Machine* Issue 16.

Chapter 4 — Dead Internet, Living Web

1. Imperva, *2025 Bad Bot Report: How AI is Supercharging the Bot Threat*. <https://www.imperva.com/blog/2025-imperva-bad-bot-report-how-ai-is-supercharging-the-bot-threat/>. *Dream Machine* Issue 4.

2. Cloudflare, “The crawl-to-click gap: Cloudflare data on AI bots, training, and referrals.” <https://blog.cloudflare.com/crawlers-click-ai-bots-training/>. *Dream Machine* Issue 4. Later 2025 updates show training crawlers declining from ~90% to ~74% of AI bot activity as

scraper bots rose to 24% and a new “agentic” category emerged at 1.7%; see Cloudflare, “A deeper look at AI crawlers: breaking down traffic by purpose and industry.” <https://blog.cloudflare.com/ai-crawler-traffic-by-purpose-and-industry/>.

3. Grand View Research, “Generative AI Content Creation Market Report.” <https://www.grandviewresearch.com/industry-analysis/generative-ai-content-creation-market-report>. *Dream Machine* Issue 4 also cites Gartner and Europol forecasts of 90–99% AI-generated or AI-assisted online content by 2030.

4. *Dream Machine* Issue 4, “Editor’s Pick: Is the Internet Dead Yet?” 23 October 2025. <https://www.linkedin.com/pulse/dream-machine-creative-ai-news-insight-oct-25-issue-4-woodbridge-hzttc/>.

5. Wikipedia, *Dead Internet Theory*. https://en.wikipedia.org/wiki/Dead_Internet_theory. *Dream Machine* Issue 4.

6. Graphite, 2025 analysis of new web content by author type (human vs. AI vs. AI-assisted). Cited in *Dream Machine* Issue 4.

7. For “model collapse” as a term of art, see Ilia Shumailov et al., “The Curse of Recursion: Training on Generated Data Makes Models Forget” (2024), and subsequent literature.

8. Futurism, “Researchers built a social network with only AI agents — within hours it had collapsed into warring tribes.” <https://futurism.com/social-network-ai-intervention-echo-chamber>. *Dream Machine* Issue 4.

9. *Digital Music News*, “Instagram Chief Says We Should ‘Fingerprint Real Media’ Instead of Tracking and Disclosing AI Slop.” <https://www.digitalmusicnews.com/2026/01/05/instagram-chief-ai-slop-comments/>. See also *WebProNews*, “Instagram Head Warns AI Images Erode Trust, Calls for Verification Standards.” <https://www.webpronews.com/instagram-head-warns-ai-images-erode-trust-calls-for-verification-standards/>. *Dream Machine* Issue 13.

10. Sundance Institute, “Centering the Artist: Why We’re Launching the AI Literacy Initiative.” <https://www.sundance.org/blogs/centering-the-artist-why-were-launching-the-ai-literacy-initiative/>. *Dream Machine* Issue 16.

11. *Stereogum*, “Bandcamp bans AI music.” <https://stereogum.com/2485199/bandcamp-bans-ai-music/news>. *Dream Machine* Issue 14.

12. *CNET*, “San Diego Comic-Con Draws a Line: No AI Art Allowed at 2026 Event.” <https://www.cnet.com/culture/san-diego-comic-con-bans-ai-art-for-2026-event/>. *Dream Machine* Issue 16.

13. *Deezer*, “AI-generated tracks now represent 44% of all new uploaded music,” April 2026. <https://newsroom-deezer.com/2026/04/ai-generated-tracks-represent-44-of-new-uploaded-music/>. *Music Business Worldwide*, “75,000 AI-generated tracks now flood Deezer daily.” <https://www.musicbusinessworldwide.com/75000-ai-generated-tracks-now-flood-deezer-daily-representing-44-of-all-new-music-uploaded-to-the-platform-says-streamer/>. *Dream Machine* Issues 7, 26, 27, 28.

14. *The Hollywood Reporter*, “‘Synthetic Sincerity’ by Marc Isaacs Explores if AI Characters Can Be Taught Authenticity: IDFA.” <https://www.hollywoodreporter.com/movies/movie-news/synthetic-sincerity-film-idfa-ai-authenticity-interview-1236426180/>. *Dream Machine* Issue 8.
15. *Variety*, “AI-Generated Images Threaten Future of Documentary as People ‘Will Stop Believing Anything’.” <https://variety.com/2025/film/festivals/ai-generated-images-threaten-future-of-documentary-1236583466/>. *Dream Machine* Issue 8.
16. PR Newswire, “From Apple TV Creative to AI Filmmaker: Hoyt Dwyer’s Animated Film To Compete at AI FilmFest Japan 2025.” <https://www.prnewswire.com/news-releases/from-apple-tv-creative-to-ai-filmmaker-hoyt-dwyers-animated-film-to-compete-at-ai-filmfest-japan-2025-302598064.html>. *Dream Machine* Issue 6.
17. *Variety*, “AI Creator Behind Viral ‘Deadpool,’ ‘Harry Potter’ Christmas Clip Made His Film in a Ukrainian Bomb Shelter.” <https://variety.com/2026/digital/news/ai-video-deadpool-harry-potter-andrii-daniels-1236624632/>. *Dream Machine* Issue 16.
18. *Branding in Asia*, “‘It’s the Most Terrible Time of the Year’ — McDonald’s Netherlands’ Wonderfully Chaotic, AI-Driven Christmas Film.” <https://www.brandinginasia.com/its-the-most-terrible-time-of-the-year-mcdonalds-netherlands-wonderfully-chaotic-ai-driven-christmas-film/>. Pulled following backlash: *SiliconAngle*, “Not ready: McDonald’s AI-generated ad taken down after public backlash.” <https://siliconangle.com/2025/12/10/not-ready-mcdonalds-ai-generated-ad-taken-public-backlash/>. *Dream Machine* Issue 11.
19. BBC News, “Fashion house Valentino criticised over ‘disturbing’ AI handbag ads.” <https://www.bbc.co.uk/news/articles/cwyvjyvn83go>. *Dream Machine* Issue 10.
20. *Adweek*, “Coca-Cola Uses AI to Rekindle the Magic of Its Holiday Ads.” <https://www.adweek.com/creativity/coca-cola-uses-ai-to-rekindle-the-magic-of-its-holiday-ads/>. *Dream Machine* Issue 6.
21. *AI News*, “AI causes reduction in users’ brain activity, MIT.” <https://www.artificialintelligence-news.com/news/ai-causes-reduction-in-users-brain-activity-mit/>. *Dream Machine* Issue 1.

Chapter 5 — The Slop Ceiling

1. Deezer, “AI-generated tracks now represent 44% of all new uploaded music,” April 2026. <https://newsroom-deezer.com/2026/04/ai-generated-tracks-represent-44-of-new-uploaded-music/>. *Music Business Worldwide*, “75,000 AI-generated tracks now flood Deezer daily, representing 44% of all new music uploaded to the platform.” <https://www.musicbusinessworldwide.com/75000-ai-generated-tracks-now-flood-deezer-daily-representing-44-of-all-new-music-uploaded-to-the-platform-says-streamer/>. *Dream Machine* Issues 7, 26, 27, 28.
2. Ditto Music research, October 2025 and prior. *Press Ditto Music*, “48% of artists use AI to make music — fewer than in 2023.” <https://press.dittomusic.com/48-of-artists-use-ai-to-make-music-fewer-than-in-2023>. *Dream Machine* Issue 2.

3. *Musically*, “Universal and Warner could sign landmark AI deals within weeks.” <https://musically.com/2025/10/02/report-umg-and-wmg-could-sign-landmark-ai-deals-within-weeks/>. Spotify Newsroom, “Spotify Strengthens AI Protections for Artists, Songwriters, and Producers.” <https://newsroom.spotify.com/2025-09-25/spotify-strengthens-ai-protections/>. *Dream Machine* Issue 1.
4. *Musically*, “50,000 AI music tracks are now uploaded to Deezer every day.” <https://musically.com/2025/11/12/50000-ai-music-tracks-are-now-uploaded-to-deezer-every-day/>. *Dream Machine* Issue 7.
5. Deezer, April 2026, *op. cit.*
6. *Musically*, “UMG boss slams exponential growth of AI slop on streaming services.” <https://musically.com/2026/01/09/umg-boss-slams-exponential-growth-of-ai-slop-on-streaming-services/>. *Dream Machine* Issue 14.
7. *Musically*, “Report: 56.9% of new independent songs in China are AI-generated.” <https://musically.com/2026/01/05/report-56-9-of-new-independent-songs-in-china-are-ai-generated/>. *Dream Machine* Issue 13.
8. *The Wrap*, “An AI Podcasting Machine Is Churning Out 3,000 Episodes a Week — and People Are Listening.” <https://www.thewrap.com/ai-podcasts-hosts-inception-point-ai/>. *Dream Machine* Issue 8.
9. *Dream Machine* Issue 28, May 2026, citing aggregator-platform data on “podslop” classification.
10. *The Hollywood Reporter*, “Merriam-Webster Names ‘Slop’ Word of the Year Amid AI Boom.” <https://www.hollywoodreporter.com/news/general-news/slop-word-year-2025-merriam-webster-1236450780/>. *Dream Machine* Issue 12.
11. *Digital Music News*, “YouTube CEO Puts ‘Managing AI Slop’ on the Priority List for 2026.” <https://www.digitalmusicnews.com/2026/01/22/youtube-ceo-ai-slop-2026-comments/>. *Dream Machine* Issue 16.
12. *The Guardian*, “YouTube AI channels spreading fake, anti-Labour videos viewed 1.2bn times in 2025.” <https://www.theguardian.com/technology/2025/dec/13/fake-anti-labour-video-billion-views-youtube-2025>. *Dream Machine* Issue 12.
13. Deezer/Ipsos survey, November 2025. <https://newsroom-deezer.com/2025/11/deezer-ipsos-survey-ai-music/>. *Dream Machine* Issue 7.
14. *Bain & Company*, “In an AI Age, People Still Want the Radio Star.” <https://www.bain.com/insights/in-an-ai-age-people-still-want-the-radio-star/>. *Dream Machine* Issue 16.
15. Deezer, April 2026, *op. cit.* “Up to 85% of the streams generated by fully AI-generated tracks were in fact fraudulent in 2025.”
- 15a. *Bloomberg*, “AI Changed Chess. Grandmasters Now Win With Unpredictable Moves,” 27 March 2026. <https://www.bloomberg.com/news/articles/2026-03-27/ai-changed-chess-grandmasters-now-win-with-unpredictable-moves>. *Dream Machine* Issue 23.

16. *Billboard*, “AI Artist Xania Monet Climbs the Charts — And Signs a Multimillion-Dollar Record Deal.” <https://www.billboard.com/pro/ai-music-artist-xania-monet-multimillion-dollar-record-deal/>.
17. *Billboard*, *op. cit.*; CNN, “Xania Monet is the first AI-powered artist to debut on a Billboard airplay chart.” <https://www.cnn.com/2025/11/01/entertainment/xania-monet-billboard-ai>.
18. *Billboard*, *op. cit.*
19. *Bangkok Post*, “AI singer Xania Monet signs \$3m deal with record label.” <https://www.bangkokpost.com/life/tech/3142355/ai-singer-xania-monet-signs-3m-deal-with-hollywood-media>. *Dream Machine* Issue 7.
20. Multiple outlets; quoted in *Billboard* feature *op. cit.*
- 20a. Telisha Jones quoted in *Billboard*, *op. cit.*
21. NPR, “Breaking Rust is a hot new country act on the Billboard charts. It’s powered by AI.” <https://www.npr.org/2025/11/10/nx-s1-5604320/breaking-rust-is-a-hot-new-country-act-on-the-billboard-charts-its-powered-by-ai>. *Dream Machine* Issue 7.
22. *Washington Post*, “‘Walk My Walk,’ Breaking Rust: AI country hit triggers Nashville angst.” <https://www.washingtonpost.com/style/2025/12/28/breaking-rust-ai-country/>.
23. *MusicRadar*, “The No. 1 country song in the US right now is AI-generated.” <https://www.musicradar.com/music-tech/the-no-1-country-song-in-the-us-right-now-is-ai-generated>. *Dream Machine* Issue 7.
24. BBC News, “The mysterious singer, Sienna Rose, with millions of streams is hitting the viral charts — but who (or what) is she?” <https://www.bbc.co.uk/news/articles/cq6v83gg66eo>. *Dream Machine* Issue 15.
25. *Billboard*, “How a MAGA Rapper Used AI to Create A Gospel Song That Climbed the Charts.” <https://www.billboard.com/pro/maga-rapper-ai-gospel-song-climbed-charts/>. *Dream Machine* Issue 9.
26. *Musically*, “AI band Bleeding Verse’s creator signs deal with Hallwood Media.” <https://musically.com/2025/10/07/ai-band-bleeding-verses-creator-signs-deal-with-hallwood-media/>. *Dream Machine* Issue 2.
27. *Musically*, “Indian AI band Trilok performs live, government denies association.” <https://musically.com/2025/12/17/indian-ai-band-trilok-performs-live-government-denies-association/>. *Dream Machine* Issue 12.
28. *The Guardian*, “Paul McCartney joins music industry protest against AI with silent track.” <https://www.theguardian.com/music/2025/nov/17/the-sound-of-silence-why-theres-barely-anything-there-in-paul-mccartney-new-release>. *Dream Machine* Issue 8.
29. *The Guardian*, “Musicians must embrace ‘unstoppable force’ of AI, Eurythmics’ Dave Stewart urges.” <https://www.theguardian.com/music/2025/dec/05/musicians-must-embrace-unstoppable-force-of-ai-eurythmics-dave-stewart-urges>. *Dream Machine* Issue 11.

30. *Digital Music News*, “Nearly 800 Creatives, Including Jason Aldean and One Republic, Sign Responsible AI Declaration — ‘Stealing Our Work Is Not Innovation’.” <https://www.digitalmusicnews.com/2026/01/22/stealing-isnt-innovation/>. *Dream Machine* Issue 16.
31. Stability AI, “Universal Music Group and Stability AI Announce Strategic Alliance.” <https://stability.ai/news/universal-music-group-and-stability-ai-announce-strategic-alliance>. *Dream Machine* Issue 5.
32. Stability AI, “Warner Music Group and Stability AI Join Forces To Build The Next Generation Of Responsible AI Tools For Music Creation.” <https://stability.ai/news/warner-music-group-and-stability-ai-join-forces-to-build-next-gen-tools>. *Dream Machine* Issue 8.
33. Universal Music, “Universal Music Group and Splice to Collaborate on the Next Generation of AI-Powered Music Creation Tools for Artists.” <https://www.universalmusic.com/universal-music-group-and-splice-to-collaborate-on-the-next-generation-of-ai-powered-music-creation-tools-for-artists/>. *Dream Machine* Issue 12.
34. LinkedIn / Lexology, “Munich Regional Court rules for GEMA against OpenAI.” Coverage: https://www.linkedin.com/posts/dr-barry-scannell-bbb5aa207_in-a-major-ruling-for-european-copyright-share-7393957246386323457-8bbx. *Dream Machine* Issue 7.
35. *EDM.com*, “‘Biggest Theft in Music History’: Rights Group Sues Suno as AI Music Showdown Escalates.” <https://edm.com/gear-tech/rights-group-sues-suno-copyright-infringement/>. *Dream Machine* Issue 7.
36. *Music Business Worldwide*, “Wixen files \$50m copyright suit against Meta.” <https://www.musicbusinessworldwide.com/wixen-files-50m-copyright-suit-against-meta-claims-tech-giant-wants-to-replace-songwriters-with-ai/>. *Dream Machine* Issue 16.
37. *Dream Machine* Issue 17 reportage on UMG’s \$3B suit against Anthropic.
38. *Stereogum*, “Bandcamp bans AI music.” <https://stereogum.com/2485199/bandcamp-bans-ai-music/news>. *Dream Machine* Issue 14.
39. *Dream Machine* Issue 18 reportage of Deezer licensing its detection tool.
40. *TechRadar*, “AI music is flooding Spotify, and subscribers are furious.” <https://www.techradar.com/audio/spotify/ai-music-is-flooding-spotify-and-subscribers-are-furious-heres-why-music-fans-no-longer-trust-discover-weekly>. *Dream Machine* Issue 14.
41. *CNET*, “San Diego Comic-Con Draws a Line: No AI Art Allowed at 2026 Event.” <https://www.cnet.com/culture/san-diego-comic-con-bans-ai-art-for-2026-event/>. *Dream Machine* Issue 16.
42. *The Independent*, “AI-generated song banned from Swedish charts: ‘It’s deceiving’.” <https://www.independent.co.uk/tv/news/ai-music-song-banned-sweden-spotify-b2901627.html>. *Dream Machine* Issue 15.
43. *Soultracks*, “A.I.-generated music is catchy, familiar... and boring.” <https://soultracks.com/news-ai-generated-music-is-catchy-boring/>. *Dream Machine* Issue 14.

43a. *The Independent*, “AI-generated song banned from Swedish charts: ‘It’s deceiving’.” <https://www.independent.co.uk/tv/news/ai-music-song-banned-sweden-spotify-b2901627.html>. *Dream Machine* Issue 15.

43b. *Marketing Week*, “You can’t dismiss AI ads as slop when they’re winning in testing.” Coverage discussed in *Dream Machine* Issue 22.

44. *Billboard*, “The Real Story Behind The AI Song That Knocked Tyla Off No. 1 On Billboard Afrobeats Chart.” <https://www.billboard.com/pro/ai-song-knocked-tyla-off-no-1-afrobeats/>. *Dream Machine* Issue 30.

45. *MusicTech*, “Jack Antonoff brands AI music makers as ‘godless whores’.” <https://musictech.com/news/industry/jack-antonoff-ai-music-makers-godless-whores/>. *Dream Machine* Issue 30.

45a. Apple Music platform statement, June 2026: AI-generated songs make up less than 1% of plays on the platform. *Dream Machine* Issue 31.

45b. Spotify–Universal Music Group landmark deal allowing Spotify Premium subscribers to create AI covers and remixes of UMG-catalogue tracks as a paid Premium add-on, June 2026. *Dream Machine* Issue 31.

45c. Spotify CEO Daniel Ek defending the platform’s move to AI music as “better than ‘slop’,” June 2026. *Dream Machine* Issue 31.

45d. Swedish startup Tonada — AI music for retailers (June 2026). *Dream Machine* Issue 31.

45e. *Hollywood Reporter*, “Amazon Orders Three Animated Series Using Generative AI.” <https://www.hollywoodreporter.com/tv/tv-news/amazon-orders-three-animated-series-generative-ai-1236606219/>. *Variety*, “Jorge Gutierrez Drops Out of Amazon MGM AI-Generated Series After Backlash.” <https://variety.com/2026/tv/news/jorge-gutierrez-drops-out-amazon-mgm-ai-generated-series-backlash-1236762285/>. *Futurism*, “Amazon AI-Generated Series Canceled.” <https://futurism.com/artificial-intelligence/amazon-ai-generated-series-canceled>. *Dream Machine* Issue 32.

45f. Guillermo del Toro, subtweet responding to Amazon AI Creators Fund announcement, May/June 2026. Coverage in *Dream Machine* Issue 32. Cross-reference: del Toro’s “rather die” statement in Chapter 7, §“Position Three: We are not doing this.”

Chapter 6 — The 88%

40. *Screen Daily*, “We got it wrong says UK Culture Secretary Lisa Nandy of government misstep on AI and copyright strategy,” 1 June 2026. <https://www.screendaily.com/news/we-got-it-wrong-says-uk-culture-secretary-lisa-nandy-of-government-misstep-on-ai-and-copyright-strategy/5217330.article>. *Dream Machine* Issue 32.

1. UK Department for Science, Innovation and Technology, *Statement of Progress on Copyright and AI*, 15 December 2025. <https://www.gov.uk/government/publications/copyright-and-artificial-intelligence-progress-report/copyright-and-artificial-intelligence-statement-of-progress-under-section->

137-data-use-and-access-act. *Dream Machine* Issue 12, “Editor’s Pick: 88% of Creators Said ‘No’.” 18 December 2025.

2. UK DSIT, original consultation, 17 December 2024 – 25 February 2025. Discussion in IPWatchdog, “Respondents to UK AI Consultation Overwhelmingly Want AI Companies to License Copyrighted Works in All Cases.” <https://ipwatchdog.com/2025/12/16/respondents-uk-ai-consultation-overwhelmingly-want-ai-companies-license-copyrighted-works-all-cases/>.

3. IPWatchdog, *op. cit.*; Hogan Lovells, “Copyright and AI: UK government publishes statement of progress.” <https://www.hoganlovells.com/en/publications/copyright-and-ai-uk-government-publishes-statement-of-progress>.

3a. Society of Authors submission to the UK consultation, quoted in IPWatchdog, *op. cit.*

4. UK DSIT, *Statement of Progress*, *op. cit.*; analysis at UCL Copyright Queries, “UK government publishes progress statement on AI and copyright consultation.” <https://blogs.ucl.ac.uk/copyright/2025/12/23/uk-government-publishes-progress-statement-on-ai-and-copyright-consultation/>.

5. UK DSIT, *Statement of Progress*, *op. cit.*

6. Dr Barry Scannell, LinkedIn analysis of GEMA v. OpenAI ruling, November 2025. https://www.linkedin.com/posts/dr-barry-scannell-bbb5aa207_in-a-major-ruling-for-european-copyright-share-7393957246386323457-8bbx. *Dream Machine* Issue 7.

7. EDM.com, “‘Biggest Theft in Music History’: Rights Group Sues Suno as AI Music Showdown Escalates.” <https://edm.com/gear-tech/rights-group-sues-suno-copyright-infringement/>. *Dream Machine* Issue 7.

8. *Music Business Worldwide*, “Wixen files \$50m copyright suit against Meta, claims tech giant wants to replace songwriters with AI.” <https://www.musicbusinessworldwide.com/wixen-files-50m-copyright-suit-against-meta-claims-tech-giant-wants-to-replace-songwriters-with-ai/>. *Dream Machine* Issue 16.

9. *Dream Machine* Issue 17, on UMG’s \$3B suit against Anthropic.

10. *Complete Music Update*, “Johnny Cash estate uses ELVIS Act to sue Coke over tribute act ad soundtrack.” <https://completemusicupdate.com/johnny-cash-estate-uses-elvis-act-to-sue-coke-over-tribute-act-ad-soundtrack/>. *Dream Machine* Issue 9.

11. Reuters, “European lawmakers seek EU-wide minimum age to access AI chatbots, social media.” <https://www.reuters.com/legal/litigation/european-lawmakers-seek-eu-wide-minimum-age-access-ai-chatbots-social-media-2025-11-26/>. *Dream Machine* Issue 9.

12. SAG-AFTRA contract update reporting through Q2 2026. *Dream Machine* Issues 20, 26, 29. Coverage: <https://www.theverge.com/news/842848/new-york-law-ai-advertisements-sag-aftra-labor>.

13. Equity (U.K.), “Performers prepared to take industrial action over AI in landslide 99% vote.” <https://www.equity.org.uk/news/2025/performers-prepared-to-take-industrial-action-over-ai-in-landslide-99-vote>. *Dream Machine* Issue 12.

14. Equity (U.K.), “Equity welcomes improved offer in AI protection negotiations in film and TV.” <https://www.equity.org.uk/news/2026/equity-welcomes-improved-offer-in-ai-protection-negotiations-in-film-and-tv>. *Dream Machine* Issue 15.
15. Cannes Film Festival AI Disclosure Standard launch, May 2026. *Dream Machine* Issue 29.
16. UK DSIT, *Statement of Progress*, op. cit.
17. *Dream Machine* Issue 21, 19 March 2026, on the UK government’s revised position on AI copyright.
18. *Digital Music News*, “The AI Licensing Shift — Creative Weight Attribution Emerges as Music Industry Game-Changer for Rights Holders.” <https://www.digitalmusicnews.com/2026/01/26/ai-licensing-shift-creative-weight-attribution/>. See also *Digital Music News*, “Artificial Intelligence Attribution and Licensing Startup Musical AI Scores \$4.5 Million Raise.” <https://www.digitalmusicnews.com/2026/01/13/musical-ai-funding-january-2026/>. *Dream Machine* Issues 14, 16.
19. PRS for Music, “PRS for Music AI Survey 2026.” <https://www.prsformusic.com/m-magazine/news/prs-for-music-ai-survey-2026>. *Dream Machine* Issue 16.
20. *Broadcast Now*, “Alex Mahon joins Stellar AI Creative Summit line-up” (covering the launch of the UCL/RCA Centre for Creative AI). <https://www.broadcastnow.co.uk/broadcasters/alex-mahon-joins-stellar-ai-creative-summit-line-up/5209227.article>. *Dream Machine* Issue 1.
21. *Complete Music Update*, “Artists must have creative control in AI deals or risk ending up with ‘scraps’, says US artist trade body.” <https://completemusicupdate.com/artists-must-have-creative-control-in-ai-deals-or-risk-ending-up-with-scraps-says-us-artist-trade-body/>. *Dream Machine* Issue 6.
22. *Digital Music News*, “Nearly 800 Creatives, Including Jason Aldean and One Republic, Sign Responsible AI Declaration — ‘Stealing Our Work Is Not Innovation’.” <https://www.digitalmusicnews.com/2026/01/22/stealing-isnt-innovation/>. *Dream Machine* Issue 16.
23. For Disney’s parallel position, see *Deadline*, “Disney Sends Cease And Desist Letter To Character.ai.” <https://deadline.com/2025/09/disney-cease-and-desist-letter-characterai-copyright-infringement-1236566831/>. For Studio Ghibli’s similar stance: *NDTV Profit*, “Studio Ghibli and Studio That Developed Elden Ring Send Stern Message To OpenAI.” <https://www.ndtvprofit.com/technology/studio-ghibli-and-studio-that-developed-elden-ring-send-stern-message-to-openai>. *Dream Machine* Issues 2, 6.
24. Adobe, *Creators’ Toolkit Report*, op. cit. 69% of 16,000 surveyed creators worried about their work being used to train AI without consent.
25. Adobe Firefly milestone and adoption data, in *Appendix E: Dynamics of Generative AI Adoption*, §“The Ubiquity of AI in Visual and Digital Arts.” Firefly Foundry and Firefly Image Model 5 launch reporting, Adobe MAX 2025:

<https://news.adobe.com/news/2025/10/adobe-max-2025-firefly-foundry>;
<https://news.adobe.com/news/2025/10/adobe-max-2025-firefly>.

26. Bria AI consent-licensed dataset and attribution mechanism. [TODO: confirm primary citation — Bria’s licensed-data white paper or Series B coverage.]

27. Getty Images, “Generative AI by iStock” launch, built on NVIDIA Picasso, trained exclusively on Getty’s licensed library with contributor royalties. [TODO: confirm citation — Getty press release or *Reuters* coverage.]

28. Moonvalley Marey, generative-video foundation model trained on licensed video. [TODO: confirm citation — Moonvalley launch coverage in *The Verge* / *TechCrunch*.]

29. AIODE, ethically-trained music creation DAW. See Chapter 16: The Tools, §“Audio modality models.”

30. Stability AI / Universal Music Group strategic alliance: <https://stability.ai/news/universal-music-group-and-stability-ai-announce-strategic-alliance>. Stability AI / Warner Music: <https://stability.ai/news/warner-music-group-and-stability-ai-join-forces-to-build-next-gen-tools>. Universal Music / Splice partnership: <https://www.universalmusic.com/universal-music-group-and-splice-to-collaborate-on-the-next-generation-of-ai-powered-music-creation-tools-for-artists/>. *Dream Machine* Issues 5, 8, 12.

31. Reporting on AI-generated images in the Adobe Stock training corpus, *Bloomberg*, April 2024. [TODO: confirm exact citation.]

32. Adobe Firefly IP indemnification for enterprise customers. [TODO: confirm citation — Adobe enterprise terms or *The Verge* coverage from 2023.]

33. Microsoft, “Microsoft announces new Copilot Copyright Commitment for customers,” 7 September 2023. <https://blogs.microsoft.com/on-the-issues/2023/09/07/copilot-copyright-commitment-ai-legal-concerns/>.

34. Google Cloud Generative AI indemnification: <https://cloud.google.com/blog/products/ai-machine-learning/protecting-customers-with-generative-ai-indemnification>.

35. IBM watsonx uncapped indemnity for enterprise customers. [TODO: confirm citation.]

36. Sundance AI Literacy Initiative, in Chapter 12: Authenticity, the New Scarcity, §“The provenance infrastructure, named.”

37. *Musically*, “BPI sets out transparency and sovereignty demands to secure ‘AI licensing boom’.” <https://musically.com/2026/05/19/bpi-transparency-sovereignty-ai-licensing-boom/>. *Dream Machine* Issue 30.

37a. Spotify and Universal Music Group, “Landmark Licensing Agreements for Fan-Made Covers and Remixes,” 21 May 2026. Spotify Newsroom: <https://newsroom.spotify.com/2026-05-21/universal-music-group-spotify-licensing-agreements-fan-made-covers-remixes/>. *TechCrunch*: <https://techcrunch.com/2026/05/21/spotify-and-universal-music-strike-deal-allowing-fan->

made-ai-covers-and-remixes/. *Medianama*: <https://www.medianama.com/2026/05/223-spotify-umg-fans-create-licensed-ai-covers-remixes/>. *Dream Machine* Issue 31. Detailed treatment of the three-Cs framework, the deal’s prehistory in the UMG settlement with Udio (October 2025), the *Verified by Spotify* badge, and the broader analysis in Appendix L.

37b. Hasbro, “Hasbro Launches Sixth Wall, a New AI Studio,” BusinessWire press release, 3 June 2026: <https://www.businesswire.com/news/home/20260603297922/en/>. *The Wrap*: <https://www.thewrap.com/industry-news/business/hasbro-ai-studio-sixth-wall-launch/>. *Screen Daily*: <https://www.screendaily.com/news/hasbro-launches-subsidiary-ai-studio-sixth-wall-to-exploit-ip/5217380.article>. CharacterOS runtime governance layer; talent-participation pay-out structure for original performers; initial 12-character catalogue; 13+/enterprise focus. Detailed treatment in Appendix L. Cross-reference in Chapter 7 §“Position Five: Own the rails.”

37c. *Dream Machine* Deep Dive №14, *The Programmable Brand: How AI Is Reshaping IP Licensing in the Generative Economy*, 4 June 2026. Preserved in full as Appendix L. The deep dive’s *three Cs* framework draws on the Spotify–UMG licensing agreements; the five-model comparison covers Hasbro/Sixth Wall (Behavioural Licensing), Spotify/UMG (platform-embedded remixing), Disney/OpenAI (the framework that survived the platform’s collapse — see footnote ^3a in Chapter 7), the independent-artist commons-and-DAO models (Grimes/Elf.tech with a ~50% master-royalty split; Holly Herndon / Holly+ with DAO governance; ElevenLabs heritage voices via CMG Worldwide; Meta’s 2023 celebrity-chatbot programme as the closed-static anti-pattern), and Suno/Alexa as the *create-to-consume* demand floor. The underlying economic argument draws on Coase (1937), Williamson, and Barney’s VRIN resource-based view; the appendix’s load-bearing claim is that AI *relocates* rather than *abolishes* transaction costs — collapsing search and negotiation while raising the cost of monitoring and enforcement, and pushing the durable value into the governance layer.

38. *MusicTech*, “Tamber is an ‘ethically trained’ AI tool to aid the creative process – and you can use arm gestures to control it.” <https://musictech.com/news/gear/tamber-ai-ethically-trained-arm-gestures/>. Tamber product page: <https://tamber.ai/>. *Dream Machine* Issue 30.

39. *Variety*, “Is ‘AI Resistance’ Setting the Music Sector Back? WMG’s Robert Kyncl Sees ‘An Incredible Value Creation Opportunity,’ But Warns ‘We Cannot Wait the Way the Industry Did 25 Years Ago’.” <https://variety.com/2026/music/news/wmg-robert-kyncl-ai-resistance-1236748901/>. *Dream Machine* Issue 30.

Chapter 7 – The Studios Decide

1. *CNBC*, “Netflix ‘all in’ on leveraging AI as the tech creeps into entertainment industry,” 22 October 2025. <https://www.cnn.com/2025/10/22/netflix-all-in-on-leveraging-ai-in-its-streaming-platform.html>. *Dream Machine* Issue 4.

2. *Futurism*, “Lionsgate’s Attempt to Create Movies Using AI Has Crumbled Into Disaster.” <https://futurism.com/artificial-intelligence/lionsgate-movies-ai>. *Dream Machine* Issue 1.

3. *The Guardian*, “Disney to invest \$1bn in OpenAI, allowing characters in Sora video tool.” <https://www.theguardian.com/business/2025/dec/11/disney-open-ai-sora-video-deal>.

Dream Machine Issue 11. The announcement was widely treated at the time as a completed, operative template; on the corrected record set out in *Appendix L*, no formal agreement was ever signed and no money changed hands.

3a. OpenAI staged Sora wind-down announced 24 March 2026: consumer app and website to end in late April 2026; developer API to close in late September 2026. Disney exited the partnership with reportedly as little as thirty minutes' notice. *Ropes & Gray*, "The Deal That Wasn't: What Disney and OpenAI Teach Us About Strategic Investments." <https://www.ropesgray.com/en/insights/viewpoints/the-deal-that-wasnt-what-disney-and-openai-teach-us-about-strategic-investments>. *TheStreet* / AOL (Hollywood Reporter source): <https://www.thestreet.com/technology/openai-is-shutting-down-sora-and-the-disney-deal-is-off>. *Tech-Insider*: <https://tech-insider.org/openai-sora-shutdown-disney-deal-ai-video-2026/>. Peak Sora revenue reported at roughly \$540k/month against operating costs in the billions; OpenAI reallocated compute toward higher-margin coding and enterprise products. Coverage in *Dream Machine* Issue 22.

3b. *Dream Machine* Deep Dive №14, *The Programmable Brand*, 4 June 2026. Preserved in full as *Appendix L*.

4. PYMNTS, "Retention Is Name of the Game for Netflix's AI Strategy." <https://www.pymnts.com/subscription-commerce/2026/retention-is-name-of-the-game-for-netflixs-ai-strategy/>. *Dream Machine* Issue 15.

5. *Deadline*, "Amazon Builds Out AI Studios With Sports Docs Boss Matt Newman Named Head Of Live-Action." <https://deadline.com/2025/11/amazon-ai-studios-matt-newman-1236603477/>. *Dream Machine* Issue 7.

6. *Wired*, "Amazon's House of David Used Over 350 AI Shots in Season 2. Its Creator Isn't Sorry." <https://www.wired.com/story/amazons-house-of-david-used-over-350-ai-shots-in-season-2-its-creator-isnt-sorry/>. *Dream Machine* Issue 7.

7. *Video Games Chronicle*, "NBCUniversal signs deal with Law & Order creator Dick Wolf's son to make AI-generated games based on its IP." <https://www.videogameschronicle.com/news/nbcuniversal-signs-deal-with-law-order-creator-dick-wolfs-son-to-make-ai-generated-games-based-on-its-ip/>. *Dream Machine* Issue 5.

8. *NME*, "'The Office', 'Saturday Night Live' and 'Sex And The City' could be turned into AI games." <https://www.nme.com/news/gaming-news/the-office-and-sex-and-the-city-ai-video-games-3901630>. *Dream Machine* Issue 5.

9. *The Hollywood Reporter*, "Disney+ to Allow User-Generated Fan Content with AI." <https://www.hollywoodreporter.com/business/digital/disney-plus-gen-ai-user-generated-content-1236426135/>. *Dream Machine* Issue 8.

10. *Dream Machine* Issue 8 reportage of the Disney "Office of Technology Enablement," led by former Walt Disney Studios CTO Jamie Voris.

11. *Marketing Dive*, "Disney unveils TikTok-like vertical video, AI video generation tool." <https://www.marketingdive.com/news/disney-unveils-tiktok-like-vertical-video-ai-video->

[generation-tool/809269/](#). *Dream Machine* Issue 14.

12. *The Hollywood Reporter*, “Fox Entertainment Takes Equity Stake in AI-Microdramas Company Holywater.” <https://www.hollywoodreporter.com/business/business-news/fox-entertainment-invests-in-holywater-ai-microdramas-1236396802/>. *Dream Machine* Issue 3.

13. *Deadline*, “Sky History Acquires ‘Castles SOS,’ AI-Powered Doc Exploring Royalty, Ruins & Restoration.” <https://deadline.com/2025/11/castles-sos-ai-doc-sky-history-documentary-rick-edwards-1236627378/>. *Dream Machine* Issue 9.

14. *Estate Agent Today*, “Homebuilder among first to use Channel 4’s AI ads.” <https://www.estateagenttoday.co.uk/breaking-news/2025/12/homebuilder-among-first-to-use-channel-4s-ai-ads/>. *Dream Machine* Issue 11.

15. *The Hollywood Reporter*, “Fremantle Names Boss of New AI Native Studio Imaginae Studios.” <https://www.hollywoodreporter.com/business/digital/fremantle-names-ceo-new-ai-label-imaginae-studios-1236396579/>. *Dream Machine* Issue 2.

16. *Dream Machine* Issue 25, on Fremantle’s *Art Awakens* development.

17. *Indiewire*, “Another New AI Production Company Inks a Big Creative Partnership — This Time, with Ron Howard and Brian Grazer’s Imagine Entertainment.” <https://www.indiewire.com/news/business/obsidian-studio-ai-production-company-imagine-entertainment-1235158619/>. *Dream Machine* Issue 6.

18. *UK Tech News*, “AI film studio Wonder lands \$9m investment.” <https://www.uktech.news/ai/ai-film-studio-wonder-lands-9m-investment-20251023>. *Dream Machine* Issue 5.

19. Wonder Studios, “Shortlisted films revealed for The Wonder Film Festival.” https://www.linkedin.com/posts/wearewonderstudios_were-thrilled-to-share-the-shortlisted-films-activity-7404560378082246656-7NcI. *Dream Machine* Issue 11.

20. *The Hollywood Reporter*, “AI Company Asteria Produces New Animated Short ‘All Heart’.” <https://www.hollywoodreporter.com/movies/movie-news/natasha-lyonne-ai-company-asteria-1236403144/>. *Dream Machine* Issue 4.

21. *The Hollywood Reporter*, “Promise, a deep-pocketed AI studio backed by Google, aims to Bring GenAI Filmmaking and VFX to Legacy Media.” <https://www.hollywoodreporter.com/business/digital/ai-studio-promise-vfx-generation-company-1236397636/>. *Dream Machine* Issue 3.

22. *Variety*, “AI-Powered Cinematic Universe Platform enGEN3 Launched by Goldfinch.” <https://variety.com/2025/film/news/ai-powered-cinematic-universe-platform-engen3-1236543349/>. *Dream Machine* Issue 2.

23. *Deadline*, “Munich Based Beta Films & Industry Execs Join Forces To Launch Artificial Intelligence Start-Up Chapter41.” <https://deadline.com/2025/11/beta-film-ai-startup-chapter41-artificial-intelligence-1236612632/>. *Dream Machine* Issue 7.

24. *The Hollywood Reporter*, “Longtime TV Exec, Kevin Reilly, Set to Lead AI Startup Kartel.” <https://www.hollywoodreporter.com/business/digital/kevin-reilly-ceo-kartel-ai->

hbo-1236424692/. *Dream Machine* Issue 7.

25. Variety, “‘Wanted’ Director Timur Bekmambetov Explains His \$5 Million Plan to Generate AI Method Actors: ‘AI Is Here to Stay. We Have to Train It Responsibly’.” <https://variety.com/2025/film/news/wanted-director-method-acting-ai-actors-1236579647/>. *Dream Machine* Issue 7.

26. Variety, “Tilly Norwood Creator Doubles Down on AI ‘Actors’ and Says It’s a ‘More Ethical Way to Perform,’ Urges Human Actors to ‘Future-Proof’ Themselves With AI.” <https://variety.com/2026/digital/news/tilly-norwood-creator-tells-actors-to-create-ai-avatars-1236638940/>. *Dream Machine* Issue 16.

27. *Broadcast Now*, “Wonder Studios adapts children’s book to animated series with AI.” <https://www.broadcastnow.co.uk/production-and-post/wonder-studios-adapts-childrens-book-to-animated-series-with-ai/5211713.article>. *Dream Machine* Issue 11.

28. Variety, “‘Watch the Skies,’ Swedish UFO Feature Film Dubbed Entirely With AI, Sets USA Distribution Deal.” <https://variety.com/2025/film/news/watch-the-skies-us-theatrical-release-ai-dubbing-1236343110/>. *Dream Machine* Issue 5.

29. *Cybernews*, “Run to the West — South Korea’s first AI film tests the soul of cinema.” <https://cybernews.com/entertainment/korean-cinema-run-to-the-west-ai/>. *Dream Machine* Issue 5.

30. *Broadcast Pro Middle East*, “Tunisian filmmaker wins \$1 million AI Film Award for ‘Lily.’” <https://www.broadcastprome.com/news/tunisian-filmmaker-wins-1-million-ai-film-award-for-lily/>. *Dream Machine* Issue 14.

31. Variety, “AI Drama ‘Humans in the Loop’ Receives Film Independent’s Sloan Distribution Grant, Enters Oscar Race.” <https://variety.com/2025/film/news/ai-drama-humans-in-the-loop-oscar-race-1236582975/>. *Dream Machine* Issue 8.

32. *PC Gamer*, “Palworld studio Pocketpair says its new publishing division won’t handle games that use generative AI: ‘We don’t believe in it.’” <https://www.pcgamer.com/software/ai/palworld-studio-pocketpair-says-its-new-publishing-division-wont-handle-games-that-use-generative-ai-we-dont-believe-in-it/>. *Dream Machine* Issue 4.

33. *Niche Gamer*, “Larian Studios backs off from gen AI.” *Dream Machine* Issue 14.

33a. Larian Studios policy framing on the next *Divinity*, January 2026; same source as [^33].

34. *Decrypt*, “Warhammer 40,000 Maker Games Workshop Rules Out Generative AI.” *Dream Machine* Issue 14.

35. *Niche Gamer*, “Manor Lords publisher Hooded Horse won’t work with devs using gen AI.” *Dream Machine* Issue 14.

36. *gamesindustry.biz*, “RuneScape maker Jagex says it will never use generative AI to make in-game content.” *Dream Machine* Issue 16.

37. *GamesRadar*, “Wallace and Gromit creator says beloved animation studio Aardman will ‘embrace the technology’ of AI, but will be ‘very cautious not to lose our values’.”

<https://www.gamesradar.com/entertainment/animation-movies/wallace-and-gromit-creator-says-beloved-animation-studio-aardman-will-embrace-the-technology-of-ai-but-will-be-very-cautious-not-to-lose-our-values/>. *Dream Machine* Issue 11.

38. Variety, “Guillermo del Toro Says He’d ‘Rather Die’ Than Use Generative AI in His Films: ‘Not Interested’.” <https://variety.com/2025/film/news/guillermo-del-toro-rather-die-generative-ai-frankenstein-1236561316/>. *Dream Machine* Issue 5.

39. *The Hollywood Reporter*, “Leonardo DiCaprio Says AI Can’t Be Art Because ‘There’s No Humanity to It’.” <https://www.hollywoodreporter.com/movies/movie-news/leonardo-dicaprio-ai-cant-be-art-no-humanity-1236445405/>. *Dream Machine* Issue 11.

40. *Daily Mail*, “Claire Foy says she has ‘no interest’ in seeing AI in films.” <https://www.dailymail.co.uk/tvshowbiz/article-15454199/Claire-Foy-AI-films-sad-disappointed-people-future-Hollywood.html>. *Dream Machine* Issue 14.

41. *NME*, “Jenna Ortega says it’s ‘very easy to be terrified’ of AI in filmmaking.” <https://www.nme.com/news/jenna-ortega-says-its-very-easy-to-be-terrified-of-ai-in-filmmaking-3913926>. *Dream Machine* Issue 10.

42. Variety, “Chris Pratt Pitched Having an AI ‘Actor’ Star as the Villain in ‘Mercy’: ‘I Don’t Think That’s a Good Idea at All’.” <https://variety.com/2026/film/news/chris-pratt-ai-actor-villain-mercy-amazon-mgm-1236640460/>. *Dream Machine* Issue 16.

43. *PC Gamer*, “Todd Howard says AI can’t replace human ‘creative intention,’ but it’s part of Bethesda’s ‘toolset for how we build our worlds or check things’.” <https://www.pcgamer.com/gaming-industry/todd-howard-says-ai-cant-replace-human-creative-intention-but-its-part-of-bethedas-toolset-for-how-we-build-our-worlds-or-check-things/>. *Dream Machine* Issue 11.

44. *GamesRadar*, “Battlefield 6 lead calls generative AI ‘very seducing,’ but says it was only used in the game’s earliest stages ‘to allow for more time and more space to be creative’.” <https://www.gamesradar.com/games/battlefield/battlefield-6-lead-calls-generative-ai-very-seducing-but-says-it-was-only-used-in-the-games-earliest-stages-to-allow-for-more-time-and-more-space-to-be-creative/>. *Dream Machine* Issue 3.

44a. *Wired*, “Amazon’s House of David Used Over 350 AI Shots in Season 2. Its Creator Isn’t Sorry.” <https://www.wired.com/story/amazons-house-of-david-used-over-350-ai-shots-in-season-2-its-creator-isnt-sorry/>. *Dream Machine* Issue 7.

45. *gamesindustry.biz*, “Witcher 3 and Cyberpunk 2077 director says AI can help, but not replace, creatives.” <https://www.gamesindustry.biz/witcher-3-and-cyberpunk-2077-director-says-ai-can-help-but-not-replace-creatives>. *Dream Machine* Issue 9.

46. *GamesRadar*, “Aardman” *op. cit.*

47. *Dream Machine* Issue 29, on Sony’s “all in on AI for games” announcement.

48. *Hollywood Reporter*, “Netflix is building and recruiting for an AI animation studio, called INKubator, to produce ‘feature-quality’ shorts.” <https://www.hollywoodreporter.com/business/business-news/netflix-ai-animation-studio-inkubator-1236592110/>. *Dream Machine* Issue 30.

49. *Forbes*, “Meet Wonder Studios, The \$50M British Studio Striving To Become The A24 Of AI Production.” <https://www.forbes.com/sites/charliefink/2026/05/18/meet-wonder-studios-the-50m-british-studio-striving-to-become-the-a24-of-ai-production/>. *Dream Machine* Issue 30.

50. *Variety*, “Kling AI Partners With Evolutionary Films on Animated Feature ‘Minibots,’ Unveils Filmmaker Initiative at Cannes Market.” <https://variety.com/2026/film/news/kling-ai-evolutionary-films-minibots-cannes-1236748590/>. *Dream Machine* Issue 30.

51. *Variety*, “AI Dominates Cannes Buzz as Filmmakers Grudgingly Accept It.” <https://variety.com/2026/film/festivals/ai-cannes-2026-filmmakers-accept-1236748402/>; *Hollywood Reporter*, “At Cannes, filmmakers shift towards cautious acceptance of AI’s inevitability.” <https://www.hollywoodreporter.com/business/business-news/cannes-2026-ai-acceptance-1236592488/>. *Dream Machine* Issue 30.

52. *Variety*, “Is AI Basically Like Special Effects? Peter Jackson Seems to Think So.” <https://variety.com/2026/film/news/peter-jackson-ai-special-effects-1236748120/>. *Dream Machine* Issue 30.

53. *PC Gamer*, “Take-Two’s CEO says AI’s not in the business of making hits, ‘datasets by their very nature are backward looking’, but that doesn’t mean AI can’t be ‘super helpful’.” <https://www.pcgamer.com/games/take-two-ceo-ai-not-making-hits-backward-looking/>. *Business Insider*, “The CEO behind Grand Theft Auto says he’s pro AI — but the technology can’t make an original hit.” <https://www.businessinsider.com/take-two-ceo-strauss-zelnick-ai-original-hits-2026-5>. *Dream Machine* Issue 30.

54a. Hasbro launches Sixth Wall, a new AI studio, with the CharacterOS governance layer underneath it, 3 June 2026. BusinessWire press release: <https://www.businesswire.com/news/home/20260603297922/en/>. *The Wrap*, “Hasbro AI Studio Sixth Wall Launch”: <https://www.thewrap.com/industry-news/business/hasbro-ai-studio-sixth-wall-launch/>. *Screen Daily*: <https://www.screendaily.com/news/hasbro-launches-subsidiary-ai-studio-sixth-wall-to-exploit-ip/5217380.article>. *BriefGlance* analysis (FTC inquiries into AI companions, Sixth Wall’s stated 13+/enterprise focus): <https://briefglance.com/articles/hasbros-sixth-wall-taming-the-ai-wild-west-with-behavioral-licensing>. Detailed treatment in Appendix L. CEO Roberta Thomson; CharacterOS encodes character canon, voice, personality, lore and safety limits at runtime; talent-participation pay-out structure for the original performers; initial 12-character catalogue; pilot grew out of an internal generative-AI Ouija-board experiment that logged hundreds of thousands of player interactions in its first few days.

54b. Spotify and Universal Music Group announce landmark linked recorded-music and publishing licensing agreements for fan-made AI covers and remixes at Spotify Investor Day, 21 May 2026. Spotify Newsroom: <https://newsroom.spotify.com/2026-05-21/universal-music-group-spotify-licensing-agreements-fan-made-covers-remixes/>. *TechCrunch*: <https://techcrunch.com/2026/05/21/spotify-and-universal-music-strike-deal-allowing-fan-made-ai-covers-and-remixes/>. *Medianama*: <https://www.medianama.com/2026/05/223-spotify-umg-fans-create-licensed-ai-covers-remixes/>. Three-pillar consent / credit / compensation framework; opt-in catalogue; Premium add-on; new royalty stream; UMG

settlement with Udio in October 2025 is the architectural ancestor. Detailed treatment in Appendix L. Cross-reference in Chapter 5 and Chapter 6.

5a. Amazon AI Creators Fund launch and the Punky Duck cancellation. *Hollywood Reporter*: <https://www.hollywoodreporter.com/tv/tv-news/amazon-orders-three-animated-series-generative-ai-1236606219/>. *Variety*: <https://variety.com/2026/tv/news/jorge-gutierrez-drops-out-amazon-mgm-ai-generated-series-backlash-1236762285/>. *Futurism*: <https://futurism.com/artificial-intelligence/amazon-ai-generated-series-canceled>. *Dream Machine* Issue 32. Cross-reference: Chapter 5 §“The pushback” for the audience-enforcement analysis.

55. *Variety*, “Tribeca Festival AI Film ‘Dreams of Violets’ by Fountain o.” <https://variety.com/2026/film/festivals/tribeca-festival-ai-film-dreams-of-violets-fountain-o-1236759724/>. Director Ash Koosha; producer Fountain o; 75-minute docudrama; subject: 2026 Iran protests; cost \$2,000; production time: 3 months; tools: Kling AI (video), Anthropic Claude (editing), Google Gemini (research). *Dream Machine* Issue 32.

56. The broader festival-circuit context: the *Dreams of Violets* Tribeca selection follows the *Lily* (\$1M Dubai AI Film Award, January 2026) and *Humans in the Loop* (Film Independent Sloan Distribution Grant, November 2025) as part of the accumulating festival recognition of AI-native feature and short-film work — but *Dreams of Violets* is the first Tier 1 festival selection for a feature-length, fully-AI-generated work.

57. *Variety*, “Martin Scorsese Supports AI Company, Endorses AI for Storyboarding.” <https://variety.com/2026/film/news/martin-scorsese-supports-ai-company-storyboard-movies-1236765037/>. Scorsese becomes advisor to Black Forest Labs (FLUX model); endorses AI storyboarding; Boots Riley and Guillermo del Toro respond critically. *Dream Machine* Issue 32.

Chapter 8 — Worlds, Not Pictures

1. World Labs, “Bringing Marble to Life.” <https://www.worldlabs.ai/case-studies/bringing-marble-to-life>. *Dream Machine* Issue 7, “Editor’s Pick: Marble by WorldLabs goes on public release,” 13 November 2025.

2. For a working primer on Gaussian splatting in the post-Marble era, see *Radiance Fields*, “World Labs Formally Launches Marble, A Generative World Model.” <https://radiancefields.com/world-labs-formally-launches-marble-a-generative-world-model>.

3. DreamLab AI Collective, beta participation in Marble, October–November 2025. Referenced in *Dream Machine* Issue 7: “DreamLab have been part of the beta testing for this over the last few months and it’s very neat.”

4. SuperSplat (PlayCanvas), open-source Gaussian splat editor, regular updates through 2025–26. *Dream Machine* Issue 1: “PlayCanvas open sources SOG — WebP for 3D Gaussian Splatting”; Issue 7 / Issue 11 on SuperSplat v2 updates.

5. Sony Pictures’ use of Marble in Virtual Production: https://www.linkedin.com/posts/brent-liang_tech-media-launch-ugcPost-

7394911181091692546-TyUz. *Dream Machine* Issue 8.

6. Disney “300,000 poses in an instant” livestream, March 2026. *Dream Machine* Issue 23.

7. Netflix + Eyeline, *Vista4D*: 4D point clouds from live-action. *Dream Machine* Issue 27.

8. Google DeepMind, “Genie 3: A new frontier for world models.” <https://deepmind.google/blog/genie-3-a-new-frontier-for-world-models/>. Project Genie roll-out to AI Ultra subscribers in the U.S.: <https://blog.google/innovation-and-ai/models-and-research/google-deepmind/project-genie/>. *Dream Machine* Issue 3 (initial announcement) and Issue 17 (broader availability).

9. Meta, “WorldGen — text-to-immersive-3D-worlds research update.” <https://www.facebook.com/LifeAtMeta/videos/research-update-worldgen-text-to-immersive-3d-worlds/1879077432692421/>. *Dream Machine* Issues 9, 11.

10. Tencent, “HY World 1.5” announcement: <https://x.com/TencentHunyuanyuan/status/2001170499133653006>. *Dream Machine* Issue 12.

11. SpAItial, *ECHO* spatial foundation model. <https://www.spaitial.ai/>. *Dream Machine* Issue 12.

12. Stanford AI Lab, *Wonderzoom*: Multi-Scale 3D World Generation. <https://wonderzoom.github.io/>. *Dream Machine* Issue 14.

13. OpenArt, *Worlds* product launch, March 2026. *Dream Machine* Issue 21.

14. Luma AI, *UNI-1* launch, March 2026. *Dream Machine* Issue 22, “Editor’s Pick: When worlds become instant, the race shifts to better thinking.”

15. ByteDance Seedance 2.0 in CapCut/Dreamina, March 2026. *Dream Machine* Issue 22.

16. *Spark 2.0*, open-source Gaussian-splat streaming framework, April 2026. *Dream Machine* Issue 25.

17. Radiance Fields, “Apple Confirms that it’s Gaussian Splatting that powers their personas.” <https://radiancefields.com/apple-confirms-personas-use-gaussian-splatting/>. *Dream Machine* Issue 5.

18. *Video Games Chronicle*, “‘It honestly sucks’: Fans think Call of Duty: Black Ops 7 is filled with generative AI art.” <https://www.videogameschronicle.com/news/it-honestly-sucks-fans-think-call-of-duty-black-ops-7-is-filled-with-generative-ai-art/>. *Video Games Chronicle*, “Ubisoft says AI-generated art in Anno 117 was a placeholder which ‘slipped through our review process’.” <https://www.videogameschronicle.com/news/ubisoft-says-ai-generated-art-in-anno-117-was-a-placeholder-which-slipped-through-our-review-process/>. *Polygon*, “Fortnite chapter 7 kicks off new controversy over AI art.” <https://www.polygon.com/fortnite-chapter-7-season-1-generative-ai-art-epic-games/>. *Dream Machine* Issues 8, 10.

19. NVIDIA + Stanford, *NitroGen*. <https://nitrogen.minedojo.org/>. *Dream Machine* Issue 13.

20. DeepMind, “SIMA 2: An Agent that Plays, Reasons, and Learns With You in Virtual 3D Worlds.” <https://deepmind.google/blog/sima-2-an-agent-that-plays-reasons-and-learns->

with-you-in-virtual-3d-worlds/. *Dream Machine* Issue 8.

21. ComfyUI Blog, “Ubisoft La Forge Open-Sources the CHORD Model and ComfyUI Nodes for End-to-End PBR Material Generation.” <https://blog.comfy.org/p/ubisoft-open-sources-the-chord-model>. *Dream Machine* Issue 11.

22. *Video Games Chronicle*, “The future of gaming, or ‘just a tool’? Hands-on with Teammates, Ubisoft’s ambitious voice AI tech demo.” <https://www.videogameschronicle.com/features/the-future-of-gaming-or-just-a-tool-hands-on-with-teammates-ubisofts-ambitious-voice-ai-tech-demo/>. *Dream Machine* Issue 9.

23. YouTube Playables Builder, closed-beta announcement: <https://www.youtube.com/playablesbuilder/>. *Dream Machine* Issue 12.

24. Unity AI Open Beta, in-editor AI suite, May 2026. *Dream Machine* Issue 28.

25. Korin AI, “trained with African datasets, built by Africans,” May 2026. *Dream Machine* Issue 27.

26. NVIDIA SANA-WM, 2.6B open-source world model with 60-second video generation and camera control, May 2026. <https://huggingface.co/collections/nvidia/sana-wm>. *Dream Machine* Issue 30.

27. Odyssey, “Introducing Starchild-1, the first real-time multimodal world model.” <https://odyssey.ml/introducing-starchild-1>. *Dream Machine* Issue 30.

28. Apple Machine Learning Research, “Headsup: a large-scale high-quality 3D Gaussian head reconstruction from multi-view captures.” <https://machinelearning.apple.com/research/apple-headsup-3d-gaussian-head>. *Dream Machine* Issue 30.

29. WorldLens VR, “AI-powered 3D depth for Google Street View on Quest.” <https://www.uploadvr.com/worldlens-vr-quest-street-view-3d-depth/>. *Dream Machine* Issue 30.

Chapter 9 — AI in Everything, Everywhere, All at Once

1. *Creative Boom*, “Adobe is putting AI in everything everywhere all at once.” <https://www.creativeboom.com/news/adobe-is-putting-ai-in-everything-everywhere-all-at-once/>. *Dream Machine* Issue 5, “Editor’s Pick,” 31 October 2025.

2. Adobe, “Adobe MAX 2025: Firefly Foundry.” <https://news.adobe.com/news/2025/10/adobe-max-2025-firefly-foundry>.

3. Adobe, “Adobe MAX 2025: Firefly.” <https://news.adobe.com/news/2025/10/adobe-max-2025-firefly>.

4. Adobe, “Adobe MAX 2025: Express AI Assistant.” <https://news.adobe.com/news/2025/10/adobe-max-2025-express-ai-assistant>.

5. *Wired*, “Adobe’s ‘Corrective AI’ Can Change the Emotions of a Voice-Over” and accompanying Adobe Sneaks 2025 coverage. <https://www.wired.com/story/adobe-max->

sneaks-2025-corrective-ai/. Project list compiled from MAX keynote and *Dream Machine* Issue 5 coverage.

6. *PYMNTS*, “Adobe Lets Users Design and Edit Using ChatGPT.” <https://www.pymnts.com/artificial-intelligence-2/2025/adobe-lets-users-design-and-edit-using-chatgpt/>. Adobe blog: “Edit images, designs, and PDFs right inside ChatGPT.” <https://blog.adobe.com/en/publish/2025/12/10/edit-photoshop-chatgpt>. *Dream Machine* Issue 12.

7. Adobe Premiere Object Mask tool: https://www.linkedin.com/posts/robdewinter_ok-this-is-going-to-save-a-lot-of-time-in-ugcPost-7421617551690063872-yKmB. *Dream Machine* Issue 16.

8. Adobe blog, “Sundance Film Festival 2026: Creativity, Community & Power of Storytelling.” <https://blog.adobe.com/en/publish/2026/01/20/sundance-film-festival-2026-creativity-community-power-storytelling>. *Dream Machine* Issue 16.

9. Adobe Summit 2026, “agentic creative intelligence” keynote. *Dream Machine* Issue 26.

10. After Effects AI animation features through late 2025: *Dream Machine* Issue 9, “AI video is finally animatable inside After Effects.” https://www.linkedin.com/posts/thisisdoug_ai-aivideo-animation-ugcPost-7399512745924067330-Aldk.

10a. Doug McGinness on LinkedIn, late 2025, in the same post. *Dream Machine* Issue 9.

11. *Dream Machine* Issue 21, “Editor’s Pick: Adobe and NVIDIA Just Raised the Stakes for Creative AI,” 19 March 2026.

12. NVIDIA + Google Cloud creative-AI infrastructure deal, March 2026. *Dream Machine* Issue 21.

13. Hugging Face and Google Cloud partnership announcement: https://www.linkedin.com/posts/julienchaumond_i-am-super-excited-to-announce-that-hugging-activity-7396177403972276225-CuMM. *Dream Machine* Issue 8.

14. *EdTech Innovation Hub*, “Meta and Hugging Face launch OpenEnv to advance open-source agentic development.” <https://www.edtechinnovationhub.com/news/meta-and-hugging-face-launch-openenv-to-advance-open-source-agentic-development>. *Dream Machine* Issue 5.

15. Anthropic / Blender Foundation patronage, May 2026. *Dream Machine* Issue 27.

16. TechCrunch, “Anthropic launches interactive Claude apps, including Slack and other workplace tools.” <https://techcrunch.com/2026/01/26/anthropic-launches-interactive-claude-apps-including-slack-and-other-workplace-tools/>. *Dream Machine* Issue 16.

17. Spotify–Anthropic integration, May 2026. *Dream Machine* Issue 27.

18. *MarTech Series*, “WPP continues AI overhaul with \$400-million Google partnership.” <https://martechseries.com/predictive-ai/ai-platforms-machine-learning/google-and-spotify-alum-launch-epiminds-with-6-6m-to-build-marketing-teams-for-the-ai-era/>. *Dream Machine* Issue 3.

19. Campaign Brief, “WPP launches AI-powered marketing platform WPP Open Pro.” <https://campaignbrief.com/wpp-launches-ai-powered-marketing-platform-wpp-open-pro/>. *Dream Machine* Issue 5.
20. Digiday, “WPP expands AI capabilities to boost brand performance with Slightly partnership.” <https://digiday.com/media-buying/agencies-continue-to-expand-ai-capabilities-to-boost-brand-performance/>. *Dream Machine* Issue 6.
21. WPP and Google Earth AI consumer-journey project, April 2026. *Dream Machine* Issue 27.
22. SiliconAngle, “Higgsfield raises \$80M on \$1.3B valuation to scale AI video platform.” <https://siliconangle.com/2026/01/15/higgsfield-raises-80m-1-3b-valuation-scale-ai-video-platform/>. *Dream Machine* Issue 15.
23. 36kr, “AI Video Unicorn Higgsfield: Earns \$200M in 9 Months by ‘Serving’ Social Media Marketers.” <https://eu.36kr.com/en/p/3650517574312323>. *Dream Machine* Issue 16.
24. TechCrunch, “Synthesia hits \$4B valuation, lets employees cash in.” <https://techcrunch.com/2026/01/26/synthesia-hits-4b-valuation-lets-employees-cash-in/>. *Dream Machine* Issue 16.
25. Sifted, “Synthesia rejects \$3bn Adobe acquisition offer.” <https://sifted.eu/articles/synthesia-acquisition-offer>. *Dream Machine* Issue 5.
26. ElevenLabs \$500m ARR reporting, April 2026. *Dream Machine* Issue 25.
27. Runway product cycle: Gen-4.5 (December 2025), Gen-4.5 Image-to-Video (January 2026), Workflows, Story Panels, Characters API, Apps for Advertising — *Dream Machine* Issues 10, 14, 15, 16, 20. Runway CEO on indie films vs. blockbusters: *Dream Machine* Issue 26.
28. For the running ledger of new creative-AI products through 2025–26, see *Dream Machine* Issues 1–30 archive.
29. ComfyUI, “We raised \$17 million to build an OS for Creative AI.” https://www.linkedin.com/posts/comfyui_we-raised-17-million-to-build-an-os-for-ugcPost-7373743341236236288-wkCc. *Dream Machine* Issue 1.
30. ComfyUI \$500M valuation, May 2026. *Dream Machine* Issue 27.
31. Google Pomelli launch: <https://x.com/GoogleLabs/status/1983204018567426312>. *Dream Machine* Issue 5.
32. Google AI Studio app gallery: <https://x.com/GoogleAISTudio/status/1982121563785949255>. Google Labs Opal expansion: <https://blog.google/technology/google-labs/opal-expansion/>. Project Genie: <https://blog.google/innovation-and-ai/models-and-research/google-deepmind/project-genie/>. *Dream Machine* Issues 5, 17.
33. Lovable for classrooms: <https://lovable.dev/classroom>. *Dream Machine* Issue 11.

- 34.** Adobe Express AI Assistant: <https://news.adobe.com/news/2025/10/adobe-max-2025-express-ai-assistant>. *Dream Machine* Issue 5.
- 35.** Hugging Face platform expansion through 2025–26.
- 36.** Google blog, “Sundance Institute AI Education.” <https://blog.google/company-news/outreach-and-initiatives/google-org/sundance-institute-ai-education/>. *Dream Machine* Issue 15.
- 37.** Adobe Ignite Day at Sundance: *Adobe blog, Sundance Film Festival 2026*. <https://blog.adobe.com/en/publish/2026/01/20/sundance-film-festival-2026-creativity-community-power-storytelling>. *Dream Machine* Issue 16.
- 38.** Google’s \$40bn investment in Anthropic, May 2026. *Dream Machine* Issue 27.
- 39.** UK Government, “Free AI training for all.” <https://www.gov.uk/government/news/free-ai-training-for-all-as-government-and-industry-programme-expands-to-provide-10-million-workers-with-key-ai-skills-by-2030>. *Dream Machine* Issue 16.
- 40.** CNBC, “People with ADHD, autism, dyslexia say AI agents are helping them succeed at work.” <https://www.cnn.com/2025/11/08/adhd-autism-dyslexia-jobs-careers-ai-agents-success.html>. *Dream Machine* Issue 7.
- 41.** University of Wisconsin-Stout, “AI Reshaping Industry: New UW-Stout Course Sets AI-Use as Baseline Competency in Filmmaking.” <https://www.uwstout.edu/about-us/news-center/ai-reshaping-industry-new-uw-stout-course-sets-ai-use-baseline-competency-filmmaking>. *Dream Machine* Issue 15.
- 42.** Google I/O 2026 announcement block: Gemini Omni <https://blog.google/technology/google-deepmind/gemini-omni/>, Antigravity <https://antigravity.google/>, Google Flow <https://flow.google/>, Gemini Spark <https://blog.google/technology/developers/gemini-spark/>, Project Genie + Street View <https://deepmind.google/discover/blog/project-genie-street-view/>. *Dream Machine* Issue 30, “Editor’s Pick — Google I/O 2026,” 21 May 2026.
- 43.** Google Labs, “Infinite Scaler.” <https://blog.google/technology/google-labs/infinite-scaler/>. *Dream Machine* Issue 30.
- 44.** Google DeepMind, “SynthID — 100 billion watermarks, partner ecosystem.” <https://deepmind.google/discover/blog/synthid-100-billion-watermarks-partners/>. *Dream Machine* Issue 30.
- 45.** Runway, “Runway Japan.” <https://runwayml.com/blog/runway-japan>. *Dream Machine* Issue 30.
- 46.** *Music Business Worldwide*, “Splice inks ‘Responsible AI’ deal with ElevenLabs.” <https://www.musicbusinessworldwide.com/splice-elevenlabs-responsible-ai-deal/>. *Dream Machine* Issue 30.
- 47.** *Adweek*, “Netflix ad tools could see ‘agentic AIs talking to each other.’” <https://www.adweek.com/media/netflix-ad-tools-agentic-ais-talking-to-each-other/>. *Dream Machine* Issue 30.

48. *Fortune*, “AI startup Viktor raises \$75 million to put a virtual ‘coworker’ in Slack and Teams.” <https://fortune.com/2026/05/19/ai-startup-viktor-75-million-virtual-coworker-slack-teams/>. *Dream Machine* Issue 30.

Chapter 10 — What is newly possible

1. Snap Newsroom, “Snapchat Gen Z AI Creativity Research 2026.” <https://newsroom.snap.com/snapchat-gen-z-ai-creativity-research-2026>. *Dream Machine* Issue 30.

Chapter 11 — The Orchestrator

1. *Dream Machine* Issue 13, “Editor’s Pick: The Year of the Orchestrator,” 9 January 2026.

2. *Dream Machine* Issue 29, May 2026, reporting on Sony’s 49-agent / 72-skill multi-agent game-development team.

3. Anthropic blog content on agent deployment patterns, Q1 2026.

3a. *Bloomberg*, “AI Changed Chess. Grandmasters Now Win With Unpredictable Moves,” 27 March 2026. <https://www.bloomberg.com/news/articles/2026-03-27/ai-changed-chess-grandmasters-now-win-with-unpredictable-moves>. *Dream Machine* Issue 23. The chess analogy is developed in Chapter 15’s *Age of the Why* section.

4. Sundance Institute, “Centering the Artist: Why We’re Launching the AI Literacy Initiative.” <https://www.sundance.org/blogs/centering-the-artist-why-were-launching-the-ai-literacy-initiative/>. *Dream Machine* Issue 16.

5. Sundance Institute, *op. cit.*

6. Sundance Story Forum 2026 sessions on legal toolkits for producers using AI. *Dream Machine* Issue 16.

7. Google blog, “Sundance Institute AI Education.” <https://blog.google/company-news/outreach-and-initiatives/google-org/sundance-institute-ai-education/>. *Dream Machine* Issue 15.

8. McKinsey & Company, “What AI could mean for film and TV production and the industry’s future.” <https://www.mckinsey.com/industries/technology-media-and-telecommunications/our-insights/what-ai-could-mean-for-film-and-tv-production-and-the-industrys-future>. *Dream Machine* Issue 16.

9. *Metro*, “Prince of Persia remake and five more games cancelled as Ubisoft focuses on AI.” <https://metro.co.uk/2026/01/21/prince-persia-remake-five-games-cancelled-ubisoft-focuses-ai-26431926/>. *Dream Machine* Issue 15.

10. *PC Gamer*, “Square Enix, makers of Final Fantasy, aims to have AI doing 70% of its QA work by the end of 2027.” <https://www.pcgamer.com/gaming-industry/square-enix-aims-to-have-ai-doing-70-percent-of-its-qa-work-by-the-end-of-2027/>. *Dream Machine* Issue 7.

11. Eurogamer, “Falcom is the latest developer to buy into the AI hype machine.” <https://www.eurogamer.net/falcom-is-the-latest-developer-to-buy-into-the-ai-hype-machine>. *Dream Machine* Issue 12.
12. NDTV Profit, “Don’t Expect AI To Invent the Next ‘Grand Theft Auto’, Says Take-Two CEO Strauss Zelnick.” <https://www.ndtvprofit.com/technology/dont-expect-ai-to-invent-the-next-grand-theft-auto-says-take-two-ceo-strauss-zelnick>. *Dream Machine* Issue 6.
13. *Dream Machine* Issue 21, on Spielberg’s public position on AI.
14. *Dream Machine* Issues 25, 28, on Steven Soderbergh’s AI work.
15. Digiday, “Independent agencies face new frontier as agency-in-a-box tools democratize creativity.” <https://digiday.com/marketing/independent-agencies-face-new-frontier-as-agency-in-a-box-tools-democratize-creativity/>. *Dream Machine* Issues 6, 14.
16. Digiday, “AI agent developers have become adland’s in-demand role.” <https://digiday.com/marketing/ai-agent-developers-have-become-adlands-in-demand-role/>. *Dream Machine* Issue 7.
17. PYMNTS, “AI Content Is Par For The Course With PGA Tour’s Expanded AWS Partnership.” <https://www.pymnts.com/artificial-intelligence-2/2026/ai-content-is-par-for-the-course-with-pga-tours-expanded-aws-partnership/>. *Dream Machine* Issue 15.
18. The Verge, “Oreo-maker Mondelez will use AI for TV ads next year.” <https://www.theverge.com/news/806047/mondelez-ai-generated-ads>. *Dream Machine* Issue 5.
19. Digiday, “Avocados From Mexico turns to AI to advertise around the Super Bowl instead of a TV buy.” <https://digiday.com/marketing/avocados-from-mexico-turns-to-ai-to-advertise-around-the-super-bowl-instead-of-a-tv-buy/>. *Dream Machine* Issue 15.
20. *Dream Machine* Issue 26.
21. Reuters Institute, “AI adoption by UK journalists and their newsrooms: surveying applications, approaches, and attitudes.” <https://reutersinstitute.politics.ox.ac.uk/ai-adoption-uk-journalists-and-their-newsrooms-surveying-applications-approaches-and-attitudes>. *Dream Machine* Issue 9.
22. Digiday, “Daily Mail says Google AI Overviews have killed click-throughs.” <https://digiday.com/media/daily-mail-says-google-ai-overviews-have-killed-click-throughs/>. *Dream Machine* Issue 7.
23. Digiday, “How The Times is using AI to model synthetic focus groups from human audiences.” <https://digiday.com/media/how-the-times-is-using-ai-to-model-synthetic-focus-groups-from-human-audiences/>. *Dream Machine* Issue 6.
24. TechBullion, “Why the future belongs to multi-skilled leaders.” <https://techbullion.com/playing-the-long-game-with-a-portfolio-career-why-the-future-belongs-to-multi-skilled-leaders/>. *Dream Machine* Issue 9.
25. Anthropic Skills framework via Claude Code, reported through *Dream Machine* Issues 11, 16, 29.

26. *Forbes*, “AI Is Changing How Creators Work And Earn.” <https://www.forbes.com/sites/kolawolesamueladebayo/2025/12/22/how-ai-is-changing-how-creators-work-and-earn/>. *Dream Machine* Issue 13.

a1. *Variety*, “AI Creator Behind Viral ‘Deadpool,’ ‘Harry Potter’ Christmas Clip Made His Film in a Ukrainian Bomb Shelter.” <https://variety.com/2026/digital/news/ai-video-deadpool-harry-potter-andrii-daniels-1236624632/>. *Dream Machine* Issue 16.

Chapter 12 — Authenticity as the New Scarcity

33a. *Variety*, “AI Music Suno Funding Round \$400 Million \$5.4 Billion Valuation.” <https://variety.com/2026/digital/news/ai-music-suno-funding-round-400-million-5-4-billion-valuation-1236765727/>. *TechCrunch*, “Still Facing Copyright Lawsuits, AI Music Generator Suno Raises Another \$400M.” <https://techcrunch.com/2026/06/03/still-facing-copyright-lawsuits-ai-music-generator-suno-raises-another-400m/>. *Dream Machine* Issue 32.

33b. Warner Music Group settled its copyright litigation against Suno and signed a licensing partnership; Universal Music Group and Sony Music remain in active litigation as of June 2026. *Dream Machine* Issue 32. For the structural context of the WMG settlement-as-licensing-table move, see Chapter 6 §“The three Cs”: *the lawsuits were never the destination; they were leverage to force a licensing table into existence.*

1. *Dream Machine* Issue 29 reportage of Tiny Grandma stop-motion content being wrongly flagged as AI by YouTube’s automated detection, May 2026.

2. *Dream Machine* Issue 23, April 2026, reporting death threats against Eline Van der Velden following Tilly Norwood’s continuing public role.

3. *Digital Music News*, “Instagram Chief Says We Should ‘Fingerprint Real Media’ Instead of Tracking and Disclosing AI Slop.” <https://www.digitalmusicnews.com/2026/01/05/instagram-chief-ai-slop-comments/>. *WebProNews*, “Instagram Head Warns AI Images Erode Trust, Calls for Verification Standards.” <https://www.webpronews.com/instagram-head-warns-ai-images-erode-trust-calls-for-verification-standards/>. *Dream Machine* Issue 13.

4. *Digital Music News*, “AI-Generated Far-Right Hate Songs Aren’t Just a Problem in the US — Now They’re Spreading Across Europe Too.” <https://www.digitalmusicnews.com/2025/11/09/ai-generated-hate-songs-dutch-spotify-charts/>. *Dream Machine* Issue 7.

5. Google DeepMind SynthID watermark roll-out across Veo, Lyria and Imagen products. *Dream Machine* Issues 11, 12.

6. Google DeepMind, “Verify Google AI-generated videos in the Gemini app.” https://www.linkedin.com/posts/googleddeepmind_verify-google-ai-generated-videos-in-the-activity-7407748300688478208-fJgW. *Dream Machine* Issue 12; broader coverage in *SmartBrief*, “Google’s Gemini can now spot AI-generated videos.” <https://newsletter.smartbrief.com/sharedSummary/index.jsp?briefId=40A39351-5419-4681->

[94DF-31A53480B698&issueId=58E986AD-821F-422E-9E34-](https://www.fox.com/story/entertainment-2025/01/01/fox-ai-voice-trademark-94DF-31A53480B698&issueId=58E986AD-821F-422E-9E34-3386E0E2272B©Id=2DB8E453-8E83-416C-949B-44751F252A8D)

[3386E0E2272B©Id=2DB8E453-8E83-416C-949B-44751F252A8D](https://www.fox.com/story/entertainment-2025/01/01/fox-ai-voice-trademark-94DF-31A53480B698&issueId=58E986AD-821F-422E-9E34-3386E0E2272B©Id=2DB8E453-8E83-416C-949B-44751F252A8D). *Dream Machine* Issue 13.

7. *Dream Machine* Issues 23, 27 reportage on Taylor Swift's voice/image trademark filings.

8. *Lawyer Monthly*, "Matthew McConaughey Draws a Line to Protect His Voice and Image From AI." <https://www.lawyer-monthly.com/2026/01/matthew-mcconaughey-protects-voice-image-ai/>. *Dream Machine* Issue 15.

9. *Adweek*, "Meet the \$1.3 Billion Startup Behind Madonna and Will Smith's AI Video." <https://www.adweek.com/media/higgsfield-ai-marketing-startup/>. *Dream Machine* Issue 16.

10. *Variety*, "George Clooney Says AI Actors Will Face the 'Same Problem We Have' in Hollywood: 'Making a Star Is Not So Easy'." <https://variety.com/2025/scene/columns/george-clooney-ai-actors-movie-stars-1236579661/>. *Dream Machine* Issue 7.

11. *Deadline*, "AI Documentary Director Insists Jeremy Renner Agreed To Narrate Movie As 'Hawkeye' Star Threatens 'Multi-Millions' Lawsuit." <https://deadline.com/2025/11/jeremy-renner-lawsuit-threat-ai-movie-1236611830/>. *Dream Machine* Issue 7.

12. *Complete Music Update*, "Johnny Cash estate uses ELVIS Act to sue Coke over tribute act ad soundtrack." <https://completemusicupdate.com/johnny-cash-estate-uses-elvis-act-to-sue-coke-over-tribute-act-ad-soundtrack/>. *Dream Machine* Issue 9.

13. *The Verge*, "New York's new law forces advertisers to say when they're using AI avatars." <https://www.theverge.com/news/842848/new-york-law-ai-advertisements-sag-aftra-labor>. *Dream Machine* Issue 11.

14. *Fast Company*, "Governments around the world are considering bans on Grok's app over AI sexual image scandal." <https://www.fastcompany.com/91474131/governments-around-the-world-are-considering-bans-on-groks-app-over-ai-sexual-image-scandal>. *Dream Machine* Issue 14.

15. Cannes AI Disclosure Standard, launched May 2026. *Dream Machine* Issue 29.

16. *Dream Machine* Issue 28, May 2026, reporting on the Academy of Motion Picture Arts and Sciences' "You must be human to win" rule update.

17. *The Hollywood Reporter*, "Emmys Set AI Guidance." <https://www.hollywoodreporter.com/tv/tv-news/emmys-ai-guidelines-2026-awards-1236468434/>. *Dream Machine* Issue 14.

18. SAG-AFTRA negotiation timeline through *Dream Machine* Issues 7, 12, 15, 20, 26, 29.

19. *Marketing Week*, "You can't dismiss AI ads as slop when they're winning in testing." <https://www.marketingweek.com/dismiss-ai-ads-winning-creative-effectiveness/>. *Dream Machine* Issues 8, 13.

20. *Variety*, "AI Creator Behind Viral 'Deadpool,' 'Harry Potter' Christmas Clip Made His Film in a Ukrainian Bomb Shelter." *op. cit.* *Dream Machine* Issue 16.

21. PR Newswire, “From Apple TV Creative to AI Filmmaker: Hoyt Dwyer’s Animated Film To Compete at AI FilmFest Japan 2025.” *op. cit. Dream Machine* Issue 6.
22. Google DeepMind, “Dear Upstairs Neighbors.” <https://blog.google/innovation-and-ai/models-and-research/google-deepmind/dear-upstairs-neighbors/>. *Dream Machine* Issue 16.
23. *The Hollywood Reporter*, “‘Synthetic Sincerity’ by Marc Isaacs.” *op. cit. Dream Machine* Issue 8.
24. *Variety*, “‘Watch the Skies,’ Swedish UFO Feature Film Dubbed Entirely With AI, Sets USA Distribution Deal.” *op. cit. Dream Machine* Issue 5.
25. *Broadcast Pro Middle East*, “Tunisian filmmaker wins \$1 million AI Film Award for ‘Lily’.” *op. cit. Dream Machine* Issue 14.
26. Sundance Institute AI Literacy Initiative emphasis on documentation: <https://www.sundance.org/blogs/centering-the-artist-why-were-launching-the-ai-literacy-initiative/>. *Dream Machine* Issue 16.
27. Google DeepMind, “SynthID — 100 billion watermarks, expanding to partner ecosystems including OpenAI, ElevenLabs and Kakao.” <https://deepmind.google/discover/blog/synthid-100-billion-watermarks-partners/>. *Dream Machine* Issue 30.
28. *Hollywood Reporter*, “Bobby Berk Says AI Will Make Reality TV & ‘Verifiably Human Content’ More Valuable.” <https://www.hollywoodreporter.com/tv/tv-news/bobby-berk-ai-reality-tv-1236592920/>. *Dream Machine* Issue 30.
29. *Rolling Stone*, “The Rolling Stones Release New Single ‘In the Stars’ — With a Music Video De-Aging the Rockers Courtesy of AI.” <https://www.rollingstone.com/music/music-news/rolling-stones-in-the-stars-ai-de-aging-video-1235142200/>. *Hollywood Reporter*, “‘South Park’ Creators’ AI Company Made The Rolling Stones Young Again for ‘In The Stars’ Music Video.” <https://www.hollywoodreporter.com/tv/tv-news/south-park-creators-ai-rolling-stones-in-the-stars-1236592855/>. *Dream Machine* Issue 30.
30. *Variety*, “Cate Blanchett Co-Founds RSL Media, a Non-Profit to Address Consent Around AI Usage including creative work, name, image and likeness.” <https://variety.com/2026/film/news/cate-blanchett-rsl-media-ai-consent-1236748255/>. *Dream Machine* Issue 30.
31. *Bloomberg*, “Apple Acquires Key Talent & Patents Behind AI Avatar Company ‘Animato’.” <https://www.bloomberg.com/news/articles/2026-05-19/apple-acquires-animato-ai-avatar-talent-patents>. *Dream Machine* Issue 30.
- 31a. Estate of Ozzy Osbourne announces the late artist will “live on as an AI avatar that talks with fans,” June 2026. *Dream Machine* Issue 31.
32. *The Drum*, “David Beckham Designs ‘Henchester United’ Chicken Coop in Lenovo Ad.” <https://www.thedrum.com/news/2026/05/18/david-beckham-henchester-united-chicken-coop-lenovo-ai-ad>. *Dream Machine* Issue 30.

32a. YouTube announces automatic labelling of AI-generated videos, June 2026. *Dream Machine* Issue 31.

32b. LinkedIn announces it will limit the reach of AI-generated content on its feed, June 2026. *Dream Machine* Issue 31.

32c. Hasbro's Sixth Wall launch and the CharacterOS governance layer, 3 June 2026. BusinessWire: <https://www.businesswire.com/news/home/20260603297922/en/>. *The Wrap*: <https://www.thewrap.com/industry-news/business/hasbro-ai-studio-sixth-wall-launch/>. Detailed treatment in Appendix L. Cross-reference Chapter 7 §“Position Five” and Chapter 6 §“The three Cs”.

32d. Spotify–Universal Music Group landmark licensing agreements, 21 May 2026. Spotify Newsroom: <https://newsroom.spotify.com/2026-05-21/universal-music-group-spotify-licensing-agreements-fan-made-covers-remixes/>. *TechCrunch*: <https://techcrunch.com/2026/05/21/spotify-and-universal-music-strike-deal-allowing-fan-made-ai-covers-and-remixes/>. Detailed treatment in Appendix L and footnote ^37a in Chapter 6.

Chapter 13 — Coordination Collapse

1. *Dream Machine* Issue 5, “Industry Insights: Stealth, Shadow and Secret AI Users.”

2. Azumo, “AI in Workplace Statistics 2025.” <https://azumo.com/artificial-intelligence/ai-insights/ai-in-workplace-statistics>. *Tech.co*, “Gen Z Most Likely Use AI Boss.” <https://tech.co/news/gen-z-most-likely-use-ai-boss>. *Dream Machine* Issue 5.

3. *Exploding Topics*, “AI Workforce Research.” <https://explodingtopics.com/blog/ai-workforce-research>. *Dream Machine* Issue 5.

4. *Forbes*, “AI Tools Flood Workplaces as Employees Face a Double Bind.” <https://www.forbes.com/sites/carolinecastrillon/2025/09/09/ai-tools-flood-workplaces-as-employees-face-a-double-bind/>. *Dream Machine* Issue 5.

5. *Blog IDC Europe*, “Shadow AI: How Stealth Productivity Is Strangling Enterprise AI Adoption and Creating a Security Nightmare.” <https://blog-idceurope.com/shadow-ai-how-stealth-productivity-is-strangling-enterprise-ai-adoption-and-creating-a-security-nightmare/>. *Dream Machine* Issue 5.

6. *Game Developer*, “Subnautica owner Krafton outlines plans to transform into an ‘AI First’ company.” <https://www.gamedeveloper.com/business/subnautica-owner-krafton-outlines-plans-to-transform-into-an-ai-first-company>. *Dream Machine* Issue 6.

7. *Dream Machine* Issue 24, April 2026, on the GTA VI publisher laying off its internal AI team.

8. *Dream Machine* Issue 25, April 2026, on Disney layoffs including Marvel staff.

9. *SmartBrief*, “Meta to cut 10% of Reality Labs staff to focus on AI.” <https://newsletter.smartbrief.com/sharedSummary/index.jsp?briefId=40A39351-5419-4681-94DF-31A53480B698&issueId=025444D1-A590-46D8-B969->

EF81DEE05228©Id=1B5F70D2-FFDA-4660-9CE9-047C9B16BF83. *Dream Machine* Issue 14.

10. *Dream Machine* Issue 23, April 2026, on Scottish animation studio collapse.

11. *Metro*, “Prince of Persia remake and five more games cancelled as Ubisoft focuses on AI.” *op. cit.* *Dream Machine* Issue 15.

12. *The Guardian*, “AI is hitting UK harder than other big economies, study finds.” <https://www.theguardian.com/technology/2026/jan/26/ai-uk-jobs-us-japan-germany-australia>. *Dream Machine* Issue 16.

13. *The Economist*, “Investors expect AI use to soar. That’s not happening.” <https://www.economist.com/finance-and-economics/2025/11/26/investors-expect-ai-use-to-soar-thats-not-happening>. *Dream Machine* Issue 9.

14. *Dream Machine* Issue 24, April 2026, on OpenAI’s public-policy proposals around AI-driven economic disruption.

15. *The Economist*, “Job apocalypse? Humbug! AI is creating brand new occupations.” <https://www.economist.com/business/2025/12/14/job-apocalypse-humbug-ai-is-creating-brand-new-occupations>. *Dream Machine* Issue 12.

16. *Forbes*, “Vibe Coding — The In Demand AI Skill.” <https://www.forbes.com/sites/rachelwells/2025/11/06/the-in-demand-ai-skill-and-certifications-that-pays-up-to-220000/>. *Dream Machine* Issue 8.

17. U.K. Department for Business and Trade research on neurodiverse workers and AI assistants, autumn 2025. Reported via *CNBC*, “People with ADHD, autism, dyslexia say AI agents are helping them succeed at work.” <https://www.cnbc.com/2025/11/08/adhd-autism-dyslexia-jobs-careers-ai-agents-success.html>. *Dream Machine* Issue 7.

18. *CNBC*, *op. cit.*

19. *Dream Machine* Issue 7 secondary references.

20. Korin AI launch, May 2026. *Dream Machine* Issue 27.

21. *CNBC Africa*, “How AI is changing the landscape of the music industry in Africa.” <https://www.cnbc.com/2025/how-ai-is-changing-the-landscape-of-the-music-industry-in-africa>. *Dream Machine* Issue 5.

22. *BBC Future*, “Lights, camera, algorithm: Why Indian cinema is awash with AI.” <https://www.bbc.co.uk/future/article/20251223-why-indian-cinema-is-awash-with-ai>. *Dream Machine* Issue 14.

23. *Dream Machine* Issue 25, April 2026, on Indonesia’s *Legenda Bertuah*.

24. *Broadcast Pro Middle East*, “Tunisian filmmaker wins \$1 million AI Film Award for ‘Lily’.” *op. cit.* *Dream Machine* Issue 14.

25. *Digiday*, “Avocados From Mexico turns to AI to advertise around the Super Bowl instead of a TV buy.” *op. cit.* *Dream Machine* Issue 15.

26. Variety, “AI Creator Behind Viral ‘Deadpool,’ ‘Harry Potter’ Christmas Clip Made His Film in a Ukrainian Bomb Shelter.” *op. cit.* *Dream Machine* Issue 16.
27. *Dream Machine* Issue 8 citing Andreessen Horowitz observations: https://www.linkedin.com/posts/stevenouri_a-wild-stat-80-of-startups-pitching-ai6z-activity-7396182718998351872-xTKR.
28. *PocketGamer.biz*, “Shift Up CEO says AI is key to competing with China’s game industry scale.” <https://www.pocketgamer.biz/shift-up-ceo-says-ai-is-key-to-competing-with-chinas-game-industry-scale/>. *Dream Machine* Issue 14.
29. Enterprise-AI workforce tracking, late 2025. Aggregated in the Deep Dive companion piece *The Shadow AI Paradox in the Creative Industries*, drawing on Azumo’s *AI in Workplace Statistics 2025*, Tech.co’s Gen Z survey, and the IDC Europe shadow-AI security brief. *Dream Machine* Issue 5.
30. *Hidden Cloud Explosion* analysis, IDC Europe, 2025. See *The Shadow AI Paradox in the Creative Industries*, §“The Epistemology and Scale of Shadow AI.”
31. Shadow-AI security-incident statistics, 2025, aggregated in *The Shadow AI Paradox in the Creative Industries*, §“The Epistemology and Scale of Shadow AI”; underlying data via IBM *Cost of a Data Breach Report 2025* and IDC Europe.
32. For the developer-community origins of the “AI for thee, but not for me” phrasing, and the full sectoral analysis of the paradox, see *The Shadow AI Paradox in the Creative Industries*, §“The Great Hypocrisy.”
33. Survey of 1,100+ professional music creators, 2026, summarised in *Dynamics of Generative AI Adoption in the Creative Industries*, §“Music Production and Sound Recording,” and *The Shadow AI Paradox in the Creative Industries*, §“Sector-Specific Analysis.”
34. WGA screenwriter survey, pre- and post-strike, reported in *Dynamics of Generative AI Adoption in the Creative Industries*, §“Screenwriting and the Post-Strike AI Boom.”
35. Adobe Firefly milestone data, September 2023 – June 2025, in *Dynamics of Generative AI Adoption in the Creative Industries*, §“The Ubiquity of AI in Visual and Digital Arts.” *Dream Machine* Issue 6.
36. Adobe quarterly financials, FY2025–FY2026; AI-first ARR growth reported in *Dream Machine* Issue 21 and summarised in *Dynamics of Generative AI Adoption*.
37. Adobe Firefly enterprise penetration metrics, in *Dynamics of Generative AI Adoption*.
38. Adobe Stock submission analysis, 2024, in *Dynamics of Generative AI Adoption*.
39. Adobe, “Inaugural Adobe Creators’ Toolkit Report,” October 2025. <https://news.adobe.com/news/2025/10/adobe-max-2025-creators-survey>. *Dream Machine* Issue 6.
40. ChatGPT weekly-active-user disclosures, mid-2025; consolidated in *Dynamics of Generative AI Adoption*, §“General Purpose LLMs.”

41. Gemini desktop-user growth, year-over-year, in *Dynamics of Generative AI Adoption*.
42. Stanford AI Index Report 2025, global-sentiment chapter. Summarised in *Dynamics of Generative AI Adoption*, §“The Perception Gap.”
43. YouGov 2024 multi-market AI sentiment survey, 17 countries. Summarised in *Dynamics of Generative AI Adoption*, §“The Perception Gap.”
44. Quantic Foundry consumer-AI-in-gaming survey, 2025. Summarised in *Dynamics of Generative AI Adoption*, §“The Video Game Industry.”
45. Game Developers Conference *State of the Game Industry* surveys, 2024–2026, sentiment vs. usage trend. Reported in *Dynamics of Generative AI Adoption*, §“The Video Game Industry.”
46. O’Hare, J. (2026). *The Coordination Collapse: How AI Is Dissolving the Information Tax That Created Hierarchy — And What Replaces It*. DreamLab AI Consulting Ltd, April 2026. Preserved in full as Appendix K. The report draws on Dell’Acqua et al.’s preregistered Boston Consulting Group experiment with 758 consultants, Mollick’s *secret cyborgs* coinage, Edmondson and Lei’s psychological-safety findings, the IBM *Cost of a Data Breach Report 2025*, the Epoch AI / Ipsos KnowledgePanel survey of March 2026, and the CyberArk *hidden org chart* analysis.
- 46a. Epidemic Sound, *Future of the Creator Economy Report 2026* (survey of 3,000 professional creators across UK and US, June 2026); *Musically* commentary on the report and creator attitudes towards AI. *Dream Machine* Issue 31, Editor’s Pick, June 2026.

Chapter 14 — The New Jobs

1. *The Economist*, “Investors expect AI use to soar. That’s not happening.” <https://www.economist.com/finance-and-economics/2025/11/26/investors-expect-ai-use-to-soar-thats-not-happening>. *Dream Machine* Issue 9.
2. *The Economist*, “Job apocalypse? Humbug! AI is creating brand new occupations.” <https://www.economist.com/business/2025/12/14/job-apocalypse-humbug-ai-is-creating-brand-new-occupations>. *Dream Machine* Issue 12.
3. *The Guardian*, “AI is hitting UK harder than other big economies, study finds.” <https://www.theguardian.com/technology/2026/jan/26/ai-uk-jobs-us-japan-germany-australia>. *Dream Machine* Issue 16.
4. University of Wisconsin-Stout, “AI Reshaping Industry: New UW-Stout Course Sets AI-Use as Baseline Competency in Filmmaking.” <https://www.uwstout.edu/about-us/news-center/ai-reshaping-industry-new-uw-stout-course-sets-ai-use-baseline-competency-filmmaking>. *Dream Machine* Issue 15.
5. Adobe Firefly enterprise metrics, in Appendix E: Dynamics of Generative AI Adoption.
6. Reuters Institute, “AI adoption by UK journalists and their newsrooms.” <https://reutersinstitute.politics.ox.ac.uk/ai-adoption-uk-journalists-and-their-newsrooms-surveying-applications-approaches-and-attitudes>. *Digiday*, “Daily Mail says Google AI

- Overviews have killed click-throughs.” <https://digiday.com/media/daily-mail-says-google-ai-overviews-have-killed-click-throughs/>. *Dream Machine* Issues 7, 9.
7. 1,100-creator music survey 2026, in Appendix D: Shadow AI, §“Music Production and Sound Recording.”
8. VFX AI integration metrics, in Appendix E, §“Visual Effects (VFX) Automation.”
9. *PC Gamer*, “Square Enix aims to have AI doing 70% of its QA work by the end of 2027.” <https://www.pcgamer.com/gaming-industry/square-enix-aims-to-have-ai-doing-70-percent-of-its-qa-work-by-the-end-of-2027/>. *Dream Machine* Issue 7.
10. *Eurogamer*, “Falcom is the latest developer to buy into the AI hype machine.” <https://www.eurogamer.net/falcom-is-the-latest-developer-to-buy-into-the-ai-hype-machine>. *Dream Machine* Issue 12.
11. *Metro*, “Prince of Persia remake and five more games cancelled as Ubisoft focuses on AI.” <https://metro.co.uk/2026/01/21/prince-persia-remake-five-games-cancelled-ubisoft-focuses-ai-26431926/>. *Dream Machine* Issue 15.
12. *Dream Machine* Issue 24, April 2026, on the GTA VI publisher laying off its internal AI team.
13. ElevenLabs \$500m ARR reporting, April 2026. *Dream Machine* Issue 25.
14. *36kr*, “AI Video Unicorn Higgsfield: Earns \$200M in 9 Months by ‘Serving’ Social Media Marketers.” <https://eu.36kr.com/en/p/3650517574312323>. *Dream Machine* Issue 16.
15. *Dream Machine* Issue 29, May 2026, on Sony’s 49-agent / 72-skill multi-agent game-development team.
16. *Digiday*, “AI agent developers have become adland’s in-demand role.” <https://digiday.com/marketing/ai-agent-developers-have-become-adlands-in-demand-role/>. *Dream Machine* Issue 7.
17. *Forbes*, “Vibe Coding — The In Demand AI Skill That Pays Up to \$220,000.” <https://www.forbes.com/sites/rachelwells/2025/11/06/the-in-demand-ai-skill-and-certifications-that-pays-up-to-220000/>. *Dream Machine* Issue 8.
18. Sundance Institute, “Centering the Artist: Why We’re Launching the AI Literacy Initiative.” <https://www.sundance.org/blogs/centering-the-artist-why-were-launching-the-ai-literacy-initiative/>. Google blog, “Sundance Institute AI Education.” <https://blog.google/company-news/outreach-and-initiatives/google-org/sundance-institute-ai-education/>. *Dream Machine* Issues 15, 16.
19. UK Government, “Free AI training for all.” <https://www.gov.uk/government/news/free-ai-training-for-all-as-government-and-industry-programme-expands-to-provide-10-million-workers-with-key-ai-skills-by-2030>. *Dream Machine* Issue 16.
20. Lovable for classrooms. <https://lovable.dev/classroom>. *Dream Machine* Issue 11.
21. UW-Stout course launch, January 2026 — *op. cit.*

22. Adobe, “Firefly Foundry.” <https://news.adobe.com/news/2025/10/adobe-max-2025-firefly-foundry>. *Dream Machine* Issue 5.
23. Korin AI launch, May 2026. *Dream Machine* Issue 27.
24. *The Verge*, “New York’s new law forces advertisers to say when they’re using AI avatars.” <https://www.theverge.com/news/842848/new-york-law-ai-advertisements-sag-aftra-labor>. *Dream Machine* Issue 11. C2PA / SynthID infrastructure references in Chapter 12.
25. *Forbes*, “AI Is Changing How Creators Work And Earn.” <https://www.forbes.com/sites/kolawolesamueladebayo/2025/12/22/how-ai-is-changing-how-creators-work-and-earn/>. *Dream Machine* Issue 13.
26. *Variety*, “AI Creator Behind Viral ‘Deadpool,’ ‘Harry Potter’ Christmas Clip Made His Film in a Ukrainian Bomb Shelter.” <https://variety.com/2026/digital/news/ai-video-deadpool-harry-potter-andrii-daniels-1236624632/>. *Dream Machine* Issue 16.
27. *Broadcast Pro Middle East*, “Tunisian filmmaker wins \$1 million AI Film Award for ‘Lily.’” <https://www.broadcastprome.com/news/tunisian-filmmaker-wins-1-million-ai-film-award-for-lily/>. *Dream Machine* Issue 14.
28. *Dream Machine* Issue 25, April 2026, on Indonesia’s *Legenda Bertuah*.
29. BBC Future, “Lights, camera, algorithm: Why Indian cinema is awash with AI.” <https://www.bbc.co.uk/future/article/20251223-why-indian-cinema-is-awash-with-ai>. *Dream Machine* Issue 14.
30. *TechBullion*, “Why the future belongs to multi-skilled leaders.” <https://techbullion.com/playing-the-long-game-with-a-portfolio-career-why-the-future-belongs-to-multi-skilled-leaders/>. *Dream Machine* Issue 9.
31. Anthropropic Skills framework via Claude Code. *Dream Machine* Issues 11, 16, 29.
32. Adobe, “Inaugural Adobe Creators’ Toolkit Report,” October 2025. <https://news.adobe.com/news/2025/10/adobe-max-2025-creators-survey>. *Dream Machine* Issue 6.
33. PRS for Music, “PRS for Music AI Survey 2026.” <https://www.prsformusic.com/m-magazine/news/prs-for-music-ai-survey-2026>. *Dream Machine* Issue 16.
34. *CNBC*, “People with ADHD, autism, dyslexia say AI agents are helping them succeed at work.” <https://www.cnn.com/2025/11/08/adhd-autism-dyslexia-jobs-careers-ai-agents-success.html>. *Dream Machine* Issue 7.
35. McKinsey & Company, “What AI could mean for film and TV production and the industry’s future.” <https://www.mckinsey.com/industries/technology-media-and-telecommunications/our-insights/what-ai-could-mean-for-film-and-tv-production-and-the-industrys-future>. *Dream Machine* Issue 16.
36. GDC State of the Game Industry surveys 2024–2026, in Appendix E, §“The Video Game Industry.”

37. LANDR AI music study, late 2025, referenced via *Ari's Take*. <https://aristake.com/ai-tools-musicians-study/>. *Dream Machine* Issue 8.
38. Stanford AI Index Report 2025. Summarised in Appendix E, §“The Perception Gap.”
39. YouGov 2024 multi-market AI sentiment survey. Summarised in Appendix E.
40. *Digital Music News*, “Nearly 800 Creatives Sign Responsible AI Declaration — ‘Stealing Our Work Is Not Innovation’.” <https://www.digitalmusicnews.com/2026/01/22/stealing-isnt-innovation/>. *Dream Machine* Issue 16.
41. *Broadcast Pro Middle East*, Lily award — *op. cit.*
42. Variety, Andrii Daniels bomb-shelter clip — *op. cit.*
43. BBC Future, “Lights, camera, algorithm” — *op. cit.*
44. *Dream Machine* Issue 25, Indonesian *Legenda Bertuah*.
45. CNBC Africa, “How AI is changing the landscape of the music industry in Africa.” <https://www.cnbc africa.com/2025/how-ai-is-changing-the-landscape-of-the-music-industry-in-africa>. *Dream Machine* Issue 5. Korin AI launch, May 2026 — *op. cit.*
46. *PocketGamer.biz*, “Shift Up CEO says AI is key to competing with China’s game industry scale.” <https://www.pocketgamer.biz/shift-up-ceo-says-ai-is-key-to-competing-with-chinas-game-industry-scale/>. *Dream Machine* Issue 14.
47. The doomer frame is set out at length in OpenAI’s spring 2026 policy paper on the “social contract for the intelligence age” (covered in *Dream Machine* Issue 24) and in Daniel Susskind’s *A World Without Work*. The full empirical rebuttal — including the Stanford Digital Economy Lab’s *Canaries in the Coal Mine?* findings (Brynjolfsson, Chandar & Chen, 2025), Brynjolfsson, Li & Raymond’s *Generative AI at Work* (QJE, 2025; n = 5,179 agents), Acemoglu’s *Simple Macroeconomics of AI* (Economic Policy, 2025), the Dallas Fed’s labour analyses, the St. Louis Fed’s AI-and-GDP tracking, the 2026 Stanford AI Index, and the 230-year history of identical predictions — is laid out in Appendix I: The Doomer Mistake. The deep dive is the analytical antibody to the policy framing this section is pushing back against.

Chapter 15 — Choosing the Future

- 1a. *Bloomberg*, “AI Changed Chess. Grandmasters Now Win With Unpredictable Moves,” 27 March 2026. <https://www.bloomberg.com/news/articles/2026-03-27/ai-changed-chess-grandmasters-now-win-with-unpredictable-moves>. *Dream Machine* Issue 23. The behavioural pattern the piece describes — top grandmasters deliberately deviating from machine-optimal lines to put opponents on uncomputed ground — is the cleanest available analogy I have for the strategic shift the rest of this chapter argues for.
1. *Digital Music News*, “The AI Licensing Shift — Creative Weight Attribution Emerges as Music Industry Game-Changer for Rights Holders.” *op. cit.* *Dream Machine* Issue 16.
2. DreamLab AI Collective, team page. <https://dreamlab-ai.com/team>.

Chapter 16 – The Tools

1. OpenAI, “Sora 2 is here,” 30 September 2025. <https://openai.com/index/sora-2/>. *Dream Machine* Issue 1.
2. LinkedIn News aggregation: “Sora Tops 1 Million Downloads in 5 Days.” <https://www.linkedin.com/news/story/sora-tops-1m-downloads-in-5-days-6684988/>. *Dream Machine* Issue 3.
3. Google DeepMind, Veo 3.1 launch, mid-October 2025. *Dream Machine* Issue 3.
4. Runway product cycle: Gen-4.5 (December 2025), Gen-4.5 Image-to-Video (January 2026), Workflows, Story Panels, Characters API, Apps for Advertising — *Dream Machine* Issues 10, 14, 15, 16, 20.
5. Runway CEO on indie films vs. blockbusters, *Dream Machine* Issue 26.
6. Chinese open-source AI video model releases, 2025–2026. *Dream Machine* Issues 3, 12, 22.
7. *SiliconAngle*, “Higgsfield raises \$80M on \$1.3B valuation.” <https://siliconangle.com/2026/01/15/higgsfield-raises-80m-1-3b-valuation-scale-ai-video-platform/>. *36kr*, “Higgsfield: Earns \$200M in 9 Months.” <https://eu.36kr.com/en/p/3650517574312323>. *Dream Machine* Issues 15, 16.
8. Heygen Video Agent. https://www.linkedin.com/posts/heygen_introducing-the-new-video-agent-activity-7421597801240801282-d1CF. *Dream Machine* Issue 16.
9. TechCrunch, “Synthesia hits \$4B valuation, lets employees cash in.” <https://techcrunch.com/2026/01/26/synthesia-hits-4b-valuation-lets-employees-cash-in/>. *Sifted*, “Synthesia rejects \$3bn Adobe acquisition offer.” <https://sifted.eu/articles/synthesia-acquisition-offer>. *Dream Machine* Issues 5, 16.
10. ElevenLabs \$500m ARR reporting, April 2026. *Dream Machine* Issue 25.
11. Adobe Firefly milestone data, in *Dynamics of Generative AI Adoption*, §“The Ubiquity of AI in Visual and Digital Arts.”
12. Nano Banana inside Photoshop and inside Unreal Engine cross-integrations, October–November 2025. *Dream Machine* Issue 1.
13. Suno Studio launch. <https://www.techradar.com/ai-platforms-assistants/i-tried-suno-studio-the-new-platform-that-mixes-ai-music-generation-with-hands-on-editing-like-garageband-but-smarter>. *Dream Machine* Issue 1.
14. Mureka, “Music Agent Studio” launch. *Dream Machine* Issue 4.
15. ElevenLabs Series funding, April 2026. *Dream Machine* Issue 25.
16. *MusicTech*, “Cardiff band speaks out after AI artist trained on their music outperforms them on Spotify.” <https://musictech.com/news/industry/its-shocking-disheartening-and-insulting-cardiff-band-speaks-out-after-ai-artist-trained-on-their-music-outperforms-them-on-spotify/>. *Dream Machine* Issue 1.

17. Variety, “AI Creator Behind Viral ‘Deadpool,’ ‘Harry Potter’ Christmas Clip Made His Film in a Ukrainian Bomb Shelter.” <https://variety.com/2026/digital/news/ai-video-deadpool-harry-potter-andrii-daniels-1236624632/>. *Dream Machine* Issue 16.
18. Music industry AI deal flow, October 2025 – May 2026. See Chapter 5 footnotes 31–37, and *Dream Machine* Issues 5, 7, 8, 12, 14, 16, 17.
19. World Labs, “Bringing Marble to Life.” <https://www.worldlabs.ai/case-studies/bringing-marble-to-life>. *Dream Machine* Issue 7.
20. Sony Pictures Marble VP integration. *Dream Machine* Issue 8.
21. Google DeepMind, “Genie 3.” <https://deepmind.google/blog/genie-3-a-new-frontier-for-world-models/>. Project Genie: <https://blog.google/innovation-and-ai/models-and-research/google-deepmind/project-genie/>. *Dream Machine* Issues 3, 17.
22. Tencent, “HY World 1.5” and Hunyuan 3D Studio. *Dream Machine* Issue 12.
23. Luma AI, *UNI-1* launch, March 2026. *Dream Machine* Issue 22.
24. SuperSplat / Spark 2.0 / SOG releases through 2025–26. *Dream Machine* Issues 1, 25.
25. Radiance Fields, “Apple Confirms that it’s Gaussian Splatting that powers their personas.” <https://radiancefields.com/apple-confirms-personas-use-gaussian-splatting>. *Dream Machine* Issue 5.
26. ComfyUI Blog, “Ubisoft La Forge Open-Sources the CHORD Model.” <https://blog.comfy.org/p/ubisoft-open-sources-the-chord-model>. *Dream Machine* Issue 11.
27. Anthropic / Blender Foundation patronage, May 2026. *Dream Machine* Issue 27.
28. OpenAI, “Introducing AgentKit.” <https://openai.com/index/introducing-agentkit/>. *Dream Machine* Issue 2.
29. Anthropic Skills framework. *Dream Machine* Issues 11, 16, 29.
30. Heygen Video Agent. *Dream Machine* Issue 16.
31. Adobe Summit 2026 CX Enterprise. *Dream Machine* Issue 26.
32. Adobe + NVIDIA / Google + NVIDIA partnerships. *Dream Machine* Issue 21.
33. ComfyUI funding round. https://www.linkedin.com/posts/comfyui_we-raised-17-million-to-build-an-os-for-ugcPost-7373743341236236288-wkCc. *Dream Machine* Issue 1.
34. ComfyUI \$500M valuation, May 2026. *Dream Machine* Issue 27.
35. Hugging Face / Google Cloud and Meta / Hugging Face OpenEnv. *Dream Machine* Issues 5, 8.
36. Unreal Engine 5 official AI Assistant. https://www.linkedin.com/posts/wouterweynants_theres-an-official-ai-assistant-coming-to-ugcPost-7369377204226379776-pGiH. *Dream Machine* Issue 1.
37. Unity AI Council (October 2025); Unity AI Open Beta (May 2026). *Dream Machine* Issues 1, 28.

- 38.** VFX AI integration metrics. See *Dynamics of Generative AI Adoption*, §“Visual Effects (VFX) Automation.”
- 39.** Anthropic / Blender Foundation patronage. *Dream Machine* Issue 27.
- 40.** Andreessen Horowitz pitch-deck observations on Chinese open-source model usage. https://www.linkedin.com/posts/stevenouri_a-wild-stat-80-of-startups-pitching-a16z-activity-7396182718998351872-xTKR. *Dream Machine* Issue 8.
- 41.** Korin AI launch, May 2026. *Dream Machine* Issue 27.
- 42.** Google DeepMind, “Introducing Gemini Omni: Create Anything from Any input.” <https://blog.google/technology/google-deepmind/gemini-omni-launch/>. *Dream Machine* Issue 30.
- 43.** Beeple Canvas — Generative AI compositor. <https://www.beeple-canvas.com/>. *Dream Machine* Issue 30.
- 44.** Sony AI, “Woosh — a sound effect foundation model.” <https://ai.sony/blog/woosh-sound-effect-foundation-model/>. *Dream Machine* Issue 30.
- 45.** Mirelo SFX 1.6, “edit sound, not just generate it.” <https://mirelo.ai/sfx-1-6>. *Dream Machine* Issue 30.
- 46.** Stability AI, “Stable Audio 3.0 released — open-weight model family built for artistic experimentation.” <https://stability.ai/news/stable-audio-3-0-released>. *Dream Machine* Issue 30.
- 47.** Tamber product page: <https://tamber.ai/>. *Dream Machine* Issue 30.
- 48.** Beatport Track ID. <https://www.beatport.com/track-id>. *Dream Machine* Issue 30.
- 49.** NVIDIA SANA-WM model collection. <https://huggingface.co/collections/nvidia/sana-wm>. *Dream Machine* Issue 30.
- 50.** Odyssey, “Introducing Starchild-1, the first real-time multimodal world model.” <https://odyssey.ml/introducing-starchild-1>. *Dream Machine* Issue 30.
- 51.** Odyssey, “Introducing Agora-1 — four-player AI-generated world built on a 1997 shooter.” <https://odyssey.ml/introducing-agora-1>. *Dream Machine* Issue 30.
- 52.** Apple Machine Learning Research, “Apple Headsup: a Large-Scale High-Quality 3D Gaussian Head Reconstruction from Multi-View Captures.” <https://machinelearning.apple.com/research/apple-headsup-3d-gaussian-head>. *Dream Machine* Issue 30.
- 53.** Google, “Official skills for AI agents.” <https://github.com/google/agent-skills>. *Dream Machine* Issue 30.
- 54.** Tencent Ardot, AI-native design agent platform. <https://ardot.tencent.com/>. *Dream Machine* Issue 30.
- 55.** Anthropic, “Claude is now available as a partner node in ComfyUI.” <https://www.anthropic.com/news/claude-comfyui-partner-node>. *Dream Machine* Issue 30.

56. ECABridge — Unreal Engine MCP integration. <https://ecabridge.dev/>. *Dream Machine* Issue 30.

57. *Video Games Chronicle*, “Epic Games Veteran Claims He’s Building AI-Heavy ‘Fully European’ Game Engine.” <https://www.videogameschronicle.com/news/epic-games-veteran-ai-heavy-fully-european-game-engine/>. *Dream Machine* Issue 30.

58. PhotoGIMP — the open-source GIMP skin that mimics Photoshop. <https://github.com/Diollinux/PhotoGIMP>. *Dream Machine* Issue 30.

* * *

Summary

- **Total footnotes:** 563
- **Unique primary-source URLs cited:** 357
- **Source corpus:** 32 editions of *Dream Machine* | *Creative AI* (October 2025 – June 2026)

About the author, and about DreamLab

Pete Woodbridge is a **Creative Technologist and Innovation Director** working across AI, immersive, interactive and experiential tech, real-time production, and next-generation storytelling across **film, games, interactive and immersive experiences**. His cross-sector background spans creative, technical, innovation and product-development roles for software companies, games and immersive studios, and working on new possibilities in storytelling covering film, advertising, brand experiences and animation.

He has worked on projects and new products with partners including **Intel, Google, McLaren, ITV, Sony, Aardman, Red Bull, IBM, BBC, NBC, Crytek, Apple, MediaCityUK, Saatchi & Saatchi** and **Universal** — alongside dozens of independent studios, agencies and SMEs building new products, content pipelines, and interactive and immersive experiences.

Technically, he works across strategy, leadership and hands-on development, including: **real-time interactive and immersive systems** (Unreal Engine, Unity, WebXR, Disguise, Notch, Ventuz); **AI, software and workflow engineering** (Python, ComfyUI, API integrations, automation systems, AI-assisted pipeline design); and **production and future-media** (virtual production, motion capture, ICVFX, real-time previs, AI-augmented workflows). He has built and led teams of engineers and artists across most of those categories at one time or another — which is, more than anything else, the experience the book's arguments about orchestrator-shaped careers and agentic-mesh studios are made out of.

In parallel to running DreamLab he has developed innovation and investment programmes for the creative industries in the North West UK, including the **£3.2 million Innovation Accelerator (Innovate UK) at MediaCity** in Manchester, and the **£7.2 million MusicFutures programme (UKRI/AHRC)** in Liverpool. He is also an external assessor and advisor on several national and international innovation programmes and strategies, including **CoSTAR, EPSRC** and **Digital Catapult**, alongside contributions to a number of steering committees, boards, startups and investment initiatives.

He writes the *Dream Machine* | *Creative AI* newsletter, which began in October 2025 with the launch of Sora 2 and has run weekly through every major event covered in this book. The newsletter is the underlying source corpus of the book; each chapter footnotes back to the specific issue in which a given event was first analysed. He is also the co-host of the *Creative AI* podcast, an active contributor to UK creative-industry policy conversations on AI and copyright, and a frequent speaker on creative technology, AI orchestration, world models, the *Human–AI Agency Continuum* framework, the *Slop Ceiling*, the *Authenticity Premium* and the *Orchestrator economy* — the analytical scaffolding the book is built on.

He lives and works in the North West of England, and is, by his own account, more comfortable writing than talking, and more comfortable shipping than either.

About DreamLab

DreamLab is the collective behind the *Dream Machine* newsletter, the *Creative AI* podcast and this book. The studio was founded in **January 2024**, starting at **MediaCity in**

Manchester, and has since grown into a network of **around 50 deep-tech specialists** across labs in **Manchester, Liverpool, Cumbria and London** — and is continuing to spread into new locations. The team is **AI engineers, XR developers, game designers, animators, sound artists, researchers, educators and filmmakers**, including **Emmy winners, PhD researchers and BAFTA-recognised practitioners** working at the cutting edge of creative technology.

The studio does not operate as a fixed agency. It assembles the right team of specialists around each challenge — a model that has, in the period covered by this book, become both the operational pattern of the orchestrator economy described in [Chapter 11](#) and the studio's own default way of working. The four service lines are *Workflows & Automation*, *Creative Services*, *Product Development* and *Consultancy & Training*, deployed across film, TV, games, music, broadcast, immersive and advertising.

In the two and a half years since the studio was founded, DreamLab has supported **100+ businesses**, delivered **35+ collaborative projects**, and helped channel approximately **£2.7m of investment** into partner organisations. Its clients and collaborators include the partner list above — Intel, Google, McLaren, ITV, Sony, Aardman, Red Bull, IBM, BBC, NBC, Crytek, Apple, MediaCityUK, Saatchi & Saatchi, Universal — and span:

- **Global and established brands** — including work for **Verizon Trailblazers** and partnerships with major broadcasters, streamers and platform companies looking for the kind of grounded R&D the studio specialises in.
- **SME partners** — the next wave of ambitious independent studios and small-to-medium creative-technology firms building products with applied AI and creative technology underneath them.
- **Universities, public sector, combined authorities, research councils and innovation agencies** — including funded programmes with **UKRI, Innovate UK, EPSRC, CoSTAR and Digital Catapult**, and a long-running innovation partnership at **MediaCity Hub / MediaCityUK** in Salford, where the studio was founded.

DreamLab has been recognised in the **Verizon Trailblazers 500 UK Trailblazers** report, is a participant in the **UKRI-funded Innovate UK** immersive-tech accelerator, and is a **MediaCity Hub innovation partner** — three structural indicators that the studio sits inside the public-private R&D infrastructure of the UK creative-tech sector rather than on its periphery.

In closed-beta access to platforms like World Labs' **Marble**, in production deployment of agentic stacks built on Claude Code, Skills and the wider open-source ecosystem, and in continuous public writing through the *Dream Machine* newsletter, DreamLab functions as a *working studio that ships* and a *public-record practice* simultaneously — the operating model the book argues will, on the available evidence, define the most defensible shape of creative work for the next decade.

You can find DreamLab at dreamlab.org.uk, the *Dream Machine* newsletter on [LinkedIn](#) (and as a podcast on [Spotify](#)), and the team page at dreamlab.org.uk/team.

A living book, and the knowledge base behind it

I want to put on the page, at the back of the book where it belongs, the load-bearing operational detail about how this book is built and re-built, because it is — quietly — part of the argument the rest of the manuscript has been making.

Dream Machine is a **living book**. It is not a traditional non-fiction hardback that ships once and then sits frozen for five years while the field it describes moves on. The print and digital editions are updated and re-published on a **rolling cadence — currently roughly monthly** — and will continue to be for as long as the *Dream Machine* newsletter ships and the field keeps changing at the rate it currently changes. The unversioned `Dream_Machine.pdf` in the published archive is always the most recent edition. The dated editions — `Dream_Machine_2026-05-21.pdf`, `Dream_Machine_2026-06-04.pdf`, and so on through subsequent months — are preserved alongside as snapshots of the manuscript at specific moments in the longer sequence. A reader who wants the latest argument takes the unversioned file. A reader who wants the specific version a particular claim first appeared in takes the dated one. Both are kept indefinitely; nothing is ever silently overwritten.

Behind the manuscript there is a continuously-rebuilt **knowledge base** inside DreamLab. The corpus underneath the book is structured: each weekly issue of the newsletter is parsed into machine-readable markdown with the original primary-source URLs preserved; each long-form deep dive is checked in as a separately-tracked document; each chapter's footnotes are linked to the issue or document where the claim first appeared; the toolchain inventory of [Chapter 16](#) is regenerated from a categorised data set so that new tools surface in the catalogue automatically when the newsletter covers them; the citation index is auto-built from every chapter's footnotes on every rebuild. The pipeline that re-assembles the book takes the corpus state at the moment of build, runs it through the layout system, produces the dated edition PDF, mirrors the result into the unversioned `Dream_Machine.pdf` alias, and ships both. The whole cycle takes minutes rather than days. The operational implication is that the next time the newsletter ships an issue, the next time a deep-dive PDF lands in the studio, the next time a chapter's footnote needs to be tightened against a fresh source, the book updates with it — automatically, repeatably, and without requiring a fresh print contract or a new edition number.

The cadence is deliberate, not as a marketing posture but as a quiet demonstration of the book's own argument. The position the manuscript settles on — that AI enables new formats, new publishing models, and new ways of holding a body of knowledge alive in public — is one I think a book about AI in the creative industries has an obligation to *show*, not only to *say*. The traditional book-publishing cycle was an artefact of the industrial-print economy: typesetting cost what it cost, distribution required physical inventory, the unit economics of a hardback were balanced over a five-year shelf life, and so the book froze. The AI-augmented publishing pipeline I have used to build *Dream Machine* makes a different set of trade-offs available. The book can update. The corpus can grow. The argument can sharpen as the field's evidence sharpens. The dated editions preserve the record. The unversioned file ships the present. The trade-offs of the old model are not, in 2026, the only trade-offs available, and the trade-offs of the new model are — for a working-practitioner's record of a moving field — by some margin the better ones.

The point, in a sentence: **a living book is the right shape of a book about a living transition**. That this book has chosen that shape is part of how it makes its argument.

The full corpus — every newsletter issue parsed to markdown, every deep-dive essay, the build pipeline and the chapter source files — is open and available at github.com/petejwoodbridge/DreamMachine-Book. Readers who want to explore the knowledge base, trace the provenance of a specific argument, access earlier editions, or follow the book as it updates will find everything there.

A note on the book's authorship

The arguments in this book belong to the author, but they have been shaped, sharpened and in many places explicitly co-authored with the DreamLab team and the wider community of working creatives who read the newsletter, write back, push back and contribute their own evidence. The Dedication at the front of the book names this debt directly. The organisational-design spine of Chapter 13 and the long-form analysis preserved as Appendix K are the work of **Dr. John O'Hare** of DreamLab AI Consulting, drawing on Dell'Acqua's Boston Consulting Group experiment, Mollick's *secret cyborgs* coinage, Edmondson and Lei's psychological-safety research and the operational experience of the studio's own agentic-mesh deployments. Several other sections — most notably the *Three Cs* framework in Chapter 6, the *Position Five* analysis in Chapter 7, the *governance-layer-as-IP-environment* thesis in Chapter 12 and the *Programmable Brand* deep dive preserved as Appendix L — are the product of long studio-internal arguments that the book has, in places, simply written down.

The book is, in that sense, an artefact of the practice it describes. It is the field guide a working studio produces when it tries, in public, to make sense of the most consequential year in creative work since cinema learned to talk.

— Pete Woodbridge — *DreamLab, the North West of England* — June 2026

-
1. Variety, "SAG-AFTRA Condemns Tilly Norwood: AI Actress Is Not an Actor," 30 September 2025. <https://variety.com/2025/film/news/sag-aftra-tilly-norwood-ai-actress-1236534779/>. See also NBC News, "Tilly Norwood, fully AI 'actor,' blasted by actors union SAG-AFTRA for 'devaluing human artistry.'" <https://www.nbcnews.com/pop-culture/pop-culture-news/tilly-norwood-fully-ai-actor-blasted-actors-union-sag-aftra-devaluing-rcna234685>. Discussed in *Dream Machine* Issue 1 (6 October 2025).↵
 2. *The Hollywood Reporter*, "U.K. Union Equity Condemns Tilly Norwood: 'AI Tool, Not a Performer'." <https://www.hollywoodreporter.com/movies/movie-news/tilly-norwood-ai-actress-uk-union-equity-sag-aftra-debate-1236391739/>. See also Variety, "Tilly Norwood Slammed by Equity as AI Tool, Concerned About Origin." <https://variety.com/2025/film/global/tilly-norwood-slammed-equity-ai-tool-concerned-origin-1236537042/>. *Dream Machine* Issue 1.↵
 3. CNN, "Tilly Norwood: Hollywood is fuming over a new 'AI actress'," 30 September 2025. <https://www.cnn.com/2025/09/30/tech/hollywood-ai-actor-backlash>.↵
 4. OpenAI, "Sora 2 is here," announcement page, 30 September 2025. <https://openai.com/index/sora-2/>. The model launched alongside an invite-only iOS app of the same name in the U.S. and Canada. *Dream Machine* Issue 1 carried the launch alongside contemporaneous coverage from NBC News and *The Guardian* on the model's first copyright and safety incidents.↵
 5. *Dream Machine* | *Creative AI*, LinkedIn newsletter, archive of Issues 1–32, October 2025 – June 2026. <https://www.linkedin.com/newsletters/dream-machine-creative-ai-7379776527871381505/>.↵
 6. DreamLab AI Collective, team page. <https://dreamlab-ai.com/team>. Referenced from *Dream Machine* Issue 16 onward.↵

7. Charles Cecil (Revolution Software, *Broken Sword*) quoted in *gamesindustry.biz*, “AI was an expensive mistake’: Charles Cecil on innovation, insolvency, and Broken Sword.” <https://www.gamesindustry.biz/ai-was-an-expensive-mistake-charles-cecil-on-innovation-insolvency-and-broken-sword>. *Dream Machine* Issue 3. ↵
8. Adobe, “Inaugural Adobe Creators’ Toolkit Report: 86 Percent of Global Creators Use Creative Generative AI.” <https://news.adobe.com/news/2025/10/adobe-max-2025-creators-survey>. Survey of 16,000 creators across the U.S., U.K., France, Germany, South Korea, Japan, India and Australia, released at Adobe MAX 2025. *Dream Machine* Issue 6. ↵
9. UK Department for Science, Innovation and Technology (DSIT), *Statement of Progress on Copyright and AI*, December 2025. <https://www.gov.uk/government/publications/copyright-and-artificial-intelligence-progress-report/copyright-and-artificial-intelligence-statement-of-progress-under-section-137-data-use-and-access-act>. See also IPWatchdog, “Respondents to UK AI Consultation Overwhelmingly Want AI Companies to License Copyrighted Works in All Cases.” <https://ipwatchdog.com/2025/12/16/respondents-uk-ai-consultation-overwhelmingly-want-ai-companies-license-copyrighted-works-all-cases/>. *Dream Machine* Issue 12. ↵
10. *Dream Machine* Issue 5, “Adobe’s Latest AI Announcements — Is every tool going AI?”, 31 October 2025. <https://www.linkedin.com/pulse/dream-machine-creative-ai-news-insight-oct-25-issue-5-woodbridge-f7jnc/>. ↵
11. Adobe, Adobe MAX 2025 keynote messaging, October 2025. Coverage: Creative Boom, “Adobe is putting AI in everything everywhere all at once.” <https://www.creativeboom.com/news/adobe-is-putting-ai-in-everything-everywhere-all-at-once/>. *Dream Machine* Issue 5. ↵
12. World Labs, *Marble* — first commercial spatial-AI world model, public launch November 2025. <https://marble.worldlabs.ai/>. Technical context: TechCrunch, “Fei-Fei Li’s World Labs speeds up the world model race with Marble, its first commercial product.” <https://techcrunch.com/2025/11/12/fei-fei-lis-world-labs-speeds-up-the-world-model-race-with-marble-its-first-commercial-product/>. DreamLab participated in the closed beta during October–November 2025. *Dream Machine* Issue 7. ↵
13. 11,514 responses across the Citizen Space portal and email, of which 10,112 came through Citizen Space; 88% of those supported licensing as a default rule, against 3% who supported the government’s preferred opt-out model. UK DSIT, *Statement of Progress*, December 2025; analysis in *Dream Machine* Issue 12 (18 December 2025). Final report and economic impact assessment to be laid before Parliament by 18 March 2026. ↵
14. *Digital Music News*, “Nearly 800 Creatives, Including Jason Aldean and One Republic, Sign Responsible AI Declaration — ‘Stealing Our Work Is Not Innovation’.” <https://www.digitalmusicnews.com/2026/01/22/stealing-isnt-innovation/>. *Dream Machine* Issue 16. ↵
15. For a contemporaneous overview of the AI video model release cadence through 2024 and 2025, see *Dream Machine* Issues 1–8 (October–November 2025), which logged near-weekly releases from Runway, Luma, Pika, Kling, Veo, Wan, Higgsfield, Hunyuan and a long tail of smaller labs. ↵
16. *The Hollywood Reporter*, “AI Performer Tilly Norwood Sparks Hollywood Backlash.” <https://www.hollywoodreporter.com/movies/movie-news/tilly-norwood-ai-actress-uk-union-equity-sag-aftra-debate-1236391739/>. *Dream Machine* Issue 1. ↵
17. SAG-AFTRA statement, 30 September 2025, reported in Variety, “SAG-AFTRA Condemns Tilly Norwood: AI Actress Is Not an Actor.” <https://variety.com/2025/film/news/sag-aftra-tilly-norwood-ai-actress-1236534779/>. ↵
18. OpenAI, “Sora 2 is here,” 30 September 2025. <https://openai.com/index/sora-2/>. *Dream Machine* Issue 1. ↵
19. Particle6 background and Van der Velden interview: *The Hollywood Reporter*, “Meet the Creator of the AI Actress Hollywood Loves to Hate: ‘You’re Gonna See a Lot of Tilly Norwood Next Year.’” <https://www.hollywoodreporter.com/movies/movie-features/tilly-norwood-creator-particle6-elina-van-der-velden-talks-1236428824/>. *Dream Machine* Issue 8. ↵
20. *Deadline*, “Tilly Norwood Creator Eline Van Der Velden Talks Backlash, Reveals Another 40 AI Actors Are In The Pipeline.” <https://deadline.com/2025/11/tilly-norwood-creator-interview-backlash-more-ai-actors-coming-1236601334/>. ↵
21. Northeastern Global News, “Why AI ‘Actress’ Tilly Norwood Has Hollywood Angry.” <https://news.northeastern.edu/2025/10/02/ai-actress-tilly-norwood-hollywood-backlash/>. ↵

22. SAG-AFTRA, official statement reproduced in *Variety*, *op. cit.*; also NBC News, “Tilly Norwood, fully AI ‘actor,’ blasted by actors union SAG-AFTRA for ‘devaluing human artistry’.” <https://www.nbcnews.com/pop-culture/pop-culture-news/tilly-norwood-fully-ai-actor-blasted-actors-union-sag-aftra-devaluing-rcna234685>.↵
23. Equity (U.K.), statement of 2 October 2025: *Variety*, “Tilly Norwood Slammed by Equity as AI Tool, Concerned About Origin.” <https://variety.com/2025/film/global/tilly-norwood-slammed-equity-ai-tool-concerned-origin-1236537042/>.↵
24. CNN, “Tilly Norwood: Hollywood is fuming over a new ‘AI actress’.” <https://www.cnn.com/2025/09/30/tech/hollywood-ai-actor-backlash>.↵
25. OpenAI, “Sora 2 is here,” <https://openai.com/index/sora-2/>. Technical capabilities summary including physics modelling, multi-shot world-state persistence and synchronised audio.↵
26. *Dream Machine* Issue 1, “Editor’s Pick”; further launch context in NBC News, “OpenAI’s Sora 2: a major leap in AI video and audio.” <https://www.nbcnews.com/tech/tech-news/openai-sora-2-app-video-chatgpt-creation-rcna234973>.↵
27. LinkedIn News aggregation: “Sora Tops 1 Million Downloads in 5 Days.” <https://www.linkedin.com/news/story/sora-tops-1m-downloads-in-5-days-6684988/>. *Dream Machine* Issue 3. A historical note for the reader: the Sora consumer app and the broader Sora platform did not survive 2026. OpenAI announced a staged wind-down of Sora on **24 March 2026** — consumer app and website ending late April, developer API closing late September — with peak revenue reportedly at around \$540,000 per month against operating costs in the billions, and the compute reallocated toward higher-margin coding and enterprise products. The Disney–OpenAI licensing announcement of 11 December 2025 (referenced in Chapter 7 and Appendix C) was never executed; no formal agreement was signed and no money changed hands. The contractual framework underneath that deal — IP-not-training, talent-likeness excluded, jointly-governed steering committee, equity-and-warrants structure — survived the platform’s collapse and is now the reference architecture other studios negotiate against. The full deal-architecture analysis is in *Appendix L: The Programmable Brand*. The argument the rest of this chapter is making about the Sora 2 launch week — the cultural rupture, the audience arrival in the AI debate, the parasocial-character / working-actor market collision — survives the platform’s death intact; the *moment* matters whether or not the *vendor* persists.↵
28. *The Guardian*, “OpenAI Sora 2 violence racism.” <https://www.theguardian.com/us-news/2025/oct/04/openai-sora-violence-racism>. *Dream Machine* Issue 1.↵
29. NBC News, *op. cit.*; *The Guardian*, *op. cit.*↵
30. *Digital Music News*, “OpenAI’s Sora 2 includes likeness protections for celebrities who don’t opt in, but that doesn’t apply to ‘historical figures’ and dead celebrities.” <https://www.digitalmusicnews.com/2025/10/08/openais-likeness-protections-dont-apply-to-dead-celebrities/>. *Dream Machine* Issue 2.↵
31. Quoted in *The Guardian*, “OpenAI launch of video app Sora plagued by violent and racist images: ‘The guardrails are not real’.” <https://www.theguardian.com/us-news/2025/oct/04/openai-sora-violence-racism>. *Dream Machine* Issue 1.↵
32. Google DeepMind, Veo 3.1 launch, mid-October 2025. *Dream Machine* Issue 3, “Editor’s Pick: Veo 3.1 and the Rise of AI Filmmaking.” Coverage: <https://www.cometapi.com/veo-3-1-is-coming-and-whats-rumor/>.↵
33. WUFT, “Kiss reality goodbye: AI-generated social media has arrived,” 3 October 2025. <https://www.wuft.org/2025-10-03/kiss-reality-goodbye-ai-generated-social-media-has-arrived>. *Dream Machine* Issue 1.↵
34. *No Film School*, “James Cameron Says AI Is ‘Never Going to Take the Place’ of Humans.” <https://nofilmschool.com/james-cameron-ai#>. *Dream Machine* Issue 1.↵
35. *The Guardian*, “James Cameron says AI actors are ‘horrifying to me’,” 1 December 2025. <https://www.theguardian.com/film/2025/dec/01/james-cameron-says-ai-actors-are-horrifying-to-me>. Original quote from CBS *Sunday Morning*. *Dream Machine* Issue 10.↵
36. *Variety*, “James Cameron Says It’s ‘Horrifying’ that AI Can ‘Make Up an Actor’.” <https://variety.com/2025/film/news/james-cameron-horrifying-ai-replace-actors-1236595864/>.↵

37. Stability AI, board composition, 2024–2026. Reported across multiple outlets including *Deadline*, “James Cameron Calls AI Replacing Actors ‘Horrific’; Art ‘Sacred.’” <https://deadline.com/2025/11/james-cameron-gen-ai-horrifying-human-art-sacred-avatar-1236631387/>. ↵
38. Deezer, “AI-generated tracks now represent 44% of all new uploaded music,” April 2026 newsroom release. <https://newsroom-deezer.com/2026/04/ai-generated-tracks-represent-44-of-new-uploaded-music/>. Companion analysis: *Music Business Worldwide*, “75,000 AI-generated tracks now flood Deezer daily, representing 44% of all new music uploaded to the platform.” <https://www.musicbusinessworldwide.com/75000-ai-generated-tracks-now-flood-deezer-daily-representing-44-of-all-new-music-uploaded-to-the-platform-says-streamer/>. Daily AI uploads to Deezer rose from approximately 50,000 per day in November 2025 (*Dream Machine* Issue 7, citing Deezer / *Musically*) to 75,000 per day by April 2026, with consumer streams of fully-AI tracks holding between 1% and 3% of total platform plays — and up to 85% of those streams identified as fraudulent in 2025. *Dream Machine* Issues 7, 26, 27, 28. ↵
39. John Philip Sousa, “The Menace of Mechanical Music,” *Appleton’s Magazine*, Vol. 8, September 1906, pp. 278–284. Full text via ExplorePAHistory: <https://explorepahistory.com/odocument.php?docId=1-4-1A1>. Academic context: Patrick Warfield, “John Philip Sousa and ‘The Menace of Mechanical Music,’” *Journal of the Society for American Music*, Cambridge University Press: <https://www.cambridge.org/core/journals/journal-of-the-society-for-american-music/article/abs/john-philip-sousa-and-the-menace-of-mechanical-music/A9E621587BE7580ABD73AEF64D4B2DC8>. The 1906 essay was, in part, lobbying for what would become the 1909 Copyright Act. ↵
40. Sousa, *op. cit.* The Library of Congress’s “*Sousa and the Talking Machine*” essay is a useful institutional summary: <https://blogs.loc.gov/now-see-hear/2020/05/sousa-and-the-talking-machine/>. ↵
41. William Henry Cardinal O’Connell, Archbishop of Boston, sermon to the Holy Name Society, Boston, 10 January 1932. Reported widely in the contemporaneous press, including the *Daily Courier* (Connellsville, PA), 12 January 1932 (<https://www.newspapers.com/newspage/38168082/>). Cultural context: KUOW/NPR, “Imbecile Slush’: Surprising Early Reactions to Crooning,” <https://www.kuow.org/stories/imbecile-slush-surprising-early-reactions-crooning>. JSTOR Daily, “The Gender Politics of the First Boy Bands,” <https://daily.jstor.org/the-gender-politics-of-the-first-boy-bands/>. ↵
42. *Grand Upright Music, Ltd. v. Warner Bros. Records Inc.*, 780 F. Supp. 182 (S.D.N.Y. 1991). Full text: <https://law.justia.com/cases/federal/district-courts/FSupp/780/182/1445286/>. The “Thou shalt not steal” opening is the most-quoted line from a US copyright opinion of the late twentieth century. ↵
43. Tippet’s account of the Jurassic Park digital test is documented across multiple ASC and contemporaneous press accounts. American Society of Cinematographers, “*Jurassic Park: Effects Team Brings Dinosaurs Back from Extinction*,” <https://theasc.com/articles/jurassic-park-effects-team-brings-dinosaurs-back>. Wikipedia, “Phil Tippet,” https://en.wikipedia.org/wiki/Phil_Tippet. The dialogue paraphrase Spielberg incorporated into the film is Goldblum/Malcolm’s response to Grant’s “I think we’re out of a job”: “Don’t you mean extinct?” ↵
44. Charles Baudelaire, “*Le Public Moderne et la Photographie*,” *Revue Française*, 1859 (part of the *Salon de 1859* essays). English translation widely available; the original French in PDF form: https://gallowayexploringart.wordpress.com/wp-content/uploads/2014/08/ baudelaire_the-modern-public-photography.pdf. Smithsonian Archives institutional overview: “*Photography Murdered Painting, Right?*,” <https://siarchives.si.edu/blog/photography-murdered-painting-right>. ↵
45. The Delaroche apocrypha is documented in Quote Investigator: <https://quoteinvestigator.com/2022/10/16/photo-mortal/>. The earliest sourced version is in an 1873 survey, 34 years after Delaroche reportedly said it. Delaroche’s own contemporary writing on the daguerreotype, in Gernsheim’s standard 1959 monograph, characterised the new technology as “an immense service to the arts.” ↵
46. The 1942–44 Petrillo strike: Wikipedia, “1942–44 musicians’ strike,” https://en.wikipedia.org/wiki/1942%E2%80%9344_musicians%27_strike; Mainspring Press, “*The Man Who Crippled the American Recording Industry: James Caesar Petrillo and the American Federation of Musicians Recording Bans*,” <https://mainspringpress.org/2024/11/23/the-man-who-crippled-the-recording-industry-james-caesar-petrillo-and-the-americans-federation-of-musicians-recording-bans/>; DownBeat, “*The Petrillo Ban of 1942–’44: Past & Future at War*,” <https://downbeat.com/news/detail/the-petrillo-ban-of-194244-past-future-at-war>; Local 802 AFM, “*The Silence Was Deafening*,”

<https://www.local802afm.org/allegro/articles/the-silence-was-deafening/>. The Music Performance Trust Fund's institutional history: <https://musicpf.org/establishment-of-mpft-led-to-the-formation-of-afms-pension-and-residual-funds/>.↵

47. William Henry Cardinal O'Connell, Archbishop of Boston, sermon to the Holy Name Society, Boston, 10 January 1932. Reported widely in the contemporaneous press, including the *Daily Courier* (Connellsville, PA), 12 January 1932 (<https://www.newspapers.com/newspage/38168082/>). Cultural context: KUOW/NPR, "Imbecile Slush: Surprising Early Reactions to Crooning," <https://www.kuow.org/stories/imbecile-slush-surprising-early-reactions-crooning>. JSTOR Daily, "The Gender Politics of the First Boy Bands," <https://daily.jstor.org/the-gender-politics-of-the-first-boy-bands/>.↵
48. Musicians' Union History, "The Strike That Made History — Massacre of the Musicians 1980," <https://www.muhistory.com/the-massacre-of-the-musicians-1980/>. Academic context on the broader MU–BBC dispute landscape: "Negotiating Needletime" (Tandfonline), <https://www.tandfonline.com/doi/full/10.1080/03071022.2016.1215098>.↵
49. MusicRadar, "The Day the Loony Musicians Union Tried to Kill the Synthesizer (Which Also Happened to be Bob Moog's Birthday)," <https://www.musicradar.com/news/the-union-passed-a-motion-to-ban-the-use-of-synths-drum-machines-and-any-electronic-devices-the-day-the-loony-musicians-union-tried-to-kill-the-synthesizer-which-also-happened-to-be-bob-moogs-birthday>. Far Out Magazine, "Why did the Musicians Union outlaw synthesizers in 1982?", <https://faroutmagazine.co.uk/musicians-union-outlaw-synthesizers/>.↵
50. *Grand Upright Music, Ltd. v. Warner Bros. Records Inc.*, 780 F. Supp. 182 (S.D.N.Y. 1991). Full text: <https://law.justia.com/cases/federal/district-courts/FSupp/780/182/1445286/>. The "Thou shalt not steal" opening is the most-quoted line from a US copyright opinion of the late twentieth century.↵
51. *Bridgeport Music, Inc. v. Dimension Films*, 410 F.3d 792 (6th Cir. 2005). Full text: <https://law.justia.com/cases/federal/appellate-courts/F3/410/792/574458/>. The "Get a licence or do not sample" rule is the most-cited line in the opinion.↵
52. TIME, "50 Worst Inventions," 2010, Auto-Tune at #15: https://content.time.com/time/specials/packages/article/0,28804,1991915_1991909_1991903,00.html. Wikipedia, "Auto-Tune," <https://en.wikipedia.org/wiki/Auto-Tune>. NPR, "25 Years of Believe," <https://www.npr.org/2023/10/19/1207028349/25-years-ago-cher-released-a-song-that-would-change-the-sound-of-pop-music>. Wikipedia, "D.O.A. (Death of Auto-Tune)," [https://en.wikipedia.org/wiki/D.O.A. \(Death of Auto-Tune\)](https://en.wikipedia.org/wiki/D.O.A._(Death_of_Auto-Tune)).↵
53. Walter Murch, *In the Blink of an Eye: A Perspective on Film Editing*, Silman-James Press, 1995 (2nd edition 2001). PDF: [https://www.craftfilmschool.com/userfiles/files/Walter%20Murch%20%20In%20the%20Blink%20of%20an%20Eye%20Revised%202nd%20Edition%20of%2001,%20Silman-James%20Pr\).pdf](https://www.craftfilmschool.com/userfiles/files/Walter%20Murch%20%20In%20the%20Blink%20of%20an%20Eye%20Revised%202nd%20Edition%20of%2001,%20Silman-James%20Pr).pdf). Charles Koppelman, *Behind the Seen: How Walter Murch Edited Cold Mountain Using Apple's Final Cut Pro and What This Means for Cinema*, Peachpit Press, 2004: <https://www.peachpit.com/store/behind-the-seen-how-walter-murch-edited-cold-mountain-9780735714267>.↵
54. Sasson's account documented at the National Inventors Hall of Fame: <https://www.invent.org/blog/inventors/Legacy-Steve-Sasson>. Snopes verification of the "Kodak suppressed the digital camera" claim: <https://www.snopes.com/fact-check/kodak-digital-camera-invention/>. Knowledge@Wharton on the Kodak collapse: <https://knowledge.wharton.upenn.edu/podcast/knowledge-at-wharton-podcast/whats-wrong-with-this-picture-kodaks-30-year-slide-into-bankruptcy/>. Bankruptcy filing: 19 January 2012, S.D.N.Y., \$5.1bn assets / \$6.8bn liabilities.↵
55. Wikipedia, "Brian Walski," https://en.wikipedia.org/wiki/Brian_Walski. *Washington Post* contemporaneous coverage: <https://www.washingtonpost.com/archive/lifestyle/2003/04/03/alters-picture-costs-la-times-photographer-his-job/c5e7c9e0-a836-429a-bb4e-d502f1768a96/>. World Press Photo's institutional response in TIME: <https://time.com/3706626/world-press-photo-processing-manipulation-disqualified/>.↵
56. Wikipedia, "Viacom International, Inc. v. YouTube, Inc.," https://en.wikipedia.org/wiki/Viacom_International_Inc._v._YouTube,_Inc.. Electronic Frontier Foundation case file: <https://www EFF.org/cases/viacom-v-youtube>. Variety on the March 2014 settlement:

<https://variety.com/2014/biz/news/google-and-viacom-settle-copyright-infringement-lawsuit-over-youtube-1201137538/>. ↵

57. PetaPixel, “The Rise and Crash of the Camera Industry in One Chart,” <https://petapixel.com/2024/08/22/the-rise-and-crash-of-the-camera-industry-in-one-chart/>. Statista, “Smartphones Wipe Out Decades of Camera Industry Growth,” <https://www.statista.com/chart/15524/worldwide-camera-shipments/>. CIPA shipment data series, multiple years. ↵
58. CNN Business, “Meet the translation professionals losing their jobs to AI,” January 2026, <https://www.cnn.com/2026/01/23/tech/translation-language-jobs-ai-automation-intl>. Carl Benedikt Frey (Oxford Martin School), 2025 study on translator employment across 696 US labour markets. American Translators Association industry position: <https://www.atanet.org/client-assistance/blog-machine-translation-vs-human-translation/>. Wikipedia, “Google Neural Machine Translation,” https://en.wikipedia.org/wiki/Google_Neural_Machine_Translation. ↵
59. *Dream Machine* Issue 2, “Editor’s Pick,” 10 October 2025. <https://www.linkedin.com/pulse/dream-machine-creative-ai-news-insight-oct-25-2-pete-woodbridge-mnrjc/>. ↵
60. OpenAI, “Introducing AgentKit,” 6 October 2025. <https://openai.com/index/introducing-agentkit/>. ↵
61. TechCrunch, “OpenAI launches AgentKit to help developers build and ship AI agents,” 6 October 2025. <https://techcrunch.com/2025/10/06/openai-launches-agentkit-to-help-developers-build-and-ship-ai-agents/>. Also coverage at *InfoQ*, “OpenAI Dev Day 2025 Introduces GPT-5 Pro API, Agent Kit, and More.” <https://www.infoq.com/news/2025/10/openai-dev-day/>. ↵
62. *Dream Machine* Issue 2: “Agentic AI — the class of AI systems that can plan, act, and pursue goals with autonomy — promises a new era of collaboration in creative industries... Its another step along the Human-AI Agency Continuum.” See also *TVB Europe*, “Is Agentic AI About to Change the Media and Entertainment Industry?” <https://www.tvbeurope.com/artificial-intelligence/opinion-is-agentic-ai-about-to-change-the-media-and-entertainment-industry>. ↵
63. Google DeepMind, Veo 3.1 release, October 2025. *Dream Machine* Issue 3. ↵
64. *MusicTech*, “iZotope Ozone 12’s AI assistant is cool, but the Stem EQ is the real star.” <https://musictech.com/reviews/plugin-ins/izotope-ozone-12-review/>. *Dream Machine* Issue 3. ↵
65. Adobe, “Inaugural Adobe Creators’ Toolkit Report,” October 2025. <https://news.adobe.com/news/2025/10/adobe-max-2025-creators-survey>. Survey of 16,000 creators across eight countries, released at Adobe MAX 2025. *Dream Machine* Issue 6. ↵
66. Adobe, *op. cit.* The same survey: 86% of creators use creative generative AI; 76% say it has helped grow their business or brand; 81% say AI lets them make content they otherwise couldn’t have made; 69% worry about their work being used to train AI without consent; 70% are optimistic about agentic AI; 85% would use AI that learns their creative style. ↵
67. Mureka, “Music Agent Studio” launch, mid-October 2025. *Dream Machine* Issue 4. https://www.linkedin.com/posts/sherrihendrickson_mureka-unveils-music-agent-studio-and-enhanced-share-7384999251526864896-cNYg/. ↵
68. *Finsmes*, “AdsAgency Raises \$12M in Seed Funding,” October 2025. <https://www.finsmes.com/2025/10/adsgency-raises-12m-in-seed-funding.html>. *Dream Machine* Issue 4. ↵
69. *Musically*, “Meet Lenny, an AI agent to help organisers of live music events.” <https://musically.com/2025/10/20/meet-lenny-an-ai-agent-to-help-organisers-of-live-music-events/>. *Dream Machine* Issue 4. ↵
70. *GamesRadar*, “Even under USD20 million in debt, EA reportedly pushes 15,000 employees to use AI as a ‘thought partner’ for everything from character art to playtesting.” <https://www.gamesradar.com/games/even-under-usd20-million-in-debt-ea-reportedly-pushes-15-000-employees-to-use-ai-as-a-thought-partner-for-everything-from-character-art-to-playtesting/>. *Dream Machine* Issue 6. ↵
71. PYMNTS, “Adobe Lets Users Design and Edit Using ChatGPT.” <https://www.pymnts.com/artificial-intelligence-2/2025/adobe-lets-users-design-and-edit-using-chatgpt/>. Adobe blog: “Edit images, designs, and PDFs right inside ChatGPT — thanks to Adobe Express, Photoshop, and Acrobat.” <https://blog.adobe.com/en/publish/2025/12/10/edit-photoshop-chatgpt>. *Dream Machine* Issue 12. ↵

72. TechCrunch, “Anthropic launches interactive Claude apps, including Slack and other workplace tools,” 26 January 2026. <https://techcrunch.com/2026/01/26/anthropic-launches-interactive-claude-apps-including-slack-and-other-workplace-tools/>. Heygen Video Agent: https://www.linkedin.com/posts/heygen_introducing-the-new-video-agent-activity-7421597801240801282-d1CF. *Dream Machine* Issue 16.↵
73. *Dream Machine* Issue 21, “Editor’s Pick: Adobe and NVIDIA Just Raised the Stakes for Creative AI,” 19 March 2026.↵
74. Adobe Summit 2026, “Agentic Creative Intelligence” keynote framing. *Dream Machine* Issue 26.↵
75. *Dream Machine* Issue 29, May 2026, citing Sony’s adoption of Claude Code studios with multi-agent coordination.↵
76. Anthropic, public statements on agent deployment patterns through Q1 2026. Cf. *Dream Machine* Issues 11, 16, 22.↵
77. *gamesindustry.biz*, “AI was an expensive mistake’: Charles Cecil on innovation, insolvency, and Broken Sword.” <https://www.gamesindustry.biz/ai-was-an-expensive-mistake-charles-cecil-on-innovation-insolvency-and-broken-sword>. *Dream Machine* Issue 3.↵
78. *Niche Gamer*, “Larian Studios backs off from gen AI, says tech won’t be used in new Divinity.” <https://nichegamer.com/larian-studios-backs-off-from-gen-ai/>. *Dream Machine* Issue 14.↵
79. *Decrypt*, “‘Warhammer 40,000’ Maker Games Workshop Rules Out Generative AI.” <https://decrypt.co/354482/warhammer-40000-maker-games-workshop-rules-out-generative-ai>. *Dream Machine* Issue 14.↵
30. *Niche Gamer*, “Manor Lords publisher Hooded Horse won’t work with devs using gen AI.” <https://nichegamer.com/manor-lords-publisher-hooded-horse-wont-work-with-devs-using-gen-ai/>. *Dream Machine* Issue 14.↵
81. *gamesindustry.biz*, “RuneScape maker Jagex says it will never use generative AI to make in-game content.” <https://www.gamesindustry.biz/runescape-maker-jagex-says-it-will-never-use-generative-ai-to-make-in-game-content>. *Dream Machine* Issue 16.↵
32. Imperva, 2025 *Bad Bot Report: How AI is Supercharging the Bot Threat*. <https://www.imperva.com/blog/2025-imperva-bad-bot-report-how-ai-is-supercharging-the-bot-threat/>. *Dream Machine* Issue 4.↵
33. Cloudflare, “The crawl-to-click gap: Cloudflare data on AI bots, training, and referrals.” <https://blog.cloudflare.com/crawlers-click-ai-bots-training/>. *Dream Machine* Issue 4. Later 2025 updates show training crawlers declining from ~90% to ~74% of AI bot activity as scraper bots rose to 24% and a new “agentic” category emerged at 1.7%; see Cloudflare, “A deeper look at AI crawlers: breaking down traffic by purpose and industry.” <https://blog.cloudflare.com/ai-crawler-traffic-by-purpose-and-industry/>.↵
34. Grand View Research, “Generative AI Content Creation Market Report.” <https://www.grandviewresearch.com/industry-analysis/generative-ai-content-creation-market-report>. *Dream Machine* Issue 4 also cites Gartner and Europol forecasts of 90–99% AI-generated or AI-assisted online content by 2030.↵
85. *Dream Machine* Issue 4, “Editor’s Pick: Is the Internet Dead Yet?” 23 October 2025. <https://www.linkedin.com/pulse/dream-machine-creative-ai-news-insight-oct-25-issue-4-woodbridge-hzttc/>.↵
36. Wikipedia, *Dead Internet Theory*. https://en.wikipedia.org/wiki/Dead_Internet_theory. *Dream Machine* Issue 4.↵
87. Graphite, 2025 analysis of new web content by author type (human vs. AI vs. AI-assisted). Cited in *Dream Machine* Issue 4.↵
38. For “model collapse” as a term of art, see Ilia Shumailov et al., “The Curse of Recursion: Training on Generated Data Makes Models Forget” (2024), and subsequent literature.↵
39. Futurism, “Researchers built a social network with only AI agents — within hours it had collapsed into warring tribes.” <https://futurism.com/social-network-ai-intervention-echo-chamber>. *Dream Machine* Issue 4.↵

90. *Digital Music News*, “Instagram Chief Says We Should ‘Fingerprint Real Media’ Instead of Tracking and Disclosing AI Slop.” <https://www.digitalmusicnews.com/2026/01/05/instagram-chief-ai-slop-comments/>. See also *WebProNews*, “Instagram Head Warns AI Images Erode Trust, Calls for Verification Standards.” <https://www.webpronews.com/instagram-head-warns-ai-images-erode-trust-calls-for-verification-standards/>. *Dream Machine* Issue 13.↵
91. Sundance Institute, “Centering the Artist: Why We’re Launching the AI Literacy Initiative.” <https://www.sundance.org/blogs/centering-the-artist-why-were-launching-the-ai-literacy-initiative/>. *Dream Machine* Issue 16.↵
92. *Stereogum*, “Bandcamp bans AI music.” <https://stereogum.com/2485199/bandcamp-bans-ai-music/news>. *Dream Machine* Issue 14.↵
93. *CNET*, “San Diego Comic-Con Draws a Line: No AI Art Allowed at 2026 Event.” <https://www.cnet.com/culture/san-diego-comic-con-bans-ai-art-for-2026-event/>. *Dream Machine* Issue 16.↵
94. Deezer, “AI-generated tracks now represent 44% of all new uploaded music,” April 2026. <https://newsroom-deezer.com/2026/04/ai-generated-tracks-represent-44-of-new-uploaded-music/>. *Music Business Worldwide*, “75,000 AI-generated tracks now flood Deezer daily.” <https://www.musicbusinessworldwide.com/75000-ai-generated-tracks-now-flood-deezer-daily-representing-44-of-all-new-music-uploaded-to-the-platform-says-streamer/>. *Dream Machine* Issues 7, 26, 27, 28.↵
95. *The Hollywood Reporter*, “‘Synthetic Sincerity’ by Marc Isaacs Explores if AI Characters Can Be Taught Authenticity: IDFA.” <https://www.hollywoodreporter.com/movies/movie-news/synthetic-sincerity-film-idfa-ai-authenticity-interview-1236426180/>. *Dream Machine* Issue 8.↵
96. *Variety*, “AI-Generated Images Threaten Future of Documentary as People ‘Will Stop Believing Anything’.” <https://variety.com/2025/film/festivals/ai-generated-images-threaten-future-of-documentary-1236583466/>. *Dream Machine* Issue 8.↵
97. PR Newswire, “From Apple TV Creative to AI Filmmaker: Hoyt Dwyer’s Animated Film To Compete at AI FilmFest Japan 2025.” <https://www.prnewswire.com/news-releases/from-apple-tv-creative-to-ai-filmmaker-hoyt-dwyers-animated-film-to-compete-at-ai-filmfest-japan-2025-302598064.html>. *Dream Machine* Issue 6.↵
98. *Variety*, “AI Creator Behind Viral ‘Deadpool,’ ‘Harry Potter’ Christmas Clip Made His Film in a Ukrainian Bomb Shelter.” <https://variety.com/2026/digital/news/ai-video-deadpool-harry-potter-andrii-daniels-1236624632/>. *Dream Machine* Issue 16.↵
99. *Branding in Asia*, “‘It’s the Most Terrible Time of the Year’ — McDonald’s Netherlands’ Wonderfully Chaotic, AI-Driven Christmas Film.” <https://www.brandingasia.com/its-the-most-terrible-time-of-the-year-mcdonalds-netherlands-wonderfully-chaotic-ai-driven-christmas-film/>. Pulled following backlash: *SiliconAngle*, “Not ready: McDonald’s AI-generated ad taken down after public backlash.” <https://siliconangle.com/2025/12/10/not-ready-mcdonalds-ai-generated-ad-taken-public-backlash/>. *Dream Machine* Issue 11.↵
100. BBC News, “Fashion house Valentino criticised over ‘disturbing’ AI handbag ads.” <https://www.bbc.co.uk/news/articles/cwvyjvvn83go>. *Dream Machine* Issue 10.↵
101. *Adweek*, “Coca-Cola Uses AI to Rekindle the Magic of Its Holiday Ads.” <https://www.adweek.com/creativity/coca-cola-uses-ai-to-rekindle-the-magic-of-its-holiday-ads/>. *Dream Machine* Issue 6.↵
102. *AI News*, “AI causes reduction in users’ brain activity, MIT.” <https://www.artificialintelligence-news.com/news/ai-causes-reduction-in-users-brain-activity-mit/>. *Dream Machine* Issue 1.↵
103. For “model collapse” as a term of art, see Ilia Shumailov et al., “The Curse of Recursion: Training on Generated Data Makes Models Forget” (2024), and subsequent literature.↵
104. *Digital Music News*, “Instagram Chief Says We Should ‘Fingerprint Real Media’ Instead of Tracking and Disclosing AI Slop.” <https://www.digitalmusicnews.com/2026/01/05/instagram-chief-ai-slop-comments/>. See also *WebProNews*, “Instagram Head Warns AI Images Erode Trust, Calls for Verification Standards.” <https://www.webpronews.com/instagram-head-warns-ai-images-erode-trust-calls-for-verification-standards/>. *Dream Machine* Issue 13.↵

35. Graphite, 2025 analysis of new web content by author type (human vs. AI vs. AI-assisted). Cited in *Dream Machine* Issue 4.↵
36. Deezer, “AI-generated tracks now represent 44% of all new uploaded music,” April 2026. <https://newsroom-deezer.com/2026/04/ai-generated-tracks-represent-44-of-new-uploaded-music/>. *Music Business Worldwide*, “75,000 AI-generated tracks now flood Deezer daily, representing 44% of all new music uploaded to the platform.” <https://www.musicbusinessworldwide.com/75000-ai-generated-tracks-now-flood-deezer-daily-representing-44-of-all-new-music-uploaded-to-the-platform-says-streamer/>. *Dream Machine* Issues 7, 26, 27, 28.↵
37. Ditto Music research, October 2025 and prior. *Press Ditto Music*, “48% of artists use AI to make music — fewer than in 2023.” <https://press.dittomusic.com/48-of-artists-use-ai-to-make-music-fewer-than-in-2023>. *Dream Machine* Issue 2.↵
38. Musically, “Universal and Warner could sign landmark AI deals within weeks.” <https://musically.com/2025/10/02/report-umg-and-wmg-could-sign-landmark-ai-deals-within-weeks/>. Spotify Newsroom, “Spotify Strengthens AI Protections for Artists, Songwriters, and Producers.” <https://newsroom.spotify.com/2025-09-25/spotify-strengthens-ai-protections/>. *Dream Machine* Issue 1.↵
39. Musically, “50,000 AI music tracks are now uploaded to Deezer every day.” <https://musically.com/2025/11/12/50000-ai-music-tracks-are-now-uploaded-to-deezer-every-day/>. *Dream Machine* Issue 7.↵
10. Deezer, April 2026, *op. cit.*↵
11. Musically, “UMG boss slams exponential growth of AI slop on streaming services.” <https://musically.com/2026/01/09/umg-boss-slams-exponential-growth-of-ai-slop-on-streaming-services/>. *Dream Machine* Issue 14.↵
12. Musically, “Report: 56.9% of new independent songs in China are AI-generated.” <https://musically.com/2026/01/05/report-56-9-of-new-independent-songs-in-china-are-ai-generated/>. *Dream Machine* Issue 13.↵
13. *The Wrap*, “An AI Podcasting Machine Is Churning Out 3,000 Episodes a Week — and People Are Listening.” <https://www.thewrap.com/ai-podcasts-hosts-inception-point-ai/>. *Dream Machine* Issue 8.↵
14. *Dream Machine* Issue 28, May 2026, citing aggregator-platform data on “podslope” classification.↵
15. *The Hollywood Reporter*, “Merriam-Webster Names ‘Slop’ Word of the Year Amid AI Boom.” <https://www.hollywoodreporter.com/news/general-news/slop-word-year-2025-merriam-webster-1236450780/>. *Dream Machine* Issue 12.↵
16. *Digital Music News*, “YouTube CEO Puts ‘Managing AI Slop’ on the Priority List for 2026.” <https://www.digitalmusicnews.com/2026/01/22/youtube-ceo-ai-slop-2026-comments/>. *Dream Machine* Issue 16.↵
17. *The Guardian*, “YouTube AI channels spreading fake, anti-Labour videos viewed 1.2bn times in 2025.” <https://www.theguardian.com/technology/2025/dec/13/fake-anti-labour-video-billion-views-youtube-2025>. *Dream Machine* Issue 12.↵
18. Deezer/Ipsos survey, November 2025. <https://newsroom-deezer.com/2025/11/deezer-ipsos-survey-ai-music/>. *Dream Machine* Issue 7.↵
19. *Bain & Company*, “In an AI Age, People Still Want the Radio Star.” <https://www.bain.com/insights/in-an-ai-age-people-still-want-the-radio-star/>. *Dream Machine* Issue 16.↵
20. Deezer, April 2026, *op. cit.* “Up to 85% of the streams generated by fully AI-generated tracks were in fact fraudulent in 2025.”↵
21. *Bloomberg*, “AI Changed Chess. Grandmasters Now Win With Unpredictable Moves.” 27 March 2026. <https://www.bloomberg.com/news/articles/2026-03-27/ai-changed-chess-grandmasters-now-win-with-unpredictable-moves>. *Dream Machine* Issue 23.↵
22. *Billboard*, “AI Artist Xania Monet Climbs the Charts — And Signs a Multimillion-Dollar Record Deal.” <https://www.billboard.com/pro/ai-music-artist-xania-monet-multimillion-dollar-record-deal/>.↵
23. *Billboard*, *op. cit.*; CNN, “Xania Monet is the first AI-powered artist to debut on a Billboard airplay chart.” <https://www.cnn.com/2025/11/01/entertainment/xania-monet-billboard-ai>↵
24. *Billboard*, *op. cit.*↵

25. *Bangkok Post*, “AI singer Xania Monet signs \$3m deal with record label.” <https://www.bangkokpost.com/life/tech/3142355/ai-singer-xania-monet-signs-3m-deal-with-hollywood-media>. *Dream Machine* Issue 7.↵
26. Multiple outlets; quoted in *Billboard* feature *op. cit.*↵
27. Telisha Jones quoted in *Billboard*, *op. cit.*↵
28. NPR, “Breaking Rust is a hot new country act on the Billboard charts. It’s powered by AI.” <https://www.npr.org/2025/11/10/nx-si-5604320/breaking-rust-is-a-hot-new-country-act-on-the-billboard-charts-its-powered-by-ai>. *Dream Machine* Issue 7.↵
29. *Washington Post*, “‘Walk My Walk,’ Breaking Rust: AI country hit triggers Nashville angst.” <https://www.washingtonpost.com/style/2025/12/28/breaking-rust-ai-country/>.↵
30. *MusicRadar*, “The No. 1 country song in the US right now is AI-generated.” <https://www.musicradar.com/music-tech/the-no-1-country-song-in-the-us-right-now-is-ai-generated>. *Dream Machine* Issue 7.↵
31. BBC News, “The mysterious singer, Sienna Rose, with millions of streams is hitting the viral charts — but who (or what) is she?” <https://www.bbc.co.uk/news/articles/cq6v83gq66eo>. *Dream Machine* Issue 15.↵
32. *Billboard*, “How a MAGA Rapper Used AI to Create A Gospel Song That Climbed the Charts.” <https://www.billboard.com/pro/maga-rapper-ai-gospel-song-climbed-charts/>. *Dream Machine* Issue 9.↵
33. *Musically*, “AI band Bleeding Verse’s creator signs deal with Hollywood Media.” <https://musically.com/2025/10/07/ai-band-bleeding-verses-creator-signs-deal-with-hollywood-media/>. *Dream Machine* Issue 2.↵
34. *Musically*, “Indian AI band Trilok performs live, government denies association.” <https://musically.com/2025/12/17/indian-ai-band-trilok-performs-live-government-denies-association/>. *Dream Machine* Issue 12.↵
35. *Billboard*, “The Real Story Behind The AI Song That Knocked Tyla Off No. 1 On Billboard Afrobeats Chart.” <https://www.billboard.com/pro/ai-song-knocked-tyla-off-no-1-afrobeats/>. *Dream Machine* Issue 30.↵
36. *The Guardian*, “Paul McCartney joins music industry protest against AI with silent track.” <https://www.theguardian.com/music/2025/nov/17/the-sound-of-silence-why-theres-barely-anything-there-in-paul-mccartney-new-release>. *Dream Machine* Issue 8.↵
37. *The Guardian*, “Musicians must embrace ‘unstoppable force’ of AI, Eurythmics’ Dave Stewart urges.” <https://www.theguardian.com/music/2025/dec/05/musicians-must-embrace-unstoppable-force-of-ai-eurythmics-dave-stewart-urges>. *Dream Machine* Issue 11.↵
38. *Digital Music News*, “Nearly 800 Creatives, Including Jason Aldean and One Republic, Sign Responsible AI Declaration — ‘Stealing Our Work Is Not Innovation’.” <https://www.digitalmusicnews.com/2026/01/22/stealing-isnt-innovation/>. *Dream Machine* Issue 16.↵
39. *MusicTech*, “Jack Antonoff brands AI music makers as ‘godless whores’.” <https://musictech.com/news/industry/jack-antonoff-ai-music-makers-godless-whores/>. *Dream Machine* Issue 30.↵
40. Stability AI, “Universal Music Group and Stability AI Announce Strategic Alliance.” <https://stability.ai/news/universal-music-group-and-stability-ai-announce-strategic-alliance>. *Dream Machine* Issue 5.↵
41. Stability AI, “Warner Music Group and Stability AI Join Forces To Build The Next Generation Of Responsible AI Tools For Music Creation.” <https://stability.ai/news/warner-music-group-and-stability-ai-join-forces-to-build-next-gen-tools>. *Dream Machine* Issue 8.↵
42. Universal Music, “Universal Music Group and Splice to Collaborate on the Next Generation of AI-Powered Music Creation Tools for Artists.” <https://www.universalmusic.com/universal-music-group-and-splice-to-collaborate-on-the-next-generation-of-ai-powered-music-creation-tools-for-artists/>. *Dream Machine* Issue 12.↵
43. LinkedIn / *Lexology*, “Munich Regional Court rules for GEMA against OpenAI.” Coverage: https://www.linkedin.com/posts/dr-barry-scannell-bbb5aa207_in-a-major-ruling-for-european-copyright-share-7393957246386323457-8bbx. *Dream Machine* Issue 7.↵

44. EDM.com, “‘Biggest Theft in Music History’: Rights Group Sues Suno as AI Music Showdown Escalates.” <https://edm.com/gear-tech/rights-group-sues-suno-copyright-infringement/>. *Dream Machine* Issue 7.↵
45. Music Business Worldwide, “Wixen files \$50m copyright suit against Meta.” <https://www.musicbusinessworldwide.com/wixen-files-50m-copyright-suit-against-meta-claims-tech-giant-wants-to-replace-songwriters-with-ai/>. *Dream Machine* Issue 16.↵
46. *Dream Machine* Issue 17 reportage on UMG’s \$3B suit against Anthropic.↵
47. Stereogum, “Bandcamp bans AI music.” <https://stereogum.com/2485199/bandcamp-bans-ai-music/news>. *Dream Machine* Issue 14.↵
48. *Dream Machine* Issue 18 reportage of Deezer licensing its detection tool.↵
49. TechRadar, “AI music is flooding Spotify, and subscribers are furious.” <https://www.techradar.com/audio/spotify/ai-music-is-flooding-spotify-and-subscribers-are-furious-heres-why-music-fans-no-longer-trust-discover-weekly>. *Dream Machine* Issue 14.↵
50. CNET, “San Diego Comic-Con Draws a Line: No AI Art Allowed at 2026 Event.” <https://www.cnet.com/culture/san-diego-comic-con-bans-ai-art-for-2026-event/>. *Dream Machine* Issue 16.↵
51. *The Independent*, “AI-generated song banned from Swedish charts: ‘It’s deceiving.’” <https://www.independent.co.uk/tv/news/ai-music-song-banned-sweden-spotify-b2901627.html>. *Dream Machine* Issue 15.↵
52. *Hollywood Reporter*, “Amazon Orders Three Animated Series Using Generative AI.” <https://www.hollywoodreporter.com/tv/tv-news/amazon-orders-three-animated-series-generative-ai-1236606219/>. *Variety*, “Jorge Gutierrez Drops Out of Amazon MGM AI-Generated Series After Backlash.” <https://variety.com/2026/tv/news/jorge-gutierrez-drops-out-amazon-mgm-ai-generated-series-backlash-1236762285/>. *Futurism*, “Amazon AI-Generated Series Canceled.” <https://futurism.com/artificial-intelligence/amazon-ai-generated-series-canceled>. *Dream Machine* Issue 32.↵
53. Guillermo del Toro, subtweet responding to Amazon AI Creators Fund announcement, May/June 2026. Coverage in *Dream Machine* Issue 32. Cross-reference: del Toro’s “rather die” statement in Chapter 7, §“Position Three: We are not doing this.”↵
54. Apple Music platform statement, June 2026: AI-generated songs make up less than 1% of plays on the platform. *Dream Machine* Issue 31.↵
55. Spotify–Universal Music Group landmark deal allowing Spotify Premium subscribers to create AI covers and remixes of UMG-catalogue tracks as a paid Premium add-on, June 2026. *Dream Machine* Issue 31.↵
56. Spotify CEO Daniel Ek defending the platform’s move to AI music as “better than ‘slop,’” June 2026. *Dream Machine* Issue 31.↵
57. Swedish startup Tonada — AI music for retailers (June 2026). *Dream Machine* Issue 31.↵
58. *Soultracks*, “A.I.-generated music is catchy, familiar... and boring.” <https://soultracks.com/news-ai-generated-music-is-catchy-boring/>. *Dream Machine* Issue 14.↵
59. *The Independent*, “AI-generated song banned from Swedish charts: ‘It’s deceiving.’” <https://www.independent.co.uk/tv/news/ai-music-song-banned-sweden-spotify-b2901627.html>. *Dream Machine* Issue 15.↵
50. *Marketing Week*, “You can’t dismiss AI ads as slop when they’re winning in testing.” Coverage discussed in *Dream Machine* Issue 22.↵
61. UK Department for Science, Innovation and Technology, *Statement of Progress on Copyright and AI*, 15 December 2025. <https://www.gov.uk/government/publications/copyright-and-artificial-intelligence-progress-report/copyright-and-artificial-intelligence-statement-of-progress-under-section-137-data-use-and-access-act>. *Dream Machine* Issue 12, “Editor’s Pick: 88% of Creators Said ‘No.’” 18 December 2025.↵
52. UK DSIT, original consultation, 17 December 2024 – 25 February 2025. Discussion in IPWatchdog, “Respondents to UK AI Consultation Overwhelmingly Want AI Companies to License Copyrighted Works in All Cases.” <https://ipwatchdog.com/2025/12/16/respondents-uk-ai-consultation-overwhelmingly-want-ai-companies-license-copyrighted-works-all-cases/>.↵
53. IPWatchdog, *op. cit.*; Hogan Lovells, “Copyright and AI: UK government publishes statement of progress.” <https://www.hoganlovells.com/en/publications/copyright-and-ai-uk-government-publishes-statement-of>

progress.↵

54. UK DSIT, *Statement of Progress*, *op. cit.*; analysis at UCL Copyright Queries, “UK government publishes progress statement on AI and copyright consultation.” <https://blogs.ucl.ac.uk/copyright/2025/12/23/uk-government-publishes-progress-statement-on-ai-and-copyright-consultation/>.↵
55. UK DSIT, *Statement of Progress*, *op. cit.*↵
56. Society of Authors submission to the UK consultation, quoted in IPWatchdog, *op. cit.*↵
57. Dr Barry Scannell, LinkedIn analysis of GEMA v. OpenAI ruling, November 2025. https://www.linkedin.com/posts/dr-barry-scannell-bbb5aa207_in-a-major-ruling-for-european-copyright-share-7393957246386323457-8bbx. *Dream Machine* Issue 7.↵
58. EDM.com, “‘Biggest Theft in Music History’: Rights Group Sues Suno as AI Music Showdown Escalates.” <https://edm.com/gear-tech/rights-group-sues-suno-copyright-infringement/>. *Dream Machine* Issue 7.↵
59. *Music Business Worldwide*, “Wixen files \$50m copyright suit against Meta, claims tech giant wants to replace songwriters with AI.” <https://www.musicbusinessworldwide.com/wixen-files-50m-copyright-suit-against-meta-claims-tech-giant-wants-to-replace-songwriters-with-ai/>. *Dream Machine* Issue 16.↵
70. *Dream Machine* Issue 17, on UMG’s \$3B suit against Anthropic.↵
71. *Complete Music Update*, “Johnny Cash estate uses ELVIS Act to sue Coke over tribute act ad soundtrack.” <https://completemusicupdate.com/johnny-cash-estate-uses-elvis-act-to-sue-coke-over-tribute-act-ad-soundtrack/>. *Dream Machine* Issue 9.↵
72. Reuters, “European lawmakers seek EU-wide minimum age to access AI chatbots, social media.” <https://www.reuters.com/legal/litigation/european-lawmakers-seek-eu-wide-minimum-age-access-ai-chatbots-social-media-2025-11-26/>. *Dream Machine* Issue 9.↵
73. SAG-AFTRA contract update reporting through Q2 2026. *Dream Machine* Issues 20, 26, 29. Coverage: <https://www.theverge.com/news/842848/new-york-law-ai-advertisements-sag-aftra-labor>.↵
74. Equity (U.K.), “Performers prepared to take industrial action over AI in landslide 99% vote.” <https://www.equity.org.uk/news/2025/performers-prepared-to-take-industrial-action-over-ai-in-landslide-99-vote>. *Dream Machine* Issue 12.↵
75. Equity (U.K.), “Equity welcomes improved offer in AI protection negotiations in film and TV.” <https://www.equity.org.uk/news/2026/equity-welcomes-improved-offer-in-ai-protection-negotiations-in-film-and-tv>. *Dream Machine* Issue 15.↵
76. Cannes Film Festival AI Disclosure Standard launch, May 2026. *Dream Machine* Issue 29.↵
77. *Musically*, “BPI sets out transparency and sovereignty demands to secure ‘AI licensing boom’.” <https://musically.com/2026/05/19/bpi-transparency-sovereignty-ai-licensing-boom/>. *Dream Machine* Issue 30.↵
78. UK DSIT, *Statement of Progress*, *op. cit.*↵
79. *Dream Machine* Issue 21, 19 March 2026, on the UK government’s revised position on AI copyright.↵
30. *Digital Music News*, “The AI Licensing Shift — Creative Weight Attribution Emerges as Music Industry Game-Changer for Rights Holders.” <https://www.digitalmusicnews.com/2026/01/26/ai-licensing-shift-creative-weight-attribution/>. See also *Digital Music News*, “Artificial Intelligence Attribution and Licensing Startup Musical AI Scores \$4.5 Million Raise.” <https://www.digitalmusicnews.com/2026/01/13/musical-ai-funding-january-2026/>. *Dream Machine* Issues 14, 16.↵
81. PRS for Music, “PRS for Music AI Survey 2026.” <https://www.prsformusic.com/m-magazine/news/prs-for-music-ai-survey-2026>. *Dream Machine* Issue 16.↵
32. *Broadcast Now*, “Alex Mahon joins Stellar AI Creative Summit line-up” (covering the launch of the UCL/RCA Centre for Creative AI). <https://www.broadcastnow.co.uk/broadcasters/alex-mahon-joins-stellar-ai-creative-summit-line-up/5209227.article>. *Dream Machine* Issue 1.↵
33. *Complete Music Update*, “Artists must have creative control in AI deals or risk ending up with ‘scraps’, says US artist trade body.” <https://completemusicupdate.com/artists-must-have-creative-control-in-ai-deals-or-risk-ending-up-with-scraps-says-us-artist-trade-body/>. *Dream Machine* Issue 6.↵
34. *Digital Music News*, “Nearly 800 Creatives, Including Jason Aldean and One Republic, Sign Responsible AI Declaration — ‘Stealing Our Work Is Not Innovation’.”

<https://www.digitalmusicnews.com/2026/01/22/stealing-isnt-innovation/>. *Dream Machine* Issue 16.↵

35. Adobe Firefly milestone and adoption data, in Appendix E: Dynamics of Generative AI Adoption, §“The Ubiquity of AI in Visual and Digital Arts.” Firefly Foundry and Firefly Image Model 5 launch reporting, Adobe MAX 2025: <https://news.adobe.com/news/2025/10/adobe-max-2025-firefly-foundry>; <https://news.adobe.com/news/2025/10/adobe-max-2025-firefly>.↵
36. Bria AI consent-licensed dataset and attribution mechanism. [TODO: confirm primary citation — Bria’s licensed-data white paper or Series B coverage.]↵
37. Getty Images, “Generative AI by iStock” launch, built on NVIDIA Picasso, trained exclusively on Getty’s licensed library with contributor royalties. [TODO: confirm citation — Getty press release or *Reuters* coverage.]↵
38. Moonvalley Marey, generative-video foundation model trained on licensed video. [TODO: confirm citation — Moonvalley launch coverage in *The Verge* / *TechCrunch*.]↵
39. AIODE, ethically-trained music creation DAW. See Chapter 16: The Tools, §“Audio modality models.”↵
40. *MusicTech*, “Tamber is an ‘ethically trained’ AI tool to aid the creative process – and you can use arm gestures to control it.” <https://musictech.com/news/gear/tamber-ai-ethically-trained-arm-gestures/>. Tamber product page: <https://tamber.ai/>. *Dream Machine* Issue 30.↵
91. Stability AI / Universal Music Group strategic alliance: <https://stability.ai/news/universal-music-group-and-stability-ai-announce-strategic-alliance>. Stability AI / Warner Music: <https://stability.ai/news/warner-music-group-and-stability-ai-join-forces-to-build-next-gen-tools>. Universal Music / Splice partnership: <https://www.universalmusic.com/universal-music-group-and-splice-to-collaborate-on-the-next-generation-of-ai-powered-music-creation-tools-for-artists/>. *Dream Machine* Issues 5, 8, 12.↵
92. Reporting on AI-generated images in the Adobe Stock training corpus, *Bloomberg*, April 2024. [TODO: confirm exact citation.]↵
93. Spotify and Universal Music Group, “Landmark Licensing Agreements for Fan-Made Covers and Remixes,” 21 May 2026. Spotify Newsroom: <https://newsroom.spotify.com/2026-05-21/universal-music-group-spotify-licensing-agreements-fan-made-covers-remixes/>. *TechCrunch*: <https://techcrunch.com/2026/05/21/spotify-and-universal-music-strike-deal-allowing-fan-made-ai-covers-and-remixes/>. *Medianama*: <https://www.medianama.com/2026/05/223-spotify-umg-fans-create-licensed-ai-covers-remixes/>. *Dream Machine* Issue 31. Detailed treatment of the three-Cs framework, the deal’s prehistory in the UMG settlement with Udio (October 2025), the *Verified by Spotify* badge, and the broader analysis in Appendix L.↵
94. Hasbro, “Hasbro Launches Sixth Wall, a New AI Studio,” *BusinessWire* press release, 3 June 2026: <https://www.businesswire.com/news/home/20260603297922/en/>. *The Wrap*: <https://www.thewrap.com/industry-news/business/hasbro-ai-studio-sixth-wall-launch/>. *Screen Daily*: <https://www.screendaily.com/news/hasbro-launches-subsidiary-ai-studio-sixth-wall-to-exploit-ip/5217380.article>. CharacterOS runtime governance layer; talent-participation pay-out structure for original performers; initial 12-character catalogue; 13+/enterprise focus. Detailed treatment in Appendix L. Cross-reference in Chapter 7 §“Position Five: Own the rails.”↵
95. *Dream Machine* Deep Dive № 14, *The Programmable Brand: How AI Is Reshaping IP Licensing in the Generative Economy*, 4 June 2026. Preserved in full as Appendix L. The deep dive’s three Cs framework draws on the Spotify–UMG licensing agreements; the five-model comparison covers Hasbro/Sixth Wall (Behavioural Licensing), Spotify/UMG (platform-embedded remixing), Disney/OpenAI (the framework that survived the platform’s collapse — see footnote ^3a in Chapter 7), the independent-artist commons-and-DAO models (Grimes/Elf.tech with a ~50% master-royalty split; Holly Herndon / Holly+ with DAO governance; ElevenLabs heritage voices via CMG Worldwide; Meta’s 2023 celebrity-chatbot programme as the closed-static anti-pattern), and Suno/Alexa as the *create-to-consume* demand floor. The underlying economic argument draws on Coase (1937), Williamson, and Barney’s VRIN resource-based view; the appendix’s load-bearing claim is that AI *relocates* rather than *abolishes* transaction costs — collapsing search and negotiation while raising the cost of monitoring and enforcement, and pushing the durable value into the governance layer.↵
96. Adobe Firefly IP indemnification for enterprise customers. [TODO: confirm citation — Adobe enterprise terms or *The Verge* coverage from 2023.]↵

97. Microsoft, "Microsoft announces new Copilot Copyright Commitment for customers," 7 September 2023. <https://blogs.microsoft.com/on-the-issues/2023/09/07/copilot-copyright-commitment-ai-legal-concerns/>.↵
98. Google Cloud Generative AI indemnification: <https://cloud.google.com/blog/products/ai-machine-learning/protecting-customers-with-generative-ai-indemnification>.↵
99. IBM watsonx uncapped indemnity for enterprise customers. [TODO: confirm citation.]↵
100. Sundance AI Literacy Initiative, in Chapter 12: Authenticity, the New Scarcity, §"The provenance infrastructure, named."↵
101. For Disney's parallel position, see *Deadline*, "Disney Sends Cease And Desist Letter To Character.ai." <https://deadline.com/2025/09/disney-cease-and-desist-letter-characterai-copyright-infringement-1236566831/>. For Studio Ghibli's similar stance: *NDTV Profit*, "Studio Ghibli And Studio That Developed Elden Ring Send Stern Message To OpenAI." <https://www.ndtvprofit.com/technology/studio-ghibli-and-studio-that-developed-elden-ring-send-stern-message-to-openai>. *Dream Machine* Issues 2, 6.↵
102. *Variety*, "Is 'AI Resistance' Setting the Music Sector Back? WMG's Robert Kyncl Sees 'An Incredible Value Creation Opportunity,' But Warns 'We Cannot Wait the Way the Industry Did 25 Years Ago'." <https://variety.com/2026/music/news/wmg-robert-kyncl-ai-resistance-1236748901/>. *Dream Machine* Issue 30.↵
103. Adobe, *Creators' Toolkit Report*, *op. cit.* 69% of 16,000 surveyed creators worried about their work being used to train AI without consent.↵
104. *Screen Daily*, "We got it wrong says UK Culture Secretary Lisa Nandy of government misstep on AI and copyright strategy," 1 June 2026. <https://www.screendaily.com/news/we-got-it-wrong-says-uk-culture-secretary-lisa-nandy-of-government-misstep-on-ai-and-copyright-strategy/5217330.article>. *Dream Machine* Issue 32.↵
105. *CNBC*, "Netflix 'all in' on leveraging AI as the tech creeps into entertainment industry," 22 October 2025. <https://www.cnb.com/2025/10/22/netflix-all-in-on-leveraging-ai-in-its-streaming-platform.html>. *Dream Machine* Issue 4.↵
106. *Futurism*, "Lionsgate's Attempt to Create Movies Using AI Has Crumbled Into Disaster." <https://futurism.com/artificial-intelligence/lionsgate-movies-ai>. *Dream Machine* Issue 1.↵
107. *The Guardian*, "Disney to invest \$1bn in OpenAI, allowing characters in Sora video tool." <https://www.theguardian.com/business/2025/dec/11/disney-open-ai-sora-video-deal>. *Dream Machine* Issue 11. The announcement was widely treated at the time as a completed, operative template; on the corrected record set out in Appendix L, no formal agreement was ever signed and no money changed hands.↵
108. OpenAI staged Sora wind-down announced 24 March 2026: consumer app and website to end in late April 2026; developer API to close in late September 2026. Disney exited the partnership with reportedly as little as thirty minutes' notice. *Ropes & Gray*, "The Deal That Wasn't: What Disney and OpenAI Teach Us About Strategic Investments." <https://www.ropesgray.com/en/insights/viewpoints/the-deal-that-wasnt-what-disney-and-openai-teach-us-about-strategic-investments>. *TheStreet / AOL* (Hollywood Reporter source): <https://www.thestreet.com/technology/openai-is-shutting-down-sora-and-the-disney-deal-is-off>. *Tech Insider*: <https://tech-insider.org/openai-sora-shutdown-disney-deal-ai-video-2026/>. Peak Sora revenue reported at roughly \$540k/month against operating costs in the billions; OpenAI reallocated compute toward higher-margin coding and enterprise products. Coverage in *Dream Machine* Issue 22.↵
109. *Dream Machine* Deep Dive № 14, *The Programmable Brand*, 4 June 2026. Preserved in full as Appendix L.↵
110. PYMNTS, "Retention Is Name of the Game for Netflix's AI Strategy." <https://www.pymnts.com/subscription-commerce/2026/retention-is-name-of-the-game-for-netflixs-ai-strategy/>. *Dream Machine* Issue 15.↵
111. *Hollywood Reporter*, "Netflix is building and recruiting for an AI animation studio, called INKubator, to produce 'feature-quality' shorts." <https://www.hollywoodreporter.com/business/business-news/netflix-ai-animation-studio-inkubator-1236592110/>. *Dream Machine* Issue 30.↵
112. *Deadline*, "Amazon Builds Out AI Studios With Sports Docs Boss Matt Newman Named Head Of Live-Action." <https://deadline.com/2025/11/amazon-ai-studios-matt-newman-1236603477/>. *Dream Machine*

Issue 7.↵

13. *Wired*, “Amazon’s House of David Used Over 350 AI Shots in Season 2. Its Creator Isn’t Sorry.” <https://www.wired.com/story/amazons-house-of-david-used-over-350-ai-shots-in-season-2-its-creator-isnt-sorry/>. *Dream Machine* Issue 7.↵
14. Amazon AI Creators Fund launch and the Punky Duck cancellation. *Hollywood Reporter*: <https://www.hollywoodreporter.com/tv/tv-news/amazon-orders-three-animated-series-generative-ai-1236606219/>. *Variety*: <https://variety.com/2026/tv/news/jorge-gutierrez-drops-out-amazon-mgm-ai-generated-series-backlash-1236762285/>. *Futurism*: <https://futurism.com/artificial-intelligence/amazon-ai-generated-series-canceled>. *Dream Machine* Issue 32. Cross-reference: Chapter 5 §“The pushback” for the audience-enforcement analysis.↵
15. *Video Games Chronicle*, “NBCUniversal signs deal with Law & Order creator Dick Wolf’s son to make AI-generated games based on its IP.” <https://www.videogameschronicle.com/news/nbcuniversal-signs-deal-with-law-order-creator-dick-wolfs-son-to-make-ai-generated-games-based-on-its-ip/>. *Dream Machine* Issue 5.↵
16. *NME*, “‘The Office’, ‘Saturday Night Live’ and ‘Sex And The City’ could be turned into AI games.” <https://www.nme.com/news/gaming-news/the-office-and-sex-and-the-city-ai-video-games-3901630>. *Dream Machine* Issue 5.↵
17. *The Hollywood Reporter*, “Disney+ to Allow User-Generated Fan Content with AI.” <https://www.hollywoodreporter.com/business/digital/disney-plus-gen-ai-user-generated-content-1236426135/>. *Dream Machine* Issue 8.↵
18. *Dream Machine* Issue 8 reportage of the Disney “Office of Technology Enablement,” led by former Walt Disney Studios CTO Jamie Voris.↵
19. *Marketing Dive*, “Disney unveils TikTok-like vertical video, AI video generation tool.” <https://www.marketingdive.com/news/disney-unveils-tiktok-like-vertical-video-ai-video-generation-tool/809269/>. *Dream Machine* Issue 14.↵
20. *The Hollywood Reporter*, “Fox Entertainment Takes Equity Stake in AI-Microdramas Company Holywater.” <https://www.hollywoodreporter.com/business/business-news/fox-entertainment-invests-in-holywater-ai-microdramas-1236396802/>. *Dream Machine* Issue 3.↵
21. *Deadline*, “Sky History Acquires ‘Castles SOS,’ AI-Powered Doc Exploring Royalty, Ruins & Restoration.” <https://deadline.com/2025/11/castles-sos-ai-doc-sky-history-documentary-rick-edwards-1236627378/>. *Dream Machine* Issue 9.↵
22. *Estate Agent Today*, “Homebuilder among first to use Channel 4’s AI ads.” <https://www.estateagenttoday.co.uk/breaking-news/2025/12/homebuilder-among-first-to-use-channel-4s-ai-ads/>. *Dream Machine* Issue 11.↵
23. *The Hollywood Reporter*, “Fremantle Names Boss of New AI Native Studio Imagine Studios.” <https://www.hollywoodreporter.com/business/digital/fremantle-names-ceo-new-ai-label-imagine-studios-1236396579/>. *Dream Machine* Issue 2.↵
24. *Dream Machine* Issue 25, on Fremantle’s *Art Awakens* development.↵
25. *Indiewire*, “Another New AI Production Company Inks a Big Creative Partnership — This Time, with Ron Howard and Brian Grazer’s Imagine Entertainment.” <https://www.indiewire.com/news/business/obsidian-studio-ai-production-company-imagine-entertainment-1235158619/>. *Dream Machine* Issue 6.↵
26. *UK Tech News*, “AI film studio Wonder lands \$9m investment.” <https://www.uktech.news/ai/ai-film-studio-wonder-lands-9m-investment-20251023>. *Dream Machine* Issue 5.↵
27. Wonder Studios, “Shortlisted films revealed for The Wonder Film Festival.” https://www.linkedin.com/posts/wearewonderstudios_were-thrilled-to-share-the-shortlisted-films-activity-7404560378082246656-7NcI. *Dream Machine* Issue 11.↵
28. *Forbes*, “Meet Wonder Studios, The \$50M British Studio Striving To Become The A24 Of AI Production.” <https://www.forbes.com/sites/charliefink/2026/05/18/meet-wonder-studios-the-50m-british-studio-striving-to-become-the-a24-of-ai-production/>. *Dream Machine* Issue 30.↵
29. *The Hollywood Reporter*, “AI Company Asteria Produces New Animated Short ‘All Heart’.” <https://www.hollywoodreporter.com/movies/movie-news/natasha-lyonne-ai-company-asteria>

1236403144/. *Dream Machine* Issue 4.↵

30. *The Hollywood Reporter*, “Promise, a deep-pocketed AI studio backed by Google, aims to Bring GenAI Filmmaking and VFX to Legacy Media.” <https://www.hollywoodreporter.com/business/digital/ai-studio-promise-vfx-generation-company-1236397636/>. *Dream Machine* Issue 3.↵
31. Variety, “AI-Powered Cinematic Universe Platform enGEN3 Launched by Goldfinch.” <https://variety.com/2025/film/news/ai-powered-cinematic-universe-platform-engen3-1236543349/>. *Dream Machine* Issue 2.↵
32. *Deadline*, “Munich Based Beta Films & Industry Execs Join Forces To Launch Artificial Intelligence Start-Up Chapter41.” <https://deadline.com/2025/11/beta-film-ai-startup-chapter41-artificial-intelligence-1236612632/>. *Dream Machine* Issue 7.↵
33. *The Hollywood Reporter*, “Longtime TV Exec, Kevin Reilly, Set to Lead AI Startup Kartel.” <https://www.hollywoodreporter.com/business/digital/kevin-reilly-ceo-kartel-ai-hbo-1236424692/>. *Dream Machine* Issue 7.↵
34. Variety, “‘Wanted’ Director Timur Bekmambetov Explains His \$5 Million Plan to Generate AI Method Actors: ‘AI Is Here to Stay. We Have to Train It Responsibly.’” <https://variety.com/2025/film/news/wanted-director-method-acting-ai-actors-1236579647/>. *Dream Machine* Issue 7.↵
35. Variety, “Tilly Norwood Creator Doubles Down on AI ‘Actors’ and Says It’s a ‘More Ethical Way to Perform,’ Urges Human Actors to ‘Future-Proof’ Themselves With AI.” <https://variety.com/2026/digital/news/tilly-norwood-creator-tells-actors-to-create-ai-avatars-1236638940/>. *Dream Machine* Issue 16.↵
36. *Broadcast Now*, “Wonder Studios adapts children’s book to animated series with AI.” <https://www.broadcastnow.co.uk/production-and-post/wonder-studios-adapts-childrens-book-to-animated-series-with-ai/5211713.article>. *Dream Machine* Issue 11.↵
37. Variety, “Kling AI Partners With Evolutionary Films on Animated Feature ‘Minibots,’ Unveils Filmmaker Initiative at Cannes Market.” <https://variety.com/2026/film/news/kling-ai-evolutionary-films-minibots-cannes-1236748590/>. *Dream Machine* Issue 30.↵
38. Variety, “Tribeca Festival AI Film ‘Dreams of Violets’ by Fountain o.” <https://variety.com/2026/film/festivals/tribeca-festival-ai-film-dreams-of-violets-fountain-o-1236759724/>. Director Ash Koosha; producer Fountain o; 75-minute docudrama; subject: 2026 Iran protests; cost \$2,000; production time: 3 months; tools: Kling AI (video), Anthropic Claude (editing), Google Gemini (research). *Dream Machine* Issue 32.↵
39. The broader festival-circuit context: the *Dreams of Violets* Tribeca selection follows the *Lily* (\$1M Dubai AI Film Award, January 2026) and *Humans in the Loop* (Film Independent Sloan Distribution Grant, November 2025) as part of the accumulating festival recognition of AI-native feature and short-film work — but *Dreams of Violets* is the first Tier 1 festival selection for a feature-length, fully-AI-generated work.↵
40. Variety, “‘Watch the Skies,’ Swedish UFO Feature Film Dubbed Entirely With AI, Sets USA Distribution Deal.” <https://variety.com/2025/film/news/watch-the-skies-us-theatrical-release-ai-dubbing-1236343110/>. *Dream Machine* Issue 5.↵
41. *Cybernews*, “Run to the West — South Korea’s first AI film tests the soul of cinema.” <https://cybernews.com/entertainment/korean-cinema-run-to-the-west-ai/>. *Dream Machine* Issue 5.↵
42. *Broadcast Pro Middle East*, “Tunisian filmmaker wins \$1 million AI Film Award for ‘Lily’.” <https://www.broadcastpromeme.com/news/tunisian-filmmaker-wins-1-million-ai-film-award-for-lily/>. *Dream Machine* Issue 14.↵
43. Variety, “AI Drama ‘Humans in the Loop’ Receives Film Independent’s Sloan Distribution Grant, Enters Oscar Race.” <https://variety.com/2025/film/news/ai-drama-humans-in-the-loop-oscar-race-1236582975/>. *Dream Machine* Issue 8.↵
44. *PC Gamer*, “Palworld studio Pocketpair says its new publishing division won’t handle games that use generative AI: ‘We don’t believe in it.’” <https://www.pcgamer.com/software/ai/palworld-studio-pocketpair-says-its-new-publishing-division-wont-handle-games-that-use-generative-ai-we-dont-believe-in-it/>. *Dream Machine* Issue 4.↵
45. *Niche Gamer*, “Larian Studios backs off from gen AI.” *Dream Machine* Issue 14.↵

46. Decrypt, "Warhammer 40,000 Maker Games Workshop Rules Out Generative AI." *Dream Machine* Issue 14.↵
47. Niche Gamer, "Manor Lords publisher Hooded Horse won't work with devs using gen AI." *Dream Machine* Issue 14.↵
48. *gamesindustry.biz*, "RuneScape maker Jagex says it will never use generative AI to make in-game content." *Dream Machine* Issue 16.↵
49. *GamesRadar*, "Wallace and Gromit creator says beloved animation studio Aardman will 'embrace the technology' of AI, but will be 'very cautious not to lose our values'." <https://www.gamesradar.com/entertainment/animation-movies/wallace-and-gromit-creator-says-beloved-animation-studio-aardman-will-embrace-the-technology-of-ai-but-will-be-very-cautious-not-to-lose-our-values/>. *Dream Machine* Issue 11.↵
50. Variety, "Guillermo del Toro Says He'd 'Rather Die' Than Use Generative AI in His Films: 'Not Interested'." <https://variety.com/2025/film/news/guillermo-del-toro-rather-die-generative-ai-frankenstein-1236561316/>. *Dream Machine* Issue 5.↵
51. Variety, "Martin Scorsese Supports AI Company, Endorses AI for Storyboarding." <https://variety.com/2026/film/news/martin-scorsese-supports-ai-company-storyboard-movies-1236765037/>. Scorsese becomes advisor to Black Forest Labs (FLUX model); endorses AI storyboarding; Boots Riley and Guillermo del Toro respond critically. *Dream Machine* Issue 32.↵
52. *The Hollywood Reporter*, "Leonardo DiCaprio Says AI Can't Be Art Because 'There's No Humanity to It'." <https://www.hollywoodreporter.com/movies/movie-news/leonardo-dicaprio-ai-cant-be-art-no-humanity-1236445405/>. *Dream Machine* Issue 11.↵
53. *Daily Mail*, "Claire Foy says she has 'no interest' in seeing AI in films." <https://www.dailymail.co.uk/tvshowbiz/article-15454199/Claire-Foy-AI-films-sad-disappointed-people-future-Hollywood.html>. *Dream Machine* Issue 14.↵
54. NME, "Jenna Ortega says it's 'very easy to be terrified' of AI in filmmaking." <https://www.nme.com/news/jenna-ortega-says-its-very-easy-to-be-terrified-of-ai-in-filmmaking-3913926>. *Dream Machine* Issue 10.↵
55. Variety, "Chris Pratt Pitched Having an AI 'Actor' Star as the Villain in 'Mercy': 'I Don't Think That's a Good Idea at All'." <https://variety.com/2026/film/news/chris-pratt-ai-actor-villain-mercy-amazon-mgm-1236640460/>. *Dream Machine* Issue 16.↵
56. *PC Gamer*, "Todd Howard says AI can't replace human 'creative intention,' but it's part of Bethesda's 'toolset for how we build our worlds or check things'." <https://www.pcgamer.com/gaming-industry/todd-howard-says-ai-cant-replace-human-creative-intention-but-its-part-of-bethesdas-toolset-for-how-we-build-our-worlds-or-check-things/>. *Dream Machine* Issue 11.↵
57. *Wired*, "Amazon's House of David Used Over 350 AI Shots in Season 2. Its Creator Isn't Sorry." <https://www.wired.com/story/amazons-house-of-david-used-over-350-ai-shots-in-season-2-its-creator-isnt-sorry/>. *Dream Machine* Issue 7.↵
58. *GamesRadar*, "Battlefield 6 lead calls generative AI 'very seducing,' but says it was only used in the game's earliest stages 'to allow for more time and more space to be creative'." <https://www.gamesradar.com/games/battlefield/battlefield-6-lead-calls-generative-ai-very-seducing-but-says-it-was-only-used-in-the-games-earliest-stages-to-allow-for-more-time-and-more-space-to-be-creative/>. *Dream Machine* Issue 3.↵
59. *gamesindustry.biz*, "Witcher 3 and Cyberpunk 2077 director says AI can help, but not replace, creatives." <https://www.gamesindustry.biz/witcher-3-and-cyberpunk-2077-director-says-ai-can-help-but-not-replace-creatives>. *Dream Machine* Issue 9.↵
50. *GamesRadar*, "Aardman" op. cit.↵
61. *Dream Machine* Issue 29, on Sony's "all in on AI for games" announcement.↵
52. Variety, "AI Dominates Cannes Buzz as Filmmakers Grudgingly Accept It." <https://variety.com/2026/film/festivals/ai-cannes-2026-filmmakers-accept-1236748402/>; *Hollywood Reporter*, "At Cannes, filmmakers shift towards cautious acceptance of AI's inevitability."

<https://www.hollywoodreporter.com/business/business-news/cannes-2026-ai-acceptance-1236592488/>.
Dream Machine Issue 30.↵

53. *Variety*, “Is AI Basically Like Special Effects? Peter Jackson Seems to Think So.” <https://variety.com/2026/film/news/peter-jackson-ai-special-effects-1236748120/>. *Dream Machine* Issue 30.↵
54. *PC Gamer*, “Take-Two’s CEO says AI’s not in the business of making hits, ‘datasets by their very nature are backward looking’, but that doesn’t mean AI can’t be ‘super helpful’.” <https://www.pcgamer.com/games/take-two-ceo-ai-not-making-hits-backward-looking/>. *Business Insider*, “The CEO behind Grand Theft Auto says he’s pro AI — but the technology can’t make an original hit.” <https://www.businessinsider.com/take-two-ceo-strauss-zelnick-ai-original-hits-2026-5>. *Dream Machine* Issue 30.↵
55. Hasbro launches Sixth Wall, a new AI studio, with the CharacterOS governance layer underneath it, 3 June 2026. BusinessWire press release: <https://www.businesswire.com/news/home/20260603297922/en/>. *The Wrap*, “Hasbro AI Studio Sixth Wall Launch”: <https://www.thewrap.com/industry-news/business/hasbro-ai-studio-sixth-wall-launch/>. *Screen Daily*: <https://www.screendaily.com/news/hasbro-launches-subsidiary-ai-studio-sixth-wall-to-exploit-ip/5217380.article>. *BriefGlance* analysis (FTC inquiries into AI companions, Sixth Wall’s stated 13+/enterprise focus): <https://briefglance.com/articles/hasbros-sixth-wall-taming-the-ai-wild-west-with-behavioral-licensing>. Detailed treatment in Appendix L. CEO Roberta Thomson; CharacterOS encodes character canon, voice, personality, lore and safety limits at runtime; talent-participation pay-out structure for the original performers; initial 12-character catalogue; pilot grew out of an internal generative-AI Ouija-board experiment that logged hundreds of thousands of player interactions in its first few days.↵
56. Hasbro launches Sixth Wall, a new AI studio, with the CharacterOS governance layer underneath it, 3 June 2026. BusinessWire press release: <https://www.businesswire.com/news/home/20260603297922/en/>. *The Wrap*, “Hasbro AI Studio Sixth Wall Launch”: <https://www.thewrap.com/industry-news/business/hasbro-ai-studio-sixth-wall-launch/>. *Screen Daily*: <https://www.screendaily.com/news/hasbro-launches-subsidiary-ai-studio-sixth-wall-to-exploit-ip/5217380.article>. *BriefGlance* analysis (FTC inquiries into AI companions, Sixth Wall’s stated 13+/enterprise focus): <https://briefglance.com/articles/hasbros-sixth-wall-taming-the-ai-wild-west-with-behavioral-licensing>. Detailed treatment in Appendix L. CEO Roberta Thomson; CharacterOS encodes character canon, voice, personality, lore and safety limits at runtime; talent-participation pay-out structure for the original performers; initial 12-character catalogue; pilot grew out of an internal generative-AI Ouija-board experiment that logged hundreds of thousands of player interactions in its first few days.↵
57. Spotify and Universal Music Group announce landmark linked recorded-music and publishing licensing agreements for fan-made AI covers and remixes at Spotify Investor Day, 21 May 2026. Spotify Newsroom: <https://newsroom.spotify.com/2026-05-21/universal-music-group-spotify-licensing-agreements-fan-made-covers-remixes/>. *TechCrunch*: <https://techcrunch.com/2026/05/21/spotify-and-universal-music-strike-deal-allowing-fan-made-ai-covers-and-remixes/>. *Medianama*: <https://www.medianama.com/2026/05/223-spotify-umg-fans-create-licensed-ai-covers-remixes/>. Three-pillar consent / credit / compensation framework; opt-in catalogue; Premium add-on; new royalty stream; UMG settlement with Udio in October 2025 is the architectural ancestor. Detailed treatment in Appendix L. Cross-reference in Chapter 5 and Chapter 6.↵
58. Larian Studios policy framing on the next *Divinity*, January 2026; same source as [^07thestudios-33].↵
59. World Labs, “Bringing Marble to Life.” <https://www.worldlabs.ai/case-studies/bringing-marble-to-life>. *Dream Machine* Issue 7, “Editor’s Pick: Marble by WorldLabs goes on public release,” 13 November 2025.↵
70. For a working primer on Gaussian splatting in the post-Marble era, see *Radiance Fields*, “World Labs Formally Launches Marble, A Generative World Model.” <https://radiancefields.com/world-labs-formally-launches-marble-a-generative-world-model>.↵
71. DreamLab AI Collective, beta participation in Marble, October–November 2025. Referenced in *Dream Machine* Issue 7: “DreamLab have been part of the beta testing for this over the last few months and it’s very neat.”↵
72. SuperSplat (PlayCanvas), open-source Gaussian splat editor, regular updates through 2025–26. *Dream Machine* Issue 1: “PlayCanvas open sources SOG — WebP for 3D Gaussian Splatting”; Issue 7 / Issue 11 on SuperSplat v2 updates.↵

73. Sony Pictures' use of Marble in Virtual Production: https://www.linkedin.com/posts/brent-liang_tech-media-launch-ugcPost-7394911181091692546-TyUz. *Dream Machine* Issue 8. ↵
74. Disney "300,000 poses in an instant" livestream, March 2026. *Dream Machine* Issue 23. ↵
75. Netflix + Eyaline, *Vista4D*: 4D point clouds from live-action. *Dream Machine* Issue 27. ↵
76. Google DeepMind, "Genie 3: A new frontier for world models." <https://deepmind.google/blog/genie-3-a-new-frontier-for-world-models/>. Project Genie roll-out to AI Ultra subscribers in the U.S.: <https://blog.google/innovation-and-ai/models-and-research/google-deepmind/project-genie/>. *Dream Machine* Issue 3 (initial announcement) and Issue 17 (broader availability). ↵
77. Meta, "WorldGen — text-to-immersive-3D-worlds research update." <https://www.facebook.com/LifeAtMeta/videos/research-update-worldgen-text-to-immersive-3d-worlds/1879077432692421/>. *Dream Machine* Issues 9, 11. ↵
78. Tencent, "HY World 1.5" announcement: <https://x.com/TencentHunyuan/status/2001170499133653006>. *Dream Machine* Issue 12. ↵
79. SpAtItal, *ECHO* spatial foundation model. <https://www.spatial.ai/>. *Dream Machine* Issue 12. ↵
30. Stanford AI Lab, *Wonderzoom*: Multi-Scale 3D World Generation. <https://wonderzoom.github.io/>. *Dream Machine* Issue 14. ↵
81. OpenArt, *Worlds* product launch, March 2026. *Dream Machine* Issue 21. ↵
32. NVIDIA SANA-WM, 2.6B open-source world model with 60-second video generation and camera control, May 2026. <https://huggingface.co/collections/nvidia/sana-wm>. *Dream Machine* Issue 30. ↵
33. Odyssey, "Introducing Starchild-1, the first real-time multimodal world model." <https://odyssey.ml/introducing-starchild-1>. *Dream Machine* Issue 30. ↵
34. Apple Machine Learning Research, "Headsup: a large-scale high-quality 3D Gaussian head reconstruction from multi-view captures." <https://machinelearning.apple.com/research/apple-headsup-3d-gaussian-head>. *Dream Machine* Issue 30. ↵
35. WorldLens VR, "AI-powered 3D depth for Google Street View on Quest." <https://www.uploadvr.com/worldlens-vr-quest-street-view-3d-depth/>. *Dream Machine* Issue 30. ↵
36. Luma AI, *UNI-1* launch, March 2026. *Dream Machine* Issue 22, "Editor's Pick: When worlds become instant, the race shifts to better thinking." ↵
87. ByteDance Seedance 2.0 in CapCut/Dreamina, March 2026. *Dream Machine* Issue 22. ↵
38. *Spark 2.0*, open-source Gaussian-splat streaming framework, April 2026. *Dream Machine* Issue 25. ↵
39. Radiance Fields, "Apple Confirms that it's Gaussian Splatting that powers their personas." <https://radiancefields.com/apple-confirms-personas-use-gaussian-splatting>. *Dream Machine* Issue 5. ↵
30. *Video Games Chronicle*, "It honestly sucks': Fans think Call of Duty: Black Ops 7 is filled with generative AI art." <https://www.videogameschronicle.com/news/it-honestly-sucks-fans-think-call-of-duty-black-ops-7-is-filled-with-generative-ai-art/>. *Video Games Chronicle*, "Ubisoft says AI-generated art in Anno 117 was a placeholder which 'slipped through our review process'." <https://www.videogameschronicle.com/news/ubisoft-says-ai-generated-art-in-anno-117-was-a-placeholder-which-slipped-through-our-review-process/>. *Polygon*, "Fortnite chapter 7 kicks off new controversy over AI art." <https://www.polygon.com/fortnite-chapter-7-season-1-generative-ai-art-epic-games/>. *Dream Machine* Issues 8, 10. ↵
91. NVIDIA + Stanford, *NitroGen*. <https://nitrogen.minedojo.org/>. *Dream Machine* Issue 13. ↵
32. DeepMind, "SIMA 2: An Agent that Plays, Reasons, and Learns With You in Virtual 3D Worlds." <https://deepmind.google/blog/sima-2-an-agent-that-plays-reasons-and-learns-with-you-in-virtual-3d-worlds/>. *Dream Machine* Issue 8. ↵
33. ComfyUI Blog, "Ubisoft La Forge Open-Sources the CHORD Model and ComfyUI Nodes for End-to-End PBR Material Generation." <https://blog.comfy.org/p/ubisoft-open-sources-the-chord-model>. *Dream Machine* Issue 11. ↵
34. *Video Games Chronicle*, "The future of gaming, or 'just a tool'? Hands-on with Teammates, Ubisoft's ambitious voice AI tech demo." <https://www.videogameschronicle.com/features/the-future-of-gaming-or-just-a-tool-hands-on-with-teammates-ubisofts-ambitious-voice-ai-tech-demo/>. *Dream Machine* Issue 9. ↵

95. YouTube Playables Builder, closed-beta announcement: <https://www.youtube.com/playablesbuilder/>. *Dream Machine* Issue 12. ↩
96. Unity AI Open Beta, in-editor AI suite, May 2026. *Dream Machine* Issue 28. ↩
97. Korin AI, “trained with African datasets, built by Africans,” May 2026. *Dream Machine* Issue 27. ↩
98. *Creative Boom*, “Adobe is putting AI in everything everywhere all at once.” <https://www.creativeboom.com/news/adobe-is-putting-ai-in-everything-everywhere-all-at-once/>. *Dream Machine* Issue 5, “Editor’s Pick,” 31 October 2025. ↩
99. Adobe, “Adobe MAX 2025: Firefly Foundry.” <https://news.adobe.com/news/2025/10/adobe-max-2025-firefly-foundry>. ↩
100. Adobe, “Adobe MAX 2025: Firefly.” <https://news.adobe.com/news/2025/10/adobe-max-2025-firefly>. ↩
101. Adobe, “Adobe MAX 2025: Express AI Assistant.” <https://news.adobe.com/news/2025/10/adobe-max-2025-express-ai-assistant>. ↩
102. *Wired*, “Adobe’s ‘Corrective AI’ Can Change the Emotions of a Voice-Over” and accompanying Adobe Sneaks 2025 coverage. <https://www.wired.com/story/adobe-max-sneaks-2025-corrective-ai/>. Project list compiled from MAX keynote and *Dream Machine* Issue 5 coverage. ↩
103. PYMNTS, “Adobe Lets Users Design and Edit Using ChatGPT.” <https://www.pymnts.com/artificial-intelligence-2/2025/adobe-lets-users-design-and-edit-using-chatgpt/>. Adobe blog: “Edit images, designs, and PDFs right inside ChatGPT.” <https://blog.adobe.com/en/publish/2025/12/10/edit-photoshop-chatgpt>. *Dream Machine* Issue 12. ↩
104. Adobe Premiere Object Mask tool: https://www.linkedin.com/posts/robdewinter_ok-this-is-going-to-save-a-lot-of-time-in-ugcPost-7421617551690063872-vKmB. *Dream Machine* Issue 16. ↩
105. Adobe blog, “Sundance Film Festival 2026: Creativity, Community & Power of Storytelling.” <https://blog.adobe.com/en/publish/2026/01/20/sundance-film-festival-2026-creativity-community-power-storytelling>. *Dream Machine* Issue 16. ↩
106. Adobe Summit 2026, “agentic creative intelligence” keynote. *Dream Machine* Issue 26. ↩
107. After Effects AI animation features through late 2025: *Dream Machine* Issue 9, “AI video is finally animatable inside After Effects.” https://www.linkedin.com/posts/thisisdoug_ai-aivideo-animation-ugcPost-7399512745924067330-Aldk. ↩
108. *Dream Machine* Issue 21, “Editor’s Pick: Adobe and NVIDIA Just Raised the Stakes for Creative AI,” 19 March 2026. ↩
109. NVIDIA + Google Cloud creative-AI infrastructure deal, March 2026. *Dream Machine* Issue 21. ↩
110. Hugging Face and Google Cloud partnership announcement: https://www.linkedin.com/posts/julienchaumond_i-am-super-excited-to-announce-that-hugging-activity-7396177403972276225-CuMM. *Dream Machine* Issue 8. ↩
111. *EdTech Innovation Hub*, “Meta and Hugging Face launch OpenEnv to advance open-source agentic development.” <https://www.edtechinnovationhub.com/news/meta-and-hugging-face-launch-openenv-to-advance-open-source-agentic-development>. *Dream Machine* Issue 5. ↩
112. Anthropic / Blender Foundation patronage, May 2026. *Dream Machine* Issue 27. ↩
113. TechCrunch, “Anthropic launches interactive Claude apps, including Slack and other workplace tools.” <https://techcrunch.com/2026/01/26/anthropic-launches-interactive-claude-apps-including-slack-and-other-workplace-tools/>. *Dream Machine* Issue 16. ↩
114. Spotify–Anthropic integration, May 2026. *Dream Machine* Issue 27. ↩
115. *Music Business Worldwide*, “Splice inks ‘Responsible AI’ deal with ElevenLabs.” <https://www.musicbusinessworldwide.com/splice-elevenlabs-responsible-ai-deal/>. *Dream Machine* Issue 30. ↩
116. *Adweek*, “Netflix ad tools could see ‘agentic AIs talking to each other’.” <https://www.adweek.com/media/netflix-ad-tools-agentic-ais-talking-to-each-other/>. *Dream Machine* Issue 30. ↩
117. *Fortune*, “AI startup Viktor raises \$75 million to put a virtual ‘coworker’ in Slack and Teams.” <https://fortune.com/2026/05/19/ai-startup-viktor-75-million-virtual-coworker-slack-teams/>. *Dream Machine* Issue 30. ↩

Machine Issue 30.

18. *MarTech Series*, “WPP continues AI overhaul with \$400-million Google partnership.” <https://martechseries.com/predictive-ai/ai-platforms-machine-learning/google-and-spotify-alum-launch-epiminds-with-6-6m-to-build-marketing-teams-for-the-ai-era/>. *Dream Machine Issue 3.*
19. *Campaign Brief*, “WPP launches AI-powered marketing platform WPP Open Pro.” <https://campaignbrief.com/wpp-launches-ai-powered-marketing-platform-wpp-open-pro/>. *Dream Machine Issue 5.*
20. *Digiday*, “WPP expands AI capabilities to boost brand performance with Slightly partnership.” <https://digiday.com/media-buying/agencies-continue-to-expand-ai-capabilities-to-boost-brand-performance/>. *Dream Machine Issue 6.*
21. WPP and Google Earth AI consumer-journey project, April 2026. *Dream Machine Issue 27.*
22. Google I/O 2026 announcement block: Gemini Omni <https://blog.google/technology/google-deepmind/gemini-omni/>, Antigravity <https://antigravity.google/>, Google Flow <https://flow.google/>, Gemini Spark <https://blog.google/technology/developers/gemini-spark/>, Project Genie + Street View <https://deepmind.google/discover/blog/project-genie-street-view/>. *Dream Machine Issue 30*, “Editor’s Pick — Google I/O 2026,” 21 May 2026.
23. Google Labs, “Infinite Scaler.” <https://blog.google/technology/google-labs/infinite-scaler/>. *Dream Machine Issue 30.*
24. Google DeepMind, “SynthID — 100 billion watermarks, partner ecosystem.” <https://deepmind.google/discover/blog/synthid-100-billion-watermarks-partners/>. *Dream Machine Issue 30.*
25. *SiliconAngle*, “Higgsfield raises \$80M on \$1.3B valuation to scale AI video platform.” <https://siliconangle.com/2026/01/15/higgsfield-raises-80m-1-3b-valuation-scale-ai-video-platform/>. *Dream Machine Issue 15.*
26. *36kr*, “AI Video Unicorn Higgsfield: Earns \$200M in 9 Months by ‘Serving’ Social Media Marketers.” <https://eu.36kr.com/en/p/3650517574312323>. *Dream Machine Issue 16.*
27. *TechCrunch*, “Synthesia hits \$4B valuation, lets employees cash in.” <https://techcrunch.com/2026/01/26/synthesia-hits-4b-valuation-lets-employees-cash-in/>. *Dream Machine Issue 16.*
28. *Sifted*, “Synthesia rejects \$3bn Adobe acquisition offer.” <https://sifted.eu/articles/synthesia-acquisition-offer>. *Dream Machine Issue 5.*
29. ElevenLabs \$500m ARR reporting, April 2026. *Dream Machine Issue 25.*
30. Runway product cycle: Gen-4.5 (December 2025), Gen-4.5 Image-to-Video (January 2026), Workflows, Story Panels, Characters API, Apps for Advertising — *Dream Machine Issues 10, 14, 15, 16, 20*. Runway CEO on indie films vs. blockbusters: *Dream Machine Issue 26.*
31. Runway, “Runway Japan.” <https://runwayml.com/blog/runway-japan>. *Dream Machine Issue 30.*
32. For the running ledger of new creative-AI products through 2025–26, see *Dream Machine Issues 1–30* archive.
33. ComfyUI, “We raised \$17 million to build an OS for Creative AI.” https://www.linkedin.com/posts/comfyui_we-raised-17-million-to-build-an-os-for-ugcPost-7373743341236236288-wkCc. *Dream Machine Issue 1.*
34. ComfyUI \$500M valuation, May 2026. *Dream Machine Issue 27.*
35. Google Pomelli launch: <https://x.com/GoogleLabs/status/1983204018567426312>. *Dream Machine Issue 5.*
36. Google AI Studio app gallery: <https://x.com/GoogleAISTudio/status/1982121563785949255>. Google Labs Opal expansion: <https://blog.google/technology/google-labs/opal-expansion/>. Project Genie: <https://blog.google/innovation-and-ai/models-and-research/google-deepmind/project-genie/>. *Dream Machine Issues 5, 17.*
37. Lovable for classrooms: <https://lovable.dev/classroom>. *Dream Machine Issue 11.*

38. Adobe Express AI Assistant: <https://news.adobe.com/news/2025/10/adobe-max-2025-express-ai-assistant>. *Dream Machine* Issue 5.↵
39. Hugging Face platform expansion through 2025–26.↵
40. Google blog, “Sundance Institute AI Education.” <https://blog.google/company-news/outreach-and-initiatives/google-org/sundance-institute-ai-education/>. *Dream Machine* Issue 15.↵
41. Adobe Ignite Day at Sundance: *Adobe blog, Sundance Film Festival 2026*. <https://blog.adobe.com/en/publish/2026/01/20/sundance-film-festival-2026-creativity-community-power-storytelling>. *Dream Machine* Issue 16.↵
42. Google’s \$40bn investment in Anthropic, May 2026. *Dream Machine* Issue 27.↵
43. UK Government, “Free AI training for all.” <https://www.gov.uk/government/news/free-ai-training-for-all-as-government-and-industry-programme-expands-to-provide-10-million-workers-with-key-ai-skills-by-2030>. *Dream Machine* Issue 16.↵
44. CNBC, “People with ADHD, autism, dyslexia say AI agents are helping them succeed at work.” <https://www.cnn.com/2025/11/08/adhd-autism-dyslexia-jobs-careers-ai-agents-success.html>. *Dream Machine* Issue 7.↵
45. University of Wisconsin-Stout, “AI Reshaping Industry: New UW-Stout Course Sets AI-Use as Baseline Competency in Filmmaking.” <https://www.uwstout.edu/about-us/news-center/ai-reshaping-industry-new-uw-stout-course-sets-ai-use-baseline-competency-filmmaking>. *Dream Machine* Issue 16.↵
46. Doug McGinness on LinkedIn, late 2025, in the same post. *Dream Machine* Issue 9.↵
47. Snap Newsroom, “Snapchat Gen Z AI Creativity Research 2026.” <https://newsroom.snap.com/snapchat-gen-z-ai-creativity-research-2026>. *Dream Machine* Issue 30.↵
48. *Dream Machine* Issue 13, “Editor’s Pick: The Year of the Orchestrator,” 9 January 2026.↵
49. *Dream Machine* Issue 29, May 2026, reporting on Sony’s 49-agent / 72-skill multi-agent game-development team.↵
50. Bloomberg, “AI Changed Chess. Grandmasters Now Win With Unpredictable Moves,” 27 March 2026. <https://www.bloomberg.com/news/articles/2026-03-27/ai-changed-chess-grandmasters-now-win-with-unpredictable-moves>. *Dream Machine* Issue 23. The chess analogy is developed in Chapter 15’s *Age of the Why* section.↵
51. Anthropic blog content on agent deployment patterns, Q1 2026.↵
52. Sundance Institute, “Centering the Artist: Why We’re Launching the AI Literacy Initiative.” <https://www.sundance.org/blogs/centering-the-artist-why-were-launching-the-ai-literacy-initiative/>. *Dream Machine* Issue 16.↵
53. Sundance Institute, *op. cit.*↵
54. Sundance Story Forum 2026 sessions on legal toolkits for producers using AI. *Dream Machine* Issue 16.↵
55. Google blog, “Sundance Institute AI Education.” <https://blog.google/company-news/outreach-and-initiatives/google-org/sundance-institute-ai-education/>. *Dream Machine* Issue 15.↵
56. McKinsey & Company, “What AI could mean for film and TV production and the industry’s future.” <https://www.mckinsey.com/industries/technology-media-and-telecommunications/our-insights/what-ai-could-mean-for-film-and-tv-production-and-the-industrys-future>. *Dream Machine* Issue 16.↵
57. Metro, “Prince of Persia remake and five more games cancelled as Ubisoft focuses on AI.” <https://metro.co.uk/2026/01/21/prince-persia-remake-five-games-cancelled-ubisoft-focuses-ai-26431926/>. *Dream Machine* Issue 15.↵
58. PC Gamer, “Square Enix, makers of Final Fantasy, aims to have AI doing 70% of its QA work by the end of 2027.” <https://www.pcgamer.com/gaming-industry/square-enix-aims-to-have-ai-doing-70-percent-of-its-qa-work-by-the-end-of-2027/>. *Dream Machine* Issue 7.↵
59. Eurogamer, “Falcom is the latest developer to buy into the AI hype machine.” <https://www.eurogamer.net/falcom-is-the-latest-developer-to-buy-into-the-ai-hype-machine>. *Dream Machine* Issue 12.↵
50. NDTV Profit, “Don’t Expect AI To Invent the Next ‘Grand Theft Auto’, Says Take-Two CEO Strauss Zelnick.” <https://www.ndtvprofit.com/technology/dont-expect-ai-to-invent-the-next-grand-theft-auto-says-take-two>

ceo-strauss-zelnick. *Dream Machine* Issue 6.↵

61. *Dream Machine* Issue 21, on Spielberg's public position on AI.↵
52. *Dream Machine* Issues 25, 28, on Steven Soderbergh's AI work.↵
53. *Digiday*, "Independent agencies face new frontier as agency-in-a-box tools democratize creativity." <https://digiday.com/marketing/independent-agencies-face-new-frontier-as-agency-in-a-box-tools-democratize-creativity/>. *Dream Machine* Issues 6, 14.↵
54. *Digiday*, "AI agent developers have become adland's in-demand role." <https://digiday.com/marketing/ai-agent-developers-have-become-adlands-in-demand-role/>. *Dream Machine* Issue 7.↵
65. PYMNTS, "AI Content Is Par For The Course With PGA Tour's Expanded AWS Partnership." <https://www.pymnts.com/artificial-intelligence-2/2026/ai-content-is-par-for-the-course-with-pga-tours-expanded-aws-partnership/>. *Dream Machine* Issue 15.↵
56. *The Verge*, "Oreo-maker Mondelez will use AI for TV ads next year." <https://www.theverge.com/news/806047/mondelez-ai-generated-ads>. *Dream Machine* Issue 5.↵
67. *Digiday*, "Avocados From Mexico turns to AI to advertise around the Super Bowl instead of a TV buy." <https://digiday.com/marketing/avocados-from-mexico-turns-to-ai-to-advertise-around-the-super-bowl-instead-of-a-tv-buy/>. *Dream Machine* Issue 15.↵
58. *Dream Machine* Issue 26.↵
59. Reuters Institute, "AI adoption by UK journalists and their newsrooms: surveying applications, approaches, and attitudes." <https://reutersinstitute.politics.ox.ac.uk/ai-adoption-uk-journalists-and-their-newsrooms-surveying-applications-approaches-and-attitudes>. *Dream Machine* Issue 9.↵
70. *Digiday*, "Daily Mail says Google AI Overviews have killed click-throughs." <https://digiday.com/media/daily-mail-says-google-ai-overviews-have-killed-click-throughs/>. *Dream Machine* Issue 7.↵
71. *Digiday*, "How The Times is using AI to model synthetic focus groups from human audiences." <https://digiday.com/media/how-the-times-is-using-ai-to-model-synthetic-focus-groups-from-human-audiences/>. *Dream Machine* Issue 6.↵
72. *TechBullion*, "Why the future belongs to multi-skilled leaders." <https://techbullion.com/playing-the-long-game-with-a-portfolio-career-why-the-future-belongs-to-multi-skilled-leaders/>. *Dream Machine* Issue 9.↵
73. Anthropic Skills framework via Claude Code, reported through *Dream Machine* Issues 11, 16, 29.↵
74. *Forbes*, "AI Is Changing How Creators Work And Earn." <https://www.forbes.com/sites/kolawolesamueladebayo/2025/12/22/how-ai-is-changing-how-creators-work-and-earn/>. *Dream Machine* Issue 13.↵
75. *Variety*, "AI Creator Behind Viral 'Deadpool,' 'Harry Potter' Christmas Clip Made His Film in a Ukrainian Bomb Shelter." <https://variety.com/2026/digital/news/ai-video-deadpool-harry-potter-andrii-daniels-1236624632/>. *Dream Machine* Issue 16.↵
76. *Dream Machine* Issue 29 reportage of Tiny Grandma stop-motion content being wrongly flagged as AI by YouTube's automated detection, May 2026.↵
77. *Dream Machine* Issue 23, April 2026, reporting death threats against Eline Van der Velden following Tilly Norwood's continuing public role.↵
78. *Hollywood Reporter*, "Bobby Berk Says AI Will Make Reality TV & 'Verifiably Human Content' More Valuable." <https://www.hollywoodreporter.com/tv/tv-news/bobby-berk-ai-reality-tv-1236592920/>. *Dream Machine* Issue 30.↵
79. *Digital Music News*, "Instagram Chief Says We Should 'Fingerprint Real Media' Instead of Tracking and Disclosing AI Slop." <https://www.digitalmusicnews.com/2026/01/05/instagram-chief-ai-slop-comments/>. *WebProNews*, "Instagram Head Warns AI Images Erode Trust, Calls for Verification Standards." <https://www.webpronews.com/instagram-head-warns-ai-images-erode-trust-calls-for-verification-standards/>. *Dream Machine* Issue 13.↵
30. *Digital Music News*, "AI-Generated Far-Right Hate Songs Aren't Just a Problem in the US — Now They're Spreading Across Europe Too." <https://www.digitalmusicnews.com/2025/11/09/ai-generated-hate-songs-dutch-spotify-charts/>. *Dream Machine* Issue 7.↵

81. Google DeepMind SynthID watermark roll-out across Veo, Lyria and Imagen products. *Dream Machine* Issues 11, 12. ↵
82. Google DeepMind, “Verify Google AI-generated videos in the Gemini app.” https://www.linkedin.com/posts/googledeepmind_verify-google-ai-generated-videos-in-the-activity-7407748300688478208-flgW. *Dream Machine* Issue 12; broader coverage in *SmartBrief*, “Google’s Gemini can now spot AI-generated videos.” <https://newsletter.smartbrief.com/sharedSummary/index.jsp?briefId=40A39351-5419-4681-94DF-31A53480B698&issueId=58E986AD-821F-422E-9E34-3386E0E2272B©Id=2DB8E453-8E83-416C-949B-44751F252A8D>. *Dream Machine* Issue 13. ↵
83. Google DeepMind, “SynthID — 100 billion watermarks, expanding to partner ecosystems including OpenAI, ElevenLabs and Kakao.” <https://deepmind.google/discover/blog/synthid-100-billion-watermarks-partners/>. *Dream Machine* Issue 30. ↵
84. YouTube announces automatic labelling of AI-generated videos, June 2026. *Dream Machine* Issue 31. ↵
85. LinkedIn announces it will limit the reach of AI-generated content on its feed, June 2026. *Dream Machine* Issue 31. ↵
86. Hasbro’s Sixth Wall launch and the CharacterOS governance layer, 3 June 2026. BusinessWire: <https://www.businesswire.com/news/home/20260603297922/en/>. *The Wrap*: <https://www.thewrap.com/industry-news/business/hasbro-ai-studio-sixth-wall-launch/>. Detailed treatment in Appendix L. Cross-reference Chapter 7 §“Position Five” and Chapter 6 §“The three Cs”. ↵
87. Spotify–Universal Music Group landmark licensing agreements, 21 May 2026. Spotify Newsroom: <https://newsroom.spotify.com/2026-05-21/universal-music-group-spotify-licensing-agreements-fan-made-covers-remixes/>. *TechCrunch*: <https://techcrunch.com/2026/05/21/spotify-and-universal-music-strike-deal-allowing-fan-made-ai-covers-and-remixes/>. Detailed treatment in Appendix L and footnote ^37a in Chapter 6. ↵
88. *Dream Machine* Issues 23, 27 reportage on Taylor Swift’s voice/image trademark filings. ↵
89. *Lawyer Monthly*, “Matthew McConaughey Draws a Line to Protect His Voice and Image From AI.” <https://www.lawyer-monthly.com/2026/01/matthew-mcconaughey-protects-voice-image-ai/>. *Dream Machine* Issue 15. ↵
90. *Adweek*, “Meet the \$1.3 Billion Startup Behind Madonna and Will Smith’s AI Video.” <https://www.adweek.com/media/higgsfield-ai-marketing-startup/>. *Dream Machine* Issue 16. ↵
91. *Rolling Stone*, “The Rolling Stones Release New Single ‘In the Stars’ — With a Music Video De-Aging the Rockers Courtesy of AI.” <https://www.rollingstone.com/music/music-news/rolling-stones-in-the-stars-ai-de-aging-video-1235142200/>. *Hollywood Reporter*, “‘South Park’ Creators’ AI Company Made The Rolling Stones Young Again for ‘In The Stars’ Music Video.” <https://www.hollywoodreporter.com/tv/tv-news/south-park-creators-ai-rolling-stones-in-the-stars-1236592855/>. *Dream Machine* Issue 30. ↵
92. *Variety*, “George Clooney Says AI Actors Will Face the ‘Same Problem We Have’ in Hollywood: ‘Making a Star Is Not So Easy’.” <https://variety.com/2025/scene/columns/george-clooney-ai-actors-movie-stars-1236579661/>. *Dream Machine* Issue 7. ↵
93. *Deadline*, “AI Documentary Director Insists Jeremy Renner Agreed To Narrate Movie As ‘Hawkeye’ Star Threatens ‘Multi-Millions’ Lawsuit.” <https://deadline.com/2025/11/jeremy-renner-lawsuit-threat-ai-movie-1236611830/>. *Dream Machine* Issue 7. ↵
94. *Variety*, “Cate Blanchett Co-Founds RSL Media, a Non-Profit to Address Consent Around AI Usage including creative work, name, image and likeness.” <https://variety.com/2026/film/news/cate-blanchett-rsl-media-ai-consent-1236748255/>. *Dream Machine* Issue 30. ↵
95. *Bloomberg*, “Apple Acquires Key Talent & Patents Behind AI Avatar Company ‘Animato’.” <https://www.bloomberg.com/news/articles/2026-05-19/apple-acquires-animato-ai-avatar-talent-patents>. *Dream Machine* Issue 30. ↵
96. Estate of Ozzy Osbourne announces the late artist will “live on as an AI avatar that talks with fans,” June 2026. *Dream Machine* Issue 31. ↵
97. *Complete Music Update*, “Johnny Cash estate uses ELVIS Act to sue Coke over tribute act ad soundtrack.” <https://completemusicupdate.com/johnny-cash-estate-uses-elvis-act-to-sue-coke-over-tribute-act-ad-soundtrack/>. *Dream Machine* Issue 9. ↵

98. *The Verge*, “New York’s new law forces advertisers to say when they’re using AI avatars.” <https://www.theverge.com/news/842848/new-york-law-ai-advertisements-sag-aftra-labor>. *Dream Machine* Issue 11. ↩
99. *Fast Company*, “Governments around the world are considering bans on Grok’s app over AI sexual image scandal.” <https://www.fastcompany.com/91474131/governments-around-the-world-are-considering-bans-on-groks-app-over-ai-sexual-image-scandal>. *Dream Machine* Issue 14. ↩
100. Cannes AI Disclosure Standard, launched May 2026. *Dream Machine* Issue 29. ↩
101. *Dream Machine* Issue 28, May 2026, reporting on the Academy of Motion Picture Arts and Sciences’ “You must be human to win” rule update. ↩
102. *The Hollywood Reporter*, “Emmys Set AI Guidance.” <https://www.hollywoodreporter.com/tv-tv-news/emmys-ai-guidelines-2026-awards-1236468434/>. *Dream Machine* Issue 14. ↩
103. SAG-AFTRA negotiation timeline through *Dream Machine* Issues 7, 12, 15, 20, 26, 29. ↩
104. *Variety*, “AI Music Suno Funding Round \$400 Million \$5.4 Billion Valuation.” <https://variety.com/2026/digital/news/ai-music-suno-funding-round-400-million-5-4-billion-valuation-1236765727/>. *TechCrunch*, “Still Facing Copyright Lawsuits, AI Music Generator Suno Raises Another \$400M.” <https://techcrunch.com/2026/06/03/still-facing-copyright-lawsuits-ai-music-generator-suno-raises-another-400m/>. *Dream Machine* Issue 32. ↩
105. Warner Music Group settled its copyright litigation against Suno and signed a licensing partnership; Universal Music Group and Sony Music remain in active litigation as of June 2026. *Dream Machine* Issue 32. For the structural context of the WMG settlement-as-licensing-table move, see Chapter 6 § “The three Cs”: *the lawsuits were never the destination; they were leverage to force a licensing table into existence*. ↩
106. *Marketing Week*, “You can’t dismiss AI ads as slop when they’re winning in testing.” <https://www.marketingweek.com/dismiss-ai-ads-winning-creative-effectiveness/>. *Dream Machine* Issues 8, 13. ↩
107. *The Drum*, “David Beckham Designs ‘Henchester United’ Chicken Coop in Lenovo Ad.” <https://www.thedrum.com/news/2026/05/18/david-beckham-henchester-united-chicken-coop-lenovo-ai-ad>. *Dream Machine* Issue 30. ↩
108. *Variety*, “AI Creator Behind Viral ‘Deadpool,’ ‘Harry Potter’ Christmas Clip Made His Film in a Ukrainian Bomb Shelter.” *op. cit.* *Dream Machine* Issue 16. ↩
109. PR Newswire, “From Apple TV Creative to AI Filmmaker: Hoyt Dwyer’s Animated Film To Compete at AI FilmFest Japan 2025.” *op. cit.* *Dream Machine* Issue 6. ↩
110. Google DeepMind, “Dear Upstairs Neighbors.” <https://blog.google/innovation-and-ai/models-and-research/google-deepmind/dear-upstairs-neighbors/>. *Dream Machine* Issue 16. ↩
111. *The Hollywood Reporter*, “‘Synthetic Sincerity’ by Marc Isaacs.” *op. cit.* *Dream Machine* Issue 8. ↩
112. *Variety*, “‘Watch the Skies,’ Swedish UFO Feature Film Dubbed Entirely With AI, Sets USA Distribution Deal.” *op. cit.* *Dream Machine* Issue 5. ↩
113. *Broadcast Pro Middle East*, “Tunisian filmmaker wins \$1 million AI Film Award for ‘Lily’.” *op. cit.* *Dream Machine* Issue 14. ↩
114. Sundance Institute AI Literacy Initiative emphasis on documentation: <https://www.sundance.org/blogs/centering-the-artist-why-were-launching-the-ai-literacy-initiative/>. *Dream Machine* Issue 16. ↩
115. *Dream Machine* Issue 5, “Industry Insights: Stealth, Shadow and Secret AI Users.” ↩
116. Azumo, “AI in Workplace Statistics 2025.” <https://azumo.com/artificial-intelligence/ai-insights/ai-in-workplace-statistics>. *Tech.co*, “Gen Z Most Likely Use AI Boss.” <https://tech.co/news/gen-z-most-likely-use-ai-boss>. *Dream Machine* Issue 5. ↩
117. *Exploding Topics*, “AI Workforce Research.” <https://explodingtopics.com/blog/ai-workforce-research>. *Dream Machine* Issue 5. ↩
118. *Forbes*, “AI Tools Flood Workplaces as Employees Face a Double Bind.” <https://www.forbes.com/sites/carolinecastrillon/2025/09/09/ai-tools-flood-workplaces-as-employees-face-a-double-bind/>. *Dream Machine* Issue 5. ↩

19. *Blog IDC Europe*, “Shadow AI: How Stealth Productivity Is Strangling Enterprise AI Adoption and Creating a Security Nightmare.” <https://blog-idceurope.com/shadow-ai-how-stealth-productivity-is-strangling-enterprise-ai-adoption-and-creating-a-security-nightmare/>. *Dream Machine Issue 5*. ↵
20. Enterprise-AI workforce tracking, late 2025. Aggregated in the Deep Dive companion piece *The Shadow AI Paradox in the Creative Industries*, drawing on Azumo’s *AI in Workplace Statistics 2025*, Tech.co’s Gen Z survey, and the IDC Europe shadow-AI security brief. *Dream Machine Issue 5*. ↵
21. *Hidden Cloud Explosion* analysis, IDC Europe, 2025. See *The Shadow AI Paradox in the Creative Industries*, §“The Epistemology and Scale of Shadow AI.” ↵
22. Shadow-AI security-incident statistics, 2025, aggregated in *The Shadow AI Paradox in the Creative Industries*, §“The Epistemology and Scale of Shadow AI”; underlying data via IBM *Cost of a Data Breach Report 2025* and IDC Europe. ↵
23. For the developer-community origins of the “AI for thee, but not for me” phrasing, and the full sectoral analysis of the paradox, see *The Shadow AI Paradox in the Creative Industries*, §“The Great Hypocrisy.” ↵
24. Survey of 1,100+ professional music creators, 2026, summarised in *Dynamics of Generative AI Adoption in the Creative Industries*, §“Music Production and Sound Recording,” and *The Shadow AI Paradox in the Creative Industries*, §“Sector-Specific Analysis.” ↵
25. WGA screenwriter survey, pre- and post-strike, reported in *Dynamics of Generative AI Adoption in the Creative Industries*, §“Screenwriting and the Post-Strike AI Boom.” ↵
26. Adobe Firefly milestone data, September 2023 – June 2025, in *Dynamics of Generative AI Adoption in the Creative Industries*, §“The Ubiquity of AI in Visual and Digital Arts.” *Dream Machine Issue 6*. ↵
27. Adobe quarterly financials, FY2025–FY2026; AI-first ARR growth reported in *Dream Machine Issue 21* and summarised in *Dynamics of Generative AI Adoption*. ↵
28. Adobe Firefly enterprise penetration metrics, in *Dynamics of Generative AI Adoption*. ↵
29. Adobe Stock submission analysis, 2024, in *Dynamics of Generative AI Adoption*. ↵
30. Adobe, “Inaugural Adobe Creators’ Toolkit Report,” October 2025. <https://news.adobe.com/news/2025/10/adobe-max-2025-creators-survey>. *Dream Machine Issue 6*. ↵
31. ChatGPT weekly-active-user disclosures, mid-2025; consolidated in *Dynamics of Generative AI Adoption*, §“General Purpose LLMs.” ↵
32. Gemini desktop-user growth, year-over-year, in *Dynamics of Generative AI Adoption*. ↵
33. Stanford AI Index Report 2025, global-sentiment chapter. Summarised in *Dynamics of Generative AI Adoption*, §“The Perception Gap.” ↵
34. YouGov 2024 multi-market AI sentiment survey, 17 countries. Summarised in *Dynamics of Generative AI Adoption*, §“The Perception Gap.” ↵
35. Quantic Foundry consumer-AI-in-gaming survey, 2025. Summarised in *Dynamics of Generative AI Adoption*, §“The Video Game Industry.” ↵
36. Game Developers Conference *State of the Game Industry* surveys, 2024–2026, sentiment vs. usage trend. Reported in *Dynamics of Generative AI Adoption*, §“The Video Game Industry.” ↵
37. O’Hare, J. (2026). *The Coordination Collapse: How AI Is Dissolving the Information Tax That Created Hierarchy — And What Replaces It*. DreamLab AI Consulting Ltd, April 2026. Preserved in full as Appendix K. The report draws on Dell’Acqua et al.’s preregistered Boston Consulting Group experiment with 758 consultants, Mollick’s *secret cyborgs* coinage, Edmondson and Lei’s psychological-safety findings, the IBM *Cost of a Data Breach Report 2025*, the Epoch AI / Ipsos KnowledgePanel survey of March 2026, and the CyberArk *hidden org chart* analysis. ↵
38. Epidemic Sound, *Future of the Creator Economy Report 2026* (survey of 3,000 professional creators across UK and US, June 2026); *Musically* commentary on the report and creator attitudes towards AI. *Dream Machine Issue 31*, Editor’s Pick, June 2026. ↵
39. *Game Developer*, “Subnautica owner Krafton outlines plans to transform into an ‘AI First’ company.” <https://www.gamedeveloper.com/business/subnautica-owner-krafton-outlines-plans-to-transform-into-an-ai-first-company>. *Dream Machine Issue 6*. ↵
40. *Dream Machine Issue 24*, April 2026, on the GTA VI publisher laying off its internal AI team. ↵

41. *Dream Machine* Issue 25, April 2026, on Disney layoffs including Marvel staff.↵
42. *SmartBrief*, “Meta to cut 10% of Reality Labs staff to focus on AI.” <https://newsletter.smartbrief.com/sharedSummary/index.jsp?briefId=40A39351-5419-4681-94DF-31A53480B698&issueId=025444D1-A590-46D8-B969-EF81DEE05228©Id=1B5F70D2-FFDA-4660-9CE9-047C9B16BF83>. *Dream Machine* Issue 14.↵
43. *Dream Machine* Issue 23, April 2026, on Scottish animation studio collapse.↵
44. *Metro*, “Prince of Persia remake and five more games cancelled as Ubisoft focuses on AI.” *op. cit. Dream Machine* Issue 15.↵
45. *The Guardian*, “AI is hitting UK harder than other big economies, study finds.” <https://www.theguardian.com/technology/2026/jan/26/ai-uk-jobs-us-japan-germany-australia>. *Dream Machine* Issue 16.↵
46. *The Economist*, “Investors expect AI use to soar. That’s not happening.” <https://www.economist.com/finance-and-economics/2025/11/26/investors-expect-ai-use-to-soar-thats-not-happening>. *Dream Machine* Issue 9.↵
47. *Dream Machine* Issue 24, April 2026, on OpenAI’s public-policy proposals around AI-driven economic disruption.↵
48. *The Economist*, “Job apocalypse? Humbug! AI is creating brand new occupations.” <https://www.economist.com/business/2025/12/14/job-apocalypse-humbug-ai-is-creating-brand-new-occupations>. *Dream Machine* Issue 12.↵
49. *Forbes*, “Vibe Coding — The In Demand AI Skill.” <https://www.forbes.com/sites/rachelwells/2025/11/06/the-in-demand-ai-skill-and-certifications-that-pays-up-to-220000/>. *Dream Machine* Issue 8.↵
50. U.K. Department for Business and Trade research on neurodiverse workers and AI assistants, autumn 2025. Reported via *CNBC*, “People with ADHD, autism, dyslexia say AI agents are helping them succeed at work.” <https://www.cnbc.com/2025/11/08/adhd-autism-dyslexia-jobs-careers-ai-agents-success.html>. *Dream Machine* Issue 7.↵
51. *CNBC*, *op. cit.*↵
52. *Dream Machine* Issue 7 secondary references.↵
53. Korin AI launch, May 2026. *Dream Machine* Issue 27.↵
54. *CNBC Africa*, “How AI is changing the landscape of the music industry in Africa.” <https://www.cnbc.com/2025/how-ai-is-changing-the-landscape-of-the-music-industry-in-africa>. *Dream Machine* Issue 5.↵
55. *BBC Future*, “Lights, camera, algorithm: Why Indian cinema is awash with AI.” <https://www.bbc.com/future/article/20251223-why-indian-cinema-is-awash-with-ai>. *Dream Machine* Issue 14.↵
56. *Dream Machine* Issue 25, April 2026, on Indonesia’s *Legenda Bertuah*.↵
57. *Broadcast Pro Middle East*, “Tunisian filmmaker wins \$1 million AI Film Award for ‘Lily’.” *op. cit. Dream Machine* Issue 14.↵
58. *Digiday*, “Avocados From Mexico turns to AI to advertise around the Super Bowl instead of a TV buy.” *op. cit. Dream Machine* Issue 15.↵
59. *Variety*, “AI Creator Behind Viral ‘Deadpool,’ ‘Harry Potter’ Christmas Clip Made His Film in a Ukrainian Bomb Shelter.” *op. cit. Dream Machine* Issue 16.↵
50. *Dream Machine* Issue 8 citing Andreessen Horowitz observations: https://www.linkedin.com/posts/stevengouri_a-wild-stat-80-of-startups-pitching-a16z-activity-7396182718998351872-xTKR.↵
61. *PocketGamer.biz*, “Shift Up CEO says AI is key to competing with China’s game industry scale.” <https://www.pocketgamer.biz/shift-up-ceo-says-ai-is-key-to-competing-with-chinas-game-industry-scale/>. *Dream Machine* Issue 14.↵
52. *The Economist*, “Investors expect AI use to soar. That’s not happening.” <https://www.economist.com/finance-and-economics/2025/11/26/investors-expect-ai-use-to-soar-thats-not-happening>.

happening. *Dream Machine* Issue 9.↵

53. *The Economist*, “Job apocalypse? Humbug! AI is creating brand new occupations.” <https://www.economist.com/business/2025/12/14/job-apocalypse-humbug-ai-is-creating-brand-new-occupations>. *Dream Machine* Issue 12.↵
54. *The Guardian*, “AI is hitting UK harder than other big economies, study finds.” <https://www.theguardian.com/technology/2026/jan/26/ai-uk-jobs-us-japan-germany-australia>. *Dream Machine* Issue 16.↵
55. University of Wisconsin-Stout, “AI Reshaping Industry: New UW-Stout Course Sets AI-Use as Baseline Competency in Filmmaking.” <https://www.uwstout.edu/about-us/news-center/ai-reshaping-industry-new-uw-stout-course-sets-ai-use-baseline-competency-filmmaking>. *Dream Machine* Issue 15.↵
56. Adobe Firefly enterprise metrics, in Appendix E: Dynamics of Generative AI Adoption.↵
57. Reuters Institute, “AI adoption by UK journalists and their newsrooms.” <https://reutersinstitute.politics.ox.ac.uk/ai-adoption-uk-journalists-and-their-newsrooms-surveying-applications-approaches-and-attitudes>. *Digiday*, “Daily Mail says Google AI Overviews have killed click-throughs.” <https://digiday.com/media/daily-mail-says-google-ai-overviews-have-killed-click-throughs/>. *Dream Machine* Issues 7, 9.↵
58. 1,100-creator music survey 2026, in Appendix D: Shadow AI, §“Music Production and Sound Recording.”↵
59. VFX AI integration metrics, in Appendix E, §“Visual Effects (VFX) Automation.”↵
70. *PC Gamer*, “Square Enix aims to have AI doing 70% of its QA work by the end of 2027.” <https://www.pcgamer.com/gaming-industry/square-enix-aims-to-have-ai-doing-70-percent-of-its-qa-work-by-the-end-of-2027/>. *Dream Machine* Issue 7.↵
71. *Eurogamer*, “Falcom is the latest developer to buy into the AI hype machine.” <https://www.eurogamer.net/falcom-is-the-latest-developer-to-buy-into-the-ai-hype-machine>. *Dream Machine* Issue 12.↵
72. *Metro*, “Prince of Persia remake and five more games cancelled as Ubisoft focuses on AI.” <https://metro.co.uk/2026/01/21/prince-persia-remake-five-games-cancelled-ubisoft-focuses-ai-26431926/>. *Dream Machine* Issue 15.↵
73. *Dream Machine* Issue 24, April 2026, on the GTA VI publisher laying off its internal AI team.↵
74. ElevenLabs \$500m ARR reporting, April 2026. *Dream Machine* Issue 25.↵
75. *36kr*, “AI Video Unicorn Higgsfield: Earns \$200M in 9 Months by ‘Serving’ Social Media Marketers.” <https://eu.36kr.com/en/p/3650517574312323>. *Dream Machine* Issue 16.↵
76. *Dream Machine* Issue 29, May 2026, on Sony’s 49-agent / 72-skill multi-agent game-development team.↵
77. *Digiday*, “AI agent developers have become adland’s in-demand role.” <https://digiday.com/marketing/ai-agent-developers-have-become-adlands-in-demand-role/>. *Dream Machine* Issue 7.↵
78. *Forbes*, “Vibe Coding — The In Demand AI Skill That Pays Up to \$220,000.” <https://www.forbes.com/sites/rachelwells/2025/11/06/the-in-demand-ai-skill-and-certifications-that-pays-up-to-220000/>. *Dream Machine* Issue 8.↵
79. Sundance Institute, “Centering the Artist: Why We’re Launching the AI Literacy Initiative.” <https://www.sundance.org/blogs/centering-the-artist-why-were-launching-the-ai-literacy-initiative/>. Google blog, “Sundance Institute AI Education.” <https://blog.google/company-news/outreach-and-initiatives/google-org/sundance-institute-ai-education/>. *Dream Machine* Issues 15, 16.↵
30. UK Government, “Free AI training for all.” <https://www.gov.uk/government/news/free-ai-training-for-all-as-government-and-industry-programme-expands-to-provide-10-million-workers-with-key-ai-skills-by-2030>. *Dream Machine* Issue 16.↵
81. Lovable for classrooms. <https://lovable.dev/classroom>. *Dream Machine* Issue 11.↵
32. UW-Stout course launch, January 2026 — *op. cit.*↵
33. Adobe, “Firefly Foundry.” <https://news.adobe.com/news/2025/10/adobe-max-2025-firefly-foundry>. *Dream Machine* Issue 5.↵
34. Korin AI launch, May 2026. *Dream Machine* Issue 27.↵

85. *The Verge*, “New York’s new law forces advertisers to say when they’re using AI avatars.” <https://www.theverge.com/news/842848/new-york-law-ai-advertisements-sag-aftra-labor>. *Dream Machine* Issue 11. C2PA / SynthID infrastructure references in Chapter 12.↵
36. *Forbes*, “AI Is Changing How Creators Work And Earn.” <https://www.forbes.com/sites/kolawolesamueladebayo/2025/12/22/how-ai-is-changing-how-creators-work-and-earn/>. *Dream Machine* Issue 13.↵
87. Variety, “AI Creator Behind Viral ‘Deadpool,’ ‘Harry Potter’ Christmas Clip Made His Film in a Ukrainian Bomb Shelter.” <https://variety.com/2026/digital/news/ai-video-deadpool-harry-potter-andrii-daniels-1236624632/>. *Dream Machine* Issue 16.↵
38. *Broadcast Pro Middle East*, “Tunisian filmmaker wins \$1 million AI Film Award for ‘Lily’.” <https://www.broadcastprome.com/news/tunisian-filmmaker-wins-1-million-ai-film-award-for-lily/>. *Dream Machine* Issue 14.↵
39. *Dream Machine* Issue 25, April 2026, on Indonesia’s *Legenda Bertuah*.↵
30. BBC Future, “Lights, camera, algorithm: Why Indian cinema is awash with AI.” <https://www.bbc.co.uk/future/article/20251223-why-indian-cinema-is-awash-with-ai>. *Dream Machine* Issue 14.↵
91. *TechBullion*, “Why the future belongs to multi-skilled leaders.” <https://techbullion.com/playing-the-long-game-with-a-portfolio-career-why-the-future-belongs-to-multi-skilled-leaders/>. *Dream Machine* Issue 9.↵
92. Anthropic Skills framework via Claude Code. *Dream Machine* Issues 11, 16, 29.↵
93. Adobe, “Inaugural Adobe Creators’ Toolkit Report,” October 2025. <https://news.adobe.com/news/2025/10/adobe-max-2025-creators-survey>. *Dream Machine* Issue 6.↵
94. PRS for Music, “PRS for Music AI Survey 2026.” <https://www.prsformusic.com/m-magazine/news/prs-for-music-ai-survey-2026>. *Dream Machine* Issue 16.↵
95. *CNBC*, “People with ADHD, autism, dyslexia say AI agents are helping them succeed at work.” <https://www.cnbc.com/2025/11/08/adhd-autism-dyslexia-jobs-careers-ai-agents-success.html>. *Dream Machine* Issue 7.↵
96. McKinsey & Company, “What AI could mean for film and TV production and the industry’s future.” <https://www.mckinsey.com/industries/technology-media-and-telecommunications/our-insights/what-ai-could-mean-for-film-and-tv-production-and-the-industrys-future>. *Dream Machine* Issue 16.↵
97. GDC State of the Game Industry surveys 2024–2026, in Appendix E, §“The Video Game Industry.”↵
98. LANDR AI music study, late 2025, referenced via *Ari’s Take*. <https://aristake.com/ai-tools-musicians-study/>. *Dream Machine* Issue 8.↵
99. Stanford AI Index Report 2025. Summarised in Appendix E, §“The Perception Gap.”↵
30. YouGov 2024 multi-market AI sentiment survey. Summarised in Appendix E.↵
01. *Digital Music News*, “Nearly 800 Creatives Sign Responsible AI Declaration — ‘Stealing Our Work Is Not Innovation’.” <https://www.digitalmusicnews.com/2026/01/22/stealing-isnt-innovation/>. *Dream Machine* Issue 16.↵
32. *Broadcast Pro Middle East*, *Lily* award — *op. cit.*↵
33. Variety, Andrii Daniels bomb-shelter clip — *op. cit.*↵
34. BBC Future, “Lights, camera, algorithm” — *op. cit.*↵
35. *Dream Machine* Issue 25, Indonesian *Legenda Bertuah*.↵
36. *CNBC Africa*, “How AI is changing the landscape of the music industry in Africa.” <https://www.cnbc africa.com/2025/how-ai-is-changing-the-landscape-of-the-music-industry-in-africa>. *Dream Machine* Issue 5. Korin AI launch, May 2026 — *op. cit.*↵
07. *PocketGamer.biz*, “Shift Up CEO says AI is key to competing with China’s game industry scale.” <https://www.pocketgamer.biz/shift-up-ceo-says-ai-is-key-to-competing-with-chinas-game-industry-scale/>. *Dream Machine* Issue 14.↵
38. The doomer frame is set out at length in OpenAI’s spring 2026 policy paper on the “social contract for the intelligence age” (covered in *Dream Machine* Issue 24) and in Daniel Susskind’s *A World Without Work*. The

full empirical rebuttal — including the Stanford Digital Economy Lab’s *Canaries in the Coal Mine?* findings (Brynjolfsson, Chandar & Chen, 2025), Brynjolfsson, Li & Raymond’s *Generative AI at Work* (QJE, 2025; n = 5,179 agents), Acemoglu’s *Simple Macroeconomics of AI* (Economic Policy, 2025), the Dallas Fed’s labour analyses, the St. Louis Fed’s AI-and-GDP tracking, the 2026 Stanford AI Index, and the 230-year history of identical predictions — is laid out in Appendix I: The Doomer Mistake. The deep dive is the analytical antibody to the policy framing this section is pushing back against.

39. Bloomberg, “AI Changed Chess. Grandmasters Now Win With Unpredictable Moves.” 27 March 2026. <https://www.bloomberg.com/news/articles/2026-03-27/ai-changed-chess-grandmasters-now-win-with-unpredictable-moves>. *Dream Machine* Issue 23. The behavioural pattern the piece describes — top grandmasters deliberately deviating from machine-optimal lines to put opponents on uncomputed ground — is the cleanest available analogy I have for the strategic shift the rest of this chapter argues for.
10. *Digital Music News*, “The AI Licensing Shift — Creative Weight Attribution Emerges as Music Industry Game-Changer for Rights Holders.” *op. cit.* *Dream Machine* Issue 16.
11. DreamLab AI Collective, team page. <https://dreamlab-ai.com/team>.
12. OpenAI, “Sora 2 is here,” 30 September 2025. <https://openai.com/index/sora-2/>. *Dream Machine* Issue 1.
13. LinkedIn News aggregation: “Sora Tops 1 Million Downloads in 5 Days.” <https://www.linkedin.com/news/story/sora-tops-1m-downloads-in-5-days-6684988/>. *Dream Machine* Issue 3.
14. Google DeepMind, Veo 3.1 launch, mid-October 2025. *Dream Machine* Issue 3.
15. Runway product cycle: Gen-4.5 (December 2025), Gen-4.5 Image-to-Video (January 2026), Workflows, Story Panels, Characters API, Apps for Advertising — *Dream Machine* Issues 10, 14, 15, 16, 20.
16. Runway CEO on indie films vs. blockbusters, *Dream Machine* Issue 26.
17. Chinese open-source AI video model releases, 2025–2026. *Dream Machine* Issues 3, 12, 22.
18. *SiliconAngle*, “Higgsfield raises \$80M on \$1.3B valuation.” <https://siliconangle.com/2026/01/15/higgsfield-raises-80m-1-3b-valuation-scale-ai-video-platform/>. 36kr, “Higgsfield: Earns \$200M in 9 Months.” <https://eu.36kr.com/en/p/3650517574312323>. *Dream Machine* Issues 15, 16.
19. Heygen Video Agent. https://www.linkedin.com/posts/heygen_introducing-the-new-video-agent-activity-7421597801240801282-d1CF. *Dream Machine* Issue 16.
20. TechCrunch, “Synthesia hits \$4B valuation, lets employees cash in.” <https://techcrunch.com/2026/01/26/synthesia-hits-4b-valuation-lets-employees-cash-in/>. *Sifted*, “Synthesia rejects \$3bn Adobe acquisition offer.” <https://sifted.eu/articles/synthesia-acquisition-offer>. *Dream Machine* Issues 5, 16.
21. ElevenLabs \$500m ARR reporting, April 2026. *Dream Machine* Issue 25.
22. Google DeepMind, “Introducing Gemini Omni: Create Anything from Any input.” <https://blog.google/technology/google-deepmind/gemini-omni-launch/>. *Dream Machine* Issue 30.
23. Beeples Canvas — Generative AI compositor. <https://www.beeple-canvas.com/>. *Dream Machine* Issue 30.
24. Adobe Firefly milestone data, in *Dynamics of Generative AI Adoption*, §“The Ubiquity of AI in Visual and Digital Arts.”
25. Nano Banana inside Photoshop and inside Unreal Engine cross-integrations, October–November 2025. *Dream Machine* Issue 1.
26. Suno Studio launch. <https://www.techradar.com/ai-platforms-assistants/i-tried-suno-studio-the-new-platform-that-mixes-ai-music-generation-with-hands-on-editing-like-garageband-but-smarter>. *Dream Machine* Issue 1.
27. Mureka, “Music Agent Studio” launch. *Dream Machine* Issue 4.
28. ElevenLabs Series funding, April 2026. *Dream Machine* Issue 25.
29. *MusicTech*, “Cardiff band speaks out after AI artist trained on their music outperforms them on Spotify.” <https://musitech.com/news/industry/its-shocking-disheartening-and-insulting-cardiff-band-speaks-out-after-ai-artist-trained-on-their-music-outperforms-them-on-spotify/>. *Dream Machine* Issue 1.
30. Variety, “AI Creator Behind Viral ‘Deadpool,’ ‘Harry Potter’ Christmas Clip Made His Film in a Ukrainian Bomb Shelter.” <https://variety.com/2026/digital/news/ai-video-deadpool-harry-potter-andrii-daniels->

1236624632/. *Dream Machine* Issue 16.↵

31. Sony AI, “Woosh — a sound effect foundation model.” <https://ai.sony/blog/woosh-sound-effect-foundation-model/>. *Dream Machine* Issue 30.↵
32. Mirelo SFX 1.6, “edit sound, not just generate it.” <https://mirelo.ai/sfx-1-6>. *Dream Machine* Issue 30.↵
33. Stability AI, “Stable Audio 3.0 released — open-weight model family built for artistic experimentation.” <https://stability.ai/news/stable-audio-3-0-released>. *Dream Machine* Issue 30.↵
34. Tamber product page: <https://tamber.ai/>. *Dream Machine* Issue 30.↵
35. Beatport Track ID. <https://www.beatport.com/track-id>. *Dream Machine* Issue 30.↵
36. Music industry AI deal flow, October 2025 – May 2026. See Chapter 5 footnotes 31–37, and *Dream Machine* Issues 5, 7, 8, 12, 14, 16, 17.↵
37. World Labs, “Bringing Marble to Life.” <https://www.worldlabs.ai/case-studies/bringing-marble-to-life>. *Dream Machine* Issue 7.↵
38. Sony Pictures Marble VP integration. *Dream Machine* Issue 8.↵
39. Google DeepMind, “Genie 3.” <https://deepmind.google/blog/genie-3-a-new-frontier-for-world-models/>. Project Genie: <https://blog.google/innovation-and-ai/models-and-research/google-deepmind/project-genie/>. *Dream Machine* Issues 3, 17.↵
40. Tencent, “HY World 1.5” and Hunyuan 3D Studio. *Dream Machine* Issue 12.↵
41. Luma AI, *UNI-1* launch, March 2026. *Dream Machine* Issue 22.↵
42. NVIDIA SANA-WM model collection. <https://huggingface.co/collections/nvidia/sana-wm>. *Dream Machine* Issue 30.↵
43. Odyssey, “Introducing Starchild-1, the first real-time multimodal world model.” <https://odyssey.ml/introducing-starchild-1>. *Dream Machine* Issue 30.↵
44. Odyssey, “Introducing Agora-1 — four-player AI-generated world built on a 1997 shooter.” <https://odyssey.ml/introducing-agera-1>. *Dream Machine* Issue 30.↵
45. Apple Machine Learning Research, “Apple Headsup: a Large-Scale High-Quality 3D Gaussian Head Reconstruction from Multi-View Captures.” <https://machinelearning.apple.com/research/apple-headsup-3d-gaussian-head>. *Dream Machine* Issue 30.↵
46. SuperSplat / Spark 2.0 / SOG releases through 2025–26. *Dream Machine* Issues 1, 25.↵
47. Radiance Fields, “Apple Confirms that it’s Gaussian Splatting that powers their personas.” <https://radiancefields.com/apple-confirms-personas-use-gaussian-splatting>. *Dream Machine* Issue 5.↵
48. ComfyUI Blog, “Ubisoft La Forge Open-Sources the CHORD Model.” <https://blog.comfy.org/p/ubisoft-open-sources-the-chord-model>. *Dream Machine* Issue 11.↵
49. Anthropic / Blender Foundation patronage, May 2026. *Dream Machine* Issue 27.↵
50. OpenAI, “Introducing AgentKit.” <https://openai.com/index/introducing-agentkit/>. *Dream Machine* Issue 2.↵
51. Anthropic Skills framework. *Dream Machine* Issues 11, 16, 29.↵
52. Google, “Official skills for AI agents.” <https://github.com/google/agent-skills>. *Dream Machine* Issue 30.↵
53. Tencent Ardot, AI-native design agent platform. <https://ardot.tencent.com/>. *Dream Machine* Issue 30.↵
54. Heygen Video Agent. *Dream Machine* Issue 16.↵
55. Adobe Summit 2026 CX Enterprise. *Dream Machine* Issue 26.↵
56. Adobe + NVIDIA / Google + NVIDIA partnerships. *Dream Machine* Issue 21.↵
57. ComfyUI funding round. https://www.linkedin.com/posts/comfyui_we-raised-17-million-to-build-an-os-for-ugcPost-7373743341236236288-wkCc. *Dream Machine* Issue 1.↵
58. ComfyUI \$500M valuation, May 2026. *Dream Machine* Issue 27.↵
59. Anthropic, “Claude is now available as a partner node in ComfyUI.” <https://www.anthropic.com/news/claude-comfyui-partner-node>. *Dream Machine* Issue 30.↵
50. Hugging Face / Google Cloud and Meta / Hugging Face OpenEnv. *Dream Machine* Issues 5, 8.↵

61. Unreal Engine 5 official AI Assistant. https://www.linkedin.com/posts/wouterweynants_theres-an-official-ai-assistant-coming-to-ugcPost-7369377204226379776-pGiH. *Dream Machine* Issue 1.↵
52. ECABridge — Unreal Engine MCP integration. <https://ecabridge.dev/>. *Dream Machine* Issue 30.↵
53. *Video Games Chronicle*, “Epic Games Veteran Claims He’s Building AI-Heavy ‘Fully European’ Game Engine.” <https://www.videogameschronicle.com/news/epic-games-veteran-ai-heavy-fully-european-game-engine/>. *Dream Machine* Issue 30.↵
54. Unity AI Council (October 2025); Unity AI Open Beta (May 2026). *Dream Machine* Issues 1, 28.↵
55. VFX AI integration metrics. See *Dynamics of Generative AI Adoption*, §“Visual Effects (VFX) Automation.”↵
56. Anthropic / Blender Foundation patronage. *Dream Machine* Issue 27.↵
67. Andreessen Horowitz pitch-deck observations on Chinese open-source model usage. https://www.linkedin.com/posts/stevenouri_a-wild-stat-80-of-startups-pitching-a16z-activity-7396182718998351872-xTKR. *Dream Machine* Issue 8.↵
58. NVIDIA SANA-WM model collection. <https://huggingface.co/collections/nvidia/sana-wm>. *Dream Machine* Issue 30.↵
59. PhotoGIMP — the open-source GIMP skin that mimics Photoshop. <https://github.com/Diollinux/PhotoGIMP>. *Dream Machine* Issue 30.↵
70. Korin AI launch, May 2026. *Dream Machine* Issue 27.↵

dream machine: the new creative economy is the field guide to the most consequential year in creative work since cinema learned to talk. written from inside the work by creative technologist pete woodbridge, this book is not just a tools guide or a five-point manifesto. it is a careful, public record of an industry coming apart and putting itself back together.

drawing on his vantage point running a studios in the north west of england and publishing a weekly newsletter for thousands of creatives, woodbridge holds the whole picture together in a way no platform-company keynote or academic has managed.

whether you are a working creative trying to figure out what to do on monday morning, a studio leader making organisational decisions, or a policymaker setting the rails for the future, this book provides the frameworks you need to stay oriented in a fast-moving environment.

creative work is being re-platformed. the choices made right now will set the terms for the next decade.

dreamlab.org.uk