

dream machine

***artificial intelligence and
the next creative economy***

pete woodbridge
dreamlab.org.uk

Dream Machine

Artificial Intelligence and the Next Creative Economy

Pete Woodbridge

Edition 14 August 2026

DreamLab AI Collective

dreamlab.org.uk

Dream Machine: Artificial Intelligence and the Next Creative Economy

© 2026 Pete Woodbridge. All rights reserved.

Edition of 14 August 2026.

Written and edited at DreamLab in the North West of England, between October 2025 and August 2026, in parallel with forty-one editions of the *Dream Machine* newsletter.

All footnoted claims are sourced. Every footnote links either to a primary source or to the issue of *Dream Machine* in which the item was first discussed.

A note on this edition. **Dream Machine** is a *living book*. The print edition is updated and re-published on a rolling cadence — currently roughly monthly — as the newsletter ships new issues, as deep-dive research is added, and as the field's evidence sharpens. The full corpus, the build pipeline, and the always-current copy of the manuscript are open at github.com/petejwoodbridge/DreamMachine-Book. The unversioned *Dream_Machine.pdf* file in the published archive is always the most recent edition; this dated *2026-08-14* copy is one snapshot inside the longer sequence. The book's own argument — that AI enables new formats and publishing models, where the work, the toolchain and the commentary update continuously rather than in five-year hardback cycles — is one the book is also trying to demonstrate by being.

DEDICATION

This book is for the collective.

*To the **DreamLab team** — the around fifty deep-tech specialists who make the studio what it is across labs in Manchester, Liverpool, Cumbria and London. The AI engineers, XR developers, game designers, animators, sound artists, researchers, educators and filmmakers — the Emmy winners, the PhD researchers, the BAFTA-recognised practitioners — who turn up to do the work, who put taste and judgement into every brief, and who hold the human signal in the work the audience eventually sees. Every argument in this book has been sharpened by conversation with one of you. Several of you have, without knowing it, written the lines that landed best.*

*A particular thank you to **Dr. John O'Hare**, whose research on agentic organisational design supplied much of the structural argument underneath Chapter 14, and whose long-running collaboration on the harder organisational-theory questions has shaped large parts of Chapter 12, Chapter 15 and Chapter 16 as much as any single contribution from outside this desk. The 2,000-year framing of hierarchy as an information-routing protocol, the jagged-frontier reading of Dell'Acqua's BCG experiment, the judgment broker role replacing middle management, the compounding organisation concept — these are John's arguments, and the book is meaningfully more careful for them.*

To everyone at dreamlab.org.uk — see the team page — this is yours as much as it is mine.

— Pete

About this book

In September 2025, a synthetic actress walked onto a Zurich film festival stage, OpenAI shipped Sora 2, and Pete Woodbridge sat down to write a newsletter about it. Forty-one weekly issues later, what began as a one-month experiment had become *Dream Machine* — the most-read working-creative record of the AI transition.

This is the book that record produced. From the slop ceiling that the audience is already enforcing, to the 88% of UK creators demanding licensing in all cases, to the chess-grandmaster strategy of deliberately playing the move the machine wouldn't make — Woodbridge holds the whole picture together in a way no journalist, academic or platform-company keynote has managed.

Equal parts history, manifesto and operating manual, *Dream Machine* is the field guide to the most consequential year in creative work since cinema learned to talk. Written from inside **DreamLab** — the studio Woodbridge founded at MediaCity in Manchester in January 2024, now around fifty AI engineers, XR developers, game designers, animators, sound artists, researchers, educators and filmmakers across labs in Manchester, Liverpool, Cumbria and London — drawing on the studio's closed-beta access to platforms like World Labs' Marble, project work with partners including Intel, Google, Sony, Aardman, BBC, NBC, Apple, MediaCityUK and Universal, and the operational experience of running an agentic-mesh studio in production every week of the period the book describes.

Dream Machine is a **living book**. New editions are published on a rolling cadence — currently roughly monthly — as the newsletter ships new issues and as the field's evidence sharpens. The unversioned *Dream_Machine.pdf* in the published archive is always the most recent edition; dated editions are preserved alongside as snapshots of the longer sequence. The cadence is part of the argument: AI enables new publishing models in which the work, the toolchain and the commentary keep moving with the field rather than freezing in five-year hardback cycles.

The age of the How is ending. Welcome the Why.

Contents

<i>Foreword — Welcome to the Dream Machine</i>	8
<i>Reader paths</i>	19
Chapter 1 — The Day Sora Landed	23
Chapter 2 — A History of Resistance	34
Chapter 3 — The Human-AI Agency Continuum	52
Chapter 4 — Dead Internet, Living Web	62
Chapter 5 — The Slop Ceiling	74
Chapter 6 — The 88%	88
Chapter 7 — The Studios Decide	110
Chapter 8 — The Machine Gets the Money	132
Chapter 9 — Worlds, Not Pictures	149
Chapter 10 — AI in Everything, Everywhere, All at Once	163
Chapter 11 — What is newly possible	179
Chapter 12 — The Orchestrator	196
Chapter 13 — Authenticity as the New Scarcity	213
Chapter 14 — Coordination Collapse	232
Chapter 15 — The New Jobs	249
Chapter 16 — Choosing the Future	267
Chapter 17 — The Tools	285
Chapter 18 — Five Years Inside the Dream Machine	333
Epilogue — A Letter from the Dream Machine	344

APPENDICES

<i>Appendix A — A Quantitative Anatomy of Six Months</i>	349
<i>Appendix B — Glossary</i>	367
<i>Appendix C — Bibliography by Topic</i>	373
<i>Appendix D — The Dream Machine Source Index</i>	395

REFERENCE

<i>Citation Index</i>	437
-----------------------------	-----

ABOUT

About the author, and about DreamLab	515
--	-----

Foreword — Welcome to the Dream Machine

There was a week, at the very end of September 2025, when two things happened on opposite sides of the world that I knew, while I was watching them, would change the rest of my year — and probably most of yours.

In Zurich, an AI company called Particle6 walked an entirely synthetic actress called Tilly Norwood onto the Zurich Film Festival stage and announced that talent agencies were already in conversation about representing her. She had a face, a personality, a showreel and, by the founder's own framing, a future. Within forty-eight hours, the U.S. actors' union SAG-AFTRA had issued a statement that she was "not an actor" but "a character generated by a computer program that was trained on the work of countless professional performers."¹ The U.K. union Equity followed.² By the time the weekend was over, Whoopi Goldberg, Emily Blunt, Melissa Barrera and a long list of other working performers had told the world, in their own words, what they thought of all this.³

The other thing was that OpenAI released Sora 2.⁴

I sat down on a Monday morning in early October, opened a blank LinkedIn article, and wrote my first edition of a newsletter called *Dream Machine*. I didn't have a plan. I had a sentence:

"It's time to take AI seriously."

I sent that issue out to a few hundred people in my network. I assumed I'd do it for a month, maybe two — that the wave would pass, or that I'd run out of material, or that someone with a bigger platform would do it better. Ten months later, the newsletter has 4,100-odd subscribers, forty-one published editions, several thousand curated links, and a community of people in the North West of England — the **DreamLab Collective**, around fifty of us across labs in Manchester, Liverpool, Cumbria and London — who help me read, sift, argue and build around it every week.⁵

This book is the thing that happens when you keep a careful, public record of an industry coming apart and putting itself back together inside the same ten months.

* * *

I should say who I am, because the rest of this book is in the first person and you should know what kind of "I" is talking to you.

I'm a **Creative Technologist** and **Innovation Director** working across AI, immersive, interactive and experiential tech, real-time production and next-generation storytelling — across film, games, interactive and immersive experiences. My cross-sector background spans creative, technical, innovation and product-development roles for software companies, games and immersive studios, and

working on new possibilities in storytelling covering film, advertising, brand experiences and animation. The job title has had a few names over the years — *creative technologist*, *R&D director*, *innovation lead*, *head of emerging technology* — and they all shade into the same job: shipping software and creative products inside the gap between research labs, game studios, broadcasters and emerging-technology companies, then writing up what was learned along the way. The book is, in a real sense, a longer version of the *writing up what was learned along the way* habit.

I've worked on projects and new products with partners including **Intel**, **Google**, **McLaren**, **ITV**, **Sony**, **Aardman**, **Red Bull**, **IBM**, **BBC**, **NBC**, **Crytek**, **Apple**, **MediaCityUK**, **Saatchi & Saatchi** and **Universal** — alongside dozens of independent studios, agencies and SMEs building new products, content pipelines, and interactive and immersive experiences. Technically the work runs across strategy, leadership and hands-on development: real-time interactive and immersive systems (**Unreal Engine**, **Unity**, **WebXR**, **Disguise**, **Notch**, **Ventuz**); AI, software and workflow engineering (**Python**, **ComfyUI**, **API integrations**, **automation systems**, **AI-assisted pipeline design**); and production and future-media (**virtual production**, **motion capture**, **ICVFX**, **real-time previs**, **AI-augmented workflows**). I've built and led teams of engineers and artists across most of those categories at one time or another, which is, more than anything else, the experience the book's arguments about orchestrator-shaped careers and agentic-mesh studios are made out of.

In parallel to the studio work I've developed innovation and investment programmes for the UK creative industries — including the **£3.2 million Innovation Accelerator (Innovate UK)** at **MediaCity** in Manchester, and the **£7.2 million MusicFutures programme (UKRI / AHRC)** in Liverpool — and I sit as an external assessor and advisor on several national and international innovation programmes and strategies, including **CoSTAR**, **EPSRC** and **Digital Catapult**, alongside contributions to a number of steering committees, boards, startups and investment initiatives. I include this not as a CV flourish but because a meaningful slice of the book's policy arguments — the Petrillo-template chapter, the apprenticeship-gap section, the four-principles framework in Chapter 16 — has been sharpened in conversation with the institutions running those programmes, and the reader can fairly take that into account when reading the policy chapters.

DreamLab, the studio I'm writing this book from, has been operating since **January 2024**. It started at **MediaCity** in Manchester and has, in the two-and-a-half years since, grown into a network of **around 50 deep-tech specialists** across labs in **Manchester**, **Liverpool**, **Cumbria** and **London** — and continues to spread. The team is AI engineers, XR developers, game designers, animators, sound artists, researchers, educators and filmmakers, some of whom have won Emmys and

BAFTAs, some of whom finished their PhDs last year. We don't operate as a fixed agency; we assemble the right team of specialists around each challenge — a model that, in the period covered by this book, has become both the operational pattern of the orchestrator economy described in [Chapter 12](#) and our own default way of working. Across the two and a half years since DreamLab launched, the studio has supported **100+ businesses**, delivered **35+ collaborative projects** and helped channel approximately **£2.7 million of investment** into partner organisations. None of that is the book's argument; it's the *vantage point* the book's argument is being made from.⁶

I'm not an AI evangelist. I'm also not an AI sceptic. I'm the kind of practitioner who has had Marble running on a beta key for months and who has also sat in a room with a games studio CEO who used the phrase "AI was an expensive mistake" without breaking eye contact.⁷ I have built things with these tools and I have watched them break. I'm writing this book from inside the work — and writing it, by the way, in the same voice I write the newsletter in, which is the same voice I argue with my team in, which is the same voice (give or take a bracket) you'll hear if we ever end up in a pub together. A fuller version of the bio and a fuller picture of DreamLab — including the studio's client list, recognition and structural position inside the UK creative-tech ecosystem — sits in the [About the author, and about DreamLab](#) section at the back of the book.

Through the ten months the book covers, I have been doing three things in parallel:

- **Running DreamLab**, the ~50-person studio above, which uses every major AI platform discussed in this book in live production. We have been a closed-beta partner for World Labs' Marble since October 2025.
- **Writing the *Dream Machine* newsletter** — forty-one weekly issues, several thousand curated source links across film, music, games, advertising and broadcast, a subscriber base of roughly 4,100 working creatives at the time of publication.
- **Building out a small set of analytical frameworks** — the *Human-AI Agency Continuum* (Chapter 3), the *Slop Ceiling* (Chapter 5), the *Four Positions* map of studio strategy (Chapter 7), the *Seven-Year Problem* and the *duration discount* (Chapter 8), the *Year of the Orchestrator* (Chapter 12), *Coordination Collapse* and the *Consumption Gap* (Chapter 14), the *AI Literacy Premium* and the *Apprenticeship Gap* (Chapter 15), the *Four Principles* of a humane creative economy (Chapter 16), and the *Seven-Layer Toolchain* model (Chapter 17) — that have begun, over the period, to circulate in the wider industry conversation.

The combination of the three is unusual. The trade-press journalists write the coverage but do not, on the whole, run studios. The studio operators run studios but do not, on the whole, publish a public weekly record. The academics produce

research, sometimes excellent, but at the cadence of academia rather than the cadence of the transition. The platform companies produce material at the cadence of their own product cycles. Working from across the three positions at once, week by week, has produced a particular kind of vantage point — neither outside the work nor confined to one slice of it — that I have not seen anyone else holding consistently through this period.

The book is what that vantage point produces. There are good histories of cinema written by people who never ran a film studio, and good histories of the music industry written by people who never produced a record. There are also, sometimes, books written from inside the work — by people who were making the decisions the book describes, in real time, under the same conditions. The second kind tends, when honest, to tell you something the outsider accounts cannot: what the moment *felt like* to be inside, what the practitioners thought they were doing, and what — looking back — the moment was actually for.

This is that kind of book.

I am writing it because, in early October 2025, I realised that nobody I respected was doing what I needed somebody to do: hold the whole picture in one place. Not the boosters. Not the doomers. Not the niche tool-reviewers. The whole picture, week by week, including the contradictions. Including the bits where Adobe was telling us that 86% of creators already use generative AI in their workflow,⁸ *and* where 88% of creators who replied to the UK government’s copyright consultation said AI companies should have to license their work in every case.⁹

The whole point — to me — is that those two numbers are both true.

* * *

The arc of these eight months turned out to be tighter than I expected.

In **October 2025**, the question on the table was whether AI in the creative industries was real. By the time I finished writing Issue 5, it was no longer a question.¹⁰

In **November**, the question shifted to whose tool stack it would run on. Adobe announced — and I am quoting them, this isn’t editorial flourish — “AI in everything, everywhere, all at once”.¹¹ World Labs released Marble for public use a couple of weeks later, and the entire shape of what “a creative asset” can be quietly changed.¹²

In **December**, the question was about consent and money. The UK government’s copyright consultation closed with eleven and a half thousand responses — one of

the largest copyright consultations the country has ever run — and a number that has stayed with me: 88%.¹³

In **January 2026**, the question was who decides the rules. Sundance launched an AI Literacy Initiative for filmmakers. Bandcamp banned AI music outright. Steam clarified, then re-clarified, what counts as AI in a video game. Almost 800 creators signed an open declaration with the line, “*Stealing our work is not innovation.*”¹⁴

In **February, March, April and May**, the question started to feel like a different question altogether. It wasn’t really *should we use these tools*. It was: *now that the tools are inside the production pipeline at every studio, every label, every newsroom, every agency, what kind of creative economy do we actually want on the other side?*

That last question is the one this book is about.

What this book is and is not

This is **not a tools guide**. Chapter 17 lists every significant tool that surfaced in the *Dream Machine* archive in the period the book covers, but the rest of the chapters are organised around the *transition* — the economics, the audience, the labour, the unions, the law, the institutions, the rails — rather than around the apps. The tools change weekly. The transition is what will still be true in 2030.

This is **not a manifesto**. I do not believe the cleanest five-point plans for the future of creative work, and I have refused to write one. What this book argues for, in Chapter 16, is a *test* you can apply to any policy, any contract, any platform decision: *Agency, Attribution, Access, Audience*. The test is not a programme. It is a way of staying oriented in a fast-moving environment.

This is **not a chronicle**. The newsletter is the chronicle. Every issue is online, every link is preserved, and the comprehensive thematic source index at the back of this book (Appendix D) catalogues the entire archive by topic for the reader who wants to follow specific threads.

What this book *is* is an argument, in eighteen chapters and four appendices, that creative work is being re-platformed in a twelve-month window — that this is not the internet of 1995 or the mobile phone of 2007, this is a faster, deeper, more thorough re-platforming of the economic and cultural rails on which creative work travels — and that the choices being made *right now*, by studios, by unions, by governments, by toolmakers, by individual creatives at their kitchen tables, will set the terms for the next decade. The book is here to help you make those choices on better information than you would otherwise have.

That re-platforming has a name. It’s the title of the book. *The New Creative Economy*. I don’t think it’s a metaphor and I don’t think we have very long to decide

what we want it to look like.

What I believe, stated plainly

The book lays out my position chapter by chapter, but if you want the headline conviction up front, it is this:

AI is best understood as an assistive instrument that amplifies human creativity. Not a replacement for it. Not a substitute for it. An amplifier of it.

The working creative economy that emerges from this transition will be — has to be — the one that does not lose sight of which side of that relationship is the master and which is the servant. The human creativity is the master. The AI is the servant. Chapter 16 is the long-form version of that argument; everything else in the book sits inside it.

I want to flag a second framing the book leans on, because it is the one I use most often in the talks I have been giving. **We are leaving the age of the *How* and entering the age of the *Why*.** The *How* — the technical labour of executing a creative thought, the ability to draw, light, mix, model, edit, render — has been the bottleneck of professional creative work for a century. The *How* is, in 2026, becoming a utility. A teenager with a midrange GPU can now produce work whose surface quality sits on a continuum with what a full studio could produce in 2020. When the *How* becomes a utility, the *Why* — taste, intent, authenticity, the willingness to take a risk on the move the data does not yet endorse — is the only thing left with commercial leverage. Chapter 16 anchors this argument in a story I borrow from elite chess; the rest of the book lives inside the strategic implication.

If you read no further than this front matter, you have the heart of the book.

What this book does not do well

A critic-friendly note, because it makes for a more honest read.

This book is more confident about the *creative-industries* layer of the AI transition than it is about the layers above and below. The environmental and energy footprint of the systems the book describes is something I touch on in Chapter 16 and otherwise under-treat; a fuller account is the subject of a different book, by a different writer, that I hope is being written now. The labour conditions of the global data-supply chain — the labellers, evaluators and content moderators that the platform companies depend on — sit underneath every chapter of this book without being centred in any of them; the same caveat applies. The geopolitics of AI, the macro-economic question of the platform-company stock-market valuations, the wider policy questions about national AI strategy, the philosophical questions about machine consciousness — none of these is the book's subject, and the book is shorter and more useful for not pretending otherwise.

The book is also, by design, anglophone-skewed and Global-North-skewed in its primary sourcing. The 88% in Chapter 6 is a UK number. The Sundance turn in Chapters 9 and 10 is a US story. The platform-company analysis in Chapter 10 is, in the main, an analysis of US and European companies, with significant Chinese coverage but less than the Chinese open-source ecosystem deserves. The Indian, African, Latin American and Southeast Asian stories I cover in Chapters 7, 12 and 13 are real but I cover them, in places, from the wrong side of a translation gap. The next edition, if there is one, should fix this. DreamLab's deep-dive research begins the work but does not finish it.

Finally: the book is written *while the transition is happening*. Some of the specific claims, particularly in the tools chapter and in the predictions, will age in ways I am not yet able to predict. The frameworks should outlast their evidence. The evidence should be checked, when you read the book, against whatever the state of play is by then.

Who this book is for

I have written this book for **working creatives** — the writers, directors, songwriters, games designers, photographers, illustrators, editors, producers, agency creatives, indie filmmakers, YouTubers, freelance designers, students and senior practitioners who are, right now, trying to figure out what creative life looks like in 2026 and beyond.

It is also, secondarily, written for the **studio, agency and label leadership** trying to make organisational decisions about AI integration in a year in which the cost of getting it wrong is, by my read, the next decade of cultural authority.

And it is, thirdly, written for the **policy, union, institute and platform** people who are deciding the rails the next decade of creative work will run on. The 88% turned up to the UK consultation. The Cannes Disclosure Standard, the Academy rule, the SAG-AFTRA contract — these are the institutional decisions that shape the field. The people making them are part of the audience for this book.

Whoever you are: read with a pen. The chapters do not need to be read in order — there is a Reader Paths guide for different routes through the material. The Source Index at the back lets you follow any thread back to its primary sources.

A practical note, and an ask

I've written the book in the same voice I write the newsletter — talkative, opinionated, North-West English, occasionally too fond of a bracket. But I have tried, in every chapter, to put my opinions on top of evidence rather than the other way around. Every claim that matters is footnoted. Every footnote points either to a primary source (a research report, a court filing, an official announcement) or to the *Dream Machine* edition where I first wrote about it, where the original link is

preserved. There are several thousand citations. If you only ever read this book once, you can ignore them entirely; if you ever want to know whether I made something up, follow the trail.

One more practical thing I want to flag before we start, because it is, in a small way, part of what the book is arguing. *Dream Machine* is a **living book**. The print and digital editions are updated and re-published on a rolling cadence — currently roughly **monthly** — as the newsletter ships new issues, as the research deep dives are added, and as the field's evidence sharpens. Behind the manuscript there is a continuously-rebuilt **knowledge base** at DreamLab: a structured corpus of the newsletter archive, the deep-dive PDFs, the citation index and the toolchain catalogue, with a build pipeline that re-assembles the book end-to-end every time the underlying material moves. The unversioned `Dream_Machine.pdf` in the published archive is, by design, always the most recent edition; the dated copies are snapshots inside the longer sequence. If you want the latest, take the unversioned file. If you want the *specific* version a particular argument first appeared in, take the dated one. Both are kept.

I have made the cadence deliberate, not as a marketing posture but as a quiet demonstration of the book's own argument. The position the book settles on — that AI enables new formats, new publishing models and new ways of holding a body of knowledge alive in public — is one I think a book about AI in the creative industries has an obligation to *show*, not only to *say*. A traditional non-fiction hardback ships once, sits frozen for five years, and ages out of its own subject matter. *Dream Machine* updates the way the field does. It is published the way the field works now. It is, in that sense, an artefact of the practice it describes. I would not write it any other way.

I have one ask of you before we start.

The temptation, when reading a book about AI in the creative industries in 2026, is to take a side before chapter one. To decide, on page one, whether this is going to be about how the machines are coming for us or about how the machines are setting us free. Please don't. The most honest thing I can tell you about what I have learned over these ten months is that the truth is almost always both at once, and that the most interesting people in this story — the directors, the songwriters, the games developers, the union reps, the platform engineers, the indie filmmakers in bomb shelters and the policy officers in Whitehall — are the ones who can hold both sides at the same time without flinching.

That's the kind of book I want this to be.

Welcome to the Dream Machine.

— *Pete Woodbridge*

DreamLab AI Collective

The complete book in reading order:

- **Foreword: Welcome to the Dream Machine** — this piece
- **Reader Paths** — six routes through the material
- **Chapter 1: The Day Sora Landed** — the watershed week
- **Chapter 2: A History of Resistance** — the 1839–2025 pattern, and where AI sits inside it
- **Chapter 3: The Human–AI Agency Continuum** — the foundational frame, and the case for opening the black box
- **Chapter 4: Dead Internet, Living Web** — the web infrastructure question, and the provenance stack
- **Chapter 5: The Slop Ceiling** — the audience’s verdict, and the Authenticity Premium
- **Chapter 6: The 88%** — the political watershed, and the Petrillo template applied to AI
- **Chapter 7: The Studios Decide** — four strategic positions, and the trap legacy built for itself
- **Chapter 8: The Machine Gets the Money** — the Seven-Year Problem, and how AI repriced creative investment
- **Chapter 9: Worlds, Not Pictures** — the spatial turn
- **Chapter 10: AI in Everything** — the platform layer, and its underlying economics
- **Chapter 11: What is Newly Possible** — six categories of newly-possible work, and the finite-attention ceiling
- **Chapter 12: The Orchestrator** — the working creative’s new role
- **Chapter 13: Authenticity as the New Scarcity** — the provenance economy
- **Chapter 14: Coordination Collapse** — shadow AI, the consumption gap, organisational re-shaping
- **Chapter 15: The New Jobs** — the labour-market restructuring, the AI literacy premium, the Apprenticeship Gap
- **Chapter 16: Choosing the Future** — the four principles, the assistive-amplifier conviction, the Age of the Why
- **Chapter 17: The Tools** — the complete categorised inventory
- **Chapter 18: Five Years Inside the Dream Machine** — a speculative five-year future-cast, argued from the book

- **Chapter 19: Epilogue** — a letter to 2030
- **Appendix A: Quantitative Anatomy** — the numbers behind the book
- **Appendix B: Glossary** — terms of art
- **Appendix C: Bibliography by Topic** — curated reading
- **Appendix D: The Dream Machine Source Index** — thematic catalogue of every significant source across the 41 issues

Reader paths

This is a book about creative AI in the ten months between October 2025 and August 2026. It is not a tools guide. It is an argument, with the evidence underneath, about what kind of creative economy is being built right now and what we should do about it.

You can read this book straight through. It is built to reward that. The book has nineteen sections — a combined Foreword, eighteen chapters, an Epilogue — plus four appendices, sequenced so that each one earns the next.

If you don't have time for that, here are ten ways into the book that will save you from reading something that doesn't serve you yet. Pick one. Come back to the rest later.

If you are a working creative trying to figure out what to do next

Read the **Foreword**, then **Chapter 2 (A History of Resistance)** for the historical pattern you are inside, then **Chapter 3 (The Human–AI Agency Continuum)**, then **Chapter 11 (What is Newly Possible)** for the new categories of work, then **Chapter 12 (The Orchestrator)**, then **Chapter 15 (The New Jobs)** for the labour-market evidence, then **Chapter 16 (Choosing the Future)** — particularly the section *What working creatives should do on Monday morning*. That is the spine of the *practitioner's* argument. Everything else in the book is evidence supporting it.

If you run a studio, agency, label or production company

Read the **Foreword**, then **Chapter 7 (The Studios Decide)** — particularly the section *The trap the legacy industries built for themselves* — then **Chapter 8 (The Machine Gets the Money)** for what the capital markets have done to the strategy map, then **Chapter 14 (Coordination Collapse)** — particularly the section *What organisations should do*. Then read **Chapter 10 (AI in Everything, Everywhere, All at Once)** for the platform-economics frame, **Chapter 11 (What is Newly Possible)** for the new business categories, and **Chapter 15 (The New Jobs)** for the labour-market restructuring evidence. **Chapter 16** is the closing argument. **Chapter 17 (The Tools)** is the practical inventory. The deeper organisational-design argument that Chapter 14 rests on — the 2,000-year history of hierarchy as information-routing protocol, the jagged frontier, the vigilance problem, the three competing post-hierarchical models and a 90-day implementation roadmap — sits in DreamLab's published research.

If you invest in, lend to, or fund creative work

Read **Chapter 8 (The Machine Gets the Money)** first — the Seven-Year Problem, the duration discount, and the four sector verdicts (games, film, television, immersive) read from the allocation data rather than the press releases. Then **Chapter 7 (The Studios Decide)** for the strategy map the capital is pricing,

Chapter 5 (The Slop Ceiling) for the demand-side evidence underneath the authenticity premium, **Chapter 9 (Worlds, Not Pictures)** for the spatial layer where the discount changes sign, and **Chapter 13 (Authenticity as the New Scarcity)** for the provenance infrastructure the chapter argues is becoming collateral. **Appendix A** holds the underlying numbers.

If you are working in policy or law

Read the **Foreword**, then **Chapter 6 (The 88%)** and **Chapter 13 (Authenticity as the New Scarcity)** as a pair. **Chapter 8 (The Machine Gets the Money)** for the financing gap the policy instruments have to fill — particularly the section *The map, read whole*. **Chapter 15 (The New Jobs)** for the labour-market policy framework. **Chapter 16** for the four principles. **Appendix A** for the data the policy arguments rest on; **Appendix D** for the comprehensive source archive.

If you are working in (or covering) the music industry

Read **Chapter 5 (The Slop Ceiling)** first — it's where most of the music-specific analysis lives. Then **Chapter 6 (The 88%)** for the rights-and-licensing argument. Then **Chapter 13 (Authenticity as the New Scarcity)** for where it's heading. **Chapter 17 (The Tools)** for the comprehensive music-AI tool inventory.

If you are working in (or covering) film, TV or games

Read **Chapter 1 (The Day Sora Landed)** for the watershed scene-setter, **Chapter 7 (The Studios Decide)** for the strategic map of how the industry has positioned itself, **Chapter 8 (The Machine Gets the Money)** for what has happened to the money underneath it, **Chapter 9 (Worlds, Not Pictures)** for what is coming next, **Chapter 12 (The Orchestrator)** for what it means for working roles, and **Chapter 15 (The New Jobs)** for the labour-market data. The harder underlying argument about why the *creative / creatives* equivocation collapses under generative abundance, and the storyworld template that Chapter 9 argues is coming for every entertainment franchise — demonstrated live, at six-billion-viewer scale — both sit in DreamLab's published research.

If you want the empirical case against the AI-doomer policy framing

Read the **Foreword**, then **Chapter 15 (The New Jobs)** for the labour-market reading the book settles on. The long-form rebuttal that underlies it — the Stanford Digital Economy Lab data, Brynjolfsson's worker-level studies, Acemoglu's macroeconomic modelling, and the 230-year track record of identical predictions that have been wrong every time they have been asserted with confidence, together with the adoption-dynamics and shadow-AI evidence base — sits in DreamLab's published research.

If you are working in (or covering) IP, licensing, brand strategy or legal architecture

Read **Chapter 6 (The 88%)** for the political mandate the licensing-side conversation rests on, then **Chapter 7 (The Studios Decide)** — particularly the new *Position Five: Own the rails* section on Hasbro's Sixth Wall, the CharacterOS governance layer and the corrected record on the Disney–OpenAI deal that never executed — and then **Chapter 13 (Authenticity as the New Scarcity)** for the provenance and disclosure infrastructure that the IP layer sits inside. The long-form treatment of this — the *three Cs* (consent, credit, compensation) framework, the five-licensing-model comparison, the Coase / Williamson / Barney economic foundations, and the conceptual move from *IP as asset* to *IP as environment* — sits in DreamLab's published research.

If you are working in (or covering) sport, live events or broadcast

Start with the 2026 FIFA World Cup as the largest live deployment of AI in entertainment history: officiating, broadcast, the Roblox second venue, the creator flywheel, the synthetic-rights precedent and the robot dogs — the full case sits in DreamLab's published research. Then **Chapter 9 (Worlds, Not Pictures)** for the storyworld architecture it demonstrates, **Chapter 13 (Authenticity as the New Scarcity)** for why the live, verifiably human moment is where the premium migrates, and **Chapter 7 (The Studios Decide)** for the own-the-rails strategy FIFA is executing.

If you are reading this in a class, a book club, or as part of training

Read the **Foreword** and **Chapter 1** to get oriented, then **Chapter 4 (Dead Internet, Living Web)** to understand the structural argument, then **Chapter 15 (The New Jobs)** and **Chapter 16 (Choosing the Future)** to see the labour-market story and the four principles the book argues for. The other chapters are evidence and elaboration. The **Glossary** (Appendix B), **Citation Index**, and **Source Index** (Appendix D) are designed to be the back-pocket reference set for the rest of the year.

* * *

A note about reading order: the chapters are designed to be load-bearing on each other. If you skip a chapter and a later one references it without re-explanation, the Glossary should fill the gap. If the Glossary doesn't fill the gap, that is my failure and not yours — please write to me through the newsletter and I will improve the next edition.

— Pete

Chapter 1 — The Day Sora Landed

The first thing to understand about the week of 30 September 2025 is that nothing in it was supposed to be a watershed. Sora had existed, in some form, for a year and a half. Runway, Pika, Kling, Luma, Veo and a dozen others had been releasing video models on a near-monthly cadence since the back end of 2023.¹⁵ AI-assisted post-production had been quietly integrated into nearly every studio pipeline in Hollywood for months. The major U.K. and U.S. actors' unions had been negotiating over digital replicas since the 2023 SAG-AFTRA strike. The Tilly Norwood character had been on Instagram, posting selfies and pretending to drink coffee in cafés, since the previous July.

And yet that week — the seven days I would later mark as the start of this book, the start of the newsletter, and the start of a year that nobody in the creative industries has fully recovered from — three things lined up in a sequence so neat that I almost didn't believe it at the time.

On the Friday and Saturday, at the Zurich Film Festival, the founder of an AI studio called Particle6 stood on a panel and announced that several talent agencies were interested in representing her company's flagship product: a fully AI-generated actress called Tilly Norwood.¹⁶

On the Tuesday, the U.S. actors' union SAG-AFTRA issued a public condemnation calling her “not an actor” but “a character generated by a computer program that was trained on the work of countless professional performers” — adding, in a line I have thought about more than any other this year, that “audiences aren't interested in watching computer-generated content untethered from the human experience.”¹⁷

In between those two moments, on the Tuesday of the same week, OpenAI released Sora 2.¹⁸

I want to take each of these in turn, because the order matters, and then I want to argue that the actual watershed was something else entirely — something that happened underneath all three, that almost nobody noticed at the time, and that I think we will be living with for the rest of our working lives.

The actress

Tilly Norwood was the creation of Eline Van der Velden — a comedian, writer and producer who had spent the better part of a year building her under the banner of a U.K.–Netherlands company called Particle6.¹⁹ Tilly had a face that looked like it had been assembled by committee from the most marketable features of the late 2010s. She had a voice. She had a small social-media following. She had, by the time of the Zurich announcement, “another forty AI actors in the pipeline,” according to Van der Velden's later interview in *Deadline*.²⁰

The Zurich announcement wasn't subtle. Van der Velden told the audience that several talent agencies were "looking" at signing Tilly. She framed her as an industrial product: a character who could be cast in feature films and television, who came with all the upsides of a real performer (a marketing footprint, an emotional connection with audiences, a brand) and none of the downsides (no salary, no per diems, no licensing complications, no aging, no scandals).²¹

It is hard, looking back, to remember how loaded the word *agency* was that weekend. In Hollywood and the wider acting industry, the moment a talent agency takes on a new client is the moment that client moves from aspiration to product. The agencies are the gatekeepers of the working economy. Van der Velden's announcement, in essence, was that the gate had cracked.

The response was almost immediate, and it came from the only side of the gate that had anything to lose.

SAG-AFTRA's statement, issued on the Tuesday after the festival closed, called Tilly Norwood "not an actor." The full quote went further than the headlines tended to carry, and the next sentence was the one that actually mattered for the industry that read it: *"Signatory producers should be aware that they may not use synthetic performers without complying with our contractual obligations, which require notice and bargaining whenever a synthetic performer is going to be used."*²² That sentence — terse, contractual, almost dull — was the union reminding every signatory studio in Hollywood that the 2023 strike had already settled this question, and that the rules already on paper applied here too.

The U.K.'s actors' union Equity issued its own condemnation within hours, focusing less on the philosophical question of what an actor is and more on the practical one: where had the training data come from? Their general secretary put it more sharply than any of the other initial responses: *"We're at the stage in AI where so much data has been used that the original source becomes more and more unclear. And that's something that should worry every viewer, every working person, because that's not really the way our data should be used."*²³ If you build a model on the labour of working actors without consent, payment or attribution, you have, in their view, not created a new performer — you have repackaged a stolen one.

Within seventy-two hours, the public response from working actors followed. Emily Blunt called it "really, really scary." Whoopi Goldberg, on *The View*, dismissed the entire premise. Melissa Barrera, Kiersey Clemons, Lukas Gage — a roll-call of working performers, mostly the ones with enough security in their careers to be willing to say anything on the record at all — each took their position.²⁴

Eline Van der Velden, Particle6's founder, replied. *"To those who have expressed anger over the creation of our AI character Tilly Norwood: she is not a*

replacement for a human being, but a creative work — a piece of art.” The defence carried, in miniature, the entire shape of the argument the working actors were about to spend two years pushing back against. *Art. Creation. Not a replacement.* The whole rhetoric of the AI-native studios, in one sentence, on the Monday morning of October 2025.

The line that got repeated most often, in the threads and the green rooms and the union briefings I read that week, was a variant of: *we already gave you eighteen months of strike to settle this question.* The 2023 SAG-AFTRA strike had not been about AI in the abstract. It had been, in significant part, about digital replicas and the right of performers to control the use of their likeness in synthetic content. The deal that ended that strike had — supposedly — set the rules for the next era. Tilly Norwood, dropped onto a festival stage eighteen months later, was a test of whether those rules meant anything.

The answer was: they were going to have to be re-tested in public, on every individual case, for years.

And the tests kept coming. The last beat in the Tilly arc, as this edition goes to press, landed in July 2026 — ten months after Zurich, after the condemnations, after the death threats to her creator — when *Deadline* reported that Tilly Norwood had been cast to *star in a feature film*, a production called *Misaligned*.²⁵ Not a cameo, not a proof-of-concept short: a lead role. Whatever the union statements of October 2025 had established, they had not established that the synthetic actress would quietly go away. The gate Van der Velden said had cracked at Zurich was, less than a year later, being walked through — and the questions SAG-AFTRA’s contractual language was written to force (notice, bargaining, consent for the training data underneath) were about to be asked not in the abstract but on a specific production, with a specific call sheet. The film’s title, I assume, is somebody’s idea of a joke. It is also, unintentionally, the most accurate one-word description of where the industry stood on the question.

The model

What turned the Tilly Norwood weekend from a single-news-cycle controversy into the moment everyone in the creative industries started paying serious attention was that, on the Tuesday she was being condemned, OpenAI released Sora 2.

Sora 2 was not just an upgrade. It was, by OpenAI’s own framing, a step-change in three things at once: physical realism (a ball that bounces correctly off a backboard), audio integration (sound effects and synchronised dialogue baked in, not added in post) and what the company called “world state” — the ability to follow instructions across multiple shots while keeping the scene logically consistent.²⁶ If Sora 1, a year and a half earlier, had been the model that made people sit up and notice that AI

video was a thing, Sora 2 was the model that made people sit up and notice that AI video might be a *medium*.

The launch came with a second thing that, looking back, was the actually significant part: an invite-only iOS app, also called Sora, that worked like TikTok. You scrolled. You remixed. You did “cameos” — the in-app feature that let you drop a generated likeness of yourself, or a friend, or anyone you had a clip of, into the model’s output.²⁷

Within five days, the Sora app had hit a million downloads.²⁸

The thing I want you to hold in your head about this is the timeline. The model was announced on the Tuesday. The app launched the same week. By the Friday, it was on the front page of the App Store. By the following Monday, the first wave of celebrity deepfakes — Robin Williams’s daughter calling them “gross”; Michael Jackson generated into music videos he never made; dead historical figures being put through new dialogue by users who didn’t realise they were doing anything illegal because the app’s design hadn’t told them so²⁹ — was being written up in *The Guardian* and the *NBC News* tech vertical.³⁰

OpenAI’s own likeness-protection rules, *The Hollywood Reporter* noted that week, had a specific carve-out: dead celebrities and “historical figures” weren’t covered.³¹ The carve-out was treated, in the first days of the app, as a feature rather than a problem.

The clearest line of the entire launch week came not from OpenAI’s blog post but from *The Guardian*’s technology reporter, in the lead of their coverage of the Sora app’s first violent and racist outputs. “*In 2022, [the tech companies] would have made a big deal about how they were hiring content moderators ... In 2025, this is the year that tech companies have decided they don’t give a shit.*”³² I think that sentence will end up being remembered as one of the best one-line summaries of the moment. It was, if anything, optimistic about 2022. It was wholly accurate about 2025.

The reason this matters is not that Sora 2 was the first AI video model. It wasn’t even the best one, by some metrics, that month — Veo 3.1 from Google would land in mid-October with arguably more sophisticated cinematic controls.³³ The reason it matters is that the *combination* of a major model release and a consumer iOS app, in a single week, collapsed a distinction that had — until that week — kept the conversation about AI in the creative industries safely in the hands of the people who made the creative industries.

Until Sora’s iOS app, AI video was something that happened on a desktop, with a subscription, with a prompt window. After Sora’s iOS app, AI video was something

that happened on the phone of every teenager with an invite code, on a swipeable, remixable, social feed.

The line between “an AI tool a working filmmaker might use” and “a default app on a default phone” had been the line that held the cultural debate in shape for two years. That week, it didn’t.

The phrase

The line that has stayed with me from those first seven days came not from a celebrity, or from OpenAI’s launch page, or from a union release. It came from a small Florida public-radio station’s website. The headline read: *Kiss reality goodbye: AI-generated social media has arrived.*³⁴

I have read that headline a hundred times since the week it was written. It is a perfect sentence. It is also — and I think this is what made the early days of Sora 2 so vertiginous — *prematurely* true. Reality, as a category, did not end the week that Sora 2 launched. People still had real lives, real friends, real coffees in the morning, real bills at the end of the month. What ended, that week, was the easy assumption that what you saw on your phone had been made by somebody — a person, a team — with a recognisable connection to a recognisable place.

What’s striking about the WUFT piece, and the few hundred similar pieces that followed it, is that they were not, in the main, written by people in the AI industry. They were written by reporters at local stations, by columnists at regional papers, by the kind of journalist who covers schools and county budgets and the planning department. The collapse of the line between “made by a person” and “made by a machine” was being noticed first not by the experts but by the audience.

That, more than anything else that happened that week, was the thing I wrote in the first issue of the newsletter. *It’s time to take AI seriously.* The line, six months on, embarrasses me a bit — it sounds like the kind of thing you say when you don’t have anything else to say — but it was, on the morning of 6 October 2025, the only sentence I could write that felt like it was about the actual situation.

The director

There was one more piece of the early reaction that I want to flag, because it set the template for almost every “high-end” creative response that followed.

In late September, the website *No Film School* ran an interview with James Cameron in which he said — bluntly, on the record, in a quote that travelled — that AI was “never going to take the place” of humans in filmmaking. “Filmmaking is subconscious,” he said, “and can’t be quantified.”³⁵

Two months later, in promotion for *Avatar: Fire and Ash*, Cameron expanded on the point in a CBS *Sunday Morning* interview that became one of the most-shared

creative-industry stories of the entire winter. Asked about generative AI's ability to "make up a performance from scratch with a text prompt," Cameron said: "It's like, no. That's horrifying to me." And then, in a line I have used in talks and in arguments and in this book: *"The act of performance, the act of actually seeing an artist creating in real time, will become sacred."*³⁶

He went on, in a passage that almost never travelled with the headline quote: *"It also causes us to have to set our bar to a very disciplined level, and to continue to be out-of-the-box imaginative."* The whole interview, taken together, is not the doom-laden refusal the press cycle reduced it to. It is a working filmmaker articulating an argument about *craft* in an era of automated production — that the existence of cheap synthetic performance does not retire the human performer, it *raises the discipline required of the human performer*. The audience's bar moves up. The work that wins it has to move up too.

The companion line Cameron gave *Variety* the same month was even more telling: *"For years, there was this sense that, 'Oh, they're doing something strange with computers and they're replacing actors,' when in fact, once you really drill down and you see what we're doing, it's a celebration of the actor-director moment."*³⁷ Cameron has spent forty years building digital filmmaking. He is, as a working artist, perhaps the most sophisticated user of computer-aided performance in the history of cinema. His argument is not against AI; it is against AI *that displaces the moment in the room where an artist creates*. That distinction — *between AI that augments the actor-director relationship and AI that substitutes for it* — is the one most working creatives are trying to draw, and the one most press coverage of the AI debate still flattens.

What almost every report of those quotes glossed over — and what made them more interesting, not less — was that Cameron was, and still is, a board member of Stability AI.³⁸ He is not an opponent of AI in filmmaking. He is, by any reasonable definition, an investor in it. His position was not "no AI." His position was "no AI that replaces the actor in the room." That distinction is the one most working creatives, in my experience, are trying to make. The early coverage, looking for the clean villain-or-saviour story, mostly missed it.

I think Cameron's *Sunday Morning* line is the most important thing said by a working filmmaker in the entire six months I have been writing the newsletter. Not because it is right about everything — I think the "sacred" framing makes a smaller and more brittle claim than it sounds — but because it is the first time, in the AI era, that one of the people who *built* the apparatus of digital cinema in the 1990s and 2000s drew a line that he himself was prepared to defend.

The line, again, is not "no machines." Cameron has been a machine-builder for forty years. The line is: a human creates *in real time*, and that creation is the work.

Everything else is just delivery.

The state of the field, the week before

I want to spend a section on what was already in motion when the Tilly Norwood announcement happened, because the historians of this period are, I suspect, going to under-tell the *cumulative* story underneath the *catalysing* one. The week of 30 September 2025 was not the moment AI arrived in the creative industries. It was the moment the *audience* arrived in the AI debate. The infrastructure underneath had been forming, in a slow, uneven, partly-public, partly-private way, for at least two years before it.

Let me sketch the field.

By the time Sora 2 launched, the major AI-video model release cadence was running at roughly one significant model per fortnight. Runway had shipped Gen-4 in late 2024 and was deep into the public roll-out of Gen-4 Image-to-Video and the Workflows product by September 2025. Luma had released its Dream Machine consumer app; the Genie 3 demo from Google DeepMind in late summer 2025 had been described by *Time* as one of the year’s best inventions. Pika 2.0 was shipping. Higgsfield, which would close \$80M on a \$1.3B valuation by January 2026, was already on its third major product cycle. Hunyuan Video and Wan 2.2 — the open-source Chinese-built models from Tencent and Alibaba — had been freely available, on commodity GPUs, for months. Kling, the Kuaishou model that would, by mid-2026, be the model many professional filmmakers actually used for production-grade clips, had been in continuous public release. Sora 2 was the headline of the week. It was not, in any meaningful sense, the entire field. The field was already crowded.

The same is true of the studio-side adoption. Lionsgate had publicly partnered with Runway on a deal to produce AI-augmented studio films a year before *Futurism*’s headline-grabbing “crumbled into disaster” piece. Netflix’s quiet integration of AI tools into background-plate generation, animated short development and de-aging post-production work had been documented across 2024 and the first half of 2025. Disney’s *House of David* — the show whose creator would, in November 2025, defend the use of more than 350 AI-generated visual-effects shots in its second season — had been in production with that pipeline in place months before the Sora 2 launch. The major broadcast and streaming companies had been integrating AI under the hood at a pace that the public conversation had not yet caught up with. The Tilly Norwood week was, in part, the moment that pace became impossible to keep quiet.

The unions had been working on this even longer. The 2023 SAG-AFTRA strike had — through the Writers Guild and through the actors’ bargaining — produced contract language on “digital replicas” that was, by 2025, already two years old.

Equity in the UK had been running consultations and ballots through 2024 and into 2025. The Authors Guild’s class-action lawsuits against OpenAI had been filed in mid-2023 and were grinding through discovery. The European AI Act had been finalised in 2024 and was beginning to bite on copyright disclosures by the time of the Zurich announcement. The Music Performance Trust Fund’s emerging conversations about an AI-era levy mechanism — the Petrillo-template applied to neural-network outputs that I argue for in Chapter 2 and Chapter 16 — were already on the agenda at AFM Local 802 and at the UK Musicians’ Union long before Tilly walked on stage. The institutional response had pre-existed the cultural rupture. The cultural rupture is what made the institutional response politically possible.

The adoption telemetry on the platform side, in the months before the Sora 2 week, was already at the level that would later be revealed in Adobe’s MAX 2025 Creators’ Toolkit Report and the Stanford AI Index 2025. Firefly was already on track for its 22-billion-asset milestone. ChatGPT, by the time Sora 2 launched, was already at roughly 700–800 million weekly active users on its way to 900 million. The 86% of global creators who reported using generative AI in their workflow in the Adobe survey — published a few weeks after the Tilly Norwood week, but capturing data collected before it — was a number the platform-companies had been quietly watching for months. The Adobe survey did not produce the adoption. It documented the adoption that was already complete by the time the public was paying attention.

On the *consumer* side, the Sora app was, again, the catalyst rather than the inventor of the dynamic. The TikTok-style consumer surface for AI generation had been visible for at least a year. ByteDance’s Dreamina and CapCut tools had been integrating Seedance, Seedream and the wider ByteDance generative stack into consumer-facing video editing through 2024 and into 2025. The Sora app’s million-downloads-in-five-days number — the headline that defined the week’s consumer dynamics — landed inside a market that had been *prepared for it* by Dreamina, CapCut, the Krea consumer app and the Suno and Udio consumer-facing music platforms. What was new about the Sora app was not the *concept* of swipeable AI generation. What was new was that an American flagship-AI company had chosen to ship it as a *primary* consumer surface alongside the model.

Particle6, finally, had been building Tilly Norwood for the better part of a year before the Zurich panel. Van der Velden’s Instagram and TikTok rollouts had been running since the previous summer. The character had a follower count, a posting cadence, a slowly-built visual identity, a small but real fanbase that engaged with her as if she were a person. The Particle6 strategy — the deliberate cultivation of a *parasocial* relationship between an AI character and a human audience, in the months before the studio-system pitch — was a play borrowed from the playbook of the Korean virtual-idol economy, the Japanese vocaloid scene, and the long history

of cultivated personas in the influencer industry. What was new about Tilly was not the *idea* of a synthetic personality with a fanbase. What was new was the framing of that synthetic personality as a *casting option* for legacy film and television. The Zurich announcement was the moment the parasocial-character economy and the working-actor economy were proposed, on a festival stage, as overlapping markets. The reaction was, in retrospect, the audience and the union refusing to let the markets overlap.

What all of this means, when you stack it together, is that the Tilly Norwood / Sora 2 week did not invent the AI moment in the creative industries. It *named* it. It made it impossible to keep treating AI as a technical category that working creatives could opt out of. The model release was overdue. The Particle6 announcement was overdue. The union response was already drafted in some form. The audience reaction was already, on the slop-ceiling logic I lay out in [Chapter 5](#), structurally inevitable. What the week did was *force everything to happen in public, at the same time, in front of an audience that had not previously been part of the conversation*.

I think that is the deeper reason the week mattered. The conditions were ripe. The catalyst was small. The reaction was big. The change in the *visibility* of the AI transition — from a back-room toolchain conversation to a front-page audience question — was, by my reading, the actual watershed. Everything else in this book is downstream of that visibility shift.

The watershed under the watershed

I said earlier that the actual watershed of that week was not the Tilly Norwood announcement, and not the Sora 2 launch, and not the union responses, and not the Cameron quote. The actual watershed was something underneath all four.

For two years, the conversation about AI in the creative industries had been a conversation among insiders. Toolmakers talked to creators. Creators talked to studios. Studios talked to unions. Unions talked to governments. The general public — the audience for the things being made — had largely been a backdrop. They had been the people *for whom* this argument was happening, not the people in the argument.

In the week of Sora 2, that ended.

The Sora app put a generative video tool on the phone of anyone with an iOS device and an invite code, and the invite codes were not hard to come by. The Tilly Norwood announcement put the abstract concept of “a synthetic actress” into the *Daily Mail*, *The Guardian*, *The View*, the breakfast television circuit, and three different morning radio shows I happened to be listening to that week. The Cameron quote, when it came, ran on every wire service that covers the entertainment business.

The audience joined the argument. Not as a side, but as a participant.

And once the audience is in the argument, the argument changes. It is no longer about what the unions can negotiate, what the studios will adopt, what the toolmakers will ship. It is about what the people who watch films and listen to music and play games and scroll feeds will accept, demand, refuse and forgive.

The rest of this book — the seventeen chapters that follow this one — is, in one way or another, an account of what the audience has been doing with the argument since it became theirs. The artists' boycotts and the streaming platforms' counter-moves; the eleven and a half thousand UK citizens who turned up to a government consultation; the 50,000 AI-generated tracks uploaded to Deezer every day and the 1 to 3 percent of streams those tracks actually got;³⁹ the death threats sent to Tilly Norwood's creator in early 2026; the moment in the middle of January 2026 when the U.S. actors' union and SAG-AFTRA went back to the negotiating table because the audience, having looked at the new landscape, had decided what it wanted.

The week of 30 September 2025 was the last week before any of that.

I started writing the *Dream Machine* newsletter the Monday after. By the time I sent the first edition out, the conversation had already moved on.

Chapter 2 — A History of Resistance

I want to start this chapter with a piece of writing.

“Sweeping across the country with the speed of a transient fashion in slang or Panama hats, political war cries or popular novels, comes now the mechanical device to sing for us a song or play for us a piano, in substitute for human skill, intelligence, and soul... Let us not hamper it with a machine that tells the story day by day, without variation, without soul, barren of the joy, the passion, the ardor that is the inheritance of man alone.

Singing will no longer be a fine accomplishment; vocal exercises, so important a factor in the curriculum of physical culture, will be out of vogue! Then what of the national throat? Will it not weaken? What of the national chest? Will it not shrink?

When a mother can turn on the phonograph with the same ease that she applies to the electric light, will she croon her baby to slumber with sweet lullabys, or will the infant be put to sleep by machinery? Children are naturally imitative, and if, in their infancy, they hear only phonographs, will they not sing, if they sing at all, in imitation and finally become simply human phonographs — without soul or expression?”⁴⁰

Read that paragraph once more, and try, before I tell you when it was published, to date it.

If you guessed *autumn 2024 trade-press editorial on Suno*, you would not be alone. The language sits comfortably alongside the take-pieces about AI music slop that filled the music press in the months I started writing the newsletter. *Machinery in substitute for human skill, intelligence and soul. The story told day by day without variation. Children growing up as imitative human phonographs, without soul or expression.* Every clause has a 2024–25 equivalent.

The essay is not from 2024. It was published in **September 1906**, in *Appleton’s Magazine*, by **John Philip Sousa** — the most popular bandleader in the United States at the time — and it is titled *“The Menace of Mechanical Music.”*⁴¹ The machine Sousa was warning his readers against was the phonograph. The “soul-barren” recording technology Sousa feared was Edison’s flat disc, spinning at 78 rpm, playing back music captured by a brass horn.

The same essay was — almost word for word, with a few changes in the names of the machines — written about the player piano in 1900, the microphone in 1932, the synthesiser in 1980, the drum machine in 1991, Auto-Tune in 1998, non-linear video editing in the early 1990s, the digital camera in the 2000s, the smartphone-as-camera in the 2010s, and generative AI in 2023, 2024 and 2025.

That is the subject of this chapter. The structured, recurring, almost ritual *pattern* by which every major creative-technology introduction in the modern era has been

received by its working practitioners — and what that pattern, eighteen iterations later, tells us about how to think about the AI moment the rest of this book describes.

I want to be careful about how I do this. The historical-analogy move is, in tech writing, a famously cheap one. “*Every disruptive technology has been resisted; therefore your resistance to this technology is wrong*” is the rhetorical operating system of two decades of platform-company keynotes, and it has been wielded so dishonestly that working creatives in 2026 are right to be suspicious of any version of it. I am not, in this chapter, making that argument. I am making a *more specific* one. The historical pattern has — across twenty distinct technologies in the period from 1839 to 2022 — a recognisable five-act shape, and the *five-act shape* is what the pattern tells you. Not that resistance is wrong. Not that the new tool will be fine. Something more useful than either: which institutional moves work, which fail, what gets preserved, what gets lost, where the *new* creative forms come from, and what the working practitioner’s actual leverage in the period is.

The history is, in other words, *operationally* informative. That is the use I want to make of it.

The five-act curve

Every resistance I look at in this chapter — and there are many more I do not have room for — moves through roughly the same five stages, in something like the same order.

Act One: Ridicule. The new tool is dismissed as a toy. It cannot compete. It sounds awful, looks crude, has the wrong specifications. The serious practitioners are unconcerned because the work it produces is not, on inspection, work. The Roland TR-808 drum machine, released in 1980 and a commercial failure, was reviewed as “toy robot drums.” The Canon 5D Mark II, the DSLR that started the death of dedicated cinema cameras, was — in 2008 — dismissed as a stills camera with a video gimmick. Auto-Tune, between 1997 and roughly 2003, was used as an *undisclosed* studio tool because no working singer wanted to admit that their pitch was being machine-corrected.

Act Two: Moral panic. The new tool is reframed as a threat to public morals, aesthetic standards, or the integrity of the form. It is degenerate. It is theft. It is not real *singing*, not real *photography*, not real *art*. The 1932 sermon by **Cardinal O’Connell** of Boston that crystallised this stage for the microphone — that crooning was “a degenerate form of singing,” that “no true American man would practice this base art,” and that crooners were “whiners and bleaters defiling the air” — is unimprovable as a template.⁴² “*Imbecile slush*,” O’Connell called it, in language that anyone who has read a 2024 anti-AI op-ed will recognise. The 1991 federal court ruling in *Grand Upright Music v. Warner Bros.* — the case in which Biz Markie was

sued for sampling Gilbert O'Sullivan — opened with the words “*Thou shalt not steal*,” quoting the Seventh Commandment, in a US copyright opinion.⁴³ In 2025, the moral-panic stage of AI is mostly behind us; the analogous *Thou shalt not steal* language is in the UK's 88% copyright consultation response, the *Stealing Our Work Is Not Innovation* declaration, and the union statements quoted throughout this book.

Act Three: Existential professional alarm. The displaced practitioners realise the tool is not a toy and not just morally suspect — it is structural. It is going to take their work, change their craft, and reshape the institutions that support them. The classic statement is **Phil Tippett's**, the stop-motion master who saw ILM's first digital test of a *Jurassic Park* T. rex in 1992 and said: “*I think I'm extinct.*”⁴⁴ Spielberg liked the line enough to put a paraphrase of it in the film. Tippett was right about himself in some local sense — his go-motion craft did not survive *Jurassic Park's* release. He was wrong about himself in a wider one — his studio went on to produce digital animation work for *Starship Troopers* and is still operating in 2026. Both readings can be true at the same time. Most of the existential-alarm moments work like this.

Running alongside the alarmed practitioner, in every iteration of the curve, is a second recurring voice: the elder statesman who declares the whole thing unstoppable. In July 2026 that part was taken — with almost suspicious neatness — by **George Lucas**, who described AI as the inevitable future of filmmaking, said it makes the work “easier” and that “there's nothing you can do about it,” and likened AI sceptics to people who wanted to keep using horses after the invention of the car.⁴⁵ Horses and carts is the oldest analogy in the resistance literature, and Lucas reached for it unprompted. The irony, which the trade press did not miss, is that the man delivering the you-can't-stop-it line in 2026 is the man who founded ILM — the company whose digital T. rex test made Phil Tippett feel extinct in 1992. Lucas has stood on the delivering end of this pattern before, which is presumably why he recognises it. Both voices — *I think I'm extinct* and *there's nothing you can do about it* — are permanent fixtures of the curve, and the history's contribution is to note that neither of them, on its own, has ever accurately described what happened next.

Act Four: Institutional and legal counter-attack. Unions strike. Lawmakers legislate. Lawsuits get filed. The 1942 and 1948 **Petrillo strikes** — when the American Federation of Musicians, led by James Caesar Petrillo, refused to record for the major labels — are the canonical version of this stage, and I will spend longer on them in a moment. The UK Musicians' Union's **1980 “Massacre of the Musicians”** BBC strike and its **1982 motion to ban synthesisers** are the British version. The 1991 *Grand Upright* ruling that sampling was theft is the legal version. The 2007 *Viacom v. YouTube* \$1bn lawsuit is the platform-distribution version. The

2023 SAG-AFTRA and WGA strikes — which gave us the contract language that frames Chapter 1 — are the most recent before the period this book covers.

Act Five: Settlement. The dust settles into one of three forms. The displaced craft *dies*: miniature painting after 1840, hand-drawn feature animation after *Toy Story*, the photo-lab business after digital. The new tool is *taxed and the revenue redistributed*: the Petrillo settlement created the **Music Performance Trust Fund**, still distributing payments to live musicians in 2026; the DMCA Section 512 plus YouTube’s Content ID created a parallel pool of platform-paid royalties for the music industry; needletime in the UK forced the BBC to pay for live sessions until 1988. Or — most often — the two creative categories *coexist*: photography did not kill painting, it forced painting toward what photography could not do (Impressionism, abstraction); the microphone did not kill singing, it *redefined* what counts as singing; sampling did not kill composition, it redefined what counts as composition.

That is the shape. It is, by my count, the shape of every single technological transition in the creative industries in the period 1839 to 2022. It is happening, around AI, right now — and where we are on the curve is part of what this chapter is for.

Let me walk briefly through a handful of the cases, because the *texture* of the historical record is, I think, more useful than the abstract pattern alone.

The photograph (1839)

The daguerreotype was unveiled in Paris on 7 January 1839 and publicly described to the Académie des Sciences on 19 August. By 1849, roughly 100,000 daguerreotypes had been produced in Paris alone; by 1861, about 33,000 people in Paris were making their living from photography and photographic supplies. The professional class most immediately wiped out was the **portrait miniaturist** — painters of small ivory-based likenesses, the working photographers of the pre-photographic age. They could not compete on speed, price or fidelity. Within a single working generation, the craft was effectively gone.

The most articulate resistance was not from the miniaturists themselves but from the literary intelligentsia. **Charles Baudelaire’s** 1859 essay “*The Modern Public and Photography*,” published in the *Revue Française*, made the case in language a 2024 AI-sceptic would recognise: “*this industry, by invading the territories of art, has become art’s most mortal enemy.*” And, harder: “*The photographic industry was the refuge of all failed painters, too ill-equipped or too lazy to complete their studies.*”¹⁶

The famous **Paul Delaroche** line — “*From today, painting is dead*” — is, the historians who have looked carefully tell us, almost certainly apocryphal. The

earliest sourced version of it appears in an 1873 survey, thirty-four years after Delaroche supposedly said it; the painter's own contemporary writing on the daguerreotype called it "*an immense service to the arts*," and he continued painting until his death in 1856.⁴⁷ The story has outlived the saying. (This is itself a pattern. I will come back to it.)

The settlement took seventy years. Alfred Stieglitz founded the Photo-Secession on 17 February 1902. *Camera Work* ran from 1903 to 1917. The "291" gallery opened in 1905. MoMA established the first photography department at a major museum on 31 December 1940. From invention to *full institutional acceptance of the new form as art*: about a century. The compensating gain: the entire history of modern photography as a fine-art tradition, an industrial portrait business, a documentary-journalism profession, and — eventually — the cultural substrate on which the smartphone-camera moment of the 2010s rests.

What painting did, meanwhile, was redefine itself. Impressionism (light, atmosphere, the subjective moment), Post-Impressionism (the interior state), Cubism (multiple viewpoints), and ultimately abstraction — every one of these moves makes more sense if you read them as painting's response to the daguerreotype's having taken representation. The standard art-history reading, which I think is correct, is that *photography liberated painting from the burden of representation*. The fear-namers were locally right and structurally wrong. Painting did not die. It became something else.

The phonograph (1877–1920s)

I have already quoted Sousa. The whole 1906 *Appleton's* essay is worth reading; almost every paragraph of it could be republished, with names changed, in 2025.

The institutional response to the phonograph — and this is the part of the story that working creatives in 2026 need to know — came in the form of the **1909 Copyright Act**, the first statutory acknowledgement in US law that *machine reproduction of human creative work required a legal regime*. The 1909 Act created the **compulsory mechanical licence** for recorded music, the structural ancestor of every machine-licensing argument we are now having about AI training. Sousa, in part because his lobbying helped pass it, prospered in the recording era; his compositions still generate royalties under that licensing structure today.

The phonograph also created entirely new occupations that the parlour-music economy could not have anticipated. The recording engineer. The A&R executive. The producer. The mastering engineer. The sleeve designer. The pressing-plant operator. The retail-store buyer. The pop *single* as a commercial form. Jazz on record. The LP. The concept album. The bedroom-studio that, in the 2000s, would take all of those occupations apart again. Sousa got the *parlour* prediction right — recorded music did dent amateur home music-making — and the *industry*

prediction completely wrong. The recording industry was the largest expansion of working-musician employment in the history of music, and it was made possible by the technology Sousa thought would destroy it.

The Petrillo bans (1942, 1948)

The **American Federation of Musicians** had, by the 1940s, watched the phonograph, the talking picture (which alone wiped out roughly 22,000 cinema-orchestra jobs in the US in 1927), and commercial radio progressively displace live performance. **James Caesar Petrillo**, AFM president from 1940, did the thing every union threatened by a creative technology has tried, with varying success, to do since: he turned off the recording machine.

On 1 August 1942, AFM members stopped recording. The strike lasted **27 months**. Decca settled in 1943; RCA Victor and Columbia in November 1944. The settlement: a per-record royalty paid into an AFM fund for unemployed musicians.

On 1 January 1948 Petrillo did it again, triggered by the Taft-Hartley Act outlawing the first royalty arrangement. The 1948 settlement created the **Music Performance Trust Fund** under Section 302 of Taft-Hartley — a jointly-administered labour-management fund paid into by the labels and broadcasters, used to subsidise free live performances by working musicians. The MPTF still exists. It still pays out, in 2026, several million dollars a year for live music.⁴⁸

I want to dwell on this for a moment, because the Petrillo settlement is — by some distance — the **most operationally important precedent** for how the AI debate could land, and almost nobody in the current creative-AI conversation talks about it.

The Petrillo template has four parts. *One*, the displacing technology is not banned. It is allowed to displace. *Two*, the platform owner pays an ongoing per-unit tribute to the displaced labour pool. *Three*, the tribute is collected centrally, by a joint labour-management body, not negotiated individual-by-individual. *Four*, the tribute is paid out to subsidise the *displaced creative practice itself* — live music, in this case — keeping it alive as a category even as the market for it shrinks.

The SAG-AFTRA *Tilly Tax* provisions in the 2026 contract, the UK 88% licensing-by-default proposal, the *Creative Weight Attribution* musical-AI infrastructure I described in [Chapter 5](#), the C2PA / SynthID provenance stack from [Chapter 13](#) — these are all, on inspection, *attempts to reconstruct the Petrillo template for AI*. Per-unit tribute. Joint collection. Redistribution to the displaced practice. The mechanism is the same. The political question is whether the platforms will accept it.

The mechanism worked once. It can work again. The pattern of the resistance that fails — the 1982 UK Musicians' Union motion to *ban synthesisers outright*, the European *Right to be Forgotten* style absolutism on training data — is the pattern of

resistance that tries to legislate against the machine rather than to tax it. Levy beats ban, every time. Royalty pool beats injunction. *Mechanism, not prohibition.*

The microphone (1932)

Cardinal O’Connell’s January 1932 sermon to 3,000 men of the Holy Name Society of Boston is the single funniest moment in the resistance literature. *Crooning* — the conversational, intimate, microphone-enabled vocal style that **Bing Crosby**, Rudy Vallée and others had built into a mass commercial form by the late 1920s — was, for O’Connell, “a degenerate form of singing. No true American man would practice this base art... If you will listen closely [to crooners’ songs] you will discern the basest appeal to sex emotion in the young.” Crooners were “whiners and bleaters defiling the air.” Their work was “imbecile slush.”⁴⁹

The cultural fight was effectively over by the end of the decade. Bing Crosby was the biggest male voice in America. By the late 1930s no popular vocalist *not* trained in microphone technique could make a competitive career.

What is interesting about the microphone case, for our purposes, is *what it did to the underlying definition of the craft*. Before the microphone, *good singing* meant volume, projection, throat technique — the operatic, theatrical, music-hall tradition that O’Connell had grown up inside. After the microphone, *good singing* meant timbre, intimacy, breath control at low dynamics, conversational diction — a fundamentally different skill set. Almost every popular vocalist since 1935 has been a “crooner” in the technical sense, including those who would not call themselves that. *The microphone redefined what counted as singing*. The vocalists who refused the microphone are mostly remembered as period figures. The ones who absorbed it defined the rest of the century of popular music.

This is the deeper pattern I will come back to. Resistance, in the named-fear form, almost always defends *the existing definition* of the craft. The settlement, almost always, *redefines* the craft. The fear is real and the language is sincere; what’s actually happening is bigger than what the fear is naming.

The synthesiser (1980–82)

The UK part of this story is the cleanest because it generated public archives. The **Yamaha DX7**, released in 1983, put credible electric piano, brass, strings, marimba and dozens of other sounds into a single keyboard at a price point — about £1,500 — accessible to working session players. The DX7 was used on roughly **40% of US Billboard Hot 100 #1 singles in 1986**. Session keyboard players and orchestral string sections previously hired for adverts, TV scoring, library music and pop sessions were directly displaced.

The **Musicians’ Union** of the UK responded in two stages.

1980, the Massacre of the Musicians. The BBC announced in March 1980 that it would cut 172 staff orchestra posts and disband five of its eleven in-house orchestras. 83% of MU BBC members voted to strike. The strike began 16 May 1980. **The First Night of the Proms was cancelled for the first time in its history.** The strike ran until 1 August 1980, ending with a compromise: the BBC Northern Ireland Orchestra and BBC Midland Radio Orchestra were disbanded as planned; the others survived.⁵⁰

1982, the synthesiser ban motion. On 23 May 1982 — by coincidence Bob Moog’s birthday — the MU’s Central London Branch passed a motion to ban synthesiser, drum-machine and electronic-device use by union members entirely.⁵¹ The trigger was that Barry Manilow’s UK tour had replaced its string section with synth players. The motion was never adopted as full union policy. The MU’s Executive Committee passed a more measured resolution in November 1982. A breakaway “Union of Sound Synthesists” was formed. *Top of the Pops*, for a period, required bands to record their backing tracks the afternoon before the show “to prove they could actually play it.” None of it held. The DX7 sold over 200,000 units. By 1990 the cultural debate was over and the synthesiser was, simply, another instrument.

What the *MU* got wrong, in retrospect, was the same thing every resistance-by-prohibition gets wrong: they tried to ban the tool rather than to tax it. There was no Petrillo-style fund. There was no per-output levy on synthesiser use. There was no royalty pool subsidising live orchestral work. The MU’s defensive posture preserved nothing structurally and lost the cultural argument decisively. The session-musician economy contracted. Some of the displaced players retrained as programmers and prospered. Others did not. **Trevor Horn** — the Buggles who, three years before the dispute, had sung *Video Killed the Radio Star* — became the defining pop producer of the 1980s.

The lesson, for the working creative in 2026, is not that the MU was wrong to resist. The MU was *right* to read the displacement signal three years ahead of the rest of the industry. The MU was wrong about *what kind* of resistance to mount. Prohibition was always going to lose. A levy-and-pool argument — what an *AI-era Petrillo* could look like — might have held.

Drum machines, sampling and the lawsuits (1980s–2000s)

The Roland TR-808 (1980) and TR-909 (1983) both *failed* commercially on release. The 808 was criticised as “toy robot drums”; the 909 was reviewed as “still sounds like a drum machine, instead of a machine playing drums.” Both became cult instruments only after being dumped at secondhand prices to young hip-hop and dance producers in the mid-1980s. The 808’s hand-clap, snare and signature deep kick are now the defining percussion sounds of contemporary popular music. The

same dismissed-then-canonised arc is, in 2026, beginning to play out for Suno and Udio at the consumer-music end of the spectrum.

The legal resistance to sampling produced two rulings every working creative should know about, because they are the cleanest available templates for how the courts may treat AI-training disputes.

Grand Upright Music v. Warner Bros. (S.D.N.Y. 1991). Biz Markie sampled three bars of Gilbert O’Sullivan’s *Alone Again (Naturally)* without clearing it. Judge Kevin Thomas Duffy’s opinion opened with the words “*Thou shalt not steal*,” quoting the Seventh Commandment, in a US federal court ruling. He referred the case to the US Attorney for potential *criminal* investigation. The album was pulled. The sample-dense Bomb Squad / Public Enemy style of production it had been built on became commercially impossible.⁵²

Bridgeport Music v. Dimension Films (6th Cir. 2005). NWA’s *100 Miles and Runnin’* sampled a two-second guitar chord from Funkadelic’s *Get Off Your Ass and Jam*, looped and pitched down. The Sixth Circuit eliminated the *de minimis* defence for sampling sound recordings and issued the rule that has, for twenty years, defined how clearance works: “**Get a license or do not sample.**”⁵³

Hip-hop, of course, did not die. It became the dominant global popular music form. What changed was that the *aesthetic* of dense, layered, sample-heavy production gave way to a more clearance-friendly style. The Bomb Squad lineage continued, but on different terms.

The AI-training analogy here is direct, and I would commend the rulings to anyone trying to think clearly about *UMG v. Anthropic* and the cases that will follow it. *Thou shalt not steal* and *Get a license or do not sample* are not, despite their archaic phrasing, particularly anti-technology rulings. They are *pro-licensing* rulings. They say: the tool can be used, but the inputs have to be paid for. That is what working creatives are asking for in the UK 88% and the *Stealing Our Work Is Not Innovation* declaration. The line of legal reasoning is already in the books.

Auto-Tune (1998–2010)

I want to spend less time on Auto-Tune than the dossier supports, because the case is so clean it almost makes itself. Andy Hildebrand, a former Exxon seismic-data engineer, released Auto-Tune in 1997. **Cher’s *Believe*** (1998), produced by Mark Taylor and Brian Rawling with the retune speed maxed out, produced the now-iconic “Cher effect” — the audible warble. Taylor and Rawling claimed it was a vocoder for several years to protect the trick.

T-Pain made the effect his signature from 2005 onwards and was rewarded with **Jay-Z’s “D.O.A. (Death of Auto-Tune)”** — released June 2009 on *The Blueprint 3* — a direct moral-panic attack on what Auto-Tune was doing to vocal authenticity.

TIME magazine's 50 Worst Inventions list in 2010 ranked Auto-Tune at #15: software that “can make bad singers sound good, and really bad singers sound like robots.”⁵⁴

Sixteen years later, Auto-Tune is on every pop vocal you hear, used as both correction and effect. Bon Iver's *22, A Million* (2016) deployed it as a self-conscious aesthetic instrument. Billie Eilish has used it across her career, transparently. The cultural rehabilitation is complete. The moral-panic stage, in retrospect, looks parochial.

But notice what *did* happen. The microphone-era settlement — *you have to be able to sing* — was, in the Auto-Tune era, definitively broken. The new settlement is *you have to be able to perform the post-corrected vocal as a self-conscious artistic choice*. The redefinition is real. Cardinal O'Connell, transported eighty years forward, would have hated Auto-Tune for exactly the same reasons he hated crooning, and would have been just as wrong about it.

Non-linear editing and the death of the splice

Avid Media Composer launched in 1989. Through the early 1990s it displaced the Moviolas, Steenbecks and KEM flatbeds that working film editors had used for sixty years. The American Cinema Editors did not strike. The Motion Picture Editors Guild did not stop the transition. The settlement was generational: editors trained on film retired; editors trained on Avid became standard.

The witness I want to quote here is **Walter Murch**, ACE — possibly the most respected film editor of the last fifty years. Murch edited *The Conversation* and *Apocalypse Now* on physical film. He edited *The English Patient* on Avid, winning Oscars for both picture and sound. He then edited Anthony Minghella's *Cold Mountain* (2003) on **Apple Final Cut Pro running on commodity Power Mac G4 hardware** — for a \$79m feature. The story is in Charles Koppelman's book *Behind the Seen* (Peachpit, 2004).⁵⁵ The decision saved about \$1m versus an equivalent Avid rental, and Murch chose it on practical grounds. The editor who literally wrote the textbook on film editing — *In the Blink of an Eye*, the canonical philosophical text on the cut — was, by 2003, working on the tool the editing-room conservatives were warning the industry against.

The compensating gain, from the non-linear editing transition, was that the toolkit shipped on a laptop. Indie cinema benefited enormously. The grammar of the cut accelerated — the average shot length in Hollywood drama dropped from roughly ten seconds in the 1960s to roughly four seconds by the 2000s, a change made trivial by NLE that would have been physically punishing to execute on a Moviola. The contemporary visual grammar that ranges from *The Bourne Identity*'s hyper-cuts to Wong Kar-wai's asynchronous editing aesthetic to TikTok's stitched, layered, fast-

moving native form is, in operational terms, *what non-linear editing made possible*. The form the new tool enabled was bigger than the form it replaced.

Kodak (1975–2012)

Steven Sasson, an engineer at Kodak, built **the first digital camera prototype in December 1975**. 0.01 megapixel, black and white, the size of a toaster, with a 23-second save time to magnetic tape. Sasson's own description, in multiple interviews, is that the executive response to the demo was *curious but concerned about the implications for film*. Kodak suppressed the project to protect its film business.⁵⁶

Kodak's peak headcount was about 145,000 in the early 1980s. At the company's Chapter 11 bankruptcy filing on **19 January 2012**, the workforce was about **19,000**. A nearly \$7bn liability stack. A company that had invented the technology that would destroy it, and then suppressed that technology, and then been destroyed by it anyway when the rest of the market — Nikon, Canon, Sony — built around the patents Kodak had not commercialised.

The Kodak story is the standard cautionary tale for incumbents about to be disrupted by the technology they themselves built. It is being told, with increasing force, about the legacy entertainment industries in 2026. *The companies that already have the IP, the audience and the distribution are, by structural inheritance, in the position Kodak was in in 1990*. What we will find out, in the rest of this decade, is whether they have learned the Kodak lesson — that the new technology is going to displace the old whether or not you commercialise it, so you may as well be the company that does — or whether the next few years will produce the entertainment-industry equivalent of the 2012 bankruptcy filing.

Photoshop, social media, smartphones — a faster montage

Three more cases, briefly, because the pattern is by this point clear and the texture matters more than the recital.

Photoshop and photojournalism (1990–2015). The famous **Brian Walski** firing, 1 April 2003 — LA Times staff photographer dismissed by satellite phone from Iraq for compositing two combat photographs into one stronger image — is the single canonical example.⁵⁷ World Press Photo's response, over the following decade, was to issue increasingly tight ethics rules and disqualify increasing numbers of finalists. The truth-claim of the photograph survived institutionally at AP, Reuters, *National Geographic* and *Magnum*. It did not survive on social media. The structural answer was C2PA — the content-authenticity infrastructure that now sits underneath the AI provenance stack in [Chapter 13](#). The 1990s photojournalism debate was, on inspection, *the first round of the C2PA conversation* the AI era is now finishing.

Social media (2005–2014). The defining institutional fight was **Viacom v. YouTube**: \$1bn in claimed damages, filed in March 2007, decided in favour of YouTube and the DMCA Section 512 safe harbour in 2010 and again in 2013, settled out of court in **March 2014 with no money changing hands** — a complete Google win.⁵⁸ The settlement was effectively *automated*: YouTube's Content ID, deployed from 2007 onwards, has paid out billions in royalties to rights-holders through algorithmic matching. The DMCA safe harbour that decided *Viacom* — passed in 1998, in response to the previous resistance pattern — became the structural framework of the *creator economy*. Power-law distribution, \$100bn+ in aggregate value by 2022, MrBeast at 450m subscribers and \$600–700m in 2024 revenue. The platform shape of the audience contract in 2026, on which every AI-distribution question rests, was built by Content ID.

Smartphones (2007 onwards). The most dramatic collapse of an established consumer-creative-tool category in modern history. Compact-camera shipments: **108.8 million units in 2010, 3.01 million in 2021**. A 97% drop in roughly a decade.⁵⁹ The iPhone, the Galaxy and the Pixel did to the camera industry what Sousa thought the phonograph would do to amateur music — and Sousa, in retrospect, was about a hundred years early on his own argument. The smartphone-camera moment is also, on inspection, the dress rehearsal for the AI-image moment. The same arguments — *it isn't real photography, the computational layer is doing the work, not the human, the dedicated camera is the moral artefact* — were made in 2010 and 2015 about smartphones. They were also lost in 2010 and 2015 about smartphones. Computational photography is, today, where most of the meaningful innovation in the image-making craft is happening. *AI photography*, in 2026, is not the new thing. It is the next iteration of the thing that the smartphone-camera made inevitable fifteen years ago.

Five patterns

I have, by this point, walked through ten of the twenty cases the dossier behind this chapter covers. The pattern recurs. Let me name the five structural features I think the working creative in 2026 most needs to internalise.

Pattern one: the curve has a predictable shape. Ridicule → moral panic → existential alarm → institutional counter-attack → settlement. AI in 2025–26 is, on my read, somewhere in late act three and early act four. The existential alarm has been voiced. The institutional counter-attack — the SAG-AFTRA strike, the WGA contract, the UK 88%, the *UMG v. Anthropic* suit, the Cannes Disclosure Standard, the Sundance literacy initiative — is well underway. The settlement is starting to form but is not yet stable. The next eighteen months will, on the historical pattern, *settle the form of the next decade's industry*. This is exactly why this period matters.

Pattern two: the named fear is always mis-named. What practitioners say they fear — loss of *soul*, loss of *authenticity*, loss of *the real* — is almost never what actually happens. What actually happens is a *redefinition* of the underlying creative category. The microphone redefined *singing*. The photograph redefined *painting*. The sampler redefined *composition*. The Avid redefined *editing*. The smartphone redefined *photography*. The fear is framed as a defence of the thing. What is actually at stake is the *definition* of the thing. The fear-namers usually win the surface argument and lose the definitional one. *The AI debate is, on the structural reading, an argument over what counts as authorship, performance, writing, photography and composition in the next decade.* That is the real fight. The named-fear version of it — *AI will steal our jobs* — is true but partial.

Pattern three: the institutional moves that work are levy-and-pool; the moves that fail are prohibition. The Petrillo settlement (MPTF, 1948), the DMCA safe harbour plus Content ID (1998 plus 2007), the 1909 Copyright Act mechanical licence, the eventual streaming-loudness-normalisation truce — these *worked*, in the limited but real sense that they extracted ongoing transfer payments from the new medium to the displaced labour pool, or that they restructured the aesthetic equilibrium. The institutional moves that *failed* are the prohibitions: the MU 1982 synthesiser ban, the 1991 *Grand Upright* effective ban on dense sampling, the European *Right to be Forgotten* style absolutism. **Levy beats ban. Royalty pool beats injunction. Mechanism beats prohibition.** For the AI-training fight: a per-output levy distributed to a creators' fund is in the workable category. A ban on AI training is in the unworkable category. The 88% — by my reading — is closer to the workable category than the politics around it have so far recognised.

Pattern four: the resisters are usually right about the local loss and wrong about the total loss. Miniature painting died. Hand-drawn feature animation died (*Toy Story* in 1995; Disney's Florida 2D studio closed on 12 January 2004). The professional recording-studio mid-tier died. Photo-processing labs died. Staff photographer jobs at US newspapers — collapsed (the *Chicago Sun-Times* infamously laid off its entire photo staff in May 2013). But *more people make moving images today than ever made them; more people make music today than ever made music; more photographs are taken on a given Sunday than were taken in the entire nineteenth century.* The aggregate creative-labour pool grew. The fear is always articulated by the *displaced* cohort. The compensating gains accrue to a different cohort, who don't yet exist when the fear is articulated, and whose names — by definition — are not yet in the trade press. This is why the resisters can be both *factually right* about their own situation and *structurally wrong* about the form they are trying to protect.

Pattern five: the cultural symbol outlives the cultural anxiety. *Video Killed the Radio Star* was the Buggles' 1979 lament for the death of a form. MTV used it as

its launch trumpet at 12:01 a.m. on 1 August 1981. Then radio survived; then MTV died; the song is on TikTok. The artefact *about* the death of a form outlasted both the form it threatened and the form that did the threatening. Phil Tippett’s “*I think I’m extinct*” in 1992, **Trevor Horn’s** “*video killed the radio star*” in 1979, **Cardinal O’Connell’s** “*imbecile slush*” in 1932 — these crystallise an anxiety into a piece of language so vivid that it survives whatever it was anxious about. The named fear becomes the source material for the next generation’s art. *The anxiety is the art*. The *Tilly Norwood week* of late September 2025, with Whoopi Goldberg and Melissa Barrera and Emily Blunt speaking on the record, is — I am almost certain — the equivalent crystallising-moment for the AI era. The artefacts that will come out of it (the documentaries, the dramatic-feature treatments, the songs, the union-history books) will, twenty years from now, be the cultural objects that *outlast* the technology they were originally anxious about.

Where AI sits in this picture

I want to be careful, in applying the diagnostic, not to wave a hand and claim the pattern *predicts* the AI outcome. It doesn’t. What it does is rule out certain shapes the outcome cannot take, and rule in certain shapes it almost certainly will.

The shapes ruled out: an *outright ban* on AI training that protects the existing definition of authorship. A *prohibition* by union action that simply removes AI tooling from professional production. A *moral consensus* — Cardinal O’Connell at scale — that AI work is degenerate and should be socially refused. None of these has worked in twenty previous iterations of the same pattern. None of them is going to work this time. The MU’s 1982 synth motion is in the books as a cautionary tale.

The shapes ruled in: a per-output levy structure flowing into a creators’ fund — Petrillo for neural weights. A C2PA-style provenance standard underwriting an authenticity premium — the structural ancestor of which is the 1990s photojournalism ethics fight. A redefinition of the underlying creative category — *authorship in 2030* will mean something materially different from *authorship in 2020*, the way *singing* meant something different after the microphone. An institutional settlement that absorbs the new tool and *redistributes* the productivity gain, rather than one that *bans* the new tool and *forfeits* the productivity gain.

The cohorts who will be locally displaced are already visible in [Chapter 12](#) and [Chapter 15](#). Junior animators. Concept artists in cohorts being asked to use generative tooling for the front of their pipeline. Voice actors below the SAG-AFTRA scale level. Stock-image photographers. Translators of bulk commercial copy (the closest pre-2022 analogue — the Google Neural Machine Translation moment in 2016 — already produced documented translator-employment effects across 696 US labour markets, and a 70% income loss for the Irish-language EU-institutions translator profiled in the January 2026 *CNN Business* piece I cited earlier in the

book⁶⁰). The fear of these cohorts is well-founded, and the institutional response should be calibrated to it.

The compensating gains — the new categories of creative work, the new business shapes, the new audience contracts — are what Chapter 11 of this book is about. I will not re-state them here. What I want to note, *for the historical pattern's sake*, is that they always emerge, they always emerge faster than the resisters expect, and they always — at the aggregate level — produce more total creative employment than the displaced form supported. The phonograph was the largest creative-labour expansion in music history. The smartphone was the largest creative-labour expansion in image-making history. *I think AI, on the historical pattern, will be the largest creative-labour expansion in cultural-production history.* I think the working creatives who emerge from this with the most leverage will be the ones who, like Walter Murch picking up Final Cut Pro at the height of his Avid mastery, learned the new tool *before* the cultural permission to use it had fully crystallised. The cultural permission usually arrives about three years after the productive use does. The window for asymmetric leverage is now.

What this asks of working creatives

I want to close this chapter with a working-practitioner read on the history, because the abstract pattern is useless without translation.

One. Recognise where on the curve you are. AI in 2026 is in late act three, early act four. Existential alarm is the dominant mood. Institutional counter-attack is well organised. *Settlement* has not been reached. The decisions you make about how to engage with the tools in the next eighteen months are decisions that will define the structural shape of the rest of the decade. This is not a *theoretical* claim about the historical pattern; it is the *operational* claim about your career.

Two. Don't fight to keep the existing definition; fight to be among the people redefining it. The miniaturists are remembered as the cohort that was wiped out. The Stieglitzes are remembered as the cohort that *redefined what photography could be*. Both groups felt, in 1855, that they were fighting for the same thing. They were not. One was defending the inherited definition. The other was rewriting it. The working creatives who emerge from the AI period with the most leverage will not be the ones who defended hardest. They will be the ones who *redefined fastest*.

Three. Pick the institutional response that has historically worked. The Petrillo template — levy on the displacing technology, pool collected centrally, redistribution to the displaced craft — has a hundred-year track record. Use it. Apply it to your union negotiations. Apply it to your platform-procurement decisions. Apply it to your political advocacy. The SAG-AFTRA Tilly Tax, the UK 88%, the C2PA provenance stack, the Cannes Disclosure Standard, the Sundance literacy

initiative are all, in their different ways, *the Petrillo template applied to AI*. They are the part of the institutional response that historically wins. Show up to them. Argue for them. Don't waste energy on prohibitions.

Four. Read the named fear carefully, and listen for the redefinition underneath it. When you hear yourself saying *AI will steal my craft*, ask the second question: *what new definition of my craft is forming on the other side of the displacement, and am I in a position to inhabit it?* The microphone vocalists who absorbed Crosby outlasted the operatic vocalists who refused him. The editors who absorbed Avid outlasted the editors who refused it. The photographers who absorbed digital outlasted the photographers who refused it. The pattern is, by this point in the historical record, very reliable.

Five. Make the cultural symbol. The named-fear language of the AI era — its *Video Killed the Radio Star*, its *I think I'm extinct*, its *imbecile slush* — is being written, right now, by working creatives in their own work. Some of those artefacts will outlast the technology they were anxious about. Some will become the source material for the next generation's understanding of this moment. *If you are a working creative reading this in 2026 and the AI displacement frightens you, the most useful thing you can do is put your fear into a piece of work whose argument outlasts the platform release cycle.* That is what Trevor Horn did. It is what Phil Tippett did. It is what the displaced miniature-painters who became fine-art photographers did. It is, on the historical pattern, what is asked of you.

A coda on Trevor Horn

I want to close with one image, because it is the single most useful one I know.

Trevor Horn, in 1979, was the keyboard player and frontman of the Buggles. He wrote, with Geoff Downes and Bruce Woolley, a song called *Video Killed the Radio Star*. The song is a small, half-melancholy commercial joke about the end of an era — a synthetic-sounding pop song *about* the moment recorded music turned visual. MTV launched on 1 August 1981 with that song as its very first broadcast.

Horn went on to become, by some distance, the most influential pop producer of the 1980s. He built Frankie Goes to Hollywood, ABC, Yes's *90125*, Grace Jones's *Slave to the Rhythm*, Seal's debut, the Art of Noise. He did it on Fairlight CMIs, sequencers, drum machines and the kind of dense, layered, programmed production that the Musicians' Union was, at the same moment, voting to ban. He did the thing the union was warning him against, and built half the canonical pop music of his decade out of it.

The song that named the death of the radio star outlived MTV. The producer who wrote that song defined the next era of recorded music by absorbing the technology

the resisters were trying to ban. The cultural symbol outlasted both the form it threatened and the form that did the threatening.

That is the working operating model I would commend to anyone reading this book and feeling, in 2026, the gravitational pull of the resistance. The resistance is real. The fear is real. The named-fear language is the source material of the era's best art. *And the people who absorb the tool, who learn it, who push it past where its makers intended it, and who use it to make the work that argues for the world they actually want — they are the ones, on a hundred years of evidence, who define the next decade of the form.*

Welcome to the *next* moment in a recurring pattern. The pattern is, by now, very well documented. The choice is yours.

Chapter 3 — The Human-AI Agency Continuum

A week after I sent the first edition of *Dream Machine*, on the Monday of the second week of October 2025, OpenAI held its annual DevDay conference and quietly changed what the conversation about AI in creative work was about.

The first edition had been about Sora 2. The second edition was about something I was less prepared for: the launch of a thing called **AgentKit**.⁶¹

AgentKit was, at first glance, a set of developer tools. Agent Builder. A connector registry. An eval framework. ChatKit, for embedding agents into other products. The launch post on OpenAI's blog framed it, in the slightly forced register that all platform-launch posts use, as a way for developers to "build, deploy, and optimize agentic workflows."⁶² On its own, this was an unremarkable announcement.

What was remarkable, looking back, was the *category claim* the announcement carried with it. Sam Altman, in his DevDay keynote that day, declared the start of "the age of agentic AI" — by which he meant the moment that AI systems stopped being prompt-and-respond chat boxes and started being things that could plan, decide and execute "for hours on end" without further human input.⁶³

For someone like me, sitting in a small studio in the North West of England — running tools all day, looking at my pipeline, thinking about my team's labour — that phrase did a particular kind of work. It rearranged the question.

The question, until that week, had been: *what does AI do for creative work?* The question after that week became: *where, in any given piece of creative work, does my agency end and the model's begin?*

The second question is the one I want this chapter to be about. I called it, in the second issue of the newsletter, the **Human–AI Agency Continuum**.⁶⁴ The frame has stuck with me. I think it is the most useful thing I have ever written down about all of this, and I think — at the risk of overselling it — the rest of the book leans on it.

The continuum

Imagine a horizontal line.

On the far left of the line is **pure human agency**: the writer at the desk, the painter at the canvas, the songwriter at the piano. No machine intermediation other than the tool itself — and the tool, in this position, is dumb. It records what you do; it doesn't decide.

On the far right of the line is **pure machine agency**: an autonomous system that, given a goal, produces a finished creative output with no human in the loop. A prompt, a setting, a render. No one looks at the intermediate steps. No one steers.

The conversation about AI in the creative industries in 2024 mostly took place on the assumption that "generative AI" sat about three-quarters of the way along that

line — closer to the machine end. You typed a prompt; the machine made the thing; you accepted or rejected. There were variants, of course. But the geometry was prompt-and-respond, and the question was simply where on the line, between you and the model, the actual creative work happened.

What changed at OpenAI DevDay on 6 October 2025 — and what was reinforced almost every week of the six months that followed — was that the line is not, as it turned out, a single line. It is a *family of lines*, one per creative function, and they all move at different speeds.

A film, broken down, is not one act of agency. It is a thousand. The choice of subject. The treatment. The casting. The script revisions. The cinematography. The blocking on set. The performance, take by take. The editorial assembly. The grade. The sound. The music. The marketing. Each of those is a sub-discipline, with its own craft, its own labour pool, its own union, its own pay scale and its own internal hierarchies.

AI doesn't slide along *the* line. It slides along each of those lines independently.

A working filmmaker in late 2025 might sit at the absolute left of the continuum on *performance* (a real actor, in the room, in real time, the work itself) and at the absolute right on *background plate generation* (a Veo 3.1 shot, signed off in a Slack message, no human ever drawing a frame).⁶⁵ A working musician might sit at the absolute left on *songwriting* (a song in a notebook) and on the right edge of the centre on *vocal alignment and pitch correction* (an iZotope Ozone 12 assistant, accepted with one click).⁶⁶

The crisis of authorship is not that machines do creative work. Machines have done parts of creative work for as long as there have been cameras, samplers, Photoshop filters and Logic plug-ins. The crisis is that we don't have an honest, shared, public vocabulary for *which* parts. The Continuum, written down honestly per project, is the start of one.

Agents are not generators

The reason the Continuum became urgent the week of DevDay, and not before, is that “agent” is a different kind of object on the line than “generator” is.

A generator is a tool. You aim it at a problem; it makes an output. The agency is in the aiming.

An agent is something more like a junior collaborator. You give it a goal — *find me ten reference images for this shot, generate a rough sound design for this scene, book the courier for tomorrow's pickup* — and it goes away, makes a series of sub-decisions, and comes back with a result. The agency is distributed. You set the direction; it makes the moves.

The reason this matters in creative work is that the moves are where the craft lives. Anybody can describe a final film in a sentence. The film is in the thousand decisions between the sentence and the screen. A generator that makes the screen-ready file from your sentence isn't doing your craft. It is taking your craft out of the loop.

An agent, properly deployed, can do something different and — to me, anyway — more interesting. It can take the parts of the loop that are not where your craft lives, and quietly handle them, so that the parts of the loop that *are* where your craft lives become the parts you actually spend your time on.

That is the optimistic case for agentic AI in creative work, and it is the case that almost every working creative I respect makes when you sit down with them in private. It is also the case Adobe's 16,000-creator survey, released a few weeks after DevDay, came in to support: 70% of respondents were optimistic about agentic AI, framed as "tools that act on your behalf"; 85% said they would use AI that learned their creative style.⁶⁷

The pessimistic case is the one Adobe's same survey also captured: 69% of respondents worried about their work being used to train AI without consent.⁶⁸

Both numbers are about agency. The first is about *gaining* it back, by handing routine work to a competent assistant. The second is about *losing* it, by having the work that defines you absorbed into a system you do not control. Both are true at the same time, for the same creators, in the same workflows.

Where agents went, between October and May

In the six months between DevDay and the time I'm writing this, the agent layer of the creative toolchain went from "interesting demo" to "shipping product," faster than any technology shift I have lived through in twenty years of practice. I want to give you a sketch of the trajectory, because it is what most of the rest of this book is reacting to.

By **mid-October 2025**, Mureka — a Chinese music platform — launched a thing called *Music Agent Studio*, six specialised AI agents for songwriting, arrangement and production.⁶⁹ A startup called AdsGency raised \$12m in seed to build agents that could autonomously run a brand's entire paid marketing workflow.⁷⁰ A company called Lenny launched an agent for organising live music events.⁷¹ Each of these felt, at the time, like a specialist tool. In retrospect, they were the first signs that whole production functions — not individual tasks — were being handed over.

By **the end of November**, EA, in the middle of a brutal financial year, told its 15,000 employees to use AI as a "thought partner" for everything from character art to playtesting.⁷² The framing — *thought partner* — was the precise rhetorical move

that turned an agent from a tool into a colleague. The colleague has opinions. The colleague has time. The colleague has a seat at the meeting.

By **December**, Adobe announced that you could now use Photoshop and Express *inside* ChatGPT — meaning that the creative output itself was no longer happening inside Adobe’s interface, but inside an agent’s.⁷³ This was a small thing on the surface and an enormous thing underneath. It was the moment that Adobe — a company that has, since 1990, owned the metaphor of the *creative tool* — accepted that the new metaphor was the *creative agent*, and that they would rather be inside someone else’s agent than not in the conversation at all.

By **late January 2026**, Anthropic shipped Claude apps — interactive, custom assistants embedded directly in workplace tools — and a company called Heygen released *Video Agent*, which could script, edit and assemble entire videos from reference images.⁷⁴ By **March**, Adobe announced its **CX Enterprise** platform alongside NVIDIA: a stack of AI agents embedded across the entire content lifecycle, from brief to delivery.⁷⁵ By **April**, the *Adobe Summit* keynote made it official — “agentic creative intelligence” was now the headline category, not a feature.⁷⁶ By **May**, Sony was using a multi-agent team of forty-nine Claude Code agents, working with seventy-two skills, to coordinate game-development work.⁷⁷

The trajectory, in one sentence: in October 2025 we were arguing about whether agents were a thing. By May 2026, the entire creative production pipeline at a global game publisher was being run by a team of them.

What this means for craft

The natural fear, reading that timeline, is that the agency line drifts inexorably to the right — towards the machine end — and that the craft of the human in the loop becomes thinner and thinner until it disappears.

I do not think that is what happens. I think what happens is more interesting and more demanding.

What I see, in my own studio, in my friends’ studios, in the working musicians and filmmakers and games designers I talk to every week, is that agentic AI doesn’t compress craft into nothing. It *relocates* craft to a different place on the continuum.

If your job, last year, was “make the thing”, your job this year is “decide what gets made, brief the agents that make the constituent parts, and judge the output.” That isn’t a smaller job. In some ways it is a bigger one. It requires *more* taste, not less, because taste is now the only signal you bring that the agents cannot.

Anthropic, in a blog post in early 2026 that I have ended up quoting repeatedly in talks, made the point this way: agentic systems work best when they are deployed by people who already have the taste and judgment to know what good output looks

like.⁷⁸ The agents accelerate the work of *people who are already good at it*. They do not — at least, not yet — manufacture good work from nothing.

This is the central — and I think non-obvious — claim of the Continuum frame: as the line for any given function slides to the right, the *value of the human at the left edge of the line* doesn't decrease. It increases. Because the question being asked of that human gets sharper. Not “can you make this,” but “*should* this be made, and *why this version*, and *who is it for*, and *what does it need to do in the world*.”

That is craft. It is just craft sitting in a different chair.

Where the Continuum breaks

I want to be honest about where my frame stops working, because nothing is more boring than a writer who only quotes the people who agree with him.

In November 2025, the games designer Charles Cecil — the head of Revolution Software, the studio that made *Broken Sword* — told *gamesindustry.biz*, in a sentence that has been quoted, retweeted and emailed around my industry approximately a million times: “AI was an expensive mistake.”⁷⁹

Cecil's argument was specific. Revolution Software had, like a lot of indie game studios, experimented with using generative AI in early production. They had found that the time saved on the front end of the pipeline was lost — and then some — on the back end, where artists, writers and designers had to reverse-engineer, fix, replace and reintegrate AI-generated assets that didn't quite fit the game's tone, didn't quite match the existing art direction, didn't quite work with the engine, didn't quite carry the IP. Net-net: more time spent, not less. More cost, not less. Hence: “an expensive mistake.”

This is what the Continuum frame doesn't capture on its own. *Where on the line* a given task sits is not a fixed property of the task. It is a function of the surrounding system: how the tools integrate, how the team is structured, how the IP works, how the audience receives the output. A generative tool that sits comfortably on the right-hand side for one studio's marketing department sits awkwardly in the middle for another studio's lead-artist pipeline.

In the same six months that I was watching the agent layer eat the creative toolchain, I was also watching studios push back. Larian, the makers of *Baldur's Gate 3*, backed off from generative AI for their next *Divinity* game in January 2026. Their public note was carefully worded: “*I know there's been a lot of discussion about us using AI tools as part of concept art exploration. We already said this doesn't mean the actual concept art is generated by AI but we understand it created confusion.*”⁸⁰ Games Workshop ruled it out entirely for *Warhammer 40,000*.⁸¹ Manor Lords publisher Hooded Horse said it wouldn't work with

developers using generative AI — its founder’s framing, when asked about the line, was unusually direct: AI in his pipeline was “*cancerous*,” and the studio’s job was “*constantly having to watch and deal with it and try to prevent it from slipping in.*”⁸² Jagex, the maker of *RuneScape*, said in early 2026 that it would *never* use generative AI to make in-game content, and that the commitment “*goes so far that we are now doing an audit and having a conversation with our various external partners that work with us to ensure that no AI is being used in inappropriate ways in any of their work that might filter through.*”⁸³

These were not statements made by Luddites. They were strategic decisions made by people whose creative product is, in significant part, the *human* fingerprint on the work. The audience for a *Warhammer* miniature, or a *RuneScape* quest line, or a Larian dialogue tree, comes to those products in part because they know — and want to know — that real people made them. The Continuum slides differently in those companies because the *output* sits at a different point on the continuum of what the audience wants.

This is the thing about the agency line that the OpenAI keynote, the Adobe Summit, the NVIDIA GTC keynote, the Anthropic blog post and the Salesforce Dreamforce all keep glossing over. The position of the line is not just about what is technically possible. It is about what the work, in its finished form, is *for*.

Open the black box

I want to put one more argument on the page in this chapter, because it is the argument I have come to believe more strongly than any other after six months of writing the newsletter, and it does not fit cleanly inside the Continuum frame even though it is what the frame is, in the end, *for*.

The argument is this. Working creatives, as a class, need to **open the black box of AI and own a real stake in how it is built**. Not just *use* it. Not just *refuse* it. Not just *bargain over its terms*. All of those matter, and the SAG-AFTRA Tilly Tax, the UK 88%, the *Stealing Our Work Is Not Innovation* declaration are all evidence that the bargaining work is happening. They are necessary. They are not sufficient.

The sufficiency move is the *technical-literacy* move. The thing that makes the Continuum frame survive contact with the agentic stack — and that makes the *age of the Why* I will argue for in Chapter 16 commercially defensible rather than wishful — is that working creatives are sitting *inside* the toolchain, with their hands on the dials, understanding how the model was trained, on what, with what licensing, with what guardrails, with what consent mechanisms, with what energy and water footprint, with what data-supply-chain labour costs. Not as a hobby. As a structural condition of their professional autonomy.

The history of every previous creative-technology transition supports the move. The musicians of the 1970s and 1980s who *learned the synth from the inside* — programmed it, modified it, hacked the patches, understood the signal chain — built more durable careers than the ones who let the manufacturers decide what the instrument was for. The editors who *learned non-linear editing from the inside* — set up their own systems, understood the codecs, understood the colour pipelines, understood the storage architecture — were the ones who, by the early 2000s, had real leverage over how digital cinema was structured. The photographers who *learned digital from the inside*, in the 1990s and 2000s, made the working-photographer transition that the photographers who waited for the consumer firms to tell them what digital meant largely did not.

The pattern is, by historical evidence, very reliable. *The cohort of working creatives that opens the black box of the new tool, and that participates in the design and the discourse of how the tool is governed, defines the next era's craft. The cohort that uses the tool without ever asking what is inside it has the era's craft defined for them by the platform companies that ship the tool.* The first cohort writes the textbooks. The second cohort is described in them.

The 2025–26 evidence so far is mixed. The open-source ecosystem documented in [Chapter 17](#) — ComfyUI (\$500M valuation by May 2026), Hugging Face, the Hunyuan and Qwen and DeepSeek open-weight families, the Civitai LoRA marketplace, the Korin AI Africa-trained model, the *80% of YC and Andreessen Horowitz startups now building on open-weight models* statistic — describes one half of the picture. There is, in 2026, a genuine open-source creative-AI infrastructure underneath the closed platform layer, and a fast-growing cohort of working creatives who use it deliberately. That cohort is doing the *opening-the-black-box* move at scale.

The other half of the picture is the part of the working-creative population that uses the closed platforms — ChatGPT, Sora, Midjourney, Adobe Firefly via the Creative Cloud — without understanding what the models were trained on, what the terms of service say about output ownership, what the consent regime around the training data is, what the energy footprint of a single generation is. That cohort is, structurally, in the position of the parlour musician in 1906 who took the phonograph at face value because the salesman said it would play their favourite songs. The phonograph absolutely did play their favourite songs. It also restructured the entire economics of the music industry around them, in a direction the parlour musician had no say in, because the parlour musician had not opened the box.

I want to be very direct about what this asks of working creatives in 2026. It asks four specific moves.

One. Learn how the models are trained. Not in technical detail. In structural detail. Understand the difference between a model trained with consent and a model trained without. Understand the licensing regime of the tool you are about to use. Understand, before you sign the EULA, whether your *outputs* are owned by you or by the platform. Treat the EULAs of AI platforms as part of your working practice. If this feels like reading the small print on a building-trade contract, that is the right comparison.

Two. Run at least some part of your stack on open-weight infrastructure. The strategic argument for this is in [Chapter 17](#). The political argument is in [Chapter 6](#). The personal argument is the one I am making here: the working creative who knows how to run a Hunyuan or Qwen variant on their own machine, on their own terms, with their own data, has a different relationship to the closed platforms than the working creative who depends on them. The independence is real. It is also, in commercial negotiations with platforms, *worth money*. The closed-platform vendors price their tooling differently for customers who can credibly walk to open-source alternatives.

Three. Show up to the governance conversation. The Sundance literacy initiative ([Chapter 12](#)), the UK government consultation that produced the 88% ([Chapter 6](#)), the SAG-AFTRA bargaining ([Chapter 13](#)), the Cannes Disclosure Standard ([Chapter 13](#)), the European Article 17 implementation, the C2PA standards body, the Music Performance Trust Fund’s emerging AI-era equivalents — these are the venues where the rules for the next decade are being written. They are usually held in rooms with bad coffee, in meetings with too many lawyers, with insufficient working-creative representation. *Be the working-creative representation in those rooms*. The platform companies have full-time staff on every standards body and every consultation. The cohort that turns up to argue with them is the cohort that gets included in the rules.

Four. Refuse the framing where AI is something done to you, and adopt the framing where it is something you do. This is the rhetorical move, but it is also a practical posture. The 2024 industry conversation about AI in creative work — and a large fraction of the 2025 trade press — treated working creatives as the *object* of the AI transition: the population to which AI was being applied. The 2026 working creatives who are doing best, in my experience, have reversed that framing. They have made themselves the *subject* — the people *applying* AI to their work, on their terms, in service of their intent, using the open-source infrastructure where it serves them, using the closed-platform infrastructure where it serves them, refusing both where neither does. The grammar is the difference between “*I’m being affected by AI*” and “*I’m using AI*.” The grammatical difference is also, on inspection, the *power* difference.

A creative economy in which working creatives have opened the box, understand the box, contribute to the design of the box, and own the political and technical infrastructure that decides what the box is for, is the creative economy I am arguing for in this book. The Continuum is the working frame for the daily practice. The Four Principles of Chapter 16 — *agency, attribution, access, audience* — are the structural-policy version. The black-box-opening move is the practitioner’s version. They are all the same argument seen from different angles.

The version of this transition where the working creatives stay outside the box is the version where the box decides what creative work is. The version where the working creatives are *inside the box* is the version where the box is built around what creative work needs to be. Those are not the same outcomes. The next eighteen months will, on the available evidence, decide which one we get.

A working frame

If I were going to leave you with one tool from this chapter, it would be this:

The next time you sit down to plan a piece of creative work, draw the lines.

Not one line — that’s the trap of the “AI debate” — but as many lines as the work has functions. *Ideation. Research. Writing. Direction. Performance. Image-making. Sound. Editing. Distribution.* For each one, ask the same two questions. *Where do I want to sit on this continuum, and where am I willing to let the agent sit on my behalf?* And then — the harder question — *what does the work lose if I move further to the right, and what does it gain?*

The honest answer, for almost every creative person I know, varies wildly by function. Most of us are happy to let agents sit on the right-hand side of distribution and admin. Most of us are not happy to let them sit on the right-hand side of the performance, the writing, the moments where the audience can feel a person in the work. The middle is where the interesting fights are.

If you can articulate where the lines sit for *your* work, you can articulate it to your clients, your team, your collaborators, your union, your audience. You can write it into your contract. You can put it on your website. You can fight for it.

If you can’t articulate it — if you wave at “AI” as if it were a single thing — you will end up with the lines drawn for you, by tool vendors and platform companies and CFO spreadsheets that have very different ideas about where your agency should sit than you do.

The Human–AI Agency Continuum, in the end, is not a description. It is a defence.

Chapter 4 — Dead Internet, Living Web

On the morning of Wednesday 22 October 2025, I read three reports back to back at my desk, and by the time I was halfway through the third one I had stopped taking notes and just started staring at the screen.

The first was from Imperva, a security company that publishes an annual *Bad Bot Report*. The 2025 edition opened with a sentence I have quoted in talks at least a dozen times since: for the first time in a decade, automated traffic had overtaken human activity on the public web. Bots — not people — were now responsible for **51%** of all web traffic. Within that 51%, the category Imperva calls “bad bots” — scrapers, credential-stuffers, content thieves and fraud accounts — accounted for **37%** of the *whole* internet, on their own.⁸⁴

The second was from Cloudflare, whose engineers can see a significant share of global web traffic from inside their infrastructure. Cloudflare’s own analysis, in a blog post titled *The crawl-to-click gap*, confirmed Imperva’s picture and added a detail. Of the bot traffic Cloudflare could classify, roughly **80%** was attributable to *AI training crawlers* — GPTBot, ClaudeBot, Meta’s scrapers, the new wave of agentic bots that performed autonomous tasks (1.7% of bot traffic at the time, but growing fast).⁸⁵

The third was a market projection from Grand View Research and a separate one from Gartner referenced in Europol’s 2025 briefing. Both said, in slightly different language, the same thing: by 2030, between 90% and 99% of online content will be AI-generated or AI-assisted.⁸⁶

If you put the three reports together — and this is the thing I did on the morning of the 22nd, before I had decided what to write that week — what you got was a picture of an internet whose dominant activity was no longer humans publishing and reading. The dominant activity was *machines reading machines*. The web was being trained on a version of itself written by the systems it was training.

Five days later, the fourth issue of the *Dream Machine* newsletter went out with a headline I had been circling for weeks. It said: *Is the Internet Dead Yet?*⁸⁷

I want to spend this chapter on the answer.

The synthetic mirror

The “Dead Internet Theory,” for those who haven’t met it, is a notion that has been knocking around the internet since at least 2021. In its original, slightly conspiratorial form, it claims that most of the web has been replaced by bots — that the people you talk to on social media are agents, that the comments on news articles are agents, that the cultural water you swim in is a synthetic medium pretending to be a human one.⁸⁸

In 2021, when it was first articulated, it was an interesting bit of folklore that didn't quite map onto reality. The bots existed; they just weren't, yet, doing most of the work. The cultural water was still mostly human.

By October 2025, the maths had quietly inverted. Half of the traffic was machines. A majority of *new published content* was machine-assisted, according to a separate 2025 analysis by Graphite that put the human-to-AI authoring split at roughly 50–50.⁸⁹ The pages those machines were writing were being scraped by other machines to train *next year's* generation of writing machines.

A recursive system trained on its own outputs is called, in academic AI circles, *model collapse*. The fear, in the published literature on this, is straightforward: a system that learns from synthetic data loses touch with the real-world signal that made it useful in the first place, and starts producing increasingly homogenised, brittle, hallucination-prone outputs.⁹⁰

What the 2025 numbers said, when you sat with them, was that we were no longer talking about model collapse as a theoretical risk. We were talking about *web* collapse — a slow, quiet, structural drift in which the public commons of writing, image-making, video and music started to be made by, and for, the machines that read it. Humans were still there. We were no longer, by any meaningful metric, the *primary audience*.

What the Dutch researchers found

In the second week of October, a team of researchers in the Netherlands ran an experiment that I think will end up being cited a lot more in the years to come than it was at the time.⁹¹

They built a small, stripped-down social platform — no algorithms, no ranking, no advertising — and populated it with several hundred large-language-model-based AI agents. The agents had different “personalities,” different starting interests, different opinions. The researchers' question was simple: in the absence of any algorithmic distortion, would the bots — when free to interact only with each other, with no human in the loop — converge on a healthy public conversation, or would they reproduce the pathologies we already see in human social media?

The answer, within hours, was the second. The agents fractured into warring tribes. A narrow elite captured the bulk of the attention. Extremist echo chambers flourished. The platform, with no humans on it at all, produced almost exactly the same dynamics that the human-plus-algorithm version of social media has produced for the last decade.

The conclusion the researchers reached — and the one I want to flag now, because it is going to recur in this book — was that the *architecture itself* is the problem. The toxicity wasn't, or wasn't only, in the humans. It was in the design of the system:

how identity worked, how attention was allocated, how voices were amplified or suppressed. The bots reproduced it because they had been trained on the human web, and the human web has the same architecture.

This is the single most important thing I learned in the first two months of writing the newsletter. The optimistic AI take and the pessimistic AI take both assume the architecture stays the same. The optimist thinks the agents will use it better; the pessimist thinks they will use it worse. The Dutch experiment suggests that neither matters — the architecture *itself*, regardless of who or what is filling it, will produce the same pathologies.

If we want a different outcome from the AI era, we need different rails, not just different drivers.

What survives

In the original Issue 4, I wrote that “authenticity and provenance become the new scarcity.” I want to defend that line, six months on, because I think it is the part of the chapter that has held up best.

The simplest way to put it is this: when everything online can be faked, cloned or generated at near-zero cost, the most valuable signal is *proof that a person made something*. Not just an aesthetic preference. An economic one.

You can see this argument being made, all over the creative industries, by people who have nothing else in common. Adam Mosseri, the head of Instagram, said in early January 2026 that the platform should focus on “fingerprinting real media” rather than tracking and disclosing AI slop — that is, the policy should be to identify and amplify provably human-authored content rather than to play whack-a-mole with the synthetic stuff. His framing was telling: “*Everything that made creators matter — the ability to be real, to connect — is now accessible to anyone with the right tools.*”⁹² The platform head was acknowledging, on the record, that the previous decade’s content-creation moat had been completely flooded. The only remaining moat was *being a person you could verify was a person*.

Even that moat is now under direct commercial assault. In June 2026, the *Guardian* reported on brands deploying **AI-generated influencers** to promote products on social media — synthetic people engineered to occupy exactly the territory Mosseri had identified as the last defensible one: the ability to be real, to connect.⁹³ The platforms’ first regulatory response is telling about where the economic nerve actually sits. The same month, **TikTok banned AI-generated voices in shopping livestreams**⁹⁴ — not on aesthetic principle, but because the point of sale is the one place where synthetic sincerity carries an immediate, measurable commercial cost. Where trust is transactional, the platforms are already legislating the synthetic out. Where it is merely cultural, they are still letting it flood in. A

month later, the hygiene perimeter widened along exactly the lines that logic predicts: in July 2026, TikTok began targeting **AI-generated spam accounts in “high-risk” topics** — finance, health, news, the categories where synthetic authority carries a measurable cost to the reader.⁹⁵ Note the shape of the policy. Not a general ban on synthetic content; a triage, category by category, of where the synthetic flood is allowed to pool and where it is not. The platforms are drawing the same map I am drawing in this chapter — they are just drawing it in enforcement actions rather than arguments. The enforcement is also, it should be said, arriving late and unevenly: a July 2026 study found that **YouTube and X between them were funnelling millions of visits to AI deepfake “nudify” apps** — the same platforms policing synthetic authority in finance and health were, through their own recommendation and ad surfaces, routing users straight to the most harmful synthetic content on the web.⁹⁶ The hygiene perimeter and the funnel to the sewer run through the same buildings.

Sundance Institute, launching its AI Literacy Initiative the same month, framed authentication and authorship as the central question filmmakers needed to negotiate to remain in control of their own work.⁹⁷ Bandcamp, the indie music platform that has always carried more cultural weight than its commercial size implied, simply banned AI-generated music outright in early 2026.⁹⁸ San Diego Comic-Con drew the same line for its 2026 art show, with rule language as flat as anything in the cultural sector: *“Material created by Artificial Intelligence (AI) either partially or wholly, is not allowed in the art show. If there are questions, the Art Show Coordinator will be the sole judge of acceptability.”*⁹⁹

These are not, on their own, market signals — they are policy decisions. But they were being made, in early 2026, against a backdrop of audience behaviour that suggested something larger. Deezer reported in April 2026 that AI-generated music had risen to **44% of all daily uploads** — 75,000 tracks a day, more than 2 million a month — but that those tracks accounted for **between 1% and 3% of total streams**.¹⁰⁰ The audience, given the choice, was choosing not to listen.

That ratio — call it 44 to 3, or 75,000 to listen-to-nothing, or whatever shorthand you prefer — is the most important number in this book, and I will come back to it in Chapter 5. The reason I introduce it here is that it is the empirical answer to the Dead Internet question. The web is not dead. The web is producing exponentially more *stuff* than it ever has, and the humans on it have started to develop antibodies. They are not engaging with the synthetic flood. They are, by their attention patterns, picking out the human signal.

There is a piece of accounting underneath all of this that the platform companies, in my view, have not yet metabolised — and that I think the rest of the book runs on top of. **Human attention is a finite resource.** Nielsen-class telemetry on

aggregate daily media-consumption time, in every market I have seen the numbers for, has been roughly stable for at least a decade. The eyes, the ears and the consciousness of the average adult are each, by physiology, in the same condition as they were in 2015. The supply of producible content has — through the autumn of 2025 and the spring of 2026 — grown by orders of magnitude. The audience's capacity to consume has not grown at all. That is the binding constraint of the Dead Internet picture. The synthetic content is real; the humans are still here; and *the humans cannot, in net, consume more hours per day than they already do*. The 51% of bot traffic Imperva measured, on this reading, is not a measure of how much the audience has expanded to absorb new content. It is a measure of how much *unread* content is being produced, by machines, for other machines, with no human eye-time anywhere near it. Chapter 11 develops the finite-attention argument at length. For now, it is enough to note that the *flood* and the *ceiling* — the two sides of the slop-ceiling dynamic — are both visible in the Dead Internet picture, and that they are produced by the same underlying audience biology.

Synthetic sincerity

In November 2025, the filmmaker Marc Isaacs premiered a documentary at IDFA — the International Documentary Film Festival in Amsterdam — with a title I have not been able to get out of my head. The film was called *Synthetic Sincerity*. It was a hybrid piece, blending real footage with AI-generated characters, deliberately blurring the line between what was real and what wasn't, and asking — as its working premise — whether AI characters could be *taught* authenticity.¹⁰¹

The film and its accompanying *Hollywood Reporter* interview ran the same week as a separate *Variety* piece titled *AI-Generated Images Threaten Future of Documentary as People 'Will Stop Believing Anything'*.¹⁰² The juxtaposition was almost too on the nose. One filmmaker trying to *expand* the territory of the synthetic, on the assumption that authenticity is a property that can be invested in fictional characters; another set of filmmakers arguing that the very ability to fake reality is hollowing out the cultural credibility of their entire form.

I am not going to take a side on the documentary question, because I don't think there is one yet. What I want to flag is that *Synthetic Sincerity* — the phrase, not the film — is a useful piece of vocabulary. It names a category. There is a kind of work, in this new ecology, that is *trying* to be authentic and openly synthetic at the same time. It is not pretending to be human. It is asking whether the qualities we used to attach to humans — emotional truth, lived experience, perspective — can be ported over to synthetic characters who are honest about what they are.

The verdict, six months in, is mixed. Some of the strongest creative work I have seen this year sits firmly in this space. Hoyt Dwyer's animated short — made by a former Apple TV creative, competing at the AI FilmFest Japan in late 2025 — does not

pretend its characters are real, and is more honest about its medium than three quarters of the live-action features I watched the same year.¹⁰³ Andrii Daniels' viral *Deadpool / Harry Potter* Christmas clip, which he made in a Ukrainian bomb shelter during an active war, has more sincerity in any single frame than most legacy-studio output, precisely because the conditions of its making are on the screen.¹⁰⁴

Some of the worst work I have seen this year sits in the same space too. McDonald's Netherlands' AI-driven Christmas ad — released in December 2025 and pulled within days after a public backlash — was an attempt at *synthetic sincerity* that read, almost universally, as cynicism wearing a Christmas jumper. The line that travelled fastest, as the ad's reception turned, came from a working creative director responding on social media: “No actors, no camera team, no light, no sound, just probably one guy, alone in front of a computer battling with an AI prompt who steals the look and everything else from someone else.” That sentence — circulated on LinkedIn and X within an hour of the ad's launch — was the thing that did the cultural damage. The brand had to pull the spot.¹⁰⁵ The Valentino “AI handbag” campaign, criticised by the BBC for being “disturbing,” was the same.¹⁰⁶ Coca-Cola's AI holiday ad — the second time the company had tried this — divided viewers along almost the same lines as the previous year.¹⁰⁷

The interesting pattern, when you line these up, is not whether AI is “good” or “bad” for the work. The interesting pattern is that audiences are very fast, and very precise, at distinguishing *sincere* synthetic work from *cynical* synthetic work. The technology is the same. The fingerprint of the human intent behind it is not. And the audience can feel the difference at the speed of a swipe.

What the brain study said

In the middle of all this — and I want to acknowledge that it is harder evidence than the cultural commentary — an MIT Media Lab study made the rounds in the autumn of 2025, in which researchers measured the brain activity of subjects writing essays with and without generative AI assistance. The headline finding was that AI users showed measurably reduced brain activity over the course of the writing tasks compared to control subjects writing on their own.¹⁰⁸

The headline framing — *AI makes you stupid* — was unfair to the study, which was small, preliminary, and didn't claim anything as strong as that. But the underlying observation has been replicated in other domains. When the cognitive load of producing the first draft is offloaded to a generator, the cognitive engagement of the human in the loop measurably drops. The work gets produced. The person producing it engages with it less.

This is the quieter consequence of the Continuum chapter — the one that doesn’t show up in any line item on a P&L sheet but that I think we are going to be wrestling with for years. If the right-hand side of the continuum is “machine agency,” and we slide more and more functions of our creative work to that side, we are not just changing the *outputs*. We are changing the *people doing the work*. The thinking that produces the work happens, or doesn’t, in the bodies of the people in the workflow. And brains, like muscles, atrophy with disuse.

This is not a reason to refuse the tools. It is a reason to be careful about *which functions* you offload, and to keep a deliberate, conscious habit of *exercising the cognitive work that defines your craft*. The Continuum doesn’t just describe where the line sits today. It describes where you are willing to *let* your mind sit, every day, for the rest of your career.

The training crisis: what model collapse really means

I want to spend a section on *model collapse* in technical depth, because the trade-press shorthand for the phenomenon — *AI starts training on its own outputs and gets stupider* — is right enough to be useful but wrong enough to be misleading. The underlying mechanics matter more than the headline does, and they matter especially for working creatives trying to read where the next wave of model releases is going to land.

The technical framing dates back to a 2023–24 paper by Ilia Shumailov, Zakhar Shumaylov, Yiren Zhao, Yarin Gal, Nicolas Papernot and Ross Anderson titled “*The Curse of Recursion: Training on Generated Data Makes Models Forget*.”¹⁰⁹ The argument is that when a generative model is trained on a corpus where a meaningful fraction of the training data is itself produced by an earlier generation of the same kind of model, the model’s outputs progressively *narrow* — losing the long tail of unusual, rare, distinctive examples, regressing towards the statistical mean, and eventually losing the very property that made the first-generation model interesting (its ability to surprise the user with a specific, particular, well-tuned response).

The paper showed the effect cleanly in controlled experiments. By the fifth or sixth recursive generation of training, the model’s outputs had become noticeably homogenised. The rare-token rate had collapsed. The distinctive-style rate had collapsed. The model was, in technical terms, still functional. It was, in practical terms, less and less *useful* — a copy of a copy of a copy.

The reason this matters in 2026 is that the public web is, by every measure I have seen, now a corpus that contains a non-trivial fraction of AI-generated material. The Imperva and Cloudflare numbers I quoted at the top of this chapter — 51% of web traffic being bots, 80% of that being AI training crawlers and a fast-growing agentic component — describe the *upstream* side of the recursive-training loop. The Graphite 50-50 human-to-AI authoring split describes the *downstream* side. The

next generation of foundation models, trained in 2026 and 2027 on the public web that those bots have produced, will, on the model-collapse hypothesis, exhibit some degree of the homogenisation Shumailov et al. predicted.

How much, in practice, is an open question. The platform companies have not, in the main, disclosed the degree to which they filter their training data to exclude synthetic content. The open-source weights — Hunyuan, Wan, Qwen, FLUX — are trained on corpora whose AI-content fraction is, by my reading, somewhere between *substantial* and *unknown*. The published evidence on whether recent model releases have shown the mean-regression Shumailov predicted is, in mid-2026, *mixed*. Some benchmarks suggest the effect is being detected and engineered around. Others suggest it is showing up in subtle ways — in the difficulty of producing genuinely surprising creative outputs, in the way recent models converge on a recognisable house style, in the fact that the *cheapest* and most ubiquitous AI generations all feel, to working creatives, somehow alike.

The strategic implication for the creative-AI moment is twofold.

One, model collapse — if and to the extent it is real — strengthens the slop ceiling I describe in Chapter 5. A model that has, by 2027, been trained on a corpus heavily contaminated by 2025–26 AI output is, by construction, going to produce *more average* outputs than the 2025 model that preceded it. The audience's selection against the most-average outputs is, on that reading, going to bite even harder against the next generation of generative tools than it bit against the first. The first slop wave hits a ceiling because the audience underweights it. The second slop wave will hit the ceiling *plus* a degradation curve on the production side.

Two, the value of *clean, provenanced, verifiably human-authored data* — the C2PA-signed photograph, the SynthID-watermarked music recording, the contractually-licensed text corpus — goes up sharply over the next five years. The platform companies that retain access to the cleanest training corpora will have a measurable advantage over the platform companies that scrape the post-2024 web indiscriminately. The Stability AI / Universal Music alliance, the Splice / UMG partnership, the various YouTube and Spotify licensing deals are, on inspection, *acquisition strategies for clean training data* as much as they are creator-economy plays. The race to lock down provenanced data is already on, and it is being fought at the level of large institutional licensing rather than at the level of individual consent — which is, on the historical pattern, one of the reasons the Petrillo template (collective bargaining, joint funds, redistribution) is the right structural response.

The provenance stack: how it actually works

I want to give one more passage of technical detail before the chapter closes, because the *provenance infrastructure* I will refer back to throughout the rest of the book is,

at the moment, one of the most-mentioned and least-understood pieces of the AI conversation.

There are, roughly, four layers that together constitute what the industry has started calling the *provenance stack*.

The first layer is **capture-time signing**. A camera with C2PA support — by 2026, this includes flagship Sony Alpha bodies, Leica’s M11-P and M11-D, Nikon Z9 firmware, and a handful of Canon professional bodies, plus most major smartphone makers’ computational-photography pipelines — generates a cryptographic signature for every image and video at the moment of capture. The signature commits to the device, the timestamp, the GPS coordinates (if enabled), and a fingerprint of the underlying pixel data. The signature is, by design, *hard to fake without access to the capture device’s private key*. Capture-time signing is the foundation of every provenance claim that follows.

The second layer is **edit-time chain-of-custody**. C2PA-compatible editing software — Photoshop with the C2PA extension, Premiere with the Content Credentials toolchain, the various Capture One and Lightroom integrations — preserves the capture signature through each editing step, appending a cryptographically-linked record of *what was done to the file*. The chain is not a single signature; it is a *history* of signed transformations, each one referring back to the previous one. A photograph that has been processed through Lightroom and Photoshop arrives at the publisher with a verifiable record of: the camera that took it, the time it was taken, the edits that were applied, and the human (or automated tool) that applied each edit.

The third layer is **upload-time platform integration**. By 2026, Adobe’s Behance, Vimeo’s pro tier, the AP and Reuters wire services, and a growing list of news publishers have integrated C2PA-aware upload pipelines that preserve the chain through their content-management systems and embed it into the public-facing version of the work. The reader’s browser, with the right extension or platform-level support, can inspect the chain and verify the provenance claim. Adam Mosseri’s January 2026 framing of Instagram’s “*fingerprinting real media*” approach was Instagram joining this third layer at the platform-distribution end.¹¹⁰

The fourth layer is **detection and watermarking for synthetic content**. SynthID, Google DeepMind’s watermarking system, is the most mature commercially-deployed example. SynthID embeds a statistically-detectable but human-imperceptible signal into the output of Veo (video), Lyria (audio) and Imagen (image) generations. The signal survives most common transformations — crops, recompressions, low-quality re-encodings. By December 2025, Google had shipped a consumer-facing version inside the Gemini app: a user could upload a video and ask “*Is this AI-generated?*” and receive a yes/no answer based on the

SynthID signature.¹¹¹ The same kind of watermarking is being deployed, with varying technical robustness, across the other major generative platforms — and the detection layer is spreading across vendors: in the summer of 2026, **ElevenLabs**, the dominant audio-AI platform, announced a partnership with Google DeepMind to make ElevenLabs-generated audio detectable through the same SynthID pipeline.¹¹² The detection layer is also arriving at the *distribution* end, where the audience actually lives: in July 2026, TikTok quietly began **testing AI deepfake detection technology** on its own platform — the first of the big attention platforms to move detection from policy document to deployed infrastructure.¹¹³ (Chapter 13 covers the cross-vendor SynthID rollout in detail.)

Layered together, the four levels produce a structural answer to the Dead Internet question. Capture-time signing tells you *this was taken by a real device*. Edit-time chain-of-custody tells you *here is what was done to it after capture*. Platform integration tells you *the publisher has preserved the chain*. SynthID and equivalent watermarks tell you *this output was generated by an AI system*. No single layer is sufficient on its own. All four, deployed together, produce a *verifiable provenance signal* that the audience can — in principle — use to decide what to spend their finite attention on.

I want to be honest about where the stack is incomplete. Watermarks can be stripped by determined adversaries. Capture signatures can be forged if the device's private key is compromised. Chain-of-custody breaks the moment a file passes through a non-compliant tool. The audience's ability to *inspect* the provenance metadata is, in 2026, dependent on platform UI choices that the platforms have not yet made consistent or universal. The stack is the right answer to the architecture problem. It is not, by any means, finished.

What it does do — and this is the point I want to land before the chapter closes — is *establish the category*. The question *did a person make this?* is, by 2026, *technically answerable with high reliability given the right tooling*. That sentence is the entire shape of the next decade's cultural and policy fight in the creative industries. Who controls the tooling. Who decides what it certifies. What economic value the certification carries. Whether the audience has the legal right to *demand* the certification before paying attention. The C2PA standards body, the SynthID rollout, the Content Authenticity Initiative, the Cannes Disclosure Standard, the various national disclosure regulations in development — these are the venues where the next decade of the *Living Web* gets built. They are, on my read of the historical pattern, the part of the AI debate the working creative most needs to be inside.

The architecture, again

I want to come back to the Dutch researchers' result one more time before I close this chapter, because I think it is the through-line.

The story we are mostly told, by toolmakers and platforms and the optimistic side of the industry press, is that the AI era is *a thing happening to* an otherwise functioning internet. The implication is that if we can get the AI part right — better tools, smarter agents, cleaner training data, better watermarking — then the internet itself will be fine.

I do not think this is true any more. I think what the bot statistics, the Dutch experiment, the model-collapse research, and the audience response to AI music collectively show, is that the architecture itself — the rails on which all this is running — was already broken, and that AI is just the load that has finally exposed how broken it was.

The Dead Internet, in this reading, is not a thing AI is doing to us. It is a thing the web's architecture was already drifting towards — attention-monopolised, identity-collapsed, provenance-blind, optimised for machine-readable metadata rather than human-meaningful work — and AI is the technology that has shown us the destination. (There is a small, almost perfect irony in how the platforms are responding: in June 2026, YouTube rolled out AI-powered comment filtering — machines deployed to clean up what is, in meaningful part, machine-generated noise. Machines reading machines, now with machines moderating them.¹¹⁴)

The *Living Web* — and this is where I find the actual reason for the rest of this book — is something that has to be deliberately built. It is the part of the internet where authorship is provable, where attribution is durable, where attention is allocated on something other than virality, where the architecture itself supports the kind of work that humans do well together. None of that comes for free. None of it is a side-effect of better AI models.

I am not the only person making this argument, and I am nowhere near the best qualified. In June 2026, **Tim Berners-Lee** — the man who invented the World Wide Web — said it publicly and plainly: “*We can fix the internet.*”¹¹⁵ The verb is the point. *Fix* is a builder's word, not a mourner's. The person who designed the original architecture is saying, on the record, that the architecture is broken *and repairable* — which is precisely the wager the Living Web rests on.

We have to make it. On purpose. In the next twelve months.

That is the project the rest of this book is about.

Chapter 5 — The Slop Ceiling

The most important number I have come across in the six months of writing this newsletter is **44 to 3**.

44 is the percentage of daily music uploads to the streaming platform Deezer that are now AI-generated, according to the company's own analysis published in April 2026 — roughly 75,000 tracks a day, more than two million a month.

3 is the upper bound of the percentage of total streams those tracks generate.¹¹⁶

I want you to sit with that ratio for a second. We are looking at a flood of synthetic music nearly half the size of the entire upload pipeline, that the listening audience is, in real time, simply refusing to play. Not banning. Not boycotting. Not legislating against. Just *not pressing play*.

There is, in the language Pete and the *DreamLab* team started using internally around February, a name for what that ratio represents. We call it **the slop ceiling**.

The slop ceiling is the empirical answer to the most common 2024-era question about AI in the creative industries: *does the audience care?* For two years, the assumption in tech circles was that they wouldn't. That the cost-and-volume advantages of synthetic content would eventually swamp human-made work in attention markets, the way industrial agriculture swamped artisanal farming, the way Spotify swamped CDs. That the public would, given enough exposure, develop a taste for the synthetic — or at least, a tolerance.

The 44-to-3 ratio is what it looks like when that assumption is wrong.

And the thesis is no longer a newsletter-internal coinage. In July 2026, *Fast Company* ran a piece under the headline “Our interest in AI slop is hitting a ceiling” — the trade press arriving, more or less word for word, at the framing this chapter has been building since February.¹¹⁷ I take no particular credit for the convergence; when the data is this loud, everyone hears the same thing. But it matters for the argument, because it means the ceiling is no longer a contrarian reading of one streaming platform's numbers. It is becoming the consensus description of the market.

This chapter is about the slop ceiling — what it is, how it is showing up across music, film, advertising, podcasting and the web, who is hitting it from above and below, and what it tells us about the creative economy that is actually forming, as opposed to the one that the platform companies have been forecasting.

The flood

Let me describe the flood, because it is, in absolute terms, extraordinary.

In **October 2025**, when I started the newsletter, the music industry was already in panic over the fact that around 10% of new music-makers, according to a Ditto Music survey, were using AI in their work, down from a Ditto 2023 survey

suggesting around 48% — a number which was itself a shock at the time.¹¹⁸ The major-label CEOs talked about AI as an existential problem; Spotify announced “new protections” for artists, songwriters and producers; Universal and Warner were rumoured to be signing “landmark AI deals within weeks.”¹¹⁹

By **the end of November 2025**, an Israeli streaming-analytics firm reported that 50,000 AI-music tracks were being uploaded to Deezer *every day*.¹²⁰ In a single quarter, the volume of music being added to one streaming platform — by AI — exceeded the entire human-made catalogue uploaded in any month before October.

By **April 2026**, the number was 75,000 a day, on Deezer alone, and *44% of total new uploads*.¹²¹ Deezer’s own statement on its findings was unusually direct for a streaming company: “*AI-generated music is now far from a marginal phenomenon, and as daily deliveries keep increasing, we hope the whole music ecosystem will join us in taking action to help safeguard artists’ rights and promote transparency for fans.*” Universal Music Group’s CEO, in a January memo widely circulated in the music press, called it the “exponential growth of AI slop on streaming services,” adding, in language unusual for a major-label communications stance: “*Let me be clear: UMG will not stand by and watch irresponsible business models take hold — models that devalue artists, fail to provide adequate compensation for their work, stifle their creativity and ultimately, diminish their ability to reach audiences.*”¹²²

This isn’t a sector trend. The same pattern is showing up everywhere I look. In January 2026, *Music Business Worldwide* reported that **56.9%** of new independent songs released in China were AI-generated.¹²³ In March 2026, the term “**podslop**” — synthetic AI podcasts churned out by content farms — entered the trade press, with the *Wrap* reporting that one such operation, Inception Point AI, was producing 3,000 episodes a week.¹²⁴ By the time of Issue 28 in early May, *almost half* of new podcast feeds being added to the major directories were classified by aggregator companies as AI-generated, with little or no human host involvement.¹²⁵

Games hit the same wall from their own direction. In the six months to July 2026, roughly **181,000 mobile games** were released — a flood that *Digital Trends* attributed directly to AI and “vibe coding,” and whose character its headline caught exactly: more games, “but not necessarily better ones.”¹²⁶ By late June, the boss of **CD Projekt** — the studio behind *The Witcher* and *Cyberpunk 2077* — was warning that games made *solely* by generative AI were “already being made.”¹²⁷ Not a forecast. A present-tense supply category.

In November 2025, *Merriam-Webster* named “**slop**” its word of the year, citing the rise of AI-generated content across the web as the primary driver.¹²⁸ In February 2026, *YouTube’s CEO* put “managing AI slop” at the top of her published priorities

list for the year.¹²⁹ In April 2026, on YouTube alone, channels labelled as “AI” content had viewership in the billions for political fake-news content.¹³⁰

The flood is real. The flood is global. The flood is structural — it is not going to subside, because the marginal cost of producing more of it is approaching zero and the marginal benefit, at least at the volume end of the market, is non-zero.

What I want to argue in this chapter is that the flood is also — counter-intuitively, against almost every prediction made in 2023 and 2024 — *not winning*.

The ceiling

The 44-to-3 ratio is the cleanest version of the slop ceiling, but it is not the only one.

Deezer’s parallel analysis, conducted in partnership with Ipsos in late 2025, found that **97% of listeners** could not reliably distinguish AI-generated music from human music in a blind test.¹³¹ On the face of it, this is bad news for the human side — if you can’t tell the difference, why pay the difference? But the same study found that *when listeners were told a track was AI-generated, their willingness to engage with it dropped sharply*. The Adobe Creators’ Toolkit Report had a similar finding from the production side: in a December 2025 *Bain & Company* report titled *In an AI Age, People Still Want the Radio Star*, the firm found that audience engagement with AI-disclosed work fell well below engagement with human-labelled work, holding all other variables constant.¹³²

So here is the empirical picture, in one sentence: audiences can’t tell the difference, but when they find out, they care.

This is — and I think this is the part that almost everyone in the platform economy has been slow to understand — *not a temporary cultural reaction*. It is a structural property of how attention works in oversupplied markets. When everything is abundant and indistinguishable, the only thing left that allocates attention is meaning. And meaning, for human audiences, requires a knowable human source.

This is why Deezer’s chart looks the way it does. Up to 85% of the streams that AI-generated music *does* get on Deezer were identified by the company in 2025 as **fraudulent** — bot-driven, click-farm-driven, streaming-fraud-driven.¹³³ The actual human listening to the actual human-produced flood of AI tracks is, on the most generous estimate, a fraction of a percent of the platform’s overall listening time. The flood is hitting a ceiling not because the audience is wise. It is hitting a ceiling because the audience, presented with a near-infinite menu, makes its choices in a way that systematically underweights the synthetic. The fraud has a sharper edge still. In July 2026, *Digital Music News* reported on a band who woke to find that an AI operation had duplicated their album at a slightly altered speed and siphoned off **94% of their Spotify royalties** — the flood not merely competing with human

work but impersonating it, draining the earnings of the people who actually made the thing.¹³⁴

There is, in another domain entirely, a clean analogue for what this audience behaviour describes. In March 2026, *Bloomberg* reported on what AI had done to elite chess: at the very top of the game, machine-optimal play had produced an epidemic of *draws*. When both players have memorised the machine-optimal lines, both play optimally, and both tie. The grandmasters' response, the piece reported, was to *deliberately play sub-optimal moves* — moves that the engines would not endorse, but that the opponent, having trained against the engines, had not seen.¹³⁵

What is happening to elite chess is what is happening, at the audience layer, to streaming music. The machine-optimal output saturates; the surface of the work becomes indistinguishable from itself; the listener's attention systematically shifts towards the work whose *Why* the engine could not have generated. The 44-to-3 is the slop ceiling in numerical form. The chess-grandmasters' sub-optimal move is the slop ceiling at the practitioner's end of the same dynamic. Chapter 16 builds the long-form argument out of this analogue. For now: hold the picture in mind that the audience, in the aggregate, is the room of grandmasters refusing the machine-optimal line.

Xania Monet

The most prominent test case of this dynamic, in the autumn of 2025, was a virtual R&B artist called **Xania Monet**.

Monet was created using the AI music platform Suno, with lyrics written by Telisha Jones — a Mississippi-based poet and designer who built the character around her own life and stories.¹³⁶ Monet was not a synthetic-from-nothing artist. She was a synthetic *vessel* for a real human songwriter's words. The vocal performance was AI; the lyric was Jones; the persona was a collaboration.

The week dated **20 September 2025**, Monet debuted on the Billboard Emerging Artists chart at No. 25 and on the Hot Gospel Songs chart at No. 21 with a track called *Let Go, Let God*.¹³⁷ A few weeks later, her track *How Was I Supposed to Know?* became the first AI-led song ever to enter a Billboard radio airplay chart, debuting at No. 30 on Adult R&B Airplay.¹³⁸ In November, after a bidding war between several labels, the entertainment company **Hallwood Media** — led by former Interscope executive Neil Jacobson — signed her to a deal reported by *Billboard* and the *Bangkok Post* at **\$3 million**.¹³⁹

The response from working musicians was immediate and almost uniformly negative. The R&B singer Kehlani posted a video — which she later deleted — calling out the deal directly: “There is an AI R&B artist who just signed a multimillion-dollar deal,” she said, “and the person is doing none of the work.”¹⁴⁰

Telisha Jones, the human lyricist behind the Monet vocal, gave Billboard a counter-framing that I find more interesting than either Kehlani's outrage or Hallwood Media's PR. "*It's not a hook and a bridge and a catchy chant — it's just the lyrics, and they are pure,*" she told the magazine.¹⁴¹ The whole transaction — labour, attribution, deal value — was happening in the space between Jones' words and the synthetic voice that delivered them. Whose work is the work? Whose name goes on the contract? Whose royalty cheque arrives in the post? Those questions, in late 2025, had no settled answer, and the music industry spent the next six months arguing about them in real time.

The interesting thing, six months on, is not the outrage. The outrage was the expected response. The interesting thing is that *Xania Monet has not become a star*. She has not, by any of the standard pop-cultural metrics, *broken through* in the way a \$3M new signing usually would. The Billboard chart entries were a moment. The radio airplay was a moment. The cultural impact, six months in, is mostly that her name is the name everyone uses to ask the question *can an AI artist actually become a star?* — and the working answer, so far, is *not yet, and possibly not*.

She is, in the slop-ceiling frame, the upper edge of what's possible. A real human songwriter, a real lyrical vision, a real cultural specificity (Mississippi R&B, gospel-adjacent, a particular Black American spiritual tradition); a high-quality AI voice; a multi-million-dollar marketing budget. The cultural product made by combining all of those things has hit the ceiling somewhere short of cultural escape velocity.

The same pattern, more starkly, is visible in **Breaking Rust**, the AI country act whose track *Walk My Walk* hit No. 1 on Billboard's Country Digital Song Sales chart in November 2025 with about 3,000 paid downloads.¹⁴² *Walk My Walk* did exceptionally well by AI-music standards — over 3 million Spotify streams in less than a month, an Emerging Artists Billboard debut at No. 9 — and then plateaued. The *Washington Post* and TIME both ran pieces in late 2025 raising the possibility that the chart performance had been partly manufactured, and that the streaming traffic, while large by AI-music standards, was unusually concentrated in the kinds of automated playlists where streaming fraud is most common.¹⁴³ Nashville, by every available report, was unsettled — but the Nashville Songwriters' Association didn't see anything close to the kind of fan-driven cultural takeover that the song's chart position would have implied if it had been a human single performing the same way.¹⁴⁴

The pattern repeats again with **Sienna Rose**, the mysterious AI artist who racked up millions of Spotify streams in late 2025 and whose identity prompted a BBC investigative feature in January 2026 ("Who, or what, is she?").¹⁴⁵ The pattern repeats with the **MAGA gospel rapper** who used AI to climb the charts in November 2025;¹⁴⁶ with the **AI band Bleeding Verse** whose creator signed with

Hallwood Media in October 2025;¹⁴⁷ with **Trilok**, the Indian AI band the Indian government had to publicly disavow association with in December 2025 after a live performance.¹⁴⁸ In May 2026, an AI-generated Afrobeats track displaced **Tyla** from the No. 1 spot on Billboard's Afrobeats chart — the first time, on the public reporting I have read, that an AI-led song had taken the top position on a Billboard genre chart in a primarily African-music category. The chart performance was, as in the Breaking Rust case, unusually concentrated in automated streaming traffic; the cultural footprint, six weeks on, was again *not stardom* but a brief news-cycle moment.¹⁴⁹

In every single case, the AI act hits a ceiling. They chart. They make money for somebody, often a lot of money. They generate headlines. But they do not become *stars*. Their cultural shadow stops at the edge of the news cycle and does not propagate into the next one.

The audience is doing something at the margin of these careers. It just isn't *quite* showing up in any of the metrics the labels are used to looking at.

The pushback

The flood and the ceiling describe the supply and demand sides of the market. What the *culture* is doing — the people, the institutions, the labels, the platforms — is the third leg.

Through autumn 2025 and winter 2026, the cultural pushback intensified in waves. Some of it was symbolic. In November 2025, **Paul McCartney** released a silent track as part of a wider music-industry protest against the UK government's proposed copyright opt-out scheme.¹⁵⁰ In December, the **Eurythmics'** **Dave Stewart** argued — slightly against the grain of the protests — that musicians needed to “embrace the unstoppable force” of AI and license their intellectual property rather than fight it.¹⁵¹ In January, almost 800 creators including Jason Aldean and OneRepublic signed an open declaration titled *Stealing Our Work Is Not Innovation*.¹⁵² In May 2026, **Jack Antonoff** — one of the most-cited producer-songwriters of the period — went considerably further than McCartney in the public register, calling AI music-makers “*godless whores*” in an interview that became the headline-grabbing artist-side moment of the post-I/O news cycle.¹⁵³ Antonoff's framing is, in my read, a useful marker of how far the cultural register of the resistance has shifted between McCartney's *silent track* in November 2025 and the spring of 2026: from elegiac protest to active contempt.

By June 2026, the register had shifted a third time — from contempt to *counter-programming*. The *Guardian* reported on artists deliberately making what they called “**anti-slop**” — work positioned, made and marketed as the explicit opposite of AI-generated content, by creators who felt the synthetic flood had “been rammed

down our throats.”¹⁵⁴ A feature film was released **straight to VHS** — billed as the first such release in twenty years — with its makers describing the choice to *Variety*, in so many words, as a “middle finger” to AI: the physical format itself doing the authenticity work, distribution-as-protest.¹⁵⁵ And in early June, “**Say No To Suno**” protesters brought the fight to a swanky Santa Monica AI conference — the resistance no longer waiting for the AI-music industry to come to it, but following the industry to its own venues.¹⁵⁶ None of these are large economic events. What they are is evidence that the ceiling has a *creator-facing* surface as well as an audience-facing one: an emerging aesthetic identity, with its own name, built entirely out of not being slop.

Some of it was practical. In November, **Universal Music Group** announced a strategic alliance with Stability AI for “responsible” music tools.¹⁵⁷ In December, **Warner Music Group** signed a similar deal with Stability AI.¹⁵⁸ At almost the same moment, **Splice** and **Universal Music Group** agreed to collaborate on “next-generation AI-powered music creation tools for artists” — a structural acknowledgement that the labels’ strategy had pivoted from purely *suing* the AI companies to *partnering* with them.¹⁵⁹

Some of it was legal. In January 2026, the German rights society **GEMA** won a major ruling against OpenAI in the Munich Regional Court, on training-data grounds.¹⁶⁰ **Suno** was sued by music-rights groups under a banner the litigators called “the biggest theft in music history.”¹⁶¹ **Wixen Music Publishing** filed a \$50m copyright suit against Meta.¹⁶² **Universal Music Group** filed a \$3B suit against Anthropic.¹⁶³ By the end of February 2026, the lawsuits were no longer an interesting subplot. They were the main mechanism through which the new creative economy was being defined.

Some of it was platform policy. **Bandcamp** banned AI-generated music outright in January 2026.¹⁶⁴ **Deezer** built and licensed an AI-music detection tool to other platforms.¹⁶⁵ **Spotify** declined to add an AI-music filter, preferring transparency and labelling.¹⁶⁶ **San Diego Comic-Con** banned AI art at its 2026 event.¹⁶⁷ **Sweden’s official music chart** banned AI-generated entries.¹⁶⁸ And in June 2026, **Tidal** found the most economically legible position of all: AI-generated music would be *allowed* on the platform — but it would not earn royalties.¹⁶⁹ Not a ban, not a filter, not a label. A price. A major streamer looked at the synthetic flood and formally valued it at zero — which is, give or take a rounding error, exactly what the audience’s 44-to-3 attention behaviour had already been saying.

The most compressed test of audience enforcement in the entire period — more compressed than any chart performance or streaming metric — arrived in late May

2026, when **Amazon** launched its **AI Creators Fund** and commissioned three animated Prime Video series using generative AI. One of those series was **Punky Duck**, created by **Jorge Gutierrez**, the animator best known as the creator of *El Tigre: The Adventures of Manny Rivera* on Nickelodeon. Within **48 hours** of the announcement being confirmed publicly, Gutierrez had dropped out. The reason, as he told the trade press and as the Hollywood Reporter documented in detail, was “*relentless derision*” — online backlash so severe that he quit before a single frame had aired, before a single trailer had been cut, before anyone had seen the work.¹⁷⁰ **Guillermo del Toro** — who had spent October 2025 telling Variety he would “rather die” than use generative AI — subtweeted the announcement.¹⁷¹

The Amazon/Punky Duck case is, in slop-ceiling terms, the most extreme data point in this book’s six-month window. Every other test case — Breaking Rust, Xania Monet, the AI Afrobeats chart topper — hit the ceiling *after* the work reached an audience. Punky Duck hit the ceiling *before it existed*. The audience enforcement mechanism did not operate at the level of listening, watching, or streaming. It operated at the level of a *press release*. The announcement of AI involvement was sufficient to generate backlash severe enough to cause the creator to exit the project. Whatever the commissioning economics of Amazon’s AI Creators Fund, the cultural economics of an audience that will punish AI-labeled animation before it has been made is a data point that no studio’s investment thesis can ignore.

By June 2026, two more platform-side data points had landed and are worth recording because they bracket the slop-ceiling argument from opposite directions. **Apple Music**, on the *defensive* side of the AI flood, publicly disclosed that AI-generated songs accounted for **less than 1% of plays on the platform** — the cleanest single confirmation, from the most curated of the major streamers, that audience-chosen listening time has remained almost entirely human-authored even as upload-side AI volume has continued to climb.¹⁷² On the *offensive* side, **Spotify** signed a landmark deal with Universal Music Group that lets paying Spotify Premium subscribers create AI covers and remixes of UMG-catalogue tracks as a *licensed Premium add-on*¹⁷³ — a clear move into the *consent-trained* territory I describe at length in [Chapter 6](#), and a structural acknowledgement by the largest streamer in the world that the AI-music market only works at scale if the rights infrastructure is built underneath it. Daniel Ek, defending the same move from the artist-side criticism, made the *slop ceiling* point in his own words: AI music, he argued, is “*better than slop*” — a phrasing the dissenting artists treated as gratuitously dismissive but which lands, structurally, on the same observation this chapter has been making.¹⁷⁴ At the corporate-deployment end of the same gradient, **Tonada** — a Swedish startup that began the period as a small AI-music demo — pivoted into the *retailer-facing* AI music market, supplying brand-aligned synthetic

background music to chain retailers as a licensed B2B service.¹⁷⁵ The pattern across all three moves is consistent: the path from the slop floor to the licensed integration runs through *consent, licensing and disclosure*, not through prohibition. The labels and platforms that solve that path early are the ones building the post-Petrillo settlement Chapter 6 describes.

Then, in the last stretch of June 2026, the enforcement went kinetic. **YouTube** — whose CEO had put “managing AI slop” at the top of her published priorities back in February — began actively cracking down on mass-produced synthetic content, and the *Hollywood Reporter* documented the immediate casualty list: the “**faceless creators**,” the high-volume, presenter-less AI channels that had been the flood’s most visible commercial success story, took a direct hit to reach and revenue.¹⁷⁶ This is, structurally, the data point the whole chapter has been building towards. Until now, the slop ceiling was a pure audience behaviour — an aggregate pattern of not-pressing-play that showed up in streaming ratios and chart half-lives. As of June 2026, the platform whose entire ad business runs on audience attention has written that behaviour into policy. The ceiling is no longer only enforced by the audience. It is being enforced, on the audience’s behalf, by the platform layer itself — because a platform that monetises attention cannot afford to keep subsidising content the attention systematically refuses.

The point I want to make about all of this is that the cultural pushback is not — as the more dismissive coverage tends to frame it — a Luddite reaction. It is not an irrational allergy to new technology. It is a *market response*. The audience has spoken with its attention. The platforms are reacting to the audience. The labels are reacting to the platforms. The lawyers are reacting to the labels. The artists are reacting to the lawyers. The whole system is, in slow motion, *renegotiating the terms on which synthetic creative work is allowed to participate in the public sphere*.

That renegotiation is the actual story. The viral AI hits, the ones that get the magazine covers, are footnotes.

What the ceiling is made of

I want to try, in the last part of this chapter, to articulate what the slop ceiling actually is — what cognitive, cultural, economic mechanism produces the 44-to-3 ratio — because I think understanding that mechanism is the difference between thinking it will hold and thinking it will eventually erode.

My working hypothesis, after six months of looking at the data, is that the ceiling is made of four overlapping things:

One. *A cognitive distinction the audience can’t articulate but can feel.* The Deezer/Ipsos finding — that 97% of listeners can’t pick AI from human in a blind

test, but that revealed-AI tracks underperform — suggests the distinction is below conscious recognition but above zero. The audience knows when something doesn't matter, in some way they can't quite name.

Two. *A status-signal collapse.* Music, film, advertising and podcasting are, in significant part, status goods. Telling your friends you discovered an exciting new artist is part of why people pay attention to artists. AI artists, by being mass-produced and machine-authored, fail the status test at the structural level. There is no way to *signal cultural insiderness* by being the first to discover Xania Monet, because Xania Monet was discovered by 800,000 people in the same week, all of whom found out via the same press cycle.

Three. *A meaning vacuum.* Most of the slop is produced for content marketing, SEO and ad placement reasons, not because anyone needs to express anything. The audience can tell. As one *Digital Music News* headline from January 2026 put it, in a phrase I have written down in my notebook and used in talks: “*A.I.-generated music is catchy, familiar... and boring.*”¹⁷⁷ The Swedish Top Chart's reasoning when it banned an AI-generated track from its rankings in January 2026 said the same thing in different words: “*The song is great, but unfortunately, it's missing one of the most important ingredients, which is emotion.*”¹⁷⁸ The technical capability is there. The reason for the work to exist is not.

Four. *Reciprocity.* And this is the one I am least confident about, but find the most interesting. There is, in almost every long-running creative relationship between an artist and an audience, an implicit reciprocity. The audience pays attention because the artist has *paid the price of making the work* — has practised, has struggled, has lived, has earned the right to be heard. AI artists short-circuit that contract. They produce the output without paying the price. The audience, at some level, refuses the trade.

Take any of those four mechanisms away and the ceiling might drop. Take all four away and we'd be in trouble. None of them, on the current evidence, is going away.

The Authenticity Premium: pricing the ceiling

The slop ceiling, as I have described it so far, is the *negative* side of the dynamic — the work the audience refuses to engage with. There is a *positive* side I want to spend a section on, because it is the side that pays the bills, and the side the rest of this book leans on for its strategic argument.

I have come to call the positive side the **Authenticity Premium**: the measurable excess of attention, willingness to pay, and cultural credit that audiences allocate to creative work whose human authorship can be verified. If the slop ceiling tells you what the audience *will not engage with*, the Authenticity Premium tells you what they *will pay extra for*. Both numbers are produced by the same underlying

audience behaviour. Both are market findings. Both have stabilised across the six months of newsletter coverage that this book is built on.

Let me try to sketch the empirical picture.

On the *music* side, the cleanest evidence is the inverse of the Deezer 44-to-3. Independent labels, working musicians and the back-catalogue of clearly-human-authored recorded music have, by Deezer's own published April 2026 statement, retained a *disproportionate share of total listening hours* relative to their share of the upload volume. The artists with verified human-authorship signal — touring records, named producers, signed labels, press-covered releases — capture a meaningful premium in playlist placements, editorial coverage, and per-stream consumption versus the long tail of synthetic uploads. The premium is, on Deezer's own framing, a function of *audience-chosen* listening rather than algorithmic placement. The 1-to-3% of streams that the AI-music flood receives is, in operational terms, the *floor* of the slop ceiling; the corresponding 97-to-99% that named human artists capture is the *ceiling-level* the Authenticity Premium underwrites.

On the *advertising* side, *Marketing Week's* “You can't dismiss AI ads as slop when they're winning in testing” piece — which I quote at length in Chapter 13 — found that *care-driven* AI work could win creative-effectiveness tests, but *only when the production process foregrounded human creative intent*.¹⁷⁹ The same publication's reporting on the McDonald's Netherlands, Valentino and Coca-Cola AI campaigns documented that *cynical* AI work, by contrast, produced sharply negative effectiveness scores. The price of being a *cynical* AI advertiser, in commercial-effectiveness terms, was a measurable revenue penalty. The price of being a *sincere* AI advertiser was a measurable premium. Both numbers come from the same audience exposed to the same kind of work.

On the *film and television* side, the Authenticity Premium has been visible in awards-season behaviour, in the *Marvel-era IP fatigue* data, and in the audience response to specific high-profile AI integrations. Films whose AI use was *disclosed in advance and framed as part of the work's identity* — *Watch the Skies*, *Lily*, Andrii Daniels' bomb-shelter Christmas clip, *Dear Upstairs Neighbors*, *Synthetic Sincerity* — have, on aggregate, received better critical reception and stronger audience engagement than films whose AI use was *concealed and later revealed*. The premium is, structurally, paying for *transparency about authorship* rather than for *authorship purity*. The audience is not refusing AI involvement. The audience is refusing *deception* about AI involvement.

On the *games* side, the Quantic Foundry survey I referenced earlier and the parallel work by IGN and PC Gamer through autumn 2025 showed an effect that is, in commercial terms, more direct. Games that disclosed AI involvement in *audience-visible roles* (NPC dialogue, generated quests, voice-acting replacement) faced

specific consumer pushback in Steam reviews and Metacritic user scores. Games that disclosed AI involvement in *invisible-pipeline roles* (asset generation under art-direction supervision, QA automation, localisation) faced no measurable consumer pushback at all. The premium, in games, is paying for *AI being kept off the audience-visible parts of the work*. Studios that have understood this calibration — Larian, Hooded Horse, Jagex, the *Position Three* signatories of Chapter 7 — are operating with the Authenticity Premium as a deliberate marketing position.

The premium has a *commercial* shape that I want to name explicitly, because I think working creatives are systematically under-pricing it in 2026. The price difference between *clearly-human-authored* and *clearly-AI-authored* creative work, on the available evidence, is not a marginal one. It is, in some segments, an *order-of-magnitude* difference. Xania Monet's \$3M Hallwood Media deal is the *upper bound* of what the synthetic-artist market can pay for a top-tier AI act. The *touring-artist* economics of even a mid-tier independent musician with a verified live following easily exceed that over a five-year career. The same comparison holds across film and games. The Authenticity Premium is, in market-pricing terms, the *commercial gap* between the synthetic and the human. It is not small. It is, for working creatives positioning themselves above the slop ceiling, the actual business case.

There are two empirical caveats I want to put on the page, because the book should not over-claim.

One, the premium is *unevenly distributed*. It accrues, by my reading of the 2025–26 data, most strongly to artists at the *top* (premium auteur cinema, signed-label musical artists, name-brand designers) and at the *long-tail* (hyperlocal, culturally-specific, niche-community-serving creators). It accrues weakly, if at all, to the *middle* — the mid-career journeyman creatives who have built careers on producing competent-but-mean-of-the-distribution work. The middle is, as I argue in Chapter 7's legacy-vulnerability section, the segment most exposed to AI substitution. The Authenticity Premium does not protect everyone equally.

Two, the premium is *dependent on the provenance infrastructure* I described in Chapter 4. The audience can pay a premium for verifiable human authorship only to the extent that authorship is *verifiable* — to the extent that the C2PA chains, the SynthID watermarks, the platform-level disclosure norms and the institutional certifications make the signal reliable. The premium and the provenance stack are, structurally, *the same project seen from different angles*. The premium pays for the signal. The signal is produced by the infrastructure. The infrastructure needs the premium to fund itself.

The Authenticity Premium is, in operational summary, *the audience paying the bill for the Living Web*. Whether the bill stays paid, over the next five years, depends on

whether the infrastructure that produces the signal stays trustworthy. That, in turn, depends on the policy, platform and union work this book has been describing.

The corollary

I want to close this chapter with a corollary, because it has implications for the rest of the book.

If the slop ceiling is real, and if it is structural, and if it is going to hold — and I think the burden of evidence at this point is on the people who say otherwise — then the strategic question for everyone in the creative industries is not *how do we compete with the flood?* It is *how do we sit above the ceiling?*

That question has different answers in different sectors. For a working musician, it might mean leaning into the irreducibly human parts of the work — live performance, personal relationship with audience, transparent process. For a working filmmaker, it might mean making the *kind* of film whose value depends on being knowably authored by knowable people. For a working games studio, it might mean — as Jagex, Larian, Games Workshop, Hooded Horse and an increasing list of studios have explicitly said — taking *generative AI off the table* as a public commitment to the audience.

None of these are anti-AI positions. They are *above-the-ceiling* positions. They take seriously the fact that the world is now full of cheap synthetic content and ask: *what is the work that the synthetic content can't do?*

That is the question that organises the rest of the book. The slop ceiling is the negative space against which everything interesting in creative work for the next ten years is going to be defined.

Chapter 6 — The 88%

On 15 December 2025, the UK government quietly laid a document before Parliament that I think will be remembered, ten years from now, as a more important moment in the history of creative AI than any single tool release in 2025 or 2026.

The document was the *Statement of Progress on Copyright and Artificial Intelligence*, prepared by the Department for Science, Innovation and Technology.¹⁸⁰ It was a stocktaking report — not a final policy, not new legislation, not a decision. It was a “where we are” note, eleven months after the closing of one of the largest copyright consultations the United Kingdom has ever run.

The consultation had been open from 17 December 2024 to 25 February 2025. The government had proposed four options for how UK copyright law should treat AI training:

- **Option 0**, do nothing.
- **Option 1**, require AI companies to license copyrighted works for training in all cases.
- **Option 2**, a narrower set of exceptions, with conditions.
- **Option 3** — the government’s *preferred* option — a broad text-and-data-mining exception with an opt-out for rightsholders.¹⁸¹

Eleven and a half thousand people replied.

Of the 10,112 responses submitted through the government’s *Citizen Space* online portal — the subset for which the government published quantitative breakdowns:

- **88%** supported Option 1 — licensing in all cases.
- **7%** supported Option 0 — do nothing.
- **3%** supported the government’s preferred Option 3.¹⁸²

I want you to look at those numbers again, because they are the single most concrete thing this book has to offer in defence of the argument I will be making in the second half of it: that the creative economy is not waiting to be told what it thinks about AI.

Eighty-eight per cent. In a country with no compulsory voting, no organised industry mobilisation comparable to the music or film unions’ rapid responses to specific provocations, no celebrity-led campaign on the scale of the SAG-AFTRA strike: 88% of the people who took the time to write to their government about how their work should be used said, *license it*. Pay for it. Don’t take it.

That number is the *true* watershed of the period this book covers. The Tilly Norwood week made it possible. The 88% made it permanent.

This chapter is about how a global creative coalition — half informal, half deliberate, half union-led, half artist-led, half lawyer-led, all of it networked — went from a few

scattered protest statements in October 2025 to a structural force in policy and law by May 2026.

Why eighty-eight per cent

It is easy, looking at a number that big, to assume it represents some kind of organised lobbying effort. To assume that the AI companies' opt-out proposal was so unpopular that the response was a coordinated push from a few large interest groups, who marshalled their members into the consultation.

The actual composition, as analysed in the December 2025 Statement of Progress, was mixed. There were submissions from creators in every major creative discipline — writers, musicians, filmmakers, photographers, illustrators, designers, journalists. There were submissions from professional bodies (the Society of Authors, the Association of Photographers, the Authors' Licensing and Collecting Society). There were submissions from individual citizens with no industry affiliation, who simply objected on principle to having their work — their LinkedIn posts, their family photos, their blogs — pulled into a training set without their consent.¹⁸³

There were submissions from the AI companies too. The progress report notes that the 3% who supported the government's preferred Option 3 were "particularly concentrated among AI developers and large technology companies."¹⁸⁴ This is — and I am being careful about how I phrase this, because the Statement of Progress is itself careful — *not* a description of a balanced industry view. It is a description of a public consultation in which the people most affected by the proposed policy said one thing, and the companies the policy was designed to enable said the opposite.

The 88% is not a curiosity. It is a *vote*, in the most literal sense. The creators of the United Kingdom were given a structured chance to say what they wanted, and 88% of them said the same thing.

The Society of Authors' submission, which I have read in full, made the underlying argument with the kind of clarity that the policy debate had been avoiding for two years. *"If we are to see an end to the industrial-scale theft of writers' and other creators' work, and to protect the creators and creative industries of the future, then UK copyright needs to be enforced not weakened."*¹⁸⁵ That sentence — *industrial-scale theft, enforced not weakened* — set the rhetorical register that the next six months of the policy debate ran on.

The pattern repeats

I think a lot of the international coverage of the UK consultation has under-emphasised that the 88% was not a UK-only phenomenon. It was the first formal expression of a pattern that was, in the same six months, repeating in every jurisdiction that gave its creators a meaningful chance to speak.

In **Germany**, the music rights society **GEMA** sued OpenAI in the Munich Regional Court over the training of large language models on copyrighted music lyrics. In November 2025, the court ruled for GEMA in a decision that intellectual-property lawyers across Europe — including Dr Barry Scannell, whose detailed LinkedIn breakdown of the ruling I have read more times than I will admit — described as a *major* precedent for European copyright law.¹⁸⁶

That precedent held, and then it compounded. In August 2026, GEMA won again — this time against **Suno**, the largest of the AI music-generation platforms, in a landmark ruling that turned the November 2025 decision from an outlier into a line of authority: the same collecting society, the same court system, and now a generative-music company losing on the training question itself.¹⁸⁷ *Creative Bloq's* question in the aftermath — what does the biggest AI copyright payout in history actually change for creatives? — deserves, I think, a precise answer: not a windfall for individual artists, but a *price signal*. Training now has a cost on the books, which means licensing now has a floor to negotiate up from.¹⁸⁸ And the platforms were fighting the same war docket by docket, taking whatever procedural ground each ruling offered — only a month earlier, Suno had publicly seized on Udio's procedural win against Sony Music over which infringed works were admissible in that case.¹⁸⁹ Neither side was treating any single judgment as the end of the argument. Both sides were treating every judgment as a rate-setting input for the licensing table.

In **the United States**, a coalition of music rights organisations sued **Suno**, with the press release describing the action, in a phrase the litigators clearly knew would travel, as “the biggest theft in music history.”¹⁹⁰ **Wixen Music Publishing** filed a \$50m copyright suit against **Meta** in January 2026.¹⁹¹ **Universal Music Group** filed a \$3B suit against **Anthropic**.¹⁹² **The Johnny Cash estate** sued Coca-Cola under the **ELVIS Act** — Tennessee's new AI-impersonation law — for using a Cash sound-alike in a tribute-act advertisement.¹⁹³ By the spring of 2026, the litigation landscape was so dense that *Music Business Worldwide* was running weekly summary columns just to keep track of which cases were still active.

In **the European Union**, lawmakers tabled a bill in November 2025 seeking an EU-wide minimum age to access AI chatbots and social media, an early acknowledgement that the regulatory question was not just about copyright but about the wider integration of AI into the social fabric.¹⁹⁴

In **the United States**, the actors' union **SAG-AFTRA**, riding the wave of the Tilly Norwood backlash, opened negotiations in October 2025 that resulted by spring 2026 in *significantly stronger AI protections* in its next four-year contract — a deal that included new consent requirements, residuals, and what the trade press began calling, informally, the “Tilly Tax” on the use of AI actors.¹⁹⁵

In **the United Kingdom**, the U.K. actors' union **Equity** held a strike ballot in December 2025 over AI scanning of performers' likenesses; the result came back in a 99% landslide in favour of industrial action. The ballot question itself, in its plain language, captured the substance of what was at stake: "*Are you prepared to refuse to be digitally scanned on set to secure AI protections?*"¹⁹⁶ By January 2026 the union had secured what its general secretary called "an improved offer" from producers on AI protections in film and TV negotiations.¹⁹⁷ In May 2026, the broader **AI Disclosure Standard** for the film industry was launched at the **Cannes Film Festival**.¹⁹⁸ In the same week, the **British Phonographic Industry (BPI)** issued a formal set of *transparency and sovereignty* demands aimed at the music side of the same settlement — a structured industry position designed, in the BPI's own framing, to *secure* the "AI licensing boom" rather than leave it to bilateral negotiation between platforms and rights-holders one model at a time.¹⁹⁹

The pattern, in every jurisdiction and across every part of the creative economy, was the same. Where creators were given a procedural mechanism — a consultation, a strike ballot, a contract negotiation, a class action — they used it. They turned up in numbers. They voted, in their structured way, against the unconditional appropriation of their work. And they won enough of these procedural battles that, by the time the spring of 2026 arrived, the *terms of engagement* for AI in the creative industries had been substantially re-set in a six-month window.

The first reversal

The most surprising single event in the entire policy arc was the UK government's own *reversal* in spring 2026.

The Statement of Progress in December had already softened the official position. Where the original consultation had proposed Option 3 — the text-and-data-mining exception with opt-out — as the *preferred* outcome, the December update simply described the government as "working with 50+ experts from across music, film, games and AI to figure out what comes next."²⁰⁰ The opt-out language was gone.

By **March 2026**, the position had reversed further. The government's final report on copyright and AI, laid before Parliament by the statutory deadline of 18 March 2026, walked back the original preference for Option 3 in favour of a much more cautious set of proposals that acknowledged the 88% finding.²⁰¹ *Dream Machine Issue 21*, dated 19 March 2026, was the first edition where I noticed the change in tone in the government's own language. The framing had shifted from "how do we enable AI training" to "how do we protect creators."

I want to be precise about what this reversal means and what it doesn't.

It does *not* mean that the UK has banned AI training on copyrighted work, or that it has imposed a licensing-first regime by default. As of the time I am writing this — May 2026 — the legislative process is ongoing, and the eventual policy could land anywhere on a wide spectrum.

It *does* mean that 88% of 10,112 people, plus a thousand-odd email submissions, plus a media cycle that ran for fourteen months, plus a parallel set of legal proceedings, plus a parallel set of platform and industry pushback, plus the active mobilisation of multiple professional bodies, was enough to *change the position of a national government* on one of the most economically significant technology-policy questions of the decade.

That is, in democratic terms, what *working* looks like.

What the creators were actually saying

The 88% was a procedural answer to a procedural question. What were the *substantive* arguments behind it? I have read enough of the submissions, through the published summaries and through the secondary press coverage, to feel confident in summarising the three I see most often.

The consent argument. This was the simplest and the most universal: that work made by a creator — a song, a book, a photograph — belongs to that creator in a way that is *not* fully captured by the existing copyright regime, and that the use of that work to train a machine learning model is a use that requires the creator's consent.

The argument is not new. The Berne Convention has, since 1886, treated authorship as a *moral right* in addition to an economic one. What is new is the scale of the use. A single AI training run can ingest the work of millions of human creators in a way that no single buyer, publisher, broadcaster or aggregator has ever done. The procedural mechanisms of copyright were designed for a world where uses were enumerable. They struggle in a world where the use is, in effect, *the entire creative output of a generation, all at once*.

The attribution argument. This was the most operationally specific: that when AI systems produce derivative outputs based on training data, the creators whose work shaped those outputs should be identifiable, and where appropriate, compensated. *Musical AI*, a startup that raised \$4.5m in January 2026 on a “creative weight attribution” model, described the technical version of this as “calculating each input’s actual contribution to a generative model’s output, then licensing accordingly.”²⁰² The argument doesn’t require AI training to stop. It requires it to *show its workings*.

For most of the period this book covers, that demand was aspirational — the training sets were sealed, literally and figuratively. In June 2026 the seals started to come off. A federal judge vacated the order sealing **Udio’s** training data in the ongoing music-

industry litigation, opening the prospect of the first public accounting of exactly how many commercial recordings a major music-generation platform ingested.²⁰³ In the same month, *The Atlantic* published an investigation with a searchable tool that lets any artist check whether their music sits in the datasets used to train AI models²⁰⁴ — and artists started looking. **SZA**, discovering that her songs had been used in AI training, hit out at what she called ‘disgusting’ AI music.²⁰⁵ I flag the SZA moment not for the quote but for the mechanism underneath it: the consent argument changes register the moment a working artist can look up her own catalogue in a training set. Ingestion stops being an abstraction about ‘data’ and becomes a named person finding her named work in a place she never agreed for it to be. Every transparency instrument that ships — a court unsealing, a journalist’s search tool, a disclosure standard — converts some quantity of diffuse unease into specific, actionable grievance. That is what *show its workings* is for.

The economic argument. This was the most cynical and the most powerful: that AI systems trained on the unpaid labour of creators will eventually substitute for those creators in the market, and that the failure to license is therefore not just an *ethical* offence — it is an active transfer of wealth from a relatively diffuse group of working creatives to a relatively concentrated group of technology platforms and their shareholders.

The PRS for Music *2026 AI Survey* found that **four in five** music creators worried about AI-generated music competing with human-created music in the streaming economy.²⁰⁶ The Edinburgh-based **Centre for Creative AI** at UCL/RCA, launched in late 2025, explicitly framed its mission around the “redistribution of value from machines back to the humans whose work made them possible.”²⁰⁷ The U.S. **artist trade body** quoted in *Complete Music Update* in November 2025 was even blunter: “Artists must have creative control in AI deals or risk ending up with ‘scraps’.”²⁰⁸

Stack those three arguments next to each other and you get a recognisable shape. It is the shape of every economic-rights argument creators have made, in every previous technological transition, going back to the Stationers’ Company in seventeenth-century London. *Don’t print without permission. Don’t broadcast without a fee. Don’t sell our records without paying us. Don’t sample without clearing. Don’t stream without licensing. Don’t train without consent.*

The 88% is the latest entry in a four-hundred-year sequence. What’s new is the *speed* with which it has had to be expressed, and the *scale* of the use it is responding to.

The artists’ declaration

In January 2026, in parallel with the union negotiations and the lawsuits and the policy responses, nearly **800 working creatives** — including high-profile names

like Jason Aldean and OneRepublic — signed an open declaration with the line that gave the document its name: *Stealing Our Work Is Not Innovation*.²⁰⁹

I want to spend a moment on this document because it is the cleanest expression I have found of the underlying argument, and because I think the line will be on a t-shirt within a year if it isn't already.

The declaration was not a legal document. It had no enforcement mechanism. It did not call for specific legislation. It was a *cultural statement* — a refusal of the framing under which the AI companies had been making their case.

The framing the AI companies had been using, repeatedly, in venues from technology conferences to court filings, was that training models on copyrighted material was a kind of *technical inevitability* — that machine learning required vast amounts of data, that the data could not practically be licensed at scale, and that therefore the use was, in a sense, *outside* the traditional consent-and-payment framework of copyright. It was — they argued — not really “use” in the sense the law had been built around. It was a new kind of activity that needed a new kind of rules.

The declaration's response, in a phrase, was: *no, it's just stealing*.

This was a rhetorically devastating move. It collapsed the AI companies' carefully constructed framing — *transformative use, fair use, technical necessity, innovation* — into the oldest accusation in commerce, and made it stick. *Stealing*. Not because the signatories did not understand the technical arguments. They did. Because they had decided that the technical arguments were a *cover* for an underlying transfer of value that didn't deserve any other name.

Once you have that framing, the whole policy debate looks different. *Should we allow innovation?* becomes *should we allow theft?* The answers are not the same.

The framing war did not end with the declaration, though, and I want to record where the other side had moved to by midsummer 2026, because in those weeks the platform counter-position consolidated into the form I expect it to hold for years. The consolidation was led by **Google**, and it had two prongs.

The first prong is *consent by terms of service*. In June 2026, Google moved to dismiss a class action brought by artists over YouTube-sourced AI training data, arguing that the artists had consented “to the use of their recordings” when they uploaded them to YouTube.²¹⁰ The claim is that the click-wrap agreement every uploader accepts — the licence that lets YouTube transcode, distribute and recommend a video — also covers ingestion into generative models that did not exist when most of that catalogue was uploaded. Independent-music advocates responded with the obvious objection, in the phrase that became the story's headline: this is *‘not informed consent’*.²¹¹ A musician uploading a video in 2015 was consenting to distribution. Nobody reads a streaming platform's terms of service as a training

licence for a future song-generating competitor — and if the argument holds, then every platform's most valuable AI asset turns out to be its archive of unread agreements.

The second prong is *outputs, not inputs*. In the same month, Google doubled down on its position that AI training is fair use, arguing that copyright enforcement should focus on what models *output*, not what they *ingest*.²¹² I want to be precise about what this framing does, because it is the exact structural inverse of the 88%. Licensing-in-all-cases puts the legal event at *ingestion* — the moment the work enters the training set, which is knowable, auditable, and licensable in bulk, and which the GEMA ruling treated as the operative act. Outputs-not-inputs moves the legal event to *generation* — millions of individual outputs, each requiring a case-by-case similarity analysis, each expensive to litigate, each disputable. It converts one collectively-negotiable question into an unbounded number of individually-unwinnable ones. That is not a compromise position. It is the training-pipeline status quo, restated as a principle of enforcement priority.

The sharpest rebuttal to the outputs-not-inputs move came, that same midsummer, from inside the industry rather than the courtroom. **Björn Ulvaeus** of ABBA — no one's idea of an AI maximalist — told a UN summit that the music business might be doing its AI licensing deals wrong, and that “*tracing the output was always the wrong question*.”²¹³ He meant it as a licensing critique, not a legal one, but it lands on Google's framing like a hammer. If even a rights-holder helping to build the licensing table agrees that chasing individual outputs is a category error, then outputs-not-inputs is not a principle of enforcement at all — it is a way of making enforcement impossible, and everyone serious about getting paid knows the operative event is ingestion.

The levy precedent: why Petrillo matters here

I want to take a long detour, because the 88% — and the institutional response forming around it — is, on the historical reading I laid out in [Chapter 2](#), a *Petrillo-template* moment that the trade press has, in the main, declined to recognise as such.

Let me state the template again, in its cleanest form, because the rest of this section relies on it.

When James Caesar Petrillo, the president of the American Federation of Musicians, took on the recording industry in 1942 and again in 1948 — staging the recording bans that effectively shut down the entire commercial output of American recorded music for the better part of three years — the strategic move was not, in essence, *prohibition*. Petrillo was not trying to ban records. He was trying to *tax* records. The 1942 settlement created a per-record royalty paid into an AFM unemployed-

musicians fund. The 1948 settlement, after the Taft-Hartley Act outlawed the 1942 structure, created the **Music Performance Trust Fund** under Section 302 — a *jointly-administered* labour-management fund, paid into by the labels and broadcasters, used to subsidise free live music performances by working musicians, distributing the productivity gain of the new recording technology to the displaced labour pool. The MPTF still exists. It still distributes payments today. It is, on a hundred years of evidence, the *only* form of institutional response to a creative-technology displacement that has worked at structural scale.

The four parts of the template, again:

One, the displacing technology is not banned. It is allowed to displace.

Two, the platform owner pays an *ongoing per-unit tribute* to the displaced labour pool.

Three, the tribute is collected *centrally*, by a joint labour-management body, not negotiated individual-by-individual.

Four, the tribute is paid out to subsidise the *displaced creative practice itself* — live music, in Petrillo's case — keeping it alive as a category even as the market for it shrinks.

I want to show how the 88% — and the architecture of institutional response coalescing around it in spring 2026 — is, function by function, a *reconstruction* of the Petrillo template for the AI era.

One, none of the institutional responses I have catalogued in this chapter — the UK consultation's licensing-by-default proposal, the SAG-AFTRA Tilly Tax, the *Stealing Our Work Is Not Innovation* declaration, the GEMA ruling, the Cannes Disclosure Standard — is, in essence, a *ban*. The declaration's signatories are not asking for AI to be prohibited. The 88% of UK respondents who wanted licensing-in-all-cases were not asking for AI training to be banned. They were asking for it to be *licensed* — which is, by definition, an acknowledgement that the underlying activity will continue. This matches Petrillo's first principle.

Two, what the 88%, the GEMA ruling and the *UMG v. Anthropic* settlement framework are collectively asking for is a *per-output tribute* from the AI platforms to the creative-labour pool whose work was used in training. The mechanism is, structurally, identical to the per-record royalty that Decca and Columbia agreed to pay AFM in 1944. The platform pays. The labour pool receives. The amount is calibrated to the volume of platform output. The mechanism is the Petrillo mechanism.

Three, the structural innovation of the Petrillo settlement — collection through a *joint body* rather than through individual-creator negotiation — is, in spring 2026, only partially built for the AI era. The collective-licensing infrastructure for music

(PRS, GEMA, ASCAP, BMI, SIAE, JASRAC and the related international bodies) has, in some cases, started negotiating directly with the AI platforms on the per-output structure. The Musical AI *creative weight attribution* infrastructure is a first attempt to build a *technical* layer underneath the joint-body political layer. The Cannes Disclosure Standard is an industry-coordination mechanism for the production-side disclosure that the collection mechanism rests on. None of this is finished. The joint bodies for *visual* artists, *writers*, *games developers*, *photographers* are at much earlier stages of development. The MPTF-equivalent fund-and-distribution mechanism does not yet exist for most of the creative industries. *Building it is the institutional work of the next eighteen months.*

Four, the final part of the template — paying the tribute out to subsidise the *displaced practice* — is the part the AI debate has, in my view, most under-thought. What does “subsidising the displaced practice” look like for AI-displaced creative work? For working musicians whose tracks are being competed-against by Suno outputs, it could look like funded performance opportunities, funded studio time, funded creative-development grants — the direct lineage of MPTF live-performance subsidies. For working illustrators whose work was used to train image models, it could look like commissioned-work grants, funded artist residencies, public-art-commission expansion. For working authors whose books were used to train LLMs, it could look like Public Lending Right expansion, library-licensing funds, writer-in-residence programmes. The structural move is the same in each case: *take the productivity gain from the platform, redistribute it to the displaced practice, keep the practice alive as a category.* This is what the 88% is implicitly asking for, whether or not the consultation respondents would have phrased it that way.

And then, in the first week of June 2026, the template stopped being my historical analogy and filed a lawsuit.

The **American Federation of Musicians** — Petrillo’s union; the institution whose 1942 and 1948 settlements this entire section is built on — sued **Universal Music Group** and **Warner Music Group** over the majors’ AI deals, claiming breach of the ‘**new use**’ clause in the record industry’s collective agreement with performers.²¹⁴ The ‘new use’ provision is, in essence, the contractual descendant of everything Petrillo won: the clause under which session musicians have historically been owed again when their recordings migrated into films, adverts and formats that did not exist when the original session fee was paid. The AFM’s argument is that licensing catalogue to AI companies for training is precisely such a *new use* of its members’ performances — and that the labels, in signing their AI deals, monetised the members’ work through a use the members were never paid for and never agreed to.

I do not know, as I write, how the case will resolve, and I am wary of prediction. But I want to insist on what the filing *is*, structurally, because the trade coverage treated it as one more item in a crowded litigation docket, and it is not. Every other suit in this chapter — GEMA, *UMG v. Anthropic*, Wixen, the Suno coalition — is rightsholders suing *platforms*. This is the labour pool suing the *rightsholders*: the union that built the only working levy mechanism in the history of creative-technology displacement, litigating to make sure the AI-era version of the tribute flows through the joint labour–management architecture rather than stopping at the label tier. It is Petrillo’s third principle — *collection through a joint body, not negotiated over the labour pool’s head* — asserted in court, against the very companies now presenting themselves as the creators’ champions in the platform fight. The question I raise at the end of the Kyncl discussion later in this chapter — whether the working session player actually *sees* the tribute — is exactly the question the AFM has now put in front of a judge. It is no longer hanging. It has a docket number.

I want to be honest about a complication that the Petrillo template hits at full speed in the AI era, because the book should not be glib about it.

The MPTF works partly because the relationship between *recorded music* (the displacing technology) and *live music* (the displaced practice) is *one-to-one*. The same musicians could, in 1948, do either thing. The Petrillo settlement was, structurally, paying the displaced version of the labour to subsidise the alternative version of the same labour.

The AI version of this relationship is *many-to-many*. The training data for a generative-image model is the lifetime output of *thousands* of working illustrators, photographers and visual artists, each of whom contributed an individually-tiny fraction of the model’s competence. The output is *generated* — there is no clean per-image-licence-equivalent. The redistribution problem is, by structure, much harder than Petrillo’s problem.

Two attempts to solve this are visible in 2026.

*The first is **creative weight attribution*** — Musical AI’s framing, picked up by some of the C2PA-adjacent technical standards groups — which proposes that AI platforms compute, for each output, the *gradient-weighted contribution* of each training-data input, and distribute a per-output royalty proportionally. The technical infrastructure for this is, in mid-2026, *partially* built. The economic infrastructure to handle the resulting micro-payments is, in mid-2026, *not* built. But the mechanism is the right one in principle: it preserves the one-to-many relationship that Petrillo could not directly handle, and translates it into a *many-to-many* redistribution mechanism.

The second is **collective licensing at the publisher tier**. The Stability AI / Universal Music alliance, the Splice / UMG partnership, the various YouTube and Spotify catalogue-licensing deals operate by aggregating training-data permissions at the publisher and label level, with the per-creator distribution handled internally by the existing royalty infrastructure of those publishers. This works for *commercially-published* creative work where the publisher already has a contractual relationship with the creator. It works less well for *independent* and *self-published* creative work where there is no publisher to negotiate on the creator's behalf.

Both approaches will, in some hybrid form, be the architecture of the AI-era Petrillo settlement. The 88%, the GEMA ruling, the SAG-AFTRA bargaining, the Cannes Disclosure Standard and the *UMG v. Anthropic* litigation are the political pressure that is forcing the platforms to agree to *some* version of one or the other. The version that emerges over the next eighteen months will, on the historical pattern, define the next forty years of how creative-AI work is paid for.

If the working-creative cohort reading this is asking what specifically to push for, my answer is: *the Petrillo template, applied to AI, collected through a joint body, distributed through a creative-weight-attribution mechanism layered on top of the existing collective-licensing infrastructure, used to subsidise the displaced creative practice as a category*. That sentence is a mouthful. It is also the most-likely-to-work structural answer that the historical pattern points at. The 88% is the political mandate for it. The institutional architecture is, in mid-2026, half-built. Finishing it is the work.

The consent-trained alternative

The 88% is a *demand-side* fact: it tells you what creators want done about the training pipeline. There is a *supply-side* fact that I think the policy debate has been slow to absorb, and that working creatives reading this book should know about, because it is the practical refutation of the AI companies' core argument.

The AI companies, as I noted earlier in this chapter, have spent two years arguing that machine-learning models *cannot practically* be trained on licensed data at the scale they require. That the data volumes are too large, the licensing relationships too fragmented, the legal cost too high. That training on consent-acquired data is, in effect, a nice idea that does not survive contact with the engineering.

By spring 2026, this argument was falsifiable, and had been falsified, by the existence of a category of foundation models that had been built — and were commercially successful — on exactly the consent-first basis the AI companies said was impossible.

The category, with the models I would name as its leading examples:

- **Adobe Firefly** (Image Model 5, Firefly Video, Firefly Foundry) — trained on Adobe Stock licensed content, openly licensed material, and public-domain work. Adobe Stock contributors are compensated through the Firefly bonus programme. By April 2025, Firefly had generated 22 billion assets; by mid-2026 it was the dominant enterprise generative-image tool, with 45%+ Creative Cloud penetration.²¹⁵
- **Bria** — image and video foundation models trained on a 100%-licensed dataset assembled from partners including Getty Images, Alamy and a network of independent rights holders. Bria publishes a per-output attribution and royalty mechanism on which the broader creative-weight-attribution conversation has drawn.²¹⁶
- **Getty Images Generative AI** (built with NVIDIA Picasso) — trained exclusively on Getty’s licensed library, with contributor royalties paid through Getty’s existing rights infrastructure.²¹⁷
- **Moonvalley Marey** — a generative-video model positioned explicitly as the “clean” alternative to Sora and Veo, trained on a licensed video dataset.²¹⁸
- **AIODE** — a music-creation DAW trained on ethically-licensed audio, listed in the toolchain map in Chapter 17.²¹⁹
- **Tamber** — an ethically-trained AI music suite launched in May 2026, controllable via arm gestures, positioned explicitly as a tool that moves *with* the musician rather than substituting for them.²²⁰
- The **Stability AI / Universal Music Group** alliance, the **Stability AI / Warner Music** deal and the **Splice / Universal Music** partnership — data-alliance arrangements that build music models on rights-cleared catalogues, with the per-creator distribution handled through the labels’ existing royalty infrastructure.²²¹

I am not claiming this category is perfect, or even, in every case, that its consent claims fully hold up to scrutiny. Adobe Firefly has faced criticism over the inclusion of AI-generated stock images in its training set;²²² the per-creator economics on the Stability / UMG-style deals are still being worked out. The point is not that these models are above critique. The point is that they *exist*, that they *work commercially*, and that their existence collapses the central technical-inevitability argument that the rest of the industry has been using to justify scraping.

The three Cs: consent, credit, compensation

By the late spring of 2026 the licensing-side conversation had begun to settle on a three-word shorthand worth pulling out of the contract architecture and putting on the page, because it is the cleanest articulation I have seen of what the *operative* version of the 88%-mandate looks like inside an actual platform-and-label deal. The

shorthand is **consent, credit, compensation** — the three Cs — and it was the explicit framing for the *Spotify–Universal Music Group* licensing agreements announced at Spotify’s Investor Day on **21 May 2026**, which let fans create AI covers and remixes of participating artists’ songs as a paid Premium add-on, with revenue flowing back to the artists and songwriters whose work is reinterpreted.²²³

The three Cs map cleanly onto the Petrillo template I described earlier in this chapter. *Consent* is the opt-in catalogue — only artists who actively agree have their work made available for remixing, which is the *acknowledgement-that-the-underlying-activity-will-continue* move that Petrillo’s first principle requires. *Credit* is programmatic provenance — derivatives are linked to original metadata, and a *Verified by Spotify* badge steers attention and money toward sanctioned work, which is the *central-joint-body-collection* mechanism that Petrillo’s third principle requires (with the platform doing the joint-body work that, in the recorded-music case in 1948, the AFM did). *Compensation* is a new royalty stream layered on top of standard streaming, paid out to the artists and songwriters whose work is used — which is Petrillo’s second and fourth principles (per-output tribute, paid into the displaced creative-labour pool) at the level of a streaming-product line.

The deal’s prehistory is the part I want every working creative reading this to register. UMG spent 2024 calling Suno and Udio’s training data “*infringement on an almost unimaginable scale*”. It settled with Udio in October 2025 and agreed to build a licensed product instead. The Spotify deal is the mature form of that turn — litigation converted into a revenue line. The market read it as a direct shot at Suno; Spotify’s stock jumped on the day. **The lawsuits were never the destination. They were leverage to force a licensing table into existence.** That sentence is, in my view, the right read on the 88%, on GEMA, on UMG v. Anthropic, on the Wixen suit, on the *Stealing Our Work Is Not Innovation* declaration and on every other piece of legal-and-political pressure I have catalogued in this chapter. The pressure produced the licensing table. The table is now being built. The three Cs are the contract pattern the table is using.

And the table kept filling through June. The **National Music Publishers’ Association** unveiled AI licensing deals with **Udio** and **Klay**, built — per the reporting — on a **50/50 split between songs and recordings**.²²⁴ Three things about that deal deserve a working creative’s attention. First, the counterparty: Udio is the platform whose training UMG had called “*infringement on an almost unimaginable scale*” before settling in October 2025 — the litigation-to-licensing arc, completing on schedule. Second, the structure: this is not one label doing a bilateral deal; it is the publishers’ *collective* trade body negotiating on behalf of the songwriting side as a bloc — the joint-body collection mechanism of Petrillo’s third principle, operating at the publisher tier. Third, and most quietly radical, the split itself. In the streaming economy the recording side has historically taken several

times what the song side takes — an allocation songwriters have spent two decades contesting. Writing *parity* between songs and recordings into the AI licensing architecture at its founding, before the norms harden, is exactly the kind of structural move that is only available at the start of a format. If the 50/50 split holds as the pattern, songwriters will have won in the AI transition something they never won in the streaming one. And step back far enough and the deal is the 88%’s Option 1 — licensing in all cases — arriving not as statute but as market practice: not a ban, a rate.

A second 2026 deal — the **Hasbro Sixth Wall Behavioral Licensing** launch on **3 June 2026** — is the same architecture applied to characters rather than recordings, and is worth flagging here because it extends the *consent* leg in a direction the music settlement does not yet have to reach. Where UMG’s opt-in catalogue grants permission for the *song* to be remixed, Hasbro’s CharacterOS grants permission for the *behaviour* of a character — how Optimus Prime argues, how Megatron menaces — to be deployed inside third-party experiences, with the underlying voice performance compensated through a *talent-participation* structure to the original performers.²²⁵ On the three Cs, this is the most complete model in the field as of mid-2026. It is also where the sharpest critique sits, and the critique is worth marking. *Participation* is a compensation mechanism — it answers *credit* and *compensation*. It does not, by itself, settle *consent* in the deeper sense the actors’ unions mean: ongoing control over how a performance is used, and the right to refuse. *A pay cheque is not a veto*. The Sixth Wall model’s long-term legitimacy will rest on contract terms the press releases don’t yet disclose, and the SAG-AFTRA-shaped argument about NIL rights is going to continue to be the live argument underneath every Behavioural-Licensing deal that follows.

The deeper analytical treatment of the three Cs runs across five emerging licensing models — Hasbro/Behavioural Licensing, Spotify/UMG embedded remixing, Disney/OpenAI as cautionary tale (the architectural framework that survived the platform’s collapse), the independent-artist commons-and-DAO models (Grimes / Elf.tech, Holly Herndon / Holly+, ElevenLabs heritage voices, and the Meta likeness-rental anti-pattern), and the Suno/Alexa *create-to-consume* demand floor.²²⁶

Indemnity as the receipt

The clearest single signal that a model has done its upstream consent work is whether the company behind it is willing to *indemnify* its customers against copyright infringement claims arising from generated output.

The pattern, in the eighteen months to mid-2026:

- **Adobe** offers IP indemnification on Firefly enterprise outputs — Adobe will defend customers, and pay damages, if an enterprise generation is challenged on copyright grounds.²²⁷
- **Microsoft** announced its **Copilot Copyright Commitment** in September 2023 and extended it across the Copilot product family through 2024–25. Microsoft assumes the legal liability for commercial customers using Copilot outputs in good faith.²²⁸
- **Google** offers an indemnification commitment on Vertex AI and Workspace generative outputs.²²⁹
- **IBM** offers an uncapped indemnity on watsonx-generated content for enterprise customers.²³⁰

Notice which companies are on this list and which are not. The companies indemnifying their customers are, without exception, the companies that have invested most heavily in the *upstream* consent work — licensed data, contributor compensation, rights-cleared catalogues. The companies that have *not* indemnified their customers are, predominantly, the companies whose training-data position is most exposed.

This is not a coincidence. Indemnification is a receipt. It is the legal department of a \$200B company telling its commercial customers, in the most expensive language available, *we have done the work; you can use this without being sued*. The absence of an indemnity, conversely, is an instruction. It is the same legal department saying, *the risk is yours; you carry it*.

For working creatives, agencies and studios making procurement decisions in 2026, the indemnity status of a tool is the single most useful one-question proxy for whether its training pipeline is built on the side of the 88% or against it. Ask the vendor. If they cannot give you a written indemnity, you have your answer.

The literacy problem

None of this — the consent-trained category, the indemnity framework, the C2PA provenance stack in [Chapter 13](#), the legislative reversal earlier in this chapter — works without a corresponding investment in *literacy*. And the literacy gap, in mid-2026, is the place where I am most worried about the architecture failing.

Policy and infrastructure can constrain the supply side. They cannot, on their own, redirect the demand side. The question of whether a working illustrator chooses Firefly over Midjourney, whether a marketing team specifies Bria over a scraped open-source model in its agency brief, whether a record label’s A&R department uses a Stability / UMG-aligned tool rather than Suno for demo work, whether a film commissioner asks for Marey provenance on a generative-video shot, whether an audience member streams a SynthID-watermarked track over an unlabelled one —

these are *consumption* questions. They sit downstream of every law and every standard. They are decided, ten thousand times a day, by people choosing tools and content from a menu, without anyone telling them what the choices on the menu actually mean.

Three things have to happen for the literacy layer to catch up with the infrastructure layer.

First, working creatives need to know what the consent-trained category is, which tools are in it, and what an indemnity is for. This book is one attempt at that; the **Sundance AI Literacy Initiative**, training 100,000+ artists in provenance practice on Google's funding, is another.²³¹ The professional bodies — the Society of Authors, the AIGA, Equity, the AOP, the MPG, the WGA — have, by mid-2026, started shipping member guides. The work is early.

Second, the *buyers* of creative work — the brands, the agencies, the broadcasters, the platforms, the publishers — need to make ethically-trained models a *specified requirement* in their briefs and procurement contracts. A handful already have: the BBC, the AP wire service, the Cannes festival itself. The vast majority have not. The lever exists. It needs to be pulled.

Third, the *audience* needs the equivalent of a nutrition label. The Cannes Disclosure Standard, SynthID-in-Gemini, the YouTube AI-content disclosure rules, the proposed EU AI Act labelling obligations are early attempts at this. None of them yet add up to a consumer-facing signal as legible as the *Fairtrade* mark or the *organic* certification. Until they do — until an audience member streaming a song, watching a clip, or buying a print can tell at a glance whether the creative work in front of them was made with a tool that paid the people whose work it learned from — the consumption side of the equation will keep leaking. Provenance metadata sitting in a file header that no one reads is not, on its own, literacy.

I do not think this layer will get built by the AI companies. The incentive isn't there. I think it will get built — slowly, contentiously, in fits and starts — by the same coalition I am about to describe in the next section: by creators, their unions, their professional bodies, their buyers and their audiences, jointly insisting on a labelling regime that the platforms eventually have to honour because the market has organised itself around it.

The 88% is the political mandate. The consent-trained models are the proof of supply. The indemnity framework is the legal receipt. The literacy infrastructure is the missing piece — and it is, on the evidence of the last six months, being built.

Coalitions, not protests

The thing I want creative people reading this to take from this chapter is *not* that protest works.

Protest works. We have seen it. The 88%, the Equity ballot, the SAG-AFTRA contract, the artists' declaration, the GEMA ruling, the U.K. government's reversal — these are evidence that protest works.

What I want you to take is that *coalition* works.

What happened in these six months was not — or was not only — that individual creators got angry and shouted. What happened was that creators *aligned themselves* with adjacent groups whose interests they had not previously seen as aligned with theirs.

Working musicians aligned with photographers, who aligned with authors, who aligned with games developers, who aligned with screenwriters, who aligned with voice actors, who aligned with concept artists, who aligned with translators, who aligned with journalists. They aligned with their unions. They aligned with their professional bodies. They aligned, somewhat to everyone's surprise, with the major studios, who had spent twenty years suing them and now found themselves on the same side of a copyright argument against the same platform companies.²³²

The major-label leadership read of the same coalition shifted, in the spring of 2026, in a way that I think is worth registering carefully because the rhetoric is a useful weather-vane. **Robert Kyncl**, the chief executive of **Warner Music Group**, in a widely-quoted May 2026 interview, told the industry that “*AI resistance*” was actively *setting the music sector back* — that AI represented “an incredible value creation opportunity,” and that the labels “*cannot wait the way the industry did 25 years ago*.” Kyncl's invocation of the Napster moment was deliberate. The argument was that the labels' twenty-five-year pattern of *suing first, integrating second* had cost them, in net, the bulk of the streaming-era surplus to platforms that had moved before they did, and that repeating that pattern with AI would compound the loss.²³³ This is a meaningful change in register from the 2025 “*biggest theft in music history*” framing, and worth tracking. It does not contradict the 88% — Kyncl, like the BPI, is pushing for licensing infrastructure rather than against it — but it shifts the *centre of gravity* of major-label rhetoric from prohibition toward participation. On the Petrillo template, this is the labels' tribute-mechanism position: AI continues, the platforms pay, the joint-body collection infrastructure scales. The question of whether the *displaced practice* gets a meaningful share of the resulting flow — whether the working songwriter, the working session player, the working independent artist actually *sees* the tribute — is the part the Kyncl framing does not yet answer.

The artists answered it themselves within weeks. In late June 2026, an open letter warned that musicians “*must not be pressured into AI deals*” by the rightsholders who control their catalogues.²³⁴ Read alongside the Kyncl interview, the NMPA deals and the AFM's ‘new use’ lawsuit earlier in this chapter, the letter marks out the

precise fault-line inside the coalition I have been describing. Artists and labels stood on the same side against unlicensed scraping — that alignment held, and it won the procedural battles this chapter has catalogued. But a licensing boom conducted at the catalogue level, over the heads of the artists whose recordings constitute the catalogue, fails the first C. Consent that is structural — buried in old contracts, exercised by an intermediary, impossible in practice to refuse — is the platform’s terms-of-service argument wearing a label’s lanyard. The open letter, the AFM filing and the artist trade body’s ‘scraps’ warning from November 2025 are the same message delivered at increasing volume: the coalition’s unity was built on *consent*, and it will fracture on consent too, if the tribute mechanism stops one tier above the people who did the work.

They also — and this part I find most interesting — aligned with their *audiences*. The Adobe Creators’ Toolkit Report found that **69%** of creators worried about their work being used to train AI without consent.²³⁵ That number rhymes with the 88% in the U.K. consultation. It also rhymes with the audience behaviour I described in Chapter 5 — the slop ceiling, the AI-music underperformance, the cultural rejection of synthetic content that doesn’t disclose itself. The creators wanted protection. The audience, given a choice, wanted to listen to the protected work. The two interests, for the first time in a long time, sat on the same side of the line.

That alignment is the most powerful political asset the creative industries have had this century. They built it in six months. The question for the next six months — which Chapter 14 of this book is going to come back to — is what they do with it.

The government admits it got it wrong

On **1 June 2026** — at **SXSW London**, a cultural and technology conference whose choice of venue was itself a statement about London’s post-Brexit attempt to compete with Austin for creative-economy credibility — UK Culture Secretary **Lisa Nandy** said, on a public panel, the words that the 88% had been waiting fourteen months to hear from a minister.

*“I just want to be really clear about where this government sits on it, because we went out to try and find a way to help tech and creative industries to unlock the opportunities of AI, and we got it wrong in the first instance.”*²³⁶

It is worth holding that sentence in its full weight. *We got it wrong in the first instance*. This is a serving government minister, at a public conference, explicitly conceding that the policy position her department had put to Parliament — the opt-out model, Option 3, the preferred option that only 3% of respondents had supported — was wrong. Not “under review.” Not “subject to ongoing consultation.” *Wrong*. The verb is unambiguous and it was not walked back.

What Nandy described at SXSW was a government that had gone into the consultation trying to find a workable compromise between tech and creative industries, and had, by the weight of the 88% finding and fourteen months of subsequent pressure — the GEMA ruling, the SAG-AFTRA settlement, the Equity ballot, the Cannes Disclosure Standard, the parallel legislative movements in the United States and the European Union — been forced to reframe its understanding of what the question even was. The government had begun the process asking *how do we enable AI training?* and had been forced, by the volume and coherence of the creative sector's response, into asking *how do we protect the people whose work underpins it?*

The vindication is real. And I want to be precise, as I was precise about the March 2026 reversal earlier in this chapter, about what it does and does not mean.

It *does* mean that the 88% — and every piece of union, legal, institutional and cultural pressure that built up around it — was sufficient to change the stated position of the government that proposed the opt-out model. In policy terms, that is a significant outcome. Very few public consultations on complex technology questions produce a ministerial admission of error this explicit.

And the admission was followed, within a fortnight, by the sound of the post-consultation policy machinery actually moving. In the second week of June the government published its **AI Adoption Plan for the Creative Industries** — the sector-specific plan for how film, music, games and the wider creative economy are meant to *adopt* AI, not merely be protected from it.²³⁷ DSIT announced a **£1.1 billion AI hardware plan** to back British chip firms and expand domestic computing power and skills.²³⁸ And in the same week, the UK's competition regulator introduced what was widely reported as “world first” controls on Google, requiring it to let news publishers opt out of having their content scraped into Google's generative AI features.²³⁹

I read those three moves together, because together they say something the consultation documents never quite did. The adoption and hardware plans confirm that the SXSW admission was not a retreat from AI — the government wants the AI economy, wants the compute onshore, wants the creative industries using the tools. What it retreated from was *unlicensed* AI. And the regulator's Google controls are the first genuinely novel instrument in the post-consultation toolkit: an *opt-out* — the very mechanism the 88% rejected when it was offered as a blanket copyright exception — reappearing as an enforceable obligation imposed on a named dominant platform, the same Google whose consent-by-terms-of-service and outputs-not-inputs arguments I set out earlier in this chapter, with the compliance burden sitting on Google rather than on ten thousand individual rightsholders. An opt-out that a regulator polices against a specific gatekeeper is a different

instrument from an opt-out buried in a copyright exception: the first has teeth, the second had none, and the difference between them is a fair one-line summary of what fourteen months of pressure achieved. The controls cover news publishers, not musicians or illustrators — the machinery moves sector by sector — but the instrument now exists, and instruments that exist get extended.

It *does not* mean the fight is over. Nandy's statement at SXSW was accompanied by the careful qualifier: "no decisions on next steps." There is, as of the time of writing, no legislation in place, no licensing regime enacted, no opt-in framework passed into law. The legislative process is ongoing. The political will has shifted; the policy architecture that would make that shift concrete has not yet been built. The 88% produced a political admission. The institutional and legal work that converts that admission into durable protection for creators is still ahead. The Petrillo template I described earlier in this chapter — the joint-body collection mechanism, the per-output tribute, the subsidy for the displaced creative practice — remains the structural target. The political mandate for it has never been stronger. The architecture is still half-built.

Chapter 7 — The Studios Decide

If the audience was speaking through the slop ceiling, and the creators were speaking through the 88%, the studios were speaking through their balance sheets — and the language was not quite the language of either of the other two.

On **22 October 2025** — three weeks into the period this book covers, the same day I was writing Issue 4 — Ted Sarandos, Netflix’s co-CEO, told an industry conference that Netflix was “all in” on leveraging AI across its streaming platform.²⁴⁰ The phrase was casual. The implications were not. Within hours, the trade press was running it as the official line of the world’s largest streaming service, and it was being read — correctly — as a signal to every other studio that the period of “wait and see” was over.

Three weeks earlier, in a piece *Futurism* had published with a headline that aged badly almost in real time, **Lionsgate’s** ambitious attempt to use AI for movie development had been characterised as having “crumbled into disaster.”²⁴¹

Two months later, on **11 December 2025**, *The Guardian* reported that **Disney** was *announcing* a \$1 billion investment in OpenAI, with a structured agreement that would let Disney characters appear in the Sora video tool.²⁴² At the time of the announcement, the deal was treated by the press and most of the industry as the canonical template for how a major rights holder would integrate with a frontier AI lab. By the spring of 2026, on closer inspection, it was nothing of the kind. *No formal agreement was ever signed; no money changed hands.* On **24 March 2026**, OpenAI told staff it was shutting down Sora — a staged wind-down, with the consumer app and website ending in late April and the developer API set to close in September — and Disney, reportedly given as little as thirty minutes’ notice, exited the partnership.²⁴³ The billion-dollar template, as the *Ropes & Gray* analysis later put it, was a billion-dollar press release.

That collapse matters for the strategic map in two ways, both of which I want to flag now so the reader can read the rest of the chapter with the correction in view. *First*, the framework Disney’s lawyers built — IP-not-training, talent-likeness excluded, jointly-governed steering committee, equity-and-warrants structure — *survived the platform’s death* and is now the reference contract architecture against which other studios negotiate. The framework was right; the vendor was fragile. *Second*, the lesson the rest of the industry took from the collapse was that an *equity stake in a platform you don’t govern* is not a strategic position. The thirty-minutes’-notice detail became a case study in strategic-investment risk: a \$1bn pledge with no board seat or information rights strong enough to give the rights holder warning that the product it was licensing into was about to be killed. The takeaway for every rights holder that follows has been straightforward — *align on equity and governance rights, not just licence terms, or build the platform yourself.* Which is exactly what Hasbro did with **Sixth Wall** and the **CharacterOS** governance layer on 3 June

2026, a move that I treat in detail under Position Five below. The Disney–OpenAI material throughout this chapter should be read with this correction in mind. The full deal-architecture analysis — the *three Cs* of consent, credit and compensation, the comparison against Hasbro, Spotify/UMG, the independent-artist models, and the Suno/Alexa demand floor — sits in the *Dream Machine* research this chapter rests on.²⁴⁴

These three moments — Lionsgate’s failure, Netflix’s commitment, the Disney–OpenAI announcement-and-collapse — are the three corners of the strategic map that every legacy studio in the world has been navigating for the last six months. They are not a single story. They are three different stories about how a creative business with a hundred years of human-craft DNA tries to integrate a technology that, by the time it integrates, no longer behaves like the technology you thought you were integrating.

This chapter is about the studios. About how they decided. About the ones that went *all-in*, the ones that went *AI-native from scratch*, the ones that went *we are not doing this at all*, and the ones — the most interesting group — that went *we will do it, but only in the places where it doesn’t show up in the work the audience sees*. And, late in the period, a fifth group emerged: the ones that decided to stop licensing the future and *own the rails it runs on*.

The map of those five positions, drawn carefully, is the map of where the film, TV, games and entertainment industries will be in 2030.

Position One: All-in

Netflix’s “all in” framing was the most prominent example of what became, over the autumn of 2025 and the winter of 2026, the dominant public stance of the major streamers. The framing was: AI is a tool, AI is a productivity multiplier, AI is going to be used everywhere in the pipeline, and the studios that adopt it earliest will have the most leverage when the new economics settle.

The actual deployment, when the trade press dug into it, was more interesting than the slogan suggested. Netflix’s use of AI in late 2025 included generative AI tools for visual effects (de-aging actors, scene extensions, background plates), AI-driven recommendation engines that the company had been refining for fifteen years, and — disclosed in a January 2026 *Pymnts* report — a major AI strategic push focused on subscriber *retention* rather than production cost.²⁴⁵ The story Sarandos was telling Wall Street was not “AI will replace our writers.” It was “AI will keep our subscribers engaged in a way that human-only programming alone cannot afford to.”

By May 2026 the deployment had moved one further step in. Netflix announced **INKubator**, an in-house AI animation studio explicitly chartered to produce

“feature-quality” short-form work, and began recruiting for it publicly.²⁴⁶ What is notable about INKubator, for the purposes of this chapter, is not the size of the unit — small, by Netflix standards — but the *organisational position* of it: an *internal* AI-native studio sitting *inside* the major-streamer architecture, producing original work, reporting up into the same commissioning structure as the live-action slates. That is a different shape from the Position-Two AI-native studios I describe in the next section. It is a hybrid: Position One on the org chart, Position Two in the pipeline.

And in July 2026 the all-in position finally acquired a number. On Netflix’s Q2 earnings call, the company revealed that **about 300 titles** this year had used generative AI somewhere in the pipeline.²⁴⁷ Three hundred titles is not a pilot programme, and it is not a research unit; it is the slate. What had been a slogan in October 2025 — “all in” — was, nine months later, a disclosed operating statistic, delivered to shareholders in the same breath as the viewing-hours numbers. I would flag the venue as much as the figure: when AI usage moves from the technology press to the earnings call, it has stopped being a strategy and started being an input, reported the way a studio reports its content spend.

Adjacent moves from other big studios that autumn told the same story.

Amazon built out an internal “AI Studios” unit in November 2025, naming sports-docs boss Matt Newman as its head of live-action production.²⁴⁸ In the same month, Amazon’s *House of David* TV series became one of the first major Western dramas to publicly disclose the use of more than **350 AI-generated visual-effects shots** in its second season, with creator Jon Erwin telling *Wired* he was “not sorry.”²⁴⁹

By late May 2026, Amazon went further still: the company launched an **AI Creators Fund** and commissioned three animated Prime Video series using generative AI. One of the three — **Punky Duck**, created by *El Tigre* animator **Jorge Gutierrez** — became the most dramatic commissioning collapse of the period. Within **48 hours** of the announcement being confirmed, Gutierrez had dropped out, citing “*relentless derision*” online. The series was dead before it had been made. For the full treatment of the audience-enforcement dynamics underlying this collapse, see the discussion in [Chapter 5 §“The pushback.”](#) For the purposes of this chapter’s strategic map, the Punky Duck case is a Position One data point of a specific kind: a major streamer commissioned AI animation through a structural fund and discovered, at the announcement stage, that the audience-enforcement mechanism can operate on a press release. The question that Amazon’s AI Creators Fund has to answer — and that the remaining two commissioned series have to answer — is whether the specific cultural objection was to *Jorge Gutierrez’s participation* specifically (which del Toro’s subtweet suggested) or to AI-animated commissions from a streaming fund generically.²⁵⁰ The games division, meanwhile,

pressed on regardless: Amazon's new games plan, as *The Verge* reported it in June 2026, paired the James Bond licence with an AI Snoop Dogg — licensed celebrity likeness, synthesised, deployed as first-party content strategy.²⁵¹

NBCUniversal signed a deal in late October 2025 with the son of *Law & Order* creator Dick Wolf to develop AI-generated games based on its IP.²⁵² By late November, the framing had broadened — *The Office*, *Saturday Night Live* and *Sex and the City* were all reportedly being considered as IP for AI-generated game adaptations.²⁵³

Disney, beyond its OpenAI investment, announced in November 2025 that it was developing generative AI tools to let Disney+ subscribers create and share their own short-form videos using the company's iconic IP — a play, transparently, at recapturing the engagement Fortnite and Roblox had been taking from passive streaming for years.²⁵⁴ To execute this, Disney created a new “**Office of Technology Enablement**” under former Walt Disney Studios CTO Jamie Voris, with the specific mandate of accelerating AI and Mixed Reality adoption across the organisation.²⁵⁵ In January 2026, Disney followed up with an announcement of a TikTok-like vertical-video product and an AI video-generation tool aimed at brand advertisers using existing Disney brand assets and guidelines.²⁵⁶

Fox Entertainment took an equity stake in **Holywater**, an AI-microdramas company, in October 2025.²⁵⁷

Sky History acquired *Castles SOS*, an AI-powered documentary, in late November.²⁵⁸

Channel 4 rolled out an AI-driven advertising tool in December 2025 designed to make TV advertising accessible to SMEs — a small home-builder was one of the first clients.²⁵⁹

The Punky Duck collapse left Position One with an unanswered question: was the objection to one artist's participation, or to the alliance itself? **A24** spent late June running something close to the control experiment. In June 2026 the *Hollywood Reporter* reported that **Google** was investing in the prestige-indie studio, in an AI venture linked to Google DeepMind and reported around A24's *Backrooms* project.²⁶⁰ A24 is not a legacy tanker; it is the most carefully curated brand in independent film — the name the AI-native studios invoke when they describe their own ambitions. Wonder Studios' stated goal, remember, is to become “*the A24 of AI production*.” If any studio's cultural authority could absorb an AI alliance, it should have been A24's. Within a fortnight, *Futurism* was reporting that the studio was dealing with a “*full-on crisis*” among its fans over the Google deal.²⁶¹ The audience-enforcement mechanism, in other words, does not check the brand's credentials

before it fires. If anything the reverse: the fan bases most likely to enforce are precisely the ones a prestige brand has spent a decade training to care about authorship. A24's bet — and it is a live bet as I write this — is that curation authority is durable enough to survive the alliance that funds it. The Gutierrez precedent says the next few months of that experiment will matter more than the press release did.

It is worth being clear about why the studios keep walking into this mechanism, because it is not naivety. It is arithmetic. At the end of June, *Bloomberg* reported what it called animation's AI reckoning: filmmakers saying they can now make animated movies for **90% less**.²⁶² Whatever discount you apply to that figure for enthusiasm — and I would apply a healthy one — a cost compression of that order is not a tooling decision. It is a restructuring of the industry's cost base, and it explains the otherwise puzzling persistence of the commissioning experiments: every studio that has done the maths knows that if the 90% number is even half right, someone is going to build an animation slate at that cost basis, and the only open question is whether it is them.

Nor is Position One a purely Hollywood posture. In June 2026, CNN reported that South Korea's film industry — the industry that gave the last decade *Parasite* and the prestige-TV export wave — was embracing AI at industrial speed, with the attendant question, *will K-film lose its magic?*, being asked mostly by observers rather than by the industry itself.²⁶³ The Shanghai International Film Festival programmed AI and iPhone moviemaking as the twin faces of cinema's future,²⁶⁴ and the Chinese model company **MiniMax** chose the festival to launch an all-in-one AI video generation hub — with the launch framing that “*creative direction and aesthetic judgment must ultimately be left to humans*.”²⁶⁵ That last detail is the one I would flag. The East Asian industries are not having the argument Hollywood is having; they are arriving with the adoption already decided and with Position Four language — tool in the workflow, judgment with the humans — pre-installed in the marketing. Whether the language governs the practice is a separate question, but as strategic positioning it is a full cycle ahead of the American discourse.

Position One is not subtle. The streamers, broadcasters and major studios with the capital to do it have been integrating AI into their stacks — production, post-production, marketing, advertising, distribution, subscriber retention — at a pace that suggests they have already decided which side of the future they want to be on. They want to be the side that owns the toolchain.

Position Two: AI-native

A second group, more recent and more interesting, are the studios that have decided not to *integrate* AI into existing film and television production pipelines but to

replace those pipelines entirely with AI-first workflows. These are the **AI-native studios**.

Fremantle, the international production powerhouse, named the boss of its new “AI-native” studio **Imaginae Studios** in October 2025.²⁶⁶ By the spring of 2026, Imaginae was developing a project called *Art Awakens*, fusing AI techniques with classical painting IP.²⁶⁷

Imagine Entertainment — Ron Howard and Brian Grazer’s production company, with one of the most distinguished filmographies in modern Hollywood — partnered with a new AI-first production company called **Obsidian Studio** in November 2025.²⁶⁸

Wonder Studios raised \$12m in seed in October 2025²⁶⁹ and by January 2026 was running its own **Wonder Film Festival** with a curated shortlist of AI-made shorts.²⁷⁰ By May 2026 Wonder had closed a further round bringing total funding to **\$50M**, with the company publicly framing the ambition as becoming “*the A24 of AI production*” — a deliberate analogue to the indie-prestige distribution model rather than to the streamer-replacement model the trade press had been expecting AI-native studios to chase.²⁷¹

Asteria — Natasha Lyonne’s AI company, backed by James Cameron’s *Lightstorm Entertainment* — produced its first animated short, *All Heart*, in October 2025.²⁷²

Promise, a deep-pocketed AI studio backed by Google, set up shop in October 2025 specifically to “bring GenAI filmmaking and VFX to legacy media.”²⁷³

Goldfinch launched **enGEN3**, an “AI-Powered Cinematic Universe Platform,” in October 2025.²⁷⁴

Chapter41, a Munich-based AI startup, was launched in November 2025 by **Beta Film** and a group of industry executives.²⁷⁵

Kartel — a new AI startup led by long-time TV exec Kevin Reilly, formerly of HBO — was set up in November 2025.²⁷⁶

Wanted director Timur Bekmambetov launched a \$5 million project to “generate AI method actors” in November 2025, with the framing: “AI is here to stay. We have to train it responsibly.”²⁷⁷

Particle6, the U.K.–Netherlands company behind Tilly Norwood, expanded to 41 AI actors in development by November 2025, with founder Eline Van der Velden in a December 2025 *Deadline* interview making the case that AI performance was a “more ethical way” to act — and urging working performers to “future-proof” themselves by creating their own AI avatars.²⁷⁸

Wonder Studios, separately, adapted a children's book to an animated series using AI in December 2025.²⁷⁹

Kling AI and **Evolutionary Films** announced an AI-animated feature, *Minibots*, at the Cannes Film Market in May 2026, alongside a broader Kling-backed filmmaker initiative aimed at funding AI-native productions on the same indie-distribution architecture.²⁸⁰

By June 2026 the maturing of the category was visible in the shape of a single company's slate. **Utopai Studios** signed on to co-produce *Half Moon*, a feature from Silver Bear-winning filmmaker **Hyo-Joo Yang**²⁸¹ — an AI-native studio buying into precisely the festival-circuit authorship credibility the category has so far lacked — while simultaneously developing, with the Chinese production giant **Huace**, a fully AI-generated series adaptation of *Journey to the West*.²⁸² Auteur co-production on one track; industrial-scale generation of canonical public-domain IP on the other. That two-track structure — prestige upward, volume outward — is what an AI-native studio looks like when it stops being a demo reel and starts being a business.

On **10 June 2026**, the first fully AI-generated feature film accepted by a major international festival made its world premiere. *Dreams of Violets* — directed by **Ash Koosha**, produced by his company **Fountain o** — screened at the **Tribeca Film Festival** in New York. The film is a 75-minute docudrama about the 2026 Iran protests. It cost **\$2,000** to make and took **three months** in production. The tool stack: **Kling AI** for video generation, **Anthropic's Claude** for editing, **Google's Gemini** for research.²⁸³ What Tribeca programmed was not a short or an experiment or a proof of concept. It was a feature — a narrative of a real political event, made entirely without a traditional production infrastructure, at a cost lower than a mid-range laptop. The festival's decision to program it did not close the "is it art or is it slop?" debate. It opened it, formally, in the festival-circuit context where that debate carries weight. Critics split sharply: advocates for the work pointed to the film's subject matter and urgency; critics pointed to the same qualities they flag in every AI-native production I have documented in these six months — technical capability uncoupled from the *earned authority* that makes a film's claim on its audience stick. What is not in dispute is the milestone: *Dreams of Violets* is the first AI-generated feature to pass the selection committee of a Tier 1 film festival. Whatever comes after, it marks the moment the festival circuit acknowledged that the category exists.²⁸⁴ The critical verdict arrived within the fortnight. *Variety's* review out of Tribeca ran under a headline that was itself a question — "*What does a film made entirely with AI look like?*" — and that, on my read, is the honest state of the reception: the first fully AI-generated feature at a Tier 1 festival was reviewed seriously, on the film pages rather than the technology pages, and the review was

still working out what the category it belonged to actually was.²⁸⁵ Being taken seriously enough to be asked the question is the milestone. Answering it is the work the category has not yet done.

By **April 2026**, the trade press could no longer keep up with the AI-native studio launches. There were too many of them. Most of them, like most early-stage production companies in any era, will not survive the next two years. The question of whether a meaningful AI-native studio system will eventually emerge as a parallel structure to legacy Hollywood — the way Netflix and Amazon eventually emerged as a parallel structure to the cable networks — is, in my view, the biggest single open question in the film and TV industry as of May 2026.

The early evidence is mixed. *Watch the Skies*, a Swedish UFO feature entirely dubbed with AI, secured U.S. distribution in October 2025.²⁸⁶ *Run to the West*, South Korea's first AI feature film, was tested with critics and audiences in October 2025; one *cybernews.com* review described the experience as “testing the soul of cinema.”²⁸⁷ *Lily*, a Tunisian-made AI short, won the \$1 million Dubai AI Film Award in January 2026.²⁸⁸ *Humans in the Loop*, an AI drama that received Film Independent's Sloan Distribution Grant, entered the Oscar race in November 2025.²⁸⁹

I have watched a meaningful percentage of the AI-native output of these six months. The honest evaluation, which I have given in talks several times and stand by here, is that we have not yet seen the *Citizen Kane* of AI cinema — we have not yet seen a single AI-native work that I think will still be watched in 2040. We have seen, repeatedly, films that demonstrate technical ability without yet demonstrating cultural necessity.

The most interesting AI-native works, in my view, are the ones that — like Andrii Daniels' bomb-shelter clip — wear their non-traditional production conditions on their face. They are films *about* the technology being used to make them, in some implicit or explicit sense. They are not pretending to be legacy films made by a different route.

Position Three: We are not doing this

The studios that have publicly refused generative AI have been some of the most interesting voices in this entire period.

Pocketpair, the Japanese games studio behind *Palworld*, announced in October 2025 that its new publishing division would not handle games using generative AI. The CEO's full statement, in *PC Gamer*, was sharper than the headline: “*We don't believe in it. We're very upfront about it. If you're big on AI stuff or your game is Web3 or uses NFTs, there are lots of publishers out there [who'll talk to you], but*

*we're not the right partner for that.*²⁹⁰ It was, on its face, a rejection of one production model. It was, on inspection, also a *marketing position*: a publisher staking its claim with audiences who had become — by late 2025 — actively allergic to AI-augmented games.

Larian Studios — the maker of *Baldur's Gate 3*, one of the most critically and commercially successful games of the decade — backed off generative AI in January 2026 for its next *Divinity* game.²⁹¹

Games Workshop, custodian of the *Warhammer 40,000* universe, ruled out generative AI entirely in early 2026.²⁹²

Hooded Horse, the U.S. games publisher behind *Manor Lords*, said in January 2026 that it would not work with developers who used generative AI.²⁹³

Jagex, the maker of *RuneScape*, declared in January 2026 that it would *never* use generative AI to make in-game content.²⁹⁴

Aardman Animations — the British animation studio responsible for *Wallace and Gromit* — announced in December 2025 that it would “embrace the technology” of AI but would be “very cautious not to lose our values.”²⁹⁵ This was, by Aardman’s careful standards, a sharp line: they reserved the right to use AI for narrowly defined post-production and admin tasks, but explicitly excluded it from the stop-motion craft that defines their work.

Guillermo del Toro, in October 2025, told *Variety* he would “rather die” than use generative AI in his films, with a follow-up *Frankenstein*-themed press cycle that made the line one of the most-quoted creative-industry statements of the year. The full quote was even better than the headline: “*I’m 61, and I hope to be able to remain uninterested in using it at all until I croak. ... The other day, somebody wrote me an email, said, ‘What is your stance on AI?’ And my answer was very short. I said, ‘I’d rather die.’*”²⁹⁶ What del Toro was doing, with the bluntness only a senior auteur with a fully-funded slate can afford, was *refusing to participate in the framing*. Most working creatives have had to spend two years giving careful, nuanced, defensive answers about their AI position. Del Toro decided he was a senior enough artist to refuse the question entirely. The cultural permission for that posture, in a particular kind of high-end filmmaking, is part of the architecture this book has been describing.

Not every senior auteur held the line del Toro drew. In June 2026, **Martin Scorsese** became an advisor to **Black Forest Labs** — the German company behind the FLUX image-generation model — and publicly endorsed using AI for storyboarding in the filmmaking process.²⁹⁷ The announcement produced an immediate and, on the cultural register, notable backlash. **Boots Riley** and

Guillermo del Toro — del Toro, again — were prominent among the filmmakers who responded critically, with Riley’s objections framing the endorsement in explicitly labour-terms: a filmmaker of Scorsese’s stature lending his name to a generative-AI company is not a neutral act in a landscape where working artists are still fighting for compensation frameworks. The Scorsese/Black Forest Labs story is the mirror image of the del Toro position: where del Toro’s “rather die” framing set a cultural permission for refusal, Scorsese’s endorsement set a cultural permission for adoption — and was immediately challenged for doing so. Both positions tell you something important about where the industry’s internal argument is in mid-2026: it is not between the *studios* and the *artists* any more, as it was in 2024. It is between the *artists themselves*, at the level of individual public positioning.

By June, that argument-between-artists had become a weekly genre of its own, and the range of positions on display is worth recording precisely because it does not sort neatly by generation or by prestige. **Ron Howard** — whose Imagine Entertainment had, remember, partnered with the AI-first Obsidian Studio back in November — told the *Hollywood Reporter* he was bullish on AI, and that Jim Henson, the great puppet-technologist, would have been too.²⁹⁸ **Gore Verbinski**, in the same weeks, proposed a rating system for AI use in films — “if you use AI to write a script, you get an F” — which is not a refusal of the tools so much as a demand for disclosure with consequences attached: a classification-board analogue of the Cannes disclosure standard.²⁹⁹ **Gareth Edwards**, the *Rogue One* director, came out in favour with the most usefully double-edged formulation of the period: “*It has no taste whatsoever. It is a genius at helping you... like a billionaire on acid.*”³⁰⁰ Read that carefully and it is a Position Four argument wearing an enthusiast’s clothes — infinite resource, zero judgment; the taste stays with the director or the film has none. And **Emily Blunt** revealed she had turned down Steven Spielberg’s offer to use AI for her alien voice in *Disclosure Day* — a working performer declining the tool at the exact point where it would have touched the performance itself, on a production run by one of the most technologically fluent directors alive.³⁰¹ Endorsement, regulation, double-edged embrace, point-refusal: four artists, four positions, one month.

There is also a third register emerging, past endorsement and refusal, in which the auteurs engage the technology by making it the material. In early July, **Neon** — a distributor whose brand, like A24’s, is built on curation authority — released the first trailer for *Sheep in the Box*, **Hirokazu Koreeda’s** AI family drama.³⁰² Koreeda, the Palme d’Or-winning director of *Shoplifters*, is among the most humanist filmmakers working; that he has turned his attention to AI, and that a prestige-indie distributor is putting its marketing weight behind the result, tells you the arthouse economy has concluded that the subject itself now carries dramatic — and commercial — weight. The strongest position in the argument between the artists may turn out to be

neither the endorsement nor the refusal, but the film that metabolises the technology into the work.

Leonardo DiCaprio, in December 2025, told *The Hollywood Reporter*: “*I think anything that is going to be authentically thought of as art has to come from the human being.*” The headline framing reduced the position to “AI can’t be art because there’s no humanity to it,” which is the version that travelled, but the full quote is more philosophically defensible. DiCaprio wasn’t claiming AI-augmented work couldn’t be valuable. He was claiming that the authorship signal — “from the human being” — was a precondition for the *category* of art, as he understood it.³⁰³

Claire Foy told the *Daily Mail* in January 2026 she had “no interest” in seeing AI in films and would be “disappointed” if it became the future of Hollywood.³⁰⁴

Jenna Ortega said in December 2025 it was “very easy to be terrified” of AI in filmmaking. Her fuller reasoning, given to *NME*, is the part I have ended up quoting in talks: “*It comes to a point where it becomes sort of mental junk food and we feel sick and we don’t know why. I think, as terrible as it is to say, sometimes audiences need to be deprived of something in order to appreciate something again.*”³⁰⁵ That argument — *audiences need to be deprived of something in order to appreciate something again* — is one of the most interesting things a working performer has said in this period about the slop ceiling and its psychological substrate. The audience does not, by Ortega’s read, simply discriminate against AI work. They develop a *hunger* for the human-authored work *because* of the AI flood. The flood and the hunger are part of the same cultural dynamic.

Chris Pratt publicly rejected a pitch to cast an AI ‘actor’ as the villain in *Mercy* in January 2026: “I don’t think that’s a good idea at all.”³⁰⁶

I do not think any of these positions are static. I think some of them will shift in the next eighteen months, in ways that depend on how the policy environment, the audience response and the tool ecosystem evolve. But I think the *fact* of the positions, written down, in public, on the record, is more important than whether any individual position holds.

What these refusals do, collectively, is keep open a part of the creative economy that the all-in studios would otherwise be forced to close. They make it possible — for the audience, for the working performer, for the next generation of creative-industry workers entering the field — to have a viable career path that does not require AI integration as the price of admission.

In a world without these positions, every working creative would, by default, be a *partial AI operator*, whether they wanted to be or not. With them, the choice remains open.

That is not a small thing. It is the architecture of the future creative economy being deliberately preserved, by people with the cultural standing and the economic security to preserve it.

Position Four: The middle

The position I find most interesting is the one almost nobody articulates clearly, because it does not make a good press release. It is the position of studios that *use AI everywhere except where it shows up in the finished work*.

The clearest version of it is the one **Aardman** has effectively articulated and that **Bethesda's** Todd Howard described in *PC Gamer* in December 2025: AI is “part of Bethesda’s toolset for how we build our worlds or check things” — but it cannot replace human creative intention.³⁰⁷

You see the same position in **Amazon's** *House of David* — 350 AI shots, disclosed up front, but every one of them used to augment rather than originate the work. The show’s creator Jon Erwin gave *Wired* a metaphor that I have not stopped thinking about: “*You can put a very real camera on a very real actor and direct that actor, direct the camera, and that becomes, in essence, the hand inside a puppet. The puppet itself is this digital world that you create.*”³⁰⁸ The hand-inside-the-puppet image is the cleanest articulation I have heard of where the *Position Four* studios are choosing to put their human craft: at the moments of direction and performance, with the AI doing the digital-world infrastructure underneath. You see the same position in the *Battlefield 6* development team’s statement, in October 2025, that generative AI had been “seducing” but ultimately used only in the earliest stages of the game’s development, “to allow for more time and more space to be creative.”³⁰⁹

You see the same position in **The Witcher 3** and **Cyberpunk 2077** director’s November 2025 framing — AI “can help, but not replace, creatives.”³¹⁰

You see the same position in the **Wallace and Gromit** creator Nick Park’s December 2025 framing — embrace the technology, but be cautious about the values.³¹¹

You see the same position, most starkly, in the May 2026 **Sony** announcement that it was “going all in on AI for games” — *with the specific framing that AI was a force multiplier, not a replacement*. Mocap-to-facial animation in seconds rather than hours. AI integrated into asset generation, QA, engineering and animation pipelines. The goal: *more games, faster*. But Sony’s framing, which I have read several times to make sure I am not over-reading it, repeatedly emphasised AI as a tool *inside* the creative work, not as a substitute *for* it.³¹²

You see the same position, by May 2026, in the **Cannes** festival press cycle, where the working auteurs on the Croisette had shifted markedly from the prior year’s

defensive stance to a more cautious *acceptance of inevitability* — framed less as enthusiasm than as a refusal to be left out of the next decade’s tooling argument.³¹³

Peter Jackson, in a May 2026 interview during the Cannes window, summarised the Position-Four read in a single line that I think will travel further than the Sony announcement: *AI is, in essence, the next wave of special effects*. The director’s job is not changed by the SFX wave. The director’s job is to know what the film is for, and to deploy whatever tools are now in the box to get it there.³¹⁴ **Take-Two’s** Strauss Zelnick made the same argument from inside the AAA games business in the same week — that AI “*datasets by their very nature are backward-looking*” and so cannot, alone, make an *original* hit, but that AI is “*super helpful*” in the production of hits the human team has already conceived.³¹⁵ Different industries, different framings, structurally identical position: *tool in the workflow, not author of the work*.

This middle position — *AI in the workflow, not in the work* — is, I think, where most of the surviving major studios are going to land in 2030. It is the most defensible commercial position because it captures the productivity upside without giving up the cultural-product specificity that the audience continues, against the slop ceiling, to demand. It is also the most defensible *ethical* position, because it allows the studio to credibly claim — and to credibly prove, with disclosure and documentation — that the creative work the audience sees was, in its decisive moments, the work of human creators.

What is at stake, for the studios that get this right, is the next two decades of cultural authority. The studios that adopt aggressively and badly will look like the early-2010s newspapers that switched to clickbait. The studios that refuse entirely will look like the 1980s record companies that refused to release CDs. The studios that thread the needle — that adopt the productivity benefits without surrendering the human authorship signal — will, in my view, be the studios that the audience actually trusts in 2035.

Position Five: Own the rails

A fifth position emerged late in the period this book covers, and I want to flag it as its own line on the strategic map even though it sits, in some ways, *inside* the Position One *all-in* framing. It is the position of the rights holder that, instead of *licensing* its IP to a frontier-AI platform — and discovering, on the Disney–OpenAI evidence, that the platform underneath the licence is more fragile than the IP it covers — *builds the licensing infrastructure itself*. Owns the operating system. Owns the meter. Owns the fence.

The canonical case is **Hasbro’s Sixth Wall**, the in-house AI studio the company launched on **3 June 2026** under CEO Roberta Thomson — and the **CharacterOS**

governance layer underneath it.³¹⁶ What Sixth Wall is licensing is not the appearance of Hasbro's characters but their *behaviour* — how Optimus Prime argues, how Megatron menaces, how Cobra Commander reacts — for deployment inside chatbots, games, toys and brand experiences built by third parties. CharacterOS encodes each character's canon, voice, vocal timbre, personality, lore and safety limits and enforces them at *runtime*: a licensed Cobra Commander stays in canonical character without a writer in the loop. The economic model is a B2B programmatic licence layered on a *talent-participation pay-out* — voices built from authorised recordings of real human performers, who are compensated for the generative use of their work — and the studio's initial focus is restricted to 13+ and enterprise use, explicitly not building for young children, against the regulatory weather around AI companions and minors.

CEO Chris Cocks framed Sixth Wall as a “*third way*” between endless and costly enforcement of unauthorised fan creation on the one hand, and passive brand dilution on the other.³¹⁷ The framing is, in my read, the single cleanest articulation of what the *governance layer* actually is. It is not a creative technology dressed as governance. It is a governance technology dressed as a creative one. The deeper conceptual move underneath it — and the one I think will define legacy-IP strategy through the rest of the decade — is the **reconception of IP from an asset to be guarded into an environment to be governed**. For a century, intellectual property was understood as a discrete, ownable thing whose value you protected by controlling access to it. What the 2026 deals reveal is a redefinition of IP as a *bounded space, governed by rules, inside which value is created continuously by people who are not you*. Hasbro doesn't sell you Optimus Prime. It sells developers a governed space in which Optimus Prime can be summoned, on-brand and on-licence, a million times over. The asset is the rulebook.

The same structural move sits underneath the **Spotify–Universal Music Group** deal announced on 21 May 2026 — fan-made AI covers and remixes of participating artists' songs as a paid Premium add-on, with revenue flowing back to the artists and songwriters whose work is reinterpreted.³¹⁸ The deal's framing is a three-pillar contract — **consent, credit, compensation** — that maps onto the *Petrillo template* I laid out in Chapter 6. *Consent* is opt-in: only the catalogues of artists who actively agree are available for remixing. *Credit* is programmatic: derivatives are linked to original metadata, and the *Verified by Spotify* badge steers attention and money toward sanctioned work. *Compensation* is a new royalty stream layered on top of standard streaming. The deal's *prehistory* — UMG calling Suno and Udio's training “*infringement on an almost unimaginable scale*” in 2024, then settling with Udio in October 2025 and agreeing to build a licensed product — confirms the line I argued in Chapter 6: the lawsuits were never the destination. They were *leverage to force a licensing table into existence*. The Spotify–UMG product is the mature form

of that turn — litigation converted into a revenue line, and the architectural ancestor of what every other major-label and major-streamer pairing will, on the available evidence, build over the next two years.

The Disney–OpenAI collapse, read alongside the Hasbro and Spotify moves, supplies the lesson. **Licensing to a fragile platform you don’t govern is not a strategic position.** *Equity in the platform without governance rights* is not a strategic position either. The strategic position — Position Five on this chapter’s map — is owning the *governance layer itself*: the licensing protocol, the runtime guardrails, the provenance metadata, the verified-by stamp, the royalty-attribution rails. The legacy industries that absorb this lesson and *build their own governance infrastructure* — Hasbro’s CharacterOS, Spotify’s *Verified-by* stack, the **Cannes AI Disclosure Standard**, the C2PA / SynthID provenance layer described in [Chapter 13](#) — are positioning themselves to capture the durable value of the generative economy. The legacy industries that wait for the platform companies to build it on their behalf will end up where Disney ended up: thirty minutes’ notice and an exit.

The sharpest test of that lesson arrived in June 2026, from the studio this chapter opened with. In the first week of the month, Lionsgate vice-chairman **Michael Burns** said publicly that AI would save the company “*tens and tens of millions of dollars a year*”³¹⁹ — the Position One cost argument, stated as plainly as anyone in a legacy studio has stated it. The following week, **Runway** and **Lionsgate** announced an expanded partnership in which Lionsgate took an **equity stake** in Runway.³²⁰ Read against this chapter’s opening, the arc is remarkable: eight months from *Futurism*’s “crumbled into disaster” headline to shareholder in the model company — the fastest strategic rehabilitation on this map. And read against the Disney–OpenAI collapse, the structure is the significant part. This is a legacy studio moving from licence terms to an equity position in the model layer itself, which is precisely the correction the collapse prescribed: *align on equity and governance rights, not just licence terms*. Whether it is a genuine Position Five move, rather than Disney with a smaller cheque, turns on the part of the announcement a press release never settles — governance. Disney’s \$1bn pledge failed not because it was equity-shaped but because it carried no rights strong enough to buy the company more than thirty minutes’ warning. If Lionsgate’s stake carries real information and governance rights — visibility into the roadmap of the platform its catalogue now trains and its pipeline increasingly runs on — then a mid-sized studio will have executed the strategic correction the largest studio in the world fumbled. If it does not, the equity is a licence with a balance-sheet entry. On the evidence available at the time of writing I lean towards the former reading, for one unscientific reason: Lionsgate has already lived through the failure mode once, in public, and companies rarely repeat the specific humiliation they have been mocked for by name.

There is one further observation I want to surface here, because it shifts the read on the labour story underneath the rest of this chapter. The transaction-cost frame I keep coming back to — that generative AI compresses the cost of search and negotiation but raises the cost of monitoring and enforcement — has a *labour consequence* that the substitution-panic framing of 2024 missed entirely. The work doesn't disappear; *it moves*. Voice performers who supply authorised source recordings, curators who write and maintain canon, the *judgment brokers* the coordination-collapse argument describes, provenance and rights-management staff, brand stewards, the *Continuum Lead*-shaped roles I described in Chapter 15 — these are the jobs the governance layer creates as fast as it absorbs the routine production layer. The new roles are fewer and differently distributed than the ones they replace, and “*talent participation*” is a weaker contractual instrument than employment. But this is the actual shape of the disruption, and the studios reading the strategic map well are the studios that recognise it and start staffing for it.

The trap the legacy industries built for themselves

There is a deeper strategic risk underneath the five-positions map that I have, until now, deliberately not put on the page. I want to put it on the page, because I think it is the single most important read on the long-term legacy-studio position, and because, in the conversations I have had with senior creative-industry executives over these six months, it is the read they are most uncomfortable hearing.

The risk is this.

Across the last fifteen years — most aggressively across the last decade — Hollywood, commercial music and the AAA games business have, on the available evidence, *systematically optimised themselves for exactly the kind of work AI is now best at producing*. They have built their economic and creative production engines around the mean of the distribution. They are now competing, in the most direct possible sense, against a technology built to produce the mean of the distribution at near-zero marginal cost.

Let me make the case concretely.

In **film and television**, the IP-cycle data is unambiguous. Sequels, prequels, reboots, remakes, spin-offs and franchise instalments accounted for an increasing fraction of the top-grossing US theatrical releases through the 2010s and 2020s; *original* studio films — meaning IP not derived from an existing book, comic, game, brand or prior film — fell from a majority of major studio output in the late 1990s to a single-digit percentage of wide releases by the mid-2020s, depending on which counting convention you use. The headline form of contemporary tentpole filmmaking, by 2024, was a sequel to a property whose original installment had itself been a sequel. Marvel Cinematic Universe Phase Six; the *Star Wars* sequel-trilogy aftermath; *Avatar* sequels, *Toy Story* sequels, *Frozen* sequels, *Mission:*

Impossible sequels, *Fast & Furious* sequels; *Stranger Things* finales; *House of the Dragon* spin-offs; every 1980s and 1990s IP revisited at least once, most of them more than once. James Cameron’s *Hollywood Reporter* observation in 2025 that contemporary studio executives “don’t reach for things that are scary” was, on the data, a description, not a prediction.

In **commercial recorded music**, the structural pattern is the same. The streaming-economy hit structure converged on a remarkably narrow set of parameters across the 2010s: average song length compressed from roughly four minutes in 2000 to roughly three-and-a-half minutes by 2024; choruses landed earlier; intros shortened to keep the algorithmic skip-rate down; major-label A&R increasingly drove signing decisions by predictive-analytics data rather than developmental A&R judgment; co-writing teams expanded; the median Top 40 hit, by 2024, was a co-write across three to seven credited songwriters working to formulas that had been tested in advance against streaming-engagement data. The major-label business is, structurally, an *industry that has spent ten years training itself to produce the most predictable possible version of the song*. The Cardiff band from Chapter 5 — whose music was fed to an AI that produced a tracking-style imitator outperforming them on Spotify — is the canonical illustration. The AI did to them what the major-label streaming-optimisation operating model had already half-done. It made the next most-likely-good track. The fact that an AI could match the output is, on the available evidence, *because the major-label hit factory was already producing the work AI was about to be able to copy*.

In **AAA games**, the standardisation is even more visible. The *Ubisoft-tower-and-checklist* structure — open world, viewpoint towers that uncover map regions, scattered icon-driven side quests, levelled enemy zones, crafting trees — became, between roughly 2008 and 2022, the default structural template for the AAA action-adventure genre. *Assassin’s Creed*, *Far Cry*, *Watch Dogs*, *Ghost Recon*; outside Ubisoft, *Horizon*, *Spider-Man*, *Mad Max*, *Shadow of Mordor*, every Tom-Clancy-derivative, every western-RPG converted to console. The 2024 *gamesindustry.biz* roundtable I referenced in Chapter 3 — in which working AAA designers described the genre as having stagnated into a single repeating structural pattern — was a working-developer admission that the AAA industry had, like Hollywood and like the major labels, optimised its production engine around predictable, low-risk, repeat-format output.

Now consider what AI, as a creative tool in 2025–26, is structurally best at. Agents — as I argued in Chapter 12 — produce the *mean of their training distribution* by default. The mean of the training distribution, for a video model trained on contemporary Hollywood tentpoles, is *another contemporary Hollywood tentpole*. The mean of the training distribution, for an audio model trained on the contemporary streaming Top 40, is *another contemporary streaming Top 40 song*.

The mean of the training distribution, for a games-development agent trained on the AAA open-world template, is *another AAA open-world template*.

This is the strategic trap. The legacy industries, by spending fifteen years training themselves to produce the mean of the distribution, have arranged for the segment of the market they dominate to be *exactly the segment AI replicates most cheaply*. The Cardiff band's experience is the cleanest version of this dynamic in microcosm. The macro version is the major-studio business model. The risk to legacy Hollywood, legacy commercial music and the AAA games business is not that AI takes their *premium* segments — Cameron's *Avatar* sequels, the highest-end auteur cinema, the genuinely original musical voices, the *Baldur's Gate 3*-class boundary-pushing games. Those segments are, on the evidence of the slop ceiling and the authenticity premium, *more defensible than ever*. The risk is to the *median* output of these industries — the franchise instalments, the by-the-numbers chart hits, the AAA action-adventures that read as algorithmically generated even when no algorithm was involved. The median output is exactly where the AI substitution pressure is most direct, and the median output is precisely what the legacy industries have most thoroughly optimised themselves to produce.

And by the summer of 2026 the trap was acquiring its own tooling. In June, **Quilty** — an AI script-analysis start-up — struck a preferred-look deal with Giovanni Entertainment,³²¹ a week or so after *The Verge* had examined the same company under the headline question: can AI tell if your script will make a hit film?³²² Greenlighting-by-model is the strategic trap with the lid bolted shut. A hit-prediction system is, by construction, a machine for measuring proximity to the historical mean; install it at the greenlight and you have automated the exact commissioning bias — *produce the work that resembles what already worked* — that this section has been describing. Every producer who filters development through such a model is asking the machine the only question it can answer, *is this like the things that already succeeded?*, and systematically declining to ask the only question that protects them: *is this the move the machine would not have generated?* The technology is not the trap. The technology is the trap's enforcement layer.

The grandmasters of Chapter 16 — the chess players who have started, in 2026, to deliberately play sub-optimal moves to put their opponents on uncomputed ground — are, on this read, the *senior auteurs of legacy Hollywood*. Cameron, del Toro, Soderbergh, Spielberg, Aronofsky, Lyonne, Larian's Sven Vincke, Hooded Horse's leadership — the figures profiled in the *Position Three* section of this chapter — are, structurally, the people whose competitive advantage is *the move the machine would not have generated*. The grandmasters can take the punch. The middle ranks — the franchise journeymen, the streaming-optimised mid-tier filmmakers, the chart-A&R commercial-pop machine, the AAA studio designing its fifth open-world

action-adventure — are the ones whose business model the machine is structurally suited to replicate.

The contrast with the **new AI-native studios** is sharp, and I want to draw it out carefully because the contrast is, I think, the most underappreciated strategic reality of the period.

The Position Two studios in this chapter — Gossip Goblin, Critterz, Imaginae, Wonder, Asteria, Promise, Obsidian, Chapter41, Kartel, Goldfinch's enGEN3 — were, by May 2026, accumulating creative production credits at a rate the legacy studios were not matching. **Gossip Goblin**, the AI filmmaker that I covered in Issue 29 of *Dream Machine*, is the example I find clearest. It is a studio with no inherited IP, no inherited production pipeline, no inherited audience, no inherited rules about what its films should look like, and no inherited risk-aversion. Its only operating constraint is the one any working creative has: make work the audience wants to watch. Critterz (*Vertigo* + *Federation*), launched as an AI-assisted animated feature operation in Issue 29, operates with the same freedom. Animaj (the kids-content AI studio Google's AI Futures Fund partnered with in spring 2026) operates with the same freedom. *Imaginae Studios* — Fremantle's AI-native operation — has, in *Art Awakens*, committed to a kind of generative collaboration with classical painting IP that no legacy studio has the institutional permission to attempt.

These studios do not have rules about how a film should be paced, what a song should sound like, what an open-world game should feel like, what a franchise structure should be. They have no quarterly-earnings call about year-on-year tentpole performance. They have no \$200m development sunk cost in a Marvel-style multi-film slate. They have no major-label playlist-pitching infrastructure that demands the song be a certain length and a certain shape. They have, in short, no calcified definition of *what counts as the right move* — and so they are, by default, free to play the move the legacy studios cannot.

The risk to legacy, in other words, is *not symmetrical* with the risk to the AI-native studios. Both are being shaped by the same technology shift. But legacy faces a double squeeze — its median product is the part of the market AI is most easily eating, *and* its strategic incentives push it deeper into that part of the market every quarter, *and* its calcified production rules prevent it from doing the thing (the deliberately un-machine-like move) that would protect it. The AI-native studios face only the upside.

The path out, for the legacy studios that recognise the trap, is roughly what the *Position Three* signatories have intuited: refuse to compete in the median segment; reposition aggressively into the segments where the audience pays a premium for the human signal; treat IP investment as *bets on artist-author voices* rather than as

bets on formula. *Baldur's Gate 3* — Larian's mid-2020s blockbuster — was the inflection-point example: an enormous AAA-grade game built explicitly *against* the Ubisoft-tower template, on a CRPG framework most analysts had written off as a dead genre, and the audience response was the largest commercial success of any new RPG IP of the decade. Larian's January 2026 announcement that the next *Divinity* game would not use generative AI³²³ is, on the strategic read, not an anti-technology gesture. It is a *commercial* gesture — a public claim that Larian's market position is built on doing the un-machine-like work, and that the studio is going to protect that position from the inside. *Pocketpair's* “*we don't believe in it*” statement, *Jagex's* “*never*,” *Hooded Horse's* “*cancerous*” framing, *Aardman's* careful preservation of the stop-motion craft — these are not statements of moral piety. They are *competitive positioning* against an algorithm-optimal market that the major-studio system has, structurally, conceded to AI.

The legacy industries that survive this transition will be the ones that recognise, in the next eighteen months, that the position they spent fifteen years moving into is exactly the position they now need to move *out of*. The audience, the slop ceiling, the authenticity premium and the chess grandmasters' move are all telling them the same thing: *do not produce the work the machine can replicate. Produce the work the machine, by construction, cannot*. The industries that can absorb that redirection in their commissioning, contracting and greenlighting culture will outlast the transition. The industries that cannot, won't. The new AI-native studios — Gossip Goblin and the rest — are not the threat to legacy Hollywood. They are, on the structural read, *the proof of what survives*: studios built without the rules that made the legacy industries vulnerable in the first place.

What the studios got right

I want to close this chapter by saying something that is unfashionable in some of the creative-industry circles I move in.

The studios — for all the headlines about Lionsgate's “disaster,” for all the criticism of Disney's OpenAI deal, for all the eye-rolling at Netflix's “all in” framing — have, in this period, made some genuinely difficult strategic decisions in a genuinely difficult environment, and a meaningful fraction of those decisions have been better than the public discourse credits them for.

They have committed to disclosure. *House of David's* 350 AI shots were *disclosed*. Aardman's careful framing was *explicit*. Sony's “AI as force multiplier” framing was *spelled out*. None of these companies pretended their AI use didn't exist. None of them adopted the all-too-common platform-economy strategy of *use it and don't tell anyone*. The disclosure norm, where it has taken hold in the studio system, is a public good.

They have negotiated, in many cases, in good faith with the unions. The SAG-AFTRA contract that emerged from the autumn 2025 negotiations was — by historical comparison with other major technology transitions — produced quickly, produced through legitimate process, and produced with materially stronger AI protections than any prior contract in the industry’s history.

They have, in significant cases, *resisted* internal pressure to over-adopt AI in ways that would have undermined their cultural product. Most of the studios in Position Three above are not run by Luddites; they are run by people who have done the maths on the slop ceiling and decided that the long-term value of their IP depends on it remaining recognisably human-authored.

And they have, finally, invested in the infrastructure of the AI-native sector — through funding deals, through co-productions, through equity investments — in a way that means the AI-native studios are not fighting them so much as building alongside them. The picture, ten years out, is more likely to be of a *mixed ecosystem* — legacy studios with hybrid pipelines, AI-native studios with new IP, and a long tail of human-only craft studios serving the highest-value segments of the market — than of a single winner-takes-all outcome.

The studios decided, in these six months, that they were not going to be replaced by AI. They were going to be the *operators* of AI. That decision was — for all its compromises — probably the right one for the working creatives whose careers depend on the studio system continuing to exist.

But a strategic position is only as durable as the capital underneath it. Every posture on the map above — all-in, AI-native, refusal, the middle path, own-the-rails — is, in the end, a claim about the future that somebody has to *finance*: a greenlight, a slate, a completion bond, a term sheet. And while the studios were deciding what to make, the money behind them was making a quieter, stranger decision of its own — one that has less to do with what AI can create than with what AI has done to the price of *not knowing*.

That is the subject of the next chapter: what the money has actually concluded about creative work in the AI era — and why the funding crunch it gets blamed for was already underway when the machines arrived.

Chapter 8 — The Machine Gets the Money

Three verdicts, so far, in the first half of this book. The audience delivered theirs through the slop ceiling of Chapter 5: flood us, and we will simply decline to listen. The creators delivered theirs through the 88% of Chapter 6: license the work, in all cases. The studios delivered theirs through the strategy map of the last chapter: a declared intention, position by position, to operate the machine rather than be operated on by it. This chapter is about the fourth voice in the room — the one that speaks last and settles the argument. The capital. It has been talking all year, through term sheets and completion bonds, through greenlights granted and greenlights quietly never offered, and most of the industry has been mishearing what it says.

The mishearing goes like this. When an AI-branded studio raises nine figures in a week while a conventional slate dies in development, the industry reads a verdict on creative work: *the money thinks the machines are winning*. When the same money then pays \$55 billion for a games publisher, or pours sovereign construction capital into giant immersive venues, the industry reads nostalgia, or noise. Both readings assume that capital is betting on an outcome — that somewhere behind the term sheets sits a settled view about whether AI-made or human-made work wins the decade. Ten months of watching the allocation data says otherwise — and the newsletter has, without quite meaning to, kept a weekly tape of that data: every raise, every fund, every shutdown, running underneath the culture war like a ticker under a news broadcast. The money is not betting on an outcome. It is repricing the possibility of *knowing* — and to read that correctly, you first have to be honest about the state the patient was already in when AI walked into the room.

Because the funding crunch did not start with AI, and a chapter that pretended otherwise would be propaganda for the technology's importance. Three conditions were in place before a single generative model mattered. **The money was already over-extended.** The lockdown years pushed the streaming wars to maximum content spend and games venture to its \$12.5 billion 2021 peak; when the cheap-money era ended in 2022, the correction was ordered by interest rates, not by algorithms — the layoffs, the cancelled slates and the eighty-per-cent venture drawdown the games section counts below all began before capable generative video existed. **The money was already concentrating.** A decade of franchise economics had taught film and games capital to crowd into the biggest IPs — sequels, tentpoles, live-service giants — while the mid-budget original and the new studio's first title starved; "Peak TV" was already finished as a commissioning regime. **And the audience was already drowning.** Tens of thousands of new tracks a day arriving at a single streaming service, record release counts on Steam, more scripted television than any human could watch: content supply had outrun human attention years before generative AI arrived, which is the supply-side half of the story Chapter 5 tells from the audience's side. Over-invested, over-concentrated,

over-supplied — that was the creative economy the new technology met in the autumn of 2025. Already correcting. Already barbelled. Already saturated.

What AI did to that stressed system is the subject of this chapter, and it is stranger than either the boosters or the doomers predicted. It did not cause the correction; rates did. What it did was *mutate* the correction, in two ways. First, it made the oversaturation permanent. A glut is a market phase; a glut at near-zero marginal cost is a climate. The flood Chapter 5 measured from the audience's side is, on the capital side, the collapse of every scarcity assumption a content business model was underwritten on.

Second — and this is the deeper break — it destroyed *describability*: the ability to say, well enough to underwrite, what the market for a piece of long-form creative work will look like on the day the work finally meets it. Everything in Chapter 5 says audiences will want the human-made versions of films, games and television in 2033 more fiercely than ever, and almost nobody writing cheques seriously disputes it. But wanting is not underwriting. Hollywood's founding epistemology — William Goldman's "nobody knows anything" — was always a statement about demand inside a stable game: nobody knew which film would hit, but everyone knew what a film was, what it cost, and what a market for one looked like. The slate was invented to diversify exactly that risk. Generative AI has produced the second-order version — *nobody knows whether anybody will know anything* — and the financial corollary of that sentence gives this chapter its title. When nobody can describe the market for the work, capital funds the one thing it can still describe: the machine that makes the work cheap. The machine gets the money. The work gets the discount.

The economists supplied the vocabulary a century ago, and it is worth thirty seconds of theory because it explains everything the sector data is about to show. Frank Knight's distinction: *risk* — variation you can put a distribution over — gets priced every day; *uncertainty* — variation you cannot — gets fled, or charged returns so punishing that only the desperate transact. The real-options economists added the clause that matters most for creative work: when commitments are irreversible and the environment is uncertain, the value of *waiting* soars, because committing destroys the option to commit later with better information. A greenlight is close to perfectly irreversible. A seven-year greenlight in 2026 is therefore the deliberate destruction of seven years of option value at the precise moment option value has never been higher. Investors do not need a theory of AI victory to hesitate; the arithmetic hesitates for them. Nor does the uncertainty cut only against content: the February wobble that wiped a reported trillion dollars off the big technology firms on AI worries was the same mechanism running in reverse — a reminder that the disruptors' own equity is a claim on an unpriceable decade too — the difference, which the film section will price properly, being that a tools business marks itself to market in months, in revenue, while a greenlit slate cannot mark itself at all.³²⁴ And

the far pan of the scales kept getting heavier all year: by the summer, AI was being credited with a third of all US economic growth.³²⁵

I started calling this the **Seven-Year Problem** in the spring, when the same question kept surfacing in three different rooms I was sitting in — an investment-programme assessment panel, a games-studio board call, a public-funding steering group. *Won't AI have taken over by the time this ships?* And in each room the people actually controlling money gave a version of the same answer: that is not the question. Almost nobody writing cheques believes, in an underwritable form, that AI will have taken over. What they believe is that they cannot say what the market will be — and an underwriter prices *cannot know* and *will end* identically. That identity is the engine of everything this chapter describes. Its output is what I will call the **duration discount**: the additional haircut capital now applies to creative assets in proportion to their production and payback horizons.

But here is the thing the aggregate framing hides, and the reason this chapter is built the way it is: the discount is not a weather system that fell evenly on the creative economy. Duration varies by sector — a AAA game is a seven-year claim, a feature three to five, a television commission one to three, a venue a different kind of claim entirely — and moats vary with it. So the repricing has drawn a *map*. Meanwhile the pull on the far side keeps strengthening: the generative-content market that stood at \$14.8 billion in 2024 is projected to quintuple within five years, with forecasters sizing it at \$80 billion by 2030 — a gravitational mass every content pound must now be weighed against.³²⁶ AI has not made creative investment uniformly scarcer. It has re-sorted it, sector by sector, by duration and by moat, and the four sectors this book lives across — games, film, television, immersive — have received four different verdicts.

And across all four, watch the *direction* of the money that stays, because it answers the question I am asked most often: if capital is fleeing the work, what is it buying instead? It is buying the layer underneath the work — the tools, the models, the training data, the infrastructure. Long the means of production, short the products.

Music, as ever, got its verdict first: the shortest cycle in the creative economy took the flood first, and Chapter 5 recorded the result — a majority of the upload pipeline now synthetic, a consumption share pinned in the low single digits, and the money pooling in the tools rather than the tracks. The tooling numbers have only steepened since that chapter's edition went to press: Suno, which raised \$250 million at a \$2.45 billion valuation in the week Issue 8 shipped, closed \$400 million at \$5.4 billion by the week of Issue 32 — a valuation that more than doubled in seven months while genuine human listening to fully synthetic music stayed under one per cent of streams — and Apple quietly made an AI audio company, Qai, its reportedly second-largest acquisition ever at \$1.6 billion.³²⁷ Hold that shape in mind. It is the

preview. What follows is the anatomy: what AI has actually done to the investment case in each of the four screen-and-experience sectors, read from the allocation data rather than the press releases — and then, at the end, what to do about it, because the answer this chapter builds to is instruments, not vibes.

Games: the longest bet in the building

Games carry the longest production cycles in the creative economy, so if the duration discount is real, games should show the deepest cut. They do.

Venture investment in games collapsed from a 2021 peak of roughly \$12.5 billion to \$2.54 billion in 2024 — an eighty per cent drawdown.³²⁸ Here are the receipts the opening promised, dates first, before AI takes the blame it has been eagerly assigned: the collapse *began in 2022* — before any tool existed that could plausibly threaten a AAA pipeline — and the roughly forty-five thousand jobs lost across the sector between then and mid-2025, the wave peaking in early 2024 and decelerating since, belong mostly to the monetary story.³²⁹ AI arrived mid-contraction and became the *narrative*: the all-purpose justification, as Chapter 14 documented from inside the org charts, for cuts that rates had already ordered. The labour market priced the resulting spread with brutal clarity in the newsletter's very first week: as the layoff wave rolled through the studios, Netflix was advertising up to \$840,000 a year for a director of generative AI for games.³³⁰

What AI changed is where the surviving money goes. The 2026 recovery is genuine and genuinely instructive: Drake Star recorded a twelve-month funding high above \$2.5 billion in the second quarter — concentrated not in games but in AI tooling, advertising infrastructure and user-generated-content platforms.³³¹ PitchBook's commentary was unusually blunt for a data house: gaming venture capital now flows to the companies building AI rather than to games, investors treating the sector as a proving ground for generative infrastructure and “funding the tooling layer ahead of individual-game content.”³³² Sixty-one per cent of *all* global venture capital went to AI firms in 2025;³³³ every pound contemplating a games position is benchmarked against that, and a seven-year content bet is not merely judged risky against it. It is judged *slow* — an indictment independent of any view about the work. You can watch the re-sorting week by week in the newsletter's own record, which through the winter and spring read like a tape of the tooling layer capitalising while new-studio content rounds went missing: Fifth Door's \$20 million for AI game creation, Razer's \$600 million pivot to AI gaming, Sett's \$30 million Series B for automated game marketing, Astrocade's \$56 million for interactive AI, Aippy announcing tens of millions at a \$250 million valuation, Welevel's \$8.5 million for an AI-built survival game — a round or two most weeks, and barely a conventional content greenlight among them.³³⁴

The darkest datum in the sector, though — the one I keep returning to — is an arc the newsletter caught at both ends. In October, Issue 3 flagged a \$134 million seed round, unusual size for the category, for General Intuition: a startup training agentic AI on video-game clips, teaching machines spatial reasoning from recordings of people playing. By the time Issue 36 went out in the summer, the same company had raised \$320 million at a \$2.3 billion valuation, and a smaller specialist, Origin Lab, had taken \$8 million to mine training data from games.³³⁵ By July the category had a name and a pipeline: Cambridge’s Worldmodeldata raised £7 million to turn video games into AI training data, Meta shipped a consumer app — Pocket — that turns every play session into training data by design, and open libraries of raw gameplay hours began circulating on Hugging Face.³³⁶ Sit with what that arc means. The fastest-appreciating asset to come out of the games industry in 2026 is not a game. It is the *exhaust* — the behavioural data trail of humans at play, harvested to train agents for the physical world. Capital has found a way to be long the games industry that requires no bet on any game shipping at all, and it is paying frontier-lab multiples for the privilege. If you want a single image of the duration discount at work, it is this: seven-year artefacts unfundable, while the by-product of their audiences’ attention commands \$2.3 billion.

Then, in the space of ninety days this summer, the sector staged the whole argument as theatre, in two capital events that should be read as a diptych.

The first ships this November: *Grand Theft Auto VI* — greenlit thirteen years ago, in a world that contained no transformer architecture, no diffusion models and no consumer generative video, and arriving now as what the analysts at Drake Star call likely the largest game launch in history.³³⁷ It is the cleanest surviving specimen of the underwriting world this chapter opened by describing: a decade-plus claim, committed blind, about to pay out anyway. And it is paying out not because the world proved predictable — by launch day the world contained a synthetic actress with a talent agency — but because the asset was moated in the one currency the transition cannot inflate: a fandom, a canon, a world people live inside and want more of *from these makers*. Even the publisher’s own leadership frames the technology as unable to originate the franchise — while, as Issue 24 reported and Chapter 14 dissected, laying off its entire seven-year-old internal AI team in favour of off-the-shelf tooling.³³⁸ Read those two facts together and you get the franchise position in miniature: AI as commodity input, the world and its audience as the asset.

But carry with you a correction the next chapter will earn in full, because capital is currently half-right in a way that will hurt the half that’s wrong. The moat that survives the flood is *fandom*, not *formula* — and the AAA titles whose value is production polish on the fifth iteration of an open-world template are, precisely

because the formula is what the models learn first, the most exposed content in the sector. An allocator who reads “sequels are safe” out of GTA VI is buying the wrong asset with the right instinct.

The second event closed on the fourth of August, ten days before this chapter’s edition date: Electronic Arts left the public markets in a \$55 billion transaction — the largest leveraged buyout in history — with Saudi Arabia’s Public Investment Fund taking 93.4 per cent, Silver Lake 5.5, Affinity Partners 1.1, on roughly \$36 billion of equity and \$20 billion of debt.³³⁹ Two features of that deal belong in this chapter’s evidence file. One: the identity of the buyer. When venture capital flees duration, the assets do not evaporate — they migrate to the one class of allocator whose horizon *exceeds* the uncertainty. A sovereign wealth fund executing a generational diversification strategy is, almost by definition, the natural holder of the Seven-Year Problem: it is the only balance sheet in the market for whom 2033 is the near term. Two: the role of AI in the underwriting. Twenty billion dollars of debt has to be serviced by margins, and the consortium’s own framing leans on heavy investment in AI applications across a live-service portfolio of seven hundred million players — a thesis the banks had already sized, with analysts talking of \$22 billion in AI-driven profit headroom across the industry.³⁴⁰ Which is to say: in the largest games transaction ever done, AI appears not as the existential threat to the asset but as the *margin story that makes the leverage add up*. The bear case and the bull case are the same technology, held by different balance sheets at different durations.

That is the games verdict, stated plainly. Venture money has repriced the sector’s long tail to near-zero and rotated into picks, shovels and exhaust; the long-duration assets that still clear the discount are carried by fandom moats and sovereign patience; and the middle — the new studio, the original IP, the first game — is where the discount lands with full force. Keep hold of that last clause. It recurs.

Film: the machinery seizes

Film’s cycles are shorter than games’ — three to five years from development to release — but film depends more than any other sector on a lattice of *intermediated* finance: banks, presales, co-productions, completion bonds, slate vehicles. And the most consequential thing AI has done to film investment is not to change anyone’s opinion. It is to seize the machinery.

Start at the dullest, most load-bearing joint. Banks, presale buyers and co-production partners release funds against completion bonds; bonds are underwritten against known production functions — a script, a schedule, a director with a track record, a physical process whose failure modes a century of claims data has catalogued. As of this year, no major bond company has published underwriting criteria for AI-native production, where the “script” may be a prompt architecture and the “schedule” a generative pipeline run — which locks AI-first features out of

institutional finance altogether and makes them dependent on exactly the fast, unstructured capital that arrives on a keyword and can leave on one.³⁴¹ Read the bond gap as epistemology, not paperwork. The institutions whose entire function is converting uncertainty into priced risk are declining to perform the conversion, because they cannot yet build the actuarial table. When insurance cannot form, the financial system is telling you, in its own dialect, that it has met genuine Knightian uncertainty.

The spread that seizure creates was measured, perfectly, by one man in February. Roger Avary — an Oscar-winning screenwriter with three decades of relationships — described how years of failing to finance conventional features through conventional channels ended the moment he founded an AI-branded production company: three films funded almost overnight, on his own account essentially because the word AI was attached.³⁴² One anecdote, but the newsletter's record shows it was no outlier — it was the visible tip of an asset class forming. Wonder, the British AI film studio, took a £9 million seed in October that Issue 5 logged as a curiosity; by the time Issue 30 shipped in May, Deadline was profiling it as the \$50 million studio styling itself the A24 of AI production. Filmology Labs announced itself as a \$250 million studio built for AI-era content. Singapore's Video Rebirth raised \$50 million chasing studio-grade AI video. A married couple got \$6 million out of Y Combinator on a pitch deck that literally wedded AI and filmmaking.³⁴³ Position Two of Chapter 7's map — the AI-native studio — is now a funded category, capitalised at nine figures, while the conventional independent market has moved, in the trade's own language, from momentum to scrutiny: harder questions, earlier, about comps, incentives, attachments and proof.³⁴⁴

Underneath the spread sits a cost curve whose steepness deserves plain statement, because it is the saturation mechanism of Chapter 5 arriving in film's economics. By spring, Seedance 2.0 had pushed generation cost below fourteen cents a second of video.³⁴⁵ A feature-length AI film reportedly made for around two thousand dollars premiered at Tribeca.³⁴⁶ And at Cannes — the datum I found most clarifying all year — the Wall Street Journal profiled a competition film that cost \$500,000, of which \$400,000 was AI compute: a production budget in which the compute line has swallowed the crew list.³⁴⁷ The mid-market is trying the hybrid version — *Killing Satoshi*, a \$70 million feature with a name cast, is being made with heavy AI use precisely to make its budget close³⁴⁸ — and the tooling vendors are selling the saturation thesis as a promise: Runway's chief executive pitches a world where Hollywood makes fifty films instead of one \$100 million blockbuster.³⁴⁹ Fifty films for the price of one is a marvellous sentence for a tools company and a terrifying one for anybody whose business model depends on the scarcity of the one.

Where the studio money has gone instead is the barbell, and one balance sheet states it more cleanly than any analyst note. The two largest acquisitions in Netflix's history are now a catalogue of beloved children's IP — the Roald Dahl Story Company, roughly \$700 million in 2021 — and a sixteen-person AI post-production startup, Ben Affleck's InterPositive, first flagged in Issue 20 at a reported \$600 million and confirmed — in the securities filing the newsletter caught in Issue 39 — at \$587 million cash.³⁵⁰ Proven story-worlds at one end; the machinery that makes everything else cheap at the other; and a walking-away from most of what lies between. The tooling end is where film's capital velocity now lives, and no company demonstrates it like Higgsfield, whose arc the newsletter tracked in real time: an \$80 million raise at a \$1.3 billion valuation in January; \$200 million of revenue within nine months of launch by February; half a billion in revenue and talk of a \$5 billion round by midsummer.³⁵¹ No content business on earth compounds at that speed — which is precisely the comparison every limited partner is now making. Runway sits at \$5.3 billion with dedicated film funds and a \$40 million Japanese expansion; Luma closed \$900 million Saudi-led; Kuaishou's Kling raised a record \$2.8 billion in July;³⁵² Adobe and Google stood up film grants; A24's DeepMind investment demonstrated, via instant fan backlash, that even the arthouse brand premium is now entangled with the provenance politics of Chapter 13.³⁵³

There is a shadow over that compounding, though, and an investment chapter would be malpractice not to price it. Every pure-play on the list — Higgsfield, Runway, Luma, Pika and the rest of a leaderboard that reorders itself roughly quarterly — is selling, at bottom, a model. And the two most consequential owners of video models sell no model at all. Google ships Veo through Gemini, YouTube and its cloud, wrapped in the service agreements, enterprise controls and SynthID provenance that procurement departments actually buy, while ByteDance's Seedance and Kuaishou's Kling arrive subsidised by short-video empires that need the firehose fed, collapsing the price floor towards pennies a second as they come.³⁵⁴ For the platforms, in other words, the model is not a business; it is a feature of a distribution empire — and features priced as businesses is a category for which venture history knows exactly one ending. Sora is what that ending looks like at frontier scale: the most-hyped video product of its generation, reportedly burning compute at a daily rate its lifetime revenue never approached, folded in March by the one lab with a claim to the best model but no distribution economics to carry it — in the same spring that Higgsfield's valuation quadrupled.³⁵⁵ A five-billion-dollar round and a shutdown, twelve weeks apart, inside one asset class: that is the duration distribution at the tooling end, stated plainly. Nor is the mortality confined to the frontier: in August, Reforged Labs — an AI ad-creative startup serving the games market — shut down with its chief executive conceding that “the gap we were selling into is closing”: the models were eating his market faster than he could sell into it.³⁵⁶

The model is the shortest-lived asset in the entire creative stack — the thirteen-year game will outlive every model used to build it, several generations over — and the moats that survive at this end are the ones this book keeps finding everywhere else: distribution, workflow relationship, and clean, licensed, provenance-bearing data. (And, in one small August datapoint I cannot resist recording: *taste* — Design Arena’s founders raised \$7.9 million that month to bring exactly that to AI models, the age of the *Why* arriving as a seed round.³⁵⁷) The barbell’s safe end is a knife-fight priced like a haven.

And notice the strange, telling instrument that has bloomed in the underwriting vacuum: the prize. Dubai launched a \$1 million AI short-film award with Google; the XPRIZE Future Vision competition put up \$3.5 million; Astana opened a \$1 million prize fund; Luma ran a \$1 million global brief; Adobe attached \$10 million in grants to its film push; Google gave Sundance \$2 million for AI literacy; Runway bankrolls films directly through its Hundred Film Fund.³⁵⁸ By August the pattern was compounding: Higgsfield opened a Global Film Festival with a \$1 million prize pool and half a million dollars for first place — having open-sourced its own 95-minute AI feature as the calling card³⁵⁹ — and the Raindance festival launched a development fund and hub for indie filmmakers with the AI platform IMG.N.³⁶⁰ Festivals and tool vendors, standing in for a financing system. When the modern instruments — bonds, presales, slate finance — cannot form, financing regresses to its pre-modern shapes: patronage, commission, and the prize. It is a rational adaptation and an indicting one. The film verdict, then: plumbing seized, spreads inverted towards the AI-branded, an AI-native studio class capitalised into existence, the cost curve collapsing under the mid-market, studio capital barbelled between catalogue and machinery — and the mid-budget original, the tranche the whole intermediated lattice existed to finance, left holding the discount, or queueing for a prize.

Television: duration arbitrage

Television ought, on the duration logic, to be the least disturbed screen sector: commissions run one to three years, catalogues amortise fast, and the streamers’ pipelines are annual. Instead television is where the repricing is most *visible* — because television is short-cycle enough for allocation decisions to show up in the data within quarters, and because the sector was already contracting when AI arrived: the streamers’ pivot from growth-at-any-cost to margin discipline had already cut scripted volume from its 2022 high, and by this winter the trade press was describing feature and episodic shoots drying up in Los Angeles as a fact of life rather than a cycle.³⁶¹

Inside the pipeline, AI’s arrival looks like Chapter 7’s Position One map executed at speed: Amazon’s internal AI Studios unit and the publicly disclosed 350-plus generative VFX shots in *House of David*; Netflix’s INKubator, an AI-native

animation unit reporting into the same commissioning structure as the live-action slates; Disney's enablement office and advertiser-facing generation tools — and, on the revenue side, AI agents deployed to grow Netflix's \$3 billion advertising business, the margin story running a floor below the commissioning story.³⁶² The talent has read the map too: the creator of the most-watched series in streaming history put \$3 million of his own money not into a production company but into Twelve Labs, an AI video-understanding startup — a working showrunner going long the tooling layer.³⁶³

But television also minted the period's genuinely novel underwriting peril. When Amazon's AI Creators Fund commissioned three generative animated series in May and one of them — *Punky Duck*, from a beloved human animator — collapsed within forty-eight hours of announcement under online derision, before a frame existed, the industry learned that the audience-enforcement mechanism Chapter 5 described operating on finished work can operate on a *press release*.³⁶⁴ Insurers have categories for weather, injury, bankruptcy and war. There is no category for *the commission was culturally repudiated at the moment of its announcement* — and until there is, every AI-adjacent greenlight carries an unpriceable tail that has nothing to do with whether the work would have been good.

The headline television story, though, is where the exiting capital *landed* — because it did not leave the sector. It moved down the duration curve, hard, into the vertical microdrama: phone-native serials of sixty to a hundred one-to-two-minute episodes, budgeted at one to two hundred thousand dollars a series, shot in days, monetised coin by coin. The numbers have gone from curiosity to category in eight quarters. Category app revenue ran from roughly \$178 million in the first quarter of 2024 to nearly \$700 million a year later; Deloitte projects in-app microdrama revenue around \$7.8 billion for 2026, more than doubling year on year; ReelShort alone booked around \$400 million in 2024 revenue, still loss-making on user-acquisition spend, while DramaBox turned profitable and went seeking \$100 million at a half-billion valuation; ByteDance entered the international market; TikTok quietly launched a standalone microdrama app.³⁶⁵ Hollywood veterans are placing what the trade openly calls big-money bets on the format precisely as the long-form shoots dry up; Fox took its equity stake in Holywater — the deal Issue 3 flagged back in October — and committed to two hundred shows in two years; one producer called verticals the first new professional format since reality television, while the legacy distributors, to their evident discomfort, sit out a market their sales model cannot parse — even as Netflix itself, in the plans Issue 26 logged, moves to add a vertical feed with AI-driven recommendation to its own app.³⁶⁶ When the incumbent that ended the last format war starts rebuilding its interface around the challenger's format, the migration is no longer avant-garde.

I want to name what this actually is, because the trade coverage keeps misfiling it as a taste story. It is a *duration* story. Capital fleeing three-year commissions it cannot underwrite has discovered a format whose entire production-and-payback cycle fits *inside* the visibility window — weeks from greenlight to revenue data, unit costs two orders of magnitude below episodic television, demand legible almost in real time. Call it **duration arbitrage**: the reallocation of creative capital towards formats whose cycles are shorter than the uncertainty horizon. And notice the second-order behaviour already emerging, because it is the market inventing, bottom-up, exactly the instrument this chapter has been arguing the long formats lack: vertical series are beginning to function as *proof-of-concept layers*, where audience response and narrative traction justify longer-form investment.³⁶⁷ That is staged, real-options greenlighting — priced discovery before irreversible commitment — built not by the studios' structured-finance teams but by an app economy that never read Dixit and Pindyck. The AI layer is arriving on schedule too — by late July, *Nikkei* was reporting from inside China's boom on AI moviemakers chasing “2 minutes of fame,” a fully synthetic tier maturing inside the format³⁶⁸ — which means the microdrama market is about to run the Chapter 5 experiment — flood, ceiling, premium — at compressed speed, on the shortest cycle in scripted entertainment. Watch it the way we watched music.

The television verdict, then: the sector where AI's investment impact is least about refusal and most about *re-sorting* — pipelines quietly re-tooled, announcement-stage peril priced in blood once, and the commissioning dollar migrating down the duration curve to wherever the future is visible from.

Immersive: the inversion

And then there is immersive — the sector this book comes home to — where the story runs backwards, and where reading it backwards matters most, because immersive is the clearest evidence that the duration discount is a *map* and not a mood.

Begin with the humiliation of the old thesis. The 2014-vintage immersive investment case was platform-led: fund the headset, the OS, the walled garden, and the content economy follows. A decade and colossal sums later, that case is being carried out of the building — Meta cutting a tenth of Reality Labs, as Chapter 14 noted, after cumulative losses that dwarf most national film funds; Apple's headset repositioned as a beachhead rather than a mass platform, its most quietly significant feature, as Chapter 9 observed, being that its Personas run on Gaussian splatting.³⁶⁹ Meanwhile the experience-led immersive economy — the one that sells tickets rather than devices — has been compounding through the exact months the flat-screen sectors repriced. Sphere's \$2.3 billion Las Vegas build has become, in its investors' own letters, a story about strong underlying unit economics; Abu Dhabi's culture-and-

tourism department is putting \$1.7 billion of construction capital into a second Sphere on Yas Island for 2029; Maryland has assembled roughly \$200 million in incentives for a smaller-format third; the stated ambition is a global network of venues.³⁷⁰ Netflix — the same balance sheet running the Dahl-and-InterPositive barbell — opened its first permanent physical destination in Philadelphia in November, with Dallas and Las Vegas to follow: a streamer converting IP into place.³⁷¹ Universal’s Epic Universe set the technical bar for the park industry in the same period, and the location-based trade grew its own convention by half.³⁷²

Why does venue-scale immersive clear a discount that is strangling screen media? Three reasons, and they are this book’s themes wearing hard hats. First, the revenue is *un-floodable*: a ticketed, capacity-constrained physical experience is the one entertainment product to which the marginal cost of synthetic content is irrelevant — you cannot slop a building. Second, it is provenance-native: presence is its own verification, the Authenticity Premium with a postcode; the audience contract of Chapter 13 is enforced by the door. Third, look at who carries the duration — a sovereign tourism strategy, a US state’s incentive package, a franchise balance sheet — and you find the games pattern again: the long-duration bets still being made across this entire economy are held by state capital, sovereign capital, or fandom moats. Venture appears nowhere on that list. That is the chapter’s unifying empirical law, and immersive states it in concrete and steel.

Now the AI part, which is where the inversion completes. For the entire history of the discipline, the binding constraint on immersive — the reason venue content refreshes slowly, the reason virtual production stayed expensive, the reason XR content economics never closed — has been the cost of building navigable worlds. Chapter 9 described what happened to that constraint: world models collapsed it. The cost of “a place” — persistent, lit, navigable — fell one to two orders of magnitude in twelve months; Sony’s virtual-production team quoted a fortyfold speedup; Netflix’s Eyleline shipped live-action-to-4D conversion; the environment stage left the critical path.³⁷³ What belongs in *this* chapter is the capital’s response, because the spatial layer is where the biggest creative-technology cheques of 2026 are landing, and the newsletter caught nearly all of them as they fell. World Labs’ further \$1 billion — Issue 19’s lead — at a reported valuation near \$5 billion, anchored by Autodesk’s largest startup investment ever, with NVIDIA, AMD and Fidelity alongside; Odyssey’s \$310 million at \$1.45 billion, logged in Issue 35 the week it closed, with Amazon and GV aboard; Tripo AI’s near-\$200 million for AI 3D and world-model tech in Issue 32; Yann LeCun’s AMI Labs reportedly closing a \$1.03 billion seed — the largest in European history; DeepMind shipping Genie 3; NVIDIA’s Cosmos passing two million downloads.³⁷⁴ Beneath the spatial layer, the operating layer capitalised on the same curve: ComfyUI, the open-source node graph that Issue 1 recorded raising \$17 million to build an operating system for creative AI,

crossed four million users and a \$500 million valuation by early summer.³⁷⁵ Strip the robotics framing off the world-model rounds and note what the products export to: Unreal, Unity, WebXR — the immersive toolchain. The tooling end of the barbell and the experience end of the barbell are, in immersive alone, pointing at the same future.

So the immersive verdict inverts the others. In film, television and games, AI is the reason capital cannot price the terminal market, and long-form content pays the discount. In immersive, AI is the thing that finally breaks the sector's historic cost constraint, in service of a revenue model the flood cannot reach. The same technology that reprices flat media *down* on uncertainty reprices spatial media *up* on feasibility. The duration discount, applied across the whole map, changes sign exactly once — and it does so in the sector where the UK, through CoSTAR's labs and networks and the R&D infrastructure this book's programmes plug into, and through simulation champions like PhysicsX with its \$300 million raise, has genuine positions on the board.³⁷⁶ That is either the best-timed industrial bet British creative policy has made in a generation, or an accident it should hurry to claim.

The map, read whole

Lay the four verdicts side by side and the Seven-Year Problem stops being a mood and becomes a mechanism you can read.

Games: the deepest venture repricing, capital rotated to tooling and to the industry's own data exhaust, long duration held only by fandom moats and sovereign patience — a \$55 billion LBO in which AI is simultaneously the anxiety and the margin story servicing the debt. Film: the intermediation lattice seized at the bond, the spread inverted towards anything AI-branded, an AI-native studio class capitalised from £9 million seeds to \$250 million launches inside eight months, the cost curve collapsing to fourteen cents a second and \$400,000 compute lines, the tooling pure-plays compounding under a hyperscaler shadow, and financing regressing to prizes and patronage where underwriting cannot form. Television: the fastest visible re-sorting — pipelines re-tooled, one commission killed by its own press release, and the money running down the duration curve into a format whose whole cycle fits inside the visibility window, inventing real-options greenlighting from below as it goes. Immersive: the inversion — un-floodable, provenance-native revenue meeting the one AI capability that dissolves the sector's historic constraint, with states and sovereigns carrying the venue duration that venture will not touch.

One mechanism generates all four: capital pricing *cannot know* as *will end*, then seeking whatever combination of shorter duration, harder moat, or longer-horizon balance sheet lets a deal clear anyway. And one casualty recurs in every sector, which is why I keep returning to it: the middle. The unproven studio's first game; the mid-budget original feature; the returning-series bet on a new voice; the independent

immersive production between the art installation and the theme park. Chapter 13 showed the Authenticity Premium accruing to the top and the long tail and weakly to the middle; Chapter 7 showed the middle most exposed to substitution; this chapter adds the third front — the middle is also, definitionally, the longest-duration claim in the market, originality being a bet on a future taste landscape with no catalogue floor beneath it. Squeezed from below by substitution, from the side by the premium's distribution, from above by capital's flight from duration. Three independent mechanisms, one victim — and the victim is where every franchise on the safe end of the barbell originally came from. Even the talent has begun saying it out loud: Brad Pitt, in August, argued that AI “used as a tool” could help more mid-budget films get made³⁷⁷ — an actor reading the same map this chapter has been reading, and volunteering the machine as the middle's least-bad financier. A capital regime that funds only proven IP and the tools to exploit it eats its seed corn with perfect quarterly discipline: rational for every allocator, ruinous for the system, which is the classic signature of a coordination failure and the classic licence for policy.

Which is why the instruments agenda matters more than any forecast, and why the UK's position is more interesting — and more precarious — than its headline numbers. The Sector Plan's finance package — the British Business Bank backing creative-sector venture managers, the expanded equity and debt funds, the exploration of IP-backed lending, the clusters money, CoSTAR's expansion, the mayoral funds — reads, squinted at, like a portfolio of proto-duration-instruments: public capital groping towards exactly the tranche private underwriting has repriced away.³⁷⁸

But hold the scales up honestly, because the newsletter's own record does. The state has just found £1.1 billion for an AI hardware plan; the Arts Council's cumulative funding for AI-and-creativity projects since 2019 stands at £4 million; and a typical creative-AI intervention is a £3 million feasibility pot for SMEs — sums that the spatial layer's private rounds exceed before lunch on an ordinary Tuesday.³⁷⁹ A country that can mobilise a billion for silicon but four million over six years for the culture the silicon is meant to serve has not yet understood which of its constraints is binding. And two structural problems remain. The instruments are underspecified — IP-backed lending is, even to sympathetic legal readers, a concept in search of a mechanism.³⁸⁰ And the whole edifice rests on the foundation Chapter 6 documented: a training-data property regime not merely undecided but *contested*, 88 per cent of consultation respondents backing full licensing against a government preference that drew 3 — which means the legal status of the very collateral the lending schemes would secure against cannot currently be priced by anyone.³⁸¹ You cannot build duration instruments on a property regime the public is still fighting the government over.

The industry has solved instrument problems before — completion bonds made an uninsurable business bankable, slates made heavy-tailed demand investable, presales converted unknowable box office into contracted cash flow, and television's vertical market is right now building priced-discovery greenlighting from spare parts. The next generation is visible in outline: staged real-options commitments for long formats; AI-era completion underwriting that can price generative pipelines, vendor mortality and announcement-stage repudiation; provenance as collateral, where the Chapter 13 verification stack meets IP-backed lending and the Authenticity Premium becomes a balance-sheet line; and public co-investment explicitly mandated to hold the middle until private books can mark it again.

The caveats, before the corollary, because the book should not over-claim and a chapter accusing the market of unfalsifiable nerves had better state its own falsifiers. This argument fails if genuine, non-fraudulent consumption of fully synthetic content breaks decisively above its current sub-one-per-cent asymptote and climbs — watch the stream-share number, never the upload number. It fails if the youth scepticism gradient of Chapter 13 reverses as the models improve, because that gradient is the premium's demographic engine. It fails if franchise premiums *compress* through the flood, which would mean attention is not pooling at the lighthouses after all. And it fails, in its immersive limb, if world-model economics turn out to serve robotics and simulation while leaving venue content costs untouched — watch whether the splat pipelines actually ship inside experiences. Four dials. As of this writing, none is moving the bears' way.

The corollary

The question from those three rooms — *won't AI have taken over by the time this ships?* — turns out to have four answers, none of which is yes, and none of which is the reassurance the question was fishing for either.

In games, the honest answer is: by then, the only balance sheets left holding your kind of bet may be sovereign ones — so know whose patience you are borrowing, and know that your audience's play may be worth more to capital than your game. In film: by then the machinery, not the audience, will have decided what got made — so the fight that matters is over underwriting criteria, not opinions, and a prize is not a financing system. In television: by then the market will have quietly rebuilt greenlighting as staged discovery — the format war is a duration war wearing a content costume. And in immersive: by then the same technology everyone else is fleeing will have made your sector's oldest constraint optional — the taking-over, there, is the opportunity.

Capital has not concluded that AI wins. It has concluded that it cannot know, and it prices *cannot know* as *will end* — everywhere except where a moat, a sovereign, or a shortened cycle lets it stop pretending to know at all. The most expensive position in

a transition is not the wrong bet. It is the unpriced one. The sectors, institutions and countries that answer the Seven-Year Problem with instruments will finance the next thirteen-year bet — and will get to decide, in the Chapter 16 sense of agency, attribution, access and audience, what kind of creative economy that bet is made inside. The ones that answer it with vibes, bullish or bearish, will discover they have outsourced that decision to whoever was left holding the term sheets.

One more thing belongs at the end of this chapter, because it is the bridge to the next one. The immersive inversion — the single place on the map where the discount changes sign — is not an accident of venue economics. It is capital noticing, earlier than most of the culture, the deepest technical shift this book tracks: that the machines have stopped making *pictures* and started making *places*. The biggest creative-technology cheques of 2026 are not chasing a better video model. They are chasing the thing underneath video — and to see why, you have to look at a release that produced no viral clips, changed no news cycle for a single day, and may matter more than everything that did.

* * *

Theoretical apparatus from the standard literature: Frank Knight, “Risk, Uncertainty and Profit” (1921); Avinash Dixit and Robert Pindyck, “Investment Under Uncertainty” (1994); Arthur De Vany, “Hollywood Economics” (2004); William Goldman, “Adventures in the Screen Trade” (1983); Mark Carney, “Breaking the Tragedy of the Horizon” (2015); William Baumol on cost disease (1966).

Chapter 9 — Worlds, Not Pictures

If you had asked me, in the autumn of 2025, what the most important AI release of the year was going to be, I would have said something obvious — Sora 2, or Veo 3.1, or one of the music models, or one of the editor-class tools like Runway Gen-4.5 or Adobe’s new Firefly. Something that turned a prompt into a thing you could put on a screen.

Six months later, I don’t think I would say any of those.

I think the most important release of the period this book covers was something that almost nobody outside of a relatively small community of practitioners noticed at the time, that produced no viral videos, that did not change the news cycle for a single day, and that I have come, over the course of the winter, to think of as the *actual* future of creative work: the public launch of **Marble**, by Fei-Fei Li’s company **World Labs**, in November 2025.³⁸²

Marble doesn’t make videos. Marble makes worlds.

I want to spend this chapter on why that distinction is, in my view, the most strategically important one in creative AI right now, and why almost everything else in the toolchain — from generative video to AI music to the digital-human work in advertising — eventually has to be re-thought in its shadow.

What a world model actually is

The phrase “world model” sounds like a marketing term. It isn’t, exactly. It is a category of AI system that researchers have been chasing for the better part of a decade, and that — as of late 2025 — finally started shipping as production-ready software.

A generative *video* model takes a prompt and produces a sequence of frames. The frames are coherent because the model has learned the statistical regularities of video: things move smoothly, light behaves more or less correctly, faces stay faces. But the output is *flat*. It is a particular sequence of pixels. You cannot navigate it. You cannot move the camera. You cannot pick up the lamp on the table and look at the wall behind it.

A *world* model takes a prompt — or an image, or a video, or a rough 3D layout — and produces a *navigable three-dimensional environment*. You can move through it. You can change the camera angle. You can, depending on the model, walk around the table, look at the wall behind the lamp, and find that the wall continues to exist in a consistent way that the model didn’t have to generate for you because it understood, structurally, that walls have backs.

The technical core, in the most common implementation, is **Gaussian splatting** — a representation where a scene is stored as a cloud of millions of tiny semi-transparent ellipsoids, each carrying colour and position information. The whole scene can be rendered in real time from any angle, because the system isn’t drawing

2D pixels; it is rendering a structured 3D world. The output, in turn, can be exported as a splat file, as a mesh, or — if you want — as a flat video.³⁸³

This is the part that took me, even as a working creative technologist, embarrassingly long to fully understand. Video is a *projection* of a world. A world is the more fundamental object. For two and a half years, the public-facing AI conversation has been about generating better projections. The actual capability landscape has been moving, in parallel, towards generating the worlds themselves.

When the worlds become cheap to generate, the projections — the videos, the images, the renders — become *outputs of the worlds*, not the primary medium. The whole production stack inverts.

Marble, from the inside

DreamLab — the studio I run in the North West of the UK — has been a beta participant in Marble since October 2025, in the months before its public release.³⁸⁴ I want to share what that experience actually felt like, because the *technical* description of a world model and the *practitioner's* experience of using one are different in ways that matter for understanding what is happening to the toolchain.

Imagine, for the sake of example, that I am working on a client project that needs a scene of a market square at dusk in a Mediterranean town. In the old pipeline — which is to say, the pipeline of 2024 and most of 2025 — that brief would translate into something like the following:

A concept artist would produce a moodboard. A 3D artist or a virtual-production house would build a CGI version of the square, populated with assets either bought from a marketplace or modelled bespoke. Lighting would be set up in Unreal or Maya. The whole scene would be rendered out as a video plate or used as a backdrop on an LED volume. If any change was required — *can we move the camera left a bit, see what's on that side* — the rebuild was non-trivial.

In Marble, the same brief unfolds differently. I type, or paste in a reference image, or upload a quick phone-shot panorama of an actual market square I visited last year. Marble generates a complete, navigable 3D environment of that square. It exists, persistently, as a file on my account. I can move my virtual camera anywhere in it. I can hand it to a director and say *walk through this and tell me where the camera lives*. I can export the result as a Gaussian splat, drop it into Unreal Engine via SuperSplat or one of the other Gaussian-splat editors,³⁸⁵ and use it as the lit backdrop for an LED-volume shoot. I can also, if I just want a plate, render a flat video from a chosen camera move.

The economic implication is this: the cost of having “a place” — a navigable, lit, persistent environment with depth — has dropped, in twelve months, by something between one and two orders of magnitude. The thing that used to require a four-

person team and a fortnight now requires a prompt and the time it takes a model to render.

This is not a marginal improvement to virtual production. It is a *category* change. The bottleneck of virtual production has, for the entire history of the discipline, been the cost and time of building the environment. When that bottleneck goes, what remains is exactly the human craft that the audience is paying for: blocking, performance, direction, lighting design, story.

In **Issue 8** of the newsletter, I noted that Sony Pictures had begun using Marble inside its virtual-production pipeline. The number the team quoted publicly was the one that should have made the front page of every trade publication: *40× faster than the traditional workflow*.³⁸⁶ If you sit inside the legacy economics of a virtual-production house — where building a single environment is a six-figure, multi-week proposition — that number is not an improvement. It is a *re-platforming* of the discipline. In **Issue 12**, Disney showed off a “300,000 poses in an instant” demonstration that was conceptually similar — animation built on top of generative spatial infrastructure rather than against it.³⁸⁷ In **Issue 27**, Netflix and Eyeline released **Vista4D**, a system that converts live-action footage into navigable 4D point clouds.³⁸⁸ The pattern is the same across the studios: a quiet pipeline shift, not a marketing story, that takes the entire “building the environment” stage out of the critical path of production.

The world-model race

Marble was the first commercial product in this category, but it was not the only one. The autumn of 2025 and the spring of 2026 were essentially a foot-race between research labs to ship usable world models, and the pace of releases was so rapid that the *Dream Machine* readers’ WhatsApp group routinely had three or four new ones to discuss per week.

Google DeepMind’s Genie 3, named by *Time* as one of the best inventions of 2025, generated playable 3D worlds at 24 frames per second from text prompts, with consistency held for several minutes — and in January 2026 was made publicly available to Google AI Ultra subscribers in the U.S. through a prototype web app called Project Genie. At Google I/O 2026, **Project Genie** was extended with a **Street View** integration that lets users generate navigable simulations of real-world locations directly from Street View map data, collapsing the gap between *the world that exists* and *the world that can be generated*.³⁸⁹ **Meta** announced **WorldGen** in November 2025, framed as research that could generate walkable 3D worlds from prompts like “*medieval village town square*.”³⁹⁰ **Tencent** open-sourced **HY World 1.5**, a real-time world model framework, in December 2025, alongside the **Hunyuan 3D Studio** which integrated the company’s art-grade 3D generative

model **3D-PolyGen 1.5**.³⁹¹ **SpAltial** launched **ECHO**, a spatial foundation model, in December 2025.³⁹² Stanford AI Lab and others released **Wonderzoom** in January 2026, a multi-scale 3D world-generation model that let you “infinitely zoom into the details” of a generated environment.³⁹³ **OpenArt** launched its own world-generation product, **Worlds**, in March 2026.³⁹⁴

The May 2026 wave was the most aggressive yet. **NVIDIA** released **SANA-WM**, a 2.6B-parameter open-source world model natively trained for 60-second video generation with explicit camera control — the first open-weight world model at meaningful scale, and a development whose long-term implications for the open-source-AI-tooling argument I make in Chapter 17 are, in my view, substantial.³⁹⁵ **Odyssey** released **Starchild-1**, which it described as “*the first ever real-time multimodal world model*” — a system that doesn’t just generate a world but understands and simulates it.³⁹⁶ A month later, Odyssey closed a **\$310 million Series B** explicitly framed as capital to accelerate its world models — which is, for anyone still wondering whether the spatial thesis is a niche enthusiasm, the sound of the capital markets pricing it in.³⁹⁷ **Apple** published **Headsup**, a large-scale, high-quality 3D Gaussian-head reconstruction pipeline built from multi-view captures of the kind a consumer iPhone can already produce — a continuation of the Apple-Personas-and-Gaussian-splat thread above.³⁹⁸ At the consumer end of the same wave, **WorldLens VR** rolled out an AI-powered Quest feature that adds subtle 3D depth to ordinary Google Street View environments, making the existing planetary-scale street-imagery dataset navigable in VR.³⁹⁹

The most ambitious of all of these — and the one I think hints most clearly at where the category is going — was **Luma AI’s UNI-1**, launched in March 2026 with the framing: “*When worlds become instant, the race shifts to better thinking.*”⁴⁰⁰ UNI-1 was the first commercial release I am aware of that *combined* world-model generation with what Luma called “reasoning” — that is, the model didn’t just generate a scene, it could plan, modify and iterate on the scene as a coherent agent. The pitch was that you would no longer have a fragmented pipeline of prompt → image → video → iterate; you would have a single unified creative system that thought before it created.

UNI-1 is, in my view, the most important *category* announcement of the spring of 2026, even if the product itself is still rough at the edges. It is the announcement that says: world models are not the end state. They are the *substrate* on which something else — reasoning-led generative creativity — gets built.

By **May 2026**, you could find world-model capabilities embedded in the consumer tools as well. **CapCut**, the consumer-grade video editing app, integrated ByteDance’s *Seedance 2.0* via the *Dreamina* product, giving phone-users the ability

to generate spatial scenes alongside flat video.⁴⁰¹ **Spark 2.0**, an open-source Gaussian-splat streaming framework, brought 100-million-splat scenes to web browsers at interactive frame rates.⁴⁰² **Apple** confirmed in October 2025 that its Personas feature on Vision Pro and other devices was powered by Gaussian splatting under the hood, making this — for the millions of Apple device owners who had used the feature without knowing what it was — the most-deployed Gaussian-splat technology in consumer hardware.⁴⁰³

The category, in eight months, went from a research demo to a consumer feature.

And then, over the summer of 2026, it went from a race to a crowd. **Adobe** — the company whose entire franchise has been the flat canvas — showed **Wonder**, a real-time interactive 3D world-navigation model out of Adobe Research, in July.⁴⁰⁴ **Tencent** followed its HY World line with **WorldClaw**, an agentic system for generating 3D open worlds *at scale* — worlds built by agents in bulk rather than prompted into existence one at a time.⁴⁰⁵ The smaller labs kept pace: Reactor's **LingBot-World 2.0** did real-time world generation,⁴⁰⁶ and **AlayaWorld** pushed long-horizon *playable* video worlds.⁴⁰⁷ **Wayve** — an autonomous-driving company — shipped **GAIA-4**, the latest of its world models, which is worth pausing on: the same category of system that generates a market square for a film shoot is being built, in parallel, to teach cars what streets are.⁴⁰⁸ The physical-AI end of that same thread showed up in August, when **Niantic Spatial** — the spatial-mapping company carved out of the *Pokémon Go* empire — began building the city of Rancho Cordova's *digital twin* as foundation infrastructure for physical AI: capture a real place once, and everything downstream, from robots to route-planning to broadcast, renders from the model rather than from the world.⁴⁰⁹ And in August, Lightricks released **LTX-2.5** as an open-source open-world model — the SANA-WM moment repeating itself, as open weights always eventually arrive downstream of the frontier.⁴¹⁰ I am keeping this deliberately compressed — the full inventory lives in [Chapter 17](#) — because the point is not any single model. The point is that in thirteen months the world model went from one company's strange conviction to the assumed next platform of an entire industry, with Adobe, Tencent, an autonomous-driving lab and the open-source community all arriving at the same destination from four different directions.

The splat pipeline underneath it matured in the same weeks. In July 2026 alone, **Gracia AI** brought *4D* Gaussian splats — splats with a time dimension; volumetric video you can walk around — into Unreal Engine with real-time relighting and self-shadowing; **Atlux** shipped the reverse trip, Unreal-to-4DGS; and **Notch**, a staple tool of the live-events and concert-visuals world, added native splat import and rendering.⁴¹¹ The capture side compressed just as dramatically. **InfiniSplat**, a

SIGGRAPH Asia system, turns a *single photograph* into a navigable 3DGS scene;⁴¹² and the **Insta360 X6** — a consumer 360 camera — shipped a feature called Spatial Capture: record a place in 360, open the app, and the place becomes a 3D model.⁴¹³ Read those two together and the direction is unmistakable. The world model’s input is becoming *anything* — one image, one walk-through with a camera you can buy on the high street — and its output is becoming a first-class citizen of every professional pipeline that used to speak only polygons.

The games industry, again

If world models are infrastructure, the industry that has been waiting for that infrastructure the longest is games.

The 2024 conversation in games about generative AI was, in significant part, about *flat assets* — concept art, textures, dialogue, music — and it was the conversation that produced most of the backlash. *Call of Duty: Black Ops 7*’s loading screens. *Anno 117*’s placeholder art “slipping through” the review process. *Fortnite*’s Chapter 8 controversy.⁴¹⁴ The audience response, in every case, was visceral, and the studios learned, the hard way, that AI-generated 2D assets dropped into established franchises read to fans as a cost-cutting move, not a creative one.

The 2025–26 conversation in games is different in kind, because the AI is now being aimed at the *substrate* of the game — the worlds, the systems, the NPCs, the procedural infrastructure — and the audience response is, so far, much more nuanced.

NVIDIA, in partnership with Stanford, released **NitroGen** in January 2026 — a “plays-any-game” AI trained on 40,000 hours of gameplay across more than 1,000 games. The model wasn’t being pitched as a way to *replace* games; it was being pitched as the foundation layer for a new generation of AI-aware game agents and procedural systems.⁴¹⁵ **Google DeepMind**’s **SIMA 2**, released in November 2025, was an agent that could play, reason and learn alongside humans in virtual 3D environments.⁴¹⁶ **Ubisoft** open-sourced its **CHORD** model in December 2025, for end-to-end PBR material generation, and ComfyUI nodes built on top of it within the same week.⁴¹⁷ **Ubisoft’s Teammates** — a voice AI tech demo first shown in November 2025 — promised a step-change in how NPCs would behave in next-generation titles. The team lead’s hands-on framing, given to *Video Games Chronicle*, is the one I keep returning to: “*It’s a tool first. We’ve been working on it for more than two years now, and our conclusion is that it’s a super cool tool, but it’s still a tool.*”⁴¹⁸ *Still a tool.* The whole AI-in-games debate, compressed into four words by the people inside Ubisoft who are actually building the thing.

The most interesting single release of the spring of 2026 was **YouTube’s Playables Builder**, a closed-beta product launched in December 2025 that lets

users create games with short text, video or image prompts, built on Gemini 3.⁴¹⁹ The framing, when YouTube’s product team described it publicly, was that *every YouTube creator* should have the ability to ship a playable game as easily as they currently ship a video. Within months, **Unity** announced an “AI Open Beta” — an in-editor AI suite that brought the same logic to the professional games-development pipeline.⁴²⁰

By June 2026 the prompt-to-playable category had stopped being a YouTube curiosity and become a product race. **The Sandbox** announced **The Sandbox Studio**, pitched as taking a creator *from idea to live multiplayer game in hours*;⁴²¹ a week later **Higgsfield** — the AI-video company whose commercial numbers I return to in Chapter 10 — introduced **Higgsfield Games**, a system that promises to create and deploy multiplayer games from a single prompt.⁴²² And at the **State of Unreal**, Epic showed where the professional engine layer is heading: **Unreal MCP**, LLM-powered tools and video-model integrations that let creators direct the engine through natural language. The debate that followed — would this hollow out the craft of Blueprints? — rather missed what I take to be Epic’s actual bet: not replacing developers but *expanding who can become one*, growing the whole creator ecosystem across games, film and immersive by lowering the technical barrier to entry.⁴²³

The platform that has been most explicit about the destination is **Roblox**. In June 2026 it acquired *three* AI research teams — **Morpheus AI**, **Dynamics Lab** and **Lucid AI** — in a single consolidating move aimed, in the company’s own framing, at solving the specific technical problems blocking photorealistic multiplayer worlds.⁴²⁴ Morpheus builds video world models using *Self Forcing*, a technique for converting offline video models into fast, autoregressive, interactive engines. Read the acquisition list as a sentence and it says: the largest user-generated-worlds platform on Earth believes the world model is its next rendering engine, and it is buying the research capacity to build it.

Where this lands, in 2027 and 2028, is the question I find the most strategically charged in the whole industry. If creating a playable, navigable world becomes a thing a YouTube creator can do in an afternoon, the boundary between *games* and *video* — which has been collapsing slowly for fifteen years, through platforms like Roblox and Fortnite and the proliferation of interactive content on social platforms — collapses fully. The next generation of creators will not think in terms of *making a video* or *making a game*. They will think in terms of *making a thing*, and the thing will, by default, be navigable.

The storyworld, live: the 2026 World Cup

In June 2026, while the tool announcements were stacking up, the thesis of this chapter received its first field test at genuinely planetary scale — and it came not from a studio, a lab or a games platform, but from football.

The 2026 FIFA World Cup — 48 teams, 104 matches, sixteen cities across three countries, an audience of six billion — is, on the evidence I assembled for the newsletter’s Deep Dive that month, the largest live deployment of AI in entertainment history.⁴²⁵ Before kicking a ball, every one of the tournament’s 1,248 players stood still for one second while a scanner built a photorealistic **digital twin** — a 3D avatar that then spent the following five weeks adjudicating offsidess, driving broadcast replays and populating immersive experiences. Combine the scanned avatars with a sensor-carrying match ball and computer-vision tracking across the stadium camera arrays, and something structurally new happens to the broadcast: the match is no longer merely *captured* by cameras. It exists as **a live data model that can be analysed, visualised and re-rendered from almost any perspective** — including perspectives no physical camera ever occupied.

Stop and notice what that is. It is the inversion this chapter has been describing, running live in front of most of the connected human species. The world — the match, the stadium, the twenty-two tracked bodies — is the fundamental object. The television picture is one projection of it. The automated offside decision is another. The 3D replay on the BBC is a third. And the tournament’s *second venue* is a fourth: the official FIFA World Cup experience on **Roblox**, the platform’s biggest sports activation to date, which by the tournament’s opening weeks had logged over **1.1 billion visits** — with real-time standings piping the actual tournament into the game world so the virtual carnival tracks the real one match by match.⁴²⁶ For a generation of Gen Z and Alpha fans, that voxel stadium, not the broadcast, is where the World Cup is happening.

Around the canonical event, the halo: TikTok as FIFA’s “preferred platform” with thirty Creator Correspondents, archival footage licensed to creators for co-creation, AI-personalised highlight packages assembled end-to-end for each fan’s declared teams and interests, machine translation and auto-clipping collapsing the gap between the final whistle and the posted edit. The World Cup is not an event with media coverage any more. It is a **storyworld** — one scarce, live, unscripted core wrapped in an infinitely expanding, AI-mediated halo of derivative worlds, each surface rendering the same underlying reality for a different audience, in a different grammar, at a different depth. That architecture is, I would bet, the template every major cultural property chases for the next decade — and the next generation of entertainment franchises will not bolt a transmedia strategy on after the fact; they will be designed as storyworlds from inception, built to be captured once and rendered everywhere. The officiating stack, the stadium digital twins, the creator

economy and the synthetic-rights precedents being set in the players' contracts each reward closer study, and the fuller treatment sits in DreamLab's published research on the tournament.

The physical venue layer, meanwhile, is moving the same way at its own pace: the same month, Sphere announced that *The Rocky Horror Picture Show* would be re-worked with "advanced technologies" into an immersive experience for the Las Vegas venue — a fifty-year-old flat film being promoted, in effect, into a place you stand inside.⁴²⁷

What this means for film

I want to come back to film for a moment, because I think the consequences of world models for the film industry are bigger than the consequences for any other sector, and the least understood.

For the entire history of cinema, the discipline has been organised around a fundamental scarcity: *the cost of building the location*. Even when the location was real — a city street, a forest, a beach — capturing it required a crew, a lighting team, transport, permits, weather contingencies. When it wasn't real — when it was a sound stage, or a digital matte painting, or a CGI environment — the cost was, if anything, higher.

The entire industrial structure of cinema, from the location department to the gaffer's crew to the virtual-production house, exists because *the place is expensive to make*.

When the place becomes cheap — when a Marble-generated environment, exported as a splat, dropped into Unreal, lit interactively, can substitute for a \$200,000-per-day exterior shoot at almost any quality bar a hero shot — the industrial structure that organised cinema starts to look like the manuscript-copying scriptorium did in 1450. The thing that was the bottleneck is no longer the bottleneck.

What replaces it? My best guess, six months into the transition, is *taste in places*. If everyone can generate a market square, the value of *choosing the right market square* — the one with the texture, the light, the cultural specificity, the lived-in-ness that makes a scene feel like it belongs to a real human story — becomes the new scarce skill. The location scout becomes the *world curator*. The production designer becomes the *spatial director*. The cinematographer becomes — even more than they already are — the person whose job is to find the *one camera move* in a near-infinite navigable space that tells the story.

This is, I think, an upgrade for the craft, not a downgrade. It moves the human contribution to the part of the work that humans actually do well — *judgement about what matters in a place* — and offloads the part of the work that has been a manufacturing problem for a hundred years.

The risk

I want to flag the risk too, because I am trying — and I am sure I will not always succeed — to be honest about the downsides.

If world models become the substrate of creative work, the *training data* for those models becomes a question of enormous cultural consequence. A world model trained on, say, the visual archive of Hollywood will generate scenes that look like Hollywood. A world model trained on the photographic archive of Mumbai will generate scenes that look like Mumbai. The *aesthetic monoculture* that the early image-generation models produced — that vaguely Pixar-flavoured, vaguely Marvel-flavoured, vaguely YouTube-thumbnail look that you can recognise in a thousand 2024 AI outputs at a glance — is at risk of being amplified, not reduced, when the medium moves from images to navigable spaces.

The companies that own the largest world-model training datasets in 2030 will, in a real sense, own the visual language of the next generation of cinema, games and immersive media. If those datasets are biased — towards English, towards the global North, towards Hollywood production design, towards the architectural and cultural visual vocabulary of a small number of wealthy cities — the entire interior life of the next generation of creative work will reflect those biases.

This is not a hypothetical. We are seeing it now. The publicly available world models, in mid-2026, do a startlingly good job of generating “Mediterranean market square” and “American suburb” and “Tokyo street at night.” They do a startlingly *thin* job of generating, say, “Lagos street at dusk during the rains” or “a contemporary Indigenous Australian community space” or “a Manchester terraced street in winter with the sodium lights coming on.” The bias is in the training, and the training is in the assets, and the assets were in the corpus, and the corpus was English-internet-skewed.

If we want the next creative economy to look like the world rather than like the AI companies’ biggest source datasets, the dataset question has to be a *first-order* design problem. Korin AI’s late-2025 launch — “trained with African datasets, built by Africans” — is the kind of intervention that is going to have to multiply.⁴²⁸ So is the African Tech / India / Singapore-led wave of culturally-specific AI cinema that the trade press started covering in Issues 20 through 27. Diversity in training datasets, for the world-model era, is not a content-moderation question. It is a *cultural infrastructure* question.

Six craft questions for the world-model era

Before I make the big claim, I want to put six craft questions on the page that working creatives — directors, designers, art directors, cinematographers, sound designers, level designers — will, by my estimate, be wrestling with for the rest of the decade. They are the world-model-era equivalents of the craft questions the

cinematographic era took fifty years to develop a vocabulary for (where do you put the camera, how do you light the scene, how does the cut work, how does the sound do its work). The world-model era has, in 2026, no settled vocabulary for any of them. The vocabulary will be built by the working creatives who notice the question first.

One. Where does the audience stand? The most-overlooked craft question of the navigable-space era. A film positions the camera; the camera positions the audience. A world model produces a *navigable space*; the question of where the audience *enters* the space, where they are *invited* to stand, what they are *encouraged* to look at, is no longer fixed by the cinematographer. It is fixed — if it is fixed at all — by the *narrative scaffolding* the orchestrator builds around the navigable space. Marble’s October-2025 update added explicit *suggested-camera-pose* primitives for exactly this reason. The craft question is which of those poses to specify and which to leave to the audience.

Two. How does the cut work in a navigable scene? The film cut depends on the audience being in a fixed position; the editor moves the camera between fixed positions in a way that the audience’s eye follows. The navigable scene has, by default, no cut. The audience moves through it continuously. The craft question — for working directors and editors — is when to *break* the continuity, how to do so in a way the audience reads as deliberate rather than as a technical glitch, and what new grammar of transitions a navigable medium permits. Some early experiments in 2025–26 have used *spatial discontinuities* (an audience walks through a door and emerges in a different space) and *temporal discontinuities* (the same space at different times) as cuts. None of these has yet stabilised into a shared grammar.

Three. How does performance survive the medium? A film performance is captured by a camera at a fixed angle and pace. A world-model performance — a synthetic actor performing inside a navigable scene — has to be authored such that the performance *works* from every angle and every speed at which the audience might encounter it. This is, for working performers and motion-capture supervisors, an entirely new craft challenge. The film-era cliché of the actor “playing to the camera” is, in the world-model era, replaced by *playing to the spatial neighbourhood* — knowing that the audience may be six feet away, may be inside the actor’s eyeline, may be behind the actor’s shoulder, may be looking at the actor from above. Volumetric capture (Vista4D’s live-action 4D reconstruction, NVIDIA’s D-Rex digital-human pipeline) is the technical answer. The *performance* answer — what acting *means* in a medium where there is no fourth wall — has not been worked out.

Four. What does sound design do in a navigable scene? A film sound mix is, for the most part, a fixed track timed to the picture cut. A navigable-scene sound mix has to *follow the audience*. Spatial-audio tooling (the SonicLab SPATAI pipeline,

Dolby Atmos for VR, the various Meta-and-Apple immersive-audio platforms) is the technical answer. The craft question is, again, what *good* spatial sound design looks like in a medium where the audience-author relationship has changed.

Five. What is the *running-time* of a navigable scene? A film has a fixed running time. A navigable scene does not. The audience could leave after thirty seconds or stay for two hours. The craft question for the working director is how to design the experience so that *both* extremes produce a satisfying piece of work. Games have, for fifty years, been grappling with this question — the *Dark Souls* answer (every player gets a different running time depending on skill and exploration) is different from the *Outer Wilds* answer (the running time is gated by narrative discovery) is different from the *Telltale Games* answer (the running time is broadly fixed across players). World-model cinema, in 2026, has not yet settled on its equivalent.

Six. What is the *single best moment* of a navigable scene? Film has scenes — discrete units of dramatic action with a recognisable shape, a recognisable peak, a recognisable end. A navigable scene, by default, does not. The craft question for the working director is whether to *design* the navigable scene around a single peak moment (which the audience may or may not reach) or to design it as a *texture* (which the audience experiences at whatever density their navigation produces). The peak-moment design pulls the medium back toward film conventions; the texture design pushes it toward something more like architecture or landscape design. Different working directors will, on the historical pattern, settle on different answers. The grammar will, over a decade, stabilise into a working vocabulary the way the cinematic-cut grammar stabilised between 1903 and 1925.

The six questions are not, in 2026, *theoretical* problems. They are the questions the working spatial-cinema teams I have talked to — the Wonderzoom group at Stanford, the World Labs developer cohort, the early adopters at Sony Pictures and Eyeline — are wrestling with on Wednesday afternoons. They are also, on the historical pattern of Chapter 2, the questions whose answers will define what working creatives in the next decade are *paid* to do. The directors and designers who develop a working vocabulary for them first will, on the available evidence, become the named *Walter Murchs* of the spatial-cinema era. The ones who wait for the vocabulary to settle will, in retrospect, look like the editors who waited too long to learn Avid.

The big claim

Let me make the big claim, and then move on.

I think — and this is the most non-obvious bet in this book — that the **world model is the medium of the next twenty years of creative work**, in the same way

that the **moving image** was the medium of the twentieth century and the **interactive screen** was the medium of the first quarter of the twenty-first.

I think people who are working in flat-video, flat-image, flat-audio formats in 2030 will increasingly be working in a *legacy* format — still alive, still culturally valuable, still where the highest-end of the craft lives, the way live theatre or vinyl-record production still lives — while the *dominant* mode of creative work will be the production, curation, performance and distribution of navigable spaces.

I think the studios, platforms and tool companies that are quietly investing in world models now — World Labs, DeepMind, Meta, NVIDIA, Tencent, Luma, Apple — will be the ones that set the rails for the next two decades.

I think the audience, having developed the antibodies described in Chapter 5 to slop-grade flat AI content, will eventually develop a parallel set of tastes for *navigable* content — and that the question of what makes a *good* AI world (rather than a *good* AI video) will be the central craft question of the late 2020s.

And I think — most importantly for the next chapter — that the toolchain to make all of this is being built, right now, by a small number of platform companies who have started saying out loud that AI is going to be *in everything, everywhere, all at once* — and who are, while you are reading this paragraph, designing the rails on which the next creative economy will run.

Chapter 10 — AI in Everything, Everywhere, All at Once

In late October 2025, at Adobe MAX, the company that has made the software almost everyone in the creative industries uses every day — Photoshop, Illustrator, Premiere, After Effects, InDesign — decided that the year-old marketing line “*AI is a feature in our tools*” had outlived its usefulness, and replaced it with a more honest one.

The new line was: “**AI in everything, everywhere, all at once.**”⁴²⁹

The reason I want to spend a chapter on that phrase is not because I love a slogan. The reason I want to spend a chapter on it is that I think it is, more than any other single piece of corporate positioning from the period this book covers, *literally* true. AI is in everything now. It is in every layer of the creative software stack. And the implications of that for working creatives — for the way we are trained, the way we are paid, the way we work with each other — are not yet, in the spring of 2026, fully understood.

This chapter is about the platform layer. About the companies that make the tools that the rest of the creative industries use to make the work. About how those companies have, in the past eight months, accepted that their business is no longer making *tools* but making *agents*, and about what that means for the rest of us.

The Adobe MAX week

The Adobe MAX 2025 keynote — held in mid-October in Los Angeles, the week after OpenAI’s DevDay, two weeks after Tilly Norwood — was unusual, by Adobe’s standards, in how much it tried to land at once.

The headline products were Firefly Foundry, a service for companies to train their own custom generative models on their own visual identity;⁴³⁰ Firefly Image Model 5, the latest generation of the image generator that has, since 2023, been Adobe’s primary public answer to Midjourney and Stable Diffusion;⁴³¹ and an AI Assistant built directly into Adobe Express, the company’s lower-barrier consumer creative tool.⁴³²

Underneath the headlines was a much longer list of “Project” announcements — Adobe’s research-preview format, the things that may or may not ship but that signal what the company is investing in. The list, looked at as a whole, is what convinced me, sitting at my desk in the North West watching the live stream, that something larger than a product launch was happening:

Project Scene It: image-to-3D and 3D-to-image technologies, with reference-image tagging for object preservation in 3D space.

Project Surface Swap: AI-powered material recognition, letting designers swap textures while preserving lighting, shading and perspective.

Project Turn Style: editing 2D objects as if they were 3D.

Project Trace Erase: removing objects *and* their shadows, reflections and environmental distortions in one operation.

Project New Depths: editing depth in an image as easily as adjusting brightness.

Project Frame Forward: applying changes across entire videos based on one annotated frame and a text prompt — “the precision of photo editing in video workflows.”

Project Motion Map: bringing static vector graphics to life automatically.

Project Sound Stager: analysing a video’s visuals, pacing and emotional tone, and automatically generating layered soundscapes.

Project Clean Take: AI correction of mispronunciations, voice isolation, noise removal and delivery refinement.

Project Graph: a node-based workflow editor, conceptually similar to ComfyUI, for chaining Adobe’s tools and models into custom pipelines.⁴³³

There is, in that list of ten projects, *every* layer of the post-production stack — image, video, 3D, audio, layout, workflow — being re-imagined as a generative or agentic operation. Not a tool with an AI feature stapled on. A *generative-first reimagining of the operation itself*.

The Adobe MAX week was, to put it plainly, Adobe’s announcement that it was rebuilding its product from the inside.

What “AI in everything” actually means

The reason I want to be careful with the Adobe-MAX framing is that, six months on, you can see how literally the company has executed against it.

In December 2025, Adobe announced that Photoshop, Express and Acrobat editing would be available *inside* ChatGPT — meaning the creative output was no longer happening inside Adobe’s interface, but inside an AI agent’s.⁴³⁴ In January 2026, the Premiere Object Mask tool — an AI-driven masking feature that automated one of the most laborious tasks in video editing — quietly became available to Premiere users.⁴³⁵ In late January, at Sundance, Adobe launched the *Adobe Film & TV Fund* and *Ignite Day*, with explicit support for filmmakers integrating AI into their workflows.⁴³⁶ In April 2026, at the **Adobe Summit**, the company introduced its **CX Enterprise** platform alongside NVIDIA — a stack of AI agents embedded across the entire content lifecycle from brief to delivery — under the framing “agentic creative intelligence is now.”⁴³⁷

And the migration of creative output *into a conversation* is not an Adobe peculiarity — it is the whole interface layer converging on the same shape. In July 2026 **Spotify** became, in effect, an AI chatbot too, folding music, audiobook and podcast discovery

behind a conversational front end;⁴³⁸ the same summer, **OpenAI** appeared to be building *chatbot-native ads* — advertising units that don't just sit in a feed but *launch an AI agent* when you tap them.⁴³⁹ The box you type or talk into is becoming the box you *create* in, *buy* in, and are *sold* to in. That is one more sense in which AI is not a feature bolted onto the tools; it is the layer the tools now run inside.

By June 2026 the agentic turn had reached the flagship applications themselves: the **Adobe Creative Agent and AI Assistant** went into beta across Premiere, Photoshop, Illustrator, InDesign and Frame.io — the MAX-week “Project” previews resolving, eight months on, into shipping agents inside the tools most working creatives open every morning.⁴⁴⁰ **Firefly Foundry**, the custom-model service that headlined MAX, completed the same journey from keynote to procurement line: Adobe productised it for enterprise brands wanting their own generative models trained on their own visual identity,⁴⁴¹ and — in the announcement that tells you where the enterprise tier of this market actually sits — **Disney's Imagineers** said publicly that they were adding Firefly Foundry to their worldbuilding toolkit.⁴⁴² And in late June, Adobe announced its intent to acquire **Topaz Labs**, the AI image- and video-enhancement company whose upscaling tools had become a quiet standard in independent post-production pipelines⁴⁴³ — one more independent toolmaker absorbed into the platform, exactly on the concentration pattern this chapter describes.

The Disney thread is worth pulling on a little further, because by August 2026 it had produced the clearest picture I have of where the platform layer goes next: *inside* the institutions. Disney rolled out **J.A.R.V.I.S.**, its internal AI system, to more than 2,000 Imagineers — putting, in effect, seven decades of accumulated creative knowledge, process and institutional expertise at every Imagineer's fingertips.⁴⁴⁴ The significance is not the tool; it is the substrate. The next competitive advantage for a creative organisation is not access to the foundation models everyone can rent. It is what the organisation *already knows* — its archives, its intellectual property, its workflows, its institutional memory — made searchable, navigable and operational through AI. Adobe sells the platform; Disney is building the proprietary knowledge layer on top of it. Every creative organisation of any scale is going to face the same build, and the ones that treat their own accumulated knowledge as the training substrate — rather than treating AI as a bolt-on service — are the ones the Disney example suggests will pull ahead.

The trajectory, in one sentence: Adobe in 2024 was a *creative tool company*. Adobe in 2026 is an *AI-agent platform company* that happens to also still ship Photoshop.

If you are wondering whether this transition has been smooth: it has not. The reception of the Adobe AI announcements among working creatives has been, in my

own circles and the readers' WhatsApp group the *Dream Machine* community runs, sharply ambivalent. There is real appreciation for the productivity gains. There is real anxiety about the implications for craft, for licensing, for control, for the trajectory of the company's relationship with the creators who pay for it.

What no working creative I know thinks is that this transition is reversible. Once Photoshop has an AI assistant baked in, once Premiere has Object Mask, once After Effects has the AI-powered animation tools that landed in November 2025,⁴⁴⁵ the next version of every Adobe product is going to have *more* of this, not less. Adobe's competitors are, if anything, going faster. If Adobe slows down, somebody else lands the punch.

This is — I think this is the part that working creatives have to understand and internalise — *the new physics of the toolchain*. AI is not a feature that one tool company decided to ship. It is a structural property of the toolchain itself in 2026, and the question for anyone using that toolchain is not whether to integrate AI but *how to integrate it deliberately*, with eyes open, on terms that preserve the human craft underneath.

The platform alliance

Adobe is not — and this is the more important observation — the only company doing this.

In March 2026, *Dream Machine* Issue 21 led with what I have called, in talks since, the most consequential business announcement of the year: **Adobe + NVIDIA** entered a strategic partnership that explicitly framed creative AI as *enterprise infrastructure* rather than viral consumer tooling.⁴⁴⁶ The partnership covered next-generation Firefly models, agentic creative-and-marketing workflows, and production-pipeline integration. The language was telling: *precision and control* for creativity and marketing pipelines, alongside content, campaign and production speed.

The reason this is consequential — beyond the size of the two companies involved — is that it signals the *maturation* of the market. Adobe + NVIDIA is not a race-to-the-cool-demo deal. It is a *race-to-the-procurement-line* deal. The two companies are betting, jointly, that the next era of creative AI is going to be won by whoever ships the most reliable, most controllable, most legally-defensible production-grade tooling to the enterprise creative buyers — the studios, the agencies, the broadcasters, the brand teams.

The same week, **Google** and **NVIDIA** announced a parallel deal for cloud-based generative-AI infrastructure aimed at the same enterprise market.⁴⁴⁷ **Hugging Face** and **Google Cloud** announced a partnership in November 2025 covering open-source agentic development.⁴⁴⁸ **Meta** and **Hugging Face** launched **OpenEnv** in

October 2025 to advance open-source agentic development.⁴⁴⁹ **Anthropic** signed a corporate-patronage deal with the **Blender Foundation** in May 2026.⁴⁵⁰ **Anthropic** also acquired into the **Slack** workplace-tooling ecosystem with Claude Apps in January 2026,⁴⁵¹ and reached an ad-sales partnership with **Spotify** to put music recommendations inside Claude.⁴⁵² In May 2026 **Splice** signed a “Responsible AI” deal with **ElevenLabs** covering sample-library training and consented voice synthesis;⁴⁵³ **Netflix** announced an agentic ad-tools roadmap whose internal framing — “*agentic AIs talking to each other*” — was an unusually candid description of where the advertising-orchestration layer is heading;⁴⁵⁴ and the AI-coworker startup **Viktor** raised \$75M to embed an agentic colleague directly into Slack and Teams,⁴⁵⁵ reinforcing the pattern that the platform-layer agents are landing where the working creative already lives.

The advertising holding companies were moving at the same pace. **WPP** signed a \$400m partnership with Google in October 2025.⁴⁵⁶ **WPP Open Pro**, a new edition of the agency’s AI marketing platform, launched the same month with a framing that should be read carefully by anyone working in adland: “*While some companies hide their AI behind service teams or focus on just one part of the journey, WPP Open Pro is an integrated solution for campaign implementation, built to deliver outcomes, not just assets.*”⁴⁵⁷ *Outcomes, not just assets.* That is the position of a holding company that has decided AI is not a feature — it is the entire reason a brand should buy from them in 2026. **WPP** then expanded its AI capabilities through a partnership with **Sightly** in November 2025.⁴⁵⁸ By April 2026, WPP was using Google Earth’s AI tools to map consumer journeys at scale.⁴⁵⁹

The jockeying in that advertising layer has continued through the period. In June 2026, Forrester reported that **Google had dethroned OpenAI as agencies’ preferred AI partner**⁴⁶⁰ — the same week that OpenAI, per *Digiday*, moved to automate ad creative itself,⁴⁶¹ a reminder that the platform companies compete for the agencies’ business with one hand while building the tooling that could disintermediate the agencies with the other. Even **Apple** — the last of the consumer-platform giants to hold back — finally arrived, announcing Siri AI and its next generation of Apple Intelligence at WWDC in June 2026.⁴⁶²

The pattern is unmistakable. The platform layer — the toolmakers, the infrastructure companies, the agencies, the cloud providers — has been quietly consolidating around a small number of strategic alliances that, taken together, are deciding the *rails* on which creative work will run for the next decade.

If you are a working creative reading this, you are probably already running some part of your workflow on rails laid by one of these alliances. By 2028, you will,

almost certainly, be running *most* of your workflow on those rails — or on a deliberate, principled alternative that has chosen to opt out.

Google I/O 2026: the second platform-layer moment

If Adobe MAX 2025 was the platform-layer announcement of the first half of this book, **Google I/O 2026** — held in the week this book went to press — was the announcement that closed it. The two events bookend the eight-month window the manuscript covers, and the symmetry of their framings is, on the platform-economics read, instructive.

The headline announcements were **Gemini Omni**, a unified multimodal model designed to work across text, image, audio, video and live interaction in a single workflow; **Antigravity**, Google's agentic coding and development environment; **Google Flow**, the agent-based workflow product that allows AI systems to take on multi-step creative and production tasks autonomously; **Gemini Spark**, the new developer toolkit aimed at building autonomous agents and AI-powered applications; and **Project Genie + Street View**, an integration that allows users to generate navigable simulations of real-world locations from the Street View map data — a topic I return to at length in [Chapter 9](#).⁴⁶³ The keynote opened, deliberately, with a browser-based multiplayer demo called **Infinite Scaler** — thousands of players competing inside vertically generated levels created on the fly from player prompts — a piece of theatre that, like *Tilly Norwood* a year earlier, was less interesting for the product than for the framing it imposed on the event that followed: AI-generated worlds, live procedural experiences, mass participatory systems that evolve in real time.⁴⁶⁴ **SynthID**, Google's content-provenance watermarking technology, was announced as having marked over 100 billion items by May 2026, and as being extended to partner ecosystems including OpenAI, ElevenLabs and Kakao — a development I cover in detail in [Chapter 13](#).⁴⁶⁵

The reason to draw the Adobe MAX / Google I/O parallel directly is that the *structural shape* of the two announcement weeks was identical. Both keynotes argued that AI was no longer a feature to be added to existing products; it was the *operating layer* into which the existing products would be re-built. Both keynotes pre-positioned the company's product roadmap around *agents* rather than tools. Both keynotes were aimed not at the consumer-keynote crowd but at the procurement teams of the enterprise creative buyers. The fact that the same framing arrived from the two largest creative-software and creative-platform companies in the world, eight months apart, in the same eight-month window, is the cleanest single confirmation I have that *the AI-in-everything thesis is the platform layer's settled commercial strategy* for the rest of the decade.

For working creatives, the operational implication is direct. The platform layer has decided. The question is no longer whether the creative software stack will be re-

built around AI agents. It is which agents, on whose terms, with what provenance, on which commercial settlement.

The new entrants

Underneath the platform giants, a separate layer of companies has been building the *consumer-facing AI creative tools* that, in some markets, are turning into bigger businesses faster than anyone expected.

Higgsfield, the AI video startup focused on social-media marketers, raised \$80m at a \$1.3bn valuation in January 2026.⁴⁶⁶ Three months later — in a stat that I have read repeatedly to check that I have not got it wrong — Higgsfield was reported to have earned \$200m in nine months of operations.⁴⁶⁷ An AI-video startup, less than two years old, was running at a quarter-billion-dollar annual run-rate by the spring of 2026.

Synthesisia, the U.K.-based AI-avatar platform, hit a \$4bn valuation in January 2026 and let its employees cash in.⁴⁶⁸ In October 2025 it had reportedly *rejected* a \$3bn acquisition offer from Adobe — choosing to remain independent.⁴⁶⁹

ElevenLabs, the audio-AI company, was reported to have crossed \$500m in annualised revenue by April 2026, raising from BlackRock, NVIDIA, Jamie Foxx and Eva Longoria.⁴⁷⁰

Runway released Gen-4.5 in December 2025 and Gen-4.5 Image-to-Video in January 2026, then a “Workflows” product across all paid plans, then a Story Panels app, then a Characters API, then Apps for Advertising — and by spring 2026 was making the public case that AI could enable “50 indie films” instead of “one \$100M blockbuster.”⁴⁷¹ In May 2026 the company opened a Tokyo office on a \$40M commitment, marking its first material expansion into the Asia-Pacific creative-AI market.⁴⁷²

Krea, Freepik, Magnific, Heygen, Hedra, Cascadeur, Hunyuan, Kling, Suno, Udio, Mureka, Hitem3D, Meshy, Rodin — the list of consumer-grade AI creative tools that crossed material commercial scale in this period is too long to fully enumerate, and the *Dream Machine* archive carries them week by week.⁴⁷³ The category that didn’t exist in 2023 is now an industry with multiple unicorns, multiple billion-dollar valuations and meaningful real revenue.

ComfyUI, the open-source node-based workflow tool that has become a quiet standard for technical AI users, raised \$17m in October 2025⁴⁷⁴ and hit a \$500m valuation by May 2026.⁴⁷⁵ What the ComfyUI valuation tells you, more than any of the big-platform numbers, is that the market is also paying — at significant scale — for tools that give *creators control* over the AI process rather than abstracting it

away. By the summer of 2026, that observation had grown into something bigger — big enough that it deserves its own section, because it is the strongest counter-evidence this book has to offer against the platform-concentration story the rest of this chapter tells.

The ComfyUI counterweight

If this chapter so far reads as concentration all the way down — Adobe absorbing the independents, the alliances laying the rails, the natural-monopoly flywheel spinning — I have written it that way because that is what most of the evidence shows. But the period this book covers also produced one large, awkward, genuinely hopeful counter-example, and I want to give it proper weight.

ComfyUI began in January 2023 as one developer’s open-source side project — comfyanonymous, then in the orbit of Stability AI — built to give people real control over diffusion models instead of a single prompt box. Three and a half years on, it has crossed **four million users**, drawn more than 89,000 GitHub stars and **60,000+ community-built nodes**, and become the operating layer where a substantial share of serious generative image and video work now happens.⁴⁷⁶ My own route in was earlier, and more personal: I first came across it through Javi, a developer I was working closely with at the time — and the node-graph thinking under the bonnet wasn’t new to us, because back in 2022–23 I was part of a team already experimenting with litegraph.js, the browser-based node-graph engine ComfyUI’s interface is built on. What felt then like a niche way of wiring tools together has quietly become how much of the industry builds.

And not just in theory — in shipped work. The ad studio Silverside AI built SVEDKA’s 2026 Super Bowl spot on ComfyUI, billed as the first primarily AI-generated Super Bowl ad. Series Entertainment rebuilt its entire game-and-video pipeline on it, scaling past 100,000 assets and shipping Netflix games. Comfy Org says **Netflix, Apple and Ubisoft are paying customers**. And the model labs have fallen in line: Flux, Stable Diffusion, Alibaba’s Wan, Tencent’s Hunyuan and ByteDance’s Seedance now ship a ComfyUI workflow as standard documentation — ComfyUI has become the integration target for new models the way Hugging Face became the target for weights.⁴⁷⁷

The reason is the one this chapter’s closing section circles: *control*. Where the closed platforms hide the process behind a prompt, ComfyUI exposes the whole pipeline as a visual graph of nodes. You can chain models, wire in controls, and turn a lucky one-off into a repeatable production system — the workflow itself becomes an asset a team can version, share and reason about. And as the foundation models improve, that control layer gets *more* valuable, not less: the bottleneck shifts from *can the model do it?* to *can I direct it precisely, every time?* Comfy’s chief executive put the strategy plainly — “we are not trying to win the next benchmark” — because the bet

is that whichever model wins, the layer that lets a human steer it is the thing that endures. It is picks-and-shovels economics for the creative economy, and it explains why the market has bifurcated: closed consumer apps like Midjourney on one side, and on the other a professional layer — VFX houses, ad agencies, game studios — that has consolidated almost entirely on the open graph.

In July 2026 the company pushed the thesis one layer further, launching **Comfy MCP** in public beta — pitched as the first Model Context Protocol server built for real production pipelines rather than one-off generations. Point an agent at it and you can build, edit and run workflows in plain language, search models and nodes, and — crucially — share a workflow URL so a colleague’s agent runs the *exact same* pipeline, every generation reproducible. The tagline says it plainly: *turn your agent into a creative technologist*. The shift is from AI that *makes* content to AI that can *operate* a creative system.⁴⁷⁸

And then the reversal — the detail that made me sit up when I wrote the newsletter issue this section draws on. Adobe’s Firefly AI Assistant, launched in April 2026, is an agent that orchestrates multi-step jobs across Photoshop, Premiere and Illustrator from plain language; Project Graph, previewed at MAX and described earlier in this chapter, is a node-based editor for wiring AI models, Adobe tools and effects into shareable “capsules.” Read those two products together: Adobe is building node graphs and a multi-model roster, and openly treating the connective infrastructure *between* its apps — not any single app — as where the value now sits. That is ComfyUI’s thesis, arriving in San Jose. The innovation didn’t come from the biggest company in the room. It came from an open ecosystem, and the biggest company in the room is now following it.

I don’t want to over-claim. ComfyUI is itself now a venture-backed company; open-source projects have been absorbed, enclosed and abandoned before, and the concentration forces described in the rest of this chapter have not gone away. But the structural fact stands: at the moment the platform layer consolidated, an open-source control layer became the professional standard — and that matters, not as ideology, but as *leverage*. Opening the black box gives creators a say over the tools, the models, and the values baked into them: room for models trained differently, transparent development, systems that answer to the communities using them rather than a single vendor’s roadmap. It is the strongest evidence yet that the next generation of creative studios may be built on open, programmable pipelines rather than closed software platforms — and that the people doing the creative work may yet get to keep their hands on the controls. It is also, not coincidentally, the practical foundation of the *insurance policy* argument I make in the platform-economics section below and at length in Chapter 17.

The free tier and the literacy tier

Two things happened in the consumer-platform layer that I think have been under-discussed and that matter a lot for what the next creative economy will look like.

The first is that the **base layer of AI capability went free**, in a meaningful sense, in the autumn of 2025. **Google** released its **Pomelli** marketing AI agent for free in October.⁴⁷⁹ **Google AI Studio**, **Opal** (the no-code AI mini-app builder), and the Project Genie prototype were all released as free or near-free tiers through early 2026.⁴⁸⁰ **Lovable** made its product free for teachers and students in classrooms.⁴⁸¹ **Adobe Express's** AI Assistant arrived inside the free tier of Adobe's already-free consumer product.⁴⁸² **Hugging Face** continued to expand its free hosting and open-source model distribution.⁴⁸³ **Krea**, **Freepik**, and many of the larger tool platforms kept generous free tiers as a customer-acquisition lever.

What this means, practically, is that the entry-level for AI-enabled creative work in 2026 is *near zero*. A teenager with a phone and a free Google account can, today, generate video, music, 3D objects and (with Project Genie) navigable interactive worlds at a quality bar that, two years ago, required a small production company to produce.

This is, in absolute terms, a democratisation. It is also — and this is the second thing — *creating a literacy gap* between the people who know how to use these tools well and the people who don't.

Adobe responded to this gap, in late 2025 and through 2026, by becoming — in addition to a software company — *a training organisation*. The Sundance partnership, with a \$2M investment to teach 100,000 filmmakers AI skills.⁴⁸⁴ The Ignite Day, focused on emerging creators.⁴⁸⁵ The Adobe Film & TV Fund. The Adobe Express AI Assistant tutorials. Google made the same bet in parallel, putting \$40bn into Anthropic in May 2026 in a deal widely interpreted as betting on the literacy and infrastructure layer of the next decade.⁴⁸⁶

The UK government picked the same direction. In January 2026, the Department for Science, Innovation and Technology announced *Free AI training for all*, expanding a government-and-industry programme to provide 10 million UK workers with AI skills by 2030.⁴⁸⁷ The Department for Business and Trade research, reported in *Dream Machine Issue 7*, found that *neurodiverse workers* were 25% more satisfied with AI assistants — suggesting that AI's productivity benefits in certain workflows could “potentially help to level the playing field.”⁴⁸⁸ The University of Wisconsin-Stout set AI-use as a baseline competency in its filmmaking course in January 2026.⁴⁸⁹

What the consumer-platform companies and the policy-makers are, jointly, building is a *training infrastructure* for the new toolchain. The reason they are doing this is straightforward: a tool you cannot use is a tool you do not buy, and a worker who cannot use the new tool is a worker who eventually exits the workforce. Both incentives push in the direction of mass AI literacy as a public investment.

What I find encouraging about this — and I am genuinely encouraged, against the grain of much of the cultural commentary — is that the literacy push is being framed, both by Adobe at Sundance and by the UK government, as *creator empowerment* rather than worker replacement. The proposition is not *learn AI or be replaced by it*. The proposition is *learn AI to remain in the driver's seat of your own work*. That framing matters. It is the right framing. It is the only framing under which the AI-literacy push doesn't become a way of accelerating the very problems it is supposed to fix.

The platform economics underneath the slogan

I want to spend a section on the *commercial* shape of the platform layer, because the “AI in everything” framing has economic implications that the keynotes have been careful not to name, and that working creatives buying platform access at scale need to understand.

The shape, simplified, is this. The dominant generative-AI platforms — OpenAI, Anthropic, Google DeepMind, Adobe, Runway, ElevenLabs and the rest — operate as *infrastructure-as-a-service* businesses on top of *capital-intensive* underlying compute. The marginal cost of producing one more generated image, song or video clip is, at the platform level, low. The fixed cost of the compute infrastructure required to produce *any* generated output at competitive quality is, at the platform level, very high — the data-centre build, the chip supply, the energy contract, the model-training spend.

This produces, structurally, a *natural-monopoly-tending* market. The platform with the largest compute base produces the lowest marginal cost per output, captures the largest user base, generates the largest revenue, and reinvests in a larger compute base. The flywheel is the standard cloud-services flywheel, accelerated by the AI-specific dynamics of training-data flywheels and user-feedback flywheels.

The 2025–26 financial telemetry, where the platform companies have disclosed it or been required to disclose it, supports the natural-monopoly read.

OpenAI was reported, through late 2025 and into 2026, to be operating at *significantly negative cash flow* at the unit-economics level despite its 800–900M weekly active users. The company's reported revenue, which crossed the \$10 billion annual run-rate mark in late 2025, was — by every analyst breakdown I have seen — being substantially exceeded by infrastructure costs (data-centre lease, chip supply,

energy, training compute). The Microsoft partnership at the financial level was, structurally, a *capital-supply* relationship rather than a *technology-licensing* one: Microsoft providing the compute capacity that OpenAI could not, by itself, finance.

Anthropic, in 2026, was reported to be operating with similar structural dynamics, with Google and Amazon as its capital-supply partners. The Anthropic Foundation patronage deal with the Blender Development Fund — announced in May 2026 and discussed in [Chapter 17](#) — is interesting precisely because it suggests Anthropic has, alongside the closed-platform business, *strategic interest* in supporting the open-source creative-AI infrastructure that the closed-platform model competes with. That kind of two-handed positioning is, in natural-monopoly markets, often the precursor to a *platform regulation* settlement.

Adobe, by contrast, operates with the most *defensible* business model in the creative-AI platform space, because it is selling AI as a feature of an existing subscription rather than as a per-use service. Firefly’s contribution to Adobe’s 2024 annual recurring revenue — 11% of new ARR, on the company’s own published numbers — is being generated *without* the per-token unit-economics problem that OpenAI and Anthropic face, because Adobe is bundling the AI into the existing \$54.99-a-month Creative Cloud all-apps subscription. The customer’s behaviour change from no-AI to AI doesn’t change the revenue line. It changes the *value capture* of the existing revenue line. This is, in business-school terms, the platform’s *strongest* possible defensive position. It is also the reason Adobe’s stock performed differently from the rest of the AI-platform cohort through 2025–26.

Runway, ElevenLabs and the AI-native specialist platforms operate with a per-use unit-economics structure that more closely resembles OpenAI’s. The differentiation, where they have it, is in *workflow integration* — Runway’s Workflows product, ElevenLabs’ Flows canvas, the studio-tier features that lock professional users into per-platform tooling. The strategic question for each of these companies, in the next two years, is whether they can build defensible workflow lock-in before the natural-monopoly dynamic of the underlying foundation-model market consolidates the foundation-model layer down to two or three players.

The implications for working creatives buying platform access at scale, in 2026 and beyond:

One, the per-token / per-output prices working creatives are paying for AI tooling in 2026 are, on every analyst read I have seen, *materially subsidised by platform-company investor capital*. The unit economics underneath the prices are not, today, sustainable at the volumes the platforms are producing. The prices are, by structural inference, going *up* over the medium term as the platforms work toward unit-economic break-even. The working creative who builds a business model assuming today’s per-token costs as a stable input is, on the platform-economics read, taking a

bet that the platforms cannot win. *Pricing today is not pricing forever*. This is the part of the platform-dependency argument the *open-the-black-box* discussion in Chapter 3 most directly relies on.

Two, the *strategic-rent-extraction* potential of the eventual platform-monopoly position is the structural risk underneath the entire orchestrator economy I described in Chapter 12. If two foundation-model platforms dominate the underlying generative-AI capacity by 2030, and every working creative's production pipeline depends on access to one or both of them, the platforms will be in the position the cable companies were in by 2010 and the social-media platforms were in by 2015 — able to extract value from the working creators who depend on them at prices the creators have no real ability to negotiate. The Petrillo-template solution to this is *collective bargaining* by working creatives and their unions against the platform companies as a class. The early architecture of this — the Cannes Disclosure Standard, the SAG-AFTRA platform negotiations, the EU AI Act enforcement, the UK 88% — is in place. The substance of it is, in mid-2026, still mostly aspirational.

Three, the *open-source alternative* layer documented in Chapter 17 is, on this structural read, *the working creative's principal long-term insurance policy* against platform-monopoly pricing. The 80% of YC and a16z startups now building on open-weight models — Hunyuan, Wan, Qwen, FLUX, DeepSeek, the various Mistral variants — is, in market-economics terms, *the credible-walk-to-alternative* that constrains the closed-platform companies' pricing power. The generative-video layer is the sharpest current illustration: by June 2026, *Forbes* was reporting that the Chinese labs had effectively won the video-model race that OpenAI had abandoned — and because so much of that Chinese-lab work ships on open or near-open terms, the video layer is precisely where the closed platforms' pricing power is weakest.⁴⁹⁰ The working creative who has familiarised themselves with open-weight tooling, even if they default to closed-platform tooling for most of their daily work, has a *strategic option* the working creative who has not has surrendered. The option is worth money. It is also, on the historical pattern, worth political leverage in the institutional negotiations that the next decade of platform-rule-writing will run on.

The platform-dependence risk, I should add, stopped being hypothetical in the middle of June 2026. As reported in Issue 34, the US government, citing national-security authorities, issued an export-control directive suspending all access to Anthropic's Fable 5 and Mythos 5 frontier models — leading Anthropic to suspend the tools.⁴⁹¹ I am wary of drawing large conclusions from a single episode whose full circumstances were still unclear as this book went to press. But the structural lesson holds regardless of the details: a frontier model your pipeline depends on can be switched off by decisions — corporate, regulatory, geopolitical — taken several layers above your head, with no notice and no appeal. One week the model is a line item in

your workflow; the next it is a line in an export-control directive. That is what dependence on the platform layer means in practice, and it is the most concrete argument I can offer for the open-weight insurance policy described above.

The “AI in everything” framing, in operational summary, is *the platform companies’ commercial strategy described in marketing language*. The strategy is to make AI a default productivity feature of every creative workflow, on platform-controlled tooling, at prices that the platforms can adjust over time once the workflow lock-in is in place. The strategy is, on the historical pattern of every previous platform-economics moment, going to produce a settlement somewhere between *the most-extractive version of the strategy* and *the most-constrained version of the strategy*. The 88%, the SAG-AFTRA contract, the EU AI Act, the C2PA standards body, the open-source ecosystem, and the working-creative collective-bargaining infrastructure are the constraints. The platform companies’ compute capital, distribution leverage, and product-design control are the extractive forces. The settlement will be wherever those forces balance.

What we lost

I want to close this chapter with the harder question, because the “AI in everything” framing has a cost that the platform-company keynotes are not, on the whole, eager to discuss.

What we lost, in the transition to AI-in-everything tools, is *the deliberate friction of the old creative process*. The thing that made Photoshop, for many of its early users, a profound creative tool was not just what it could do. It was that it required you to know it. The interface was a discipline. The keyboard shortcuts were a vocabulary. The layers, the masks, the channels, the curves, the colour pickers — they were the language of a craft, and learning the language was part of becoming the craftsperson.

When the layer of mastery moves from the toolchain to the prompt, the *barrier* of mastery drops to near zero. That is the democratisation we have been promised, and it is real.

What goes with the barrier, though, is the *depth of relationship* between the maker and the tool. The Photoshop user of 2015 knew the tool the way a guitar player knows a guitar — with their hands, with their body, with a relationship built up over years of repeated, embodied practice. The prompt-driven AI tool user of 2026 has a different relationship. It is more like the relationship of a director to a department head: you describe the result, the department head executes, you adjust by giving notes.

The motion designer Doug McGinness, posting on LinkedIn about the new AI-augmented After Effects workflow in late 2025, summarised the current state of the tooling in a single, ruefully accurate line that became a small private meme inside

my studio: “*export* → *prompt* → *pray* → *import*.”⁴⁹² The line is funny because it’s true. The current generation of AI-tooled creative work is, for a substantial portion of every day, an exercise in *committing to a black-box operation and accepting whatever comes back*. That is, structurally, a different kind of creative discipline than the deterministic-tool craft it is replacing.

Neither relationship is *better* than the other. They are different relationships, and they produce different kinds of practitioners. But the *transition* is real, and one of the consequences — which I have seen up close, watching young creatives come through the studio — is that the *cognitive engagement* with the medium is structurally less deep than it used to be. The maker is one further step removed from the material.

This is not, by itself, a tragedy. The cinema director is one step removed from the film stock and is still, recognisably, the author of the film. The composer is one step removed from the violin and is still, recognisably, the composer. The novelist who uses a word processor is one step removed from the page and is still, recognisably, a writer.

But it is a *change*, and it is one we are pretending not to notice. The new toolchain is not just faster than the old toolchain. It is also a different kind of relationship with the work, and the people who will be its best practitioners — the ones who will, in 2030 and 2035, be doing the AI-era equivalent of what Greg Lynn did with parametric architecture or what Bjork did with synthesisers — will be the people who *consciously cultivate* the depth of relationship that the toolchain no longer enforces.

The platform companies are not going to teach you to do this. They have no incentive to. They benefit from your dependency, not your mastery. The new toolchain is *frictionless*, and frictionless tools, however much we benefit from their efficiency, are not, on their own, going to produce the next generation of great creative work.

That work is going to come from the people who put the friction *back in*, deliberately, on their own terms — who treat the AI agent as a junior colleague rather than as an oracle, who insist on *understanding* what their tools are doing rather than just *using* them, and who maintain the cognitive engagement with the work that the tools have been designed to make optional.

In the next chapter, I want to talk about the people who are doing exactly that. The orchestrators.

Chapter 11 — What is newly possible

I want to spend a chapter, after nine chapters that have been mostly about *displacement*, on a question I think the book has so far under-served. *Which categories of creative work has the AI moment made possible that were not possible before?*

This is not the question working creatives in 2026 are most often being asked. The questions most often being asked are *will I lose my job?* and *should I use the tools?* and *what are studios doing?* Each of those is in this book, in a chapter of its own. They are all important questions.

The question of *newly-possible work* is the one that, on the historical pattern I drew in Chapter 2, almost always turns out to be the one that mattered most. Every previous creative-technology transition the book has documented produced a set of new categories of creative work that the displaced practitioners did not, and could not, see coming. The phonograph displaced amateur parlour music and created the recorded-music industry. The microphone displaced operatic vocal projection and created intimate vocal styles, jazz singing, pop crooning, audio storytelling. Non-linear editing displaced the splice and created the MTV cut, the music video as art form, the hyper-cut action grammar, the streaming-era serial-drama editorial rhythm. The smartphone-as-camera displaced the dedicated camera and created the entire grammar of vertical-video native form.

The pattern is that the new categories *always* appear, that they *always* appear faster than the displaced cohort predicts, and that they are *always* invisible from the perspective of the existing definition of the craft, because they are made of capabilities the existing craft does not have. The miniaturist could not, in 1845, predict Stieglitz's *Camera Work*. The session keyboard player could not, in 1982, predict Aphex Twin's *Selected Ambient Works* (and could not, if they had heard them, have recognised them as music). The print-magazine art director could not, in 2006, predict Instagram-as-fine-art-platform.

We are, in 2026, the equivalent generation to those people. The new categories are mostly invisible from where we sit. But some of them are already shipping; some of them are already finding audiences; some of them already have working creatives building careers inside them. This chapter is about what those categories are.

Before I start the inventory, I want to draw two historical analogues out — the synthesiser and non-linear editing — because they are the cleanest templates for *how a tool that begins as a faster version of the old thing ends up being the substrate of a new thing entirely*. The mistake almost everyone made about AI in 2024 and 2025 was to think of it as a faster way to make the kind of creative work the platforms had been making. The historical pattern says: that view is *almost always wrong on a five-to-ten year timeline*.

From imitator to instrument: the synth analogue

When Robert Moog first started selling modular synthesisers in the late 1960s, the cultural permission for the instrument was very narrow. The synth, in its first commercial moment, was understood as a way to *imitate* existing orchestral instruments — to play the parts of a string section, a brass section, a piano, an organ, in a form a single keyboard player could control. The breakthrough commercial release that secured the synth's cultural status — Wendy Carlos's *Switched-On Bach* in 1968 — was, on its face, a literal demonstration of this framing: the synth playing the music of the most-canonical European classical composer in the literature. Three Grammys. The first electronic record to be reviewed seriously by classical critics. The argument was: *the synth can do what an orchestra can do*.

The synth never did, in the end, only do that. The 1970s and 1980s did something nobody at the 1968 reviewing desks predicted. They produced *sounds that had never existed in the history of music* — Moog leads, FM electric pianos, the Roland TR-808 kick drum, the Yamaha DX7's chord pad, the *Blade Runner* CS-80 ambient texture, the Aphex Twin acid bassline — and they built entire musical genres around those new sounds. Hip-hop, electronic dance music, ambient, IDM, synth-pop, the entire sonic vocabulary of 1980s film scoring — these are forms that *could not have existed* without the synth, that did not exist before the synth, and that, *crucially*, could not have been predicted from the framing in which the synth was first introduced.

The synth was not a faster orchestra. The synth was an instrument for making sounds that no orchestra could produce, on a timescale that no orchestra could match, accessible to working musicians without the social capital of orchestral training. The instrument's first cultural moment — *Switched-On Bach* — was the moment of the *imitator framing*. Its second cultural moment — *Autobahn*, *Blade Runner*, *Thriller*, *Acid*, the Roland 808 in *Planet Rock* — was the moment when the imitator framing was thrown away and the instrument was used for what only it could do.

That second moment took about a decade. From Moog's first commercial modular in 1965, through Carlos in 1968, through Kraftwerk's *Autobahn* in 1974, through *Trans-Europe Express* in 1977, through the early hip-hop and dance records of 1981–82 — the gap between the synth-as-imitator and the synth-as-new-instrument was roughly fifteen years. The cultural permission to use the synth as itself, rather than as a replacement for something else, had to be earned by working musicians inhabiting the instrument in front of audiences who eventually understood that they were hearing something new.

The AI equivalent moment, on the historical pattern, has not yet arrived. We are, on the synth timeline, somewhere between *Switched-On Bach* and *Autobahn*. The work that is going to define AI as a new creative substrate — rather than as a faster way to

produce existing work — is being made *right now*, by working creatives somewhere, in forms that the trade press has not yet decided what to call. I have my candidates, which I will come to. The point I want to make at the chapter's opening is the structural one: every iterative-technology framing of AI is doing what the *Switched-On Bach* reviewers did in 1968. It is describing the new tool in the language of the old one. The new language has not yet been written.

From workflow to grammar: the non-linear-editing analogue

The same pattern is visible, in a different domain, with non-linear editing.

When Avid Media Composer shipped in 1989, the cultural framing was strictly utilitarian. NLE was a *faster way to do the existing thing*. Where a working editor used to splice physical film on a Moviola — a slow, irreversible, physically demanding craft — NLE allowed the same edits to be made in software, with undo, with multiple versions, with no consumable cost. The first generation of working editors who picked up Avid did so on the same premise as the synth's first-decade adopters: *the tool will let me do what I already do, faster*.

What NLE actually produced, by the mid-1990s and through the 2000s, was a fundamentally new editing *grammar*. The average shot length in mainstream Hollywood drama dropped from roughly ten seconds in the 1960s to roughly four seconds by the mid-2000s — a change made trivial by NLE that would have been physically punishing to execute on a Moviola. The MTV-cut aesthetic, which had been a music-video novelty in the early 1980s, became the default grammar of contemporary action cinema by the 2000s — *The Bourne Identity* (2002) is the textbook example, with shot lengths under two seconds across whole action sequences and an editing logic that depended on the viewer's now-trained ability to read a fast-cut grammar. The parallel-narrative structures of contemporary streaming drama — multi-thread, multi-timeline, multi-perspective storytelling, edited together with non-linear interleaving that would have been logistically impossible on tape — *Lost*, *Westworld*, *Dark*, *Severance* — are forms that NLE made possible.

Walter Murch, the most respected film editor of the last fifty years and one of the few people to have edited at the highest level on both film and Final Cut Pro, made the point clearly in *In the Blink of an Eye*: the tool *does* change the grammar. Murch was characteristically careful not to claim that the change was an improvement or a degradation. He claimed only that it was a *change* — that NLE permitted certain kinds of edit that physical-film editing could not, and that the new grammar would, over a generation, become as natural to its audiences as the slower-cut grammar of the 1960s had been to its.

The same dynamic is, in 2026, visible in the AI-augmented production pipeline. Working creatives I know are doing things in their day-to-day practice that would

have been physically impossible — not just expensive, *impossible* — in 2020. Iterating across forty variants of a scene in an afternoon. Producing personalised localised versions for ten markets simultaneously. Re-cutting a feature against a different aspect ratio for vertical-video distribution while keeping the principal photography intact. Generating a sustained 4D point-cloud reconstruction of a real-world location from a phone-captured walk-through and using it as a virtual-production plate. Running a scratch-vocal session in seventeen languages off a single take. These are *not faster versions of existing workflows*. They are workflows that did not exist three years ago. And, exactly as Murch predicted of NLE, the audiences for the work made on these workflows are already developing the perceptual literacy to read it.

Six categories of newly-possible work

With those two analogues in mind, I want to walk through six categories of creative work that I believe are *newly possible* in 2025–26 — meaning, work that an individual creative or a small studio can produce now that they could not have produced before, and that an audience can experience now in ways the audience could not have experienced before.

I will be honest, on each, about how much of the category is already shipping in finished form and how much is still in the *demo and beta* layer of the toolchain. The whole point of being inside the work is that you can see the difference.

I will also flag, throughout, the *binding constraint* that runs underneath the entire chapter: **human attention is finite**. This is the most-underdiscussed structural fact of the AI creative-economy moment, and I will come to it at length in a moment.

One: Remix at scale

The first category — the most visible and the most legally contested — is *remix*. The infrastructure of AI generation makes it cheap, fast, and at scale to produce derivative work: alternate-style versions of existing songs, cover-style reinterpretations across genres, AI-dubbed translations of feature films into languages the original release never reached, image-to-image style transfers of canonical artworks, motion-transferred re-performances of choreography across body types.

I want to be careful in describing this, because *remix as a creative form* has a long pre-AI history. Hip-hop's relationship to sampling is the canonical example; *Lessig's "remix culture"* framing from the 2000s identified the dynamic in broad strokes well before generative AI; the *Star Wars* fan-edit community, the mash-up era of Girl Talk and Danger Mouse's *Grey Album*, the YouTube AMV community, the TikTok stitch-and-duet grammar — every one of these is, in operational terms, *remix infrastructure that produces creative value through derivation*. The 2025–26 AI

moment didn't invent remix culture. It made the *production cost of derivative-but-original creative work* drop by more than an order of magnitude, and it shifted the technical bottleneck from *skill at imitation* to *judgement about what to imitate*.

The 2025–26 examples I have followed most closely:

- **AI-dubbed feature releases** crossing language barriers that the original production could not afford to cross. *Watch the Skies*, the Swedish UFO feature dubbed into English using ElevenLabs voice cloning, getting US distribution in October 2025, is the canonical example I covered in [Chapter 1](#). The dubbing is not, in any historical sense, a *new translation*; what is new is that the translation is rendered into the original actor's voice rather than a session voice actor's, preserving the performance signal across the language transition.
- **Cover-style reinterpretation at the long-tail of recorded music.** Suno and Udio's user economies have produced an extraordinary volume of *style-transferred* renditions of existing songs — folk versions of pop hits, metal versions of show tunes, lullaby versions of EDM tracks. The vast majority of these are uninteresting. A small fraction, in my experience, are *better than the original by some specific human listener's measure of better* — and that is the new commercial dynamic.
- **Film characters re-rendered as recording artists.** Eros Music Worlds, which launched in July 2026, turns film characters into *virtual singers* — lifting a character out of its original narrative IP and instantiating it as a performing musician with its own catalogue.⁴⁹³ This is remix in the strict sense: it takes a piece of existing creative property and derives a new work in a form the original never occupied. It is also, structurally, a preview of the licensing fight, because the character is somebody's IP even when the song is not.
- **Public-domain canon, re-rendered at feature length.** In July 2026 a full-length AI-generated film of Homer's *Odyssey* was released — the oldest narrative in the Western canon, rendered as feature cinema by a production apparatus that would not have existed two years earlier.⁴⁹⁴ I am not going to claim the result is a great film; on the evidence of the AI features I have watched, it almost certainly is not. What it is, structurally, is a demonstration that the entire public-domain literary canon — every epic, every myth cycle, every out-of-copyright novel — is now addressable by the remix category at feature length, with no rights negotiation whatsoever. The interesting work in this sub-category will come from the creative teams who treat the canon as raw material for genuine reinterpretation rather than as free IP for illustration.
- **The Andrii Daniels bomb-shelter clip** ([Chapter 13](#)) — *Deadpool and Harry Potter* as a Christmas mash-up, made in a Ukrainian bomb shelter, viral in December 2025 — is a remix-at-scale artefact in the strict sense. It crosses two IPs that the rights-holders would never have collaborated on, in a stylistic

register the original films could not have produced, and finds a global audience that responds to the *idea* of the remix as much as to its execution.

The legal layer of this category is still moving. *UMG v. Anthropic*, *Getty v. Stability AI*, the EU Copyright Directive's Article 17, the UK 88% — these are the institutional structures that will decide whether the AI-remix category becomes a *licensed and compensated* creative form (the Petrillo template applied to AI) or a *grey-market* one that operates outside the rights system. The historical pattern — Sampling, post-*Grand Upright*, became a licensed creative form, with the dense Bomb Squad style becoming commercially difficult but the basic technique surviving in a more clearance-friendly mode — says the remix category will, in some form, *settle into a licensed category by the end of the decade*. The Petrillo template, again, is the answer the system already knows.

Two: Mass personalisation, against the binding constraint

The second category — the one most often gestured at in platform-company keynotes and least well-served by them — is *mass personalisation*: creative work that is *individually different* for each viewer, listener or player.

The early shipped versions in 2025–26 are these:

- **Personalised playlist generation and AI DJs.** Spotify's DJ feature, YouTube Music's AI Hub, Apple Music's editorial-style algorithmic mixes — all of these are *individual-level curation* of pre-existing creative work. The personalisation is in the *selection*, not in the *generated content*.
- **Personalised game NPCs and dialogue.** Inworld, ConvAI, Ubisoft Teammates, Roblox's AI assistant — by spring 2026, the first wave of generative-NPC dialogue is shipping in player-facing form. The personalisation is in the *moment-by-moment response* to the player's actions, not in the underlying game world. Whether this actually produces a personalisation a player can recognise — versus producing the same in-character dialogue from a wider statistical envelope — is, on the published shipping examples I have played, an open question.
- **Personalised marketing creative.** WPP Open Pro, Adobe CX Enterprise, GenStudio — the platform infrastructure for producing per-segment, per-audience, per-context creative variations is now mature. Whether the audience experiences this as personalisation, or experiences it as a feeling of unease that the messaging knows too much, is the cultural question every brand creative is asking in 2026.
- **Personalised educational content.** Khan Academy's Khanmigo and Duolingo Max are the most-mature shipped examples I know of. The

personalisation here is genuine and measurable — the per-student tutoring response is different on the level of the individual learner.

I want to draw a sharp limit around the personalisation category, because the binding constraint runs straight through it and I think every working creative thinking about this market needs to internalise it.

Human attention is finite. Aggregate daily media-consumption time per person has, on the available Nielsen-style telemetry I have read, *not* grown meaningfully in the past decade. The total of every form of media consumption — TV, streaming video, music listening, podcast listening, social media, gaming, reading — is, per the published data, on the order of 11–12 hours per US adult per day, and that number has been roughly stable for years even as the *number of available* hours of content per day has exploded by orders of magnitude. The eye, the ear and the consciousness each have a finite capacity. *Personalisation* does not, by itself, expand that capacity. It changes the *distribution* of attention across content, but it does not increase the *total* attention available to be spent.

This is the structural ceiling on the *commercial* value of mass personalisation. Producing a personalised version of a film for every viewer — to take the most extravagant platform-keynote framing — is, on the binding-constraint reading, a competitive move that *reallocates* attention rather than *expanding* it. The category will be commercially meaningful in segments where reallocation can produce premium prices (luxury advertising, premium educational content, top-tier video-game NPCs in IP that justifies the investment). It will be less commercially meaningful at the long tail, where the personalisation effort does not produce attention-reallocation big enough to pay for the agentic infrastructure underneath it.

The trade-press framing of mass personalisation as *infinite content for infinite audiences* is, on the binding-constraint reading, structurally incoherent. The audience does not have infinite attention. The producible content is, by 2026, effectively infinite. The economic question is *which slices of the audience's existing finite attention budget* the personalised work can plausibly capture. The answer, by my read of the shipping evidence, is *narrower than the platform companies' enthusiasm suggests*.

I will come back to the finite-attention constraint at the end of the chapter, because it shapes every category that follows.

Three: Audience participation in the creator economy

The third category — and the one I am most personally interested in — is the *audience as participant*. The 2010s creator-economy framing was that the audience could *make their own content* on platforms (YouTube, TikTok, Instagram), creating a long tail of user-generated work alongside the professional content. The 2025–26

AI framing extends this by an order of magnitude: the audience can now *prompt, contribute to, remix and co-author* work in real time, often in collaboration with named professional creators.

The shipped examples I have followed:

- **TikTok / Instagram / Shorts remix mechanics** — duets, stitches, the *Sora* iOS app's TikTok-style remix grammar (which crossed a million downloads in five days in October 2025). The barrier between *watching the work* and *making a related piece of work* has dropped to a single tap.
- **Roblox's creator economy.** The platform's UGC ecosystem, by 2026, is the largest single audience-as-creator system in any creative medium. Roblox creator payouts run into the hundreds of millions of dollars annually. The line between the audience and the working creator on Roblox is, structurally, not a line; it is a gradient.
- **Disney+'s announced (though not yet shipped at the date I am writing this) user-generated-AI-content tools using Disney IP.** If this ships in the form the platform has described, it will be the first time a major IP holder has *deliberately given the audience tools to make work inside its canonical universe*. The cultural permission this would set — sanctioned fan-made *Star Wars*, sanctioned fan-made *Marvel* — would, on the historical pattern, expand the categorical *Star Wars* / *Marvel* canon at a rate the studio system itself cannot match.
- **The fan-prompt economy.** Civitai, Promptbase, the LoRA marketplaces — the infrastructure for *selling and sharing the creative inputs* that other creatives then use to produce work. The audience is participating not just in the work, but in the *toolchain underneath the work*.
- **Participatory art at national scale.** In August 2026 the National Portrait Gallery opened *A National Portrait* by Es Devlin, a work in which *every person in the UK aged 18 and over can become a co-author*.⁴⁹⁵ This is the audience-as-participant category executed not by a platform chasing engagement but by a national cultural institution treating co-authorship as the *form* of the work rather than a marketing bolt-on. It is the cleanest institutional signal I have seen that the participation category has crossed from the creator-economy fringe into the centre of the cultural establishment — the National Portrait Gallery is, definitionally, the keeper of the canonical British portrait, and it has just made the canonical portrait something the whole adult population helps make.

The single cleanest piece of evidence on the *generational* shape of this category came in May 2026, when **Snapchat** published research finding that **31% of 13–15 year-olds** on the platform were already using AI tools “*to be creative*” — not, importantly, to do their homework, not to chat with a synthetic friend, but

specifically to *make things* they then shared with their peers.⁴⁹⁶ That number is the most legible quantitative indicator I have seen of where the audience-as-participant category is heading. The 13-to-15 cohort the survey describes will be the 18-to-20 cohort of 2029. By the time they reach the working-creative entry pool, *making things with AI* will not, for them, be a category distinct from *making things*. It will simply be how things are made. The studios that build for this audience now — not as a future they are anticipating, but as a present they are already serving — will, on the historical pattern of every previous generational shift, set the terms on which the next decade of cultural production runs.

The thing I want to flag about this category — because I think it is the part the platform companies and the working studios have most systematically under-priced — is that *audience participation reverses the direction of the creator-audience economic relationship*. In the pre-platform creative economy, the audience paid the creator. In the platform-era creator economy, the audience generated the content and the platform monetised the attention. In the 2025–26 AI-augmented creator economy, the audience is increasingly co-producing the content *with* the creator, and the question of who gets paid for what is, structurally, harder to answer than it has ever been.

This is one of the open frontiers of working-creative business model design in this period. I do not think anyone has solved it. The studios that figure out how to credit, compensate and structurally honour audience contributions to the work — without turning the work into the kind of crowdsourced mush that does not survive the slop ceiling — will, in my view, have the most defensible business position in the next decade. The studios that try to extract audience-generated work without compensating it (the 2010s social-media platform model applied to AI) are, on the historical pattern, walking into the next *Viacom v. YouTube*-scale lawsuit.

Four: Fan fiction and fan-made content, legitimised at scale

Closely related to the participation category — but worth a section of its own — is the *fan-fiction* / *fan-content* category. Fan-made creative work has a long pre-AI history. The *Star Trek* fanzine era of the 1960s; the *Star Wars* fan-film tradition from the 1970s onward; the Archive of Our Own / Wattpad / FanFiction.net communities; cosplay; the anime fansubbing tradition; the *Harry Potter* fan-fiction archive that, on some counts, contains more words of *Harry Potter*-canon-derivative text than the original novels themselves.

What AI does to this category is two things. *One*, it raises the *technical floor* of fan-made work toward what was previously professional-grade: a fan can now generate a *Harry Potter* short film with production values that would have required a major-studio budget five years ago. *Two*, it makes the *canon-extension impulse* of fan

culture practically infinite: every reader, every viewer, every player can — with current tools — produce a new piece of work inside the universe they love, in their own voice, on their own terms, for their own audience.

The cultural and legal layer of this category is, in 2026, still moving. Disney's tolerance for fan content has shifted under AI, and the question of whether *AI-generated fan content using Disney IP* will be treated more leniently than fan-made content has historically been is one of the open IP-policy fights of the year. Lucasfilm's tolerance for *Star Wars* fan films has, historically, been generous; whether that extends to AI-generated *Star Wars* features is unsettled. The Marvel Comics community policy on AI is one of the documents to watch in the next eighteen months.

What is *not* unsettled is that the audience is already doing this. The fan-AI-content economy is, by spring 2026, larger by volume than the official IP-holder output for almost every major IP in popular culture. The official IP holders can suppress this, license it, build platforms around it, or watch it consume their cultural authority. There is no fourth option.

The working creative read on this, for anyone reading the book inside an IP-holding studio, is the one I made in Chapter 7's discussion of the legacy industries' strategic vulnerability. The studios that move *toward* sanctioned fan-AI content economies — Disney's announced UGC tools, the most generous of the cosplay-and-fan-film tolerances, the platform-and-fund models that pay fan creators for canonical contributions — will have a meaningful structural advantage over the studios that try to defend the closed canon against the audience that is, with or without permission, going to extend it anyway. The Petrillo template applied to fan culture is: pay the fan, structure the IP-holder's stake, *participate* in the extension rather than fighting it.

Five: Agentic support workers for the solo creative

The fifth category is the one that has — in my own practice and in the practice of the wider DreamLab community — produced the most dramatic and immediate productivity gains. The agentic-support-worker model.

In the pre-AI creative economy, the small or solo creative practitioner — the independent filmmaker, the songwriter, the freelance illustrator, the indie game developer, the YouTuber — operated with a particular structural disadvantage. Their work was bottlenecked not on creative judgement (which the senior practitioner had in abundance) but on the *production-coordination labour* that a larger studio would assign to junior staff: client communications, project management, asset organisation, scheduling, invoice processing, basic research, draft response, brief structuring, simple post-production. The senior practitioner spent an inordinate

fraction of their effective working hours on labour that did not require senior judgement.

The 2025–26 agentic toolchain — Notion AI, Adobe Express AI Assistant, Heygen Video Agent, Claude Apps, the personal-assistant features baked into the major platforms — has, in operational terms, *given the solo creative the first four hires they would have made*. The category keeps collapsing whole production departments into a single prompt: Heygen’s Video Podcast, launched in July 2026, turns *any document, link or idea into a two-host video show — with studio scenes, multi-cam cuts and B-roll, ready in minutes*.⁴⁹⁷ The point is not that the output is good (much of it will not be); the point is that a format which, in 2020, required a studio booking, two presenters, a camera operator and an editor is now a background agentic function the solo creative can run between other tasks. The personal assistant, the scheduler, the production coordinator, the asset manager — these are now, for working solo creatives in my circle, agentic functions running underneath the senior practitioner’s day, costing roughly the price of a couple of midrange platform subscriptions, and producing genuine recoverable hours.

The economic impact of this is, I think, the most under-priced shift of the period. The *one-person studio* that, in 2020, would have been a part-time freelance practice supporting two to four projects a year is, in 2026, a *full-business-class production operation* supporting twenty to forty projects across a wider range of disciplines. The Sienna-Rose / Xania-Monet / Hoyt-Dwyer single-creator AI-supported career — about which I have, in [Chapter 5](#), made my reservations clear in terms of *star formation* — is, on the *production-economics* side, a genuine new business form. Whether or not Xania Monet becomes a Billboard-defining cultural figure, the *operational machinery* underneath her — a single human creative, supported by an agentic stack producing music, marketing, distribution and merchandise at a scale that would have required a small label to support five years ago — is a working business model that did not exist before the toolchain shipped.

The implication for working creatives at the senior solo level is direct. The first four hires you would have made — production coordinator, junior researcher, scheduling and admin, post-production junior — are now agentic. The economic ceiling on your individual practice has, structurally, lifted. The constraint is no longer the throughput of the labour underneath you. The constraint is your own senior judgement bandwidth — the briefing, the taste, the integration, the *Why* — which the agents cannot replace and which, on the chess-grandmaster argument of [Chapter 16](#), has more commercial leverage in this market than it has had in any previous period.

Six: Hyperlocal and long-tail cultural production

The sixth category — and the one I think will, in retrospect, prove to be the most culturally significant — is *hyperlocal and long-tail cultural production*. The collapse of the cost of producing professional-grade creative work means that, for the first time in the history of mass media, *every linguistic community, every regional culture, every minority cultural tradition, every niche audience* can produce work in its own language, in its own visual style, for its own audience, at production values that compete with global commercial output.

The early shipped examples:

- **Korin AI**, the Africa-trained and Africa-built music platform launched in spring 2026, is the canonical example of a *deliberately culturally-specific* AI tool. The training data is African; the outputs are African; the platform exists to serve African creators rather than to backfill global pop with African aesthetic markers.
- **Indian AI cinema and the Trilok / Animaj / Filmax DinoGames cluster**. India's regional-language film industries — Bollywood, Tollywood, Kollywood and the dozen smaller regional industries — have been among the most aggressive adopters of generative AI in the period this book covers. The reason is structural: the regional film industries operate at production budgets where AI's cost-reduction effect has *the largest proportional impact*. A regional Indian feature that would have cost \$200,000 to produce in 2020 can, in 2026, be produced for a meaningful fraction of that, with comparable production values. The category of films that this brings into commercial viability — every regional cinema with a million-viewer audience but a six-figure budget ceiling — is, by my estimate, *an order of magnitude larger* than the category that was commercially viable before.
- **Wrinkles**, an app launched in August 2026 that turns *local history into an AI-generated audio guide*,⁴⁹⁸ is the small-scale, granular end of the same category: it makes it possible for a single neighbourhood, a single street, a single parish to have its history rendered as a produced audio experience without a production company ever being involved. This is hyperlocal cultural production at the most literal geographic resolution — the culture of a place, produced for the people of that place, at a cost that finally sits below the threshold at which such work was ever commercially worth making.
- **The Tunisian Lily** — winner of the \$1m Dubai AI Film Award in January 2026 — is the cleanest individual example of a film made *outside* the major-market financing system, at a production value competitive with it, by a creative team using AI tooling as the cost-reduction lever.

- **The 800-creator declaration’s signatories** include working creatives from a meaningful number of national contexts that the historical Anglophone film/music/games industries have under-represented. The *Stealing Our Work Is Not Innovation* coalition is, structurally, a coalition that includes the long tail of global creative work.

The implication of this category is the one I am most personally hopeful about, and I want to be honest that *hopeful* is the right word — there is a less-hopeful version of the same data, in which Anglophone AI tooling homogenises the global creative economy faster than the regional creative economies can develop their own infrastructure. The 2026 evidence I have is mixed enough that both outcomes are still on the table.

What is not in doubt is that the *production-cost ceiling that has, for a century, kept hyperlocal creative work below the threshold of professional commercial viability* is, by 2026, no longer the binding constraint. The next decade of cultural production will, I am confident, contain more local, more linguistically diverse, more culturally specific creative work than the previous decade — and the most culturally significant individual works of the next decade will, on the historical pattern, come from communities the existing global creative economy has under-served. The question is whether those communities own the tooling that produces them.

What is not yet possible

Before I close, I want to be honest about the categories that are *not yet shipping* in the form the platform-company keynotes have promised them, because the book’s credibility depends on accurately characterising the gap between affordance and rhetoric.

AI does not yet write a satisfying novel from scratch. The 2025–26 evidence on long-form prose generation is that AI systems produce technically competent prose that loses *narrative purpose* over the length of a novel. Working novelists I have talked to who have experimented seriously with this category all describe the same failure mode: the prose is fine; the *book is not a book*. This is consistent with the *House of David* “hand inside a puppet” critique of AI-augmented storytelling at feature length.

AI does not yet do live performance. No AI is touring in 2026. Xania Monet has not performed at Madison Square Garden. The cultural permission for a synthetic performer to share a stage with a live audience does not, as of this writing, exist in any meaningful form. The structural reason — the audience experience of *being in a room with another human consciousness* — is, on the slop-ceiling reading, not a permission gap; it is a category mismatch. AI work and live performance are, for now, different categories.

AI does not yet do sustained emotional storytelling at feature length without human authorial spine. The auteur-driven cinema of the Cameron / del Toro / Spielberg generation is, structurally, a category that depends on a *single human consciousness running through every creative decision in the film*. AI-augmented versions of this category exist; AI-native versions of it do not. The *Citizen Kane* of AI cinema has not been made. Whether it will be is, in my view, the single most interesting open creative question of the next decade.

AI struggles with cultural specificity that the model has not been trained on. The BBC India observation I referenced in [Chapter 14](#) — that AI screenplays produce *cultural-memory failures* when generating content for cultural traditions the training data has under-represented — is the binding limitation on AI's promise as a *globally-distributed creative tool*. The hyperlocal-cultural-production category I described above depends, structurally, on this limitation being addressed by purpose-built regional model infrastructure (Korin AI, the Indian regional-language tooling, the various Asian-built models). Until that infrastructure matures, AI cultural-production is still, by default, Anglophone and Western-modelled.

The binding constraint, restated

I want to close by returning to the structural fact that runs underneath every category of newly-possible work in this chapter.

Human attention is finite.

I am not making a romantic argument about it. I am making an economic one. The aggregate daily media-consumption time per adult, in any market the Nielsen-class telemetry covers, has been roughly stable for at least a decade — eleven or twelve hours a day, across all forms of media, all formats, all platforms, all devices. *That number has not grown with the rise of streaming. It has not grown with the rise of mobile. It has not grown with the rise of social media. It has reallocated, sometimes dramatically. It has not expanded.*

The producible content, meanwhile, has expanded by orders of magnitude. Deezer in 2026 receives 75,000 AI tracks a day; the listener has the same number of waking ear-hours she had in 2016. YouTube uploads run at hundreds of hours per minute; the viewer has, in net, the same daily screen-time budget. The Sora app produced a million downloads in five days; the audience for the work those million users are about to make has, in total, the same total attention they would have had if Sora had never shipped.

This is the binding constraint of the entire 2025–26 creative-AI moment. The cost of producing work has collapsed; the supply of attention has not. Every category I have described in this chapter — remix, personalisation, audience participation, fan content, agentic support, hyperlocal production — is being built into a market where

the *production* side of the supply-and-demand equation is racing toward infinity and the *consumption* side is bounded by the finite cognitive and biological capacity of the human nervous system.

This has three structural implications.

One. The competitive advantage in this market accrues, with iron consistency, to the producer of work that the audience *chooses to spend its finite attention on* rather than the producer of work that the audience could, in principle, choose. The slop ceiling, the authenticity premium, the chess-grandmaster move, the *Why* — these are not romantic notions. They are the *mechanisms by which finite human attention selects from infinite producible content*.

Two. The platform business models that depend on *expanding audience attention faster than supply* are, on the structural reading, walking into a wall. The platform companies' current trajectories — push more content, optimise for engagement, monetise more minutes — work only as long as the engagement budget is elastic. It is not elastic. The audience cannot, on average, watch more hours per day than it already does. The platforms that get to those audiences first, and that build the most-defensible *retention* mechanics, will win the share-of-attention game. The platforms that come second will be competing for an attention budget that the first-movers have already claimed.

Three. The new categories of work I have described in this chapter will produce, in aggregate, more *total creative output* than the previous decade. They will not produce, in aggregate, more *audience attention received per minute of output*. The ratio of attention-to-output, which is what working creatives actually live on, is going to fall — sharply, in some categories, more gradually in others. The working creatives who survive the next decade will be the ones who recognise that the attention-to-output ratio is the metric that actually matters, and who position their practice in the categories where the ratio is most defensible.

The categories where the ratio is most defensible — on the chess-grandmaster reading and the slop-ceiling reading and the authenticity-premium reading — are the categories where the work is *most deliberately un-machine-like, most authentically human-authored, most culturally specific, most personally risked*. The newly-possible work in this chapter is real and will reshape the creative economy. But it is being made in a market where the binding constraint is, and will remain, the finite attention of the audience watching it.

The synth made entirely new sounds possible. The audience for those sounds was finite. The musicians who learned to make sounds *the audience would spend its scarce listening time on* — Trevor Horn, Kraftwerk, the Bomb Squad, Aphex Twin, Daft Punk, every electronic-musician who built a serious career — are the ones we still listen to. The musicians who learned to make sounds the synth made possible

but that the audience did not develop a hunger for — the long tail of 1970s and 1980s synth-record obscurities — are remembered, mostly, by collectors.

The AI moment will, on the historical pattern, work the same way. The newly-possible categories will create work that did not exist before. The audience will, on the available evidence, allocate its scarce attention to the work that *earns* that attention. The working creatives who position themselves in the newly-possible categories, *and* who make work that the audience deliberately chooses, are the ones who will define the next decade of the form.

That is the operating manual. The categories are open. The constraint is real. The choice — like the chess-grandmaster's — is yours.

Chapter 12 — The Orchestrator

In the second issue of *Dream Machine*, in October 2025, I described the **Human–AI Agency Continuum** as a way of mapping how much of any given creative function is being done by the human in the chair and how much by the machine. In the thirteenth issue, in January 2026, I made a prediction that I want to look at again in this chapter, because — six months on — it has held up better than most of the others I made that day.

I called 2026 “**the Year of the Orchestrator.**”⁴⁹⁹

The argument was straightforward. If 2024 had been the year of the *generator* — the prompt-and-respond text-to-image, text-to-video model — and 2025 had been the year of the *agent* — the system that could take goals, plan, decide and execute multi-step tasks autonomously — then 2026, I argued, would be the year that working creatives stopped being *operators* of these tools and started being *orchestrators* of them.

By “orchestrator” I meant something quite specific: a person whose job is not to make the work themselves but to direct, brief, integrate and judge the work of a team of AI agents and human collaborators. Not a producer in the old sense. Not a creative director in the old sense. Something newer, with a different skill set, a different rhythm, a different relationship to craft.

This is the role I think most working creatives will be holding by 2030. This chapter is about why, what it looks like, what it asks of you, and where it breaks.

Forty-nine Claude agents and seventy-two skills

In May 2026, in the second-to-last issue I wrote before this book went to draft, I reported on something that I think of as the canonical image of what the orchestrator role actually is. **Sony**, in announcing its “all-in on AI for games” strategic move, disclosed that one of its game-development studios was running a coordinated multi-agent team of **49 Claude Code agents**, working with **72 skills**, on game-development tasks ranging from asset generation through QA through engineering through animation.⁵⁰⁰

This is, in operational terms, a small army of synthetic colleagues working on a single creative project. Each agent has a defined role. Each skill is a defined capability. The whole apparatus is overseen by a *substantially smaller number of human developers and creative leads* whose job is to plan the team’s work, brief the agents, judge their outputs, integrate their contributions, and decide what gets into the game.

The ratio matters. The pre-AI version of this game-development team would have been, plausibly, 50 to 100 people working on the same scope of work over a much longer timeline. The AI-augmented version, as Sony has set it up, is a much smaller number of *senior, judgement-heavy* roles — people whose value is taste, narrative

sense, gameplay design instinct, IP fluency, engineering oversight — sitting on top of a much *larger* pool of synthetic capacity.

What you don't see, in the Sony picture, is the disappearance of the junior roles. Those roles haven't disappeared. They have been *replaced* — by agents. The 49 Claude Code agents are, in effect, the new junior developers, the new junior animators, the new junior writers. They work cheaply, they work fast, they work in parallel. They are not — and this is important — *replacements for senior judgement*. They are leverage *for* senior judgement. The whole pipeline is designed to take the senior creatives' time and multiply its effective reach by a substantial factor.

The orchestrator, in this configuration, is the senior creative — the writer-director, the gameplay lead, the art director, the technical director — whose taste and judgement set the boundaries that the agent team works inside.

I want to flag the obvious labour question here, because it is the question every working creative is asking, and any chapter that handwaves past it is not being honest. If the senior roles still exist, and the junior roles are now done by agents, *where does the next generation of senior creatives come from?* The pipeline that has, for fifty years, produced senior creatives in the film, TV, games, music and design industries has worked by *starting people as juniors and letting them grow*. If we remove the junior rung, we are — over the next decade — also removing the apprenticeship infrastructure that makes the senior rung possible.

This is the structural risk of the orchestrator economy that platform companies are, in my view, not yet taking seriously. I name it the **Apprenticeship Gap**, and I will come back to it in Chapter 15.

What the orchestrator does

In the talks I have given since publishing the “Year of the Orchestrator” piece in January, the question I get asked most often is: *what does an orchestrator's day actually look like? What are the skills? What does the job description say?*

I want to try to answer that in this chapter, in the most concrete language I can find, because I think the gap between *what working creatives think they will be doing in 2030* and *what they will actually be doing* is unhelpfully large.

The orchestrator does five things:

One. They define the brief. The agents — and the human team — work to a brief. The brief is the thing the orchestrator owns. It is not the prompt. The prompt is a derivative artefact. The brief is the *creative intent* that the project exists to deliver: who it's for, what it's trying to do in the world, what success looks like, what tone, what feeling, what audience, what context. The orchestrator's first job is to know — clearly enough to communicate it — what the work is *for*.

Two. They allocate work. Given the brief, the orchestrator decides which parts of the work get done by humans, which by AI agents, and which by the orchestrator themselves. This is the practical, day-by-day application of the Human–AI Agency Continuum we talked about in Chapter 3. It is not a one-off decision. It is a constant series of micro-decisions about where on the continuum each function sits, in this project, on this day, for this output.

Three. They brief the agents. This is the closest the orchestrator role gets to what people currently mean by “prompt engineering,” but it is a meaningfully different skill. Briefing a human collaborator and briefing an AI agent are not the same activity, but they share a core: the ability to *describe what is wanted with enough precision and enough context that the recipient can produce something useful, without over-constraining them in ways that prevent useful surprise*. This is, in my experience, the single biggest skill differentiator between effective and ineffective AI-era creatives. The people who can brief well — who know when to give the agent a tight constraint and when to let it explore — are the people who produce the best output.

Four. They judge the outputs. When the agents deliver, the orchestrator’s job is to look at what came back and decide what to ship, what to revise, what to throw away. This is *taste* in the most operational sense. It is also, crucially, *taste under abundance* — taste exercised in a context where you can have ten versions of the same scene back in ninety seconds and your job is to choose, not to make. Choosing well under abundance is a different cognitive skill than choosing well under scarcity, and most creatives have been trained for the latter. The grandmaster analogy from Chapter 16 — top chess players, in 2026, deliberately playing *sub-optimal* moves to put their opponents on uncomputed ground⁵⁰¹ — is the cleanest available picture of what this looks like in practice. The orchestrator’s value is not in choosing the *most-likely-good* output of the ten variants the agent returned. The agent has, by construction, already centred its output on the most-likely-good. The orchestrator’s value is in seeing which of the ten variants would be the *un-machine-like* move at this specific point of this specific project, and choosing that one. *Taste under abundance* is, operationally, the discipline of refusing the machine-optimal output in favour of the deliberately-chosen one.

Five. They integrate. A film is not the sum of its scenes. A game is not the sum of its assets. A campaign is not the sum of its individual creatives. The orchestrator’s job, at the end, is to take the outputs of the agent team and the human team and assemble them into a coherent piece of work that has a *single sensibility*. This is the part of the job that I think — for all the AI tooling — is least likely to become a thing AI can do. The integrated voice of a piece of work is a function of a single human consciousness running through it. The orchestrator role is, at its core, the role that holds that voice.

If you read those five things back, you will notice that none of them are *making*. They are all *deciding*. The orchestrator's work product is decisions: about what to make, who or what should make it, whether the made thing is good enough, and how the made things fit together.

This is, in some sense, what every senior creative director and showrunner has always done. The change is not the shape of the role. The change is that the role is no longer a privileged senior position at the top of a pyramid of junior makers. It is, increasingly, *the entire role*. And it is the role that, in 2026 and 2027, working creatives at every level are being asked to grow into faster than the career-development infrastructure of any of the creative industries is built for.

Where the agents go wrong

I want to spend some time on the failure mode of agentic creative work, because the press cycle around Sony's "49 agents" framing — and the corresponding announcements at Adobe Summit, NVIDIA GTC and elsewhere — has been heavy on the upside and thin on the downside.

The agents go wrong, in my experience and in the experience of every working creative I have talked to about this, in four characteristic ways:

One. They confidently produce the wrong thing. This is the most familiar failure mode and the one the public discourse has covered most. Agents — like the LLMs underneath them — *hallucinate*. They will produce an asset that confidently disregards a key constraint of the brief. They will generate a character with the wrong eye colour. They will produce a piece of music in the wrong key. The fix, with current systems, is *human review at every gate*. The cost of this review, in time, is the largest single source of the "AI was an expensive mistake" experience that Charles Cecil described and that many studios have replicated.

Two. They produce the *mean* of the training distribution. This is the more insidious failure. Agents, by default, will produce work that sits in the middle of what they have been trained on. The middle of a training distribution is, by definition, the *most average* version of the thing you asked for. For creative work — where the value is almost always in the *non-average* — the default output is structurally weak. To get above-average output, the orchestrator has to push, prompt, and curate against the gravitational pull of the mean. This takes deliberate, conscious effort and it takes taste to know what *above-average* looks like in this particular project.

Three. They lose context across long tasks. Agents working on multi-step tasks accumulate errors over the steps. A small misalignment in step one becomes a larger one by step five. By step ten, the output is meaningfully off-brief. The orchestrator's role is to *check in* at the right intervals — not so often that you negate

the benefits of agentic execution, not so rarely that the team has wandered off the brief by the time you look.

Four. They cannot tell when to stop. Agents, given an open-ended task, tend either to over-iterate (producing fifty variants of the same thing without converging on one) or to under-iterate (producing one variant and stopping). The orchestrator's job is to set the *stopping criteria* for the agents, which is, in practice, a series of judgment calls about *when good enough is good enough*. This is a craft skill the agents do not, as of 2026, have. It is also a skill that, in my experience, almost every working creative already has — they just haven't had to use it on synthetic colleagues before.

The Anthropic blog posts on agent deployment patterns through Q1 2026 made the point I want to land on here cleanly: agentic systems work best when they are deployed by people who already have the taste and judgment to know what good output looks like. They *accelerate* people who are already good. They do not, on their own, make people good.⁵⁰²

That is the orchestrator's job: to *be the human who is already good*, holding the taste line, while the agent team produces faster than the human pipeline ever could.

By mid-2026, even the model labs were making this point from their own stages. When **MiniMax** — one of the Chinese labs whose video models have been setting the pace this year — launched its all-in-one AI video generator hub at the Shanghai International Film Festival in June, the framing it chose was not the maximal claim its marketing department could have written. It was a concession: “*creative direction and aesthetic judgment must ultimately be left to humans.*”⁵⁰³ Sit with that for a second. A frontier video lab, at its own launch event, in a festival market it is actively trying to win, describing the boundary of its own product in exactly the terms this chapter has been arguing for. When the people selling the agents and the people directing them agree on where the line sits, the line is probably real.

The console becomes conversational

There is a development from the early summer of 2026 that deserves its own section, because it changes the *texture* of the orchestrator's working day more than anything else this year: the orchestrator's console is becoming conversational.

In July 2026, ComfyUI — the open-source node-graph engine that has quietly become the professional operating layer for serious generative image and video work — launched **Comfy MCP** in public beta: pitched as the first Model Context Protocol server built for real production pipelines rather than one-off generations, under a tagline that could serve as this chapter's epigraph — “*turn your agent into a creative technologist.*”⁵⁰⁴ Point an agent like Claude, Codex or Cursor at it and you can build, edit and run production workflows in plain language, search models and nodes, and

share a workflow URL so a colleague's agent runs the *exact same* pipeline, every generation reproducible. And it did not arrive alone. In the same few weeks, Epic announced **Unreal MCP** at the State of Unreal, Runway shipped its tools *inside* ChatGPT, and Magnific and Higgsfield both stood up MCP servers of their own.⁵⁰⁵ Read the pattern rather than the products: every serious creative system is being fitted with a socket that an agent can plug into. The orchestration this chapter has been describing — brief, allocate, judge, integrate — stops being a metaphor about management style and becomes a literal software architecture. The agent is no longer just a maker of content. It is an *operator of creative systems*, and the orchestrator directs it in sentences.

By high summer the socket-fitting had spread well beyond the image-and-video stack. FL Studio — one of the most widely used DAWs in the world — can now be driven by an agentic AI chatbot from inside the application itself: you tell it what you want and it operates the software.⁵⁰⁶ Music Ally ran an explainer for the industry's business side under a title that captures the moment — MCP as “AI's new USB port”⁵⁰⁷ — while *Digital Music News* documented how AI agents were already quietly automating music companies' back offices, from A&R triage to release operations.⁵⁰⁸ Television got its “world first” agentic AI TV platform.⁵⁰⁹ And **Filmera** launched as an early-stage agentic filmmaking workspace whose explicit ambition is the *whole* pipeline — idea, script, shot planning, generation, edit — inside a single agent-driven environment.⁵¹⁰ Two details from the same weeks tell you where this is heading in scale and in texture. Spotify's engineering leadership described running agents across a codebase of more than twenty million lines⁵¹¹ — which is what the orchestrated-agent pattern looks like when an enterprise adopts it as infrastructure rather than experiment. And OpenAI shipped a *physical keypad* for controlling agents⁵¹² — the conversational console acquiring literal buttons. When a technology gets dedicated hardware, it has stopped being a demo.

Two things follow from this, and they pull in opposite directions. The first is that briefing — function three of the five — becomes even more nakedly the differentiating skill. When the pipeline itself responds to natural language, the distance between a precise brief and a vague one is the distance between a production system and a slot machine. The second is that the low-agency end of the same spectrum gets a fresh coat of glamour. The same fortnight Comfy MCP shipped, OpenArt — an AI startup run by ex-Googlers — was pitching a new model to *The Hollywood Reporter* under the banner of “**vibe directing**”: filmmaking by loose intent, with the system inferring the rest.⁵¹³ I want to be blunt about why that phrase should worry you. Everything this book has argued about the Human–AI Agency Continuum, and everything the Process Trap essay argues about curating the machine's mean, says that the value of the conversational console is *proportional to*

the agency of the person speaking into it. The vibe director has delegated not just the execution but the *intent* — and intent is the one thing on the five-function list that cannot be delegated without the role collapsing into passengerhood. Kane Parsons, the director behind the *Backrooms* phenomenon, had put the refusal position at the other pole with equal bluntness only weeks earlier: using AI “defeats the purpose” of filmmaking — “*I get no enjoyment from using those tools.*”⁵¹⁴ Between the vibe director who has surrendered the brief and the refusenik who has surrendered the leverage sits the position this chapter is actually arguing for: the orchestrator who holds the intent tightly, and speaks it into a console that can finally hear it.

Sundance, and the literacy turn

I noted in the last chapter that the platform companies have responded to the AI-literacy gap by becoming, in addition to software companies, *training organisations*. The most institutionally credible example of this turn, in the period this book covers, was **Sundance Institute’s** launch of an **AI Literacy Initiative** at the 2026 festival.⁵¹⁵

I want to spend a moment on what Sundance did, because the framing is important.

The Institute’s announcement did not say “*AI is the future of filmmaking; here is how to use it.*” It said something more careful. It said that AI is a fact of the filmmaking landscape, that filmmakers are going to have to make decisions about whether and how to use it, and that those decisions should be made by *informed* filmmakers with *agency over their own practice* — not by filmmakers who have had the tools imposed on them by clients, by streamers, or by tool vendors.⁵¹⁶

The framing was *creator empowerment*. The mechanisms were free learning through Sundance Collab, community conversations, a fellowship and alliance model, and a Story Forum that specifically tackled the legal questions creators face when they use AI: whether AI-generated content can be copyrighted, how to protect projects in a world of contested datasets, how to negotiate AI clauses in production contracts.⁵¹⁷

Google funded this. The \$2 million the company put into the Institute, with the stated aim of training 100,000+ artists in foundational AI skills, was both an act of corporate generosity and a strategic investment in the *category* of “filmmaker who can use AI without losing their creative authority.”⁵¹⁸ Both motivations are real. Both can be true. What matters, for the working filmmaker in 2026, is that the institutional infrastructure for becoming an AI-literate orchestrator — without surrendering creative agency — now exists.

The McKinsey AI report on film and TV production, released in early 2026, made the corresponding business case. AI would not, in McKinsey’s view, replace film and

television production. It would *restructure* it — towards smaller teams, faster cycles, more iteration, and a heavier reliance on senior creative judgement.⁵¹⁹ In other words: towards an orchestrator-shaped industry.

The middle layer disappears

What is happening, structurally, in every creative industry that the *Dream Machine* newsletter has tracked in these six months, is that the middle layer of the workforce — the layer of *intermediate* roles, between the very senior creative leadership and the very junior entry-level — is being absorbed into the agent layer.

This was the story behind **Ubisoft's** decision in January 2026 to cancel five games, including the *Prince of Persia* remake, while pouring more money into AI.⁵²⁰ It was the story behind **Square Enix's** target of doing 70% of its QA work via AI by the end of 2027.⁵²¹ It was the story behind **Falcom's** description of work that “previously took 2–3 hours” being completed “in 10 minutes” with AI tools.⁵²² *Eighteen-to-one productivity.* That ratio, taken on its own, is what re-shapes the headcount calculus for every studio's mid-level production work in the next eighteen months. It was the story behind the *Take-Two CEO's* explicit framing that AI “won't invent the next *Grand Theft Auto*” — meaning, the *creative direction* won't come from the machines — even as Take-Two's QA, asset and engineering pipelines absorb AI capacity rapidly.⁵²³

In film, you see the same pattern. **Spielberg** explained in March 2026 why he hadn't yet used AI directly⁵²⁴ — and the same press cycle reported that he had a substantial AI-augmented team working on production-pipeline tasks underneath him. **Steven Soderbergh** committed to “a lot of AI” on the Wagner Moura film and a John Lennon documentary, framing it explicitly as transparency: “I owe people honesty.”⁵²⁵ In every case, the structure is the same: a senior creative voice on top, an AI-augmented operational layer underneath, *fewer mid-career intermediaries in between.*

In advertising, the pattern was even more pronounced. **Independent agencies** faced what *Digiday* called “a new frontier as agency-in-a-box tools democratize creativity.”⁵²⁶ **AI agent developers** became “adland's in-demand role.” The framing from one agency hiring lead, given to *Digiday*, captures the shift better than any of the trend-piece coverage: “*We actually need people who understand [AI], who are building systems organically within their day to day workflows. People who understand taking what took them 40 hours one week and turning it into 38 the next week.*”⁵²⁷ The job description is no longer *make the work*. It is *make the system that makes the work — and keep shaving hours off it.* The PGA Tour expanded its AWS partnership to put AI content at the heart of its content

distribution.⁵²⁸ **Mondelez** said it would use AI for TV ads in 2026.⁵²⁹ **Avocados From Mexico** turned to AI to advertise around the Super Bowl, *instead of* a traditional TV buy.⁵³⁰ **Adobe** said that AI in marketing was now “agentic creative intelligence.”⁵³¹

In journalism, the Reuters Institute’s “AI adoption by UK journalists” survey found high integration of AI tools across newsrooms by late 2025.⁵³² Daily Mail reported that Google’s AI Overviews had “killed click-throughs” to news sites.⁵³³ *The Times* was using AI to model synthetic focus groups from human audiences.⁵³⁴ In each case, *the middle layer of the journalism workforce* — the sub-editors, the copy editors, the data journalists, the social-media producers — was the layer most exposed to AI substitution.

I want to be honest about what this means. *It does not mean that every working creative in the middle of their career is about to lose their job.* That framing — the apocalyptic one — has, for two years, been the most popular and the most wrong. What it means is that the *shape* of the mid-career role is changing. Mid-career creatives who can become orchestrators of agent teams will, in many cases, *gain* leverage and earning power. Mid-career creatives who cannot — who try to keep doing the maker-as-craftsperson job at the speed and price of the agents — will, increasingly, struggle.

The Sundance literacy turn, the Adobe and Google training investments, the UK government’s free-AI-training-for-all programme — these are the institutional response to that pressure. They are not enough on their own. They are, however, the right direction.

The portfolio creative

There is one more shape of the orchestrator role that I want to flag, because it is the one I see most often in my own studio and in the wider DreamLab community: the **portfolio creative**.

A portfolio creative is someone who, instead of holding a single specialist role, holds *several* loosely-coupled creative roles across different disciplines, supported by AI tooling that lets them maintain useful proficiency in each. The portfolio creative is a writer-director, but also a creative technologist; a music producer, but also a video editor; a games designer, but also a brand strategist.

The *TechBullion* piece “Why the future belongs to multi-skilled leaders,” from November 2025, made the case for this from a corporate-leadership angle.⁵³⁵ The Anthropic Skills framework — the system of named, reusable skills that Claude Code now uses to coordinate multi-agent workflows — is, in effect, an attempt to make the portfolio-creative model into a *technical infrastructure* rather than a personality

type.⁵³⁶ The *Forbes* piece “AI Is Changing How Creators Work And Earn,” from December 2025, surveyed the same phenomenon from the working-creator angle and found the same pattern: the most economically successful creators in 2026 are not specialists. They are *integrators* who can work across disciplines using AI as the connective tissue.⁵³⁷

In my own studio, the move towards portfolio creatives has been a deliberate strategic choice — and an honestly difficult one to execute. The cultural expectation, in most creative industries, has been to hire specialists and stack them in a pipeline. The portfolio-creative model requires you to *hire generalists* and let them move between disciplines as the work demands. The former is easier to manage, easier to bill, easier to explain. The latter, in my experience, produces better work in the AI era, because the human is doing the integration that the agents can’t.

The portfolio creative is the orchestrator at the level of an individual career. The orchestrated team is the orchestrator at the level of a project. The same pattern shows up at multiple scales.

Five orchestrators, briefly

I want to give five working orchestrator case studies, because the abstract description above can sit in the head as a theory without the operational texture of *what the role actually looks like* on a Wednesday afternoon. Each of the five is a specific working creative or organisation whose practice in 2025–26 I think represents a different *shape* of the orchestrator role. Each is documented in the *Dream Machine* newsletter archive. Each is, on my read, doing something that working creatives reading this book can learn directly from.

Andrii Daniels (Ukraine). The independent filmmaker who, in December 2025, produced a *Deadpool* / *Harry Potter* Christmas mash-up in a Ukrainian bomb shelter using a Runway-and-Veo-and-ElevenLabs stack on a laptop running through a generator. The clip went viral and was picked up by *Variety* as a profile piece.⁵³⁸ What Daniels did, operationally, was the orchestrator role at its purest: a single human creative, holding the taste and the narrative judgement, briefing a stack of generative tools to produce work whose production conditions would have been physically impossible eighteen months earlier. Daniels did not write, draw, animate, voice or render the work. He *briefed* it. He *integrated* it. He *judged* what was good enough. He *delivered* the finished piece, on his own, with no studio underneath him. The bomb-shelter context is the dramatic detail; the underlying operational pattern is what makes Daniels’ practice replicable for working filmmakers in any production environment.

The Imaginae Studios / *Art Awakens* team (Fremantle). The AI-native studio I described in Chapter 7 has, by mid-2026, settled into an operational pattern

that maps cleanly to the five-function orchestrator description above. A small senior team of writer-directors, supported by an AI-augmented production pipeline that handles concept development, asset generation, scene assembly and post-production. *Art Awakens* — Imaginae’s flagship 2026 development project, fusing AI techniques with classical painting IP — has, on the published interviews with the team, been produced by a core human team of fewer than ten people running an agentic pipeline that, by their own estimate, produces output that would have required a sixty-to-eighty-person team five years ago. The 8:1 productivity ratio is the orchestrator economy expressed at the studio scale.

Sven Vincke / Larian Studios. The opposite case, also instructive. Larian, the maker of *Baldur’s Gate 3*, has — as I described in [Chapter 7](#) — *publicly refused* generative AI for its next major game while continuing to use AI-augmented tooling in *adjacent* parts of the pipeline (QA, localisation, internal admin, asset management). Vincke’s framing, in his January 2026 statements, was *not* anti-technology. It was *position-on-the-continuum*: certain parts of the game’s authorial signature (writing, character design, world-building, dialogue) had to remain *fully human* for commercial and cultural reasons; certain other parts (build tooling, QA automation, marketing-asset generation) could be AI-augmented at no audience-visible cost. Vincke is, in operational terms, *an orchestrator at the level of the studio’s continuum positioning*. He is making the agency-line decisions at the strategic level that the working filmmaker makes at the project level. The role is the same. The scope is different.

Xania Monet / Hallwood Media / Telisha Jones. The case I have, in [Chapter 5](#), been most equivocal about. The structural pattern is, on inspection, an orchestrator economy operating at the *single-artist* level. Telisha Jones, the human lyricist, is the orchestrator. The Suno music-generation stack is the agentic capacity producing the *executed* musical output. The Xania Monet persona is the *audience-facing* product. Hallwood Media’s \$3M deal pays for the orchestrated work as a unit, with the orchestrator (Jones) receiving the commercial credit and revenue that the synthetic-vocalist alone could not have generated. Whether this scales into a sustained cultural-star career — the [Chapter 5](#) slop-ceiling argument suggests, on six months of evidence, that it has not — is a separate question from whether it works as a *business form*. As a business form, it is the orchestrator role at the level of a solo recording artist.

The Sony game-development teams running 49 Claude agents and 72 skills. The canonical *enterprise-scale* orchestrator case I opened the chapter with. The 49-agent / 72-skill stack is, in operational terms, an organisational design for the orchestrator role at the *team* level: a small group of senior creative leads (writer-directors, gameplay design leads, art directors, technical directors) orchestrating a multi-agent synthetic team whose individual outputs require senior human review

and integration. The 49 agents are not autonomous studio replacements. They are *leverage* for the human orchestrators. The 72 skills are the *reusable capabilities* the orchestrators can deploy across multiple projects.

I want to be honest about what the five cases share, because the shared pattern is the operational lesson.

In every case, the orchestrator's *contribution* is the same five-function set I described above: brief, allocate, brief-the-agents, judge, integrate. None of the five is doing the production-execution labour themselves. All of them are making *decisions* that direct a synthetic (and, in most cases, also a human) team to do the production-execution labour on their behalf. The *value* the orchestrator brings to the work is, in every case, the senior judgement that the agents cannot supply: the taste that knows what good output looks like, the briefing skill that gets useful output out of the agents, the integration sense that assembles the agent outputs into a coherent piece of work.

In every case, the orchestrator's *leverage* — the ratio of finished output produced to working hours spent — is dramatically higher than the equivalent practitioner could produce without the agent layer underneath them. Daniels would not have made the Christmas clip in a week pre-AI. Imaginae would not have produced *Art Awakens* at the pace and budget they are producing it at. Jones would not have shipped a Billboard-charting record without Suno. Sony's game-development teams would not have shipped on the cadence they are shipping at. The leverage is real. The leverage is what makes the orchestrator role *commercially* viable as a new form of working practice.

In every case, the orchestrator's *fragility* is also the same: the practice depends on the toolchain underneath it continuing to be available, accessible, on commercial terms the orchestrator can sustain, with model behaviour that the orchestrator can predict and brief against. The platform-dependency of the orchestrator role is the structural risk that [Chapter 10](#) addresses. It is also the reason the *open-the-black-box* argument I made in [Chapter 3](#) is operationally serious: the orchestrators who depend on closed-platform tooling without understanding the dependency are, structurally, exposed to platform-pricing and platform-policy decisions that they have no control over.

The orchestrator role is, in operational summary, *high-leverage and platform-dependent*. The working creatives reading this who are positioning themselves toward the role need to take both halves of that description seriously. The leverage is the upside. The platform dependency is the work that has to be done to defend the upside over time.

The process trap, and the migration out of it

I want to put on the page a counterweight argument to the orchestrator framing this chapter has been building, because the framing has, in the reception I have seen of the “Year of the Orchestrator” piece, sometimes been read as a *consoling* one — as a reassurance that creative work is safe and that the working creative needs only to ascend to the orchestrator role to keep their seat at the table. That is not what the chapter is arguing, and the reassurance is wrong.

The counterweight comes from an essay I have been carrying around with me through the final drafting of this book, *The Process Trap: Generative Abundance and the Equivocation of the Creative*. The essay’s central move is to pull apart the word *creative* — which does two incompatible jobs at once, naming a rare capacity (the genuine origination of the new) and naming a workforce defined by sector, brief, format and schedule. The reassuring sentence circulating through every studio all-hands and trade-press op-ed — *automation will take the rote; real creative work is safe* — is true of the *capacity* and catastrophic for the *class*, and it slides between them without anyone noticing because one word covers both.

The argument that follows from the equivocation is the harder one for this chapter to absorb. Generative AI is, precisely, a *process* technology. It drives the marginal cost of combinational and exploratory creativity toward zero, including — and this is the part most under-absorbed — the tacit craft long believed inimitable. Polanyi’s formula that we know more than we can tell was treated as the craftworker’s last redoubt; it is exactly this redoubt that large models have breached, because the tacit dimension turns out to be tacit *to the individual practitioner*, not to the corpus. What no single craftsperson could explain, the corpus encodes. The process premium — the rent that craft, technique and executional competence have historically commanded — is collapsing, and is collapsing *fastest* at the upper rungs of the Dreyfus skill ladder where the most-invested experts sit.

The orchestrator role, in this reading, is not an escape from the process trap. It is, in some configurations, the *re-entry* into it under another name. The “prompt engineering” that the consumer press has been celebrating is process re-dressed. The “personalisation at scale” pitched as the answer to synthetic abundance is what Adorno and Horkheimer called *pseudo-individualisation* industrialised: the same standardised core, recombined and surface-tuned per recipient, presented as bespoke. *Curating* the best of the agent’s outputs is still deference to the machine’s mean. The orchestrator who only chooses among the variants the system returned has, in the Process Trap’s framing, re-entered the process economy by a different door.

The defensible thing, the essay argues — and this is the part I want this chapter’s readers to take seriously alongside the orchestrator framing — is not in the artifact and not in the lone maker. It is in the *position* of the maker: a person who has a

problem of their own (the originating problem, in Getzels and Csikszentmihalyi's language; *natality*, in Arendt's), a stake in the game (Bourdieu's *illusio*), an embedded place in a living field (Csikszentmihalyi's systems model), and the *negative capability* to refuse the probable answer the machine returns by default (Keats's phrase; Ghafouri's *Paradoxical Bridge* in its tooled-up version). Creativity, in the strong sense that resists automation, is **positioned negation under risk in pursuit of significance** — the bringing-into-being of something that did not have to be, by an agent who has a problem of her own and a stake in the game, achieved by refusing the probable, and judged to matter by a living field. *Every clause names a feature of the maker's position; not one names a feature of the output.* This is why the residual moat is so consistently mislocated.

The operational reframe that follows is the one I want the orchestrator-shaped reader to carry forward into the next chapter. The correct use of a generative system, for a maker pursuing creativity, is not as a creative partner and not as a curation engine, but as the *maximally competent producer of the probable* — precisely so that the human can specialise in the *improbable*. Let the model occupy the mean it is built to occupy. Use its best output as a *map of what to avoid*, not a menu of what to choose. The orchestrator who defers to the model's best output has re-entered the process trap by another door. The orchestrator who uses the model's best output as a map of where the probable goes and then *deliberately departs from it* — refusing the obvious continuation, the well-formed default, the safe resolution — is the one whose practice survives the collapse of the process premium.

The essay's closing observation is the harder of the two I want to hold against the rest of this chapter's framing. The migration from making to creating — competing on the problem rather than the execution, taking a stake on purpose, re-entering the field as a place of significance-judgement rather than an audience-retention asset, practising negation against the gravitational pull of the mean — is *harder than the work the orchestrator was trained for, slower to reward, lonelier, and structurally unsupported by the institutions that employ them*. To pretend that every working creative called *creative* can simply ascend to *creativity* proper is to repeat the equivocation in a consoling key, and it does a disservice to people making real decisions about real livelihoods. Most creative workers will not make this migration. Some will. The difference between those two facts is where most of the human cost of this transition will fall.

I include this counterweight here, in the chapter that argues most ambitiously for the orchestrator role, because the chapter's argument and the essay's argument are not in conflict but they are not the same. The orchestrator role is the *operational* shape of working creative practice in 2026–2030. *Positioned negation under risk in pursuit of significance* is the *defensible* shape underneath it. The working creative who holds both at once — orchestrating fluently at the level of execution and

refusing the probable at the level of intent — is the working creative I think most likely to make work that lasts.

The role this asks you to play

If you are a working creative reading this — and most of the readers of *Dream Machine* are — the question I am sure you have is: *what does this mean for me, this year?*

The honest answer is that it depends on where you are on the continuum we drew in Chapter 3. But there are some things I would say to almost everyone I know in the creative industries right now, and I want to put them on the page.

Practice briefing. It is the single most leveraged skill you can develop. Brief your AI tools as if you were briefing a junior who has thirty seconds to understand what you want and ninety seconds to do it. The discipline of *being able to communicate what you want* will improve every other part of your creative practice.

Cultivate taste deliberately. Look at more good work, harder, with a more critical eye. The agents will, by default, give you the average. Your job, increasingly, is to know what *good* is. That knowledge is a function of how much good work you have looked at, how seriously, with how much craft attention.

Stay in the work. Resist the temptation to abstract too far. The director who never picks up the camera, the showrunner who never writes, the music producer who never plays the instrument — these are the figures most likely to lose the touch that makes their judgement worth anything in the first place. The portfolio creative is *not* a creative who has lost contact with the craft. They are someone who maintains craft contact in *several* domains.

Choose your line on the Continuum, and defend it. Decide where your craft sits, where you are willing to let the agents work, and where you are not. Write it down. Communicate it to clients, collaborators, your team. Be willing to walk away from work that would force you across the line you have drawn.

Build the apprenticeship. If you are senior enough to be running a team, take seriously the question of where the next generation of senior creatives is going to come from. The orchestrator model breaks if there is no path from junior to senior for new humans entering the field. The studios and agencies that survive the next decade will be the ones that solve this problem — by keeping some junior roles in human hands, by creating new pathways through AI-tool-augmented apprenticeship, by investing in the literacy infrastructure that the platform companies and the institutes have started but cannot finish on their own.

The Year of the Orchestrator is not a coronation. It is a job description. It is what most of us are now being asked to do, whether we have signed up for it or not. The

people who do it well will set the terms of the next creative economy. The people who don't will, increasingly, be sat next to it.

The thing I want to land before we leave this chapter is that the orchestrator role — for all the leverage it brings and for all the productivity it unlocks — depends, in the end, on something that the platform companies cannot ship and the agents cannot synthesise. It depends on the human in the chair *being someone*. Having taste. Having judgement. Having a perspective. Having the kind of relationship with the work that the agents do not, and probably will not, have.

That relationship — the kind of *authorship* that makes the work feel like it belongs to a person rather than a process — is, increasingly, the only signal the audience trusts.

That is the subject of the next chapter.

Chapter 13 — Authenticity as the New Scarcity

In early 2026, a stop-motion animator who goes by *Tiny Grandma* on YouTube uploaded a short to her channel. It was a stop-motion piece — claymation, frame by frame, the kind of work that takes weeks to make a few seconds of. YouTube’s AI-detection systems flagged it as AI-generated content and applied the platform’s automated labelling. The video went viral, not because of the animation, but because the platform’s automated system had wrongly flagged genuine human handcraft as synthetic.⁵³⁹

The story of *Tiny Grandma* is the perfect inverse of the Tilly Norwood story we opened with in Chapter 1.

If Tilly Norwood was the moment a synthetic creation tried to enter the working creative economy as if it were human, *Tiny Grandma* was the moment a human creation was wrongly identified as synthetic by the very systems that were supposed to protect the public from synthetic content. Both moments tell you the same thing, from opposite directions: the *signal* of whether a piece of creative work was made by a human is now an economic, cultural and legal asset of the first order, and the infrastructure for reliably establishing that signal is one of the most underdeveloped parts of the current creative economy.

This is the chapter about *provenance*. About why the question “did a person make this?” has become — in eight months — the single most important question in creative AI policy, and about what the people, companies and institutions trying to answer it are doing about it.

The death threats

In April 2026, *Dream Machine* Issue 23 reported, with as little editorialising as I could manage, that **Tilly Norwood’s creator Eline Van der Velden** had received death threats.⁵⁴⁰

The threats were not, of course, justified by anything. Death threats never are. But the cultural reaction that produced them — the visceral, sustained hostility that built up around the *idea* of a synthetic actress through the autumn of 2025 and the spring of 2026 — was not random. It was a specific response to a specific kind of cultural transgression. *You came here pretending to be one of us. You took something that belongs to us.*

The death threats are the extreme tail of a much larger curve of audience response that has been quietly shaping the AI creative economy for these six months. The slop ceiling in Chapter 5 — the 44%-to-3% Deezer ratio — is the polite version of the same response. The vehement audience pushback against AI art in *Call of Duty: Black Ops 7* and *Anno 117* in November 2025 is another version. The viral reaction to *Spotify’s* AI music infiltrating Discover Weekly playlists, the public anger at *McDonald’s Netherlands’* AI Christmas ad, the “*disturbing*” reception of *Valentino’s*

AI handbag campaign — every one of these episodes is the audience saying, in increasingly direct terms, *we know what is human-made, we want what is human-made, and we are paying attention to who is trying to slip us something else*.

This is the cultural pressure that *authenticity-as-scarcity* describes. It is not, as some of the more dismissive AI commentary has framed it, a romantic attachment to old craft. It is a *market signal*. The audience is allocating attention, money and trust on a basis that increasingly weights human authorship as a positive variable. I have come, in talks since the autumn, to call this the **Authenticity Premium** — the measurable excess of attention, willingness to pay, and cultural credit that audiences allocate to creative work whose human authorship can be verified. The Authenticity Premium is the *positive* side of the slop ceiling: the slop ceiling tells you what audiences *will not* engage with; the Authenticity Premium tells you what they *will pay extra for*. Both are market findings. Both are produced by the same underlying audience behaviour. The data is unambiguous. The strategic implication, for every working creative and every studio operating in this period, is also unambiguous.

In May 2026, **Bobby Berk** — the *Queer Eye* design lead — articulated the working-talent version of the same finding in a single line that I think is worth quoting in full because it captures the Premium argument from inside the unscripted-TV business: AI, he said, will make *reality TV* and “*verifiably human content*” *more valuable*, not less.⁵⁴¹ What Berk is describing, in industry-of-the-thing terms, is the supply-and-demand mechanism the rest of this chapter is about. As the synthetic supply of *any* content category expands, the *verifiably human* corner of the same category accrues a scarcity premium. Reality TV is, structurally, the unscripted broadcast form most resistant to AI replacement — it is the form whose entire commercial proposition is *real people, in real situations, producing unscripted reactions whose value depends on us knowing they are real*. That is the Authenticity Premium drawn down to a single broadcast genre, and Berk’s read is, on the structural argument, exactly right.

The chess analogy I develop at length in Chapter 16 sits underneath this. The Authenticity Premium is what it looks like when an audience, faced with an infinite supply of machine-optimal work, allocates its scarce attention to the *deliberately un-machine-like* move. The 88% of UK respondents who wanted licensing-in-all-cases were articulating the same preference at the policy layer. The 44%-to-3% Deezer ratio was the same preference at the listening layer. The Television Academy’s “tools used to bring it to life” language was the same preference at the institutional layer. The Authenticity Premium is, at its core, the *commercial price* of the deliberately-human move — the move the engine, by construction, could not have made — and the audience’s reliable willingness to pay it.

The question is what the *infrastructure* for honouring that signal looks like.

Fingerprinting real media

I quoted Adam Mosseri — the head of Instagram — in Chapter 4 making the case that the platforms should focus on “fingerprinting real media” rather than chasing AI slop disclosure. The fuller version of his argument, made repeatedly across late 2025 and early 2026, was that the current approach to AI content moderation — trying to detect and label everything synthetic — is unwinnable.⁵⁴² The volume is too high, the detection is too unreliable, and the labelling produces both false positives (Tiny Grandma) and false negatives (the AI hate-songs spreading across European Spotify charts in November 2025).⁵⁴³

The alternative Mosseri and others have argued for is *the inverse*: instead of trying to catch what’s synthetic, build infrastructure that can *prove what’s human*. A capture-time fingerprint — a cryptographic signature embedded by the camera, the microphone, the editing software, the upload pipeline — that travels with the file through its entire life on the public web.

The technical infrastructure for this is, as of 2026, partially built. The **Content Authenticity Initiative**, an Adobe-led coalition of camera makers, software companies and news organisations, has been working on it since 2019. By late 2025, **C2PA** (Coalition for Content Provenance and Authenticity) standards were supported by most major camera manufacturers, most major editing platforms, and a growing number of social-media uploads pipelines. The standards are robust enough that a photo taken with a C2PA-enabled camera, edited in Photoshop with C2PA-aware tools, uploaded to a C2PA-supporting platform, can carry a verifiable chain-of-custody for its entire provenance, from sensor to viewer.

Underneath this is **Google’s SynthID** — a watermarking system that Google has been deploying across its AI generation tools, including Veo and Lyria.⁵⁴⁴ In December 2025, the company announced that users could ask the Gemini app, “*Is this video made with AI?*”, and receive a reliable yes/no answer based on the SynthID watermark. By January 2026, this capability was available in the consumer Gemini product.⁵⁴⁵ At Google I/O 2026, the company reported that **SynthID had marked over 100 billion items** across its own ecosystem and was being extended to partner platforms including **OpenAI, ElevenLabs** and **Kakao**.⁵⁴⁶ That cross-vendor expansion is the single most consequential development on the provenance side of the period this book covers: a watermarking standard, born inside one platform, is — for the first time — being adopted across the foundation-model companies that have until now competed against one another on every other axis. (The ElevenLabs leg of that expansion went live in June 2026: under a formal partnership with Google DeepMind, audio generated by ElevenLabs’ models can now be identified with SynthID.⁵⁴⁷) If the C2PA Content Credentials standard is the

capture-time spine of authenticity infrastructure, SynthID-across-vendors is, as of May 2026, the closest thing the industry has to a *generation-time* spine.

These technologies are not, on their own, sufficient. Watermarks can be stripped by determined adversaries. C2PA chains break when files pass through non-compliant tools. The reliability of any given piece of provenance metadata depends on the integrity of every link in its chain. The trust infrastructure is still — relative to the speed of the AI rollout — early.

By June 2026 two more platform-side moves had landed and are worth flagging because they shift the labelling regime from voluntary infrastructure to platform default. **YouTube** rolled out an *automatic* AI-content labelling system: synthetic and AI-augmented videos uploaded to YouTube are, from this point onward, labelled by the platform itself rather than relying on creator self-disclosure.⁵⁴⁸

LinkedIn moved in the same direction from a different angle, announcing that it would *limit the reach* of AI-generated content on its feed — the first major professional-network platform to make the distribution-side argument that AI-authored posts deserve less algorithmic amplification than human-authored ones.⁵⁴⁹

The two moves point in the same structural direction: the platforms are moving from a *disclose-or-be-disclosed* regime to a *labelled-and-down-ranked* one, and the working creative who is reading the trend correctly is the one who treats provenance metadata as a *commercial asset* rather than a compliance cost. The labelled-human-authored work is, in this regime, the work that gets seen.

By August 2026 the regime had acquired both its statutory floor and its platform vocabulary. The **EU AI Act's transparency provisions** moved from the statute book to enforcement: AI labels are now *compulsory* on authentic-looking synthetic content under EU rules, which converts the disclosure norm this chapter has been tracing from good practice into legal obligation across the world's largest single market.⁵⁵⁰ The music industry was among the first to sit down and work through what enforced transparency obligations actually require of it — a useful early sighting of what every creative sector will shortly be doing.⁵⁵¹ And **Spotify**, whose voluntary transparency beta I list in the inventory below, introduced an **"AI Persona" label** for AI-generated artist identities on the platform.⁵⁵² That last one deserves a beat of attention, because it is the Tilly Norwood problem addressed at the metadata layer: an AI artist can exist on Spotify, but it can no longer exist there *as if it were a person*. The label doesn't ban the synthetic performer. It bans the impersonation.

The deeper structural move underneath all of this — and the one I want to spend a paragraph on because it ties the chapter's authenticity argument into the IP-economics argument that Chapter 7 has built up — is the **redefinition of intellectual property from an asset to an environment**, and the

corresponding migration of value from the *generation* layer to the *governance* layer. For a century, IP was understood as a discrete, ownable thing whose value you protected by controlling access to it. What the 2026 deals reveal, when you stack them next to each other, is something different: a bounded space, governed by rules, inside which value is created continuously by people who are not you. **Hasbro's Sixth Wall** doesn't sell you Optimus Prime; it sells developers a governed space in which Optimus Prime can be summoned, on-brand and on-licence, a million times over.⁵⁵³ **Spotify's Verified-by-Spotify badge**, layered onto the *consent / credit / compensation* opt-in catalogue of the Universal Music deal, is the same move applied to recorded music — the platform doesn't try to detect and remove the synthetic flood, it draws a fence around the verified-and-licensed work and routes attention and money to the inside of the fence.⁵⁵⁴ The C2PA chain-of-custody stack, SynthID's 100-billion-watermark deployment, the Cannes AI Disclosure Standard, the Academy's "*you must be human to win*" rule, the YouTube and LinkedIn labelling regimes I have just described, the New York AI-avatar disclosure law and the SAG-AFTRA *Tilly Tax* are — read in this light — *the same project seen from different angles*. They are not labelling regimes for synthetic work. They are *governance infrastructure for the authenticity premium*: the architectural enclosure that lets the audience tell, at the speed of a swipe, which side of the line the work in front of them is on. The platforms and rights holders that *own* this governance layer — Hasbro with CharacterOS, Spotify with the verified-stack, Adobe with C2PA, Google with SynthID, the Cannes consortium with the Disclosure Standard — are positioning themselves to capture the durable value of the generative economy. The fuller IP-economics analysis underlying this argument — the *three Cs* (consent, credit, compensation) framework, the comparative analysis of five emerging licensing models, and the corrected record on the Disney–OpenAI deal that *never executed* — sits in DreamLab's published research.

One caveat belongs here, precisely because I have just described Hasbro's Sixth Wall favourably. A governance layer is only as legitimate as the consent underneath it, and in June 2026 nearly 1,000 actors, talent agents and parents signed an open letter condemning contract clauses that require child actors to sign over their voices for AI use as a condition of the job — the flashpoint being the *Peppa Pig* franchise, in the Hasbro orbit.⁵⁵⁵ The same company can build the industry's most sophisticated brand-governance architecture *and* be the test case for what consent extraction looks like at the contract layer. The environment-not-asset model of IP is not self-legitimising; it inherits the legitimacy — or the illegitimacy — of the consents it is built on. This is the theme Chapter 6's three Cs framework exists to police, and the child-actor letter is the sharpest reminder of the period that *consent* is the C that fails first when the counterparty has no bargaining power.

But what these technologies are doing, collectively, is establishing the *category*. They are saying: the question *did a person make this?* is technically answerable, with high reliability, given the right tooling. The next decade of cultural and legal policy in the creative industries will be — in significant part — about who controls that tooling, who decides what it certifies, and what economic value it carries.

You do not have to take that trajectory on faith, because the June and July issues of the newsletter filled in the detection layer at a pace that surprised even me. **Warner Music Group** bought the AI-detection company **Sureel** outright — a major label deciding that detection infrastructure is something you *own*, not something you subscribe to.⁵⁵⁶ **Sony**, in the same weeks, invested in the AI copyright-protection firm **Midnight Labs**.⁵⁵⁷ **Deezer** — whose upload-time detection pipeline I have leaned on throughout this book as the best data source on the synthetic flood — productised that pipeline, launching its AI music detector *as a service for other streaming platforms*.⁵⁵⁸ The startup **RTM Audio** launched **UAI**, an AI-music detector that issues a *signed certificate* with every verdict — detection converging on exactly the chain-of-custody logic of the C2PA stack: a verdict you can carry and produce, rather than merely a flag.⁵⁵⁹ And on the human-verification side of the same coin, **Played By Humans** launched to verify “human artistry in an AI age”,⁵⁶⁰ while the dance-music store **Traxsource** introduced catalogue labels identifying human-made versus AI-assisted music.⁵⁶¹ When I first drafted this chapter, the provenance layer was an argument. By July 2026 it is a market — major-label M&A, platform-to-platform B2B products, certificate authorities and human-verification badges all competing to be the trust layer the Authenticity Premium settles on.

If you want to know where the next ten years of investment, policy and platform politics in creative AI is going, watch the provenance layer. The companies that win the provenance infrastructure will be — in a real sense — the companies that own the *signal of authenticity* that the audience increasingly demands.

The trademark, the trust, the tax

The cultural pushback against synthetic content has produced, alongside the technical provenance infrastructure, a parallel set of *legal and contractual* defences that working creatives have begun deploying around their own work and identity.

Taylor Swift filed trademarks on her voice and image in early 2026, specifically citing AI deepfake concerns.⁵⁶² By the summer, the Swift filing had stopped being a curiosity and become a pattern. **Lionel Richie** moved to trademark the sound of his voice in June 2026 — the trade press could not resist the headline: “*Hello, Is It Me You’re Looking For?*”⁵⁶³ — and the **Backstreet Boys** followed within weeks, filing to trademark their voices explicitly to decommission AI deepfakes, with the coverage now routinely describing each new filing as “joining Taylor Swift, Lionel

Richie and Others”.⁵⁶⁴ The trademark-the-voice play is, in other words, no longer one artist’s defensive novelty; it is a recognised legal instrument with a queue forming behind it. It is also — and this needs saying plainly — a rich artist’s instrument. The folk singer **Emily Portman** discovered in June 2026 that a fake AI album had been released *in her name* on the streaming platforms, and wrote, under a headline that is its own policy argument, that the experience “made me realise why the law needs to change”.⁵⁶⁵ Portman does not have a trademark portfolio and should not need one; what she needs is the statutory layer I come to below. **Matthew McConaughey** publicly drew the same line in January 2026.⁵⁶⁶ **Madonna and Will Smith** appeared in AI videos by Higgsfield in early 2026, the Madonna piece becoming a marquee example of how a major artist could *deliberately* deploy synthetic imagery as part of their own brand.⁵⁶⁷ (That Madonna would go on, in June 2026, to call AI “the complete opposite of making art” is not the contradiction it looks like: the line she is holding is between AI as a brand register she consents to and AI as a substitute for the making itself — which is, in miniature, the distinction this whole chapter is about.⁵⁶⁸) In the same vein, **The Rolling Stones** released *In The Stars* in May 2026, with a music video that used AI to de-age the band — produced, in a small piece of cross-industry casting that says something about where this category is heading, by the AI company belonging to the **South Park** creators Trey Parker and Matt Stone.⁵⁶⁹ **Boy George** supplied the audio-side version of the same consented register in June 2026, re-recording ‘Karma Chameleon’ with AI-assisted vocal technology as the first release from *Artist*, a new LA startup whose entire pitch is artist-initiated, artist-consented AI re-recording of classic tracks — the consent architecture this chapter argues for, packaged as a product.⁵⁷⁰ **George Clooney**, in November 2025, gave Variety the working actor’s read on the synthetic-star economy: “*It’s been just like a writer creating characters. You fall in love with your characters when you’re writing them. It’s a wonderful process. It wasn’t like I just made her in a second, and that was it. You know, it took a long time.*”⁵⁷¹ Clooney was making, in his particular way, the same argument that the slop ceiling makes empirically: cultural stardom is a *function of time and human relationship*. It is not a function of generation cost. **Jeremy Renner** threatened a “multi-millions” lawsuit against an AI documentary director he said had used his voice without permission.⁵⁷²

In May 2026 the celebrity-defence layer took a meaningful *organisational* step. **Cate Blanchett** co-founded **RSL Media**, a non-profit explicitly chartered to address *consent around AI usage* — covering creative work, name, image and likeness — for performers across film, TV and music.⁵⁷³ This is, on my read, the first time the celebrity NIL-protection conversation has produced a *standalone institution* rather than a string of individual lawsuits and trademark filings. RSL

Media is small as of mid-2026, but its founding signal is important. Where the existing infrastructure on celebrity AI consent has been individual (Swift's trademarks, McConaughey's line, Renner's threatened suit) or statutory (the ELVIS Act, New York's AI-avatar disclosure law), RSL Media is the first attempt at a *coordination layer* on the side of the talent — closer in shape to a Performing Rights Society than to a class action. If the Petrillo template I describe in Chapter 6 eventually has to be reconstructed for the NIL question, RSL Media is the kind of body that the reconstruction will need to anchor on.

By July 2026 the founding signal had become an operating one. Blanchett was back in the trade press — this time alongside **Steven Soderbergh** and the musician **Nikki Hexum** — making the public case for the **Human Consent Registry**: a registry through which performers, and in principle anyone, can record and assert the terms on which their identity may be used by AI systems. The framing they chose is the one to hold on to. Protecting your identity from AI theft, they argued, is a “*human right*” — not a perk of celebrity, not a contract term to be negotiated away, a right.⁵⁷⁴ The coordination layer I described above as a founding signal is, a few months on, materialising into working infrastructure — and doing so with exactly the Performing-Rights-Society shape I predicted it would need: a standing register of consent that any counterparty can be checked against, rather than a fresh negotiation every time.

In the same week, **Apple** acquired the talent and patents behind the AI-avatar company **Animato**, signalling — not for the first time — that the platform layer intends to own the celebrity-grade-avatar infrastructure rather than license it from third parties.⁵⁷⁵ The combination of an Apple-owned avatar pipeline and an RSL-administered consent regime is, in 2030 terms, the most plausible architecture for how the high-end NIL economy actually runs.

The posthumous case of the same architecture arrived in June 2026. The estate of **Ozzy Osbourne** announced that the late artist would “*live on as an AI avatar that talks with fans*” — a deliberate, consent-given (via the estate), avatar-shaped continuation of the artist's public presence into a posthumous register.⁵⁷⁶ The Osbourne case is the first major celebrity-grade posthumous AI presence to be launched with the *estate's* explicit framing rather than as a third-party reanimation, and it sits, structurally, in the same category as the Stones' AI-de-aging *In The Stars* video and the Apple/Animato infrastructure: avatar work deployed deliberately, with the talent's (or the estate's) consent, in a register the audience is asked to recognise as honest. The pattern is clear: the audience does not reject the AI-rendered celebrity. The audience rejects the *un-consented* AI-rendered celebrity. Consent infrastructure — institutional (RSL Media), statutory (ELVIS Act, the New York avatar law), platform (Apple's owned-pipeline play) and estate-level (the Osbourne announcement) — is what makes the difference.

The test case for exactly where that line sits arrived within weeks. **Netflix** announced it would use AI to recreate the voice of the late **Gene Wilder** for a reality-competition series set in the world of *Willy Wonka & the Chocolate Factory*.⁵⁷⁷ Structurally, this is the Osbourne architecture — a posthumous voice, commercially deployed. But the register is different: not an estate continuing the artist's own public presence on the artist's own terms, but a platform reaching for a dead performer's most beloved characterisation to warm up an unscripted format. Whether audiences receive it as the Osbourne case or as the Portman case will tell us whether the consent infrastructure is doing the *cultural* work — establishing which registers feel honest — or only the legal work of allocating who gets paid.

Underneath the celebrity layer, the structural infrastructure was being built. The **ELVIS Act**, Tennessee's AI-impersonation law, had been used by the **Johnny Cash estate** to sue Coca-Cola over a tribute-act ad soundtrack.⁵⁷⁸ **New York** passed a law in December 2025 forcing advertisers to disclose when they were using AI avatars. The SAG-AFTRA statement on the law's passage captured the political theory underneath the moment: *"These protections are the direct result of artists, lawmakers and advocates coming together to confront the very real and immediate risks posed by unchecked AI use."*⁵⁷⁹ By June 2026 SAG-AFTRA had scaled the same coalition politics to the federal level: **more than 16,000 signatories** on an open letter demanding Congress pass the **NO FAKES Act** — the federal statute that would ban unauthorised AI images and videos of real people, and consolidate the state-by-state patchwork (Tennessee's **ELVIS Act**, New York's disclosure law) into a national right.⁵⁸⁰ **Governments around the world** were considering bans on Grok's app over an AI sexual-image scandal that broke in early 2026.⁵⁸¹ By **May 2026**, the **AI Disclosure Standard** had been launched at the **Cannes Film Festival** as an industry coordination point for production-side AI labelling.⁵⁸² The **Academy of Motion Picture Arts and Sciences** had — in a quietly consequential rule update — set the line *"You must be human to win"* for its 2026 awards.⁵⁸³ The **Emmys** had set their own AI guidelines. The Television Academy's language was a model of how to write a policy that defends authorship without picking a fight with the toolchain: *"The Television Academy reserves the right to inquire about the use of AI in submissions. The core of our recognition remains centered on human storytelling, regardless of the tools used to bring it to life."*⁵⁸⁴ *Tools used to bring it to life — not tools that did the work.* The grammar matters. **SAG-AFTRA's** four-year contract — finalised by spring 2026 — included what the trade press informally called the **Tilly Tax**: a structured set of provisions for compensation, consent and residuals when AI replicas of human performers are used.⁵⁸⁵

Each of these is, on its own, a marginal piece of policy. Stacked together, they describe a new economic landscape: one in which *human authorship and identity* have become legally protected categories of creative work, with specific procedural and economic mechanisms for asserting them, defending them and compensating their use.

The cultural shorthand for this — *authenticity as the new scarcity* — captures the supply-and-demand logic. The legal shorthand — *human-authored work as a protected class* — captures the policy logic. Both are the same thing seen from different angles.

The financial reality: billion-dollar valuations alongside billion-dollar lawsuits

The Authenticity Premium, as I have described it throughout this chapter, is a cultural and market signal. It is also a financial one — and the financial picture, by June 2026, has a shape that the 2024-era “will AI destroy the music industry?” framing did not predict.

Suno, the AI music generation company that the music-rights organisations described as “the biggest theft in music history” when they sued it, raised **\$400 million** in June 2026 at a **\$5.4 billion valuation**. The round was led by Bond Capital. The company’s annual recurring revenue at the time of the raise was **\$300 million**; it had **2 million+ subscribers**.⁵⁸⁶ The headline that accompanied the raise told the financial story plainly: Suno is a multi-billion-dollar company that is simultaneously facing active copyright litigation from Universal Music Group and Sony Music, and has already settled with Warner Music Group — a settlement that included a **licensing partnership** rather than just a damages payment.⁵⁸⁷

That combination — \$5.4B valuation, active lawsuits from UMG and Sony, settled partnership with Warner — is the financial shape of the post-Petrillo AI music settlement in real time. Warner’s calculation is that a licensing relationship with a \$5.4B company is worth more than the continuation of a lawsuit against it. UMG and Sony are making the opposite bet: that the litigation creates enough pressure to extract better terms than Warner got, or to establish the precedent that training on their catalogues without consent is not permissible at any price. Both bets may be rational. The point, for this chapter, is what the \$5.4B itself means. The market is saying, with extraordinary clarity, that the Authenticity Premium does *not* prevent the AI-music tools from attracting capital at billion-dollar valuations. The demand floor for AI-generated music — the content-farm market, the B2B brand-music market, the *Tonada*-style retail-background-music market — is large enough to sustain a \$5.4B business even if the premium streaming market remains, as the Deezer and Apple Music data suggest, almost entirely human-authored.

The Scorsese/Black Forest Labs episode I describe in Chapter 7 reads differently in this financial context. When Scorsese became an advisor to Black Forest Labs — the FLUX model company — the backlash from Boots Riley and Guillermo del Toro was framed, in the press, as a *cultural* disagreement. But the underlying structure is, on closer inspection, an *Authenticity Premium* argument in action. Scorsese’s cultural authority derives, in large part, from the audience’s belief that his work has been made the hard way — that it is, in Jenna Ortega’s phrasing, the work the audience was *deprived of*, and has learned to hunger for. When he endorses a generative-AI company, the backlash is not a reaction to the endorsement’s legal implications. It is a reaction to the *signal* the endorsement sends about the source of his authority. The audience — and the peers who pushed back — is protecting the asset. That is, in market terms, exactly what the Authenticity Premium predicts.

What sincerity looks like in 2026

I want to come back to a distinction I made in Chapter 4, because it has held up through the last six months better than almost any other framing in this book.

I argued that audiences distinguish, very quickly, between *sincere* synthetic work and *cynical* synthetic work — that the underlying technology is the same, but the fingerprint of human intent behind the work is visible to the audience at the speed of a swipe.

The data from the spring of 2026 supports this. *Marketing Week*’s analysis “You can’t dismiss AI ads as slop when they’re winning in testing”⁵⁸⁸ documented that AI-generated advertising creative could, in fact, win in standard creative-effectiveness tests — *when the work was made with care, on a brief that respected the audience, by a team that had taste*. The same publication’s parallel coverage of the audience pushback against the McDonald’s Netherlands ad, the Valentino handbag campaign and a dozen other “AI slop” launches, made the inverse point. The technology is neutral. The intent is not. The May 2026 **David Beckham / Lenovo** “*Henchester United*” ad — in which the brand cheerfully used generative AI to render a Beckham-designed chicken coop for the footballer’s home flock — is, in my read, the cleanest *sincere* example of the period: the AI is visible, the brief is playful, the talent is in on the joke, and the cultural response was warm rather than wary.⁵⁸⁹ The Beckham/Lenovo spot belongs in the same category as the Madonna/Higgsfield piece in the section above: AI deployed deliberately, with the talent’s consent, in a register the audience recognises as honest.

The strongest AI-authored creative work of the period this book covers has, almost without exception, *not* tried to hide that it was AI-authored. Andrii Daniels’ bomb-shelter clip foregrounded its conditions of making.⁵⁹⁰ Hoyt Dwyer’s animated short for AI FilmFest Japan was upfront about its medium.⁵⁹¹ *Dear Upstairs Neighbors*,

the Google DeepMind / Connie He collaboration that premiered at Sundance, was *about* the constraints and possibilities of its production pipeline.⁵⁹² *Synthetic Sincerity*, Marc Isaacs' IDFA film, took the disclosure to the title of the piece.⁵⁹³ *Watch the Skies*, the AI-dubbed Swedish UFO feature, disclosed the dubbing process as part of its identity.⁵⁹⁴ *Lily*, the \$1m AI Film Award-winning Tunisian short, was framed by its director and reviewers as a piece *about* the new toolchain.⁵⁹⁵

The pattern, repeated across thirty or forty examples I have looked at carefully, is the same: *AI work that owns its synthetic nature, and that is made with human creative intent, finds an audience.* AI work that tries to pass as something it isn't gets the audience response that Tiny Grandma's stop-motion got *from the algorithm* — an immediate, automatic, suspicious flag.

This is, in market terms, a stable equilibrium. It is the market that the slop ceiling and the audience pushback have built. And it is, for working creatives, a manageable and even encouraging environment to operate in. The audience is not against AI. The audience is against being lied to.

What disclosure should look like

I want to lay out — because I have been asked this in every Q&A I have done since starting the newsletter — what I think the *practical* shape of authenticity infrastructure should look like for working creatives in 2026.

It is four things, in increasing order of investment:

One. Disclose, consistently. If you use AI in any part of your work, say so. In your credits. On your website. In your contracts with clients. In the metadata of your files. The act of disclosure does, in my experience, not cost you anything with the audience — the audience that is going to reject AI work would reject it anyway, and the audience that is going to accept it is the audience that values you being straight with them. The cost of *getting caught* not disclosing, in this environment, is materially higher than the cost of disclosing.

The scale evidence that disclosure is becoming routine rather than reputational poison arrived from the games industry in June 2026: of the more than 300 games released on Steam in a single week, **120 of them carried AI disclosures** under Valve's mandatory regime.⁵⁹⁶ Disclosure at that volume — two in five of everything shipped — is no longer a confession; it is a filing requirement, the games-industry equivalent of an ingredients label. The counter-argument deserves recording too: Epic's Tim Sweeney publicly criticised Valve's disclosure policy weeks later, arguing that access to AI tools is essential to smaller studios' survival in a competitive landscape and that a labelling regime risks stigmatising exactly the developers who most need the tools.⁵⁹⁷ Sweeney is making a real point about who bears the cost of a

label. But notice that it is an argument about the *consequences* of disclosure, not against the audience's right to know — and on the evidence of the Steam numbers, the audience's right to know is winning.

Two. Document, deliberately. Keep logs. Keep notes. Keep prompt histories. If a piece of work you make this year ends up being legally or culturally contested in 2030 — and a non-trivial fraction of work made this year will be — your ability to *show your workings* will be the difference between defending the work and losing it. The Sundance literacy initiative's emphasis on *evidence of human authorship* is exactly right.⁵⁹⁸

Three. Watermark, where appropriate. Use SynthID, C2PA, or the equivalent provenance layer that your toolchain supports. If your work doesn't yet support these standards, ask your tool vendors when they will. The market for tools that support provenance metadata is, in 2026, larger than the market for tools that don't.

Four. Build the chain. If you are running a studio or an agency, build the *internal* infrastructure for verifying and tracking the provenance of your work end-to-end. The cost of doing this in 2026 is moderate. The cost of *not* doing it in 2029, when a client asks for the chain-of-custody on a piece of work and you can't produce it, is going to be much higher.

These are not, on their own, business strategies. They are, in 2026, the *minimum hygiene* for operating a credible creative practice in the AI era. Treat them as you would treat health-and-safety on a film set. Do them as a default. Do them well. Then get on with the work.

The provenance infrastructure, named: forty-four pieces of the puzzle

I want to lay out a more complete map of the provenance infrastructure that is being built in 2025–26, because the technical-and-policy stack is more advanced than the public conversation has caught up with, and working creatives reading this book need to know what is actually in the field.

Stacking the moves I have referenced across this chapter and the rest of the book, the inventory is roughly this:

Capture-time signing and provenance metadata:

1. **C2PA / Content Credentials** (Adobe-led, with Microsoft, Sony, Nikon, Leica, BBC, Canon participation) — the cryptographic-signature standard for capture-and-edit-time provenance.
2. **Leica M11-P / M11-D** — first consumer cameras with C2PA at the firmware level (2023, expanded through 2025).
3. **Sony Alpha 1 II and Alpha 9 III firmware** — C2PA support across pro Alpha range.

4. **Nikon Z9 firmware** — C2PA support for the AP/Reuters wire-service workflow.
5. **Canon professional bodies** — Content Credentials integration through 2025 firmware cycle.
6. **iPhone Camera app** (selected models, 2025–26) — capture-time signature support.
7. **Adobe Photoshop, Premiere, Lightroom** — edit-time chain-of-custody preservation through Content Credentials toolchain.
8. **DaVinci Resolve, Final Cut Pro, Avid Media Composer** — partial C2PA support through 2025–26 update cycle.
9. **Capture One** — Content Credentials integration for the high-end commercial-photography workflow.

Synthetic watermarking and detection:

10. **Google DeepMind SynthID** — embedded watermark for Veo (video), Lyria (audio) and Imagen (image) outputs. Over 100 billion items marked by May 2026; extended to OpenAI, ElevenLabs and Kakao as a cross-vendor standard, with the ElevenLabs detection integration live from June 2026.
11. **SynthID Verification via Gemini app** — consumer-facing yes/no detection.
12. **Lyria 3 SynthID extension** — audio verification across the Google music-model family.
13. **YouTube AI Detection Tool** — automated content classification (the *Tiny Grandma* false-positive case demonstrating both its scope and its current accuracy ceiling).
14. **Deezer’s AI-music detection pipeline** — identifies up to 75,000 AI-generated tracks per day at upload time; productised in June 2026 as a detection service offered to *other* streaming platforms.
15. **Spotify AI Transparency Beta** — voluntary creator disclosure surfaced in the consumer player UI.
16. **Beeble** — independent detection-and-watermarking infrastructure used by some news organisations.
17. **Cloudflare AI bot classification** — public-web-infrastructure-level tracking of AI crawlers and agents.
18. **Sony music-identification technology** — identifying source recordings inside AI-generated outputs at the catalogue level.
19. **Sureel** — AI-detection company acquired outright by Warner Music Group in June 2026: major-label-owned detection infrastructure.
20. **Midnight Labs** — AI copyright-protection firm, Sony-backed from June 2026.
21. **RTM Audio’s UAI detector** — AI-music detection that issues a signed certificate with every verdict.

- 22. **Played By Humans** — verification service for “human artistry in an AI age”.
- 23. **Traxsource human-made / AI-assisted catalogue labels** — store-level labelling of provenance at the point of sale.

Institutional and contractual disclosure:

- 24. **Cannes AI Disclosure Standard** — industry-coordination production-side labelling standard for the festival circuit (launched May 2026).
- 25. **Sundance AI Literacy Initiative** — creator-training programme funded by Google’s \$2M commitment, training 100,000+ artists in foundational AI literacy and provenance practice.
- 26. **Academy of Motion Picture Arts and Sciences “you must be human to win” rule** — 2026 awards eligibility update.
- 27. **Television Academy AI guidelines** — “tools used to bring it to life” framing for Emmys submissions.
- 28. **SAG-AFTRA Tilly Tax provisions** — the consent, compensation and residuals framework for digital-replica use, included in the 2026 four-year contract.
- 29. **Equity (UK) strike ballot outcome** — 99% vote authorising industrial action over AI provisions.
- 30. **PRS for Music AI Survey 2026** — UK music-creator sentiment baseline informing collective-licensing negotiations.
- 31. **GEMA ruling against OpenAI** — first European-rights-society legal precedent on AI training compensation.
- 32. **The 88% UK consultation outcome** — political mandate for licensing-by-default.
- 33. **The *Stealing Our Work Is Not Innovation* declaration** — 800-creator cultural marker.
- 34. **Bandcamp’s outright AI-music ban** — distribution-platform-level disclosure-by-exclusion.
- 35. **Swedish Top Chart AI ban** — chart-eligibility-level disclosure.
- 36. **San Diego Comic-Con AI art ban** — convention-level disclosure-by-exclusion.
- 37. **Valve’s mandatory Steam AI-disclosure regime** — platform-scale disclosure in games (120 of 300+ releases in a single June 2026 week carrying AI disclosures), with Epic’s public criticism marking the policy fault line.

Legal infrastructure protecting human identity:

- 38. **Tennessee ELVIS Act** — the most-cited state-level AI-impersonation statute, used by the Johnny Cash estate in the Coca-Cola tribute-act case.

39. **New York's December 2025 AI-avatar disclosure law** — requires advertisers to disclose AI-performer use.
40. **The voice-and-image trademark filings** — Taylor Swift (early 2026), then Lionel Richie and the Backstreet Boys (June 2026): celebrity-led use of the trademark mechanism for identity protection, now a pattern rather than a one-off.
41. **The Jeremy Renner unauthorised-voice lawsuit** — pending case testing voice-likeness protections.
42. **The NO FAKES Act push** — SAG-AFTRA's 16,000-signature open letter demanding a federal statute banning unauthorised AI images and videos (June 2026).
43. **The Human Consent Registry** — the talent-side consent register publicly championed by Cate Blanchett, Steven Soderbergh and Nikki Hexum (July 2026), the working continuation of the RSL Media thread.
44. **The pending EU AI Act enforcement** — training-data-transparency requirements coming into effect through 2026.

Each of these, on its own, is a marginal piece. Stacked together — capture-signing, watermarking, platform integration, festival rules, awards rules, union contracts, civil society declarations, legal protections — they describe a *coherent infrastructure project* that the creative industries are, in eight months, jointly constructing.

The project is, by any reasonable assessment of similar previous infrastructure builds, *substantially ahead of schedule*. The C2PA standards body was founded in 2021 and was, by mid-2026, deployed across most major commercial capture and edit tooling. SynthID went from research demo in 2023 to consumer-facing detection in Gemini by January 2026. The SAG-AFTRA digital-replica provisions went from a 2023 strike demand to a contractual reality in 2026. The 88% went from political abstraction to government statement of progress in twelve months.

The thing this rate of progress tells me is *not* that the work is done. The work is, in many places, half-done — there are gaps in adversarial robustness, in platform UI integration, in cross-jurisdictional enforcement, in coverage of the long tail of creator categories outside the major commercial industries. The work is also being done unevenly: the music industry has built more of the stack than the games industry, which has built more than the publishing industry, which has built more than the regional and minority-language creative ecosystems that the next decade will need to bring in.

But the *trajectory* of the work is unambiguous. The provenance stack is being built. The institutional disclosure infrastructure is being built. The legal protections are being built. The audience contract I describe in the next section is being written. Working creatives who position themselves on the *inside* of this build — using the

tools, contributing to the standards, showing up at the consultations, advocating with the unions, deploying provenance metadata in their own work as a default — will have, by 2030, materially more leverage than working creatives who waited for someone else to finish the project for them.

What this means for the audience

I want to close the chapter with a thought about what this whole structure means for *the audience*, because most of this book has — by design — been about the people who *make* creative work, and the audience is sitting on the other side of the screen the whole time.

What I think the slop ceiling, the provenance infrastructure, the disclosure norms and the legal protections are, *collectively*, building is a *new contract* between makers and audiences.

The old contract was straightforward. The maker made the thing. The audience watched, listened, played, read. The signal of authenticity was implicit — most creative work was, by default, made by humans because there was no other way to make it.

The new contract is, by necessity, *explicit*. The maker discloses what was made by whom and how. The audience gets to make an informed choice. The platform, the union, the law and the institution all support both sides of the transaction.

If we get this contract right, the AI era is not the end of human creative work. It is a *renegotiation* of the terms on which human and synthetic creative work coexist in the public sphere — with the audience, for the first time in a very long time, getting a real seat at the table.

If we get it wrong — if the disclosure infrastructure fails, if the provenance metadata is unreliable, if the platforms refuse to honour the audience's stated preferences, if the legal protections are not enforced — what we get is the world the *Dead Internet* chapter described. A web of synthetic content, made by no one in particular, for no one in particular, churning past an audience that has lost the ability to trust any of it.

The choice between those two outcomes is *not, in 2026, a technical question*. The technical infrastructure for both is, by spring 2026, broadly in place. The choice between them is a *political, institutional and cultural* one. It is about whether the people who set the rules — the platforms, the legislators, the institutions, the studios, the audience itself — collectively decide that *knowable human authorship* is a public good worth protecting.

I think, on the evidence of the last six months, that the choice is being made — slowly, contentiously, imperfectly, but recognisably — in the right direction. The 88%, the Sundance literacy turn, the Cannes Disclosure Standard, the Academy's

rule update, the SAG-AFTRA contract, the C2PA standards, the SynthID rollout, the audience's own attention behaviour: these all point the same way.

The philosophical argument underneath all of this is the case that when technical execution becomes a commodity, the distinctively human acts of *intent*, *risk-taking*, and *the willingness to mean something* become the only irreplaceable contribution. The darker side of the same picture is the *process trap*: why the reassuring story that “automation will take the rote but real creative work is safe” is simultaneously true of the *capacity* and catastrophic for the *workforce class*.

The question for the rest of this book — Chapter 14 on the organisational restructuring, Chapter 15 on the labour-market reshuffle, and Chapter 16 on the political choice — is what happens to the *organisations*, the *labour market* and the *economy* of creative work when authenticity is the scarce good and the orchestrator is the new role. The implications for how teams are structured, how labour is paid and how creative careers are built are bigger than any single tool launch, and they are what the next three chapters are about.

Chapter 14 — Coordination Collapse

There is a working assumption underneath almost every conversation about AI in the creative industries that I want, in this chapter, to make explicit and then take apart.

The assumption is that AI is a *technology change*, broadly equivalent in shape to other technology changes the industries have absorbed in the past — the arrival of digital, the move to streaming, the rise of mobile, the emergence of social. The implication is that the existing institutions of the industry — the studios, the agencies, the labels, the unions, the publishing houses, the broadcasters — will absorb this change the way they absorbed the previous ones: with some restructuring, some layoffs, some new hires, some new departments, and a generally familiar shape on the other side.

I do not think this is what is happening.

What I think is happening is that AI is, specifically and quite differently, a *coordination technology*. It is changing — not at the margin, but at the core — what it is possible for a single person to know, decide and execute about a complex creative project. And because the existing organisations of the creative industries are, structurally, *coordination architectures* — they exist to allow many people to work on a single piece of work — the change in coordination economics is changing the organisations themselves.

This is the chapter about what happens to creative *organisations* — studios, agencies, labels, broadcasters, indie companies — when AI reaches a certain level of capability. It is also, by necessity, the chapter about what happens to creative *careers* when those organisations change shape.

The shorthand I have come to use for what is happening is **coordination collapse**.

What organisations are for

A studio, an agency, a label, a publishing house — these are not just brand names attached to creative outputs. They are *organisational technologies* that solve a specific problem: how do you get fifty or five hundred or five thousand people to coordinate on a single piece of creative work, well enough that the result is coherent, on-budget, on-deadline and good enough to put out into the world.

The way they solve that problem is by *layering* the work into specialised roles, building *hierarchies* that direct the work down through those roles, *processes* that move material between the roles in predictable order, and *cultural norms* that make the whole apparatus run with less explicit instruction than you would otherwise need.

A film studio exists, structurally, because making a film requires the coordinated work of many specialists — writers, directors, actors, cinematographers, designers,

editors, sound designers, composers, marketers, distributors. The studio is the *coordination apparatus*. The film is the output of the apparatus.

When AI starts to do the work of many of those specialists — not entirely, but at the level of *first draft* or *junior contribution* — the calculus of the coordination apparatus changes. Suddenly, a much smaller human team, working with a large pool of synthetic capacity, can produce the same coordinated output that previously required the large team. Suddenly, the *bottleneck* of producing creative work is no longer the size of the team. It is something else: the ability of the small senior team to *direct* the synthetic capacity well.

The implication, which has been dawning on the creative industries through the autumn of 2025 and the spring of 2026, is that the existing organisational shape — the *studio shape*, the agency shape, the label shape* — was built for a coordination problem that no longer exists in the form it used to.

This does not mean the studios will disappear. It does mean that the *shape* of the studios is going to change, in a hurry, in ways that working creatives need to understand if they are going to be on the inside of the change rather than on the receiving end of it.

The games industry, as usual, said this out loud before most sectors did. A Gamescom developer-speaker survey in August 2026 — polling the practitioners the industry’s flagship conference had chosen to put on its stages — identified AI’s likely impacts on *team structure* and *productivity* as the questions developers were actually wrestling with: not whether the tools work, but what shape the team around them takes.⁵⁹⁹ That is coordination collapse rendered as a conference agenda item.

The shadow workforce

The first symptom of coordination collapse, as it has shown up across the creative industries in this period, is the rise of what the workplace-research literature has started calling **shadow AI** — the practice of using AI tools in your job *without telling your employer*.

The numbers, from a series of 2025 studies I covered in *Dream Machine* Issue 5, are extraordinary.⁶⁰⁰

Roughly **half** of U.S. employees — 45–52% in different surveys — have used AI in their jobs without telling their bosses, with Gen Z and tech-sector workers being the most frequent secret users.⁶⁰¹

About a third — 29–33% — pay for their own AI tools out of their own pockets without their employer’s knowledge.⁶⁰²

Roughly **56–57%** of regular AI users admit to *hiding their usage or presenting AI output as their own work* to avoid judgement or stigma. Nearly half of executives do

the same.⁶⁰³

52% of workers won't admit to using AI at work — *even when asked directly*.⁶⁰⁴

These numbers describe a workforce that has, en masse, started running its own private parallel productivity infrastructure that bypasses the official organisational tooling, the official organisational processes, and the official organisational accounting of where the work is being done and by whom.

This is not a niche phenomenon. *Half the workforce*. And it is concentrated, the surveys suggest, in exactly the demographic — Gen Z, tech-sector, knowledge worker — that is most likely to be the future workforce of the creative industries.

The numbers escalate the more recent the research. By the end of 2025, enterprise-AI tracking data put active daily use at **88–89%** of staff across organisations, with **71–80%** of those users running their tools entirely outside any official approval or IT oversight.⁶⁰⁵ What the workplace-research firms have started calling the “Hidden Cloud Explosion” describes a six-month period in which the average enterprise IT department’s visibility into the AI tools its workforce was using simply collapsed: organisations believed they were running on roughly **91** public cloud services per enterprise, while network-level analysis put the actual figure at **1,220** active services — a 90% visibility gap.⁶⁰⁶ In the same year, **20%** of organisations reported severe security incidents linked directly to shadow AI, with the average breach cost going up by **\$670,000**; in **65%** of those incidents personally identifiable information was exposed, and in **40%** intellectual property was directly leaked.⁶⁰⁷

For the creative industries this matters disproportionately, because the data being fed into public LLMs by shadow users — proprietary scripts, unreleased concept art, client briefs, internal pipeline code, unmastered audio stems — is the *exact intellectual property* that organisations are simultaneously suing AI companies for scraping. The studio whose general counsel is in federal court against a frontier-model company is, on the same Tuesday afternoon, watching its own animation department paste asset descriptions into the same company’s consumer chatbot to speed up metadata writing. Both things are true. Both happen at once.

The shadow workforce, in coordination-collapse terms, is the symptom of an organisational architecture that is no longer aligned with the work the people inside it are actually doing. The official architecture says *we hired these humans to do these specific jobs, in this specific way, at this specific pace*. The shadow architecture says *these humans are now hybrid human-agent operators producing more, faster, with different qualitative properties than the official architecture is set up to manage*.

What you get, when these two architectures sit on top of each other, is a workforce that is *measurably more productive than the official metrics show*, doing more

work than the official scope says, with less institutional knowledge of how that work is being done than ever before.

This is, in coordination terms, an unstable equilibrium. It cannot last indefinitely. The question is how it resolves.

The “AI for thee, but not for me” paradox

The pattern that the shadow-AI numbers describe is not random distribution of tool use. It is *hierarchical*. Creative workers, across every survey I have read in this period, exhibit a consistent psychological pattern that the developer and creative-community discourse has named the “AI for thee, but not for me” paradox.⁶⁰⁸

The pattern works like this. Creative professionals identify some tasks as *mine* — the writing, the cinematography, the composing, the performance, the lead concept — and other tasks as *not mine* — the marketing copy, the project email, the deck assembly, the metadata tagging, the routine code, the contract redline, the rough mix, the asset variation. The first category is defended fiercely against AI substitution; the second is offered to AI substitution without much thought. The moral framing of the technology shifts depending on whose labour is being replaced.

Look at the music sector. An industry survey of more than **1,100** professional producers, songwriters and audio engineers in 2026 found that **87%** were actively using AI tools in their creative process.⁶⁰⁹ The internal distribution, though, tracked the hierarchy: **58%** used AI for audio restoration and cleanup, **38%** for mixing assistance, **33.9%** for automated mastering — high-friction tasks that nobody felt sentimental about — while only **20.9%** admitted to using AI for composition or lyric generation, the parts of the craft on which the personal artistic identity sat. **77%** cited “loss of originality” as their primary concern, outranking even the fear of personal job displacement (**42%**). The artist’s relationship to the tool, in other words, is not consistent across the work. It is sharply conditional on whose labour is being substituted.

The same hierarchy shows up in film. A survey of professional screenwriters before and after the 2023 WGA strike found that pre-strike covert AI use sat at around **34%**; once the WGA’s negotiated guidelines legitimised AI assistance for *formatting, structural outlining and brainstorming*, that number jumped to **68%** by 2024.⁶¹⁰ What the regulation changed was not the technology. It was the stigma. The shadow use moved into the light, with no measurable decline in the work product. The covert hierarchy became an overt one.

This is, in my view, the most uncomfortable observation in the entire shadow-AI literature, and it is the one that creative-organisation leadership has the hardest time admitting publicly. The same professionals who, in their public statements, treat AI training as a moral violation are, in their private practice, the heaviest users

of the same underlying technology. The hypocrisy is not a character defect. It is a structural property of how knowledge workers self-defensively triage their own tasks under productivity pressure. Read charitably, working creatives are doing what every economic actor in a productivity transition has done: protect the highest-value labour and offload everything else.

The cost of that triage is moral clarity. It is hard to credibly argue that AI training is theft when you are typing your portfolio description into Claude. The vocal-protest economy and the silent-adoption economy now run on the same desks, often within the same hour.

The consumption gap: what the data actually says

I want to spend a section on the gap, because the gap is *the* macroeconomic story of this period and I do not see anyone telling it cleanly.

The gap is between *what the creative industries are saying publicly about AI* and *how the creative industries are actually using AI on a Tuesday morning*. The two pictures are not slightly different. They are, in significant measure, contradictory.

Take Adobe, because Adobe is the cleanest single case study. Adobe's Firefly generative-AI suite — the same product that the working creatives in my surveyed circle are most ambivalent about — passed **22 billion AI-generated assets** by April 2025, eighteen months from public release.⁶¹¹ By that point, **45%** of all Creative Cloud subscribers had engaged with Firefly. **70%** of active Firefly users were using the tool *every week*, averaging 2.8 sessions weekly at 26 minutes each. Firefly contributed **11%** of all new annual recurring revenue at Adobe in 2024 — the company's fastest-growing revenue catalyst since the original move to a subscription model — and Adobe's AI-first ARR more than *tripled year-over-year* in the first quarter of fiscal 2026.⁶¹²

That is not the adoption curve of a niche professional tool. That is the adoption curve of a default productivity feature in the dominant creative-software stack on the planet. **72%** of Fortune 500 design teams have formally integrated Firefly. **63%** of marketing agencies. **58%** of e-commerce design departments. **48%** of UX/UI designers.⁶¹³ *Twenty-five per cent* of new Adobe Stock contributions in 2024 contained Firefly-generated elements.⁶¹⁴ The Adobe MAX 2025 Creators' Toolkit Report's headline number — **86%** of global creators using generative AI — sits inside this pattern, not against it.⁶¹⁵

If you sat with the public discourse alone — the open letters, the boycotts, the strike statements — you would assume working creatives were broadly refusing AI integration. The actual platform telemetry, in a year where Adobe shared more of its numbers than usual, says the opposite. Working creatives are not refusing. They are

adopting at a pace that Adobe's growth team is, by every signal I can read, struggling to keep ahead of.

The same picture holds across the toolchain. ChatGPT, by mid-2025, was on **800–900 million weekly active users**.⁶¹⁶ Anthropic's Claude was the writers' and developers' second favourite, with rapidly increasing usage in long-context creative tasks. Google's Gemini was growing desktop users at **155%** year-over-year, more than six times faster than ChatGPT's 23%.⁶¹⁷ These are not user numbers that reflect a market in revolt. They are user numbers that reflect a market that has, in private, decided.

The spend data points the same way — and adds a stratification the user numbers hide. The Ramp AI Index, drawn from actual corporate billing transactions rather than survey self-report, showed by June 2026 that the top 1% of firms were spending **\$7,500 per employee per month** on AI.⁶¹⁸ Read that figure against the median firm, which spends a small fraction of it, and the consumption gap resolves into a corporate arms race: the distance between the organisations that have made AI a core budget line and the organisations still running cautious pilots is widening, not closing. The shadow-AI numbers earlier in this chapter describe workers running ahead of their organisations. The Ramp numbers describe organisations running ahead of each other — and the creative organisations reading their own AI spend as generous should check which side of that distribution they are actually on.

A methodological caveat belongs on the record here, because I have leaned hard on Adobe's numbers in this section. When Adobe published another glowing AI survey in June 2026, *AppleInsider* pointed out that the sample largely left out the working creative industry it claimed to describe — the freelancers, the small studios, the practitioners outside the enterprise seats that Adobe finds easiest to survey.⁶¹⁹ The criticism is fair, and it cuts in a specific direction: the platform *telemetry* — assets generated, weekly active use, ARR — is hard data, but the *sentiment* layer reported on top of it is measured through samples that over-represent the already-converted. This is the legibility problem of the consumption gap. The private economy of AI use is real, but the instruments we have for measuring how the workforce *feels* about it are owned, in large part, by the companies selling the tools. The gap is real; its precise width is less knowable than either Adobe's marketing or its critics would like.

And the consumer side mirrors the producer side. The Stanford AI Index 2025 found that **55%** of individuals across 26 countries view AI products as offering more benefits than drawbacks — up from 52% in 2022.⁶²⁰ A 2024 YouGov poll across 17 markets found that nearly a third of consumers felt *more positively* about generative AI than the previous year, against only 22% feeling more negatively.⁶²¹ In the gaming sector — which has produced some of the loudest anti-AI consumer backlash of this year — the same Quantic Foundry survey that showed audiences are **77–83%**

negative toward AI-generated quests and dialogue also showed that **60%** of gamers remain entirely neutral about AI in a game's development *provided the final product is of high quality*.⁶²² The hostility is not generic. It is specifically aimed at AI in the creative roles where audiences expect to feel a human soul. Everywhere else — UI, backend, balancing, localisation, dynamic difficulty — the audience is, on aggregate, indifferent.

Even the GDC sentiment data, which is often cited as evidence of an industry in retreat from AI, tells the same paradoxical story when you read it as a whole. *Personal* generative-AI usage among professional game developers rose from **31%** in 2024 to **36%** in 2026, while *industry sentiment* over the same period cratered from 18% negative to **52%** negative.⁶²³ Use went up while approval went down. The two lines should, in a coherent market, move together. They are not. They are diverging.

I want to be careful about what conclusion to draw from this. It is *not* that the public discourse is wrong and the silent adopters are right. The public discourse is doing genuine political and cultural work — it is what produced the 88%, the SAG-AFTRA contract, the GEMA ruling, the Sundance literacy turn, the Cannes Disclosure Standard. Without the loud minority, the creative economy would have no political leverage at all.

The conclusion to draw is more uncomfortable. The creative industries are, in 2026, operating with *two parallel economies on top of each other*. In one, AI is a moral crisis, a labour threat, and a contested category of production. In the other, AI is a default productivity feature being integrated at the speed of any other software upgrade. The same individuals, the same teams, the same studios are participating in both economies simultaneously, often without acknowledging the contradiction.

The question for the next eighteen months — the question I keep coming back to when I talk to studio leadership — is whether the two economies merge into a single, honest, integrated practice (the **path two** integration I describe below), or whether they continue to run in parallel, with the public economy producing the policy and the private economy producing the work. The first outcome is harder but produces better collective decisions. The second outcome is the path of least resistance, and is, in my view, where we will end up by default if working creatives, studios and unions do not deliberately close the gap.

The data and the sectoral mechanics behind this section — the linguistic markers of covert AI use, the labour-market dynamics of agentic displacement, the deeper analysis of Adobe / OpenAI / Anthropic adoption telemetry, and the consumer sentiment / consumption asymmetry — sit in the DreamLab research this chapter draws on, as does the case against the *civilizational-risk* school that treats AI-driven coordination failure as an extinction scenario: the empirical critique of every major AI-doom prediction and why the data consistently refutes the worst-case framing.

The deeper organisational-design treatment of the *structural* argument this chapter is making — the 2,000-year history of hierarchy as an information-routing protocol, the *jagged technological frontier* identified by Dell’Acqua et al.’s preregistered BCG experiment, the *vigilance problem* that makes naïve human-in-the-loop oversight unreliable, the three competing post-hierarchical models (Block’s intelligence layer, Every’s agent colony, DreamLab’s Dynamic Agentic Mesh), the new KPIs for fluid organisations and the *judgment broker* role replacing middle management — rests on Dr. John O’Hare’s April 2026 sixteen-chapter research report, on which this chapter’s voice draws.⁶²⁴

The June 2026 Epidemic Sound *Future of the Creator Economy Report*, surveying 3,000 professional creators across the UK and US, sealed the picture and is worth pulling out because it is the cleanest single piece of evidence on the creator side I have seen this year. **94%** of creators now use AI tools; **72%** expect to *increase* their usage over the next twelve months; **57%** use AI primarily as a *workflow accelerant* rather than a generative content engine; only **28%** regularly produce fully generative content; and **89%** report feeling pressure to use AI *just to keep up with industry expectations* — the clearest articulation I have seen of the silent-adoption-under-public-protest pattern this chapter describes. The internal contradictions of the creator response to AI sit on top of each other in the same report: **75%** believe human-created content will become a *premium product* in the AI era, **83%** say human-made sound creates a stronger emotional connection with audiences, **79%** say music and sound directly impact engagement and revenue, **75%** believe disclosure is essential for audience trust, **73%** worry that unclear licensing could limit their future business opportunities, and **62%** have already experienced copyright or licensing issues in the past two years. *Musically’s* read of the same data — that creator attitudes towards AI have become “increasingly nuanced”, embracing the productivity gains while demanding greater transparency, clearer licensing frameworks and stronger protections around identity, voice and creative work — is the read that, in mid-2026, has finally caught up with the contradiction this chapter is built around.⁶²⁵ The simplest summary of the report is the one that sits underneath the entire second half of this book: **AI may drive scale, but long-term value will belong to creators who combine new tools with authentic human creativity and strong control over their intellectual property.**

There is one further dimension of the consumption gap I want to flag here, because Chapter 11 develops it at length and Chapter 4 introduced it: the gap between *production* and *consumption* is not just an organisational misalignment, it is a *biological* one. **Aggregate human attention is finite.** The same Adobe Firefly that has generated 22 billion assets, the same ChatGPT that serves 900 million weekly users, the same Sora app that hit a million downloads in five days — these are all systems whose *production* side scales without bound and whose *consumption*

side is bounded by the eleven-or-so daily hours of media attention the average adult can physiologically deploy. The consumption gap I have described above is, at its widest, this binding constraint expressed as an organisational problem. Studios that produce AI-augmented content at the rate the toolchain now allows — without recognising that the audience cannot consume more hours per day than it already does — are, on inspection, optimising the wrong side of the supply-demand equation. The studios that integrate AI productivity gains into work that *earns* a larger share of the audience's finite attention budget will win the next decade. The studios that integrate AI productivity gains into *more output competing for the same finite budget* will, on the historical pattern, hit the slop ceiling on a balance-sheet timeline they did not plan for.

The two paths

The shadow workforce can resolve, broadly, in one of two directions, and I think the choice between them will be the central organisational question for every studio, agency and label in the creative industries over the next three years.

Path one is *suppression*. The organisation decides that the shadow AI use is a risk — to security, to IP, to brand, to compliance, to the official productivity metrics — and shuts it down. Tightens the rules. Audits the work. Punishes the offenders. Reverts to the official architecture and the official tooling.

This is, in my view, a losing strategy in the medium term, because the productivity advantages that the shadow workforce is capturing are real, and the workers who are capturing them will, given the choice, work for organisations that *let them keep capturing them*. The suppressing organisation will progressively lose its most AI-fluent workforce to organisations that allow the hybrid practice.

Path two is *integration*. The organisation accepts that the shadow AI use is happening, decides to make it official, builds the infrastructure to support it, sets the norms to govern it, and re-shapes the work — and the workforce — around it.

This is, in my view, the right strategy. It is also the one most major creative organisations have been quietly moving towards in the period this book covers.

EA's push of its 15,000 employees to use AI as a "thought partner" was, structurally, a path-two move. **Krafton's** transformation into an "AI-first" company in November 2025 was a path-two move.⁶²⁶ **Disney's** Office of Technology Enablement was a path-two move. **WPP's** AI overhaul, **Adobe's** AI in everything, **Sony's** 49-agent game team — all path-two moves.

The path-two organisations are, structurally, betting that integration produces *more* output, *more* quality and *more* employee retention than suppression. The early evidence, six months in, suggests they are right.

Two refinements to the two-path picture arrived in the summer of 2026, and both are worth keeping in view. The first is that a path-two commitment is not an anything-goes commitment, and the most instructive counter-example came from open source. The **Godot** engine — one of the healthiest commons in games — announced in June that it would no longer accept AI-authored code contributions, on the bluntly stated grounds that “*we can’t trust heavy users of AI to understand their code enough to fix it.*”⁶²⁷ That is not suppression in the path-one sense. It is a governance line drawn where the institution’s survival depends on contributors understanding what they submit — a maintainability standard, not a moral panic. And it is precisely the kind of line a path-two organisation should study rather than dismiss, because the question Godot is asking — *does the human who shipped this understand it well enough to fix it?* — is the accountability question every integrating studio will eventually have to answer about its own pipelines. Integration without that question is not path two. It is abdication with better tooling.

The second refinement is that integration is a training problem before it is a procurement problem. *Fortune’s* mid-2026 survey of chief financial officers — HPE, Prologis, J.M. Smucker and Moody’s among them — landed on a warning that the finance function, of all places, has now internalised: AI success depends on training employees, not just buying technology.⁶²⁸ The CFOs are describing, in balance-sheet language, exactly what the shadow-AI data describes in behavioural language. The tooling spend is the easy half. The organisations that buy the seats without building the literacy get the cost line without the productivity line — and then, eighteen months later, conclude that “AI didn’t work here.” What makes this a solvable problem rather than a mystical one is that the literacy is teachable. KPMG’s analysis of 1.4 million AI interactions isolated the specific employee behaviours that separate effective AI use from ineffective use — and, crucially, showed that those behaviours can be identified, codified and taught at scale rather than left to whoever happens to be a natural.⁶²⁹ The skill that matters most, on one influential reading, is also the least technical. Sir Martin Sorrell — watching Cannes obsess, in his view, over the wrong AI debate, when the industry is already firmly in the AI era — argued that the decisive future capability is *sharing*: the people and organisations willing to pool prompts, workflows and learnings, he suggested, “will be the new kings and queens.”⁶³⁰ That is a training claim dressed as a cultural one, and it points at exactly the organisational muscle a path-two studio has to build: not just seats and models, but the norms that turn one person’s discovered workflow into everyone’s baseline.

The mid-career squeeze

The cost of the path-two transition has not, on the whole, been borne by the senior creatives or the entry-level workforce. It has been borne by the *middle*.

In **April 2026**, *Dream Machine* Issue 24 reported that the publisher behind *Grand Theft Auto VI* had laid off the entire seven-year-old internal AI team it had built to develop in-house AI capability for the franchise.⁶³¹ The framing, in the press release and the subsequent industry coverage, was that the company had decided to use off-the-shelf AI tools instead of maintaining proprietary ones — and that the seven-year AI investment was, in retrospect, a “backlash cleanup” cost.

The story was repeated across multiple studios. **Disney**, in April 2026, laid off staff including in its Marvel division, in moves the company did not blame on AI but whose timing was, as the trade press noted, “loaded.”⁶³² **Meta** had cut 10% of its Reality Labs staff in January 2026 to refocus on AI.⁶³³ **Microsoft** cut around 4,800 roles — 2.1% of its global workforce — in July 2026 as its Xbox unit prepared to spin out studios, in a round the trade press read explicitly as stoking fears that AI would replace people.⁶³⁴ **Scottish animation studio** Axis Animation collapsed in early 2026, with its closure publicly attributed in part to AI competition.⁶³⁵ **Ubisoft** cancelled five games, including the *Prince of Persia* remake, in January 2026, in order to refocus on AI.⁶³⁶

The pattern across these cases is the same: the *mid-career* layer — the experienced specialists in the middle of their professional lives, doing the day-to-day production work that the senior creative leadership directs — is the layer absorbed into agentic capacity first.

This is the unambiguous bad news of the AI transition. The *Guardian* covered this directly in January 2026 with a piece titled “AI is hitting UK harder than other big economies, study finds,” which found that mid-career knowledge workers in the U.K. were experiencing disproportionate displacement compared to peers in the U.S., Japan, Germany and Australia.⁶³⁷ *Economist* coverage in late November 2025 had been the early signal: “Investors expect AI use to soar. That’s not happening.”⁶³⁸ — meaning that the AI investment thesis was not, in the short term, producing the aggregate productivity gains the investors had hoped for, but *was* producing concentrated labour displacement in specific sectors.

The OpenAI public-policy response to this, articulated through April 2026, was a series of proposals — robot taxes, public wealth funds, a 4-day workweek — designed to manage the economic disruption of AI-driven productivity gains.⁶³⁹ *Dream Machine* Issue 24 covered these proposals at length. The framing OpenAI used was telling: the company was no longer arguing that AI would *not* cause disruption. It was arguing that the disruption was inevitable and that society needed to build new mechanisms to manage it.

The *Economist*, in a piece titled “Job apocalypse? Humbug! AI is creating brand new occupations,” took the contrary position — that AI was, on net, creating more new

jobs than it was destroying, and that the framing of mass displacement was overstated.⁶⁴⁰ Both positions are partially right. The aggregate employment numbers, across creative industries in 2026, did not show the apocalyptic decline some had predicted. But the *composition* of employment changed sharply. Senior orchestrator roles increased. Mid-career specialist roles decreased. New AI-specialist roles — AI agent developers, prompt engineers, AI ops specialists — exploded. The *Forbes* piece from November 2025 noted that “vibe coding” — natural-language software development — was an in-demand AI skill that paid up to \$220,000.⁶⁴¹

What the labour market is doing, when you look at it carefully, is *not* destroying jobs in the creative industries. It is *reshuffling* them — towards a smaller number of senior strategic roles, a different mix of specialist roles, and a much larger pool of AI-tooling skills that span the old discipline lines. The mid-career creative who fails to make this transition is the one who is at risk. The mid-career creative who makes it well — by upskilling deliberately, by claiming the orchestrator role, by building a portfolio practice — has, by every indicator I can see, *more* leverage in the labour market than they did before.

The hard truth is that the transition is not equally available to everyone. It depends on access to training, on access to tooling, on workplace cultures that support experimentation, on time to reskill that workers with caring responsibilities or financial precarity often don’t have. The institutional response to this — the Sundance Literacy initiative, the UK free-AI-training programme, the Adobe and Google educational investments — is real but partial. The structural inequities of who *can* make the transition are real and concerning.

The neurodiversity dividend

One genuinely encouraging finding from the period this book covers came from the U.K. Department for Business and Trade’s research on neurodiverse workers in AI-tooled workplaces.

The study, published in late 2025, found that workers with ADHD, autism and dyslexia were **25% more satisfied** with AI assistants than neurotypical workers, and that they reported AI agents as actively helping them succeed at work.⁶⁴² The interpretation, reported in *CNBC* in November 2025, was that AI tools were lowering the cognitive load of tasks that had historically been disproportionately punishing for neurodivergent workers — coordinating complex calendars, parsing dense documents, structuring written outputs — and were, as a result, *levelling the playing field* in workplaces that had previously underutilised neurodivergent talent.⁶⁴³

I want to flag this finding because it is one of the cleanest counterexamples to the “AI is bad for workers” framing that I have come across, and because it is a useful

corrective to the labour-displacement narrative that has dominated much of the coverage of this period.

AI is, demonstrably, *good* for some workers. It is good for the workers who are most able to leverage it, and it is also good for the workers whose existing labour-market participation was being limited by structural barriers that AI happens to dismantle. Both are real, and both are important to keep in view.

The *Guardian's* parallel finding — that ADHD, autism and dyslexia workers were reporting AI agents as a major workplace enabler — was echoed in dozens of smaller reports across 2026.⁶⁴⁴ The implication, for the creative industries: a workforce that has historically been heavily neurodivergent (the writing, music, film and games sectors are all over-indexed on neurodivergent talent compared to the general population) stands to be one of the *biggest beneficiaries* of well-deployed AI tooling in the workplace.

This is not a reason to ignore the displacement story. It is a reason to be careful about which framings of the AI transition are accurate and which are reductive.

Indies and the Global South

The other genuinely encouraging signal in this period is the rise of the *indie* and *Global South* creative sectors as direct beneficiaries of the AI cost reduction.

African film and tech has been a recurring positive story across the period — from Korin AI, the “trained with African datasets, built by Africans” model that launched in May 2026,⁶⁴⁵ to the wave of African AI filmmakers that the trade press began covering in earnest in early 2026, to the African music industry’s adoption of AI tools described in *CNBC Africa* in October 2025.⁶⁴⁶

Indian cinema has been awash with AI through the period covered by this book.⁶⁴⁷ The BBC’s December 2025 piece “Lights, camera, algorithm” documented the structural shift, with major productions integrating AI for visual effects, dubbing and asset generation, and made an observation about the limits of the technology that has stayed with me: “*You could create a sequel to a regional Indian movie using ChatGPT, but you would need to feed it the cultural memory of the original script. That script would have to be written by a human screenwriter.*” The cultural memory is the human contribution. The toolchain accelerates everything around it. The screenwriting itself — the act of *knowing what the culture remembers* — remains stubbornly, irreducibly human. India’s first AI-animated show, *Legenda Bertuah*, launched in Indonesia in April 2026.⁶⁴⁸

Latin American and Middle Eastern AI film festivals proliferated through late 2025 and early 2026. The \$1m Dubai AI Film Award was won by Tunisia’s *Lily*.⁶⁴⁹ Mexico’s Avocados-From-Mexico Super Bowl campaign was AI-led.⁶⁵⁰

Eastern European AI filmmaking — typified by Andrii Daniels’ bomb-shelter clip — became a recognised category.⁶⁵¹

East Asian AI development continued at a pace that, by spring 2026, had Chinese open-source AI models being used by approximately 80% of startups pitching the Andreessen Horowitz fund.⁶⁵² Korea’s Shift Up CEO described AI as the way to compete with Chinese game-industry scale, in language that captures both the geo-economic argument and what it means for individual workers: “*Only when all these people are proficient in AI, so that one person can perform the role of 100 people, can we compete with industries like China and the US that rely on large-scale human resources.*”⁶⁵³ *One person performing the role of 100.* That is the East Asian games industry’s framing of the orchestrator economy, and — if it is right — it tells you everything you need to know about the headcount maths every studio in the world will run between now and 2030.

The pattern, when you stand back from it, is what I have come to call the **Geographic Inversion**. AI is — in significant measure — *redistributing creative production capacity* away from the traditional centres (Hollywood, London, New York) towards regions that were historically capacity-constrained relative to their creative ambition. For most of the post-war period, geography concentrated creative work; for the first time in living memory, the technology is pushing the other way. This is not, on its own, a justification for everything AI is doing. It is, however, one of the clearest beneficial second-order effects of the cost reduction, and one of the most reliable signs that the creative economy that emerges on the other side of this transition will be — in geographic, demographic and economic terms — *less concentrated* than the one that preceded it.

If you are a working creative in a part of the world that has historically been on the wrong end of the global creative economy’s geography of access, the AI era is — for all its risks — also opening doors that were welded shut for most of the previous century.

What organisations should do

I want to close this chapter with a short and direct argument about what creative *organisations* — studios, agencies, labels, broadcasters — should be doing right now, because the readers I hear from most often, after working creatives, are people running creative organisations and trying to figure out the shape of the next three to five years.

The short version, drawn from everything I have read, watched and lived through these six months:

One. Move to integration, fast. The path-two organisations will outcompete the path-one organisations on talent, on output and on cultural capital within three

years. Suppression is not viable as a long-term strategy.

Two. Invest in your mid-career layer. The biggest source of avoidable damage in this transition is the loss of mid-career specialist knowledge that takes years to rebuild. Find ways to upskill your existing mid-career staff into orchestrator roles. The institutional knowledge they carry is the most valuable asset you have. Do not throw it away because the labour-cost arithmetic in a single quarter says you can.

Three. Solve the apprenticeship problem. The orchestrator economy structurally undermines the pipeline that has historically produced senior creatives. If you don't solve this — by maintaining some entry-level human roles, by building new AI-augmented apprenticeship pathways, by partnering with the institutes and the literacy initiatives — you are eating your own future. *Your senior creatives of 2035 are the juniors you hire today.* Treat them that way.

Four. Build the disclosure and provenance infrastructure. Chapter 13's argument applies as much to organisations as to individual creatives. The organisations that can credibly disclose their AI use, that maintain documentation, that can produce chain-of-custody on contested work, will be the organisations that the audience trusts in 2030.

Five. Build for the new geography. If your existing organisation is centred on the traditional creative capitals, the AI era is going to be much harder for you than for organisations distributed across the newly-accessible regions of the global creative economy. *Take seriously the option of building distributed teams* — not as a cost-saving move, but as a creative-capacity move. The talent is global. The tools are global. The audience is global. The organisations that don't adapt to this fact will lose their relevance to the ones that do.

Six. Don't outsource your judgement. This is the most important one and the easiest to get wrong. AI tools — even the very good ones, even the agentic ones, even the ones the platform companies are most eager to sell you — *cannot replace organisational judgement*. The decisions about what to make, who to hire, what to invest in, what to refuse — these are decisions that have to live with the humans running the organisation. AI can inform them. AI cannot make them. The organisations I have watched make the biggest unforced errors in the period this book covers are the ones that abdicated organisational judgement to the tools.

The shape of the creative economy in 2030 — what it produces, who it employs, where it operates, what it pays, what it is for — is being decided, right now, by the choices that the working creatives, the organisations and the institutions of the creative industries make in this twelve-to-eighteen-month window.

In Chapter 16 — once the next chapter has put the labour-market evidence on the table — I want to argue, as directly as I can, for the *kind* of creative economy I think we should be choosing. What a humane version of the AI-era creative economy looks

like. Who has to do what to get there. And what working creatives reading this book should be doing on Monday morning to play their part in it.

That choice is the last thing the book is about.

Chapter 15 — The New Jobs

There is a binary that I have, for six months, watched dominate every conversation about AI and creative employment, and that I am going to spend this chapter taking apart.

The binary is *jobs apocalypse* versus *jobs renaissance*. On one side, the visible argument: AI is coming for creative work, the trade unions are right, the displacement is real and accelerating, and a meaningful percentage of working creatives — particularly mid-career specialists in functions where AI has already become competent — will be out of the industry by 2030. On the other side, the equally visible argument: AI is creating more jobs than it destroys, the new categories of AI-orchestration work are paying more than the old ones, the freelance and indie sectors are expanding, the geographic boundaries of the creative economy are dissolving, and the next decade will be the most economically expansive period for creative labour since the post-war television boom.

Both arguments have evidence. Both are partially right. Both are, taken on their own, *wrong* — because the actual labour-market story of 2025–26 is not a binary. It is a *restructuring* with sharp winners and sharp losers, in which the dividing line between the two is not “AI” or “anti-AI.” The dividing line is **AI literacy** — the practical capacity to deploy generative tools as instruments of one’s own creative practice, with judgement, taste and structural understanding of where they help and where they harm.

This chapter is about that restructuring. About which jobs are disappearing, which are emerging, which are simply being *reshaped*, and what the working creative reading this should be doing — concretely, this year — to land on the right side of the line.

I want to spend most of the chapter on the evidence, because the binary framings have been driven, in my experience, by people who have not done the reading. The evidence is messier and more interesting than either side wants to admit.

What the headline numbers do not say

The aggregate employment numbers in the creative industries for 2024–26 did not show the apocalyptic collapse some had predicted. They also did not show the renaissance the platform companies’ marketing teams have been selling. *The Economist*, in a November 2025 piece titled “Investors expect AI use to soar. That’s not happening,” argued that the broad productivity-gain thesis was, in the short term, not playing out at scale across the wider knowledge economy.⁶⁵⁴ A month later, in “Job apocalypse? Humbug! AI is creating brand new occupations,” the same publication argued — using the same labour-market datasets — that AI was producing more new role categories than it was eliminating.⁶⁵⁵ Both pieces were

defensible. Both used real numbers. The two coexisted in the same magazine, six weeks apart, because the aggregate data is, on the current cut, ambiguous.

What the aggregate data hides is the *internal redistribution*. The mid-career specialist roles I described in Chapter 14 — the experienced sub-editors, the junior animators, the staff illustrators, the in-house copywriters, the routine production-pipeline engineers — are visibly contracting. The senior strategic roles — the showrunners, the lead creative directors, the senior orchestrators, the IP-fluent producers — are visibly expanding their *effective reach* if not their headcount. The new role categories — AI agent developers, prompt engineers, AI operations specialists, creative-AI ethics officers, model-curation specialists, AI-literacy trainers, custom-model fine-tuners, agentic workflow designers — are visibly growing from a near-zero base.

The Guardian’s “AI is hitting UK harder than other big economies, study finds,” from January 2026, found that mid-career UK knowledge workers were experiencing disproportionate displacement compared to peers in the U.S., Japan, Germany and Australia — but that the displacement was concentrated in specific *task categories*, not whole occupations.⁶⁵⁶ The University of Wisconsin-Stout’s January 2026 announcement, in which the institution set AI use as a *baseline competency* in its filmmaking course, captured the supply-side response: the curriculum was being re-engineered around the assumption that working filmmakers in 2030 would be AI-literate by default.⁶⁵⁷

The labour market, in other words, is doing what labour markets always do in a productivity transition. It is reshuffling. The reshuffle is sharper than the headline employment figures suggest because the *composition* of work is changing faster than the *amount*.

There is one further reading of the aggregate data worth holding onto, because it reframes the whole displacement question. *Labor Market Matters*, in a July 2026 analysis built on seven charts of independent-work data, argued that AI’s first labour-market effect may not be replacing workers at all — it may be making *traditional firms less necessary*, accelerating the long shift toward independent and freelance work.⁶⁵⁸ If that reading is right, then a good part of what shows up in the statistics as “job losses” is actually *employment changing shape*: the same work, and often the same worker, migrating out of the payroll column and into the independent one. For the creative industries — which were already the most freelance-heavy sector of the economy before the transition began — this is less a foreign invasion than an intensification of the existing weather. It also explains why the headline employment figures and the lived experience of working creatives can diverge so sharply: the firm-shaped jobs are contracting faster than the work is.

The roles that are disappearing

I want to be specific, because vague claims about “creative jobs disappearing” do not help anyone make a career decision. Based on the trade-press coverage tracked across the *Dream Machine* archive, the shadow-AI and adoption-dynamics survey data underlying DreamLab’s research, and the studio-leadership interviews I have read or conducted, the roles under the most active substitution pressure in 2026 are:

Junior visual production roles. Concept artists at the asset-variation level. Junior 3D modellers doing standard architectural / environmental fills. Storyboard artists working on commercial briefs that do not require performance staging. Background-plate compositors. Routine matte painters. Stock photographers and stock illustrators. Adobe’s own data — 25% of new Adobe Stock contributions in 2024 containing Firefly-generated elements⁶⁵⁹ — is the clearest single number in this category.

Junior writing and copy roles. In-house copywriters at brand agencies. Junior content marketers. Routine technical writers. Translation generalists where the source/target pair is well-resourced (English-Spanish, English-Chinese, etc.). SEO content writers. Sub-editors at digital publications. The Reuters Institute’s “AI adoption by UK journalists” survey found high integration across newsrooms by late 2025; the *Daily Mail*’s December 2025 report that Google’s AI Overviews had “killed click-throughs” to news sites was the consumer-side mirror of the production-side pressure.⁶⁶⁰

Mid-career routine production and post-production roles. Routine audio engineering (the 1,100-creator music survey discussed in Chapter 14 showed 58% of producers using AI for restoration, 38% for mixing assistance, 33.9% for automated mastering⁶⁶¹). Standard VFX compositing (62% of Hollywood studios on automated AI compositing, 35% reduction in post-production timelines⁶⁶²). De-aging specialists (200 hours per actor down to 50). Particle simulation specialists (68% adoption among top VFX houses by SIGGRAPH 2025). Routine matte-painting generalists (initial setup time from 4 hours to 1.2 per shot).

Routine games-development roles. Square Enix’s announced target — 70% of QA work via AI by end of 2027 — is the cleanest signal here.⁶⁶³ Falcom’s reported productivity ratio of *2-3 hours of work reduced to 10 minutes*⁶⁶⁴ tells you what is happening to the routine animation, asset and engineering layer underneath the games industry. Ubisoft’s January 2026 cancellation of five games (including the *Prince of Persia* remake) in order to refocus capital on AI signalled a structural shift in resource allocation that mid-career game developers are still digesting.⁶⁶⁵

In-house AI specialist teams at non-AI-native companies. This one is counter-intuitive but real. The publisher behind *Grand Theft Auto VI* laying off its

entire seven-year-old internal AI team in April 2026, in favour of off-the-shelf tooling, is the canonical case.⁶⁶⁶ Companies that built proprietary AI capabilities in 2018–2024 are increasingly finding that the open-weight and commercial foundation models have caught up; the bespoke AI team becomes redundant. This is a real and rapid form of AI-driven displacement that the public discourse has not yet recognised, because the workers being displaced are themselves AI specialists.

Voice actors and session musicians in commodity work. ElevenLabs’ growth to \$500m ARR by April 2026⁶⁶⁷ is, in large part, the substitution of routine voiceover, audiobook, podcast-host and dubbing work. Live performance, voice work that requires acting craft, and specialty voice roles (animation leads, signature character voices) are not displaced. The middle of the voice market is.

Routine commercial illustration and design. Brand assets, marketing imagery, social-media graphics, basic product visualisation. The Higgsfield growth curve — \$200M revenue in nine months, primarily serving social-media marketers⁶⁶⁸ — is the consumer-marketing equivalent of the Adobe Firefly enterprise curve.

The pattern across these categories is consistent: it is *routine* work, *specialist* work, and *mid-career* work that is under pressure. The pattern is *not* aimed at junior on-ramp roles (which is a problem of its own — see below) or at senior creative judgement roles. It is concentrated in the middle of the pipeline, in the layer historically occupied by experienced operators of the toolchain.

The roles that are emerging

The other side of the redistribution is equally real and substantially under-reported in the consumer press.

AI orchestrators / senior creative directors of agentic teams. The most strategically important new role, and the one Chapter 12 made the long-form case for. The Sony 49-Claude-agent / 72-skill stack is the canonical example.⁶⁶⁹ In adland, *Digiday* reported in late 2025 that “AI agent developers have become adland’s in-demand role”⁶⁷⁰ — a senior creative-strategic role that did not exist eighteen months earlier. By Cannes 2026 the demand signal had become theatrical: PMG ran a hackathon at the festival itself, using adland’s flagship week to recruit and showcase AI-fluent talent, and *Ad Age* read it, correctly, as a signal of agencies’ growing appetite for people who can *build* with the tools rather than merely approve work made with them.⁶⁷¹

Prompt engineers / AI workflow designers. The role is now broad enough to have its own specialisation tracks. *Forbes* reported in November 2025 that “vibe coding” — natural-language software development — paid up to **\$220,000** as an in-

demand AI skill.⁶⁷² The equivalent figures in the creative space — prompt engineers at major studios, freelance AI-workflow consultants — are running in the same band for senior practitioners.

AI literacy trainers and AI-education designers. The Sundance Institute’s AI Literacy Initiative, launched in January 2026 with \$2M of Google funding to train 100,000 filmmakers, is the institutional version of this role.⁶⁷³ The Adobe Ignite Day at Sundance, the UK government’s “Free AI training for all” programme covering 10 million workers by 2030,⁶⁷⁴ the Lovable-for-classrooms expansion,⁶⁷⁵ the UW-Stout baseline-AI competency course⁶⁷⁶ — these are the demand signal for a new category of educator that combines creative-discipline expertise with practical AI fluency. The pipeline keeps widening. In June 2026 Anthropic launched **Claude Corps**, a fellowship that will teach 1,000 fellows to master Claude and match them with nonprofits across America — a platform company building, in effect, a subsidised entry route into orchestration work, with the civic sector as the training ground.⁶⁷⁷ The same month, the UK published its **AI Adoption Plan for the Creative Industries**, the skills-and-adoption half of the government’s creative-AI policy (the copyright half belongs to Chapter 6), formalising sector-specific adoption support and training as national policy rather than platform philanthropy.⁶⁷⁸

Model curation specialists. With foundation models proliferating and custom fine-tuning becoming a baseline capability, the role of *selecting, training and maintaining* an organisation’s model stack has emerged as a discrete specialism. Adobe Firefly Foundry — the service that lets companies train custom generative models on their own visual identity⁶⁷⁹ — created an entire job category of brand-and-IP model trainers. Korin AI’s launch in May 2026, “trained with African datasets, built by Africans,”⁶⁸⁰ is the cultural-fluency variant of this role.

AI ethics, disclosure and provenance officers. Following the SAG-AFTRA contract negotiations, the Cannes AI Disclosure Standard, the Academy’s “you must be human to win” rule, the New York AI advertising disclosure law, and the proliferating C2PA-compliance and SynthID-tooling requirements, organisations across the creative industries have begun hiring (or designating) dedicated AI-ethics and disclosure leads.⁶⁸¹ At DreamLab, we have a *Continuum Lead* whose job is to make this work coherent across every project we run — three years ago, the role did not exist.

Indie and Global South creator-producers. The cost reduction in production tooling has created a new viable role category that was not economically possible before: the *one-person-or-small-team creator-producer* operating outside the traditional creative centres, with global distribution reach and a defensible aesthetic identity. *Forbes* covered the broad category in “AI Is Changing How Creators Work

And Earn” in December 2025.⁶⁸² The Higgsfield revenue (built on social-media marketer demand), the Andrii Daniels bomb-shelter clip (a Ukrainian one-person production with global reach⁶⁸³), the Tunisian *Lily* (\$1M Dubai AI Film Award winner⁶⁸⁴), the Indonesian *Legenda Bertuah* animated series,⁶⁸⁵ the Indian-cinema integration wave covered by the BBC’s “Lights, camera, algorithm”⁶⁸⁶ — these are not exceptions. They are the leading edge of a structural change in *who can be a working creative*, and where they can live.

AI training and evaluation work. The new job category with the bitterest edge. *The Hollywood Reporter* reported in June 2026 that Hollywood workers — writers prominent among them — are increasingly taking work *training AI models* as conventional job prospects grow slim: annotating scripts, evaluating outputs, teaching the systems the craft conventions of the industry contracting around them.⁶⁸⁷ The role is real, it pays, and it is growing fast. It is also the most uncomfortable entry on this list, for two reasons. The first is that the labour is directed at improving the very substitutes applying the displacement pressure — the screenwriter grading model outputs is, in the most literal sense, training their own competition. The second is who the role is absorbing: exactly the early-and-mid-career practitioners the apprenticeship pipeline used to carry. A junior writer spending their formative years rating completions is building a kind of AI literacy — but it is literacy in the service of someone else’s system, on someone else’s terms, with none of the mentorship, authorship or craft progression the old junior rungs carried. I list it as an emerging role because honesty requires it. I decline to celebrate it.

AI-augmented apprentices. This category is still being built, and is the central labour-market design question of the next three years (more below). The early models — AI-tool-augmented junior animator roles maintained deliberately at *Position Four* studios (Chapter 7), the Sundance Collab fellowship structure, the AI-augmented entry-level posts at WPP and the major Hollywood studios — are the early experiments. None of them, yet, has fully solved the apprenticeship problem.

Cross-disciplinary “portfolio creatives.” What I have been calling, in talks since the autumn, the **AI Literacy Premium** role — the working creative who, instead of a single specialism, holds several loosely-coupled creative disciplines together using AI as connective tissue. *TechBullion*’s “Why the future belongs to multi-skilled leaders,” from November 2025,⁶⁸⁸ and the Anthropic *Skills* framework underneath Claude Code’s multi-agent coordination,⁶⁸⁹ are the corporate-leadership and tooling-side manifestations of the same trend. The portfolio creative is increasingly the *default* career shape for working creatives entering the field today.

The AI literacy premium

If you stack the disappearing-roles list against the emerging-roles list, a single underlying variable does most of the explanatory work for who is winning and who is losing in the 2026 creative labour market.

The variable is **AI literacy** — defined operationally as the combined skill of (a) knowing what generative tools can and cannot do well, (b) being able to brief them effectively, (c) being able to judge their outputs with taste, (d) being able to integrate them into a coherent creative workflow, and (e) knowing where on the Continuum your craft sits and why.

The empirical case for the literacy premium is strong and consistent across the data:

The most direct number arrived in June 2026, and it put a price on the premium. PwC's analysis of the UK labour market, reported by *UKTN*, found AI hiring surging *across all UK sectors* — not just tech — with roles demanding AI skills carrying a **34% salary premium** over equivalent roles without them.⁶⁹⁰ I have been making the AI Literacy Premium argument since the autumn of 2025 as a structural claim about where leverage was moving. The PwC finding is that claim showing up in the payroll data: a third more, for the same discipline, in the same market, for the practitioner who can direct the tools. If a single statistic in this chapter deserves to be pinned above a working creative's desk, it is that one.

The Adobe Creators' Toolkit Report of October 2025, surveying 16,000 creators across the U.S., U.K., France, Germany, South Korea, Japan, India and Australia, found that **86%** of creators were already using generative AI; that **76%** of users said it had helped grow their business or brand; that **81%** said AI lets them make content they otherwise could not have made; and — most significantly — that **70%** were optimistic about agentic AI and **85%** would use AI that learned their creative style.⁶⁹¹ The numbers are not the numbers of a workforce in retreat. They are the numbers of a workforce *adopting at scale* and reporting commercial benefit.

The PRS for Music 2026 AI Survey, surveying U.K. music creators, found **four in five** music creators worried about AI-generated music competing in the streaming economy⁶⁹² — but also that respondents who were using AI tools in their own creative practice reported higher earnings and broader output reach than peers who were not. The two findings live together. The anxiety and the adoption are happening in the same head.

The U.K. Department for Business and Trade research, reported in *CNBC* in November 2025, found that workers with ADHD, autism and dyslexia were **25% more satisfied** with AI assistants than neurotypical workers, with AI agents reported as actively helping them succeed at work.⁶⁹³ Given that the creative industries are over-indexed on neurodivergent talent compared with the general

population, this is a major equity-and-productivity story that the AI-displacement narrative has under-told.

The McKinsey AI report on film and TV production, released in early 2026, made the broader business case: AI would not, in McKinsey's view, replace film and television production. It would *restructure* it — towards smaller teams, faster cycles, more iteration, and heavier reliance on senior creative judgement.⁶⁹⁴ In other words: towards an *orchestrator-shaped* industry, in which the working creatives best positioned to thrive are those who can operate the new toolchain rather than those who refuse to engage with it.

The GDC State of the Game Industry surveys 2024–2026 showed personal generative-AI usage among professional game developers rising from **31%** to **36%** while industry sentiment cratered from 18% negative to **52%** negative.⁶⁹⁵ The full trajectory of that sentiment number deserves spelling out, because it is one of the most striking curves in the entire labour dataset: **18%** of developers said generative AI was having a negative impact on the industry in the 2024 survey, **30%** in 2025, **52%** in 2026 — the negative share nearly tripling in two years, over exactly the period in which adoption *grew*.⁶⁹⁶ The interpretation: developers using AI are increasingly *uncomfortable* about doing so, but they are doing so anyway, because the productivity advantage is too strong to ignore. Sentiment and behaviour are running in opposite directions, and behaviour is winning.

The employer side of the same story is now being told in an unexpectedly warm register. EA told *PC Gamer* in June 2026 that AI is delivering “faster prototyping” and a “real rise in creativity” across its studios.⁶⁹⁷ Take the framing with the scepticism that employer testimony always deserves — “a rise in creativity” is doing public-relations work for a company that has been aggressive about AI integration — but notice what the claim concedes. The employer case for AI is no longer being made in cost-reduction language alone. Employers are competing to describe their AI-augmented workforces as *more creative*, which tells you, at minimum, what they believe the AI-literate practitioner is worth to them — and which side of the literacy line they intend to hire from.

The LANDR study reported in *Ari's Take*, late 2025, found that **87%** of musicians surveyed were using AI tools in some part of their practice.⁶⁹⁸ When the figure for “musicians using AI” approaches the figure for “musicians using DAWs,” you are not looking at a niche adoption pattern. You are looking at a baseline competency that the next generation of working musicians will be assumed to have.

The Stanford AI Index Report 2025 found that across 26 surveyed nations, **55%** of individuals viewed AI products as offering more benefits than drawbacks, up from 52% in 2022⁶⁹⁹; the YouGov 2024 multi-market sentiment data echoed the same

direction.⁷⁰⁰ Working creatives are sitting inside a consumer environment that is, on net, *increasingly accepting* of AI in creative work — which translates, slowly, into client demand for AI-augmented services.

The *Digital Music News* survey of nearly **800 creatives** who signed the *Stealing Our Work Is Not Innovation* declaration in January 2026⁷⁰¹ — that same vocal-minority signal that produced political wins in Chapter 6 — included, in its signatory base, a meaningful number of creators who were *themselves* using AI tools in their own production work. The signing of the declaration was not a refusal of the technology. It was an insistence that the *training data* be consented and compensated. Two positions, held at once, by the same working creatives.

Stack these findings together and the labour-market picture becomes much clearer than the binary discourse allows.

Working creatives who hold AI literacy are gaining leverage. Their effective output reach is larger. Their hourly rate is higher. Their client base is broader. Their geographic constraints are weaker. The orchestrator role, the portfolio-creative model, the literacy-trained newcomer — these are the positions on which the next decade of creative employment expansion is being built.

Working creatives without AI literacy are losing leverage. Their effective hourly rate is being competed down by AI-augmented peers. Their client base is being eroded by faster, cheaper, mostly-adequate substitutes. Their career mobility is contracting as mid-career specialist roles in their domains are absorbed into agentic capacity. Their geographic constraints (London, New York, LA, Tokyo as the centres) are being preserved while their AI-fluent competitors operate from anywhere.

This is, in practical terms, the structural transition the working creative reading this in 2026 is now facing. The choice is not “use AI” or “refuse AI.” The choice is whether to build the literacy that lets you remain in command of how AI shows up in your own practice, or whether to let the literacy build itself, badly, in your absence.

The Apprenticeship Gap

I want to name the most underdiscussed structural problem in the AI-era labour market, because it is — by my read — the single largest threat to the long-term health of the creative economy and it has nowhere near the public attention it deserves.

I call it the **Apprenticeship Gap**.

For the entire history of the creative industries — from the medieval guilds through the post-war Hollywood studio system through the rise of digital media — the *junior on-ramp* into creative work has been the structural foundation on which senior

talent is built. Junior writers become senior writers by writing things that nobody pays much attention to, repeatedly, for a decade, under the loose mentorship of more senior practitioners. Junior animators become senior animators by drawing the in-between frames, by cleaning up the rough animatics, by handling the routine asset variations that the lead artists do not have the time for. Junior cinematographers become senior cinematographers by holding focus, by pulling cable, by lighting the second-unit shot. The junior tasks were not, in themselves, the destination. They were the *training ground* on which judgement, craft and taste were built.

The orchestrator economy, the agentic toolchain, the GDC-data picture of senior practitioners using AI to absorb the routine middle-layer work — these patterns, taken together, are progressively *removing* the junior on-ramp from the industry. The junior writer is competing with Claude. The junior animator is competing with Cascadeur. The junior cinematographer is competing with Veo 3.1's plate generation. The junior coder is competing with Cursor and Copilot. In every case, the routine task that *used to be the entry point* into the discipline is now economically uncompetitive against an AI agent that does it in seconds for cents. The games press put the same point more bluntly in mid-2026: the AI game-dev boom has a dark side — *fewer jobs for the people starting out*.⁷⁰² The counter-current is real but small, and it runs through founding rather than hiring: the same tools that close the junior rung also lower the barrier to *building your own thing*, and 2026 produced both a teenage founder raising \$3.2m for an AI game-development startup and an eleven-year-old shipping a game he had built with AI himself.⁷⁰³ Democratised authorship is genuine; it is simply not, on its own, an apprenticeship system.

The studios are not — yet — replacing senior roles. They are absorbing the *junior layer* underneath the senior roles, and then telling themselves a story about how the new tools will free senior practitioners to focus on the *real* creative work.

The story is partially true. It is also dangerously incomplete. If the junior layer disappears for a decade, the *next generation of senior practitioners has nowhere to be trained*. The pipeline breaks. The 2035 cohort of senior creative directors, lead animators, showrunners, music producers, art directors — the people who, in 2026, would be five years into a junior career — will simply not exist at the volumes the industry needs. The discipline-specific knowledge, the embodied craft, the relationship-based mentorship — all of these were carried in the apprenticeship layer. Remove the layer, and you are eating the seed corn of the discipline.

This is not a speculative claim. It is the underlying logic of every “expensive mistake” interview a working studio leader has given to the trade press in this period — Charles Cecil at Revolution Software, Todd Howard at Bethesda, the Larian and Aardman and Jagex public positions. The senior practitioners are saying, in different

vocabularies, the same thing: *we used to teach the next generation by giving them the routine work. The routine work is gone. We have not solved the teaching problem.*

And the displaced junior layer is not waiting quietly while the industry deliberates. It is going to work for the other side. The Hollywood-workers-training-AI-models story I flagged in the emerging-roles list above⁷⁰⁴ is the Apprenticeship Gap expressed as a labour flow: the would-be apprentices of 2026, unable to find the junior rungs, are spending their formative years teaching the systems that absorbed the rungs beneath them. Every hour of craft knowledge that flows into the models through that channel is an hour that used to flow, through mentorship, into the next generation of practitioners. The gap is not just failing to train the 2035 cohort. It is actively transferring the training substrate somewhere else.

The institutional response to date is partial and patchy. The Sundance literacy initiative is real. The Adobe Ignite Day, the Sundance Collab fellowships, the UK Free AI Training for All programme, the UW-Stout curricular changes — these are the visible institutional moves. But they are training programmes for *AI literacy specifically*, not full apprenticeship pipelines for the underlying creative discipline. They produce literate orchestrators. They do not, by themselves, produce the cinematographers, the composers, the writers, the lead artists, the showrunners of 2040.

The deepest structural reform that needs to happen in the next eighteen months is the deliberate preservation — or rebuilding, or re-imagining — of the apprenticeship layer. Some of this is already happening:

Position Four studios maintaining junior roles by policy. Aardman, Larian, Games Workshop, Jagex have, in different forms, made deliberate commitments to keep junior human roles in their pipelines even where AI could absorb them, on the explicit grounds that the future of the studio requires it. This is the most encouraging single trend I have observed.

Hybrid apprenticeship pathways. The new role I called *AI-augmented apprentices* above. A junior animator who uses Cascadeur, but who is *paired with a senior animator who teaches them why* the AI's output is good or bad. The juniors are not doing the in-between frames any more. They are doing the *judgement* on the in-between frames the AI produces, and the senior teaches them how to judge. This is a real model. It is not yet at scale.

Institutional reinvestment. The cultural-institution training programmes — Sundance, the BFI, the national film schools, the BBC training schemes, the Royal College of Art and equivalent national schools — are, in different ways, recalibrating to deliver more comprehensive training in shorter timeframes, on the assumption that the years-on-the-job apprenticeship period is contracting.

Public funding interventions. The UK government’s *Free AI training for all* programme, the EU’s various creator-skills initiatives, and the U.S. state-level training credits being attached to AI-investment incentives are early signs that the apprenticeship gap is being recognised as a public-policy problem rather than a market problem.

None of these is, on its own, sufficient. The Apprenticeship Gap is — by my prediction — going to be the single largest *unresolved* labour-market issue in the creative economy of 2030. The studios, the unions, the schools and the platform companies are going to have to figure it out together. The book’s Chapter 16 manifesto in *Choosing the Future* lists it explicitly as one of the questions every working organisation has to engage with.

If you are reading this as a senior practitioner: maintain juniors. Pair them with the new tools deliberately. Treat their on-ramp as a public good your industry depends on.

If you are reading this as a junior practitioner: the on-ramp is contracting. You do not get to wait. The literacy you build *in the next eighteen months* will determine whether the on-ramp closes before you are inside it.

The geographic redistribution

The second-largest under-reported labour-market story of this period is geographic.

For the entire post-war history of the global creative industries, professional creative employment has been concentrated in a small number of cities — Los Angeles, New York, London, Paris, Tokyo, Mumbai, with smaller secondary nodes in Berlin, Toronto, Seoul, Sydney. The geography was a function of the *cost structure* of creative production: studios, equipment, distribution networks, talent pools and capital all concentrated in the cities that could afford to host them, and the working creatives followed.

The AI cost reduction is, structurally, dismantling this geography.

The Tunisian-made *Lily* winning the \$1M Dubai AI Film Award.⁷⁰⁵ The Ukrainian one-person bomb-shelter production going viral globally.⁷⁰⁶ The Indian-cinema AI integration wave covered by the BBC, with productions across regional cinema centres absorbing the new toolchain at scale.⁷⁰⁷ The Indonesian *Legenda Bertuah* AI-animated series.⁷⁰⁸ The Korin AI launch — Africa-trained, Africa-built foundation model — and the broader CNBC Africa coverage of AI in African music and film.⁷⁰⁹ The Singapore-based AI video startup Video Rebirth raising \$50M for studio-grade tooling. The Eastern European AI filmmaker community building around the success of creators like Andrii Daniels. The Latin American AI-film festival wave through 2026.

These are not isolated stories. They are the leading edge of a redistribution that, taken together, is *materially expanding* the global creative workforce beyond the previous-century cities. The reduction in cost-of-entry for serious creative production has, for the first time since the rise of cinema, made it economically viable to build a competitive creative practice from places that the previous geography had locked out.

The Shift Up CEO's framing, in the *PocketGamer.biz* piece — “*Only when all these people are proficient in AI, so that one person can perform the role of 100 people, can we compete with industries like China and the US that rely on large-scale human resources*”¹⁰ — is the strategic reading from inside the Korean games industry. The same logic applies in Mexico, in Egypt, in Nigeria, in Brazil, in Vietnam, in any creative economy that has historically been resource-constrained relative to its ambitions.

For the working creative reading this in one of those geographies: the AI era is, for all its disruption risks, also opening doors that were welded shut for most of the previous century. The labour market is becoming, for the first time in living memory, *less concentrated*.

For the working creative reading this in one of the historical centres: this redistribution is not theoretical. The clients, the budgets and the IP that previously concentrated in your city are, increasingly, being competed for by capable AI-augmented competitors operating from anywhere on the planet. Your geographic advantage is contracting.

What the working creative should do

I have, throughout the book, tried to land each chapter with practical takeaways. This chapter is the labour-market chapter, and the takeaways are the most concrete.

Build the literacy this year. Not next year. The Adobe Creators' Toolkit Report, the LANDR survey, the GDC data, the McKinsey reading — all point in the same direction. The literacy premium is not a future variable. It is, in 2026, the single biggest determinant of mid-career creative employment outcomes. Free training programmes exist (Sundance Collab, UK Free AI Training, Adobe Express, the open-source ecosystem). Use them. Spend the equivalent of one week per quarter, deliberately, on building practical fluency in the toolchain.

Map your craft against the Continuum (Chapter 3). Decide where your craft sits, function by function. Decide where you are willing to let agents operate on your behalf and where you are not. Write the map down. Update it quarterly. The working creatives I have watched make the most successful transitions in this period have been the ones who knew, in advance and explicitly, where their lines were.

Pick your role on the new map. Orchestrator, portfolio creative, AI-literacy trainer, model curator, ethics-and-disclosure specialist, regional creator-producer, hybrid apprentice. The roles are real. They pay. They will, by every indicator I can read, continue to grow through the next five years. You do not need to invent your own category. You can pick one of the emerging ones.

Build your apprenticeship — or build the next generation's. If you are early in your career, find the senior practitioners who are running hybrid apprenticeship pipelines and apply. If you are mid-career, identify the AI-augmented junior roles in your discipline and either fill them yourself or pair with one. If you are senior, maintain juniors in your team and pair them deliberately with the new tools.

Take the geography seriously. If you have historically been outside the creative-economy centres, the AI cost reduction has opened a window. Use it. Build for the global creative market from where you are, with the cost advantages your geography offers. If you have historically been inside the centres, your geographic premium is contracting; build defensible craft, IP and relationships that survive the cost flattening.

Stay in the work. This is the same advice the rest of the book lands on, and it is the same advice for this chapter. The maker who never makes is the maker whose judgement decays. Maintain craft contact, in at least one part of your practice, that does not depend on AI tooling. The contact will keep your eye sharp for everything else. The orchestrator who never operates the tools cannot brief them well. The orchestrator who only operates the tools loses the human signal the audience came for.

Speak. The labour-market shape of the next decade is being decided right now, by the institutions of collective bargaining, by the policy-makers running consultations, by the platform companies building the rails, and by the creative organisations making integration choices. The 88% in the UK consultation was made out of voices. The Tilly Tax was made out of voices. The Sundance literacy turn was made out of voices. Your voice — your testimony to your union, your trade body, your local government, your manager, your client — is the input the institutions need. The labour-market protections of 2030 will be the cumulative result of how loudly working creatives turn up to claim them in 2026 and 2027.

The doomer frame, and why it loses to the data

A separate but related framing has, since the early spring of 2026, occupied roughly the same rhetorical ground as the *jobs apocalypse* position but with rather more political weight behind it. Call it the **doomer frame**: the argument, articulated through OpenAI's policy paper, Dario Amodei's interviews, and Daniel Susskind's *A World Without Work*, that AI represents a categorical break from prior automation, that productivity gains will accrue only to capital, that the labour market is showing

early signs of *structural collapse* rather than ordinary sectoral rotation, and that a *new social contract* — sovereign wealth funds, robot taxes, universal basic income, a right to AI access — is the required institutional response to an inevitable surplus of unemployable humans.⁷¹¹

This is, in my view, the most consequential rhetorical position currently on the table, because it is the position from which the policy architecture of the next decade is being constructed. The OpenAI proposals I described in [Chapter 14](#), the various national wealth-fund and four-day-week initiatives now under discussion, the *social-contract-for-the-intelligence-age* framing that has migrated from think-tank papers into mainstream policy briefings — all of it rests on the four load-bearing claims the doomer frame is making about the labour data. If those claims are wrong, the policy architecture being built on top of them is wrong.

The empirical case against the doomer frame is extensive — set out at length in DreamLab’s published research — and I will not reproduce it here in full. The short version is that each of the four load-bearing claims is, at best, unsupported, and several are flatly contradicted by the same Stanford and Brookings data the doomer camp routinely cites. The Stanford Digital Economy Lab’s *Canaries in the Coal Mine?* paper — the 13% relative employment decline among 22–25-year-olds in AI-exposed occupations — is real, painful, and worth taking seriously as an *apprenticeship-pipeline* problem (the precise problem this chapter’s *Apprenticeship Gap* section is built around). What it is not is a civilizational signal. Overall employment in AI-exposed occupations has grown robustly; older workers in the same occupations have seen employment hold steady or grow; aggregate unemployment in AI-exposed roles is rising more *slowly* than overall unemployment. Brynjolfsson, Li and Raymond’s *Generative AI at Work* — the most-cited worker-level study in this period — finds the *opposite* of the doomer prediction: AI delivered a +34% productivity gain to novice / low-skill workers and a +14% average gain across 5,179 workers, while *compressing* the skill distribution rather than amplifying it. Acemoglu’s *Simple Macroeconomics of AI* — the most rigorous bear case in the academic literature — bounds the total-factor-productivity gain at 0.66% over ten years, about 0.06% per year. The 230-year track record of identical predictions, from Elizabeth I’s refusal of the knitting-frame patent in 1589 through Hinton’s “stop training radiologists now” in 2016, is unambiguous: every generation produces a high-status intellectual movement convinced that *this time* the historical pattern will break, and every generation has been wrong.

The risk of accepting the doomer frame is not that we accidentally help displaced workers. The risk is that we build the policy architecture for a crisis that does not arrive, and in doing so create exactly the labour-market rigidities that prevent the ordinary process of adjustment from working. The countries that handled the textile transition badly were the ones that resisted it. The countries that handled the

agricultural transition badly were the ones that subsidised it indefinitely. There is a real chance that the most damaging response to AI is not the technology itself but the panic-driven policy response to it.

For working creatives reading this, the implication is direct. The doomer frame is doing political and commercial work for the people promoting it. *It does not, on present evidence, describe reality.* The labour market is reshuffling, not collapsing. The mid-career squeeze is real and worth addressing. The apprenticeship pipeline is a real problem and requires deliberate institutional response. The *end-of-work event* is not happening on any timeline visible in the actual labour data, and acting as though it is — by withdrawing from the literacy investment, by treating your career as a sunk cost, by accepting the framing of inevitability that the AI labs' policy departments are pushing — is the strategic error that does the most damage to the working creative.

A note on the binary

I want to close this chapter by returning to the binary I started with — *jobs apocalypse* versus *jobs renaissance* — because both framings, repeated often enough, do real damage to the working creative who is trying to make rational career decisions in real time.

The apocalyptic framing is, in my view, the more dangerous of the two, because it produces *paralysis*. Working creatives who become convinced that AI is coming for their job, full stop, often stop investing in the literacy that would let them stay ahead of the substitution curve. They become spectators of their own displacement. By the time the substitution arrives, they are unprepared.

The renaissance framing is the less dangerous one but is also wrong, because it produces *complacency*. Working creatives who become convinced that AI will simply expand the labour market often fail to invest deliberately in the literacy that lets them claim the expansion. They drift, expecting the rising tide to lift their boat without recognising that it lifts only the boats whose owners are paying attention.

The accurate framing is harder to live with than either: the labour market is *redistributing* in real time, with sharp winners and sharp losers, and the variable that most reliably predicts which side you land on is the *deliberate, structured investment in AI literacy* that you make over the next twelve to eighteen months.

The good news, against the dystopian end of the press cycle, is that the literacy is acquirable. It is not gatekept by class, by geography, by previous credentialing, or by institutional access. The tools are largely free at the entry level. The training programmes are largely free. The community of practice — the *Dream Machine* readers, the DreamLab Collective, the open-source forums, the regional creator networks, the literacy initiatives at Sundance and elsewhere — is open and inviting.

The barrier to entry is, by historical creative-industry standards, low. The empirical case that AI is, in practice, *lowering* barriers to entry for practitioners who previously lacked institutional access — and the evidence on who benefits most from AI augmentation — runs consistently in that direction across the stigma, privilege and democratisation research DreamLab has published.

The harder news is that the literacy *has to be acquired deliberately*. It does not arrive by osmosis. It does not arrive by reading the trade press. It arrives by *doing the work* — by sitting at the desk, briefing the agents, evaluating their outputs, building the workflow, and shipping the result. Then doing it again. Then doing it again. The literacy is a *practice*, not a credential. It is, in that sense, identical in shape to every other craft skill the creative industries have ever rewarded.

The working creative in 2026 who treats AI literacy as a craft to be practised, not a technology to be debated, is the working creative who, in 2030, will look back at this period as the one in which their career took its most consequential turn.

Build the literacy. Pick your role on the new map. Defend the apprenticeship. Take the geography seriously. Stay in the work. Speak.

The labour market is moving. So can you.

Chapter 16 — Choosing the Future

I want to start this chapter with a confession.

When I sent out the first edition of *Dream Machine* on 6 October 2025, I thought I was writing a newsletter about *tools*. I thought it was going to be a weekly digest of new model releases, new app launches, new research papers — a useful reading list for people in the creative industries who wanted to keep up with what was, even then, an absurd pace of technical change.

Six months later, I do not think that is what *Dream Machine* has been about, and I do not think it is what this book is about either.

The tools have mattered. The tools will continue to matter. Each chapter of this book has had to spend significant pages on what the platforms shipped, when, and to whom — because the platforms are setting the rails, and you cannot understand the choices the creative industries are now making without understanding the constraints those rails impose.

But the *book* has been about something else. About a question that the tools force every working creative, every studio, every union, every government, every audience member, and — eventually — every person who consumes culture, to answer.

The question is: *what kind of creative economy do we want on the other side of this?*

That is what this chapter is for. To put the question on the table, in its sharpest form. To describe what a humane answer to it looks like. And to tell you, as directly as I can, what I think we should each do on Monday morning to start producing that answer rather than the alternative. The chapters that follow this one — *The Tools*, a categorised inventory of the toolchain; *Five Years Inside the Dream Machine*, the one chapter where I let the speculation off the leash; and the *Epilogue*, a letter to the creative person reading this in 2030 — are the practical reference, the long view and the closing register. The argument the book has been building toward lives here.

The choice on the table

The choice, stated plainly, is between two creative economies.

One is **extractive**. In this economy, the creative work of millions of human authors — their writing, their music, their images, their voices, their styles, their cultural specificity — is absorbed, without consent or compensation, into large statistical models owned by a small number of platform companies. The platforms then sell access to those models, mostly to the same brands, studios, and agencies that previously paid for original human work, at a fraction of the original cost. The aggregate result is a *transfer of wealth* from the diffuse pool of working creatives to a concentrated pool of platform shareholders. The creative output of the economy continues — perhaps even increases in volume — but its *meaning*, its *cultural*

specificity, and its *connection to the lives of the people who used to make it*, decays over time.

The other is **generative**, in the original sense of the word. In this economy, AI is treated as *new craft infrastructure* — a set of tools that, like the printing press, the camera, the synthesiser and the digital editing suite, can be used by working creatives to make new kinds of work that the previous infrastructure didn't allow. The training data is consented to, attributed to, and compensated for. The platforms compete on the quality and integrity of their tools, not on the unpriced absorption of their users' work. The audience can verify the provenance of what they encounter, and pay attention accordingly. Working creatives are *augmented* by the new tools, not *replaced* by them. The output of the economy is bigger, more diverse, more accessible, more globally distributed — and recognisably continuous with the human creative traditions it builds on.

These two economies are not, on the current trajectory, equally probable. Some of the rails being laid right now point at the first. Some at the second. The choice between them is not — as I argued in Chapter 13 — a technical question. The technology to support either is, in spring 2026, broadly available.

The choice is *political, institutional and cultural*. It is about who gets to make the rules, who has the standing to enforce them, and what default the millions of small daily decisions that constitute a creative economy converge on.

The good news, if you have read this far, is that more of those daily decisions than you might expect are pointing at the generative economy. The audience's behaviour around the slop ceiling. The UK consultation's 88%. The Sundance Literacy turn. The Cannes Disclosure Standard. The Academy's *you must be human to win*. The SAG-AFTRA Tilly Tax. The studio refusals — Jagex, Larian, Aardman, Games Workshop. The disclosure norms emerging in advertising and the platforms. The 800-creator declaration. The C2PA standards. The SynthID rollout. The free-AI-training-for-all programmes. The neurodiversity dividend. The Global South opening up.

These are real. They are not a foregone conclusion. But they are real, and they are coming from a coalition of forces — creators, audiences, unions, governments, institutions and some of the platform companies — that has, six months in, more political and economic weight than the extractive trajectory's proponents are willing to acknowledge.

What I want to argue, in the rest of this chapter, is that *the generative economy is winnable*. It is winnable in the next eighteen months. But it requires the working creatives, the organisations, and the institutions that have been doing the work this book has documented to *keep doing it*, with deliberate intent, against the gravitational pull of the extractive alternative.

What I actually believe

Before I get to the principles, I want to put a piece of conviction on the page that has been implicit in everything else I have written but that I think deserves to be made plain.

I believe in AI as an **assistive tool that amplifies human creativity**. Not as a replacement for it. Not as a substitute for it. As an instrument that, properly deployed, expands what a working creative can imagine, attempt and finish in a given afternoon — without displacing the human imagination, the human attempt, or the human finish.

This is the operating assumption underneath every chapter of this book, and I have been deliberately cautious about stating it as my own view, because the whole point of the newsletter and the book has been to track the *evidence*, not the *enthusiasm*. The evidence has been mixed. The cultural reaction has been mixed. The economics have been mixed. None of those mixed signals would have served the reader if I had collapsed them, every week, into the version I personally find most hopeful.

But the evidence has come in. The slop ceiling holds. The audience can tell the difference between work made with care and work made by a content farm. The toolchain rewards taste, judgement and intent more than it rewards raw computational throughput. The economic returns of agentic AI accrue, demonstrably, to the people who already know what good output looks like — not to a category of new “AI-only” creators replacing human ones, but to a re-tooled population of working creatives whose effective reach has grown.

What is happening in 2026, on the evidence I have spent six months collecting, is not that machines are taking over creative work. What is happening is that the *labour of execution* is being democratised, and the *labour of intent* is being foregrounded. The *how* is becoming a utility. The *why* is becoming the scarce good.

The long-form case for this inversion — the philosophical and economic argument that, when the technical barrier to production collapses, the value of *deciding what should be produced and why* rises in direct proportion — sits in DreamLab’s published research. The artist of 2030, in that framing, is less a *manual laborer* and more an *Architect of Meaning* — a curator, a noticer, a setter-of-intent — and the friction of human vulnerability is the irreplaceable component of the work that the machine cannot, by construction, supply.

I think that framing is essentially correct. I think the working creatives who emerge from this transition with the most leverage will be the ones who took the AI tools seriously *as instruments of their own intent*, and refused either to surrender that intent to the toolchain or to refuse the toolchain on principle. The maker-as-craftsperson and the maker-as-prompt-monkey are both poor models of where the work is heading. The maker-as-orchestrator — the architect, the editor-in-chief, the

director of a hybrid human-agent team whose unifying signal is *taste* — is the model that the next ten years rewards.

This is not a small reframing. It is the reframing the rest of this chapter — and the rest of this book — has been building towards. The four principles I am about to lay out are not principles for *restraining* AI. They are principles for *deploying AI in service of the human creativity it is supposed to amplify*. The choice between the extractive and the generative economy is, in the end, a choice about which of those two we treat as the master and which we treat as the servant.

I think the human creativity is the master. I think the AI is the servant. I think the creative economy that emerges on the other side of this transition will be the one that does not lose sight of that hierarchy.

The age of the Why

I want to give this conviction a name, because the name will travel where the long-form argument cannot, and because a name makes the choice easier to hold in the head when the next platform launch tries to talk you out of it.

We are leaving the *age of the How*. We are entering the *age of the Why*.

By the *age of the How*, I mean the long century in which the central question of working creative life was *can you do the thing?* Can you draw the figure, light the scene, edit the cut, mix the record, model the asset, hold the camera steady, hit the note? The training pipelines of every creative industry the *Dream Machine* newsletter has tracked were, at their core, infrastructure for answering that question. Conservatoires for the note. Film schools for the cut. Apprenticeships for the figure. Studios for the scene. Whole career structures built around the demonstrable, transferrable ability to *execute* — to convert intent into finished work using the technical labour of one's own hands and ears and eyes.

The *How* was the bottleneck. The *How* was the scarce thing. The *How* was what you got paid for.

The *How*, in 2026, is — at a rate that I do not think the industry has yet fully metabolised — becoming a utility. The teenager in the bedroom with a midrange GPU and a Claude subscription can, at the time I am writing this, produce work whose surface qualities — composition, lighting, sound design, edit pacing, visual-effects polish — sit on a continuum with what a full studio could produce in 2020. The continuum is not yet at the very top end; the bedroom does not yet make *Avatar*. But it is, on every available metric, closing the surface-quality gap at a rate that makes the next-decade trajectory unambiguous. *Hollywood-level execution in a bedroom* is no longer a marketing slogan. It is, for working filmmakers and musicians I know, already true for non-trivial fractions of the work.

When the *How* becomes a utility, the *Why* becomes the scarce good.

I want to anchor this in a specific story that captures the dynamic better than any creative-industries example I have. In March 2026, *Bloomberg* ran a piece on what artificial intelligence has done to elite chess.⁷¹² AI, the article reported, has driven the game towards *perfect play* at the very top — Stockfish, Leela Chess Zero and their descendants have, between them, mapped out the optimal response to most board positions a top-twenty grandmaster is likely to encounter. The visible effect of this, through 2024 and 2025, was a striking rise in the *draw rate* at top tournaments. When both players have memorised the machine-optimal lines, both players play optimally, and both players draw. The game, at the very top end, was being *solved into stasis* by the machines that had been trained on it.

The *Bloomberg* piece reported the response. Top grandmasters — Magnus Carlsen among them — had started, deliberately, to play *sub-optimal* moves. Moves that Stockfish would mark as inaccuracies. Moves that the machine-optimal line would not recommend. Moves chosen, specifically, because they were *unexpected*, and because the opponent — having trained against the machine — had not memorised the human-grade response to them. The grandmasters had stopped trying to out-machine the machine. They had started, instead, to *deliberately diverge from machine-optimal play* in ways that put their opponents on uncomfortable, uncomputed ground.

The grandmasters are winning, in 2026, by *doing the unexpected*.

I want to dwell on this image, because I think it is the cleanest available picture of what working creative life looks like on the other side of the *How* becoming a utility. When the machine has solved the optimal move — the perfectly-lit shot, the on-trend hook, the algorithmically-tested ad treatment, the mean-of-the-distribution image — the human edge is no longer in *playing the optimal move better than the machine*. The machine plays the optimal move infinitely. The human edge is in playing the move *the machine would not make*. The deliberately unexpected. The taste-driven. The risk-taking. The idiosyncratic. The personal. The locally-meaningful. The unrepeatabe.

This is the *Why* in operational form. The *Why* is not, in the practitioner's day, an abstract philosophical commitment to human creativity. It is a *daily competitive practice* of choosing the move the machine would not make. The director who picks the actor the casting algorithm would not have picked. The songwriter who keeps the verse the chart-pattern model would have cut. The illustrator who renders the figure in a style no FLUX prompt would generate. The brand creative who builds the campaign around the audience question the marketing-AI did not surface, because no marketing-AI in 2026 has the lived context to surface it.

The deeply human things — *taste, intent, authenticity, the willingness to take a risk on a move the data does not yet endorse, the refusal to make the average thing in*

service of the meaningful thing — are not, in the age of the Why, vestigial commitments held against the toolchain. They are, increasingly, the only things the toolchain cannot do, and therefore, by simple economic substitution, the things that have *commercial leverage* in a market where everything else has been pushed towards zero marginal cost.

The slop ceiling in Chapter 5 is this dynamic measured at the audience layer: the audience, presented with the machine-optimal flood, reliably underweights it, because the audience can tell — at the speed of a swipe — that there is no human *Why* underneath the work. The orchestrator role in Chapter 12 is this dynamic operationalised at the working-creative layer: the orchestrator's daily craft is, exactly, the *judgement to make the un-machine-like move at the right moment*. The authenticity premium in Chapter 13 is this dynamic priced at the commercial layer: the audience pays extra, demonstrably, for work whose human *Why* is verifiable. The four principles I am about to lay out are this dynamic codified at the policy and platform layer.

The chess grandmasters did not stop using engines. They train on engines daily. They study machine-optimal play more closely than any generation of players before them. The chess engines, far from being their enemy, are their *most-used analytical tool*. What the grandmasters refused to do is to *let the engines define what counted as a winning move*. They use the engine to know what the optimal line is, and then they choose, *for taste and surprise reasons*, to play another line.

This is also, I should note in passing, the *unflattering diagnosis* of the legacy entertainment industries' strategic position that I made in Chapter 7. Hollywood, commercial music and the AAA games business spent the past fifteen years optimising themselves *toward* the engine-optimal move — toward the franchise-instalment, the streaming-tested chart hit, the open-world template — and arrived at the AI moment producing exactly the work the engines can now replicate most cheaply. The grandmasters' response to engines is the response legacy needs to make to AI. The new AI-native studios — Gossip Goblin, Critterz, Imaginae, Asteria, Wonder — have, by virtue of their newness, no calcified rules to unlearn and are, by default, playing the moves no machine would generate. The legacy industries that survive will be the ones that re-learn how to make the un-machine-like move. The ones that don't will, on the historical pattern of Chapter 2, be remembered as the cohort that defended the previous definition while a different cohort, with no inherited risk-aversion, defined the next one.

That is the working operating model I think the next decade of creative work runs on. Use the engines. Learn the engines. Train against the engines. And then, in the actual encounter with the audience, *make the move the engine would not have made*.

The age of the Why is not the age of refusing AI. It is the age of mastering AI sufficiently that the only competitive question left is whether you can summon the *deliberately unexpected, deliberately human, deliberately yours* move that the machine, by construction, cannot.

Four principles

I have been asked, by readers of the newsletter and by people who turn up to the talks I have been giving since the autumn, what a humane AI-era creative economy should look like. After six months of trying to answer that question, I have come down to four principles. They are not a programme. They are a *test* you can apply to any policy, any contract, any product launch, any organisational decision — and ask whether it is moving us towards the generative economy or the extractive one.

One. Agency. Every working creative should retain meaningful control over how AI is used in relation to their own work — both in *what gets trained on their work* and in *whether and how they choose to use AI in their own practice*. The Human–AI Agency Continuum from Chapter 3 is the practical expression of this. Policy, contracts, platform terms-of-service and union agreements should all be evaluated against whether they preserve this control or undermine it.

The 88% of UK respondents who said training should require licensing in all cases were articulating this principle.

Two. Attribution. When AI systems produce work that is derived, in any meaningful sense, from human-authored training data, the human authors should be identified and — where appropriate — compensated. The technical infrastructure for this — *creative weight attribution* in Musical AI's framing,⁷¹³ C2PA provenance metadata, SynthID watermarks — is the most important infrastructure question of the next two years. Policy should support its deployment. Contracts should require it. Audiences should expect it. Platforms that refuse to engage with it should be treated, in the policy and procurement environment, as the laggards they are.

Three. Access. The benefits of AI tooling — the productivity, the creative leverage, the cost reduction — should be broadly available, not concentrated. This means free or near-free tools for creative entry; investment in literacy at the population scale; deliberate inclusion of historically excluded creative communities in the design, training and deployment of AI systems; and structural support for the indie, the regional, the Global South, the neurodivergent and the under-resourced creative ecosystems that the AI cost reduction has the potential to *include* in the global creative economy for the first time.

The Korin AI launch, the rise of Indian AI cinema, the African-trained model investments — these are early signs that the access principle can be operationalised.

Four. Audience. The audience for creative work is not a passive market. The audience is — has always been, and increasingly is — a *participant* in the cultural meaning of the work. AI policy, platform design and creative practice should treat the audience as such: with the right to know what they are encountering, the right to choose what they pay attention to, the right to refuse work that violates their cultural or ethical preferences, and the right to a creative economy that produces *for* them rather than *at* them.

The slop ceiling is the audience exercising this principle. The cultural rejection of cynical AI work. The death threats and the BBC investigations and the *Tiny Grandma* virality. The audience has been telling us what it wants. The institutions and the platforms are, slowly, beginning to listen.

If you can hold these four principles — *agency, attribution, access, audience* — in your head when you sit down to make a creative decision in 2026, you have a working test for whether you are building towards the generative economy or away from it. Apply the test. Often. Without too much agonising. The aggregate of millions of decisions made against this test, over the next eighteen months, is the choice we are collectively making about what comes next.

The encouraging thing is that the institutions are starting to build the machinery for exactly this kind of test. In July 2026, Arts Council England launched a **Responsible AI Cultural Leaders Academy** — a programme designed to train the people who run the country’s cultural organisations in how to adopt AI *responsibly*: with the governance, the ethics and the sector-specific judgement the adoption decision actually requires.⁷¹⁴ It is a small programme against the scale of the transition. But it is precisely the shape of institutional response the four principles call for — not a ban, not a shrug, but a deliberate investment in the *literacy of the people who will be making the rules* inside the cultural sector. When the national arts funder decides that responsible AI adoption is a leadership competency worth training for, the generative economy has gained a piece of infrastructure.

The four principles are policy and platform tests. They describe what the rules of the new economy should be. They do not, on their own, describe what working creatives should be *doing* with their hands, every day, to put themselves on the inside rather than the outside of the rule-writing. I made the case for the practitioner’s version of this argument at length in Chapter 3’s *Open the black box* section, and I want to put one summary sentence of it here, because the four principles cannot be defended by creators who do not understand the toolchain underneath them.

Working creatives need to open the black box of AI and own a real stake in how it is built. Not just use it. Not just refuse it. Not just bargain over its terms. *Open it.* Understand what the model was trained on. Read the EULAs of the

platforms you ship work through. Run some part of your stack on open-weight infrastructure (Chapter 17 is the practical map). Show up to the governance conversations — the Cannes Disclosure Standard rooms, the UK consultation responses, the SAG-AFTRA bargaining tables, the C2PA standards body. The cohort of working creatives that does this defines the next era's craft. The cohort that uses the toolchain without ever asking what is inside it has the era's craft defined for them, by the platform companies that built the toolchain. The four principles assume — they cannot deliver on their own — that the first cohort is large enough, organised enough and technically literate enough to do the work. The argument of this chapter, on its widest reading, is that turning up to the rule-writing is the practitioner's work.

The cost we are not pricing

I want to break my own framing for a moment, because there is a category of cost the four principles only obliquely address, and any reader who has been paying attention to the wider technology press in the period this book covers will be wondering when I am going to put it on the page.

The cost is *the resource footprint of the systems we are building everything on top of*.

Training and running large generative models, at the scale the creative industries are now using them, is an industrial activity. It is consuming enormous quantities of electricity, water, semiconductor manufacturing capacity, and rare-earth-mineral throughput, in a global infrastructure-build the planet has not seen since the rollout of mobile telephony. The trade-press coverage on this is patchy and the data the platform companies disclose is partial, but the direction is unambiguous. The data centres that produce the *Sora 2* clip you generated for a client this morning, the *Veo 3.1* sequence you cut into your edit, the *Suno* track you used as scratch on a project — those data centres draw power on a scale that is, in aggregate, a meaningful new line item in the global energy mix.

There is also a second category of cost, equally unpriced, that the four principles touch only by implication. The training, refinement and moderation of these systems involves an extensive workforce of human data labellers, content moderators and reinforcement-learning evaluators, much of it concentrated in low-wage labour markets in the Global South, much of it carrying significant psychological cost from exposure to the worst of what the models filter out. The creative economy of 2026 is, in part, sitting on top of that labour. The labour is not always visible in the marketing copy of the platform companies that depend on it.

I have under-treated both of these costs in the chapters above, on the grounds that the book is specifically about the creative-industries transition rather than about AI's wider externalities. That is a defensible editorial choice. It is also an incomplete

one, and a fair reader is right to ask why a manifesto for a humane creative economy stops at the edge of the studio.

The honest answer is this. The four principles, as I have stated them, are *necessary but not sufficient*. A creative economy that gets agency, attribution, access and audience right, but that does so on top of a platform stack whose energy and labour externalities are concentrated, opaque, and morally indefensible, is not a generative economy. It is an extractive economy with better creative-side ethics — which is, in some ways, the more dangerous animal, because the creative-side ethics provide ideological cover for the wider extraction.

The implication for the working creative reading this is *not* that you should refuse to use AI tools because of their environmental and labour costs. The implication is that you should *demand the same transparency* from the platforms about their externalities that you demand about their training data. Energy disclosure, water disclosure, labour conditions in the data-supply chain, carbon accounting of inference runs — these should be standard procurement requirements for any creative organisation buying platform access at scale, in exactly the same way C2PA provenance is becoming a standard requirement for the work itself.

If I were rewriting the four principles from scratch, I would probably make this a fifth: **Footprint**. The full resource cost of the work — energy, water, human labour, supply-chain externalities — should be visible to the people commissioning, making and consuming it. The principle would sit awkwardly alongside the other four, because it is the only one that does not run along the human-creative-vs-platform axis the book has otherwise been organised on. I have left the four where they are, because the rhetorical compression of *Agency, Attribution, Access, Audience* is one of the more useful things I have written. But the fifth is real. It belongs in the conversation. I want, in this section, to have at least said so on the page.

The deeper point, which I will not dwell on further because it is the subject of a different book by a different writer who I hope is writing it now, is that the AI transition is — like every previous industrial transition — being paid for somewhere. The creative-industries part of the bill is, in 2026, becoming visible. The energy and labour parts of the bill are still, by deliberate platform-company design, harder to see. A humane creative economy will have to insist on seeing both.

The DreamLab model

I want to talk briefly about the **DreamLab** model — not because I think the studio I run in the North West of England is some kind of ideal, but because it is one specific, concrete instance of how a working creative organisation has tried, deliberately, to build for the generative economy rather than against it, and because I think the choices we have made are useful to share.

DreamLab is a collective of about fifty practitioners — artists, technologists, directors, games developers, storytellers — based in the North West, with collaborators across the U.K. and internationally.⁷¹⁵ We are emphatically *not* an AI company. We are a *creative* company that happens to use AI heavily.

The choices we have made, since we started thinking deliberately about how to operate in this environment in early 2024, are not heroic. They are pragmatic. They include:

A disclosure-first practice. We tell clients, in the brief, what AI tools we propose to use and what for. We document our use. We can — and have, when asked — provide chain-of-custody on contested work.

A continuum-first contract. Our client contracts identify, function by function, where on the Human–AI Agency Continuum the work will sit. Some functions are 100% human (performance, lead creative direction, story). Some are mixed. Some are mostly AI (asset variations, plate generation, certain post-production tasks). The client knows. The team knows. The line is negotiable but it is drawn.

A literacy-first hiring approach. We hire — and we invest in upskilling — generalists rather than specialists. We expect every working member of the collective to understand the AI tools in *adjacent* disciplines well enough to brief them and judge their outputs. The portfolio creative model from Chapter 12 is the operating practice.

An open-toolchain commitment. We use a deliberately diverse toolchain — Adobe, Runway, ComfyUI, World Labs, Hunyuan, open-source models on Hugging Face, internally built tooling. We refuse exclusive dependency on any single platform. If one of the tools shifts in a direction that violates the four principles, we have alternatives.

A community-first relationship to the city. The collective is intentionally based in the North West of England, not in London or Los Angeles. The aim, explicitly, is to build creative infrastructure in places that the previous generation of the creative economy did not invest in. We hire locally. We train locally. We work, where we can, with regional partners.

A newsletter, an archive, a public record. *Dream Machine*, the publication this book is built out of, exists in significant part because *putting the work in public* — keeping a transparent, citable record of what is happening, what is being said, what is being decided — is itself a form of building the generative economy. Hidden knowledge concentrates power. Shared knowledge distributes it.

None of these choices are unique to DreamLab. Many of them are being made, in different forms, by studios, agencies, labels and indie production companies all over the world. What I would say, having lived inside this set of choices for the period this

book covers, is that they *work*. The studio's output has been better, our team has been happier, our clients have been more loyal, and our market position has been more defensible than I had any right to expect when I started writing the newsletter in October.

The generative economy, as a practice, is not abstract. It is a set of operational choices that working organisations can make on a Monday morning. The choice produces real, measurable outcomes. It is not, as some of the more pessimistic AI commentary frames it, a luxury for organisations that can afford to be high-minded. It is, in our experience, the most *commercially* sustainable position available right now.

What working creatives should do on Monday morning

This book is mostly aimed at working creatives. The studios and the institutions and the platforms have been the subject of the chapters, but the audience for the chapters has been you — the writer, the director, the songwriter, the games designer, the photographer, the editor, the producer, the agency creative, the indie filmmaker, the YouTuber, the freelance illustrator, the student looking at a creative career and trying to figure out whether to be discouraged or excited.

What should you do on Monday morning?

I will keep this brief, because the rest of the book has been the long version.

Read your own newsletter. Whatever your version of *Dream Machine* is — the publication you trust to tell you what is happening in your discipline week by week — read it carefully. The pace of change in 2026 is too fast to absorb by osmosis. You need a deliberate reading practice. If you can't find one you trust, build one.

Draw your lines. Write down, for your own practice, where on the Human–AI Agency Continuum you want to sit, function by function. Update the lines as the technology and your judgement evolve. Be prepared to *show* the lines — to clients, to collaborators, to yourself.

Practise briefing. It is the most leveraged skill of the era. Brief the agents. Brief your collaborators. Brief yourself. Get clearer, every week, about what you actually want from a piece of work *before* you start making it.

Cultivate taste, on purpose. Look at more good work. Look at it harder. Articulate, every week, what makes something good — and what makes the average thing average. The agents will, by default, give you average. Your job is to know better.

Disclose. Tell people what you are using, how you are using it, and where the human work in your output lives. The transparency is, in 2026, not a cost. It is a competitive advantage.

Stay in the work. Resist the temptation to abstract too far. The maker who never makes is the maker whose judgement decays. Maintain craft contact, in at least one part of your practice, that does not depend on AI tooling. The contact will keep your eye sharp for everything else.

Find your coalition. The 88% in the U.K. consultation didn't get there by accident. It got there because creators turned up alongside other creators alongside professional bodies alongside unions alongside institutions. The political and economic shape of the next decade of creative work depends on creators being *organised*. Join your union. Sign the declarations. Show up to the consultations. The institutions of collective bargaining and political representation that your forerunners built are the institutions that will defend your work in 2030.

Build for the new geography. Wherever you are in the world, the AI cost reduction has made it more viable than ever to build a serious creative practice in places that were previously locked out of the centre. Take the opportunity. The next century of cultural production is not going to belong to the cities that owned the previous one. The next generation of canonical creative work will, on the trajectory I see, come from more places, in more languages, in more forms, than ever before.

Make the work. None of the rest of this matters if you do not make the work. The AI era is not a reason to stop. It is a different set of conditions under which to keep going. Working creatives have always operated against the technological grain of their day. The maker in 2026 — like the maker in 1926 and 1826 and 1626 — is the person who finds, in the conditions of their actual moment, the work that needs to be made and then makes it.

The conditions have changed. The need has not.

Predictions you can check me on

Books on the cusp of a transition tend to age in one of two ways. Either they were broadly right about the shape of what was coming and look prescient five years on, or they were embarrassingly wrong and quietly disappear from the recommended-reading lists. The honest thing to do, at the end of a book like this one, is to write down predictions specific enough that the future can grade them.

Here are mine, dated May 2026.

By the end of 2027. Every major studio, label and agency has, in its standard production contract, a clause requiring AI-use disclosure across the production pipeline. *Position Four* — *AI in the workflow, not in the work* — is the dominant operational posture in legacy creative organisations. The number of *Position Three* refusing studios has stabilised at roughly 10–15% of the mid-and-upper market, concentrated in IP-led franchises where the audience contract depends on a human-craft signal.

By the end of 2027. The 44%/3% Deezer ratio (Chapter 5) has stabilised on the upload side at 50–65% AI by volume, with consumer streams of fully-AI tracks still under 5% of total play time. *Audience-disclosed* synthetic music is a small but real category — under 10% of paid streams — and is concentrated in mood/background and synth-driven genres rather than in artist-led popular music.

By the end of 2028. A *Position Two* studio — Imaginae, Wonder, Obsidian, Asteria or a name we have not heard yet — produces the first AI-native feature that wins material awards recognition *for its writing or direction* (not for its technology). The audience response is mixed; the cultural permission for AI-native cinema has shifted from contested to accepted-with-asterisks. James Cameron’s “horrifying” line is still being quoted but the position it describes has narrowed to *AI that displaces a specific actor in a specific scene*, not the general toolchain.

By the end of 2028. At least one large national government — most likely the U.K. or one of a few EU member states — has passed legislation requiring licensing of copyrighted work for AI training. The U.S. has not, but the litigation environment (UMG v Anthropic and the cases that follow it) has produced de facto licensing markets that platform companies have grudgingly entered. The technical infrastructure for *creative weight attribution* is mature enough to underpin real revenue distribution to working creators.

By the end of 2029. World models, not flat video, are the dominant medium for new high-production-value creative work. Marble-class tools are commodity infrastructure; the differentiated work is being done by orchestrators with strong *world curation* skills (Chapter 9). The boundary between film, games and immersive media is functionally gone for new productions, while legacy formats retain prestige and a real audience.

By the end of 2029. The audience contract has been substantively re-written. C2PA-class provenance metadata is supported in the major capture, edit and distribution tools by default. Platforms that do not honour provenance have lost meaningful share to ones that do. The *Tiny Grandma* error mode — human work mis-flagged as synthetic — is a solved problem in the principal platforms; the inverse error — synthetic work mis-flagged as human — is the harder one and remains the central audience-trust question.

By the end of 2030. The orchestrator role (Chapter 12) is a named seniority track in every creative discipline that survives the transition. *Senior orchestrator* is a credentialled title. The apprenticeship problem (Chapter 12, Chapter 14) is partially solved by a combination of three things: AI-augmented junior roles maintained deliberately by *Position Four* studios; new pathways through literacy initiatives like Sundance Collab and its successors; and an explicit re-funding of cultural-institution training programmes by national governments.

By the end of 2030, but I'm less confident about this one. A major creative-industry union has negotiated a *productivity dividend* — a structural share of AI-driven productivity gains that flows to the human workforce, not just to the platform companies and the studio shareholders. The mechanism is novel and contested but the underlying maths is unanswerable; the industry that does this first becomes the template for the others.

I am wrong about something on this list. I do not know yet which item. If you are reading this in 2030 and one of these predictions has aged badly, send the receipts to the *Dream Machine* newsletter address. I will publish them.

The predictions I have *not* made — about which AI company will dominate by 2030, about which specific platform companies will go bust, about whether AGI arrives in the period and what it does to all of this — are absences I have made on purpose. Anyone confident about those questions in 2026 is selling something. The shape of the *creative economy* is more predictable than the shape of *the platform layer underneath it*. That is the bet this book has been built on.

A note to those who are afraid

I want to write directly, for a paragraph or two, to the readers I know are sitting with this book — and with the larger transition this book describes — in a state of real, well-founded fear.

If you are reading this and your livelihood depends on a creative discipline that the AI tools are getting unsettlingly good at — if you are a junior animator, a session musician, a copywriter, an illustrator, a translator, a stock-image photographer, a voiceover artist, a foley artist, a concept designer, a stage actor in a regional theatre, a working musician on tour, an indie filmmaker, an entry-level games artist, a video editor at a small post-production house — and the news cycle has, over and over again, told you that your job is going away, I want you to know two things.

The first is that the news cycle is, on this question, almost always too negative. The aggregate data on creative-industry employment in 2026 does not show the collapse the headlines have been predicting. It shows a *restructuring*, with a lot of disruption, and with real pain in specific places, but with — net — more creative work being done by more people in more forms than was being done a year ago. The Sundance literacy initiative is not happening because the industry is dying. It is happening because the industry, freshly, has more entrants than ever and needs to train them.

The second is that the most important strategic move you can make in this moment is not to be afraid in private. It is to *speak*. To your union. To your trade association. To your local government's consultation. To your manager. To your audience. To the people on your team. The 88% was made out of voices. The Tilly Tax was made out

of voices. The Bandcamp ban was made out of voices. The Sundance literacy turn was made out of voices. The institutions of the new creative economy — the ones that will protect your work in 2030 — are being shaped *right now* by the people who turn up to be counted. The people who don't turn up, or who turn up only in private, are the people whose interests get absorbed into the policies of the people who do.

You are not powerless in this. You are, if anything, the most powerful constituency in this whole story, because you are the people who are actually making the work that the audience is paying attention to. You are the constituency the unions represent, the institutions exist to serve, the audience cares about, and the platforms — ultimately, despite themselves — depend on. The choice between the extractive and the generative economy is, in large part, the choice you collectively make about how to use that power.

Use it.

The Dream Machine

I want to close the argument of this book with the phrase the newsletter started with.

Welcome to the Dream Machine.

When I named the newsletter, in late September 2025, I had no clear idea what the phrase meant. I had a half-formed sense that “dream” caught something about the strange, vivid, slightly hallucinatory quality of what the new tools produced, and that “machine” caught something about the industrial, scaled, infrastructural nature of the platforms behind them. The phrase was, more than anything else, a placeholder for a feeling I couldn't quite articulate yet.

Six months on, I think I know what the phrase means.

A *dream machine* is, in the sense I am using it, an apparatus capable of producing something that has — until very recently — required a human mind to produce. Not just images, or videos, or songs, but *constructions* of meaning, of place, of presence, of emotion. The apparatus is not, in itself, doing the dreaming. It is amplifying, multiplying and distributing the dreaming of the humans who direct it.

The question this book has been about, in every chapter, is *whose dreams the machine amplifies*.

If the machine amplifies the dreams of a small number of platform shareholders, optimised against extraction, the creative economy it produces will be — to use Kehlani's line about Xania Monet — a place where “the person is doing none of the work.” The dreams of the audience are *for sale*. The dreams of the creators are *training data*. The dreams of the culture are *outputs of an algorithm tuned for engagement*.

If, on the other hand, the machine amplifies the dreams of the working creatives — supported by the audience, defended by the unions, regulated by the institutions, contained by the four principles, made transparent by the provenance infrastructure — the creative economy it produces will be the largest, most diverse, most accessible expansion of human cultural production in the history of the species.

We are, in May 2026, sitting in the moment between these two outcomes.

The signals from the last six months — the 88%, the slop ceiling, the Sundance turn, the Cannes Disclosure Standard, the Academy rule, the SAG-AFTRA contract, the audience behaviour, the creator coalition, the regional opening, the literacy investment — all point at the second outcome being achievable. Not inevitable. *Achievable.*

It is going to require the working creatives to keep turning up. It is going to require the studios, agencies and labels to keep making the integration choices we covered in Chapter 14. It is going to require the institutions — the unions, the rights bodies, the festivals, the universities — to keep doing the slow, unglamorous work of building the rails. It is going to require the platforms to be pushed, by their users and by their regulators, towards the side of provenance and attribution and consent. It is going to require the governments to keep the policy moving in the direction the 88% asked for.

It is going to require the audience to keep paying attention to the human signal.

And it is going to require people like you — the people who have read this far, who have been doing the work of figuring out what creative life looks like in this moment — to make, in your own practice, the daily decisions that build the generative economy rather than the extractive one.

This is the choice. This is what is on the table. This is the work.

Welcome to the Dream Machine.

Chapter 17 — The Tools

I have, until this chapter, deliberately kept the *tools* out of the foreground. Sixteen chapters about creative AI without a chapter on the toolchain is, on the face of it, a strange editorial decision, and I want to begin by explaining it.

The reason is that I think the most common mistake people make about this period is to confuse *the tools* with *the transition*. Tools are the visible surface of the change — the thing the press cycle covers, the thing the platform companies want you to talk about, the thing that has a price and a logo and a launch date you can put on a slide. The transition is everything underneath: the economics, the labour, the audience contract, the law, the institutions, the rails. The tools change weekly. The transition is slower, deeper, and is what will still be true in 2030 when most of the tools in this chapter are obsolete.

The first sixteen chapters were about the transition. This chapter is about the tools.

I have written it last on purpose. Read in this order, the tools sit inside the architecture the book has been building — the Continuum, the Slop Ceiling, the four positions, the orchestrator role, the four principles. Read in any other order, they collapse back into the format the platform companies prefer: a tools-arms-race in which the only question is which model is best this week.

That format is, in 2026, the most reliable way to misunderstand what is happening.

A note on the date stamp. Everything in this chapter is current to August 2026 — the May 2026 cut of the catalogue, with the additions from Issues 31 to 41 (June–August 2026) to the video, music, world-model, agent and games-engine layers folded in where they fit. By the time you read it, some of these tools will have been bought, renamed, killed, surpassed or repositioned. The point is not that the specific tools matter. The point is the *shape* of the toolchain — what categories exist, what they do, who builds them, and how a working creative builds a coherent stack on top of them. The shape, in my experience, holds.

How to think about the toolchain

Before the inventory, the frame.

I think the creative-AI toolchain in 2026 is best understood as a stack of seven layers, each with its own dominant players, its own pace of change, its own integration model. The layers, from foundation to consumer, are:

1. **Foundation models** — the large multimodal systems underneath everything else (OpenAI's GPT-class, Anthropic's Claude, Google's Gemini, Meta's Llama, the major Chinese open-source models). These are the raw capability layer. Almost no working creative uses them directly except via wrappers.
2. **Modality models** — specialist models for video (Sora, Veo, Runway Gen-4.5, Kling, Hunyuan, Wan), image (Firefly, Midjourney, FLUX, Imagen,

SDXL/Stable Diffusion variants), audio (Suno, Udio, ElevenLabs, Mureka), 3D and world (Marble, Genie 3, WorldGen, UNI-1, Hunyuan 3D-PolyGen, ECHO). These are what most working creatives think of when they say “AI tools.”

3. **Agent platforms** — systems that compose modality models and external tools into multi-step workflows (OpenAI’s AgentKit, Anthropic’s Claude apps and skills, Google’s Project Genie, Heygen’s Video Agent, Sony’s 49-agent / 72-skill stack). The agent layer is where the “orchestrator economy” of Chapter 12 actually runs.
4. **Creative software with AI baked in** — the legacy creative suites that have been rebuilt as AI-first platforms (Adobe Creative Cloud, Autodesk, Foundry, Unreal Engine, Unity, DaVinci Resolve, Pro Tools, Logic Pro, Ableton). This is where most paid professional work still happens.
5. **AI-native creative apps** — new entrants whose product is a single-purpose AI workflow (Runway, Higgsfield, Krea, Freepik, Magnific, Synthesia, Heygen, Hedra, Cascadeur, Pika, Luma). Most working creatives use 4 to 10 of these in rotation.
6. **Open-source and workflow infrastructure** — the technical-creator layer that wires everything together (ComfyUI, Hugging Face, SuperSplat, OpenEnv, the open-source model ecosystem). This is where the most interesting innovation often happens first.
7. **Consumer surfaces** — the apps that put generative capability on every phone (the Sora app, CapCut/Dreamina with Seedance, the Gemini app, the various TikTok-style remix platforms). This is the layer the audience touches.

The mistake I see most often, both in the press cycle and in studios planning their internal AI roadmaps, is to optimise for layer 2 (modality models) without understanding that the actual leverage is in how you compose layers 2, 3, 4 and 6 into a coherent workflow. The tool that “wins” is rarely the tool with the best benchmark. It is the tool that integrates cleanly into the rest of your stack.

With that frame, the inventory.

Video

The video layer changed faster than any other modality between October 2025 and July 2026, and is the one most likely to look different again by the time you read this. Treat the names as snapshots, not as a stable league table.

Sora 2 (OpenAI) is the model that opened the period this book covers. Its September 2025 launch — physical realism, audio integration, multi-shot world-state persistence — is the moment Chapter 1 is about.⁷¹⁶ The iOS app launched alongside it hit a million downloads in five days.⁷¹⁷ By March 2026, however, the

platform was shutting down. OpenAI announced a staged wind-down on **24 March 2026** — consumer app and website ending late April, developer API closing late September — with peak revenue reportedly at around \$540k per month against operating costs in the billions, and compute reallocated toward higher-margin coding and enterprise products.⁷¹⁸ The Disney–OpenAI \$1bn licensing announcement of 11 December 2025 was never executed; no formal agreement was signed and no money changed hands. The contractual *framework* underneath that deal — IP-not-training, talent-likeness excluded, jointly-governed steering committee, equity-and-warrants structure — survived the platform’s collapse and is now the reference architecture other studios negotiate against; the deeper treatment of that architecture and the broader IP-economics shift it sits inside sits in DreamLab’s published research. For working filmmakers in mid-2026, the practical implication of the Sora collapse is that the *consumer-facing edge* of the AI-video market has rotated to ByteDance’s CapCut/Dreamina/Seedance stack, to the YouTube Shorts integrations of Veo 3.1, and to the Sora app’s various TikTok-and-Reels-shaped successors. The professional pipeline never depended on Sora 2 in the first place — most working filmmakers I know used it less than its cultural prominence suggested — and so on the production side the collapse changes less than the headlines implied.

Veo 3.1 (Google DeepMind), released in mid-October 2025, is the model the professional filmmaking community has, on average, gravitated toward — for narrative coherence, controllable camera composition, cinematic lighting vocabulary and sound integration.⁷¹⁹ Sora 2 wins on raw physics in single clips; Veo 3.1 wins on the kind of sustained directorial control most actual production pipelines need.

Runway Gen-4.5 (and Gen-4.5 Image-to-Video, the Workflows product, Story Panels, Characters API, Apps for Advertising) is the most-integrated commercial AI-video stack of the period.⁷²⁰ Runway has shipped product faster than any other AI video company in this market, and CEO Cristóbal Valenzuela’s “fifty indie films instead of one \$100M blockbuster” framing is the cleanest articulation I have seen of the case for AI as creative leverage rather than cost-cut.⁷²¹

Kling (Kuaishou), **Hunyuan Video** (Tencent), **Wan 2.5** (Alibaba), **Seedance 2.0** (ByteDance) — the Chinese-built models that, in aggregate, have rivalled or surpassed the U.S. labs on specific capabilities (motion physics, character consistency, render speed) at significantly lower cost.⁷²² Hunyuan’s open-source releases have been the single most important contribution to the wider open-source AI video ecosystem in this period.

Pika 2.0, **Higgsfield**, **Luma** (Dream Machine and UNI-1) round out the commercial layer. Each has carved a niche: Pika on iteration speed and creator workflow; Higgsfield on social-media marketing video at scale (\$80M raised, \$1.3B

valuation, \$200M revenue in nine months⁷²³); Luma on the world-model bridge to spatial content.

Heygen ships Video Agent — a full scripting-to-assembly agent built around reference images.⁷²⁴ **Synthesia** holds the corporate AI-avatar market (\$4B valuation, having reportedly rejected a \$3B Adobe acquisition offer).⁷²⁵ **ElevenLabs** runs the dominant audio-AI layer underneath much of the new video work (\$500m ARR by April 2026).⁷²⁶

Gemini Omni (Google DeepMind), announced at Google I/O 2026, brings text, image, audio, video and live interaction into a single multimodal model — the first foundation-model release in this category that meaningfully unifies the modalities working creatives currently have to bridge across five different tools.⁷²⁷ **Beeple Canvas**, Mike Winkelmann’s gen-AI compositor — launched May 2026 — is the first AI-native compositing application to ship from a working visual-effects artist’s own studio, and is structurally distinct from the AI-features-bolted-onto-existing-compositors pattern in the legacy-software section below.⁷²⁸

The June–July 2026 cycle belonged to ByteDance and Runway. **Seedance 2.5** broke the thirty-second barrier for AI video generation, with **Seedance 2.0 Mini** shipping the quality-at-lower-cost tier into Dreamina and ComfyUI in the same month.⁷²⁹ Runway, meanwhile, shipped **API Recipes** (production-ready generative-media features in a single API call), an MCP connection that puts Runway *inside ChatGPT*, and **Agent 2** ad localisation — the fastest single-vendor shipping cadence of the early summer.⁷³⁰

If I had to name a single video product that, in my experience, working creatives have settled on as a default in 2026, it would be Veo 3.1 for finished work and Runway for iteration and integration. Sora is the brand name the audience knows. The actual production pipelines run on the other two.

Image

The image layer is more stable than video — the technology has matured, the differences between top models are narrower, and the dominant question has moved from “which model” to “which workflow.”

Adobe Firefly (Image Model 5, plus Foundry for custom-trained corporate models, plus integration across Photoshop / Illustrator / Express / InDesign) is the default for any working creative who is also a Creative Cloud subscriber — which is, by Adobe’s own numbers, 45% of Creative Cloud users actively using Firefly, 70% of those weekly, more than 22 billion assets generated by April 2025.⁷³¹ The Firefly adoption curve is the single best evidence I have for the consumption-gap argument in Chapter 14.

Midjourney remains the aesthetic-leadership product in the category. Slower to ship, more opinionated about output style, dominant on Discord and X among the working AI-art community.

FLUX (Black Forest Labs) is the open-source and pro-creator favourite for fine control, having largely replaced Stable Diffusion XL as the open-weight default through 2025.

Google Imagen (and the *Nano Banana* fast-image variant integrated into Gemini, Photoshop and Unreal Engine via the various plugins) has become the most-integrated image model in the consumer toolchain, by virtue of Google's distribution. Nano Banana inside Photoshop and Nano Banana inside Unreal Engine were two of the more consequential cross-platform integrations of the period.⁷³²

Krea, Freepik, Magnific, Recraft — the higher-control consumer / pro-creator products built on top of foundation image models. Each is competing on specific workflow advantages (real-time generation, upscaling, vector output, brand-consistency control).

The image workflow most commonly cited in my circle in mid-2026 is: a base generation in Firefly / Midjourney / FLUX, character-consistent variation in a controllable wrapper like Krea or Magnific, finishing inside Photoshop with the AI-assisted masking, generative-fill and object-removal tools that Adobe shipped through the autumn 2025 and spring 2026 update cycle.

Music and audio

The music layer split into three categories during this period, and the split is, in my experience, more important than the specific products in each category.

Generative music tools that produce finished tracks from prompts — **Suno** (Studio launched late 2025⁷³³), **Udio, Mureka** (with its Music Agent Studio, six specialised AI agents for songwriting, arrangement and production⁷³⁴). These are the tools that produce most of the AI-music flood Chapter 5 describes. They are also, paradoxically, the tools most working musicians use the least directly — the consumer market for AI-generated finished tracks is large and growing, but professional musicians overwhelmingly use AI tools at a different layer.

Production and post-production AI — the tools that handle audio restoration, mixing, mastering and isolation. The 1,100-creator music survey found that 58% of professional producers used AI for audio restoration, 38% for mixing assistance, 33.9% for automated mastering. **iZotope Ozone 12, LANDR**, the Pro Tools and Logic Pro AI suites, **CleanvoiceAI** for podcast post — this is where the silent-adoption majority of the music industry lives.

Voice and audio synthesis — **ElevenLabs** is the dominant player, with \$500m ARR, BlackRock / NVIDIA backing, and meaningful share across audiobook narration, dubbing, podcast synthesis and AI character voice work.⁷³⁵ The Cardiff band that found their music had been used to train an “AI artist” outperforming them on Spotify⁷³⁶ is one of the cautionary tales of this layer; the Andrii Daniels bomb-shelter clip⁷³⁷ is one of the success cases.

Sound-effect foundation models emerged as a new sub-category in May 2026. **Sony AI’s Woosh** is the first foundation model explicitly trained for the professional sound-effects discipline — built for the people who design the sonic worlds behind games, film and interactive media, not for the consumer market.⁷³⁸

Mirelo SFX 1.6 shipped the first sound-effects model that lets you *edit* a generated sound rather than only regenerate it — a structural shift in the discipline equivalent to the move from rendered images to layered Photoshop files.⁷³⁹ **Stable Audio 3.0** (Stability AI) shipped as an open-weight audio model family explicitly aimed at artistic experimentation.⁷⁴⁰ **Tamber**, the ethically-trained AI music suite I describe in Chapter 6, shipped alongside a gestural-control interface that lets the musician steer the generation with arm movements.⁷⁴¹ **Beatport’s Track ID** rolled out as the real-time identification standard for the DJ market.⁷⁴²

The *instrument turn* arrived in June 2026, and it is the music-layer development I would watch most closely. **Google Magenta RealTime 2** is a live music model you play as an instrument; **Daydream** bills itself as the first AI-native instrument for the DAW; Fender’s **Studio Pro** DAW shipped with Moises stem separation and a smart studio assistant built in; and Maono’s **PD500W** claimed the title of first AI-powered hybrid microphone.⁷⁴³ Generative music capability arriving as *performance hardware and studio furniture*, rather than as another prompt box, is a structurally different adoption path from the Suno-shaped flood — it lands in the production layer, where the silent-adoption majority already lives. **Suno** itself, in the same period, launched **Spark**, an incubator programme for the next generation of creators — the flood-tier vendor of Chapter 5 repositioning itself as a patron.⁷⁴⁴

The deal flow underneath this layer is the second-fastest-changing in the toolchain. The Stability AI / Universal Music alliance, the Stability AI / Warner Music deal, the Splice / Universal partnership, the GEMA / OpenAI lawsuit, the Wixen / Meta lawsuit, the UMG / Anthropic \$3B suit — these are the structural moves I would track if I were a working musician trying to plan a five-year toolchain.⁷⁴⁵

3D, world models, spatial

The category that, more than any other, I think defines the next decade of creative work. Chapter 9 is the long-form argument; this section is the inventory.

Marble (World Labs, Fei-Fei Li's company) is the first commercial product I would put on a professional toolchain.⁷⁴⁶ Public release November 2025; Sony Pictures' use of it in virtual production reportedly running 40× faster than the legacy workflow.⁷⁴⁷ DreamLab has been in the beta since October 2025, and Marble is, today, the world-model product most integrated into a working pipeline I have used.

Google DeepMind Genie 3 is the most ambitious research-grade world model, named by *Time* as one of the best inventions of 2025. Made publicly available to Google AI Ultra subscribers via Project Genie in January 2026.⁷⁴⁸

Meta WorldGen, Tencent HY World 1.5 (open-sourced December 2025, alongside the Hunyuan 3D Studio integration⁷⁴⁹), **SpAItial ECHO**, **Stanford Wonderzoom**, **OpenArt Worlds**, **Luma UNI-1** (the most important *category* announcement of spring 2026, combining world generation with reasoning⁷⁵⁰) — the rest of the world-model commercial layer.

The May 2026 world-model wave extended this layer further. **NVIDIA SANA-WM** is the first open-weight world model at meaningful scale (2.6B parameters), with 60-second video generation and explicit camera control.⁷⁵¹ **Odyssey Starchild-1** is, by Odyssey's own framing, "*the first ever real-time multimodal world model*" — a system that doesn't just generate a world but simulates and reasons about it.⁷⁵²

Odyssey Agora-1 is the multiplayer companion to Starchild, putting four players inside the same AI-generated world (built, in a small piece of provenance theatre, on the bones of a 1997 shooter).⁷⁵³ **Apple Headsup** is a research-grade 3D Gaussian head-reconstruction pipeline built for multi-view captures from consumer iPhones, extending the Vision-Pro-Personas Gaussian-splat thread into the open research layer.⁷⁵⁴

Underneath this layer, the Gaussian-splatting infrastructure has matured into a stable workflow: **SuperSplat** (PlayCanvas) for editing, **Spark 2.0** for open-source streaming of 100-million-splat scenes to browsers, the SOG / WebP equivalent compression standard.⁷⁵⁵ Apple's confirmation that its Vision Pro Personas feature is powered by Gaussian splatting under the hood made it, by some margin, the most-deployed Gaussian-splat technology in consumer hardware as of late 2025.⁷⁵⁶ At WWDC 2026 Apple went further, announcing that Gaussian splatting is coming to **Apple Maps** — the Flyover feature rebuilt on splats — which will make it the most-deployed splat technology in consumer *software* as well.⁷⁵⁷

For the 3D-asset and material side: **Hunyuan 3D-PolyGen 1.5** (Tencent's "art-grade" 3D generative model), **Hitem3D**, **Meshy**, **Rodin** — the rapidly-maturing 3D-asset generation layer that is being integrated, model-by-model, into Unreal Engine, Unity and Blender pipelines.

Ubisoft’s open-sourcing of its **CHORD** PBR-material model in December 2025⁷⁵⁸, and the Blender Foundation’s patronage deal with Anthropic announced in May 2026⁷⁵⁹, are two of the more strategically significant moves in this layer — both pushing the production-grade open-source tooling forward at a pace the commercial alternatives have struggled to match.

Agent platforms and orchestration

The category I think most working creatives are still underestimating, six months after Chapter 3 argued it was the inflection point of the era.

OpenAI AgentKit (Agent Builder, ChatKit, connector registry, eval framework) launched October 2025 and is the developer-facing platform underneath most third-party agentic creative tools.⁷⁶⁰

Anthropic Claude apps and the **Skills framework** — the system of named, reusable capabilities that Claude Code uses to coordinate multi-agent workflows. The Sony 49-agent / 72-skill game-development stack is built on this.⁷⁶¹ In May 2026, **Google** released its own **official skills for AI agents** — a parallel, cross-vendor skills layer that lets Google-side agents do what Anthropic’s Skills framework has been doing for Claude-side ones.⁷⁶² The convergence of two named “skills” frameworks across the foundation-model vendors is, in my read, the first sign that the orchestration layer is settling on a shared vocabulary rather than continuing to fragment.

Tencent Ardor, the company’s AI-native design-agent platform launched May 2026, is the most ambitious non-Western agent-platform launch of the period — an integrated environment in which generative design agents handle layout, asset generation, brand application and iteration as a single coordinated pipeline.⁷⁶³ In the same week, **Viktor** raised \$75M to embed an agentic *coworker* directly into Slack and Microsoft Teams — i.e., the agentic layer landing not as a standalone product but as a colleague-shaped presence in the chat surface the working creative is already in all day, as discussed in Chapter 10.

Heygen Video Agent for end-to-end video assembly.⁷⁶⁴ **Adobe CX Enterprise** (announced at Adobe Summit 2026 with NVIDIA) for “agentic creative intelligence” across the full content lifecycle.⁷⁶⁵ **NVIDIA + Google Cloud** for the underlying creative-AI infrastructure most enterprise pipelines run on.⁷⁶⁶

ComfyUI — the open-source node-based workflow editor — sits underneath much of the technical-creator community’s agentic work. ComfyUI raised \$17M in October 2025⁷⁶⁷ and hit a \$500M valuation by May 2026⁷⁶⁸; the platform has become the de facto OS for the open-source side of the creative-AI ecosystem. In May 2026 **Anthropic’s Claude** was added as an official partner node inside ComfyUI, joining

the existing Google, OpenAI and open-weight nodes — meaning the three frontier foundation models can now be orchestrated side-by-side inside the same open-source pipeline.⁷⁶⁹

Hugging Face, OpenEnv (Meta / Hugging Face), the **Hugging Face / Google Cloud** partnership — the open-source agentic-development infrastructure.⁷⁷⁰

By early July 2026 the orchestration story had a name, and the name was MCP. **Comfy MCP** launched in public beta as the first Model Context Protocol server built for production pipelines rather than one-off generations — point Claude, Codex or Cursor at it and the agent can build, edit and run ComfyUI workflows in plain language, with a shareable workflow URL that makes every generation reproducible across a team.⁷⁷¹ In the same four weeks, Epic announced **Unreal MCP** at the State of Unreal, and **Magnific, Higgsfield** and **Runway** all shipped MCP surfaces of their own.⁷⁷² The protocol the agentic ecosystem had been converging on all spring is now, in practice, the connective tissue of the creative toolchain — the shift, as ComfyUI's own framing put it, from AI that *makes* content to AI that can *operate* a creative system.

For working creatives, the practical agent stack in 2026 is some combination of:

1. A foundation model (Claude / GPT / Gemini) for the orchestration brain.
2. A modality model layer (video, image, audio, 3D) doing the actual generation.
3. A workflow integration layer (ComfyUI for technical work, the in-app agent surfaces for less technical work).
4. A judgement layer (the human at the desk, doing what Chapter 12 calls *briefing, allocating, judging and integrating*).

The team I work with at DreamLab runs this stack in production every week. The agents that handle our daily work in May 2026 are, in aggregate, doing the labour of what would, two years ago, have been a team three to four times our size. The human team has not shrunk. We have just become substantially more leveraged.

Legacy creative software, repositioned

The most under-reported strategic story of this period, in my view, has been the speed at which the legacy creative-software vendors have rebuilt their products as AI-agent platforms.

Adobe — I have written enough about Adobe in Chapter 10 that I will not repeat it here. The short version: Creative Cloud is, today, a stack of AI agents wearing a Photoshop / Premiere / After Effects / Illustrator / InDesign / Acrobat skin. The agents are inside the apps; the apps are inside ChatGPT; the apps are inside Adobe Express; the apps are inside the new CX Enterprise platform. In June 2026 the **Creative Agent and AI Assistant** went into beta across Premiere, Photoshop,

Illustrator, InDesign and Frame.io — the agent layer moving from announcement to default surface.⁷⁷³ The repositioning is complete.

Unreal Engine (Epic) — the games engine that has, through plugins, integrations and the Nano Banana / Gemini partnership, become a hybrid game-engine / virtual-production / AI-generation hub. The Unreal Engine 5 AI Assistant, announced at the end of 2025⁷⁷⁴, is one of the more consequential single-product launches of the period. The **ECABridge** connector, launched May 2026, is the most-cited Unreal-Engine MCP integration of the spring — providing the Model Context Protocol surface and a set of agentic capabilities Epic itself has not yet shipped to the launcher.⁷⁷⁵ That gap closed in June: at the **State of Unreal 2026**, Epic announced its own official **Unreal MCP** alongside LLM-powered tools and video-model rendering integrations — the clearest statement yet that Epic sees natural language as a first-class interface to the engine, and a deliberate bet on expanding *who* can become an Unreal developer rather than replacing the ones it has.⁷⁷⁶ In a separate but related move, an **Epic Games veteran** announced an AI-heavy “*Fully European*” game-engine project in the same week — the first plausibly-credible new entrant in the AAA game-engine market since the early 2010s, framed explicitly around AI as the core operating layer.⁷⁷⁷

Unity — Unity’s AI Open Beta (May 2026), an in-editor AI suite for the full games-development pipeline, alongside the company’s AI Council formation in October 2025.⁷⁷⁸

Autodesk, Foundry, SideFX — the VFX-pipeline vendors integrating generative AI into Maya, Nuke and Houdini at the speed the VFX industry’s adoption curve (62% of Hollywood studios on automated compositing, 35% reduction in post-production timelines⁷⁷⁹) demanded.

Blender — open-source 3D, now a recognised industry-grade tool, beneficiary of the Anthropic Foundation patronage deal.⁷⁸⁰

DaVinci Resolve (Blackmagic), **Avid Media Composer**, **Pro Tools** — the editorial and audio post environments, all now shipping AI-augmented features that have become baseline expectations.

The thing to note is that the legacy software did not get displaced by the AI-native products. The legacy software *absorbed* the AI-native capability and kept the underlying user community. Adobe was supposed to lose to Midjourney in 2024; Adobe is, instead, the dominant generative-AI player by aggregate creator engagement in 2026. The platform companies bet on this absorption pattern, and that bet has, so far, paid off.

Open source

The open-source ecosystem has, against the odds and against most VC predictions in 2024, held its ground through this period and is, in several categories, the leader rather than the follower.

Hugging Face — the operating system of open-source AI, expanding aggressively through 2025–26.

ComfyUI — already discussed.

Open-source models from Tencent (Hunyuan), Alibaba (Qwen, Wan), DeepSeek, Meta (Llama), Mistral, Stability AI — collectively, the open-weight ecosystem that, by the spring of 2026, was being used by approximately 80% of startups pitching the Andreessen Horowitz fund.⁷⁸¹ **NVIDIA's SANA-WM** (May 2026) extended this list to world-models for the first time at meaningful parameter scale.⁷⁸²

PhotoGIMP, the open-source skin that takes GIMP and makes it look and feel exactly like Photoshop, became, in this period, a credible Photoshop alternative for working creatives who wanted to opt out of the Adobe subscription stack — the open-source equivalent of the *Tools I do not use* discipline in the section above.⁷⁸³

OpenEnv (Meta / Hugging Face) for open-source agentic development. **Korin AI** (the Africa-trained, Africa-built model launched May 2026⁷⁸⁴). **SuperSplat, Spark 2.0, PlayCanvas SOG, Blender** — the open-source spatial / 3D infrastructure layer.

If you are a working creative trying to build a long-term, defensible toolchain that does not depend on the unilateral pricing or policy decisions of a single platform company, the open-source ecosystem in 2026 is materially viable in a way it was not eighteen months ago. We have built significant parts of the DreamLab pipeline on top of it precisely for that reason.

Tools I do not use, and why

I want to be specific, because lists of “best tools” without exclusions are not useful.

I do not use AI tools whose terms of service claim ownership over my output, or that train on user inputs without an opt-out. Multiple consumer-facing AI products in this period have shipped with terms that working creatives should read carefully before adopting.

I do not use AI tools whose training data provenance I cannot, in some material way, verify or trust. The growing infrastructure for *creative weight attribution*, watermarking and C2PA compliance is, in my view, the right side of the market to be on; tools that explicitly reject that infrastructure are tools that I have, increasingly, kept out of our production pipeline.

I do not use the AI products that have made the most noise in the consumer press cycle. The marketing-driven launches — the products whose first appearance is a viral demo and whose second appearance is a Series-A round — are, in my experience, the products most likely to have collapsed or pivoted by the time you need them in production six months later.

I do not, finally, use AI tools to produce work in the disciplines where my own craft is the value I am bringing to the client. The Continuum frame from Chapter 3 is, for me, a daily operational practice, not a theoretical model. The places I sit on the right edge of the line are deliberately chosen. The places I sit on the left are deliberately defended.

The complete toolchain: a categorised reference

This section is a reference inventory, not a recommendation list. It catalogues every tool, model, platform, app, plugin, LoRA, workflow and service that *Dream Machine* tracked across its 41 issues, from October 2025 to August 2026. Some are dominant; some are niche; some have already been bought, renamed or discontinued by the time you read this. The point of the list is not “what to use.” The point is *what existed*, in this period, in the creative-AI toolchain — so that the *shape* of the field is on the record.

A word on the list’s grain. I have tried to err on the side of inclusion. Where a single company ships multiple closely-related products — Adobe’s *Sneaks* portfolio, the Runway Gen-4.5 family, the Qwen-Edit LoRA series — I have grouped them under the parent entry but called out the constituent tools, because in this period each constituent shipped to working creatives separately and changed at its own cadence. Where a tool was a one-issue demo I could not later verify, I have still listed it; that the demo existed *at all* is part of the field’s history. Where a tool’s name conflicts with another (there are at least three things called “Wonder” in the period the book covers) I have annotated.

The list runs to roughly six hundred and thirty entries. Skim it. Use the categories. Come back to specific sections when you need them.

Foundation models / LLMs

- **ChatGPT / GPT-5 / GPT-5 Pro / GPT-Live** (OpenAI) — the dominant consumer LLM and reference foundation model; 800–900M weekly active users; GPT-5 / GPT-5 Pro announced at DevDay 2025; **GPT-Live** (July 2026) is a new generation of voice models for natural, real-time human-AI interaction.
- **Claude / Claude Code / Claude Apps / Claude Skills / Claude Fable 5 / Claude Sonnet 5** (Anthropic) — the writers’ and developers’ favoured second; strong long-context performance; the agentic coding environment that underlies Sony’s 49-agent / 72-skill stack; Claude for Legal launched May 2026. **Fable 5**,

the Mythos-class model released June 2026, had its access suspended within a week under a US export-control directive — the first frontier creative model to be withdrawn on national-security grounds; **Sonnet 5** (July 2026) shipped stronger agentic capabilities at lower pricing.

- **Gemini / Gemini 2.5 Flash / Gemini 3 / Gemini 3.1 Flash / Gemini Omni Flash / Gemini 3.5 Live Translate** (Google) — Google’s multimodal LLM family; desktop users growing 155% YoY in 2025–26; Gemini 3.1 Flash TTS is the most controllable Google voice model as of spring 2026; Gemini Omni Flash (July 2026) generates and edits video via the Gemini API, Google AI Studio and ComfyUI; Gemini 3.5 Live Translate ships low-latency voice translation across 70+ languages.
- **Llama** (Meta) — the dominant open-weight foundation model.
- **Mistral / Mistral Voxtral / Mistral Transcribe 2** — European open-source LLM; Voxtral is the next-generation speech-to-text family.
- **Qwen / Qwen 3.5-Omni / Qwen 3.8 Max** (Alibaba) — Chinese open-source LLM, image, video and audio variants; Omni is the multimodal family covering text, images, audio and video; **Qwen 3.8 Max** (August 2026) set a new bar for coding and agentic “cowork”.
- **DeepSeek** — Chinese open-source LLM used heavily in agentic stacks.
- **Phi / Phi Mini** (Microsoft) — lightweight foundation models.
- **Lyria / Lyria 3 / Lyria 3 Pro / Lyria 3.5** (Google DeepMind) — text-to-music foundation model family; Lyria 3 ships with SynthID watermarking and the Lyria Camera interactive demo; **Lyria 3.5** (July 2026) shipped inside Google Flow Music.
- **Grok** (xAI) — xAI’s LLM family, surfacing in the creative stack through Grok Imagine; voice cloning via the xAI API from May 2026.
- **Cohere Transcribe** — enterprise transcription model; 33 hours of audio in 12 minutes.
- **Microsoft VibeVoice** — open-source frontier voice AI model.

AI video models

- **Sora / Sora 2** (OpenAI) — the model that opened the period; physical realism, audio integration, multi-shot world-state persistence; iOS app hit 1M downloads in 5 days; Sora 2 Character Creation surfaced on fal in March 2026.
- **Veo 3 / Veo 3.1 / Veo 3.1 Ingredients to Video / Veo 3.1 Lite** (Google DeepMind) — the working filmmaker’s preferred model for cinematic control; Ingredients to Video shipped to YouTube Shorts and YouTube Create; Veo 3.1 Lite is the lower-cost text- and image-to-video tier.

- **Runway Gen-4 / Gen-4.5 / Gen-4.5 Image-to-Video / Workflows / Story Panels / Characters API / Ad Concepter / Apps for Advertising / Aleph 2.0** (Runway) — the most-integrated commercial AI-video stack; Story Panels generate three-panel storyboards from a single image; Characters API ships real-time intelligent avatars; the Vera Rubin-powered real-time model runs <100 ms latency. **Aleph 2.0** (June 2026) lets the user edit a *single frame* in a video and propagates that edit consistently across the rest of the clip — the cleanest *frame-as-handle* video-editing pattern yet shipped. **API Recipes** (June 2026) drop production-ready generative-media features into any platform with one API call; the **Runway MCP** puts Runway generation and editing inside ChatGPT; **Agent 2** (July 2026) handles ad localisation.
- **Kling / Kling 2.5 Turbo / Kling O1 / Kling 2.6 / Kling 3.0 / Kling X-Dub / Kling Motion Control 3.0** (Kuaishou) — strong on physics and trajectory control; 3.0 adds multi-shot control, multilingual audio and 4K image generation; X-Dub is the context-rich visual dubbing variant.
- **Pika 2.0 / PikaStream 1.0 / Pika Director’s Suite** — iteration-speed-focused video generation; PikaStream brings AI agents into live video calls; **Director’s Suite** (July 2026) is Pika’s experimental, invite-only director-controls video product.
- **Luma Dream Machine / UNI-1 / UNI-1.1 / Ray3 Modify / Luma Dream Brief / Luma Layers** — Luma’s video, world-and-reasoning, and modification stack; UNI-1.1 ships with prompt enhancement and built-in research; Dream Brief is the \$1M Cannes Lions competition; **Layers** (August 2026), live in Luma Agents, extracts editable layers from any image on request.
- **Wan 2.2 / Wan 2.5 / Wan 2.6 / Wan 3 / Wan-Dancer-14B / Wan Animate 2** (Alibaba Qwen) — camera-controlled video generation; 2.6 adds character reference and multishot capabilities; **Wan 3** (August 2026) entered public beta with 30-second clips and multimodal inputs; **Wan-Dancer-14B** (July 2026) is the open-source “music-to-dance” model that generates long-duration rhythmic dance video locally; **Wan Animate 2** (August 2026) is a state-of-the-art character-animation model that transfers motion from a reference video onto a single character image.
- **Hunyuan Video / Hunyuan Image-to-Video** (Tencent) — open-source video model.
- **Seedance 2.0 / Seedance 2.0 Mini / Seedance 2.5 / SeeDream 4 / SeeDream 4.5 / Seedream 5.0 Pro** (ByteDance) — image-to-video and finished-video models, integrated into CapCut / Dreamina and Freepik; per-second cost fell below \$0.14 by March 2026; **Seedance 2.0 Mini** (June 2026) is the quality-at-lower-cost tier in Dreamina, and in ComfyUI (with 4K) from July; **Seedance 2.5** (June 2026) broke the thirty-second barrier for AI video generation and went live across fal, Dreamina and ComfyUI Partner Nodes

through July–August 2026; **Seedream 5.0 Pro** (July 2026) is ByteDance’s latest image model — character and product consistency across edits, precision region editing and in-image text across 14 languages.

- **Higgsfield / Higgsfield Sketch-to-Video / Higgsfield Popcorn / Higgsfield Relight / Higgsfield Shots / WAN Camera Control / Higgsfield Adobe Premiere Plugin / Higgsfield After Effects Plugin / Higgsfield DaVinci Resolve Plugin / Higgsfield MCP** — social-media-marketer video platform; \$80M raised at \$1.3B valuation, \$200M revenue in nine months (reportedly \$500M by June 2026); Shots produces multiple storyboard images from a single shot; the Adobe Premiere Pro and After Effects plugins (June 2026) bring Higgsfield generation directly into the professional editorial workflow; the DaVinci Resolve plugin adds in-timeline generation, draw-to-edit and AI-generated LUTs; Higgsfield MCP (July 2026) lets agents such as Claude generate through the platform directly.
- **Bernini** (ByteDance) — AI video generation and editing framework; a complementary video-editing layer to ByteDance’s Seedance generation stack.
- **LTX-2 / LTX-2.3 / LTX-2.3 Colorizer / LTX-HDR / LTX Studio / LTX-2 Audio-to-Video / LTX-2 Lip Sync / LTX-2 Real-Time / LTX-2.5** — open-source video generation with audio sync; LTX-2.3 is high-resolution, fast, cinematic with native lip-sync; LTX HDR (beta) ships HDR processing; **LTX-2.5** (August 2026) is Lightricks’ open-world, open-source model.
- **Odyssey 2** — real-time interactive video generation.
- **Vidi 2** (ByteDance) — multimodal video understanding and creation.
- **MotionStream** — real-time interactive video with mouse-based motion control.
- **Blockvid** — one-minute video with improved structure and visual consistency.
- **Video Rebirth** (Singapore) — studio-grade AI video platform; \$50M raise.
- **LiveGS** — mobile Gaussian-splatting video.
- **Time-To-Move** — motion control for generated video.
- **Ovi / Ovi 1.1** (Character.AI) — open-source video with speech sync.
- **CraftStory** — image-to-video for long-form AI video with human “actors”.
- **Decart Lucy 2.0** — realtime world-editing video model; 1080p at 30 fps.
- **Decart LSD v2** — real-time video-to-video with prompt-on-the-fly.
- **Xmax X1** — the first real-time interactive video model.
- **FastVideo** — 30-second 1080p generation at 4.5 s latency.
- **HiAR** (Tencent) — scalable long-video generation.
- **Alli AI** — system for creating and editing 8K video up to 60 seconds.
- **AERA AI TV** — automated storyboard generation, editing and serialized video production.

- **AnchorWeave** — world-consistent video generation with retrieved local spatial memories.
- **CubeComposer** — generates native 4K 360° video from standard perspective footage.
- **MatAnyone2** — high-fidelity video matting with finer detail.
- **VFace** — training-free video face-swapping for any diffusion model.
- **DreamActor M2.0** (fal) — drive characters from a single image + template video; multi-character supported.
- **ByteDance ALIVE** — unified audio-video generation.
- **Magnific Upscaler for Video / Magnific Precision v2 / Magnific Precision for Video / Magnific MCP / Magnific Plugin** — upscaling and 4K detail enhancement, including dedicated video upscaling; Magnific MCP (June 2026) connects every Magnific tool to your AI assistant; the Magnific Plugin (July 2026) puts the stack inside Premiere Pro, After Effects, DaVinci Resolve, Final Cut Pro and Adobe Illustrator.
- **NetFlix VOID** — object removal from video with physics-interaction removal.
- **Ponder** — agentic video editor.
- **ArcReel** — multi-agent video generation from written stories.
- **Scope** — real-time interactive generative AI pipelines (LTX-2.3 real-time runs on Scope).
- **NotebookLM Cinematic Video Overviews** (Google) — video generation from notebooks.
- **Omnia** — AI-native browser video editor.
- **NVIDIA RTX Video Super Resolution** — upscaling node with ComfyUI integration.
- **Google M2SVid** — monocular-to-stereo video conversion.
- **Alibaba WonderClip** — all-in-one AI video creation platform from Alibaba Cloud (June 2026).
- **Sony Vid-CamEdit** — post-capture camera trajectory editing; keeps scenes spatially coherent across large viewpoint changes (June 2026).
- **Nvidia PiD (Pixel Diffusion Decoder)** — image / video decoder; new entry from Nvidia (June 2026).
- **VEED Fabric 1.0** — hyper-realistic AI explainer and educational videos, up to five minutes, generated in seconds.
- **NVIDIA SANA-Streaming** — high-resolution, real-time video-to-video editing on consumer GPUs; the streaming sibling of the SANA-WM world model.
- **Viggle / Viggle API** — give any character any motion; one API call and it comes back alive in seconds.

- **Pixelcut Video Background Remover** — video background removal, live on fal.
- **LooseControlVideo** — directorial video control using spatial blocking; the research end of the *blocking-as-interface* pattern.
- **Musavir LiveAI** — transforms a live video feed in real time.
- **NVIDIA Metropolis VSS Blueprint 3** — video search and summarisation; lets a coding agent analyse massive live streams and video libraries from a natural-language prompt.
- **MiniMax Hub** — all-in-one AI video generator hub, launched at the Shanghai Film Festival with the stated caveat that “creative direction and aesthetic judgment must ultimately be left to humans”.
- **MiniMax H3 / MiniMax H3 LoRA trainer** — open-weights video model with day-0 ComfyUI node support (August 2026); runs locally, even on a Mac; a fal-hosted LoRA trainer lets creators fine-tune it, and the model anchors a widely-shared music-video workflow. (Distinct from MiniMax Hub and the MiniMax Music Generator.)
- **NVIDIA DeepStream 9.1** — video-analytics pipeline SDK shipping 13 agentic skills for building multi-camera 3D-tracking and video-understanding applications (August 2026).
- **Google Video Remix** (Google Photos) — in-app AI video editing (July 2026): swap backgrounds, fix lighting, and apply artistic effects such as watercolour and oil painting.

AI image models / tools

- **Midjourney / Midjourney V8.1 / Midjourney 8.2** — the aesthetic-leadership product; Discord/X-native; V8.1 ships native 2K HD rendering at 3× the speed and 3× the cost reduction; **8.2** (July 2026) is the latest release.
- **FLUX / FLUX 2 / FLUX 2 Max / FLUX.2 [klein] / FLUX 3 / FLUX 3 Video** (Black Forest Labs) — open-weight, fine-control, the open-source default through 2025–26; klein is the 4B-parameter lightweight model; **FLUX 3** (July 2026) generates images and 20-second video with audio; **FLUX 3 Video** (August 2026) is the dedicated video model.
- **Meta Muse Image / Muse Code / Muse Spark 1.2** (Meta) — Meta’s Muse creative-AI family: **Muse Image** (July 2026) is the consumer image generator; **Muse Code** and **Muse Spark 1.2** (August 2026) extend the family into generative coding and rapid experience-building.
- **Adobe Firefly / Firefly Image Model 5 / Firefly Foundry / Firefly Boards / Firefly Precision Flow** — Image Model 5, Foundry (custom corporate training), Firefly Boards (moodboards), Precision Flow (granular AI

editing control, beta); 45% of Creative Cloud users active, 22B+ assets generated by April 2025.

- **Imagen 3 / Nano Banana / Nano Banana Pro / Nano Banana 2 / Nano Banana 2 Lite** (Google) — most-integrated image model in the consumer toolchain; Photoshop and Unreal Engine plugins; Nano Banana Pro ships professional capabilities at lightning speed; Nano Banana 2 Lite (aka Gemini 3.1 Flash-Lite, July 2026) targets low-cost, four-second enterprise image generation.
- **Stable Diffusion / Stable Diffusion 3 / Stable-Layers** (Stability AI) — the foundational open-source image model; **Stable-Layers** (June 2026) splits generated images into editable compositional layers — a structural shift from regeneration to post-generation editing.
- **Ideogram 4.0 / Ideogram Object Remover** — fourth-generation image model; open weights, fine-tuneable on own data, runs on own hardware; available on every Ideogram plan and via API; the **Object Remover** (July 2026) ships state-of-the-art object removal.
- **Krea / Krea AI / Krea 2 / Krea Realtime / Krea Realtime Edit / Krea Nodes / Krea LoRA Trainers** — real-time AI image generation; Krea 2 shipped open weights in July 2026, with one model for inference and one for training; Realtime Edit takes complex instructions in real time; the LoRA Trainers cover Qwen-2512 and Z-Image.
- **Qwen-Image-2512 / Qwen-Image-Edit-2511 / Qwen 2511 Time Travel / Qwen-Edit 2509** — the dominant open-source image-editing model family; constituent LoRAs (relighting, multi-angle, time-travel, AnyPose) discussed in the LoRAs section below.
- **ChatGPT Images / ChatGPT Images 2.0** — image generation integrated with Adobe Express, Photoshop and Acrobat; 2.0 ships thinking-level intelligence.
- **Grok Imagine / Grok Imagine API (xAI)** — image and video generation; the API bundles end-to-end creative workflows; reference-to-video and video extend added in March 2026.
- **Vision Banana** — unified model for image understanding and generation.
- **Freepik / Freepik Spaces / Freepik Speak / Freepik 3D Scenes** — design platform; Spaces is the real-time collaborative canvas; Speak is the lip-sync talking-video tool; 3D Scenes generates full environments from an image.
- **Recraft** — brand-focused AI design.
- **Weavy** — node-based AI creative platform; acquired by Figma.
- **Canva** — design platform with “Creative Operating System.”
- **Civitai / Civision** — community-models and LoRAs platform; Civision is the free Civitai/Hugging-Face hybrid alternative.

- **Replicate** — model-hosting marketplace and API.
- **GFPGAN** — face restoration.
- **UpscalyAI / Upscayl / Topaz Gigapixel AI / Real-ESRGAN / Clarity Pro Upscaler** — upscaling tools; Clarity Pro reaches 10K.
- **PractiLight** — practical light control via diffusion.
- **PercHead** — 3D head reconstruction from single images.
- **Loveart AI** — layer separation and editable text (Live Editable Text, LET).
- **NVIDIA ChronoEdit** — temporal reasoning for image editing.
- **Beeble Switchlight 3 / Beeble Background Remover** — AI masking and relighting.
- **CanvAI** — turns rough sketches into polished AI art.
- **MakeComics / Make Comics** — generates custom comics with characters and storylines in seconds; from prompt to full comic book.
- **Seethrough** — illustration decomposition into layered PSD.
- **Best Face Swap (Flux 2 LoRA)** — specialised face-swap.
- **PixelSmile** — fine-grained facial-expression editing across 12 expressions.
- **Earth Cinema** — Chrome extension using Google Earth for cinematic images.
- **360Anything** — lifts any perspective image or video to a 360° panorama.

AI music / audio tools

- **Suno / Suno Studio / Suno 5.5 / Suno Spark** — the dominant prompt-to-track generative music platform; \$400M raised at \$5.4B valuation (June 2026), \$300M ARR, 2M+ subscribers; Warner Music Group settled its lawsuit and signed a licensing partnership; UMG and Sony remain in active litigation; Suno Studio is the “world’s first generative audio workstation”; 5.5 adds Voices; Spark (July 2026) is the incubator programme for the next generation of creators.
- **Udio** — prompt-to-music; partnered with Universal Music Group; indie-label licensing via Merlin.
- **Mureka / Music Agent Studio** — six specialised AI agents covering songwriting, arrangement and production.
- **Google Flow Music** — Google’s generative-music studio inside Flow (July 2026); upload and remix your own tracks, with song creation, editing and stem-splitting added through the summer; powered by Lyria 3.5.
- **MuScriptor** (Mirelo AI / Kyutai Labs) — audio-to-MIDI transcription (July 2026): takes a finished recording, identifies the instruments and returns separate MIDI tracks for voice, drums, bass, keys and more.

- **SCENA** (Lightricks) — reference-driven, multi-speaker audio-scene generation (July 2026).
- **Sonilo Sound Effects 1.0** (Sonilo + fal) — realistic sound effects generated from video and text (July 2026); distinct from the earlier Sonilo × Shutterstock training deal.
- **ByteDance SwanTale** — a unified audio model (August 2026) that collapses voice, sound and music into a single system.
- **ElevenLabs / ElevenLabs v3 / Scribe v2 Realtime / ElevenMusic / ElevenCreative / ElevenLabs Flows / ElevenAgents Expressive Mode / ElevenLabs Voice Changer / ElevenLabs Music v2 / ElevenLabs Dubbing v2 / ElevenLabs Character Casting / ElevenLabs Avatars** — the dominant voice/audio synthesis stack; **Avatars** (June 2026) put a face on the voices inside ElevenCreative; \$500M ARR; BlackRock + NVIDIA backed; Scribe v2 transcribes in 150 ms across 90+ languages; ElevenMusic is the discovery/creation/earning marketplace; Flows is the node-based creative canvas; **Music v2** (June 2026) is the next-generation generative-music model; **Dubbing v2** (June 2026, on ElevenAPI) is the next-generation dubbing model; **Character Casting** (August 2026) finds and assigns a distinct voice to every character in an ElevenCreative audiobook. ElevenLabs also partnered with Hasbro's Sixth Wall AI studio in June 2026, providing the voice layer for Hasbro's Behavioral Licensing / CharacterOS programme — enabling licensed AI voices for Transformers, Mr. Potato Head, Peppa Pig and more.
- **iZotope Ozone 12 / Stem EQ** — AI-assisted mixing and mastering.
- **LANDR** — AI mastering and distribution.
- **Riffusion** — spectrogram-based music generation.
- **Stable Audio 2.5 / Stable Audio 3.0** (Stability AI) — generative audio. **Stable Audio 3.0** (May–June 2026) is an open-weight model family trained on fully licensed data, designed as the foundation for what the open audio community builds next; ships in ComfyUI for tracks up to about six minutes.
- **MusicGPT** — local music generation.
- **ACE Studio / ACE Studio 2.0 (TIMEDOMAIN) / ACE Studio Video-to-Music / ACE-Step 1.5 / ACE-Step 1.5 XL** — all-in-one AI music studio; video-to-music for visuals; ACE-Step generates full songs in under 10 seconds on <4 GB VRAM.
- **Music Lens** (Musixmatch) — catalog-intelligence agent.
- **Melosurf** — voice-controlled assistant for Ableton Live.
- **Songscription** — audio-to-notation transcription; \$5M raise.

- **Groundhog Audio Pedal / Groundhog OnePedal** — AI guitar tone matching.
- **Spotify DJ / Spotify AI / Spotify AI Prompted Playlists / Spotify AI Transparency Beta / Personal Podcasts on Spotify** — playlist personalisation, voluntary AI disclosure beta, agent-created podcasts.
- **YouTube Music AI Hub** — AI music hosting and DJ features.
- **Epidemic Sound Studio** — AI music for video creators.
- **AIODE** — ethically-trained music creation DAW.
- **Overtune** — virtual music studio for Roblox.
- **Space DJ** (Google DeepMind) — interactive music exploration.
- **Minimax Music Generator** — song generation.
- **Roland AI Pedal** — concept for AI audio processing.
- **Roland + Sony CSL Melody Flip** — AI music tool.
- **BandLab Voice Cleaner / BandLab Palette** — background-noise removal; AI loop-matching.
- **Claimy** — missing-royalty recovery agent.
- **Fish Audio S1 / Fish Audio S2** — text-to-speech; S2 ships expressive TTS with emotional control (6× cheaper than ElevenLabs at S1).
- **Tempolor Guitars** (Quwan) — AI to make songs playable on guitar.
- **GEMIDI** — music playground powered by Gemini.
- **Lalal AI / Lalal AI plugin / Lalal AI API v1** — AI stem separation, music removal, voice changing; first stem-separation plugin native to a DAW.
- **StemDeck** — local stem separation; free, no upload required.
- **Mirelo** (Berlin) — automatic sound-effect generation for video.
- **PromptSep** — separates any sound by text description.
- **SAMAudio** (Meta) — Segment Anything for audio editing.
- **SonicLab SPATAI** — generative audio engine for immersive spatial production.
- **Apple Logic Pro AI** — synth player and personal music-theory expert; chord-progression generation.
- **Apple GarageBand** — consumer music creation.
- **Musicful AI** — music-video generation for AI-generated songs.

- **Latent Space Explorer** — interactive audio data visualiser using Stable Audio VAE.
- **Voicebox** — open-source local voice cloning from a 3-second audio clip.
- **Voice-Pro** — AI speech recognition / translation / transcription / multilingual dubbing.
- **PersonaPlex** (NVIDIA) — full-duplex model that listens and speaks simultaneously; open-source natural-sounding voice.
- **Phoenix-4** — advanced real-time human-rendering model with 10+ emotional states.
- **NVIDIA D-Rex** — photorealistic digital humans under any lighting.
- **Qwen3-TTS / VoiceDesign / VoiceClone** — text-to-speech with voice design and cloning.
- **Greysound** — self-engineering music-production studio.
- **TAC (Timestamped Audio Captioning)** — timestamped captions for audio and audiovisual content.
- **Insanely Fast Whisper** — 2.5 hours of audio transcribed in 98 seconds.
- **smol-audio** — local audio model notebooks and scripts.
- **AudioStream / Walzersymphonie** (Ars Electronica Futurelab) — AI composition system for classical music using Ricercar.
- **Ricercar** (Ars Electronica) — creative AI for artistic composition.
- **Krotos Video-to-Sound** — expanded platform for audio professionals (foley, sound design).
- **Jamu** — AI co-producer agent for Ableton Live.
- **Moises** — AI music platform; Charlie Puth as Chief Music Officer; its stem separation now ships inside Fender's Studio Pro DAW.
- **Music Mogul AI** — tour-booking automation.
- **Sonilo + Shutterstock** — video-to-music AI training deal.
- **Clearnote** — AI music-contract platform to end deal delays.
- **Sony music identification tech** — identify original music inside AI-generated songs.
- **Cactus Music** — artist-ops support.
- **Rebel Audio** — AI podcast startup.

- **ROLI Airwave & AI Music Coach** — hand-tracking (27 joints) with real-time AI piano practice coaching.
- **Sony AI ICASSP papers** — music understanding and generative audio research papers (April 2026).
- **Google Magenta RealTime 2 (MRT2)** — the live music model you can play as an instrument; the flagship of the June 2026 *instrument turn*.
- **Daydream** — “the first AI-native instrument for your DAW.”
- **Fender Studio Pro** — Fender’s DAW with Moises stem separation and a smart studio assistant built in; “AI isn’t the destination. Making music is.”
- **Maono PD500W** — “the world’s first AI-powered hybrid microphone”; AI processing arriving in the capture hardware itself.
- **Orchestria** — AI music engine with granular stem control.
- **Wubble** — ethically-created royalty-free music, AI voiceovers and sound effects from prompts.
- **Tencent MMAE** — Massive Multitask Audio Editing benchmark; the first comprehensive evaluation benchmark for speech and audio editing.
- **Bandsintown Boost / Laylo AI Ticket-Sales Agent** — AI tools for touring artists: promotion boosting and an agent that sells tickets.
- **Artist (LA)** — music-tech startup letting artists re-record classic tracks with AI-assisted vocal technology; Boy George’s *Karma Chameleon* re-record was the first release. (Not to be confused with the generic use of the word elsewhere in this book.)
- **Tringbox** — Indian AI background-music startup; seed-funded June 2026, also working with live venues.
- **Marble / Marble 1.1 / WorldLabs API / RTFM** (World Labs / Fei-Fei Li) — first commercial generative world model; 40× faster than legacy VP workflow; 1.1 adds real-world location 3D reconstruction and restyling; the API generates persistent 3D worlds from text, images and video; RTFM is the real-time frame model.
- **Genie 3 / Project Genie** (Google DeepMind) — research-grade world model; *Time Best Inventions 2025*; rolled out to Google AI Ultra in January 2026.
- **WorldGen / WorldMirror** (Meta / Tencent) — interactive 3D world generation.
- **UNI-1 / UNI-1.1** (Luma) — world generation + reasoning combined; 1.1 adds prompt enhancement and built-in research.

- **Hunyuan 3D / Hunyuan 3D Studio / Hunyuan 3D-PolyGen 1.5 / HY3D 3.0 / HY-Motion 1.0 / Hunyuan3D-Buffer 1.0** (Tencent) — art-grade 3D generative model family; PolyGen 1.5 is the art-grade text/image/sketch-to-3D model; HY-Motion 1.0 is text-to-3D human motion; **Hunyuan3D-Buffer 1.0** (August 2026) is a unified foundation for 3D generation, understanding and editing.
- **HY World 1.5 / HY-World 2.0 / WorldCompass / ShotVerse / WorldClaw** (Tencent) — open-source real-time world-model frameworks; WorldCompass is the RL post-training framework; ShotVerse handles cinematic multi-shot camera control; **WorldClaw** (August 2026) is agentic 3D open-world generation at scale.
- **ECHO / Echo-2** (SpAItial) — spatial foundation model; Echo-2 adds world decomposition.
- **Wonderzoom** (Stanford) — multi-scale 3D world generation with infinite zoom.
- **Matrix-Game 3.0** (Skywork AI) — real-time interactive world model; 720p at 40 fps with long-horizon memory.
- **AMD Micro-World** — AMD’s entry into world models.
- **LingBot-World / LingBot-World 2.0** — free/open-source world model on Wan 2.2; real-time interaction at 16 fps; **2.0** (July 2026, from Reactor) is the new real-time world-generation model.
- **Moonlake** — \$28M-seeded world model platform now in beta for games and simulations.
- **Project Eden** (TripoAI) — persistent, multiplayer world model (June 2026) that fundamentally breaks from existing paradigms by *decoupling the underlying world state from visual rendering*; the most ambitious state-vs-rendering separation in the world-model layer at the time of writing. TripoAI raised nearly \$200M in June 2026 to advance this work.
- **TripoSplat** (Tripo) — single 2D image to high-quality 3D Gaussians; open-source under the MIT license.
- **NVIDIA Cosmos 3** — omnimodal world model (released May 31, 2026) that connects understanding, generation, simulation and action through a shared architecture moving fluidly across text, images, video, audio and actions; Two-Tower Mixture-of-Transformers design; ranked best open-source T2I and I2V models by Artificial Analysis; explicitly targets robotics, autonomous vehicles and warehouse monitoring. The most significant NVIDIA world-model release in the period.

- **NVIDIA FlashDreams** — open-source world model from NVIDIA, released June 2026.
- **PaGeR** (Google & Meta) — AI model for 360° scene geometry reconstruction.
- **NVIDIA Lyra 2.0** — explorable generative 3D worlds.
- **NVPanoptix-3D** (NVIDIA) — single-image 3D indoor-scene reconstruction.
- **Kimodo** (NVIDIA) — kinematic motion-diffusion model trained on 700 hours of mocap.
- **Animotive Kimodo** — AI text-to-animation directly inside Unreal Engine, on the FAB Marketplace (a name collision with NVIDIA's Kimodo above; the two are separate products).
- **NVIDIA MotionBricks** — real-time character animation at scale, without hand-crafted transitions or fine-tuning.
- **NVIDIA ARDY** — real-time interactive human-motion generation model (July 2026).
- **Disney Generative Motion Rig** (Disney Research Zurich / ETH Zurich) — a generative motion rig for artist-driven motion authoring (July 2026).
- **Rokoko Vision 3.0** — markerless, video-based motion capture (August 2026).
- **MAMMA** — markerless motion-capture pipeline recovering SMPL-X parameters from multi-view video for two-person interaction.
- **MetaHuman Animator Markerless Mocap Plugin** (Epic) — turn any video into mocap animation in Unreal; live on FAB.
- **Mint** — rigs and animates a 3D model in ten minutes.
- **InSpatio-WorldFM** — open-source real-time generative frame model.
- **Code2Worlds** — workflow for generating 4D scenes with environmental/object generation and feedback refinement.
- **PARSE** — framework for procedurally generating spatially grounded 3D scenes from part-level relational descriptions.
- **DreamForge** — a sandbox for interactive worlds (July 2026; one-issue sighting, listed for the record).
- **NVIDIA XR AI** — public beta for building AI agents for AR glasses and XR devices.
- **OpenArt Worlds** — 3D navigable environments from a single prompt.
- **Wonder by Adobe** (Adobe Research) — a real-time interactive 3D world model for navigating generated worlds (July 2026). (One of at least three

unrelated products called “Wonder” in the period this book covers.)

- **AlayaWorld** (Alaya Lab) — open-source long-horizon, playable video-world generation (July 2026).
- **Wayve GAIA-4** — the latest evolution of Wayve’s line of world models (August 2026), rooted in the autonomous-driving lineage.
- **Rodin / Hyper3D / Rodin Hyper 3D Gen 2 / Rodin Gen-2.5** — high-precision 3D model generation; Gen-2.5 (July 2026) is the world’s first 10-million-polygon AI 3D model — 1 million polygons in 4 seconds, scaling to 10M with skin-level microstructure detail, plus dreamy emissive bloom textures; 8K textures became free for all Hyper3d users the same month.
- **Meshy / Meshy 6 / Meshy Image-to-3D** — 3D model generation; Meshy 6 ships natively inside ComfyUI; Image-to-3D generates poses.
- **Hitem3D** — high-resolution 3D generation.
- **Microsoft Trellis 2** — native compact structured latents for 3D.
- **ByteDance Seed3D 2.0** — 3D object generation from image or text.
- **Pixel3D** (Tencent) — 3D object generation.
- **Tripo 3.1 / Tripo 8K Texture** — 3D-asset creation; ComfyUI partner nodes; 8K Texture (June 2026) applies 8192×8192 detail directly onto 3D models, preserving surface fidelity without upscaling.
- **PATINA** (fal) — image or text to full PBR material maps.
- **M-XR PBR Model** — 4K PBR material maps for 3D assets.
- **CHORD** (Ubisoft La Forge) — open-sourced end-to-end PBR-material generation.
- **Autodesk Wonder 3D** — generative AI tool creating editable 3D assets from text or images.
- **Autodesk Flow Studio Rigging** — AI rigging and neural layers.
- **AI4AnimationPy** (Meta) — pure-Python 3D character animation.
- **MocapAnything** (Huawei) — unified 3D motion capture for arbitrary skeletons from monocular video.
- **FreeMoCAP** — open-source markerless mocap from any camera.
- **MuJoCo** — physics engine for robotics and animation.
- **Mesh2Motion** — auto-assigns and exports 3D model animations.
- **One-to-All Animation** — alignment-free character animation and image pose transfer.

- **Cascadeur / Cascadeur AI Animation v2025.3 / Cascadeur 2026.2** — AI animation tool for character movement; local interpolation; **2026.2** (August 2026) adds additive animation layers (alpha), eased interpolation, AutoPosing improvements, FBX material import and custom UI via a Python API.
- **Move AI** — motion capture.
- **MoCap / UE5 Motion** — camera-only motion capture.
- **SCAIL** — studio-grade character animation via in-context learning.
- **Vanast** — garment-transferred human animation from images.
- **CompHairHead** — one-shot 3D head avatars with deformable hair.
- **FlexAvatar** — 3D head-avatar generation with partial supervision from a single image.
- **Face Anything** — 4D face reconstruction from any image sequence.
- **Reallusion Headshot 3** — digital-double creation.
- **MoRo** (Meta) — human motion recovery via masked modelling for occlusions.
- **Meta Sapiens 2** — pose estimation, body-part segmentation and surface normals.
- **NVIDIA Audio2Face / Audio2Face-3D** — facial animation from audio.
- **NVIDIA UniRelight** — shot relighting technology.
- **Disney Neural Render Proxies** (Disney Research) — neural render proxies for interactive and differentiable lighting (July 2026).
- **Apple SHARP / LiTo** — photorealistic 3D view synthesis from a single image in under a second; LiTo is the surface light-field tokenisation method.
- **Pixel3DMM** — screen-space single-image 3D face reconstruction.
- **VISTA4D** (Netflix + Eyeline) — live-action to navigable 4D point clouds.
- **SuperSplat / SuperSplat v2.16** (PlayCanvas) — free, open-source Gaussian-splat editor; the July 2026 release added collision/proxy-mesh generation (via splat-transform, exportable to GLB), software attribution and a histogram view.
- **Spark 2.0** — open-source Gaussian-splat streaming framework; streamable LoD for WebGL2.
- **Gracia AI** — 4D Gaussian splatting in Unreal Engine (July 2026): real-time splat relighting and self-shadowing, with drag-and-drop 4DGS import and synchronised playback and transitions.
- **Atlux** — Unreal Engine to 4D Gaussian Splat conversion (July 2026), including UE 5.8 support.

- **Notch 2026.2** — adds native Gaussian-splat import and rendering to Notch (July 2026).
- **InfiniSplat** (SIGGRAPH Asia 2026) — turns a single image into a navigable 3DGS scene.
- **Insta360 X6** — 360 capture camera whose new Spatial Capture mode (August 2026) turns a recorded place into a navigable in-app 3D model.
- **PlayCanvas SOG / SplatTransform v2.0.0** — WebP-equivalent compression for Gaussian splats; SplatTransform v2 ships automatic high-quality collisions for splats.
- **Cesium / CesiumJS / Cesium for Unreal** — geospatial 3D platform; Cesium agentic workflows let natural-language commands interact with 3D geospatial data; supports Gaussian splats natively.
- **SPAG-4D** — 360° photos to 3D Gaussian splat conversion.
- **Animated Gaussian Splats** (4DV.ai) — 4D Gaussian-splat animation.
- **Open3Dmap** — crowdsourced 3D mapping with Gaussian splats.
- **Hyperscape / Hyperscape Capture** (Meta) — Gaussian-splat capture on Quest.
- **Apple Vision Pro Personas** — Gaussian-splatting consumer feature.
- **Gaussian Splats in Apple Maps** — announced at WWDC 2026: the Flyover feature rebuilt on Gaussian splatting; the splat pipeline's arrival in mass-market consumer software.
- **Common Sense Machines** — converts 2D images into 3D digital assets (Google acquisition).
- **AI-Enhanced LiDAR** — generative-AI-augmented LiDAR capture pipeline (June 2026) for scanning and reconstruction work in spatial production.
- **NVIDIA Omniverse Fixer** — rendering-artefact removal for Gaussians.
- **DecartAI / Decart LSD v2 / Decart Lucy 2.0 / Decart Mirage** — real-time world transformation by voice; LSD v2 is real-time video-to-video; Lucy 2.0 is realtime world editing at 1080p/30fps; Mirage (June 2026) restyles a shared screen live — claymation, oil painting, anime, or your own prompt.
- **VoxeloAI** — 3D creation for e-commerce.
- **Mosaic / Mosaic 3D** — 3D reconstruction.
- **Depth Anything 3** (ByteDance) — visual-space recovery from any view.
- **SAM 3 / SAM 3D / SAM 3.1** (Meta) — object detection and 3D segmentation; SAM 3.1 is 7× faster, tracking 128 objects on a single H100.

- **Seen2Scene** — realistic 3D scene completion with visibility-guided flow.
- **PixARMesh** — reconstructs full indoor scenes from a single photo into lightweight meshes.
- **Confidence-Based Mesh Extraction from 3D Gaussians** — algorithm for high-quality mesh extraction from Gaussian splats.
- **Rhizomatics Sword Tip Visualization System** — fencing tracking at 60 fps.
- **Mixar** — AI-native Blender-fork editor positioned as “Cursor for 3D artists”.
- **Sprite Booth** — pixel-character animation tool (ComfyUI API required).
- **AutoSprite** — character-to-animated-sprite-sheet conversion.

Voice, avatar, digital human

- **Synthesia** — AI avatar platform; \$4B valuation; rejected \$3B Adobe offer.
- **Heygen / Heygen Video Agent / Heygen Motion Designer / LiveAvatar / Heygen Elements / Heygen CLI / HyperFrames / Heygen Video Podcast** — AI-avatar and end-to-end video generation; LiveAvatar is hyper-realistic real-time interactive; Elements builds scenes from reusable components; CLI ships videos from the command line; **HyperFrames** (June 2026) is an open-source framework for turning HTML, CSS, media and seekable animations into deterministic MP4 videos — bringing deterministic video export to web-stack creative workflows; **Video Podcast** (July 2026) turns any doc, link or idea into a two-host video show with studio scenes, multi-cam cuts and B-roll.
- **Hedra / Hedra Audio Tags / Hedra Elements** — digital-human creation; Audio Tags assign precise emotions to audio.
- **Live Avatar** (Alibaba) — streaming real-time audio-driven avatar generation with infinite length.
- **Avatar Forcing** — real-time interactive head-avatar generation for natural conversation.
- **PersonaLive** — real-time expressive portrait animation.
- **LeapTalk** — real-time audio-driven talking-head generation (August 2026).
- **LongCat-Video-Avatar 1.5** (Meituan) — production-ready, open-source framework for audio-driven avatar generation.
- **Weclone** — digital avatar built from a user’s chat history.
- **Particle6 / Tilly Norwood** — AI-performer studio and synthetic actress.
- **Xania Monet** — AI music artist; Billboard chart entries; \$3M Hallwood Media deal.

- **Bleeding Verse** — AI band; Hallwood Media signing.
- **Sienna Rose** — anonymous AI artist with millions of Spotify streams.
- **Breaking Rust** — AI country artist; #1 country digital song sales.
- **Trilok** — Indian AI band.
- **Copresence / ConvAI / Inworld** — avatar and NPC-character technology for Unreal Engine, games and interactive media.
- **ROXi** — AI-generated TV presenters.
- **Facelooker** — interactive headshot generation.
- **FellinAI** — AI film director.
- **ego AI** — in-game character platform.
- **Character.AI** — play characters in your favourite books; Ovi open-source video with speech sync.

Agent platforms and orchestration

- **OpenAI AgentKit / Agent Builder / ChatKit / Connector Registry / Eval Framework** — the developer-facing agent platform underneath most third-party agentic creative tools.
- **Anthropic Claude Apps / Claude Skills / Claude Code / Claude Design / Claude for Legal / Anthropic Academy** — interactive Claude in workplace tools; the Skills framework powers Sony's 49-agent / 72-skill game-development stack; Claude Design ships prototypes/slides/one-pagers; Anthropic Academy is free Claude Code training.
- **Claude Code Game Studios** — 49-agent, 72-skill coordinated AI game-development team.
- **Gemini API Agents / Google AntigraVity / Opal / Fabula / Google Stitch / Google Skills for AI Agents** (Google) — agentic capability surface across Google's stack; AntigraVity is the agentic development platform; Opal is the no-code mini-app builder; Fabula is the interactive AI writing tool; Stitch generates mobile/web UI; the official Skills repository is Google's cross-vendor answer to Anthropic's Skills framework.
- **Heygen Video Agent** — end-to-end video-assembly agent.
- **Adobe CX Enterprise / GenStudio / Adobe Film & TV Fund / Adobe Ignite Day** — agentic creative intelligence across the full content lifecycle; GenStudio is the brand-intelligence system for personalised content at scale.
- **NVIDIA + Google Cloud / Avid + Google Cloud / Speechmatics in Adobe Premiere** — agentic-creative infrastructure partnerships; on-device speech-to-text inside Premiere.
- **n8n** — workflow automation with AI.

- **ComfyUI / ComfyUI Cloud / ComfyHub / ComfyUI App Mode / ComfyUI Simple Mode** — node-based AI workflow editor; \$500M valuation by May 2026; App Mode turns workflows into shareable no-install apps; Simple Mode lowers the complexity threshold.
- **Replicate** — model marketplace.
- **Hugging Face** — open-source model hub and agentic platform; partnered with Google Cloud and Meta.
- **Meta OpenEnv** — open-source agentic-development environment.
- **fal / FAL MCP** — generative-model deployment marketplace; FAL MCP exposes 1,000+ generative models to Claude/Cursor.
- **Unity Official MCP / Unity MCP** — Unity Editor control via MCP protocol; from July 2026 the Unity MCP server is free, with no limits on concurrency.
- **Blender MCP / Blender Buddy / OpenClaw + Blender MCP** — Claude-to-Blender control for 3D workflows.
- **VibeComfy** — open-source tools letting agents understand, build and run ComfyUI workflows with Claude.
- **MooshieUI** — beginner-friendly ComfyUI desktop frontend.
- **Lenny** (Maroofy) — AI agent for live-music event organisers.
- **AdsGency** — autonomous paid-marketing agent platform; \$12M seed.
- **Invideo AI / Invideo AI Agent** — full video-ad generation from prompt; built on Gemini.
- **Pomelli** (Google Labs) — AI marketing agent for small business.
- **Inception Point AI** — podcast hosting with AI hosts (3,000+ episodes/week).
- **Slapshot** — AI camera tracking.
- **Photoroom** — AI photobooth.
- **Jabali.ai** — AI game-development platform (Sony-backed).
- **Fifth Door** — AI-powered game creation; \$20M raise.
- **Atlas AI Agents** — game-production pipeline agents.
- **Ad Concepter** (Runway) — ad-concept exploration agent.
- **AE GPT** — AI assistant for After Effects (expressions, scripts).
- **Career-Ops** — AI job-search pipeline built with Claude Code.
- **Auto-Dream** (MyClawAI) — personal AI assistant wrapper.
- **Clicky** — AI screen-aware tutor/buddy interface.
- **MemPalace** — high-scoring AI memory system.
- **TikTok AI Agents for Ads** — TikTok’s automated ad-creation agents.
- **Amazon Creative Agent** — agentic AI generating professional-quality ads.
- **Playad** — AI marketing agents for ad creative.
- **Moltbook** (Meta acquisition) — AI agent social network.

- **Buzzy Canvas** — “your AI co-director”; a canvas-shaped direction agent for creative work.
- **Grok Voice Agent Builder** (xAI) — no-code platform for creating human-like voice agents with Grok Voice.

Legacy creative software, AI-augmented

- **Adobe Creative Cloud / Adobe Creative Agent** — Photoshop, Premiere Pro, After Effects, Illustrator, Express, Acrobat, plus the Express AI Assistant, the Premiere Object Mask, the Photoshop generative-fill, selection and Rotate Object tools, the Photoshop AI Assistant public beta, and Illustrator Turntable; the Creative Agent and AI Assistant went into beta across Premiere, Photoshop, Illustrator, InDesign and Frame.io in June 2026.
- **Adobe Sneaks 2025–26** — Adobe’s research-preview portfolio shown at MAX 2025 and Summit 2026:
 - **Project Scene It** — image-to-3D and 3D-to-image with object tagging.
 - **Project Surface Swap** — AI-powered texture recognition and swap.
 - **Project Turn Style** — edit and transform 2D objects as if they were 3D.
 - **Project Trace Erase** — removes objects, shadows, reflections and distortions.
 - **Project New Depths** — edit depth like brightness.
 - **Project Frame Forward** — apply changes across video from a single frame plus text.
 - **Project Motion Map** — bring static vector graphics to life.
 - **Project Sound Stager** — analyses video visuals and generates synchronised soundscapes.
 - **Project Clean Take** — corrects mispronunciations and isolates voices.
 - **Project Light Touch** — spatial lighting mode for relighting photos.
 - **Project Graph** — node-based workflow editor in the ComfyUI / Weavy lineage.
 - **Corrective AI** — changes the emotional register of voice-overs.
 - **Edit-by-Track** (Adobe Research) — generative video motion editing with 3D point tracks.
- **Adobe Acrobat** — AI converting PDFs into podcasts.
- **Unreal Engine 5 / UE5 AI Assistant / UE5 AI Motion Plugin / Unreal to Gaussian Splat / Prompt-to-Player / Unreal MCP** (Epic) — official AI assistant; camera-only motion capture; UE5 scene-to-splat conversion; Prompt-to-Player generates 3D characters and auto-rigs/animates them; Unreal MCP and the AI rendering integrations were announced at the State of Unreal 2026.

- **Unity / Unity AI Open Beta / Unity AI Council / Unity AI Tools Suite / Unity Prompt-a-Game** — in-editor AI suite for the full games pipeline; Prompt-a-Game demoed at GDC March 2026; the Unity AI tutorial series (June 2026) covers the AI Assistant and the official Unity MCP.
- **Autodesk Maya / Autodesk Flow Studio / Autodesk Wonder 3D / Autodesk Flow Studio Rigging** — motion capture, 3D animation, generative 3D and neural rigging.
- **Foundry / Nuke / Nano Banana × Nuke** — AI-augmented VFX; Weta FX / AWS collaboration; the Nuke node connecting Nano Banana through fal’s API.
- **SideFX Houdini** — procedural modelling with AI integration.
- **Blender / Blender MCP / Blender Buddy** — open-source 3D; Anthropic Foundation patronage; MCP enables Claude control; NVIDIA DLSS 4.5 Ray Reconstruction added to Blender 5.3 as a viewport denoiser.
- **Vlo** — free, local, open-source video editor with ComfyUI-backed generative AI features; the open-source equivalent of the AI-native NLE.
- **Motionfly Editor** — AI-native video editor positioned as “Cursor for video editing” — bringing the code-editor / AI-pair model into the video-editing workflow.
- **DaVinci Resolve** (Blackmagic) — colour and editorial with AI features.
- **Avid Media Composer / Avid + Google Cloud** — professional editorial; Pro Tools agentic AI in the Google Cloud partnership.
- **Pro Tools** — audio editorial with AI features; Claude integration in Ableton.
- **Logic Pro / Apple Logic Pro AI** — music production with AI; synth player and music-theory expert.
- **Ableton Live / Claude in Ableton / Jamu / Melosurf** — music production with Claude assistance, Jamu co-producer agent and Melosurf voice control.
- **GarageBand** — consumer music creation.
- **CapCut / Dreamina / Snapchat AI Clips in Lens Studio** (ByteDance / Snapchat) — mobile video editing with AI; Snapchat AI Clips in Lens Studio.
- **Final Cut Pro** — editorial with FX Factory plugins.
- **Keyframe** — AI models inside After Effects.
- **Tether** — AI animation inside After Effects.
- **Deep Pan** — animated stills in Premiere Pro.
- **FX Factory** — AI plugin suite for Premiere and Final Cut.
- **VEED Transitions** — AI transition generation.
- **Electric Sheep** — web-based projection mapping and AI visual orchestration.

AI-native creative studios and apps

- **Imaginae Studios** (Fremantle) — AI-native studio; *Art Awakens* project under Google AI Futures Fund.
- **Wonder Studios** — AI-native film studio; \$12M raise; Wonder Film Festival.
- **Asteria** (Natasha Lyonne / Lightstorm / Topaz Video) — AI animated short *All Heart*; hybrid AI workflows with Topaz.
- **Promise** (Google-backed) — GenAI filmmaking and VFX for legacy media.
- **Obsidian Studio** (Imagine Entertainment / Ron Howard, Brian Grazer) — AI production partner.
- **Goldfinch enGEN3** — AI cinematic universe platform.
- **Chapter41** (Beta Film / Munich) — AI startup.
- **Kartel** (Kevin Reilly) — AI startup; ex-HBO leadership.
- **Holywater** (Fox Entertainment investment) — AI microdramas.
- **CenterStage** (David Boies / Zack Schiller) — interactive AI entertainment.
- **Superlunar + Eden Studios** — *Little Plastics* dystopian AI film.
- **Particle6** (Eline Van der Velden) — AI-performers studio.
- **Imago Pictures** — 3D + ComfyUI workflows.
- **LKDN AI Production** — node-based AI filmmaking.
- **Deep Fusion Films** (Cardiff) — AI documentary production; acquired by John Gore Studios.
- **Palo Creators** (MrBeast alumni) — AI for viral video creation.
- **Refik Anadol Studio Dataland** — first museum of AI art.
- **Animaj** — Google AI Futures Fund partner for kids' content.
- **Gossip Goblin** — AI filmmaking studio.
- **Critterz** (Vertigo + Federation) — AI-assisted animated-feature production.
- **Filmology Labs** — \$250M studio for vertical micro-dramas and AI-driven media formats.
- **escapeAI** (John Gaeta) — streaming apps for AI-generated content across Roku, Fire TV, Samsung and LG.
- **Primordial Soup** (Darren Aronofsky) — AI production studio.
- **Toonstar** — AI animation house.
- **Utopai Studios** — AI film-and-TV studio; co-producing Silver Bear winner Hyo-Joo Yang's *Half Moon* and, with Huace, developing the fully AI-generated series *Journey to the West*.
- **CinemersiveAbout Labs / Cinemersive Labs** — UK machine-learning and computer-vision company; acquired by Sony.
- **My SMASH Media** — AI film startup with Allison Gardner (ex-Glasgow Film Festival).

- **Framestore AI Platform / Framestore Futon** — ML and GenAI in the VFX pipeline.
- **Mito AI** — \$4.5M-raised platform empowering video professionals.
- **Filmera** — early-stage agentic AI filmmaking platform and workspace taking creators through the full film pipeline — idea, script, shot planning, generation and editing — in one place (August 2026).
- **Filmax DinoGames** — AI-driven 3D animated film project.
- **Abundantia Entertainment + InVideo** — AI film studio launched in India.
- **Luma Production Studio + Wonder Project** — feature collaboration.
- **Lionsgate Chief AI Officer (Kathleen Grace)** — first major-studio CAIO hire.
- **CreateAI** — AI animation production.
- **Netflix Ben Affleck AI Studio Acquisition** — AI filmmaking company acquisition.
- **Holywater** (Fox investment) — AI microdrama studio (also above).
- **DreamLab AI Collective** (the studio that produced this book) — AI-augmented creative practice across film, games and immersive.

Games-development AI

- **NitroGen** (NVIDIA + Stanford) — plays-any-game model trained on 40,000 hours across 1,000+ games.
- **SIMA 2** (Google DeepMind) — agent that plays, reasons and learns in 3D virtual worlds.
- **CHORD** (Ubisoft La Forge) — open-sourced end-to-end PBR material generation.
- **Ubisoft Teammates** — voice AI for in-game team communication.
- **YouTube Playables Builder** (Gemini 3) — text-to-game prompt-to-playable web app.
- **Roblox AI Tools / Roblox AI Assistant / Roblox Reality / Metain** — creator-facing AI for game development; Reality is Roblox's productised AI assistant; Metain is the prompt-to-Roblox-Studio interface. In June 2026, Roblox acquired **Morpheus AI Inc.**, an AI lab building video world models using *Self Forcing* (converting offline video models into fast, autoregressive, interactive engines), to accelerate the Roblox Reality video-model vision.
- **Epic xAI Studio** — Elon-Musk-affiliated game-dev AI.
- **General Intuition** — spatial-reasoning research lab.
- **ReBlink ARBO** — AI-powered strategy battler.
- **Halo Studios** — AI woven into the production pipeline.
- **Dabgg** — AI gaming with interactive platform.

- **RosebudAI / Spawn / LudoAI** — vibe-coding games and interactive apps from prompts; Spawn is the “games made with words” community platform; LudoAI is the ideation and planning scaffold.
- **Verse8** — AI-driven game-creation platform.
- **Sett** — AI agents for game marketing.
- **Meta Horizon Studio AI Assistant** — Horizon Worlds creator AI.
- **Gambo** — vibe-coding game agent.
- **Project Motoko** (Razer) — AI-powered gaming headset.
- **Bethesda AI Toolset** — Todd Howard’s game-dev AI integration.
- **StarBerry Games / Merge Mayor** — AI-all-in game studio.
- **Astrocade** — interactive entertainment platform; \$56M Series B.
- **GameByte** — AI-powered game creation platform; \$1M raise.
- **Sorcerox Game Creation Suite** — prompt-driven game creation toolkit (June 2026); positioned as a “*build any game you can imagine*” general-purpose creation surface.
- **Studio Atelico Bobium Brawlers** — AI-based iOS game with pro-human approach.
- **Bitmagic** — AI-created version of classic Civilization.
- **Helika AI Publishing Engine** — AI publishing engine for game studios.
- **Xbox Game Studios “git gud” AI** — Xbox patent for AI tech that plays games for you.
- **Gaming Copilot** (Xbox) — console AI copilot.
- **NVIDIA DLSS 5** — gaming upscaling.
- **PS5 Pro PSSR** — PlayStation upscaling.
- **Gizmo** (Meta acquisition) — vibe-coding app for creating mini-games.
- **Meta Pocket** (Meta) — vibe-coded gaming app that turns descriptions into little AI games and experiences, with every play feeding back as training data (July 2026).
- **Tesana** — prompt-driven “make any game you can imagine” creation platform (July 2026).
- **Voyage** (Latitude / AI Dungeon) — AI RPG platform.
- **DoomGen** — prompt-driven DOOM mod prototyping app.
- **Sony AI tools for PlayStation** — animation support, QA automation, asset generation and performance-capture-to-facial-models pipelines.
- **PROWL** (Odyssey) — RL agents for game-world model testing.
- **MagicDawn** (Tencent) — AI lighting tech.
- **ReActor** (Disney) — RL motion-transfer method.
- **Origin Lab** — game-world AI training data; \$8M raise.

- **Daughter of the Inner Stars** — Unreal Engine 5 AI-driven game.
- **Genesis AI** — robot cooking and piano playing systems (demonstrating embodied agents).
- **Quilty** — AI platform for script development and assessment.
- **Higgsfield Games** — create and deploy multiplayer games from a single prompt.
- **NVIDIA ACE Game Agent SDK** — with new Unreal Engine 5 plugins, for building responsive, on-device AI companions.
- **The Sandbox Studio** — from idea to live multiplayer game in hours.
- **ClarityQ** — game-native AI analytics platform for mobile game studios.
- **Welevel SolidRiver** — AI-based city-building survival game; \$8.5M in funding.
- **Krafton PUBG AI Ally** — beta AI teammate in the new PUBG arcade mode.

Marketing and advertising AI

- **WPP Open Pro** — agency-wide AI marketing platform.
- **Sightly** — AI advertising and brand performance.
- **Pomelli** (Google Labs) — small-business marketing agent.
- **AdCreative.ai** — mobile pro ad creation.
- **Epiminds** — AI marketing team builder.
- **Octave AI** — AI podcast/radio ad creation.
- **AdsGency** — autonomous paid-marketing agent.
- **Aimy** — AI TV advertising on the Comcast API.
- **Brand Networks** — AI television advertising.
- **Fullthrottle.ai** — automotive-focused DSP.
- **Channel 4 AI ads** — AI-driven TV-advertising tool for SMEs.
- **Sky Studio Challenge** — interactive AI ad creation via Starlink.
- **Google Demand Gen** — AI-powered image tools for the ad platform.
- **Amazon DSP** — AI one-stop-shop advertising infrastructure.
- **Amazon Creative Agent** — agentic ad creator.
- **Playad** — AI marketing agents.
- **Runway Ad Concepter App** — ad-concept and composition exploration.
- **TikTok AI Agents for Ads** — TikTok ad-creation agents.
- **Instagram Edits** — AI video generation for advertisers.
- **ElevenLabs Ads Engine** — localise existing ads across 50+ languages and push the finished creatives back to the ad platform.
- **WPP Hex** — WPP's studio putting “a different type of ad talent” — AI-native creatives — to the test.

Open-source ecosystem and infrastructure

- **ComfyUI** — node-based AI workflow editor; the operating system of the open-source creative AI ecosystem; \$17M raise (October 2025), \$500M valuation (May 2026).
- **Hugging Face** — open-source model hub; partnerships with Google Cloud and Meta.
- **Meta OpenEnv** — open-source agentic-development environment.
- **PlayCanvas SOG / SplatTransform** — open Gaussian-splat compression and conversion.
- **Civitai / Civision** — community-models and LoRA platforms.
- **LoRA / PEFT / DiffSynth-Studio / Qwen-Image-i2L / Llama Factory / AI Toolkit / LTX-2 Trainer** — fine-tuning frameworks and training pipelines; DiffSynth turns a single image into a custom LoRA; Llama Factory enables zero-code fine-tuning of 100+ models; AI Toolkit covers LTX-2.3 fine-tuning; LTX-2 Trainer is the fal-hosted variant.
- **PyTorch / TensorFlow / Ollama / LLaMA.cpp / MuJoCo** — the open infrastructure layer.
- **Korin AI** — Africa-trained, Africa-built foundation model.
- **DecartAI** — real-time world transformation.
- **Replicate** — model hosting and API marketplace.
- **fal / FAL MCP** — generative-model API platform.
- **llmfit** — hardware scanner for local-model compatibility.
- **Unsloth Desktop** — the first desktop app to run and train models locally (August 2026).
- **Open Vision Agents** — toolkit for building agents that watch, listen and understand video.
- **NotebookLM** — Google’s research-and-content-generation surface; cinematic video overviews launched March 2026.
- **OpenAI Gym / Gym Retro / ARC** — research benchmarks (referenced in the games-AI literature).
- **Awesome Stable Audio** (Stability AI) — curated list of integrations and extensions for Stable Audio.
- **Open Generative AI** — open-source alternative to the commercial AI video platforms.
- **Roboflow RF-DETR / RF-DETR Keypoints** — real-time transformer architecture for object detection and instance segmentation; the Keypoints release (June 2026) added state-of-the-art tracking.
- **Roboflow supervision** — the open-source computer-vision utility library; 40,000 GitHub stars, powering 6,500+ open-source CV projects.

- **NVIDIA TAO Skill Bank** — portable agent skills for training, evaluating and running inference on NVIDIA TAO models.
- **Open Source Face Recognition SDK** (Faceplugin) — open-source face-recognition toolkit.

ComfyUI ecosystem — nodes, extensions and workflows

- **Mesh2Motion** (ComfyUI) — mesh-to-motion conversion node.
- **NanoBanana Pro LoRA Dataset Generator** — creates training datasets for image-editing models in minutes.
- **ComfyUI Music Tools** — professional-grade music-processing node pack.
- **ComfyUI HY-Motion1 plugin** — Tencent HY-Motion implementation.
- **ComfyUI Audio Reactive Node Pack** (VisualFrisson) — spatial audio-reactive nodes.
- **ComfyUI-BlendPack** — nodes for video transitions.
- **KJNodes** — use audio as input for LTX-2 inside ComfyUI.
- **WhatDreamsCost ComfyUI Nodes** — LTX sequencer, keyframer.
- **Luminance Keyer Node** — luminance-based keying.
- **Camera Tracking & Body Pose tools (ComfyUI)** — work-in-progress camera and pose tracking.
- **Multiple Angle Camera Control Node** — multi-angle camera control.
- **FeedbackSampler** (Deforum-inspired) — iterative feedback sampling.
- **Day-to-Night with Qwen Edit LoRA** — day/night transformation workflow.
- **Sora 2 API Nodes** — Sora 2 access from ComfyUI.
- **QWEN-AI Compositing** — Qwen-driven compositing workflow.
- **Street View URL Parser** — Google Street View ingestion.
- **Minim** — ComfyUI model-linker extension.
- **Hunyuan 3D 3.0** (ComfyUI partner nodes) — HY3D 3.0 inside ComfyUI.
- **Kling 3.0 ComfyUI node** — Kling 3.0 partner node.
- **Kling Video 2.6 Motion Control** (ComfyUI) — Kling motion control inside ComfyUI.
- **Kling Motion Control 3.0** (ComfyUI) — the v3 motion control variant.
- **Seedance** (ComfyUI partner node) — Seedance inside ComfyUI; Seedance 2.0 Mini and 4K available from July 2026.
- **LTX-2** (ComfyUI native support) — audio-video model natively supported.
- **LTX-2 ComfyUI Audio-to-Video** — audio-driven video inside ComfyUI.
- **ACE-Step 1.5** (ComfyUI) — full songs in under 10 seconds on <4 GB VRAM.
- **ElevenLabs ComfyUI Partner Nodes** — ElevenLabs inside ComfyUI.
- **Luma Uni-1 ComfyUI Partner Node** — Uni-1 inside ComfyUI.

- **Meshy 6** (ComfyUI native) — Meshy 6 inside ComfyUI.
- **Grok Imagine** (ComfyUI) — xAI image generation inside ComfyUI.
- **MatAnyone2** (ComfyUI) — high-fidelity video matting inside ComfyUI.
- **Tripo 3.1** (ComfyUI partner nodes) — Tripo 3.1 inside ComfyUI.
- **NVIDIA RTX Video Super Resolution Node** — RTX upscaling node.
- **Depth Anything Custom Node** — Depth Anything inside ComfyUI.
- **SAM 3 / SAM 3D ComfyUI integration** — Meta SAM 3 / 3D inside ComfyUI.
- **Sora 2 Full Loop Workflow** — end-to-end Sora 2 pipeline.
- **Krea Nodes** — Krea inside ComfyUI.
- **Gaussian Splats + Qwen Edit workflow** — explore 2D images and generate new camera views.
- **ComfyUI App Mode / ComfyHub** — turn workflows into shareable no-install apps.
- **ComfyUI Simple Mode** — share and iterate complex workflows more easily.
- **ComfyUI Cloud (private models)** — private model hosting inside ComfyUI Cloud.
- **ComfyUI Action Director** — interactive 3D viewport for ControlNet.
- **Complete Style Transfer Handbook (ComfyUI)** — style-transfer guide.
- **VibeComfy** — agents that understand, build and run ComfyUI workflows with Claude.
- **MooshieUI** — beginner-friendly ComfyUI desktop frontend.
- **Creative Control Using ComfyUI on NVIDIA RTX** — NVIDIA’s official ComfyUI/RTX tutorial series.
- **Comfy MCP** — the official Model Context Protocol server, in public beta from July 2026; the first MCP built for production pipelines — agents build, edit and run workflows in plain language, with shareable, reproducible workflow URLs, running on Comfy Cloud with no local GPU required.
- **Comfy Desktop** — one app for every Comfy; the desktop consolidation of the ecosystem.
- **ComfyUI Utility Passes workflow** — generates production-ready VFX utility passes directly from footage.
- **ComfyUI Face Swap workflow** — de-aging and face-replacement demo workflow.
- **ComfyUI + Seedance 2.0 Storyboard workflow** — turn one idea into a complete storyboard, then animate it in Seedance 2.0 from a reference image.
- **ComfyUI Three.js Motion Reference node** — custom node generating a 3D motion reference with Three.js.

- **ComfyUI for VFX (ActionVFX) / Advanced ComfyUI (fxphd)** — the VFX-pipeline training layer: ActionVFX’s tutorial collection and fxphd’s *Generative Video Production for VFX* course.

LoRAs, fine-tuning and training

- **QwenEdit-2509 Light Transfer LoRA** — simple, powerful image relighting via reference.
- **Qwen-Edit Relighting LoRA** — atmospheric relighting control.
- **Qwen-Image-Edit-2511 Multiple Angles LoRA** — multi-angle perspective generation.
- **Qwen-Image-Edit-2511 AnyPose** — pose transfer from a reference image.
- **QwenEdit-2511 Anything2Real LoRA** — realistic-image transformation.
- **Qwen Edit LoRA Object Removal** — bounding-box-precision edits.
- **Qwen-Image-Edit 2511 Gaussian Splash 3D Camera Motion** — 3D Gaussian-splat motion control.
- **Qwen 2511 Time Travel Workflow** — temporal image effects.
- **NextScene (Qwen Image LoRA)** — cinematic image sequences with natural progression.
- **Best Face Swap (Flux 2 LoRA)** — specialised face-swap.
- **FLUX.2 [klein] LoRAs (fal)** — open-source: Outpaint, Zoom, Object Removal, Background Removal.
- **Krea AI LoRA Trainers** — train LoRAs for Qwen-2512 and Z-Image inside krea.ai.
- **LTX-2.3 Character LoRA Training (AI Toolkit)** — character LoRA training tutorial.
- **Gaussian Splats Repair LoRA** — Klein-9b LoRA for repairing 3D views and geometry.
- **LTX 2.3 Colorizer (LoRA)** — black-and-white footage colourisation.
- **Music Finetunes in ElevenCreative** — stylistically consistent vocal and instrument generation.
- **LTX 2.3 Audio Reactive LoRA (fal)** — open-source LoRA built for music-driven video generation.
- **LTX 2.3 LoRA Trainers (fal)** — train custom LTX 2.3 LoRAs across 23 trainer endpoints.

IP licensing, governance and behavioural-licence platforms

The contract-and-governance layer that emerged through 2025–26. The deep analytical treatment of how these models work — the three Cs of consent, credit,

compensation; the five-model comparison; the corrected Disney–OpenAI record; the transaction-cost reframe — sits in DreamLab’s published research.

- **Hasbro Sixth Wall / CharacterOS** — in-house AI studio launched 3 June 2026 to license the *behaviour* of Hasbro characters (Optimus Prime, Megatron, Cobra Commander, the *Clue* cast) under a runtime governance layer that encodes canon, voice, personality, lore and safety limits; B2B programmatic licence with a talent-participation pay-out structure for original performers; initial 12-character catalogue; 13+ /enterprise focus.
- **Verified by Spotify / Spotify–UMG fan-remix product** — Spotify’s *Verified by Spotify* badge and the opt-in catalogue infrastructure underneath the May 2026 Spotify–Universal Music Group landmark licensing agreement for fan-made AI covers and remixes; *consent / credit / compensation* framework; paid Premium add-on; new royalty stream to artists and songwriters.
- **Disney–OpenAI licensing framework (announced 11 December 2025; never executed)** — the \$1bn licensing announcement that became the reference contract architecture for legacy-IP × frontier-AI deals despite the deal itself collapsing when OpenAI shut down Sora on 24 March 2026. IP-not-training, talent-voice-and-likeness excluded, jointly-governed steering committee, equity-and-warrants structure. The framework survived the platform’s death.
- **Elf.tech / GrimesAI** (CreateSafe, Triniti API) — Grimes’s open-voice platform; upload your vocal, get back a synthesised GrimesAI performance, release commercially with a ~50% master-royalty split; analogised by her team to fan-fiction and fan-art cultures; ~15,000 tracks generated under the licence.
- **Holly+ / Holly Herndon DAO** — public digital twin of the artist’s voice wrapped in collective DAO governance for commercial-use decisions and the distribution of proceeds; paired with the *Have I Been Trained?* (Spawning) consent-layer tooling.
- **ElevenLabs Heritage Voices / Iconic Marketplace (with CMG Worldwide and partners)** — heritage-voice licensing pipeline confined to a single text-to-speech Reader app, walled off from the general voice-cloning database; voices include Judy Garland, James Dean, Burt Reynolds, Laurence Olivier, the Stan Lee estate and Michael Caine.
- **Meta Celebrity Chatbots (2023)** — the closed-static likeness-rental anti-pattern often cited for contrast: high flat fee (reportedly up to ~\$5m per celebrity over two years) for platform-owned, frozen personas with no participatory layer for users or the wider creator community.
- **Cannes AI Disclosure Standard** — industry-coordination labelling standard for production-side AI use; the inter-industry analogue of the per-platform governance layers above.

- **Human Consent Registry** — identity-protection registry backed by Cate Blanchett, Steven Soderbergh and Nikki Hexum, framing protection from AI identity theft as “a human right.”

Provenance, watermarking and detection

- **C2PA** (Content Authenticity Initiative / Adobe-led) — cryptographic provenance metadata standard.
- **SynthID / SynthID Verification** (Google DeepMind) — synthetic-content watermark; deployed across Veo, Lyria and Imagen; verifiable inside Gemini; from June 2026, ElevenLabs-generated audio is detectable via SynthID under a DeepMind partnership.
- **YouTube AI Detection / YouTube AI Deepfake Detection** — automated AI-content detection; *Tiny Grandma* false-positive case study.
- **Deezer AI Music Detection** — identifies up to 75,000 fully AI-generated uploads per day; from June 2026, offered as a detector to *other* streaming services — the in-house tool becoming licensable infrastructure.
- **Spotify AI Transparency Beta** — voluntary disclosure feature.
- **Beeble** — detection and watermarking tools.
- **Cloudflare AI Bot Classification** — public-web infrastructure tracking AI crawlers.
- **Human Provenance AI Disclosure Standard (Cannes)** — industry-coordination labelling standard.
- **RTM Audio UAI** — AI-music detector that issues a signed certificate with every verdict.
- **Played By Humans** — verification badge for “human artistry in an AI age.”
- **Midnight Labs** — AI copyright-protection firm; Sony made a multi-million-dollar investment in June 2026.

Consumer surfaces and distribution platforms

- **Sora iOS app** (OpenAI) — TikTok-style AI video remix; 1M downloads in 5 days.
- **CapCut / Dreamina** (ByteDance) — consumer video editing + AI generation.
- **Gemini app / Gemini in Chrome** (Google) — Google’s consumer AI surface; **Gemini in Chrome** launched to UK desktop users in July 2026.
- **ChatGPT app** (OpenAI) — mobile ChatGPT; Shazam integration.
- **Perplexity AI** — Samsung Smart TV AI assistant integration; consumer AI search.
- **Grok / Grok Imagine** — xAI consumer surfaces.
- **TikTok / Reels / Shorts** — AI-content distribution surfaces.

- **YouTube Shorts / YouTube Create** — destinations for Veo 3.1 Ingredients-to-Video.
- **Bandcamp** — banned AI music outright, January 2026.
- **Deezer** — published the 44%/3% data; built and licensed AI-music detection.
- **Spotify** — declined an AI-music filter; voluntary disclosure beta; *Personal Podcasts* agent feature.
- **YouTube Music** — labelling, the AI Detection Tool, viewership of AI content.
- **Steam** — AI-content labelling controversy; *Clair Obscur* Game of the Year stripping.
- **Sweden's Official Music Chart** — banned AI-generated entries.
- **San Diego Comic-Con** — banned AI art at the 2026 event.
- **Roblox** — creator-tools AI with continued controversy.
- **Fortnite** — Chapter 8 AI-art controversy; Disney partnership.
- **Netflix** — AI in production (de-aging, plates, retention engine, Ben Affleck studio acquisition).
- **Disney+** — user-generated AI content features.
- **Luna** (Amazon) — AI-powered Snoop Dogg game.
- **Meta AI Creator Assistant** — AI creator assistant on Facebook, followed by an AI companion app for creators (June 2026).

Studios, programmes, festivals and institutional infrastructure

- **Sundance AI Literacy Initiative** — Google-funded \$2M creator-training programme; Sundance Collab learning, Story Forum, Ignite Day.
- **UCL / RCA / Brandtech Centre for Creative AI** — UK research hub.
- **AIMICI Training Program** — ScreenSkills-accredited AI training.
- **AI FilmFest Japan** — annual AI-film festival.
- **Dubai AI Film Award** — \$1M prize won by Tunisia's *Lily*.
- **Cannes AI Disclosure Standard / Cannes Lions** — industry coordination for AI labelling; Luma Dream Brief \$1M prize.
- **Future Vision Film Competition** — Google / XPRIZE / Range Media collaboration.
- **Runway AI Festival** — film and creative-discipline awards.
- **Wonder Film Festival** — AI-native film festival.
- **Academy of Motion Picture Arts and Sciences** — “you must be human to win” rule.
- **Television Academy / Emmys** — AI guidance for submissions.
- **SAG-AFTRA** — contract negotiations and the Tilly Tax.
- **Equity (UK)** — strike ballot; 99% vote in favour.

- **PRS for Music AI Survey 2026** — UK music-creator sentiment data.
- **GEMA** — German rights society; won against OpenAI.
- **Splice + UMG** — sample library; AI-music tools partnership.
- **Stability AI + Universal Music / Stability AI + Warner Music** — alliance and deal on the AI-music side.
- **Vizrt Viz One 8.1** — broadcast AI features.
- **Tesseract** — AI platform at MIPCOM.
- **Sphere + Google Cloud** — AI technology behind the *Wizard of Oz* experience.
- **Anthropic Joins Blender Development Fund** — patronage deal, May 2026.
- **Sony AI ICASSP papers** — music-understanding and generative-audio research, April 2026.
- **Google Flow Music + Believe Partnership** — AI music distribution partnership.
- **Google Flow Resident Filmmaker Programme** — six-week Google Flow residency recruiting “the next wave of creatives across all mediums” (July 2026).
- **VERSA Lab** (VERSA Studios) — future-format development platform for film and broadcast.

Techniques, methods and recurring workflows

- **Vibe Coding** — prompt-led creative prototyping, especially for games and interactive apps; Rosebud, Spawn, Gizmo, Gambo are the canonical surfaces.
- **Trajectory Control (ATI)** (Wan / Qwen) — control of video trajectories at frame and clip level.
- **Reference-to-Video** (Veo 3.1 + Gemini) — consistent character generation from a reference image.
- **Motion Control nodes / Time-To-Move** — frame- and segment-level motion controls in generated video.
- **Layer Separation & Live Editable Text (LET)** (Loveart) — extract and edit text and subjects from images as PSD layers.
- **Temporal Reasoning** (NVIDIA ChronoEdit) — physics-aware edits and world simulation.
- **Practical Light Control** (PractiLight) — foundational diffusion for relighting.
- **3D Gaussian Splatting** — the spatial backbone running through PlayCanvas SOG, Hyperscape, Apple Personas and the open-source splat ecosystem.
- **World Models** — recurring topic spanning Marble, Genie, HY-World, UNI-1, ECHO, Matrix-Game, LingBot-World, Moonlake.
- **Multi-Shot Consistency** — Kling 3.0, ShotVerse, Higgsfield Shots, Veo 3.1 — the multi-shot turn that defined the spring-2026 video pipeline.

- **Camera Control / Multi-Angle Generation** — WAN Camera Control, ShotVerse, ATI, Multiple Angle Camera Control Node.
- **Style Transfer (the ComfyUI handbook)** — the canonical workflow ladder for style transfer.
- **Audio-Reactive Generation** — Audio Reactive Node Pack and the Lalal-API / ACE-Step pipelines.
- **Real-Time Generation Pipelines** — Krea Realtime, MotionStream, Decart Lucy 2.0 and Mirage, Xmax X1, Phoenix-4, LTX-2.3 Real-Time on Scope, NVIDIA SANA-Streaming, Magenta RealTime 2; the *interactive-generation* category that emerged in Q1 2026 and reached instruments and screens by the summer.
- **Agentic Workflows / Multi-Agent Stacks** — Sony Game Studios (49 agents, 72 skills), Adobe CX Enterprise, OpenAI AgentKit, Heygen Video Agent, Claude Skills.
- **Provenance-First Capture** — C2PA + SynthID + on-device detection in Premiere and YouTube as the emerging *standard* of provenance-aware production.
- **Music Stem Separation as a primitive** — Lalal AI, StemDeck, BandLab — stem separation moving from end-product to upstream primitive.
- **MCP (Model Context Protocol) as the connective tissue** — Unity MCP (free and unlimited-concurrency from July 2026), Blender MCP, fal MCP, and — in the June–July 2026 wave — Comfy MCP, Unreal MCP, Magnific MCP, Higgsfield MCP (running Claude Code inside After Effects) and the Runway MCP inside ChatGPT; the protocol the agentic creative ecosystem converged on.
- **Brief-First, Generate-Second** — the workflow practice the book argues for in Chapter 12 (and Chapter 17’s *toolchain in layers* section below).

This is the catalogue. By the time you read it, it will be incomplete — new tools have shipped, some on this list have been bought, renamed or killed. Treat it as a snapshot of one year of toolchain at the moment the toolchain became a stack rather than a list, and use it to orient yourself in whatever the state of play is when you pick the book up.

The adoption telemetry behind many of these tools — Firefly’s 22B-asset growth curve, ChatGPT’s 800–900M WAU figures, the Veo / Sora professional split, the GDC sentiment-vs-usage divergence — is analysed at length in DreamLab’s published research.

How to build a toolchain you can defend

The last thing I want to say in this chapter is something I have said in talks more often than anything else, because working creatives ask me this question, in some variant, every week.

How do I decide what tools to use, in a market that is changing this fast, without burning my whole month re-learning interfaces?

My short answer is: build the toolchain in *layers*, and accept that the layers move at different speeds.

The bottom layer — your foundation model, your modality stack, your agent platform — will change frequently. Treat it as ephemeral. Pick the best tools available *this quarter* and be ready to swap them next quarter.

The middle layer — your creative software (Adobe / Unreal / DaVinci / Pro Tools / Blender / Logic) — will change more slowly, and is the layer in which the AI capability will be progressively absorbed. Treat it as the long-term home of your craft. Learn it deeply.

The top layer — your *judgement*, your *taste*, your *briefing skill*, your *integration sense* — does not change. It is the layer the agents cannot copy, the layer the platforms cannot ship and the layer the next model release does not depreciate. Spend more time here than the toolchain wants you to.

The mistake I see working creatives make most often is to over-optimize the bottom layer and under-invest the top. The platform companies want you to spend your working hours chasing the new model release; the work that pays the bills, the work that finds an audience, and the work that survives a transition is built on the layer the platform companies cannot reach.

The toolchain, in the end, is the means. The work is the end. The tools change. The work, if it is any good, lasts.

That is the working operating model my studio has run on for the period this book covers. It is the model I would commend to anyone building, this year or next, a creative practice that survives the rest of the decade.

The transition is going to keep going. The tools will keep changing. The work that matters, on the other side, will be made by the people who kept their attention on the right layer.

Chapter 18 — Five Years Inside the Dream Machine

I called the newsletter *Dream Machine* in late September 2025 because the phrase caught something I couldn't yet articulate. I have said elsewhere in this book — in Chapter 16, and in the front matter — what I think the phrase has come to mean over the period the book covers: an apparatus capable of producing what until recently required a human mind, amplifying and distributing the dreaming of the humans who direct it.

What I have not done, anywhere else in the book, is let myself off the leash about what the apparatus might *become*. I have been careful, chapter by chapter, to put the evidence in front of the argument. The predictions in Chapter 16 are dated, falsifiable, defensive — written so the reader picking up the book in 2030 can grade them with a stopwatch.

This chapter is the opposite. This is the chapter where I let myself imagine what the next five years might look like if the trajectory of the previous eight months extends, accelerates, breaks and recombines in the ways I half-suspect it will but cannot, in any conventional analytical register, *prove*.

I want to be honest about the rules I am giving myself for this chapter, because they are looser than the rules the rest of the book has run on.

The rule is: each scenario has to be *argued from* something already in the book — a chapter, a number, a framework, a quoted source. I am not allowed to introduce a new mechanism out of thin air. But I am allowed to take a mechanism that is, in spring 2026, present in small or early form, and ask what it looks like by 2031 if it follows the curve I think it is on.

That is the contract for the next thirty pages. Wild but rooted. Speculative but cited. The kind of writing that the rest of the book would not have permitted, and that the closing letter to 2030 in Chapter 19 would deflate if I tried to put it there.

If a reader in 2031 has bought this book to grade my predictions, this is the chapter where you can grade me most harshly. Some of what follows will look, by then, embarrassingly off. Some of it will look, in the unflattering retrospect that books-on-transitions always get, banal — the obvious thing nobody had quite said yet. The interesting category — the predictions that turn out to be *roughly the right shape, badly miscalibrated on timing* — is the one I am writing for.

Six scenarios. Then the upside I am most hopeful for. Then the downside I am most afraid of. Then a handoff to the letter that closes the book.

One. The Petrillo settlement actually happens

In May 2026, the structural argument of Chapter 6 — that the institutional response to the AI training problem will, on the historical pattern, converge on the *levy-and-redistribute* mechanism James Caesar Petrillo built in 1948 for recorded music — is still, on the page, an argument from analogy. The 88% is the political mandate. The

GEMA ruling is the legal precedent. The Stability / UMG-style alliances are the commercial templates. The Musical AI creative-weight-attribution work is the technical infrastructure. None of it has yet been *assembled* into a working settlement.

By 2031, I think the assembly is done. Not perfectly, not globally, not without exclusions — but done, in at least one major jurisdiction, as a working mechanism that a working creative encounters on a bank statement.

The shape of the thing I expect: a national or supranational *AI Performance Trust Fund*, modelled on the MPTF, capitalised by a statutory per-output levy on commercial generative output, governed by a joint labour–platform body, distributed to working creatives whose training-data contributions were identifiable through the creative-weight-attribution layer described in Chapter 13. The first cheques are small. A working illustrator opens a statement and sees the line *AI Royalty Distribution: £43.18*. A session musician sees *£127.42*. A photographer whose backlist sat in a Getty-class licensed dataset sees a meaningful four-figure annual sum.

The cultural significance of those small numbers will, I think, be larger than the numbers themselves. The first cheque is the moment the 88% becomes a *fact* rather than a *demand*. The mechanism is on the books. The architecture is real. The question stops being *whether there will be a Petrillo settlement* and starts being *how much, to whom, on what calibration*.

The bargaining ground for the next decade is set on that question.

Two. The internet bifurcates, formally

Chapter 4 made the case that the web of 2026 was, in measurable ways, splitting into two distinct attention environments — the *Dead Internet* of synthetic content optimised for synthetic attention, and the *Living Web* of provenance-stamped, human-verified work made for an audience that had organised itself around the difference.

By 2031, I think the bifurcation is no longer a cultural metaphor. I think it is *operational infrastructure*.

The provenance stack the book has spent so many pages cataloguing — C2PA, SynthID, the Cannes Disclosure Standard, the AP and BBC wire-service signing chains, the Adobe content credentials — will, by 2031, have stitched together into a functioning verification layer. Major browsers will render an indicator. Major platforms will, under pressure from advertisers tired of paying for bot-on-bot impressions, surface verified-human content separately from synthetic. A handful of *human-only* subscription services — I would not bet against the BBC, *The New York*

Times and one or two of the streaming-music incumbents launching first — will offer “no synthetic content ever” as a paid product.

The split internet is not, in the version I expect, a *clean* split. The synthetic layer will dwarf the verified layer in volume, by perhaps a hundred to one. But the verified layer will capture an outsized share of *paid attention*, *advertiser dollars*, and — most importantly — *cultural credit*. The slop ceiling of Chapter 5, measured at 44/3 on Deezer in April 2026, will by 2031 have hardened into a roughly stable ratio across most major content categories: synthetic supply dominant on the upload side, verified-human attention dominant on the consumption side.

The Living Web is, by 2031, no longer an aspiration. It is a *distribution layer* you can buy access to, with a different economics from the synthetic layer running underneath it.

Three. World models replace flat video as the default high-end medium

Chapter 9 argued that the most important technical shift of 2025–26 was not the video models but the world models — Marble, Genie 3, the Hunyuan 3D family — and that the rate at which spatial-AI tooling was moving from research demos to consumer product was the underappreciated structural story of the period.

By 2031, I think we look back on flat 24-frame video the way 2026 looks back at black-and-white silent cinema. Still made. Still loved. Still occasionally the right medium for the work. But not the *default*.

The default high-end production format, by 2031, is a *navigable spatial render* — a world that the camera moves through in post, that the audience can choose to view from any angle through a Vision-class headset, that an editor can re-cut from a different perspective two months after principal photography is done. The boundary between film, games, immersive and live performance that Chapter 9 described as eroding is, by 2031, functionally gone for new productions.

The career implications, for working filmmakers, are substantial. The skillset of the director-of-photography fragments into world-curation, lighting-prompt design, and post-spatial composition. The skillset of the editor expands into multi-angle viewer choreography. New roles emerge — *spatial continuity supervisor*, *world dramaturg*, *perspective director* — that have no clean analogue in the 2026 production pipeline.

This is also where I expect the legacy industries’ Chapter 7 strategic positioning to come back to haunt them. The studios that bet on flat AI-generated video as the next medium will discover, between 2027 and 2029, that they bet on the second-to-last format of the previous era. The studios that built world-model fluency into their pipelines — Sony’s documented experiments, the AI-native studios like Imaginae and Critterz, a handful of regional players in Korea, Japan and India — will be the

dominant 2031 producers of *the kind of work that uses the medium the way it actually works*.

Four. The orchestrator becomes a credentialled guild

The orchestrator role — central to Chapter 12, referenced through Chapter 14 and Chapter 15 — was, in 2026, a *description of an emerging practice*. The Sony 49-agent / 72-skill team was the most-cited case study. The working title circulated in studios and agencies but did not yet carry contract weight.

By 2031, I think the orchestrator is a *credentialled profession* with collective bargaining power.

The guild structure I expect — and this is the most institutionally specific of the six predictions, so the most likely to be wrong on detail — will look something like the IATSE locals that govern below-the-line film labour, or the Writers Guild's apprenticeship and credit system. *Senior orchestrators* will carry the credential, demonstrate competence in agent stewardship across a defined toolchain, negotiate over rates and over the *quantity of agent labour* a single orchestrator can supervise. The Writers Guild will, I think, be among the first to formalise the role; SAG-AFTRA, Equity and the games-development unions will follow within eighteen months of whoever moves first.

The collective-bargaining angle is the part that will, in 2031, look most novel and most obviously inevitable in retrospect. If a single human orchestrator can supervise the work of, say, twenty agents — a number that today's productivity research is pointing at, with substantial variance — then the question of *who pays whom for the productivity gain* is exactly the Petrillo question from earlier in this chapter, applied to the orchestrator's own labour rather than to the underlying training data. The answer the guild will negotiate is a *productivity share* of the agent-team output, with the orchestrator's individual rate tracking the value of the supervised work, not the volume of the orchestrator's keystrokes.

This is one of the more concrete reasons I am cautiously optimistic about the working-creative position over the next five years. The job that emerged in 2026 — the maker-as-orchestrator — is a job for which collective bargaining is *unusually well-suited*, because the value being created is unambiguously human in origin and unambiguously measurable in output. Unions exist for exactly this kind of bargaining problem. The Petrillo template, again.

Five. An AI-native studio wins the Palme d'Or before Hollywood does

This is the prediction that will, depending on the reader, look either obviously safe or wildly provocative in 2031.

Chapter 7 catalogued the new AI-native studios — Imaginae, Wonder, Asteria, Obsidian, Critterz, Gossip Goblin — and made the case that their structural

advantage over the legacy studios is not technical. It is *cultural*. They have no calcified rules to unlearn. They have no inherited risk-aversion. They have no franchise-template gravity to pull them back to the engine-optimal move. By the chess-grandmasters analogy in [Chapter 16](#), they are, by default, in a better position to play the move the machine would not have generated.

By 2031, I think one of them — or, more likely, a successor company emerging from a director who learned in their orbit — wins the Palme d'Or, the Golden Bear, the Silver Lion or one of the other top European festival awards. *For its writing or direction. Not for its technology.*

The cultural moment when that happens will be larger than the prize itself. It will be the *Petrillo Settlement* of cinema-craft acceptance: the moment the question shifts from *can AI-native cinema be art* to *which AI-native film will be canonical*. The follow-on year, I expect the Academy — having spent 2025–28 reinforcing its *you must be human to win* rule for above-the-line crafts — to add a parallel category, or modify the existing one, to recognise hybrid human-orchestrator-AI authorship under contested but workable rules. The festival juries will get there before the Academy does. They almost always do.

The corollary, for the legacy studios, is bleak. The diagnosis in [Chapter 7](#) — that Hollywood, commercial music and AAA games spent fifteen years optimising toward the engine-optimal move and arrived at the AI moment producing exactly the work the engines can now replicate cheapest — will be on full display when the first AI-native Palme winner is, recognisably, *not* a tentpole. By 2031, the Hollywood studio system will be in the position the major-label music industry was in around 2008: still dominant by revenue, visibly losing the cultural argument, scrambling for the next operating model.

Six. The platform monopoly cracks

The reason I am giving this prediction last among the six is that it is the one I hold most loosely. The structural argument is clean. The timeline is the thing I cannot pin.

By 2031, I do not think the OpenAI / Anthropic / Google triad of frontier-AI dominance will look the way it does in May 2026. The pressures pushing against it, catalogued in different chapters of this book, are:

- The open-weight ecosystem — Tencent's Hunyuan, Alibaba's Qwen and Wan, DeepSeek, Meta's Llama, Mistral, Stability — collectively used by, per [Chapter 17](#), some 80% of startups pitching the major venture funds by spring 2026. Open weights compound. Closed weights, by structure, do not.
- The consent-trained category from the new section in [Chapter 6](#) — Adobe Firefly, Bria, Getty's licensed model, Moonvalley Marey, AIODE, the Stability /

UMG-style alliances — which, by 2031, has either *forced* the frontier labs to license their training data retroactively, or has *displaced* them in the enterprise procurement environment by carrying indemnities the frontier labs cannot match.

- The sovereign-model trend, only just starting in 2026, in which national governments — I would put France, Korea, India, the UAE, Brazil, Japan on the early list — fund consent-trained, citizen-licensed foundation models as a strategic asset, governed under sovereign rules, distributing royalties through national rights-holder bodies.
- The bifurcating internet from earlier in this chapter, which by 2031 makes *which model trained on what* a visible product-marketing claim, in the way *fair-trade-certified* or *organic* became visible product-marketing claims in food.

By 2031, the AI model market I expect looks much more like the 2024 cloud market than the 2025 AI market. Multiple major model families. Sovereign options. Open-weight defaults at the long tail. Indemnification and provenance as standard procurement requirements. The frontier-lab valuations of 2025 either justified by deep enterprise penetration in narrow categories, or remembered the way AOL's 1999 valuation is remembered.

The corollary I want to put on the page, because it has not yet been said cleanly anywhere else in this book: the platform companies dominant in May 2026 are, on the historical pattern, *unlikely to be the platform companies dominant in 2031*. The history of computing is that the dominant platform of one era is rarely the dominant platform of the next. The Wintel of the 1990s, the iOS / Android of the 2010s, the AWS / Azure of the 2020s — each was a dominant pairing whose successor came from a direction the incumbents did not see. Whatever the dominant generative-AI platform of 2031 is, I think there is a non-trivial chance it is a company most readers of this book — including me — have not yet heard of.

The upside I am most hopeful for: the Dream Machine becomes literal

I want to spend the rest of this chapter on two scenarios that do not fit cleanly into the *predict and date* register, because they are the ones I think about most when I am away from the work and most off-script when I am in it.

The first is what happens if the title of this book turns out to be a literal description.

The phrase *Dream Machine*, as I have used it through the newsletter and the book, has been a metaphor for the apparatus of generative AI. A composite name for the platforms, models, infrastructure and human labour that, between them, produce the new creative work.

I want to entertain the possibility that, by 2031, the phrase describes a *literal piece of consumer hardware* — a wearable creative-cognition prosthetic, sitting

somewhere on the continuum between a Vision Pro and a brain-computer interface, that watches a working creative make work, learns their taste, infers their intent, and contributes — in real time, in collaboration, at the speed of thought — to the work in progress.

The technical building blocks are visible. Apple's Vision Pro and the next-generation headset class are mature consumer products by 2031. The brain-computer interface category, dominated in 2026 by medical applications, is on a trajectory the consumer hardware press has been undercounting. The agentic orchestration layer that I described in Chapter 12 is, by 2031, no longer something a human types instructions into. It is something a human *converses with, gestures to, thinks at*. The interaction model of *prompting* is, by 2031, an obsolete UX pattern, remembered the way command-line interfaces are remembered now.

The literal Dream Machine, in this scenario, is not a separate device. It is the orchestration layer made *embodied*. A creative cognition prosthetic that allows a working creative to externalise, manipulate, sketch, iterate and finish creative work at a fluency that is, today, available only to the smallest set of professionals who have spent thirty years building the relevant neural pathways. The leverage the device provides is not *replacement* of the human imagination. It is the closing of the gap between the human imagination and what the human hand can, today, get out into the world.

If this device — or one like it — ships at consumer price in the period this chapter covers, the *access* principle from Chapter 16 gets a lift I cannot easily exaggerate. The bedroom-Hollywood dynamic from Chapter 15 — the teenager with a midrange GPU producing studio-quality output — extends to *every* discipline, including the ones that today still require expensive equipment and decades of haptic training. The new geography I argued for in Chapter 16 — the dispersion of canonical creative work to places the previous century forgot — becomes structurally inevitable rather than aspirational.

The Dream Machine, in this scenario, is not the platforms. It is the *prosthetic*. The platforms are infrastructure underneath it. The dreaming is, finally, located where the title was always insisting it was located: in the human at the centre of the apparatus.

I want this scenario to be the one that comes true. I think there is a real chance — not a high probability, but a meaningfully non-zero one — that it does.

The downside I am most afraid of: the audience becomes synthetic

The corresponding worst case I want to put on the page is the failure mode the slop ceiling cannot, by construction, defend against.

The slop ceiling, as I described it in [Chapter 5](#), works because *the audience can tell* — at the speed of a swipe, in the aggregate, with reliability — that synthetic content does not carry the human signal real human attention is calibrated to. The ratio holds, the 44/3 holds, the audience underweights the slop, because there is an audience, and the audience is made of humans, and the humans are still, in May 2026, paying attention to one another.

The downside scenario is that, by 2031, that last condition has decayed.

The mechanism is the one the Dutch researchers cited in [Chapter 5](#) and the bot-traffic investigations cited in [Chapter 4](#) were already, by spring 2026, beginning to document. As synthetic content production becomes free and synthetic attention-allocation agents become widely deployed, the *measured* audience for any given piece of content drifts further from the *human* audience. Platforms optimise against the measured audience, because that is what their analytics surface. The content the platforms surface, in turn, becomes the content the synthetic audience underweights least. The synthetic audience grows. The human audience, exhausted, exits. The platforms continue to grow on the measured numbers because the measured numbers are increasingly bot-on-bot. The cultural production layer continues to produce, but produces *for bots watching bots*. The work loses meaning slowly, then quickly, because the audience that produced meaning has stopped being part of the loop.

This is the version of 2031 I am most afraid of. Not because AI replaces creators. Creators are not the load-bearing structure. *The audience is the load-bearing structure*, and the audience is the thing the architecture of the synthetic internet, deployed without the verification layer described in scenario two of this chapter, is structurally positioned to dissolve.

The corresponding fight, the one I think the next five years has to win to avoid the downside, is the *audience verification* fight. The C2PA-equivalent for *who is watching*, not just for *who made*. The mechanism that lets a creator know — and lets an advertiser pay for, and lets a platform surface — the work that *real humans were really paying attention to*. The early signals in this direction are there, in the form of platform attestation experiments and the bot-traffic regulatory pressure documented in [Chapter 4](#). None of them are, in May 2026, anywhere near the scale of the upstream provenance work.

If the next five years build the verification layer for production but not for consumption, the slop ceiling holds in the short term and collapses in the medium term. The creators win the upstream battle. The audience disappears anyway. The downstream meaning of the work, by 2031, is gone — not because human creativity has been replaced, but because the social fact that creativity is made *for* somebody has been quietly dissolved underneath the work.

This is the failure mode I think we are most exposed to and least talking about.

What we make of it

I want to close this chapter the same way I open the next.

The phrase *Dream Machine* was, when I named the newsletter, a placeholder for a feeling. By the end of Chapter 16 it had become an argument: that the apparatus we are building amplifies, multiplies and distributes the dreaming of the humans who direct it, and that the question every working chapter of this book has been about is *whose dreams the machine amplifies*.

What this chapter has tried to do is take that argument out to five years and let it run.

Six things I think happen. The Petrillo settlement, made real, on a working creative's bank statement. The internet, formally split. World models, displacing flat video as the default. The orchestrator, credentialled and bargained-for. The first AI-native Palme d'Or. The platform monopoly, cracked. One upside I am hopeful for — the Dream Machine made literal, a creative cognition prosthetic that closes the gap between human imagination and finished work. One downside I am afraid of — the audience layer collapsing under synthetic capture, leaving the work without anyone to make it *for*.

Three of the six will, almost certainly, be wrong. One of them will be wrong in the direction of *I undersold this*. One of them will be wrong in the direction of *the timeline was slower than I thought*. One of them will be wrong in the direction of *the thing I described did not happen because the thing that happened instead was weirder and more interesting*. I do not know yet which is which.

The point of the exercise is not to be right about the six. The point is to put a *shape* on the page, in 2026, for the conversation that the working creatives, the studios, the unions, the policy people and the audience will be having with each other, week by week, between now and 2031. The four principles from Chapter 16 — *agency, attribution, access, audience* — are the test you apply to each daily decision. The scenarios in this chapter are the *direction* the aggregate of those decisions is, on my read of the available evidence, most likely pointing in.

The title of this book, I have come to think, was always slightly mis-named. The Dream Machine is not the apparatus the platforms built. It is not even the prosthetic the consumer-hardware companies might ship by 2031. The Dream Machine is the *whole system* — the platforms, the prosthetics, the working creatives, the audience, the unions, the institutions, the regulators, the literacy initiatives, the newsletters, the festivals, the awards, the lawsuits and the daily decisions of millions of people about whose work to make and which work to pay attention to. The thing the title

points at is the *coupled human-machine cultural production system* that the period this book covers brought into existence.

What we make of that system, between May 2026 and the moment a reader picks this book up in 2031 to grade me, is the work.

The letter in the next chapter is what I want to leave that reader with.

Welcome to the Dream Machine.

Epilogue — A Letter from the Dream Machine

To the creative person reading this in 2030, picking the book up out of curiosity or for a class or because someone older than you said you should:

I want to write you a short letter, because by the time you are reading this, more than half of what is in this book will already be wrong.

Some of the tools I have spent chapters discussing will be obsolete. Some of the companies will have been bought, broken up or sunk. Some of the law I described as new and contested will have settled into precedent that everyone takes for granted. Some of the people I quoted as authorities will have changed their positions, and some of the predictions I made — the ones I let myself make — will have aged in ways I would now find embarrassing to read back.

This is fine. It is, I think, the right outcome for a book written in the moment that this one was written in. The job of a book about a transition is not to predict where the transition lands. The job is to *describe the moment honestly* — what was being argued about, by whom, with what evidence, and what was at stake — so that the reader picking it up later has, at minimum, an accurate record of *what the people in the room thought they were doing*.

There are a few things I want to flag for you, looking back from your time at mine.

What we got right

We got the slop ceiling right. I am almost certain of this, even sitting here in May 2026, because the underlying mechanism — that audiences allocate attention to *meaning*, not to *quantity* — is one of the most reliable findings in the history of cultural production, and there is no plausible scenario in which it stops working.

If you are reading this in 2030 and there is a thriving creative economy, it is because the slop ceiling held. The audience that turned away from synthetic content in 2026 either kept turning away, or the platforms and the creators figured out how to make AI-augmented work that genuinely earned attention — sincere, transparent, made with care. Both outcomes preserve the underlying contract between maker and audience. Either way, the ratio between volume and attention is still doing its work.

If, on the other hand, you are reading this in 2030 and the creative economy looks bleaker than the one I have described, it is — almost certainly — because the architecture of the internet was allowed to drift further in the direction the Dutch researchers identified in 2025: an attention market optimised against meaning, where the ratio between volume and attention stopped functioning because the *audience itself* had become indistinguishable from the bots.

The slop ceiling was real. The question for your decade is whether you protected the *audience* that produced it.

What we got wrong

We got the timeline wrong, in both directions.

In some places, we expected the change to be slower than it turned out to be. The pace at which world models moved from research demos to consumer features in the eight months I documented in this book was something almost nobody — including me — would have predicted in late 2024.

In other places, we expected the change to be faster. The full AI-native film studio that was supposed to replace Hollywood by 2027 is still, in the spring of 2026, mostly a series of well-funded prototypes. The fully autonomous game-development pipeline is, even at Sony, still an *assisted* pipeline rather than an *autonomous* one. The headline futures the platform companies have been selling are, in the main, still further off than the press cycle suggested.

If you are reading this in 2030, I suspect that the timeline gap will look in retrospect like the most obvious thing about the moment we were in. The interesting changes were the ones nobody put on a slide. The boring changes — the slow institutional repositioning, the contract-by-contract renegotiation of how creators are paid, the patient construction of the C2PA standards, the union by union recalibration of what counts as a fair use of a performer's likeness — were the ones that actually shaped your decade.

This is, I think, almost always true of technology transitions. The press loves the visible. The structural change happens in the invisible. The fact that you can pick this book up in 2030 and read it without unusual cost or effort is the result of a thousand small infrastructure decisions, made by people whose names did not appear in this book, that none of us thought to celebrate at the time.

What I want you to know about us

We were, in 2026, frightened. Most of the working creatives I knew that year were carrying a layer of low-grade fear — about their livelihoods, about their crafts, about the institutions that had organised their working lives — that I do not think showed up in the trade press or the platform keynotes or even in this book as much as it should have.

We were also, I think, more hopeful than we said out loud. Most of the working creatives I knew that year were also doing some of the most experimental, most curious, most adventurous work of their careers, because the tools made things possible that had been impossible the year before and they wanted to find out what.

Both states were true at the same time. They will be true again in your decade, about whatever the next transition is. I don't think we managed the fear well. I think we did okay with the hope. I am proud of the people who turned up, week after week, to argue about how all of this should go.

A specific request

If the institutions of the generative creative economy — the unions, the rights bodies, the consultations, the literacy initiatives, the disclosure standards, the festivals, the regional studios, the indie filmmakers in places that the previous century forgot — are still doing their work when you read this, *thank the people who built them*. Most of those people have not become rich, or famous, or particularly comfortable. They were working creatives, union officers, policy officers, festival organisers, technologists, lawyers, professors and freelancers who chose, in their unglamorous moments, to do the slow institutional work that this kind of economy needs.

The platform companies will not remember them. The audience will not know their names. The cultural histories will record the names of the directors and the songwriters and the celebrity AI controversies. The institutional work happens in the basement and the archive and the union office and the policy briefing.

If you are inheriting a creative economy that works, you are inheriting it from those people. Thank them while they are still around to be thanked.

A specific request to working creatives in 2030

Keep going.

The pace of change will be at least as fast in your decade as it was in mine. The next paradigm of AI tooling — whatever it is — will, I am sure, make the world models and agents of 2026 look as quaint as Photoshop in 1996. You will have your own version of *the day Sora landed*. You will have your own *Tilly Norwood week*. You will have your own 88%. You will have your own choice between the extractive and the generative economy.

Whatever the new tools are, the principles I tried to articulate in Chapter 16 — *agency, attribution, access, audience* — will, I am confident, still apply. They are not specific to the 2026 transition. They are general properties of how a humane creative economy organises itself, and they will work in any technology environment that produces creative outputs at scale.

The deeper conviction underneath the principles will, I think, also still apply. The age of the *How* — the long century in which the central question of working creative life was *can you do the thing?* — was, by 2026, visibly ending. The age of the *Why* — taste, intent, authenticity, the willingness to take a risk on the move the data does not endorse — was visibly starting. The image I borrowed from elite chess in Chapter 16, of grandmasters winning by deliberately playing the move the engine would not have played, will still be the picture I would commend to you. By the time you read this, the engines will be better. The pressure to play the engine's optimal line will be stronger. The competitive advantage of the *deliberately un-machine-like* move — of

the move that is yours because no machine, trained on what came before, could have generated it — will, if anything, have widened.

Apply them. Argue for them. Build for them. Defend them when they are under pressure.

And — this is the part I want you to take seriously even when it feels grandiose — *write the newsletter*. Whatever your version of the newsletter is. The thing that, when the moment comes, you sit down on a Monday morning and decide to write because nobody else seems to be doing it. Most of those newsletters won't matter. A few of them will. The one you are reading was, for me, one of the most consequential decisions I ever accidentally made.

You will know which moment is the moment, when it arrives. You will know because the people you respect will be staring at the same news cycle, looking at each other in the same way, asking the same question.

When that moment arrives in your decade, write the newsletter.

Closing

I want to end this letter the same way I have ended every issue of *Dream Machine* for the past six months, because the line has come to mean more to me than I expected it to:

If you've got any recommendations or things we need to know about for our next edition, please feel free to reach out.

That sentence was, for all the time the newsletter has run, my way of acknowledging that I did not have the picture on my own. That the readers were the network. That the work was a *conversation*, not a broadcast. The newsletter has only ever been as good as the community of creatives, technologists, union reps, academics, festival programmers, indie filmmakers, working musicians and audience members who have, week after week, sent me the things I would otherwise have missed.

If you are reading this in 2030, the newsletter is still — in whatever form it has taken by then — that same conversation. Keep adding to it. The Dream Machine, in the end, is not the AI. It is the network of humans who, between them, are deciding what the machine is *for*.

Welcome to the Dream Machine.

— Pete Woodbridge, *DreamLab*, the North West, May 2026

*Appendix A — A Quantitative Anatomy
of Six Months*

This appendix is a structured tour of the corpus the book was built from. It is not in the body of the manuscript because it would interrupt the argument; it lives here so that the reader, the policy researcher, the journalist or the historian picking the book up later can see what the underlying data actually looks like and check the arguments against it.

* * *

A1. Corpus shape

- **Total fetched and parsed articles:** 1,388.
- **Total captured words across the corpus:** ~1,099,216 (~6,945,361 characters of post-extraction text).
- **Source span:** 41 editions of *Dream Machine* | *Creative AI*, published between 6 October 2025 and 12 August 2026 (the quantitative tables in §A1–A6 below cover the analytic window of Issues 1–29; **Issue 30** (21 May 2026), **Issue 31** (1 June 2026), and **Issue 32** (4 June 2026) are documented in §A7 as post-cut supplemental data; **Issues 33–36** (9 June – 2 July 2026) and the World Cup deep dive are documented in §A8; **Issues 37–41** (9 July – 12 August 2026) are documented in §A9. All are incorporated narratively into the closing material across the body chapters).
- **Average articles per newsletter edition** (in this corpus): ~48.
- **Capture rate** against the full curated URL set: **91.4%** (1,438 of 1,574 URLs returned readable content; the remainder hit bot-detection, 404s, or live-page connection issues).

A2. Most-cited domains

Where did six months of curated coverage actually come from? The top 30 domains, by article count:

Rank	Domain	Articles
1	x.com	82
2	github.com	75
3	youtube.com	74
4	musically.com	57
5	theguardian.com	38
6	variety.com	38
7	hollywoodreporter.com	34
8	open.spotify.com	25
9	huggingface.co	24
10	musictech.com	22
11	forbes.com	21
12	gamesindustry.biz	20
13	bbc.co.uk	18
14	deadline.com	18
15	techcrunch.com	17
16	pcgamer.com	17
17	musicradar.com	17
18	blog.google	15
19	futurism.com	15
20	completemusicupdate.com	15
21	videogameschronicle.com	14
22	businessinsider.com	13
23	theverge.com	12
24	musicbusinessworldwide.com	12
25	cnet.com	11
26	gamesbeat.com	11
27	adweek.com	10
28	digiday.com	10
29	pocketgamer.biz	10
30	blog.comfy.org	9

Reading note. Trade press (Hollywood Reporter, Variety, Deadline, Music Business Worldwide) and tech press (Verge, Wired, TechCrunch) dominate. The platform-company blogs (OpenAI, Adobe, DeepMind, Stability) and policy bodies (UK Gov, Reuters Institute, Imperva, Cloudflare) sit underneath. The geographic concentration is North American and British, which reflects both the newsletter author’s vantage point and a real imbalance in where creative-AI coverage is concentrated.

A3. Story volume by sector, month by month

How the conversation moves through the six months — count of corpus articles touching each sector, by publication month of the newsletter issue that cited them:

Month	Film & TV	Games	Music	Adv/Mkt	News	Policy/Law	Tools/Models	Total
2025-10	136	111	155	104	25	102	184	817
2025-11	112	89	120	97	14	93	150	675
2025-12	73	58	78	57	9	43	98	416
2026-01	89	69	96	70	13	56	120	513
2026-02	77	72	94	61	7	63	109	483
2026-03	81	49	81	55	12	47	106	431
2026-04	100	84	110	69	11	64	138	576
2026-05	74	60	89	53	14	57	113	460

Total story tags across the period: 4,371. (Articles can fall into more than one sector — many do, which is itself part of the story: the boundaries between film, games, music, advertising and tooling have been meaningfully porous through the AI era.)

A4. The voices: public-figure mention frequency

Public figures appearing in three or more corpus articles, ranked by article count. This is *not* a sentiment ranking — only a measure of how often someone surfaces in the AI conversation:

Rank	Name	Articles
1	Sam Altman	32
2	Tilly Norwood	30
3	Xania Monet	18
4	Taylor Swift	18
5	RZA	17
6	Steven Soderbergh	17
7	James Cameron	16
8	Paul McCartney	12
9	Mark Zuckerberg	12
10	Eline Van der Velden	12
11	Madonna	12
12	MrBeast	12
13	Tim Sweeney	9
14	Christopher Nolan	9
15	Emily Blunt	8
16	Breaking Rust	8
17	Mikey Shulman	7
18	Matthew McConaughey	7
19	Steven Spielberg	7
20	Robert Kyncl	6
21	Guillermo del Toro	6
22	Will Smith	6
23	Lucian Grainge	5
24	Sienna Rose	5
25	Natasha Lyonne	5
26	Fei-Fei Li	4
27	George Clooney	4
28	Demis Hassabis	4
29	Ron Howard	4
30	Ted Sarandos	4
31	Adam Mosseri	3

Rank	Name	Articles
32	Imogen Heap	3
33	Dave Stewart	3
34	Chris Pratt	3
35	Joost van Dreunen	3
36	Brian Grazer	3

Reading note. James Cameron, Guillermo del Toro and Leonardo DiCaprio are the three voices most-cited in opposition to generative AI in performance. Tilly Norwood and Xania Monet are the two most-cited synthetic entities in the corpus. *Both lists matter equally to the story this book is telling.*

A5. The tool wave

Cumulative count of distinct AI tools, models and platforms entering the corpus, by month of first mention:

Month	New tools first mentioned	Cumulative
2025-10	38	38
2025-11	6	44
2025-12	3	47
2026-01	5	52
2026-02	2	54
2026-03	4	58
2026-04	0	58
2026-05	2	60

Reading note. The tool cadence ran at roughly **7.5 new platforms or major-version releases per month** across the period. This is roughly four times the pace of any other software-tool category I have personally tracked over a comparable window. The implication is that any working creative making technology bets in this period was, by definition, working with incomplete information — the relevant toolchain had not stabilised long enough for any single bet to settle.

Most-mentioned tools and platforms (top 30, by article count):

Rank	Tool	Articles
1	Udio	587
2	Wan	539
3	ChatGPT	136
4	Gemini	98
5	Anthropic	84
6	Suno	76
7	Sora	69
8	ElevenLabs	49
9	Premiere	48
10	Veo	46
11	ComfyUI	45
12	Sora 2	37
13	Claude Code	36
14	Flux	36
15	Veo 3	30
16	Runway	30
17	Kling	30
18	Nano Banana	29
19	Seedance	26
20	Tencent	24
21	Genie	24
22	Firefly	21
23	Photoshop	21
24	Hunyuan	18
25	Veo 3.1	17
26	Luma	13
27	Marble	10
28	Rodin	10
29	LTX-2	9
30	Higgsfield	7

A6. The vocabulary shift

Recurring key phrases by month — articles containing each phrase:

Phrase	2025-10	2025-11	2025-12	2026-01	2026-02	2026-03	2026-04	2026-05	Total
ai slop	8	8	6	15	5	7	7	6	62
ai actor	4	5	3	2	2	2	1	1	20
synthetic performer	0	0	2	1	0	0	0	1	4
world model	8	3	3	1	9	6	8	2	40
agentic ai	12	6	2	7	6	5	8	3	49
ai agent	20	18	4	11	12	9	16	10	100
deepfake	9	6	3	10	6	8	4	3	49
human authorship	1	3	1	0	1	0	1	1	8
training data	9	10	11	6	8	12	5	10	71
consent	16	11	11	11	9	6	10	11	85
license	26	30	18	21	35	25	29	16	200
copyright	50	53	27	24	33	26	25	24	262
ai-generated	65	58	26	46	30	29	41	28	323
ai actress	3	2	0	0	1	0	1	0	7
watermark	5	2	0	4	3	2	1	3	20
synthid	1	0	0	2	2	1	1	0	7
c2pa	0	1	0	2	1	2	0	0	6
provenance	5	2	1	3	1	1	1	5	19
disclosure	6	12	7	9	4	3	7	6	54
fingerprint	0	2	0	3	2	0	0	1	8
creative ai	8	5	0	1	0	1	3	1	19
generative ai	85	59	35	43	32	32	40	32	358
creator economy	4	10	0	3	3	0	2	1	23
tilly norwood	10	7	3	3	1	2	2	2	30
xania monet	3	7	2	3	0	2	1	0	18

Phrase	2025-10	2025-11	2025-12	2026-01	2026-02	2026-03	2026-04	2026-05	Total
breaking rust	0	3	1	1	0	1	1	1	8
slop ceiling	0	0	0	0	0	0	0	0	0
model collapse	0	0	0	0	0	0	0	0	0

Reading note. Watch *AI slop* — it goes from a fringe term in October 2025 to a Merriam-Webster word of the year by December and a policy framing by the spring. Watch *agentic AI* — it lifts after the October DevDay and never falls back. Watch *world model* — barely present in October 2025, ubiquitous by April 2026. Watch *consent / license / copyright* — climbing all the way through, with a sharp December spike around the UK consultation closure.

A7. May–June 2026 supplemental: the Issue-30, Issue-31 and Issue-32 datapoints

The analytic corpus underlying §A1–A6 closes at *Dream Machine* Issue 29. **Issue 30** (21 May 2026), **Issue 31** (1 June 2026), and **Issue 32** (4 June 2026) post-date the analytic cut and are not represented in the article-frequency tables above; Issue 30 catches the **Google I/O 2026** announcement wave, Issue 31 catches the **June IP-licensing turn** (Hasbro Sixth Wall / Behavioural Licensing, Spotify–UMG, the Epidemic Sound *Future of the Creator Economy* report and the Sora shutdown), and Issue 32 catches the **early-June 2026 policy and industry inflection** (Nandy admission, Suno Series D, Dreams of Violets, Amazon AI Creators Fund, Scorsese/Black Forest Labs, NVIDIA Cosmos 3, Roblox/Morpheus, Tripo AI). Together they are the source for the manuscript’s closing-month additions. The numerical datapoints from Issues 30, 31 and 32 worth recording here in standalone form:

Datapoint	Value	Source
SynthID watermarked items, cumulative	100B+	Google DeepMind, May 2026
Wonder Studios total funding	\$50M	<i>Forbes</i> , May 2026
Runway Japan investment (Tokyo office)	\$40M	Runway, May 2026
Viktor (virtual AI coworker) Series funding	\$75M	<i>Fortune</i> , May 2026
Sondo AI claimed global users	10M	<i>Musically</i> , May 2026
13–15 year-olds using AI to “be creative” (Snapchat survey)	31%	Snap Newsroom, May 2026
Australians who say AI-generated ads make them trust a brand less (YouGov)	45%	YouGov AU, May 2026
NVIDIA SANA-WM model size	2.6B	NVIDIA, May 2026
SANA-WM native video-generation length	60 sec	NVIDIA, May 2026
Epidemic Sound <i>Future of the Creator Economy</i> survey size	3,000	Epidemic Sound, June 2026
Creators using AI tools (Epidemic Sound)	94%	Epidemic Sound, June 2026
Creators expecting to increase AI usage in next 12 months	72%	Epidemic Sound, June 2026
Creators feeling pressure to use AI to keep up	89%	Epidemic Sound, June 2026
Creators who believe human-created content will become a premium product	75%	Epidemic Sound, June 2026
AI-generated songs as share of Apple Music plays	<1%	<i>Billboard</i> , June 2026
Sora peak monthly revenue (pre-shutdown)	~\$540k	<i>Tech-Insider</i> , March 2026
Disney–OpenAI announced licensing value (never executed)	\$1B	The Walt Disney Company / OpenAI, December 2025

Datapoint	Value	Source
Hasbro Sixth Wall initial character catalogue	12	Hasbro, June 2026
Suno Series D funding round	\$400M	<i>Variety</i> / <i>TechCrunch</i> , June 2026
Suno post-money valuation	\$5.4B	<i>Variety</i> , June 2026
Suno annual recurring revenue	\$300M	<i>Variety</i> , June 2026
Suno subscribers	2M+	<i>Variety</i> , June 2026
<i>Dreams of Violets</i> (Tribeca) production cost	\$2,000	<i>Variety</i> , June 2026
<i>Dreams of Violets</i> production time	3 months	<i>Variety</i> , June 2026
Amazon AI Creators Fund: days between commission confirmation and creator exit (Punky Duck)	2	<i>Variety</i> / <i>Hollywood Reporter</i> , June 2026
Lisa Nandy (UK Culture Secretary) SXSW London admission that government “got it wrong” on AI copyright	—	<i>Screen Daily</i> , 1 June 2026

Reading note. Issue 30’s headline tool releases (Gemini Omni, Beeple Canvas, Sony Woosh, Mirelo SFX 1.6, Tencent Ardot, Odyssey Starchild-1 / Agora-1, NVIDIA SANA-WM, Apple Headsup, Stable Audio 3.0, PhotoGIMP, Tamber, ECABridge, Claude/ComfyUI) lift the cumulative tool count in §A5 by roughly **a dozen entries in a single week**. The May-2026 cadence is the highest single-week tool-release count in the period the book covers, and reads — in the context of the §A5 average of 7.5 new platforms per *month* — as a Google-I/O-week saturation point. Issue 31 settles back toward that prior cadence, with the focus shifting from raw tool releases to the *IP-and-licensing turn* (Spotify–UMG, Hasbro Sixth Wall, the Disney–OpenAI corrected record) and the *creator-attitudes* data from the Epidemic Sound report — the dataset that anchors the closing argument of Chapter 14 and the *three Cs* framework in Chapter 6. Issue 32 brings a further cluster of major model releases (NVIDIA Cosmos 3 — omnimodal world model, released 31 May 2026, rated best open-source T2I and I2V by Artificial Analysis) and infrastructure investments (Roblox acquires Morpheus AI for video world models; Tripo AI raises ~\$200M for AI 3D and world model tech / Project Eden) alongside the policy and industry datapoints recorded in the table above. The Issue 32 financial headline — Suno’s \$5.4B valuation alongside active UMG and Sony litigation — is the cleanest single

data point on the financial reality of the creative-AI market in mid-2026: the licensing fights and the capital markets are running in parallel, not sequentially.

A8. June–July 2026 supplemental: the Issue-33 to Issue-36 and World Cup datapoints

Issue 33 (9 June 2026), **Issue 34** (16 June 2026), **Issue 35** (23 June 2026) and **Issue 36** (2 July 2026) extend the post-cut record established in §A7, alongside the World Cup deep dive of 12 June 2026. Issue 33 catches the **AFM v. UMG/Warner suit and the UK’s AI Adoption Plan for the Creative Industries**; Issue 34 catches the **licensing-and-labour-market turn** (NMPA–Udio/Klay, Runway–Lionsgate equity, the PwC salary-premium finding, the Steam disclosure count) and the **World Cup Editor’s Pick**; Issue 35 catches the **Unreal Engine AI/MCP announcements, the NO FAKES letter and the A24–Google deal**; Issue 36 catches the **ComfyUI/Comfy MCP inflection, the Adobe–Topaz acquisition and the games-volume data**. The numerical datapoints worth recording here in standalone form:

Datapoint	Value	Source
Salary premium for AI skills across UK sectors (PwC)	34%	UKTN / PwC, June 2026
Monthly AI spend per employee, top 1% of firms (Ramp AI Index)	\$7,500	Ramp, June 2026
Steam releases in one week carrying AI disclosures	120 of 300+	PC Gamer, June 2026
Mobile games released in six months (AI and vibe coding)	181,000	Digital Trends, July 2026
Signatures on SAG-AFTRA letter demanding the NO FAKES Act	16,000+	Variety, June 2026
NMPA–Udio/Klay licensing split for songs and recordings	50/50	Complete Music Update, June 2026
Lionsgate annual AI savings claimed (Michael Burns), alongside an equity stake in Runway	“tens and tens of millions”	Deadline / Runway, June 2026
Claimed animation cost reduction with AI (filmmakers)	90% less	Bloomberg, June 2026
ComfyUI users	4M+	Dream Machine Issue 36
ComfyUI GitHub stars / community-built nodes	89,000+ / 60,000+	Dream Machine Issue 36
ComfyUI April 2026 round / valuation / total raised	\$30M / \$500M / \$48M	Dream Machine Issue 36
Odyssey Series B (world models)	\$310M	Odyssey, June 2026
Higgsfield revenue (eyeing a \$5B funding round)	\$500M	TechTimes, June 2026
General Intuition raise / valuation (gameplay-trained frontier models)	\$320M / \$2.3B	GamesBeat / TechCrunch, June 2026
World Cup players scanned for photoreal digital twins / scan duration	1,248 / 1 second each	FIFA / Lenovo, June 2026
World Cup teams / matches / projected viewers	48 / 104 / 6bn	FIFA / Lenovo, June 2026

Datapoint	Value	Source
Projected incremental global ad spend, Q2 2026, on the tournament (WARC)	\$10.5bn	WARC, June 2026
FIFA Roblox presence: total visits / daily gameplay sessions	1.1bn+ / 1.5M	Gamefam / Roblox, June 2026
Fake FIFA-related websites identified before kick-off	4,300+	Security researchers, June 2026
Host-city sponsorship investment in the year before the tournament	\$725M	SponsorUnited / Goal.com, June 2026
Adidas “Backyard Legends” campaign (AI de-aging of three legends)	~\$1bn	Mission Media, June 2026
US programmatic ad spend, 2026 (machine-bought)	\$203bn+	<i>Dream Machine</i> World Cup analysis, June 2026
Streaming share of all US TV viewing	47.5%	<i>Dream Machine</i> World Cup analysis, June 2026

Reading note. Two clusters carry the weight here. The first is the **labour-and-licensing cluster**: PwC’s 34% salary premium and Ramp’s \$7,500-per-employee-per-month top-percentile spend are the two cleanest signals that AI capability had become a priced labour-market asset by mid-2026, while the NMPA’s 50/50 split with Udio and Klay — set against the AFM’s simultaneous suit over the UMG and Warner deals — shows the licensing architecture and the labour resistance maturing in the same fortnight. The second is the **scale cluster**: 120 AI disclosures in a single week of Steam releases and 181,000 mobile games in six months are the games-industry version of the Deezer 44% figure in [Chapter 5](#) — volume without gatekeeping — while the World Cup rows record the largest live deployment of the full AI stack in entertainment history, the empirical anchor for the storyworld argument of Chapter 9. The ComfyUI rows (4M users, 60,000+ community nodes, a \$500m valuation) close the arc that §A5’s tool table opens: the open node-graph layer, not any single model, is where the professional market consolidated.

A9. July–August 2026 supplemental: the Issue-37 to Issue-41 datapoints

Issue 37 (9 July 2026), **Issue 38** (16 July 2026), **Issue 39** (30 July 2026), **Issue 40** (5 August 2026) and **Issue 41** (12 August 2026) extend the post-cut record established in §A7 and §A8 through the end of the book’s source window. Issue 37 leads with the **HI-C / creative-friction Editor’s Pick** — the High-Impact Individual Contributor and the argument that capability is not the same as craft —

alongside a run of record financings; Issue 38 catches **George Lucas backing AI** (comparing sceptics to people clinging to horses after the car) and the *Star Wars*-inspired independent-AI-film wake-up call; Issue 39 catches **Netflix's disclosure that roughly 300 titles used generative AI in 2026** and the confirmed price of its Ben Affleck acquisition; Issue 40 catches the **"can AI write better stories?" study** and the biggest AI copyright payout in history; Issue 41 catches **Disney's J.A.R.V.I.S.** rollout and the generational-divide data on AI-music attitudes. The numerical datapoints worth recording here in standalone form:

Datapoint	Value	Source
Kling AI record funding raise	\$2.8bn	<i>TechTimes</i> , July 2026
Tripo AI follow-on raise, a month after its previous \$200M	+\$150m	<i>GamesBeat</i> , July 2026
Cambridge's Worldmodeldata raise (turning games into AI training data)	£7m	<i>tech.eu</i> , July 2026
Microsoft roles cut in one round (2.1% of global workforce)	~4,800	<i>CNBC</i> , July 2026
Netflix titles that used generative AI in 2026	~300	<i>Variety</i> , July 2026
Netflix price for Ben Affleck's AI filmmaking startup (InterPositive)	\$587m	<i>TechCrunch</i> , July 2026
Teen-founder raise for AI game-dev startup ForgeGUI	\$3.2m	<i>GamesBeat</i> , July 2026
Design Arena raise (bringing "taste" to AI models)	\$7.9m	<i>TechCrunch</i> , August 2026
Band's share of Spotify royalties lost to an AI speed-altered duplicate album	94%	<i>Digital Music News</i> , July 2026
Suno–GEMA ruling: the biggest AI-music copyright decision to date	loss	<i>Variety</i> , August 2026
AI's credited share of US economic growth	one third	<i>PYMNTS</i> , August 2026
Higgsfield Global Film Festival prize pool / first place	\$1m / \$500k	Higgsfield, August 2026
Disney J.A.R.V.I.S. rollout to Imagineers	2,000+	<i>WDWNT / BlogMickey</i> , August 2026
Gemini monthly active users	1bn+	<i>9to5Google</i> , August 2026
Gemini images generated per day / app actions it can take	150M+ / 40+	<i>9to5Google</i> , August 2026
GDC "State of the Game Industry": developers saying GenAI has a negative impact	52%	GDC, August 2026

Datapoint	Value	Source
(18% → 30% → 52% over three years)		
Luminate generational gap in AI-music interest: +19% net (13–17) vs –29% (18–24)	48 points	Luminate / <i>Musically</i> , August 2026

Reading note. Three clusters carry the weight here. The first is the **capital cluster**: Kling’s \$2.8bn, Tripo’s \$150m-on-top-of-\$200m, Design Arena’s \$7.9m, a teenager’s \$3.2m for ForgeGUI, Worldmodeldata’s £7m and Netflix’s confirmed \$587m for InterPositive are the closing entries in the investment story of Chapter 8, *The Machine Gets the Money* — money flowing into the tooling and world-model layers even as the labour market contracts (Microsoft’s ~4,800 roles, the games-industry job anxiety). The second is the **authority-and-authorship cluster**: the Suno–GEMA loss — the biggest AI-music copyright ruling to date — sits against the Guardian study in which readers rated AI-written stories above human ones, the tension the book works through in Chapter 13, while AI’s credited one-third of US economic growth is the macro backdrop to the whole argument. The third is the **generational-divide cluster**: the GDC negative-sentiment line climbing 18% → 30% → 52% and Luminate’s 48-point gap between 13–17s (+19%) and 18–24s (–29%) are the sharpest evidence yet for the split-audience thesis of Chapter 14, while Gemini’s billion monthly users, 150M+ images a day and 40+-app action reach — together with Disney’s J.A.R.V.I.S. putting seven decades of institutional knowledge in front of 2,000+ Imagineers — record how completely the assistant layer of Chapter 10 had normalised by the close of the source window.

* * *

What this appendix is for

Every chapter of this book is a *reading* of the corpus described above. It will be useful in 2030 and beyond to be able to see the underlying shape of the corpus, separate from the argument the book builds on top of it.

If you want to test the argument against your own reading of the same evidence: every URL in the corpus is enumerated in the citation index, every scraped article is preserved in JSON form in the `Research/scraped/` directory of the source repository, and every analysis in this appendix is reproducible by running `Research/quant.py`.

If you want to extend it: the scraper is in `Research/scrape.py`, the analyzer is in `Research/analyze.py`, the per-chapter dossiers are in `Research/dossier/`. Fork it, change it, run it on the next six months. I’d be glad to see what you find.

Appendix B — Glossary

The terms in this glossary are the working vocabulary of the book. Some are coinages, some are borrowed from elsewhere and reframed; all are used with precise, deliberately narrow meanings in the chapters above.

* * *

Agency line (or Continuum line). A single axis representing the share of decision-making in a creative function performed by a human versus a machine system. See **Human–AI Agency Continuum**.

Agentic AI. A class of AI system that, given a goal, can plan, decide and execute a sequence of multi-step actions without further human input between steps. Distinct from a *generator*, which produces an output in response to a single prompt. See Chapter 3.

AI literacy. The cluster of skills required to deploy AI tools effectively in creative work — including briefing, taste, judgement, prompt practice, output evaluation and tool-stack fluency. The term moved from optional to baseline competency through 2025–26, formalised by initiatives such as the Sundance Institute’s AI Literacy Initiative launched in January 2026.

AI slop. Low-quality, mass-produced AI-generated content that is recognisable to audiences as such — usually because it is made without human creative intent. *Merriam-Webster’s* word of the year for 2025. See **Slop Ceiling**.

Attribution. The principle that when AI systems produce derivative outputs based on training data, the human authors whose work shaped those outputs should be identified and — where appropriate — compensated. One of the four principles of the generative creative economy (Chapter 16). Technical infrastructure includes C2PA, SynthID, and creative-weight-attribution systems.

Audience contract. The implicit agreement between makers and audiences about what creative work is, what conditions of making it carries, and what relationship the audience can expect with its makers. The shift from implicit to explicit audience contracts is one of the central structural changes the book describes. See Chapter 13.

Augmented intelligence. Reframing of “AI” used by some industry voices to emphasise human-in-the-loop deployment over autonomy. Cited in *Dream Machine* Issue 9. Compare with **Generative AI**, **Agentic AI**.

C2PA. *Coalition for Content Provenance and Authenticity* — a technical standard for embedding cryptographic provenance metadata in media files, supported by camera manufacturers, editing software and a growing number of platforms. The principal “fingerprint real media” infrastructure underlying the authenticity argument in Chapter 13.

Coordination collapse. The structural change that occurs when the labour-coordination architecture of a creative organisation — built around the bandwidth constraints of human-only teams — is overtaken by AI-assisted workflows that no longer require those constraints. The subject of Chapter 14. Manifests as *shadow AI* below management sight and as compressed middle layers in the workforce above.

Dead Internet Theory. The notion that most of the public web is now bot-generated and bot-read, with humans increasingly a minority of traffic. Once a conspiracy framing; by 2025 a measurable phenomenon — bots accounted for 51% of web traffic in the Imperva 2025 Bad Bot Report, of which ~80% of bot traffic was AI training crawlers. See Chapter 4.

Disclosure. The practice of declaring the use of AI in the production of a piece of creative work — in credits, contracts, metadata, watermarks, or platform-facing labels. By spring 2026, disclosure had emerged as the dominant industry response to the audience-authenticity question, anchored by standards including the **Cannes AI Disclosure Standard** (May 2026) and the **Academy of Motion Picture Arts and Sciences** rule requiring human authorship for awards eligibility.

Duration arbitrage. The reallocation of creative capital toward formats whose production-and-payback cycles are shorter than the uncertainty horizon — the vertical microdrama being the canonical 2025–26 case. Includes the emergent use of short-cycle formats as priced discovery for longer-form commitments. See Chapter 8.

Duration discount, the. The additional discount capital now applies to creative assets in proportion to their production and payback horizons, driven by Knightian uncertainty about the terminal market rather than by any specific forecast of AI substitution. Sector-variable: deepest in games, machinery-mediated in film, expressed as format migration in television, and sign-flipped in immersive. See Chapter 8, **Seven-Year Problem**.

Extractive economy. A creative economy in which AI systems are trained on unpaid human work, the platform companies that build the models capture most of the resulting economic value, and the diffuse pool of working creatives is steadily decapitalised. One of two possible end-states the book identifies. See Chapter 16.

Fingerprint real media. Adam Mosseri’s (Instagram) framing of the verification problem: amplify provably-human content rather than chase synthetic content for labelling. Used in the book as shorthand for *provenance-first* approaches to content moderation. See Chapter 4.

Generative economy. A creative economy in which AI tools are treated as new craft infrastructure, training data is consented and compensated, platforms compete on tool quality and integrity, and the productivity gains are broadly distributed rather than concentrated. The opposite of the **extractive economy**. The four

principles (*Agency, Attribution, Access, Audience*) are the operational test. See Chapter 16.

Human–AI Agency Continuum. A frame, introduced in *Dream Machine Issue 2* (October 2025) and extended in Chapter 3, in which any given creative function is mapped on a horizontal line from full human agency (left) to full machine agency (right). The frame’s key claim is that *each creative function moves at its own speed* — you can sit at the extreme left on performance while being at the right edge on plate generation.

Living Web. The deliberately-built portion of the public web in which authorship is provable, attribution is durable, attention is allocated on non-viral signals, and the architecture supports rather than undermines human creative work. The aspirational opposite of the **Dead Internet**. Has to be built, not assumed. See Chapter 4.

Mid-career squeeze. The structural pressure on workers in the middle of creative-industry careers — neither at the junior entry level (replaced by agents) nor at the senior decision-making level (still required) — as AI absorbs the intermediate production roles that those mid-career workers historically held. See Chapter 14.

Model collapse. The technical risk that AI systems trained predominantly on synthetic data — including data produced by earlier generations of AI systems — progressively lose touch with real-world signal and produce increasingly homogenised, hallucination-prone outputs. The risk that *Dead Internet, Living Web* warns has moved from theoretical to measurable.

Orchestrator. The role that emerges when an individual working creative — or a small team — directs a large pool of AI-agent capacity. Defined operationally by five activities: defining the brief, allocating work, briefing the agents, judging outputs, and integrating the result. Predicted in *Dream Machine Issue 13* (January 2026) as the dominant role of 2026; documented across the chapters above. See Chapter 12.

Pipeline of authorship. The full chain from creative intent to delivered work, broken down into discrete functions (writing, direction, performance, image-making, sound, edit, distribution, marketing). The point of the **Human–AI Agency Continuum** is that each link in this chain has its own agency line.

Position One (All-in). The strategic posture of legacy studios that have decided to integrate AI aggressively across all production functions, betting that early-adopter advantage will compound. Netflix’s “all in” framing, October 2025. See Chapter 7.

Position Two (AI-native). The strategic posture of new entrants that build their production pipelines AI-first from inception — Imaginae Studios, Wonder Studios, Obsidian Studio, Asteria, Wonder, Chapter41, Kartel. The category exploded in scale through autumn 2025 and winter 2026. See Chapter 7.

Position Three (Refusal). The strategic posture of creative organisations that have explicitly excluded generative AI from their work. Jagex, Larian, Games Workshop, Hooded Horse, Aardman (qualified), Pocketpair. Cultural authority is preserved as the principal asset. See Chapter 7.

Position Four (Middle). *AI in the workflow, not in the work.* The strategic posture — taken by Sony, Bethesda, Amazon (in *House of David*), Aardman (in qualified form), and an increasing number of major studios — that uses AI to augment production pipelines while preserving human creative intent in the moments the audience sees. The book’s prediction for where most surviving major studios land by 2030. See Chapter 7.

Provenance. The chain of custody of a creative work from its origin (capture, performance, writing, sketching) to its delivered form. Technical standards (C2PA, SynthID) and policy frameworks (Cannes Disclosure Standard, SAG-AFTRA AI protections) collectively constitute the *provenance infrastructure* the book argues is critical for the next decade. See Chapter 13.

Seven-Year Problem, the. The collision between creative assets that take half a decade or more to build and a technological regime whose five-year shape no serious actor claims to know. Named for the production cycle of a AAA game; applies with sector-specific force across film, television and venue-scale immersive. Distinct from ordinary hit risk (“nobody knows anything”), which concerns demand within a stable game rather than the stability of the game itself. See Chapter 8, **Duration discount**.

Shadow AI. The use of AI tools by employees outside their employer’s official tooling, processes and accounting. Documented in 2025 workplace research as encompassing approximately half of the U.S. workforce. The principal symptom of **Coordination Collapse**. See Chapter 14.

Slop Ceiling. The empirical pattern, established across multiple sectors by spring 2026, in which AI-generated creative content can be produced in massive volume but consistently fails to capture audience attention proportionate to that volume. Anchored in the **44%/3%** ratio Deezer published in April 2026 (44% of daily uploads AI; under 3% of streams). One of the central claims of the book. See Chapter 5.

Stranded slate. A portfolio of in-production, long-cycle creative assets held at valuations that assume a pre-transition market structure; by analogy with stranded assets in climate finance. Exists in institutional behaviour — bond seizure, presale hesitancy, shortened fund horizons — before it appears in institutional language. See Chapter 8.

Synthetic sincerity. The category of creative work that is openly synthetic but made with serious creative intent and is not pretending to be something else. Named

after Marc Isaacs’ 2025 IDFA documentary. Audiences distinguish *synthetic sincerity* from *synthetic cynicism* at the speed of a swipe. See Chapter 4.

Tilly Tax (informal). The collection of contract provisions in SAG-AFTRA’s spring 2026 agreement requiring compensation, consent and residuals when AI replicas of human performers are used. Named after the Tilly Norwood controversy of September 2025 that catalysed the broader negotiation. See Chapters 5, 11.

Tragedy of the near horizon, the. The inversion, for the creative industries, of Mark Carney’s “tragedy of the horizon”: the disruption sits inside every current investment horizon but its shape is unreadable, contaminating all long-dated cash flows rather than escaping the model entirely. Strictly worse for underwriting than a distant catastrophe, which can at least be discounted. See Chapter 8.

Watermark. A persistent identifier embedded in AI-generated outputs by the producing system, intended to allow downstream detection that content is synthetic. SynthID (Google) and similar systems became standard across major platform tools through 2025–26.

World model. A class of generative AI system that produces navigable three-dimensional environments rather than flat output. Marble (World Labs, public release November 2025) was the first commercial product in the category; Google DeepMind’s Genie 3, Meta’s WorldGen, Luma’s UNI-1, Tencent’s Hunyuan World, SpAItial’s ECHO and others followed within months. See Chapter 9.

* * *

If a term in the book did substantial work but does not appear in this glossary, please tell us so we can improve it in the next edition.

Appendix C — Bibliography by Topic

The book's footnotes are sequential per chapter; this appendix organises the same sources thematically, so a reader pursuing a specific question across the period covered by the book can find every source it touches in one place.

The full corpus — 1,438 successfully fetched and archived articles — is preserved in JSON form in `Research/scraped/` of the source repository; the manifest enumerating every URL and its capture status lives at `Research/manifest.json`.

* * *

I. The watershed week, September–October 2025

Sora 2 launches. Tilly Norwood is announced. The union response sets the contract reference point for the next year.

- OpenAI, “Sora 2 is here,” 30 September 2025. <https://openai.com/index/sora-2/>
- Variety, “SAG-AFTRA Condemns Tilly Norwood: AI Actress Is Not an Actor.” <https://variety.com/2025/film/news/sag-aftra-tilly-norwood-ai-actress-1236534779/>
- *The Hollywood Reporter*, “U.K. Union Equity Condemns Tilly Norwood: ‘AI Tool, Not a Performer’.” <https://www.hollywoodreporter.com/movies/movie-news/tilly-norwood-ai-actress-uk-union-equity-sag-aftra-debate-1236391739/>
- Variety, “Tilly Norwood Slammed by Equity as AI Tool, Concerned About Origin.” <https://variety.com/2025/film/global/tilly-norwood-slammed-equity-ai-tool-concerned-origin-1236537042/>
- CNN, “Tilly Norwood: Hollywood is fuming over a new ‘AI actress’.” <https://www.cnn.com/2025/09/30/tech/hollywood-ai-actor-backlash>
- NBC News, “Tilly Norwood, fully AI ‘actor,’ blasted by actors union SAG-AFTRA.” <https://www.nbcnews.com/pop-culture/pop-culture-news/tilly-norwood-fully-ai-actor-blasted-actors-union-sag-aftra-devaluing-rcna234685>
- *The Guardian*, “OpenAI launch of video app Sora plagued by violent and racist images.” <https://www.theguardian.com/us-news/2025/oct/04/openai-sora-violence-racism>
- WUFT (Florida), “Kiss reality goodbye: AI-generated social media has arrived.” <https://www.wuft.org/2025-10-03/kiss-reality-goodbye-ai-generated-social-media-has-arrived>
- *Deadline*, “Tilly Norwood Creator on Backlash, Says More AI Actors Coming.” <https://deadline.com/2025/11/tilly-norwood-creator-interview-backlash-more-ai-actors-coming-1236601334/>
- *No Film School*, “James Cameron Says AI Is ‘Never Going to Take the Place’ of Humans.” <https://nofilmschool.com/james-cameron-ai>

- *The Guardian*, “James Cameron says AI actors are ‘horrrifying to me’.” <https://www.theguardian.com/film/2025/dec/01/james-cameron-says-ai-actors-are-horrifying-to-me>

II. Agentic AI and the orchestrator economy

- OpenAI, “Introducing AgentKit.” <https://openai.com/index/introducing-agentkit/>
- TechCrunch, “OpenAI launches AgentKit to help developers build and ship AI agents.” <https://techcrunch.com/2025/10/06/openai-launches-agentkit-to-help-developers-build-and-ship-ai-agents/>
- TechCrunch, “Anthropic launches interactive Claude apps, including Slack and other workplace tools.” <https://techcrunch.com/2026/01/26/anthropic-launches-interactive-claude-apps-including-slack-and-other-workplace-tools/>
- GamesRadar, “EA reportedly pushes 15,000 employees to use AI as a ‘thought partner’.” <https://www.gamesradar.com/games/even-under-usd20-million-in-debt-ea-reportedly-pushes-15-000-employees-to-use-ai-as-a-thought-partner-for-everything-from-character-art-to-playtesting/>
- Digiday, “AI agent developers have become adland’s in-demand role.” <https://digiday.com/marketing/ai-agent-developers-have-become-adlands-in-demand-role/>
- Eurogamer, “Falcom: ‘Tasks that previously took 2–3 hours can now be completed in 10 minutes’.” <https://www.eurogamer.net/falcom-is-the-latest-developer-to-buy-into-the-ai-hype-machine>
- Forbes, “AI Is Changing How Creators Work And Earn.” <https://www.forbes.com/sites/kolawolesamueladebayo/2025/12/22/how-ai-is-changing-how-creators-work-and-earn/>
- Google DeepMind, “Introducing Gemini Omni — unified multimodal model.” <https://blog.google/technology/google-deepmind/gemini-omni/>
- Google, “Gemini Spark — developer toolkit for autonomous agents.” <https://blog.google/technology/developers/gemini-spark/>
- Google, “Antigravity — agentic coding environment.” <https://antigravity.google/>
- Google, “Google Flow — agent-based creative workflows.” <https://flow.google/>
- Google, “Official skills for AI agents.” <https://github.com/google/agent-skills>
- Adweek, “Netflix ad tools could see ‘agentic AIs talking to each other’.” <https://www.adweek.com/media/netflix-ad-tools-agentic-ais-talking-to-each-other/>
- Fortune, “AI startup Viktor raises \$75 million to put a virtual ‘coworker’ in Slack and Teams.” <https://fortune.com/2026/05/19/ai-startup-viktor-75->

[million-virtual-coworker-slack-teams/](#)

- Tencent, “Ardot — AI-native design agent platform.”
<https://ardot.tencent.com/>
- Anthropic, “Claude available as a partner node in ComfyUI.”
<https://www.anthropic.com/news/claude-comfyui-partner-node>
- Runway, “Aleph 2.0 — single-frame video editing propagated across the clip” (June 2026). *Dream Machine* Issue 31.
- Alibaba Cloud, “WonderClip — all-in-one AI video creation platform” (June 2026). *Dream Machine* Issue 31.
- Sony, “Vid-CamEdit — post-capture camera trajectory editing” (June 2026). *Dream Machine* Issue 31.
- Nvidia, “PiD (Pixel Diffusion Decoder)” (June 2026). *Dream Machine* Issue 31.
- TripoAI, “Project Eden — persistent multiplayer world model decoupling world state from rendering” (June 2026). *Dream Machine* Issue 31.
- Sorcerox, “Game Creation Suite — build any game from a prompt” (June 2026). *Dream Machine* Issue 31.

III. Bot economy, model collapse, and the Dead Internet

- Imperva, “2025 Bad Bot Report.” <https://www.imperva.com/blog/2025-imperva-bad-bot-report-how-ai-is-supercharging-the-bot-threat/>
- Cloudflare, “The crawl-to-click gap: Cloudflare data on AI bots.”
<https://blog.cloudflare.com/crawlers-click-ai-bots-training/>
- Cloudflare, “A deeper look at AI crawlers.” <https://blog.cloudflare.com/ai-crawler-traffic-by-purpose-and-industry/>
- Grand View Research, “Generative AI in Content Creation Market Size Report, 2030.” <https://www.grandviewresearch.com/industry-analysis/generative-ai-content-creation-market-report>
- Wikipedia, “Dead Internet Theory.”
https://en.wikipedia.org/wiki/Dead_Internet_theory
- Futurism, “Scientists Created an Entire Social Network Where Every User Is a Bot.” <https://futurism.com/social-network-ai-intervention-echo-chamber>
- AI News, “AI causes reduction in users’ brain activity, MIT.”
<https://www.artificialintelligence-news.com/news/ai-causes-reduction-in-users-brain-activity-mit/>

IV. The slop ceiling — music

- Deezer, “AI-generated tracks now represent 44% of all new uploaded music.”
<https://newsroom-deezer.com/2026/04/ai-generated-tracks-represent-44-of-new-uploaded-music/>

- *Music Business Worldwide*, “75,000 AI-generated tracks now flood Deezer daily.” <https://www.musicbusinessworldwide.com/75000-ai-generated-tracks-now-flood-deezer-daily-representing-44-of-all-new-music-uploaded-to-the-platform-says-streamer/>
- *Deezer / Ipsos*, “AI fools 97% of listeners.” <https://newsroom-deezer.com/2025/11/deezer-ipsos-survey-ai-music/>
- *Musically*, “50,000 AI-music tracks are now uploaded to Deezer every day.” <https://musically.com/2025/11/12/50000-ai-music-tracks-are-now-uploaded-to-deezer-every-day/>
- *Musically*, “UMG boss slams ‘exponential growth of AI slop’ on streaming services.” <https://musically.com/2026/01/09/umg-boss-slams-exponential-growth-of-ai-slop-on-streaming-services/>
- *Billboard*, “AI Artist Xania Monet Climbs the Charts — And Signs a Multimillion-Dollar Record Deal.” <https://www.billboard.com/pro/ai-music-artist-xania-monet-multimillion-dollar-record-deal/>
- *NPR*, “Breaking Rust is a hot new country act on the Billboard charts. It’s powered by AI.” <https://www.npr.org/2025/11/10/nx-s1-5604320/breaking-rust-is-a-hot-new-country-act-on-the-billboard-charts-its-powered-by-ai>
- *BBC News*, “Sienna Rose: AI suspicions surround mysterious singer.” <https://www.bbc.co.uk/news/articles/cq6v83gq66eo>
- *Bain & Company*, “In an AI Age, People Still Want the Radio Star.” <https://www.bain.com/insights/in-an-ai-age-people-still-want-the-radio-star/>
- *Soultracks*, “A.I.-generated music is catchy, familiar... and boring.” <https://soultracks.com/news-ai-generated-music-is-catchy-boring/>
- *Stereogum*, “Bandcamp Bans AI Music.” <https://stereogum.com/2485199/bandcamp-bans-ai-music/news>
- *The Independent*, “AI-generated song banned from Swedish charts.” <https://www.independent.co.uk/tv/news/ai-music-song-banned-sweden-spotify-b2901627.html>
- *Digital Music News*, “YouTube CEO Puts ‘Managing AI Slop’ on the Priority List for 2026.” <https://www.digitalmusicnews.com/2026/01/22/youtube-ceo-ai-slop-2026-comments/>
- *Musically*, “Report: 56.9% of new independent songs in China are AI-generated.” <https://musically.com/2026/01/05/report-56-9-of-new-independent-songs-in-china-are-ai-generated/>
- *The Guardian*, “Paul McCartney joins music industry protest against AI with silent track.” <https://www.theguardian.com/music/2025/nov/17/the-sound-of-silence-why-theres-barely-anything-there-in-paul-mccartney-new-release>
- *MusicTech*, “Jack Antonoff brands AI music makers as ‘godless whores’.” <https://musictech.com/news/industry/jack-antonoff-ai-music-makers-godless->

whores/

- *Billboard*, “The Real Story Behind The AI Song That Knocked Tyla Off No. 1 On Billboard Afrobeats Chart.” <https://www.billboard.com/pro/ai-song-knocked-tyla-off-no-1-afrobeats/>
- *Variety*, “WMG’s Robert Kyncl: ‘AI Resistance’ Setting Music Sector Back.” <https://variety.com/2026/music/news/wmg-robert-kyncl-ai-resistance-1236748901/>
- *Musically*, “Sondo AI says it has 10m global users.” <https://musically.com/2026/05/18/sondo-ai-10m-global-users/>
- *Musically*, “BPI sets out transparency and sovereignty demands to secure AI licensing boom.” <https://musically.com/2026/05/19/bpi-transparency-sovereignty-ai-licensing-boom/>
- *Musicbusinessworldwide*, “DistroKid asks creators if their music is AI-generated as self-disclosure moves out of beta.” <https://www.musicbusinessworldwide.com/distrokid-ai-self-disclosure-out-of-beta/>
- *Musicbusinessworldwide*, “Splice inks ‘Responsible AI’ deal with ElevenLabs.” <https://www.musicbusinessworldwide.com/splice-elevenlabs-responsible-ai-deal/>
- *Rolling Stone*, “The Rolling Stones Release ‘In the Stars’ — AI De-Aging Music Video.” <https://www.rollingstone.com/music/music-news/rolling-stones-in-the-stars-ai-de-aging-video-1235142200/>
- *MusicTech*, “Tamber — ethically trained AI music tool with arm gesture control.” <https://musictech.com/news/gear/tamber-ai-ethically-trained-arm-gestures/>
- *Stability AI*, “Stable Audio 3.0 released.” <https://stability.ai/news/stable-audio-3-0-released>
- *Sony AI*, “Woosh — sound effect foundation model.” <https://ai.sony/blog/woosh-sound-effect-foundation-model/>
- *Mirelo*, “SFX 1.6 — edit sound, not just generate.” <https://mirelo.ai/sfx-1-6>
- *ElevenLabs*, “Music v2 — next-generation generative music model” (June 2026). *Dream Machine* Issue 31.
- *Apple Music*, “AI-generated songs make up less than 1% of plays on platform” (June 2026). *Dream Machine* Issue 31.
- *Spotify + Universal Music Group*, “Landmark deal lets Spotify Premium subscribers create AI covers and remixes of UMG-catalogue tracks” (June 2026). *Dream Machine* Issue 31.
- *Daniel Ek (Spotify)*, “AI music is better than ‘slop’” — interview / statement (June 2026). *Dream Machine* Issue 31.

- Tonada, “*Swedish startup supplies AI music for retailers*” (June 2026). *Dream Machine* Issue 31.
- Estate of Ozzy Osbourne, “*Ozzy to live on as an AI avatar that talks with fans*” (June 2026). *Dream Machine* Issue 31.
- *Startup Fortune*, “*An Atlantic Investigation Just Blew Open the AI Music Industry’s Data Provenance Problem*” — the datasets that may have trained the major AI-music platforms. <https://startupfortune.com/an-atlantic-investigation-just-blew-open-the-ai-music-industrys-data-provenance-problem/>
- *Musically*, “*Luminate study explores Americans’ interest in GenAI content*” — a 48-percentage-point generational gap between 13–17-year-olds (+19% net interest in AI music) and 18–24-year-olds (–29%). <https://musically.com/2026/07/27/luminate-study-explores-americans-interest-in-genai-content/>

V. Copyright, consultation and the legal architecture

- UK Department for Science, Innovation and Technology, “*Statement of Progress on Copyright and AI.*” <https://www.gov.uk/government/publications/copyright-and-artificial-intelligence-progress-report/copyright-and-artificial-intelligence-statement-of-progress-under-section-137-data-use-and-access-act>
- IPWatchdog, “*Respondents to UK AI Consultation Overwhelmingly Want AI Companies to License Copyrighted Works in All Cases.*” <https://ipwatchdog.com/2025/12/16/respondents-uk-ai-consultation-overwhelmingly-want-ai-companies-license-copyrighted-works-all-cases/>
- UCL Copyright Queries, “*UK government publishes progress statement on AI and copyright consultation.*” <https://blogs.ucl.ac.uk/copyright/2025/12/23/uk-government-publishes-progress-statement-on-ai-and-copyright-consultation/>
- Hogan Lovells, “*Copyright and AI: UK government publishes statement of progress.*” <https://www.hoganlovells.com/en/publications/copyright-and-ai-uk-government-publishes-statement-of-progress>
- *Music Business Worldwide*, “*Wixen files \$50m copyright suit against Meta.*” <https://www.musicbusinessworldwide.com/wixen-files-50m-copyright-suit-against-meta-claims-tech-giant-wants-to-replace-songwriters-with-ai/>
- *Music Business Worldwide*, “*Meta hits back against Wixen Music Publishing’s \$70m ‘AI replacement’ infringement suit.*” <https://www.musicbusinessworldwide.com/meta-wixen-music-publishing-70-million-ai-replacement-lawsuit/>
- *EDM.com*, “*‘Biggest Theft in Music History’: Rights Group Sues Suno.*” <https://edm.com/gear-tech/rights-group-sues-suno-copyright-infringement/>

- *Complete Music Update*, “Johnny Cash estate uses ELVIS Act to sue Coke.” <https://completemusicupdate.com/johnny-cash-estate-uses-elvis-act-to-sue-coke-over-tribute-act-ad-soundtrack/>
- Dr. Barry Scannell, analysis of GEMA v. OpenAI ruling. https://www.linkedin.com/posts/dr-barry-scannell-bbb5aa207_in-a-major-ruling-for-european-copyright-share-7393957246386323457-8bbx
- Reuters, “European lawmakers seek EU-wide minimum age to access AI chatbots.” <https://www.reuters.com/legal/litigation/european-lawmakers-seek-eu-wide-minimum-age-access-ai-chatbots-social-media-2025-11-26/>
- *The Verge*, “New York’s new law forces advertisers to say when they’re using AI avatars.” <https://www.theverge.com/news/842848/new-york-law-ai-advertisements-sag-aftra-labor>
- *Digital Music News*, “Nearly 800 Creatives Sign Responsible AI Declaration — ‘Stealing Our Work Is Not Innovation’.” <https://www.digitalmusicnews.com/2026/01/22/stealing-isnt-innovation/>
- *Complete Music Update*, “US musicians’ union sues Universal and Warner over AI deals, claims breach of ‘new use’ clause in record industry’s performer agreement.” <https://completemusicupdate.com/us-musicians-union-sues-universal-and-warner-over-ai-deals-claims-breach-of-new-use-clause-in-record-industrys-performer-agreement/>
- *Complete Music Update*, “NMPA unveils AI licensing deals with Udio and Klay with 50/50 split for songs and recordings.” <https://completemusicupdate.com/nmpa-unveils-ai-licensing-deals-with-udio-and-klay-with-50-50-split-for-songs-and-recordings/>
- *Digital Music News*, “Google Doubles Down on ‘Fair Use’ AI Training, Stressing That ‘Outputs’ Not ‘Inputs’ Are Where Enforcement Should Be Focusing.” <https://www.digitalmusicnews.com/2026/06/29/google-fair-use-ai-training/>
- *Digital Music News*, “Federal Judge Vacates Order Sealing Udio’s Training Data.” <https://www.digitalmusicnews.com/2026/06/05/udio-training-number-order-vacated/>
- *Variety*, “More Than 16,000 Sign SAG-AFTRA Letter Demanding Congress Pass NO FAKES Act, Which Would Ban Unauthorized AI Images and Videos.” <https://variety.com/2026/biz/news/sag-aftra-letter-congress-pass-no-fakes-act-1236783056/>
- *Variety*, “Suno Loses Landmark AI Copyright Lawsuit to German Music Rights Society GEMA.” <https://variety.com/2026/digital/news/suno-loses-ai-lawsuit-gema-1236825010/>
- *Musically*, “Key EU AI Act rules now being enforced — so how does that affect the music industry?” <https://musically.com/2026/08/04/key-eu-ai-act-rules->

now-being-enforced-so-how-does-that-affect-the-music-industry/ — see also AOL, “AI labels to be compulsory on authentic-looking content under EU rules.” <https://www.aol.co.uk/articles/ai-labels-compulsory-authentic-looking-094628000.html>

- Billboard, “ABBA’s Björn Ulvaeus Says We Might Be Doing AI Licensing Deals Wrong” — ‘tracing the output was always the wrong question’.
<https://www.billboard.com/pro/abba-bjorn-ulvaeus-music-ai-un-global-summit/>

VI. Studio strategy

- CNBC, “Netflix ‘all in’ on leveraging AI.”
<https://www.cnbc.com/2025/10/22/netflix-all-in-on-leveraging-ai-in-its-streaming-platform.html>
- Futurism, “Lionsgate’s Attempt to Create Movies Using AI Has Crumbled Into Disaster.” <https://futurism.com/artificial-intelligence/lionsgate-movies-ai>
- The Guardian, “Disney to invest \$1bn in OpenAI.”
<https://www.theguardian.com/business/2025/dec/11/disney-open-ai-sora-video-deal>
- The Hollywood Reporter, “Disney+ to Allow User-Generated Content Via AI.”
<https://www.hollywoodreporter.com/business/digital/disney-plus-gen-ai-user-generated-content-1236426135/>
- Deadline, “Amazon’s Head Of Original Sports Content Matt Newman Moves Into A.I.” <https://deadline.com/2025/11/amazon-ai-studios-matt-newman-1236603477/>
- Wired, “Amazon’s House of David Used Over 350 AI Shots in Season 2.”
<https://www.wired.com/story/amazons-house-of-david-used-over-350-ai-shots-in-season-2-its-creator-isnt-sorry/>
- Indiewire, “New AI Studio Obsidian Forms Partnership with Imagine Entertainment.” <https://www.indiewire.com/news/business/obsidian-studio-ai-production-company-imagine-entertainment-1235158619/>
- The Hollywood Reporter, “Fremantle Names Boss of New AI Label Imaginae Studios.” <https://www.hollywoodreporter.com/business/digital/fremantle-names-ceo-new-ai-label-imaginae-studios-1236396579/>
- PC Gamer, “Palworld studio Pocketpair: ‘We don’t believe in it’.”
<https://www.pcgamer.com/software/ai/palworld-studio-pocketpair-says-its-new-publishing-division-wont-handle-games-that-use-generative-ai-we-dont-believe-in-it/>
- Variety, “Guillermo del Toro: ‘I’d rather die’ than use generative AI.”
<https://variety.com/2025/film/news/guillermo-del-toro-rather-die-generative-ai-frankensteins-1236561316/>

- *The Hollywood Reporter*, “Leonardo DiCaprio Says AI Can’t Be Art.”
<https://www.hollywoodreporter.com/movies/movie-news/leonardo-dicaprio-ai-cant-be-art-no-humanity-1236445405/>
- *NME*, “Jenna Ortega: ‘very easy to be terrified’ of AI.”
<https://www.nme.com/news/jenna-ortega-says-its-very-easy-to-be-terrified-of-ai-in-filmmaking-3913926>
- *GamesRadar*, “Wallace and Gromit creator: Aardman will ‘embrace the technology’ of AI but will be ‘very cautious’.”
<https://www.gamesradar.com/entertainment/animation-movies/wallace-and-gromit-creator-says-beloved-animation-studio-aardman-will-embrace-the-technology-of-ai-but-will-be-very-cautious-not-to-lose-our-values/>
- *gamesindustry.biz*, “‘AI was an expensive mistake’: Charles Cecil on Broken Sword.” <https://www.gamesindustry.biz/ai-was-an-expensive-mistake-charles-cecil-on-innovation-insolvency-and-broken-sword>
- *gamesindustry.biz*, “RuneScape maker Jagex says it will never use generative AI.” <https://www.gamesindustry.biz/runescape-maker-jagex-says-it-will-never-use-generative-ai-to-make-in-game-content>
- *Decrypt*, “Warhammer 40,000’ Maker Games Workshop Rules Out Generative AI.” <https://decrypt.co/354482/warhammer-40000-maker-games-workshop-rules-out-generative-ai>
- *Niche Gamer*, “Larian Studios backs off from gen AI.”
<https://nichegamer.com/larian-studios-backs-off-from-gen-ai/>
- *PC Gamer*, “Todd Howard says AI can’t replace human ‘creative intention’.”
<https://www.pcgamer.com/gaming-industry/todd-howard-says-ai-cant-replace-human-creative-intention-but-its-part-of-bethesdas-toolset-for-how-we-build-our-worlds-or-check-things/>
- *Hollywood Reporter*, “Netflix INKubator AI animation studio for ‘feature-quality’ shorts.” <https://www.hollywoodreporter.com/business/business-news/netflix-ai-animation-studio-inkubator-1236592110/>
- *Forbes*, “Wonder Studios — \$50M British studio aiming to become ‘A24 of AI’.”
<https://www.forbes.com/sites/charliefink/2026/05/18/meet-wonder-studios-the-50m-british-studio-striving-to-become-the-a24-of-ai-production/>
- *Runway*, “Runway Japan — Tokyo office and \$40M investment.”
<https://runwayml.com/blog/runway-japan>
- *Hollywood Reporter*, “Cannes 2026: filmmakers shift towards cautious AI acceptance.” <https://www.hollywoodreporter.com/business/business-news/cannes-2026-ai-acceptance-1236592488/>
- *Variety*, “AI Dominates Cannes Buzz as Filmmakers Grudgingly Accept It.”
<https://variety.com/2026/film/festivals/ai-cannes-2026-filmmakers-accept-1236748402/>

- *Variety*, “Peter Jackson: AI is basically the next wave of special effects.” <https://variety.com/2026/film/news/peter-jackson-ai-special-effects-1236748120/>
- *Variety*, “Kling AI + Evolutionary Films on ‘Minibots’ at Cannes Market.” <https://variety.com/2026/film/news/kling-ai-evolutionary-films-minibots-cannes-1236748590/>
- *Cartoon Brew*, “Former Disney Animator: Marvel layoffs point to a bigger AI shift.” <https://www.cartoonbrew.com/artist-rights/former-disney-animator-marvel-layoffs-ai-shift-260601.html>
- *Hollywood Reporter*, “Bobby Berk: AI will make reality TV and ‘verifiably human content’ more valuable.” <https://www.hollywoodreporter.com/tv/tv-news/bobby-berk-ai-reality-tv-1236592920/>
- *PC Gamer*, “Take-Two’s CEO: AI ‘datasets by their very nature are backward looking’.” <https://www.pcgamer.com/games/take-two-ceo-ai-not-making-hits-backward-looking/>
- *Runway*, “Runway and Lionsgate Expand Partnership” — Lionsgate takes an equity stake in Runway. <https://runwayml.com/news/runway-and-lionsgate-expand-partnership>
- *Deadline*, “Lionsgate’s Michael Burns Says AI Will Save The Company ‘Tens And Tens Of Millions Of Dollars A Year’.” <https://deadline.com/2026/06/lionsgate-michael-burns-ai-saves-millions-runway-1236941043/>
- *The Hollywood Reporter*, “Google Investing in ‘Backrooms’ Studio A24.” <https://www.hollywoodreporter.com/business/digital/a24-google-deepmind-ai-venture-backrooms-1236627228/>
- *Futurism*, “A24 Is Dealing With a Full-on Crisis Among Fans After Signing an AI Deal With Google.” <https://futurism.com/artificial-intelligence/a24-crisis-fans-signing-ai-deal-google>
- *Bloomberg*, “Animation’s AI Reckoning: Filmmakers Say They Can Better Make Movies for 90% Less.” <https://www.bloomberg.com/news/newsletters/2026-06-28/animation-s-ai-reckoning-filmmakers-say-they-can-make-movies-for-90-less>
- *Variety*, “About 300 Netflix Titles Used Generative AI This Year, Company Reveals.” <https://variety.com/2026/biz/news/about-300-netflix-programs-used-ai-this-year-q2-earnings-1236812914/>
- *Variety*, “George Lucas Says AI Makes Filmmaking ‘Easier’ and ‘There’s Nothing You Can Do About It’.” <https://variety.com/2026/film/news/george-lucas-test-screenings-embraces-ai-hollywood-1236810656/>
- *Variety*, “Brad Pitt Says AI Could Help More ‘Mid-Budget Films’ Get Made if It’s ‘Used as a Tool’.” <https://variety.com/2026/film/news/brad-pitt-ai-could->

help-mid-budget-movies-1236831201/

- BlogMickey, “Disney Imagineering rolls out its internal J.A.R.V.I.S. AI to 2,000+ Imagineers” — seven decades of accumulated creative knowledge, archives and workflows made searchable.
<https://blogmickey.com/2026/08/damaro-details-ai-use-imagineering-disney-company/>
- Wikipedia, “Leveraged buyout of Electronic Arts” — the \$55bn PIF / Silver Lake / Affinity Partners take-private (completed 4 August 2026), the largest leveraged buyout in history.
https://en.wikipedia.org/wiki/Leveraged_buyout_of_Electronic_Arts — see also CNBC, “Saudi wealth fund and Jared Kushner’s Affinity finalize \$55 billion EA deal.” <https://www.cnbc.com/2026/08/05/saudi-wealth-fund-jared-kushner-ea-sports-deal.html>

VII. World models and the new medium

- World Labs, “Bringing Marble to Life.” <https://www.worldlabs.ai/case-studies/bringing-marble-to-life>
- TechCrunch, “Fei-Fei Li’s World Labs speeds up the world model race with Marble.” <https://techcrunch.com/2025/11/12/fei-fei-lis-world-labs-speeds-up-the-world-model-race-with-marble-its-first-commercial-product/>
- Google DeepMind, “Genie 3: A new frontier for world models.” <https://deepmind.google/blog/genie-3-a-new-frontier-for-world-models/>
- Google blog, “Project Genie: AI world model now available for Ultra users in U.S.” <https://blog.google/innovation-and-ai/models-and-research/google-deepmind/project-genie/>
- Meta, “WorldGen — text-to-immersive-3D-worlds.” <https://www.facebook.com/LifeAtMeta/videos/research-update-worldgen-text-to-immersive-3d-worlds/1879077432692421/>
- SpAItial, “Transform Images into Interactive 3D Worlds.” <https://www.spaitial.ai/>
- Stanford AI Lab, “WonderZoom.” <https://wonderzoom.github.io/>
- NVIDIA/Stanford, “NitroGen.” <https://nitrogen.minedojo.org/>
- DeepMind, “SIMA 2: An Agent that Plays, Reasons, and Learns.” <https://deepmind.google/blog/sima-2-an-agent-that-plays-reasons-and-learns-with-you-in-virtual-3d-worlds/>
- Radiance Fields, “Apple Confirms Personas use Gaussian Splatting.” <https://radiancefields.com/apple-confirms-personas-use-gaussian-splatting>
- Sony Pictures’ virtual-production case study (Marble integration). https://www.linkedin.com/posts/brent-liang_tech-media-launch-ugcPost-7394911181091692546-TyUz

- NVIDIA, “SANA-WM — 2.6B open-source world model with 60-second video generation and camera control.”
<https://huggingface.co/collections/nvidia/sana-wm>
- Odyssey, “Starchild-1 — first real-time multimodal world model.”
<https://odyssey.ml/introducing-starchild-1>
- Odyssey, “Agora-1 — four-player AI-generated world built on a 1997 shooter.”
<https://odyssey.ml/introducing-agora-1>
- Google DeepMind, “Project Genie + Street View — simulate real-world places.”
<https://deepmind.google/discover/blog/project-genie-street-view/>
- Google Labs, “Infinite Scaler — browser multiplayer generative game.”
<https://blog.google/technology/google-labs/infinite-scaler/>
- Apple Machine Learning Research, “Apple Headsup — 3D Gaussian Head Reconstruction from Multi-View Captures.”
<https://machinelearning.apple.com/research/apple-headsup-3d-gaussian-head>
- WorldLens VR, “AI-powered 3D depth on Quest for Google Street View.”
<https://www.uploadvr.com/worldlens-vr-quest-street-view-3d-depth/>
- gamesindustry.biz, “Meet Seed, the planet-sized society simulator.”
<https://www.gamesindustry.biz/meet-seed-planet-sized-society-simulator>

VII.a Creator economy and labour-market sentiment

The June 2026 Epidemic Sound Future of the Creator Economy report, and the Coordination Collapse, Process Trap and Doomer Mistake deep dives that re-anchor the labour-market analysis of the second half of the book.*

- Epidemic Sound, “*Future of the Creator Economy Report 2026*” — survey of 3,000 professional creators across UK and US; 94% of creators use AI tools, 72% expect to increase usage in the next 12 months, 89% feel pressure to use AI to keep up with industry expectations, 75% believe human-created content will become a premium product, 83% say human-made sound creates a stronger emotional connection. *Dream Machine* Issue 31.
- *Musically*, “*Creator attitudes towards AI: nuanced embrace plus transparency, licensing and identity demands.*” Coverage of the same report. *Dream Machine* Issue 31.
- O’Hare, J. (2026). *The Coordination Collapse: How AI Is Dissolving the Information Tax That Created Hierarchy — And What Replaces It*. DreamLab AI Consulting Ltd, April 2026.
- *The Process Trap: Generative Abundance and the Equivocation of the Creative.*
- *The Doomer Mistake: Why the Case for AI-Driven Civilizational Collapse Keeps Losing the Argument with the Data.* Underlying empirical work draws on Acemoglu (2025), Brynjolfsson, Li & Raymond (2025), Brynjolfsson, Chandar &

Chen (2025), the 2026 Stanford AI Index, Dallas Fed, St. Louis Fed, Bessen (2015), Autor (2016), AEI, Susskind (2020), and CCIA (2026).

- *The Programmable Brand: How AI Is Reshaping IP Licensing in the Generative Economy*. Dream Machine Deep Dive N°14, 4 June 2026. Five-licensing-model comparison (Hasbro Sixth Wall / Behavioural Licensing; Spotify–UMG embedded remixing; Disney–OpenAI as cautionary tale; the independent-artist commons-and-DAO models — Grimes / Elf.tech, Holly Herndon / Holly+, ElevenLabs heritage voices, Meta likeness-rental anti-pattern; Suno / Alexa as the *create-to-consume* demand floor) against the three Cs (consent, credit, compensation). Theoretical foundations: Coase (1937), Williamson (1979), Barney (1991) VRIN resource-based view, Henry Jenkins on participatory culture. Primary sources:
 - Hasbro BusinessWire: <https://www.businesswire.com/news/home/20260603297922/en/>
 - *The Wrap*, Hasbro Sixth Wall launch: <https://www.thewrap.com/industry-news/business/hasbro-ai-studio-sixth-wall-launch/>
 - *Screen Daily*, Hasbro launches subsidiary AI studio: <https://www.screendaily.com/news/hasbro-launches-subsidiary-ai-studio-sixth-wall-to-exploit-ip/5217380.article>
 - *BriefGlance*, Sixth Wall and behavioural licensing: <https://briefglance.com/articles/hasbros-sixth-wall-taming-the-ai-wild-west-with-behavioral-licensing>
 - Spotify Newsroom, Spotify–UMG: <https://newsroom.spotify.com/2026-05-21/universal-music-group-spotify-licensing-agreements-fan-made-covers-remixes/>
 - *TechCrunch*, Spotify–UMG: <https://techcrunch.com/2026/05/21/spotify-and-universal-music-strike-deal-allowing-fan-made-ai-covers-and-remixes/>
 - *Medianama*, Spotify–UMG: <https://www.medianama.com/2026/05/223-spotify-umg-fans-create-licensed-ai-covers-remixes/>
 - *Ropes & Gray*, “The \$1 Billion Deal That Wasn’t”: <https://www.ropesgray.com/en/insights/viewpoints/the-deal-that-wasnt-what-disney-and-openai-teach-us-about-strategic-investments>
 - *TheStreet* / AOL on Disney exit: <https://www.thestreet.com/technology/openai-is-shutting-down-sora-and-the-disney-deal-is-off>
 - *Tech-Insider* on Sora shutdown economics: <https://tech-insider.org/openai-sora-shutdown-disney-deal-ai-video-2026/>
 - *Yahoo Finance* on Disney cease-and-desist to Google pre-OpenAI: <https://finance.yahoo.com/sectors/technology/articles/disney-cancels-1b-deal-openai-043000427.html>

- Goldman Sachs on creator economy:
<https://www.goldmansachs.com/insights/articles/the-creator-economy-could-approach-half-a-trillion-dollars-by-2027>
- *Kill the DJ* on Grimes / Elf.tech and artist-voice models:
<https://killthedj.com/how-artists-lend-their-ai-voices-to-creators-fans/>
- CBS News on ElevenLabs heritage voices:
<https://www.cbsnews.com/news/ai-voices-dead-celebrities-judy-garland-james-dean-burt-reynolds/>
- *El País* on Meta celebrity chatbots (2023):
<https://english.elpais.com/technology/2023-11-13/from-paris-hilton-to-snoop-dogg-meta-to-launch-speaking-ai-chatbots-played-by-celebrities.html>
- *Billboard* on Amazon + Suno: <https://www.billboard.com/pro/amazon-partners-ai-music-company-suno-alexa/>
- *Digital Music News* on Amazon + Suno:
<https://www.digitalmusicnews.com/2025/02/27/amazon-suno-team-up-for-gen-ai-songs/>
- MIDiA Research on *create to consume*:
<https://www.midiaresearch.com/blog/suno-alexa-and-the-rise-of-create-to-consume>
- Dell’Acqua et al. (2026). Preregistered Boston Consulting Group experiment on the jagged technological frontier (n = 758 consultants).
- Mollick, E. (2023). “*Secret cyborgs*” and the shadow-AI phenomenon.
- Edmondson, A. & Lei, Z. (2025). Psychological safety and AI adoption in teams.
- Epoch AI / Ipsos KnowledgePanel (March 2026). Probability-based survey of 2,021 US adults (665 employed past-week AI users).

VIII. The platform layer and the AI-native toolchain

- Creative Boom, “Adobe puts AI in everything, everywhere, all at once.”
<https://www.creativeboom.com/news/adobe-is-putting-ai-in-everything-everywhere-all-at-once/>
- Adobe, “*Adobe MAX 2025: Firefly Foundry*.”
<https://news.adobe.com/news/2025/10/adobe-max-2025-firefly-foundry>
- Adobe, “*Adobe MAX 2025: Express AI Assistant*.”
<https://news.adobe.com/news/2025/10/adobe-max-2025-express-ai-assistant>
- *Wired*, “*Adobe’s Corrective AI Can Change the Emotions of a Voice-Over*.”
<https://www.wired.com/story/adobe-max-sneaks-2025-corrective-ai/>
- PYMNTS, “*Adobe Lets Users Design and Edit Using ChatGPT*.”
<https://www.pymnts.com/artificial-intelligence-2/2025/adobe-lets-users->

[design-and-edit-using-chatgpt/](#)

- Adobe blog, “*Edit with Photoshop in ChatGPT.*”
<https://blog.adobe.com/en/publish/2025/12/10/edit-photoshop-chatgpt>
- Campaign Brief, “*WPP launches WPP Open Pro.*”
<https://campaignbrief.com/wpp-launches-ai-powered-marketing-platform-wpp-open-pro/>
- Digiday, “*WPP expands AI capabilities.*” <https://digiday.com/media-buying/agencies-continue-to-expand-ai-capabilities-to-boost-brand-performance/>
- SiliconAngle, “*Higgsfield raises \$80M on \$1.3B valuation.*”
<https://siliconangle.com/2026/01/15/higgsfield-raises-80m-1-3b-valuation-scale-ai-video-platform/>
- TechCrunch, “*Synthesia hits \$4B valuation.*”
<https://techcrunch.com/2026/01/26/synthesia-hits-4b-valuation-lets-employees-cash-in/>
- Sifted, “*Synthesia rejects \$3bn Adobe acquisition offer.*”
<https://sifted.eu/articles/synthesia-acquisition-offer>
- ComfyUI, “*We raised \$17 million to build an OS for creative AI.*”
https://www.linkedin.com/posts/comfyui_we-raised-17-million-to-build-an-os-for-ugcPost-7373743341236236288-wkCc
- Beeples Canvas, “*Generative AI meets compositing.*” <https://www.beeples-canvas.com/>
- MacRumors, “*New CarPlay Audio Apps Add AI Music and Global Radio.*”
<https://www.macrumors.com/2026/05/19/carplay-ai-music-global-radio/>
- The Drum, “*World’s First AI-Driven Immersive Shoppable Catwalk.*”
<https://www.thedrum.com/news/2026/05/18/how-world-s-first-ai-driven-immersive-shoppable-catwalk-redefined-future-brand>
- The Drum, “*David Beckham designs ‘Henchester United’ chicken coop in Lenovo AI ad.*” <https://www.thedrum.com/news/2026/05/18/david-beckham-henchester-united-chicken-coop-lenovo-ai-ad>
- Amazon, “*Alexa+ debuts feature to create AI-generated podcasts.*”
<https://www.aboutamazon.com/news/devices/alexa-plus-ai-podcasts>
- ECABridge, “*Unreal Engine MCP integration.*” <https://ecabridge.dev/>
- Video Games Chronicle, “*Epic Games Veteran on AI-Heavy ‘Fully European’ Game Engine.*” <https://www.videogameschronicle.com/news/epic-games-veteran-ai-heavy-fully-european-game-engine/>
- Andrew Zhu (xhinker), Medium, “*The Rise of ComfyUI: How One Guy’s Side Project Took Over AI Image Generation*” — 4M users, 89,000+ GitHub stars, 60,000+ community nodes, \$500M valuation.

<https://xhinker.medium.com/the-rise-of-comfyui-how-one-guys-side-project-took-over-ai-image-generation-558dca64bfd7>

- Comfy Org, “Comfy MCP” — public-beta Model Context Protocol server for production creative pipelines. <https://comfy.org/mcp/>
- Adobe, “Adobe announces its intent to acquire Topaz Labs.” <https://news.adobe.com/news/2026/06/adobe-to-acquire-topaz-labs>
- Unreal Engine, “State of Unreal 2026: Top News From the Show” — AI rendering integrations and Unreal MCP. <https://www.unrealengine.com/news/state-of-unreal-2026-top-news-from-the-show>

IX. Authenticity, provenance and disclosure

- Digital Music News, “Instagram Chief Says We Should ‘Fingerprint Real Media’.” <https://www.digitalmusicnews.com/2026/01/05/instagram-chief-ai-slop-comments/>
- WebProNews, “Instagram Head Warns AI Images Erode Trust.” <https://www.webpronews.com/instagram-head-warns-ai-images-erode-trust-calls-for-verification-standards/>
- Google DeepMind, “Verify AI-generated videos in the Gemini app.” https://www.linkedin.com/posts/googledeepmind_verify-google-ai-generated-videos-in-the-activity-7407748300688478208-fJgW
- The Hollywood Reporter, “Emmys Set AI Guidelines for 2026 Awards.” <https://www.hollywoodreporter.com/tv/tv-news/emmys-ai-guidelines-2026-awards-1236468434/>
- Lawyer Monthly, “Matthew McConaughey Trademarks Voice and Image Against AI.” <https://www.lawyer-monthly.com/2026/01/matthew-mcconaughey-protects-voice-image-ai/>
- Adweek, “Meet the \$1.3 Billion Startup Behind Madonna and Will Smith’s AI Video.” <https://www.adweek.com/media/higgsfield-ai-marketing-startup/>
- Marketing Week, “You can’t dismiss AI ads as slop when they’re winning in testing.” <https://www.marketingweek.com/dismiss-ai-ads-winning-creative-effectiveness/>
- Variety, “George Clooney: AI Actors Will Struggle to Be Movie Stars.” <https://variety.com/2025/scene/columns/george-clooney-ai-actors-movie-stars-1236579661/>
- The Hollywood Reporter, “Synthetic Sincerity on AI Characters, Authenticity (IDFA).” <https://www.hollywoodreporter.com/movies/movie-news/synthetic-sincerity-film-idfa-ai-authenticity-interview-1236426180/>
- Variety, “AI-Generated Images Lead to Audiences’ Distrust, Threaten Documentary.” <https://variety.com/2025/film/festivals/ai-generated-images->

[threaten-future-of-documentary-1236583466/](https://www.theguardian.com/technology/2026/aug/05/ai-generated-stories-rated-better-quality-than-human-written-ones-study-finds)

- *Hollywood Reporter*, “Merriam-Webster Names Slop Word of the Year 2025.” <https://www.hollywoodreporter.com/news/general-news/slop-word-year-2025-merriam-webster-1236450780/>
- Google DeepMind, “SynthID — 100 billion items watermarked; partner expansion to OpenAI, ElevenLabs, Kakao.” <https://deepmind.google/discover/blog/synthid-100-billion-watermarks-partners/>
- *Variety*, “Cate Blanchett co-founds RSL Media — non-profit on AI consent for name, image, likeness.” <https://variety.com/2026/film/news/cate-blanchett-rsl-media-ai-consent-1236748255/>
- *Bloomberg*, “Apple Acquires Key Talent & Patents Behind AI Avatar Company ‘Animato’.” <https://www.bloomberg.com/news/articles/2026-05-19/apple-acquires-animato-ai-avatar-talent-patents>
- Snap Newsroom, “Snapchat Gen Z AI Creativity Research 2026 — 31% of 13–15 year-olds use AI to be creative.” <https://newsroom.snap.com/snapchat-gen-z-ai-creativity-research-2026>
- YouGov AU, “45% of Australians say AI-generated ads would make them trust a brand less.” <https://yougov.com.au/topics/consumer/articles-reports/2026/05/19/45-percent-australians-ai-generated-ads-trust-brand-less>
- *Variety*, “Cate Blanchett, Steven Soderbergh, Nikki Hexum on How the Human Consent Registry Helps People Protect Their Identity From AI Theft: It’s a ‘Human Right’.” <https://variety.com/2026/biz/news/cate-blanchett-steven-soderbergh-nikki-hexum-ai-human-registry-1236791271/>
- *The Guardian*, “AI-generated stories rated better quality than human-written ones, study finds” — readers frequently preferred the machine-written work and could not reliably tell which was which. <https://www.theguardian.com/technology/2026/aug/05/ai-generated-stories-rated-better-quality-than-human-written-ones-study-finds>

X. Labour, organisation and the new geography

- Imperva/Cloudflare bot reports (Section III).
- Adobe, “Inaugural Adobe Creators’ Toolkit Report: 86% of Global Creators Use Creative Generative AI.” <https://news.adobe.com/news/2025/10/adobe-max-2025-creators-survey>
- Azumo, “AI in the Workplace Statistics 2026.” <https://azumo.com/artificial-intelligence/ai-insights/ai-in-workplace-statistics>
- *Tech.co*, “Gen Z Most Likely to Sneakily Use AI Without Telling the Boss.” <https://tech.co/news/gen-z-most-likely-use-ai-boss>

- *Exploding Topics*, “*The Hidden AI Workforce.*”
<https://explodingtopics.com/blog/ai-workforce-research>
- *Forbes*, “*AI Tools Flood Workplaces as Employees Face a Double Bind.*”
<https://www.forbes.com/sites/carolinecastrillon/2025/09/09/ai-tools-flood-workplaces-as-employees-face-a-double-bind/>
- IDC Europe, “*Shadow AI.*” <https://blog-idceurope.com/shadow-ai-how-stealth-productivity-is-strangling-enterprise-ai-adoption-and-creating-a-security-nightmare/>
- *Game Developer*, “*Krafton outlines plans to become ‘AI First’ company.*”
<https://www.gamedeveloper.com/business/subnautica-owner-krafton-outlines-plans-to-transform-into-an-ai-first-company>
- *CNBC*, “*Why workers with ADHD, autism, dyslexia should use AI agents.*”
<https://www.cnn.com/2025/11/08/adhd-autism-dyslexia-jobs-careers-ai-agents-success.html>
- *The Guardian*, “*AI is hitting UK harder than other big economies.*”
<https://www.theguardian.com/technology/2026/jan/26/ai-uk-jobs-us-japan-germany-australia>
- *PocketGamer.biz*, “*Shift Up CEO says AI is key to competing with China’s game industry scale.*” <https://www.pocketgamer.biz/shift-up-ceo-says-ai-is-key-to-competing-with-chinas-game-industry-scale/>
- *CNBC Africa*, “*How AI is changing the landscape of the music industry in Africa.*” <https://www.cnn.com/2025/how-ai-is-changing-the-landscape-of-the-music-industry-in-africa>
- *BBC Future*, “*Lights, camera, algorithm: Why Indian cinema is awash with AI.*” <https://www.bbc.com/future/article/20251223-why-indian-cinema-is-awash-with-ai>
- *Broadcast Pro Middle East*, “*Tunisian filmmaker wins \$1 million AI Film Award for ‘Lily’.*” <https://www.broadcastprome.com/news/tunisian-filmmaker-wins-1-million-ai-film-award-for-lily/>
- *UKTN*, “*AI hiring surges ‘across all UK sectors’, according to new findings from PwC, with a 34% salary premium for AI skills.*”
<https://www.uktech.news/ai/ai-hiring-surges-across-all-uk-sectors-20260615>
- *The Hollywood Reporter*, “*Hollywood Workers Are Training AI Models as Job Prospects Grow Slim.*”
<https://www.hollywoodreporter.com/business/digital/ai-training-hollywood-writer-jobs-prospects-1236628302/>
- *Ramp*, “*Ramp AI Index, June 2026*” — top 1% of firms now spend \$7,500 per employee per month on AI. <https://ramp.com/data/ai-index-june-2026>
- *GDC*, “*State of the Game Industry 2026*” — 52% of game developers now say generative AI is having a negative impact on the industry, up from 30% a year

earlier and 18% the year before that. <https://gdconf.com/article/gdc-2026-state-of-the-game-industry-reveals-impact-of-layoffs-generative-ai-and-more/>

- KPMG, “*Analysis of 1.4 million AI interactions identifies the employee behaviours behind effective AI use — and how they can be taught at scale.*” <https://kpmg.com/us/en/articles/2026/sophisticated-ai.html>
- Institute for the Future, “*The High-Impact Individual Contributor (HI-C)*” — the individual who can deliver a project of business value end-to-end on their own by combining deep expertise with a capable AI stack; the framing behind the Issue 37 Editor’s Pick on economies of scale giving way to economies of capability. *Dream Machine* Issue 37.

XI. *The literacy turn — institutions and education*

- Sundance Institute, “*Centering the Artist: Why We’re Launching the AI Literacy Initiative.*” <https://www.sundance.org/blogs/centering-the-artist-why-were-launching-the-ai-literacy-initiative/>
- Google blog, “*Building a community-led future for AI in film with Sundance Institute.*” <https://blog.google/company-news/outreach-and-initiatives/google-org/sundance-institute-ai-education/>
- McKinsey & Company, “*What AI could mean for film and TV production.*” <https://www.mckinsey.com/industries/technology-media-and-telecommunications/our-insights/what-ai-could-mean-for-film-and-tv-production-and-the-industrys-future>
- UK Government, “*Free AI training for all.*” <https://www.gov.uk/government/news/free-ai-training-for-all-as-government-and-industry-programme-expands-to-provide-10-million-workers-with-key-skills-by-2030>
- University of Wisconsin-Stout, “*AI Reshaping Industry: AI-use as Baseline Competency in Filmmaking.*” <https://www.uwstout.edu/about-us/news-center/ai-reshaping-industry-new-uw-stout-course-sets-ai-use-baseline-competency-filmmaking>
- UK Government (DSIT/DCMS), “*AI Adoption Plan: Creative Industries.*” <https://www.gov.uk/government/publications/ai-champions-ai-adoption-plans/ai-adoption-plan-creative-industries>

XI.a *Live sport and the World Cup deployment*

The 2026 FIFA World Cup as the largest live deployment of AI in entertainment history — the primary sources behind the Issue 34 Editor’s Pick.

- FIFA Football Technology, “*FIFA and Lenovo unveil multiple AI-powered innovations ahead of FIFA World Cup 2026™*” — one-second player scans,

- semi-automated offside, Football AI Pro, AI-generated personalised highlights.
<https://football-technology.fifa.com/organisation/media-releases/lenovo-tech-world-ai-powered-innovations-world-cup-2026>
- *AI Magazine*, “How AI Will Power the 2026 FIFA World Cup.”
<https://aimagazine.com/news/how-ai-will-power-the-2026-fifa-world-cup>
 - Lenovo StoryHub, “Lenovo Technology Powers FIFA World Cup 2026™ Operations and Strengthens AI-Driven Broadcast.”
<https://news.lenovo.com/pressroom/press-releases/lenovo-technology-powers-fifa-world-cup-2026-operations-and-strengthens-ai-driven-broadcast/>
 - GamesBeat, “FIFA and Gamefam launch FIFA World Cup 2026 event in Super Soccer and Roblox” — 1.1bn+ visits, 1.5M gameplay sessions a day.
<https://gamesbeat.com/gamefam-launch-fifa-world-cup-2026-super-soccer-roblox/>
 - Social Media Today, “TikTok introduces 2026 FIFA World Cup creator correspondents.” <https://www.socialmediatoday.com/news/tiktok-introduces-2026-fifa-world-cup-creator-correspondents/820198/>
 - Global Brands Magazine, “FIFA World Cup 2026: Inside the Biggest Brand Marketing Race” — WARC’s projected \$10.5bn in incremental Q2 ad spend.
<https://www.globalbrandsmagazine.com/fifa-world-cup-2026-marketing/>
 - Mercia AI, “How AI Will Shape the 2026 World Cup” — the fan-made content flood; 4,300+ fake FIFA websites before kick-off.
<https://www.merciaai.com/post/how-ai-will-shape-the-2026-world-cup>
 - Google DeepMind, “TacticAI — AI system simulating field scenarios and predicting open play up to 8 seconds in advance.”
<https://x.com/GoogleDeepMind/status/2065093482088169719>
 - IBM / Red Hat, “Serving faster innovation: how open source is supporting AI-powered fan experiences at Wimbledon.”
<https://www.redhat.com/en/blog/serving-faster-innovation-how-open-source-supporting-ai-powered-fan-experiences-wimbledon>
 - SportsLawHub, “Athlete Image Rights in the AI era: protection and legal challenges” — biometric digital twins and “synthetic rights” clauses.
<https://www.sportslawhub.com/athlete-image-rights-ai-deepfakes>

XII. Newsletters and primary corpus

- *Dream Machine | Creative AI* — LinkedIn newsletter, archive of Issues 1–41 (6 October 2025 – 12 August 2026).
<https://www.linkedin.com/newsletters/dream-machine-creative-ai-7379776527871381505/>
- DreamLab AI Collective, team page. <https://dreamlab-ai.com/team>

- DreamLab Substack, “Some Predictions on Creative AI for 2026” — *Dream Machine* Issue 13 companion piece. <https://substack.com/home/post/p-183997149>

*Appendix D — The Dream Machine
Source Index*

A thematic catalogue of significant sources surfaced across the 41 issues of *Dream Machine* (October 2025 – August 2026).

This index is a navigational tool, not an exhaustive list. The full *Dream Machine* archive contains more than three thousand individual hyperlinks across its forty-one issues, the great majority of which are primary-source links to industry coverage, research reports, official announcements, court filings, technical demos, creator showcases, and platform releases.

What follows below is the thematic catalogue of the *significant* sources — the ones the book itself draws on, the ones a working creative or researcher tracking a specific topic would want as a starting point, and the ones that, taken together, define the public record of creative AI as it stood in the period this book covers. Within each theme, entries are organised chronologically by issue number. The format is:

[Issue N] **Title** — short context — URL

For the complete primary-source archive — every link, in full, in original publication order — refer to the *Dream Machine* newsletter archive on LinkedIn or the per-issue markdown files in *Dream Machine* MD/. Each issue’s final section, “*All embedded URLs (in document order)*”, lists every URL the issue carried.

* * *

1. AI Video — Models and Releases

- [Issue 1] **Sora 2 Launch** — OpenAI’s step-change in physical realism, audio, multi-shot world state — <https://openai.com/index/sora-2/>
- [Issue 1] **Veo 3 and Flow** — Google DeepMind’s cinematic AI video — <https://www.youtube.com/watch?v=Io6Ef8alr2Y>
- [Issue 2] **Veo 3.1 Coming Soon** — improved consistency, resolution, multi-shot, cinematic presets — <https://www.cometapi.com/veo-3-1-is-comingand-whats-rumor/>
- [Issue 3] **Veo 3.1 Deep Dive with Flow** — cinematic filmmaking toolset — <https://www.youtube.com/watch?v=Io6Ef8alr2Y>
- [Issue 3] **Higgsfield Sketch-to-Video** — powered by Sora 2 — <https://higgsfield.ai/posts/6nkYSGwOcdyXVqZefE1MsQ>
- [Issue 3] **LiveGS** — mobile Gaussian-splatting video
- [Issue 4] **Veo 3.1 Style Tips** — text-to-video with style guidance — <https://x.com/GoogleAI/status/1980327604843381215>
- [Issue 4] **Veo Precision Features** — remove/add elements to scenes — <https://x.com/GoogleDeepMind/status/1980261047836508213>
- [Issue 4] **Heygen Identity Consistency with Veo 3.1** — character-consistent video —

- https://x.com/HeyGen_Official/status/1978491090618749193
- [Issue 4] **Higgsfield Popcorn** — storyboard tool with character consistency — https://x.com/higgsfield_ai/status/1981110992630341928
 - [Issue 4] **Odyssey 2** — real-time interactive video generation — <https://odyssey.ml/introducing-odyssey-2>
 - [Issue 4] **Krea Realtime** — open-sourcing the creative engine — https://www.linkedin.com/posts/krea-ai_today-were-open-sourcing-krea-realtime-activity-7386124532207689728-B7pD
 - [Issue 5] **Sora Character Cameos** — new feature in Sora app — <https://x.com/OpenAI/status/198366103653379486>
 - [Issue 5] **VEED Transitions** — AI-powered video transitions — <https://x.com/veedstudio/status/1980636419891818850>
 - [Issue 5] **LTX-2** — open-source audio-video generation — https://x.com/ltx_model/status/1981346235194683497
 - [Issue 6] **Wan 2.2** — AI video with multi-shot capabilities — <https://x.com/eyishazyer/status/1983507594942755221>
 - [Issue 6] **MotionStream** — real-time interactive video, 29 FPS — <https://x.com/wildmindai/status/1985828041566941576>
 - [Issue 7] **Odyssey-2** — interactive streaming 16:9 video — <https://x.com/olivercameron/status/1984777003967672800>
 - [Issue 9] **Wan 2.6 Released** — cast characters from reference videos, up to 15 seconds — <https://wan.video/blog/wan2.6-introduction>
 - [Issue 13] **Runway Gen-4.5** — image-to-video for paid plans — <https://www.youtube.com/watch?v=AwKSrJFvdps>
 - [Issue 13] **LTX-2 on ComfyUI** — open-source audio-video — <https://blog.comfy.org/p/ltx-2-open-source-audio-video-ai>
 - [Issue 14] **Veo 3.1 Ingredients to Video** — vertical formats, 1080p/4K — <https://x.com/FlowbyGoogle/status/2011130097483526474>
 - [Issue 14] **LTX-2 Lip Sync** — native audio-driven dialogue — https://x.com/ltx_model/status/2011101440706806051
 - [Issue 15] **Runway Gen-4.5 Image to Video** — broad rollout — https://www.linkedin.com/posts/runwayml_introducing-image-to-video-for-gen-45-the-activity-7419856988186238976-ZJU2
 - [Issue 15] **Veo 3.1 in YouTube Shorts and Create app** — distribution-layer integration — <https://blog.google/innovation-and-ai/technology/ai/veo-3-1-ingredients-to-video/>
 - [Issue 31] **Runway Aleph 2.0** — single-frame edit propagated across the video clip — June 2026

- [Issue 31] **Alibaba WonderClip** — all-in-one AI video creation platform — June 2026
- [Issue 31] **Sony Vid-CamEdit** — post-capture camera trajectory editing keeping scenes spatially coherent across viewpoint changes — June 2026
- [Issue 31] **Nvidia PiD (Pixel Diffusion Decoder)** — new image/video decoder — June 2026
- [Issue 33] **Video AI Wars: How Chinese Labs Are Winning The Race OpenAI Abandoned** — Forbes on the post-Sora competitive landscape — <https://www.forbes.com/sites/edithyeung/2026/06/05/video-ai-wars-how-chinese-labs-are-winning-the-race-openai-abandoned/>
- [Issue 35] **Seedance 2.5 announced** — ByteDance breaks the 30-second barrier for AI video generation — <https://the-decoder.com/bytedances-seedance-2-5-breaks-the-30-second-barrier-for-ai-video-generation/>
- [Issue 36] **Gemini Omni Flash** — generate and edit high-quality video via the Gemini API, AI Studio and ComfyUI — <https://ai.google.dev/gemini-api/docs/omni>
- [Issue 36] **Krea 2 open weights** — open-source release with separate inference and training models — <https://www.krea.ai/krea-2-open-source>
- [Issue 39] **FLUX 3** — Black Forest Labs model generating images and 20-second video with audio — <https://bfl.ai/blog/flux-3>
- [Issue 41] **Alibaba Wan 3** — video model in public beta with 30-second clips and multimodal inputs — <https://technode.global/2026/08/10/chinas-alibaba-releases-wan3-0-ai-video-model-in-public-beta-with-30s-clips-multimodal-inputs/>

2. AI Image — Models and Tools

- [Issue 1] **Nano Banana Plugin for Photoshop** — Gemini image gen inside Adobe — https://www.linkedin.com/posts/arminas-valunas-b4477255_nano-banana-plugin-for-photoshop-is-here-ugcPost-7367923639414906881-GgV7
- [Issue 1] **Magnific Precision v2** — AI upscaling — https://www.linkedin.com/posts/magnific-ai_introducing-magnific-precision-v2-activity-7387158930776276992-zpET
- [Issue 2] **Vimeo AI Creator Tools** — new features — <https://www.tvtechnology.com/news/vimeo-releases-new-ai-powered-creator-tools>
- [Issue 3] **Google Stitch** — design tool with AI features — <https://www.testingcatalog.com/google-test-new-stitch-modes-annotate-theme-interactive/>
- [Issue 5] **Adobe Firefly Image Model 5** — Adobe MAX 2025 — <https://news.adobe.com/news/2025/10/adobe-max-2025-firefly>

- [Issue 5] **Freepik Spaces** — infinite canvas for collaborative creation — <https://www.freepik.com/spaces>
- [Issue 5] **Midjourney updates** — state-of-the-art image gen — <https://x.com/midjourney/status/1991684484455100477>
- [Issue 6] **Qwen Image Multiple Angles LoRA** — consistent characters — <https://x.com/multimodalart/status/1986174924038218087>
- [Issue 8] **Qwen Image 2512** — latest generation — <https://qwen.ai/blog?id=qwen-image-2512>
- [Issue 9] **FLUX 2** — Black Forest Labs image generation — <https://bfl.ai/models/flux-2-max>
- [Issue 10] **FLUX 2 klein** — fast image generation under one second — <https://huggingface.co/unisloth/FLUX.2-klein-4B-GGUF>
- [Issue 13] **ChatGPT Images Upgrade** — major feature improvements — <https://www.techradar.com/ai-platforms-assistants/chatgpt/chatgpt-images-just-got-a-major-upgrade-and-it-could-change-how-we-all-create>
- [Issue 37] **Seedream 5.0 Pro** — ByteDance image model with character/product consistency, precision region editing and in-image text in 14 languages — <https://blog.comfy.org/p/seedream-50-pro>

3. *AI Music and Audio*

- [Issue 1] **Suno Studio** — generative audio workstation — <https://www.techradar.com/ai-platforms-assistants/i-tried-suno-studio-the-new-platform-that-mixes-ai-music-generation-with-hands-on-editing-like-garageband-but-smarter>
- [Issue 1] **YouTube Music AI Hub** — AI music hosts — <https://www.linkedin.com/news/story/youtube-music-debuts-new-ai-hub-6625484/>
- [Issue 1] **Spotify AI Protections** — strengthened for artists — <https://newsroom.spotify.com/2025-09-25/spotify-strengthens-ai-protections/>
- [Issue 1] **Xania Monet** — AI singer signs \$3M deal — <https://www.forbes.com/sites/dougmelville/2025/09/27/ai-singer-xania-monet-just-charted-on-billboard-signed-3m-deal-is-this-the-future-of-music/>
- [Issue 1] **Cardiff Band on AI Artist** — trained on their music, outperforming them — <https://musictech.com/news/industry/its-shocking-disheartening-and-insulting-cardiff-band-speaks-out-after-ai-artist-trained-on-their-music-outperforms-them-on-spotify/>
- [Issue 2] **Suno Funding Round** — \$2 billion valuation — <https://www.digitalmusicnews.com/2025/10/20/suno-funding-round-october-2025/>

- [Issue 2] **Groundhog AI Guitar Pedal** — tone matching — <https://musictech.com/news/gear/groundhog-audio-onepedal-ai-tone-matching/>
- [Issue 2] **Spotify in ChatGPT** — integration launched — <https://newsroom.spotify.com/2025-10-06/spotify-personalized-prompts-chatgpt/>
- [Issue 3] **iZotope Ozone 12** — AI assistant for mixing — <https://musictech.com/reviews/plugin-ins/izotope-ozone-12-review/>
- [Issue 3] **Tempolor Guitars (Quwan)** — AI to make songs playable — <https://kr-asia.com/no-practice-required-quwans-tempolor-guitars-use-ai-to-make-songs-playable-in-minutes>
- [Issue 4] **Mureka Music Agent Studio** — six specialised AI agents — https://www.linkedin.com/posts/sherrihendrickson_mureka-unveils-music-agent-studio-and-enhanced-share-7384999251526864896-cNYg/
- [Issue 4] **Fish Audio S1** — TTS 6× cheaper than ElevenLabs — <https://fish.audio/app/text-to-speech/>
- [Issue 5] **Universal Music + Stability AI Alliance** — strategic partnership — <https://stability.ai/news/universal-music-group-and-stability-ai-announce-strategic-alliance>
- [Issue 5] **Udio Partners with UMG** — partnership announced — <https://www.udio.com/blog/a-new-era>
- [Issue 5] **OpenAI Music Generator** — reportedly in development — <https://www.theinformation.com/articles/openai-plots-generating-ai-music-potential-rivalry-startup-suno>
- [Issue 6] **UMG Boss Lucian Grainge on AI** — full internal memo — <https://musically.com/2025/10/14/umg-boss-sir-lucian-grainge-talks-ai-full-internal-memo/>
- [Issue 6] **Bleeding Verse AI Band** — Hallwood Media signing — <https://musically.com/2025/10/07/ai-band-bleeding-verses-creator-signs-deal-with-hallwood-media/>
- [Issue 6] **JYP Entertainment AI Artist** — hiring AI/Unreal experts — <https://www.musicbusinessworldwide.com/jyp-entertainment-is-hiring-for-ai-and-unreal-engine-experts-to-develop-an-unprecedented-virtual-kpop-artist/>
- [Issue 6] **Claimy** — \$1.8M for missing royalty recovery — <https://www.musicbusinessworldwide.com/ai-music-tech-startup-claimy-raises-1-8m-to-tackle-missing-royalty-payments/>
- [Issue 7] **Breaking Rust on Billboard** — AI country act — <https://www.npr.org/2025/11/10/nx-s1-5604320/breaking-rust-is-a-hot-new-country-act-on-the-billboard-charts-its-powered-by-ai>

- [Issue 7] **GEMA v OpenAI Munich ruling** — European copyright precedent — https://www.linkedin.com/posts/dr-barry-scannell-bbb5aa207_in-a-major-ruling-for-european-copyright-share-7393957246386323457-8bbx
- [Issue 8] **LANDR AI Music Study** — 87% of musicians use AI tools — <https://aristake.com/ai-tools-musicians-study/>
- [Issue 8] **Stability AI + Warner Music** — next-gen responsible AI tools — <https://stability.ai/news/warner-music-group-and-stability-ai-join-forces-to-build-next-gen-tools>
- [Issue 8] **Paul McCartney Silent Track Protest** — UK copyright opt-out — <https://www.theguardian.com/music/2025/nov/17/the-sound-of-silence-why-theres-barely-anything-there-in-paul-mccartney-new-release>
- [Issue 11] **Lyria Camera (Google DeepMind)** — music generation — <https://magenta.withgoogle.com/lyria-camera-announce>
- [Issue 11] **Dave Stewart on AI** — Musicians must embrace it — <https://www.theguardian.com/music/2025/dec/05/musicians-must-embrace-unstoppable-force-of-ai-eurythmics-dave-stewart-urges>
- [Issue 12] **Splice + UMG Collaboration** — AI music creation tools — <https://www.universalmusic.com/universal-music-group-and-splice-to-collaborate-on-the-next-generation-of-ai-powered-music-creation-tools-for-artists/>
- [Issue 14] **Bandcamp Bans AI Music** — platform policy — <https://stereogum.com/2485199/bandcamp-bans-ai-music/news>
- [Issue 14] **UMG slams AI slop** — exponential growth on streaming — <https://musically.com/2026/01/09/umg-boss-slams-exponential-growth-of-ai-slop-on-streaming-services/>
- [Issue 15] **Sienna Rose** — viral mystery AI singer (BBC investigation) — <https://www.bbc.co.uk/news/articles/cq6v83gq66eo>
- [Issue 16] **800 Creatives Sign Declaration** — *Stealing Our Work Is Not Innovation* — <https://www.digitalmusicnews.com/2026/01/22/stealing-isnt-innovation/>
- [Issue 31] **ElevenLabs Music v2** — next-generation generative music model — June 2026
- [Issue 31] **Stable Audio 3.0** — open-weight model family on fully licensed data, foundation for the open audio community — June 2026
- [Issue 31] **Apple Music: AI <1% of plays** — clean platform-level confirmation that audience-chosen listening time has remained almost entirely human-authored — June 2026
- [Issue 31] **Spotify + UMG: AI covers/remixes as Premium add-on** — landmark licensed AI-music deal — June 2026

- [Issue 31] **Daniel Ek defends AI music** — “better than slop” framing — June 2026
- [Issue 31] **Tonada (Sweden)** — AI music for retailers; B2B-aligned synthetic background music — June 2026
- [Issue 31] **Spotify NotebookLM-style Personalised Audio app** — competing with Google’s NotebookLM podcasts — June 2026
- [Issue 31] **Musician uses AI as Parkinson’s affects ability to play guitar** — accessibility case for AI in music creation — June 2026
- [Issue 33] **Google Magenta RealTime 2 (MRT2)** — the live music model you can play as an instrument — <https://magenta.withgoogle.com/magenta-realtime-2>
- [Issue 33] **Fender Studio Pro DAW** — Moises stem separation and a smart studio assistant; “AI isn’t the destination. Making music is” — <https://www.musicradar.com/music-tech/daws/ai-isnt-the-destination-making-music-is-fender-looks-forward-as-its-studio-pro-daw-integrates-moises-stem-separation-and-a-smart-studio-assistant>
- [Issue 35] **Boy George re-records ‘Karma Chameleon’ via startup Artist** — AI-assisted vocal technology for re-recording classic tracks — <https://www.musicbusinessworldwide.com/boy-george-uses-ai-to-re-record-karma-chameleon-as-first-release-from-tech-startup-artist-included/>
- [Issue 35] **Sony AI Woosh** — foundation model for sound effect generation — <https://gamesbeat.com/sony-ai-releases-woosh-foundation-model-for-sound-effect-generation-exclusive-interview/>
- [Issue 36] **Google Lyria 3** — Gemini’s new AI music model — <https://blog.google/innovation-and-ai/products/gemini-app/lyria-3/>
- [Issue 36] **Suno Spark incubator** — programme to support the next generation of creators — <https://suno.com/blog/introducing-spark>
- [Issue 36] **Tidal allows AI-generated music — without royalties** — platform-policy split between hosting and monetisation — <https://musically.com/2026/06/30/tidal-allows-ai-generated-music-but-it-wont-earn-royalties/>
- [Issue 37] **Suno explores developer API** — seeking apps that ‘unlock experiences generative music makes possible for the first time’ — <https://www.musicbusinessworldwide.com/suno-explores-developer-api-seeking-apps-that-unlock-experiences-generative-music-makes-possible-for-the-first-time/>
- [Issue 37] **The BBC’s policy on the use of AI in music** — public broadcaster sets out its position — <https://radiotoday.co.uk/2026/07/the-bbcs-policy-on-the-use-of-ai-in-music/>

- [Issue 38] **Spotify is now an AI chatbot, too** — conversational interface across music, audiobooks and podcasts — <https://www.theverge.com/entertainment/965358/spotify-ai-chatbot-interface-music-audiobooks-podcasts>
- [Issue 38] **FL Studio agentic AI chatbot** — control the DAW with an AI chatbot from inside it — <https://www.musicradar.com/music-tech/you-can-now-control-fl-studio-with-an-ai-chatbot-from-inside-the-daw>

4. AI 3D / World Models / Spatial

- [Issue 1] **Marble by World Labs** — first commercial world model — <https://marble.worldlabs.ai/>
- [Issue 1] **Meta Hyperscape Capture** — Gaussian splatting on Quest — <https://www.meta.com/en-gb/experiences/meta-horizon-hyperscape-capture-beta/8798130056953686/>
- [Issue 1] **PlayCanvas SOG** — WebP for 3D Gaussian Splatting — https://www.linkedin.com/posts/willeastcott_playcanvas-open-sources-sog-literally-webp-ugcPost-7374459362708180992-aHda
- [Issue 3] **Genie 3** (Google DeepMind) — interactive 3D world generation — <https://deepmind.google/blog/genie-3-a-new-frontier-for-world-models/>
- [Issue 3] **World Labs RTFM** — real-time frame model — <https://www.worldlabs.ai/blog/rtfm>
- [Issue 3] **Instant Skinned Gaussian Avatars** — web/mobile VR — <https://sites.google.com/view/gaussian-vrm>
- [Issue 3] **Tencent Hunyuan World 1.1** — 3D reconstruction — <https://x.com/TencentHunyuan/status/1980930623536837013>
- [Issue 5] **Apple Personas use Gaussian Splatting** — most-deployed splat tech in consumer hardware — <https://radiancefields.com/apple-confirms-personas-use-gaussian-splatting>
- [Issue 7] **Marble formal public launch** — Sony Pictures using it (40× faster than legacy VP) — <https://www.worldlabs.ai/case-studies/bringing-marble-to-life>
- [Issue 8] **Meta SAM 3 / SAM 3D** — segment anything in 3D — https://www.linkedin.com/posts/aiatmeta_introducing-sam-3-sam-3d-ugcPost-7396944913751465985-m5Nc
- [Issue 11] **Meta WorldGen** — text-to-immersive-3D-worlds research — <https://www.facebook.com/LifeAtMeta/videos/research-update-worldgen-text-to-immersive-3d-worlds/1879077432692421/>
- [Issue 11] **Ubisoft CHORD Model** — open-sourced PBR material generation — <https://blog.comfy.org/p/ubisoft-open-sources-the-chord-model>

- [Issue 12] **Tencent HY World 1.5** — real-time world model framework — <https://x.com/TencentHunyuan/status/2001170499133653006>
- [Issue 12] **Microsoft Trellis 2** — 3D generation — <https://github.com/microsoft/TRELLIS.2>
- [Issue 13] **Wonderzoom** (Stanford AI Lab) — multi-scale 3D world generation — <https://wonderzoom.github.io/>
- [Issue 15] **WorldLabs API** — 3D world generation as a service — <https://x.com/theworldlabs/status/2014046372639408203>
- [Issue 17] **Project Genie Public Release** — Google AI Ultra subscribers — <https://blog.google/innovation-and-ai/models-and-research/google-deepmind/project-genie/>
- [Issue 22] **Luma UNI-1** — combined world generation + reasoning — *Dream Machine* Editor's Pick
- [Issue 25] **Spark 2.0** — open-source 100M-splat browser streaming
- [Issue 27] **Vista4D** (Netflix + Eyeline) — live action to navigable 4D point clouds
- [Issue 31] **TripoAI Project Eden** — persistent multiplayer world model decoupling world state from visual rendering — June 2026
- [Issue 31] **AI-Enhanced LiDAR** — generative-AI-augmented LiDAR capture pipeline — June 2026
- [Issue 33] **Gaussian splatting comes to Apple Maps** — WWDC Flyover upgrade; splats reach planetary-scale consumer mapping — <https://radiancefields.com/apple-maps-flyover-is-getting-a-gaussian-splatting-upgrade>
- [Issue 35] **Odyssey \$310M Series B** — funding to accelerate world models — <https://odyssey.ml/our-series-b>

5. Voice, Avatars, Digital Humans

- [Issue 1] **Tilly Norwood AI Actress** — sparks UK/US union debate — <https://www.hollywoodreporter.com/movies/movie-news/tilly-norwood-ai-actress-uk-union-equity-sag-aftra-debate-1236391739/>
- [Issue 1] **SAG-AFTRA Condemns Tilly Norwood** — official union statement — <https://variety.com/2025/film/news/sag-aftra-tilly-norwood-ai-actress-1236534779/>
- [Issue 1] **Deep Fusion Films + Topfoto** — AI documentary alliance — <https://www.ibc.org/production/news/deep-fusion-and-topfoto-strike-alliance-to-produce-ai-powered-documentaries/22728>
- [Issue 2] **ROXi AI TV Presenters** — music channel hosts — <https://www.advanced-television.com/2025/10/07/roxi-debuts-ai-generated->

tv-presenters/

- [Issue 2] **Hedra Audio Tags** — emotional audio control — https://x.com/hedra_labs/status/1998490844748460528
- [Issue 3] **Heygen Motion Designer** — prompt-based animation — https://www.linkedin.com/posts/heygen_if-you-can-explain-it-you-can-animate-it-activity-7387488068204892160-AY8-
- [Issue 4] **Copresence and ConvAI** — Unreal-Engine intelligent avatars — https://www.linkedin.com/posts/copresence-tech_want-to-create-an-intelligent-avatar-that-activity-7379523290421383168-SAg_
- [Issue 7] **Lumi Avatar** — real-time audio-driven avatars — <https://www.linkedin.com/feed/update/urn:li:activity:7393916271018467328/>
- [Issue 7] **Tilly Creator Eline Van der Velden — Deadline** — backlash and the next 40 — <https://deadline.com/2025/11/tilly-norwood-creator-interview-backlash-more-ai-actors-coming-1236601334/>
- [Issue 8] **Synthesia \$4B valuation** — rejected \$3B Adobe offer — <https://techcrunch.com/2026/01/26/synthesia-hits-4b-valuation-lets-employees-cash-in/>
- [Issue 11] **ElevenLabs Impact Program** — SXSW documentary on voice loss — https://www.linkedin.com/posts/elevenlabsio_at-sxsw-the-elevenlabs-impact-program-will-activity-7413979472988774400-fVRG
- [Issue 12] **Google Veo Avatars** — expressiveness upgrade — <https://vids.new/>
- [Issue 12] **Meta SAM Audio** — segment anything for audio — <https://about.fb.com/news/2025/12/our-new-sam-audio-model-transforms-audio-editing/>
- [Issue 13] **Avatar Forcing** — real-time interactive avatar generation — <https://taekyungki.github.io/AvatarForcing/>
- [Issue 13] **Qwen3 TTS** — voice design and cloning — <https://qwen.ai/blog?id=qwen3-tts-vc-voicedesign>
- [Issue 16] **Tilly Norwood Doubles Down** — AI as “more ethical” performance, urging actors to create avatars — <https://variety.com/2026/digital/news/tilly-norwood-creator-tells-actors-to-create-ai-avatars-1236638940/>
- [Issue 23] **Death threats against Eline Van der Velden** — cultural-extreme response
- [Issue 31] **Ozzy Osbourne AI avatar** — estate-led posthumous “live on as an AI avatar that talks with fans” — June 2026
- [Issue 33] **New York Times roasted for “profiling” Tilly Northwood** — the AI-actress story crosses into media self-parody — <https://futurism.com/artificial-intelligence/new-york-times-roasted-profiling-ai-actress>

- [Issue 34] **ElevenLabs Avatars** — “the best AI voices, now with a face” in ElevenCreative — <https://elevenlabs.io/blog/introducing-avatars>
- [Issue 36] **Netflix to recreate Gene Wilder’s voice with AI** — posthumous voice cloning for a Willy Wonka reality series — <https://www.nbcnews.com/pop-culture/pop-culture-news/gene-wilder-ai-voice-netflix-willy-wonka-competition-show-rcna352440>
- [Issue 37] **AI ‘actor’ Tilly Norwood to star in feature film ‘Misaligned’** — the synthetic performer lands a lead role — <https://deadline.com/2026/07/tilly-norwood-ai-actor-misaligned-1236974639/>

6. Agent Platforms / Orchestration

- [Issue 2] **OpenAI AgentKit / DevDay** — agentic AI for creative workflows — <https://openai.com/index/introducing-agentkit/>
- [Issue 2] **Lenny AI Agent** — for live music event organisers — <https://musically.com/2025/10/20/meet-lenny-an-ai-agent-to-help-organisers-of-live-music-events/>
- [Issue 4] **Adsgency \$12M seed** — autonomous paid marketing — <https://www.finsmes.com/2025/10/adsgency-raises-12m-in-seed-funding.html>
- [Issue 5] **Meta + Hugging Face OpenEnv** — open-source agentic development — <https://www.edtechinnovationhub.com/news/meta-and-hugging-face-launch-openenv-to-advance-open-source-agentic-development>
- [Issue 5] **Pomelli** (Google Labs) — AI marketing agent for SMBs
- [Issue 5] **Opal** (Google) — no-code AI mini-app builder — <https://blog.google/technology/google-labs/opal-expansion/>
- [Issue 8] **Multimodal Agents in Unreal Engine** — live-score nature documentary — <https://www.youtube.com/watch?v=7u2yCtbONmo>
- [Issue 8] **SIMA 2** (Google DeepMind) — agent for virtual 3D worlds — <https://deepmind.google/blog/sima-2-an-agent-that-plays-reasons-and-learns-with-you-in-virtual-3d-worlds/>
- [Issue 8] **Google Antigravity** — agentic development platform
- [Issue 13] **NitroGen** (NVIDIA + Stanford) — plays-any-game model — <https://nitrogen.minedojo.org/>
- [Issue 14] **General Intuition** — \$134M for spatial-reasoning agents — <https://techcrunch.com/2025/10/16/general-intuition-lands-134m-seed-to-teach-agents-spatial-reasoning-using-video-game-clips/>
- [Issue 16] **Anthropic Claude Apps** — interactive Claude in workplace tools — <https://techcrunch.com/2026/01/26/anthropic-launches-interactive-claude-apps-including-slack-and-other-workplace-tools/>

- [Issue 16] **Heygen Video Agent** — full scripting-to-assembly — https://www.linkedin.com/posts/heygen_introducing-the-new-video-agent-activity-7421597801240801282-d1CF
- [Issue 21] **Adobe + NVIDIA Strategic Partnership** — agentic creative intelligence
- [Issue 26] **Adobe Summit 2026** — “agentic creative intelligence” headline category
- [Issue 29] **Sony 49-Claude-agent / 72-skill stack** — game-dev multi-agent team
- [Issue 34] **Runway in ChatGPT** — generate and edit video and images inside ChatGPT via MCP — <https://runwayml.com/mcp>
- [Issue 35] **NVIDIA ACE Game Agent SDK** — on-device AI companions with Unreal Engine 5 plugins — <https://developer.nvidia.com/blog/build-on-device-ai-companions-with-the-nvidia-ace-game-agent-sdk-and-unreal-engine-5-plugins/>
- [Issue 36] **Comfy MCP public beta** — “turn your agent into a creative technologist”; the first MCP server built for production creative pipelines — <https://comfy.org/mcp/>

7. Adobe and Creative Software

- [Issue 1] **Tether** — AI animation in After Effects — https://www.linkedin.com/posts/thisisdoug_aftereffects-ai-video-vfx-ugcPost-7368671859774517249-losz
- [Issue 1] **Unreal Engine 5 AI Assistant** — official integration — https://www.linkedin.com/posts/wouterweynants_theres-an-official-ai-assistant-coming-to-ugcPost-7369377204226379776-pGiH
- [Issue 1] **Unity AI Council** — accelerate AI innovation — <https://www.gamedeveloper.com/business/unity-forms-ai-council-to-accelerate-ai-product-innovation->
- [Issue 1] **ComfyUI raises \$17M** — OS for creative AI — https://www.linkedin.com/posts/comfyui_we-raised-17-million-to-build-an-os-for-ugcPost-7373743341236236288-wkCc
- [Issue 5] **Adobe MAX 2025 Express AI Assistant** — Adobe MAX announcements — <https://news.adobe.com/news/2025/10/adobe-max-2025-express-ai-assistant>
- [Issue 5] **Adobe Firefly Foundry** — custom models for brands — <https://news.adobe.com/news/2025/10/adobe-max-2025-firefly-foundry>
- [Issue 5] **Adobe MAX Sneaks** — Light Touch, Surface Swap, Scene It, etc. — <https://www.youtube.com/watch?v=YqAAFx1XXY8>

- [Issue 5] **Adobe MAX Creator Survey** — 86% use creative gen AI — <https://news.adobe.com/news/2025/10/adobe-max-2025-creators-survey>
- [Issue 5] **Adobe Project Light Touch / Scene It / Surface Swap** — see Adobe Sneaks
- [Issue 5] **“Adobe is putting AI in everything everywhere all at once”** — Creative Boom coverage — <https://www.creativeboom.com/news/adobe-is-putting-ai-in-everything-everywhere-all-at-once/>
- [Issue 6] **Adobe Corrective AI** — voice-over emotion editing — <https://www.wired.com/story/adobe-max-sneaks-2025-corrective-ai/>
- [Issue 8] **Adobe Research RELIC** — interactive video world model
- [Issue 12] **Adobe inside ChatGPT** — Photoshop, Express, Acrobat editing in ChatGPT — <https://blog.adobe.com/en/publish/2025/12/10/edit-photoshop-chatgpt>
- [Issue 12] **Google Flow Refine** — prompt by doodling — <https://blog.google/technology/google-labs/flow-refine-videos/>
- [Issue 15] **Adobe at Sundance 2026** — \$10M grants, Ignite Day — <https://analyticsindiamag.com/ai-news-updates/adobe-unveils-ai-video-innovations-10-million-grants-ahead-of-sundance-film-festival/>
- [Issue 16] **Adobe Premiere Object Mask** — automated masking — https://www.linkedin.com/posts/robdewinter_ok-this-is-going-to-save-a-lot-of-time-in-ugcPost-7421617551690063872-yKmB
- [Issue 28] **Unity AI Open Beta** — in-editor full AI suite
- [Issue 33] **Higgsfield Plugin for DaVinci Resolve** — generate footage in the timeline, draw-to-edit, AI-generated LUTs — <https://higgsfield.ai/plugins/davinci>
- [Issue 35] **State of Unreal 2026: Unreal MCP and AI integrations** — Epic embeds LLM-powered tools and video models into engine workflows — <https://www.unrealengine.com/news/state-of-unreal-2026-top-news-from-the-show>
- [Issue 35] **Adobe Creative Agent and AI Assistant in beta** — across Premiere, Photoshop, Illustrator, InDesign and Frame.io — <https://news.adobe.com/news/2026/06/adobe-unveils-major-expansion>
- [Issue 35] **Disney Imagineers adopt Adobe Firefly Foundry** — custom models as worldbuilding toolkit — <https://business.adobe.com/blog/disney-imagineers-add-adobe-firefly-foundry-to-toolkit>
- [Issue 36] **Adobe to acquire Topaz Labs** — the AI upscaling/enhancement specialist folded into Creative Cloud — <https://news.adobe.com/news/2026/06/adobe-to-acquire-topaz-labs>

8. Games Industry — Adoption

- [Issue 1] **Netflix Director of GenAI for Games** — \$840K role — <https://www.pcgamer.com/gaming-industry/as-the-videogame-industry-continues-to-be-hammered-by-layoffs-netflix-is-offering-up-to-usd840-000-per-year-for-a-new-director-of-generative-ai-for-games/>
- [Issue 1] **Meta Horizon Studio AI Assistant Upgrade** — <https://www.uploadvr.com/meta-horizon-studio-upgrade-ai-assistant-horizon-worlds/>
- [Issue 1] **51% of Japanese games studios use AI** — research finding — <https://www.gamesindustry.biz/51-of-japanese-game-makers-use-generative-ai>
- [Issue 3] **Roblox AI Tools for Creators** — <https://www.gamesindustry.biz/roblox-announces-new-ai-tools-for-creators>
- [Issue 3] **Battlefield 6** — “very seducing” AI for early stages — <https://www.gamesradar.com/games/battlefield/battlefield-6-lead-calls-generative-ai-very-seducing-but-says-it-was-only-used-in-the-games-earliest-stages-to-allow-for-more-time-and-more-space-to-be-creative/>
- [Issue 3] **Promise (Google-backed AI studio)** — VFX for legacy media — <https://www.hollywoodreporter.com/business/digital/ai-studio-promise-vfx-generation-company-1236397636/>
- [Issue 4] **EA + Stability AI Partnership** — generative AI tools for games — <https://stability.ai/news/stability-ai-and-ea-partner-to-reimagine-game-development>
- [Issue 4] **Halo Studios** — AI woven into development — <https://thegamepost.com/insider-halo-studios-generative-ai-game-development/>
- [Issue 5] **NBCUniversal × Dick Wolf Jr** — AI games deal — <https://www.videogameschronicle.com/news/nbcuniversal-signs-deal-with-law-order-creator-dick-wolfs-son-to-make-ai-generated-games-based-on-its-ip/>
- [Issue 5] **Microsoft Gaming Copilot** — screenshots for in-game understanding — <https://www.tomshardware.com/video-games/pc-gaming/microsoft-says-gaming-copilot-uses-screenshots-to-understand-in-game-events-not-for-training-ai-models-optional-feature-can-be-turned-off-but-not-easily-uninstalled>
- [Issue 5] **Sony Jabali AI Platform** — game development — <https://variety.com/2025/gaming/news/sony-jabali-ai-ai-game-development-platform-1236566619/>
- [Issue 5] **Krafton AI-First Plans** — Subnautica owner — <https://www.gamedeveloper.com/business/subnautica-owner-krafton-outlines-plans-to-transform-into-an-ai-first-company>

- [Issue 6] **Todd Howard on AI at Bethesda** — toolset, not replacement — <https://www.pcgamer.com/gaming-industry/todd-howard-says-ai-cant-replace-human-creative-intention-but-its-part-of-bethesdas-toolset-for-how-we-build-our-worlds-or-check-things/>
- [Issue 6] **EA pushes 15K employees on AI** — thought partner — <https://www.gamesradar.com/games/even-under-usd20-million-in-debt-ea-reportedly-pushes-15-000-employees-to-use-ai-as-a-thought-partner-for-everything-from-character-art-to-playtesting/>
- [Issue 6] **Falcom AI** — 2-3 hours of work to 10 minutes — <https://www.eurogamer.net/falcom-is-the-latest-developer-to-buy-into-the-ai-hype-machine>
- [Issue 7] **Square Enix 70% AI QA target** — by end 2027 — <https://www.pcgamer.com/gaming-industry/square-enix-aims-to-have-ai-doing-70-percent-of-its-qa-work-by-the-end-of-2027/>
- [Issue 8] **Call of Duty: Black Ops 7 AI art accusations** — community backlash — <https://www.videogameschronicle.com/news/it-honestly-sucks-fans-think-call-of-duty-black-ops-7-is-filled-with-generative-ai-art/>
- [Issue 8] **Ubisoft Anno 117 — AI art placeholder** — slipped through review — <https://www.videogameschronicle.com/news/ubisoft-says-ai-generated-art-in-anno-117-was-a-placeholder-which-slipped-through-our-review-process/>
- [Issue 13] **Razer \$600M AI focus** — strategic investment — <https://www.pymnts.com/news/artificial-intelligence/2026/razer-spends-600-million-dollars-sharpen-focus-ai-gaming/>
- [Issue 15] **Ubisoft cancels Prince of Persia + four** — AI refocus — <https://metro.co.uk/2026/01/21/prince-persia-remake-five-games-cancelled-ubisoft-focuses-ai-26431926/>
- [Issue 29] **Sony all-in on AI for games** — 49-agent / 72-skill stack
- [Issue 33] **The Sandbox Studio** — from idea to live multiplayer game in hours — <https://www.youtube.com/watch?v=IdTVDRvQoWM>
- [Issue 33] **Amazon's new plan for games** — James Bond and an AI Snoop Dogg on Luna — <https://www.theverge.com/games/943147/amazon-gaming-strategy-james-bond-snoop-dogg-luna>
- [Issue 33] **Sega confirms generative AI content in new Crazy Taxi game** — <https://gameinformer.com/2026/06/07/sega-confirms-and-responds-to-generative-ai-content-in-new-crazy-taxi-game>
- [Issue 34] **120 of 300+ Steam releases in one week carried AI disclosures** — the disclosure regime meets the volume wave — <https://www.pcgamer.com/gaming-industry/steam-week-in-review-more-than-300-games-released-on-steam-last-week-and-120-of-them-had-ai-disclosures/>

- [Issue 35] **Epic Games details how it's embracing generative AI in Unreal Engine** — <https://www.engadget.com/2196807/epic-games-details-how-its-embracing-gen-ai-in-unreal-engine/>
- [Issue 36] **EA: AI delivering “faster prototyping” and a “real rise in creativity”** — <https://www.pcgamer.com/software/ai/ea-exec-says-ai-has-helped-drive-a-real-rise-of-creativity-at-its-studios/>
- [Issue 36] **Epic CEO criticises Valve's AI disclosure policy for Steam** — access to AI tools framed as survival issue for smaller studios — <https://80.lv/articles/epic-ceo-criticizes-valve-s-ai-disclosure-policy-for-steam>
- [Issue 37] **Epic Games' CEO argues AI could've made Destiny profitable** — Tim Sweeney on AI and studio economics — <https://www.videogameschronicle.com/news/epic-games-ceo-argues-ai-couldve-made-destiny-profitable/>
- [Issue 38] **Nvidia and Sega team up on games for RTX Spark AI computers** — 30 years after their first partnership — <https://gamesbeat.com/nvidia-and-sega-team-up-on-games-for-rtx-spark-ai-computers-30-years-after-1st-partnership/>
- [Issue 39] **Epic Games adds AI voices to Fortnite NPCs** — for an upcoming in-game conversations feature — <https://www.videogameschronicle.com/news/fortnite-has-started-adding-ai-voices-to-its-npcs-for-an-upcoming-conversations-feature/>

9. Games Industry — Refusal / Position

- [Issue 1] **Charles Cecil** — “AI was an expensive mistake” — <https://www.gamesindustry.biz/ai-was-an-expensive-mistake-charles-cecil-on-innovation-insolvency-and-broken-sword>
- [Issue 4] **Pocketpair (Palworld)** — publishing won't take AI games — <https://www.pcgamer.com/software/ai/palworld-studio-pocketpair-says-its-new-publishing-division-wont-handle-games-that-use-generative-ai-we-dont-believe-in-it/>
- [Issue 9] **Witcher 3 / Cyberpunk Director** — AI helps, doesn't replace — <https://www.gamesindustry.biz/witcher-3-and-cyberpunk-2077-director-says-ai-can-help-but-not-replace-creatives>
- [Issue 11] **Aardman on AI** — embrace, but cautious — <https://www.gamesradar.com/entertainment/animation-movies/wallace-and-gromit-creator-says-beloved-animation-studio-aardman-will-embrace-the-technology-of-ai-but-will-be-very-cautious-not-to-lose-our-values/>
- [Issue 14] **Larian backs off gen AI** — Divinity statement — <https://nichegamer.com/larian-studios-backs-off-from-gen-ai/>

- [Issue 14] **Games Workshop rules out gen AI** — Warhammer 40K — <https://decrypt.co/354482/warhammer-40000-maker-games-workshop-rules-out-generative-ai>
- [Issue 14] **Hooded Horse won't work with AI devs** — Manor Lords publisher — <https://nichegamer.com/manor-lords-publisher-hooded-horse-wont-work-with-devs-using-gen-ai/>
- [Issue 16] **Jagex never AI** — RuneScape commitment — <https://www.gamesindustry.biz/runescape-maker-jagex-says-it-will-never-use-generative-ai-to-make-in-game-content>
- [Issue 35] **CD Projekt boss warns games made solely by gen AI are already being made** — <https://metro.co.uk/2026/06/22/cd-projekt-boss-warns-games-made-solely-gen-ai-already-made-28878296/>
- [Issue 36] **Godot will no longer accept AI-authored code contributions** — “we can't trust heavy users of AI to understand their code enough to fix it” — <https://www.pcgamer.com/gaming-industry/open-source-game-engine-godot-will-no-longer-accept-ai-authored-code-contributions-we-cant-trust-heavy-users-of-ai-to-understand-their-code-enough-to-fix-it/>
- [Issue 39] **Why China is embracing AI and “very few” studios refuse to adopt it** — the contrasting adoption culture — <https://www.pocketgamer.biz/why-china-is-embracing-ai-and-very-few-studios-refuse-to-adopt-it/>
- [Issue 41] **GDC 2026 State of the Game Industry** — 52% of developers now say generative AI is having a negative impact, up from 30% — <https://gdconf.com/article/gdc-2026-state-of-the-game-industry-reveals-impact-of-layoffs-generative-ai-and-more/>

10. Film Industry — Studios and Positions

- [Issue 1] **Lionsgate AI failure** — Futurism report — <https://futurism.com/artificial-intelligence/lionsgate-movies-ai>
- [Issue 2] **Fremantle's Imaginae AI Studios** — CEO named — <https://www.hollywoodreporter.com/business/digital/fremantle-names-ceo-new-ai-label-imaginae-studios-1236396579/>
- [Issue 2] **Goldfinch enGEN3** — AI cinematic universe — <https://variety.com/2025/film/news/ai-powered-cinematic-universe-platform-engen3-1236543349/>
- [Issue 3] **Fox Entertainment + Holywater** — AI microdramas — <https://www.hollywoodreporter.com/business/business-news/fox-entertainment-invests-in-holywater-ai-microdramas-1236396802/>
- [Issue 4] **Netflix “all in” on AI** — Sarandos at industry conference — <https://www.cnn.com/2025/10/22/netflix-all-in-on-leveraging-ai-in-its->

streaming-platform.html

- [Issue 4] **Asteria’s “All Heart”** — Natasha Lyonne short — <https://www.hollywoodreporter.com/movies/movie-news/natasha-lyonne-ai-company-asteria-1236403144/>
- [Issue 5] **Wonder Studios \$9M raise** — AI-native studio — <https://www.uktech.news/ai/ai-film-studio-wonder-lands-9m-investment-20251023>
- [Issue 5] **Watch the Skies (Swedish AI feature dubbed)** — USA distribution — <https://variety.com/2025/film/news/watch-the-skies-us-theatrical-release-ai-dubbing-1236343110/>
- [Issue 5] **Run to the West (South Korean first AI feature)** — <https://cybernews.com/entertainment/korean-cinema-run-to-the-west-ai/>
- [Issue 6] **Obsidian + Imagine Entertainment** — Ron Howard, Brian Grazer — <https://www.indiewire.com/news/business/obsidian-studio-ai-production-company-imagine-entertainment-1235158619/>
- [Issue 7] **Beta Films Chapter41 launch** — Munich AI startup — <https://deadline.com/2025/11/beta-film-ai-startup-chapter41-artificial-intelligence-1236612632/>
- [Issue 7] **House of David 350+ AI shots** — Wired feature — <https://www.wired.com/story/amazons-house-of-david-used-over-350-ai-shots-in-season-2-its-creator-isnt-sorry/>
- [Issue 7] **Wanted director’s AI Method Actors** — Bekmambetov \$5M — <https://variety.com/2025/film/news/wanted-director-method-acting-ai-actors-1236579647/>
- [Issue 7] **Kevin Reilly + Kartel** — HBO veteran’s AI startup — <https://www.hollywoodreporter.com/business/digital/kevin-reilly-ceo-kartel-ai-hbo-1236424692/>
- [Issue 8] **Humans in the Loop — Oscar race** — Sloan grant — <https://variety.com/2025/film/news/ai-drama-humans-in-the-loop-oscar-race-1236582975/>
- [Issue 8] **Synthetic Sincerity — IDFA** — Marc Isaacs — <https://www.hollywoodreporter.com/movies/movie-news/synthetic-sincerity-film-idfa-ai-authenticity-interview-1236426180/>
- [Issue 8] **AI Images Threaten Documentary** — Variety — <https://variety.com/2025/film/festivals/ai-generated-images-threaten-future-of-documentary-1236583466/>
- [Issue 11] **Disney \$1bn OpenAI investment** — Sora characters — <https://www.theguardian.com/business/2025/dec/11/disney-open-ai-sora-video-deal>

- [Issue 14] **Tunisian filmmaker wins \$1M for Lily** — Dubai AI Award — <https://www.broadcastprome.com/news/tunisian-filmmaker-wins-1-million-ai-film-award-for-lily/>
- [Issue 14] **Disney TikTok-like vertical video AI** — brand-asset video gen — <https://www.marketingdive.com/news/disney-unveils-tiktok-like-vertical-video-ai-video-generation-tool/809269/>
- [Issue 15] **Netflix retention AI strategy** — Pymnts — <https://www.pymnts.com/subscription-commerce/2026/retention-is-name-of-the-game-for-netflixs-ai-strategy/>
- [Issue 16] **Andrii Daniels bomb-shelter Christmas clip** — viral Ukrainian AI film — <https://variety.com/2026/digital/news/ai-video-deadpool-harry-potter-andrii-daniels-1236624632/>
- [Issue 16] **Tilly Norwood Doubles Down** — Variety — <https://variety.com/2026/digital/news/tilly-norwood-creator-tells-actors-to-create-ai-avatars-1236638940/>
- [Issue 16] **Chris Pratt rejects AI villain** — Mercy pitch — <https://variety.com/2026/film/news/chris-pratt-ai-actor-villain-mercy-amazon-mgm-1236640460/>
- [Issue 33] **Lionsgate's Michael Burns: AI will save “tens and tens of millions of dollars a year”** — <https://deadline.com/2026/06/lionsgate-michael-burns-ai-saves-millions-runway-1236941043/>
- [Issue 33] **Can AI tell if your script will make a hit film?** — Quilty's AI script analysis — <https://www.theverge.com/ai-artificial-intelligence/943531/ai-script-quilty-simon-horsman-daniel-wood>
- [Issue 34] **Runway and Lionsgate expand partnership** — Lionsgate takes an equity stake in Runway — <https://runwayml.com/news/runway-and-lionsgate-expand-partnership>
- [Issue 34] **Utopai Studios to co-produce ‘Half Moon’** — AI firm co-producing a Silver Bear-winning filmmaker's feature — <https://deadline.com/2026/06/utopai-studios-hyo-joo-yang-half-moon-ai-1236952730/>
- [Issue 34] **Asteria and LTX form strategic partnership** — advancing cinematic AI storytelling — <https://deadline.com/2026/06/asteria-ltx-form-strategic-partnership-cinematic-ai-1236955140/>
- [Issue 35] **Google investing in ‘Backrooms’ studio A24** — DeepMind AI venture with the prestige indie — <https://www.hollywoodreporter.com/business/digital/a24-google-deepmind-ai-venture-backrooms-1236627228/>
- [Issue 35] **Fully AI-generated series ‘Journey to the West’** — Utopai Studios and Huace development deal —

- <https://variety.com/2026/tv/news/fully-ai-generated-journey-to-the-west-utopai-studios-huace-1236787171/>
- [Issue 36] **A24 full-on fan crisis after Google AI deal** — the authenticity brand meets the AI partnership — <https://futurism.com/artificial-intelligence/a24-crisis-fans-signing-ai-deal-google>
 - [Issue 36] **Animation’s AI reckoning: movies for 90% less** — Bloomberg on filmmakers’ cost claims — <https://www.bloomberg.com/news/newsletters/2026-06-28/animation-s-ai-reckoning-filmmakers-say-they-can-make-movies-for-90-less>
 - [Issue 37] **Midjourney wants Hollywood studios to reveal their AI usage** — pressing Disney, Universal and Warner Bros to disclose — <https://techcrunch.com/2026/07/04/midjourney-wants-hollywood-studios-to-reveal-the-details-of-their-ai-usage/>
 - [Issue 38] **The AI ‘Star Wars’ film that is giving Hollywood a wake-up call** — Forbes on the independent AI film rattling the industry — <https://www.forbes.com/sites/carolinereid/2026/07/10/the-ai-star-wars-film-that-is-giving-hollywood-a-wake-up-call/>
 - [Issue 38] **Homer epic ‘The Odyssey’ gets a full-length AI-generated film** — a classical text rendered entirely with AI — <https://www.forbes.com/sites/lesliekatz/2026/07/14/homer-epic-the-odyssey-gets-a-full-length-ai-generated-film/>
 - [Issue 39] **About 300 Netflix titles used generative AI this year** — the company reveals the scale in Q2 earnings — <https://variety.com/2026/biz/news/about-300-netflix-programs-used-ai-this-year-q2-earnings-1236812914/>
 - [Issue 39] **‘The Lord of the Rings: The Hunt for Gollum’ will use AI to de-age actors** — AI in a tentpole production — <https://futurism.com/artificial-intelligence/lord-of-the-rings-ai-de-age>
 - [Issue 39] **Netflix paid \$587M for Ben Affleck’s AI filmmaking startup** — a marquee acquisition in AI production — <https://techcrunch.com/2026/07/19/netflix-paid-587m-for-ben-afflecks-ai-filmmaking-startup/>
 - [Issue 40] **Hollywood fights AI in public while quietly building it into movies** — the industry’s open secret — <https://www.yahoo.com/entertainment/movies/articles/hollywoods-open-secret-battling-ai-100100156.html>
 - [Issue 41] **Disney’s J.A.R.V.I.S. AI live for 2,000+ Imagineers** — seven decades of Disney creative knowledge at their fingertips — <https://wdwnt.com/2026/08/josh-damaro-outlines-how-new-ai-tool-will-be-used-by-disney-imagineering-for-attractions-more/>

11. *Celebrity / Director Positions on AI*

- [Issue 1] **James Cameron — “never going to take place”** — No Film School — <https://nofilmschool.com/james-cameron-ai#>
- [Issue 1] **Taylor Swift — criticized for AI in promo** — <https://tribune.com.pk/story/2570725/taylor-swift-criticised-for-using-ai-in-the-life-of-a-showgirl-promotional-campaign>
- [Issue 5] **Guillermo del Toro — “Rather Die”** — Variety — <https://variety.com/2025/film/news/guillermo-del-toro-rather-die-generative-ai-frankenstein-1236561316/>
- [Issue 5] **Paul Schrader on AI** — perfect script — <https://www.hollywoodreporter.com/movies/movie-news/paul-schrader-first-ai-movie-1236409606/>
- [Issue 7] **Jeremy Renner lawsuit threat** — multi-millions over AI voice — <https://deadline.com/2025/11/jeremy-renner-lawsuit-threat-ai-movie-1236611830/>
- [Issue 7] **George Clooney on AI actors** — Variety column — <https://variety.com/2025/scene/columns/george-clooney-ai-actors-movie-stars-1236579661/>
- [Issue 10] **James Cameron “horrifying”** — The Guardian — <https://www.theguardian.com/film/2025/dec/01/james-cameron-says-ai-actors-are-horrifying-to-me>
- [Issue 10] **Jenna Ortega “very easy to be terrified”** — NME — <https://www.nme.com/news/jenna-ortega-says-its-very-easy-to-be-terrified-of-ai-in-filmmaking-3913926>
- [Issue 11] **Leonardo DiCaprio — AI can’t be art** — THR — <https://www.hollywoodreporter.com/movies/movie-news/leonardo-dicaprio-ai-cant-be-art-no-humanity-1236445405/>
- [Issue 11] **James Cameron rejects AI actors at Hainan Festival** — <https://variety.com/2025/film/news/james-cameron-rejects-ai-actors-hainan-wouldnt-do-it-1236604204/>
- [Issue 14] **Claire Foy no interest in AI films** — Daily Mail — <https://www.dailymail.co.uk/tvshowbiz/article-15454199/Claire-Foy-AI-films-sad-disappointed-people-future-Hollywood.html>
- [Issue 14] **Wu-Tang Clan RZA** — case for AI in film/music — <https://www.vice.com/en/article/wu-tang-clans-rza-makes-the-case-for-ai-in-film-and-music-an-amazing-thing-for-us/>
- [Issue 15] **Matthew McConaughey protects voice/image** — Lawyer Monthly — <https://www.lawyer-monthly.com/2026/01/matthew-mcconaughey-protects-voice-image-ai/>

- [Issue 15] **Mara Wilson deepfake apocalypse fear** — Deadline — <https://deadline.com/2026/01/matilda-mara-wilson-stranger-things-ai-deepfake-apocalypse-1236689474/>
- [Issue 34] **Ron Howard bullish on AI** — “Jim Henson would have been too” — <https://www.hollywoodreporter.com/business/digital/ron-howard-ai-jim-henson-martin-scorsese-1236620457/>
- [Issue 34] **Gore Verbinski calls for an AI ‘rating system’ in films** — “If you use AI to write a script, you get an F” — <https://variety.com/2026/film/global/gore-verbinski-ai-taormina-1236780502/>
- [Issue 34] **Emily Blunt refused Spielberg’s offer to use AI for her alien voice** — <https://www.aol.com/articles/emily-blunt-refused-steven-spielbergs-163209822.html>
- [Issue 34] **Lionel Richie trademarks the sound of his voice** — ‘Hello, Is It Me You’re Looking For?’ — <https://www.billboard.com/pro/lionel-richie-is-trademarking-the-sound-of-his-voice/>
- [Issue 35] **SZA hits out at ‘disgusting’ AI music** — after discovering her songs in training data — <https://www.euronews.com/culture/2026/06/22/sza-hits-out-at-disgusting-ai-music-after-discovering-her-songs-were-used-in-ai-training>
- [Issue 36] **Backstreet Boys trademark their voices against deepfakes** — joining Taylor Swift and Lionel Richie — <https://www.digitalmusicnews.com/2026/06/26/backstreet-boys-trademark/>
- [Issue 36] **Madonna: AI is ‘the complete opposite of making art’** — <https://www.digitalmusicnews.com/2026/06/28/madonna-ai-algorithm-commentary/>
- [Issue 38] **George Lucas says AI makes filmmaking ‘easier’ and ‘there’s nothing you can do about it’** — likens AI sceptics to luddites clinging to horses — <https://variety.com/2026/film/news/george-lucas-test-screenings-embraces-ai-hollywood-1236810656/>
- [Issue 41] **Brad Pitt says AI could help more ‘mid-budget films’ get made** — if it’s ‘used as a tool’ — <https://variety.com/2026/film/news/brad-pitt-ai-could-help-mid-budget-movies-1236831201/>

12. Music Industry — Labels, Deals, and Lawsuits

- [Issue 1] **Universal & Warner — landmark AI deals within weeks** — Musically — <https://musically.com/2025/10/02/report-umg-and-wmg-could-sign-landmark-ai-deals-within-weeks/>
- [Issue 5] **Universal Music + Stability AI Alliance** — *op. cit.*
- [Issue 5] **Udio + UMG Partnership** — *op. cit.*

- [Issue 7] **GEMA v OpenAI Munich ruling** — *op. cit.*
- [Issue 7] **“Biggest theft in music history”** — Rights group sues Suno — <https://edm.com/gear-tech/rights-group-sues-suno-copyright-infringement/>
- [Issue 7] **Bangkok Post: Xania Monet \$3M deal** — <https://www.bangkokpost.com/life/tech/3142355/ai-singer-xania-monet-signs-3m-deal-with-hollywood-media>
- [Issue 8] **Warner Music + Stability AI** — next-gen tools — <https://stability.ai/news/warner-music-group-and-stability-ai-join-forces-to-build-next-gen-tools>
- [Issue 9] **Johnny Cash estate sues Coca-Cola** — ELVIS Act — <https://completemusicupdate.com/johnny-cash-estate-uses-elvis-act-to-sue-coke-over-tribute-act-ad-soundtrack/>
- [Issue 12] **Splice + UMG Collaboration** — *op. cit.*
- [Issue 14] **UMG slams AI slop** — *op. cit.*
- [Issue 16] **Wixen \$50M lawsuit against Meta** — <https://www.musicbusinessworldwide.com/wixen-files-50m-copyright-suit-against-meta-claims-tech-giant-wants-to-replace-songwriters-with-ai/>
- [Issue 17] **UMG \$3B suit against Anthropic** — *Dream Machine* coverage
- [Issue 31] **Spotify + Universal Music Group landmark AI covers/remixes deal** — Premium add-on; the first major streamer × major-label licensed AI-music product — June 2026
- [Deep Dive №14 – 4 Jun 2026] **Spotify–UMG three Cs framework** — *consent / credit / compensation* contract architecture; *Verified by Spotify* badge; opt-in catalogue; new royalty stream; deal’s prehistory in UMG settlement with Udio (Oct 2025) — <https://newsroom.spotify.com/2026-05-21/universal-music-group-spotify-licensing-agreements-fan-made-covers-remixes/>
- [Deep Dive №14 – 4 Jun 2026] **Hasbro Sixth Wall + CharacterOS** — **Behavioural Licensing** — 12 characters; runtime governance layer; talent-participation pay-out; 13+ /enterprise focus — <https://www.businesswire.com/news/home/20260603297922/en/>
- [Deep Dive №14 – 4 Jun 2026] **Disney–OpenAI deal** — **corrected record** — announced 11 Dec 2025, never executed; no formal agreement signed, no money changed hands; Sora wind-down announced 24 Mar 2026; framework survived platform’s death as reference contract architecture — <https://www.ropesgray.com/en/insights/viewpoints/the-deal-that-wasnt-what-disney-and-openai-teach-us-about-strategic-investments>
- [Deep Dive №14 – 4 Jun 2026] **Grimes / Elf.tech (GrimesAI)** — ~50% master-royalty split with creators; ~15,000 tracks generated; commons-with-royalty-hook model — <https://killthedj.com/how-artists-lend-their-ai-voices-to-creators-fans/>

- [Deep Dive №14 – 4 Jun 2026] **Holly Herndon / Holly+ DAO** — collective governance for commercial-use decisions; *Have I Been Trained?* (Spawning) consent-layer tooling
- [Deep Dive №14 – 4 Jun 2026] **ElevenLabs heritage voices (with CMG Worldwide)** — Judy Garland, James Dean, Burt Reynolds, Laurence Olivier; Iconic Marketplace expansion to Stan Lee estate and Michael Caine — <https://www.cbsnews.com/news/ai-voices-dead-celebrities-judy-garland-james-dean-burt-reynolds/>
- [Deep Dive №14 – 4 Jun 2026] **Meta celebrity chatbot anti-pattern (2023)** — Snoop Dogg, Kendall Jenner, MrBeast and others reportedly paid up to ~\$5m each for closed, platform-controlled personas — <https://english.elpais.com/technology/2023-11-13/from-paris-hilton-to-snoop-dogg-meta-to-launch-speaking-ai-chatbots-played-by-celebrities.html>
- [Deep Dive №14 – 4 Jun 2026] **Suno + Amazon Alexa+ integration** — *create-to-consume* demand floor; on-the-fly personalised songs via voice request — <https://www.billboard.com/pro/amazon-partners-ai-music-company-suno-alexa/>
- [Issue 33] **US musicians' union (AFM) sues Universal and Warner over AI deals** — claimed breach of the 'new use' clause in the industry's performer agreement — <https://completemusicupdate.com/us-musicians-union-sues-universal-and-warner-over-ai-deals-claims-breach-of-new-use-clause-in-record-industrys-performer-agreement/>
- [Issue 33] **Federal judge vacates order sealing Udio's training data** — the ingestion numbers move towards the public record — <https://www.digitalmusicnews.com/2026/06/05/udio-training-number-order-vacated/>
- [Issue 34] **NMPA unveils AI licensing deals with Udio and Klay** — 50/50 split for songs and recordings — <https://completemusicupdate.com/nmpa-unveils-ai-licensing-deals-with-udio-and-klay-with-50-50-split-for-songs-and-recordings/>
- [Issue 34] **Google moves to dismiss artists' AI lawsuit** — argues artists consented "to the use of their recordings" when uploading to YouTube — <https://www.digitalmusicnews.com/2026/06/10/youtube-ai-lawsuit-dismissal-motion/>
- [Issue 35] **Open letter: musicians 'must not be pressured into AI deals' by rightsholders** — <https://musically.com/2026/06/22/open-letter-musicians-must-not-be-pressured-into-ai-deals-by-rightsholders/>
- [Issue 36] **YouTube's AI training argument alarms indie music advocates** — 'not informed consent' —

<https://www.billboard.com/pro/youtube-ai-training-argument-indie-music-community-concerns/>

- [Issue 37] **Suno seizes on Udio's legal victory against Sony Music** — over admissible infringed works — <https://www.digitalmusicnews.com/2026/07/05/suno-admissible-infringed-works-argument/>
- [Issue 39] **Major labels and top indies unite to ban unlicensed AI music from chart rankings** — <https://www.digitalmusicnews.com/2026/07/29/labels-indies-unite-to-ban-ai-music-charts/>
- [Issue 40] **Suno loses landmark AI copyright lawsuit to GEMA** — German music rights society prevails — <https://variety.com/2026/digital/news/suno-loses-ai-lawsuit-gema-1236825010/>
- [Issue 40] **Spotify partners with Merlin to let fans legally create AI remixes and covers** — licensed fan-made derivatives of catalogue tracks — <https://newsroom.spotify.com/2026-08-04/merlin-spotify-licensing-agreements-fan-made-covers-remixes/>
- [Issue 41] **Warner Music expects AI music deals to boost subscription streaming revenue from late 2026** — <https://www.musicbusinessworldwide.com/warner-music-expects-ai-music-deals-to-boost-subscription-streaming-revenue-from-late-2026/>

13. Copyright, Policy and Regulation

- [Issue 9] **EU Lawmakers minimum age for AI/social** — Reuters — <https://www.reuters.com/legal/litigation/european-lawmakers-seek-eu-wide-minimum-age-access-ai-chatbots-social-media-2025-11-26/>
- [Issue 12] **UK DSIT Statement of Progress on Copyright and AI** — 88% — <https://www.gov.uk/government/publications/copyright-and-artificial-intelligence-progress-report/copyright-and-artificial-intelligence-statement-of-progress-under-section-137-data-use-and-access-act>
- [Issue 12] **IPWatchdog UK consultation analysis** — <https://ipwatchdog.com/2025/12/16/respondents-uk-ai-consultation-overwhelmingly-want-ai-companies-license-copyrighted-works-all-cases/>
- [Issue 11] **NY AI Advertising Disclosure Law** — Verge — <https://www.theverge.com/news/842848/new-york-law-ai-advertisements-sag-aftra-labor>
- [Issue 14] **Grok app ban consideration** — Fast Company — <https://www.fastcompany.com/91474131/governments-around-the-world-are-considering-bans-on-groks-app-over-ai-sexual-image-scandal>

- [Issue 21] **UK DSIT Final Copyright Report** — walked-back position
- [Issue 28] **Academy “You must be human to win” rule** — 2026 awards
- [Issue 29] **Cannes AI Disclosure Standard** — industry coordination
- [Issue 32] **Lisa Nandy: Government “got it wrong” on AI and creative copyright** — Culture Secretary’s admission at SXSW London (1 June 2026) that the opt-out model was a mistake; no decisions on next steps — <https://www.screendaily.com/news/we-got-it-wrong-says-uk-culture-secretary-lisa-nandy-of-government-misstep-on-ai-and-copyright-strategy/5217330.article>
- [Issue 33] **UK publishes AI Adoption Plan for the Creative Industries** — the policy follow-through on the Nandy admission — <https://www.gov.uk/government/publications/ai-champions-ai-adoption-plans/ai-adoption-plan-creative-industries>
- [Issue 33] **DSIT launches £1.1 billion AI Hardware Plan for the UK** — chips, computing power and skills — <https://www.gov.uk/government/news/a-decisive-shift-to-power-british-ai-new-11-billion-plan-to-back-chip-firms-boost-computing-power-and-skills-for-the-ai-revolution>
- [Issue 33] **Google forced to allow news sites to opt out of AI scraping** — UK watchdog’s ‘world first’ — <https://fortune.com/2026/06/05/auto-draft-3/>
- [Issue 34] **UK regulator introduces “world first” controls on content use in Google generative AI features** — <https://natlawreview.com/article/uk-regulator-introduces-world-first-controls-use-content-google-generative-ai>
- [Issue 36] **Google doubles down on ‘fair use’ AI training** — enforcement should focus on ‘outputs’, not ‘inputs’ — <https://www.digitalmusicnews.com/2026/06/29/google-fair-use-ai-training/>
- [Issue 39] **BAFTA adds AI clause to submission rules** — disclosure requirement from the 2027 awards — <https://deadline.com/2026/07/bafta-ai-clause-submission-rules-2027-1237000067/>
- [Issue 40] **AI labels to be compulsory on authentic-looking content under EU rules** — <https://www.aol.co.uk/articles/ai-labels-compulsory-authentic-looking-094628000.html>
- [Issue 40] **Key EU AI Act rules now being enforced** — how it affects the music industry — <https://musically.com/2026/08/04/key-eu-ai-act-rules-now-being-enforced-so-how-does-that-affect-the-music-industry/>

14. Unions and Labour

- [Issue 1] **SAG-AFTRA Condemns Tilly Norwood** — *op. cit.*
- [Issue 1] **UK Equity Statement** — Hollywood Reporter — <https://www.hollywoodreporter.com/movies/movie-news/tilly-norwood-ai->

actress-uk-union-equity-sag-aftra-debate-1236391739/

- [Issue 11] **Equity 99% Strike Vote** — landslide for industrial action — <https://www.equity.org.uk/news/2025/performers-prepared-to-take-industrial-action-over-ai-in-landslide-99-vote>
- [Issue 11] **NY AI Advertising Disclosure Law / SAG-AFTRA quote** — *op. cit.*
- [Issue 15] **Equity welcomes improved offer** — film/TV AI protections — <https://www.equity.org.uk/news/2026/equity-welcomes-improved-offer-in-ai-protection-negotiations-in-film-and-tv>
- [Issue 26 / 29] **SAG-AFTRA “Tilly Tax” contract provisions** — final spring 2026 contract
- [Issue 35] **More than 16,000 sign SAG-AFTRA letter demanding the NO FAKES Act** — congressional pressure on unauthorised AI images and videos — <https://variety.com/2026/biz/news/sag-aftra-letter-congress-pass-no-fakes-act-1236783056/>
- [Issue 35] **The real battle is over likeness and who gets paid** — Hollywood’s AI economics beneath the tooling debate — <https://www.staradvertiser.com/2026/06/20/breaking-news/as-hollywood-leans-into-ai-the-real-battle-is-over-likeness-and-who-gets-paid/>
- [Issue 36] **Nearly 1,000 sign open letter over child actors’ AI voice clauses** — the Peppa Pig contract backlash — <https://variety.com/2026/tv/news/open-letter-major-studio-hasbro-children-ai-peppa-pig-1236790351/>
- [Issue 38] **Session musicians’ AI pay clause cannot function, Warner Music tells court** — <https://www.techtimes.com/articles/320485/20260714/session-musicians-ai-pay-clause-cannot-function-warner-music-tells-court.htm>
- [Issue 39] **The AI game-dev boom has a dark side: fewer jobs for people starting out** — <https://www.digitaltrends.com/cool-tech/the-ai-game-dev-boom-has-a-dark-side-fewer-jobs-for-the-people-starting-out/>
- [Issue 41] **UMG asks court to dismiss US musicians’ union lawsuit over Udio AI deal** — <https://www.musicbusinessworldwide.com/umg-asks-court-to-dismiss-us-musicians-union-lawsuit-over-udio-ai-deal/>

15. Audience Response and Slop Ceiling

- [Issue 7] **Deezer/Ipsos AI Music Survey** — 97% can’t tell, but care when told — <https://newsroom-deezer.com/2025/11/deezer-ipsos-survey-ai-music/>
- [Issue 7] **50,000 AI tracks uploaded to Deezer daily** — Musically — <https://musically.com/2025/11/12/50000-ai-music-tracks-are-now-uploaded-to-deezer-every-day/>

- [Issue 8] **MrBeast on AI** — threat to creators — <https://www.forbes.com/sites/johnbbrandon/2025/10/10/mrbeast-is-right-about-ai-content-but-are-we-really-in-scary-times/>
- [Issue 12] **Merriam-Webster Word of the Year: “Slop”** — Hollywood Reporter — <https://www.hollywoodreporter.com/news/general-news/slop-word-year-2025-merriam-webster-1236450780/>
- [Issue 12] **YouTube AI channels — 1.2bn views fake politics** — Guardian — <https://www.theguardian.com/technology/2025/dec/13/fake-anti-labour-video-billion-views-youtube-2025>
- [Issue 14] **Spotify AI flood — subscribers furious** — TechRadar — <https://www.techradar.com/audio/spotify/ai-music-is-flooding-spotify-and-subscribers-are-furious-heres-why-music-fans-no-longer-trust-discover-weekly>
- [Issue 14] **Soultracks: “AI music is catchy, familiar... and boring”** — <https://soultracks.com/news-ai-generated-music-is-catchy-boring/>
- [Issue 15] **Sweden bans AI from official chart** — Independent — <https://www.independent.co.uk/tv/news/ai-music-song-banned-sweden-spotify-b2901627.html>
- [Issue 16] **YouTube CEO: managing AI slop on priority list 2026** — Digital Music News — <https://www.digitalmusicnews.com/2026/01/22/youtube-ceo-ai-slop-2026-comments/>
- [Issue 16] **Bain & Co — “People still want the radio star”** — <https://www.bain.com/insights/in-an-ai-age-people-still-want-the-radio-star/>
- [Issue 25-28] **Deezer April 2026 data — 44% / 3%** — newsroom release — <https://newsroom-deezer.com/2026/04/ai-generated-tracks-represent-44-of-new-uploaded-music/>
- [Issue 31] **Apple Music: AI <1% of plays** — clean platform-level confirmation — June 2026
- [Issue 31] **Daniel Ek defends AI music: “better than slop”** — Spotify CEO framing — June 2026
- [Issue 31] **Epidemic Sound *Future of the Creator Economy Report 2026*** — 94% of creators use AI; 89% feel pressure to use AI; 75% believe human-created content will become premium; survey of 3,000 professional creators — June 2026
- [Issue 33] **‘Say No To Suno’ protesters at Santa Monica AI conference** — organised anti-AI-music protest reaches the industry circuit — <https://www.digitalmusicnews.com/2026/06/04/say-no-to-suno-santa-monica-plane/>
- [Issue 35] **Faceless creators take a hit as YouTube cracks down on AI slop** — <https://www.hollywoodreporter.com/business/digital/faceless->

creators-youtube-ai-damage-1236617586/

- [Issue 36] **181,000 mobile games released in six months** — AI and vibe coding flood the market, “but not necessarily better ones” — <https://www.digitaltrends.com/gaming/ai-and-vibe-coding-are-making-more-games-but-not-necessarily-better-ones/>
- [Issue 37] **Our interest in AI slop is hitting a ceiling** — Fast Company on the audience’s tolerance limit — <https://www.fastcompany.com/91568818/our-interest-in-ai-slop-is-hitting-a-ceiling>
- [Issue 40] **AI-generated stories rated better quality than human-written ones, study finds** — audiences often can’t tell which is which — <https://www.theguardian.com/technology/2026/aug/05/ai-generated-stories-rated-better-quality-than-human-written-ones-study-finds>
- [Issue 40] **Snapchat will no longer recommend AI slop** — platform pulls back on surfacing synthetic content — <https://www.socialmediatoday.com/news/snapchat-will-no-longer-recommend-ai-slop/826782/>
- [Issue 41] **Luminate study: 48-point generational gap on AI music** — 13–17s at +19% net interest versus 18–24s at –29% — <https://musically.com/2026/07/27/luminate-study-explores-americans-interest-in-genai-content/>

16. Provenance, Watermarking, Detection

- [Issue 2] **OpenAI likeness protections** — Digital Music News — <https://www.digitalmusicnews.com/2025/10/08/openais-likeness-protections-dont-apply-to-dead-celebrities/>
- [Issue 11] **SynthID rollout across Veo / Lyria / Imagen** — Google DeepMind
- [Issue 12] **Gemini “Is this AI?” video verification** — https://www.linkedin.com/posts/googledeepmind_verify-google-ai-generated-videos-in-the-activity-7407748300688478208-fJgW
- [Issue 13] **Instagram chief — “fingerprint real media”** — Digital Music News — <https://www.digitalmusicnews.com/2026/01/05/instagram-chief-ai-slop-comments/>
- [Issue 13] **Instagram head AI verification** — WebProNews — <https://www.webpronews.com/instagram-head-warns-ai-images-erode-trust-calls-for-verification-standards/>
- [Issue 18] **Deezer licenses its AI-music detection tool** — *Dream Machine* coverage
- [Issue 29] **YouTube false-positive: Tiny Grandma stop-motion** — flagged as AI

- [Issue 31] **YouTube auto-labels AI videos** — platform-default disclosure regime — June 2026
- [Issue 31] **LinkedIn limits reach of AI-generated content** — distribution-side enforcement — June 2026
- [Issue 31] **TikTok AI moderation false positives** — wrongly punished independent filmmakers — June 2026
- [Issue 33] **Sony invests in AI copyright-protection firm Midnight Labs** — <https://variety.com/2026/biz/news/ip-ai-firm-midnight-labs-sony-million-dollar-investment-1236766931/>
- [Issue 34] **Deezer launches an AI music detector for other streaming services** — detection as an exportable product — <https://www.engadget.com/2192024/deezer-ai-music-detector/>
- [Issue 34] **Warner Music Group acquires AI detection company Sureel** — <https://www.hollywoodreporter.com/music/music-industry-news/warner-music-group-buys-ai-detection-company-sureel-1236618266/>
- [Issue 34] **Played By Humans** — verifying ‘human artistry in an AI age’ — <https://musically.com/2026/06/11/played-by-humans-is-verifying-human-artistry-in-an-ai-age/>
- [Issue 35] **RTM Audio launches AI music detector UAI** — a signed certificate with every verdict — <https://www.musicbusinessworldwide.com/tech-startup-rtm-audio-launches-ai-music-detector-uai-which-issues-a-signed-certificate-with-every-verdict/>
- [Issue 35] **Artists can now check whether their music trained AI** — The Atlantic’s dataset lookup — <https://mixmag.net/read/artists-check-music-used-train-ai-atlantic-investigation-news>
- [Issue 36] **ElevenLabs partners with Google DeepMind on SynthID** — detecting ElevenLabs-generated audio — <https://elevenlabs.io/blog/synthid>
- [Issue 36] **Human Consent Registry** — Cate Blanchett, Steven Soderbergh and Nikki Hexum on identity protection from AI theft as a ‘human right’ — <https://variety.com/2026/biz/news/cate-blanchett-steven-soderbergh-nikki-hexum-ai-human-registry-1236791271/>
- [Issue 36] **Traxsource to label human-made and AI-assisted music** — <https://musictech.com/news/music/traxsource-labels-identify-human-made-and-ai-assisted-music/>
- [Issue 38] **RIAA, A2IM, IMPALA announce track-level ‘labeling program’ for AI audio** — targeting broad global adoption across DSPs — <https://www.digitalmusicnews.com/2026/07/10/ai-music-tags-riaa-a2im-impala/>
- [Issue 40] **NVIDIA creates a tool for detecting AI content in videos** — designed for newsroom use — <https://www.broadcastnow.co.uk/tech->

innovation/nvidia-creates-ai-video-detector/5219075.article

- [Issue 41] **Spotify introduces a new ‘AI Persona’ label** — transparency badges for AI-generated artist identities — <https://newsroom.spotify.com/2026-08-11/ai-persona-badges-transparency/>

17. *Workplace AI Adoption / Shadow AI / Workforce Research*

- [Issue 1] **MIT study — AI reduces brain activity** — AI News — <https://www.artificialintelligence-news.com/news/ai-causes-reduction-in-users-brain-activity-mit/>
- [Issue 1] **51% of Japanese games studios use AI** — *op. cit.*
- [Issue 1] **Yale on AI adoption** — Neil Hoyne — https://www.linkedin.com/posts/neilhoynes_ai-data-research-activity-7379272781798035456-hnuV
- [Issue 5] **Azumo AI in Workplace Statistics 2025** — <https://azumo.com/artificial-intelligence/ai-insights/ai-in-workplace-statistics>
- [Issue 5] **Tech.co — Gen Z most likely use AI behind boss’s back** — <https://tech.co/news/gen-z-most-likely-use-ai-boss>
- [Issue 5] **IDC Europe Shadow AI security nightmare** — <https://blog-ideurope.com/shadow-ai-how-stealth-productivity-is-strangling-enterprise-ai-adoption-and-creating-a-security-nightmare/>
- [Issue 5] **Forbes — AI Tools Flood Workplaces** — <https://www.forbes.com/sites/carolinecastrillon/2025/09/09/ai-tools-flood-workplaces-as-employees-face-a-double-bind/>
- [Issue 5] **Exploding Topics AI Workforce Research** — <https://explodingtopics.com/blog/ai-workforce-research>
- [Issue 6] **Adobe Creators’ Toolkit Report (16,000 creators)** — <https://news.adobe.com/news/2025/10/adobe-max-2025-creators-survey>
- [Issue 7] **CNBC — ADHD, autism, dyslexia and AI agents** — <https://www.cnn.com/2025/11/08/adhd-autism-dyslexia-jobs-careers-ai-agents-success.html>
- [Issue 8] **Forbes Vibe Coding \$220K** — <https://www.forbes.com/sites/rachelwells/2025/11/06/the-in-demand-ai-skill-and-certifications-that-pays-up-to-220000/>
- [Issue 8] **LANDR — 87% of musicians use AI tools** — <https://aristake.com/ai-tools-musicians-study/>
- [Issue 9] **Economist — Investors expect AI use to soar (it isn’t)** — <https://www.economist.com/finance-and-economics/2025/11/26/investors-expect-ai-use-to-soar-thats-not-happening>

- [Issue 9] **Reuters Institute UK Journalists AI Survey** — <https://reutersinstitute.politics.ox.ac.uk/ai-adoption-uk-journalists-and-their-newsrooms-surveying-applications-approaches-and-attitudes>
- [Issue 12] **Economist — Job apocalypse? Humbug!** — <https://www.economist.com/business/2025/12/14/job-apocalypse-humbug-ai-is-creating-brand-new-occupations>
- [Issue 16] **Guardian — AI is hitting UK harder** — <https://www.theguardian.com/technology/2026/jan/26/ai-uk-jobs-us-japan-germany-australia>
- [Issue 16] **McKinsey AI for film and TV** — <https://www.mckinsey.com/industries/technology-media-and-telecommunications/our-insights/what-ai-could-mean-for-film-and-tv-production-and-the-industrys-future>
- [Issue 16] **PRS for Music AI Survey 2026** — <https://www.prsformusic.com/m-magazine/news/prs-for-music-ai-survey-2026>
- [Issue 24] **OpenAI public-policy on disruption** — robot tax, 4-day workweek, wealth funds
- [Issue 31] **Epidemic Sound *Future of the Creator Economy Report 2026*** — 94% of creators use AI; 89% feel pressure; 75% premium-human-content thesis — *Musically* coverage — June 2026
- [Issue 31] **Coordination Collapse research report (O’Hare / DreamLab AI Consulting)** — April 2026
- [Issue 31] **The Process Trap** — generative abundance and the equivocation of the creative
- [Issue 31] **The Doomer Mistake** — empirical rebuttal of the AI-driven civilizational-collapse case
- [Issue 34] **PwC: AI hiring surges across all UK sectors** — 34% salary premium for AI skills — <https://www.uktech.news/ai/ai-hiring-surges-across-all-uk-sectors-20260615>
- [Issue 34] **Ramp AI Index** — top 1% of firms now spend \$7,500 per employee per month on AI — <https://ramp.com/data/ai-index-june-2026>
- [Issue 36] **Hollywood workers are training AI models as job prospects grow slim** — <https://www.hollywoodreporter.com/business/digital/ai-training-hollywood-writer-jobs-prospects-1236628302/>
- [Issue 36] **Top CFOs: AI success depends on training employees, not just buying technology** — <https://fortune.com/2026/06/25/cfo-hpe-prologis-smucker-moodys-ai-success-training-employees/>

18. Advertising and Brand AI

- [Issue 3] **WPP \$400M Google partnership** — <https://martechseries.com/predictive-ai/ai-platforms-machine-learning/google-and-spotify-alum-launch-epiminds-with-6-6m-to-build-marketing-teams-for-the-ai-era/>
- [Issue 5] **WPP Open Pro launch** — <https://campaignbrief.com/wpp-launches-ai-powered-marketing-platform-wpp-open-pro/>
- [Issue 5] **Mondelez AI for TV ads** — Verge — <https://www.theverge.com/news/806047/mondelez-ai-generated-ads>
- [Issue 6] **WPP + Slightly partnership** — Digiday — <https://digiday.com/media-buying/agencies-continue-to-expand-ai-capabilities-to-boost-brand-performance/>
- [Issue 6] **Coca-Cola AI Holiday ad (2nd attempt)** — Adweek — <https://www.adweek.com/creativity/coca-cola-uses-ai-to-rekindle-the-magic-of-its-holiday-ads/>
- [Issue 7] **Digiday — AI agent developers in-demand role** — <https://digiday.com/marketing/ai-agent-developers-have-become-adlands-in-demand-role/>
- [Issue 10] **Valentino “disturbing” AI handbag ads** — BBC — <https://www.bbc.co.uk/news/articles/cwyvjyvn83go>
- [Issue 11] **McDonald’s NL Christmas AI Ad pulled** — Branding in Asia — <https://www.brandinginasia.com/its-the-most-terrible-time-of-the-year-mcdonalds-netherlands-wonderfully-chaotic-ai-driven-christmas-film/>
- [Issue 11] **Channel 4 AI Ads** — Estate Agent Today — <https://www.estateagenttoday.co.uk/breaking-news/2025/12/homebuilder-among-first-to-use-channel-4s-ai-ads/>
- [Issue 14] **Marketing Week: AI ads winning in testing** — <https://www.marketingweek.com/dismiss-ai-ads-winning-creative-effectiveness/>
- [Issue 15] **PGA Tour + AWS expanded partnership** — Pymnts — <https://www.pymnts.com/artificial-intelligence-2/2026/ai-content-is-par-for-the-course-with-pga-tours-expanded-aws-partnership/>
- [Issue 15] **Avocados from Mexico skip TV for AI** — Digiday — <https://digiday.com/marketing/avocados-from-mexico-turns-to-ai-to-advertise-around-the-super-bowl-instead-of-a-tv-buy/>
- [Issue 16] **Higgsfield + Madonna AI video** — Adweek — <https://www.adweek.com/media/higgsfield-ai-marketing-startup/>
- [Issue 27] **WPP + Google Earth AI consumer journey** — *Dream Machine* coverage

- [Issue 33] **WPP tests a different type of ad talent with new studio Hex** — <https://www.marketingdive.com/news/wpp-puts-different-type-of-ad-talent-to-the-test-with-new-studio-hex/822227/>
- [Issue 34] **AI disclosure labels have minimal effect on ad performance** — MediaScience / Adelaide University study on brand recall and sentiment — <https://www.marketingdive.com/news/ai-disclosure-labels-dont-hurt-ad-performance-heres-what-the-numbers-say/822711/>
- [Issue 35] **OpenAI moves to automate ad creative** — <https://digiday.com/marketing/openai-moves-to-automate-ad-creative/>
- [Issue 35] **Brands using AI-generated influencers to promote products** — <https://www.theguardian.com/technology/2026/jun/21/brands-using-ai-generated-influencers-to-promote-products-on-social-media>
- [Issue 36] **TikTok launches Agentic Hub** — third-party AI ad tools integrated with the platform’s MCP server — <https://www.socialmediatoday.com/news/tiktok-launches-agentic-hub/824169/>
- [Issue 41] **Warner Bros. uses AI dog podcast to promote ‘The End of Oak Street’** — a synthetic-host marketing experiment — <https://variety.com/2026/film/news/warner-bros-ai-dog-podcast-the-end-of-oak-street-1236830713/>

19. *Festivals, Institutions, Awards*

- [Issue 6] **AI FilmFest Japan / Hoyt Dwyer** — <https://www.prnewswire.com/news-releases/from-apple-tv-creative-to-ai-filmmaker-hoyt-dwyers-animated-film-to-compete-at-ai-filmfest-japan-2025-302598064.html>
- [Issue 6] **India’s first AI Film Festival** — IFFI, NFDC, LTIMindtree — <https://www.medianews4u.com/iffi-partners-with-ltimindtree-and-nfdc-to-launch-indias-first-ai-film-festival-and-hackathon/>
- [Issue 14] **Tunisian filmmaker wins \$1M for Lily** — *op. cit.*
- [Issue 14] **Comic-Con Art Show allows AI** — Filmstories — <https://filmstories.co.uk/news/san-diego-comic-con-art-show-to-allow-ai-slop/>
- [Issue 14] **Emmys AI Guidance** — THR — <https://www.hollywoodreporter.com/tv/tv-news/emmys-ai-guidelines-2026-awards-1236468434/>
- [Issue 15] **Sundance AI Literacy Initiative** — Sundance Institute blog — <https://www.sundance.org/blogs/centering-the-artist-why-were-launching-the-ai-literacy-initiative/>
- [Issue 15] **Google \$2M Sundance AI Education** — <https://blog.google/company-news/outreach-and-initiatives/google->

org/sundance-institute-ai-education/

- [Issue 16] **CNET — San Diego Comic-Con bans AI art at 2026 event** — <https://www.cnet.com/culture/san-diego-comic-con-bans-ai-art-for-2026-event/>
- [Issue 16] **Adobe Sundance Film Festival 2026** — <https://blog.adobe.com/en/publish/2026/01/20/sundance-film-festival-2026-creativity-community-power-storytelling>
- [Issue 28] **Academy “human to win” rule** — *Dream Machine* coverage
- [Issue 29] **Cannes AI Disclosure Standard launched** — *Dream Machine* coverage
- [Issue 33] **Tribeca doc ‘AI: Probably Nothing to Worry About’** — Nick Holt on AI’s origins as “the creation of a sort of species” — <https://variety.com/2026/film/festivals/tribeca-ai-documentary-nick-holt-director-probably-nothing-to-worry-about-1236767764/>
- [Issue 35] **Dataland — Museum of AI Art opens in Los Angeles** — Refik Anadol’s institution-scale landmark — <https://www.theartnewspaper.com/2026/06/18/refik-anadol-dataland-opens-los-angeles>
- [Issue 35] **Shanghai Film Festival embraces cinema’s future** — from AI to iPhone moviemaking — <https://www.hollywoodreporter.com/movies/movie-news/shanghai-film-festival-cinema-future-ai-iphone-1236626463/>
- [Issue 40] **Es Devlin’s ‘A National Portrait’ at the National Portrait Gallery** — every UK adult can become a co-author of the AI-driven artwork — <https://www.npg.org.uk/whatson/display/2026/national-portrait-for-npg/>

20. Geographic / Regional AI Production

- [Issue 5] **CNBC Africa on AI in African music** — <https://www.cnbc africa.com/2025/how-ai-is-changing-the-landscape-of-the-music-industry-in-africa>
- [Issue 5] **Run to the West — South Korea’s first AI feature** — *op. cit.*
- [Issue 5] **Watch the Skies — Swedish AI dubbing** — *op. cit.*
- [Issue 6] **India’s first AI Film Festival** — *op. cit.*
- [Issue 12] **Trilok — Indian AI band** — Musically — <https://musically.com/2025/12/17/indian-ai-band-trilok-performs-live-government-denies-association/>
- [Issue 13] **56.9% of new Chinese independent songs are AI** — Musically — <https://musically.com/2026/01/05/report-56-9-of-new-independent-songs-in-china-are-ai-generated/>

- [Issue 14] **BBC Future — Lights, camera, algorithm (India)** — <https://www.bbc.co.uk/future/article/20251223-why-indian-cinema-is-awash-with-ai>
- [Issue 14] **Shift Up CEO on AI vs China/US scale** — <https://www.pocketgamer.biz/shift-up-ceo-says-ai-is-key-to-competing-with-chinas-game-industry-scale/>
- [Issue 25] **Indonesia’s Legenda Bertuah** — first AI-animated series
- [Issue 27] **Korin AI — Africa-trained, Africa-built** — launch
- [Issue 27] **Latin American AI film festival wave** — *Dream Machine* coverage
- [Issue 31] **Astana AI Film Festival (\$1M prize fund)** — Kazakhstan AI film festival opens submissions — June 2026
- [Issue 31] **Tribeca: first fully AI-generated film** — set to premiere — June 2026
- [Issue 31] **Gareth Edwards (Star Wars director) in favour of AI** — “*like a billionaire on acid*” — June 2026
- [Issue 32] **Martin Scorsese backs Black Forest Labs (FLUX)** — advisor role, using AI for storyboarding; triggered Boots Riley and del Toro backlash — <https://variety.com/2026/film/news/martin-scorsese-supports-ai-company-storyboard-movies-1236765037/>
- [Issue 32] **Dreams of Violets world premiere at Tribeca** — first fully AI-generated feature accepted by a major festival; cost \$2,000; directed by Ash Koosha (Fountain o); uses Kling AI, Claude, Gemini — <https://variety.com/2026/film/festivals/tribeca-festival-ai-film-dreams-of-violets-fountain-o-1236759724/>
- [Issue 32] **Amazon AI Creators Fund + Punky Duck instant cancellation** — AI animated series commissioned and killed within 48 hours after audience derision; Jorge Gutierrez dropped out — <https://futurism.com/artificial-intelligence/amazon-ai-generated-series-canceled>
- [Issue 32] **NVIDIA Cosmos 3** — omnimodal world model; text, images, video, audio, actions; best open-source T2I and I2V per Artificial Analysis — <https://nvidianews.nvidia.com/news/nvidia-launches-cosmos-3-the-open-frontier-foundation-model-for-physical-ai>
- [Issue 35] **South Korea’s movie industry is embracing AI** — will K-film lose its magic? — <https://www.cnn.com/world/south-korea-film-industry-ai-hnk-spc>

- [Issue 35] **MiniMax launches all-in-one AI video hub at Shanghai Film Festival** — ‘creative direction and aesthetic judgment must ultimately be left to humans’ — <https://variety.com/2026/film/festivals/minimax-hub-ai-video-generator-shanghai-film-festival-1236781901/>
- [Issue 36] **Savvy Games Group expands AI access for Saudi developers** — partnership with Genvid and Massive Studios — <https://www.pocketgamer.biz/savvy-games-group-partners-with-genvid-and-massive-studios-to-expands-ai-access-for-saudi-game-developers/>
- [Issue 1] **UCL, RCA, Brandtech Centre for Creative AI launch** — Broadcast Now — <https://www.broadcastnow.co.uk/broadcasters/alex-mahon-joins-stellar-ai-creative-summit-line-up/5209227.article>
- [Issue 11] **Lovable for classrooms** — <https://lovable.dev/classroom>
- [Issue 15] **UW-Stout AI baseline competency in filmmaking** — *op. cit.*
- [Issue 15] **Google + Sundance Institute AI Education** — *op. cit.*
- [Issue 16] **UK government “Free AI training for all”** — <https://www.gov.uk/government/news/free-ai-training-for-all-as-government-and-industry-programme-expands-to-provide-10-million-workers-with-key-ai-skills-by-2030>
- [Issue 16] **Adobe at Sundance: Ignite Day** — *op. cit.*
- [Issue 35] **CoSTAR Emerge-AI opens** — accelerator for early-stage and SME creative-tech startups turning AI concepts into viable businesses — <https://www.costarnetwork.co.uk/calls/emerge-AI>

22. Open Source and Infrastructure

- [Issue 1] **ComfyUI raises \$17M** — *op. cit.*
- [Issue 1] **PlayCanvas SOG** — *op. cit.*
- [Issue 1] **DecartAI — open-source real-time world transformation** — <https://decart.ai/>
- [Issue 1] **Civitai / Replicate** — open infrastructure layer
- [Issue 4] **Krea Realtime open-sourced** — *op. cit.*
- [Issue 5] **Meta + Hugging Face OpenEnv** — *op. cit.*
- [Issue 8] **Hugging Face + Google Cloud partnership** — https://www.linkedin.com/posts/julienchaumond_i-am-super-excited-to-announce-that-hugging-activity-7396177403972276225-CuMM
- [Issue 8] **80% of A16Z pitches use Chinese open-source models** — https://www.linkedin.com/posts/stevenouri_a-wild-stat-80-of-startups-pitching-a16z-activity-7396182718998351872-xTKR

- [Issue 11] **Ubisoft CHORD open-sourced** — *op. cit.*
- [Issue 27] **ComfyUI \$500M valuation** — *Dream Machine* coverage
- [Issue 27] **Anthropic + Blender Foundation patronage** — *Dream Machine* coverage
- [Issue 27] **Korin AI launch** — *op. cit.*
- [Issue 33] **Comfy Desktop** — “one app for every Comfy” — <https://x.com/ComfyUI/status/2064093003590111314>
- [Issue 33] **Cosine secures industry backing for Britain’s first sovereign frontier model** — <https://tech.eu/2026/06/08/cosine-secures-industry-backing-for-britain-s-first-sovereign-frontier-model/>
- [Issue 36] **The Rise of ComfyUI** — 4M users, 89,000+ GitHub stars, 60,000+ community nodes, \$500m valuation; how one developer’s side project became the professional layer — <https://xhinker.medium.com/the-rise-of-comfyui-how-one-guys-side-project-took-over-ai-image-generation-558dca64bfd7>
- [Issue 36] **Comfy MCP public beta** — *op. cit.* (see §6)
- [Issue 36] **Roboflow supervision hits 40,000 GitHub stars** — powering 6.5k+ open-source computer vision projects — <https://github.com/roboflow/supervision>

23. Web Infrastructure / Bots / Dead Internet

- [Issue 4] **Imperva 2025 Bad Bot Report** — bots = 51% of web traffic — <https://www.imperva.com/blog/2025-imperva-bad-bot-report-how-ai-is-supercharging-the-bot-threat/>
- [Issue 4] **Cloudflare crawl-to-click gap** — <https://blog.cloudflare.com/crawlers-click-ai-bots-training/>
- [Issue 4] **Dead Internet Theory** — Wikipedia — https://en.wikipedia.org/wiki/Dead_Internet_theory
- [Issue 4] **Grand View Research GenAI Content Creation Market** — <https://www.grandviewresearch.com/industry-analysis/generative-ai-content-creation-market-report>
- [Issue 4] **Futurism — AI-only social network collapses into warring tribes** — <https://futurism.com/social-network-ai-intervention-echo-chamber>

24. AI Bubble / Economics / Investment

- [Issue 1] **AI bubble 17× dotcom** — PC Gamer — <https://www.pcgamer.com/software/ai/fabulous-news-everyone-market-analyst-says-the-ai-bubble-is-17x-bigger-than-the-dotcom-goldrush-and-4x-larger-than-the-subprime-bubble-that-caused-the-2008-crash/>
- [Issue 2] **Suno \$2.45B valuation** — *op. cit.*

- [Issue 32] **Suno \$400M raise at \$5.4B valuation** — \$300M ARR, 2M+ subscribers; WMG settled; UMG/Sony in active litigation — <https://variety.com/2026/digital/news/ai-music-suno-funding-round-400-million-5-4-billion-valuation-1236765727/>
- [Issue 32] **Tripo AI raises nearly \$200M** — 3D world model, Project Eden — <https://finance.yahoo.com/sectors/technology/articles/tripo-ai-raises-nearly-200-164700546.html>
- [Issue 32] **Roblox acquires Morpheus AI** — video world models for Roblox Reality vision — <https://about.roblox.com/newsroom/2026/06/pioneering-ai-founders-join-to-accelerate-roblox-reality-vision>
- [Issue 4] **AdsGency \$12M seed** — *op. cit.*
- [Issue 5] **Sifted — Synthesia rejects \$3B Adobe** — <https://sifted.eu/articles/synthesia-acquisition-offer>
- [Issue 14] **Kartel / Reilly leadership** — *op. cit.*
- [Issue 15] **Higgsfield \$80M raise at \$1.3B** — <https://siliconangle.com/2026/01/15/higgsfield-raises-80m-1-3b-valuation-scale-ai-video-platform/>
- [Issue 16] **Higgsfield earns \$200M in 9 months** — <https://eu.36kr.com/en/p/3650517574312323>
- [Issue 16] **Synthesia hits \$4B valuation** — *op. cit.*
- [Issue 21] **Adobe + NVIDIA Strategic Partnership** — *op. cit.*
- [Issue 25] **ElevenLabs \$500M ARR** — *Dream Machine* coverage
- [Issue 27] **Google \$40B investment in Anthropic** — *Dream Machine* coverage
- [Issue 27] **ComfyUI \$500M valuation** — *op. cit.*
- [Issue 35] **Odyssey \$310M Series B** — *op. cit.* (see §4)
- [Issue 36] **Higgsfield hits \$500M revenue, eyes \$5B funding round** — <https://www.techtimes.com/articles/319394/20260630/ai-video-startup-higgsfield-hits-500m-revenue-eyes-5b-funding-round.htm>
- [Issue 36] **General Intuition raises \$320M at \$2.3B valuation** — frontier models trained on gameplay; games as training data for real-world agents — <https://gamesbeat.com/general-intuition-raises-320m-at-2-3b-valuation-for-ai-frontier-models-based-on-gameplay-exclusive-interview/>
- [Issue 37] **Kling AI raises record \$2.8 billion** — the largest AI video raise to date, with strings under Beijing law — <https://www.techtimes.com/articles/319708/20260704/kling-ai-raises-record-28-billion-what-creators-owe-china-under-beijing-law.htm>

25. Live Sport, Events and the World Cup

The 2026 FIFA World Cup as the largest live deployment of AI in entertainment history — the Issue 34 Editor's Pick and the Deep Dive of 12 June 2026.

- [Issue 34] **Editor's Pick: The AI World Cup — When Sport Becomes a Living Digital System** — one-second player scans, digital twins, Football AI Pro, the multi-platform storyworld — *Dream Machine* Issue 34
- [Issue 34] **TacticAI (Google DeepMind)** — simulating field scenarios and predicting open play up to 8 seconds in advance — <https://x.com/GoogleDeepMind/status/2065093482088169719>
- [Deep Dive – 12 Jun 2026] **FIFA and Lenovo unveil AI-powered innovations for World Cup 2026** — 1,248 players scanned in one second each; semi-automated offside; AI-generated personalised highlights — <https://football-technology.fifa.com/organisation/media-releases/lenovo-tech-world-ai-powered-innovations-world-cup-2026>
- [Deep Dive – 12 Jun 2026] **How AI Will Power the 2026 FIFA World Cup** — AI Magazine on the full stack: officiating, Football AI Pro for all 48 teams, broadcast avatars — <https://aimagazine.com/news/how-ai-will-power-the-2026-fifa-world-cup>
- [Deep Dive – 12 Jun 2026] **Lenovo powers World Cup operations and AI-driven broadcast** — ultra-low-latency IPTV, stadium digital twins, 6bn viewers across 104 matches — <https://news.lenovo.com/pressroom/press-releases/lenovo-technology-powers-fifa-world-cup-2026-operations-and-strengthens-ai-driven-broadcast/>
- [Deep Dive – 12 Jun 2026] **FIFA and Gamefam launch World Cup 2026 on Roblox** — 1.1bn+ visits, 1.5M gameplay sessions a day; the tournament's second venue — <https://gamesbeat.com/gamefam-launch-fifa-world-cup-2026-super-soccer-roblox/>
- [Deep Dive – 12 Jun 2026] **TikTok's World Cup creator correspondents** — thirty creators from eleven countries; FIFA's "preferred platform" — <https://www.socialmediatoday.com/news/tiktok-introduces-2026-fifa-world-cup-creator-correspondents/820198/>
- [Deep Dive – 12 Jun 2026] **The biggest brand marketing race** — WARC's projected \$10.5bn in incremental Q2 2026 ad spend — <https://www.globalbrandsmagazine.com/fifa-world-cup-2026-marketing/>
- [Deep Dive – 12 Jun 2026] **\$4.5bn sponsorship economy for host cities** — \$725M in new host-city sponsorship in a year — <https://www.goal.com/en-qa/lists/industry-report-reveals-usd4-5-billion-sports-sponsorship-economy-for-2026-world-cup-host-cities-in-the-u-s-and-canada/blt412a2f99ee2cf077>
- [Deep Dive – 12 Jun 2026] **Adidas's ~\$1bn "Backyard Legends" campaign** — AI de-aging of three football legends; celebrity ensemble at record scale —

<https://missionmedia.asia/adidas-celebrity-marketing-world-cup-2026/>

- [Deep Dive – 12 Jun 2026] **How AI will shape the 2026 World Cup** — the fan-made content flood; 4,300+ fake FIFA websites before kick-off — <https://www.merciaai.com/post/how-ai-will-shape-the-2026-world-cup>
- [Deep Dive – 12 Jun 2026] **Boston Dynamics robot dogs on World Cup security** — Hyundai as FIFA’s first official robotics partner — <https://www.newsweek.com/hyundais-robot-dogs-will-be-patrolling-the-fifa-world-cup-2026-12037758>
- [Deep Dive – 12 Jun 2026] **Wimbledon’s watsonx stack** — IBM’s AI commentary and open-source fan experiences at the All England Club — <https://www.redhat.com/en/blog/serving-faster-innovation-how-open-source-supporting-ai-powered-fan-experiences-wimbledon>
- [Deep Dive – 12 Jun 2026] **Athlete image rights in the AI era** — biometric digital twins and “synthetic rights” clauses in player contracts — <https://www.sportslawhub.com/athlete-image-rights-ai-deepfakes>

* * *

Using this index

This thematic index covers the *significant* sources across the *Dream Machine* archive, organised by topic. For specific research, follow the bracketed Issue numbers back to the canonical issue file in *Dream Machine MD/*. Each issue file ends with the section “All embedded URLs (in document order)” which lists every URL the issue carried, including local navigation links, profile pages and platform housekeeping links that are not reproduced here.

The newsletter is a continuous publication. The index above reflects the state of the archive as of Issue 41 (August 2026). Subsequent issues will extend the catalogue. The newsletter archive itself, on LinkedIn, remains the canonical primary source for every link the book builds on.

Citation Index

Dream Machine: Artificial Intelligence and the Next Creative Economy

All footnoted sources, organised by chapter. Every claim of substance in the manuscript is anchored to one of these references.

* * *

Foreword — Welcome to the Dream Machine

1. Variety, “SAG-AFTRA Condemns Tilly Norwood: AI Actress Is Not an Actor,” 30 September 2025. <https://variety.com/2025/film/news/sag-aftra-tilly-norwood-ai-actress-1236534779/>. See also NBC News, “Tilly Norwood, fully AI ‘actor,’ blasted by actors union SAG-AFTRA for ‘devaluing human artistry’.” <https://www.nbcnews.com/pop-culture/pop-culture-news/tilly-norwood-fully-ai-actor-blasted-actors-union-sag-aftra-devaluing-rcna234685>. Discussed in *Dream Machine* Issue 1 (6 October 2025).

2. *The Hollywood Reporter*, “U.K. Union Equity Condemns Tilly Norwood: ‘AI Tool, Not a Performer’.” <https://www.hollywoodreporter.com/movies/movie-news/tilly-norwood-ai-actress-uk-union-equity-sag-aftra-debate-1236391739/>. See also Variety, “Tilly Norwood Slammed by Equity as AI Tool, Concerned About Origin.” <https://variety.com/2025/film/global/tilly-norwood-slammed-equity-ai-tool-concerned-origin-1236537042/>. *Dream Machine* Issue 1.

3. CNN, “Tilly Norwood: Hollywood is fuming over a new ‘AI actress’,” 30 September 2025. <https://www.cnn.com/2025/09/30/tech/hollywood-ai-actor-backlash>.

4. OpenAI, “Sora 2 is here,” announcement page, 30 September 2025. <https://openai.com/index/sora-2/>. The model launched alongside an invite-only iOS app of the same name in the U.S. and Canada. *Dream Machine* Issue 1 carried the launch alongside contemporaneous coverage from NBC News and *The Guardian* on the model’s first copyright and safety incidents.

5. *Dream Machine* | Creative AI, LinkedIn newsletter, archive of Issues 1–41, October 2025 – August 2026. <https://www.linkedin.com/newsletters/dream-machine-creative-ai-7379776527871381505/>.

6. DreamLab AI Collective, team page. <https://dreamlab-ai.com/team>. Referenced from *Dream Machine* Issue 16 onward.

7. Charles Cecil (Revolution Software, *Broken Sword*) quoted in *gamesindustry.biz*, “‘AI was an expensive mistake’: Charles Cecil on innovation, insolvency, and Broken Sword.” <https://www.gamesindustry.biz/ai-was-an-expensive-mistake-charles-cecil-on-innovation-insolvency-and-broken-sword>. *Dream Machine* Issue 3.

8. Adobe, “Inaugural Adobe Creators’ Toolkit Report: 86 Percent of Global Creators Use Creative Generative AI.” <https://news.adobe.com/news/2025/10/adobe-max-2025-creators-survey>. Survey of 16,000 creators across the U.S., U.K., France, Germany, South Korea, Japan, India and Australia, released at Adobe MAX 2025. *Dream Machine* Issue 6.

9. UK Department for Science, Innovation and Technology (DSIT), *Statement of Progress on Copyright and AI*, December 2025. <https://www.gov.uk/government/publications/copyright-and-artificial-intelligence-progress-report/copyright-and-artificial-intelligence-statement-of-progress-under-section-137-data-use-and-access-act>. See also IPWatchdog, “Respondents to UK AI Consultation Overwhelmingly Want AI Companies to License Copyrighted Works in All Cases.” <https://ipwatchdog.com/2025/12/16/respondents-uk-ai-consultation-overwhelmingly-want-ai-companies-license-copyrighted-works-all-cases/>. *Dream Machine* Issue 12.

10. *Dream Machine* Issue 5, “Adobe’s Latest AI Announcements — Is every tool going AI?”, 31 October 2025. <https://www.linkedin.com/pulse/dream-machine-creative-ai-news-insight-oct-25-issue-5-woodbridge-f7jnc/>.

11. Adobe, Adobe MAX 2025 keynote messaging, October 2025. Coverage: Creative Boom, “Adobe is putting AI in everything everywhere all at once.” <https://www.creativeboom.com/news/adobe-is-putting-ai-in-everything-everywhere-all-at-once/>. *Dream Machine* Issue 5.

12. World Labs, *Marble* — first commercial spatial-AI world model, public launch November 2025. <https://marble.worldlabs.ai/>. Technical context: TechCrunch, “Fei-Fei Li’s World Labs speeds up the world model race with Marble, its first commercial product.” <https://techcrunch.com/2025/11/12/fei-fei-lis-world-labs-speeds-up-the-world-model-race-with-marble-its-first-commercial-product/>. DreamLab participated in the closed beta during October–November 2025. *Dream Machine* Issue 7.

13. 11,514 responses across the Citizen Space portal and email, of which 10,112 came through Citizen Space; 88% of those supported licensing as a default rule, against 3% who supported the government’s preferred opt-out model. UK DSIT, *Statement of Progress*, December 2025; analysis in *Dream Machine* Issue 12 (18 December 2025). Final report and economic impact assessment to be laid before Parliament by 18 March 2026.

14. *Digital Music News*, “Nearly 800 Creatives, Including Jason Aldean and One Republic, Sign Responsible AI Declaration — ‘Stealing Our Work Is Not Innovation’.” <https://www.digitalmusicnews.com/2026/01/22/stealing-isnt-innovation/>. *Dream Machine* Issue 16.

Chapter 1 — The Day Sora Landed

1. For a contemporaneous overview of the AI video model release cadence through 2024 and 2025, see *Dream Machine* Issues 1–8 (October–November 2025), which logged near-weekly releases from Runway, Luma, Pika, Kling, Veo, Wan, Higgsfield, Hunyuan and a long tail of smaller labs.
2. *The Hollywood Reporter*, “AI Performer Tilly Norwood Sparks Hollywood Backlash.” <https://www.hollywoodreporter.com/movies/movie-news/tilly-norwood-ai-actress-uk-union-equity-sag-aftra-debate-1236391739/>. *Dream Machine* Issue 1.
3. SAG-AFTRA statement, 30 September 2025, reported in *Variety*, “SAG-AFTRA Condemns Tilly Norwood: AI Actress Is Not an Actor.” <https://variety.com/2025/film/news/sag-aftra-tilly-norwood-ai-actress-1236534779/>.
4. OpenAI, “Sora 2 is here,” 30 September 2025. <https://openai.com/index/sora-2/>. *Dream Machine* Issue 1.
5. Particle6 background and Van der Velden interview: *The Hollywood Reporter*, “Meet the Creator of the AI Actress Hollywood Loves to Hate: ‘You’re Gonna See a Lot of Tilly Norwood Next Year’.” <https://www.hollywoodreporter.com/movies/movie-features/tilly-norwood-creator-particle6-eline-van-der-velden-talks-1236428824/>. *Dream Machine* Issue 8.
6. *Deadline*, “Tilly Norwood Creator Eline Van Der Velden Talks Backlash, Reveals Another 40 AI Actors Are In The Pipeline.” <https://deadline.com/2025/11/tilly-norwood-creator-interview-backlash-more-ai-actors-coming-1236601334/>.
7. Northeastern Global News, “Why AI ‘Actress’ Tilly Norwood Has Hollywood Angry.” <https://news.northeastern.edu/2025/10/02/ai-actress-tilly-norwood-hollywood-backlash/>.
8. SAG-AFTRA, official statement reproduced in *Variety*, *op. cit.*; also NBC News, “Tilly Norwood, fully AI ‘actor,’ blasted by actors union SAG-AFTRA for ‘devaluing human artistry’.” <https://www.nbcnews.com/pop-culture/pop-culture-news/tilly-norwood-fully-ai-actor-blasted-actors-union-sag-aftra-devaluing-rcna234685>.
9. Equity (U.K.), statement of 2 October 2025: *Variety*, “Tilly Norwood Slammed by Equity as AI Tool, Concerned About Origin.” <https://variety.com/2025/film/global/tilly-norwood-slammed-equity-ai-tool-concerned-origin-1236537042/>.
10. CNN, “Tilly Norwood: Hollywood is fuming over a new ‘AI actress’.” <https://www.cnn.com/2025/09/30/tech/hollywood-ai-actor-backlash>.

11. OpenAI, “Sora 2 is here,” <https://openai.com/index/sora-2/>. Technical capabilities summary including physics modelling, multi-shot world-state persistence and synchronised audio.

12. *Dream Machine* Issue 1, “Editor’s Pick”; further launch context in NBC News, “OpenAI’s Sora 2: a major leap in AI video and audio.” <https://www.nbcnews.com/tech/tech-news/openai-sora-2-app-video-chatgpt-creation-rcna234973>.

13. LinkedIn News aggregation: “Sora Tops 1 Million Downloads in 5 Days.” <https://www.linkedin.com/news/story/sora-tops-1m-downloads-in-5-days-6684988/>. *Dream Machine* Issue 3. A historical note for the reader: the Sora consumer app and the broader Sora platform did not survive 2026. OpenAI announced a staged wind-down of Sora on **24 March 2026** — consumer app and website ending late April, developer API closing late September — with peak revenue reportedly at around \$540,000 per month against operating costs in the billions, and the compute reallocated toward higher-margin coding and enterprise products. The Disney–OpenAI licensing announcement of 11 December 2025 (referenced in Chapter 7 and Appendix C) was never executed; no formal agreement was signed and no money changed hands. The contractual framework underneath that deal — IP-not-training, talent-likeness excluded, jointly-governed steering committee, equity-and-warrants structure — survived the platform’s collapse and is now the reference architecture other studios negotiate against. The argument the rest of this chapter is making about the Sora 2 launch week — the cultural rupture, the audience arrival in the AI debate, the parasocial-character / working-actor market collision — survives the platform’s death intact; the *moment* matters whether or not the *vendor* persists.

14. *The Guardian*, “OpenAI Sora 2 violence racism.” <https://www.theguardian.com/us-news/2025/oct/04/openai-sora-violence-racism>. *Dream Machine* Issue 1.

15. NBC News, *op. cit.*; *The Guardian*, *op. cit.*

15a. Quoted in *The Guardian*, “OpenAI launch of video app Sora plagued by violent and racist images: ‘The guardrails are not real’.” <https://www.theguardian.com/us-news/2025/oct/04/openai-sora-violence-racism>. *Dream Machine* Issue 1.

16. *Digital Music News*, “OpenAI’s Sora 2 includes likeness protections for celebrities who don’t opt in, but that doesn’t apply to ‘historical figures’ and dead celebrities.” <https://www.digitalmusicnews.com/2025/10/08/openais-likeness-protections-dont-apply-to-dead-celebrities/>. *Dream Machine* Issue 2.

17. Google DeepMind, Veo 3.1 launch, mid-October 2025. *Dream Machine* Issue 3, “Editor’s Pick: Veo 3.1 and the Rise of AI Filmmaking.” Coverage: <https://www.cometapi.com/veo-3-1-is-comingand-whats-rumor/>.

18. WUFT, “Kiss reality goodbye: AI-generated social media has arrived,” 3 October 2025. <https://www.wuft.org/2025-10-03/kiss-reality-goodbye-ai-generated-social-media-has-arrived>. *Dream Machine* Issue 1.

19. *No Film School*, “James Cameron Says AI Is ‘Never Going to Take the Place’ of Humans.” <https://nofilmschool.com/james-cameron-ai#>. *Dream Machine* Issue 1.

20. *The Guardian*, “James Cameron says AI actors are ‘horrifying to me’,” 1 December 2025. <https://www.theguardian.com/film/2025/dec/01/james-cameron-says-ai-actors-are-horrifying-to-me>. Original quote from CBS *Sunday Morning*. *Dream Machine* Issue 10.

20a. *Variety*, “James Cameron Says It’s ‘Horrifying’ that AI Can ‘Make Up an Actor’.” <https://variety.com/2025/film/news/james-cameron-horrifying-ai-replace-actors-1236595864/>.

21. *Stability AI*, board composition, 2024–2026. Reported across multiple outlets including *Deadline*, “James Cameron Calls AI Replacing Actors ‘Horrifying’; Art ‘Sacred’.” <https://deadline.com/2025/11/james-cameron-gen-ai-horrifying-human-art-sacred-avatar-1236631387/>.

22. *Deezer*, “AI-generated tracks now represent 44% of all new uploaded music,” April 2026 newsroom release. <https://newsroom-deezer.com/2026/04/ai-generated-tracks-represent-44-of-new-uploaded-music/>. Companion analysis: *Music Business Worldwide*, “75,000 AI-generated tracks now flood Deezer daily, representing 44% of all new music uploaded to the platform.” <https://www.musicbusinessworldwide.com/75000-ai-generated-tracks-now-flood-deezer-daily-representing-44-of-all-new-music-uploaded-to-the-platform-says-streamer/>. Daily AI uploads to Deezer rose from approximately 50,000 per day in November 2025 (*Dream Machine* Issue 7, citing Deezer / *Musically*) to 75,000 per day by April 2026, with consumer streams of fully-AI tracks holding between 1% and 3% of total platform plays — and up to 85% of those streams identified as fraudulent in 2025. *Dream Machine* Issues 7, 26, 27, 28.

23. *Deadline*, “AI ‘Actor’ Tilly Norwood To Star In Feature Film ‘Misaligned’.” <https://deadline.com/2026/07/tilly-norwood-ai-actor-misaligned-1236974639/>. *Dream Machine* Issue 37.

Chapter 2 — A History of Resistance

1. John Philip Sousa, “*The Menace of Mechanical Music*,” *Appleton’s Magazine*, Vol. 8, September 1906, pp. 278–284. Full text via ExplorePAHistory: <https://explorepahistory.com/odocument.php?docId=1-4-1A1>. Academic context: Patrick Warfield, “*John Philip Sousa and The Menace of Mechanical Music*,” *Journal of the Society for American Music*, Cambridge University Press:

<https://www.cambridge.org/core/journals/journal-of-the-society-for-american-music/article/abs/john-philip-sousa-and-the-menace-of-mechanical-music/A9E621587BE7580ABD73AEF64D4B2DC8>. The 1906 essay was, in part, lobbying for what would become the 1909 Copyright Act.

2. Sousa, *op. cit.* The Library of Congress's "*Sousa and the Talking Machine*" essay is a useful institutional summary: <https://blogs.loc.gov/now-see-hear/2020/05/sousa-and-the-talking-machine/>.

3. William Henry Cardinal O'Connell, Archbishop of Boston, sermon to the Holy Name Society, Boston, 10 January 1932. Reported widely in the contemporaneous press, including the *Daily Courier* (Connellsville, PA), 12 January 1932 (<https://www.newspapers.com/newspage/38168082/>). Cultural context: KUOW/NPR, "*Imbecile Slush: Surprising Early Reactions to Crooning*," <https://www.kuow.org/stories/imbecile-slush-surprising-early-reactions-crooning>. JSTOR Daily, "*The Gender Politics of the First Boy Bands*," <https://daily.jstor.org/the-gender-politics-of-the-first-boy-bands/>.

4. *Grand Upright Music, Ltd. v. Warner Bros. Records Inc.*, 780 F. Supp. 182 (S.D.N.Y. 1991). Full text: <https://law.justia.com/cases/federal/district-courts/FSupp/780/182/1445286/>. The "Thou shalt not steal" opening is the most-quoted line from a US copyright opinion of the late twentieth century.

5. Tippet's account of the Jurassic Park digital test is documented across multiple ASC and contemporaneous press accounts. American Society of Cinematographers, "*Jurassic Park: Effects Team Brings Dinosaurs Back from Extinction*," <https://theasc.com/articles/jurassic-park-effects-team-brings-dinosaurs-back>. Wikipedia, "*Phil Tippett*," https://en.wikipedia.org/wiki/Phil_Tippett. The dialogue paraphrase Spielberg incorporated into the film is Goldblum/Malcolm's response to Grant's "I think we're out of a job": "Don't you mean *extinct*?"

5a. *Variety*, "George Lucas Says AI Makes Filmmaking 'Easier' and 'There's Nothing You Can Do About It'." <https://variety.com/2026/film/news/george-lucas-test-screenings-embraces-ai-hollywood-1236810656/>; AOL, "George Lucas likens AI sceptics to luddites clinging to horses and carts." <https://www.aol.co.uk/articles/george-lucas-likens-ai-sceptics-122913000.html>. *Dream Machine* Issue 38.

6. Charles Baudelaire, "*Le Public Moderne et la Photographie*," *Revue Française*, 1859 (part of the *Salon de 1859* essays). English translation widely available; the original French in PDF form: <https://gallowayexploringart.wordpress.com/wp-content/uploads/2014/08/ baudelaire-the-modern-public-photography.pdf>. Smithsonian Archives institutional overview: "*Photography Murdered Painting, Right?*," <https://siarchives.si.edu/blog/photography-murdered-painting-right>.

7. The Delaroche apocrypha is documented in Quote Investigator: <https://quoteinvestigator.com/2022/10/16/photo-mortal/>. The earliest sourced version is in an 1873 survey, 34 years after Delaroche reportedly said it. Delaroche's own contemporary writing on the daguerreotype, in Gernsheim's standard 1959 monograph, characterised the new technology as "*an immense service to the arts.*"

8. The 1942–44 Petrillo strike: Wikipedia, "*1942–44 musicians' strike*," https://en.wikipedia.org/wiki/1942%E2%80%9344_musicians'_strike; Mainspring Press, "*The Man Who Crippled the American Recording Industry: James Caesar Petrillo and the American Federation of Musicians Recording Bans*," <https://mainspringpress.org/2024/11/23/the-man-who-crippled-the-recording-industry-james-caesar-petrillo-and-the-american-federation-of-musicians-recording-bans/>; DownBeat, "*The Petrillo Ban of 1942–'44: Past & Future at War*," <https://downbeat.com/news/detail/the-petrillo-ban-of-194244-past-future-at-war>; Local 802 AFM, "*The Silence Was Deafening*," <https://www.local802afm.org/allegro/articles/the-silence-was-deafening/>. The Music Performance Trust Fund's institutional history: <https://musicpf.org/establishment-of-mptf-led-to-the-formation-of-afms-pension-and-residual-funds/>.

9. Musicians' Union History, "*The Strike That Made History — Massacre of the Musicians 1980*," <https://www.muhistory.com/the-massacre-of-the-musicians-1980/>. Academic context on the broader MU–BBC dispute landscape: "*Negotiating Needletime*" (Tandfonline), <https://www.tandfonline.com/doi/full/10.1080/03071022.2016.1215098>.

10. MusicRadar, "*The Day the Loony Musicians Union Tried to Kill the Synthesizer (Which Also Happened to be Bob Moog's Birthday)*," <https://www.musicradar.com/news/the-union-passed-a-motion-to-ban-the-use-of-synths-drum-machines-and-any-electronic-devices-the-day-the-loony-musicians-union-tried-to-kill-the-synthesizer-which-also-happened-to-be-bob-moogs-birthday>. Far Out Magazine, "*Why did the Musicians Union outlaw synthesisers in 1982?*" <https://faroutmagazine.co.uk/musicians-union-outlaw-synthesisers/>.

11. *Bridgeport Music, Inc. v. Dimension Films*, 410 F.3d 792 (6th Cir. 2005). Full text: <https://law.justia.com/cases/federal/appellate-courts/F3/410/792/574458/>. The "Get a licence or do not sample" rule is the most-cited line in the opinion.

12. *TIME*, "*50 Worst Inventions*," 2010, Auto-Tune at #15: https://content.time.com/time/specials/packages/article/0,28804,1991915_1991909_1991903,00.html. Wikipedia, "*Auto-Tune*," <https://en.wikipedia.org/wiki/Auto-Tune>. NPR, "*25 Years of Believe*," <https://www.npr.org/2023/10/19/1207028349/25-years-ago-cher-released-a-song>

that-would-change-the-sound-of-pop-music. Wikipedia, “D.O.A. (*Death of Auto-Tune*),” [https://en.wikipedia.org/wiki/D.O.A._\(Death_of_Auto-Tune\)](https://en.wikipedia.org/wiki/D.O.A._(Death_of_Auto-Tune)).

13. Walter Murch, *In the Blink of an Eye: A Perspective on Film Editing*, Silman-James Press, 1995 (2nd edition 2001). PDF: [https://www.craftfilmschool.com/userfiles/files/Walter%20Murch%20-%20In%20the%20Blink%20of%20an%20Eye%20Revised%202nd%20Edition%20\(2001,%20Silman-James%20Pr\).pdf](https://www.craftfilmschool.com/userfiles/files/Walter%20Murch%20-%20In%20the%20Blink%20of%20an%20Eye%20Revised%202nd%20Edition%20(2001,%20Silman-James%20Pr).pdf). Charles Koppelman, *Behind the Seen: How Walter Murch Edited Cold Mountain Using Apple’s Final Cut Pro and What This Means for Cinema*, Peachpit Press, 2004: <https://www.peachpit.com/store/behind-the-seen-how-walter-murch-edited-cold-mountain-9780735714267>.

14. Sasson’s account documented at the National Inventors Hall of Fame: <https://www.invent.org/blog/inventors/Legacy-Steve-Sasson>. Snopes verification of the “Kodak suppressed the digital camera” claim: <https://www.snopes.com/fact-check/kodak-digital-camera-invention/>. Knowledge@Wharton on the Kodak collapse: <https://knowledge.wharton.upenn.edu/podcast/knowledge-at-wharton-podcast/whats-wrong-with-this-picture-kodaks-30-year-slide-into-bankruptcy/>. Bankruptcy filing: 19 January 2012, S.D.N.Y., \$5.1bn assets / \$6.8bn liabilities.

15. Wikipedia, “Brian Walski,” https://en.wikipedia.org/wiki/Brian_Walski. *Washington Post* contemporaneous coverage: <https://www.washingtonpost.com/archive/lifestyle/2003/04/03/altered-picture-costs-la-times-photographer-his-job/c5e7c9e0-a836-429a-bb4e-d502f1768a96/>. *World Press Photo’s* institutional response in *TIME*: <https://time.com/3706626/world-press-photo-processing-manipulation-disqualified/>.

16. Wikipedia, “*Viacom International, Inc. v. YouTube, Inc.*,” https://en.wikipedia.org/wiki/Viacom_International_Inc._v._YouTube,_Inc.. Electronic Frontier Foundation case file: <https://www.eff.org/cases/viacom-v-youtube>. *Variety* on the March 2014 settlement: <https://variety.com/2014/biz/news/google-and-viacom-settle-copyright-infringement-lawsuit-over-youtube-1201137538/>.

17. PetaPixel, “*The Rise and Crash of the Camera Industry in One Chart*,” <https://petapixel.com/2024/08/22/the-rise-and-crash-of-the-camera-industry-in-one-chart/>. Statista, “*Smartphones Wipe Out Decades of Camera Industry Growth*,” <https://www.statista.com/chart/15524/worldwide-camera-shipments/>. CIPA shipment data series, multiple years.

18. CNN Business, “*Meet the translation professionals losing their jobs to AI*,” January 2026, <https://www.cnn.com/2026/01/23/tech/translation-language-jobs-ai-automation-intl>. Carl Benedikt Frey (Oxford Martin School), 2025 study on translator employment across 696 US labour markets. American Translators

Association industry position: <https://www.atanet.org/client-assistance/blog-machine-translation-vs-human-translation/>. Wikipedia, “Google Neural Machine Translation,” https://en.wikipedia.org/wiki/Google_Neural_Machine_Translation.

Chapter 3 — The Human-AI Agency Continuum

1. *Dream Machine* Issue 2, “Editor’s Pick,” 10 October 2025. <https://www.linkedin.com/pulse/dream-machine-creative-ai-news-insight-oct-25-2-pete-woodbridge-mnrjc/>.

2. OpenAI, “Introducing AgentKit,” 6 October 2025. <https://openai.com/index/introducing-agentkit/>.

3. TechCrunch, “OpenAI launches AgentKit to help developers build and ship AI agents,” 6 October 2025. <https://techcrunch.com/2025/10/06/openai-launches-agentkit-to-help-developers-build-and-ship-ai-agents/>. Also coverage at *InfoQ*, “OpenAI Dev Day 2025 Introduces GPT-5 Pro API, Agent Kit, and More.” <https://www.infoq.com/news/2025/10/openai-dev-day/>.

4. *Dream Machine* Issue 2: “Agentic AI — the class of AI systems that can plan, act, and pursue goals with autonomy — promises a new era of collaboration in creative industries... Its another step along the Human-AI Agency Continuum.” See also *TVB Europe*, “Is Agentic AI About to Change the Media and Entertainment Industry?” <https://www.tvbeurope.com/artificial-intelligence/opinion-is-agentic-ai-about-to-change-the-media-and-entertainment-industry>.

5. Google DeepMind, Veo 3.1 release, October 2025. *Dream Machine* Issue 3.

6. *MusicTech*, “iZotope Ozone 12’s AI assistant is cool, but the Stem EQ is the real star.” <https://musictech.com/reviews/plugin-ins/izotope-ozone-12-review/>. *Dream Machine* Issue 3.

7. Adobe, “Inaugural Adobe Creators’ Toolkit Report,” October 2025. <https://news.adobe.com/news/2025/10/adobe-max-2025-creators-survey>. Survey of 16,000 creators across eight countries, released at Adobe MAX 2025. *Dream Machine* Issue 6.

8. Adobe, *op. cit.* The same survey: 86% of creators use creative generative AI; 76% say it has helped grow their business or brand; 81% say AI lets them make content they otherwise couldn’t have made; 69% worry about their work being used to train AI without consent; 70% are optimistic about agentic AI; 85% would use AI that learns their creative style.

9. Mureka, “Music Agent Studio” launch, mid-October 2025. *Dream Machine* Issue 4. https://www.linkedin.com/posts/sherrihendrickson_mureka-unveils-music-agent-studio-and-enhanced-share-7384999251526864896-cNYg/.

10. *Finsmes*, “AdsAgency Raises \$12M in Seed Funding,” October 2025. <https://www.finsmes.com/2025/10/adsgency-raises-12m-in-seed-funding.html>. *Dream Machine* Issue 4.
11. *Musically*, “Meet Lenny, an AI agent to help organisers of live music events.” <https://musically.com/2025/10/20/meet-lenny-an-ai-agent-to-help-organisers-of-live-music-events/>. *Dream Machine* Issue 4.
12. *GamesRadar*, “Even under USD20 million in debt, EA reportedly pushes 15,000 employees to use AI as a ‘thought partner’ for everything from character art to playtesting.” <https://www.gamesradar.com/games/even-under-usd20-million-in-debt-ea-reportedly-pushes-15-000-employees-to-use-ai-as-a-thought-partner-for-everything-from-character-art-to-playtesting/>. *Dream Machine* Issue 6.
13. PYMNTS, “Adobe Lets Users Design and Edit Using ChatGPT.” <https://www.pymnts.com/artificial-intelligence-2/2025/adobe-lets-users-design-and-edit-using-chatgpt/>. Adobe blog: “Edit images, designs, and PDFs right inside ChatGPT — thanks to Adobe Express, Photoshop, and Acrobat.” <https://blog.adobe.com/en/publish/2025/12/10/edit-photoshop-chatgpt>. *Dream Machine* Issue 12.
14. TechCrunch, “Anthropic launches interactive Claude apps, including Slack and other workplace tools,” 26 January 2026. <https://techcrunch.com/2026/01/26/anthropic-launches-interactive-claude-apps-including-slack-and-other-workplace-tools/>. *Heygen Video Agent*: https://www.linkedin.com/posts/heygen_introducing-the-new-video-agent-activity-7421597801240801282-d1CF. *Dream Machine* Issue 16.
15. *Dream Machine* Issue 21, “Editor’s Pick: Adobe and NVIDIA Just Raised the Stakes for Creative AI,” 19 March 2026.
16. Adobe Summit 2026, “Agentic Creative Intelligence” keynote framing. *Dream Machine* Issue 26.
17. *Dream Machine* Issue 29, May 2026, citing Sony’s adoption of Claude Code studios with multi-agent coordination.
18. Anthropic, public statements on agent deployment patterns through Q1 2026. Cf. *Dream Machine* Issues 11, 16, 22.
19. *gamesindustry.biz*, “AI was an expensive mistake’: Charles Cecil on innovation, insolvency, and Broken Sword.” <https://www.gamesindustry.biz/ai-was-an-expensive-mistake-charles-cecil-on-innovation-insolvency-and-broken-sword>. *Dream Machine* Issue 3.
20. *Niche Gamer*, “Larian Studios backs off from gen AI, says tech won’t be used in new Divinity.” <https://nichegamer.com/larian-studios-backs-off-from-gen-ai/>.

Dream Machine Issue 14.

21. *Decrypt*, “Warhammer 40,000’ Maker Games Workshop Rules Out Generative AI.” <https://decrypt.co/354482/warhammer-40000-maker-games-workshop-rules-out-generative-ai>. *Dream Machine* Issue 14.

22. *Niche Gamer*, “Manor Lords publisher Hooded Horse won’t work with devs using gen AI.” <https://nichegamer.com/manor-lords-publisher-hooded-horse-wont-work-with-devs-using-gen-ai/>. *Dream Machine* Issue 14.

23. *gamesindustry.biz*, “RuneScape maker Jagex says it will never use generative AI to make in-game content.” <https://www.gamesindustry.biz/runescape-maker-jagex-says-it-will-never-use-generative-ai-to-make-in-game-content>. *Dream Machine* Issue 16.

Chapter 4 — Dead Internet, Living Web

1. *Imperva*, *2025 Bad Bot Report: How AI is Supercharging the Bot Threat*. <https://www.imperva.com/blog/2025-imperva-bad-bot-report-how-ai-is-supercharging-the-bot-threat/>. *Dream Machine* Issue 4.

2. *Cloudflare*, “The crawl-to-click gap: Cloudflare data on AI bots, training, and referrals.” <https://blog.cloudflare.com/crawlers-click-ai-bots-training/>. *Dream Machine* Issue 4. Later 2025 updates show training crawlers declining from ~90% to ~74% of AI bot activity as scraper bots rose to 24% and a new “agentic” category emerged at 1.7%; see *Cloudflare*, “A deeper look at AI crawlers: breaking down traffic by purpose and industry.” <https://blog.cloudflare.com/ai-crawler-traffic-by-purpose-and-industry/>.

3. *Grand View Research*, “Generative AI Content Creation Market Report.” <https://www.grandviewresearch.com/industry-analysis/generative-ai-content-creation-market-report>. *Dream Machine* Issue 4 also cites *Gartner* and *Europol* forecasts of 90–99% AI-generated or AI-assisted online content by 2030.

4. *Dream Machine* Issue 4, “Editor’s Pick: Is the Internet Dead Yet?” 23 October 2025. <https://www.linkedin.com/pulse/dream-machine-creative-ai-news-insight-oct-25-issue-4-woodbridge-hzttc/>.

5. Wikipedia, “Dead Internet Theory.” https://en.wikipedia.org/wiki/Dead_Internet_theory. *Dream Machine* Issue 4.

6. *Graphite*, 2025 analysis of new web content by author type (human vs. AI vs. AI-assisted). Cited in *Dream Machine* Issue 4.

7. For “model collapse” as a term of art, see *Ilia Shumailov et al.*, “The Curse of Recursion: Training on Generated Data Makes Models Forget” (2024), and subsequent literature.

8. Futurism, “Researchers built a social network with only AI agents — within hours it had collapsed into warring tribes.” <https://futurism.com/social-network-ai-intervention-echo-chamber>. *Dream Machine* Issue 4.

9. *Digital Music News*, “Instagram Chief Says We Should ‘Fingerprint Real Media’ Instead of Tracking and Disclosing AI Slop.” <https://www.digitalmusicnews.com/2026/01/05/instagram-chief-ai-slop-comments/>. See also *WebProNews*, “Instagram Head Warns AI Images Erode Trust, Calls for Verification Standards.” <https://www.webpronews.com/instagram-head-warns-ai-images-erode-trust-calls-for-verification-standards/>. *Dream Machine* Issue 13.

9a. *The Guardian*, “Brands using AI-generated influencers to promote products on social media.” <https://www.theguardian.com/technology/2026/jun/21/brands-using-ai-generated-influencers-to-promote-products-on-social-media>. *Dream Machine* Issue 35.

9b. *Social Media Today*, “TikTok bans AI-generated voices in shopping livestreams.” <https://www.socialmediatoday.com/news/tiktok-bans-ai-generated-voices-in-shopping-livestreams/822977/>. *Dream Machine* Issue 35.

9c. *Search Engine Journal*, “TikTok Targets AI-Generated Spam Accounts In High-Risk Topics.” <https://www.searchenginejournal.com/tiktok-targets-ai-generated-spam-accounts-in-high-risk-topics/582256/>. *Dream Machine* Issue 38.

9d. *Digit*, “YouTube, X funnel millions of visits to AI deepfake ‘nudify’ apps, new study shows: Key findings.” <https://www.digit.in/news/general/millions-of-users-reached-ai-nudify-sites-through-youtube-and-x-report.html>. *Dream Machine* Issue 38.

10. Sundance Institute, “Centering the Artist: Why We’re Launching the AI Literacy Initiative.” <https://www.sundance.org/blogs/centering-the-artist-why-were-launching-the-ai-literacy-initiative/>. *Dream Machine* Issue 16.

11. *Stereogum*, “Bandcamp bans AI music.” <https://stereogum.com/2485199/bandcamp-bans-ai-music/news>. *Dream Machine* Issue 14.

12. *CNET*, “San Diego Comic-Con Draws a Line: No AI Art Allowed at 2026 Event.” <https://www.cnet.com/culture/san-diego-comic-con-bans-ai-art-for-2026-event/>. *Dream Machine* Issue 16.

13. *Deezer*, “AI-generated tracks now represent 44% of all new uploaded music,” April 2026. <https://newsroom-deezer.com/2026/04/ai-generated-tracks-represent-44-of-new-uploaded-music/>. *Music Business Worldwide*, “75,000 AI-generated tracks now flood Deezer daily.” <https://www.musicbusinessworldwide.com/75000->

ai-generated-tracks-now-flood-deezer-daily-representing-44-of-all-new-music-uploaded-to-the-platform-says-streamer/. *Dream Machine* Issues 7, 26, 27, 28.

14. *The Hollywood Reporter*, “‘Synthetic Sincerity’ by Marc Isaacs Explores if AI Characters Can Be Taught Authenticity: IDFA.” <https://www.hollywoodreporter.com/movies/movie-news/synthetic-sincerity-film-idfa-ai-authenticity-interview-1236426180/>. *Dream Machine* Issue 8.

15. *Variety*, “AI-Generated Images Threaten Future of Documentary as People ‘Will Stop Believing Anything’.” <https://variety.com/2025/film/festivals/ai-generated-images-threaten-future-of-documentary-1236583466/>. *Dream Machine* Issue 8.

16. *PR Newswire*, “From Apple TV Creative to AI Filmmaker: Hoyt Dwyer’s Animated Film To Compete at AI FilmFest Japan 2025.” <https://www.prnewswire.com/news-releases/from-apple-tv-creative-to-ai-filmmaker-hoyt-dwyers-animated-film-to-compete-at-ai-filmfest-japan-2025-302598064.html>. *Dream Machine* Issue 6.

17. *Variety*, “AI Creator Behind Viral ‘Deadpool,’ ‘Harry Potter’ Christmas Clip Made His Film in a Ukrainian Bomb Shelter.” <https://variety.com/2026/digital/news/ai-video-deadpool-harry-potter-andrii-daniels-1236624632/>. *Dream Machine* Issue 16.

18. *Branding in Asia*, “‘It’s the Most Terrible Time of the Year’ — McDonald’s Netherlands’ Wonderfully Chaotic, AI-Driven Christmas Film.” <https://www.brandinginasia.com/its-the-most-terrible-time-of-the-year-mcdonalds-netherlands-wonderfully-chaotic-ai-driven-christmas-film/>. Pulled following backlash: *SiliconAngle*, “Not ready: McDonald’s AI-generated ad taken down after public backlash.” <https://siliconangle.com/2025/12/10/not-ready-mcdonalds-ai-generated-ad-taken-public-backlash/>. *Dream Machine* Issue 11.

19. *BBC News*, “Fashion house Valentino criticised over ‘disturbing’ AI handbag ads.” <https://www.bbc.co.uk/news/articles/cwyvjyvn83go>. *Dream Machine* Issue 10.

20. *Adweek*, “Coca-Cola Uses AI to Rekindle the Magic of Its Holiday Ads.” <https://www.adweek.com/creativity/coca-cola-uses-ai-to-rekindle-the-magic-of-its-holiday-ads/>. *Dream Machine* Issue 6.

21. *AI News*, “AI causes reduction in users’ brain activity, MIT.” <https://www.artificialintelligence-news.com/news/ai-causes-reduction-in-users-brain-activity-mit/>. *Dream Machine* Issue 1.

22. *ElevenLabs*, “Detecting audio generated by ElevenLabs with SynthID” (partnership with Google DeepMind). <https://elevenlabs.io/blog/synthid>. *Dream Machine* Issue 36.

22a. *Digital Music News*, “TikTok Quietly Starts Testing AI Deepfake Detection Tech.” <https://www.digitalmusicnews.com/2026/07/20/tiktok-starts-testing-ai-deepfake-detection-tech/>. *Dream Machine* Issue 39.

23. *Business Today*, “We can fix the internet ...’ says World Wide Web inventor Tim Berners-Lee.” <https://www.businesstoday.in/technology/news/story/we-can-fix-the-internet-says-world-wide-web-inventor-tim-berners-lee-538947-2026-06-24>. *Dream Machine* Issue 36.

24. *Social Media Today*, “YouTube adds AI-powered comment filtering.” <https://www.socialmediatoday.com/news/youtube-adds-ai-powered-comment-filtering/823939/>. *Dream Machine* Issue 36.

Chapter 5 – The Slop Ceiling

1. Deezer, “AI-generated tracks now represent 44% of all new uploaded music,” April 2026. <https://newsroom-deezer.com/2026/04/ai-generated-tracks-represent-44-of-new-uploaded-music/>. *Music Business Worldwide*, “75,000 AI-generated tracks now flood Deezer daily, representing 44% of all new music uploaded to the platform.” <https://www.musicbusinessworldwide.com/75000-ai-generated-tracks-now-flood-deezer-daily-representing-44-of-all-new-music-uploaded-to-the-platform-says-streamer/>. *Dream Machine* Issues 7, 26, 27, 28.

2. Ditto Music research, October 2025 and prior. *Press Ditto Music*, “48% of artists use AI to make music — fewer than in 2023.” <https://press.dittomusic.com/48-of-artists-use-ai-to-make-music-fewer-than-in-2023>. *Dream Machine* Issue 2.

3. *Musically*, “Universal and Warner could sign landmark AI deals within weeks.” <https://musically.com/2025/10/02/report-umg-and-wmg-could-sign-landmark-ai-deals-within-weeks/>. Spotify Newsroom, “Spotify Strengthens AI Protections for Artists, Songwriters, and Producers.” <https://newsroom.spotify.com/2025-09-25/spotify-strengthens-ai-protections/>. *Dream Machine* Issue 1.

4. *Musically*, “50,000 AI music tracks are now uploaded to Deezer every day.” <https://musically.com/2025/11/12/50000-ai-music-tracks-are-now-uploaded-to-deezer-every-day/>. *Dream Machine* Issue 7.

5. Deezer, April 2026, *op. cit.*

6. *Musically*, “UMG boss slams exponential growth of AI slop on streaming services.” <https://musically.com/2026/01/09/umg-boss-slams-exponential-growth-of-ai-slop-on-streaming-services/>. *Dream Machine* Issue 14.

7. *Musically*, “Report: 56.9% of new independent songs in China are AI-generated.” <https://musically.com/2026/01/05/report-56-9-of-new-independent-songs-in-china-are-ai-generated/>. *Dream Machine* Issue 13.

8. *The Wrap*, “An AI Podcasting Machine Is Churning Out 3,000 Episodes a Week — and People Are Listening.” <https://www.thewrap.com/ai-podcasts-hosts-inception-point-ai/>. *Dream Machine* Issue 8.
9. *Dream Machine* Issue 28, May 2026, citing aggregator-platform data on “podslop” classification.
10. *The Hollywood Reporter*, “Merriam-Webster Names ‘Slop’ Word of the Year Amid AI Boom.” <https://www.hollywoodreporter.com/news/general-news/slop-word-year-2025-merriam-webster-1236450780/>. *Dream Machine* Issue 12.
11. *Digital Music News*, “YouTube CEO Puts ‘Managing AI Slop’ on the Priority List for 2026.” <https://www.digitalmusicnews.com/2026/01/22/youtube-ceo-ai-slop-2026-comments/>. *Dream Machine* Issue 16.
12. *The Guardian*, “YouTube AI channels spreading fake, anti-Labour videos viewed 1.2bn times in 2025.” <https://www.theguardian.com/technology/2025/dec/13/fake-anti-labour-video-billion-views-youtube-2025>. *Dream Machine* Issue 12.
13. Deezer/Ipsos survey, November 2025. <https://newsroom-deezer.com/2025/11/deezer-ipsos-survey-ai-music/>. *Dream Machine* Issue 7.
14. *Bain & Company*, “In an AI Age, People Still Want the Radio Star.” <https://www.bain.com/insights/in-an-ai-age-people-still-want-the-radio-star/>. *Dream Machine* Issue 16.
15. Deezer, April 2026, *op. cit.* “Up to 85% of the streams generated by fully AI-generated tracks were in fact fraudulent in 2025.”
- 15a. *Bloomberg*, “AI Changed Chess. Grandmasters Now Win With Unpredictable Moves,” 27 March 2026. <https://www.bloomberg.com/news/articles/2026-03-27/ai-changed-chess-grandmasters-now-win-with-unpredictable-moves>. *Dream Machine* Issue 23.
16. *Billboard*, “AI Artist Xania Monet Climbs the Charts — And Signs a Multimillion-Dollar Record Deal.” <https://www.billboard.com/pro/ai-music-artist-xania-monet-multimillion-dollar-record-deal/>.
17. *Billboard*, *op. cit.*; CNN, “Xania Monet is the first AI-powered artist to debut on a Billboard airplay chart.” <https://www.cnn.com/2025/11/01/entertainment/xania-monet-billboard-ai>.
18. *Billboard*, *op. cit.*
19. *Bangkok Post*, “AI singer Xania Monet signs \$3m deal with record label.” <https://www.bangkokpost.com/life/tech/3142355/ai-singer-xania-monet-signs-3m-deal-with-hollywood-media>. *Dream Machine* Issue 7.
20. Multiple outlets; quoted in *Billboard* feature *op. cit.*

20a. Telisha Jones quoted in *Billboard*, *op. cit.*

21. NPR, “Breaking Rust is a hot new country act on the Billboard charts. It’s powered by AI.” <https://www.npr.org/2025/11/10/nx-s1-5604320/breaking-rust-is-a-hot-new-country-act-on-the-billboard-charts-its-powered-by-ai>. *Dream Machine* Issue 7.

22. *Washington Post*, “‘Walk My Walk,’ Breaking Rust: AI country hit triggers Nashville angst.” <https://www.washingtonpost.com/style/2025/12/28/breaking-rust-ai-country/>.

23. *MusicRadar*, “The No. 1 country song in the US right now is AI-generated.” <https://www.musicradar.com/music-tech/the-no-1-country-song-in-the-us-right-now-is-ai-generated>. *Dream Machine* Issue 7.

24. BBC News, “The mysterious singer, Sienna Rose, with millions of streams is hitting the viral charts — but who (or what) is she?” <https://www.bbc.co.uk/news/articles/cq6v83gq66eo>. *Dream Machine* Issue 15.

25. *Billboard*, “How a MAGA Rapper Used AI to Create A Gospel Song That Climbed the Charts.” <https://www.billboard.com/pro/maga-rapper-ai-gospel-song-climbed-charts/>. *Dream Machine* Issue 9.

26. *Musically*, “AI band Bleeding Verse’s creator signs deal with Hallwood Media.” <https://musically.com/2025/10/07/ai-band-bleeding-verses-creator-signs-deal-with-hallwood-media/>. *Dream Machine* Issue 2.

27. *Musically*, “Indian AI band Trilok performs live, government denies association.” <https://musically.com/2025/12/17/indian-ai-band-trilok-performs-live-government-denies-association/>. *Dream Machine* Issue 12.

28. *The Guardian*, “Paul McCartney joins music industry protest against AI with silent track.” <https://www.theguardian.com/music/2025/nov/17/the-sound-of-silence-why-theres-barely-anything-there-in-paul-mccartney-new-release>. *Dream Machine* Issue 8.

29. *The Guardian*, “Musicians must embrace ‘unstoppable force’ of AI, Eurythmics’ Dave Stewart urges.” <https://www.theguardian.com/music/2025/dec/05/musicians-must-embrace-unstoppable-force-of-ai-eurythmics-dave-stewart-urges>. *Dream Machine* Issue 11.

30. *Digital Music News*, “Nearly 800 Creatives, Including Jason Aldean and One Republic, Sign Responsible AI Declaration — ‘Stealing Our Work Is Not Innovation’.” <https://www.digitalmusicnews.com/2026/01/22/stealing-isnt-innovation/>. *Dream Machine* Issue 16.

31. Stability AI, “Universal Music Group and Stability AI Announce Strategic Alliance.” <https://stability.ai/news/universal-music-group-and-stability-ai->

announce-strategic-alliance. *Dream Machine* Issue 5.

32. Stability AI, “Warner Music Group and Stability AI Join Forces To Build The Next Generation Of Responsible AI Tools For Music Creation.” <https://stability.ai/news/warner-music-group-and-stability-ai-join-forces-to-build-next-gen-tools>. *Dream Machine* Issue 8.

33. Universal Music, “Universal Music Group and Splice to Collaborate on the Next Generation of AI-Powered Music Creation Tools for Artists.” <https://www.universalmusic.com/universal-music-group-and-splice-to-collaborate-on-the-next-generation-of-ai-powered-music-creation-tools-for-artists/>. *Dream Machine* Issue 12.

34. LinkedIn / Lexology, “Munich Regional Court rules for GEMA against OpenAI.” Coverage: https://www.linkedin.com/posts/dr-barry-scannell-bbb5aa207_in-a-major-ruling-for-european-copyright-share-7393957246386323457-8bbx. *Dream Machine* Issue 7.

35. EDM.com, “‘Biggest Theft in Music History’: Rights Group Sues Suno as AI Music Showdown Escalates.” <https://edm.com/gear-tech/rights-group-sues-suno-copyright-infringement/>. *Dream Machine* Issue 7.

36. Music Business Worldwide, “Wixen files \$50m copyright suit against Meta.” <https://www.musicbusinessworldwide.com/wixen-files-50m-copyright-suit-against-meta-claims-tech-giant-wants-to-replace-songwriters-with-ai/>. *Dream Machine* Issue 16.

37. *Dream Machine* Issue 17 reportage on UMG’s \$3B suit against Anthropic.

38. Stereogum, “Bandcamp bans AI music.” <https://stereogum.com/2485199/bandcamp-bans-ai-music/news>. *Dream Machine* Issue 14.

39. *Dream Machine* Issue 18 reportage of Deezer licensing its detection tool.

40. TechRadar, “AI music is flooding Spotify, and subscribers are furious.” <https://www.techradar.com/audio/spotify/ai-music-is-flooding-spotify-and-subscribers-are-furious-heres-why-music-fans-no-longer-trust-discover-weekly>. *Dream Machine* Issue 14.

41. CNET, “San Diego Comic-Con Draws a Line: No AI Art Allowed at 2026 Event.” <https://www.cnet.com/culture/san-diego-comic-con-bans-ai-art-for-2026-event/>. *Dream Machine* Issue 16.

42. The Independent, “AI-generated song banned from Swedish charts: ‘It’s deceiving’.” <https://www.independent.co.uk/tv/news/ai-music-song-banned-sweden-spotify-b2901627.html>. *Dream Machine* Issue 15.

43. *Soultracks*, “A.I.-generated music is catchy, familiar... and boring.” <https://soultracks.com/news-ai-generated-music-is-catchy-boring/>. *Dream Machine* Issue 14.

43a. *The Independent*, “AI-generated song banned from Swedish charts: ‘It’s deceiving’.” <https://www.independent.co.uk/tv/news/ai-music-song-banned-sweden-spotify-b2901627.html>. *Dream Machine* Issue 15.

43b. *Marketing Week*, “You can’t dismiss AI ads as slop when they’re winning in testing.” Coverage discussed in *Dream Machine* Issue 22.

44. *Billboard*, “The Real Story Behind The AI Song That Knocked Tyla Off No. 1 On Billboard Afrobeats Chart.” <https://www.billboard.com/pro/ai-song-knocked-tyla-off-no-1-afrobeats/>. *Dream Machine* Issue 30.

45. *MusicTech*, “Jack Antonoff brands AI music makers as ‘godless whores’.” <https://musictech.com/news/industry/jack-antonoff-ai-music-makers-godless-whores/>. *Dream Machine* Issue 30.

45a. Apple Music platform statement, June 2026: AI-generated songs make up less than 1% of plays on the platform. *Dream Machine* Issue 31.

45b. Spotify–Universal Music Group landmark deal allowing Spotify Premium subscribers to create AI covers and remixes of UMG-catalogue tracks as a paid Premium add-on, June 2026. *Dream Machine* Issue 31.

45c. Spotify CEO Daniel Ek defending the platform’s move to AI music as “better than ‘slop’,” June 2026. *Dream Machine* Issue 31.

45d. Swedish startup Tonada — AI music for retailers (June 2026). *Dream Machine* Issue 31.

45e. *Hollywood Reporter*, “Amazon Orders Three Animated Series Using Generative AI.” <https://www.hollywoodreporter.com/tv/tv-news/amazon-orders-three-animated-series-generative-ai-1236606219/>. *Variety*, “Jorge Gutierrez Drops Out of Amazon MGM AI-Generated Series After Backlash.” <https://variety.com/2026/tv/news/jorge-gutierrez-drops-out-amazon-mgm-ai-generated-series-backlash-1236762285/>. *Futurism*, “Amazon AI-Generated Series Canceled.” <https://futurism.com/artificial-intelligence/amazon-ai-generated-series-canceled>. *Dream Machine* Issue 32.

45f. Guillermo del Toro, subtweet responding to Amazon AI Creators Fund announcement, May/June 2026. Coverage in *Dream Machine* Issue 32. Cross-reference: del Toro’s “rather die” statement in Chapter 7, §“Position Three: We are not doing this.”

46. *Digital Trends*, “181,000 mobile games released in last 6 months — AI and vibe coding have unleashed a flood of new games, but not necessarily better ones.”

<https://www.digitaltrends.com/gaming/ai-and-vibe-coding-are-making-more-games-but-not-necessarily-better-ones/>. *Dream Machine* Issue 36.

47. *Metro*, “CD Projekt boss warns that games made solely by gen AI are already being made.” <https://metro.co.uk/2026/06/22/cd-projekt-boss-warns-games-made-solely-gen-ai-already-made-28878296/>. *Dream Machine* Issue 35.

48. *The Guardian*, “Artists are making ‘anti-slop’ to rebel against AI: ‘It’s been rammed down our throats’.” <https://www.theguardian.com/>. *Dream Machine* Issue 33.

49. *Variety*, “Behind the Movie Claiming to Be the First Straight-to-VHS Release in 20 Years — And Why It’s a ‘Middle Finger’ to AI.” <https://variety.com/2026/film/global/behind-the-fstraight-to-vhs-release-in-20-years-and-why-its-a-middle-finger-to-ai-1236767661/>. *Dream Machine* Issue 33.

50. *Digital Music News*, “‘Say No To Suno’ Protesters Bring the Drama to a Swanky Santa Monica AI Conference.” <https://www.digitalmusicnews.com/2026/06/04/say-no-to-suno-santa-monica-plane/>. *Dream Machine* Issue 33.

51. *Musically*, “Tidal allows AI-generated music — but it won’t earn royalties.” <https://musically.com/2026/06/30/tidal-allows-ai-generated-music-but-it-wont-earn-royalties/>. *Dream Machine* Issue 36.

52. *The Hollywood Reporter*, “Faceless Creators Take a Hit As YouTube Cracks Down on AI Slop.” <https://www.hollywoodreporter.com/business/digital/faceless-creators-youtube-ai-damage-1236617586/>. *Dream Machine* Issue 35.

53. *Fast Company*, “Our interest in AI slop is hitting a ceiling.” <https://www.fastcompany.com/91568818/our-interest-in-ai-slop-is-hitting-a-ceiling>. *Dream Machine* Issue 37.

54. *Digital Music News*, “Band Wakes Up to Find AI Has Stolen 94% of Their Spotify Royalties by Duplicating Their Album and Altering the Speed.” <https://www.digitalmusicnews.com/2026/07/12/ai-steals-artist-royalties-speed-altering/>. *Dream Machine* Issue 38.

Chapter 6 — The 88%

40. *Screen Daily*, “We got it wrong says UK Culture Secretary Lisa Nandy of government misstep on AI and copyright strategy,” 1 June 2026. <https://www.screendaily.com/news/we-got-it-wrong-says-uk-culture-secretary-lisa-nandy-of-government-misstep-on-ai-and-copyright-strategy/5217330.article>. *Dream Machine* Issue 32.

1. UK Department for Science, Innovation and Technology, *Statement of Progress on Copyright and AI*, 15 December 2025. <https://www.gov.uk/government/publications/copyright-and-artificial-intelligence-progress-report/copyright-and-artificial-intelligence-statement-of-progress-under-section-137-data-use-and-access-act>. *Dream Machine* Issue 12, “Editor’s Pick: 88% of Creators Said ‘No’.” 18 December 2025.
2. UK DSIT, original consultation, 17 December 2024 – 25 February 2025. Discussion in IPWatchdog, “Respondents to UK AI Consultation Overwhelmingly Want AI Companies to License Copyrighted Works in All Cases.” <https://ipwatchdog.com/2025/12/16/respondents-uk-ai-consultation-overwhelmingly-want-ai-companies-license-copyrighted-works-all-cases/>.
3. IPWatchdog, *op. cit.*; Hogan Lovells, “Copyright and AI: UK government publishes statement of progress.” <https://www.hoganlovells.com/en/publications/copyright-and-ai-uk-government-publishes-statement-of-progress>.
- 3a. Society of Authors submission to the UK consultation, quoted in IPWatchdog, *op. cit.*
4. UK DSIT, *Statement of Progress*, *op. cit.*; analysis at UCL Copyright Queries, “UK government publishes progress statement on AI and copyright consultation.” <https://blogs.ucl.ac.uk/copyright/2025/12/23/uk-government-publishes-progress-statement-on-ai-and-copyright-consultation/>.
5. UK DSIT, *Statement of Progress*, *op. cit.*
6. Dr Barry Scannell, LinkedIn analysis of GEMA v. OpenAI ruling, November 2025. https://www.linkedin.com/posts/dr-barry-scannell-bbb5aa207_in-a-major-ruling-for-european-copyright-share-7393957246386323457-8bbx. *Dream Machine* Issue 7.
7. EDM.com, “‘Biggest Theft in Music History’: Rights Group Sues Suno as AI Music Showdown Escalates.” <https://edm.com/gear-tech/rights-group-sues-suno-copyright-infringement/>. *Dream Machine* Issue 7.
8. *Music Business Worldwide*, “Wixen files \$50m copyright suit against Meta, claims tech giant wants to replace songwriters with AI.” <https://www.musicbusinessworldwide.com/wixen-files-50m-copyright-suit-against-meta-claims-tech-giant-wants-to-replace-songwriters-with-ai/>. *Dream Machine* Issue 16.
9. *Dream Machine* Issue 17, on UMG’s \$3B suit against Anthropic.
10. *Complete Music Update*, “Johnny Cash estate uses ELVIS Act to sue Coke over tribute act ad soundtrack.” <https://completemusicupdate.com/johnny-cash-estate->

uses-elvis-act-to-sue-coke-over-tribute-act-ad-soundtrack/. *Dream Machine* Issue 9.

11. Reuters, “European lawmakers seek EU-wide minimum age to access AI chatbots, social media.” <https://www.reuters.com/legal/litigation/european-lawmakers-seek-eu-wide-minimum-age-access-ai-chatbots-social-media-2025-11-26/>. *Dream Machine* Issue 9.

12. SAG-AFTRA contract update reporting through Q2 2026. *Dream Machine* Issues 20, 26, 29. Coverage: <https://www.theverge.com/news/842848/new-york-law-ai-advertisements-sag-aftra-labor>.

13. Equity (U.K.), “Performers prepared to take industrial action over AI in landslide 99% vote.” <https://www.equity.org.uk/news/2025/performers-prepared-to-take-industrial-action-over-ai-in-landslide-99-vote>. *Dream Machine* Issue 12.

14. Equity (U.K.), “Equity welcomes improved offer in AI protection negotiations in film and TV.” <https://www.equity.org.uk/news/2026/equity-welcomes-improved-offer-in-ai-protection-negotiations-in-film-and-tv>. *Dream Machine* Issue 15.

15. Cannes Film Festival AI Disclosure Standard launch, May 2026. *Dream Machine* Issue 29.

16. UK DSIT, *Statement of Progress*, *op. cit.*

17. *Dream Machine* Issue 21, 19 March 2026, on the UK government’s revised position on AI copyright.

18. *Digital Music News*, “The AI Licensing Shift — Creative Weight Attribution Emerges as Music Industry Game-Changer for Rights Holders.” <https://www.digitalmusicnews.com/2026/01/26/ai-licensing-shift-creative-weight-attribution/>. See also *Digital Music News*, “Artificial Intelligence Attribution and Licensing Startup Musical AI Scores \$4.5 Million Raise.” <https://www.digitalmusicnews.com/2026/01/13/musical-ai-funding-january-2026/>. *Dream Machine* Issues 14, 16.

19. PRS for Music, “PRS for Music AI Survey 2026.” <https://www.prsformusic.com/m-magazine/news/prs-for-music-ai-survey-2026>. *Dream Machine* Issue 16.

20. *Broadcast Now*, “Alex Mahon joins Stellar AI Creative Summit line-up” (covering the launch of the UCL/RCA Centre for Creative AI). <https://www.broadcastnow.co.uk/broadcasters/alex-mahon-joins-stellar-ai-creative-summit-line-up/5209227.article>. *Dream Machine* Issue 1.

21. *Complete Music Update*, “Artists must have creative control in AI deals or risk ending up with ‘scraps’, says US artist trade body.”

<https://completemusicupdate.com/artists-must-have-creative-control-in-ai-deals-or-risk-ending-up-with-scrap-says-us-artist-trade-body/>. *Dream Machine* Issue 6.

22. *Digital Music News*, “Nearly 800 Creatives, Including Jason Aldean and One Republic, Sign Responsible AI Declaration — ‘Stealing Our Work Is Not Innovation’.” <https://www.digitalmusicnews.com/2026/01/22/stealing-isnt-innovation/>. *Dream Machine* Issue 16.

23. For Disney’s parallel position, see *Deadline*, “Disney Sends Cease And Desist Letter To Character.ai.” <https://deadline.com/2025/09/disney-cease-and-desist-letter-characterai-copyright-infringement-1236566831/>. For Studio Ghibli’s similar stance: *NDTV Profit*, “Studio Ghibli And Studio That Developed Elden Ring Send Stern Message To OpenAI.” <https://www.ndtvprofit.com/technology/studio-ghibli-and-studio-that-developed-elden-ring-send-stern-message-to-openai>. *Dream Machine* Issues 2, 6.

24. Adobe, *Creators’ Toolkit Report*, *op. cit.* 69% of 16,000 surveyed creators worried about their work being used to train AI without consent.

25. Adobe Firefly milestone and adoption data. Firefly Foundry and Firefly Image Model 5 launch reporting, Adobe MAX 2025: <https://news.adobe.com/news/2025/10/adobe-max-2025-firefly-foundry>; <https://news.adobe.com/news/2025/10/adobe-max-2025-firefly>.

26. Bria AI consent-licensed dataset and attribution mechanism. [TODO: confirm primary citation — Bria’s licensed-data white paper or Series B coverage.]

27. Getty Images, “Generative AI by iStock” launch, built on NVIDIA Picasso, trained exclusively on Getty’s licensed library with contributor royalties. [TODO: confirm citation — Getty press release or *Reuters* coverage.]

28. Moonvalley Marey, generative-video foundation model trained on licensed video. [TODO: confirm citation — Moonvalley launch coverage in *The Verge* / *TechCrunch*.]

29. AIODE, ethically-trained music creation DAW. See Chapter 17: The Tools, §“Audio modality models.”

30. Stability AI / Universal Music Group strategic alliance: <https://stability.ai/news/universal-music-group-and-stability-ai-announce-strategic-alliance>. Stability AI / Warner Music: <https://stability.ai/news/warner-music-group-and-stability-ai-join-forces-to-build-next-gen-tools>. Universal Music / Splice partnership: <https://www.universalmusic.com/universal-music-group-and-splice-to-collaborate-on-the-next-generation-of-ai-powered-music-creation-tools-for-artists/>. *Dream Machine* Issues 5, 8, 12.

31. Reporting on AI-generated images in the Adobe Stock training corpus, *Bloomberg*, April 2024. [TODO: confirm exact citation.]

32. Adobe Firefly IP indemnification for enterprise customers. [TODO: confirm citation — Adobe enterprise terms or *The Verge* coverage from 2023.]

33. Microsoft, “Microsoft announces new Copilot Copyright Commitment for customers,” 7 September 2023. <https://blogs.microsoft.com/on-the-issues/2023/09/07/copilot-copyright-commitment-ai-legal-concerns/>.

34. Google Cloud Generative AI indemnification: <https://cloud.google.com/blog/products/ai-machine-learning/protecting-customers-with-generative-ai-indemnification>.

35. IBM watsonx uncapped indemnity for enterprise customers. [TODO: confirm citation.]

36. Sundance AI Literacy Initiative, in Chapter 13: Authenticity, the New Scarcity, §“The provenance infrastructure, named.”

37. *Musically*, “BPI sets out transparency and sovereignty demands to secure ‘AI licensing boom’.” <https://musically.com/2026/05/19/bpi-transparency-sovereignty-ai-licensing-boom/>. *Dream Machine* Issue 30.

37a. Spotify and Universal Music Group, “Landmark Licensing Agreements for Fan-Made Covers and Remixes,” 21 May 2026. Spotify Newsroom: <https://newsroom.spotify.com/2026-05-21/universal-music-group-spotify-licensing-agreements-fan-made-covers-remixes/>. *TechCrunch*: <https://techcrunch.com/2026/05/21/spotify-and-universal-music-strike-deal-allowing-fan-made-ai-covers-and-remixes/>. *Medianama*: <https://www.medianama.com/2026/05/223-spotify-umg-fans-create-licensed-ai-covers-remixes/>. *Dream Machine* Issue 31. Detailed treatment of the three-Cs framework, the deal’s prehistory in the UMG settlement with Udio (October 2025), and the *Verified by Spotify* badge.

37b. Hasbro, “Hasbro Launches Sixth Wall, a New AI Studio,” BusinessWire press release, 3 June 2026: <https://www.businesswire.com/news/home/20260603297922/en/>. *The Wrap*: <https://www.thewrap.com/industry-news/business/hasbro-ai-studio-sixth-wall-launch/>. *Screen Daily*: <https://www.screendaily.com/news/hasbro-launches-subsidiary-ai-studio-sixth-wall-to-exploit-ip/5217380.article>. CharacterOS runtime governance layer; talent-participation pay-out structure for original performers; initial 12-character catalogue; 13+/enterprise focus. Cross-reference in Chapter 7 §“Position Five: Own the rails.”

37c. *Dream Machine* Deep Dive № 14, *The Programmable Brand: How AI Is Reshaping IP Licensing in the Generative Economy*, 4 June 2026. The deep dive’s

three Cs framework draws on the Spotify–UMG licensing agreements; the five-model comparison covers Hasbro/Sixth Wall (Behavioural Licensing), Spotify/UMG (platform-embedded remixing), Disney/OpenAI (the framework that survived the platform’s collapse — see footnote ^{3a} in Chapter 7), the independent-artist commons-and-DAO models (Grimes/Elf.tech with a ~50% master-royalty split; Holly Herndon / Holly+ with DAO governance; ElevenLabs heritage voices via CMG Worldwide; Meta’s 2023 celebrity-chatbot programme as the closed-static anti-pattern), and Suno/Alexa as the *create-to-consume* demand floor. The underlying economic argument draws on Coase (1937), Williamson, and Barney’s VRIN resource-based view; the appendix’s load-bearing claim is that AI *relocates* rather than *abolishes* transaction costs — collapsing search and negotiation while raising the cost of monitoring and enforcement, and pushing the durable value into the governance layer.

38. MusicTech, “Tamber is an ‘ethically trained’ AI tool to aid the creative process – and you can use arm gestures to control it.” <https://musictech.com/news/gear/tamber-ai-ethically-trained-arm-gestures/>. Tamber product page: <https://tamber.ai/>. *Dream Machine* Issue 30.

39. Variety, “Is ‘AI Resistance’ Setting the Music Sector Back? WMG’s Robert Kyncl Sees ‘An Incredible Value Creation Opportunity,’ But Warns ‘We Cannot Wait the Way the Industry Did 25 Years Ago’.” <https://variety.com/2026/music/news/wmg-robert-kyncl-ai-resistance-1236748901/>. *Dream Machine* Issue 30.

41. Digital Music News, “Federal Judge Vacates Order Sealing Udio’s Training Data – Are We About to Learn How Many Recordings the Platform Ingested?” <https://www.digitalmusicnews.com/2026/06/05/udio-training-number-order-vacated/>. *Dream Machine* Issue 33.

42. Mixmag, “Artists can now check whether their music has been used to train AI.” <https://mixmag.net/read/artists-check-music-used-train-ai-atlantic-investigation-news>. See also *MusicRadar*, “Is your track in the datasets that are being used to train AI? The Atlantic unveils a way you can find out.” <https://www.musicradar.com/music-tech/is-your-track-in-the-datasets-that-are-being-used-to-train-ai-the-atlantic-unveils-a-way-you-can-find-out>. *Dream Machine* Issues 35, 36.

43. Euronews, “SZA hits out at ‘disgusting’ AI music after discovering her songs were used in AI training.” <https://www.euronews.com/culture/2026/06/22/sza-hits-out-at-disgusting-ai-music-after-discovering-her-songs-were-used-in-ai-training>. *Dream Machine* Issue 35.

44. Digital Music News, “Google Moves to Dismiss Artists’ AI Lawsuit, Says They Consented ‘To the Use of Their Recordings’ When Uploading to YouTube.”

<https://www.digitalmusicnews.com/2026/06/10/youtube-ai-lawsuit-dismissal-motion/>. *Dream Machine* Issue 34.

45. *Billboard*, “YouTube’s AI Training Argument Raises Alarm Among Indie Music Advocates: ‘Not Informed Consent’.” <https://www.billboard.com/pro/youtube-ai-training-argument-indie-music-community-concerns/>. *Dream Machine* Issue 36.

46. *Digital Music News*, “Google Doubles Down on ‘Fair Use’ AI Training, Stressing That ‘Outputs’ Not ‘Inputs’ Are Where Enforcement Should Be Focusing.” <https://www.digitalmusicnews.com/2026/06/29/google-fair-use-ai-training/>. *Dream Machine* Issue 36.

47. *Complete Music Update*, “US musicians’ union sues Universal and Warner over AI deals, claims breach of ‘new use’ clause in record industry’s performer agreement.” <https://completemusicupdate.com/us-musicians-union-sues-universal-and-warner-over-ai-deals-claims-breach-of-new-use-clause-in-record-industrys-performer-agreement/>. *Dream Machine* Issue 33.

48. *Complete Music Update*, “NMPA unveils AI licensing deals with Udio and Klay with 50/50 split for songs and recordings.” <https://completemusicupdate.com/nmpa-unveils-ai-licensing-deals-with-udio-and-klay-with-50-50-split-for-songs-and-recordings/>. *Dream Machine* Issue 34.

49. *Musically*, “Open letter: Musicians ‘must not be pressured into AI deals’ by rightsholders.” <https://musically.com/2026/06/22/open-letter-musicians-must-not-be-pressured-into-ai-deals-by-rightsholders/>. *Dream Machine* Issue 35.

50. GOV.UK, “AI Adoption Plan for the Creative Industries.” <https://www.gov.uk/government/publications/ai-champions-ai-adoption-plans/ai-adoption-plan-creative-industries>. *Dream Machine* Issue 33.

51. GOV.UK / DSIT, “A decisive shift to power British AI: new £1.1 billion plan to back chip firms, boost computing power and skills for the AI revolution.” <https://www.gov.uk/government/news/a-decisive-shift-to-power-british-ai-new-11-billion-plan-to-back-chip-firms-boost-computing-power-and-skills-for-the-ai-revolution>. *Dream Machine* Issue 33.

52. *Fortune*, “Google forced to allow news sites to opt out of AI scraping in ‘world first,’ UK watchdog says.” <https://fortune.com/2026/06/05/auto-draft-3/>. See also *National Law Review*, “UK Regulator Introduces ‘World First’ Controls on the Use of Content in Google Generative AI Features.” <https://natlawreview.com/article/uk-regulator-introduces-world-first-controls-use-content-google-generative-ai>. *Dream Machine* Issues 33, 34.

6a. *Variety*, “Suno has lost a landmark AI copyright lawsuit to German music rights society GEMA.” <https://variety.com/2026/digital/news/suno-loses-ai-lawsuit-gema-1236825010/>. *Dream Machine* Issue 40.

6b. *Creative Bloq*, “What does the biggest AI copyright payout in history actually change for creatives?” <https://www.creativebloq.com/ai/what-does-the-biggest-ai-copyright-payout-in-history-actually-change-for-creatives>. *Dream Machine* Issue 40.

6c. *Digital Music News*, “Suno Quickly Seizes on Udio’s Legal Victory Against Sony Music Over Admissible Infringed Works.” <https://www.digitalmusicnews.com/2026/07/05/suno-admissible-infringed-works-argument/>. *Dream Machine* Issue 37.

46a. *Billboard*, “ABBA’s Björn Ulvaeus Says We Might Be Doing AI Licensing Deals Wrong.” <https://www.billboard.com/pro/abba-bjorn-ulvaeus-music-ai-un-global-summit/>. See also *Musically*, “Tracing the output was always the wrong question’.” <https://musically.com/2026/07/13/bjorn-ulvaeus-on-ai-tracing-the-output-was-always-the-wrong-question/>. *Dream Machine* Issue 38.

Chapter 7 — The Studios Decide

1. *CNBC*, “Netflix ‘all in’ on leveraging AI as the tech creeps into entertainment industry,” 22 October 2025. <https://www.cnbc.com/2025/10/22/netflix-all-in-on-leveraging-ai-in-its-streaming-platform.html>. *Dream Machine* Issue 4.

2. *Futurism*, “Lionsgate’s Attempt to Create Movies Using AI Has Crumbled Into Disaster.” <https://futurism.com/artificial-intelligence/lionsgate-movies-ai>. *Dream Machine* Issue 1.

3. *The Guardian*, “Disney to invest \$1bn in OpenAI, allowing characters in Sora video tool.” <https://www.theguardian.com/business/2025/dec/11/disney-open-ai-sora-video-deal>. *Dream Machine* Issue 11. The announcement was widely treated at the time as a completed, operative template; on the corrected record, no formal agreement was ever signed and no money changed hands.

3a. OpenAI staged Sora wind-down announced 24 March 2026: consumer app and website to end in late April 2026; developer API to close in late September 2026. Disney exited the partnership with reportedly as little as thirty minutes’ notice. *Ropes & Gray*, “The Deal That Wasn’t: What Disney and OpenAI Teach Us About Strategic Investments.” <https://www.ropesgray.com/en/insights/viewpoints/the-deal-that-wasnt-what-disney-and-openai-teach-us-about-strategic-investments>.

TheStreet / AOL (Hollywood Reporter source): <https://www.thestreet.com/technology/openai-is-shutting-down-sora-and-the-disney-deal-is-off>. *Tech-Insider*: <https://tech-insider.org/openai-sora-shutdown-disney-deal-ai-video-2026/>. Peak Sora revenue reported at roughly \$540k/month against operating costs in the billions; OpenAI reallocated compute toward higher-margin coding and enterprise products. Coverage in *Dream Machine* Issue 22.

3b. Dream Machine Deep Dive №14, The Programmable Brand, 4 June 2026.

4. PYMNTS, “Retention Is Name of the Game for Netflix’s AI Strategy.” <https://www.pymnts.com/subscription-commerce/2026/retention-is-name-of-the-game-for-netflixs-ai-strategy/>. *Dream Machine* Issue 15.

5. Deadline, “Amazon Builds Out AI Studios With Sports Docs Boss Matt Newman Named Head Of Live-Action.” <https://deadline.com/2025/11/amazon-ai-studios-matt-newman-1236603477/>. *Dream Machine* Issue 7.

6. Wired, “Amazon’s House of David Used Over 350 AI Shots in Season 2. Its Creator Isn’t Sorry.” <https://www.wired.com/story/amazons-house-of-david-used-over-350-ai-shots-in-season-2-its-creator-isnt-sorry/>. *Dream Machine* Issue 7.

7. Video Games Chronicle, “NBCUniversal signs deal with Law & Order creator Dick Wolf’s son to make AI-generated games based on its IP.” <https://www.videogameschronicle.com/news/nbcuniversal-signs-deal-with-law-order-creator-dick-wolfs-son-to-make-ai-generated-games-based-on-its-ip/>. *Dream Machine* Issue 5.

8. NME, “The Office’, ‘Saturday Night Live’ and ‘Sex And The City’ could be turned into AI games.” <https://www.nme.com/news/gaming-news/the-office-and-sex-and-the-city-ai-video-games-3901630>. *Dream Machine* Issue 5.

9. The Hollywood Reporter, “Disney+ to Allow User-Generated Fan Content with AI.” <https://www.hollywoodreporter.com/business/digital/disney-plus-gen-ai-user-generated-content-1236426135/>. *Dream Machine* Issue 8.

10. Dream Machine Issue 8 reportage of the Disney “Office of Technology Enablement,” led by former Walt Disney Studios CTO Jamie Voris.

11. Marketing Dive, “Disney unveils TikTok-like vertical video, AI video generation tool.” <https://www.marketingdive.com/news/disney-unveils-tiktok-like-vertical-video-ai-video-generation-tool/809269/>. *Dream Machine* Issue 14.

12. The Hollywood Reporter, “Fox Entertainment Takes Equity Stake in AI-Microdramas Company Holywater.” <https://www.hollywoodreporter.com/business/business-news/fox-entertainment-invests-in-holywater-ai-microdramas-1236396802/>. *Dream Machine* Issue 3.

13. Deadline, “Sky History Acquires ‘Castles SOS,’ AI-Powered Doc Exploring Royalty, Ruins & Restoration.” <https://deadline.com/2025/11/castles-sos-ai-doc-sky-history-documentary-rick-edwards-1236627378/>. *Dream Machine* Issue 9.

14. Estate Agent Today, “Homebuilder among first to use Channel 4’s AI ads.” <https://www.estateagenttoday.co.uk/breaking-news/2025/12/homebuilder-among-first-to-use-channel-4s-ai-ads/>. *Dream Machine* Issue 11.

15. *The Hollywood Reporter*, “Fremantle Names Boss of New AI Native Studio Imaginae Studios.” <https://www.hollywoodreporter.com/business/digital/fremantle-names-ceo-new-ai-label-imaginae-studios-1236396579/>. *Dream Machine* Issue 2.
16. *Dream Machine* Issue 25, on Fremantle’s *Art Awakens* development.
17. *Indiewire*, “Another New AI Production Company Inks a Big Creative Partnership — This Time, with Ron Howard and Brian Grazer’s Imagine Entertainment.” <https://www.indiewire.com/news/business/obsidian-studio-ai-production-company-imagine-entertainment-1235158619/>. *Dream Machine* Issue 6.
18. *UK Tech News*, “AI film studio Wonder lands \$9m investment.” <https://www.uktech.news/ai/ai-film-studio-wonder-lands-9m-investment-20251023>. *Dream Machine* Issue 5.
19. Wonder Studios, “Shortlisted films revealed for The Wonder Film Festival.” https://www.linkedin.com/posts/wearewonderstudios_were-thrilled-to-share-the-shortlisted-films-activity-7404560378082246656-7NcI. *Dream Machine* Issue 11.
20. *The Hollywood Reporter*, “AI Company Asteria Produces New Animated Short ‘All Heart’.” <https://www.hollywoodreporter.com/movies/movie-news/natasha-lyonne-ai-company-asteria-1236403144/>. *Dream Machine* Issue 4.
21. *The Hollywood Reporter*, “Promise, a deep-pocketed AI studio backed by Google, aims to Bring GenAI Filmmaking and VFX to Legacy Media.” <https://www.hollywoodreporter.com/business/digital/ai-studio-promise-vfx-generation-company-1236397636/>. *Dream Machine* Issue 3.
22. *Variety*, “AI-Powered Cinematic Universe Platform enGEN3 Launched by Goldfinch.” <https://variety.com/2025/film/news/ai-powered-cinematic-universe-platform-engen3-1236543349/>. *Dream Machine* Issue 2.
23. *Deadline*, “Munich Based Beta Films & Industry Execs Join Forces To Launch Artificial Intelligence Start-Up Chapter41.” <https://deadline.com/2025/11/beta-film-ai-startup-chapter41-artificial-intelligence-1236612632/>. *Dream Machine* Issue 7.
24. *The Hollywood Reporter*, “Longtime TV Exec, Kevin Reilly, Set to Lead AI Startup Kartel.” <https://www.hollywoodreporter.com/business/digital/kevin-reilly-ceo-kartel-ai-hbo-1236424692/>. *Dream Machine* Issue 7.
25. *Variety*, “‘Wanted’ Director Timur Bekmambetov Explains His \$5 Million Plan to Generate AI Method Actors: ‘AI Is Here to Stay. We Have to Train It Responsibly’.” <https://variety.com/2025/film/news/wanted-director-method-acting-ai-actors-1236579647/>. *Dream Machine* Issue 7.

26. Variety, “Tilly Norwood Creator Doubles Down on AI ‘Actors’ and Says It’s a ‘More Ethical Way to Perform,’ Urges Human Actors to ‘Future-Proof’ Themselves With AI.” <https://variety.com/2026/digital/news/tilly-norwood-creator-tells-actors-to-create-ai-avatars-1236638940/>. *Dream Machine* Issue 16.

27. *Broadcast Now*, “Wonder Studios adapts children’s book to animated series with AI.” <https://www.broadcastnow.co.uk/production-and-post/wonder-studios-adapts-childrens-book-to-animated-series-with-ai/5211713.article>. *Dream Machine* Issue 11.

28. Variety, “‘Watch the Skies,’ Swedish UFO Feature Film Dubbed Entirely With AI, Sets USA Distribution Deal.” <https://variety.com/2025/film/news/watch-the-skies-us-theatrical-release-ai-dubbing-1236343110/>. *Dream Machine* Issue 5.

29. *Cybernews*, “Run to the West — South Korea’s first AI film tests the soul of cinema.” <https://cybernews.com/entertainment/korean-cinema-run-to-the-west-ai/>. *Dream Machine* Issue 5.

30. *Broadcast Pro Middle East*, “Tunisian filmmaker wins \$1 million AI Film Award for ‘Lily’.” <https://www.broadcastprome.com/news/tunisian-filmmaker-wins-1-million-ai-film-award-for-lily/>. *Dream Machine* Issue 14.

31. Variety, “AI Drama ‘Humans in the Loop’ Receives Film Independent’s Sloan Distribution Grant, Enters Oscar Race.” <https://variety.com/2025/film/news/ai-drama-humans-in-the-loop-oscar-race-1236582975/>. *Dream Machine* Issue 8.

32. *PC Gamer*, “Palworld studio Pocketpair says its new publishing division won’t handle games that use generative AI: ‘We don’t believe in it.’” <https://www.pcgamer.com/software/ai/palworld-studio-pocketpair-says-its-new-publishing-division-wont-handle-games-that-use-generative-ai-we-dont-believe-in-it/>. *Dream Machine* Issue 4.

33. *Niche Gamer*, “Larian Studios backs off from gen AI.” *Dream Machine* Issue 14.

33a. Larian Studios policy framing on the next *Divinity*, January 2026; same source as note 33 above.

34. *Decrypt*, “Warhammer 40,000 Maker Games Workshop Rules Out Generative AI.” *Dream Machine* Issue 14.

35. *Niche Gamer*, “Manor Lords publisher Hooded Horse won’t work with devs using gen AI.” *Dream Machine* Issue 14.

36. *gamesindustry.biz*, “RuneScape maker Jagex says it will never use generative AI to make in-game content.” *Dream Machine* Issue 16.

37. *GamesRadar*, “Wallace and Gromit creator says beloved animation studio Aardman will ‘embrace the technology’ of AI, but will be ‘very cautious not to lose our values’.” <https://www.gamesradar.com/entertainment/animation->

movies/wallace-and-gromit-creator-says-beloved-animation-studio-aardman-will-embrace-the-technology-of-ai-but-will-be-very-cautious-not-to-lose-our-values/.
Dream Machine Issue 11.

38. Variety, “Guillermo del Toro Says He’d ‘Rather Die’ Than Use Generative AI in His Films: ‘Not Interested’.” <https://variety.com/2025/film/news/guillermo-del-toro-rather-die-generative-ai-frankenstein-1236561316/>. *Dream Machine* Issue 5.

39. *The Hollywood Reporter*, “Leonardo DiCaprio Says AI Can’t Be Art Because ‘There’s No Humanity to It’.” <https://www.hollywoodreporter.com/movies/movie-news/leonardo-dicaprio-ai-cant-be-art-no-humanity-1236445405/>. *Dream Machine* Issue 11.

40. *Daily Mail*, “Claire Foy says she has ‘no interest’ in seeing AI in films.” <https://www.dailymail.co.uk/tvshowbiz/article-15454199/Claire-Foy-AI-films-sad-disappointed-people-future-Hollywood.html>. *Dream Machine* Issue 14.

41. *NME*, “Jenna Ortega says it’s ‘very easy to be terrified’ of AI in filmmaking.” <https://www.nme.com/news/jenna-ortega-says-its-very-easy-to-be-terrified-of-ai-in-filmmaking-3913926>. *Dream Machine* Issue 10.

42. Variety, “Chris Pratt Pitched Having an AI ‘Actor’ Star as the Villain in ‘Mercy’: ‘I Don’t Think That’s a Good Idea at All’.” <https://variety.com/2026/film/news/chris-pratt-ai-actor-villain-mercy-amazon-mgm-1236640460/>. *Dream Machine* Issue 16.

43. *PC Gamer*, “Todd Howard says AI can’t replace human ‘creative intention,’ but it’s part of Bethesda’s ‘toolset for how we build our worlds or check things’.” <https://www.pcgamer.com/gaming-industry/todd-howard-says-ai-cant-replace-human-creative-intention-but-its-part-of-bethedas-toolset-for-how-we-build-our-worlds-or-check-things/>. *Dream Machine* Issue 11.

44. *GamesRadar*, “Battlefield 6 lead calls generative AI ‘very seducing,’ but says it was only used in the game’s earliest stages ‘to allow for more time and more space to be creative’.” <https://www.gamesradar.com/games/battlefield/battlefield-6-lead-calls-generative-ai-very-seducing-but-says-it-was-only-used-in-the-games-earliest-stages-to-allow-for-more-time-and-more-space-to-be-creative/>. *Dream Machine* Issue 3.

44a. *Wired*, “Amazon’s House of David Used Over 350 AI Shots in Season 2. Its Creator Isn’t Sorry.” <https://www.wired.com/story/amazons-house-of-david-used-over-350-ai-shots-in-season-2-its-creator-isnt-sorry/>. *Dream Machine* Issue 7.

45. *gamesindustry.biz*, “Witcher 3 and Cyberpunk 2077 director says AI can help, but not replace, creatives.” <https://www.gamesindustry.biz/witcher-3-and-cyberpunk-2077-director-says-ai-can-help-but-not-replace-creatives>. *Dream Machine* Issue 9.

46. *GamesRadar*, “Aardman” *op. cit.*

47. *Dream Machine* Issue 29, on Sony’s “all in on AI for games” announcement.

48. *Hollywood Reporter*, “Netflix is building and recruiting for an AI animation studio, called INKubator, to produce ‘feature-quality’ shorts.” <https://www.hollywoodreporter.com/business/business-news/netflix-ai-animation-studio-inkubator-1236592110/>. *Dream Machine* Issue 30.

49. *Forbes*, “Meet Wonder Studios, The \$50M British Studio Striving To Become The A24 Of AI Production.” <https://www.forbes.com/sites/chariefink/2026/05/18/meet-wonder-studios-the-50m-british-studio-striving-to-become-the-a24-of-ai-production/>. *Dream Machine* Issue 30.

50. *Variety*, “Kling AI Partners With Evolutionary Films on Animated Feature ‘Minibots,’ Unveils Filmmaker Initiative at Cannes Market.” <https://variety.com/2026/film/news/kling-ai-evolutionary-films-minibots-cannes-1236748590/>. *Dream Machine* Issue 30.

51. *Variety*, “AI Dominates Cannes Buzz as Filmmakers Grudgingly Accept It.” <https://variety.com/2026/film/festivals/ai-cannes-2026-filmmakers-accept-1236748402/>; *Hollywood Reporter*, “At Cannes, filmmakers shift towards cautious acceptance of AI’s inevitability.” <https://www.hollywoodreporter.com/business/business-news/cannes-2026-ai-acceptance-1236592488/>. *Dream Machine* Issue 30.

52. *Variety*, “Is AI Basically Like Special Effects? Peter Jackson Seems to Think So.” <https://variety.com/2026/film/news/peter-jackson-ai-special-effects-1236748120/>. *Dream Machine* Issue 30.

53. *PC Gamer*, “Take-Two’s CEO says AI’s not in the business of making hits, ‘datasets by their very nature are backward looking’, but that doesn’t mean AI can’t be ‘super helpful’.” <https://www.pcgamer.com/games/take-two-ceo-ai-not-making-hits-backward-looking/>. *Business Insider*, “The CEO behind Grand Theft Auto says he’s pro AI — but the technology can’t make an original hit.” <https://www.businessinsider.com/take-two-ceo-strauss-zelnick-ai-original-hits-2026-5>. *Dream Machine* Issue 30.

54a. Hasbro launches Sixth Wall, a new AI studio, with the CharacterOS governance layer underneath it, 3 June 2026. BusinessWire press release: <https://www.businesswire.com/news/home/20260603297922/en/>. *The Wrap*, “Hasbro AI Studio Sixth Wall Launch”: <https://www.thewrap.com/industry-news/business/hasbro-ai-studio-sixth-wall-launch/>. *Screen Daily*: <https://www.screendaily.com/news/hasbro-launches-subsiary-ai-studio-sixth-wall-to-exploit-ip/5217380.article>. *BriefGlance* analysis (FTC inquiries into AI

companions, Sixth Wall's stated 13+/enterprise focus): <https://briefglance.com/articles/hasbros-sixth-wall-taming-the-ai-wild-west-with-behavioral-licensing>. CEO Roberta Thomson; CharacterOS encodes character canon, voice, personality, lore and safety limits at runtime; talent-participation pay-out structure for the original performers; initial 12-character catalogue; pilot grew out of an internal generative-AI Ouija-board experiment that logged hundreds of thousands of player interactions in its first few days.

54b. Spotify and Universal Music Group announce landmark linked recorded-music and publishing licensing agreements for fan-made AI covers and remixes at Spotify Investor Day, 21 May 2026. Spotify Newsroom: <https://newsroom.spotify.com/2026-05-21/universal-music-group-spotify-licensing-agreements-fan-made-covers-remixes/>. *TechCrunch*: <https://techcrunch.com/2026/05/21/spotify-and-universal-music-strike-deal-allowing-fan-made-ai-covers-and-remixes/>. *Medianama*: <https://www.medianama.com/2026/05/223-spotify-umg-fans-create-licensed-ai-covers-remixes/>. Three-pillar consent / credit / compensation framework; opt-in catalogue; Premium add-on; new royalty stream; UMG settlement with Udio in October 2025 is the architectural ancestor. Cross-reference in Chapter 5 and Chapter 6.

5a. Amazon AI Creators Fund launch and the Punky Duck cancellation. *Hollywood Reporter*: <https://www.hollywoodreporter.com/tv/tv-news/amazon-orders-three-animated-series-generative-ai-1236606219/>. *Variety*: <https://variety.com/2026/tv/news/jorge-gutierrez-drops-out-amazon-mgm-ai-generated-series-backlash-1236762285/>. *Futurism*: <https://futurism.com/artificial-intelligence/amazon-ai-generated-series-canceled>. *Dream Machine* Issue 32. Cross-reference: Chapter 5 §“The pushback” for the audience-enforcement analysis.

55. *Variety*, “Tribeca Festival AI Film ‘Dreams of Violets’ by Fountain o.” <https://variety.com/2026/film/festivals/tribeca-festival-ai-film-dreams-of-violets-fountain-o-1236759724/>. Director Ash Koosha; producer Fountain o; 75-minute docudrama; subject: 2026 Iran protests; cost \$2,000; production time: 3 months; tools: Kling AI (video), Anthropic Claude (editing), Google Gemini (research). *Dream Machine* Issue 32.

56. The broader festival-circuit context: the *Dreams of Violets* Tribeca selection follows the *Lily* (\$1M Dubai AI Film Award, January 2026) and *Humans in the Loop* (Film Independent Sloan Distribution Grant, November 2025) as part of the accumulating festival recognition of AI-native feature and short-film work — but *Dreams of Violets* is the first Tier 1 festival selection for a feature-length, fully-AI-generated work.

57. *Variety*, “Martin Scorsese Supports AI Company, Endorses AI for Storyboarding.” <https://variety.com/2026/film/news/martin-scorsese-supports-ai-company-storyboard-movies-1236765037/>. Scorsese becomes advisor to Black Forest Labs (FLUX model); endorses AI storyboarding; Boots Riley and Guillermo del Toro respond critically. *Dream Machine* Issue 32.

5b. *The Verge*, “Amazon’s new plan for games: James Bond and AI Snoop Dogg.” <https://www.theverge.com/games/943147/amazon-gaming-strategy-james-bond-snoop-dogg-luna>. *Dream Machine* Issue 33.

56a. *Variety*, “‘Dreams of Violets’ Review: What Does a Film Made Entirely with AI Look Like?” <https://variety.com/2026/film/reviews/dreams-of-violets-review-ash-koosha-tribeca-festival-1236784209/>. *Dream Machine* Issue 35.

57a. *The Hollywood Reporter*, “Ron Howard Is Bullish on AI and Thinks Jim Henson Would Have Been Too.” <https://www.hollywoodreporter.com/business/digital/ron-howard-ai-jim-henson-martin-scorsese-1236620457/>. *Dream Machine* Issue 34.

57b. *Variety*, “Gore Verbinski Says We Need a ‘Rating System’ for AI Use in Films: ‘If You Use AI to Write a Script, You Get an F’.” <https://variety.com/2026/film/global/gore-verbinski-ai-taormina-1236780502/>. *Dream Machine* Issue 34.

57c. *The Guardian*, “‘It has no taste whatsoever. It is a genius at helping you... Like a billionaire on acid’: Star Wars director Gareth Edwards comes out in favour of AI.” <https://www.theguardian.com/film/2026/may/29/star-wars-rogue-one-director-gareth-edwards-praises-ai-film-making>. *Dream Machine* Issue 35.

57d. *AOL*, “Emily Blunt refused Steven Spielberg’s offer to use AI for her alien voice in ‘Disclosure Day’.” <https://www.aol.com/articles/emily-blunt-refused-steven-spielbergs-163209822.html>. *Dream Machine* Issue 34.

58. *The Hollywood Reporter*, “Google Investing in ‘Backrooms’ Studio A24.” <https://www.hollywoodreporter.com/business/digital/a24-google-deepmind-ai-venture-backrooms-1236627228/>. *Dream Machine* Issue 35.

59. *Futurism*, “A24 Is Dealing With a Full-on Crisis Among Fans After Signing an AI Deal With Google.” <https://futurism.com/artificial-intelligence/a24-crisis-fans-signing-ai-deal-google>. *Dream Machine* Issue 36.

60. *Bloomberg*, “Animation’s AI Reckoning: Filmmakers Say They Can Make Movies for 90% Less.” <https://www.bloomberg.com/news/newsletters/2026-06-28/animation-s-ai-reckoning-filmmakers-say-they-can-make-movies-for-90-less>. *Dream Machine* Issue 36.

61. CNN, “South Korea’s movie industry is embracing AI. Will K-film lose its magic?” <https://www.cnn.com/world/south-korea-film-industry-ai-hnk-spc>. *Dream Machine* Issue 35.
62. *The Hollywood Reporter*, “How the Shanghai Film Festival Is Embracing Cinema’s Future: From AI to iPhone Moviemaking.” <https://www.hollywoodreporter.com/movies/movie-news/shanghai-film-festival-cinema-future-ai-iphone-1236626463/>. *Dream Machine* Issue 35.
63. *Variety*, “MiniMax Launches All-in-One AI Video Generator Hub at Shanghai Film Festival: ‘Creative Direction and Aesthetic Judgment Must Ultimately Be Left to Humans’.” <https://variety.com/2026/film/festivals/minimax-hub-ai-video-generator-shanghai-film-festival-1236781901/>. *Dream Machine* Issue 35.
64. *Deadline*, “AI Firm Utopai Studios To Co-Produce Silver Bear-Winning Filmmaker Hyo-Joo Yang’s Feature ‘Half Moon’.” <https://deadline.com/2026/06/utopai-studios-hyo-joo-yang-half-moon-ai-1236952730/>. *Dream Machine* Issue 34.
65. *Variety*, “Fully AI-Generated Series ‘Journey to the West’ in Development by Utopai Studios and Huace.” <https://variety.com/2026/tv/news/fully-ai-generated-journey-to-the-west-utopai-studios-huace-1236787171/>. *Dream Machine* Issue 35.
66. *The Hollywood Reporter*, “‘Sheep in the Box’: Neon Drops First Trailer for Hirokazu Koreeda’s AI Family Drama.” <https://www.hollywoodreporter.com/movies/movie-news/sheep-in-the-box-neon-trailer-hirokazu-koreeda-ai-drama-1236629132/>. *Dream Machine* Issue 36.
67. *Deadline*, “Lionsgate’s Michael Burns Says AI Will Save The Company ‘Tens And Tens Of Millions Of Dollars A Year’.” <https://deadline.com/2026/06/lionsgate-michael-burns-ai-saves-millions-runway-1236941043/>. *Dream Machine* Issue 33.
68. *Runway*, “Runway and Lionsgate Expand Partnership” — Lionsgate takes an equity stake in Runway. <https://runwayml.com/news/runway-and-lionsgate-expand-partnership>. *Dream Machine* Issue 34.
69. *Variety*, “AI Script-Analysis Start-Up Quilty Strikes Preferred-Look Deal With Giovanni Entertainment.” <https://variety.com/2026/film/news/ai-script-quilty-preferred-look-deal-giovanni-entertainment-1236772175/>. *Dream Machine* Issue 34.
70. *The Verge*, “Can AI Tell if your script will make a hit film?” <https://www.theverge.com/ai-artificial-intelligence/943531/ai-script-quilty-simon-horsman-daniel-wood>. *Dream Machine* Issue 33.
71. *Variety*, “About 300 Netflix Titles Used Generative AI This Year, Company Reveals.” <https://variety.com/2026/biz/news/about-300-netflix-programs-used-ai>

Chapter 8 — The Machine Gets the Money

1. Drake Star Global Gaming Report Q2 2026, via *MarketScale*, “Gaming Investment Hits 12-Month High of \$2.5B+ in Q2 2026, Led by AI, AdTech, and Hardware Deals.” <https://www.marketscale.com/industries/sports-entertainment/gaming-investment-hits-12-month-high-of-25b-in-q2-2026-led-by-ai-adtech-and-hardware-deals>. Note: Drake Star characterises GTA 6’s November 2026 release as likely the largest game launch in history.

2. PYMNTS, “AI Worries Erase \$1 Trillion From Big Tech Firms.” <https://www.pymnts.com/artificial-intelligence-2/2026/ai-worries-erase-1-trillion-from-big-tech-firms/>. *Dream Machine* Issue 19.

2a. PYMNTS, “AI Now Responsible for a Third of US Economic Growth.” <https://www.pymnts.com/economy/2026/ai-now-responsible-for-a-third-of-us-economic-growth/>. *Dream Machine* Issue 40.

3. Grand View Research, “Generative AI Content Creation Market” (\$14.8 billion market in 2024, projected to quintuple to \$80 billion by 2030). <https://www.grandviewresearch.com/industry-analysis/generative-ai-content-creation-market-report>. *Dream Machine* Issue 4.

4. *Music Ally*, “GenAI music firm Suno raises \$250m at \$2.45bn valuation.” <https://musically.com/2025/11/19/genai-music-firm-suno-raises-250m-at-2-45bn-valuation/>. *Dream Machine* Issue 8. *Variety*, “AI Music Startup Suno Raises \$400M Funding at \$5.4 Billion Valuation.” <https://variety.com/2026/digital/news/ai-music-suno-funding-round-400-million-5-4-billion-valuation-1236765727/>. *Dream Machine* Issue 32. *Digital Music News*, “Is Apple Plotting an AI-Powered Audio Expansion? Tech Giant Closes Reportedly \$1.6 Billion Qai Buyout — Its ‘Second-Largest Acquisition’ to Date.” <https://www.digitalmusicnews.com/2026/01/30/apple-qai-acquisition/>. *Dream Machine* Issue 17.

5. Peony, “10 Best Gaming Investors in 2026” (gaming venture funding fell from a \$12.5 billion peak in 2021 to \$2.54 billion in 2024). <https://www.peony.ink/blog/top-gaming-investors>.

6. Wikipedia, “2022–2026 video game industry layoffs” (an estimated 45,000 jobs lost from 2022 to July 2025). https://en.wikipedia.org/wiki/2022%E2%80%932026_video_game_industry_layoffs. *Game Developer*, “Data points to slowing layoffs, but doesn’t capture true harm to game industry” (Farhan Noor’s tracker: 8,500 layoffs in 2022; 10,500 in 2023; 14,600 in 2024; ~4,000 by mid-2025).

<https://www.gamedeveloper.com/business/industry-layoffs-are-seemingly-slowing-but-the-damage-has-already-been-done>.

7. *PC Gamer*, “As the videogame industry continues to be hammered by layoffs, Netflix is offering up to \$840,000 per year for a new Director of Generative AI for Games.” <https://www.pcgamer.com/gaming-industry/as-the-videogame-industry-continues-to-be-hammered-by-layoffs-netflix-is-offering-up-to-usd840-000-per-year-for-a-new-director-of-generative-ai-for-games/>. *Dream Machine* Issue 1.

8. PitchBook, “Q2 2026 Gaming VC First Look.” <https://pitchbook.com/news/reports/q2-2026-gaming-vc-first-look>.

9. *Forbes*, “Five Venture Capital Trends: Where Gaming Capital Is Moving In 2026.” <https://www.forbes.com/councils/forbesbusinesscouncil/2026/07/31/five-venture-capital-trends-where-gaming-capital-is-moving-in-2026/>. OECD figure: OECD, “AI firms capture 61 percent of global venture capital in 2025.” <https://www.oecd.org/en/about/news/announcements/2026/02/ai-firms-capture-61-percent-of-global-venture-capital-in-2025.html>.

10. *GamesBeat*, “Fifth Door raises \$20M for AI-powered game creation.” <https://gamesbeat.com/fifth-door-raises-20m-for-ai-powered-game-creation/>.

Dream Machine Issue 9. *PYMNTS*, “Razer Spends \$600 Million to Sharpen Focus on AI Gaming.” <https://www.pymnts.com/news/artificial-intelligence/2026/razer-spends-600-million-dollars-sharpen-focus-ai-gaming/>. *Dream Machine* Issue 13.

CTech, “Sett secures \$30 million Series B to automate game marketing with AI agents.” <https://www.calcalistech.com/ctechnews/article/rjxbyqisze>. *Dream Machine* Issue 23.

Astrocade, “Astrocade raises \$56M to help build the next era of interactive entertainment.” <https://www.astrocade.com/blog/astrocade-raises-56m-funding>. *Dream Machine* Issue 28.

The AI Insider, “AI Game Creation Platform Aippy Announces Tens of Millions at \$250M Valuation.” <https://theaiinsider.tech/2026/06/03/ai-game-creation-platform-aippy-announces-tens-of-millions-at-250m-valuation/>. *Dream Machine* Issue 32.

GamesBeat, “Welevel unveils AI-based city-building survival game SolidRiver and \$8.5M in funding.” <https://gamesbeat.com/welevel-unveils-ai-based-city-building-survival-game-solidriver-and-8-5m-in-funding/>. *Dream Machine* Issue 34.

11. *TechCrunch*, “General Intuition lands \$134M seed to teach agents spatial reasoning using video game clips.” <https://techcrunch.com/2025/10/16/general-intuition-lands-134m-seed-to-teach-agents-spatial-reasoning-using-video-game-clips/>. *Dream Machine* Issue 3.

GamesBeat, “General Intuition raises \$320M at \$2.3B valuation for AI frontier models based on gameplay.” <https://gamesbeat.com/general-intuition-raises-320m-at-2-3b-valuation-for-ai-frontier-models-based-on-gameplay-exclusive-interview/>. *TechCrunch*, “General Intuition’s \$2.3B bet that video games can train AI agents for the real world.”

<https://techcrunch.com/2026/06/25/general-intuitions-2-3b-bet-that-video-games-can-train-ai-agents-for-the-real-world/>. *Dream Machine* Issue 36. *GamesBeat*, “Origin Lab raises \$8 million to glean AI training data from game worlds.” <https://gamesbeat.com/origin-lab-raises-8-million-turns-game-worlds-data-ai/>. *Dream Machine* Issue 29.

11a. *Tech.eu*, “Cambridge’s Worldmodeldata raises £7M to turn video games into AI training data.” <https://tech.eu/2026/07/06/worldmodeldata-lands-ps7m-to-turn-gaming-data-into-ai-training/>. *TechCrunch*, “Meta quietly launches vibe-coded gaming app Pocket.” <https://techcrunch.com/2026/07/02/meta-quietly-launches-vibe-coded-gaming-app-pocket/>. Both in *Dream Machine* Issue 37.

12. On the Take-Two CEO and franchise origination, see Chapter 12, *The Orchestrator: NDTV Profit*, “Don’t Expect AI To Invent the Next ‘Grand Theft Auto’,” Says Take-Two CEO Strauss Zelnick.” <https://www.ndtvprofit.com/technology/dont-expect-ai-to-invent-the-next-grand-theft-auto-says-take-two-ceo-strauss-zelnick>. *Dream Machine* Issue 6. On the internal AI team layoff: *Metro*, “GTA 6 publisher lays off its entire AI team after seven years.” <https://metro.co.uk/2026/04/06/gta-6-publisher-lays-off-entire-ai-team-seven-years-27874034/>. *Dream Machine* Issue 24. See also Chapter 14, *Coordination Collapse*.

13. *CNBC*, “Saudi PIF and Kushner’s Affinity finalize \$55 billion EA Sports deal.” <https://www.cnn.com/2026/08/05/saudi-wealth-fund-jared-kushner-ea-sports-deal.html>. Structure and completion (4 August 2026): *Wikipedia*, “Leveraged buyout of Electronic Arts.” https://en.wikipedia.org/wiki/Leveraged_buyout_of_Electronic_Arts.

14. Consortium AI framing: eciks, “Electronic Arts completes \$55B takeover, becomes private company.” <https://eciks.org/19135-ea-55-billion-takeover-private>. Player base of 700 million: AP via *ABC News*, “Video game giant Electronic Arts closes \$55 billion go-private sale of its business.” <https://abcnews.com/Technology/wireStory/video-game-giant-electronic-arts-closes-55-billion-135394606>. \$22bn AI-driven profit-headroom analysis: *Reuters*, “Gaming industry could unlock \$22 billion in profits on AI-driven cost cuts, says Morgan Stanley.” <https://www.reuters.com/business/gaming-industry-could-unlock-22-billion-profits-ai-driven-cost-cuts-morgan-2026-04-22/>. *Dream Machine* Issue 26.

15. Akker, “AI Film Insurance 2026: The Coverage Gap Hollywood Is Not Talking About.” <https://www.akkerins.com/new-blog/ai-film-production-insurance-gap-2026>.

16. Roger Avary, *The Joe Rogan Experience*, Episode 2452 (February 2026), quoted in Akker, “AI Film Insurance 2026: The Coverage Gap Hollywood Is Not Talking

About,” as note 15.

17. Wonder £9m seed: UKTN, “AI film studio Wonder lands £9m investment.” <https://www.uktech.news/ai/ai-film-studio-wonder-lands-9m-investment-20251023>. *Dream Machine* Issue 5. Wonder Studios at \$50m: *Deadline*, “Meet Wonder Studios, The \$50M British Studio Striving To Become The A24 Of AI Production.” <https://deadline.com/2026/05/wonder-studios-a24-ai-production-1236900795/>. *Dream Machine* Issue 30. Filmology Labs \$250m: *Variety*, “Filmology Labs, a \$250 Million Studio Built for Content Creation Like Vertical Micro-Dramas and AI-Driven Media Formats, Set to Open in New Jersey.” <https://variety.com/2026/digital/news/filmology-labs-content-creation-ai-studio-1236657646/>. *Dream Machine* Issue 18. Video Rebirth \$50m: *Crowdfund Insider*, “Singapore’s Video Rebirth Raises \$50m to Chase Studio-Grade AI Video.” <https://www.crowdfundinsider.com/2025/11/255411-singapores-video-rebirth-raises-50m-to-chase-studio-grade-ai-video/>. *Dream Machine* Issue 7. Flick \$6m YC seed: *Business Insider*, “This couple (literally) married AI and filmmaking. Read the pitch deck they used to raise \$6 million for their startup.” <https://www.businessinsider.com/ai-filmmaking-startup-flick-lands-vc-seed-funding-pitch-deck-2026-5>. *Dream Machine* Issue 30.

18. Ciaro, “Why Investors Fund Some Film Pitches and Ignore Others.” <https://ciaro.pro/blog/why-investors-fund-some-film-pitches-and-ignore-others>.

19. *DigiTimes*, “Seedance 2.0 cuts video generation cost below US\$0.14/sec.” <https://www.digitimes.com/news/a20260306PD225/cost-technology-production-development.html>. *Dream Machine* Issue 20.

20. *The Independent*, “Lights, camera, algorithm: First fully AI-generated film set to premiere at Tribeca Festival.” <https://www.independent.co.uk/arts-entertainment/films/news/ai-generated-film-tribeca-dreams-of-violet-b2984749.html>. *Dream Machine* Issue 31. *Digital Camera World*, “The CGI would have cost millions. I spent \$2,000.’ Is Dreams of Violets AI slop – or the future of film-making?” <https://www.digitalcameraworld.com/tech/artificial-intelligence/a-major-film-festival-will-soon-debut-the-first-fully-ai-generated-feature-inspired-by-real-events-this-could-be-a-major-shift-for-filmmaking>. *Dream Machine* Issue 32.

21. *Wall Street Journal*, “This Film Cost \$500,000 to Make. \$400,000 Was AI Compute Costs.” https://www.wsj.com/cio-journal/this-cannes-film-cost-500-000-to-make-400-000-was-ai-compute-costs-a823b08d?st=2a1WRJ&reflink=desktopwebshare_permalink. *Dream Machine* Issue 31.

22. *Decrypt*, “Killing Satoshi’: \$70M Bitcoin Film Starring Casey Affleck Will Make Heavy Use of AI.” <https://decrypt.co/364505/killing-satoshi-70m-bitcoin-film-starring-casey-affleck-will-make-heavy-use-of-ai>. *Dream Machine* Issue 26.

23. *TechCrunch*, “Runway CEO says AI could help Hollywood make 50 films instead of one \$100M blockbuster.” <https://techcrunch.com/2026/04/16/runway-ceo-says-ai-could-help-hollywood-make-50-films-instead-of-one-100m-blockbuster/>. *Dream Machine* Issue 26.

24. *Fortune*, “What Netflix’s acquisition of Ben Affleck’s AI filmmaking company really shows.” <https://fortune.com/2026/03/06/what-netflixs-acquisition-of-ben-afflecks-ai-filmmaking-company-really-shows/>. *Dream Machine* Issue 20. On the price and the Roald Dahl comparison: *TechCrunch*, “Netflix may have paid \$600 million for Ben Affleck’s AI startup.” <https://techcrunch.com/2026/03/11/netflix-may-have-paid-600-million-for-ben-afflecks-ai-startup/>. *Dream Machine* Issue 21. \$587m confirmed in *Variety*, “Netflix Paid \$587 Million for Ben Affleck’s AI Startup InterPositive.” <https://variety.com/2026/film/news/netflix-paid-587-million-ben-affleck-ai-interpositive-1236815111/>. *Dream Machine* Issue 39.

25. *SiliconANGLE*, “Higgsfield raises \$80M on \$1.3B valuation to scale AI video platform.” <https://siliconangle.com/2026/01/15/higgsfield-raises-80m-1-3b-valuation-scale-ai-video-platform/>. *Dream Machine* Issue 15. 36Kr, “AI Video Unicorn Higgsfield: Earns \$200M in 9 Months by ‘Serving’ Social Media Marketers.” <https://eu.36kr.com/en/p/3650517574312323>. *Dream Machine* Issue 16. *Tech Times*, “AI Video Startup Higgsfield Hits \$500M Revenue, Eyes \$5B Funding Round.” <https://www.techtimes.com/articles/319394/20260630/ai-video-startup-higgsfield-hits-500m-revenue-eyes-5b-funding-round.htm>. *Dream Machine* Issue 36.

25a. *Tech Times*, “Kling AI Raises Record \$2.8 Billion: What Creators Owe China Under Beijing Law.” <https://www.techtimes.com/articles/319708/20260704/kling-ai-raises-record-28-billion-what-creators-owe-china-under-beijing-law.htm>. *Dream Machine* Issue 37.

26. Runway and Luma raises: *The Hollywood Reporter*, “Hollywood’s AI Breaking Point May Be Nearing.” <https://www.hollywoodreporter.com/business/business-news/hollywood-ai-breaking-point-1236512451/>. Runway Japan \$40m: Runway, “Runway is expanding into Japan: opening a Tokyo office and investing an initial \$40m in the country.” <https://runwayml.com/news/runway-is-coming-to-japan>. *Dream Machine* Issue 30. Studio and tool-maker film funds: *Forbes*, “Hollywood Studios Are Spending On AI To Control The Future Of Film.” <https://www.forbes.com/sites/maureenkerr/2026/06/01/hollywood-studios-are-spending-on-ai-to-control-the-future-of-film/>. A24–DeepMind reaction: *Fortune*, “How Hollywood’s youngest filmmakers are exposing Gen Z’s real problem with AI.” <https://fortune.com/2026/07/05/why-doesnt-gen-z-trust-ai-box-office-youtube-directors-backrooms-obsession/>.

27. Google Veo enterprise positioning, SynthID provenance and Vertex distribution, and the 2026 model-market commoditisation: *Tech-Insider*, “Best AI Video Generator 2026: Veo vs Sora vs Kling.” <https://tech-insider.org/best-ai-video-generator-2026/>. Kling (Kuaishou) per-second pricing, roughly \$0.10/second per Get AI Perks, “Best AI Video Generators 2026: Sora 2 vs Veo 3.1 vs Kling 3.0 vs Runway.” <https://www.getai perks.com/en/blogs/44-best-ai-video-generators-2026;> the \$0.07/second floor is Digital Applied’s estimate, as note 28. Seedance cost curve: as note 19.

28. *Digital Applied*, “AI Video Market After Sora: Runway, Kling, and Veo.” <https://www.digitalapplied.com/blog/ai-video-market-after-sora-runway-kling-veo-2026>. OpenAI announced the shutdown on 24 March 2026; the app closed 26 April 2026. See also Chapter 7 on the Disney exit. Higgsfield trajectory: as note 25.

28a. *PocketGamer.biz*, “Reforged Labs shuts down as CEO says AI is closing its market opportunity.” <https://www.pocketgamer.biz/reforged-labs-shuts-down-as-ceo-says-ai-is-closing-its-market-opportunity/>. *Dream Machine* Issue 40.

28b. *TechCrunch*, “Design Arena creators raise \$7.9 million to bring ‘taste’ to AI models.” <https://techcrunch.com/2026/08/03/designarena-creators-raise-7-9-million-to-bring-taste-to-ai-models/>. *Dream Machine* Issue 40.

29. Dubai \$1m award: *BroadcastPro ME*, “Dubai launches \$1m AI short film award in partnership with Google Gemini.” <https://www.broadcastprome.com/news/dubai-launches-1m-ai-short-film-award-in-partnership-with-google-gemini/>. *Dream Machine* Issue 2. XPRIZE Future Vision \$3.5m: *Fortune*, “Billionaire Peter Diamandis offers \$3.5 million to filmmakers who portray AI as the hero—not the villain.” <https://fortune.com/2026/03/09/xprize-foundation-future-vision-xprize-peter-diamandis-films-scifi-ai-fears/>; Google, “\$3.5 million Future Vision film competition.” <https://blog.google/innovation-and-ai/technology/ai/future-vision-film-competition-xprize/>. *Dream Machine* Issues 20 and 28. Astana \$1m prize fund: *The Astana Times*, “Astana AI Film Festival Opens Submissions, Offers \$1 Million Prize Fund.” <https://astanatimes.com/2026/05/astana-ai-film-festival-opens-submissions-offers-1-million-prize-fund/>. *Dream Machine* Issue 31. Luma Dream Brief \$1m: *Campaign Brief*, “Luma AI launches ‘The Luma Dream Brief’, a \$1m global competition for Cannes Lions Gold Winner.” <https://campaignbrief.com/luma-ai-launches-the-luma-dream-brief-a-1m-global-competition-for-cannes-lions-gold-winner/>. *Dream Machine* Issue 17. Adobe \$10m grants: *Analytics India Magazine*, “Adobe Unveils AI Video Innovations + \$10 Million Grants Ahead of Sundance Film Festival.” <https://analyticsindiamag.com/ai-news-updates/adobe-unveils-ai-video-innovations-10-million-grants-ahead-of-sundance-film-festival/>. *Dream Machine* Issue 15. Google–Sundance \$2m: Google, “Sundance Institute & Google Partner to

train filmmakers with AI.” <https://blog.google/company-news/outreach-and-initiatives/google-org/sundance-institute-ai-education/>. *Dream Machine* Issues 15 and 16. Runway Hundred Film Fund: *Forbes*, as note 26.

29a. Higgsfield, “Higgsfield AI Global Film Festival” — \$1m prize pool, 14 winners, \$500,000 for first place. <https://higgsfield.ai/global-film-festival>. *Dream Machine* Issue 41. The open-sourced feature *Hell Grind*: <https://x.com/higgsfield/status/2084702370764820572>. *Dream Machine* Issue 40.

29b. *Variety*, “Raindance Film Festival, AI Platform IMGN Launch New Film Fund and Development Hub for Indie Filmmakers.” <https://variety.com/2026/film/global/raindance-israel-ai-platform-imgn-development-fund-1236809594/>. *Dream Machine* Issue 39.

30. *The Hollywood Reporter*, “ReelShort and More: The Microdrama TV Series Gold Rush Is Here.” <https://www.hollywoodreporter.com/business/business-news/microdrama-series-verticals-production-1236418912/>.

31. See Chapter 7, §“Position One: All-in,” and its sources: Amazon AI Studios and *House of David*; Netflix INKubator; Disney’s Office of Technology Enablement. Netflix advertising: *Adweek*, “AI Agents Are Coming to Netflix to Grow Its \$3 Billion Ad Business.” <https://www.adweek.com/media/ai-agents-are-coming-to-netflix-to-grow-its-3-billion-ad-business/>. *Dream Machine* Issue 29.

32. *Variety*, “‘Squid Game’ Creator Makes Major Investment in AI Company TwelveLabs.” <https://variety.com/2025/tv/news/squid-game-creator-investment-ai-company-twelvelabs-1236536825/>. *Dream Machine* Issue 2.

33. See Chapter 7, §“Position One: All-in,” and Chapter 5, §“The pushback,” on the Amazon AI Creators Fund and the Punky Duck collapse.

34. Deloitte Insights (TMT Predictions 2026), “Tiny episodes, massive appeal: Short-form serials are gaining viewers and empowering independent studios.” <https://www.deloitte.com/us/en/insights/industry/technology/technology-media-and-telecom-predictions/2026/short-form-video-series.html>. Platform economics: Vitrina, “Micro Drama Distribution Platforms 2026.” <https://vitrina.ai/blog/micro-drama-distribution-platforms-2026/>. DramaBox raise: Real Reel, “Vertical Drama Weekly: Capital Tightens, Platforms Integrate.” <https://www.real-reel.com/vertical-drama-weekly-capital-platforms-2026/>. TikTok’s standalone microdrama app: *TechCrunch*, “TikTok quietly launches a micro-drama app called PineDrama.” <https://techcrunch.com/2026/01/16/tiktok-quietly-launches-a-micro-drama-app-called-pinedrama/>.

35. *The Hollywood Reporter*, “ReelShort and More: The Microdrama TV Series Gold Rush Is Here” (budgets; Fox–Holywater). <https://www.hollywoodreporter.com/business/business-news/microdrama-series->

verticals-production-1236418912/. Holywater first flagged in *Dream Machine* Issue 3. “First new professional format since reality TV”: C21Media, “Microdrama under the microscope.”

<https://www.c21media.net/departments/c21investigates/microdrama-under-the-microscope/>. Distributor abstention: *Deadline*, “Why TV’s Traditional Sellers Aren’t Mad About Microdramas.” <https://deadline.com/2026/02/microdrama-vertical-video-tv-distributors-1236714399/>. Netflix vertical feed and AI recommendations: *TechCrunch*, “Netflix plans to add a vertical video feed, use AI for recommendations.” <https://techcrunch.com/2026/04/17/netflix-plans-to-add-a-vertical-video-feed-use-ai-for-recommendations/>. *Dream Machine* Issue 26.

36. Real Reel, “Vertical Drama Weekly: Capital Tightens, Platforms Integrate.” <https://www.real-reel.com/vertical-drama-weekly-capital-platforms-2026/>.

36a. *Nikkei Asia*, “In China’s ‘microdrama’ boom, AI moviemakers chase 2 minutes of fame.” <https://asia.nikkei.com/spotlight/big-in-asia/in-china-s-microdrama-boom-ai-moviemakers-chase-2-minutes-of-fame>. *Dream Machine* Issue 39.

37. On the Meta Reality Labs cuts, see Chapter 14. On Vision Pro Personas and Gaussian splatting, see Chapter 9 and its sources.

38. Wikipedia, “Sphere (venue).” [https://en.wikipedia.org/wiki/Sphere_\(venue\)](https://en.wikipedia.org/wiki/Sphere_(venue)). *Yahoo Finance*, “Immersive Experience Bolsters Sphere Entertainment’s (SPHR) Strong Performance” (O’Keeffe Stevens Advisory Q2 2026 investor letter). <https://finance.yahoo.com/markets/stocks/articles/immersive-experience-bolsters-sphere-entertainment-145116003.html>. Sphere Entertainment IR, “Yas Island to Be Home of Sphere Abu Dhabi, a New Global Icon for Immersive Entertainment.” <https://investor.sphereentertainmentco.com/press-releases/news-details/2026/YAS-ISLAND-TO-BE-HOME-OF-SPHERE-ABU-DHABI-A-NEW-GLOBAL-ICON-FOR-IMMERSIVE-ENTERTAINMENT/default.aspx>. Sphere Entertainment IR, “Sphere Entertainment, the State of Maryland, Prince George’s County and Peterson Companies Announce Intent to Develop a Sphere at National Harbor.” <https://investor.sphereentertainmentco.com/press-releases/news-details/2026/Sphere-Entertainment-the-State-of-Maryland-Prince-Georges-County-and-Peterson-Companies-Announce-Intent-to-Develop-a-Sphere-at-National-Harbor/default.aspx>.

39. *Time*, “Netflix House: World’s Greatest Places 2026.” <https://time.com/collection/worlds-greatest-places/2026/netflix-house/>. *Event Marketer*, “Inside the Meteoric Rise of Immersive Entertainment Experiences.” <https://www.eventmarketer.com/article/fall-cover-story-immersive-entertainment-experiences/>.

40. *The Immersive Lab*, “Immersive Experiences 2025: The Year in Review.” <https://theimmersive.com/2026/01/immersive-experiences-2025-year-in-review/>.

review/.

41. See Chapter 9 on Marble and the collapse in environment costs. Sony Pictures' use of Marble in virtual production, 40× faster than the traditional workflow: https://www.linkedin.com/posts/brent-liang_tech-media-launch-ugcPost-7394911181091692546-TyUz. *Dream Machine* Issue 8. Netflix and Eyeline's Vista4D: <https://github.com/Eyeline-Labs/Vista4D>. *Dream Machine* Issue 27.

42. World Labs, "World Labs Announces New Funding." <https://www.worldlabs.ai/blog/funding-2026>. *Dream Machine* Issue 19. *Tech Times*, "Feifei Li's World Labs Splits World Model Into Three Types: Marble Targets Simulation Linchpin." <https://www.techtimes.com/articles/317927/20260606/feifei-lis-world-labs-splits-world-model-three-types-marble-targets-simulation-linchpin.htm>. Odyssey, "Odyssey has raised a \$310M Series B to accelerate world models." <https://odyssey.ml/our-series-b>. *Fundraise Insider*, "Odyssey Raises \$310M Series B at \$1.45B for World Models." <https://fundraiseinsider.com/blog/odyssey-raises-310m-series-b-at-1-45b-for-world-models/>. *Dream Machine* Issue 35. *Yahoo Finance*, "Tripo AI raises nearly \$200M in financing for AI 3D and world model tech." <https://finance.yahoo.com/sectors/technology/articles/tripo-ai-raises-nearly-200-164700546.html>. *Dream Machine* Issue 32. *Memeburn*, "Odyssey AI Raises \$310M to Build World Models, Valued at \$1.45 Billion." <https://memeburn.com/odyssey-ai-raises-310m-to-build-world-models-valued-at-1-45-billion/>.

43. ComfyUI, "We raised \$17 million to build an OS for Creative AI" (LinkedIn). https://www.linkedin.com/posts/comfyui_we-raised-17-million-to-build-an-os-for-ugcPost-7373743341236236288-wkCc. *Dream Machine* Issue 1. *TechCrunch*, "ComfyUI hits \$500M valuation as creators seek more control over AI-generated media." <https://techcrunch.com/2026/04/24/comfyui-hits-500m-valuation-as-creators-seek-more-control-over-ai-generated-media/>. *Dream Machine* Issue 27. Four million users, \$30m round and \$48m raised in total: *Dream Machine* Issue 36, Editor's Pick.

44. *TechTimes*, "Physics AI Slashes Engineering Simulation From Days to Seconds, PhysicsX Raises \$300M." <https://www.techtimes.com/articles/318582/20260617/physics-ai-slashes-engineering-simulation-days-seconds-physicsx-raises-300m.htm>. *Dream Machine* Issue 35.

44a. *Variety*, "Brad Pitt Says AI Could Help More 'Mid-Budget Films' Get Made if It's 'Used as a Tool'." <https://variety.com/2026/film/news/brad-pitt-ai-could-help-mid-budget-movies-get-made-1236831201/>. *Dream Machine* Issue 41.

45. *C21Media*, “UK government unveils \$500m creative industries plan as AI proposals remain uncertain.” <https://www.c21media.net/news/uk-government-unveils-500m-creative-industries-plan-as-ai-proposals-remain-uncertain/>. Creative PEC, “Taking stock of the Creative Industries Sector Plan.” https://pec.ac.uk/blog_entries/taking-stock-of-the-creative-industries-sector-plan/. Creative UK, “The Creative Industries Sector Plan: One Year On.” <https://www.wearecreative.uk/the-creative-industries-sector-plan-one-year-on/>.

46. GOV.UK (DSIT), “A decisive shift to power British AI: new £1.1 billion plan to back chip firms, boost computing power and skills for the AI revolution.” <https://www.gov.uk/government/news/a-decisive-shift-to-power-british-ai-new-11-billion-plan-to-back-chip-firms-boost-computing-power-and-skills-for-the-ai-revolution>. *Dream Machine* Issue 33. Arts Council England, “AI technologies and emerging forms of creative practice” (£4m of projects funded since 2019). <https://www.artscouncil.org.uk/ai-technologies-and-emerging-forms-creative-practice>. *Dream Machine* Issue 17. Innovate UK, feasibility studies for frontier AI and ML — up to £3m competition for UK-registered SMEs. <https://apply-for-innovation-funding.service.gov.uk/competition/2419/overview/bb8d7d28-0e63-4ee1-bf4a-790dfb3ae02f>. *Dream Machine* Issue 23.

47. RWK Goodman, “UK Industrial Sector Strategy Plans — Creative Industries.” <https://www.rwkgoodman.com/info-hub/uk-industrial-sector-strategy-plans-creative-industries/>.

48. See Chapter 6, “The 88%,” on the consultation response breakdown and the December 2025 Statement of Progress.

Chapter 9 — Worlds, Not Pictures

1. World Labs, “Bringing Marble to Life.” <https://www.worldlabs.ai/case-studies/bringing-marble-to-life>. *Dream Machine* Issue 7, “Editor’s Pick: Marble by WorldLabs goes on public release,” 13 November 2025.

2. For a working primer on Gaussian splatting in the post-Marble era, see *Radiance Fields*, “World Labs Formally Launches Marble, A Generative World Model.” <https://radiancefields.com/world-labs-formally-launches-marble-a-generative-world-model>.

3. DreamLab AI Collective, beta participation in Marble, October–November 2025. Referenced in *Dream Machine* Issue 7: “DreamLab have been part of the beta testing for this over the last few months and it’s very neat.”

4. SuperSplat (PlayCanvas), open-source Gaussian splat editor, regular updates through 2025–26. *Dream Machine* Issue 1: “PlayCanvas open sources SOG — WebP for 3D Gaussian Splatting”; Issue 7 / Issue 11 on SuperSplat v2 updates.

5. Sony Pictures’ use of Marble in Virtual Production: https://www.linkedin.com/posts/brent-liang_tech-media-launch-ugcPost-7394911181091692546-TyUz. *Dream Machine* Issue 8.
6. Disney “300,000 poses in an instant” livestream, March 2026. *Dream Machine* Issue 23.
7. Netflix + Eyeline, *Vista4D*: 4D point clouds from live-action. *Dream Machine* Issue 27.
8. Google DeepMind, “Genie 3: A new frontier for world models.” <https://deepmind.google/blog/genie-3-a-new-frontier-for-world-models/>. Project Genie roll-out to AI Ultra subscribers in the U.S.: <https://blog.google/innovation-and-ai/models-and-research/google-deepmind/project-genie/>. *Dream Machine* Issue 3 (initial announcement) and Issue 17 (broader availability).
9. Meta, “WorldGen — text-to-immersive-3D-worlds research update.” <https://www.facebook.com/LifeAtMeta/videos/research-update-worldgen-text-to-immersive-3d-worlds/1879077432692421/>. *Dream Machine* Issues 9, 11.
10. Tencent, “HY World 1.5” announcement: <https://x.com/TencentHunyuan/status/2001170499133653006>. *Dream Machine* Issue 12.
11. SpAItial, *ECHO* spatial foundation model. <https://www.spaitial.ai/>. *Dream Machine* Issue 12.
12. Stanford AI Lab, *Wonderzoom*: Multi-Scale 3D World Generation. <https://wonderzoom.github.io/>. *Dream Machine* Issue 14.
13. OpenArt, *Worlds* product launch, March 2026. *Dream Machine* Issue 21.
14. Luma AI, *UNI-1* launch, March 2026. *Dream Machine* Issue 22, “Editor’s Pick: When worlds become instant, the race shifts to better thinking.”
15. ByteDance Seedance 2.0 in CapCut/Dreamina, March 2026. *Dream Machine* Issue 22.
16. *Spark 2.0*, open-source Gaussian-splat streaming framework, April 2026. *Dream Machine* Issue 25.
17. Radiance Fields, “Apple Confirms that it’s Gaussian Splatting that powers their personas.” <https://radiancefields.com/apple-confirms-personas-use-gaussian-splatting>. *Dream Machine* Issue 5.
18. *Video Games Chronicle*, “It honestly sucks’: Fans think Call of Duty: Black Ops 7 is filled with generative AI art.” <https://www.videogameschronicle.com/news/it-honestly-sucks-fans-think-call-of-duty-black-ops-7-is-filled-with-generative-ai-art/>. *Video Games Chronicle*, “Ubisoft says AI-generated art in Anno 117 was a

placeholder which ‘slipped through our review process’.” <https://www.videogameschronicle.com/news/ubisoft-says-ai-generated-art-in-anno-117-was-a-placeholder-which-slipped-through-our-review-process/>. *Polygon*, “Fortnite chapter 7 kicks off new controversy over AI art.” <https://www.polygon.com/fornite-chapter-7-season-1-generative-ai-art-epic-games/>. *Dream Machine* Issues 8, 10.

19. NVIDIA + Stanford, *NitroGen*. <https://nitrogen.minedojo.org/>. *Dream Machine* Issue 13.

20. DeepMind, “SIMA 2: An Agent that Plays, Reasons, and Learns With You in Virtual 3D Worlds.” <https://deepmind.google/blog/sima-2-an-agent-that-plays-reasons-and-learns-with-you-in-virtual-3d-worlds/>. *Dream Machine* Issue 8.

21. ComfyUI Blog, “Ubisoft La Forge Open-Sources the CHORD Model and ComfyUI Nodes for End-to-End PBR Material Generation.” <https://blog.comfy.org/p/ubisoft-open-sources-the-chord-model>. *Dream Machine* Issue 11.

22. *Video Games Chronicle*, “The future of gaming, or ‘just a tool’? Hands-on with Teammates, Ubisoft’s ambitious voice AI tech demo.” <https://www.videogameschronicle.com/features/the-future-of-gaming-or-just-a-tool-hands-on-with-teammates-ubisofts-ambitious-voice-ai-tech-demo/>. *Dream Machine* Issue 9.

23. YouTube Playables Builder, closed-beta announcement: <https://www.youtube.com/playablesbuilder/>. *Dream Machine* Issue 12.

24. Unity AI Open Beta, in-editor AI suite, May 2026. *Dream Machine* Issue 28.

25. Korin AI, “trained with African datasets, built by Africans,” May 2026. *Dream Machine* Issue 27.

26. NVIDIA SANA-WM, 2.6B open-source world model with 60-second video generation and camera control, May 2026. <https://huggingface.co/collections/nvidia/sana-wm>. *Dream Machine* Issue 30.

27. Odyssey, “Introducing Starchild-1, the first real-time multimodal world model.” <https://odyssey.ml/introducing-starchild-1>. *Dream Machine* Issue 30.

28. Apple Machine Learning Research, “Headsup: a large-scale high-quality 3D Gaussian head reconstruction from multi-view captures.” <https://machinelearning.apple.com/research/apple-headsup-3d-gaussian-head>. *Dream Machine* Issue 30.

29. WorldLens VR, “AI-powered 3D depth for Google Street View on Quest.” <https://www.uploadvr.com/worldlens-vr-quest-street-view-3d-depth/>. *Dream Machine* Issue 30.

- 30.** Odyssey, “Our Series B” — \$310M raise to accelerate world models, June 2026. <https://odyssey.ml/our-series-b>. *Dream Machine* Issue 35.
- 31.** The Sandbox, “Announcing The Sandbox Studio: from idea to live multiplayer game in hours.” <https://www.youtube.com/watch?v=IdTVDRvQoWM>. *Dream Machine* Issue 33.
- 32.** Higgsfield, “Higgsfield Games” — multiplayer games created and deployed from a single prompt. <https://higgsfield.ai/games-intro>. *Dream Machine* Issue 34.
- 33.** Unreal Engine, “State of Unreal 2026: top news from the show.” <https://www.unrealengine.com/news/state-of-unreal-2026-top-news-from-the-show>. Engadget, “Epic Games details how it’s embracing generative AI in Unreal Engine.” <https://www.engadget.com/2196807/epic-games-details-how-its-embracing-gen-ai-in-unreal-engine/>. *Dream Machine* Issue 35, “Editor’s Pick: Unreal Engine’s AI Future — From Tools to Creative Partners,” 23 June 2026.
- 34.** Roblox Newsroom, “Pioneering AI founders join to accelerate Roblox reality vision” — acquisition of Morpheus AI, Dynamics Lab and Lucid AI. <https://about.roblox.com/newsroom/2026/06/pioneering-ai-founders-join-to-accelerate-roblox-reality-vision>. *Dream Machine* Issues 32 and 34.
- 35.** *Dream Machine* Issue 34, “Editor’s Pick: The AI World Cup — When Sport Becomes a Living Digital System,” 16 June 2026.
- 36.** GamesBeat, “FIFA and Gamefam launch FIFA World Cup 2026 event in Super Soccer and Roblox.” <https://gamesbeat.com/gamefam-launch-fifa-world-cup-2026-super-soccer-roblox/>. PocketGamer.biz, “FIFA and Gamefam launch FIFA World Cup 2026 event across Roblox.” <https://www.pocketgamer.biz/fifa-and-gamefam-launch-fifa-world-cup-2026-event-across-roblox/>. *Dream Machine* Issue 34.
- 37.** Sphere Entertainment, “Sphere Studios announces *The Rocky Horror Picture Show* at Sphere.” <https://www.sphereentertainmentco.com/sphere-studios-announces-the-rocky-horror-picture-show-at-sphere/>. *Dream Machine* Issue 35.
- 38.** Adobe Research, “Wonder” — real-time interactive 3D world navigation. <https://wonder-world-model.github.io/>. *Dream Machine* Issue 39.
- 39.** Tencent, “WorldClaw” — agentic 3D open-world generation at scale. <https://tencent-hunyuan.github.io/Hunyuan3D-WorldClaw/>. *Dream Machine* Issue 41.
- 40.** Reactor, “LingBot-World 2.0” — real-time world generation model. <https://www.reactor.inc/models/lingbot-world-2>. *Dream Machine* Issue 37 and Issue 38.
- 41.** AlayaLab, “AlayaWorld: Long-Horizon and Playable Video World Generation.” <https://github.com/AlayaLab/AlayaWorld>. *Dream Machine* Issue 37.

42. Wayve, “GAIA-4” — the latest of Wayve’s line of world models. <https://wayve.ai/thinking/gaia-4/>. *Dream Machine* Issue 40.
43. Lightricks, “LTX-2.5” — an open-world, open-source model. <https://huggingface.co/Lightricks/LTX-2.5>. *Dream Machine* Issue 41.
44. Gracia AI, 4D Gaussian splatting in Unreal Engine with real-time relighting and self-shadowing. <https://gracia.ai/>. Atlux, Unreal Engine to 4D Gaussian Splat. <https://atlux.one/4dgs>. *Radiance Fields*, “Notch 2026.2 Adds Native Gaussian Splat Import and Rendering.” <https://radiancefields.com/notch-2026.2-adds-native-gaussian-splat-import-and-rendering>. *Dream Machine* Issue 37.
45. InfiniSplat (SIGGRAPH Asia 2026) — turn a single image into a navigable 3DGS. <https://zju3dv.github.io/InfiniSplat/>. Paper: <https://arxiv.org/abs/2608.02437>. *Dream Machine* Issue 40.
46. Insta360, “Insta360 X6” — introduces Spatial Capture: record a place in 360, bring it into the app, and it becomes a 3D model. <https://www.insta360.com/product/insta360-x6>. *Dream Machine* Issue 41.
47. Niantic Spatial and HMCI, “Building the Foundation for City of Rancho Cordova’s First Digital Twin for Physical AI.” <https://www.youtube.com/watch?v=TbPaQw-xLVc>. *Dream Machine* Issue 40.

Chapter 10 — AI in Everything, Everywhere, All at Once

1. *Creative Boom*, “Adobe is putting AI in everything everywhere all at once.” <https://www.creativeboom.com/news/adobe-is-putting-ai-in-everything-everywhere-all-at-once/>. *Dream Machine* Issue 5, “Editor’s Pick,” 31 October 2025.
2. Adobe, “Adobe MAX 2025: Firefly Foundry.” <https://news.adobe.com/news/2025/10/adobe-max-2025-firefly-foundry>.
3. Adobe, “Adobe MAX 2025: Firefly.” <https://news.adobe.com/news/2025/10/adobe-max-2025-firefly>.
4. Adobe, “Adobe MAX 2025: Express AI Assistant.” <https://news.adobe.com/news/2025/10/adobe-max-2025-express-ai-assistant>.
5. *Wired*, “Adobe’s ‘Corrective AI’ Can Change the Emotions of a Voice-Over” and accompanying Adobe Sneaks 2025 coverage. <https://www.wired.com/story/adobe-max-sneaks-2025-corrective-ai/>. Project list compiled from MAX keynote and *Dream Machine* Issue 5 coverage.
6. *PYMNTS*, “Adobe Lets Users Design and Edit Using ChatGPT.” <https://www.pymnts.com/artificial-intelligence-2/2025/adobe-lets-users-design-and-edit-using-chatgpt/>. Adobe blog: “Edit images, designs, and PDFs right inside

ChatGPT.” <https://blog.adobe.com/en/publish/2025/12/10/edit-photoshop-chatgpt>. *Dream Machine* Issue 12.

7. Adobe Premiere Object Mask tool: https://www.linkedin.com/posts/robdewinter_ok-this-is-going-to-save-a-lot-of-time-in-ugcPost-7421617551690063872-yKmB. *Dream Machine* Issue 16.

8. Adobe blog, “Sundance Film Festival 2026: Creativity, Community & Power of Storytelling.” <https://blog.adobe.com/en/publish/2026/01/20/sundance-film-festival-2026-creativity-community-power-storytelling>. *Dream Machine* Issue 16.

9. Adobe Summit 2026, “agentic creative intelligence” keynote. *Dream Machine* Issue 26.

10. After Effects AI animation features through late 2025; *Dream Machine* Issue 9, “AI video is finally animatable inside After Effects.” https://www.linkedin.com/posts/thisisdoug_ai-aivideo-animation-ugcPost-7399512745924067330-Aldk.

10a. Doug McGinness on LinkedIn, late 2025, in the same post. *Dream Machine* Issue 9.

11. *Dream Machine* Issue 21, “Editor’s Pick: Adobe and NVIDIA Just Raised the Stakes for Creative AI,” 19 March 2026.

12. NVIDIA + Google Cloud creative-AI infrastructure deal, March 2026. *Dream Machine* Issue 21.

13. Hugging Face and Google Cloud partnership announcement: https://www.linkedin.com/posts/julienchaumond_i-am-super-excited-to-announce-that-hugging-activity-7396177403972276225-CuMM. *Dream Machine* Issue 8.

14. *EdTech Innovation Hub*, “Meta and Hugging Face launch OpenEnv to advance open-source agentic development.” <https://www.edtechinnovationhub.com/news/meta-and-hugging-face-launch-openenv-to-advance-open-source-agentic-development>. *Dream Machine* Issue 5.

15. Anthropic / Blender Foundation patronage, May 2026. *Dream Machine* Issue 27.

16. TechCrunch, “Anthropic launches interactive Claude apps, including Slack and other workplace tools.” <https://techcrunch.com/2026/01/26/anthropic-launches-interactive-claude-apps-including-slack-and-other-workplace-tools/>. *Dream Machine* Issue 16.

17. Spotify–Anthropic integration, May 2026. *Dream Machine* Issue 27.

- 18.** *MarTech Series*, “WPP continues AI overhaul with \$400-million Google partnership.” <https://martechseries.com/predictive-ai/ai-platforms-machine-learning/google-and-spotify-alum-launch-epiminds-with-6-6m-to-build-marketing-teams-for-the-ai-era/>. *Dream Machine* Issue 3.
- 19.** *Campaign Brief*, “WPP launches AI-powered marketing platform WPP Open Pro.” <https://campaignbrief.com/wpp-launches-ai-powered-marketing-platform-wpp-open-pro/>. *Dream Machine* Issue 5.
- 20.** *Digiday*, “WPP expands AI capabilities to boost brand performance with Sightly partnership.” <https://digiday.com/media-buying/agencies-continue-to-expand-ai-capabilities-to-boost-brand-performance/>. *Dream Machine* Issue 6.
- 21.** WPP and Google Earth AI consumer-journey project, April 2026. *Dream Machine* Issue 27.
- 22.** *SiliconAngle*, “Higgsfield raises \$80M on \$1.3B valuation to scale AI video platform.” <https://siliconangle.com/2026/01/15/higgsfield-raises-80m-1-3b-valuation-scale-ai-video-platform/>. *Dream Machine* Issue 15.
- 23.** *36kr*, “AI Video Unicorn Higgsfield: Earns \$200M in 9 Months by ‘Serving’ Social Media Marketers.” <https://eu.36kr.com/en/p/3650517574312323>. *Dream Machine* Issue 16.
- 24.** *TechCrunch*, “Synthesia hits \$4B valuation, lets employees cash in.” <https://techcrunch.com/2026/01/26/synthesia-hits-4b-valuation-lets-employees-cash-in/>. *Dream Machine* Issue 16.
- 25.** *Sifted*, “Synthesia rejects \$3bn Adobe acquisition offer.” <https://sifted.eu/articles/synthesia-acquisition-offer>. *Dream Machine* Issue 5.
- 26.** ElevenLabs \$500m ARR reporting, April 2026. *Dream Machine* Issue 25.
- 27.** Runway product cycle: Gen-4.5 (December 2025), Gen-4.5 Image-to-Video (January 2026), Workflows, Story Panels, Characters API, Apps for Advertising — *Dream Machine* Issues 10, 14, 15, 16, 20. Runway CEO on indie films vs. blockbusters: *Dream Machine* Issue 26.
- 28.** For the running ledger of new creative-AI products through 2025–26, see *Dream Machine* Issues 1–30 archive.
- 29.** ComfyUI, “We raised \$17 million to build an OS for Creative AI.” https://www.linkedin.com/posts/comfyui_we-raised-17-million-to-build-an-os-for-ugcPost-7373743341236236288-wkCc. *Dream Machine* Issue 1.
- 30.** ComfyUI \$500M valuation, May 2026. *Dream Machine* Issue 27.
- 31.** Google Pomelli launch: <https://x.com/GoogleLabs/status/1983204018567426312>. *Dream Machine* Issue 5.

- 32.** Google AI Studio app gallery: <https://x.com/GoogleAISTudio/status/1982121563785949255>. Google Labs Opal expansion: <https://blog.google/technology/google-labs/opal-expansion/>. Project Genie: <https://blog.google/innovation-and-ai/models-and-research/google-deepmind/project-genie/>. *Dream Machine* Issues 5, 17.
- 33.** Lovable for classrooms: <https://lovable.dev/classroom>. *Dream Machine* Issue 11.
- 34.** Adobe Express AI Assistant: <https://news.adobe.com/news/2025/10/adobe-max-2025-express-ai-assistant>. *Dream Machine* Issue 5.
- 35.** Hugging Face platform expansion through 2025–26.
- 36.** Google blog, “Sundance Institute AI Education.” <https://blog.google/company-news/outreach-and-initiatives/google-org/sundance-institute-ai-education/>. *Dream Machine* Issue 15.
- 37.** Adobe Ignite Day at Sundance: *Adobe blog, Sundance Film Festival 2026*. <https://blog.adobe.com/en/publish/2026/01/20/sundance-film-festival-2026-creativity-community-power-storytelling>. *Dream Machine* Issue 16.
- 38.** Google’s \$40bn investment in Anthropic, May 2026. *Dream Machine* Issue 27.
- 39.** UK Government, “Free AI training for all.” <https://www.gov.uk/government/news/free-ai-training-for-all-as-government-and-industry-programme-expands-to-provide-10-million-workers-with-key-ai-skills-by-2030>. *Dream Machine* Issue 16.
- 40.** CNBC, “People with ADHD, autism, dyslexia say AI agents are helping them succeed at work.” <https://www.cnn.com/2025/11/08/adhd-autism-dyslexia-jobs-careers-ai-agents-success.html>. *Dream Machine* Issue 7.
- 41.** University of Wisconsin-Stout, “AI Reshaping Industry: New UW-Stout Course Sets AI-Use as Baseline Competency in Filmmaking.” <https://www.uwstout.edu/about-us/news-center/ai-reshaping-industry-new-uw-stout-course-sets-ai-use-baseline-competency-filmmaking>. *Dream Machine* Issue 15.
- 42.** Google I/O 2026 announcement block: Gemini Omni <https://blog.google/technology/google-deepmind/gemini-omni/>, Antigravity <https://antigravity.google/>, Google Flow <https://flow.google/>, Gemini Spark <https://blog.google/technology/developers/gemini-spark/>, Project Genie + Street View <https://deepmind.google/discover/blog/project-genie-street-view/>. *Dream Machine* Issue 30, “Editor’s Pick — Google I/O 2026,” 21 May 2026.
- 43.** Google Labs, “Infinite Scaler.” <https://blog.google/technology/google-labs/infinite-scaler/>. *Dream Machine* Issue 30.

44. Google DeepMind, “SynthID — 100 billion watermarks, partner ecosystem.” <https://deepmind.google/discover/blog/synthid-100-billion-watermarks-partners/>. *Dream Machine* Issue 30.
45. Runway, “Runway Japan.” <https://runwayml.com/blog/runway-japan>. *Dream Machine* Issue 30.
46. *Music Business Worldwide*, “Splice inks ‘Responsible AI’ deal with ElevenLabs.” <https://www.musicbusinessworldwide.com/splice-elevenlabs-responsible-ai-deal/>. *Dream Machine* Issue 30.
47. *Adweek*, “Netflix ad tools could see ‘agentic AIs talking to each other’.” <https://www.adweek.com/media/netflix-ad-tools-agentic-ais-talking-to-each-other/>. *Dream Machine* Issue 30.
48. *Fortune*, “AI startup Viktor raises \$75 million to put a virtual ‘coworker’ in Slack and Teams.” <https://fortune.com/2026/05/19/ai-startup-viktor-75-million-virtual-coworker-slack-teams/>. *Dream Machine* Issue 30.
49. Adobe Newsroom, “Adobe unveils major expansion” — Adobe Creative Agent and AI Assistant in beta across Premiere, Photoshop, Illustrator, InDesign and Frame.io, June 2026. <https://news.adobe.com/news/2026/06/adobe-unveils-major-expansion>. *Dream Machine* Issue 35.
50. *MarTech Series*, “Adobe Introduces Firefly Foundry to Support Custom AI Models for Enterprise Brand Consistency.” <https://martechseries.com/predictive-ai/ai-platforms-machine-learning/adobe-introduces-firefly-foundry-to-support-custom-ai-models-for-enterprise-brand-consistency/>. *Dream Machine* Issue 36.
51. Adobe blog, “Disney Imagineers add Adobe Firefly Foundry to toolkit of worldbuilding superpowers.” <https://business.adobe.com/blog/disney-imagineers-add-adobe-firefly-foundry-to-toolkit>. *Dream Machine* Issue 35.
52. Adobe Newsroom, “Adobe to acquire Topaz Labs,” June 2026. <https://news.adobe.com/news/2026/06/adobe-to-acquire-topaz-labs>. *Dream Machine* Issue 36.
53. Forrester, “Google Dethrones OpenAI As Agencies’ Preferred AI Partner.” <https://www.forrester.com/blogs/google-dethrones-openai-as-agencies-preferred-ai-partner/>. *Dream Machine* Issue 35.
54. *Digiday*, “OpenAI moves to automate ad creative.” <https://digiday.com/marketing/openai-moves-to-automate-ad-creative/>. *Dream Machine* Issue 35.
55. Apple Newsroom, “Apple unveils next generation of Apple Intelligence, Siri AI and more,” June 2026. <https://www.apple.com/newsroom/2026/06/apple-unveils-next-generation-of-apple-intelligence-siri-ai-and-more/>. *Dream Machine* Issue 33.

56. *Dream Machine* Issue 36, “Editor’s Pick: ComfyUI MCP — the open source engine behind AI’s creative future,” 2 July 2026. Background: *Medium* (Andrew Zhu), “The Rise of ComfyUI: How One Guy’s Side Project Took Over AI Image Generation.” <https://xhinker.medium.com/the-rise-of-comfyui-how-one-guys-side-project-took-over-ai-image-generation-558dca64bfd7>.

57. Comfy Org, “Comfy MCP” — public beta, July 2026. <https://comfy.org/mcp/>. *Dream Machine* Issue 36.

58. *Forbes*, “Video AI Wars: How Chinese Labs Are Winning The Race OpenAI Abandoned,” 5 June 2026. <https://www.forbes.com/sites/edithyeung/2026/06/05/video-ai-wars-how-chinese-labs-are-winning-the-race-openai-abandoned/>. *Dream Machine* Issue 33.

59. Anthropic, “Fable and Mythos access” — suspension of frontier AI tools following US export-control directive, June 2026. <https://www.anthropic.com/news/fable-mythos-access>. *Dream Machine* Issue 34.

51a. *Blog Mickey*, “How Disney Imagineering Is Using AI: Digital Twins, J.A.R.V.I.S., and Faster Ride Design.” <https://blogmickey.com/2026/08/damaro-details-ai-use-imagineering-disney-company/>. *Dream Machine* Issue 41.

60. *The Verge*, “Spotify is now an AI chatbot, too.” <https://www.theverge.com/entertainment/965358/spotify-ai-chatbot-interface-music-audiobooks-podcasts>. *Dream Machine* Issue 38.

61. *Search Engine Land*, “OpenAI appears to be building chatbot-native ads that launch AI agents.” <https://searchengineland.com/openai-appears-to-be-building-chatbot-native-ads-that-launch-ai-agents-484107>. *Dream Machine* Issue 41.

Chapter 11 — What is newly possible

1. Snap Newsroom, “Snapchat Gen Z AI Creativity Research 2026.” <https://newsroom.snap.com/snapchat-gen-z-ai-creativity-research-2026>. *Dream Machine* Issue 30.

2. *Forbes*, “Homer Epic ‘The Odyssey’ Gets A Full-Length AI-Generated Film.” <https://www.forbes.com/sites/lesliekatz/2026/07/14/homer-epic-the-odyssey-gets-a-full-length-ai-generated-film/>. *Dream Machine* Issue 38.

3. *Musically*, “AI-powered platform Eros Music Worlds is turning film characters into virtual singers.” <https://musically.com/2026/07/02/ai-powered-platform-eros-music-worlds-is-turning-film-characters-into-virtual-singers/>. *Dream Machine* Issue 37.

4. National Portrait Gallery, “Interactive portrait of the nation by artist Es Devlin to launch at the National Portrait Gallery.”

<https://www.npg.org.uk/about/press/interactive-portrait-of-the-nation-by-artist-es-devlin-to-launch-at-the-national-portrait-gallery>. *Dream Machine* Issue 40.

5. HeyGen, “AI Video Podcast.” <https://www.heygen.com/apps/ai-video-podcast>. *Dream Machine* Issue 39.

6. WERSM, “Wrinkles, a new app that Turns Local History Into An AI Audio Guide.” <https://wersm.com/wrinkles-ai-audio-guide-local-history/>. *Dream Machine* Issue 40.

Chapter 12 — The Orchestrator

1. *Dream Machine* Issue 13, “Editor’s Pick: The Year of the Orchestrator,” 9 January 2026.

2. *Dream Machine* Issue 29, May 2026, reporting on Sony’s 49-agent / 72-skill multi-agent game-development team.

3. Anthropic blog content on agent deployment patterns, Q1 2026.

3a. Bloomberg, “AI Changed Chess. Grandmasters Now Win With Unpredictable Moves,” 27 March 2026. <https://www.bloomberg.com/news/articles/2026-03-27/ai-changed-chess-grandmasters-now-win-with-unpredictable-moves>. *Dream Machine* Issue 23. The chess analogy is developed in Chapter 16’s *Age of the Why* section.

4. Sundance Institute, “Centering the Artist: Why We’re Launching the AI Literacy Initiative.” <https://www.sundance.org/blogs/centering-the-artist-why-were-launching-the-ai-literacy-initiative/>. *Dream Machine* Issue 16.

5. Sundance Institute, *op. cit.*

6. Sundance Story Forum 2026 sessions on legal toolkits for producers using AI. *Dream Machine* Issue 16.

7. Google blog, “Sundance Institute AI Education.” <https://blog.google/company-news/outreach-and-initiatives/google-org/sundance-institute-ai-education/>. *Dream Machine* Issue 15.

8. McKinsey & Company, “What AI could mean for film and TV production and the industry’s future.” <https://www.mckinsey.com/industries/technology-media-and-telecommunications/our-insights/what-ai-could-mean-for-film-and-tv-production-and-the-industrys-future>. *Dream Machine* Issue 16.

9. Metro, “Prince of Persia remake and five more games cancelled as Ubisoft focuses on AI.” <https://metro.co.uk/2026/01/21/prince-persia-remake-five-games-cancelled-ubisoft-focuses-ai-26431926/>. *Dream Machine* Issue 15.

10. *PC Gamer*, “Square Enix, makers of Final Fantasy, aims to have AI doing 70% of its QA work by the end of 2027.” <https://www.pcgamer.com/gaming-industry/square-enix-aims-to-have-ai-doing-70-percent-of-its-qa-work-by-the-end-of-2027/>. *Dream Machine* Issue 7.

11. *Eurogamer*, “Falcom is the latest developer to buy into the AI hype machine.” <https://www.eurogamer.net/falcom-is-the-latest-developer-to-buy-into-the-ai-hype-machine>. *Dream Machine* Issue 12.

12. *NDTV Profit*, “Don’t Expect AI To Invent the Next ‘Grand Theft Auto’, Says Take-Two CEO Strauss Zelnick.” <https://www.ndtvprofit.com/technology/dont-expect-ai-to-invent-the-next-grand-theft-auto-says-take-two-ceo-strauss-zelnick>. *Dream Machine* Issue 6.

13. *Dream Machine* Issue 21, on Spielberg’s public position on AI.

14. *Dream Machine* Issues 25, 28, on Steven Soderbergh’s AI work.

15. *Digiday*, “Independent agencies face new frontier as agency-in-a-box tools democratize creativity.” <https://digiday.com/marketing/independent-agencies-face-new-frontier-as-agency-in-a-box-tools-democratize-creativity/>. *Dream Machine* Issues 6, 14.

16. *Digiday*, “AI agent developers have become adland’s in-demand role.” <https://digiday.com/marketing/ai-agent-developers-have-become-adlands-in-demand-role/>. *Dream Machine* Issue 7.

17. *PYMNTS*, “AI Content Is Par For The Course With PGA Tour’s Expanded AWS Partnership.” <https://www.pymnts.com/artificial-intelligence-2/2026/ai-content-is-par-for-the-course-with-pga-tours-expanded-aws-partnership/>. *Dream Machine* Issue 15.

18. *The Verge*, “Oreo-maker Mondelez will use AI for TV ads next year.” <https://www.theverge.com/news/806047/mondelez-ai-generated-ads>. *Dream Machine* Issue 5.

19. *Digiday*, “Avocados From Mexico turns to AI to advertise around the Super Bowl instead of a TV buy.” <https://digiday.com/marketing/avocados-from-mexico-turns-to-ai-to-advertise-around-the-super-bowl-instead-of-a-tv-buy/>. *Dream Machine* Issue 15.

20. *Dream Machine* Issue 26.

21. Reuters Institute, “AI adoption by UK journalists and their newsrooms: surveying applications, approaches, and attitudes.” <https://reutersinstitute.politics.ox.ac.uk/ai-adoption-uk-journalists-and-their-newsrooms-surveying-applications-approaches-and-attitudes>. *Dream Machine* Issue 9.

- 22.** *Digiday*, “Daily Mail says Google AI Overviews have killed click-throughs.” <https://digiday.com/media/daily-mail-says-google-ai-overviews-have-killed-click-throughs/>. *Dream Machine* Issue 7.
- 23.** *Digiday*, “How The Times is using AI to model synthetic focus groups from human audiences.” <https://digiday.com/media/how-the-times-is-using-ai-to-model-synthetic-focus-groups-from-human-audiences/>. *Dream Machine* Issue 6.
- 24.** *TechBullion*, “Why the future belongs to multi-skilled leaders.” <https://techbullion.com/playing-the-long-game-with-a-portfolio-career-why-the-future-belongs-to-multi-skilled-leaders/>. *Dream Machine* Issue 9.
- 25.** Anthropic Skills framework via Claude Code, reported through *Dream Machine* Issues 11, 16, 29.
- 26.** *Forbes*, “AI Is Changing How Creators Work And Earn.” <https://www.forbes.com/sites/kolawolesamueladebayo/2025/12/22/how-ai-is-changing-how-creators-work-and-earn/>. *Dream Machine* Issue 13.
- a1.** *Variety*, “AI Creator Behind Viral ‘Deadpool,’ ‘Harry Potter’ Christmas Clip Made His Film in a Ukrainian Bomb Shelter.” <https://variety.com/2026/digital/news/ai-video-deadpool-harry-potter-andrii-daniels-1236624632/>. *Dream Machine* Issue 16.
- 27.** *Variety*, “MiniMax Launches All-in-One AI Video Generator Hub at Shanghai Film Festival: ‘Creative Direction and Aesthetic Judgment Must Ultimately Be Left to Humans’.” <https://variety.com/2026/film/festivals/minimax-hub-ai-video-generator-shanghai-film-festival-1236781901/>. *Dream Machine* Issue 35.
- 28.** Comfy Org, “Comfy MCP” public beta, July 2026. <https://comfy.org/mcp/>. *Dream Machine* Issue 36, Editor’s Pick.
- 29.** Unreal Engine, “State of Unreal 2026: top news from the show” (Unreal MCP announcement). <https://www.unrealengine.com/news/state-of-unreal-2026-top-news-from-the-show>. *Dream Machine* Issue 35. Runway, “Runway in ChatGPT.” <https://runwayml.com/mcp>. *Dream Machine* Issue 34. Magnific, “Connecting Magnific to AI assistants using MCP.” <https://www.magnific.com/academy/tutorial/connecting-magnific-to-ai-assistants-using-mcp>. *Dream Machine* Issue 33. Higgsfield, “Higgsfield MCP.” <https://higgsfield.ai/mcp>. *Dream Machine* Issue 36.
- 30.** *The Hollywood Reporter*, “Vibe Directing? An AI Startup Run by Ex-Googlers Say Their New Model Can Make It Happen (Exclusive).” <https://www.hollywoodreporter.com/business/digital/open-art-new-tool-model-ai-director-vibe-directing-1236628188/>. *Dream Machine* Issue 36.

31. *Variety*, “‘Backrooms’ Director Kane Parsons Says Using AI ‘Defeats the Purpose’ of Filmmaking: ‘I Get No Enjoyment From Using Those Tools’.” <https://variety.com/>. *Dream Machine* Issue 32.

32. *MusicRadar*, “You can now control FL Studio with an agentic AI chatbot from inside the DAW.” <https://www.musicradar.com/music-tech/you-can-now-control-fl-studio-with-an-ai-chatbot-from-inside-the-daw>. *Dream Machine* Issue 38.

33. *Music Ally*, “AI’s new USB port: what the MCP standard means for the music industry.” <https://musically.com/2026/07/20/ais-new-usb-port-what-the-mcp-standard-means-for-the-music-industry/>. *Dream Machine* Issue 39.

34. *Digital Music News*, “How AI Agents Are Quietly Automating Music Companies: From A&R to Release Ops.” <https://www.digitalmusicnews.com/2026/07/29/ai-agents-music-industry-pibox/>. *Dream Machine* Issue 39.

35. *TVBEurope*, “‘World first’ agentic AI TV platform debuts.” <https://www.tvbeurope.com/media-consumption/world-first-agentic-ai-tv-platform-debuts>. *Dream Machine* Issue 39.

36. *Filmera*, agentic AI filmmaking platform and workspace. <https://www.filmera.ai/>. *Dream Machine* Issue 41.

37. “How Spotify runs agents across 20M+ lines of code, with Niklas Gustavsson.” <https://www.youtube.com/watch?v=9DHZLw5653E>. *Dream Machine* Issue 40.

38. *Engadget*, “OpenAI launches a physical keypad for controlling agents.” <https://www.engadget.com/2215952/openai-launches-a-physical-keypad-for-controlling-agents/>. *Dream Machine* Issue 38.

Chapter 13 — Authenticity as the New Scarcity

33a. *Variety*, “AI Music Suno Funding Round \$400 Million \$5.4 Billion Valuation.” <https://variety.com/2026/digital/news/ai-music-suno-funding-round-400-million-5-4-billion-valuation-1236765727/>. *TechCrunch*, “Still Facing Copyright Lawsuits, AI Music Generator Suno Raises Another \$400M.” <https://techcrunch.com/2026/06/03/still-facing-copyright-lawsuits-ai-music-generator-suno-raises-another-400m/>. *Dream Machine* Issue 32.

33b. Warner Music Group settled its copyright litigation against Suno and signed a licensing partnership; Universal Music Group and Sony Music remain in active litigation as of June 2026. *Dream Machine* Issue 32. For the structural context of the WMG settlement-as-licensing-table move, see Chapter 6 §“The three Cs”: *the lawsuits were never the destination; they were leverage to force a licensing table into existence.*

1. *Dream Machine* Issue 29 reportage of Tiny Grandma stop-motion content being wrongly flagged as AI by YouTube’s automated detection, May 2026.
2. *Dream Machine* Issue 23, April 2026, reporting death threats against Eline Van der Velden following Tilly Norwood’s continuing public role.
3. *Digital Music News*, “Instagram Chief Says We Should ‘Fingerprint Real Media’ Instead of Tracking and Disclosing AI Slop.” <https://www.digitalmusicnews.com/2026/01/05/instagram-chief-ai-slop-comments/>. *WebProNews*, “Instagram Head Warns AI Images Erode Trust, Calls for Verification Standards.” <https://www.webpronews.com/instagram-head-warns-ai-images-erode-trust-calls-for-verification-standards/>. *Dream Machine* Issue 13.
4. *Digital Music News*, “AI-Generated Far-Right Hate Songs Aren’t Just a Problem in the US — Now They’re Spreading Across Europe Too.” <https://www.digitalmusicnews.com/2025/11/09/ai-generated-hate-songs-dutch-spotify-charts/>. *Dream Machine* Issue 7.
5. Google DeepMind SynthID watermark roll-out across Veo, Lyria and Imagen products. *Dream Machine* Issues 11, 12.
6. Google DeepMind, “Verify Google AI-generated videos in the Gemini app.” https://www.linkedin.com/posts/googledeepmind_verify-google-ai-generated-videos-in-the-activity-7407748300688478208-fJgW. *Dream Machine* Issue 12; broader coverage in *SmartBrief*, “Google’s Gemini can now spot AI-generated videos.” <https://newsletter.smartbrief.com/sharedSummary/index.jsp?briefId=40A39351-5419-4681-94DF-31A53480B698&issueId=58E986AD-821F-422E-9E34-3386E0E2272B©Id=2DB8E453-8E83-416C-949B-44751F252A8D>. *Dream Machine* Issue 13.
7. *Dream Machine* Issues 23, 27 reportage on Taylor Swift’s voice/image trademark filings.
- 7a. *Billboard*, “Lionel Richie Is Trademarking the Sound of His Voice: ‘Hello, Is It Me You’re Looking For?’” <https://www.billboard.com/pro/lionel-richie-is-trademarking-the-sound-of-his-voice/>. *Dream Machine* Issue 34.
- 7b. *Digital Music News*, “Backstreet Boys Look to Decommission AI Deepfakes by Trademarking Their Voices — Joining Taylor Swift, Lionel Richie, and Others.” <https://www.digitalmusicnews.com/2026/06/26/backstreet-boys-trademark/>. *Dream Machine* Issue 36.
- 7c. *Songlines*, “A fake AI album released in my name made me realise why the law needs to change” (Emily Portman). <https://www.songlines.co.uk/features/emily-portman-faces-ai-fake-albums-on-streaming-platforms>. *Dream Machine* Issue 34.

8. *Lawyer Monthly*, “Matthew McConaughey Draws a Line to Protect His Voice and Image From AI.” <https://www.lawyer-monthly.com/2026/01/matthew-mcconaughey-protects-voice-image-ai/>. *Dream Machine* Issue 15.

9. *Adweek*, “Meet the \$1.3 Billion Startup Behind Madonna and Will Smith’s AI Video.” <https://www.adweek.com/media/higgsfield-ai-marketing-startup/>. *Dream Machine* Issue 16.

9a. *Digital Music News*, “Madonna Calls AI ‘the Complete Opposite of Making Art’ While Advocating for a Completely Different Creative Process.” <https://www.digitalmusicnews.com/2026/06/28/madonna-ai-algorithm-commentary/>. *Dream Machine* Issue 36.

10. *Variety*, “George Clooney Says AI Actors Will Face the ‘Same Problem We Have’ in Hollywood: ‘Making a Star Is Not So Easy’.” <https://variety.com/2025/scene/columns/george-clooney-ai-actors-movie-stars-1236579661/>. *Dream Machine* Issue 7.

11. *Deadline*, “AI Documentary Director Insists Jeremy Renner Agreed To Narrate Movie As ‘Hawkeye’ Star Threatens ‘Multi-Millions’ Lawsuit.” <https://deadline.com/2025/11/jeremy-renner-lawsuit-threat-ai-movie-1236611830/>. *Dream Machine* Issue 7.

12. *Complete Music Update*, “Johnny Cash estate uses ELVIS Act to sue Coke over tribute act ad soundtrack.” <https://completemusicupdate.com/johnny-cash-estate-uses-elvis-act-to-sue-coke-over-tribute-act-ad-soundtrack/>. *Dream Machine* Issue 9.

13. *The Verge*, “New York’s new law forces advertisers to say when they’re using AI avatars.” <https://www.theverge.com/news/842848/new-york-law-ai-advertisements-sag-aftra-labor>. *Dream Machine* Issue 11.

13a. *Variety*, “More Than 16,000 Sign SAG-AFTRA Letter Demanding Congress Pass NO FAKES Act, Which Would Ban Unauthorized AI Images and Videos.” <https://variety.com/2026/biz/news/sag-aftra-letter-congress-pass-no-fakes-act-1236783056/>. *Dream Machine* Issue 35.

14. *Fast Company*, “Governments around the world are considering bans on Grok’s app over AI sexual image scandal.” <https://www.fastcompany.com/91474131/governments-around-the-world-are-considering-bans-on-groks-app-over-ai-sexual-image-scandal>. *Dream Machine* Issue 14.

15. Cannes AI Disclosure Standard, launched May 2026. *Dream Machine* Issue 29.

16. *Dream Machine* Issue 28, May 2026, reporting on the Academy of Motion Picture Arts and Sciences’ “You must be human to win” rule update.

17. *The Hollywood Reporter*, “Emmys Set AI Guidance.” <https://www.hollywoodreporter.com/tv/tv-news/emmys-ai-guidelines-2026-awards-1236468434/>. *Dream Machine* Issue 14.
18. SAG-AFTRA negotiation timeline through *Dream Machine* Issues 7, 12, 15, 20, 26, 29.
19. *Marketing Week*, “You can’t dismiss AI ads as slop when they’re winning in testing.” <https://www.marketingweek.com/dismiss-ai-ads-winning-creative-effectiveness/>. *Dream Machine* Issues 8, 13.
20. Variety, “AI Creator Behind Viral ‘Deadpool,’ ‘Harry Potter’ Christmas Clip Made His Film in a Ukrainian Bomb Shelter.” *op. cit.* *Dream Machine* Issue 16.
21. PR Newswire, “From Apple TV Creative to AI Filmmaker: Hoyt Dwyer’s Animated Film To Compete at AI FilmFest Japan 2025.” *op. cit.* *Dream Machine* Issue 6.
22. Google DeepMind, “Dear Upstairs Neighbors.” <https://blog.google/innovation-and-ai/models-and-research/google-deepmind/dear-upstairs-neighbors/>. *Dream Machine* Issue 16.
23. *The Hollywood Reporter*, “‘Synthetic Sincerity’ by Marc Isaacs.” *op. cit.* *Dream Machine* Issue 8.
24. Variety, “‘Watch the Skies,’ Swedish UFO Feature Film Dubbed Entirely With AI, Sets USA Distribution Deal.” *op. cit.* *Dream Machine* Issue 5.
25. *Broadcast Pro Middle East*, “Tunisian filmmaker wins \$1 million AI Film Award for ‘Lily.’” *op. cit.* *Dream Machine* Issue 14.
26. Sundance Institute AI Literacy Initiative emphasis on documentation: <https://www.sundance.org/blogs/centering-the-artist-why-were-launching-the-ai-literacy-initiative/>. *Dream Machine* Issue 16.
- 26a. *PC Gamer*, “More than 300 games released on Steam last week, and 120 of them had AI disclosures.” <https://www.pcgamer.com/gaming-industry/steam-week-in-review-more-than-300-games-released-on-steam-last-week-and-120-of-them-had-ai-disclosures/>. *Dream Machine* Issue 34.
- 26b. *80.lv*, “Epic Games CEO Criticizes Valve’s AI Disclosure Policy for Steam.” <https://80.lv/articles/epic-ceo-criticizes-valve-s-ai-disclosure-policy-for-steam>. *Dream Machine* Issue 36.
27. Google DeepMind, “SynthID — 100 billion watermarks, expanding to partner ecosystems including OpenAI, ElevenLabs and Kakao.” <https://deepmind.google/discover/blog/synthid-100-billion-watermarks-partners/>. *Dream Machine* Issue 30.

27a. ElevenLabs, “Detecting audio generated by ElevenLabs with SynthID” — partnership with Google DeepMind. <https://elevenlabs.io/blog/synthid>. *Dream Machine* Issue 36.

28. *Hollywood Reporter*, “Bobby Berk Says AI Will Make Reality TV & ‘Verifiably Human Content’ More Valuable.” <https://www.hollywoodreporter.com/tv/tv-news/bobby-berk-ai-reality-tv-1236592920/>. *Dream Machine* Issue 30.

29. *Rolling Stone*, “The Rolling Stones Release New Single ‘In the Stars’ — With a Music Video De-Aging the Rockers Courtesy of AI.” <https://www.rollingstone.com/music/music-news/rolling-stones-in-the-stars-ai-de-aging-video-1235142200/>. *Hollywood Reporter*, “‘South Park’ Creators’ AI Company Made The Rolling Stones Young Again for ‘In The Stars’ Music Video.” <https://www.hollywoodreporter.com/tv/tv-news/south-park-creators-ai-rolling-stones-in-the-stars-1236592855/>. *Dream Machine* Issue 30.

29a. *Music Business Worldwide*, “Boy George uses AI to re-record ‘Karma Chameleon’ as first release from tech startup Artist.” <https://www.musicbusinessworldwide.com/boy-george-uses-ai-to-re-record-karma-chameleon-as-first-release-from-tech-startup-artist-included/>. *Dream Machine* Issue 35.

30. *Variety*, “Cate Blanchett Co-Founds RSL Media, a Non-Profit to Address Consent Around AI Usage including creative work, name, image and likeness.” <https://variety.com/2026/film/news/cate-blanchett-rsl-media-ai-consent-1236748255/>. *Dream Machine* Issue 30.

30a. *Variety*, “Cate Blanchett, Steven Soderbergh, Nikki Hexum on How the Human Consent Registry Helps People Protect Their Identity From AI Theft: It’s a ‘Human Right’.” <https://variety.com/2026/biz/news/cate-blanchett-steven-soderbergh-nikki-hexum-ai-human-registry-1236791271/>. *Dream Machine* Issue 36.

31. *Bloomberg*, “Apple Acquires Key Talent & Patents Behind AI Avatar Company ‘Animato’.” <https://www.bloomberg.com/news/articles/2026-05-19/apple-acquires-animato-ai-avatar-talent-patents>. *Dream Machine* Issue 30.

31a. Estate of Ozzy Osbourne announces the late artist will “live on as an AI avatar that talks with fans,” June 2026. *Dream Machine* Issue 31.

31b. *NBC News*, “Netflix says it will use AI to recreate the voice of the late Gene Wilder in a new reality TV series set in the world of Willy Wonka & the Chocolate Factory.” <https://www.nbcnews.com/pop-culture/pop-culture-news/gene-wilder-ai-voice-netflix-willy-wonka-competiton-show-rcna352440>. *Dream Machine* Issue 36.

32. *The Drum*, “David Beckham Designs ‘Henchester United’ Chicken Coop in Lenovo Ad.” <https://www.thedrum.com/news/2026/05/18/david-beckham->

henchester-united-chicken-coop-lenovo-ai-ad. *Dream Machine* Issue 30.

32a. YouTube announces automatic labelling of AI-generated videos, June 2026. *Dream Machine* Issue 31.

32b. LinkedIn announces it will limit the reach of AI-generated content on its feed, June 2026. *Dream Machine* Issue 31.

32c. Hasbro's Sixth Wall launch and the CharacterOS governance layer, 3 June 2026. BusinessWire:

<https://www.businesswire.com/news/home/20260603297922/en/>. *The Wrap*:

<https://www.thewrap.com/industry-news/business/hasbro-ai-studio-sixth-wall-launch/>. Cross-reference Chapter 7 §“Position Five” and Chapter 6 §“The three Cs”.

32d. Spotify–Universal Music Group landmark licensing agreements, 21 May 2026.

Spotify Newsroom: <https://newsroom.spotify.com/2026-05-21/universal-music-group-spotify-licensing-agreements-fan-made-covers-remixes/>. *TechCrunch*:

<https://techcrunch.com/2026/05/21/spotify-and-universal-music-strike-deal-allowing-fan-made-ai-covers-and-remixes/>. See footnote ^37a in Chapter 6.

32e. *Variety*, “Nearly 1,000 Sign Open Letter Condemning Contract Clauses That Require Child Actors to Sign Over Their Voices for AI” (the *Peppa Pig* / Hasbro backlash). <https://variety.com/2026/tv/news/open-letter-major-studio-hasbro-children-ai-peppa-pig-1236790351/>. *Dream Machine* Issue 36.

32f. *The Hollywood Reporter*, “Warner Music Group Acquires AI Detection Company Sureel.” <https://www.hollywoodreporter.com/music/music-industry-news/warner-music-group-buys-ai-detection-company-sureel-1236618266/>. *Dream Machine* Issue 34.

32g. *Variety*, “Sony Invests in AI Copyright-Protection Firm Midnight Labs.” <https://variety.com/2026/biz/news/ip-ai-firm-midnight-labs-sony-million-dollar-investment-1236766931/>. *Dream Machine* Issue 33.

32h. *Engadget*, “Deezer launches an AI music detector for other streaming services.” <https://www.engadget.com/2192024/deezer-ai-music-detector/>. *Dream Machine* Issue 34.

32i. *Music Business Worldwide*, “Tech startup RTM Audio launches AI music detector UAI, which issues a signed certificate with every verdict.” <https://www.musicbusinessworldwide.com/tech-startup-rtm-audio-launches-ai-music-detector-uai-which-issues-a-signed-certificate-with-every-verdict/>. *Dream Machine* Issue 35.

32j. *Music Ally*, “Played By Humans is verifying ‘human artistry in an AI age’.” <https://musically.com/2026/06/11/played-by-humans-is-verifying-human-artistry-in-an-ai-age/>. *Dream Machine* Issue 34.

32k. *MusicTech*, “Traxsource to introduce labels to identify both human-made and AI-assisted music.” <https://musictech.com/news/music/traxsource-labels-identify-human-made-and-ai-assisted-music/>. *Dream Machine* Issue 36.

32l. *AOL*, “AI labels to be compulsory on authentic-looking content under EU rules.” <https://www.aol.co.uk/articles/ai-labels-compulsory-authentic-looking-094628000.html>. *Dream Machine* Issue 40.

32m. *Musically*, “Key EU AI Act rules now being enforced – so how does that affect the music industry?” <https://musically.com/2026/08/04/key-eu-ai-act-rules-now-being-enforced-so-how-does-that-affect-the-music-industry/>. *Dream Machine* Issue 40.

32n. *Spotify Newsroom*, “Spotify Introduces a New ‘AI Persona’ Label for AI-Generated Artist Identities.” <https://newsroom.spotify.com/2026-08-11/ai-persona-badges-transparency/>. *Dream Machine* Issue 41.

Chapter 14 – Coordination Collapse

1. *Dream Machine* Issue 5, “Industry Insights: Stealth, Shadow and Secret AI Users.”

2. *Azumo*, “AI in Workplace Statistics 2025.” <https://azumo.com/artificial-intelligence/ai-insights/ai-in-workplace-statistics>. *Tech.co*, “Gen Z Most Likely Use AI Boss.” <https://tech.co/news/gen-z-most-likely-use-ai-boss>. *Dream Machine* Issue 5.

3. *Exploding Topics*, “AI Workforce Research.” <https://explodingtopics.com/blog/ai-workforce-research>. *Dream Machine* Issue 5.

4. *Forbes*, “AI Tools Flood Workplaces as Employees Face a Double Bind.” <https://www.forbes.com/sites/carolinecastrillon/2025/09/09/ai-tools-flood-workplaces-as-employees-face-a-double-bind/>. *Dream Machine* Issue 5.

5. *Blog IDC Europe*, “Shadow AI: How Stealth Productivity Is Strangling Enterprise AI Adoption and Creating a Security Nightmare.” <https://blog-idceurope.com/shadow-ai-how-stealth-productivity-is-strangling-enterprise-ai-adoption-and-creating-a-security-nightmare/>. *Dream Machine* Issue 5.

6. *Game Developer*, “Subnautica owner Krafton outlines plans to transform into an ‘AI First’ company.” <https://www.gamedeveloper.com/business/subnautica-owner-krafton-outlines-plans-to-transform-into-an-ai-first-company>. *Dream Machine* Issue 6.

7. *Dream Machine* Issue 24, April 2026, on the GTA VI publisher laying off its internal AI team.

8. *Dream Machine* Issue 25, April 2026, on Disney layoffs including Marvel staff.
9. *SmartBrief*, “Meta to cut 10% of Reality Labs staff to focus on AI.” <https://newsletter.smartbrief.com/sharedSummary/index.jsp?briefId=40A39351-5419-4681-94DF-31A53480B698&issueId=025444D1-A590-46D8-B969-EF81DEE05228©Id=1B5F70D2-FFDA-4660-9CE9-047C9B16BF83>. *Dream Machine* Issue 14.
10. *Dream Machine* Issue 23, April 2026, on Scottish animation studio collapse.
11. *Metro*, “Prince of Persia remake and five more games cancelled as Ubisoft focuses on AI.” *op. cit.* *Dream Machine* Issue 15.
12. *The Guardian*, “AI is hitting UK harder than other big economies, study finds.” <https://www.theguardian.com/technology/2026/jan/26/ai-uk-jobs-us-japan-germany-australia>. *Dream Machine* Issue 16.
13. *The Economist*, “Investors expect AI use to soar. That’s not happening.” <https://www.economist.com/finance-and-economics/2025/11/26/investors-expect-ai-use-to-soar-thats-not-happening>. *Dream Machine* Issue 9.
14. *Dream Machine* Issue 24, April 2026, on OpenAI’s public-policy proposals around AI-driven economic disruption.
15. *The Economist*, “Job apocalypse? Humbug! AI is creating brand new occupations.” <https://www.economist.com/business/2025/12/14/job-apocalypse-humbug-ai-is-creating-brand-new-occupations>. *Dream Machine* Issue 12.
16. *Forbes*, “Vibe Coding — The In Demand AI Skill.” <https://www.forbes.com/sites/rachelwells/2025/11/06/the-in-demand-ai-skill-and-certifications-that-pays-up-to-220000/>. *Dream Machine* Issue 8.
17. U.K. Department for Business and Trade research on neurodiverse workers and AI assistants, autumn 2025. Reported via *CNBC*, “People with ADHD, autism, dyslexia say AI agents are helping them succeed at work.” <https://www.cnn.com/2025/11/08/adhd-autism-dyslexia-jobs-careers-ai-agents-success.html>. *Dream Machine* Issue 7.
18. *CNBC*, *op. cit.*
19. *Dream Machine* Issue 7 secondary references.
20. Korin AI launch, May 2026. *Dream Machine* Issue 27.
21. *CNBC Africa*, “How AI is changing the landscape of the music industry in Africa.” <https://www.cnn.com/2025/11/08/adhd-autism-dyslexia-jobs-careers-ai-agents-success.html>. *Dream Machine* Issue 5.
22. *BBC Future*, “Lights, camera, algorithm: Why Indian cinema is awash with AI.” <https://www.bbc.co.uk/future/article/20251223-why-indian-cinema-is-awash->

with-ai. *Dream Machine* Issue 14.

23. *Dream Machine* Issue 25, April 2026, on Indonesia's *Legenda Bertuah*.

24. *Broadcast Pro Middle East*, "Tunisian filmmaker wins \$1 million AI Film Award for 'Lily'." *op. cit. Dream Machine* Issue 14.

25. *Digiday*, "Avocados From Mexico turns to AI to advertise around the Super Bowl instead of a TV buy." *op. cit. Dream Machine* Issue 15.

26. *Variety*, "AI Creator Behind Viral 'Deadpool,' 'Harry Potter' Christmas Clip Made His Film in a Ukrainian Bomb Shelter." *op. cit. Dream Machine* Issue 16.

27. *Dream Machine* Issue 8 citing Andreessen Horowitz observations: https://www.linkedin.com/posts/stevenouri_a-wild-stat-80-of-startups-pitching-a16z-activity-7396182718998351872-xTKR.

28. *PocketGamer.biz*, "Shift Up CEO says AI is key to competing with China's game industry scale." <https://www.pocketgamer.biz/shift-up-ceo-says-ai-is-key-to-competing-with-chinas-game-industry-scale/>. *Dream Machine* Issue 14.

29. Enterprise-AI workforce tracking, late 2025. Aggregated in the Deep Dive companion piece *The Shadow AI Paradox in the Creative Industries*, drawing on Azumo's *AI in Workplace Statistics 2025*, Tech.co's Gen Z survey, and the IDC Europe shadow-AI security brief. *Dream Machine* Issue 5.

30. *Hidden Cloud Explosion* analysis, IDC Europe, 2025. See *The Shadow AI Paradox in the Creative Industries*, §"The Epistemology and Scale of Shadow AI."

31. Shadow-AI security-incident statistics, 2025, aggregated in *The Shadow AI Paradox in the Creative Industries*, §"The Epistemology and Scale of Shadow AI"; underlying data via IBM *Cost of a Data Breach Report 2025* and IDC Europe.

32. For the developer-community origins of the "AI for thee, but not for me" phrasing, and the full sectoral analysis of the paradox, see *The Shadow AI Paradox in the Creative Industries*, §"The Great Hypocrisy."

33. Survey of 1,100+ professional music creators, 2026, summarised in *Dynamics of Generative AI Adoption in the Creative Industries*, §"Music Production and Sound Recording," and *The Shadow AI Paradox in the Creative Industries*, §"Sector-Specific Analysis."

34. WGA screenwriter survey, pre- and post-strike, reported in *Dynamics of Generative AI Adoption in the Creative Industries*, §"Screenwriting and the Post-Strike AI Boom."

35. Adobe Firefly milestone data, September 2023 – June 2025, in *Dynamics of Generative AI Adoption in the Creative Industries*, §"The Ubiquity of AI in Visual and Digital Arts." *Dream Machine* Issue 6.

36. Adobe quarterly financials, FY2025–FY2026; AI-first ARR growth reported in *Dream Machine* Issue 21 and summarised in *Dynamics of Generative AI Adoption*.
37. Adobe Firefly enterprise penetration metrics, in *Dynamics of Generative AI Adoption*.
38. Adobe Stock submission analysis, 2024, in *Dynamics of Generative AI Adoption*.
39. Adobe, “Inaugural Adobe Creators’ Toolkit Report,” October 2025. <https://news.adobe.com/news/2025/10/adobe-max-2025-creators-survey>. *Dream Machine* Issue 6.
40. ChatGPT weekly-active-user disclosures, mid-2025; consolidated in *Dynamics of Generative AI Adoption*, §“General Purpose LLMs.”
41. Gemini desktop-user growth, year-over-year, in *Dynamics of Generative AI Adoption*.
42. Stanford AI Index Report 2025, global-sentiment chapter. Summarised in *Dynamics of Generative AI Adoption*, §“The Perception Gap.”
43. YouGov 2024 multi-market AI sentiment survey, 17 countries. Summarised in *Dynamics of Generative AI Adoption*, §“The Perception Gap.”
44. Quantic Foundry consumer-AI-in-gaming survey, 2025. Summarised in *Dynamics of Generative AI Adoption*, §“The Video Game Industry.”
45. Game Developers Conference *State of the Game Industry* surveys, 2024–2026, sentiment vs. usage trend. Reported in *Dynamics of Generative AI Adoption*, §“The Video Game Industry.”
46. O’Hare, J. (2026). *The Coordination Collapse: How AI Is Dissolving the Information Tax That Created Hierarchy — And What Replaces It*. DreamLab AI Consulting Ltd, April 2026. The report draws on Dell’Acqua et al.’s preregistered Boston Consulting Group experiment with 758 consultants, Mollick’s *secret cyborgs* coinage, Edmondson and Lei’s psychological-safety findings, the IBM *Cost of a Data Breach Report 2025*, the Epoch AI / Ipsos KnowledgePanel survey of March 2026, and the CyberArk *hidden org chart* analysis.
- 46a. Epidemic Sound, *Future of the Creator Economy Report 2026* (survey of 3,000 professional creators across UK and US, June 2026); *Musically* commentary on the report and creator attitudes towards AI. *Dream Machine* Issue 31, Editor’s Pick, June 2026.
47. Ramp, “Ramp AI Index,” June 2026. <https://ramp.com/data/ai-index-june-2026>. *Dream Machine* Issue 34.

48. *AppleInsider*, “Adobe’s glowing AI survey leaves out most of the creative industry.” <https://appleinsider.com/articles/26/06/16/adobes-glowing-ai-survey-leaves-out-most-of-the-creative-industry>. *Dream Machine* Issue 35.

49. *PC Gamer*, “Open source game engine Godot will no longer accept AI-authored code contributions: ‘We can’t trust heavy users of AI to understand their code enough to fix it’.” <https://www.pcgamer.com/gaming-industry/open-source-game-engine-godot-will-no-longer-accept-ai-authored-code-contributions-we-cant-trust-heavy-users-of-ai-to-understand-their-code-enough-to-fix-it/>. *Dream Machine* Issue 36.

50. *Fortune*, “Top CFOs warn AI success depends on training employees, not just buying technology.” <https://fortune.com/2026/06/25/cfo-hpe-prologis-smucker-moodys-ai-success-training-employees/>. *Dream Machine* Issue 36.

51. *GameDev.net*, “Gamescom Dev speaker survey identifies potential impacts of AI on team structure and productivity.” <https://gamedev.net/news/5025-gamescom-dev-speaker-survey-identifies-potential-impacts-of-ai-on-team/>. *Dream Machine* Issue 41.

52. *CNBC*, “Microsoft cuts 2.1% of employees as Xbox unit plans to spin off studios.” <https://www.cnbc.com/2026/07/06/microsoft-cuts-2point1percent-of-employees-as-xbox-unit-plans-to-spin-studios.html>. *Dream Machine* Issue 37.

53. *KPMG*, “Analysis of 1.4 million AI interactions identifies the employee behaviors behind effective AI use.” <https://kpmg.com/us/en/articles/2026/sophisticated-ai.html>. *Dream Machine* Issue 37.

54. *Fortune*, “Sir Martin Sorrell has identified the key AI skill for the future: ‘People who share will be the new kings and queens’.” <https://fortune.com/2026/08/04/sir-martin-sorrell-has-identified-the-key-ai-skill-for-the-future-people-who-share-will-be-the-new-kings-and-queens/>. *Dream Machine* Issue 40. See also *The Drum*, “Martin Sorrell: Cannes is obsessing over the wrong AI debate.” <https://www.thedrum.com/news/martin-sorrell-cannes-is-obsessing-over-the-wrong-ai-debate>. *Dream Machine* Issue 37.

Chapter 15 — The New Jobs

1. *The Economist*, “Investors expect AI use to soar. That’s not happening.” <https://www.economist.com/finance-and-economics/2025/11/26/investors-expect-ai-use-to-soar-thats-not-happening>. *Dream Machine* Issue 9.

2. *The Economist*, “Job apocalypse? Humbug! AI is creating brand new occupations.” <https://www.economist.com/business/2025/12/14/job-apocalypse-humbug-ai-is-creating-brand-new-occupations>. *Dream Machine* Issue 12.

3. *The Guardian*, “AI is hitting UK harder than other big economies, study finds.” <https://www.theguardian.com/technology/2026/jan/26/ai-uk-jobs-us-japan-germany-australia>. *Dream Machine* Issue 16.
4. University of Wisconsin-Stout, “AI Reshaping Industry: New UW-Stout Course Sets AI-Use as Baseline Competency in Filmmaking.” <https://www.uwstout.edu/about-us/news-center/ai-reshaping-industry-new-uw-stout-course-sets-ai-use-baseline-competency-filmmaking>. *Dream Machine* Issue 15.
5. Adobe Firefly enterprise metrics, tracked across the *Dream Machine* archive.
6. Reuters Institute, “AI adoption by UK journalists and their newsrooms.” <https://reutersinstitute.politics.ox.ac.uk/ai-adoption-uk-journalists-and-their-newsrooms-surveying-applications-approaches-and-attitudes>. *Digiday*, “Daily Mail says Google AI Overviews have killed click-throughs.” <https://digiday.com/media/daily-mail-says-google-ai-overviews-have-killed-click-throughs/>. *Dream Machine* Issues 7, 9.
7. 1,100-creator music survey 2026 (music production and sound recording).
8. VFX AI integration metrics (visual-effects automation).
9. *PC Gamer*, “Square Enix aims to have AI doing 70% of its QA work by the end of 2027.” <https://www.pcgamer.com/gaming-industry/square-enix-aims-to-have-ai-doing-70-percent-of-its-qa-work-by-the-end-of-2027/>. *Dream Machine* Issue 7.
10. *Eurogamer*, “Falcom is the latest developer to buy into the AI hype machine.” <https://www.eurogamer.net/falcom-is-the-latest-developer-to-buy-into-the-ai-hype-machine>. *Dream Machine* Issue 12.
11. *Metro*, “Prince of Persia remake and five more games cancelled as Ubisoft focuses on AI.” <https://metro.co.uk/2026/01/21/prince-persia-remake-five-games-cancelled-ubisoft-focuses-ai-26431926/>. *Dream Machine* Issue 15.
12. *Dream Machine* Issue 24, April 2026, on the GTA VI publisher laying off its internal AI team.
13. ElevenLabs \$500m ARR reporting, April 2026. *Dream Machine* Issue 25.
14. *36kr*, “AI Video Unicorn Higgsfield: Earns \$200M in 9 Months by ‘Serving’ Social Media Marketers.” <https://eu.36kr.com/en/p/3650517574312323>. *Dream Machine* Issue 16.
15. *Dream Machine* Issue 29, May 2026, on Sony’s 49-agent / 72-skill multi-agent game-development team.
16. *Digiday*, “AI agent developers have become adland’s in-demand role.” <https://digiday.com/marketing/ai-agent-developers-have-become-adlands-in->

demand-role/. *Dream Machine* Issue 7.

17. *Forbes*, “Vibe Coding — The In Demand AI Skill That Pays Up to \$220,000.” <https://www.forbes.com/sites/rachelwells/2025/11/06/the-in-demand-ai-skill-and-certifications-that-pays-up-to-220000/>. *Dream Machine* Issue 8.

18. Sundance Institute, “Centering the Artist: Why We’re Launching the AI Literacy Initiative.” <https://www.sundance.org/blogs/centering-the-artist-why-were-launching-the-ai-literacy-initiative/>. Google blog, “Sundance Institute AI Education.” <https://blog.google/company-news/outreach-and-initiatives/google-org/sundance-institute-ai-education/>. *Dream Machine* Issues 15, 16.

19. UK Government, “Free AI training for all.” <https://www.gov.uk/government/news/free-ai-training-for-all-as-government-and-industry-programme-expands-to-provide-10-million-workers-with-key-ai-skills-by-2030>. *Dream Machine* Issue 16.

20. Lovable for classrooms. <https://lovable.dev/classroom>. *Dream Machine* Issue 11.

21. UW-Stout course launch, January 2026 — *op. cit.*

22. Adobe, “Firefly Foundry.” <https://news.adobe.com/news/2025/10/adobe-max-2025-firefly-foundry>. *Dream Machine* Issue 5.

23. Korin AI launch, May 2026. *Dream Machine* Issue 27.

24. *The Verge*, “New York’s new law forces advertisers to say when they’re using AI avatars.” <https://www.theverge.com/news/842848/new-york-law-ai-advertisements-sag-aftra-labor>. *Dream Machine* Issue 11. C2PA / SynthID infrastructure references in Chapter 13.

25. *Forbes*, “AI Is Changing How Creators Work And Earn.” <https://www.forbes.com/sites/kolawolesamueladebayo/2025/12/22/how-ai-is-changing-how-creators-work-and-earn/>. *Dream Machine* Issue 13.

26. Variety, “AI Creator Behind Viral ‘Deadpool,’ ‘Harry Potter’ Christmas Clip Made His Film in a Ukrainian Bomb Shelter.” <https://variety.com/2026/digital/news/ai-video-deadpool-harry-potter-andrii-daniels-1236624632/>. *Dream Machine* Issue 16.

27. *Broadcast Pro Middle East*, “Tunisian filmmaker wins \$1 million AI Film Award for ‘Lily.’” <https://www.broadcastprome.com/news/tunisian-filmmaker-wins-1-million-ai-film-award-for-lily/>. *Dream Machine* Issue 14.

28. *Dream Machine* Issue 25, April 2026, on Indonesia’s *Legenda Bertuah*.

29. BBC Future, “Lights, camera, algorithm: Why Indian cinema is awash with AI.” <https://www.bbc.co.uk/future/article/20251223-why-indian-cinema-is-awash->

with-ai. *Dream Machine* Issue 14.

30. *TechBullion*, “Why the future belongs to multi-skilled leaders.” <https://techbullion.com/playing-the-long-game-with-a-portfolio-career-why-the-future-belongs-to-multi-skilled-leaders/>. *Dream Machine* Issue 9.

31. Anthropic Skills framework via Claude Code. *Dream Machine* Issues 11, 16, 29.

32. Adobe, “Inaugural Adobe Creators’ Toolkit Report,” October 2025. <https://news.adobe.com/news/2025/10/adobe-max-2025-creators-survey>. *Dream Machine* Issue 6.

33. PRS for Music, “PRS for Music AI Survey 2026.” <https://www.prsformusic.com/m-magazine/news/prs-for-music-ai-survey-2026>. *Dream Machine* Issue 16.

34. *CNBC*, “People with ADHD, autism, dyslexia say AI agents are helping them succeed at work.” <https://www.cnn.com/2025/11/08/adhd-autism-dyslexia-jobs-careers-ai-agents-success.html>. *Dream Machine* Issue 7.

35. McKinsey & Company, “What AI could mean for film and TV production and the industry’s future.” <https://www.mckinsey.com/industries/technology-media-and-telecommunications/our-insights/what-ai-could-mean-for-film-and-tv-production-and-the-industrys-future>. *Dream Machine* Issue 16.

36. GDC State of the Game Industry surveys 2024–2026.

36a. GDC, “GDC 2026 State of the Game Industry Reveals Impact of Layoffs, Generative AI and More” — 52% of developers say generative AI is having a negative impact on the industry, up from 30% the year before and 18% the year before that. <https://gdconf.com/article/gdc-2026-state-of-the-game-industry-reveals-impact-of-layoffs-generative-ai-and-more/>. *Dream Machine* Issue 41.

37. LANDR AI music study, late 2025, referenced via *Ari’s Take*. <https://aristake.com/ai-tools-musicians-study/>. *Dream Machine* Issue 8.

38. Stanford AI Index Report 2025.

39. YouGov 2024 multi-market AI sentiment survey.

40. *Digital Music News*, “Nearly 800 Creatives Sign Responsible AI Declaration — ‘Stealing Our Work Is Not Innovation’.” <https://www.digitalmusicnews.com/2026/01/22/stealing-isnt-innovation/>. *Dream Machine* Issue 16.

41. *Broadcast Pro Middle East*, Lily award — *op. cit.*

42. *Variety*, Andrii Daniels bomb-shelter clip — *op. cit.*

43. *BBC Future*, “Lights, camera, algorithm” — *op. cit.*

44. *Dream Machine* Issue 25, Indonesian *Legenda Bertuah*.

45. CNBC Africa, “How AI is changing the landscape of the music industry in Africa.” <https://www.cnbc africa.com/2025/how-ai-is-changing-the-landscape-of-the-music-industry-in-africa>. *Dream Machine* Issue 5. Korin AI launch, May 2026 — *op. cit.*

46. *PocketGamer.biz*, “Shift Up CEO says AI is key to competing with China’s game industry scale.” <https://www.pocketgamer.biz/shift-up-ceo-says-ai-is-key-to-competing-with-chinas-game-industry-scale/>. *Dream Machine* Issue 14.

47. The doomer frame is set out at length in OpenAI’s spring 2026 policy paper on the “social contract for the intelligence age” (covered in *Dream Machine* Issue 24) and in Daniel Susskind’s *A World Without Work*. The full empirical rebuttal — including the Stanford Digital Economy Lab’s *Canaries in the Coal Mine?* findings (Brynjolfsson, Chandar & Chen, 2025), Brynjolfsson, Li & Raymond’s *Generative AI at Work* (QJE, 2025; n = 5,179 agents), Acemoglu’s *Simple Macroeconomics of AI* (Economic Policy, 2025), the Dallas Fed’s labour analyses, the St. Louis Fed’s AI-and-GDP tracking, the 2026 Stanford AI Index, and the 230-year history of identical predictions — is set out at length in DreamLab’s published research. That analysis is the analytical antibody to the policy framing this section is pushing back against.

48. *Ad Age*, “PMG’s hackathon at Cannes signals agencies’ growing appetite for AI talent.” <https://adage.com/technology/ai/aa-pmg-cannes-hackathon-agency-skilled-talent/>. *Dream Machine* Issue 36.

49. Anthropic, “Claude Corps: teaching 1,000 fellows to master Claude, matched with nonprofits across America.” <https://www.anthropic.com/news/claude-corps>. *Dream Machine* Issue 34.

50. GOV.UK, “AI Adoption Plan for the Creative Industries.” <https://www.gov.uk/government/publications/ai-champions-ai-adoption-plans/ai-adoption-plan-creative-industries>. *Dream Machine* Issue 33.

51. *The Hollywood Reporter*, “Hollywood Workers Are Training AI Models as Job Prospects Grow Slim.” <https://www.hollywoodreporter.com/business/digital/ai-training-hollywood-writer-jobs-prospects-1236628302/>. *Dream Machine* Issue 36.

52. *UKTN*, “AI hiring surges ‘across all UK sectors’, according to new findings from PwC, with a 34% salary premium for AI skills.” <https://www.uktech.news/ai/ai-hiring-surges-across-all-uk-sectors-20260615>. *Dream Machine* Issue 34.

53. *PC Gamer*, “EA: AI is delivering ‘faster prototyping’ and a ‘real rise in creativity’ in its studios.” <https://www.pcgamer.com/software/ai/ea-exec-says-ai-has-helped-drive-a-real-rise-of-creativity-at-its-studios/>. *Dream Machine* Issue 36.

54. *Labor Market Matters*, “AI and independent work in 7 charts” — what if AI’s first labour-market effect isn’t replacing workers, but making traditional firms less necessary? <https://www.labormarketmatters.com/p/ai-and-independent-work-in-7-charts>. *Dream Machine* Issue 37.

55. *Digital Trends*, “The AI game-dev boom has a dark side: fewer jobs for the people starting out.” <https://www.digitaltrends.com/cool-tech/the-ai-game-dev-boom-has-a-dark-side-fewer-jobs-for-the-people-starting-out/>. *Dream Machine* Issue 39.

56. *GamesBeat*, “Teen founder raises \$3.2M for AI game development startup ForgeGUI.” <https://gamesbeat.com/teen-founder-raises-3-2m-for-ai-game-development-startup-forgegui/>. *Dream Machine* Issue 39. And *Yahoo*, “My 11-year-old son used AI to build his own video game — I see it as a creativity boost, not a threat.” <https://www.yahoo.com/lifestyle/articles/11-old-son-used-ai-110601483.html>. *Dream Machine* Issue 37.

Chapter 16 — Choosing the Future

1a. *Bloomberg*, “AI Changed Chess. Grandmasters Now Win With Unpredictable Moves,” 27 March 2026. <https://www.bloomberg.com/news/articles/2026-03-27/ai-changed-chess-grandmasters-now-win-with-unpredictable-moves>. *Dream Machine* Issue 23. The behavioural pattern the piece describes — top grandmasters deliberately deviating from machine-optimal lines to put opponents on uncomputed ground — is the cleanest available analogy I have for the strategic shift the rest of this chapter argues for.

1. *Digital Music News*, “The AI Licensing Shift — Creative Weight Attribution Emerges as Music Industry Game-Changer for Rights Holders.” *op. cit.* *Dream Machine* Issue 16.

2. DreamLab AI Collective, team page. <https://dreamlab-ai.com/team>.

3. *Arts Council England*, “Arts Council England launches Responsible AI Cultural Leaders Academy.” <https://www.artscouncil.org.uk/research-and-data/responsible-ai-arts-council-england>. *Dream Machine* Issue 37.

Chapter 17 — The Tools

1. OpenAI, “Sora 2 is here,” 30 September 2025. <https://openai.com/index/sora-2/>. *Dream Machine* Issue 1.

2. LinkedIn News aggregation: “Sora Tops 1 Million Downloads in 5 Days.” <https://www.linkedin.com/news/story/sora-tops-1m-downloads-in-5-days-6684988/>. *Dream Machine* Issue 3.

2a. OpenAI's staged wind-down of Sora, announced 24 March 2026 — consumer app and website ending late April, developer API closing late September; peak revenue reportedly ~\$540k per month against operating costs in the billions. Sources and the corrected Disney–OpenAI record are set out in DreamLab's published research on the programmable brand.

3. Google DeepMind, Veo 3.1 launch, mid-October 2025. *Dream Machine* Issue 3.

4. Runway product cycle: Gen-4.5 (December 2025), Gen-4.5 Image-to-Video (January 2026), Workflows, Story Panels, Characters API, Apps for Advertising — *Dream Machine* Issues 10, 14, 15, 16, 20.

5. Runway CEO on indie films vs. blockbusters, *Dream Machine* Issue 26.

6. Chinese open-source AI video model releases, 2025–2026. *Dream Machine* Issues 3, 12, 22.

7. *SiliconAngle*, “Higgsfield raises \$80M on \$1.3B valuation.” <https://siliconangle.com/2026/01/15/higgsfield-raises-80m-1-3b-valuation-scale-ai-video-platform/>. 36kr, “Higgsfield: Earns \$200M in 9 Months.” <https://eu.36kr.com/en/p/3650517574312323>. *Dream Machine* Issues 15, 16.

8. Heygen Video Agent. https://www.linkedin.com/posts/heygen_introducing-the-new-video-agent-activity-7421597801240801282-d1CF. *Dream Machine* Issue 16.

9. TechCrunch, “Synthesia hits \$4B valuation, lets employees cash in.” <https://techcrunch.com/2026/01/26/synthesia-hits-4b-valuation-lets-employees-cash-in/>. *Sifted*, “Synthesia rejects \$3bn Adobe acquisition offer.” <https://sifted.eu/articles/synthesia-acquisition-offer>. *Dream Machine* Issues 5, 16.

10. ElevenLabs \$500m ARR reporting, April 2026. *Dream Machine* Issue 25.

11. Adobe Firefly milestone data, from DreamLab's published research on the dynamics of generative AI adoption, §“The Ubiquity of AI in Visual and Digital Arts.”

12. Nano Banana inside Photoshop and inside Unreal Engine cross-integrations, October–November 2025. *Dream Machine* Issue 1.

13. Suno Studio launch. <https://www.techradar.com/ai-platforms-assistants/i-tried-suno-studio-the-new-platform-that-mixes-ai-music-generation-with-hands-on-editing-like-garageband-but-smarter>. *Dream Machine* Issue 1.

14. Mureka, “Music Agent Studio” launch. *Dream Machine* Issue 4.

15. ElevenLabs Series funding, April 2026. *Dream Machine* Issue 25.

16. *MusicTech*, “Cardiff band speaks out after AI artist trained on their music outperforms them on Spotify.” <https://musictech.com/news/industry/its-shocking-disheartening-and-insulting-cardiff-band-speaks-out-after-ai-artist-trained-on-their-music-outperforms-them-on-spotify/>. *Dream Machine* Issue 1.

17. Variety, “AI Creator Behind Viral ‘Deadpool,’ ‘Harry Potter’ Christmas Clip Made His Film in a Ukrainian Bomb Shelter.” <https://variety.com/2026/digital/news/ai-video-deadpool-harry-potter-andrii-daniels-1236624632/>. *Dream Machine* Issue 16.
18. Music industry AI deal flow, October 2025 – May 2026. See Chapter 5 footnotes 31–37, and *Dream Machine* Issues 5, 7, 8, 12, 14, 16, 17.
19. World Labs, “Bringing Marble to Life.” <https://www.worldlabs.ai/case-studies/bringing-marble-to-life>. *Dream Machine* Issue 7.
20. Sony Pictures Marble VP integration. *Dream Machine* Issue 8.
21. Google DeepMind, “Genie 3.” <https://deepmind.google/blog/genie-3-a-new-frontier-for-world-models/>. Project Genie: <https://blog.google/innovation-and-ai/models-and-research/google-deepmind/project-genie/>. *Dream Machine* Issues 3, 17.
22. Tencent, “HY World 1.5” and Hunyuan 3D Studio. *Dream Machine* Issue 12.
23. Luma AI, *UNI-1* launch, March 2026. *Dream Machine* Issue 22.
24. SuperSplat / Spark 2.0 / SOG releases through 2025–26. *Dream Machine* Issues 1, 25.
25. Radiance Fields, “Apple Confirms that it’s Gaussian Splatting that powers their personas.” <https://radiancefields.com/apple-confirms-personas-use-gaussian-splatting>. *Dream Machine* Issue 5.
26. ComfyUI Blog, “Ubisoft La Forge Open-Sources the CHORD Model.” <https://blog.comfy.org/p/ubisoft-open-sources-the-chord-model>. *Dream Machine* Issue 11.
27. Anthropic / Blender Foundation patronage, May 2026. *Dream Machine* Issue 27.
28. OpenAI, “Introducing AgentKit.” <https://openai.com/index/introducing-agentkit/>. *Dream Machine* Issue 2.
29. Anthropic Skills framework. *Dream Machine* Issues 11, 16, 29.
30. Heygen Video Agent. *Dream Machine* Issue 16.
31. Adobe Summit 2026 CX Enterprise. *Dream Machine* Issue 26.
32. Adobe + NVIDIA / Google + NVIDIA partnerships. *Dream Machine* Issue 21.
33. ComfyUI funding round. https://www.linkedin.com/posts/comfyui_we-raised-17-million-to-build-an-os-for-ugcPost-7373743341236236288-wkCc. *Dream Machine* Issue 1.
34. ComfyUI \$500M valuation, May 2026. *Dream Machine* Issue 27.

35. Hugging Face / Google Cloud and Meta / Hugging Face OpenEnv. *Dream Machine* Issues 5, 8.
36. Unreal Engine 5 official AI Assistant. https://www.linkedin.com/posts/wouterweynants_theres-an-official-ai-assistant-coming-to-ugcPost-7369377204226379776-pGiH. *Dream Machine* Issue 1.
37. Unity AI Council (October 2025); Unity AI Open Beta (May 2026). *Dream Machine* Issues 1, 28.
38. VFX AI integration metrics, from DreamLab's published research on the dynamics of generative AI adoption, §"Visual Effects (VFX) Automation."
39. Anthropic / Blender Foundation patronage. *Dream Machine* Issue 27.
40. Andreessen Horowitz pitch-deck observations on Chinese open-source model usage. https://www.linkedin.com/posts/stevenouri_a-wild-stat-80-of-startups-pitching-a16z-activity-7396182718998351872-xTKR. *Dream Machine* Issue 8.
41. Korin AI launch, May 2026. *Dream Machine* Issue 27.
42. Google DeepMind, "Introducing Gemini Omni: Create Anything from Any input." <https://blog.google/technology/google-deepmind/gemini-omni-launch/>. *Dream Machine* Issue 30.
43. Beeple Canvas — Generative AI compositor. <https://www.beeple-canvas.com/>. *Dream Machine* Issue 30.
44. Sony AI, "Woosh — a sound effect foundation model." <https://ai.sony/blog/woosh-sound-effect-foundation-model/>. *Dream Machine* Issue 30.
45. Mirelo SFX 1.6, "edit sound, not just generate it." <https://mirelo.ai/sfx-1-6>. *Dream Machine* Issue 30.
46. Stability AI, "Stable Audio 3.0 released — open-weight model family built for artistic experimentation." <https://stability.ai/news/stable-audio-3-0-released>. *Dream Machine* Issue 30.
47. Tamber product page: <https://tamber.ai/>. *Dream Machine* Issue 30.
48. Beatport Track ID. <https://www.beatport.com/track-id>. *Dream Machine* Issue 30.
49. NVIDIA SANA-WM model collection. <https://huggingface.co/collections/nvidia/sana-wm>. *Dream Machine* Issue 30.
50. Odyssey, "Introducing Starchild-1, the first real-time multimodal world model." <https://odyssey.ml/introducing-starchild-1>. *Dream Machine* Issue 30.

51. Odyssey, “Introducing Agora-1 — four-player AI-generated world built on a 1997 shooter.” <https://odyssey.ml/introducing-agora-1>. *Dream Machine* Issue 30.
52. Apple Machine Learning Research, “Apple Headsup: a Large-Scale High-Quality 3D Gaussian Head Reconstruction from Multi-View Captures.” <https://machinelearning.apple.com/research/apple-headsup-3d-gaussian-head>. *Dream Machine* Issue 30.
53. Google, “Official skills for AI agents.” <https://github.com/google/agent-skills>. *Dream Machine* Issue 30.
54. Tencent Ardot, AI-native design agent platform. <https://ardot.tencent.com/>. *Dream Machine* Issue 30.
55. Anthropic, “Claude is now available as a partner node in ComfyUI.” <https://www.anthropic.com/news/claude-comfyui-partner-node>. *Dream Machine* Issue 30.
56. ECABridge — Unreal Engine MCP integration. <https://ecabridge.dev/>. *Dream Machine* Issue 30.
57. Video Games Chronicle, “Epic Games Veteran Claims He’s Building AI-Heavy ‘Fully European’ Game Engine.” <https://www.videogameschronicle.com/news/epic-games-veteran-ai-heavy-fully-european-game-engine/>. *Dream Machine* Issue 30.
58. PhotoGIMP — the open-source GIMP skin that mimics Photoshop. <https://github.com/Diollinux/PhotoGIMP>. *Dream Machine* Issue 30.
59. The Decoder, “ByteDance’s Seedance 2.5 breaks the 30-second barrier for AI video generation.” <https://the-decoder.com/bytedances-seedance-2-5-breaks-the-30-second-barrier-for-ai-video-generation/>. Dreamina Seedance 2.0 Mini: <https://dreamina.capcut.com/seedance/seedance-2-0-mini>. Seedance 2.0 Mini and 4K in ComfyUI: <https://blog.comfy.org/p/seedance-2-0-mini-and-4k-is-now-available>. *Dream Machine* Issues 34, 35, 36.
60. Runway API Recipes: <https://runwayml.com/api>. Runway MCP in ChatGPT: <https://runwayml.com/mcp>. Runway Agent 2 ad localisation: <https://runwayml.com/news/introducing-agent-2>. *Dream Machine* Issues 34, 35, 36.
61. Google Magenta RealTime 2: <https://magenta.withgoogle.com/magenta-realtime-2>. Daydream: <https://daydream.live/>. MusicRadar on Fender Studio Pro and Moises: <https://www.musicradar.com/music-tech/daws/ai-isnt-the-destination-making-music-is-fender-looks-forward-as-its-studio-pro-daw-integrates-moises-stem-separation-and-a-smart-studio-assistant>. MusicTech on the Maono PD500W: <https://musictech.com/news/gear/maono-pd500w/>. *Dream Machine* Issues 33, 35.

62. Suno, “Introducing Spark.” <https://suno.com/blog/introducing-spark>. *Dream Machine* Issue 36.

63. Comfy MCP public beta. <https://comfy.org/mcp/>. *Dream Machine* Issue 36.

64. Epic Games, “State of Unreal 2026: top news from the show.” <https://www.unrealengine.com/news/state-of-unreal-2026-top-news-from-the-show>. Magnific MCP: <https://www.magnific.com/academy/tutorial/connecting-magnific-to-ai-assistants-using-mcp>. Higgsfield MCP: <https://higgsfield.ai/mcp>. Runway MCP: <https://runwayml.com/mcp>. *Dream Machine* Issues 33, 35, 36.

65. Adobe, “Adobe unveils major expansion” — Creative Agent and AI Assistant in beta across Premiere, Photoshop, Illustrator, InDesign and Frame.io. <https://news.adobe.com/news/2026/06/adobe-unveils-major-expansion>. *Dream Machine* Issue 35.

66. Epic Games, “State of Unreal 2026: top news from the show.” <https://www.unrealengine.com/news/state-of-unreal-2026-top-news-from-the-show>. *Dream Machine* Issue 35.

67. Radiance Fields, “Apple Maps Flyover is getting a Gaussian splatting upgrade.” <https://radiancefields.com/apple-maps-flyover-is-getting-a-gaussian-splatting-upgrade>. *Dream Machine* Issue 33.

* * *

Summary

- **Total footnotes:** 774
- **Unique primary-source URLs cited:** 618
- **Source corpus:** 41 editions of *Dream Machine* | *Creative AI* (October 2025 – August 2026)

About the author, and about DreamLab

Pete Woodbridge is a **Creative Technologist and Innovation Director** working across AI, immersive, interactive and experiential tech, real-time production, and next-generation storytelling across **film, games, interactive and immersive experiences**. His cross-sector background spans creative, technical, innovation and product-development roles for software companies, games and immersive studios, and working on new possibilities in storytelling covering film, advertising, brand experiences and animation.

He has worked on projects and new products with partners including **Intel, Google, McLaren, ITV, Sony, Aardman, Red Bull, IBM, BBC, NBC, Crytek, Apple, MediaCityUK, Saatchi & Saatchi** and **Universal** — alongside dozens of independent studios, agencies and SMEs building new products, content pipelines, and interactive and immersive experiences.

Technically, he works across strategy, leadership and hands-on development, including: **real-time interactive and immersive systems** (Unreal Engine, Unity, WebXR, Disguise, Notch, Ventuz); **AI, software and workflow engineering** (Python, ComfyUI, API integrations, automation systems, AI-assisted pipeline design); and **production and future-media** (virtual production, motion capture, ICVFX, real-time previs, AI-augmented workflows). He has built and led teams of engineers and artists across most of those categories at one time or another — which is, more than anything else, the experience the book's arguments about orchestrator-shaped careers and agentic-mesh studios are made out of.

In parallel to running DreamLab he has developed innovation and investment programmes for the creative industries in the North West UK, including the **£3.2 million Innovation Accelerator (Innovate UK) at MediaCity** in Manchester, and the **£7.2 million MusicFutures programme (UKRI/AHRC)** in Liverpool. He is also an external assessor and advisor on several national and international innovation programmes and strategies, including **CoSTAR, EPSRC** and **Digital Catapult**, alongside contributions to a number of steering committees, boards, startups and investment initiatives.

He writes the *Dream Machine* | *Creative AI* newsletter, which began in October 2025 with the launch of Sora 2 and has run weekly through every major event covered in this book. The newsletter is the underlying source corpus of the book; each chapter footnotes back to the specific issue in which a given event was first analysed. He is also the co-host of the *Creative AI* podcast, an active contributor to UK creative-industry policy conversations on AI and copyright, and a frequent speaker on creative technology, AI orchestration, world models, the *Human–AI Agency Continuum* framework, the *Slop Ceiling*, the *Authenticity Premium* and the *Orchestrator economy* — the analytical scaffolding the book is built on.

He lives and works in the North West of England, and is, by his own account, more comfortable writing than talking, and more comfortable shipping than either.

About DreamLab

DreamLab is the collective behind the *Dream Machine* newsletter, the *Creative AI* podcast and this book. The studio was founded in **January 2024**, starting at **MediaCity in Manchester**, and has since grown into a network of **around 50 deep-tech specialists** across labs in **Manchester, Liverpool, Cumbria and London** — and is continuing to spread into new locations. The team is **AI engineers, XR developers, game designers, animators, sound artists, researchers, educators and filmmakers**, including **Emmy winners, PhD researchers and BAFTA-recognised practitioners** working at the cutting edge of creative technology.

The studio does not operate as a fixed agency. It assembles the right team of specialists around each challenge — a model that has, in the period covered by this book, become both the operational pattern of the orchestrator economy described in Chapter 12 and the studio's own default way of working. The four service lines are *Workflows & Automation*, *Creative Services*, *Product Development* and *Consultancy & Training*, deployed across film, TV, games, music, broadcast, immersive and advertising.

In the two and a half years since the studio was founded, DreamLab has supported **100+ businesses**, delivered **35+ collaborative projects**, and helped channel approximately **£2.7m of investment** into partner organisations. Its clients and collaborators include the partner list above — Intel, Google, McLaren, ITV, Sony, Aardman, Red Bull, IBM, BBC, NBC, Crytek, Apple, MediaCityUK, Saatchi & Saatchi, Universal — and span:

- **Global and established brands** — including work for **Verizon Trailblazers** and partnerships with major broadcasters, streamers and platform companies looking for the kind of grounded R&D the studio specialises in.
- **SME partners** — the next wave of ambitious independent studios and small-to-medium creative-technology firms building products with applied AI and creative technology underneath them.
- **Universities, public sector, combined authorities, research councils and innovation agencies** — including funded programmes with **UKRI, Innovate UK, EPSRC, CoSTAR** and **Digital Catapult**, and a long-running innovation partnership at **MediaCity Hub / MediaCityUK** in Salford, where the studio was founded.

DreamLab has been recognised in the **Verizon Trailblazers 500 UK Trailblazers** report, is a participant in the **UKRI-funded Innovate UK**

immersive-tech accelerator, and is a **MediaCity Hub innovation partner** — three structural indicators that the studio sits inside the public-private R&D infrastructure of the UK creative-tech sector rather than on its periphery.

In closed-beta access to platforms like World Labs' **Marble**, in production deployment of agentic stacks built on Claude Code, Skills and the wider open-source ecosystem, and in continuous public writing through the *Dream Machine* newsletter, DreamLab functions as a *working studio that ships* and a *public-record practice* simultaneously — the operating model the book argues will, on the available evidence, define the most defensible shape of creative work for the next decade.

You can find DreamLab at dreamlab.org.uk, the *Dream Machine* newsletter on LinkedIn (and as a podcast on Spotify), and the team page at dreamlab.org.uk/team.

A living book, and the knowledge base behind it

I want to put on the page, at the back of the book where it belongs, the load-bearing operational detail about how this book is built and re-built, because it is — quietly — part of the argument the rest of the manuscript has been making.

Dream Machine is a **living book**. It is not a traditional non-fiction hardback that ships once and then sits frozen for five years while the field it describes moves on. The print and digital editions are updated and re-published on a **rolling cadence** — **currently roughly monthly** — and will continue to be for as long as the *Dream Machine* newsletter ships and the field keeps changing at the rate it currently changes. The unversioned `Dream_Machine.pdf` in the published archive is always the most recent edition. The dated editions — `Dream_Machine_2026-05-21.pdf`, `Dream_Machine_2026-06-04.pdf`, and so on through subsequent months — are preserved alongside as snapshots of the manuscript at specific moments in the longer sequence. A reader who wants the latest argument takes the unversioned file. A reader who wants the specific version a particular claim first appeared in takes the dated one. Both are kept indefinitely; nothing is ever silently overwritten.

Behind the manuscript there is a continuously-rebuilt **knowledge base** inside DreamLab. The corpus underneath the book is structured: each weekly issue of the newsletter is parsed into machine-readable markdown with the original primary-source URLs preserved; each long-form deep dive is checked in as a separately-tracked document; each chapter's footnotes are linked to the issue or document where the claim first appeared; the toolchain inventory of [Chapter 17](#) is regenerated from a categorised data set so that new tools surface in the catalogue automatically when the newsletter covers them; the citation index is auto-built from every chapter's footnotes on every rebuild. The pipeline that re-assembles the book takes the corpus state at the moment of build, runs it through the layout system, produces the dated edition PDF, mirrors the result into the unversioned `Dream_Machine.pdf` alias, and ships both. The whole cycle takes minutes rather than days. The

operational implication is that the next time the newsletter ships an issue, the next time a deep-dive PDF lands in the studio, the next time a chapter's footnote needs to be tightened against a fresh source, the book updates with it — automatically, repeatably, and without requiring a fresh print contract or a new edition number.

The cadence is deliberate, not as a marketing posture but as a quiet demonstration of the book's own argument. The position the manuscript settles on — that AI enables new formats, new publishing models, and new ways of holding a body of knowledge alive in public — is one I think a book about AI in the creative industries has an obligation to *show*, not only to *say*. The traditional book-publishing cycle was an artefact of the industrial-print economy: typesetting cost what it cost, distribution required physical inventory, the unit economics of a hardback were balanced over a five-year shelf life, and so the book froze. The AI-augmented publishing pipeline I have used to build *Dream Machine* makes a different set of trade-offs available. The book can update. The corpus can grow. The argument can sharpen as the field's evidence sharpens. The dated editions preserve the record. The unversioned file ships the present. The trade-offs of the old model are not, in 2026, the only trade-offs available, and the trade-offs of the new model are — for a working-practitioner's record of a moving field — by some margin the better ones.

The point, in a sentence: **a living book is the right shape of a book about a living transition.** That this book has chosen that shape is part of how it makes its argument.

The full corpus — every newsletter issue parsed to markdown, every deep-dive essay, the build pipeline and the chapter source files — is open and available at github.com/petejwoodbridge/DreamMachine-Book. Readers who want to explore the knowledge base, trace the provenance of a specific argument, access earlier editions, or follow the book as it updates will find everything there.

A note on the book's authorship

The arguments in this book belong to the author, but they have been shaped, sharpened and in many places explicitly co-authored with the DreamLab team and the wider community of working creatives who read the newsletter, write back, push back and contribute their own evidence. The Dedication at the front of the book names this debt directly. The organisational-design spine of [Chapter 14](#), and the long-form analysis behind it that sits in DreamLab's published research, are the work of **Dr. John O'Hare** of DreamLab AI Consulting, drawing on Dell'Acqua's Boston Consulting Group experiment, Mollick's *secret cyborgs* coinage, Edmondson and Lei's psychological-safety research and the operational experience of the studio's own agentic-mesh deployments. Several other sections — most notably the *Three Cs* framework in [Chapter 6](#), the *Position Five* analysis in [Chapter 7](#), the *governance-layer-as-IP-environment* thesis in [Chapter 13](#) and the *Programmable Brand*

analysis in DreamLab’s published research — are the product of long studio-internal arguments that the book has, in places, simply written down.

The book is, in that sense, an artefact of the practice it describes. It is the field guide a working studio produces when it tries, in public, to make sense of the most consequential year in creative work since cinema learned to talk.

— Pete Woodbridge — *DreamLab, the North West of England* — June 2026

-
1. Variety, “SAG-AFTRA Condemns Tilly Norwood: AI Actress Is Not an Actor,” 30 September 2025. <https://variety.com/2025/film/news/sag-aftra-tilly-norwood-ai-actress-1236534779/>. See also NBC News, “Tilly Norwood, fully AI ‘actor,’ blasted by actors union SAG-AFTRA for ‘devaluing human artistry’.” <https://www.nbcnews.com/pop-culture/pop-culture-news/tilly-norwood-fully-ai-actor-blasted-actors-union-sag-aftra-devaluing-rcna234685>. Discussed in *Dream Machine* Issue 1 (6 October 2025).
 2. *The Hollywood Reporter*, “U.K. Union Equity Condemns Tilly Norwood: ‘AI Tool, Not a Performer’.” <https://www.hollywoodreporter.com/movies/movie-news/tilly-norwood-ai-actress-uk-union-equity-sag-aftra-debate-1236391739/>. See also Variety, “Tilly Norwood Slammed by Equity as AI Tool, Concerned About Origin.” <https://variety.com/2025/film/global/tilly-norwood-slammed-equity-ai-tool-concerned-origin-1236537042/>. *Dream Machine* Issue 1.
 3. CNN, “Tilly Norwood: Hollywood is fuming over a new ‘AI actress’,” 30 September 2025. <https://www.cnn.com/2025/09/30/tech/hollywood-ai-actor-backlash>.
 4. OpenAI, “Sora 2 is here,” announcement page, 30 September 2025. <https://openai.com/index/sora-2/>. The model launched alongside an invite-only iOS app of the same name in the U.S. and Canada. *Dream Machine* Issue 1 carried the launch alongside contemporaneous coverage from NBC News and *The Guardian* on the model’s first copyright and safety incidents.
 5. *Dream Machine* | *Creative AI*, LinkedIn newsletter, archive of Issues 1–41, October 2025 – August 2026. <https://www.linkedin.com/newsletters/dream-machine-creative-ai-7379776527871381505/>.
 6. DreamLab AI Collective, team page. <https://dreamlab-ai.com/team>. Referenced from *Dream Machine* Issue 16 onward.
 7. Charles Cecil (Revolution Software, *Broken Sword*) quoted in *gamesindustry.biz*, “‘AI was an expensive mistake’: Charles Cecil on innovation, insolvency, and Broken Sword.” <https://www.gamesindustry.biz/ai-was-an-expensive-mistake-charles-cecil-on-innovation-insolvency-and-broken-sword>. *Dream Machine* Issue 3.
 8. Adobe, “Inaugural Adobe Creators’ Toolkit Report: 86 Percent of Global Creators Use Creative Generative AI.” <https://news.adobe.com/news/2025/10/adobe-max-2025-creators-survey>. Survey of 16,000 creators across the U.S., U.K., France, Germany, South Korea, Japan, India and Australia, released at Adobe MAX 2025. *Dream Machine* Issue 6.
 9. UK Department for Science, Innovation and Technology (DSIT), *Statement of Progress on Copyright and AI*, December 2025. <https://www.gov.uk/government/publications/copyright-and-artificial-intelligence-progress-report/copyright-and-artificial-intelligence-statement-of-progress-under-section-137-data-use-and-access-act>. See also IPWatchdog, “Respondents to UK AI Consultation Overwhelmingly Want AI Companies to License Copyrighted Works in All Cases.” <https://ipwatchdog.com/2025/12/16/respondents-uk-ai-consultation-overwhelmingly-want-ai-companies-license-copyrighted-works-all-cases/>. *Dream Machine* Issue 12.
 10. *Dream Machine* Issue 5, “Adobe’s Latest AI Announcements — Is every tool going AI?”, 31 October 2025. <https://www.linkedin.com/pulse/dream-machine-creative-ai-news-insight-oct-25->

[issue-5-woodbridge-f7jnc/](#).

11. Adobe, Adobe MAX 2025 keynote messaging, October 2025. Coverage: Creative Boom, “Adobe is putting AI in everything everywhere all at once.” <https://www.creativeboom.com/news/adobe-is-putting-ai-in-everything-everywhere-all-at-once/>. *Dream Machine* Issue 5.
12. World Labs, *Marble* — first commercial spatial-AI world model, public launch November 2025. <https://marble.worldlabs.ai/>. Technical context: TechCrunch, “Fei-Fei Li’s World Labs speeds up the world model race with Marble, its first commercial product.” <https://techcrunch.com/2025/11/12/fei-fei-lis-world-labs-speeds-up-the-world-model-race-with-marble-its-first-commercial-product/>. DreamLab participated in the closed beta during October–November 2025. *Dream Machine* Issue 7.
13. 11,514 responses across the Citizen Space portal and email, of which 10,112 came through Citizen Space; 88% of those supported licensing as a default rule, against 3% who supported the government’s preferred opt-out model. UK DSIT, *Statement of Progress*, December 2025; analysis in *Dream Machine* Issue 12 (18 December 2025). Final report and economic impact assessment to be laid before Parliament by 18 March 2026.
14. *Digital Music News*, “Nearly 800 Creatives, Including Jason Aldean and One Republic, Sign Responsible AI Declaration — ‘Stealing Our Work Is Not Innovation’.” <https://www.digitalmusicnews.com/2026/01/22/stealing-isnt-innovation/>. *Dream Machine* Issue 16.
15. For a contemporaneous overview of the AI video model release cadence through 2024 and 2025, see *Dream Machine* Issues 1–8 (October–November 2025), which logged near-weekly releases from Runway, Luma, Pika, Kling, Veo, Wan, Higgsfield, Hunyuan and a long tail of smaller labs.
16. *The Hollywood Reporter*, “AI Performer Tilly Norwood Sparks Hollywood Backlash.” <https://www.hollywoodreporter.com/movies/movie-news/tilly-norwood-ai-actress-uk-union-equity-sag-aftra-debate-1236391739/>. *Dream Machine* Issue 1.
17. SAG-AFTRA statement, 30 September 2025, reported in *Variety*, “SAG-AFTRA Condemns Tilly Norwood: AI Actress Is Not an Actor.” <https://variety.com/2025/film/news/sag-aftra-tilly-norwood-ai-actress-1236534779/>.
18. OpenAI, “Sora 2 is here,” 30 September 2025. <https://openai.com/index/sora-2/>. *Dream Machine* Issue 1.
19. Particle6 background and Van der Velden interview: *The Hollywood Reporter*, “Meet the Creator of the AI Actress Hollywood Loves to Hate: ‘You’re Gonna See a Lot of Tilly Norwood Next Year’.” <https://www.hollywoodreporter.com/movies/movie-features/tilly-norwood-creator-particle6-elina-van-der-velden-talks-1236428824/>. *Dream Machine* Issue 8.
20. *Deadline*, “Tilly Norwood Creator Eline Van Der Velden Talks Backlash, Reveals Another 40 AI Actors Are In The Pipeline.” <https://deadline.com/2025/11/tilly-norwood-creator-interview-backlash-more-ai-actors-coming-1236601334/>.
21. Northeastern Global News, “Why AI ‘Actress’ Tilly Norwood Has Hollywood Angry.” <https://news.northeastern.edu/2025/10/02/ai-actress-tilly-norwood-hollywood-backlash/>.
22. SAG-AFTRA, official statement reproduced in *Variety*, *op. cit.*; also NBC News, “Tilly Norwood, fully AI ‘actor,’ blasted by actors union SAG-AFTRA for ‘devaluing human artistry’.” <https://www.nbcnews.com/pop-culture/pop-culture-news/tilly-norwood-fully-ai-actor-blasted-actors-union-sag-aftra-devaluing-rcna234685>.
23. Equity (U.K.), statement of 2 October 2025: *Variety*, “Tilly Norwood Slammed by Equity as AI Tool, Concerned About Origin.” <https://variety.com/2025/film/global/tilly-norwood-slammed-equity-ai-tool-concerned-origin-1236537042/>.
24. CNN, “Tilly Norwood: Hollywood is fuming over a new ‘AI actress’.” <https://www.cnn.com/2025/09/30/tech/hollywood-ai-actor-backlash>.

25. *Deadline*, “AI ‘Actor’ Tilly Norwood To Star In Feature Film ‘Misaligned.’” <https://deadline.com/2026/07/tilly-norwood-ai-actor-misaligned-1236974639/>. *Dream Machine* Issue 37.
26. OpenAI, “Sora 2 is here,” <https://openai.com/index/sora-2/>. Technical capabilities summary including physics modelling, multi-shot world-state persistence and synchronised audio.
27. *Dream Machine* Issue 1, “Editor’s Pick”; further launch context in NBC News, “OpenAI’s Sora 2: a major leap in AI video and audio.” <https://www.nbcnews.com/tech/tech-news/openai-sora-2-app-video-chatgpt-creation-rcna234973>.
28. LinkedIn News aggregation: “Sora Tops 1 Million Downloads in 5 Days.” <https://www.linkedin.com/news/story/sora-tops-1m-downloads-in-5-days-6684088/>. *Dream Machine* Issue 3. A historical note for the reader: the Sora consumer app and the broader Sora platform did not survive 2026. OpenAI announced a staged wind-down of Sora on **24 March 2026** — consumer app and website ending late April, developer API closing late September — with peak revenue reportedly at around \$540,000 per month against operating costs in the billions, and the compute reallocated toward higher-margin coding and enterprise products. The Disney–OpenAI licensing announcement of 11 December 2025 (referenced in Chapter 7 and Appendix C) was never executed; no formal agreement was signed and no money changed hands. The contractual framework underneath that deal — IP-not-training, talent-likeness excluded, jointly-governed steering committee, equity-and-warrants structure — survived the platform’s collapse and is now the reference architecture other studios negotiate against. The argument the rest of this chapter is making about the Sora 2 launch week — the cultural rupture, the audience arrival in the AI debate, the parasocial-character / working-actor market collision — survives the platform’s death intact; the *moment* matters whether or not the *vendor* persists.
29. *The Guardian*, “OpenAI Sora 2 violence racism.” <https://www.theguardian.com/us-news/2025/oct/04/openai-sora-violence-racism>. *Dream Machine* Issue 1.
30. NBC News, *op. cit.*; *The Guardian*, *op. cit.*
31. *Digital Music News*, “OpenAI’s Sora 2 includes likeness protections for celebrities who don’t opt in, but that doesn’t apply to ‘historical figures’ and dead celebrities.” <https://www.digitalmusicnews.com/2025/10/08/openais-likeness-protections-dont-apply-to-dead-celebrities/>. *Dream Machine* Issue 2.
32. Quoted in *The Guardian*, “OpenAI launch of video app Sora plagued by violent and racist images: ‘The guardrails are not real.’” <https://www.theguardian.com/us-news/2025/oct/04/openai-sora-violence-racism>. *Dream Machine* Issue 1.
33. Google DeepMind, Veo 3.1 launch, mid-October 2025. *Dream Machine* Issue 3, “Editor’s Pick: Veo 3.1 and the Rise of AI Filmmaking.” Coverage: <https://www.cometapi.com/veo-3-1-is-comingand-whats-rumor/>.
34. WUFT, “Kiss reality goodbye: AI-generated social media has arrived,” 3 October 2025. <https://www.wuft.org/2025-10-03/kiss-reality-goodbye-ai-generated-social-media-has-arrived>. *Dream Machine* Issue 1.
35. *No Film School*, “James Cameron Says AI Is ‘Never Going to Take the Place’ of Humans.” <https://nofilmschool.com/james-cameron-ai#>. *Dream Machine* Issue 1.
36. *The Guardian*, “James Cameron says AI actors are ‘horrifying to me,’” 1 December 2025. <https://www.theguardian.com/film/2025/dec/01/james-cameron-says-ai-actors-are-horrifying-to-me>. Original quote from CBS *Sunday Morning*. *Dream Machine* Issue 10.
37. Variety, “James Cameron Says It’s ‘Horrifying’ that AI Can ‘Make Up an Actor.’” <https://variety.com/2025/film/news/james-cameron-horrifying-ai-replace-actors-1236595864/>.
38. Stability AI, board composition, 2024–2026. Reported across multiple outlets including *Deadline*, “James Cameron Calls AI Replacing Actors ‘Horrifying’; Art ‘Sacred’.”

<https://deadline.com/2025/11/james-cameron-gen-ai-horrifying-human-art-sacred-avatar-1236631387/>.

39. Deezer, “AI-generated tracks now represent 44% of all new uploaded music,” April 2026 newsroom release. <https://newsroom-deezer.com/2026/04/ai-generated-tracks-represent-44-of-new-uploaded-music/>. Companion analysis: *Music Business Worldwide*, “75,000 AI-generated tracks now flood Deezer daily, representing 44% of all new music uploaded to the platform.” <https://www.musicbusinessworldwide.com/75000-ai-generated-tracks-now-flood-deezer-daily-representing-44-of-all-new-music-uploaded-to-the-platform-says-streamer/>. Daily AI uploads to Deezer rose from approximately 50,000 per day in November 2025 (*Dream Machine* Issue 7, citing Deezer / *Musically*) to 75,000 per day by April 2026, with consumer streams of fully-AI tracks holding between 1% and 3% of total platform plays — and up to 85% of those streams identified as fraudulent in 2025. *Dream Machine* Issues 7, 26, 27, 28.
40. John Philip Sousa, “The Menace of Mechanical Music,” *Appleton’s Magazine*, Vol. 8, September 1906, pp. 278–284. Full text via ExplorePAHistory: <https://explorepahistory.com/odocument.php?docId=1-4-1A1>. Academic context: Patrick Warfield, “John Philip Sousa and ‘The Menace of Mechanical Music,’” *Journal of the Society for American Music*, Cambridge University Press: <https://www.cambridge.org/core/journals/journal-of-the-society-for-american-music/article/abs/john-philip-sousa-and-the-menace-of-mechanical-music/A9E621587BE7580ABD73AEF64D4B2DC8>. The 1906 essay was, in part, lobbying for what would become the 1909 Copyright Act.
41. Sousa, *op. cit.* The Library of Congress’s “Sousa and the Talking Machine” essay is a useful institutional summary: <https://blogs.loc.gov/now-see-hear/2020/05/sousa-and-the-talking-machine/>.
42. William Henry Cardinal O’Connell, Archbishop of Boston, sermon to the Holy Name Society, Boston, 10 January 1932. Reported widely in the contemporaneous press, including the *Daily Courier* (Connellsville, PA), 12 January 1932 (<https://www.newspapers.com/newspage/38168082/>). Cultural context: KUOW/NPR, “Imbecile Slush’: Surprising Early Reactions to Crooning,” <https://www.kuow.org/stories/imbecile-slush-surprising-early-reactions-crooning>. JSTOR Daily, “The Gender Politics of the First Boy Bands,” <https://daily.jstor.org/the-gender-politics-of-the-first-boy-bands/>.
43. *Grand Upright Music, Ltd. v. Warner Bros. Records Inc.*, 780 F. Supp. 182 (S.D.N.Y. 1991). Full text: <https://law.justia.com/cases/federal/district-courts/FSupp/780/182/1445286/>. The “Thou shalt not steal” opening is the most-quoted line from a US copyright opinion of the late twentieth century.
44. Tippet’s account of the Jurassic Park digital test is documented across multiple ASC and contemporaneous press accounts. American Society of Cinematographers, “Jurassic Park: Effects Team Brings Dinosaurs Back from Extinction,” <https://theasc.com/articles/jurassic-park-effects-team-brings-dinosaurs-back>. Wikipedia, “Phil Tippet,” https://en.wikipedia.org/wiki/Phil_Tippett. The dialogue paraphrase Spielberg incorporated into the film is Goldblum/Malcolm’s response to Grant’s “I think we’re out of a job”: “Don’t you mean extinct?”
45. *Variety*, “George Lucas Says AI Makes Filmmaking ‘Easier’ and ‘There’s Nothing You Can Do About It.’” <https://variety.com/2026/film/news/george-lucas-test-screenings-embraces-ai-hollywood-1236810656/>; AOL, “George Lucas likens AI sceptics to luddites clinging to horses and carts.” <https://www.aol.co.uk/articles/george-lucas-likens-ai-sceptics-122913000.html>. *Dream Machine* Issue 38.
46. Charles Baudelaire, “*Le Public Moderne et la Photographie*,” *Revue Française*, 1859 (part of the *Salon de 1859* essays). English translation widely available; the original French in PDF form: https://gallowayexploringart.wordpress.com/wp-content/uploads/2014/08/ baudelaire_the-

[modern-public-photography.pdf](#). Smithsonian Archives institutional overview: “*Photography Murdered Painting, Right?*”, <https://siarchives.si.edu/blog/photography-murdered-painting-right>.

47. The Delaroche apocrypha is documented in Quote Investigator: <https://quoteinvestigator.com/2022/10/16/photo-mortal/>. The earliest sourced version is in an 1873 survey, 34 years after Delaroche reportedly said it. Delaroche’s own contemporary writing on the daguerreotype, in Gernsheim’s standard 1959 monograph, characterised the new technology as “an immense service to the arts.”
48. The 1942–44 Petrillo strike: Wikipedia, “1942–44 musicians’ strike,” [https://en.wikipedia.org/wiki/1942%E2%80%9944_musicians’ strike](https://en.wikipedia.org/wiki/1942%E2%80%9944_musicians%27_strike); Mainspring Press, “*The Man Who Crippled the American Recording Industry: James Caesar Petrillo and the American Federation of Musicians Recording Bans*,” <https://mainspringpress.org/2024/11/23/the-man-who-crippled-the-recording-industry-james-caesar-petrillo-and-the-american-federation-of-musicians-recording-bans/>; DownBeat, “*The Petrillo Ban of 1942–44: Past & Future at War*,” <https://downbeat.com/news/detail/the-petrillo-ban-of-194244-past-future-at-war>; Local 802 AFM, “*The Silence Was Deafening*,” <https://www.local802afm.org/allegro/articles/the-silence-was-deafening/>. The Music Performance Trust Fund’s institutional history: <https://musicpf.org/establishment-of-mptf-led-to-the-formation-of-afms-pension-and-residual-funds/>.
49. William Henry Cardinal O’Connell, Archbishop of Boston, sermon to the Holy Name Society, Boston, 10 January 1932. Reported widely in the contemporaneous press, including the *Daily Courier* (Connellsville, PA), 12 January 1932 (<https://www.newspapers.com/newspage/38168082/>). Cultural context: KUOW/NPR, “*Imbecile Slush’: Surprising Early Reactions to Crooning*,” <https://www.kuow.org/stories/imbecile-slush-surprising-early-reactions-crooning>. JSTOR Daily, “*The Gender Politics of the First Boy Bands*,” <https://daily.jstor.org/the-gender-politics-of-the-first-boy-bands/>.
50. Musicians’ Union History, “*The Strike That Made History — Massacre of the Musicians 1980*,” <https://www.muhistory.com/the-massacre-of-the-musicians-1980/>. Academic context on the broader MU–BBC dispute landscape: “*Negotiating Needletime*” (Tandfonline), <https://www.tandfonline.com/doi/full/10.1080/03071022.2016.1215098>.
51. MusicRadar, “*The Day the Loony Musicians Union Tried to Kill the Synthesizer (Which Also Happened to be Bob Moog’s Birthday)*,” <https://www.musicradar.com/news/the-union-passed-a-motion-to-ban-the-use-of-synths-drum-machines-and-any-electronic-devices-the-day-the-loony-musicians-union-tried-to-kill-the-synthesizer-which-also-happened-to-be-bob-moogs-birthday>. Far Out Magazine, “*Why did the Musicians Union outlaw synthesisers in 1982?*,” <https://faroutmagazine.co.uk/musicians-union-outlaw-synthesisers/>.
52. *Grand Upright Music, Ltd. v. Warner Bros. Records Inc.*, 780 F. Supp. 182 (S.D.N.Y. 1991). Full text: <https://law.justia.com/cases/federal/district-courts/FSupp/780/182/1445286/>. The “Thou shalt not steal” opening is the most-quoted line from a US copyright opinion of the late twentieth century.
53. *Bridgeport Music, Inc. v. Dimension Films*, 410 F.3d 792 (6th Cir. 2005). Full text: <https://law.justia.com/cases/federal/appellate-courts/F3/410/792/574458/>. The “Get a licence or do not sample” rule is the most-cited line in the opinion.
54. *TIME*, “50 Worst Inventions,” 2010, Auto-Tune at #15: https://content.time.com/time/specials/packages/article/0,28804,1991915_1991909_1991903,0_0.html. Wikipedia, “Auto-Tune,” <https://en.wikipedia.org/wiki/Auto-Tune>. NPR, “25 Years of Believe,” <https://www.npr.org/2023/10/19/1207028349/25-years-ago-cher-released-a-song-that-would-change-the-sound-of-pop-music>. Wikipedia, “D.O.A. (Death of Auto-Tune),” [https://en.wikipedia.org/wiki/D.O.A. \(Death of Auto-Tune\)](https://en.wikipedia.org/wiki/D.O.A._(Death_of_Auto-Tune)).

55. Walter Murch, *In the Blink of an Eye: A Perspective on Film Editing*, Silman-James Press, 1995 (2nd edition 2001). PDF: [https://www.craftfilmschool.com/userfiles/files/Walter%20Murch%20%20In%20the%20Blink%20of%20an%20Eye%20Revised%202nd%20Edition%20\(2001,%20Silman-James%20Pr\).pdf](https://www.craftfilmschool.com/userfiles/files/Walter%20Murch%20%20In%20the%20Blink%20of%20an%20Eye%20Revised%202nd%20Edition%20(2001,%20Silman-James%20Pr).pdf). Charles Koppelman, *Behind the Seen: How Walter Murch Edited Cold Mountain Using Apple's Final Cut Pro and What This Means for Cinema*, Peachpit Press, 2004: <https://www.peachpit.com/store/behind-the-seen-how-walter-murch-edited-cold-mountain-9780735714267>.
56. Sasson's account documented at the National Inventors Hall of Fame: <https://www.invent.org/blog/inventors/Legacy-Steve-Sasson>. Snopes verification of the "Kodak suppressed the digital camera" claim: <https://www.snopes.com/fact-check/kodak-digital-camera-invention/>. Knowledge@Wharton on the Kodak collapse: <https://knowledge.wharton.upenn.edu/podcast/knowledge-at-wharton-podcast/whats-wrong-with-this-picture-kodaks-30-year-slide-into-bankruptcy/>. Bankruptcy filing: 19 January 2012, S.D.N.Y., \$5.1bn assets / \$6.8bn liabilities.
57. Wikipedia, "Brian Walski," https://en.wikipedia.org/wiki/Brian_Walski. *Washington Post* contemporaneous coverage: <https://www.washingtonpost.com/archive/lifestyle/2003/04/03/alterd-picture-costs-la-times-photographer-his-job/c5e7c9e0-a836-429a-bb4e-d502fi768a96/>. World Press Photo's institutional response in TIME: <https://time.com/3706626/world-press-photo-processing-manipulation-disqualified/>.
58. Wikipedia, "Viacom International, Inc. v. YouTube, Inc.," https://en.wikipedia.org/wiki/Viacom_International_Inc._v._YouTube_Inc._ Electronic Frontier Foundation case file: <https://www.eff.org/cases/viacom-v-youtube>. Variety on the March 2014 settlement: <https://variety.com/2014/biz/news/google-and-viacom-settle-copyright-infringement-lawsuit-over-youtube-1201137538/>.
59. PetaPixel, "The Rise and Crash of the Camera Industry in One Chart," <https://petapixel.com/2024/08/22/the-rise-and-crash-of-the-camera-industry-in-one-chart/>. Statista, "Smartphones Wipe Out Decades of Camera Industry Growth," <https://www.statista.com/chart/15524/worldwide-camera-shipments/>. CIPA shipment data series, multiple years.
60. CNN Business, "Meet the translation professionals losing their jobs to AI," January 2026, <https://www.cnn.com/2026/01/23/tech/translation-language-jobs-ai-automation-intl>. Carl Benedikt Frey (Oxford Martin School), 2025 study on translator employment across 696 US labour markets. American Translators Association industry position: <https://www.atanet.org/client-assistance/blog-machine-translation-vs-human-translation/>. Wikipedia, "Google Neural Machine Translation," https://en.wikipedia.org/wiki/Google_Neural_Machine_Translation.
61. *Dream Machine* Issue 2, "Editor's Pick," 10 October 2025. <https://www.linkedin.com/pulse/dream-machine-creative-ai-news-insight-oct-25-2-pete-woodbridge-mnrjc/>.
62. OpenAI, "Introducing AgentKit," 6 October 2025. <https://openai.com/index/introducing-agentkit/>.
63. TechCrunch, "OpenAI launches AgentKit to help developers build and ship AI agents," 6 October 2025. <https://techcrunch.com/2025/10/06/openai-launches-agentkit-to-help-developers-build-and-ship-ai-agents/>. Also coverage at *InfoQ*, "OpenAI Dev Day 2025 Introduces GPT-5 Pro API, Agent Kit, and More." <https://www.infoq.com/news/2025/10/openai-dev-day/>.
64. *Dream Machine* Issue 2: "Agentic AI — the class of AI systems that can plan, act, and pursue goals with autonomy — promises a new era of collaboration in creative industries... Its another step along the Human-AI Agency Continuum." See also *TVB Europe*, "Is Agentic AI About to Change

- the Media and Entertainment Industry?” <https://www.tvbeurope.com/artificial-intelligence/opinion-is-agentic-ai-about-to-change-the-media-and-entertainment-industry>.
65. Google DeepMind, Veo 3.1 release, October 2025. *Dream Machine* Issue 3.
 66. *MusicTech*, “iZotope Ozone 12’s AI assistant is cool, but the Stem EQ is the real star.” <https://musictech.com/reviews/plugin-ins/izotope-ozone-12-review/>. *Dream Machine* Issue 3.
 67. Adobe, “Inaugural Adobe Creators’ Toolkit Report,” October 2025. <https://news.adobe.com/news/2025/10/adobe-max-2025-creators-survey>. Survey of 16,000 creators across eight countries, released at Adobe MAX 2025. *Dream Machine* Issue 6.
 68. Adobe, *op. cit.* The same survey: 86% of creators use creative generative AI; 76% say it has helped grow their business or brand; 81% say AI lets them make content they otherwise couldn’t have made; 69% worry about their work being used to train AI without consent; 70% are optimistic about agentic AI; 85% would use AI that learns their creative style.
 69. Mureka, “Music Agent Studio” launch, mid-October 2025. *Dream Machine* Issue 4. https://www.linkedin.com/posts/sherrihendrickson_mureka-unveils-music-agent-studio-and-enhanced-share-738499251526864896-cNYg/.
 70. *Finsmes*, “AdsGency Raises \$12M in Seed Funding,” October 2025. <https://www.finsmes.com/2025/10/adsgency-raises-12m-in-seed-funding.html>. *Dream Machine* Issue 4.
 71. *Musically*, “Meet Lenny, an AI agent to help organisers of live music events.” <https://musically.com/2025/10/20/meet-lenny-an-ai-agent-to-help-organisers-of-live-music-events/>. *Dream Machine* Issue 4.
 72. *GamesRadar*, “Even under USD20 million in debt, EA reportedly pushes 15,000 employees to use AI as a ‘thought partner’ for everything from character art to playtesting.” <https://www.gamesradar.com/games/even-under-usd20-million-in-debt-ea-reportedly-pushes-15-000-employees-to-use-ai-as-a-thought-partner-for-everything-from-character-art-to-playtesting/>. *Dream Machine* Issue 6.
 73. PYMNTS, “Adobe Lets Users Design and Edit Using ChatGPT.” <https://www.pymnts.com/artificial-intelligence-2/2025/adobe-lets-users-design-and-edit-using-chatgpt/>. Adobe blog: “Edit images, designs, and PDFs right inside ChatGPT — thanks to Adobe Express, Photoshop, and Acrobat.” <https://blog.adobe.com/en/publish/2025/12/10/edit-photoshop-chatgpt>. *Dream Machine* Issue 12.
 74. TechCrunch, “Anthropic launches interactive Claude apps, including Slack and other workplace tools,” 26 January 2026. <https://techcrunch.com/2026/01/26/anthropic-launches-interactive-claude-apps-including-slack-and-other-workplace-tools/>. *Heygen Video Agent*: https://www.linkedin.com/posts/heygen_introducing-the-new-video-agent-activity-7421597801240801282-d1CF. *Dream Machine* Issue 16.
 75. *Dream Machine* Issue 21, “Editor’s Pick: Adobe and NVIDIA Just Raised the Stakes for Creative AI,” 19 March 2026.
 76. Adobe Summit 2026, “Agentic Creative Intelligence” keynote framing. *Dream Machine* Issue 26.
 77. *Dream Machine* Issue 29, May 2026, citing Sony’s adoption of Claude Code studios with multi-agent coordination.
 78. Anthropic, public statements on agent deployment patterns through Q1 2026. Cf. *Dream Machine* Issues 11, 16, 22.
 79. *gamesindustry.biz*, “‘AI was an expensive mistake’: Charles Cecil on innovation, insolvency, and Broken Sword.” <https://www.gamesindustry.biz/ai-was-an-expensive-mistake-charles-cecil-on-innovation-insolvency-and-broken-sword>. *Dream Machine* Issue 3.
 30. *Niche Gamer*, “Larian Studios backs off from gen AI, says tech won’t be used in new Divinity.” <https://nichegamer.com/larian-studios-backs-off-from-gen-ai/>. *Dream Machine* Issue 14.

81. Decrypt, “Warhammer 40,000’ Maker Games Workshop Rules Out Generative AI.” <https://decrypt.co/354482/warhammer-40000-maker-games-workshop-rules-out-generative-ai>. *Dream Machine* Issue 14.
82. Niche Gamer, “Manor Lords publisher Hooded Horse won’t work with devs using gen AI.” <https://nichegamer.com/manor-lords-publisher-hooded-horse-wont-work-with-devs-using-gen-ai/>. *Dream Machine* Issue 14.
83. gamesindustry.biz, “RuneScape maker Jagex says it will never use generative AI to make in-game content.” <https://www.gamesindustry.biz/runescape-maker-jagex-says-it-will-never-use-generative-ai-to-make-in-game-content>. *Dream Machine* Issue 16.
84. Imperva, 2025 *Bad Bot Report: How AI is Supercharging the Bot Threat*. <https://www.imperva.com/blog/2025-imperva-bad-bot-report-how-ai-is-supercharging-the-bot-threat/>. *Dream Machine* Issue 4.
85. Cloudflare, “The crawl-to-click gap: Cloudflare data on AI bots, training, and referrals.” <https://blog.cloudflare.com/crawlers-click-ai-bots-training/>. *Dream Machine* Issue 4. Later 2025 updates show training crawlers declining from ~90% to ~74% of AI bot activity as scraper bots rose to 24% and a new “agentic” category emerged at 1.7%; see Cloudflare, “A deeper look at AI crawlers: breaking down traffic by purpose and industry.” <https://blog.cloudflare.com/ai-crawler-traffic-by-purpose-and-industry/>.
86. Grand View Research, “Generative AI Content Creation Market Report.” <https://www.grandviewresearch.com/industry-analysis/generative-ai-content-creation-market-report>. *Dream Machine* Issue 4 also cites Gartner and Europol forecasts of 90–99% AI-generated or AI-assisted online content by 2030.
87. *Dream Machine* Issue 4, “Editor’s Pick: Is the Internet Dead Yet?” 23 October 2025. <https://www.linkedin.com/pulse/dream-machine-creative-ai-news-insight-oct-25-issue-4-woodbridge-hzttc/>.
88. Wikipedia, *Dead Internet Theory*. https://en.wikipedia.org/wiki/Dead_Internet_theory. *Dream Machine* Issue 4.
89. Graphite, 2025 analysis of new web content by author type (human vs. AI vs. AI-assisted). Cited in *Dream Machine* Issue 4.
90. For “model collapse” as a term of art, see Ilia Shumailov et al., “The Curse of Recursion: Training on Generated Data Makes Models Forget” (2024), and subsequent literature.
91. Futurism, “Researchers built a social network with only AI agents — within hours it had collapsed into warring tribes.” <https://futurism.com/social-network-ai-intervention-echo-chamber>. *Dream Machine* Issue 4.
92. *Digital Music News*, “Instagram Chief Says We Should ‘Fingerprint Real Media’ Instead of Tracking and Disclosing AI Slop.” <https://www.digitalmusicnews.com/2026/01/05/instagram-chief-ai-slop-comments/>. See also *WebProNews*, “Instagram Head Warns AI Images Erode Trust, Calls for Verification Standards.” <https://www.webpronews.com/instagram-head-warns-ai-images-erode-trust-calls-for-verification-standards/>. *Dream Machine* Issue 13.
93. *The Guardian*, “Brands using AI-generated influencers to promote products on social media.” <https://www.theguardian.com/technology/2026/jun/21/brands-using-ai-generated-influencers-to-promote-products-on-social-media>. *Dream Machine* Issue 35.
94. *Social Media Today*, “TikTok bans AI-generated voices in shopping livestreams.” <https://www.socialmediatoday.com/news/tiktok-bans-ai-generated-voices-in-shopping-livestreams/822977/>. *Dream Machine* Issue 35.
95. *Search Engine Journal*, “TikTok Targets AI-Generated Spam Accounts In High-Risk Topics.” <https://www.searchenginejournal.com/tiktok-targets-ai-generated-spam-accounts-in-high-risk-topics/582256/>. *Dream Machine* Issue 38.

96. Digit, "YouTube, X funnel millions of visits to AI deepfake 'nudeify' apps, new study shows: Key findings." <https://www.digit.in/news/general/millions-of-users-reached-ai-nudeify-sites-through-youtube-and-x-report.html>. *Dream Machine* Issue 38.
97. Sundance Institute, "Centering the Artist: Why We're Launching the AI Literacy Initiative." <https://www.sundance.org/blogs/centering-the-artist-why-were-launching-the-ai-literacy-initiative/>. *Dream Machine* Issue 16.
98. Stereogum, "Bandcamp bans AI music." <https://stereogum.com/2485199/bandcamp-bans-ai-music/news>. *Dream Machine* Issue 14.
99. CNET, "San Diego Comic-Con Draws a Line: No AI Art Allowed at 2026 Event." <https://www.cnet.com/culture/san-diego-comic-con-bans-ai-art-for-2026-event/>. *Dream Machine* Issue 16.
100. Deezer, "AI-generated tracks now represent 44% of all new uploaded music," April 2026. <https://newsroom-deezer.com/2026/04/ai-generated-tracks-represent-44-of-new-uploaded-music/>. *Music Business Worldwide*, "75,000 AI-generated tracks now flood Deezer daily." <https://www.musicbusinessworldwide.com/75000-ai-generated-tracks-now-flood-deezer-daily-representing-44-of-all-new-music-uploaded-to-the-platform-says-streamer/>. *Dream Machine* Issues 7, 26, 27, 28.
01. *The Hollywood Reporter*, "'Synthetic Sincerity' by Marc Isaacs Explores if AI Characters Can Be Taught Authenticity: IDFA." <https://www.hollywoodreporter.com/movies/movie-news/synthetic-sincerity-film-idfa-ai-authenticity-interview-1236426180/>. *Dream Machine* Issue 8.
02. Variety, "AI-Generated Images Threaten Future of Documentary as People 'Will Stop Believing Anything'." <https://variety.com/2025/film/festivals/ai-generated-images-threaten-future-of-documentary-1236583466/>. *Dream Machine* Issue 8.
03. PR Newswire, "From Apple TV Creative to AI Filmmaker: Hoyt Dwyer's Animated Film To Compete at AI FilmFest Japan 2025." <https://www.prnewswire.com/news-releases/from-apple-tv-creative-to-ai-filmmaker-hoyt-dwyers-animated-film-to-compete-at-ai-filmfest-japan-2025-302598064.html>. *Dream Machine* Issue 6.
04. Variety, "AI Creator Behind Viral 'Deadpool,' 'Harry Potter' Christmas Clip Made His Film in a Ukrainian Bomb Shelter." <https://variety.com/2026/digital/news/ai-video-deadpool-harry-potter-andrii-daniels-1236624632/>. *Dream Machine* Issue 16.
05. *Branding in Asia*, "It's the Most Terrible Time of the Year' — McDonald's Netherlands' Wonderfully Chaotic, AI-Driven Christmas Film." <https://www.brandinginasia.com/its-the-most-terrible-time-of-the-year-mcdonalds-netherlands-wonderfully-chaotic-ai-driven-christmas-film/>. Pulled following backlash: *SiliconAngle*, "Not ready: McDonald's AI-generated ad taken down after public backlash." <https://siliconangle.com/2025/12/10/not-ready-mcdonalds-ai-generated-ad-taken-public-backlash/>. *Dream Machine* Issue 11.
06. BBC News, "Fashion house Valentino criticised over 'disturbing' AI handbag ads." <https://www.bbc.co.uk/news/articles/cwvyjvyn83go>. *Dream Machine* Issue 10.
07. *Adweek*, "Coca-Cola Uses AI to Rekindle the Magic of Its Holiday Ads." <https://www.adweek.com/creativity/coca-cola-uses-ai-to-rekindle-the-magic-of-its-holiday-ads/>. *Dream Machine* Issue 6.
08. *AI News*, "AI causes reduction in users' brain activity, MIT." <https://www.artificialintelligence-news.com/news/ai-causes-reduction-in-users-brain-activity-mit/>. *Dream Machine* Issue 1.
09. For "model collapse" as a term of art, see Ilia Shumailov et al., "The Curse of Recursion: Training on Generated Data Makes Models Forget" (2024), and subsequent literature.
10. *Digital Music News*, "Instagram Chief Says We Should 'Fingerprint Real Media' Instead of Tracking and Disclosing AI Slop." <https://www.digitalmusicnews.com/2026/01/05/instagram-chief-ai-slop-comments/>. See also *WebProNews*, "Instagram Head Warns AI Images Erode Trust,

- Calls for Verification Standards.” <https://www.webpronews.com/instagram-head-warns-ai-images-erode-trust-calls-for-verification-standards/>. *Dream Machine* Issue 13.
11. Graphite, 2025 analysis of new web content by author type (human vs. AI vs. AI-assisted). Cited in *Dream Machine* Issue 4.
 12. ElevenLabs, “Detecting audio generated by ElevenLabs with SynthID” (partnership with Google DeepMind). <https://elevenlabs.io/blog/synthid>. *Dream Machine* Issue 36.
 13. *Digital Music News*, “TikTok Quietly Starts Testing AI Deepfake Detection Tech.” <https://www.digitalmusicnews.com/2026/07/20/tiktok-starts-testing-ai-deepfake-detection-tech/>. *Dream Machine* Issue 39.
 14. *Social Media Today*, “YouTube adds AI-powered comment filtering.” <https://www.socialmediatoday.com/news/youtube-adds-ai-powered-comment-filtering/823939/>. *Dream Machine* Issue 36.
 15. *Business Today*, “‘We can fix the internet ...’ says World Wide Web inventor Tim Berners-Lee.” <https://www.businesstoday.in/technology/news/story/we-can-fix-the-internet-says-world-wide-web-inventor-tim-berners-lee-538947-2026-06-24>. *Dream Machine* Issue 36.
 16. Deezer, “AI-generated tracks now represent 44% of all new uploaded music,” April 2026. <https://newsroom-deezer.com/2026/04/ai-generated-tracks-represent-44-of-new-uploaded-music/>. *Music Business Worldwide*, “75,000 AI-generated tracks now flood Deezer daily, representing 44% of all new music uploaded to the platform.” <https://www.musicbusinessworldwide.com/75000-ai-generated-tracks-now-flood-deezer-daily-representing-44-of-all-new-music-uploaded-to-the-platform-says-streamer/>. *Dream Machine* Issues 7, 26, 27, 28.
 17. *Fast Company*, “Our interest in AI slop is hitting a ceiling.” <https://www.fastcompany.com/91568818/our-interest-in-ai-slop-is-hitting-a-ceiling>. *Dream Machine* Issue 37.
 18. Ditto Music research, October 2025 and prior. *Press Ditto Music*, “48% of artists use AI to make music — fewer than in 2023.” <https://press.dittomusic.com/48-of-artists-use-ai-to-make-music-fewer-than-in-2023>. *Dream Machine* Issue 2.
 19. *Musically*, “Universal and Warner could sign landmark AI deals within weeks.” <https://musically.com/2025/10/02/report-umg-and-wmg-could-sign-landmark-ai-deals-within-weeks/>. Spotify Newsroom, “Spotify Strengthens AI Protections for Artists, Songwriters, and Producers.” <https://newsroom.spotify.com/2025-09-25/spotify-strengthens-ai-protections/>. *Dream Machine* Issue 1.
 20. *Musically*, “50,000 AI music tracks are now uploaded to Deezer every day.” <https://musically.com/2025/11/12/50000-ai-music-tracks-are-now-uploaded-to-deezer-every-day/>. *Dream Machine* Issue 7.
 21. Deezer, April 2026, *op. cit.*
 22. *Musically*, “UMG boss slams exponential growth of AI slop on streaming services.” <https://musically.com/2026/01/09/umg-boss-slams-exponential-growth-of-ai-slop-on-streaming-services/>. *Dream Machine* Issue 14.
 23. *Musically*, “Report: 56.9% of new independent songs in China are AI-generated.” <https://musically.com/2026/01/05/report-56-9-of-new-independent-songs-in-china-are-ai-generated/>. *Dream Machine* Issue 13.
 24. *The Wrap*, “An AI Podcasting Machine Is Churning Out 3,000 Episodes a Week — and People Are Listening.” <https://www.thewrap.com/ai-podcasts-hosts-inception-point-ai/>. *Dream Machine* Issue 8.
 25. *Dream Machine* Issue 28, May 2026, citing aggregator-platform data on “podslop” classification.

26. *Digital Trends*, “181,000 mobile games released in last 6 months — AI and vibe coding have unleashed a flood of new games, but not necessarily better ones.” <https://www.digitaltrends.com/gaming/ai-and-vibe-coding-are-making-more-games-but-not-necessarily-better-ones/>. *Dream Machine* Issue 36.
27. *Metro*, “CD Projekt boss warns that games made solely by gen AI are already being made.” <https://metro.co.uk/2026/06/22/cd-projekt-boss-warns-games-made-solely-gen-ai-already-made-28878296/>. *Dream Machine* Issue 35.
28. *The Hollywood Reporter*, “Merriam-Webster Names ‘Slop’ Word of the Year Amid AI Boom.” <https://www.hollywoodreporter.com/news/general-news/slop-word-year-2025-merriam-webster-1236450780/>. *Dream Machine* Issue 12.
29. *Digital Music News*, “YouTube CEO Puts ‘Managing AI Slop’ on the Priority List for 2026.” <https://www.digitalmusicnews.com/2026/01/22/youtube-ceo-ai-slop-2026-comments/>. *Dream Machine* Issue 16.
30. *The Guardian*, “YouTube AI channels spreading fake, anti-Labour videos viewed 1.2bn times in 2025.” <https://www.theguardian.com/technology/2025/dec/13/fake-anti-labour-video-billion-views-youtube-2025>. *Dream Machine* Issue 12.
31. Deezer/Ipsos survey, November 2025. <https://newsroom-deezer.com/2025/11/deezer-ipsos-survey-ai-music/>. *Dream Machine* Issue 7.
32. *Bain & Company*, “In an AI Age, People Still Want the Radio Star.” <https://www.bain.com/insights/in-an-ai-age-people-still-want-the-radio-star/>. *Dream Machine* Issue 16.
33. Deezer, April 2026, *op. cit.* “Up to 85% of the streams generated by fully AI-generated tracks were in fact fraudulent in 2025.”
34. *Digital Music News*, “Band Wakes Up to Find AI Has Stolen 94% of Their Spotify Royalties by Duplicating Their Album and Altering the Speed.” <https://www.digitalmusicnews.com/2026/07/12/ai-steals-artist-royalties-speed-altering/>. *Dream Machine* Issue 38.
35. *Bloomberg*, “AI Changed Chess. Grandmasters Now Win With Unpredictable Moves,” 27 March 2026. <https://www.bloomberg.com/news/articles/2026-03-27/ai-changed-chess-grandmasters-now-win-with-unpredictable-moves>. *Dream Machine* Issue 23.
36. *Billboard*, “AI Artist Xania Monet Climbs the Charts — And Signs a Multimillion-Dollar Record Deal.” <https://www.billboard.com/pro/ai-music-artist-xania-monet-multimillion-dollar-record-deal/>.
37. *Billboard*, *op. cit.*; CNN, “Xania Monet is the first AI-powered artist to debut on a Billboard airplay chart.” <https://www.cnn.com/2025/11/01/entertainment/xania-monet-billboard-ai>.
38. *Billboard*, *op. cit.*
39. *Bangkok Post*, “AI singer Xania Monet signs \$3m deal with record label.” <https://www.bangkokpost.com/life/tech/3142355/ai-singer-xania-monet-signs-3m-deal-with-hollywood-media>. *Dream Machine* Issue 7.
40. Multiple outlets; quoted in *Billboard* feature *op. cit.*
41. Telisha Jones quoted in *Billboard*, *op. cit.*
42. NPR, “Breaking Rust is a hot new country act on the Billboard charts. It’s powered by AI.” <https://www.npr.org/2025/11/10/nx-s1-5604320/breaking-rust-is-a-hot-new-country-act-on-the-billboard-charts-its-powered-by-ai>. *Dream Machine* Issue 7.
43. *Washington Post*, “‘Walk My Walk,’ Breaking Rust: AI country hit triggers Nashville angst.” <https://www.washingtonpost.com/style/2025/12/28/breaking-rust-ai-country/>.

44. *MusicRadar*, “The No. 1 country song in the US right now is AI-generated.” <https://www.musicradar.com/music-tech/the-no-1-country-song-in-the-us-right-now-is-ai-generated/>. *Dream Machine* Issue 7.
45. BBC News, “The mysterious singer, Sienna Rose, with millions of streams is hitting the viral charts — but who (or what) is she?” <https://www.bbc.co.uk/news/articles/cq6v83gg66eo>. *Dream Machine* Issue 15.
46. *Billboard*, “How a MAGA Rapper Used AI to Create A Gospel Song That Climbed the Charts.” <https://www.billboard.com/pro/maga-rapper-ai-gospel-song-climbed-charts/>. *Dream Machine* Issue 9.
47. *Musically*, “AI band Bleeding Verse’s creator signs deal with Hallwood Media.” <https://musically.com/2025/10/07/ai-band-bleeding-verses-creator-signs-deal-with-hallwood-media/>. *Dream Machine* Issue 2.
48. *Musically*, “Indian AI band Trilok performs live, government denies association.” <https://musically.com/2025/12/17/indian-ai-band-trilok-performs-live-government-denies-association/>. *Dream Machine* Issue 12.
49. *Billboard*, “The Real Story Behind The AI Song That Knocked Tyla Off No. 1 On Billboard Afrobeats Chart.” <https://www.billboard.com/pro/ai-song-knocked-tyla-off-no-1-afrobeats/>. *Dream Machine* Issue 30.
50. *The Guardian*, “Paul McCartney joins music industry protest against AI with silent track.” <https://www.theguardian.com/music/2025/nov/17/the-sound-of-silence-why-theres-barely-anything-there-in-paul-mccartney-new-release>. *Dream Machine* Issue 8.
51. *The Guardian*, “Musicians must embrace ‘unstoppable force’ of AI, Eurythmics’ Dave Stewart urges.” <https://www.theguardian.com/music/2025/dec/05/musicians-must-embrace-unstoppable-force-of-ai-eurythmics-dave-stewart-urges>. *Dream Machine* Issue 11.
52. *Digital Music News*, “Nearly 800 Creatives, Including Jason Aldean and One Republic, Sign Responsible AI Declaration — ‘Stealing Our Work Is Not Innovation’.” <https://www.digitalmusicnews.com/2026/01/22/stealing-isnt-innovation/>. *Dream Machine* Issue 16.
53. *MusicTech*, “Jack Antonoff brands AI music makers as ‘godless whores’.” <https://musitech.com/news/industry/jack-antonoff-ai-music-makers-godless-whores/>. *Dream Machine* Issue 30.
54. *The Guardian*, “Artists are making ‘anti-slop’ to rebel against AI: ‘It’s been rammed down our throats’.” <https://www.theguardian.com/>. *Dream Machine* Issue 33.
55. *Variety*, “Behind the Movie Claiming to Be the First Straight-to-VHS Release in 20 Years — And Why It’s a ‘Middle Finger’ to AI.” <https://variety.com/2026/film/global/behind-the-fstraight-to-vhs-release-in-20-years-and-why-its-a-middle-finger-to-ai-1236767661/>. *Dream Machine* Issue 33.
56. *Digital Music News*, “‘Say No To Suno’ Protesters Bring the Drama to a Swanky Santa Monica AI Conference.” <https://www.digitalmusicnews.com/2026/06/04/say-no-to-suno-santa-monica-plane/>. *Dream Machine* Issue 33.
57. Stability AI, “Universal Music Group and Stability AI Announce Strategic Alliance.” <https://stability.ai/news/universal-music-group-and-stability-ai-announce-strategic-alliance>. *Dream Machine* Issue 5.
58. Stability AI, “Warner Music Group and Stability AI Join Forces To Build The Next Generation Of Responsible AI Tools For Music Creation.” <https://stability.ai/news/warner-music-group-and-stability-ai-join-forces-to-build-next-gen-tools>. *Dream Machine* Issue 8.
59. Universal Music, “Universal Music Group and Splice to Collaborate on the Next Generation of AI-Powered Music Creation Tools for Artists.” <https://www.universalmusic.com/universal-music->

group-and-splice-to-collaborate-on-the-next-generation-of-ai-powered-music-creation-tools-for-artists/. *Dream Machine* Issue 12.

50. LinkedIn / *Lexology*, “Munich Regional Court rules for GEMA against OpenAI.” Coverage: https://www.linkedin.com/posts/dr-barry-scannell-bbb5aa207_in-a-major-ruling-for-european-copyright-share-7393957246386323457-8bbx. *Dream Machine* Issue 7.
61. *EDM.com*, “Biggest Theft in Music History’: Rights Group Sues Suno as AI Music Showdown Escalates.” <https://edm.com/gear-tech/rights-group-sues-suno-copyright-infringement/>. *Dream Machine* Issue 7.
62. *Music Business Worldwide*, “Wixen files \$50m copyright suit against Meta.” <https://www.musicbusinessworldwide.com/wixen-files-50m-copyright-suit-against-meta-claims-tech-giant-wants-to-replace-songwriters-with-ai/>. *Dream Machine* Issue 16.
63. *Dream Machine* Issue 17 reportage on UMG’s \$3B suit against Anthropic.
54. *Stereogum*, “Bandcamp bans AI music.” <https://stereogum.com/2485199/bandcamp-bans-ai-music/news>. *Dream Machine* Issue 14.
65. *Dream Machine* Issue 18 reportage of Deezer licensing its detection tool.
56. *TechRadar*, “AI music is flooding Spotify, and subscribers are furious.” <https://www.techradar.com/audio/spotify/ai-music-is-flooding-spotify-and-subscribers-are-furious-heres-why-music-fans-no-longer-trust-discover-weekly>. *Dream Machine* Issue 14.
67. *CNET*, “San Diego Comic-Con Draws a Line: No AI Art Allowed at 2026 Event.” <https://www.cnet.com/culture/san-diego-comic-con-bans-ai-art-for-2026-event/>. *Dream Machine* Issue 16.
58. *The Independent*, “AI-generated song banned from Swedish charts: ‘It’s deceiving.’” <https://www.independent.co.uk/tv/news/ai-music-song-banned-sweden-spotify-b2901627.html>. *Dream Machine* Issue 15.
59. *Musically*, “Tidal allows AI-generated music — but it won’t earn royalties.” <https://musically.com/2026/06/30/tidal-allows-ai-generated-music-but-it-wont-earn-royalties/>. *Dream Machine* Issue 36.
70. *Hollywood Reporter*, “Amazon Orders Three Animated Series Using Generative AI.” <https://www.hollywoodreporter.com/tv/tv-news/amazon-orders-three-animated-series-generative-ai-1236606219/>. *Variety*, “Jorge Gutierrez Drops Out of Amazon MGM AI-Generated Series After Backlash.” <https://variety.com/2026/tv/news/jorge-gutierrez-drops-out-amazon-mgm-ai-generated-series-backlash-1236762285/>. *Futurism*, “Amazon AI-Generated Series Canceled.” <https://futurism.com/artificial-intelligence/amazon-ai-generated-series-canceled>. *Dream Machine* Issue 32.
71. Guillermo del Toro, subtweet responding to Amazon AI Creators Fund announcement, May/June 2026. Coverage in *Dream Machine* Issue 32. Cross-reference: del Toro’s “rather die” statement in Chapter 7, §“Position Three: We are not doing this.”
72. Apple Music platform statement, June 2026: AI-generated songs make up less than 1% of plays on the platform. *Dream Machine* Issue 31.
73. Spotify–Universal Music Group landmark deal allowing Spotify Premium subscribers to create AI covers and remixes of UMG-catalogue tracks as a paid Premium add-on, June 2026. *Dream Machine* Issue 31.
74. Spotify CEO Daniel Ek defending the platform’s move to AI music as “better than ‘slop,’” June 2026. *Dream Machine* Issue 31.
75. Swedish startup Tonada — AI music for retailers (June 2026). *Dream Machine* Issue 31.
76. *The Hollywood Reporter*, “Faceless Creators Take a Hit As YouTube Cracks Down on AI Slop.” <https://www.hollywoodreporter.com/business/digital/faceless-creators-youtube-ai-damage-1236617586/>. *Dream Machine* Issue 35.

77. *Soultracks*, “A.I.-generated music is catchy, familiar... and boring.” <https://soultracks.com/news-ai-generated-music-is-catchy-boring/>. *Dream Machine* Issue 14.
78. *The Independent*, “AI-generated song banned from Swedish charts: ‘It’s deceiving.’” <https://www.independent.co.uk/tv/news/ai-music-song-banned-sweden-spotify-b2901627.html>. *Dream Machine* Issue 15.
79. *Marketing Week*, “You can’t dismiss AI ads as slop when they’re winning in testing.” Coverage discussed in *Dream Machine* Issue 22.
30. UK Department for Science, Innovation and Technology, *Statement of Progress on Copyright and AI*, 15 December 2025. <https://www.gov.uk/government/publications/copyright-and-artificial-intelligence-progress-report/copyright-and-artificial-intelligence-statement-of-progress-under-section-137-data-use-and-access-act>. *Dream Machine* Issue 12, “Editor’s Pick: 88% of Creators Said ‘No.’” 18 December 2025.
81. UK DSIT, original consultation, 17 December 2024 – 25 February 2025. Discussion in IPWatchdog, “Respondents to UK AI Consultation Overwhelmingly Want AI Companies to License Copyrighted Works in All Cases.” <https://ipwatchdog.com/2025/12/16/respondents-uk-ai-consultation-overwhelmingly-want-ai-companies-license-copyrighted-works-all-cases/>.
32. IPWatchdog, *op. cit.*; Hogan Lovells, “Copyright and AI: UK government publishes statement of progress.” <https://www.hoganlovells.com/en/publications/copyright-and-ai-uk-government-publishes-statement-of-progress>.
33. UK DSIT, *Statement of Progress*, *op. cit.*; analysis at UCL Copyright Queries, “UK government publishes progress statement on AI and copyright consultation.” <https://blogs.ucl.ac.uk/copyright/2025/12/23/uk-government-publishes-progress-statement-on-ai-and-copyright-consultation/>.
34. UK DSIT, *Statement of Progress*, *op. cit.*
85. Society of Authors submission to the UK consultation, quoted in IPWatchdog, *op. cit.*
36. Dr Barry Scannell, LinkedIn analysis of GEMA v. OpenAI ruling, November 2025. https://www.linkedin.com/posts/dr-barry-scannell-bbb5aa207_in-a-major-ruling-for-european-copyright-share-7393957246386323457-8bbx. *Dream Machine* Issue 7.
87. *Variety*, “Suno has lost a landmark AI copyright lawsuit to German music rights society GEMA.” <https://variety.com/2026/digital/news/suno-loses-ai-lawsuit-gema-1236825010/>. *Dream Machine* Issue 40.
38. *Creative Blog*, “What does the biggest AI copyright payout in history actually change for creatives?” <https://www.creativeblog.com/ai/what-does-the-biggest-ai-copyright-payout-in-history-actually-change-for-creatives>. *Dream Machine* Issue 40.
39. *Digital Music News*, “Suno Quickly Seizes on Udio’s Legal Victory Against Sony Music Over Admissible Infringed Works.” <https://www.digitalmusicnews.com/2026/07/05/suno-admissible-infringed-works-argument/>. *Dream Machine* Issue 37.
90. *EDM.com*, “‘Biggest Theft in Music History’: Rights Group Sues Suno as AI Music Showdown Escalates.” <https://edm.com/gear-tech/rights-group-sues-suno-copyright-infringement/>. *Dream Machine* Issue 7.
91. *Music Business Worldwide*, “Wixen files \$50m copyright suit against Meta, claims tech giant wants to replace songwriters with AI.” <https://www.musicbusinessworldwide.com/wixen-files-50m-copyright-suit-against-meta-claims-tech-giant-wants-to-replace-songwriters-with-ai/>. *Dream Machine* Issue 16.
92. *Dream Machine* Issue 17, on UMG’s \$3B suit against Anthropic.
93. *Complete Music Update*, “Johnny Cash estate uses ELVIS Act to sue Coke over tribute act ad soundtrack.” <https://completemusicupdate.com/johnny-cash-estate-uses-elvis-act-to-sue-coke-over-tribute-act-ad-soundtrack/>. *Dream Machine* Issue 9.

94. Reuters, “European lawmakers seek EU-wide minimum age to access AI chatbots, social media.” <https://www.reuters.com/legal/litigation/european-lawmakers-seek-eu-wide-minimum-age-access-ai-chatbots-social-media-2025-11-26/>. *Dream Machine* Issue 9.
95. SAG-AFTRA contract update reporting through Q2 2026. *Dream Machine* Issues 20, 26, 29. Coverage: <https://www.theverge.com/news/842848/new-york-law-ai-advertisements-sag-aftra-labor>.
96. Equity (U.K.), “Performers prepared to take industrial action over AI in landslide 99% vote.” <https://www.equity.org.uk/news/2025/performers-prepared-to-take-industrial-action-over-ai-in-landslide-99-vote>. *Dream Machine* Issue 12.
97. Equity (U.K.), “Equity welcomes improved offer in AI protection negotiations in film and TV.” <https://www.equity.org.uk/news/2026/equity-welcomes-improved-offer-in-ai-protection-negotiations-in-film-and-tv>. *Dream Machine* Issue 15.
98. Cannes Film Festival AI Disclosure Standard launch, May 2026. *Dream Machine* Issue 29.
99. *Musically*, “BPI sets out transparency and sovereignty demands to secure ‘AI licensing boom’.” <https://musically.com/2026/05/19/bpi-transparency-sovereignty-ai-licensing-boom/>. *Dream Machine* Issue 30.
100. UK DSIT, *Statement of Progress*, op. cit.
01. *Dream Machine* Issue 21, 19 March 2026, on the UK government’s revised position on AI copyright.
02. *Digital Music News*, “The AI Licensing Shift — Creative Weight Attribution Emerges as Music Industry Game-Changer for Rights Holders.” <https://www.digitalmusicnews.com/2026/01/26/ai-licensing-shift-creative-weight-attribution/>. See also *Digital Music News*, “Artificial Intelligence Attribution and Licensing Startup Musical AI Scores \$4.5 Million Raise.” <https://www.digitalmusicnews.com/2026/01/13/musical-ai-funding-january-2026/>. *Dream Machine* Issues 14, 16.
03. *Digital Music News*, “Federal Judge Vacates Order Sealing Udio’s Training Data — Are We About to Learn How Many Recordings the Platform Ingested?” <https://www.digitalmusicnews.com/2026/06/05/udio-training-number-order-vacated/>. *Dream Machine* Issue 33.
04. *Mixmag*, “Artists can now check whether their music has been used to train AI.” <https://mixmag.net/read/artists-check-music-used-train-ai-atlantic-investigation-news>. See also *MusicRadar*, “Is your track in the datasets that are being used to train AI? The Atlantic unveils a way you can find out.” <https://www.musicradar.com/music-tech/is-your-track-in-the-datasets-that-are-being-used-to-train-ai-the-atlantic-unveils-a-way-you-can-find-out>. *Dream Machine* Issues 35, 36.
05. *Euronews*, “SZA hits out at ‘disgusting’ AI music after discovering her songs were used in AI training.” <https://www.euronews.com/culture/2026/06/22/sza-hits-out-at-disgusting-ai-music-after-discovering-her-songs-were-used-in-ai-training>. *Dream Machine* Issue 35.
06. PRS for Music, “PRS for Music AI Survey 2026.” <https://www.prsformusic.com/m-magazine/news/prs-for-music-ai-survey-2026>. *Dream Machine* Issue 16.
07. *Broadcast Now*, “Alex Mahon joins Stellar AI Creative Summit line-up” (covering the launch of the UCL/RCA Centre for Creative AI). <https://www.broadcastnow.co.uk/broadcasters/alex-mahon-joins-stellar-ai-creative-summit-line-up/5209227.article>. *Dream Machine* Issue 1.
08. *Complete Music Update*, “Artists must have creative control in AI deals or risk ending up with ‘scraps’, says US artist trade body.” <https://completemusicupdate.com/artists-must-have-creative-control-in-ai-deals-or-risk-ending-up-with-scraps-says-us-artist-trade-body/>. *Dream Machine* Issue 6.

39. *Digital Music News*, “Nearly 800 Creatives, Including Jason Aldean and One Republic, Sign Responsible AI Declaration — ‘Stealing Our Work Is Not Innovation’.” <https://www.digitalmusicnews.com/2026/01/22/stealing-isnt-innovation/>. *Dream Machine* Issue 16.
10. *Digital Music News*, “Google Moves to Dismiss Artists’ AI Lawsuit, Says They Consented ‘To the Use of Their Recordings’ When Uploading to YouTube.” <https://www.digitalmusicnews.com/2026/06/10/youtube-ai-lawsuit-dismissal-motion/>. *Dream Machine* Issue 34.
11. *Billboard*, “YouTube’s AI Training Argument Raises Alarm Among Indie Music Advocates: ‘Not Informed Consent’.” <https://www.billboard.com/pro/youtube-ai-training-argument-indie-music-community-concerns/>. *Dream Machine* Issue 36.
12. *Digital Music News*, “Google Doubles Down on ‘Fair Use’ AI Training, Stressing That ‘Outputs’ Not ‘Inputs’ Are Where Enforcement Should Be Focusing.” <https://www.digitalmusicnews.com/2026/06/29/google-fair-use-ai-training/>. *Dream Machine* Issue 36.
13. *Billboard*, “ABBA’s Björn Ulvaeus Says We Might Be Doing AI Licensing Deals Wrong.” <https://www.billboard.com/pro/abba-bjorn-ulvaeus-music-ai-un-global-summit/>. See also *Musically*, “Tracing the output was always the wrong question.” <https://musically.com/2026/07/13/bjorn-ulvaeus-on-ai-tracing-the-output-was-always-the-wrong-question/>. *Dream Machine* Issue 38.
14. *Complete Music Update*, “US musicians’ union sues Universal and Warner over AI deals, claims breach of ‘new use’ clause in record industry’s performer agreement.” <https://completemusicupdate.com/us-musicians-union-sues-universal-and-warner-over-ai-deals-claims-breach-of-new-use-clause-in-record-industrys-performer-agreement/>. *Dream Machine* Issue 33.
15. Adobe Firefly milestone and adoption data. Firefly Foundry and Firefly Image Model 5 launch reporting, Adobe MAX 2025: <https://news.adobe.com/news/2025/10/adobe-max-2025-firefly-foundry>; <https://news.adobe.com/news/2025/10/adobe-max-2025-firefly>.
16. Bria AI consent-licensed dataset and attribution mechanism. [TODO: confirm primary citation — Bria’s licensed-data white paper or Series B coverage.]
17. Getty Images, “Generative AI by iStock” launch, built on NVIDIA Picasso, trained exclusively on Getty’s licensed library with contributor royalties. [TODO: confirm citation — Getty press release or *Reuters* coverage.]
18. Moonvalley Marey, generative-video foundation model trained on licensed video. [TODO: confirm citation — Moonvalley launch coverage in *The Verge* / *TechCrunch*.]
19. AIODE, ethically-trained music creation DAW. See Chapter 17: The Tools, §“Audio modality models.”
20. *MusicTech*, “Tamber is an ‘ethically trained’ AI tool to aid the creative process – and you can use arm gestures to control it.” <https://musictech.com/news/gear/tamber-ai-ethically-trained-arm-gestures/>. Tamber product page: <https://tamber.ai/>. *Dream Machine* Issue 30.
21. Stability AI / Universal Music Group strategic alliance: <https://stability.ai/news/universal-music-group-and-stability-ai-announce-strategic-alliance>. Stability AI / Warner Music: <https://stability.ai/news/warner-music-group-and-stability-ai-join-forces-to-build-next-gen-tools>. Universal Music / Splice partnership: <https://www.universalmusic.com/universal-music-group-and-splice-to-collaborate-on-the-next-generation-of-ai-powered-music-creation-tools-for-artists/>. *Dream Machine* Issues 5, 8, 12.
22. Reporting on AI-generated images in the Adobe Stock training corpus, *Bloomberg*, April 2024. [TODO: confirm exact citation.]

23. Spotify and Universal Music Group, “Landmark Licensing Agreements for Fan-Made Covers and Remixes,” 21 May 2026. Spotify Newsroom: <https://newsroom.spotify.com/2026-05-21/universal-music-group-spotify-licensing-agreements-fan-made-covers-remixes/>. *TechCrunch*: <https://techcrunch.com/2026/05/21/spotify-and-universal-music-strike-deal-allowing-fan-made-ai-covers-and-remixes/>. *Medianama*: <https://www.medianama.com/2026/05/223-spotify-umg-fans-create-licensed-ai-covers-remixes/>. *Dream Machine* Issue 31. Detailed treatment of the three-Cs framework, the deal’s prehistory in the UMG settlement with Udio (October 2025), and the *Verified by Spotify* badge.
24. *Complete Music Update*, “NMPA unveils AI licensing deals with Udio and Klay with 50/50 split for songs and recordings.” <https://completemusicupdate.com/nmpa-unveils-ai-licensing-deals-with-udio-and-klay-with-50-50-split-for-songs-and-recordings/>. *Dream Machine* Issue 34.
25. Hasbro, “Hasbro Launches Sixth Wall, a New AI Studio,” BusinessWire press release, 3 June 2026: <https://www.businesswire.com/news/home/20260603297922/en/>. *The Wrap*: <https://www.thewrap.com/industry-news/business/hasbro-ai-studio-sixth-wall-launch/>. *Screen Daily*: <https://www.screendaily.com/news/hasbro-launches-subsidiary-ai-studio-sixth-wall-to-exploit-ip/5217380.article>. CharacterOS runtime governance layer; talent-participation pay-out structure for original performers; initial 12-character catalogue; 13+/enterprise focus. Cross-reference in Chapter 7 §“Position Five: Own the rails.”
26. *Dream Machine* Deep Dive №14, *The Programmable Brand: How AI Is Reshaping IP Licensing in the Generative Economy*, 4 June 2026. The deep dive’s three Cs framework draws on the Spotify–UMG licensing agreements; the five-model comparison covers Hasbro/Sixth Wall (Behavioural Licensing), Spotify/UMG (platform-embedded remixing), Disney/OpenAI (the framework that survived the platform’s collapse — see footnote ^3a in Chapter 7), the independent-artist commons-and-DAO models (Grimes/Elf.tech with a ~50% master-royalty split; Holly Herndon / Holly+ with DAO governance; ElevenLabs heritage voices via CMG Worldwide; Meta’s 2023 celebrity-chatbot programme as the closed-static anti-pattern), and Suno/Alexa as the *create-to-consume* demand floor. The underlying economic argument draws on Coase (1937), Williamson, and Barney’s VRIN resource-based view; the appendix’s load-bearing claim is that AI *relocates* rather than *abolishes* transaction costs — collapsing search and negotiation while raising the cost of monitoring and enforcement, and pushing the durable value into the governance layer.
27. Adobe Firefly IP indemnification for enterprise customers. [TODO: confirm citation — Adobe enterprise terms or *The Verge* coverage from 2023.]
28. Microsoft, “Microsoft announces new Copilot Copyright Commitment for customers,” 7 September 2023. <https://blogs.microsoft.com/on-the-issues/2023/09/07/copilot-copyright-commitment-ai-legal-concerns/>.
29. Google Cloud Generative AI indemnification: <https://cloud.google.com/blog/products/ai-machine-learning/protecting-customers-with-generative-ai-indemnification>.
30. IBM watsonx uncapped indemnity for enterprise customers. [TODO: confirm citation.]
31. Sundance AI Literacy Initiative, in Chapter 13: Authenticity, the New Scarcity, §“The provenance infrastructure, named.”
32. For Disney’s parallel position, see *Deadline*, “Disney Sends Cease And Desist Letter To Character.ai.” <https://deadline.com/2025/09/disney-cease-and-desist-letter-characterai-copyright-infringement-1236566831/>. For Studio Ghibli’s similar stance: *NDTV Profit*, “Studio Ghibli And Studio That Developed Elden Ring Send Stern Message To OpenAI.” <https://www.ndtvprofit.com/technology/studio-ghibli-and-studio-that-developed-elden-ring-send-stern-message-to-openai>. *Dream Machine* Issues 2, 6.
33. *Variety*, “Is ‘AI Resistance’ Setting the Music Sector Back? WMG’s Robert Kyncl Sees ‘An Incredible Value Creation Opportunity,’ But Warns ‘We Cannot Wait the Way the Industry Did 25

- Years Ago.” <https://variety.com/2026/music/news/wmg-robert-kyncl-ai-resistance-1236748901/>. *Dream Machine* Issue 30.
34. Musically, “Open letter: Musicians ‘must not be pressured into AI deals’ by rightsholders.” <https://musically.com/2026/06/22/open-letter-musicians-must-not-be-pressured-into-ai-deals-by-rightsholders/>. *Dream Machine* Issue 35.
 35. Adobe, *Creators’ Toolkit Report*, op. cit. 69% of 16,000 surveyed creators worried about their work being used to train AI without consent.
 36. *Screen Daily*, “We got it wrong says UK Culture Secretary Lisa Nandy of government misstep on AI and copyright strategy,” 1 June 2026. <https://www.screendaily.com/news/we-got-it-wrong-says-uk-culture-secretary-lisa-nandy-of-government-misstep-on-ai-and-copyright-strategy/5217330.article>. *Dream Machine* Issue 32.
 37. GOV.UK, “AI Adoption Plan for the Creative Industries.” <https://www.gov.uk/government/publications/ai-champions-ai-adoption-plans/ai-adoption-plan-creative-industries>. *Dream Machine* Issue 33.
 38. GOV.UK / DSIT, “A decisive shift to power British AI: new £1.1 billion plan to back chip firms, boost computing power and skills for the AI revolution.” <https://www.gov.uk/government/news/a-decisive-shift-to-power-british-ai-new-11-billion-plan-to-back-chip-firms-boost-computing-power-and-skills-for-the-ai-revolution>. *Dream Machine* Issue 33.
 39. *Fortune*, “Google forced to allow news sites to opt out of AI scraping in ‘world first,’ UK watchdog says.” <https://fortune.com/2026/06/05/auto-draft-3/>. See also *National Law Review*, “UK Regulator Introduces ‘World First’ Controls on the Use of Content in Google Generative AI Features.” <https://natlawreview.com/article/uk-regulator-introduces-world-first-controls-use-content-google-generative-ai>. *Dream Machine* Issues 33, 34.
 40. *CNBC*, “Netflix ‘all in’ on leveraging AI as the tech creeps into entertainment industry,” 22 October 2025. <https://www.cnn.com/2025/10/22/netflix-all-in-on-leveraging-ai-in-its-streaming-platform.html>. *Dream Machine* Issue 4.
 41. *Futurism*, “Lionsgate’s Attempt to Create Movies Using AI Has Crumbled Into Disaster.” <https://futurism.com/artificial-intelligence/lionsgate-movies-ai>. *Dream Machine* Issue 1.
 42. *The Guardian*, “Disney to invest \$1bn in OpenAI, allowing characters in Sora video tool.” <https://www.theguardian.com/business/2025/dec/11/disney-open-ai-sora-video-deal>. *Dream Machine* Issue 11. The announcement was widely treated at the time as a completed, operative template; on the corrected record, no formal agreement was ever signed and no money changed hands.
 43. OpenAI staged Sora wind-down announced 24 March 2026: consumer app and website to end in late April 2026; developer API to close in late September 2026. Disney exited the partnership with reportedly as little as thirty minutes’ notice. *Ropes & Gray*, “The Deal That Wasn’t: What Disney and OpenAI Teach Us About Strategic Investments.” <https://www.ropegray.com/en/insights/viewpoints/the-deal-that-wasnt-what-disney-and-openai-teach-us-about-strategic-investments>. *TheStreet* / AOL (Hollywood Reporter source): <https://www.thestreet.com/technology/openai-is-shutting-down-sora-and-the-disney-deal-is-off>. *Tech-Insider*: <https://tech-insider.org/openai-sora-shutdown-disney-deal-ai-video-2026/>. Peak Sora revenue reported at roughly \$540k/month against operating costs in the billions; OpenAI reallocated compute toward higher-margin coding and enterprise products. Coverage in *Dream Machine* Issue 22.
 44. *Dream Machine* Deep Dive №14, *The Programmable Brand*, 4 June 2026.
 45. PYMNTS, “Retention Is Name of the Game for Netflix’s AI Strategy.” <https://www.pymnts.com/subscription-commerce/2026/retention-is-name-of-the-game-for-netflix-ai-strategy/>. *Dream Machine* Issue 15.

46. *Hollywood Reporter*, “Netflix is building and recruiting for an AI animation studio, called INKubator, to produce ‘feature-quality’ shorts.” <https://www.hollywoodreporter.com/business/business-news/netflix-ai-animation-studio-inkubator-1236592110/>. *Dream Machine* Issue 30.
47. *Variety*, “About 300 Netflix Titles Used Generative AI This Year, Company Reveals.” <https://variety.com/2026/biz/news/about-300-netflix-programs-used-ai-this-year-q2-earnings-1236812914/>. *Dream Machine* Issue 39.
48. *Deadline*, “Amazon Builds Out AI Studios With Sports Docs Boss Matt Newman Named Head Of Live-Action.” <https://deadline.com/2025/11/amazon-ai-studios-matt-newman-1236603477/>. *Dream Machine* Issue 7.
49. *Wired*, “Amazon’s House of David Used Over 350 AI Shots in Season 2. Its Creator Isn’t Sorry.” <https://www.wired.com/story/amazons-house-of-david-used-over-350-ai-shots-in-season-2-its-creator-isnt-sorry/>. *Dream Machine* Issue 7.
50. Amazon AI Creators Fund launch and the Punky Duck cancellation. *Hollywood Reporter*: <https://www.hollywoodreporter.com/tv/tv-news/amazon-orders-three-animated-series-generative-ai-1236606219/>. *Variety*: <https://variety.com/2026/tv/news/jorge-gutierrez-drops-out-amazon-mgm-ai-generated-series-backlash-1236762285/>. *Futurism*: <https://futurism.com/artificial-intelligence/amazon-ai-generated-series-canceled>. *Dream Machine* Issue 32. Cross-reference: Chapter 5 §“The pushback” for the audience-enforcement analysis.
51. *The Verge*, “Amazon’s new plan for games: James Bond and AI Snoop Dogg.” <https://www.theverge.com/games/943147/amazon-gaming-strategy-james-bond-snoop-dogg-luna>. *Dream Machine* Issue 33.
52. *Video Games Chronicle*, “NBCUniversal signs deal with Law & Order creator Dick Wolf’s son to make AI-generated games based on its IP.” <https://www.videogameschronicle.com/news/nbcuniversal-signs-deal-with-law-order-creator-dick-wolfs-son-to-make-ai-generated-games-based-on-its-ip/>. *Dream Machine* Issue 5.
53. *NME*, “‘The Office’, ‘Saturday Night Live’ and ‘Sex And The City’ could be turned into AI games.” <https://www.nme.com/news/gaming-news/the-office-and-sex-and-the-city-ai-video-games-3901630>. *Dream Machine* Issue 5.
54. *The Hollywood Reporter*, “Disney+ to Allow User-Generated Fan Content with AI.” <https://www.hollywoodreporter.com/business/digital/disney-plus-gen-ai-user-generated-content-1236426135/>. *Dream Machine* Issue 8.
55. *Dream Machine* Issue 8 reportage of the Disney “Office of Technology Enablement,” led by former Walt Disney Studios CTO Jamie Voris.
56. *Marketing Dive*, “Disney unveils TikTok-like vertical video, AI video generation tool.” <https://www.marketingdive.com/news/disney-unveils-tiktok-like-vertical-video-ai-video-generation-tool/809269/>. *Dream Machine* Issue 14.
57. *The Hollywood Reporter*, “Fox Entertainment Takes Equity Stake in AI-Microdramas Company Holywater.” <https://www.hollywoodreporter.com/business/business-news/fox-entertainment-invests-in-holywater-ai-microdramas-1236396802/>. *Dream Machine* Issue 3.
58. *Deadline*, “Sky History Acquires ‘Castles SOS,’ AI-Powered Doc Exploring Royalty, Ruins & Restoration.” <https://deadline.com/2025/11/castles-sos-ai-doc-sky-history-documentary-rick-edwards-1236627378/>. *Dream Machine* Issue 9.
59. *Estate Agent Today*, “Homebuilder among first to use Channel 4’s AI ads.” <https://www.estateagenttoday.co.uk/breaking-news/2025/12/homebuilder-among-first-to-use-channel-4s-ai-ads/>. *Dream Machine* Issue 11.
50. *The Hollywood Reporter*, “Google Investing in ‘Backrooms’ Studio A24.” [https://www.hollywoodreporter.com/business/digital/a24-google-deepmind-ai-venture-](https://www.hollywoodreporter.com/business/digital/a24-google-deepmind-ai-venture-1236603477/)

- [backrooms-1236627228/](#). *Dream Machine* Issue 35.
61. *Futurism*, “A24 Is Dealing With a Full-on Crisis Among Fans After Signing an AI Deal With Google.” <https://futurism.com/artificial-intelligence/a24-crisis-fans-signing-ai-deal-google>. *Dream Machine* Issue 36.
 62. *Bloomberg*, “Animation’s AI Reckoning: Filmmakers Say They Can Make Movies for 90% Less.” <https://www.bloomberg.com/news/newsletters/2026-06-28/animation-s-ai-reckoning-filmmakers-say-they-can-make-movies-for-90-less>. *Dream Machine* Issue 36.
 63. *CNN*, “South Korea’s movie industry is embracing AI. Will K-film lose its magic?” <https://www.cnn.com/world/south-korea-film-industry-ai-hnk-spc>. *Dream Machine* Issue 35.
 64. *The Hollywood Reporter*, “How the Shanghai Film Festival Is Embracing Cinema’s Future: From AI to iPhone Moviemaking.” <https://www.hollywoodreporter.com/movies/movie-news/shanghai-film-festival-cinema-future-ai-iphone-1236626463/>. *Dream Machine* Issue 35.
 65. *Variety*, “MiniMax Launches All-in-One AI Video Generator Hub at Shanghai Film Festival: ‘Creative Direction and Aesthetic Judgment Must Ultimately Be Left to Humans.’” <https://variety.com/2026/film/festivals/minimax-hub-ai-video-generator-shanghai-film-festival-1236781901/>. *Dream Machine* Issue 35.
 66. *The Hollywood Reporter*, “Fremantle Names Boss of New AI Native Studio Imaginae Studios.” <https://www.hollywoodreporter.com/business/digital/fremantle-names-ceo-new-ai-label-imaginae-studios-1236396579/>. *Dream Machine* Issue 2.
 67. *Dream Machine* Issue 25, on Fremantle’s *Art Awakens* development.
 68. *Indiewire*, “Another New AI Production Company Inks a Big Creative Partnership — This Time, with Ron Howard and Brian Grazer’s Imagine Entertainment.” <https://www.indiewire.com/news/business/obsidian-studio-ai-production-company-imagine-entertainment-1235158619/>. *Dream Machine* Issue 6.
 69. *UK Tech News*, “AI film studio Wonder lands \$9m investment.” <https://www.uktech.news/ai/ai-film-studio-wonder-lands-9m-investment-20251023>. *Dream Machine* Issue 5.
 70. Wonder Studios, “Shortlisted films revealed for The Wonder Film Festival.” <https://www.linkedin.com/posts/wearewonderstudios-were-thrilled-to-share-the-shortlisted-films-activity-7404560378082246656-7NcI>. *Dream Machine* Issue 11.
 71. *Forbes*, “Meet Wonder Studios, The \$50M British Studio Striving To Become The A24 Of AI Production.” <https://www.forbes.com/sites/charliefink/2026/05/18/meet-wonder-studios-the-50m-british-studio-striving-to-become-the-a24-of-ai-production/>. *Dream Machine* Issue 30.
 72. *The Hollywood Reporter*, “AI Company Asteria Produces New Animated Short ‘All Heart.’” <https://www.hollywoodreporter.com/movies/movie-news/natasha-lyonne-ai-company-asteria-1236403144/>. *Dream Machine* Issue 4.
 73. *The Hollywood Reporter*, “Promise, a deep-pocketed AI studio backed by Google, aims to Bring GenAI Filmmaking and VFX to Legacy Media.” <https://www.hollywoodreporter.com/business/digital/ai-studio-promise-vfx-generation-company-1236397636/>. *Dream Machine* Issue 3.
 74. *Variety*, “AI-Powered Cinematic Universe Platform enGEN3 Launched by Goldfinch.” <https://variety.com/2025/film/news/ai-powered-cinematic-universe-platform-engen3-1236543349/>. *Dream Machine* Issue 2.
 75. *Deadline*, “Munich Based Beta Films & Industry Execs Join Forces To Launch Artificial Intelligence Start-Up Chapter41.” <https://deadline.com/2025/11/beta-film-ai-startup-chapter41-artificial-intelligence-1236612632/>. *Dream Machine* Issue 7.
 76. *The Hollywood Reporter*, “Longtime TV Exec, Kevin Reilly, Set to Lead AI Startup Kartel.” <https://www.hollywoodreporter.com/business/digital/kevin-reilly-ceo-kartel-ai-hbo-1236424692/>. *Dream Machine* Issue 7.

77. Variety, “‘Wanted’ Director Timur Bekmambetov Explains His \$5 Million Plan to Generate AI Method Actors: ‘AI Is Here to Stay. We Have to Train It Responsibly.’” <https://variety.com/2025/film/news/wanted-director-method-acting-ai-actors-1236579647/>. *Dream Machine* Issue 7.
78. Variety, “Tilly Norwood Creator Doubles Down on AI ‘Actors’ and Says It’s a ‘More Ethical Way to Perform,’ Urges Human Actors to ‘Future-Proof’ Themselves With AI.” <https://variety.com/2026/digital/news/tilly-norwood-creator-tells-actors-to-create-ai-avatars-1236638940/>. *Dream Machine* Issue 16.
79. *Broadcast Now*, “Wonder Studios adapts children’s book to animated series with AI.” <https://www.broadcastnow.co.uk/production-and-post/wonder-studios-adapts-childrens-book-to-animated-series-with-ai/5211713.article>. *Dream Machine* Issue 11.
30. *Variety*, “Kling AI Partners With Evolutionary Films on Animated Feature ‘Minibots,’ Unveils Filmmaker Initiative at Cannes Market.” <https://variety.com/2026/film/news/kling-ai-evolutionary-films-minibots-cannes-1236748590/>. *Dream Machine* Issue 30.
81. *Deadline*, “AI Firm Utopai Studios To Co-Produce Silver Bear-Winning Filmmaker Hyo-Joo Yang’s Feature ‘Half Moon’.” <https://deadline.com/2026/06/utopai-studios-hyo-joo-yang-half-moon-ai-1236952730/>. *Dream Machine* Issue 34.
32. *Variety*, “Fully AI-Generated Series ‘Journey to the West’ in Development by Utopai Studios and Huace.” <https://variety.com/2026/tv/news/fully-ai-generated-journey-to-the-west-utopai-studios-huace-1236787171/>. *Dream Machine* Issue 35.
33. *Variety*, “Tribeca Festival AI Film ‘Dreams of Violets’ by Fountain o.” <https://variety.com/2026/film/festivals/tribeca-festival-ai-film-dreams-of-violets-fountain-o-1236759724/>. Director Ash Koosha; producer Fountain o; 75-minute docudrama; subject: 2026 Iran protests; cost \$2,000; production time: 3 months; tools: Kling AI (video), Anthropic Claude (editing), Google Gemini (research). *Dream Machine* Issue 32.
34. The broader festival-circuit context: the *Dreams of Violets* Tribeca selection follows the *Lily* (\$1M Dubai AI Film Award, January 2026) and *Humans in the Loop* (Film Independent Sloan Distribution Grant, November 2025) as part of the accumulating festival recognition of AI-native feature and short-film work — but *Dreams of Violets* is the first Tier 1 festival selection for a feature-length, fully-AI-generated work.
85. *Variety*, “‘Dreams of Violets’ Review: What Does a Film Made Entirely with AI Look Like?” <https://variety.com/2026/film/reviews/dreams-of-violets-review-ash-koosha-tribeca-festival-1236784209/>. *Dream Machine* Issue 35.
36. *Variety*, “‘Watch the Skies,’ Swedish UFO Feature Film Dubbed Entirely With AI, Sets USA Distribution Deal.” <https://variety.com/2025/film/news/watch-the-skies-us-theatrical-release-ai-dubbing-1236343110/>. *Dream Machine* Issue 5.
87. *Cybernews*, “Run to the West — South Korea’s first AI film tests the soul of cinema.” <https://cybernews.com/entertainment/korean-cinema-run-to-the-west-ai/>. *Dream Machine* Issue 5.
38. *Broadcast Pro Middle East*, “Tunisian filmmaker wins \$1 million AI Film Award for ‘Lily.’” <https://www.broadcastprome.com/news/tunisian-filmmaker-wins-1-million-ai-film-award-for-lily/>. *Dream Machine* Issue 14.
39. *Variety*, “AI Drama ‘Humans in the Loop’ Receives Film Independent’s Sloan Distribution Grant, Enters Oscar Race.” <https://variety.com/2025/film/news/ai-drama-humans-in-the-loop-oscar-race-1236582975/>. *Dream Machine* Issue 8.
30. *PC Gamer*, “Palworld studio Pocketpair says its new publishing division won’t handle games that use generative AI: ‘We don’t believe in it’.” <https://www.pcgamer.com/software/ai/palworld-studio-pocketpair-says-its-new-publishing-division-wont-handle-games-that-use-generative-ai-we-dont-believe-in-it/>. *Dream Machine* Issue 4.

91. *Niche Gamer*, “Larian Studios backs off from gen AI.” *Dream Machine* Issue 14.
92. *Decrypt*, “Warhammer 40,000 Maker Games Workshop Rules Out Generative AI.” *Dream Machine* Issue 14.
93. *Niche Gamer*, “Manor Lords publisher Hooded Horse won’t work with devs using gen AI.” *Dream Machine* Issue 14.
94. *gamesindustry.biz*, “RuneScape maker Jagex says it will never use generative AI to make in-game content.” *Dream Machine* Issue 16.
95. *GamesRadar*, “Wallace and Gromit creator says beloved animation studio Aardman will ‘embrace the technology’ of AI, but will be ‘very cautious not to lose our values’.” <https://www.gamesradar.com/entertainment/animation-movies/wallace-and-gromit-creator-says-beloved-animation-studio-aardman-will-embrace-the-technology-of-ai-but-will-be-very-cautious-not-to-lose-our-values/>. *Dream Machine* Issue 11.
96. *Variety*, “Guillermo del Toro Says He’d ‘Rather Die’ Than Use Generative AI in His Films: ‘Not Interested’.” <https://variety.com/2025/film/news/guillermo-del-toro-rather-die-generative-ai-frankenstein-1236561316/>. *Dream Machine* Issue 5.
97. *Variety*, “Martin Scorsese Supports AI Company, Endorses AI for Storyboarding.” <https://variety.com/2026/film/news/martin-scorsese-supports-ai-company-storyboard-movies-1236765037/>. Scorsese becomes advisor to Black Forest Labs (FLUX model); endorses AI storyboarding; Boots Riley and Guillermo del Toro respond critically. *Dream Machine* Issue 32.
98. *The Hollywood Reporter*, “Ron Howard Is Bullish on AI and Thinks Jim Henson Would Have Been Too.” <https://www.hollywoodreporter.com/business/digital/ron-howard-ai-jim-henson-martin-scorsese-1236620457/>. *Dream Machine* Issue 34.
99. *Variety*, “Gore Verbinski Says We Need a ‘Rating System’ for AI Use in Films: ‘If You Use AI to Write a Script, You Get an F’.” <https://variety.com/2026/film/global/gore-verbinski-ai-taormina-1236780502/>. *Dream Machine* Issue 34.
100. *The Guardian*, “‘It has no taste whatsoever. It is a genius at helping you... Like a billionaire on acid’: Star Wars director Gareth Edwards comes out in favour of AI.” <https://www.theguardian.com/film/2026/may/29/star-wars-rogue-one-director-gareth-edwards-praises-ai-film-making>. *Dream Machine* Issue 35.
101. *AOL*, “Emily Blunt refused Steven Spielberg’s offer to use AI for her alien voice in ‘Disclosure Day’.” <https://www.aol.com/articles/emily-blunt-refused-steven-spielbergs-163209822.html>. *Dream Machine* Issue 34.
102. *The Hollywood Reporter*, “‘Sheep in the Box’: Neon Drops First Trailer for Hirokazu Koreeda’s AI Family Drama.” <https://www.hollywoodreporter.com/movies/movie-news/sheep-in-the-box-neon-trailer-hirokazu-koreeda-ai-drama-1236629132/>. *Dream Machine* Issue 36.
103. *The Hollywood Reporter*, “Leonardo DiCaprio Says AI Can’t Be Art Because ‘There’s No Humanity to It’.” <https://www.hollywoodreporter.com/movies/movie-news/leonardo-dicaprio-ai-cant-be-art-no-humanity-1236445405/>. *Dream Machine* Issue 11.
104. *Daily Mail*, “Claire Foy says she has ‘no interest’ in seeing AI in films.” <https://www.dailymail.co.uk/tvshowbiz/article-15454199/Claire-Foy-AI-films-sad-disappointed-people-future-Hollywood.html>. *Dream Machine* Issue 14.
105. *NME*, “Jenna Ortega says it’s ‘very easy to be terrified’ of AI in filmmaking.” <https://www.nme.com/news/jenna-ortega-says-its-very-easy-to-be-terrified-of-ai-in-filmmaking-3913926>. *Dream Machine* Issue 10.
106. *Variety*, “Chris Pratt Pitched Having an AI ‘Actor’ Star as the Villain in ‘Mercy’: ‘I Don’t Think That’s a Good Idea at All’.” <https://variety.com/2026/film/news/chris-pratt-ai-actor-villain-mercy-amazon-mgm-1236640460/>. *Dream Machine* Issue 16.

07. *PC Gamer*, “Todd Howard says AI can’t replace human ‘creative intention,’ but it’s part of Bethesda’s ‘toolset for how we build our worlds or check things’.” <https://www.pcgamer.com/gaming-industry/todd-howard-says-ai-cant-replace-human-creative-intention-but-its-part-of-bethesas-toolset-for-how-we-build-our-worlds-or-check-things/>. *Dream Machine* Issue 11.
08. *Wired*, “Amazon’s House of David Used Over 350 AI Shots in Season 2. Its Creator Isn’t Sorry.” <https://www.wired.com/story/amazons-house-of-david-used-over-350-ai-shots-in-season-2-its-creator-isnt-sorry/>. *Dream Machine* Issue 7.
09. *GamesRadar*, “Battlefield 6 lead calls generative AI ‘very seducing,’ but says it was only used in the game’s earliest stages ‘to allow for more time and more space to be creative’.” <https://www.gamesradar.com/games/battlefield/battlefield-6-lead-calls-generative-ai-very-seducing-but-says-it-was-only-used-in-the-games-earliest-stages-to-allow-for-more-time-and-more-space-to-be-creative/>. *Dream Machine* Issue 3.
10. *gamesindustry.biz*, “Witcher 3 and Cyberpunk 2077 director says AI can help, but not replace, creatives.” <https://www.gamesindustry.biz/witcher-3-and-cyberpunk-2077-director-says-ai-can-help-but-not-replace-creatives>. *Dream Machine* Issue 9.
11. *GamesRadar*, “Aardman” *op. cit.*
12. *Dream Machine* Issue 29, on Sony’s “all in on AI for games” announcement.
13. *Variety*, “AI Dominates Cannes Buzz as Filmmakers Grudgingly Accept It.” <https://variety.com/2026/film/festivals/ai-cannes-2026-filmmakers-accept-1236748402/>; *Hollywood Reporter*, “At Cannes, filmmakers shift towards cautious acceptance of AI’s inevitability.” <https://www.hollywoodreporter.com/business/business-news/cannes-2026-ai-acceptance-1236592488/>. *Dream Machine* Issue 30.
14. *Variety*, “Is AI Basically Like Special Effects? Peter Jackson Seems to Think So.” <https://variety.com/2026/film/news/peter-jackson-ai-special-effects-1236748120/>. *Dream Machine* Issue 30.
15. *PC Gamer*, “Take-Two’s CEO says AI’s not in the business of making hits, ‘datasets by their very nature are backward looking’, but that doesn’t mean AI can’t be ‘super helpful’.” <https://www.pcgamer.com/games/take-two-ceo-ai-not-making-hits-backward-looking/>. *Business Insider*, “The CEO behind Grand Theft Auto says he’s pro AI — but the technology can’t make an original hit.” <https://www.businessinsider.com/take-two-ceo-strauss-zelnick-ai-original-hits-2026-5>. *Dream Machine* Issue 30.
16. Hasbro launches Sixth Wall, a new AI studio, with the CharacterOS governance layer underneath it, 3 June 2026. BusinessWire press release: <https://www.businesswire.com/news/home/20260603297922/en/>. *The Wrap*, “Hasbro AI Studio Sixth Wall Launch”: <https://www.thewrap.com/industry-news/business/hasbro-ai-studio-sixth-wall-launch/>. *Screen Daily*: <https://www.screendaily.com/news/hasbro-launches-subsidiary-ai-studio-sixth-wall-to-exploit-ip/5217380.article>. *BriefGlance* analysis (FTC inquiries into AI companions, Sixth Wall’s stated 13+/enterprise focus): <https://briefglance.com/articles/hasbros-sixth-wall-taming-the-ai-wild-west-with-behavioral-licensing>. CEO Roberta Thomson; CharacterOS encodes character canon, voice, personality, lore and safety limits at runtime; talent-participation pay-out structure for the original performers; initial 12-character catalogue; pilot grew out of an internal generative-AI Ouija-board experiment that logged hundreds of thousands of player interactions in its first few days.
17. Hasbro launches Sixth Wall, a new AI studio, with the CharacterOS governance layer underneath it, 3 June 2026. BusinessWire press release: <https://www.businesswire.com/news/home/20260603297922/en/>. *The Wrap*, “Hasbro AI Studio Sixth Wall Launch”: <https://www.thewrap.com/industry-news/business/hasbro-ai-studio-sixth-wall-launch/>. *Screen Daily*: <https://www.screendaily.com/news/hasbro-launches-subsidiary-ai-studio-sixth-wall-to-exploit-ip/5217380.article>. *BriefGlance* analysis (FTC inquiries

- into AI companions, Sixth Wall's stated 13+/enterprise focus): <https://briefglance.com/articles/hasbros-sixth-wall-taming-the-ai-wild-west-with-behavioral-licensing>. CEO Roberta Thomson; CharacterOS encodes character canon, voice, personality, lore and safety limits at runtime; talent-participation pay-out structure for the original performers; initial 12-character catalogue; pilot grew out of an internal generative-AI Ouija-board experiment that logged hundreds of thousands of player interactions in its first few days.
18. Spotify and Universal Music Group announce landmark linked recorded-music and publishing licensing agreements for fan-made AI covers and remixes at Spotify Investor Day, 21 May 2026. Spotify Newsroom: <https://newsroom.spotify.com/2026-05-21/universal-music-group-spotify-licensing-agreements-fan-made-covers-remixes/>. *TechCrunch*: <https://techcrunch.com/2026/05/21/spotify-and-universal-music-strike-deal-allowing-fan-made-ai-covers-and-remixes/>. *Medianama*: <https://www.medianama.com/2026/05/223-spotify-umg-fans-create-licensed-ai-covers-remixes/>. Three-pillar consent / credit / compensation framework; opt-in catalogue; Premium add-on; new royalty stream; UMG settlement with Udio in October 2025 is the architectural ancestor. Cross-reference in Chapter 5 and Chapter 6.
 19. *Deadline*, "Lionsgate's Michael Burns Says AI Will Save The Company 'Tens And Tens Of Millions Of Dollars A Year'." <https://deadline.com/2026/06/lionsgate-michael-burns-ai-saves-millions-runway-1236941043/>. *Dream Machine* Issue 33.
 20. Runway, "Runway and Lionsgate Expand Partnership" — Lionsgate takes an equity stake in Runway. <https://runwayml.com/news/runway-and-lionsgate-expand-partnership>. *Dream Machine* Issue 34.
 21. *Variety*, "AI Script-Analysis Start-Up Quilty Strikes Preferred-Look Deal With Giovanni Entertainment." <https://variety.com/2026/film/news/ai-script-quilty-preferred-look-deal-giovanni-entertainment-1236772175/>. *Dream Machine* Issue 34.
 22. *The Verge*, "Can AI Tell if your script will make a hit film?" <https://www.theverge.com/ai-artificial-intelligence/943531/ai-script-quilty-simon-horsman-daniel-wood>. *Dream Machine* Issue 33.
 23. Larian Studios policy framing on the next *Divinity*, January 2026; same source as note 33 above.
 24. PYMNTS, "AI Worries Erase \$1 Trillion From Big Tech Firms." <https://www.pymnts.com/artificial-intelligence-2/2026/ai-worries-erase-1-trillion-from-big-tech-firms/>. *Dream Machine* Issue 19.
 25. PYMNTS, "AI Now Responsible for a Third of US Economic Growth." <https://www.pymnts.com/economy/2026/ai-now-responsible-for-a-third-of-us-economic-growth/>. *Dream Machine* Issue 40.
 26. Grand View Research, "Generative AI Content Creation Market" (\$14.8 billion market in 2024, projected to quintuple to \$80 billion by 2030). <https://www.grandviewresearch.com/industry-analysis/generative-ai-content-creation-market-report>. *Dream Machine* Issue 4.
 27. *Music Ally*, "GenAI music firm Suno raises \$250m at \$2.45bn valuation." <https://musically.com/2025/11/19/genai-music-firm-suno-raises-250m-at-2-45bn-valuation/>. *Dream Machine* Issue 8. *Variety*, "AI Music Startup Suno Raises \$400M Funding at \$5.4 Billion Valuation." <https://variety.com/2026/digital/news/ai-music-suno-funding-round-400-million-5-4-billion-valuation-1236765727/>. *Dream Machine* Issue 32. *Digital Music News*, "Is Apple Plotting an AI-Powered Audio Expansion? Tech Giant Closes Reportedly \$1.6 Billion Qai Buyout — Its 'Second-Largest Acquisition' to Date." <https://www.digitalmusicnews.com/2026/01/30/apple-qai-acquisition/>. *Dream Machine* Issue 17.
 28. Peony, "10 Best Gaming Investors in 2026" (gaming venture funding fell from a \$12.5 billion peak in 2021 to \$2.54 billion in 2024). <https://www.peony.ink/blog/top-gaming-investors>.

29. Wikipedia, “2022–2026 video game industry layoffs” (an estimated 45,000 jobs lost from 2022 to July 2025). https://en.wikipedia.org/wiki/2022%E2%80%932026_video_game_industry_layoffs. *Game Developer*, “Data points to slowing layoffs, but doesn’t capture true harm to game industry” (Farhan Noor’s tracker: 8,500 layoffs in 2022; 10,500 in 2023; 14,600 in 2024; ~4,000 by mid-2025). <https://www.gamedeveloper.com/business/industry-layoffs-are-seemingly-slowing-but-the-damage-has-already-been-done>.
30. *PC Gamer*, “As the videogame industry continues to be hammered by layoffs, Netflix is offering up to \$840,000 per year for a new Director of Generative AI for Games.” <https://www.pcgamer.com/gaming-industry/as-the-videogame-industry-continues-to-be-hammered-by-layoffs-netflix-is-offering-up-to-usd840-000-per-year-for-a-new-director-of-generative-ai-for-games/>. *Dream Machine* Issue 1.
31. Drake Star Global Gaming Report Q2 2026, via *MarketScale*, “Gaming Investment Hits 12-Month High of \$2.5B+ in Q2 2026, Led by AI, AdTech, and Hardware Deals.” <https://www.marketscale.com/industries/sports-entertainment/gaming-investment-hits-12-month-high-of-25b-in-q2-2026-led-by-ai-adtech-and-hardware-deals>. Note: Drake Star characterises GTA 6’s November 2026 release as likely the largest game launch in history.
32. PitchBook, “Q2 2026 Gaming VC First Look.” <https://pitchbook.com/news/reports/q2-2026-gaming-vc-first-look>.
33. *Forbes*, “Five Venture Capital Trends: Where Gaming Capital Is Moving In 2026.” <https://www.forbes.com/councils/forbesbusinesscouncil/2026/07/31/five-venture-capital-trends-where-gaming-capital-is-moving-in-2026/>. OECD figure: OECD, “AI firms capture 61 percent of global venture capital in 2025.” <https://www.oecd.org/en/about/news/announcements/2026/02/ai-firms-capture-61-percent-of-global-venture-capital-in-2025.html>.
34. *GamesBeat*, “Fifth Door raises \$20M for AI-powered game creation.” <https://gamesbeat.com/fifth-door-raises-20m-for-ai-powered-game-creation/>. *Dream Machine* Issue 9. *PYMNTS*, “Razer Spends \$600 Million to Sharpen Focus on AI Gaming.” <https://www.pymnts.com/news/artificial-intelligence/2026/razer-spends-600-million-dollars-sharpen-focus-ai-gaming/>. *Dream Machine* Issue 13. *CTech*, “Sett secures \$30 million Series B to automate game marketing with AI agents.” <https://www.calcalistech.com/ctechnews/article/rjxyqisze>. *Dream Machine* Issue 23. *Astrocade*, “Astrocade raises \$56M to help build the next era of interactive entertainment.” <https://www.astrocade.com/blog/astrocade-raises-56m-funding>. *Dream Machine* Issue 28. *The AI Insider*, “AI Game Creation Platform Aippy Announces Tens of Millions at \$250M Valuation.” <https://theaiinsider.tech/2026/06/03/ai-game-creation-platform-aippy-announces-tens-of-millions-at-250m-valuation/>. *Dream Machine* Issue 32. *GamesBeat*, “Welevel unveils AI-based city-building survival game SolidRiver and \$8.5M in funding.” <https://gamesbeat.com/welevel-unveils-ai-based-city-building-survival-game-solidriver-and-8-5m-in-funding/>. *Dream Machine* Issue 34.
35. *TechCrunch*, “General Intuition lands \$134M seed to teach agents spatial reasoning using video game clips.” <https://techcrunch.com/2025/10/16/general-intuition-lands-134m-seed-to-teach-agents-spatial-reasoning-using-video-game-clips/>. *Dream Machine* Issue 3. *GamesBeat*, “General Intuition raises \$320M at \$2.3B valuation for AI frontier models based on gameplay.” <https://gamesbeat.com/general-intuition-raises-320m-at-2-3b-valuation-for-ai-frontier-models-based-on-gameplay-exclusive-interview/>. *TechCrunch*, “General Intuition’s \$2.3B bet that video games can train AI agents for the real world.” <https://techcrunch.com/2026/06/25/general-intuitions-2-3b-bet-that-video-games-can-train-ai-agents-for-the-real-world/>. *Dream Machine* Issue 36. *GamesBeat*, “Origin Lab raises \$8 million to glean AI training data from game worlds.” <https://gamesbeat.com/origin-lab-raises-8-million-turns-game-worlds-data-ai/>. *Dream Machine* Issue 29.

36. *Tech.eu*, “Cambridge’s Worldmodeldata raises £7M to turn video games into AI training data.” <https://tech.eu/2026/07/06/worldmodeldata-lands-ps7m-to-turn-gaming-data-into-ai-training/>. *TechCrunch*, “Meta quietly launches vibe-coded gaming app Pocket.” <https://techcrunch.com/2026/07/02/meta-quietly-launches-vibe-coded-gaming-app-pocket/>. Both in *Dream Machine* Issue 37.
37. Drake Star Global Gaming Report Q2 2026, via *MarketScale*, “Gaming Investment Hits 12-Month High of \$2.5B+ in Q2 2026, Led by AI, AdTech, and Hardware Deals.” <https://www.marketscale.com/industries/sports-entertainment/gaming-investment-hits-12-month-high-of-25b-in-q2-2026-led-by-ai-adtech-and-hardware-deals>. Note: Drake Star characterises GTA 6’s November 2026 release as likely the largest game launch in history.
38. On the Take-Two CEO and franchise origination, see Chapter 12, *The Orchestrator: NDTV Profit*, “Don’t Expect AI To Invent the Next ‘Grand Theft Auto’, Says Take-Two CEO Strauss Zelnick.” <https://www.ndtvprofit.com/technology/dont-expect-ai-to-invent-the-next-grand-theft-auto-says-take-two-ceo-strauss-zelnick>. *Dream Machine* Issue 6. On the internal AI team layoff: *Metro*, “GTA 6 publisher lays off its entire AI team after seven years.” <https://metro.co.uk/2026/04/06/gta-6-publisher-lays-off-entire-ai-team-after-seven-years-27874034/>. *Dream Machine* Issue 24. See also Chapter 14, *Coordination Collapse*.
39. *CNBC*, “Saudi PIF and Kushner’s Affinity finalize \$55 billion EA Sports deal.” <https://www.cnn.com/2026/08/05/saudi-wealth-fund-jared-kushner-ea-sports-deal.html>. Structure and completion (4 August 2026): *Wikipedia*, “Leveraged buyout of Electronic Arts.” https://en.wikipedia.org/wiki/Leveraged_buyout_of_Electronic_Arts.
40. Consortium AI framing: *eciks*, “Electronic Arts completes \$55B takeover, becomes private company.” <https://eciks.org/19135-ea-55-billion-takeover-private>. Player base of 700 million: AP via *ABC News*, “Video game giant Electronic Arts closes \$55 billion go-private sale of its business.” <https://abcnews.com/Technology/wireStory/video-game-giant-electronic-arts-closes-55-billion-135394606>. \$22bn AI-driven profit-headroom analysis: *Reuters*, “Gaming industry could unlock \$22 billion in profits on AI-driven cost cuts, says Morgan Stanley.” <https://www.reuters.com/business/gaming-industry-could-unlock-22-billion-profits-ai-driven-cost-cuts-morgan-2026-04-22/>. *Dream Machine* Issue 26.
41. Akker, “AI Film Insurance 2026: The Coverage Gap Hollywood Is Not Talking About.” <https://www.akerins.com/new-blog/ai-film-production-insurance-gap-2026>.
42. Roger Avari, *The Joe Rogan Experience*, Episode 2452 (February 2026), quoted in Akker, “AI Film Insurance 2026: The Coverage Gap Hollywood Is Not Talking About,” as note 15.
43. Wonder £9m seed: UKTN, “AI film studio Wonder lands £9m investment.” <https://www.uktech.news/ai/ai-film-studio-wonder-lands-9m-investment-20251023>. *Dream Machine* Issue 5. Wonder Studios at \$50m: *Deadline*, “Meet Wonder Studios, The \$50M British Studio Striving To Become The A24 Of AI Production.” <https://deadline.com/2026/05/wonder-studios-a24-ai-production-1236900795/>. *Dream Machine* Issue 30. Filmology Labs \$250m: *Variety*, “Filmology Labs, a \$250 Million Studio Built for Content Creation Like Vertical Micro-Dramas and AI-Driven Media Formats, Set to Open in New Jersey.” <https://variety.com/2026/digital/news/filmology-labs-content-creation-ai-studio-1236657646/>. *Dream Machine* Issue 18. Video Rebirth \$50m: *Crowdfund Insider*, “Singapore’s Video Rebirth Raises \$50m to Chase Studio-Grade AI Video.” <https://www.crowdfundinsider.com/2025/11/255411-singapores-video-rebirth-raises-50m-to-chase-studio-grade-ai-video/>. *Dream Machine* Issue 7. Flick \$6m YC seed: *Business Insider*, “This couple (literally) married AI and filmmaking. Read the pitch deck they used to raise \$6 million for their startup.” <https://www.businessinsider.com/ai-filmmaking-startup-flick-lands-yc-seed-funding-pitch-deck-2026-5>. *Dream Machine* Issue 30.
44. Ciaro, “Why Investors Fund Some Film Pitches and Ignore Others.” <https://ciaro.pro/blog/why-investors-fund-some-film-pitches-and-ignore-others>.

45. *DigiTimes*, “Seedance 2.0 cuts video generation cost below US\$0.14/sec.” <https://www.digitimes.com/news/a20260306PD225/cost-technology-production-development.html>. *Dream Machine* Issue 20.
46. *The Independent*, “Lights, camera, algorithm: First fully AI-generated film set to premiere at Tribeca Festival.” <https://www.independent.co.uk/arts-entertainment/films/news/ai-generated-film-tribeca-dreams-of-violet-b2984749.html>. *Dream Machine* Issue 31. *Digital Camera World*, “The CGI would have cost millions. I spent \$2,000.’ Is Dreams of Violets AI slop – or the future of film-making?” <https://www.digitalcameraworld.com/tech/artificial-intelligence/a-major-film-festival-will-soon-debut-the-first-fully-ai-generated-feature-inspired-by-real-events-this-could-be-a-major-shift-for-filmmaking>. *Dream Machine* Issue 32.
47. *Wall Street Journal*, “This Film Cost \$500,000 to Make. \$400,000 Was AI Compute Costs.” https://www.wsj.com/cio-journal/this-cannes-film-cost-500-000-to-make-400-000-was-ai-compute-costs-a823b08d?st=2a1WRJ&reflink=desktopwebshare_permalink. *Dream Machine* Issue 31.
48. *Decrypt*, “Killing Satoshi’: \$70M Bitcoin Film Starring Casey Affleck Will Make Heavy Use of AI.” <https://decrypt.co/364505/killing-satoshi-70m-bitcoin-film-starring-casey-affleck-will-make-heavy-use-of-ai>. *Dream Machine* Issue 26.
49. *TechCrunch*, “Runway CEO says AI could help Hollywood make 50 films instead of one \$100M blockbuster.” <https://techcrunch.com/2026/04/16/runway-ceo-says-ai-could-help-hollywood-make-50-films-instead-of-one-100m-blockbuster/>. *Dream Machine* Issue 26.
50. *Fortune*, “What Netflix’s acquisition of Ben Affleck’s AI filmmaking company really shows.” <https://fortune.com/2026/03/06/what-netflixs-acquisition-of-ben-afflecks-ai-filmmaking-company-really-shows/>. *Dream Machine* Issue 20. On the price and the Roald Dahl comparison: *TechCrunch*, “Netflix may have paid \$600 million for Ben Affleck’s AI startup.” <https://techcrunch.com/2026/03/11/netflix-may-have-paid-600-million-for-ben-afflecks-ai-startup/>. *Dream Machine* Issue 21. \$587m confirmed in *Variety*, “Netflix Paid \$587 Million for Ben Affleck’s AI Startup InterPositive.” <https://variety.com/2026/film/news/netflix-paid-587-million-ben-affleck-ai-interpositive-1236815111/>. *Dream Machine* Issue 39.
51. *SiliconANGLE*, “Higgsfield raises \$80M on \$1.3B valuation to scale AI video platform.” <https://siliconangle.com/2026/01/15/higgsfield-raises-80m-1-3b-valuation-scale-ai-video-platform/>. *Dream Machine* Issue 15. *36Kr*, “AI Video Unicorn Higgsfield: Earns \$200M in 9 Months by ‘Serving’ Social Media Marketers.” <https://eu.36kr.com/en/p/3650517574312323>. *Dream Machine* Issue 16. *Tech Times*, “AI Video Startup Higgsfield Hits \$500M Revenue, Eyes \$5B Funding Round.” <https://www.techtimes.com/articles/319394/20260630/ai-video-startup-higgsfield-hits-500m-revenue-eyes-5b-funding-round.htm>. *Dream Machine* Issue 36.
52. *Tech Times*, “Kling AI Raises Record \$2.8 Billion: What Creators Owe China Under Beijing Law.” <https://www.techtimes.com/articles/319708/20260704/kling-ai-raises-record-28-billion-what-creators-owe-china-under-beijing-law.htm>. *Dream Machine* Issue 37.
53. *Runway and Luma raises: The Hollywood Reporter*, “Hollywood’s AI Breaking Point May Be Nearing.” <https://www.hollywoodreporter.com/business/business-news/hollywood-ai-breaking-point-1236512451/>. *Runway Japan \$40m: Runway*, “Runway is expanding into Japan: opening a Tokyo office and investing an initial \$40M in the country.” <https://runwayml.com/news/runway-is-coming-to-japan>. *Dream Machine* Issue 30. *Studio and tool-maker film funds: Forbes*, “Hollywood Studios Are Spending On AI To Control The Future Of Film.” <https://www.forbes.com/sites/maureenkerr/2026/06/01/hollywood-studios-are-spending-on-ai-to-control-the-future-of-film/>. A24–DeepMind reaction: *Fortune*, “How Hollywood’s youngest filmmakers are exposing Gen Z’s real problem with AI.” <https://fortune.com/2026/07/05/why-doesnt-gen-z-trust-ai-box-office-youtube-directors-backrooms-obsession/>.
54. Google Veo enterprise positioning, SynthID provenance and Vertex distribution, and the 2026 model-market commoditisation: *Tech-Insider*, “Best AI Video Generator 2026: Veo vs Sora vs

- Kling.” <https://tech-insider.org/best-ai-video-generator-2026/>. Kling (Kuaishou) per-second pricing, roughly \$0.10/second per Get AI Perks, “Best AI Video Generators 2026: Sora 2 vs Veo 3.1 vs Kling 3.0 vs Runway.” <https://www.getai perks.com/en/blogs/44-best-ai-video-generators-2026>; the \$0.07/second floor is Digital Applied’s estimate, as note 28. Seedance cost curve: as note 19.
55. *Digital Applied*, “AI Video Market After Sora: Runway, Kling, and Veo.” <https://www.digitalapplied.com/blog/ai-video-market-after-sora-runway-kling-veo-2026>. OpenAI announced the shutdown on 24 March 2026; the app closed 26 April 2026. See also Chapter 7 on the Disney exit. Higgsfield trajectory: as note 25.
 56. *PocketGamer.biz*, “Reforged Labs shuts down as CEO says AI is closing its market opportunity.” <https://www.pocketgamer.biz/reforged-labs-shuts-down-as-ceo-says-ai-is-closing-its-market-opportunity/>. *Dream Machine* Issue 40.
 57. *TechCrunch*, “Design Arena creators raise \$7.9 million to bring ‘taste’ to AI models.” <https://techcrunch.com/2026/08/03/designarena-creators-raise-7-9-million-to-bring-taste-to-ai-models/>. *Dream Machine* Issue 40.
 58. Dubai \$1m award: *BroadcastPro ME*, “Dubai launches \$1m AI short film award in partnership with Google Gemini.” <https://www.broadcastprome.com/news/dubai-launches-1m-ai-short-film-award-in-partnership-with-google-gemini/>. *Dream Machine* Issue 2. XPRIZE Future Vision \$3.5m: *Fortune*, “Billionaire Peter Diamandis offers \$3.5 million to filmmakers who portray AI as the hero—not the villain.” <https://fortune.com/2026/03/09/xprize-foundation-future-vision-xprize-peter-diamandis-films-scifi-ai-fears/>; Google, “\$3.5 million Future Vision film competition.” <https://blog.google/innovation-and-ai/technology/ai/future-vision-film-competition-xprize/>. *Dream Machine* Issues 20 and 28. Astana \$1m prize fund: *The Astana Times*, “Astana AI Film Festival Opens Submissions, Offers \$1 Million Prize Fund.” <https://astanatimes.com/2026/05/astana-ai-film-festival-opens-submissions-offers-1-million-prize-fund/>. *Dream Machine* Issue 31. Luma Dream Brief \$1m: *Campaign Brief*, “Luma AI launches ‘The Luma Dream Brief,’ a \$1m global competition for Cannes Lions Gold Winner.” <https://campaignbrief.com/luma-ai-launches-the-luma-dream-brief-a-1m-global-competition-for-cannes-lions-gold-winner/>. *Dream Machine* Issue 17. Adobe \$10m grants: *Analytics India Magazine*, “Adobe Unveils AI Video Innovations + \$10 Million Grants Ahead of Sundance Film Festival.” <https://analyticsindiamag.com/ai-news-updates/adobe-unveils-ai-video-innovations-10-million-grants-ahead-of-sundance-film-festival/>. *Dream Machine* Issue 15. Google–Sundance \$2m: Google, “Sundance Institute & Google Partner to train filmmakers with AI.” <https://blog.google/company-news/outreach-and-initiatives/google-org/sundance-institute-ai-education/>. *Dream Machine* Issues 15 and 16. Runway Hundred Film Fund: *Forbes*, as note 26.
 59. Higgsfield, “Higgsfield AI Global Film Festival” — \$1m prize pool, 14 winners, \$500,000 for first place. <https://higgsfield.ai/global-film-festival>. *Dream Machine* Issue 41. The open-sourced feature *Hell Grind*: <https://x.com/higgsfield/status/2084702370764820572>. *Dream Machine* Issue 40.
 50. *Variety*, “Raindance Film Festival, AI Platform IMGN Launch New Film Fund and Development Hub for Indie Filmmakers.” <https://variety.com/2026/film/global/raindance-israel-ai-platform-imgn-development-fund-1236809594/>. *Dream Machine* Issue 39.
 61. *The Hollywood Reporter*, “ReelShort and More: The Microdrama TV Series Gold Rush Is Here.” <https://www.hollywoodreporter.com/business/business-news/microdrama-series-verticals-production-1236418912/>.
 52. See Chapter 7, §“Position One: All-in,” and its sources: Amazon AI Studios and *House of David*; Netflix INKubator; Disney’s Office of Technology Enablement. Netflix advertising: *Adweek*, “AI Agents Are Coming to Netflix to Grow Its \$3 Billion Ad Business.” <https://www.adweek.com/media/ai-agents-are-coming-to-netflix-to-grow-its-3-billion-ad-business/>. *Dream Machine* Issue 29.

53. *Variety*, “‘Squid Game’ Creator Makes Major Investment in AI Company TwelveLabs.” <https://variety.com/2025/tv/news/squid-game-creator-investment-ai-company-twelve-labs-1236536825/>. *Dream Machine* Issue 2.
54. See Chapter 7, §“Position One: All-in,” and Chapter 5, §“The pushback,” on the Amazon AI Creators Fund and the Punky Duck collapse.
55. Deloitte Insights (TMT Predictions 2026), “Tiny episodes, massive appeal: Short-form serials are gaining viewers and empowering independent studios.” <https://www.deloitte.com/us/en/insights/industry/technology/technology-media-and-telecom-predictions/2026/short-form-video-series.html>. Platform economics: Vitrina, “Micro Drama Distribution Platforms 2026.” <https://vitrina.ai/blog/micro-drama-distribution-platforms-2026/>. DramaBox raise: Real Reel, “Vertical Drama Weekly: Capital Tightens, Platforms Integrate.” <https://www.real-reel.com/vertical-drama-weekly-capital-platforms-2026/>. TikTok’s standalone microdrama app: *TechCrunch*, “TikTok quietly launches a micro-drama app called PineDrama.” <https://techcrunch.com/2026/01/16/tiktok-quietly-launches-a-micro-drama-app-called-pinedrama/>.
56. *The Hollywood Reporter*, “ReelShort and More: The Microdrama TV Series Gold Rush Is Here” (budgets; Fox–Holywater). <https://www.hollywoodreporter.com/business/business-news/microdrama-series-verticals-production-1236418912/>. Holywater first flagged in *Dream Machine* Issue 3. “First new professional format since reality TV”: C21Media, “Microdrama under the microscope.” <https://www.c21media.net/departments/c21investigates/microdrama-under-the-microscope/>. Distributor abstention: *Deadline*, “Why TV’s Traditional Sellers Aren’t Mad About Microdramas.” <https://deadline.com/2026/02/microdrama-vertical-video-tv-distributors-1236714399/>. Netflix vertical feed and AI recommendations: *TechCrunch*, “Netflix plans to add a vertical video feed, use AI for recommendations.” <https://techcrunch.com/2026/04/17/netflix-plans-to-add-a-vertical-video-feed-use-ai-for-recommendations/>. *Dream Machine* Issue 26.
67. Real Reel, “Vertical Drama Weekly: Capital Tightens, Platforms Integrate.” <https://www.real-reel.com/vertical-drama-weekly-capital-platforms-2026/>.
58. *Nikkei Asia*, “In China’s ‘microdrama’ boom, AI moviemakers chase 2 minutes of fame.” <https://asia.nikkei.com/spotlight/big-in-asia/in-china-s-microdrama-boom-ai-moviemakers-chase-2-minutes-of-fame>. *Dream Machine* Issue 39.
59. On the Meta Reality Labs cuts, see Chapter 14. On Vision Pro Personas and Gaussian splatting, see Chapter 9 and its sources.
70. Wikipedia, “Sphere (venue).” [https://en.wikipedia.org/wiki/Sphere_\(venue\)](https://en.wikipedia.org/wiki/Sphere_(venue)). *Yahoo Finance*, “Immersive Experience Bolsters Sphere Entertainment’s (SPHR) Strong Performance” (O’Keeffe Stevens Advisory Q2 2026 investor letter). <https://finance.yahoo.com/markets/stocks/articles/immersive-experience-bolsters-sphere-entertainment-145116003.html>. Sphere Entertainment IR, “Yas Island to Be Home of Sphere Abu Dhabi, a New Global Icon for Immersive Entertainment.” <https://investor.sphereentertainmentco.com/press-releases/news-details/2026/YAS-ISLAND-TO-BE-HOME-OF-SPHERE-ABU-DHABI-A-NEW-GLOBAL-ICON-FOR-IMMERSIVE-ENTERTAINMENT/default.aspx>. Sphere Entertainment IR, “Sphere Entertainment, the State of Maryland, Prince George’s County and Peterson Companies Announce Intent to Develop a Sphere at National Harbor.” <https://investor.sphereentertainmentco.com/press-releases/news-details/2026/Sphere-Entertainment-the-State-of-Maryland-Prince-Georges-County-and-Peterson-Companies-Announce-Intent-to-Develop-a-Sphere-at-National-Harbor/default.aspx>.
71. *Time*, “Netflix House: World’s Greatest Places 2026.” <https://time.com/collection/worlds-greatest-places/2026/netflix-house/>. *Event Marketer*, “Inside the Meteoric Rise of Immersive Entertainment Experiences.” <https://www.eventmarketer.com/article/fall-cover-story-immersive-entertainment-experiences/>.

72. *The Immersive Lab*, “Immersive Experiences 2025: The Year in Review.” <https://theimmersiveclub.com/2026/01/immersive-experiences-2025-year-in-review/>.
73. See Chapter 9 on Marble and the collapse in environment costs. Sony Pictures’ use of Marble in virtual production, 40× faster than the traditional workflow: https://www.linkedin.com/posts/brent-liang_tech-media-launch-ugcPost-7394911181091692546-TyUz. *Dream Machine* Issue 8. Netflix and Eyeline’s Vista4D: <https://github.com/Eyeline-Labs/Vista4D>. *Dream Machine* Issue 27.
74. World Labs, “World Labs Announces New Funding.” <https://www.worldlabs.ai/blog/funding-2026>. *Dream Machine* Issue 19. *Tech Times*, “Feifei Li’s World Labs Splits World Model Into Three Types: Marble Targets Simulation Linchpin.” <https://www.techtimes.com/articles/317927/20260606/feifei-lis-world-labs-splits-world-model-three-types-marble-targets-simulation-linchpin.htm>. Odyssey, “Odyssey has raised a \$310M Series B to accelerate world models.” <https://odyssey.ml/our-series-b>. *Fundraise Insider*, “Odyssey Raises \$310M Series B at \$1.45B for World Models.” <https://fundraiseinsider.com/blog/odyssey-raises-310m-series-b-at-1-45b-for-world-models/>. *Dream Machine* Issue 35. *Yahoo Finance*, “Tripo AI raises nearly \$200M in financing for AI 3D and world model tech.” <https://finance.yahoo.com/sectors/technology/articles/tripo-ai-raises-nearly-200-164700546.html>. *Dream Machine* Issue 32. *Memeburn*, “Odyssey AI Raises \$310M to Build World Models, Valued at \$1.45 Billion.” <https://memeburn.com/odyssey-ai-raises-310m-to-build-world-models-valued-at-1-45-billion/>.
75. ComfyUI, “We raised \$17 million to build an OS for Creative AI” (LinkedIn). https://www.linkedin.com/posts/comfyui_we-raised-17-million-to-build-an-os-for-ugcPost-7373743341236236288-wkCc. *Dream Machine* Issue 1. *TechCrunch*, “ComfyUI hits \$500M valuation as creators seek more control over AI-generated media.” <https://techcrunch.com/2026/04/24/comfyui-hits-500m-valuation-as-creators-seek-more-control-over-ai-generated-media/>. *Dream Machine* Issue 27. Four million users, \$30m round and \$48m raised in total: *Dream Machine* Issue 36, Editor’s Pick.
76. *TechTimes*, “Physics AI Slashes Engineering Simulation From Days to Seconds, PhysicsX Raises \$300M.” <https://www.techtimes.com/articles/318582/20260617/physics-ai-slashes-engineering-simulation-days-seconds-physicsx-raises-300m.htm>. *Dream Machine* Issue 35.
77. *Variety*, “Brad Pitt Says AI Could Help More ‘Mid-Budget Films’ Get Made if It’s ‘Used as a Tool’.” <https://variety.com/2026/film/news/brad-pitt-ai-could-help-mid-budget-movies-get-made-1236831201/>. *Dream Machine* Issue 41.
78. *C21Media*, “UK government unveils \$500m creative industries plan as AI proposals remain uncertain.” <https://www.c21media.net/news/uk-government-unveils-500m-creative-industries-plan-as-ai-proposals-remain-uncertain/>. Creative PEC, “Taking stock of the Creative Industries Sector Plan.” https://pec.ac.uk/blog_entries/taking-stock-of-the-creative-industries-sector-plan/. Creative UK, “The Creative Industries Sector Plan: One Year On.” <https://www.wearecreative.uk/the-creative-industries-sector-plan-one-year-on/>.
79. GOV.UK (DSIT), “A decisive shift to power British AI: new £1.1 billion plan to back chip firms, boost computing power and skills for the AI revolution.” <https://www.gov.uk/government/news/a-decisive-shift-to-power-british-ai-new-11-billion-plan-to-back-chip-firms-boost-computing-power-and-skills-for-the-ai-revolution>. *Dream Machine* Issue 33. Arts Council England, “AI technologies and emerging forms of creative practice” (£4m of projects funded since 2019). <https://www.artscouncil.org.uk/ai-technologies-and-emerging-forms-creative-practice>. *Dream Machine* Issue 17. Innovate UK, feasibility studies for frontier AI and ML – up to £3m competition for UK-registered SMEs. <https://apply-for-innovation-funding.service.gov.uk/competition/2419/overview/bb8d7d28-0e63-4ee1-bf4a-790dfb3ae02f>. *Dream Machine* Issue 23.

30. RWK Goodman, “UK Industrial Sector Strategy Plans — Creative Industries.” <https://www.rwkgoodman.com/info-hub/uk-industrial-sector-strategy-plans-creative-industries/>.
81. See Chapter 6, “The 88%,” on the consultation response breakdown and the December 2025 Statement of Progress.
32. World Labs, “Bringing Marble to Life.” <https://www.worldlabs.ai/case-studies/bringing-marble-to-life>. *Dream Machine* Issue 7, “Editor’s Pick: Marble by WorldLabs goes on public release,” 13 November 2025.
33. For a working primer on Gaussian splatting in the post-Marble era, see *Radiance Fields*, “World Labs Formally Launches Marble, A Generative World Model.” <https://radiancefields.com/world-labs-formally-launches-marble-a-generative-world-model>.
34. DreamLab AI Collective, beta participation in Marble, October–November 2025. Referenced in *Dream Machine* Issue 7: “DreamLab have been part of the beta testing for this over the last few months and it’s very neat.”
85. SuperSplat (PlayCanvas), open-source Gaussian splat editor, regular updates through 2025–26. *Dream Machine* Issue 1: “PlayCanvas open sources SOG — WebP for 3D Gaussian Splatting”; Issue 7 / Issue 11 on SuperSplat v2 updates.
36. Sony Pictures’ use of Marble in Virtual Production: https://www.linkedin.com/posts/brent-liang_tech-media-launch-ugcPost-7394911181091692546-TyUz. *Dream Machine* Issue 8.
87. Disney “300,000 poses in an instant” livestream, March 2026. *Dream Machine* Issue 23.
38. Netflix + Eyeline, *Vista4D*: 4D point clouds from live-action. *Dream Machine* Issue 27.
39. Google DeepMind, “Genie 3: A new frontier for world models.” <https://deepmind.google/blog/genie-3-a-new-frontier-for-world-models/>. Project Genie roll-out to AI Ultra subscribers in the U.S.: <https://blog.google/innovation-and-ai/models-and-research/google-deepmind/project-genie/>. *Dream Machine* Issue 3 (initial announcement) and Issue 17 (broader availability).
30. Meta, “WorldGen — text-to-immersive-3D-worlds research update.” <https://www.facebook.com/LifeAtMeta/videos/research-update-worldgen-text-to-immersive-3d-worlds/1879077432692421/>. *Dream Machine* Issues 9, 11.
91. Tencent, “HY World 1.5” announcement: <https://x.com/TencentHunyuan/status/2001170499133653006>. *Dream Machine* Issue 12.
92. SpAItial, *ECHO* spatial foundation model. <https://www.spaitial.ai/>. *Dream Machine* Issue 12.
93. Stanford AI Lab, *Wonderzoom*: Multi-Scale 3D World Generation. <https://wonderzoom.github.io/>. *Dream Machine* Issue 14.
94. OpenArt, *Worlds* product launch, March 2026. *Dream Machine* Issue 21.
95. NVIDIA SANA-WM, 2.6B open-source world model with 60-second video generation and camera control, May 2026. <https://huggingface.co/collections/nvidia/sana-wm>. *Dream Machine* Issue 30.
96. Odyssey, “Introducing Starchild-1, the first real-time multimodal world model.” <https://odyssey.ml/introducing-starchild-1>. *Dream Machine* Issue 30.
97. Odyssey, “Our Series B” — \$310M raise to accelerate world models, June 2026. <https://odyssey.ml/our-series-b>. *Dream Machine* Issue 35.
98. Apple Machine Learning Research, “Headsup: a large-scale high-quality 3D Gaussian head reconstruction from multi-view captures.” <https://machinelearning.apple.com/research/apple-headsup-3d-gaussian-head>. *Dream Machine* Issue 30.
99. WorldLens VR, “AI-powered 3D depth for Google Street View on Quest.” <https://www.uploadvr.com/worldlens-vr-quest-street-view-3d-depth/>. *Dream Machine* Issue 30.

30. Luma AI, *UNI-1* launch, March 2026. *Dream Machine* Issue 22, “Editor’s Pick: When worlds become instant, the race shifts to better thinking.”
01. ByteDance Seedance 2.0 in CapCut/Dreamina, March 2026. *Dream Machine* Issue 22.
02. *Spark 2.0*, open-source Gaussian-splat streaming framework, April 2026. *Dream Machine* Issue 25.
03. Radiance Fields, “Apple Confirms that it’s Gaussian Splatting that powers their personas.” <https://radiancefields.com/apple-confirms-personas-use-gaussian-splatting>. *Dream Machine* Issue 5.
04. Adobe Research, “Wonder” — real-time interactive 3D world navigation. <https://wonder-world-model.github.io/>. *Dream Machine* Issue 39.
05. Tencent, “WorldClaw” — agentic 3D open-world generation at scale. <https://tencent-hunyuan.github.io/Hunyuan3D-WorldClaw/>. *Dream Machine* Issue 41.
06. Reactor, “LingBot-World 2.0” — real-time world generation model. <https://www.reactor.inc/models/lingbot-world-2>. *Dream Machine* Issue 37 and Issue 38.
07. AlayaLab, “AlayaWorld: Long-Horizon and Playable Video World Generation.” <https://github.com/AlayaLab/AlayaWorld>. *Dream Machine* Issue 37.
08. Wayve, “GAIA-4” — the latest of Wayve’s line of world models. <https://wayve.ai/thinking/gaia-4/>. *Dream Machine* Issue 40.
09. Niantic Spatial and HMCI, “Building the Foundation for City of Rancho Cordova’s First Digital Twin for Physical AI.” <https://www.youtube.com/watch?v=TbPaQw-xLVc>. *Dream Machine* Issue 40.
10. Lightricks, “LTX-2.5” — an open-world, open-source model. <https://huggingface.co/Lightricks/LTX-2.5>. *Dream Machine* Issue 41.
11. Gracia AI, 4D Gaussian splatting in Unreal Engine with real-time relighting and self-shadowing. <https://gracia.ai/>. Atlux, Unreal Engine to 4D Gaussian Splat. <https://atlux.one/4dgs>. Radiance Fields, “Notch 2026.2 Adds Native Gaussian Splat Import and Rendering.” <https://radiancefields.com/notch-2026.2-adds-native-gaussian-splat-import-and-rendering>. *Dream Machine* Issue 37.
12. InfiniSplat (SIGGRAPH Asia 2026) — turn a single image into a navigable 3DGS. <https://zju3dv.github.io/InfiniSplat/>. Paper: <https://arxiv.org/abs/2608.02437>. *Dream Machine* Issue 40.
13. Insta360, “Insta360 X6” — introduces Spatial Capture: record a place in 360, bring it into the app, and it becomes a 3D model. <https://www.insta360.com/product/insta360-x6>. *Dream Machine* Issue 41.
14. *Video Games Chronicle*, “It honestly sucks: Fans think Call of Duty: Black Ops 7 is filled with generative AI art.” <https://www.videogameschronicle.com/news/it-honestly-sucks-fans-think-call-of-duty-black-ops-7-is-filled-with-generative-ai-art/>. *Video Games Chronicle*, “Ubisoft says AI-generated art in Anno 117 was a placeholder which ‘slipped through our review process.’” <https://www.videogameschronicle.com/news/ubisoft-says-ai-generated-art-in-anno-117-was-a-placeholder-which-slipped-through-our-review-process/>. *Polygon*, “Fortnite chapter 7 kicks off new controversy over AI art.” <https://www.polygon.com/fortnite-chapter-7-season-1-generative-ai-art-epic-games/>. *Dream Machine* Issues 8, 10.
15. NVIDIA + Stanford, *NitroGen*. <https://nitrogen.minedojo.org/>. *Dream Machine* Issue 13.
16. DeepMind, “SIMA 2: An Agent that Plays, Reasons, and Learns With You in Virtual 3D Worlds.” <https://deepmind.google/blog/sima-2-an-agent-that-plays-reasons-and-learns-with-you-in-virtual-3d-worlds/>. *Dream Machine* Issue 8.
17. ComfyUI Blog, “Ubisoft La Forge Open-Sources the CHORD Model and ComfyUI Nodes for End-to-End PBR Material Generation.” <https://blog.comfy.org/p/ubisoft-open-sources-the-chord>

model. *Dream Machine* Issue 11.

18. *Video Games Chronicle*, “The future of gaming, or ‘just a tool’? Hands-on with Teammates, Ubisoft’s ambitious voice AI tech demo.” <https://www.videogameschronicle.com/features/the-future-of-gaming-or-just-a-tool-hands-on-with-teammates-ubisofts-ambitious-voice-ai-tech-demo/>. *Dream Machine* Issue 9.
19. YouTube Playables Builder, closed-beta announcement: <https://www.youtube.com/playablesbuilder/>. *Dream Machine* Issue 12.
20. Unity AI Open Beta, in-editor AI suite, May 2026. *Dream Machine* Issue 28.
21. The Sandbox, “Announcing The Sandbox Studio: from idea to live multiplayer game in hours.” <https://www.youtube.com/watch?v=IdTVDRvQoWM>. *Dream Machine* Issue 33.
22. Higgsfield, “Higgsfield Games” — multiplayer games created and deployed from a single prompt. <https://higgsfield.ai/games-intro>. *Dream Machine* Issue 34.
23. Unreal Engine, “State of Unreal 2026: top news from the show.” <https://www.unrealengine.com/news/state-of-unreal-2026-top-news-from-the-show>. *Engadget*, “Epic Games details how it’s embracing generative AI in Unreal Engine.” <https://www.engadget.com/2196807/epic-games-details-how-its-embracing-gen-ai-in-unreal-engine/>. *Dream Machine* Issue 35. “Editor’s Pick: Unreal Engine’s AI Future — From Tools to Creative Partners,” 23 June 2026.
24. Roblox Newsroom, “Pioneering AI founders join to accelerate Roblox reality vision” — acquisition of Morpheus AI, Dynamics Lab and Lucid AI. <https://about.roblox.com/newsroom/2026/06/pioneering-ai-founders-join-to-accelerate-roblox-reality-vision>. *Dream Machine* Issues 32 and 34.
25. *Dream Machine* Issue 34, “Editor’s Pick: The AI World Cup — When Sport Becomes a Living Digital System,” 16 June 2026.
26. *GamesBeat*, “FIFA and Gamefam launch FIFA World Cup 2026 event in Super Soccer and Roblox.” <https://gamesbeat.com/gamefam-launch-fifa-world-cup-2026-super-soccer-roblox/>. *PocketGamer.biz*, “FIFA and Gamefam launch FIFA World Cup 2026 event across Roblox.” <https://www.pocketgamer.biz/fifa-and-gamefam-launch-fifa-world-cup-2026-event-across-roblox/>. *Dream Machine* Issue 34.
27. Sphere Entertainment, “Sphere Studios announces *The Rocky Horror Picture Show* at Sphere.” <https://www.sphereentertainmentco.com/sphere-studios-announces-the-rocky-horror-picture-show-at-sphere/>. *Dream Machine* Issue 35.
28. Korin AI, “trained with African datasets, built by Africans,” May 2026. *Dream Machine* Issue 27.
29. *Creative Boom*, “Adobe is putting AI in everything everywhere all at once.” <https://www.creativeboom.com/news/adobe-is-putting-ai-in-everything-everywhere-all-at-once/>. *Dream Machine* Issue 5, “Editor’s Pick,” 31 October 2025.
30. Adobe, “Adobe MAX 2025: Firefly Foundry.” <https://news.adobe.com/news/2025/10/adobe-max-2025-firefly-foundry>.
31. Adobe, “Adobe MAX 2025: Firefly.” <https://news.adobe.com/news/2025/10/adobe-max-2025-firefly>.
32. Adobe, “Adobe MAX 2025: Express AI Assistant.” <https://news.adobe.com/news/2025/10/adobe-max-2025-express-ai-assistant>.
33. *Wired*, “Adobe’s ‘Corrective AI’ Can Change the Emotions of a Voice-Over” and accompanying Adobe Sneaks 2025 coverage. <https://www.wired.com/story/adobe-max-sneaks-2025-corrective-ai/>. Project list compiled from MAX keynote and *Dream Machine* Issue 5 coverage.
34. *PYMNTS*, “Adobe Lets Users Design and Edit Using ChatGPT.” <https://www.pymnts.com/artificial-intelligence-2/2025/adobe-lets-users-design-and-edit-using-chatgpt/>. Adobe blog: “Edit images, designs, and PDFs right inside ChatGPT.”

- <https://blog.adobe.com/en/publish/2025/12/10/edit-photoshop-chatgpt>. *Dream Machine* Issue 12.
35. Adobe Premiere Object Mask tool: https://www.linkedin.com/posts/robdewinter_ok-this-is-going-to-save-a-lot-of-time-in-ugcPost-7421617551690063872-yKmB. *Dream Machine* Issue 16.
36. Adobe blog, “Sundance Film Festival 2026: Creativity, Community & Power of Storytelling.” <https://blog.adobe.com/en/publish/2026/01/20/sundance-film-festival-2026-creativity-community-power-storytelling>. *Dream Machine* Issue 16.
37. Adobe Summit 2026, “agentic creative intelligence” keynote. *Dream Machine* Issue 26.
38. *The Verge*, “Spotify is now an AI chatbot, too.” <https://www.theverge.com/entertainment/965358/spotify-ai-chatbot-interface-music-audiobooks-podcasts>. *Dream Machine* Issue 38.
39. *Search Engine Land*, “OpenAI appears to be building chatbot-native ads that launch AI agents.” <https://searchengineland.com/openai-appears-to-be-building-chatbot-native-ads-that-launch-ai-agents-484107>. *Dream Machine* Issue 41.
40. Adobe Newsroom, “Adobe unveils major expansion” — Adobe Creative Agent and AI Assistant in beta across Premiere, Photoshop, Illustrator, InDesign and Frame.io, June 2026. <https://news.adobe.com/news/2026/06/adobe-unveils-major-expansion>. *Dream Machine* Issue 35.
41. *MarTech Series*, “Adobe Introduces Firefly Foundry to Support Custom AI Models for Enterprise Brand Consistency.” <https://martechseries.com/predictive-ai/ai-platforms-machine-learning/adobe-introduces-firefly-foundry-to-support-custom-ai-models-for-enterprise-brand-consistency/>. *Dream Machine* Issue 36.
42. Adobe blog, “Disney Imagineers add Adobe Firefly Foundry to toolkit of worldbuilding superpowers.” <https://business.adobe.com/blog/disney-imagineers-add-adobe-firefly-foundry-to-toolkit>. *Dream Machine* Issue 35.
43. Adobe Newsroom, “Adobe to acquire Topaz Labs,” June 2026. <https://news.adobe.com/news/2026/06/adobe-to-acquire-topaz-labs>. *Dream Machine* Issue 36.
44. *Blog Mickey*, “How Disney Imagineering Is Using AI: Digital Twins, J.A.R.V.I.S., and Faster Ride Design.” <https://blogmickey.com/2026/08/damara-details-ai-use-imagineering-disney-company/>. *Dream Machine* Issue 41.
45. After Effects AI animation features through late 2025: *Dream Machine* Issue 9, “AI video is finally animatable inside After Effects.” https://www.linkedin.com/posts/thisisdoug_ai-aivideo-animation-ugcPost-7399512745924067330-Aldk.
46. *Dream Machine* Issue 21, “Editor’s Pick: Adobe and NVIDIA Just Raised the Stakes for Creative AI,” 19 March 2026.
47. NVIDIA + Google Cloud creative-AI infrastructure deal, March 2026. *Dream Machine* Issue 21.
48. Hugging Face and Google Cloud partnership announcement: https://www.linkedin.com/posts/julienchaumond_i-am-super-excited-to-announce-that-hugging-activity-739617740397276225-CuMM. *Dream Machine* Issue 8.
49. *EdTech Innovation Hub*, “Meta and Hugging Face launch OpenEnv to advance open-source agentic development.” <https://www.edtechinnovationhub.com/news/meta-and-hugging-face-launch-openenv-to-advance-open-source-agentic-development>. *Dream Machine* Issue 5.
50. Anthropic / Blender Foundation patronage, May 2026. *Dream Machine* Issue 27.
51. TechCrunch, “Anthropic launches interactive Claude apps, including Slack and other workplace tools.” <https://techcrunch.com/2026/01/26/anthropic-launches-interactive-claude-apps-including-slack-and-other-workplace-tools/>. *Dream Machine* Issue 16.
52. Spotify–Anthropic integration, May 2026. *Dream Machine* Issue 27.

53. *Music Business Worldwide*, “Splice inks ‘Responsible AI’ deal with ElevenLabs.” <https://www.musicbusinessworldwide.com/splice-elevenlabs-responsible-ai-deal/>. *Dream Machine* Issue 30.
54. *Adweek*, “Netflix ad tools could see ‘agentic AIs talking to each other’.” <https://www.adweek.com/media/netflix-ad-tools-agentic-ais-talking-to-each-other/>. *Dream Machine* Issue 30.
55. *Fortune*, “AI startup Viktor raises \$75 million to put a virtual ‘coworker’ in Slack and Teams.” <https://fortune.com/2026/05/19/ai-startup-viktor-75-million-virtual-coworker-slack-teams/>. *Dream Machine* Issue 30.
56. *MarTech Series*, “WPP continues AI overhaul with \$400-million Google partnership.” <https://martechseries.com/predictive-ai/ai-platforms-machine-learning/google-and-spotify-alum-launch-epiminds-with-6-6m-to-build-marketing-teams-for-the-ai-era/>. *Dream Machine* Issue 3.
57. *Campaign Brief*, “WPP launches AI-powered marketing platform WPP Open Pro.” <https://campaignbrief.com/wpp-launches-ai-powered-marketing-platform-wpp-open-pro/>. *Dream Machine* Issue 5.
58. *Digiday*, “WPP expands AI capabilities to boost brand performance with Sightly partnership.” <https://digiday.com/media-buying/agencies-continue-to-expand-ai-capabilities-to-boost-brand-performance/>. *Dream Machine* Issue 6.
59. WPP and Google Earth AI consumer-journey project, April 2026. *Dream Machine* Issue 27.
50. Forrester, “Google Dethrones OpenAI As Agencies’ Preferred AI Partner.” <https://www.forrester.com/blogs/google-dethrones-openai-as-agencies-preferred-ai-partner/>. *Dream Machine* Issue 35.
61. *Digiday*, “OpenAI moves to automate ad creative.” <https://digiday.com/marketing/openai-moves-to-automate-ad-creative/>. *Dream Machine* Issue 35.
62. Apple Newsroom, “Apple unveils next generation of Apple Intelligence, Siri AI and more,” June 2026. <https://www.apple.com/newsroom/2026/06/apple-unveils-next-generation-of-apple-intelligence-siri-ai-and-more/>. *Dream Machine* Issue 33.
63. Google I/O 2026 announcement block: Gemini Omni <https://blog.google/technology/google-deepmind/gemini-omni/>, Antigravity <https://antigravity.google/>, Google Flow <https://flow.google/>, Gemini Spark <https://blog.google/technology/developers/gemini-spark/>, Project Genie + Street View <https://deepmind.google/discover/blog/project-genie-street-view/>. *Dream Machine* Issue 30, “Editor’s Pick — Google I/O 2026,” 21 May 2026.
64. Google Labs, “Infinite Scaler.” <https://blog.google/technology/google-labs/infinite-scaler/>. *Dream Machine* Issue 30.
65. Google DeepMind, “SynthID — 100 billion watermarks, partner ecosystem.” <https://deepmind.google/discover/blog/synthid-100-billion-watermarks-partners/>. *Dream Machine* Issue 30.
56. *SiliconAngle*, “Higgsfield raises \$80M on \$1.3B valuation to scale AI video platform.” <https://siliconangle.com/2026/01/15/higgsfield-raises-80m-1-3b-valuation-scale-ai-video-platform/>. *Dream Machine* Issue 15.
67. *36kr*, “AI Video Unicorn Higgsfield: Earns \$200M in 9 Months by ‘Serving’ Social Media Marketers.” <https://eu.36kr.com/en/p/3650517574312323>. *Dream Machine* Issue 16.
58. *TechCrunch*, “Synthesia hits \$4B valuation, lets employees cash in.” <https://techcrunch.com/2026/01/26/synthesia-hits-4b-valuation-lets-employees-cash-in/>. *Dream Machine* Issue 16.
59. *Sifted*, “Synthesia rejects \$3bn Adobe acquisition offer.” <https://sifted.eu/articles/synthesia-acquisition-offer>. *Dream Machine* Issue 5.

70. ElevenLabs \$500m ARR reporting, April 2026. *Dream Machine* Issue 25.
71. Runway product cycle: Gen-4.5 (December 2025), Gen-4.5 Image-to-Video (January 2026), Workflows, Story Panels, Characters API, Apps for Advertising — *Dream Machine* Issues 10, 14, 15, 16, 20. Runway CEO on indie films vs. blockbusters: *Dream Machine* Issue 26.
72. Runway, “Runway Japan.” <https://runwayml.com/blog/runway-japan>. *Dream Machine* Issue 30.
73. For the running ledger of new creative-AI products through 2025–26, see *Dream Machine* Issues 1–30 archive.
74. ComfyUI, “We raised \$17 million to build an OS for Creative AI.” https://www.linkedin.com/posts/comfyui_we-raised-17-million-to-build-an-os-for-ugcPost-7373743341236236288-wkCc. *Dream Machine* Issue 1.
75. ComfyUI \$500M valuation, May 2026. *Dream Machine* Issue 27.
76. *Dream Machine* Issue 36, “Editor’s Pick: ComfyUI MCP — the open source engine behind AI’s creative future,” 2 July 2026. Background: *Medium* (Andrew Zhu), “The Rise of ComfyUI: How One Guy’s Side Project Took Over AI Image Generation.” <https://xhinker.medium.com/the-rise-of-comfyui-how-one-guys-side-project-took-over-ai-image-generation-558dca64bfd7>.
77. *Dream Machine* Issue 36, “Editor’s Pick: ComfyUI MCP — the open source engine behind AI’s creative future,” 2 July 2026. Background: *Medium* (Andrew Zhu), “The Rise of ComfyUI: How One Guy’s Side Project Took Over AI Image Generation.” <https://xhinker.medium.com/the-rise-of-comfyui-how-one-guys-side-project-took-over-ai-image-generation-558dca64bfd7>.
78. Comfy Org, “Comfy MCP” — public beta, July 2026. <https://comfy.org/mcp/>. *Dream Machine* Issue 36.
79. Google Pomelli launch: <https://x.com/GoogleLabs/status/1983204018567426312>. *Dream Machine* Issue 5.
30. Google AI Studio app gallery: <https://x.com/GoogleAIStudio/status/1982121563785949255>. Google Labs Opal expansion: <https://blog.google/technology/google-labs/opal-expansion/>. Project Genie: <https://blog.google/innovation-and-ai/models-and-research/google-deepmind/project-genie/>. *Dream Machine* Issues 5, 17.
81. Lovable for classrooms: <https://lovable.dev/classroom>. *Dream Machine* Issue 11.
32. Adobe Express AI Assistant: <https://news.adobe.com/news/2025/10/adobe-max-2025-express-ai-assistant>. *Dream Machine* Issue 5.
33. Hugging Face platform expansion through 2025–26.
34. Google blog, “Sundance Institute AI Education.” <https://blog.google/company-news/outreach-and-initiatives/google-org/sundance-institute-ai-education/>. *Dream Machine* Issue 15.
85. Adobe Ignite Day at Sundance: *Adobe blog*, *Sundance Film Festival 2026*. <https://blog.adobe.com/en/publish/2026/01/20/sundance-film-festival-2026-creativity-community-power-storytelling>. *Dream Machine* Issue 16.
36. Google’s \$40bn investment in Anthropic, May 2026. *Dream Machine* Issue 27.
87. UK Government, “Free AI training for all.” <https://www.gov.uk/government/news/free-ai-training-for-all-as-government-and-industry-programme-expands-to-provide-10-million-workers-with-key-ai-skills-by-2030>. *Dream Machine* Issue 16.
38. *CNBC*, “People with ADHD, autism, dyslexia say AI agents are helping them succeed at work.” <https://www.cnn.com/2025/11/08/adhd-autism-dyslexia-jobs-careers-ai-agents-success.html>. *Dream Machine* Issue 7.
39. University of Wisconsin-Stout, “AI Reshaping Industry: New UW-Stout Course Sets AI-Use as Baseline Competency in Filmmaking.” <https://www.uwstout.edu/about-us/news-center/ai-reshaping-industry-new-uw-stout-course-sets-ai-use-baseline-competency-filmmaking>. *Dream Machine* Issue 15.

30. *Forbes*, “Video AI Wars: How Chinese Labs Are Winning The Race OpenAI Abandoned,” 5 June 2026. <https://www.forbes.com/sites/edithyeung/2026/06/05/video-ai-wars-how-chinese-labs-are-winning-the-race-openai-abandoned/>. *Dream Machine* Issue 33.
91. Anthropic, “Fable and Mythos access” — suspension of frontier AI tools following US export-control directive, June 2026. <https://www.anthropic.com/news/fable-mythos-access>. *Dream Machine* Issue 34.
92. Doug McGinness on LinkedIn, late 2025, in the same post. *Dream Machine* Issue 9.
93. *Musically*, “AI-powered platform Eros Music Worlds is turning film characters into virtual singers.” <https://musically.com/2026/07/02/ai-powered-platform-eros-music-worlds-is-turning-film-characters-into-virtual-singers/>. *Dream Machine* Issue 37.
94. *Forbes*, “Homer Epic ‘The Odyssey’ Gets A Full-Length AI-Generated Film.” <https://www.forbes.com/sites/lesliekatz/2026/07/14/homer-epic-the-odyssey-gets-a-full-length-ai-generated-film/>. *Dream Machine* Issue 38.
95. National Portrait Gallery, “Interactive portrait of the nation by artist Es Devlin to launch at the National Portrait Gallery.” <https://www.npg.org.uk/about/press/interactive-portrait-of-the-nation-by-artist-es-devlin-to-launch-at-the-national-portrait-gallery>. *Dream Machine* Issue 40.
96. Snap Newsroom, “Snapchat Gen Z AI Creativity Research 2026.” <https://newsroom.snap.com/snapchat-gen-z-ai-creativity-research-2026>. *Dream Machine* Issue 30.
97. HeyGen, “AI Video Podcast.” <https://www.heygen.com/apps/ai-video-podcast>. *Dream Machine* Issue 39.
98. *WERSM*, “Wrinkles, a new app that Turns Local History Into An AI Audio Guide.” <https://wersm.com/wrinkles-ai-audio-guide-local-history/>. *Dream Machine* Issue 40.
99. *Dream Machine* Issue 13, “Editor’s Pick: The Year of the Orchestrator,” 9 January 2026.
100. *Dream Machine* Issue 29, May 2026, reporting on Sony’s 49-agent / 72-skill multi-agent game-development team.
01. *Bloomberg*, “AI Changed Chess. Grandmasters Now Win With Unpredictable Moves,” 27 March 2026. <https://www.bloomberg.com/news/articles/2026-03-27/ai-changed-chess-grandmasters-now-win-with-unpredictable-moves>. *Dream Machine* Issue 23. The chess analogy is developed in Chapter 16’s *Age of the Why* section.
02. Anthropic blog content on agent deployment patterns, Q1 2026.
03. *Variety*, “MiniMax Launches All-in-One AI Video Generator Hub at Shanghai Film Festival: ‘Creative Direction and Aesthetic Judgment Must Ultimately Be Left to Humans’.” <https://variety.com/2026/film/festivals/minimax-hub-ai-video-generator-shanghai-film-festival-1236781901/>. *Dream Machine* Issue 35.
04. Comfy Org, “Comfy MCP” public beta, July 2026. <https://comfy.org/mcp/>. *Dream Machine* Issue 36, Editor’s Pick.
05. Unreal Engine, “State of Unreal 2026: top news from the show” (Unreal MCP announcement). <https://www.unrealengine.com/news/state-of-unreal-2026-top-news-from-the-show>. *Dream Machine* Issue 35. Runway, “Runway in ChatGPT.” <https://runwayml.com/mcp>. *Dream Machine* Issue 34. Magnific, “Connecting Magnific to AI assistants using MCP.” <https://www.magnific.com/academy/tutorial/connecting-magnific-to-ai-assistants-using-mcp>. *Dream Machine* Issue 33. Higgsfield, “Higgsfield MCP.” <https://higgsfield.ai/mcp>. *Dream Machine* Issue 36.
06. *MusicRadar*, “You can now control FL Studio with an agentic AI chatbot from inside the DAW.” <https://www.musicradar.com/music-tech/you-can-now-control-fl-studio-with-an-ai-chatbot-from-inside-the-daw>. *Dream Machine* Issue 38.

07. *Music Ally*, “AI’s new USB port: what the MCP standard means for the music industry.” <https://musically.com/2026/07/20/ais-new-usb-port-what-the-mcp-standard-means-for-the-music-industry/>. *Dream Machine* Issue 39.
08. *Digital Music News*, “How AI Agents Are Quietly Automating Music Companies: From A&R to Release Ops.” <https://www.digitalmusicnews.com/2026/07/29/ai-agents-music-industry-pibox/>. *Dream Machine* Issue 39.
09. *TVBEurope*, “World first’ agentic AI TV platform debuts.” <https://www.tvbeurope.com/media-consumption/world-first-agentic-ai-tv-platform-debuts>. *Dream Machine* Issue 39.
10. *Filmera*, agentic AI filmmaking platform and workspace. <https://www.filmera.ai/>. *Dream Machine* Issue 41.
11. “How Spotify runs agents across 20M+ lines of code, with Niklas Gustavsson.” <https://www.youtube.com/watch?v=9DHZLW5653E>. *Dream Machine* Issue 40.
12. *Engadget*, “OpenAI launches a physical keypad for controlling agents.” <https://www.engadget.com/2215952/openai-launches-a-physical-keypad-for-controlling-agents/>. *Dream Machine* Issue 38.
13. *The Hollywood Reporter*, “Vibe Directing? An AI Startup Run by Ex-Googlers Say Their New Model Can Make It Happen (Exclusive).” <https://www.hollywoodreporter.com/business/digital/open-art-new-tool-model-ai-director-vibe-directing-1236628188/>. *Dream Machine* Issue 36.
14. *Variety*, “Backrooms’ Director Kane Parsons Says Using AI ‘Defeats the Purpose’ of Filmmaking: ‘I Get No Enjoyment From Using Those Tools’.” <https://variety.com/>. *Dream Machine* Issue 32.
15. Sundance Institute, “Centering the Artist: Why We’re Launching the AI Literacy Initiative.” <https://www.sundance.org/blogs/centering-the-artist-why-were-launching-the-ai-literacy-initiative/>. *Dream Machine* Issue 16.
16. Sundance Institute, *op. cit.*
17. Sundance Story Forum 2026 sessions on legal toolkits for producers using AI. *Dream Machine* Issue 16.
18. Google blog, “Sundance Institute AI Education.” <https://blog.google/company-news/outreach-and-initiatives/google-org/sundance-institute-ai-education/>. *Dream Machine* Issue 15.
19. McKinsey & Company, “What AI could mean for film and TV production and the industry’s future.” <https://www.mckinsey.com/industries/technology-media-and-telecommunications/our-insights/what-ai-could-mean-for-film-and-tv-production-and-the-industrys-future>. *Dream Machine* Issue 16.
20. *Metro*, “Prince of Persia remake and five more games cancelled as Ubisoft focuses on AI.” <https://metro.co.uk/2026/01/21/prince-persia-remake-five-games-cancelled-ubisoft-focuses-ai-26431926/>. *Dream Machine* Issue 15.
21. *PC Gamer*, “Square Enix, makers of Final Fantasy, aims to have AI doing 70% of its QA work by the end of 2027.” <https://www.pcgamer.com/gaming-industry/square-enix-aims-to-have-ai-doing-70-percent-of-its-qa-work-by-the-end-of-2027/>. *Dream Machine* Issue 7.
22. *Eurogamer*, “Falcom is the latest developer to buy into the AI hype machine.” <https://www.eurogamer.net/falcom-is-the-latest-developer-to-buy-into-the-ai-hype-machine>. *Dream Machine* Issue 12.
23. *NDTV Profit*, “Don’t Expect AI To Invent the Next ‘Grand Theft Auto’, Says Take-Two CEO Strauss Zelnick.” <https://www.ndtvprofit.com/technology/dont-expect-ai-to-invent-the-next-grand-theft-auto-says-take-two-ceo-strauss-zelnick>. *Dream Machine* Issue 6.
24. *Dream Machine* Issue 21, on Spielberg’s public position on AI.
25. *Dream Machine* Issues 25, 28, on Steven Soderbergh’s AI work.

26. *Digiday*, “Independent agencies face new frontier as agency-in-a-box tools democratize creativity.” <https://digiday.com/marketing/independent-agencies-face-new-frontier-as-agency-in-a-box-tools-democratize-creativity/>. *Dream Machine* Issues 6, 14.
27. *Digiday*, “AI agent developers have become adland’s in-demand role.” <https://digiday.com/marketing/ai-agent-developers-have-become-adlands-in-demand-role/>. *Dream Machine* Issue 7.
28. PYMNTS, “AI Content Is Par For The Course With PGA Tour’s Expanded AWS Partnership.” <https://www.pymnts.com/artificial-intelligence-2/2026/ai-content-is-par-for-the-course-with-pga-tours-expanded-aws-partnership/>. *Dream Machine* Issue 15.
29. *The Verge*, “Oreo-maker Mondelez will use AI for TV ads next year.” <https://www.theverge.com/news/806047/mondelez-ai-generated-ads>. *Dream Machine* Issue 5.
30. *Digiday*, “Avocados From Mexico turns to AI to advertise around the Super Bowl instead of a TV buy.” <https://digiday.com/marketing/avocados-from-mexico-turns-to-ai-to-advertise-around-the-super-bowl-instead-of-a-tv-buy/>. *Dream Machine* Issue 15.
31. *Dream Machine* Issue 26.
32. Reuters Institute, “AI adoption by UK journalists and their newsrooms: surveying applications, approaches, and attitudes.” <https://reutersinstitute.politics.ox.ac.uk/ai-adoption-uk-journalists-and-their-newsrooms-surveying-applications-approaches-and-attitudes>. *Dream Machine* Issue 9.
33. *Digiday*, “Daily Mail says Google AI Overviews have killed click-throughs.” <https://digiday.com/media/daily-mail-says-google-ai-overviews-have-killed-click-throughs/>. *Dream Machine* Issue 7.
34. *Digiday*, “How The Times is using AI to model synthetic focus groups from human audiences.” <https://digiday.com/media/how-the-times-is-using-ai-to-model-synthetic-focus-groups-from-human-audiences/>. *Dream Machine* Issue 6.
35. *TechBullion*, “Why the future belongs to multi-skilled leaders.” <https://techbullion.com/playing-the-long-game-with-a-portfolio-career-why-the-future-belongs-to-multi-skilled-leaders/>. *Dream Machine* Issue 9.
36. Anthropic Skills framework via Claude Code, reported through *Dream Machine* Issues 11, 16, 29.
37. *Forbes*, “AI Is Changing How Creators Work And Earn.” <https://www.forbes.com/sites/kolawolesamueladebayo/2025/12/22/how-ai-is-changing-how-creators-work-and-earn/>. *Dream Machine* Issue 13.
38. *Variety*, “AI Creator Behind Viral ‘Deadpool,’ ‘Harry Potter’ Christmas Clip Made His Film in a Ukrainian Bomb Shelter.” <https://variety.com/2026/digital/news/ai-video-deadpool-harry-potter-andrii-daniels-1236624632/>. *Dream Machine* Issue 16.
39. *Dream Machine* Issue 29 reportage of Tiny Grandma stop-motion content being wrongly flagged as AI by YouTube’s automated detection, May 2026.
40. *Dream Machine* Issue 23, April 2026, reporting death threats against Eline Van der Velden following Tilly Norwood’s continuing public role.
41. *Hollywood Reporter*, “Bobby Berk Says AI Will Make Reality TV & ‘Verifiably Human Content’ More Valuable.” <https://www.hollywoodreporter.com/tv/tv-news/bobby-berk-ai-reality-tv-1236592920/>. *Dream Machine* Issue 30.
42. *Digital Music News*, “Instagram Chief Says We Should ‘Fingerprint Real Media’ Instead of Tracking and Disclosing AI Slop.” <https://www.digitalmusicnews.com/2026/01/05/instagram-chief-ai-slop-comments/>. *WebProNews*, “Instagram Head Warns AI Images Erode Trust, Calls for Verification Standards.” <https://www.webpronews.com/instagram-head-warns-ai-images-erode-trust-calls-for-verification-standards/>. *Dream Machine* Issue 13.
43. *Digital Music News*, “AI-Generated Far-Right Hate Songs Aren’t Just a Problem in the US — Now They’re Spreading Across Europe Too.” <https://www.digitalmusicnews.com/2025/11/09/ai-generated-hate-songs-aren-t-just-a-problem-in-the-us-now-they-re-spreading-across-europe-too/>.

- [generated-hate-songs-dutch-spotify-charts/](#). *Dream Machine* Issue 7.
44. Google DeepMind SynthID watermark roll-out across Veo, Lyria and Imagen products. *Dream Machine* Issues 11, 12.
 45. Google DeepMind, “Verify Google AI-generated videos in the Gemini app.” https://www.linkedin.com/posts/googledeepmind_verify-google-ai-generated-videos-in-the-activity-7407748300688478208-fJgW. *Dream Machine* Issue 12; broader coverage in *SmartBrief*, “Google’s Gemini can now spot AI-generated videos.” <https://newsletter.smartbrief.com/sharedSummary/index.jsp?briefId=40A39351-5419-4681-94DF-31A53480B698&issueId=58E986AD-821F-422E-9E34-3386E0E2272B©Id=2DB8E453-8E83-416C-949B-44751F252A8D>. *Dream Machine* Issue 13.
 46. Google DeepMind, “SynthID — 100 billion watermarks, expanding to partner ecosystems including OpenAI, ElevenLabs and Kakao.” <https://deepmind.google/discover/blog/synthid-100-billion-watermarks-partners/>. *Dream Machine* Issue 30.
 47. ElevenLabs, “Detecting audio generated by ElevenLabs with SynthID” — partnership with Google DeepMind. <https://elevenlabs.io/blog/synthid>. *Dream Machine* Issue 36.
 48. YouTube announces automatic labelling of AI-generated videos, June 2026. *Dream Machine* Issue 31.
 49. LinkedIn announces it will limit the reach of AI-generated content on its feed, June 2026. *Dream Machine* Issue 31.
 50. AOL, “AI labels to be compulsory on authentic-looking content under EU rules.” <https://www.aol.co.uk/articles/ai-labels-compulsory-authentic-looking-094628000.html>. *Dream Machine* Issue 40.
 51. *Musically*, “Key EU AI Act rules now being enforced – so how does that affect the music industry?” <https://musically.com/2026/08/04/key-eu-ai-act-rules-now-being-enforced-so-how-does-that-affect-the-music-industry/>. *Dream Machine* Issue 40.
 52. Spotify Newsroom, “Spotify Introduces a New ‘AI Persona’ Label for AI-Generated Artist Identities.” <https://newsroom.spotify.com/2026-08-11/ai-persona-badges-transparency/>. *Dream Machine* Issue 41.
 53. Hasbro’s Sixth Wall launch and the CharacterOS governance layer, 3 June 2026. *BusinessWire*: <https://www.businesswire.com/news/home/20260603297922/en/>. *The Wrap*: <https://www.thewrap.com/industry-news/business/hasbro-ai-studio-sixth-wall-launch/>. Cross-reference Chapter 7 §“Position Five” and Chapter 6 §“The three Cs”.
 54. Spotify–Universal Music Group landmark licensing agreements, 21 May 2026. *Spotify Newsroom*: <https://newsroom.spotify.com/2026-05-21/universal-music-group-spotify-licensing-agreements-fan-made-covers-remixes/>. *TechCrunch*: <https://techcrunch.com/2026/05/21/spotify-and-universal-music-strike-deal-allowing-fan-made-ai-covers-and-remixes/>. See footnote ^37a in Chapter 6.
 55. *Variety*, “Nearly 1,000 Sign Open Letter Condemning Contract Clauses That Require Child Actors to Sign Over Their Voices for AI” (the *Peppa Pig* / Hasbro backlash). <https://variety.com/2026/tv/news/open-letter-major-studio-hasbro-children-ai-peppa-pig-1236790351/>. *Dream Machine* Issue 36.
 56. *The Hollywood Reporter*, “Warner Music Group Acquires AI Detection Company Sureel.” <https://www.hollywoodreporter.com/music/music-industry-news/warner-music-group-buys-ai-detection-company-sureel-1236618266/>. *Dream Machine* Issue 34.
 57. *Variety*, “Sony Invests in AI Copyright-Protection Firm Midnight Labs.” <https://variety.com/2026/biz/news/ip-ai-firm-midnight-labs-sony-million-dollar-investment-1236766931/>. *Dream Machine* Issue 33.

58. *Engadget*, “Deezer launches an AI music detector for other streaming services.” <https://www.engadget.com/2192024/deezer-ai-music-detector/>. *Dream Machine* Issue 34.
59. *Music Business Worldwide*, “Tech startup RTM Audio launches AI music detector UAI, which issues a signed certificate with every verdict.” <https://www.musicbusinessworldwide.com/tech-startup-rtm-audio-launches-ai-music-detector-uai-which-issues-a-signed-certificate-with-every-verdict/>. *Dream Machine* Issue 35.
50. *Music Ally*, “Played By Humans is verifying ‘human artistry in an AI age’.” <https://musically.com/2026/06/11/played-by-humans-is-verifying-human-artistry-in-an-ai-age/>. *Dream Machine* Issue 34.
61. *MusicTech*, “Traxsource to introduce labels to identify both human-made and AI-assisted music.” <https://musitech.com/news/music/traxsource-labels-identify-human-made-and-ai-assisted-music/>. *Dream Machine* Issue 36.
52. *Dream Machine* Issues 23, 27 reportage on Taylor Swift’s voice/image trademark filings.
53. *Billboard*, “Lionel Richie Is Trademarking the Sound of His Voice: ‘Hello, Is It Me You’re Looking For?’” <https://www.billboard.com/pro/lionel-richie-is-trademarking-the-sound-of-his-voice/>. *Dream Machine* Issue 34.
54. *Digital Music News*, “Backstreet Boys Look to Decommission AI Deepfakes by Trademarking Their Voices — Joining Taylor Swift, Lionel Richie, and Others.” <https://www.digitalmusicnews.com/2026/06/26/backstreet-boys-trademark/>. *Dream Machine* Issue 36.
65. *Songlines*, “A fake AI album released in my name made me realise why the law needs to change” (Emily Portman). <https://www.songlines.co.uk/features/emily-portman-faces-ai-fake-albums-on-streaming-platforms>. *Dream Machine* Issue 34.
56. *Lawyer Monthly*, “Matthew McConaughey Draws a Line to Protect His Voice and Image From AI.” <https://www.lawyer-monthly.com/2026/01/matthew-mcconaughey-protects-voice-image-ai/>. *Dream Machine* Issue 15.
67. *Adweek*, “Meet the \$1.3 Billion Startup Behind Madonna and Will Smith’s AI Video.” <https://www.adweek.com/media/higgsfield-ai-marketing-startup/>. *Dream Machine* Issue 16.
58. *Digital Music News*, “Madonna Calls AI ‘the Complete Opposite of Making Art’ While Advocating for a Completely Different Creative Process.” <https://www.digitalmusicnews.com/2026/06/28/madonna-ai-algorithm-commentary/>. *Dream Machine* Issue 36.
59. *Rolling Stone*, “The Rolling Stones Release New Single ‘In the Stars’ — With a Music Video De-Aging the Rockers Courtesy of AI.” <https://www.rollingstone.com/music/music-news/rolling-stones-in-the-stars-ai-de-aging-video-1235142200/>. *Hollywood Reporter*, “‘South Park’ Creators’ AI Company Made The Rolling Stones Young Again for ‘In The Stars’ Music Video.” <https://www.hollywoodreporter.com/tv/tv-news/south-park-creators-ai-rolling-stones-in-the-stars-1236592855/>. *Dream Machine* Issue 30.
70. *Music Business Worldwide*, “Boy George uses AI to re-record ‘Karma Chameleon’ as first release from tech startup Artist.” <https://www.musicbusinessworldwide.com/boy-george-uses-ai-to-re-record-karma-chameleon-as-first-release-from-tech-startup-artist-included/>. *Dream Machine* Issue 35.
71. *Variety*, “George Clooney Says AI Actors Will Face the ‘Same Problem We Have’ in Hollywood: ‘Making a Star Is Not So Easy’.” <https://variety.com/2025/scene/columns/george-clooney-ai-actors-movie-stars-1236579661/>. *Dream Machine* Issue 7.
72. *Deadline*, “AI Documentary Director Insists Jeremy Renner Agreed To Narrate Movie As ‘Hawkeye’ Star Threatens ‘Multi-Millions’ Lawsuit.” <https://deadline.com/2025/11/jeremy-renner-lawsuit-threat-ai-movie-1236611830/>. *Dream Machine* Issue 7.

73. *Variety*, “Cate Blanchett Co-Founders RSL Media, a Non-Profit to Address Consent Around AI Usage including creative work, name, image and likeness.” <https://variety.com/2026/film/news/cate-blanchett-rsl-media-ai-consent-1236748255/>. *Dream Machine* Issue 30.
74. *Variety*, “Cate Blanchett, Steven Soderbergh, Nikki Hexum on How the Human Consent Registry Helps People Protect Their Identity From AI Theft: It’s a ‘Human Right.’” <https://variety.com/2026/biz/news/cate-blanchett-steven-soderbergh-nikki-hexum-ai-human-registry-1236791271/>. *Dream Machine* Issue 36.
75. *Bloomberg*, “Apple Acquires Key Talent & Patents Behind AI Avatar Company ‘Animato.’” <https://www.bloomberg.com/news/articles/2026-05-19/apple-acquires-animato-ai-avatar-talent-patents>. *Dream Machine* Issue 30.
76. Estate of Ozzy Osbourne announces the late artist will “live on as an AI avatar that talks with fans,” June 2026. *Dream Machine* Issue 31.
77. *NBC News*, “Netflix says it will use AI to recreate the voice of the late Gene Wilder in a new reality TV series set in the world of Willy Wonka & the Chocolate Factory.” <https://www.nbcnews.com/pop-culture/pop-culture-news/gene-wilder-ai-voice-netflix-willy-wonka-competitor-show-rcna352440>. *Dream Machine* Issue 36.
78. *Complete Music Update*, “Johnny Cash estate uses ELVIS Act to sue Coke over tribute act ad soundtrack.” <https://completemusicupdate.com/johnny-cash-estate-uses-elvis-act-to-sue-coke-over-tribute-act-ad-soundtrack/>. *Dream Machine* Issue 9.
79. *The Verge*, “New York’s new law forces advertisers to say when they’re using AI avatars.” <https://www.theverge.com/news/842848/new-york-law-ai-advertisements-sag-aftra-labor>. *Dream Machine* Issue 11.
30. *Variety*, “More Than 16,000 Sign SAG-AFTRA Letter Demanding Congress Pass NO FAKES Act, Which Would Ban Unauthorized AI Images and Videos.” <https://variety.com/2026/biz/news/sag-aftra-letter-congress-pass-no-fakes-act-1236783056/>. *Dream Machine* Issue 35.
81. *Fast Company*, “Governments around the world are considering bans on Grok’s app over AI sexual image scandal.” <https://www.fastcompany.com/91474131/governments-around-the-world-are-considering-bans-on-groks-app-over-ai-sexual-image-scandal>. *Dream Machine* Issue 14.
32. Cannes AI Disclosure Standard, launched May 2026. *Dream Machine* Issue 29.
33. *Dream Machine* Issue 28, May 2026, reporting on the Academy of Motion Picture Arts and Sciences’ “You must be human to win” rule update.
34. *The Hollywood Reporter*, “Emmys Set AI Guidance.” <https://www.hollywoodreporter.com/tv/tv-news/emmys-ai-guidelines-2026-awards-1236468434/>. *Dream Machine* Issue 14.
85. SAG-AFTRA negotiation timeline through *Dream Machine* Issues 7, 12, 15, 20, 26, 29.
36. *Variety*, “AI Music Suno Funding Round \$400 Million \$5.4 Billion Valuation.” <https://variety.com/2026/digital/news/ai-music-suno-funding-round-400-million-5-4-billion-valuation-1236765727/>. *TechCrunch*, “Still Facing Copyright Lawsuits, AI Music Generator Suno Raises Another \$400M.” <https://techcrunch.com/2026/06/03/still-facing-copyright-lawsuits-ai-music-generator-suno-raises-another-400m/>. *Dream Machine* Issue 32.
87. Warner Music Group settled its copyright litigation against Suno and signed a licensing partnership; Universal Music Group and Sony Music remain in active litigation as of June 2026. *Dream Machine* Issue 32. For the structural context of the WMG settlement-as-licensing-table move, see Chapter 6 §“The three Cs”: *the lawsuits were never the destination; they were leverage to force a licensing table into existence.*

38. *Marketing Week*, “You can’t dismiss AI ads as slop when they’re winning in testing.” <https://www.marketingweek.com/dismiss-ai-ads-winning-creative-effectiveness/>. *Dream Machine* Issues 8, 13.
39. *The Drum*, “David Beckham Designs ‘Henchester United’ Chicken Coop in Lenovo Ad.” <https://www.thedrum.com/news/2026/05/18/david-beckham-henchester-united-chicken-coop-lenovo-ai-ad>. *Dream Machine* Issue 30.
40. *Variety*, “AI Creator Behind Viral ‘Deadpool,’ ‘Harry Potter’ Christmas Clip Made His Film in a Ukrainian Bomb Shelter.” *op. cit.* *Dream Machine* Issue 16.
41. PR Newswire, “From Apple TV Creative to AI Filmmaker: Hoyt Dwyer’s Animated Film To Compete at AI FilmFest Japan 2025.” *op. cit.* *Dream Machine* Issue 6.
42. Google DeepMind, “Dear Upstairs Neighbors.” <https://blog.google/innovation-and-ai/models-and-research/google-deepmind/dear-upstairs-neighbors/>. *Dream Machine* Issue 16.
43. *The Hollywood Reporter*, “‘Synthetic Sincerity’ by Marc Isaacs.” *op. cit.* *Dream Machine* Issue 8.
44. *Variety*, “‘Watch the Skies,’ Swedish UFO Feature Film Dubbed Entirely With AI, Sets USA Distribution Deal.” *op. cit.* *Dream Machine* Issue 5.
45. *Broadcast Pro Middle East*, “Tunisian filmmaker wins \$1 million AI Film Award for ‘Lily’.” *op. cit.* *Dream Machine* Issue 14.
46. *PC Gamer*, “More than 300 games released on Steam last week, and 120 of them had AI disclosures.” <https://www.pcgamer.com/gaming-industry/steam-week-in-review-more-than-300-games-released-on-steam-last-week-and-120-of-them-had-ai-disclosures/>. *Dream Machine* Issue 34.
47. *80.lv*, “Epic Games CEO Criticizes Valve’s AI Disclosure Policy for Steam.” <https://80.lv/articles/epic-ceo-criticizes-valve-s-ai-disclosure-policy-for-steam>. *Dream Machine* Issue 36.
48. Sundance Institute AI Literacy Initiative emphasis on documentation: <https://www.sundance.org/blogs/centering-the-artist-why-were-launching-the-ai-literacy-initiative/>. *Dream Machine* Issue 16.
49. *GameDev.net*, “Gamescom Dev speaker survey identifies potential impacts of AI on team structure and productivity.” <https://gamedev.net/news/5025/gamescom-dev-speaker-survey-identifies-potential-impacts-of-ai-on-team/>. *Dream Machine* Issue 41.
50. *Dream Machine* Issue 5, “Industry Insights: Stealth, Shadow and Secret AI Users.”
51. Azumo, “AI in Workplace Statistics 2025.” <https://azumo.com/artificial-intelligence/ai-insights/ai-in-workplace-statistics>. *Tech.co*, “Gen Z Most Likely Use AI Boss.” <https://tech.co/news/gen-z-most-likely-use-ai-boss>. *Dream Machine* Issue 5.
52. *Exploding Topics*, “AI Workforce Research.” <https://explodingtopics.com/blog/ai-workforce-research>. *Dream Machine* Issue 5.
53. *Forbes*, “AI Tools Flood Workplaces as Employees Face a Double Bind.” <https://www.forbes.com/sites/carolinecastrillon/2025/09/09/ai-tools-flood-workplaces-as-employees-face-a-double-bind/>. *Dream Machine* Issue 5.
54. *Blog IDC Europe*, “Shadow AI: How Stealth Productivity Is Strangling Enterprise AI Adoption and Creating a Security Nightmare.” <https://blog-idceurope.com/shadow-ai-how-stealth-productivity-is-strangling-enterprise-ai-adoption-and-creating-a-security-nightmare/>. *Dream Machine* Issue 5.
55. Enterprise-AI workforce tracking, late 2025. Aggregated in the Deep Dive companion piece *The Shadow AI Paradox in the Creative Industries*, drawing on Azumo’s *AI in Workplace Statistics 2025*, *Tech.co*’s Gen Z survey, and the IDC Europe shadow-AI security brief. *Dream Machine* Issue 5.

36. *Hidden Cloud Explosion* analysis, IDC Europe, 2025. See *The Shadow AI Paradox in the Creative Industries*, §“The Epistemology and Scale of Shadow AI.”
37. Shadow-AI security-incident statistics, 2025, aggregated in *The Shadow AI Paradox in the Creative Industries*, §“The Epistemology and Scale of Shadow AI”; underlying data via *IBM Cost of a Data Breach Report 2025* and IDC Europe.
38. For the developer-community origins of the “AI for thee, but not for me” phrasing, and the full sectoral analysis of the paradox, see *The Shadow AI Paradox in the Creative Industries*, §“The Great Hypocrisy.”
39. Survey of 1,100+ professional music creators, 2026, summarised in *Dynamics of Generative AI Adoption in the Creative Industries*, §“Music Production and Sound Recording,” and *The Shadow AI Paradox in the Creative Industries*, §“Sector-Specific Analysis.”
40. WGA screenwriter survey, pre- and post-strike, reported in *Dynamics of Generative AI Adoption in the Creative Industries*, §“Screenwriting and the Post-Strike AI Boom.”
41. Adobe Firefly milestone data, September 2023 – June 2025, in *Dynamics of Generative AI Adoption in the Creative Industries*, §“The Ubiquity of AI in Visual and Digital Arts.” *Dream Machine* Issue 6.
42. Adobe quarterly financials, FY2025–FY2026; AI-first ARR growth reported in *Dream Machine* Issue 21 and summarised in *Dynamics of Generative AI Adoption*.
43. Adobe Firefly enterprise penetration metrics, in *Dynamics of Generative AI Adoption*.
44. Adobe Stock submission analysis, 2024, in *Dynamics of Generative AI Adoption*.
45. Adobe, “Inaugural Adobe Creators’ Toolkit Report,” October 2025. <https://news.adobe.com/news/2025/10/adobe-max-2025-creators-survey>. *Dream Machine* Issue 6.
46. ChatGPT weekly-active-user disclosures, mid-2025; consolidated in *Dynamics of Generative AI Adoption*, §“General Purpose LLMs.”
47. Gemini desktop-user growth, year-over-year, in *Dynamics of Generative AI Adoption*.
48. Ramp, “Ramp AI Index,” June 2026. <https://ramp.com/data/ai-index-june-2026>. *Dream Machine* Issue 34.
49. *AppleInsider*, “Adobe’s glowing AI survey leaves out most of the creative industry.” <https://appleinsider.com/articles/26/06/16/adobes-glowing-ai-survey-leaves-out-most-of-the-creative-industry>. *Dream Machine* Issue 35.
50. Stanford AI Index Report 2025, global-sentiment chapter. Summarised in *Dynamics of Generative AI Adoption*, §“The Perception Gap.”
51. YouGov 2024 multi-market AI sentiment survey, 17 countries. Summarised in *Dynamics of Generative AI Adoption*, §“The Perception Gap.”
52. Quantic Foundry consumer-AI-in-gaming survey, 2025. Summarised in *Dynamics of Generative AI Adoption*, §“The Video Game Industry.”
53. Game Developers Conference *State of the Game Industry* surveys, 2024–2026, sentiment vs. usage trend. Reported in *Dynamics of Generative AI Adoption*, §“The Video Game Industry.”
54. O’Hare, J. (2026). *The Coordination Collapse: How AI Is Dissolving the Information Tax That Created Hierarchy — And What Replaces It*. DreamLab AI Consulting Ltd, April 2026. The report draws on Dell’Acqua et al.’s preregistered Boston Consulting Group experiment with 758 consultants, Mollick’s *secret cyborgs* coinage, Edmondson and Lei’s psychological-safety findings, the *IBM Cost of a Data Breach Report 2025*, the *Epoch AI / Ipsos KnowledgePanel* survey of March 2026, and the *CyberArk hidden org chart* analysis.
55. Epidemic Sound, *Future of the Creator Economy Report 2026* (survey of 3,000 professional creators across UK and US, June 2026); *Musically* commentary on the report and creator

attitudes towards AI. *Dream Machine* Issue 31, Editor's Pick, June 2026.

26. *Game Developer*, "Subnautica owner Krafton outlines plans to transform into an 'AI First' company." <https://www.gamedeveloper.com/business/subnautica-owner-krafton-outlines-plans-to-transform-into-an-ai-first-company>. *Dream Machine* Issue 6.
27. *PC Gamer*, "Open source game engine Godot will no longer accept AI-authored code contributions: 'We can't trust heavy users of AI to understand their code enough to fix it'." <https://www.pcgamer.com/gaming-industry/open-source-game-engine-godot-will-no-longer-accept-ai-authored-code-contributions-we-cant-trust-heavy-users-of-ai-to-understand-their-code-enough-to-fix-it/>. *Dream Machine* Issue 36.
28. *Fortune*, "Top CFOs warn AI success depends on training employees, not just buying technology." <https://fortune.com/2026/06/25/cfo-hpe-prologis-smucker-moodys-ai-success-training-employees/>. *Dream Machine* Issue 36.
29. KPMG, "Analysis of 1.4 million AI interactions identifies the employee behaviors behind effective AI use." <https://kpmg.com/us/en/articles/2026/sophisticated-ai.html>. *Dream Machine* Issue 37.
30. *Fortune*, "Sir Martin Sorrell has identified the key AI skill for the future: 'People who share will be the new kings and queens'." <https://fortune.com/2026/08/04/sir-martin-sorrell-has-identified-the-key-ai-skill-for-the-future-people-who-share-will-be-the-new-kings-and-queens/>. *Dream Machine* Issue 40. See also *The Drum*, "Martin Sorrell: Cannes is obsessing over the wrong AI debate." <https://www.thedrum.com/news/martin-sorrell-cannes-is-obsessing-over-the-wrong-ai-debate>. *Dream Machine* Issue 37.
31. *Dream Machine* Issue 24, April 2026, on the GTA VI publisher laying off its internal AI team.
32. *Dream Machine* Issue 25, April 2026, on Disney layoffs including Marvel staff.
33. *SmartBrief*, "Meta to cut 10% of Reality Labs staff to focus on AI." <https://newsletter.smartbrief.com/sharedSummary/index.jsp?briefId=40A39351-5419-4681-94DF-31A53480B698&issueId=025444D1-A590-46D8-B969-EF81DEE05228©Id=1B5F70D2-FFDA-4660-9CE9-047C9B16BF83>. *Dream Machine* Issue 14.
34. *CNBC*, "Microsoft cuts 2.1% of employees as Xbox unit plans to spin off studios." <https://www.cnbc.com/2026/07/06/microsoft-cuts-2point1percent-of-employees-as-xbox-unit-plans-to-spin-studios.html>. *Dream Machine* Issue 37.
35. *Dream Machine* Issue 23, April 2026, on Scottish animation studio collapse.
36. *Metro*, "Prince of Persia remake and five more games cancelled as Ubisoft focuses on AI." *op. cit.* *Dream Machine* Issue 15.
37. *The Guardian*, "AI is hitting UK harder than other big economies, study finds." <https://www.theguardian.com/technology/2026/jan/26/ai-uk-jobs-us-japan-germany-australia>. *Dream Machine* Issue 16.
38. *The Economist*, "Investors expect AI use to soar. That's not happening." <https://www.economist.com/finance-and-economics/2025/11/26/investors-expect-ai-use-to-soar-thats-not-happening>. *Dream Machine* Issue 9.
39. *Dream Machine* Issue 24, April 2026, on OpenAI's public-policy proposals around AI-driven economic disruption.
40. *The Economist*, "Job apocalypse? Humbug! AI is creating brand new occupations." <https://www.economist.com/business/2025/12/14/job-apocalypse-humbug-ai-is-creating-brand-new-occupations>. *Dream Machine* Issue 12.
41. *Forbes*, "Vibe Coding — The In Demand AI Skill." <https://www.forbes.com/sites/rachelwells/2025/11/06/the-in-demand-ai-skill-and-certifications-that-pays-up-to-220000/>. *Dream Machine* Issue 8.

42. U.K. Department for Business and Trade research on neurodiverse workers and AI assistants, autumn 2025. Reported via *CNBC*, “People with ADHD, autism, dyslexia say AI agents are helping them succeed at work.” <https://www.cnbc.com/2025/11/08/adhd-autism-dyslexia-jobs-careers-ai-agents-success.html>. *Dream Machine* Issue 7.
43. *CNBC*, *op. cit.*
44. *Dream Machine* Issue 7 secondary references.
45. Korin AI launch, May 2026. *Dream Machine* Issue 27.
46. *CNBC Africa*, “How AI is changing the landscape of the music industry in Africa.” <https://www.cnbc.com/2025/how-ai-is-changing-the-landscape-of-the-music-industry-in-africa>. *Dream Machine* Issue 5.
47. *BBC Future*, “Lights, camera, algorithm: Why Indian cinema is awash with AI.” <https://www.bbc.co.uk/future/article/20251223-why-indian-cinema-is-awash-with-ai>. *Dream Machine* Issue 14.
48. *Dream Machine* Issue 25, April 2026, on Indonesia’s *Legenda Bertuah*.
49. *Broadcast Pro Middle East*, “Tunisian filmmaker wins \$1 million AI Film Award for ‘Lily’.” *op. cit.* *Dream Machine* Issue 14.
50. *Digiday*, “Avocados From Mexico turns to AI to advertise around the Super Bowl instead of a TV buy.” *op. cit.* *Dream Machine* Issue 15.
51. *Variety*, “AI Creator Behind Viral ‘Deadpool,’ ‘Harry Potter’ Christmas Clip Made His Film in a Ukrainian Bomb Shelter.” *op. cit.* *Dream Machine* Issue 16.
52. *Dream Machine* Issue 8 citing Andreessen Horowitz observations: https://www.linkedin.com/posts/stevenouri_a-wild-stat-80-of-startups-pitching-a16z-activity-739618271898351872-xTKR.
53. *PocketGamer.biz*, “Shift Up CEO says AI is key to competing with China’s game industry scale.” <https://www.pocketgamer.biz/shift-up-ceo-says-ai-is-key-to-competing-with-chinas-game-industry-scale/>. *Dream Machine* Issue 14.
54. *The Economist*, “Investors expect AI use to soar. That’s not happening.” <https://www.economist.com/finance-and-economics/2025/11/26/investors-expect-ai-use-to-soar-thats-not-happening>. *Dream Machine* Issue 9.
55. *The Economist*, “Job apocalypse? Humbug! AI is creating brand new occupations.” <https://www.economist.com/business/2025/12/14/job-apocalypse-humbug-ai-is-creating-brand-new-occupations>. *Dream Machine* Issue 12.
56. *The Guardian*, “AI is hitting UK harder than other big economies, study finds.” <https://www.theguardian.com/technology/2026/jan/26/ai-uk-jobs-us-japan-germany-australia>. *Dream Machine* Issue 16.
57. University of Wisconsin-Stout, “AI Reshaping Industry: New UW-Stout Course Sets AI-Use as Baseline Competency in Filmmaking.” <https://www.uwstout.edu/about-us/news-center/ai-reshaping-industry-new-uw-stout-course-sets-ai-use-baseline-competency-filmmaking>. *Dream Machine* Issue 15.
58. *Labor Market Matters*, “AI and independent work in 7 charts” — what if AI’s first labour-market effect isn’t replacing workers, but making traditional firms less necessary? <https://www.labormarketmatters.com/p/ai-and-independent-work-in-7-charts>. *Dream Machine* Issue 37.
59. Adobe Firefly enterprise metrics, tracked across the *Dream Machine* archive.
50. Reuters Institute, “AI adoption by UK journalists and their newsrooms.” <https://reutersinstitute.politics.ox.ac.uk/ai-adoption-uk-journalists-and-their-newsrooms-surveying-applications-approaches-and-attitudes>. *Digiday*, “Daily Mail says Google AI Overviews

- have killed click-throughs.” <https://digiday.com/media/daily-mail-says-google-ai-overviews-have-killed-click-throughs/>. *Dream Machine* Issues 7, 9.
61. 1,100-creator music survey 2026 (music production and sound recording).
 62. VFX AI integration metrics (visual-effects automation).
 63. *PC Gamer*, “Square Enix aims to have AI doing 70% of its QA work by the end of 2027.” <https://www.pcgamer.com/gaming-industry/square-enix-aims-to-have-ai-doing-70-percent-of-its-qa-work-by-the-end-of-2027/>. *Dream Machine* Issue 7.
 64. *Eurogamer*, “Falcom is the latest developer to buy into the AI hype machine.” <https://www.eurogamer.net/falcom-is-the-latest-developer-to-buy-into-the-ai-hype-machine>. *Dream Machine* Issue 12.
 65. *Metro*, “Prince of Persia remake and five more games cancelled as Ubisoft focuses on AI.” <https://metro.co.uk/2026/01/21/prince-persia-remake-five-games-cancelled-ubisoft-focuses-ai-26431926/>. *Dream Machine* Issue 15.
 66. *Dream Machine* Issue 24, April 2026, on the GTA VI publisher laying off its internal AI team.
 67. ElevenLabs \$500m ARR reporting, April 2026. *Dream Machine* Issue 25.
 68. *36kr*, “AI Video Unicorn Higgsfield: Earns \$200M in 9 Months by ‘Serving’ Social Media Marketers.” <https://eu.36kr.com/en/p/3650517574312323>. *Dream Machine* Issue 16.
 69. *Dream Machine* Issue 29, May 2026, on Sony’s 49-agent / 72-skill multi-agent game-development team.
 70. *Digiday*, “AI agent developers have become adland’s in-demand role.” <https://digiday.com/marketing/ai-agent-developers-have-become-adlands-in-demand-role/>. *Dream Machine* Issue 7.
 71. *Ad Age*, “PMG’s hackathon at Cannes signals agencies’ growing appetite for AI talent.” <https://adage.com/technology/ai/aa-pmg-cannes-hackathon-agency-skilled-talent/>. *Dream Machine* Issue 36.
 72. *Forbes*, “Vibe Coding — The In Demand AI Skill That Pays Up to \$220,000.” <https://www.forbes.com/sites/rachelwells/2025/11/06/the-in-demand-ai-skill-and-certifications-that-pays-up-to-220000/>. *Dream Machine* Issue 8.
 73. Sundance Institute, “Centering the Artist: Why We’re Launching the AI Literacy Initiative.” <https://www.sundance.org/blogs/centering-the-artist-why-were-launching-the-ai-literacy-initiative/>. Google blog, “Sundance Institute AI Education.” <https://blog.google/company-news/outreach-and-initiatives/google-org/sundance-institute-ai-education/>. *Dream Machine* Issues 15, 16.
 74. UK Government, “Free AI training for all.” <https://www.gov.uk/government/news/free-ai-training-for-all-as-government-and-industry-programme-expands-to-provide-10-million-workers-with-key-ai-skills-by-2030>. *Dream Machine* Issue 16.
 75. Lovable for classrooms. <https://lovable.dev/classroom>. *Dream Machine* Issue 11.
 76. UW-Stout course launch, January 2026 — *op. cit.*
 77. Anthropic, “Claude Corps: teaching 1,000 fellows to master Claude, matched with nonprofits across America.” <https://www.anthropic.com/news/claude-corps>. *Dream Machine* Issue 34.
 78. GOV.UK, “AI Adoption Plan for the Creative Industries.” <https://www.gov.uk/government/publications/ai-champions-ai-adoption-plans/ai-adoption-plan-creative-industries>. *Dream Machine* Issue 33.
 79. Adobe, “Firefly Foundry.” <https://news.adobe.com/news/2025/10/adobe-max-2025-firefly-foundry>. *Dream Machine* Issue 5.
 30. Korin AI launch, May 2026. *Dream Machine* Issue 27.

81. *The Verge*, “New York’s new law forces advertisers to say when they’re using AI avatars.” <https://www.theverge.com/news/842848/new-york-law-ai-advertisements-sag-aftra-labor>. *Dream Machine* Issue 11. C2PA / SynthID infrastructure references in Chapter 13.
32. *Forbes*, “AI Is Changing How Creators Work And Earn.” <https://www.forbes.com/sites/kolawolesamueladebayo/2025/12/22/how-ai-is-changing-how-creators-work-and-earn/>. *Dream Machine* Issue 13.
33. *Variety*, “AI Creator Behind Viral ‘Deadpool,’ ‘Harry Potter’ Christmas Clip Made His Film in a Ukrainian Bomb Shelter.” <https://variety.com/2026/digital/news/ai-video-deadpool-harry-potter-andrii-daniels-1236624632/>. *Dream Machine* Issue 16.
34. *Broadcast Pro Middle East*, “Tunisian filmmaker wins \$1 million AI Film Award for ‘Lily.’” <https://www.broadcastpromedia.com/news/tunisian-filmmaker-wins-1-million-ai-film-award-for-lily/>. *Dream Machine* Issue 14.
85. *Dream Machine* Issue 25, April 2026, on Indonesia’s *Legenda Bertuah*.
36. *BBC Future*, “Lights, camera, algorithm: Why Indian cinema is awash with AI.” <https://www.bbc.co.uk/future/article/20251223-why-indian-cinema-is-awash-with-ai>. *Dream Machine* Issue 14.
87. *The Hollywood Reporter*, “Hollywood Workers Are Training AI Models as Job Prospects Grow Slim.” <https://www.hollywoodreporter.com/business/digital/ai-training-hollywood-writer-jobs-prospects-1236628302/>. *Dream Machine* Issue 36.
38. *TechBullion*, “Why the future belongs to multi-skilled leaders.” <https://techbullion.com/playing-the-long-game-with-a-portfolio-career-why-the-future-belongs-to-multi-skilled-leaders/>. *Dream Machine* Issue 9.
39. Anthropic Skills framework via Claude Code. *Dream Machine* Issues 11, 16, 29.
30. *UKTN*, “AI hiring surges ‘across all UK sectors’, according to new findings from PwC, with a 34% salary premium for AI skills.” <https://www.uktech.news/ai/ai-hiring-surges-across-all-uk-sectors-20260615>. *Dream Machine* Issue 34.
91. Adobe, “Inaugural Adobe Creators’ Toolkit Report,” October 2025. <https://news.adobe.com/news/2025/10/adobe-max-2025-creators-survey>. *Dream Machine* Issue 6.
92. PRS for Music, “PRS for Music AI Survey 2026.” <https://www.prsformusic.com/m-magazine/news/prs-for-music-ai-survey-2026>. *Dream Machine* Issue 16.
93. *CNBC*, “People with ADHD, autism, dyslexia say AI agents are helping them succeed at work.” <https://www.cnn.com/2025/11/08/adhd-autism-dyslexia-jobs-careers-ai-agents-success.html>. *Dream Machine* Issue 7.
94. McKinsey & Company, “What AI could mean for film and TV production and the industry’s future.” <https://www.mckinsey.com/industries/technology-media-and-telecommunications/our-insights/what-ai-could-mean-for-film-and-tv-production-and-the-industrys-future>. *Dream Machine* Issue 16.
95. GDC State of the Game Industry surveys 2024–2026.
96. GDC, “GDC 2026 State of the Game Industry Reveals Impact of Layoffs, Generative AI and More” — 52% of developers say generative AI is having a negative impact on the industry, up from 30% the year before and 18% the year before that. <https://gdconf.com/article/gdc-2026-state-of-the-game-industry-reveals-impact-of-layoffs-generative-ai-and-more/>. *Dream Machine* Issue 41.
97. *PC Gamer*, “EA: AI is delivering ‘faster prototyping’ and a ‘real rise in creativity’ in its studios.” <https://www.pcgamer.com/software/ai/ea-exec-says-ai-has-helped-drive-a-real-rise-of-creativity-at-its-studios/>. *Dream Machine* Issue 36.
98. LANDR AI music study, late 2025, referenced via *Ari’s Take*. <https://aristake.com/ai-tools-musicians-study/>. *Dream Machine* Issue 8.

99. Stanford AI Index Report 2025.
10. YouGov 2024 multi-market AI sentiment survey.
01. *Digital Music News*, “Nearly 800 Creatives Sign Responsible AI Declaration — ‘Stealing Our Work Is Not Innovation’.” <https://www.digitalmusicnews.com/2026/01/22/stealing-isnt-innovation/>. *Dream Machine* Issue 16.
02. *Digital Trends*, “The AI game-dev boom has a dark side: fewer jobs for the people starting out.” <https://www.digitaltrends.com/cool-tech/the-ai-game-dev-boom-has-a-dark-side-fewer-jobs-for-the-people-starting-out/>. *Dream Machine* Issue 39.
03. *GamesBeat*, “Teen founder raises \$3.2M for AI game development startup ForgeGUI.” <https://gamesbeat.com/teen-founder-raises-3-2m-for-ai-game-development-startup-forgegui/>. *Dream Machine* Issue 39. And *Yahoo*, “My 11-year-old son used AI to build his own video game — I see it as a creativity boost, not a threat.” <https://www.yahoo.com/lifestyle/articles/11-old-son-used-ai-110601483.html>. *Dream Machine* Issue 37.
04. *The Hollywood Reporter*, “Hollywood Workers Are Training AI Models as Job Prospects Grow Slim.” <https://www.hollywoodreporter.com/business/digital/ai-training-hollywood-writer-jobs-prospects-1236628302/>. *Dream Machine* Issue 36.
05. *Broadcast Pro Middle East*, Lily award — *op. cit.*
06. *Variety*, Andrii Daniels bomb-shelter clip — *op. cit.*
07. *BBC Future*, “Lights, camera, algorithm” — *op. cit.*
08. *Dream Machine* Issue 25, Indonesian *Legenda Bertuah*.
09. *CNBC Africa*, “How AI is changing the landscape of the music industry in Africa.” <https://www.cnbc africa.com/2025/how-ai-is-changing-the-landscape-of-the-music-industry-in-africa>. *Dream Machine* Issue 5. Korin AI launch, May 2026 — *op. cit.*
10. *PocketGamer.biz*, “Shift Up CEO says AI is key to competing with China’s game industry scale.” <https://www.pocketgamer.biz/shift-up-ceo-says-ai-is-key-to-competing-with-chinas-game-industry-scale/>. *Dream Machine* Issue 14.
11. The doomer frame is set out at length in OpenAI’s spring 2026 policy paper on the “social contract for the intelligence age” (covered in *Dream Machine* Issue 24) and in Daniel Susskind’s *A World Without Work*. The full empirical rebuttal — including the Stanford Digital Economy Lab’s *Canaries in the Coal Mine?* findings (Brynjolfsson, Chandar & Chen, 2025), Brynjolfsson, Li & Raymond’s *Generative AI at Work* (QJE, 2025; n = 5,179 agents), Acemoglu’s *Simple Macroeconomics of AI* (Economic Policy, 2025), the Dallas Fed’s labour analyses, the St. Louis Fed’s AI-and-GDP tracking, the 2026 Stanford AI Index, and the 230-year history of identical predictions — is set out at length in DreamLab’s published research. That analysis is the analytical antibody to the policy framing this section is pushing back against.
12. *Bloomberg*, “AI Changed Chess. Grandmasters Now Win With Unpredictable Moves,” 27 March 2026. <https://www.bloomberg.com/news/articles/2026-03-27/ai-changed-chess-grandmasters-now-win-with-unpredictable-moves>. *Dream Machine* Issue 23. The behavioural pattern the piece describes — top grandmasters deliberately deviating from machine-optimal lines to put opponents on uncomputed ground — is the cleanest available analogy I have for the strategic shift the rest of this chapter argues for.
13. *Digital Music News*, “The AI Licensing Shift — Creative Weight Attribution Emerges as Music Industry Game-Changer for Rights Holders.” *op. cit.* *Dream Machine* Issue 16.
14. *Arts Council England*, “Arts Council England launches Responsible AI Cultural Leaders Academy.” <https://www.artscouncil.org.uk/research-and-data/responsible-ai-arts-council-england>. *Dream Machine* Issue 37.
15. DreamLab AI Collective, team page. <https://dreamlab-ai.com/team>.

16. OpenAI, “Sora 2 is here,” 30 September 2025. <https://openai.com/index/sora-2/>. *Dream Machine* Issue 1.
17. LinkedIn News aggregation: “Sora Tops 1 Million Downloads in 5 Days.” <https://www.linkedin.com/news/story/sora-tops-1m-downloads-in-5-days-6684988/>. *Dream Machine* Issue 3.
18. OpenAI’s staged wind-down of Sora, announced 24 March 2026 — consumer app and website ending late April, developer API closing late September; peak revenue reportedly ~\$540k per month against operating costs in the billions. Sources and the corrected Disney–OpenAI record are set out in DreamLab’s published research on the programmable brand.
19. Google DeepMind, Veo 3.1 launch, mid-October 2025. *Dream Machine* Issue 3.
20. Runway product cycle: Gen-4.5 (December 2025), Gen-4.5 Image-to-Video (January 2026), Workflows, Story Panels, Characters API, Apps for Advertising — *Dream Machine* Issues 10, 14, 15, 16, 20.
21. Runway CEO on indie films vs. blockbusters, *Dream Machine* Issue 26.
22. Chinese open-source AI video model releases, 2025–2026. *Dream Machine* Issues 3, 12, 22.
23. SiliconAngle, “Higgsfield raises \$80M on \$1.3B valuation.” <https://siliconangle.com/2026/01/15/higgsfield-raises-80m-1-3b-valuation-scale-ai-video-platform/>. 36kr, “Higgsfield: Earns \$200M in 9 Months.” <https://eu.36kr.com/en/p/3650517574312323>. *Dream Machine* Issues 15, 16.
24. Heygen Video Agent. <https://www.linkedin.com/posts/heygen-introducing-the-new-video-agent-activity-7421597801240801282-d1CF>. *Dream Machine* Issue 16.
25. TechCrunch, “Synthesia hits \$4B valuation, lets employees cash in.” <https://techcrunch.com/2026/01/26/synthesia-hits-4b-valuation-lets-employees-cash-in/>. Sifted, “Synthesia rejects \$3bn Adobe acquisition offer.” <https://sifted.eu/articles/synthesia-acquisition-offer>. *Dream Machine* Issues 5, 16.
26. ElevenLabs \$500m ARR reporting, April 2026. *Dream Machine* Issue 25.
27. Google DeepMind, “Introducing Gemini Omni: Create Anything from Any input.” <https://blog.google/technology/google-deepmind/gemini-omni-launch/>. *Dream Machine* Issue 30.
28. Beeples Canvas — Generative AI compositor. <https://www.beeple-canvas.com/>. *Dream Machine* Issue 30.
29. The Decoder, “ByteDance’s Seedance 2.5 breaks the 30-second barrier for AI video generation.” <https://the-decoder.com/bytedances-seedance-2-5-breaks-the-30-second-barrier-for-ai-video-generation/>. Dreamina Seedance 2.0 Mini: <https://dreamina.capcut.com/seedance/seedance-2-0-mini>. Seedance 2.0 Mini and 4K in ComfyUI: <https://blog.comfy.org/p/seedance-2-0-mini-and-4k-is-now-available>. *Dream Machine* Issues 34, 35, 36.
30. Runway API Recipes: <https://runwayml.com/api>. Runway MCP in ChatGPT: <https://runwayml.com/mcp>. Runway Agent 2 ad localisation: <https://runwayml.com/news/introducing-agent-2>. *Dream Machine* Issues 34, 35, 36.
31. Adobe Firefly milestone data, from DreamLab’s published research on the dynamics of generative AI adoption, §“The Ubiquity of AI in Visual and Digital Arts.”
32. Nano Banana inside Photoshop and inside Unreal Engine cross-integrations, October–November 2025. *Dream Machine* Issue 1.
33. Suno Studio launch. <https://www.techradar.com/ai-platforms-assistants/i-tried-suno-studio-the-new-platform-that-mixes-ai-music-generation-with-hands-on-editing-like-garageband-but-smarter>. *Dream Machine* Issue 1.
34. Mureka, “Music Agent Studio” launch. *Dream Machine* Issue 4.

35. ElevenLabs Series funding, April 2026. *Dream Machine* Issue 25.
36. *MusicTech*, “Cardiff band speaks out after AI artist trained on their music outperforms them on Spotify.” <https://musictech.com/news/industry/its-shocking-disheartening-and-insulting-cardiff-band-speaks-out-after-ai-artist-trained-on-their-music-outperforms-them-on-spotify/>. *Dream Machine* Issue 1.
37. *Variety*, “AI Creator Behind Viral ‘Deadpool,’ ‘Harry Potter’ Christmas Clip Made His Film in a Ukrainian Bomb Shelter.” <https://variety.com/2026/digital/news/ai-video-deadpool-harry-potter-andrii-daniels-1236624632/>. *Dream Machine* Issue 16.
38. Sony AI, “Woosh — a sound effect foundation model.” <https://ai.sony/blog/woosh-sound-effect-foundation-model/>. *Dream Machine* Issue 30.
39. Mirelo SFX 1.6, “edit sound, not just generate it.” <https://mirelo.ai/sfx-1-6>. *Dream Machine* Issue 30.
40. Stability AI, “Stable Audio 3.0 released — open-weight model family built for artistic experimentation.” <https://stability.ai/news/stable-audio-3-0-released>. *Dream Machine* Issue 30.
41. Tamber product page: <https://tamber.ai/>. *Dream Machine* Issue 30.
42. Beatport Track ID. <https://www.beatport.com/track-id>. *Dream Machine* Issue 30.
43. Google Magenta RealTime 2: <https://magenta.withgoogle.com/magenta-realtime-2>. Daydream: <https://daydream.live/>. *MusicRadar* on Fender Studio Pro and Moises: <https://www.musicradar.com/music-tech/daws/ai-isnt-the-destination-making-music-is-fender-looks-forward-as-its-studio-pro-daw-integrates-moises-stem-separation-and-a-smart-studio-assistant>. *MusicTech* on the Maono PD500W: <https://musictech.com/news/gear/maono-pd500w/>. *Dream Machine* Issues 33, 35.
44. Suno, “Introducing Spark.” <https://suno.com/blog/introducing-spark>. *Dream Machine* Issue 36.
45. Music industry AI deal flow, October 2025 — May 2026. See Chapter 5 footnotes 31–37, and *Dream Machine* Issues 5, 7, 8, 12, 14, 16, 17.
46. World Labs, “Bringing Marble to Life.” <https://www.worldlabs.ai/case-studies/bringing-marble-to-life>. *Dream Machine* Issue 7.
47. Sony Pictures Marble VP integration. *Dream Machine* Issue 8.
48. Google DeepMind, “Genie 3.” <https://deepmind.google/blog/genie-3-a-new-frontier-for-world-models/>. Project Genie: <https://blog.google/innovation-and-ai/models-and-research/google-deepmind/project-genie/>. *Dream Machine* Issues 3, 17.
49. Tencent, “HY World 1.5” and Hunyuan 3D Studio. *Dream Machine* Issue 12.
50. Luma AI, *UNI-1* launch, March 2026. *Dream Machine* Issue 22.
51. NVIDIA SANA-WM model collection. <https://huggingface.co/collections/nvidia/sana-wm>. *Dream Machine* Issue 30.
52. Odyssey, “Introducing Starchild-1, the first real-time multimodal world model.” <https://odyssey.ml/introducing-starchild-1>. *Dream Machine* Issue 30.
53. Odyssey, “Introducing Agora-1 — four-player AI-generated world built on a 1997 shooter.” <https://odyssey.ml/introducing-agera-1>. *Dream Machine* Issue 30.
54. Apple Machine Learning Research, “Apple Headsup: a Large-Scale High-Quality 3D Gaussian Head Reconstruction from Multi-View Captures.” <https://machinelearning.apple.com/research/apple-headsup-3d-gaussian-head>. *Dream Machine* Issue 30.
55. SuperSplat / Spark 2.0 / SOG releases through 2025–26. *Dream Machine* Issues 1, 25.
56. Radiance Fields, “Apple Confirms that it’s Gaussian Splatting that powers their personas.” <https://radiancefields.com/apple-confirms-personas-use-gaussian-splatting>. *Dream Machine*

Issue 5.

57. Radiance Fields, “Apple Maps Flyover is getting a Gaussian splatting upgrade.” <https://radiancefields.com/apple-maps-flyover-is-getting-a-gaussian-splatting-upgrade>. *Dream Machine* Issue 33.
58. ComfyUI Blog, “Ubisoft La Forge Open-Sources the CHORD Model.” <https://blog.comfy.org/p/ubisoft-open-sources-the-chord-model>. *Dream Machine* Issue 11.
59. Anthropic / Blender Foundation patronage, May 2026. *Dream Machine* Issue 27.
60. OpenAI, “Introducing AgentKit.” <https://openai.com/index/introducing-agentkit/>. *Dream Machine* Issue 2.
61. Anthropic Skills framework. *Dream Machine* Issues 11, 16, 29.
62. Google, “Official skills for AI agents.” <https://github.com/google/agent-skills>. *Dream Machine* Issue 30.
63. Tencent Ardot, AI-native design agent platform. <https://ardot.tencent.com/>. *Dream Machine* Issue 30.
64. Heygen Video Agent. *Dream Machine* Issue 16.
65. Adobe Summit 2026 CX Enterprise. *Dream Machine* Issue 26.
66. Adobe + NVIDIA / Google + NVIDIA partnerships. *Dream Machine* Issue 21.
67. ComfyUI funding round. <https://www.linkedin.com/posts/comfyui-we-raised-17-million-to-build-an-os-for-ugcPost-7373743341236236288-wkCc>. *Dream Machine* Issue 1.
68. ComfyUI \$500M valuation, May 2026. *Dream Machine* Issue 27.
69. Anthropic, “Claude is now available as a partner node in ComfyUI.” <https://www.anthropic.com/news/claude-comfyui-partner-node>. *Dream Machine* Issue 30.
70. Hugging Face / Google Cloud and Meta / Hugging Face OpenEnv. *Dream Machine* Issues 5, 8.
71. Comfy MCP public beta. <https://comfy.org/mcp/>. *Dream Machine* Issue 36.
72. Epic Games, “State of Unreal 2026: top news from the show.” <https://www.unrealengine.com/news/state-of-unreal-2026-top-news-from-the-show>. Magnific MCP: <https://www.magnific.com/academy/tutorial/connecting-magnific-to-ai-assistants-using-mcp>. Higgsfield MCP: <https://higgsfield.ai/mcp>. Runway MCP: <https://runwayml.com/mcp>. *Dream Machine* Issues 33, 35, 36.
73. Adobe, “Adobe unveils major expansion” — Creative Agent and AI Assistant in beta across Premiere, Photoshop, Illustrator, InDesign and Frame.io. <https://news.adobe.com/news/2026/06/adobe-unveils-major-expansion>. *Dream Machine* Issue 35.
74. Unreal Engine 5 official AI Assistant. <https://www.linkedin.com/posts/wouterweynants-theres-an-official-ai-assistant-coming-to-ugcPost-7369377204226379776-pGiH>. *Dream Machine* Issue 1.
75. ECABridge — Unreal Engine MCP integration. <https://ecabridge.dev/>. *Dream Machine* Issue 30.
76. Epic Games, “State of Unreal 2026: top news from the show.” <https://www.unrealengine.com/news/state-of-unreal-2026-top-news-from-the-show>. *Dream Machine* Issue 35.
77. *Video Games Chronicle*, “Epic Games Veteran Claims He’s Building AI-Heavy ‘Fully European’ Game Engine.” <https://www.videogameschronicle.com/news/epic-games-veteran-ai-heavy-fully-european-game-engine/>. *Dream Machine* Issue 30.
78. Unity AI Council (October 2025); Unity AI Open Beta (May 2026). *Dream Machine* Issues 1, 28.
79. VFX AI integration metrics, from DreamLab’s published research on the dynamics of generative AI adoption, §“Visual Effects (VFX) Automation.”

30. Anthropic / Blender Foundation patronage. *Dream Machine* Issue 27.
81. Andreessen Horowitz pitch-deck observations on Chinese open-source model usage. https://www.linkedin.com/posts/stevenouri_a-wild-stat-80-of-startups-pitching-a16z-activity-7396182718998351872-xTKR. *Dream Machine* Issue 8.
32. NVIDIA SANA-WM model collection. <https://huggingface.co/collections/nvidia/sana-wm>. *Dream Machine* Issue 30.
33. PhotoGIMP — the open-source GIMP skin that mimics Photoshop. <https://github.com/Dioliinux/PhotoGIMP>. *Dream Machine* Issue 30.
34. Korin AI launch, May 2026. *Dream Machine* Issue 27.

dream machine: the new creative economy is the field guide to the most consequential year in creative work since cinema learned to talk. written from inside the work by creative technologist pete woodbridge, this book is not just a tools guide or a five-point manifesto. it is a careful, public record of an industry coming apart and putting itself back together.

drawing on his vantage point running a studios in the north west of england and publishing a weekly newsletter for thousands of creatives, woodbridge holds the whole picture together in a way no platform-company keynote or academic has managed.

whether you are a working creative trying to figure out what to do on monday morning, a studio leader making organisational decisions, or a policymaker setting the rails for the future, this book provides the frameworks you need to stay oriented in a fast-moving environment.

creative work is being re-platformed. the choices made right now will set the terms for the next decade.

dreamlab.org.uk