DREAM MACHINE / DEEP ANALYSIS

The Doomer Mistake

Why the case for AI-driven civilizational collapse keeps losing the argument with the data

An analysis grounded in the labour data, the productivity record, and three centuries of failed technology panics

Executive summary

AI doomerism — the broad family of arguments holding that artificial intelligence will hollow out the labour market, dissolve the social contract, and leave humanity economically redundant — has become the consensus mood of 2026. Sam Altman is *publicly negotiating* a post-work settlement. Dario Amodei tells interviewers the current economic model may need to be replaced. Newspaper headlines treat a 13% drop in employment for AI-exposed 22-to-25-year-olds as a leading indicator of civilizational replacement.

This analysis argues that the consensus is wrong. Not in every particular — the labour market is genuinely being reshaped, and the disruption to entry-level cognitive work is real — but in its central frame. The doomer position rests on four claims that do not survive contact with the evidence:

- That AI represents a categorical break from prior automation, making historical analogies invalid
- That productivity gains from AI will accrue only to capital, not to workers
- That the labour market is showing early signs of structural collapse rather than ordinary sectoral rotation
- That a new social contract is required to manage an inevitable surplus of unemployable humans

Each of these is a strong claim. Each is, at best, unsupported. At worst, several are flatly contradicted by the same Stanford and Brookings data the doomer camp routinely cites. The most rigorous recent work — Acemoglu's macroeconomic modelling, Brynjolfsson's worker-level studies, the Dallas Fed's labour analyses — points to a far more boring conclusion: AI is a general-purpose technology following the same pattern as steam, electricity, and the personal computer. It is reshaping tasks, not eliminating work. It is compressing skill premiums, not abolishing them. And it is doing so on a timeline that gives institutions room to adapt.

The doomer narrative serves political and commercial functions for the people promoting it. It does not, on present evidence, describe reality.

1. The starting point: what doomers actually believe

Before dismantling the position, it has to be stated fairly. The strongest version of the doomer case — articulated by OpenAI's recent policy document calling for a new "social contract," by Dario Amodei's public commentary, and by academic voices such as Daniel Susskind — runs roughly as follows.

> *Productivity growth and wealth creation are being divorced from jobs and income. Previous technologies augmented human labour; AI substitutes for it. When machines become more economical than humans at both physical and cognitive work, there is simply no role left for humans to escape forward into.* —
> Composite of OpenAI policy paper and Susskind, A World Without Work

This argument has three load-bearing components. First, an empirical claim: that AI is already cutting employment in exposed occupations. Second, a theoretical claim: that this time the *lump of labour fallacy* — the long-standing economic rebuttal to technology panics — does not apply, because AI eliminates the cognitive frontier workers would otherwise migrate toward. Third, a normative claim: that the appropriate response is a redistributive social contract — sovereign wealth funds, robot taxes, universal basic income, a right to AI access.

Each component deserves separate treatment. The empirical claim is partially true but dramatically overstated. The theoretical claim is the oldest mistake in labour economics dressed in new clothes. The normative claim builds an emergency infrastructure for a crisis that the evidence does not show is happening.

2. The data that gets quoted — and what it actually shows

The doomer position leans on a small set of recurring data points. They are real. They are also routinely stripped of the context that makes them legible. Here is the inventory.

2.1 The Stanford entry-level finding

Erik Brynjolfsson's Stanford Digital Economy Lab published *Canaries in the Coal Mine?* in August 2025 — the most-cited piece of empirical work in the doomer canon. Using ADP payroll data covering millions of workers, the paper found that workers aged 22 to 25 in the most AI-exposed occupations had experienced a 13% relative decline in employment since late 2022. The 2026 Stanford AI Index extended this to a nearly 20% decline in developer employment for the same age cohort.

What gets dropped when this is recycled: the same paper finds that **overall employment in AI-exposed occupations has grown robustly**. The 13% figure is for one cohort — early-career — in one tail of occupations. Older workers in the same occupations have seen employment hold steady or grow. Workers in less-exposed roles have grown faster than the overall economy. The Dallas Fed, summarising the same data, was explicit: this "suggests only a slight impact on the aggregate unemployment rate so far," driven primarily by fewer people transitioning into employment rather than by layoffs.

This is not the signature of a labour market in collapse. It is the signature of a labour market where the rungs at the bottom of one ladder have been kicked off while every other ladder remains intact. Painful for the people on that ladder. Not civilizational.

2.2 The McEntarfer caveat

Erika McEntarfer, until recently head of the US Bureau of Labor Statistics, made the point cleanly at Stanford's 2026 SIEPR Economic Summit: "Unemployment is edging up for those most AI-exposed occupations, but much more slowly than it is rising for everyone else." The hiring slowdown over the prior 18 months had, she noted, mostly impacted manual labour jobs — exactly the opposite of what the AI-displacement thesis predicts.

If AI were the dominant force restructuring labour, we would expect to see AI-exposed workers losing ground *faster* than the rest of the economy. We are seeing the reverse.

2.3 The productivity data nobody wants to talk about

Brynjolfsson, Li, and Raymond's *Generative AI at Work* — published in the Quarterly Journal of Economics — studied 5,179 customer support agents at a Fortune 500 firm. The findings:

Worker segment	Productivity gain from AI access
Novice / low-skill workers	**+34% issues resolved per hour**
Average across all workers	**+14%**
Experienced / high-skill workers	Minimal effect on speed; small decline in quality
New agents at 2 months tenure (with AI)	Matched performance of 6-month-tenure agents without AI

Source: Brynjolfsson, Li & Raymond, Generative AI at Work, QJE (2025), n = 5,179 agents.

The headline finding inverts the doomer prediction. AI was supposed to widen inequality by amplifying the gap between superstar knowledge workers and everyone else — the standard "skill-biased technical change" pattern that defined the 1990s and 2000s. What Brynjolfsson found is the opposite: **generative**

AI compresses the skill distribution. The novice gets the apprenticeship that previously took years, delivered in real time, free of charge. Customer sentiment improved. Worker retention improved. English fluency among international agents improved. None of this is what civilizational replacement looks like.

2.4 Adoption is broad, not concentrated

The St. Louis Fed reported in January 2026 that 55% of Americans and 37% of US workers were using generative AI tools. A separate Gallup Workforce Panel covering 30,000 employees from 2023 to 2026 showed workplace AI adoption rising from 9% to 26%. These are diffusion curves that look like the early personal computer, not like an extinction-level technology being deployed by a small cabal.

Two facts follow. First, the benefits are spreading widely — this is consumer technology, not a hedge fund's secret edge. Second, if AI were producing the catastrophic concentration the doomer position requires, we would expect adoption to look like the smartphone (a few firms dominate, everyone else consumes) or like industrial robotics (a few rich countries dominate). Instead, it looks like Microsoft Office.

3. The oldest mistake in labour economics, restated for 2026

Every generation of technological disruption produces a high-status intellectual movement convinced that *this time* the historical pattern will break. Each movement makes essentially the same argument: previous technologies augmented narrow human capacities, but the new technology substitutes for human work as such, so the displaced have nowhere to escape forward to. The movement is always confident, always credentialled, and — over a span of roughly 230 years of available evidence — always wrong.

3.1 The historical record

Era	Technology	Prediction	What actually happened
1589	Mechanical knitting frame	Elizabeth I refused the patent; predicted mass unemployment of hand knitters	Knitting industry expanded; textile employment grew for two centuries
1810s	Power loom	Luddites: weaving as a trade will be destroyed	98% labour reduction per yard of cloth; cloth consumption surged; textile employment quadrupled
1900	Mechanised agriculture	Permanent rural unemployment	US farm employment fell from 41% to 2% of workforce by 2000; total employment grew by ~150 million
1930	General industrial automation	Keynes warned of "technological unemployment" lasting generations	Postwar boom; lowest unemployment of the 20th century
1980s–90s	ATM machines	Bank tellers will be extinct	Tellers per branch fell from ~21 to ~13; branches multiplied; total teller employment rose ~10–15%
1980s	Spreadsheet software	Two million bookkeeping jobs destroyed	Bookkeeping declined; financial analyst, auditor, and accountant roles grew by millions

Era	Technology	Prediction	What actually happened
2016	Deep learning in radiology	Geoffrey Hinton: "Stop training radiologists now"	Radiology employment grew; the field is currently short-staffed

Sources: St. Louis Fed; Bessen (IMF F&D, 2015); Autor (TED 2016); AEI; Slate Magazine; CCIA labour analysis (2026).

The pattern is consistent enough to constitute a law: when a technology reduces the cost of performing a task, demand for the output of that task expands, and the residual human tasks become more valuable. ATMs did not destroy bank tellers because banks responded to cheaper branches by opening more of them, and because the remaining teller tasks — relationship banking, product sales — were higher-value than cash counting. Spreadsheets did not destroy bookkeepers because lower-cost financial analysis meant more demand for financial analysis. Power looms did not destroy weavers because the price elasticity of cloth was much higher than the technology's pessimists assumed.

3.2 The "this time is different" claim

Doomers anticipate this rebuttal and reach for the same counter-argument every generation reaches for: previous technologies replaced *some* tasks but left the cognitive frontier untouched. AI, the argument goes, replaces the cognitive frontier itself, so there is no "escape forward."

This is the version Susskind makes in *A World Without Work*, and the version Dario Amodei has restated in interviews. It is also exactly the argument made in 1589, 1810, 1900, 1930, and 2016 — adapted each time to whatever cognitive territory the contemporary doomer believed could not be conceded. Elizabeth I's advisors did not think knitters could become factory operators. Keynes did not think factory operators could become software engineers. Hinton did not think radiologists could become anything else. The argument's structure is invariant; only the territory shifts.

The argument also misreads what AI actually does. Current systems are extraordinarily good at *codified* cognitive work — the kind of work that can be specified, examined, and trained on. They are markedly worse at tacit knowledge, novel-context judgement, embodied skill, and what David Autor calls "expert intuition." The Stanford data confirms this directly: the entry-level cohort is being displaced **precisely because** entry-level work is codified work. Experienced workers are not being displaced because their value is in the tacit dimension AI cannot yet touch.

This does not mean AI will never reach those frontiers. It means the doomer prediction is making a claim about the next two decades that is not supported by the current technology's actual capabilities or trajectory. The model that successfully passed the bar exam still cannot reliably remember what it agreed to in the previous paragraph of a contract negotiation.

4. Where the productivity actually goes

The doomer position holds that AI's productivity gains flow to capital and starve labour. This is also testable.

4.1 Acemoglu's skeptical baseline

Daron Acemoglu's *The Simple Macroeconomics of AI* (Economic Policy, 2025) is the most rigorous attempt to bound the productivity question. Acemoglu uses a Hulten-theorem-based calculation: total productivity gain equals share of tasks affected multiplied by average cost saving per task. Plugging in current estimates, he gets a **0.66% increase in total factor productivity over ten years** — about 0.06% per year. He thinks even this is probably too high.

This matters for two reasons. First, Acemoglu is a careful, cautious labour economist who has spent his career documenting how technology can hurt workers — he is not a hype-skeptic from the AI-optimist camp. When the most credible bear in the academy says the productivity gains over a decade are likely under 1%, the doomer claim that AI is about to dissolve the labour market is making a claim that even AI's harshest economic critic does not endorse.

Second, productivity gains that small *cannot* produce mass unemployment. Mass unemployment requires that capital substitute for labour faster than demand expands. Acemoglu's modelling shows that under any plausible scenario, AI moves labour shares modestly — comparable to a single decade of normal technological change.

4.2 The PwC concentration finding

PwC's 2026 AI Performance Study surveyed 1,217 senior executives and found that 74% of AI's measurable economic value was being captured by 20% of firms. The doomer reading: AI consolidates wealth at the top. The actually-correct reading: **the technology is not yet reshaping aggregate productivity** because most firms have not figured out how to deploy it.

This is precisely the pattern Brynjolfsson and Hitt documented for IT investment in the 1990s: a long J-curve, with most firms experiencing negative initial returns as they fumble organizational redesign, before the diffusion phase produces broad productivity gains. We are in the fumbling phase. The doomer interpretation treats the early-phase concentration as a permanent feature.

4.3 Wages, not just jobs

Recent analysis from Phys.org and the Brookings AI labour reports finds that workers with AI-related skills are commanding wage premiums of more than 50%. Industries most exposed to AI are seeing not just productivity gains but employment growth *and* wage growth simultaneously. The pattern is closer to the early-internet wage premium than to a labour-share collapse.

None of this is the signature of capital strip-mining labour. It is the signature of a complement: workers who use the tool earn more; firms that use the tool produce more; the diffusion is uneven but the direction is unambiguous.

5. The social contract trap

This brings us back to the framing the original *Narrative Goldmine* note collects: that AI requires a "new social contract." Altman's OpenAI policy document — robot taxes, sovereign wealth funds, four-day workweek, right-to-AI — has become the dominant frame for how this debate is conducted at the policy level.

It is worth being clear about what is happening here. Altman is the CEO of an $852 billion company asking governments to prepare for a future in which his own product breaks the economic system. There are two ways to read this. Charitable: he genuinely believes the disruption is that large and is being responsible. Less charitable: he is constructing a political environment in which the most powerful AI companies become the indispensable partners of a new redistributive state, and in which the only acceptable critique of their products is downstream of accepting their framing of inevitability.

> *You don't make that pitch unless you believe it's actually coming — or unless making it serves you regardless of whether it's coming.* — Adaptation of The Rundown's commentary

The empirical case for the doomer-driven social contract is, on the evidence reviewed above, thin. The political case is strong, and the two are easy to confuse. A reasonable response to AI's actual labour effects — entry-level displacement in some occupations, productivity compression of skill premiums, sectoral rotation — would look more like targeted reskilling, education reform, and a serious conversation about apprenticeship pathways for under-25s. It would not look like sovereign wealth funds and a robot tax.

The risk of accepting the doomer frame is not that we accidentally help displaced workers. The risk is that we build the policy architecture for a crisis that does not arrive, and in doing so create exactly the labour-market rigidities that prevent the ordinary process of adjustment from working. The countries that handled the textile transition badly were the ones that resisted it. The countries that handled the agricultural transition badly were the ones that subsidised it indefinitely. There is a real chance that the most damaging response to AI is not the technology itself but the panic-driven policy response to it.

6. What is actually happening

Strip out the doomer overlay and what does the evidence actually describe?

- AI is being adopted faster than any prior general-purpose technology — 26% of US workers using it within roughly three years of ChatGPT's release. This is faster than the PC and faster than the internet.
- Productivity gains in early-adopter firms are real but heavily concentrated; the diffusion curve is in its messy phase. Aggregate measured productivity will likely lag for another two to five years before showing the gains.
- The entry-level segment of cognitively-codified occupations — software, customer service, basic marketing copy — is being squeezed. Workers aged 22-25 in those roles are seeing 13-20% employment declines relative to peers. This is real and worth taking seriously as an apprenticeship-pipeline problem, not as a civilizational signal.
- Older workers, manual workers, and workers in tacit-knowledge-intensive roles are unaffected or are gaining. Aggregate unemployment in AI-exposed occupations is rising more slowly than overall unemployment.
- The productivity gains observed at the worker level disproportionately benefit novices and low performers — the opposite of the skill-biased pattern of the 1990s and 2000s. AI looks more like a skill leveller than a skill polariser.
- Wages for AI-using workers are rising, not falling. The premium for AI fluency is currently above 50%, comparable to the early-internet skill premium.
- There is no evidence of structural collapse in the labour market. There is evidence of structural reshuffling, which is what general-purpose technologies always produce.

None of this denies that the transition will hurt specific people in specific places. It always does. The doomer mistake is to confuse the visible pain of the affected with a coming systemic failure. The visible pain is real. The systemic failure is not.

7. Where doomerism comes from

It is worth asking why a position so weakly supported by the evidence is so dominant in the discourse. Three factors are worth naming.

Commercial incentive

AI labs benefit from the narrative that their technology is so powerful it requires a redesign of civilization. It justifies valuations, regulatory capture, and the displacement of accountability onto governments. "Our product will break society, so society should pay us to break it carefully" is a remarkable business model. It is also the explicit pitch.

Cognitive availability

Displaced workers are visible; created jobs are not. The bookkeeper losing her job to spreadsheets in 1985 was a story; the financial analyst hired in 1992 because spreadsheet-enabled analysis became newly affordable was just "a financial analyst." The doomer narrative has a constant supply of vivid anecdotes; the rebuttal has only diffuse statistics. This is structural, not evidential.

Status defence

AI is uniquely threatening to the high-credentialed knowledge class — the journalists, lawyers, consultants, and academics who shape the discourse — in a way previous technologies were not. The agricultural mechanisation of 1900 was a story about somebody else. The current automation wave is a story about the people writing the story. The framing reflects that, and the framing should be read with that in mind.

8. Conclusion

The doomer position is not entirely empty. There is a serious case to be made about entry-level cognitive work, about apprenticeship pipelines, about the speed of adaptation in education systems that move on twenty-year cycles. There is a genuine policy challenge in helping the 22-to-25 cohort transition into a labour market where the bottom rungs of certain ladders have been removed. These are real problems and deserve real attention.

But the claim that AI represents an end-of-work event, that the social contract must be rewritten, that humans are about to become economically redundant — these claims rest on a misreading of the labour data, a misunderstanding of how previous technologies actually worked, and a misuse of the historical record. The Stanford findings the doomer camp cites show **sectoral disruption, not systemic collapse**. Brynjolfsson's worker-level studies show **skill compression, not capital triumph**. Acemoglu's macroeconomic modelling shows **modest productivity gains, not labour-share extinction**. The 230-year track record of identical predictions shows that "this time is different" has been wrong every single time it has been asserted with confidence.

The most likely future is the boring one. AI joins the personal computer, the internet, and electricity as a general-purpose technology that reshapes work, expands output, raises living standards unevenly, and produces a generation of policy disputes about distribution. The disputes are worth having. The civilizational frame is not.

The doomer position will not survive the data. It is already failing to.

Sources

Acemoglu, D. (2025). The Simple Macroeconomics of AI. Economic Policy, 40(121), 13–58.

Brynjolfsson, E., Chandar, B., & Chen, R. (2025). Canaries in the Coal Mine? Six Facts about the Recent Employment Effects of Artificial Intelligence. Stanford Digital Economy Lab.

Brynjolfsson, E., Li, D., & Raymond, L. (2025). Generative AI at Work. Quarterly Journal of Economics, 140(2), 889–942.

Stanford Institute for Human-Centered AI (2026). AI Index Report 2026.

Dallas Federal Reserve (2026). Young workers' employment drops in occupations with high AI exposure.

Bessen, J. (2015). Toil and Technology. Finance & Development, IMF.

St. Louis Federal Reserve (2026). Tracking AI's Contribution to GDP Growth.

PwC (2026). AI Performance Study. Survey of 1,217 senior executives across 25 sectors.

OpenAI (2026). Policy paper on the social contract for the intelligence age.

Bureau of Economic Analysis (2026). Q1 GDP estimate.

Susskind, D. (2020). A World Without Work.

Autor, D. (2016). Will Automation Take Away All Our Jobs? TED.

American Enterprise Institute (2022). What the Story of ATMs and Bank Tellers Reveals About the Rise of the Robots.

CCIA (2026). What Bank Tellers and Radiologists Can Tell Us about Our Job Security in the AI Era.